

Backup and Recovery Tool Version 2.6

1	EDB Postgres Backup and Recovery Installation Guide	4
1.1	Installing BART	4
1.2	Configuring BART	11
1.3	Upgrading BART	27
2	Quick Start Guide for RHEL/CentOS 7	29
3	Quick Start Guide for RHEL/CentOS 8	35
4	EDB Postgres Backup and Recovery Reference Guide	40
4.1	BART Subcommand Syntax and Examples	41
4.1.1	BACKUP	43
4.1.2	CHECK-CONFIG	48
4.1.3	DELETE	49
4.1.4	INIT	52
4.1.5	MANAGE	57
4.1.6	RESTORE	62
4.1.7	SHOW-SERVERS	66
4.1.8	SHOW-BACKUPS	67
4.1.9	VERIFY-CHKSUM	69
4.1.10	Running the BART WAL Scanner	70
4.2	Additional Examples	73
4.3	Sample BART System with Local and Remote Database Servers	100
5	EDB Postgres Backup and Recovery User Guide	124
5.1	Introduction	124
5.2	Overview	125
5.2.1	Block-Level Incremental Backup	127
5.2.1.1	1 Incremental Backup Limitations and Requirements	128
5.2.1.2	2 Concept Overview	129
5.2.1.3		130
5.2.1.4	4 Performing an Incremental Backup	131
5.2.1.5	5 Restoring an Incremental Backup	132
5.2.2	Creating a Backup Chain	134
5.3	Using BART	134
5.3.1	BART Management Overview	134
5.3.1.1		135
5.3.1.2		137
5.3.2	Managing Backups Using a Retention Policy	138
5.3.2.1		139
5.3.2.2		139
5.3.2.3		140
5.3.2.4		141
5.3.2.5		144
5.3.3	Basic BART Subcommand Usage	145
5.3.3.1		146
5.3.3.2		147
5.3.3.3		148
5.3.3.4		152
5.3.3.5		152
5.3.3.6		153
5.3.3.7	7 MANAGE	154

5.3.3.8	RESTORE	156
5.3.3.9	DELETE	158
5.3.4	Running the BART WAL Scanner	159
5.4	Using Tablespaces	161

1 EDB Postgres Backup and Recovery Installation Guide

This guide provides information about how to install and configure the EDB Backup and Recovery Tool (BART) 2.6.

Requirements Overview

Supported Platforms and Database Versions

• To view a complete list of supported platforms, visit the EDB website.

!!! Note BART 2.6 does not support CentOS/RHEL/OEL 6.x platforms. It is strongly recommended that EDB products running on these platforms be migrated to a supported platform.

- BART supports the following database versions:
 - Advanced Server versions 9.6, 10, 11, 12, and 13.
 - PostgreSQL versions 9.6, 10, 11, 12, and 13.

Software Requirements

You require the following components to install BART.

- BART Host Components Use EDB packages to add BART host components; see Installing BART for information about how to install these components.
- Additional Components In addition to the BART host components, the following components are required:
 - The Secure Shell (SSH) server daemon and Secure Copy (SCP) client programs must be enabled and activated on the BART host as well as on the remote database server hosts on which BART will be managing backup and recovery.
 - BART uses the pg_basebackup utility program when taking full backups.

Limitation

BART supports taking only a full backup of standby servers; it does not support taking incremental or parallel backups of standby servers.

1.1 Installing BART

This section will walk you through performing a fresh installation of BART on a host. Installation instructions are organized into the following platform/installer specific sections:

- Installing BART on a CentOS
- Installing BART on a RHEL Host
- Installing BART on a CentOS or RHEL Host
- Installing BART on a Debian or Ubuntu Host
- Installing BART on an SLES 12 Host

!!! Note If you are using the pdf version of this document, using cut/paste to copy command may result in extra spaces or carriage returns in the pasted command. If a command fails, check the command carefully for additional characters.

Installing BART on a CentOS Host

The following section demonstrates installing BART on a CentOS host using an RPM package. This section assumes that the user has some knowledge of installation and system administration procedures, and has administrative privileges on the host.

1. To install the repository configuration, assume superuser privileges and invoke one of the following platform-specific commands:

On CentOS 7:

```
yum -y install https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

On CentOS 8:

```
dnf -y install https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

2. Replace the USERNAME: PASSWORD in the following command with the username and password of a registered EnterpriseDB user:

```
sed -i "s@<username>:<password>@USERNAME:PASSWORD@" /etc/yum.repos.d/edb.repo
```

To request credentials for the repository, visit the EDB website.

3. Before installing BART, execute the following command to install the Extra Packages for Enterprise Linux (EPEL) release package:

On CentOS 7:

```
yum -y install https://dl.fedoraproject.org/pub/epel/epel-release-latest-
7.noarch.rpm
```

On CentOS 8:

```
dnf -y install epel-release
```

4. For CentOS 8, enable the PowerTools repository to satisfy EPEL package dependencies:

```
dnf config-manager --set-enabled PowerTools
```

5. For CentOS 8, disable the built-in PostgreSQL module:

```
dnf -qy module disable postgresql
```

6. Optionally, install the pg_basebackup utility program using the server client package. If you do not already have

the pg_basebackup program installed on the BART host, you can install a limited number of files that include the pg_basebackup program by invoking the following command:

On CentOS 7:

yum install edb-as<xx>-server-client

On CentOS 8:

dnf install edb-as<xx>-server-client

In the above command, replace <xx> with the required Advanced Server version. The pg_basebackup version must be the same or more recent than the database server to be backed up. For example, pg_basebackup version 10 can be used to back up database server version 10, but cannot be used to back up database server version 11.

7. Use the following command to install BART:

On CentOS 7:

yum -y install edb-bart

On CentOS 8:

dnf -y install edb-bart

Repeat the installation process described in this section to install BART on each remote host on which an incremental backup is to be restored.

To verify the BART installation, navigate to the /usr/edb/bart/bin directory and execute the following command:

bart --version

The bart --version command should return the current BART version. If the bart --version command returns an error stating the PATH is not available after switching from the root user to another BART user account, adjust the setting of the PATH environment variable to include the directory location of the BART bin subdirectory in the ~/.bashrc or ~/.bash_profile files of the following user accounts:

- The BART user account on the BART host. See Configuring BART for details.
- The remote user account on the remote host to which incremental backups are to be restored. For details, see the *EDB Backup and Recovery User Guide* available at the *EDB website*.

Upon successful installation, BART is installed in the BART_HOME directory:

/usr/edb/bart

The installation includes the following files:

File Name	Location	Description
bart	<bart_home>/bin</bart_home>	BART command line, executable program
bart-scanner	<bart_home>/bin</bart_home>	BART WAL scanner program
bart.cfg.sample	<bart_home> /etc</bart_home>	Sample BART configuration file
xlogreader_ident.so	<bart_home>/lib</bart_home>	Libraries supporting WAL versions

File Name	Location	Description
bart_license.txt	<bart_home></bart_home>	License agreement

After BART is installed successfully, you need to configure the installation.

Installing BART on a RHEL Host

The following section demonstrates installing BART on a RHEL host using an RPM package. This section assumes that the user has some knowledge of installation and system administration procedures, and has administrative privileges on the host.

1. To install the repository configuration, assume superuser privileges and invoke one of the following platform-specific commands:

On RHEL 7:

```
yum -y install https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

On RHEL 8:

```
dnf -y install https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

2. Replace the USERNAME: PASSWORD in the following command with the username and password of a registered EnterpriseDB user:

```
sed -i "s@<username>:<password>@USERNAME:PASSWORD@" /etc/yum.repos.d/edb.repo
```

To request credentials for the repository, visit the EDB website.

3. Before installing BART, execute the following command to install the Extra Packages for Enterprise Linux (EPEL) release package:

On RHEL 7:

```
yum -y install https://dl.fedoraproject.org/pub/epel/epel-release-latest-
7.noarch.rpm
```

On RHEL 8:

```
dnf -y install https://dl.fedoraproject.org/pub/epel/epel-release-latest-
8.noarch.rpm
```

4. Enable the repository:

On RHEL 7, enable the optional, extras, and HA repositories to satisfy EPEL package dependencies:

```
subscription-manager repos --enable "rhel-*-optional-rpms" --enable "rhel-*-
extras-rpms" --enable "rhel-ha-for-rhel-*-server-rpms"
```

On RHEL 8, enable the codeready-builder-for-rhel-8-*-rpms repository to satisfy EPEL packages dependency:

```
ARCH=$( /bin/arch )
```

subscription-manager repos --enable "codeready-builder-for-rhel-8-\${ARCH}-rpms"

5. For RHEL 8, disable the built-in PostgreSQL module:

dnf -qy module disable postgresql

6. Optionally, install the pg_basebackup utility program using the server client package. If you do not already have the pg_basebackup program installed on the BART host, you can install a limited number of files that include the pg_basebackup program by invoking the following command:

On RHEL 7:

yum install edb-as<xx>-server-client

On RHEL 8:

dnf install edb-as<xx>-server-client

In the above command, replace <xx> with the required Advanced Server version. The pg_basebackup version must be the same or more recent than the database server to be backed up. For example, pg_basebackup version 10 can be used to back up database server version 10, but cannot be used to back up database server version 11.

7. Use the following command to install the BART:

On RHEL 7:

yum -y install edb-bart

On RHEL 8:

dnf -y install edb-bart

Repeat the installation process described in this section to install BART on each remote host on which an incremental backup is to be restored.

To verify the BART installation, navigate to the /usr/edb/bart/bin directory and execute the following command:

bart --version

The bart --version command should return the current BART version. If the bart --version command returns an error stating the PATH is not available after switching from the root user to another BART user account, adjust the setting of the PATH environment variable to include the directory location of the BART bin subdirectory in the ~/.bashrc or ~/.bash_profile files of the following user accounts:

- The BART user account on the BART host. See Configuring BART for details.
- The remote user account on the remote host to which incremental backups are to be restored. For details, see the *EDB Backup and Recovery User Guide* available at the *EDB website*.

Upon successful installation, BART is installed in the BART_HOME directory:

/usr/edb/bart

The installation includes the following files:

File Name	Location	Description
bart	<bart_home>/bin</bart_home>	BART command line, executable program
bart-scanner	<bart_home>/bin</bart_home>	BART WAL scanner program
bart.cfg.sample	<bart_home> /etc</bart_home>	Sample BART configuration file
xlogreader_ident.so	<bart_home> /lib</bart_home>	Libraries supporting WAL versions
bart_license.txt	<bart_home></bart_home>	License agreement

After BART is installed successfully, you need to configure the installation.

Installing BART on a RHEL/CentOS 7 PPCLE Host

The following section demonstrates installing BART on a RHEL host using an RPM package. This section assumes that the user has some knowledge of installation and system administration procedures, and has administrative privileges on the host.

1. Install Advance Toolchain:

```
rpm --import
https://public.dhe.ibm.com/software/server/POWER/Linux/toolchain/at/redhat/RHEL7/g|
-pubkey-6976a827-5164221b

cat > /etc/yum.repos.d/advance-toolchain.repo <<EOF

# Beginning of the configuration file
[advance-toolchain]
name=Advance Toolchain IBM FTP
baseurl=https://public.dhe.ibm.com/software/server/POWER/Linux/toolchain/at/redhat,7
failovermethod=priority
enabled=1
gpgcheck=1
gpgkey=ftp://public.dhe.ibm.com/software/server/POWER/Linux/toolchain/at/redhat/RHI-pubkey-6976a827-5164221b
# End of the configuration file</pre>
```

2. To install the repository configuration, assume superuser privileges and invoke the following command:

```
yum -y install https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

3. Replace the USERNAME: PASSWORD in the following command with the username and password of a registered EnterpriseDB user:

```
sed -i "s@<username>:<password>@USERNAME:PASSWORD@" /etc/yum.repos.d/edb.repo
```

To request credentials for the repository, visit the EDB website.

4. Before installing BART, execute the following command to install the Extra Packages for Enterprise Linux (EPEL) release package:

```
yum -y install https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm
```

5. On RHEL 7, enable the optional, extras, and HA repositories to satisfy EPEL package dependencies:

```
subscription-manager repos --enable "rhel-*-optional-rpms" --enable "rhel-*-
extras-rpms" --enable "rhel-ha-for-rhel-*-server-rpms"
```

6. Invoke the following command to install BART:

```
yum -y install edb-bart
```

Installing BART on a Debian or Ubuntu Host

Perform the following steps to install a Debian package using the EnterpriseDB apt repository.

To request credentials for the repository, visit the EDB website.

1. Assume the superuser privileges.

```
sudo su -
```

2. To configure the EnterpriseDB repository on Debian 9, Ubuntu 18, and Ubuntu 20:

```
sh -c 'echo "deb https://username:password@apt.enterprisedb.com/$(lsb_release -
cs)-edb/ $(lsb_release -cs) main" > /etc/apt/sources.list.d/edb-$(lsb_release -
cs).list'
```

On Debian 10:

a. Set up the EnterpriseDB repository:

```
sh -c 'echo "deb [arch=amd64] https://apt.enterprisedb.com/$(lsb_release -cs)-
edb/ $(lsb_release -cs) main" > /etc/apt/sources.list.d/edb-$(lsb_release -
cs).list'
```

b. Substitute your EnterpriseDB credentials for the username and password placeholders in the following command:

```
sh -c 'echo "machine apt.enterprisedb.com login <username> password <password>" >
/etc/apt/auth.conf.d/edb.conf'
```

3. Add support to your system for secure APT repositories.

```
apt-get install apt-transport-https
```

4. Add the EDB signing key; When invoking the command, replace the username and password with the credentials provided by EnterpriseDB.

```
wget -q -0 - https://apt.enterprisedb.com/edb-deb.gpg.key | apt-key add -
```

5. Update the repository metadata.

```
apt-get update
```

6. Install the Debian package.

```
apt-get install edb-bart
```

Installing BART on an SLES 12 Host

This section provides instructions for installing BART on an SLES 12 SP4 host using the zypper package manager. BART is supported on SLES SP4 and SP5 versions.

1. Assume superuser privileges.

```
sudo su -
```

2. Use the following command to add the EDB repository to your SLES host:

```
zypper addrepo https://zypp.enterprisedb.com/suse/edb-sles.repo
```

3. Invoke the following command to refresh the metadata:

```
zypper refresh
```

4. Install SUSEConnect to register the host with SUSE to allow access to SUSE repositories:

```
zypper install SUSEConnect
```

5. Register the host with SUSE to allow access to SUSE repositories and replace 'REGISTRATION_CODE' and 'EMAIL' with your SUSE registration information:

```
SUSEConnect -r 'REGISTRATION_CODE' -e 'EMAIL'
```

SUSEConnect -p PackageHub/12.4/x86_64

```
SUSEConnect -p sle-sdk/12.4/x86_64
```

6. Install the following repository for PEM dependencies:

```
zypper addrepo
```

https://download.opensuse.org/repositories/Apache:/Modules/SLE_12_SP4/Apache:Module
o

7. Refresh the metadata:

```
zypper refresh
```

8. Then, use the zypper utility to install BART:

```
zypper -n install edb-bart
```

1.2 Configuring BART

To configure BART, you must establish the BART user account, configure the BART host, and configure the database server that will be backed up.

Establishing the BART User Account

The BART user account is an operating system user that will run the BART command line program. The BART user account must:

- own the BART backup catalog.
- be able to run the bart program and the bart-scanner program.
- be able to establish a SSH/SCP connection to and from each database server managed by BART.

You can optionally use the enterprised database user as the BART user account for an Advanced Server database and postgreSQL server. If you do not wish to use an existing database user as the BART user account, you must create an operating system user to assume the role.

Configuring BART and Database Server

As stated earlier, to configure BART, you must configure the BART host as well as the database server. The following table acts as a configuration parameter reference listing the mandatory and optional parameters with default values for [SERVER] as well as [BART] sections.

- Parameters set in the [BART] section are applicable to all BART managed database servers.
- Parameters set in the Server section are applicable only to the specific server; the Server parameter setting overrides the [BART] section setting.

For information about configuring BART host parameters, see the BART Host Parameter Reference and for information about configuring the database server parameters, see the Database Server Parameter Reference.

Parameter	Туре	Default	[SERVER]	[BART]
[BART]	Mandatory	N/A	N/A	Yes
<pre><bart_host></bart_host></pre>	Mandatory	N/A	N/A	Yes
<pre><backup_path></backup_path></pre>	Mandatory	N/A	N/A	Yes
<pg_basebackup_path></pg_basebackup_path>	Mandatory	N/A	N/A	Yes
retention_policy	Optional	<max_number>BACKUPS</max_number>	Yes	Yes
wal_compression	Optional	Disabled	Yes	Yes
<pre>copy_wals_during_restore</pre>	Optional	Disabled	Yes	Yes
xlog_method	Optional	fetch	Yes	Yes
logfile	Optional	<pre>/tmp/bart.log</pre>	N/A	Yes
scanner_logfile	Optional	<pre>/tmp/bart_scanner.log</pre>	N/A	Yes
<pre><bart_socket_directory></bart_socket_directory></pre>	Optional	/tmp	N/A	Yes
<pre><bart_socket_name></bart_socket_name></pre>	Optional	<md5 checksum=""></md5>	N/A	Yes
<thread_count></thread_count>	Optional	1	Yes	Yes
<pre><batch_size></batch_size></pre>	Optional	49152	Yes	Yes
<scan_interval></scan_interval>	Optional	0	Yes	Yes
<mbm_scan_timeout></mbm_scan_timeout>	Optional	20 seconds	Yes	Yes
<workers></workers>	Optional	1	Yes	Yes

Parameter	Туре	Default	[SERVER]	[BART]
[Server Name]	Mandatory	N/A	Yes	N/A
<pre><backup_name></backup_name></pre>	Optional	N/A	Yes	N/A
host	Mandatory	N/A	Yes	N/A
port	Mandatory	5444 for EPAS; 5432 for Postgres	Yes	N/A
user	Mandatory	N/A	Yes	N/A
<archive_path></archive_path>	Optional	BART backup catalog	Yes	N/A
<archive_command></archive_command>	Optional	N/A	Yes	N/A
		enterprisedb for EPAS		
<cluster_owner></cluster_owner>	Mandatory		Yes	N/A
		postgres for PostgreSQL		
<remote_host></remote_host>	Optional	N/A	Yes	N/A
<tablespace_path></tablespace_path>	Optional	N/A	Yes	N/A
allow_incremental_backup	Optional	Disabled	Yes	N/A
description	Optional	N/A	Yes	N/A

Configuring the BART Host

To configure the BART host, perform the following steps on the BART host as a root user:

Step 1. Navigate to the usr/edb/bart/etc directory and make a copy of the bart.cfg.sample file to create the bart.cfg file that will contain the parameter settings.

Step 2. Confirm that the Postgres pg_basebackup utility program is installed on the BART host. The pg_basebackup utility resides in the bin directory under your Postgres installation.

Step 3. Ensure the LD_LIBRARY_PATH environment variable includes the location of the libpq library. If your libpq library does not reside in the default location (POSTGRES_INSTALL_HOME/lib), you must add the library path to the LD_LIBRARY_PATH environment variable in the BART user account's profile (bash_profile) located in /home/

/home/sbart user account>:

```
### .bash_profile
### Get the aliases and functions
if [ -f ~/.bashrc ]; then
. ~/.bashrc
fi
### User specific environment and startup programs
export LD_LIBRARY_PATH=/usr/edb/as11/lib:$LD_LIBRARY_PATH
```

Step 4. Create the BART backup catalog and ensure the BART user account holds privileges on the BART backup catalog. In the following example, the BART configuration file specifies /opt/backup as the parent directory for the BART backup catalog in the
backup_path> parameter:

```
[BART]
bart_host = bartuser@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log
```

In the following example, bartuser is the BART user account. The example creates and sets the ownership and permissions on the BART backup catalog:

su root
mkdir /opt/backup
chown bartuser /opt/backup
chgrp bartuser /opt/backup
chmod 700 /opt/backup

If the subdirectory does not exist, BART creates a subdirectory for each database server listed in the configuration file when you invoke the bart command line program.

Step 5. Use your choice of editor to open the BART configuration file (located in the usr/edb/bart/etc directory) and edit the configuration as required. You must add the mandatory parameters to the [BART] section. Default values may be used for optional parameters.

The following table describes the [BART] host parameters.

Parameters/Placeholder Type		Description		
[BART]	Mandatory	Identifies the global section of the configuration file. It must be named BART.		
bart_host	Mandatory	Specify the bart user name and the IP address of the bart host on which the BART utility resides. You must specify it in the form of specify it in the form of specify it is specify in the specify of specify it is specify in the specify of specify it is specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify in the specify of specify in the specify of specify in t		
backup_path	Mandatory	Specify the path to the file system parent directory where all BART backups are stored.		
pg_basebackup_path	Mandatory	Specify the path to the pg_basebackup program that you installed on the BART host. For information about pg_basebackup version-specific restrictions, see the EDB Backup and Recovery User Guide available at the EDB website.		
wal_compression	Optional	Set this parameter to enabled to compress the archived WAL files in gzip format in the BART backup catalog when the MANAGE subcommand is invoked. By default it is set to disabled. The gzip compression program must be in the BART user account's PATH and the WAL compression setting must not be enabled for those database servers where you need to take incremental backups.		

Parameters/Placeholder	Туре	Description
copy_wals_during_restore	Optional	Set this parameter to enabled to copy the archived WAL files from the BART backup catalog to the restore_path/archived_wals directory prior to the database server archive recovery. Enabling this option helps you save time during the restore operation. Set this parameter to disabled (default) to retrieve the archived WAL files directly from the BART backup catalog during the database server archive recovery. During the restore operation, recovery settings will be saved in the postgresql.auto.conf file. The restore_command in the postgresql.auto.conf file will be determined by the value specified in the copy_wals_during_restore parameter. If the RESTORE subcommand is invoked with the -c option, the archived WAL files are copied from the BART backup catalog to the restore_path/archived_wals directory, thus overriding any setting of the copy_wals_during_restore parameter. If the RESTORE subcommand is invoked without the -c option, the value specified by the copy_wals_during_restore parameter is used.
xlog_method	Optional	Specify how the transaction log is collected during the execution of pg_basebackup through the BACKUP subcommand. Set xlog_method to fetch (default) to collect the transaction log files after the backup is completed. Set to stream to stream the transaction log in parallel with the full backup creation.
retention_policy	Optional	Set this parameter to determine when an active backup should be marked as obsolete when the MANAGE subcommand is used. You can specify the retention policy either in terms of number of backups or duration (days, weeks, or months). <max_number> BACKUPS (default), <max_number> DAYS, <max_number> WEEKS, or <max_number> MONTHS where <max_number> is a positive integer. For information about managing backups using a retention policy, see the EDB Backup and Recovery User Guide available at the EDB website.</max_number></max_number></max_number></max_number></max_number>
logfile	Optional	Use this parameter to specify the path to the BART log file. The default log file location is /tmp/bart.log. The log file will be created the first time you invoke the bart command line program using the sample configuration file value. To change the default setting, you must delete the bart.log file from the /tmp directory and create a new log file in another directory so that a new log file will be created and owned by the new BART user account. If no path to a log file is specified, BART does not create a log file.
scanner_logfile	Optional	Use this parameter to specify the path to the XLOG/WAL scanner log file. The default location is tmp/bart_scanner.log . The scanner log file will be created the first time you invoke the bart_scanner program using the sample configuration file value. To change the default setting, you must delete the bart_scanner.log file from the tmp directory and create a new log file in another directory so that a new log file will be created and owned by the new BART user account. If no path to a log file is specified, BART does not create a WAL scanner log file.

Parameters/Placeholder Type		Description
<pre><bart_socket_directory></bart_socket_directory></pre>	Optional	Specify the socket directory path where all BART sockets will be stored. The default directory is /tmp. While specifying the bart_socket_directory path, you must ensure that the directory exists and the BART user has the required access permissions to the directory.
<pre><bart_socket_name></bart_socket_name></pre>	Optional	Specify a user-friendly BART socket file name. Using this option overrides the default BART socket name generated using MD5 checksum. You must shut down the bart-scanner before setting this option. You can view the bart_socket_name in the sockPath field after starting the bart-scanner in the debug mode. This option helps in preventing the use of MD5 during the bart-scanner startup, thus making BART more compliant in FIPS mode.
<thread_count></thread_count>	Optional	Specify the number of worker threads for copying blocks (for incremental backups) or data files (for full backup) from the database server to the archive_path when the BACKUP subcommand is invoked. The default value is 1. The same set of worker threads are used for the compression operation when taking full backups in order to provide parallel, compressed backups when the BACKUP subcommand is specified with the -z or -c options. The compression operation does not apply to incremental backups. See thread count for more information.
<batch_size></batch_size>	Optional	Specify the number of blocks of memory used for copying modified blocks from the database server to the archive_path when the BACKUP subcommand is invoked for incremental backups. The default value is 49152 blocks; each block is 8192 bytes. The maximum permitted value is 131072 blocks and the minimum permitted value is 1 block. Reduce the batch_size setting if the server runs out of memory while executing the pg_read_binary_file()).
<scan_interval></scan_interval>	Optional	Specify the number of seconds after which the WAL scanner should initiate force scanning of the new WAL files. The default value is 0, which means no brute-force scanning will be started. After upgrading to the latest version of BART, users who have set this parameter to a non-default value may see increased CPU consumption on the part of bart-scanner. If this is an issue, consider increasing the configured value of scan_interval parameter, or removing the setting if it is not required.
<mbm_scan_timeout></mbm_scan_timeout>	Optional	Specify the number of seconds to wait for MBM files before timing out; this parameter is applicable only for incremental backup. You must set the <pre>scan_interval</pre> to a value significantly less than the MBM scan timeout. The default value is 20 seconds. The <pre>mbm_scan_timeout</pre> parameter value must be greater than 0. If the value is 0 or negative, then an error will be displayed during an incremental backup.
<workers></workers>	Optional	Specify the number of parallel worker processes required to stream the modified blocks of an incremental backup to the restore host. The default value is 1.

Thread Count

If the $\ensuremath{\mathsf{BACKUP}}$ subcommand is invoked with the $-{\sf thread-count}$ option, then the number of worker threads

specified by this option overrides any setting of the thread_count parameter in the BART configuration file. If the BACKUP subcommand is invoked without the --thread-count option, then the following determines the number of worker threads used:

- The setting of the thread_count parameter in the server section of the BART configuration file overrides the setting of thread_count in the global section for that particular database server.
- If omitted in the server section, the setting of thread_count in the global section is used.
- If the thread_count parameter is not specified in either section, the default is 1.
- When taking a full backup, if the thread count in effect is only 1, then the pg_basebackup utility is used to take the full backup unless the --no-pg_basebackup option is specified with the BACKUP subcommand.

<thread_count> will not be effective if the backup is taken on a standby server.

If parallel backup is run with $\frac{N}{N}$ number of worker threads, then it will initiate $\frac{N+1}{N}$ concurrent connections with the server.

Step 6 Invoke the CHECK-CONFIG subcommand, omitting the -s option to check the parameter settings in the BART configuration file. It should return the current BART version.

bart CHECK-CONFIG

The CHECK-CONFIG subcommand displays an error message if the required configuration is not properly set. You need to check the logfile to fix this.

Configuring the Database Server

To configure the database server, you must:

- 1. Authorize SSH/SCP access without a password prompt.
- 2. Create and configure a replication database user.
- 3. Adding the database server to the configuration file (server section).
- 4. Enable WAL archiving of the server.
- 5. Verify the server configuration settings.

The following section will walk you through the configuration process.

!!! Note You must authorize SSH/SCP access and set up a replication database user before restarting the database server with WAL archiving enabled.

Authorizing SSH/SCP Access

BART uses the Secure Shell (ssh) and Secure Copy (scp) Linux utility programs to copy the backup and WAL files from the BART managed database servers to the BART host as well as to restore backups.

- The client/server ssh and scp commands must not prompt for a password when establishing a connection with the target server (the server to which a passwordless connection is being made).
- A passwordless connection uses *authorized public keys* (public key of a client user account) to authenticate with the target server.
- You must add the public key of each client user account to the target user account's authorized public keys list on the target server.

For BART usage, there are two scenarios that require a passwordless SSH/SCP connection:

• When connecting from each BART managed database server (SSH/SCP client) to the BART host (target SSH/SCP server) to support WAL archiving as implemented by the archive_command parameter.

- In this case, the database server user account should generate the public key file (id_rsa.pub) with the ssh-keygen -t rsa command on the database server host.
- The public key file name should be appended to the ~/.ssh/authorized_keys file on the BART host. The authorized_keys file is in the BART user account's home directory.
- When connecting from the BART host (SSH/SCP client) to each BART managed database server (target SSH/SCP server) for taking incremental backups and for supporting restoration of the full backup, the archived WAL files, and the modified blocks, which occurs when the BART RESTORE subcommand is given.
 - In this case, the BART user account should generate the public key file (id_rsa.pub) with the ssh-keygen
 t rsa command on the BART host.
 - The public key file name should be appended to the ~/.ssh/authorized_keys file on the database server host. The authorized_keys file is in the home directory of the user account that owns the directory where the database backup is to be restored.
- If backups are to be taken from a given database server host, but restored to a different database server host, the passwordless SSH/SCP connections must be configured from the BART host to the database server host from which the backup is to be taken as well as from the BART host to the database server host to which the backup is to be restored.

See the EDB Backup and Recovery Reference Guide available at the EDB website to view examples of creating a passwordless connection.

Enabling Public Key Authentication

The following example enables SSH/SCP access on a CentOS 7.x host; similar (platform-specific) steps will apply to other platforms/versions.

- 1. In the SSH server daemon configuration file (sshd_config) located in the /etc/ssh, set the PubkeyAuthentication parameter to yes.
- 2. Reload the configuration file:

service sshd reload

If you get any SSH or SCP errors, examine the /var/log/secure log file.

Creating a Passwordless Connection

The following general instructions will walk you through generating a client's public key file, creating the target server's authorized public keys file, and creating a passwordless connection.

Step 1. On the client system, log in as the user account that will be initiating the SSH or SCP connection.

Step 2. Navigate to the user account's home directory and check for an existing .ssh subdirectory. If the .ssh directory does not exist, create one and assign the required privileges to the user.

Step 3. Generate the public key file with the following command. Accept all prompted defaults and do not specify a passphrase when prompted for one.

ssh-keygen -t rsa

The public key file named id_rsa.pub is created in the .ssh subdirectory.

Step 4. While logged into the client where you just generated the public key file, use SCP to make a temporary copy of it on the target server:

scp ~/.ssh/id_rsa.pub <target_user>@<host_address>:tmp.pub

Step 5. Navigate into the target user account's home directory and check for an existing ssh.org/ssh.org/s

Step 6. Append the temporary, client's public key file, tmp.pub, to the authorized_keys file. If an authorized keys file does not exist, create a new file, but do not completely replace any existing authorized keys file.

```
cat tmp.pub >> ~/.ssh/authorized_keys
```

Make sure the authorized_keys file is only accessible by the file owner and not by groups or other users. If the authorized_keys file does not have the required permission setting or it was newly created, change the file permissions as follows:

```
chmod 600 ~/.ssh/authorized_keys
```

Step 7. Delete the temporary public key file:

```
rm tmp.pub
```

Now, when logged into the client system as user there should be no prompt for a password when commands such as the following is given:

```
ssh target_user@host_address
```

Setting up a Replication Database User

For each database server that is to be managed by BART, a database user must be chosen to serve as the *replication database user*. The replication database user sets the Postgres archive_command configuration parameter when the INIT subcommand in invoked and creates backups when the BACKUP subcommand is invoked. The replication database user must be a superuser.

When executed with the PSQL client, the following PostgreSQL command creates a superuser to be the replication database user:

```
CREATE ROLE repuser WITH LOGIN SUPERUSER PASSWORD 'password';
```

The pg_hba.conf file must minimally permit the replication database user to have access to the database.

In the following example, the pg_hba.conf file permits the repuser (replication database user) to have access to the template1 database. The IP address from which repuser has access to template1 database is the location of the BART host:

For pg_basebackup only: If pg_basebackup is to be used for taking any backups (such as for standby servers), the replication database user must also be included in the pg_hba.conf file as a replication database connection as shown by the last entry in the following example.

```
### TYPE DATABASE USER ADDRESS METHOD
### "local" is for Unix domain socket connections only
local all all md5
### IPv4 local connections:
host template1 repuser 192.168.2.22/32 md5
host all enterprisedb 127.0.0.1/32 md5
### IPv6 local connections:
host all all ::1/128 md5
### Allow replication connections from localhost, by a user with the
### replication privilege.
host replication repuser 192.168.2.22/32 md5
```

The replication database user must be specified for the user parameter in the BART configuration file for the database server as shown in the following example:

```
[ACCTG]
host = 192.168.2.24
port = 5444
user = repuser
cluster_owner = enterprisedb
remote_host = enterprisedb@192.168.2.24
description = "Accounting"
```

There must be no password prompt when connecting to the database server with the replication database user. There are several ways to permit this; one recommended method is to use a pgpass file located in the BART user account's home directory.

For example, if bartuser is the BART user account, then the .pgpass file located in the /home/bartuser directory must contain the following entry:

```
192.168.2.24:5444::repuser:password
```

When bartuser invokes a BART backup, the password for the replication database user, repuser, is obtained from the .pgpass file of bartuser to connect to the database server running at 192.168.2.24 on port 5444.

The .pgpass file must contain an entry for each BART managed database server and its corresponding replication database user and password.

Adding a Database Server to the BART Configuration File

To manage the backup and recovery of a database server, you must add entries to the <code>[SERVER]</code> section of the BART configuration file, which is located in <code>SBART_HOME</code> directory.

Database Server Parameter Reference

Set the following parameters in the [SERVER] section of the BART configuration file. The parameter setting in the [SERVER] section overrides the setting in the global [BART] section for that particular database server. If omitted, the default value will be used.

For each cluster serviced by BART, the following parameters are mandatory:

```
host = 192.168.2.24
port = 5432
user = postgres
cluster_owner = postgres
```

!!! Note The port parameter setting is required only if the database server listens on a port other than the default (for example if Postgres listens on a port other than 5432).

The following table describes the database server parameters.

Parameters/Placeholder	Туре	Description
[ServerName]	Mandatory	Specify the server name that you want to backup using BART. It is not case-sensitive when referenced with BART subcommand options. A lowercase conversion of this name is used to create a subdirectory in the BART backup catalog for storing the backups and WAL files for this database server (for eg., epas12).
<pre><backup_name> host</backup_name></pre>	Optional	Specify a template for user-defined, friendly names that will be assigned to the backups of the database server. The maximum permitted length of backup name is 49 characters. The backup_name parameter can be overridden by thebackup_name option of the BACKUP subcommand. If this parameter is omitted from the BART configuration file, and thebackup_name option with a user-defined name is not specified with the BACKUP subcommand, then the backup can only be referenced in BART subcommands by the BART assigned, integer backup identifier. Specify the IP address of the database server to be configured for
nost	Mandatory	backup.
port	Mandatory	Specify the port number identifying the database server instance (that is, the relevant database cluster) to be backed up. The default port number for EPAS is 5444 and for Postgres it is 5432. The port parameter setting is only required if the database server listens on a port other than the default value.
User	Mandatory	Specify the replication database user name used by BART to establish the connection to the database server for full backups. See Setting up a Replication Database User for more information.
<archive_path></archive_path>	Optional	Specify the path where archived WAL files will be stored. The default location is the BART backup catalog (<backup_path>/<server_name>/archived_wals).</server_name></backup_path>
<archive_command></archive_command>	Optional	When the INIT subcommand is used, the content and variables specified in the BART <archive_command> result in the archive command string to be generated into the Postgres archive_command parameter in the postgresql.auto.conf file. To configure the BART <archive_command> parameter, enclose the command string within single quotes ('). If you do not specify the <archive_command> parameter in the configuration file, the default setting is taken as 'scp %p %h:%a/%f'. See Archive Command Auto Configuration for information about variables. The BART <archive_command> parameter in the BART configuration file, and the Postgres <archive_command> parameter in the postgresql.conf file (or the postgresql.auto.conf file) refer to two different parameters that are to be set in different manner.</archive_command></archive_command></archive_command></archive_command></archive_command>
<cluster_owner></cluster_owner>	Mandatory	Specify the Linux operating system user account that owns the database cluster. This is typically <code>enterprisedb</code> for Advanced Server database clusters installed in the Oracle compatible mode, or <code>postgres</code> for Advanced Server database clusters installed in the PostgreSQL compatible mode and PostgreSQL database clusters.

Parameters/Placeholder	Туре	Description
<remote_host></remote_host>	Optional	Specify the IP address of the remote server to which a backup is to be restored. Specify this parameter in the form of <remote_user>@<remote_host_address>.</remote_host_address></remote_user> is the user account on the target database server host that accepts a passwordless SSH/SCP login connection and owns the directory where the backup is to be restored. <remote_host_address> is the IP address of the remote host. For restoring a backup to a remote host or for restoring any backup where <remote_user> and the BART user account are not the same users, either this parameter must be set or it may be specified with the -r option with the BART RESTORE subcommand.</remote_user></remote_host_address>
<tablespace_path></tablespace_path>	Optional	Specify path to which tablespaces are to be restored in the format OID = <tablespace_path>; If the backup is to be restored to a remote host specified by the <remote_host> parameter, then the tablespace paths must exist on the remote host.</remote_host></tablespace_path>
allow_incremental_backups	Optional	Set this parameter to enabled to enable use of the WAL scanner and permit taking incremental backups when the BACKUP subcommand is invoked with theparent option. Set it to disabled (default) to disallow incremental backups and thus permit only full backups. For information about using the BACKUP subcommand and running the WAL scanner, please see the EDB Backup and Recovery User Guide available at the EDB website.
Description	Optional	Specify the description that will be used to identify the database server.

For information regarding the following parameters, see configuring the BART host.

- retention_policy
- xlog_method
- wal_compression
- copy_wals_during_restore.
- thread_count.
- batch_size.
- scan_interval.
- mbm_scan_timeout.
- workers

Backup Name Template

- The template is an alphanumeric string that may include the following variables that will be replaced with the timestamp values when the backup is taken:
 - %year to be replaced by 4-digit year
 - %month to be replaced by 2-digit month
 - %day to be replaced by 2-digit day
 - %hour to be replaced by 2-digit hour
 - %minute to be replaced by 2-digit minute
 - %second to be replaced by 2-digit second
- To include a percent sign (%) as a character in the backup name, specify ** in the template.
- Do not enclose the template string in quotes even if you want the template to include space characters, otherwise the enclosing quotes are stored as part of the backup name. However, when referenced with the -i option by BART

subcommands use of space characters in the backup name requires enclosing the backup name in quotes.

The following example shows the configuration settings of three database servers:

```
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
cluster_owner = enterprisedb
backup_name = acctg_%year-%month-%dayT%hour:%minute:%second
archive_command = 'cp %p %a/%f'
allow_incremental_backups = enabled
retention_policy = 8 BACKUPS
description = "Accounting"
[MKTG]
host = 192.168.2.24
port = 5444
user = repuser
cluster_owner = enterprisedb
remote_host = enterprisedb@192.168.2.24
allow_incremental_backups = enabled
description = "Marketing"
[HR]
host = 127.0.0.1
port = 5432
user = postgres
cluster_owner = postgres
retention_policy = 4 DAYS
description = "Human Resources"
```

Enabling WAL Archiving

WAL archiving must be enabled for the database server for which BART is to perform backup and recovery management.

- The WAL Archiving Configuration section describes the manual WAL archiving configuration process.
- The Archive Command Auto Configuration section describes an automated WAL archiving process.

WAL Archiving Configuration

Set the following configuration parameters in the postgresql.conf file to enable WAL archiving

- Set wal level to replica or higher.
- Set archive_mode to on.
- Set the PostgreSQL archive_command parameter to copy the WAL files to the archive_path. The archive_command configuration parameter mentioned here is located in the postgresql.conf file; the PostgreSQL archive_command parameter is used in a different manner than the BART archive_command.
- Set max_wal_senders to a value high enough to leave at least one session available for the backup. If the xlog_method=stream parameter setting is to be used by this database server, the max_wal_senders setting must account for an additional session for the transaction log streaming (the setting must be a minimum of 2). See Configuring the BART host for information about the xlog_method parameter.

For detailed information about WAL archiving, see the PostgreSQL Core Documentation.

The ARCHIVE PATH field displayed by the BART SHOW-SERVERS subcommand displays the full directory path where the WAL files should be copied as specified in the postgresql.conf file:

```
-bash-4.1$ bart SHOW-SERVERS -s acctg
SERVER NAME : acctg
HOST NAME: 192.168.2.24
USER NAME: repuser
PORT: 5444
REMOTE HOST:
RETENTION POLICY: none
DISK UTILIZATION: 0.00 bytes
NUMBER OF ARCHIVES: 0
ARCHIVE PATH : /opt/backup/acctg/archived_wals
ARCHIVE COMMAND : (disabled)
XLOG METHOD : fetch
WAL COMPRESSION: disabled
TABLESPACE PATH(s):
INCREMENTAL BACKUP : DISABLED
DESCRIPTION: "Accounting"
```

The parameter settings in the following example will copy the WAL files to a directory named /opt/backup/acctg/archived_wals on the BART host located at 192.168.2.22 as the bartuser user account. Using the bartuser account ensures that the operation will have sufficient permissions to copy to the BART backup catalog owned by bartuser.

The database server must be restarted in order to initiate WAL archiving, but do not do so until you have verified that the full path of the BART backup catalog has been created by a prior BART subcommand or the archive operation will fail.

Start the WAL scanner by executing the following command:

```
./bart-scanner
```

Archive Command Auto Configuration

To enable WAL archiving:

• In the postgresql.conf file, set the wal_level to replica or higher, archive_mode to on, and max_wal_senders to a value high enough to leave at least one session available for the backup. If the xlog_method=stream parameter setting is to be used by this database server as determined in the BART configuration file, the max_wal_senders setting must account for an additional session for the transaction log

streaming (that is, the setting must be a minimum of 2). See Configuring the BART host for information on the xlog_method parameter.

• Configure the Postgres archive_command parameter automatically with the INIT subcommand and restart the database server when you are ready to initiate WAL archiving. The INIT subcommand invokes the Postgres ALTER SYSTEM command to set the Postgres archive_command configuration parameter in the postgresql.auto.conf file located in the managed database server's POSTGRES_INSTALL_HOME data directory. For additional information about the INIT subcommand, see the EDB Backup and Recovery User Guide available at the EDB website.

The archive command string that the INIT subcommand generates into the postgresql.auto.conf file is determined by the parameter setting of the BART archive_command parameter in the server section of the BART configuration file. If the BART archive_command parameter is not set in the server section for a given database server, the command string that is configured uses the following default format:

```
'scp %p %h:%a/%f'
```

The following table describes these variables:

Variable Description

	·					
%p	The path of the file to archive used by the Postgres archiving process.					
%h	Will be replaced by the @<bart_host_address">bart_user>@<bart_host_address< a=""> as specified in the <bart_host> parameter setting.</bart_host></bart_host_address<>					
%a	Will be replaced by the BART archived_wals directory as specified in the archive path parameter setting. If the <archive_path> is not specified, then the default directory is</archive_path>					
	<pre><backup_path>/<server_name>/archived_wals. <server_name> is the lowercase conversion of the database server name.</server_name></server_name></backup_path></pre>					
%f	The archived file name used by the Postgres archiving process.					

The placeholders %h and %a are replaced by the INIT subcommand when creating the archive command string. The placeholders %p and %f are not replaced by the INIT subcommand, but are kept as given to be used by the Postgres archiving process.

For example, to use the default archive command format, the BART configuration file contains the following settings where the BART archive_command parameter is omitted from the server section for ACCTG:

```
[BART]
bart_host= bartuser@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log

[ACCTG]
host = 127.0.0.1
port = 5444
user = repuser
cluster_owner = enterprisedb
description = "Accounting"
```

The INIT subcommand is invoked by BART user account bartuser as follows:

```
[bartuser@localhost ~]$ bart INIT -s acctg -o
INFO: setting archive_command for server 'acctg'
WARNING: archive_command is set. server restart is required
```

If the BART backup catalog directory is not already complete, it will be completed.

The resulting archive command string in the postgresql.auto.conf file located in the managed database server's POSTGRES_INSTALL_HOME/data directory appears as follows:

```
### Do not edit this file manually!
### It will be overwritten by ALTER SYSTEM command.
archive_command = 'scp %p
bartuser@192.168.2.22:/opt/backup/acctg/archived_wals/%f'
```

Run the INIT subcommand with the -o option to override any existing archive_command setting in the postgresql.conf or the postgresql.auto.conf file. In addition, the -o option must be used to generate the command string if the archive_mode is set to off even if there are no existing settings of the archive_command in the postgresql.conf or postgresql.auto.conf files.

In this example, the following BART configuration file is used with an explicit setting of the BART archive_command parameter:

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log

[ACCTG]
host = 127.0.0.1
port = 5444
user = repuser
cluster_owner = enterprisedb
archive_command = 'cp %p %a/%f'
description = "Accounting"
```

The INIT subcommand is invoked by BART user account enterprisedb as follows:

```
-bash-4.1$ bart INIT -s acctg -o
INFO: setting archive_command for server 'acctg'
WARNING: archive_command is set. server restart is required
```

The resulting Postgres archive_command parameter in the postgresql.auto.conf file appears as follows:

```
### Do not edit this file manually!
### It will be overwritten by ALTER SYSTEM command.
archive_command = 'cp %p /opt/backup/acctg/archived_wals/%f'
```

When the database server has been restarted, the ARCHIVE COMMAND field of the SHOW-SERVERS subcommand displays the active Postgres archive command as shown by the following example:

```
-bash-4.1$ bart SHOW-SERVERS -s acctg
```

SERVER NAME : acctg HOST NAME : 127.0.0.1 USER NAME : repuser

PORT : 5444
REMOTE HOST :

RETENTION POLICY : none
DISK UTILIZATION : 48.00 MB

NUMBER OF ARCHIVES: 0

ARCHIVE PATH : /opt/backup/acctg/archived_wals

ARCHIVE SCOMMAND: `cp %p /opt/backup/acctg/archived_wals/%f`

XLOG METHOD : fetch

WAL COMPRESSION: disabled

TABLESPACE PATH(s):

INCREMENTAL BACKUP : DISABLED
DESCRIPTION : "Accounting"

Verifying Configuration Settings

To verify the parameter settings of the database server specified, execute the CHECK-CONFIG subcommand with the -s option:

```
bart CHECK-CONFIG [ -s server_name ]
```

The CHECK-CONFIG subcommand confirms the following:

- The cluster_owner parameter is set to the user account owning the database cluster directory.
- A passwordless SSH/SCP connection is set between the BART user and the user account specified by the cluster_owner parameter.
- The BART user parameter specifies a database superuser.
- The BART user has access to the backup directory catalog.
- The pg_hba.conf file contains a replication entry for the database superuser specified by the BART user parameter.
- The archive_mode parameter in the postgresql.conf file is enabled.
- The archive_command parameter in the postgresql.auto.conf or the postgresql.conf file is set.
- The allow_incremental_backups parameter in the BART configuration file is enabled for database servers for which incremental backups are to be taken.
- Archiving of WAL files to the archive_path is in process.
- The WAL scanner program is running.

After configuring the BART host and the database server(s), you can start using BART. For information about using BART, see the EDB Backup and Recovery User Guide available at the EDB website.

1.3 Upgrading BART

This section outlines the process of upgrading BART from an existing version to the latest version.

Upgrade Restrictions

The following restrictions apply with regard to previous BART versions.

• You can take incremental backups using the latest version only when the parent backup (full or incremental backup)

has also been taken with the latest version.

• Using the latest version, you can restore incremental backups that are taken only with the latest version of BART. However, using the latest version you can restore full backups that were taken with older versions.

Upgrading from Older Versions of BART

Perform the following steps to upgrade from older versions of BART to the latest version:

Step 1: Assume the identity of the BART user account and invoke the following command to stop the BART WAL scanner program (bart-scanner):

```
bart-scanner STOP
```

Step 2: As the root user, upgrade to the latest BART version with the yum upgrade command.

• To upgrade the BART RPM package directly from the EDB Yum Repository website, specify only the package name:

On CentOS 7:

```
yum upgrade edb-bart
```

You can also use a downloaded RPM package file to upgrade. To use a downloaded BART RPM package file to upgrade, use the yum command, specifying the complete RPM package file name:

```
yum upgrade edb-bart-x.y.z rhel7.x86_64.rpm
```

Where x denotes the major version of BART, and y and z denotes the minor version.

On a Debian or Ubuntu Host:

```
apt-get upgrade edb-bart
```

On a SLES Host:

```
zypper update edb-bart
```

Step 3: Repeat the process described in this section to upgrade to the latest BART version on each remote hosts where an incremental backup will be restored.

For additional information about restoration of incremental backups on remote hosts, see the EDB Backup and Recovery User Guide available at the EDB website.

Step 4: If the bart --version command returns an error stating the PATH is not available after switching from root user to another BART user account, adjust the setting of the PATH environment variable to include the location of the BART x.y.z (x denotes the major version of BART, and y and z denotes the minor version) executable (the bin subdirectory) in the ~/.bashrc or ~/.bash_profile files of the following user accounts:

- The BART user account on the BART host.
- The remote user account on the remote host to which incremental backups are to be restored. For details, see the EDB Backup and Recovery User Guide available at the EDB website.

The PATH setting should be the same as set for BART x.y.z since all versions use /usr/edb/bart/bin.

!!! Note After upgrading to the latest BART version, you must take a new full backup of your system before performing an incremental backup.

--- 1.4 Uninstalling BART ---

This section walks you through uninstalling BART.

Uninstalling BART on a RHEL/CentOS Host

To uninstall BART on a RHEL/CentOS host, assume the identity of the root user and invoke the following command:

On RHEL/CentOS 7:

```
yum remove edb-bart
```

On RHEL/CentOS 8:

```
dnf remove edb-bart
```

Uninstalling BART does not delete the backup files and archived WAL files that reside in the BART backup catalog. To permanently delete the backup files and archived WAL files in the BART backup catalog (/opt/backup), use one of the following commands:

- rm -rf /opt/backup
- BART DELETE subcommand

For information about the BART DELETE subcommand, see the *EDB Backup and Recovery User Guide* available at the EDB website.

Uninstalling BART on an SLES 12 Host

To uninstall BART on an SLES 12 host, assume the identity of the root user and invoke the following command:

```
zypper remove edb-bart
```

Uninstalling BART on a Debian/Ubuntu Host

To uninstall BART on a Debian or Ubuntu host, invoke the following command:

```
apt-get remove edb-bart
```

2 Quick Start Guide for RHEL/CentOS 7

This tutorial demonstrates using yum to install and configure Backup and Recovery Tool (BART) 2.6 on a CentOS 7 host with minimal configuration settings. The tutorial assumes that the user has some knowledge of installation and system administration procedures, and has administrative privileges on the host.

For detailed information about BART installation and configuration, see the BART Installation and Upgrade Guide,

available at the EDB website.

- BART is tested with the following database versions:
 - Advanced Server 9.6, 10, 11, 12, and 13.
 - o PostgreSQL 9.6, 10, 11, 12, and 13.

Installing BART

The following steps describe installing BART on CentOS 7.x OS using yum.

1. Assume superuser privileges and use yum to create the repository configuration file:

```
yum install -y https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

- 2. Create an EDB user account to request credentials to the EDB repository; for a user account visit the Enterprise DB website.
- 3. Use your choice of editor to open the repository configuration file (named edb.repo, located in /etc/yum.repos.d), and set the enabled parameter value to 1, and replace the username and password placeholders in the baseurl specification with the username and password of a registered EnterpriseDB user.
- 4. Update the cache:

yum makecache

5. Install an Advanced Server or PostgreSQL database server.

To install Advanced Server, execute the following command:

```
yum -y install edb-as12-server
```

Use sudo to assume the identity of the enterprisedb database superuser

```
sudo su - enterprisedb
```

Create an Advanced Server cluster named acctg on listener port 5444:

```
/usr/edb/as12/bin/initdb -D /var/lib/edb/as12/acctg
```

As the enterprisedb user, start the cluster:

```
/usr/edb/as12/bin/pg_ctl start -D /var/lib/edb/as12/acctg
```

You can check the status of the cluster with the following command:

```
/usr/edb/as12/bin/pg_ctl status -D /var/lib/edb/as12/acctg
```

!!! Note The BART host server is not required to have an Advanced Server or PostgreSQL installation, but must include a copy of the Postgres libpq library, the pg_basebackup utility program, and Boost Libraries 1.53 version for CentOS 7.

6. If you do not already have the pg_basebackup program installed on the BART host, you can use the following command to install a limited number of files that include the pg_basebackup program:

```
yum install edb-as<xx>-server-client
```

Where <xx> is the Advanced Server version.

7. As a root user, execute the following command to install BART:

```
yum install edb-bart
```

BART (the bart program and bart-scanner) is installed in the /usr/edb/bart/bin directory, referred to as <BART_HOME>. Repeat the installation process described in this section to install BART on all remote hosts where incremental backups are to be restored.

Configuring BART

Before configuring BART, establish the BART user account (the operating system user) that will run the BART command line program. Then, to configure the BART host and each database server that is to be managed by BART, perform the following steps:

1. Assume superuser privileges, create the directory that will hold the BART backup catalog, and assign its ownership (with restrictive privileges) to the BART user account:

In this example, bartuser is the BART user account and /opt/backup is the BART backup catalog.

```
su root
mkdir /opt/backup
chown bartuser /opt/backup
chgrp bartuser /opt/backup
chmod 700 /opt/backup
```

2. Navigate to the /usr/edb/bart/etc directory and copy the bart.cfg.sample file to create the BART configuration file (bart.cfg):

```
cp bart.cfg.sample bart.cfg
```

3. Open the BART configuration file (bart.cfg) using an editor of your choice and scroll through the BART configuration file to edit the file as required; sample settings are included for your reference. You must add the mandatory parameters to the [BART] and [ServerName] sections. Default values may be used for optional parameters. For detailed information about parameter settings, see the BART Installation and Upgrade Guide, available at the EDB website.

Parameters set in the <code>[BART]</code> section are applicable to all BART managed database servers, while parameters set in the <code>[ServerName]</code> section are applicable only to the specific server; <code>[ServerName]</code> settings override <code>[BART]</code> section settings.

In the following example, only mandatory parameters are set:

```
[BART]
bart_host= bartuser@192.168.169.199
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as12/bin/pg_basebackup

[EPAS12]
host = 127.0.0.1
user = repuser
cluster_owner = enterprisedb
```

The following table describes only mandatory parameters:

Parameters/Placeholder	Section	Description
bart_host	[BART]	Use this field to specify the BART user and the IP address of the host on which the BART utility is installed. Specify the value in the form of specify the value in the form of specify the value in the form of specify the BART user and the IP address of the host on which the BART user and the IP address of the host on which the BART user and the IP address of the host on which the BART user and the IP address of the host on which the BART utility is installed. Specify the value in the form of specify the specify the specif
backup_path	[BART]	Use this field to specify the path where all BART backups and archived WAL files will be stored. Ensure the BART user account holds privileges to create subdirectories and files within the location specified in the backup_path parameter. The default backup_path is BART backup catalog (/opt/backup).
pg_basebackup_path	[BART]	Use this field to specify the path to the pg_basebackup utility (/usr/edb/as <xx>/bin/pg_basebackup).</xx>
[ServerName]	[ServerName]	Specify the name of the database server to be backed up (for example, [EPAS12]).
host	[ServerName]	Specify the IP address of the database server to be configured for backup.
user	[ServerName]	Specify the replication database user name used by BART to establish the connection to the database server for full backups.
cluster_owner	[ServerName]	Specify the Linux operating system user account that owns the database cluster.

4. As a BART user, navigate to /usr/edb/bart/bin and invoke the following subcommand (omitting the -s option) to verify the [BART] section parameter settings:

bart CHECK-CONFIG

- 5. Authorize SSH/SCP access between the server and the BART host without a password prompt.
- 6. Create a replication database user for each database server that BART manages.
- 7. To enable continuous WAL archiving for any database server for which BART is to perform a backup, modify the postgresql.conf file, setting:
 - wal_level to replica or higher (for Postgres 9.6 or later)
 - archive_mode to on
 - archive_command (if it is not set in the bart.cfg file)
 - max_wal_senders to a value high enough to leave at least one session available for the backup.

After setting the parameters, restart the database server.

8. To start the WAL scanner, navigate to /usr/edb/bart/bin as a BART user and execute the following command:

```
./bart-scanner
```

9. If you are using the default archive_command, then navigate to /usr/edb/bart/bin as a BART user, run the INIT subcommand without the -o option, and restart the database server:

```
bart INIT [ -s { <server_name> | all } ]
```

Where <server_name> is the name of the database server to be backed up.

If you have customized the archive_command setting in the bart.cfg file, run the INIT subcommand with the o option to override any existing Postgresql archive_command setting in the postgresql.conf or the postgresql.auto.conf file, and restart the database server.

```
bart INIT [ -s { <server_name> | all } ] [ -o ]
```

10. To verify the database server parameter settings, as a BART user navigate to /usr/edb/bart/bin and invoke the CHECK-CONFIG subcommand with the -s option:

```
bart CHECK-CONFIG [ -s <server_name> ]
```

BART is now configured successfully. For detailed information about using BART, see the *EDB Backup and Recovery Tool User Guide*, available at the *EDB website*.

Creating a Passwordless Connection

The following example enables SSH/SCP access on a CentOS 7.x host; similar (platform-specific) steps will apply to other platforms/versions. You must create a passwordless connection between the BART host (SSH/SCP client) and the database server (target SSH/SCP server), as well as a passwordless connection between the database server (SSH/SCP client) and the BART host (target SSH/SCP server).

- 2. As a root user navigate to /usr/edb/bart, open the /etc/ssh/sshd_config file and set the PubkeyAuthentication parameter to yes.
- 3. Reload the configuration file:

```
service sshd reload
```

If you get any SSH or SCP errors, examine the log file (/var/log/secure).

4. As a BART user, use the following command to generate the public key file; you can accept the default responses:

```
ssh-keygen -t rsa
```

The public key file named id_rsa.pub is created in the .ssh subdirectory.

5. Use SCP to make a temporary copy of the public key file on the target server:

```
scp ~/.ssh/id_rsa.pub target_user@host_address:tmp.pub
```

- 6. As a target_user, log into the target server using ssh target_user@host_address command and navigate to the user account's home directory to check if there is an existing .ssh subdirectory. If it does not exist, create one with the required privileges.
- 7. Append the client's temporary public key file, tmp.pub, to the authorized_keys file:

```
cat tmp.pub >> ~/.ssh/authorized_keys
```

If an authorized_keys file does not exist, create a new file, but be careful not to completely replace any existing authorized_keys file.

8. Ensure only the file owner (and not other groups or users) has access to authorized_keys file:

```
chmod 600 ~/.ssh/authorized_keys
```

9. Delete the temporary public key file:

```
rm tmp.pub
```

Now, when logged into the BART host as a user, there should be no prompt for a password when you are connecting to the target database server:

```
ssh target_user@database_server_address
```

Creating a Passwordless Connection Between the Database Server and the BART Host

If backups are to be taken from a given database server host, but restored to a different database server host, the passwordless SSH/SCP connections must be configured from the BART host to the database server host from which the backup is to be taken as well as from the BART host to the database server host to which the backup is to be restored.

An example of how to create a passwordless connection is documented in the *EDB Backup and Recovery Reference Guide*, available at the *EDB website*.

Even when the Advanced Server database is on the same host as BART, and the Advanced Server database cluster owner is also the BART user account, a passwordless SSH/SCP connection must be established from the same user account to itself.

- 1. On the database server, navigate into the target user account's home directory to check for an existing .ssh subdirectory. If it does does not exist, create one in the user account's home directory with the required privileges.
- 2. As a database server user, generate the public key file:

```
ssh-keygen -t rsa
```

3. Create a temporary copy of the public key file:

```
scp ~/.ssh/id_rsa.pub target_user@host_address:tmp.pub
```

4. As a target user, log into the BART host and navigate to the user account's home directory to check if there is an existing .ssh subdirectory. If it does not exist, create one with the required privileges:

```
ssh target_user@host_address
```

5. Append the temporary, client's public key file to the authorized_keys file:

```
cat tmp.pub >> ~/.ssh/authorized_keys
```

If an authorized_keys file does not exist, create a new file, but do not completely replace any existing authorized_keys file.

6. Ensure only the file owner (and not other groups or users) has access to authorized_keys file:

```
chmod 600 ~/.ssh/authorized_keys
```

7. Delete the temporary public key file:

```
rm tmp.pub
```

Now, when logged into the database server as a user, there should be no prompt for a password when you are connecting to the BART host:

```
ssh bart_user@bartip_address
```

Creating a Replication Database User

1. To create a replication database user (a superuser), connect to the database server with the psql client, and invoke the following PostgreSQL command:

```
CREATE ROLE <repuser> WITH LOGIN SUPERUSER PASSWORD '<password>';
```

- 2. Specify this replication database user in the user parameter of the bart.cfg file.
- 3. The pg_hba.conf file must minimally permit the replication database user to have access to the database. The IP address from which the replication database user has access to the database is the BART host location. The replication database user must also be included in the pg_hba.conf file as a replication database connection if pg_basebackup is to be used for taking any backups.
- 4. To ensure there is no password prompt when connecting to the database server with the replication database user, a recommended method is to use the .pgpass file located in the BART user account's home directory (if it does not exist, you need to create the .pgpass file with the required privileges). The .pgpass file must contain an entry for each BART managed database server, and its corresponding replication database user and password.

3 Quick Start Guide for RHEL/CentOS 8

This tutorial demonstrates using the dnf command to install and configure the EDB Backup and Recovery Tool (BART) 2.6 on a CentOS 8 host with minimal configuration settings. The tutorial assumes that the user has some knowledge of installation and system administration procedures and has administrative privileges on the host.

For detailed information about BART installation and configuration, see the *BART Installation and Upgrade Guide* available at the EDB website.

- BART is tested with the following database versions:
 - Advanced Server 9.6, 10, 11, 12, and 13.
 - o PostgreSQL 9.6, 10, 11, 12, and 13.

Installing BART

The following steps describe installing BART on CentOS 8.x OS.

1. Assume superuser privileges and use dnf to create the repository configuration file:

```
dnf install -y https://yum.enterprisedb.com/edbrepos/edb-repo-latest.noarch.rpm
```

- 2. Create an EDB user account to request credentials to the EDB repository; for a user account visit the Enterprise DB website.
- 3. Use your choice of editor to open the repository configuration file (named edb.repo, located in /etc/yum.repos.d), set the enabled parameter value to 1, and replace the username and password placeholders in the baseurl specification with the username and password of a registered EnterpriseDB user.
- 4. Update the cache:

dnf makecache

5. Install an Advanced Server or PostgreSQL database server.

To install Advanced Server, execute the following command:

```
dnf -y install edb-as12-server
```

Use sudo to assume the identity of the enterprisedb database superuser:

```
sudo su - enterprisedb
```

Create an Advanced Server cluster named acctg on listener port 5444:

```
/usr/edb/as12/bin/initdb -D /var/lib/edb/as12/acctg
```

As the enterprisedb user, start the cluster:

```
/usr/edb/as12/bin/pg_ctl start -D /var/lib/edb/as12/acctg
```

You can check the status of the cluster with the following command:

```
/usr/edb/as12/bin/pg_ctl status -D /var/lib/edb/as12/acctg
```

!!! Note The BART host server is not required to have an Advanced Server or PostgreSQL installation, but must include a copy of the Postgres libpq library, the pg_basebackup utility program, and Boost Libraries 1.66 version for CentOS 8.

6. If you do not already have the pg_basebackup program installed on the BART host, you can use the following command to install a limited number of files that include the pg_basebackup program:

```
dnf install edb-asxx-server-client
```

7. As a root user, use the following command to install the BART RPM package:

```
dnf install edb-bart
```

BART (the bart program and bart-scanner) is installed in the /usr/edb/bart/bin directory, referred to as <BART_HOME>. Repeat the installation process described in this section to install BART on all remote hosts where incremental backups are to be restored.

Configuring BART

Before configuring BART, establish the BART user account (the operating system user) to run the BART command line program. Then, to configure the BART host and each database server that is to be managed by BART, perform the following steps:

1. Assume superuser privileges, create the directory that will hold the BART backup catalog, and assign its ownership (with restrictive privileges) to the BART user account:

In this example, bartuser is the BART user account and /opt/backup is the BART backup catalog.

```
su root
mkdir /opt/backup
chown bartuser /opt/backup
chgrp bartuser /opt/backup
chmod 700 /opt/backup
```

2. Navigate to the /usr/edb/bart/etc directory and copy the bart.cfg.sample file to create the BART configuration file (bart.cfg):

```
cp bart.cfg.sample bart.cfg
```

3. Open the BART configuration file (bart.cfg) using an editor of your choice and scroll through the BART configuration file to edit the file as required; sample settings are included for your reference. You must add the mandatory parameters to the [BART] and [ServerName] sections. Default values may be used for optional parameters. For detailed information about parameter settings, see the BART Installation and Upgrade Guide, available at the EDB website.

Parameters set in the <code>[BART]</code> section are applicable to all BART managed database servers, while parameters set in the <code>[ServerName]</code> section are applicable only to the specific server; <code>[ServerName]</code> settings override <code>[BART]</code> section settings.

In the following example, only mandatory parameters are set:

```
[BART]
bart_host= bartuser@192.168.169.199
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as12/bin/pg_basebackup

[EPAS12]
host = 127.0.0.1
user = repuser
cluster_owner = enterprisedb
```

The following table describes only mandatory parameters:

Parameters/Placeholder	Section	Description
bart_host	[BART]	Use this field to specify the BART user and the IP address of the host on which the BART utility is installed. Specify the value in the form of bart_user>@<bart_host_address></bart_host_address> .
backup_path	[BART]	Use this field to specify the path where all BART backups and archived WAL files will be stored. Ensure the BART user account holds privileges to create subdirectories and files within the location specified in the backup_path parameter. The default backup_path is BART backup catalog (/opt/backup).
pg_basebackup_path	[BART]	Use this field to specify the path to the pg_basebackup utility (/usr/edb/as <xx>/bin/pg_basebackup).</xx>
[ServerName]	[ServerName]	Specify the name of the database server to be backed up (for example, [EPAS12]).
host	[ServerName]	Specify the IP address of the database server to be configured for backup.
user	[ServerName]	Specify the replication database user name used by BART to establish the connection to the database server for full backups.
cluster_owner	[ServerName]	Specify the Linux operating system user account that owns the database cluster.

4. As a BART user, navigate to /usr/edb/bart/bin and invoke the following subcommand (omitting the -s option) to verify the [BART] section parameter settings:

bart CHECK-CONFIG

- 5. Authorize SSH/SCP access between the server and the BART host without a password prompt.
- 6. Create a replication database user for each database server that BART manages.
- 7. To enable continuous WAL archiving for any database server for which BART is to perform a backup, modify the postgresql.conf file, setting:
 - wal_level to replica or higher (for Postgres 9.6 or later)
 - o archive mode to on
 - archive_command (if it is not set in the bart.cfg file)
 - max_wal_senders to a value high enough to leave at least one session available for the backup.

After setting the parameters, restart the database server.

8. To start the WAL scanner, navigate to /usr/edb/bart/bin as a BART user and execute the following command:

```
./bart-scanner
```

9. If you are using the default archive_command, then navigate to /usr/edb/bart/bin as a BART user, run the INIT subcommand without the -o option, and restart the database server:

```
bart INIT [ -s { <server_name> | all } ]
```

Where <server_name> is the name of the database server to be backed up.

If you have customized the archive_command setting in the bart.cfg file, run the INIT subcommand with the o option to override any existing Postgresql archive_command setting in the postgresql.conf or the postgresql.auto.conf file, and restart the database server.

```
bart INIT [ -s { <server_name> | all } ] [ -o ]
```

10. To verify the database server parameter settings, as a BART user navigate to /usr/edb/bart/bin and invoke the CHECK-CONFIG subcommand with the -s option:

```
bart CHECK-CONFIG [ -s <server_name> ]
```

BART is now configured successfully. For detailed information about using BART, see the *EDB Backup and Recovery Tool User Guide* available at the *EDB website*.

Creating a Passwordless Connection

The following example enables SSH/SCP access on a CentOS 7.x host; similar (platform-specific) steps will apply to other platforms/versions. You must create a passwordless connection between the BART host (SSH/SCP client) and the database server (target SSH/SCP server), as well as a passwordless connection between the database server (SSH/SCP client) and the BART host (target SSH/SCP server).

- 1. Log in as the user account on the BART host that will be initiating the SSH or SCP connection and navigate to the user account's home directory and check for an existing .ssh subdirectory. If the .ssh directory does not exist, create one with the required privileges.
- 2. As a root user navigate to /usr/edb/bart, open the /etc/ssh/sshd_config file and set the PubkeyAuthentication parameter to yes.
- 3. Reload the configuration file:

service sshd reload

If you get any SSH or SCP errors, examine the log file (/var/log/secure).

4. As a BART user, use the following command to generate the public key file; you can accept the default responses:

```
ssh-keygen -t rsa
```

The public key file named id_rsa.pub is created in the .ssh subdirectory.

5. Use SCP to make a temporary copy of the public key file on the target server:

```
scp ~/.ssh/id_rsa.pub target_user@host_address:tmp.pub
```

- 6. As a target_user, log into the target server using ssh target_user@host_address command and navigate to the user account's home directory to check if there is an existing .ssh subdirectory. If it does not exist, create one with the required privileges.
- 7. Append the temporary client's public key file, tmp.pub, to the authorized_keys file:

```
cat tmp.pub >> ~/.ssh/authorized_keys
```

If an authorized_keys file does not exist, create a new file, but be careful not to completely replace any existing authorized_keys file.

8. Ensure only the file owner (and not other groups or users) has access to authorized_keys file:

```
chmod 600 ~/.ssh/authorized_keys
```

9. Delete the temporary public key file:

```
rm tmp.pub
```

Now, when logged into the BART host as a user, there should be no prompt for a password when you are connecting to the target database server:

```
ssh target_user@database_server_address
```

Creating a Passwordless Connection Between the Database Server and the BART Host

If backups are to be taken from a given database server host, but restored to a different database server host, the passwordless SSH/SCP connections must be configured from the BART host to the database server host from which the backup is to be taken as well as from the BART host to the database server host to which the backup is to be restored.

An example of how to create a passwordless connection is documented in the *EDB Backup and Recovery Reference Guide*, available at the *EDB website*.

Even when the Advanced Server database is on the same host as BART, and the Advanced Server database cluster owner is also the BART user account, a passwordless SSH/SCP connection must be established from the same user account to itself.

- 1. On the database server, navigate into the target user account's home directory to check for an existing .ssh subdirectory. If it does does not exist, create one in the user account's home directory with the required privileges.
- 2. As a database server user, generate the public key file:

```
ssh-keygen -t rsa
```

3. Create a temporary copy of the public key file:

```
scp ~/.ssh/id_rsa.pub target_user@host_address:tmp.pub
```

4. As a target user, log into the BART host and navigate to the user account's home directory to check if there is an existing .ssh subdirectory. If it does not exist, create one with the required privileges:

```
ssh target_user@host_address
```

5. Append the client's temporary public key file to the authorized_keys file:

```
cat tmp.pub >> ~/.ssh/authorized_keys
```

If the authorized_keys file does not exist, create a new file, but do not completely replace any existing authorized_keys file.

6. Ensure that only the file owner (and not other groups or users) has access to authorized_keys file:

```
chmod 600 ~/.ssh/authorized_keys
```

7. Delete the temporary public key file:

```
rm tmp.pub
```

Now, when logged into the database server as a user, there should be no prompt for a password when you are connecting to the BART host:

```
ssh bart_user@bartip_address
```

Creating a Replication Database User

1. To create a replication database user (a superuser), connect to the database server with the psql client, and invoke the following PostgreSQL command:

```
CREATE ROLE <repuser> WITH LOGIN SUPERUSER PASSWORD '';
```

- 2. Specify this replication database user in the user parameter of the bart.cfg file.
- 3. The pg_hba.conf file must minimally permit the replication database user to have access to the database. The IP address from which the replication database user has access to the database is the BART host location. The replication database user must also be included in the pg_hba.conf file as a replication database connection if pg_basebackup is to be used for taking any backups.
- 4. To ensure there is no password prompt when connecting to the database server with the replication database user, a recommended method is to use the .pgpass file located in the BART user account's home directory (if it does not exist, you need to create the .pgpass file with the required privileges). The .pgpass file must contain an entry for each BART managed database server, and its corresponding replication database user and password.

4 EDB Postgres Backup and Recovery Reference Guide

This guide acts as a quick reference for BART subcommands and provides comprehensive examples of the following BART operations:

- Performing a full backup of database servers
- Performing a point-in-time recovery (PITR) on a remote PostgreSQL database server
- Restoring an incremental backup
- Restoring a database cluster with tablespaces
- Evaluating, marking, and deleting backups and incremental backups
- Configuring and operating local and remote database servers

For detailed information about BART subcommands and operations, see the EDB Backup and Recovery User Guide available at the EDB website.

The document is organized as follows:

- See Subcommands to view BART subcommand examples.
- See Examples to view BART operations examples.
- See Sample BART System to view examples of both local and remote database server configuration and operation.

4.1 BART Subcommand Syntax and Examples

This section briefly describes each BART subcommand and provides an example.

Invoking BART

BART subcommands are invoked at the Linux command line as a BART user. You can invoke the bart program (located in the <BART_HOME>/bin directory) with the desired options to manage your BART installation.

The following examples demonstrate ways of invoking BART. In these examples, the BART user account is named bartuser.

```
$ su bartuser
Password:
$ export
LD_LIBRARY_PATH=/opt/PostgresPlus/9.6AS/lib/:$LD_LIBRARY_PATH
$ ./bart SHOW-SERVERS
```

To run BART from any current working directory:

```
$ su bartuser
Password:
$ export
LD_LIBRARY_PATH=/opt/PostgresPlus/9.6AS/lib/:$LD_LIBRARY_PATH
$ bart SHOW-SERVERS
```

Syntax for invoking BART

```
bart [ <general_option> ]... [ <subcommand> ] [<subcommand_option>]...
```

You can use either abbreviated or long option forms on the command line (for example -h or --help).

General Options

You can specify the following general options with bart.

```
-h or (--help)
```

- Displays general syntax and information about BART usage.
- All subcommands support a help option (-h, --help). If the help option is specified, information is displayed regarding that particular subcommand. The subcommand, itself, is not executed.

The following code sample displays the result of invoking the --help option for the BACKUP subcommand:

```
-bash-4.2$ bart BACKUP --help
bart: backup and recovery tool

Usage:
bart BACKUP [OPTION]...

Options:
-h, --help Show this help message and exit
-s, --server Name of the server or 'all' (full backups only) to specify all servers
-F, --format=p|t Backup output format (tar (default) or plain)
-z, --gzip Enables gzip compression of tar files
-c, --compress-level Specifies the compression level (1 through 9, 9 being best compression)
--backup-name Specify a friendly name for the current backup
--parent Specify parent backup for incremental backup
--check Verify checksum of required mbm files
```

```
The following code sample displays information returned by the bart --version subcommand:
```

```
[edb@localhost bin]$ bart --version
bart (EnterpriseDB) 2.5.2
[edb@localhost bin]$
```

```
-d (or --debug)
```

-v (or --version)

The following code sample displays debugging output returned by the bart MANAGE subcommand:

```
-bash-4.1$ bart -d MANAGE -n

DEBUG: Server: acctg, Now: 2015-04-17 16:34:03 EDT, RetentionWindow:

259200 (secs) ==> 72 hour(s)

DEBUG: Server: dev, Now: 2015-04-17 16:34:03 EDT, RetentionWindow:

1814400 (secs) ==> 504 hour(s)

DEBUG: Server: hr, Now: 2015-04-17 16:34:03 EDT, RetentionWindow:

7776000 (secs) ==> 2160 hour(s)
```

```
-c (or --config-path) <config_file_path>
```

The following code sample demonstrates using the -c option to specify a non-default configuration file name and installation location:

```
$ su bartuser
Password:
$ export
LD_LIBRARY_PATH=/opt/PostgresPlus/9.6AS/lib/:$LD_LIBRARY_PATH
$ bart -c /home/bartuser/bart.cfg SHOW-SERVERS
```

4.1.1 BACKUP

Use the BACKUP subcommand to create a full or incremental backup.

Syntax for a Full Backup:

```
bart BACKUP -s { <server_name> | all } [ -F { p | t } ]

[ -z ] [ -c <compression_level> ]

[ --backup-name <backup_name> ]

[ --thread-count <number_of_threads> ]

[ { --with-pg_basebackup | --no-pg_basebackup } ]
```

Syntax for an Incremental Backup:

Before performing an incremental backup, you must take a full backup. For more details about incremental backup, refer to *Block-Level Incremental Backup* in the *EDB Backup and Recovery User Guide* available at the <u>EDB website</u>.

The following table describes the BACKUP options:

Options

Description

```
-s { <server_name>
| all }
--server {
<server_name> | all
}
```

Use this option to specify the database server to be backed up.

Specify <server_name> to take a backup of the database server (as specified in the BART configuration file).

Specify all to take a backup of all servers.

Options	Description
-F { p t }format { p t }	Use this option to specify the backup file format. Specify p option to take backup in plain text format and specify t option to take backup in tar format. If the p or t option is omitted, the default is tar format. Use p option with the BACKUP subcommand when streaming is used as a backup method. An incremental backup can only be taken in plain text format (p).
-z (gzip)	This option is applicable only for full backup and tar format. Use this option to enable gzip compression of tar files using the default compression level (typically 6).
<pre>-c <compression_level>compress-level <compression_level></compression_level></compression_level></pre>	This is applicable only for full backup and tar format. Use this option to specify the gzip compression level on the tar file output. compression_level is a digit from 1 through 9, with 9 being the best compression.
backup-name <backup_name></backup_name>	Use this option to assign a user-defined, alphanumeric friendly name to the backup. The maximum permitted length of backup name is 49 characters. For detailed information about this parameter, see the EDB Backup and Recovery User Guide available at the EDB website. If the optionbackup-name is not specified and the backup_name parameter is not set for this database server in the BART configuration file, then the backup can only be referenced in other BART subcommands by the BART assigned backup identifier.
<pre>thread-count <number_of_threads></number_of_threads></pre>	Use this option to specify the number of worker threads to run in parallel to copy blocks for a backup. For detailed information about thethread-count parameter, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.
with- pg_basebackup	This is applicable only for full backup. Use this option to specify the use of pg_basebackup to take a full backup. The number of thread counts in effect is ignored as given by the thread_count parameter in the BART configuration file. When taking a full backup, if the thread count in effect is greater than 1, then the pg_basebackup utility is not used to take the full backup (parallel worker threads are used) unless thewith-pg_basebackup option is specified with the BACKUP subcommand.
no-pg_basebackup	This is applicable only for full backup. Use this option to specify that pg_basebackup is not to be used to take a full backup. When taking a full backup, if the thread count in effect is only 1, then the pg_basebackup utility is used to take the full backup unless theno-pg_basebackup option is specified with the BACKUP subcommand.
<pre>parent { <backup_id> <backup_name> }</backup_name></backup_id></pre>	Use this option to take an incremental backup. The parent backup is a backup taken prior to the incremental backup; it can be either a full backup or an incremental backup. <backup_id> is the backup identifier of a parent backup and <backup_name> is the user-defined alphanumeric name of a parent backup.</backup_name></backup_id>
check	This is applicable only for incremental backup. Use this option to verify if the required MBM files are present in the BART backup catalog before taking an incremental backup. However, an actual incremental backup is not taken when thecheck option is specified. Theparent option must be used along with thecheck option.
checksum- algorithm	While taking a backup, you can specify one of the following values with thechecksum-algorithm option:checksum-algorithm=MD5 (default) to generate MD5 checksum fileschecksum-algorithm=SHA256 to generate SHA256 checksum fileschecksum-algorithm=NONE to skip generating checksum files.

Examples

The following code sample demonstrates using variables with the BACKUP subcommand:

```
./bart backup -s ppas12 -Ft --backup-name "YEAR = %year MONTH = %month DAY = %day"

./bart backup -s ppas12 -Ft --backup-name "YEAR = %year MONTH = %month DAY = %day %%"
```

```
./bart show-backups -s ppas12 -i "test backup"
```

The following code sample displays the result of creating a full backup in the default tar format with gzip compression when the BACKUP subcommand was invoked. Note that checksums are generated for the full backup and user-defined tablespaces for the tar format backup:

```
[edb@localhost bin]$ ./bart BACKUP -s hr -z
INFO: DebugTarget - getVar(checkDiskSpace.bytesAvailable)
INFO: new backup identifier generated 1567591909098
INFO: creating 5 harvester threads
NOTICE: all required WAL segments have been archived
INFO: backup completed successfully
INFO:
BART VERSION: 2.5
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1567591909098
BACKUP NAME: none
BACKUP PARENT: none
BACKUP LOCATION: /home/edb/bkup_new/hr/1567591909098
BACKUP SIZE: 13.91 MB
BACKUP FORMAT: tar.gz
BACKUP TIMEZONE: America/New_York
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 0
TABLESPACE(s): 3
0id
        Name
               Location
16387 test1 /home/edb/tbl1
16388 test2 /home/edb/tbl2
16389
       test3
              /home/edb/tbl3
START WAL LOCATION: 000000010000000000000025
STOP WAL LOCATION: 0000000100000000000000026
BACKUP METHOD: streamed
BACKUP FROM: master
START TIME: 2019-09-04 06:11:49 EDT
STOP TIME: 2019-09-04 06:11:53 EDT
TOTAL DURATION: 4 sec(s)
```

The following code sample displays information about the directory containing the full backup:

```
[edb@localhost bin]$number_of_threads>
[edb@localhost bin]$ ls -l /home/edb/bkup_new/hr/
total 8
drwxrwxr-x. 3 edb edb 34 Aug 27 05:57 1566899819709
```

```
drwxrwxr-x. 3 edb edb 58 Aug 27 05:57 1566899827751
drwxrwxr-x. 3 edb edb 4096 Sep 4 06:11 1567591909098
drwxrwxr-x. 2 edb edb 4096 Sep 4 06:11 archived_wals
[edb@localhost bin]$
```

The following code sample displays information about the creation of a full backup while streaming the transaction log. Note that the Pop option must be specified with the BACKUP subcommand when streaming is used as a backup method.

```
[edb@localhost bin]$ ./bart BACKUP -s ACCTG -Fp
INFO: DebugTarget - getVar(checkDiskSpace.bytesAvailable)
INFO: new backup identifier generated 1566898964200
INFO: creating 5 harvester threads
NOTICE: pg_stop_backup complete, all required WAL segments have been archived
INFO: backup completed successfully
INFO:
BART VERSION: 2.5
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1566898964200
BACKUP NAME: none
BACKUP PARENT: none
BACKUP LOCATION: /home/edb/bkup_new/acctg/1566898964200
BACKUP SIZE: 46.03 MB
BACKUP FORMAT: plain
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 0
TABLESPACE(s): 0
START WAL LOCATION: 000000010000000000000017
BACKUP METHOD: streamed
BACKUP FROM: master
START TIME: 2019-08-27 05:42:44 EDT
STOP TIME: 2019-08-27 05:42:46 EDT
TOTAL DURATION: 2 sec(s)
```

The following code sample displays the assignment of a user-defined backup name with the --backup-name option:

```
[edb@localhost bin]$ ./bart BACKUP -s acctg --backup-name acctg_%year-%month-%day
INFO: DebugTarget - getVar(checkDiskSpace.bytesAvailable)
INFO: new backup identifier generated 1566899004804
INFO: creating 5 harvester threads
NOTICE: pg stop backup complete, all required WAL segments have been archived
INFO: backup completed successfully
INFO:
BART VERSION: 2.5
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1566899004804
BACKUP NAME: acctg_2019-08-27
BACKUP PARENT: none
BACKUP LOCATION: /home/edb/bkup_new/acctg/1566899004804
BACKUP SIZE: 46.86 MB
BACKUP FORMAT: tar
BACKUP TIMEZONE: US/Eastern
```

XLOG METHOD: fetch
BACKUP CHECKSUM(s): 0
TABLESPACE(s): 0

START WAL LOCATION: 00000001000000000000001A

BACKUP METHOD: streamed BACKUP FROM: master

START TIME: 2019-08-27 05:43:24 EDT STOP TIME: 2019-08-27 05:43:24 EDT

TOTAL DURATION: 0 sec(s)

The following code sample displays an incremental backup taken by specifying the --parent option. The option -Fp must be specified while taking an incremental backup as incremental backup can be taken only in plain text format.

[edb@localhost bin]\$./bart BACKUP -s hr -Fp --parent hr_full_1 --backup-name hr_incr_1 INFO: DebugTarget - getVar(checkDiskSpace.bytesAvailable) INFO: checking /home/edb/bkup_new/hr/archived_wals for MBM files from 0/20000028 to 0/22000000 INFO: new backup identifier generated 1566899827751 INFO: creating 5 harvester threads NOTICE: all required WAL segments have been archived INFO: backup completed successfully INFO: BART VERSION: 2.5 **BACKUP DETAILS:** BACKUP STATUS: active BACKUP IDENTIFIER: 1566899827751 BACKUP NAME: hr_incr_1 BACKUP PARENT: 1566899819709 BACKUP LOCATION: /home/edb/bkup_new/hr/1566899827751 BACKUP SIZE: 7.19 MB BACKUP FORMAT: plain BACKUP TIMEZONE: America/New_York XLOG METHOD: fetch BACKUP CHECKSUM(s): 0 TABLESPACE(s): 0 START WAL LOCATION: 000000010000000000000022 STOP WAL LOCATION: 0000000100000000000000023 BACKUP METHOD: streamed BACKUP FROM: master START TIME: 2019-08-27 05:57:07 EDT STOP TIME: 2019-08-27 05:57:08 EDT TOTAL DURATION: 1 sec(s)

The following code sample displays an incremental backup taken by specifying the --checksum-algorithm=NONE option to skip generating checksum files.

First, the bart-scanner is started.

```
edb@localhost bin]$
[edb@localhost bin]$ ./bart-scanner -d --checksum-algorithm=NONE

DEBUG: sockPath = /tmp/fc557c1c8853d75f1cb52a8a578f371a

INFO: process created for server 'ppas11', pid = 19012

DEBUG: could not load XLogReaderLibrary at this time, archived_wals is empty
```

Then, an incremental backup is taken with the --checksum-algorithm=NONE option to skip generating checksum files.

```
[edb@localhost bin]$ ./bart backup -s ppas11 -Fp --parent 1593506709152 --checksum-
algorithm=NONE
INFO: DebugTarget - getVar(checkDiskSpace.bytesAvailable)
INFO: checking /home/edb/bkup/ppas11/archived_wals for MBM files from 1/D3000028
to 1/D9000000
INFO: new backup identifier generated 1593507779811
INFO: creating 5 harvester threads
NOTICE: pg_stop_backup complete, all required WAL segments have been archived
INFO: backup completed successfully
INFO:
BART VERSION: 2.6devel
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1593507779811
BACKUP NAME: none
BACKUP PARENT: 1593506709152
BACKUP LOCATION: /home/edb/bkup/ppas11/1593507779811
BACKUP SIZE: 7.30 MB
BACKUP FORMAT: plain
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 0
TABLESPACE(s): 0
START WAL LOCATION: 0000000100000001000000009
BACKUP METHOD: streamed
BACKUP FROM: master
START TIME: 2020-06-30 05:02:59 EDT
STOP TIME: 2020-06-30 05:03:05 EDT
```

!!! Note To restore an incremental backup taken with the --checksum-algorithm=NONE option, you must specify --disable-checksum while restoring.

Similarly, you can specify --checksum-algorithm=MD5 or --checksum-algorithm=SHA256 while taking an incremental backup if you want to generate MD5 or SHAD256 checksum files.

4.1.2 CHECK-CONFIG

TOTAL DURATION: 6 sec(s)

The CHECK-CONFIG subcommand checks the global parameter settings as well as the database server configuration in the BART configuration file.

The following syntax is used to check the BART configuration file global section settings.

```
bart CHECK-CONFIG
```

The following syntax is used to check the database server configuration settings.

```
bart CHECK-CONFIG [ -s <server_name> ]
```

The following table describes the CHECK-CONFIG option:

Example

The following code sample demonstrates successfully checking the BART configuration file global parameters with the bart CHECK-CONFIG command:

```
bash-4.1$ bart CHECK-CONFIG
INFO: Verifying that pg_basebackup is executable
INFO: success -
INFO: success - pg_basebackup(/usr/edb/as11/bin/pg_basebackup) returns
version 11.400000
```

The following code sample demonstrates successfully checking the BART configuration file database server parameters with the bart CHECK-CONFIG command with the -s option:

```
[edb@localhost bin]$ ./bart check-config -s hr
INFO: Checking server hr
INFO: Verifying cluster_owner and ssh/scp connectivity
INFO: success
INFO: Verifying user, host, and replication connectivity
INFO: success
INFO: Verifying that user is a database superuser
INFO: success
INFO: Verifying that cluster_owner can read cluster data files
INFO: success
INFO: Verifying that you have permission to write to vault
INFO: success
INFO: /home/edb/bkup_new/hr
INFO: Verifying database server configuration
INFO: success
INFO: Verifying that WAL archiving is working
INFO: waiting 30 seconds for
/home/edb/bkup_new/hr/archived_wals/00000001000000000000001E
INFO: success
INFO: Verifying that bart-scanner is configured and running
```

4.1.3 DELETE

The **DELETE** subcommand removes the subdirectory and data files from the BART backup catalog for the specified backups along with archived WAL files.

Syntax:

```
bart DELETE -s <server_name>
-i { all | [']{ <backup_id> | <backup_name> },... }['] }
[ -n ]
```

Note that when invoking the DELETE subcommand, you must specify a database server.

For database servers under a retention policy, there are conditions where certain backups may not be deleted. For more information, see the EDB Backup and Recovery User Guide available at the EDB website.

The following table describes the **DELETE** options:

Options

Description

```
-s <server_name>
                                 <server_name> is the name of the database server whose backups are to be
                                 deleted.
--server <server_name>
-i { all | [']{
<backup_id> |
                                 <backup id> is the backup identifier of the backup to be deleted.
<backup_name> }',... }
                                 <backup_name> is the user-defined alphanumeric name for the backup.
[`]}
                                 Multiple backup identifiers and backup names may be specified in a comma-
                                 separated list. The list must be enclosed within single quotes if there is any white
--backupid { all | [']{
                                 space appearing before or after each comma (see Example).
<backup_id> |
                                 If all is specified, all backups and their archived WAL files for the specified
                                 database server are deleted.
<backup_name> }',... }
[`] }
-n
                                 Performs the test run and displays the results prior to physically removing files;
                                 no files are actually deleted.
--dry-run
```

Example

The following code sample demonstrates deleting a backup from the specified database server:

After the deletion, the BART backup catalog for the database server no longer contains the corresponding directory for the deleted backup ID. The following code sample displays information about archived_wals subdirectory that no longer contains the backup WAL files:

```
[edb@localhost acctg]$ ls -l
total 16
drwxrwxr-x. 3 edb edb 4096 Aug 27 06:03 1566900199604
drwxrwxr-x. 3 edb edb 4096 Aug 27 06:03 1566900204377
drwxrwxr-x. 3 edb edb 4096 Aug 27 06:03 1566900209087
drwxrwxr-x. 3 edb edb 4096 Aug 27 06:05 1566900321228
drwxrwxr-x. 2 edb edb 6 Aug 27 06:01 archived_wals
```

The following code sample demonstrates deleting multiple backups from the database server.

```
[edb@localhost bin]$ ./bart DELETE -s acctg -i `1566988095633,1566988100760,
acctg_2019-08-28`
INFO: deleting backup `1566988095633` of server `acctg`
INFO: deleting backup `1566988095633`
INFO: WALs of deleted backup(s) will belong to prior backup(if any), or will
be marked unused
WARNING: not marking any WALs as unused WALs, the WAL file
`/home/edb/bkup_new/acctg/archived_wals/000000010000000000000037` is required,
yet not available in archived_wals directory
INFO: backup(s) deleted
INFO: deleting backup `1566988100760` of server `acctg`
INFO: deleting backup `1566988100760`
INFO: WALs of deleted backup(s) will belong to prior backup(if any), or will
be marked unused
WARNING: not marking any WALs as unused WALs, the WAL file
`/home/edb/bkup_new/acctg/archived_wals/000000010000000000000039` is
required, yet not available in archived_wals directory
INFO: backup(s) deleted
INFO: deleting backup `acctg_2019-08-28` of server `acctg`
INFO: deleting backup `1566988115512`
INFO: WALs of deleted backup(s) will belong to prior backup(if any), or will
be marked unused
WARNING: not marking any WALs as unused WALs, the WAL file
`/home/edb/bkup_new/acctg/archived_wals/000000010000000000000003C` is required,
yet not available in archived_wals directory
INFO: backup(s) deleted
[edb@localhost bin]$
[edb@localhost bin]$
[edb@localhost bin]$
[edb@localhost acctg]$
[edb@localhost acctg]$ ls -l
total 8
drwxrwxr-x. 3 edb edb 4096 Aug 28 06:28 1566988105086
drwxrwxr-x. 3 edb edb 4096 Aug 28 06:28 1566988109477
drwxrwxr-x. 2 edb edb 6 Aug 28 06:09 archived_wals
[edb@localhost acctg]$
```

Deleting Multiple Backups with Space Characters

The following code sample demonstrates deleting multiple backups; since there are space characters in the commaseparated list, the entire list must be enclosed within single quotes:

```
INFO: WALs of deleted backup(s) will belong to prior backup(if any), or will
be marked unused
WARNING: not marking any WALs as unused WALs, the WAL file
`/home/edb/bkup_new/acctg/archived_wals/000000010000000000000002A` is required,
yet not available in archived wals directory
INFO: backup(s) deleted
INFO: deleting backup `1566900209087` of server `acctg`
INFO: deleting backup `1566900209087`
INFO: WALs of deleted backup(s) will belong to prior backup(if any), or will
be marked unused
WARNING: not marking any WALs as unused WALs, the WAL file
`/home/edb/bkup_new/acctg/archived_wals/000000010000000000000002C` is required,
yet not available in archived_wals directory
INFO: backup(s) deleted
[edb@localhost bin]$
[edb@localhost bin]$
[edb@localhost acctg]$ ls -l
total 4
drwxrwxr-x. 3 edb edb 4096 Aug 27 06:05 1566900321228
drwxrwxr-x. 2 edb edb 6 Aug 27 06:01 archived_wals
[edb@localhost acctg]$
```

4.1.4 INIT

The INIT subcommand is used to create the BART backup catalog directory, rebuild the BART backupinfo file, and set the archive_command in the server based on the archive_command setting in the bart.cfg file.

Syntax:

```
bart INIT [ -s { <server_name> | all } ] [ -o ]
[ -r [ -i { <backup_id> | <backup_name> | all } ] ]
[--no-configure]
```

The following table describes the **INIT** options:

Options Description -s { <server_name> <server_name> is the name of the database server to which the INIT actions are to be | all } applied. If all is specified or if the option is omitted, actions are applied to all servers. --server { <server_name> | all } Overrides the existing Postgres archive_command configuration parameter setting in the postgresql.conf file or the postgresql.auto.conf file using the BART -0 archive_command parameter in the BART configuration file. The INIT generated archive --override command string is written to the postgresql.auto.conf file.

Options	Description
-r rebuild	Rebuilds the backupinfo file located in each backup subdirectory. If all is specified or if the option is omitted, the backupinfo files of all backups for the database servers specified by the -s option are recreated. This option is only intended for recovering from a situation where the backupinfo file has become corrupt. If the backup was initially created with a user-defined backup name, and then the INIT -r option is invoked to rebuild that backupinfo file, the user-defined backup name is no longer available. Thus, future references to the backup must use the backup identifier.
<pre>-i {</pre>	<pre><backup_id> is an integer, backup identifier and <backup_name> is the user-defined alphanumeric name for the backup. The -i option can only be used with the -r option.</backup_name></backup_id></pre>
no- configure	Prevents the archive_command from being set in the PostgreSQL server.

Examples

In the following code sample, you can see that archive_mode = off and archive_command is not set. After invoking the BART INIT subcommand, archive_mode is set to on and archive_command is set:

```
archive_mode = off # enables archiving; off, on, or always
### (change requires restart)
archive_command = ''
### command to use to archive a logfile segment
[edb@localhost bin]$ ./bart init -s ppas11
INFO: setting archive_mode/archive_command for server 'ppas11'
WARNING: archive_mode/archive_command is set. Restart the PostgreSQL
server using 'pg_ctl restart'
[edb@localhost bin]$
### Do not edit this file manually!
### It will be overwritten by the ALTER SYSTEM command.
archive_mode = 'on'
archive_command = 'scp %p
edb@127.0.0.1:/home/edb/bkup/ppas11/archived_wals/%f'
```

In the following code sample, you can see that archive_mode = on, and archive_command is not set. After invoking the INIT subcommand, archive_command is set:

```
archive_mode = on # enables archiving; off, on, or always
### (change requires restart)
archive_command = '' # command to use to archive a logfile segment
[edb@localhost bin]$ ./bart init -s ppas11
INFO: setting archive_mode/archive_command for server 'ppas11'
WARNING: archive_command is set. Reload the configuration in the
PostgreSQL server using pg_reload_conf() or 'pg_ctl reload'
[edb@localhost bin]$
### Do not edit this file manually!
### It will be overwritten by the ALTER SYSTEM command.
archive_command = 'scp %p
edb@127.0.0.1:/home/edb/bkup/ppas11/archived_wals/%f'
```

In the following code sample, you can see that archive_command are already set. After invoking the INIT subcommand, there is no change in their settings. Note that to override the existing archive_command, you must include the -o option.

```
archive_mode = on # enables archiving; off, on, or always
### (change requires restart)
archive_command = 'scp %p
edb@127.0.0.1:/home/edb/bkup/ppas11/archived_wals/%f' # command to use
to archive a logfile segment
### placeholders: %p = path of file to archive
[edb@localhost bin]$ ./bart init -s ppas11
INFO: setting archive_mode/archive_command for server 'ppas11'
WARNING: archive_command is not set for server 'ppas11'
[edb@localhost bin]$
### Do not edit this file manually!
### It will be overwritten by the ALTER SYSTEM command.
```

In the following code sample, you can see that archive_mode = off and archive_command is already set. After invoking the INIT subcommand archive_mode is set to on:

```
archive_mode = off # enables archiving; off, on, or always
### (change requires restart)
archive_command = 'scp %p
edb@127.0.0.1:/home/edb/bkup/ppas11/archived_wals/%f' # command to use
to archive a log file segment
[edb@localhost bin]$ ./bart init -s ppas11
INFO: setting archive_mode/archive_command for server 'ppas11'
WARNING: archive_mode/archive_command is set. Restart the PostgreSQL
server using 'pg_ctl restart'
### Do not edit this file manually!
### It will be overwritten by the ALTER SYSTEM command.
archive_mode = 'on'
archive_command = 'scp %p
edb@127.0.0.1:/home/edb/bkup/ppas11/archived_wals/%f'
```

In the following code sample an existing archive command setting is overridden by resetting the archive_command in the PostgreSQL server with the archive_command = 'cp %p %a/%f' parameter from the bart.cfg file:

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup_edb
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log

[ACCTG]
host = 127.0.0.1
port = 5444
user = repuser
cluster_owner = enterprisedb
archive_command = 'cp %p %a/%f'
description = "Accounting"
```

The archive_mode and archive_command parameters in the database server are set as follows:

Invoke the INIT subcommand with the -o option to override the current archive_command setting in the PostgreSQL server:

```
-bash-4.1$ bart INIT -s acctg -o
INFO: setting archive_mode/archive_command for server 'acctg'
WARNING: archive_command is set. Reload the configuration in the
PostgreSQL server using pg_reload_conf() or 'pg_ctl reload'
```

Reload the database server configuration; a restart of the database server is not necessary to reset only the archive_command parameter:

```
[root@localhost tmp]# service ppas11 reload
```

The archive_command in the PostgreSQL server is now set as follows:

The new command string is written to the postgresql.auto.conf file:

```
### Do not edit this file manually!
### It will be overwritten by ALTER SYSTEM command.
archive_command = 'cp %p /opt/backup_edb/acctg/archived_wals/%f'
```

When you invoke the BART INIT command with the -r option, BART rebuilds the backupinfo file using the content of the backup directory for the server specified or for all servers. The BART backupinfo file is initially created by the BACKUP subcommand and contains the backup information used by BART.

!!! Note If the backup was initially created with a user-defined backup name, and then the INIT -r option is invoked to rebuild the backup info file, the user-defined backup name is no longer available. Thus, future references to the backup must use the backup identifier.

The following code sample shows the backupinfo file location in a backup subdirectory:

```
[root@localhost acctg]# pwd
/opt/backup/acctg
[root@localhost acctg]# ls -l
total 4
```

```
drwx----- 2 enterprisedb enterprisedb 38 Oct 26 10:21 1477491569966
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Oct 26 10:19 archived_wals
[root@localhost acctg]# ls -l 1477491569966
total 61144
-rw-rw-r-- 1 enterprisedb enterprisedb 703 Oct 26 10:19 backupinfo
-rw-rw-r-- 1 enterprisedb enterprisedb 62603776 Oct 26 10:19 base.tar
```

The following code sample displays the backupinfo file content:

```
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1477491569966
BACKUP NAME: none
BACKUP PARENT: none
BACKUP LOCATION: /opt/backup/acctg/1477491569966
BACKUP SIZE: 59.70 MB
BACKUP FORMAT: tar
BACKUP TIMEZONE:
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 1
ChkSum File
84b3eeb1e3f7b3e75c2f689570d04f10 base.tar
TABLESPACE(s): 0
START WAL LOCATION: 2/A5000028 (file 0000000100000002000000A5)
STOP WAL LOCATION: 2/A50000C0 (file 00000001000000002000000A5)
CHECKPOINT LOCATION: 2/A5000028
BACKUP METHOD: streamed
BACKUP FROM: master
START TIME: 2016-10-26 10:19:30 EDT
LABEL: pg_basebackup base backup
STOP TIME: 2016-10-26 10:19:30 EDT
TOTAL DURATION: 0 sec(s)
```

The following code sample displays an error message if the backupinfo file is missing when invoking a BART subcommand:

```
-bash-4.2$ bart SHOW-BACKUPS
ERROR: 'backupinfo' file does not exist for backup '1477491569966'
please use 'INIT -r' to generate the file
```

The backupinfo file may be missing if the BACKUP subcommand did not complete successfully.

The following code sample displays information about rebuilding the backupinfo file of the specified backup for database server acctg:

```
-bash-4.1$ bart INIT -s acctg -r -i 1428346620427
INFO: rebuilding BACKUPINFO for backup '1428346620427' of server 'acctg'
INFO: backup checksum: ced59b72a7846ff8fb8afb6922c70649 of base.tar
```

The following code sample displays information about how the backupinfo files of all backups are rebuilt for all database servers:

```
-bash-4.1$ bart INIT -r

INFO: rebuilding BACKUPINFO for backup '1428347191544' of server 'acctg'
```

```
INFO: backup checksum: lac5c6lf055c9l0db3l47832l2f2544f of base.tar
INFO: rebuilding BACKUPINFO for backup 'l428346620427' of server 'acctg'
INFO: backup checksum: ced59b72a7846ff8fb8afb6922c70649 of base.tar
INFO: rebuilding BACKUPINFO for backup 'l428347l98335' of server 'dev'
INFO: backup checksum: a8890dd8ab7e6be5d5bc0f38028a237b of base.tar
INFO: rebuilding BACKUPINFO for backup 'l428346957515' of server 'dev'
INFO: backup checksum: ea62549cf090573625d4adeb7d919700 of base.tar
```

The following code sample displays information about invoking BART INIT with the -r - i option:

```
edb@localhost bin]$ ./bart init -s ppas11 -i 1551778898392 -r
INFO: rebuilding BACKUPINFO for backup '1551778898392' of server
'ppas11'
[edb@localhost bin]$ ls /home/edb/bkup/ppas11/1551778898392/
backupinfo backup_label base base-1.tar base-2.tar base-3.tar
base-4.tar base-5.tar base.tar
```

The following code sample displays information about invoking the BART INIT command with the --no-configure option. You can use the --no-configure option with the INIT subcommand to prevent the archive_command option from being set in the PostgreSQL server.

```
[edb@localhost bin]$ ./bart init -s ppas11 -o --no-configure
[edb@localhost bin]$
### Do not edit this file manually!
### It will be overwritten by the ALTER SYSTEM command.
```

4.1.5 MANAGE

The MANAGE subcommand can be invoked to:

- Evaluate backups, mark their status, and delete obsolete backups based on the retention_policy parameter in the BART configuration file.
- Compress the archived WAL files based on the wal_compression parameter in the BART configuration file.

Syntax:

```
bart MANAGE [ -s { <server_name> | all} ]
[ -l ] [ -d ]
[ -c { keep | nokeep }
-i { <backup_id> | <backup_name> | all } ]
[ -n ]
```

To view detailed information about the MANAGE subcommand and retention policy management, see the EDB Backup and Recovery User Guide. For information about setting the wal_compression parameter, see the EDB Backup and Recovery Installation and Upgrade Guide. These guides are available at the EDB website.

The following table describes the MANAGE options:

Options Description

Description Options -s [<server_name> <server name> is the name of the database server to which the MANAGE actions are to be | all] applied. --server [If all is specified or if the -s option is omitted, actions are applied to all database servers. <server_name> | all] -1 --list-Lists the backups marked as obsolete. obsolete -d Deletes the backups marked as obsolete. This action physically deletes the backup along with its --deletearchived WAL files and any MBM files for incremental backups. obsolete Specify keep to change the backup status to keep to retain the backup indefinitely. -c { keep nokeep } Specify nokeep to change the backup status back to active. You can then re-evaluate and --changepossibly mark the backup as obsolete (according to the retention policy) using the MANAGE status { subcommand. keep | nokeep } The -c option can only be used with the -i option. -i {<backup_id> <backup_name> <backup_id> is a backup identifier and <backup_name> is the user-defined alphanumeric | all } name for the backup. If all is specified, actions are applied to all backups. --backupid The -i option can only be used with the -c option. {<backup_id> <backup_name> | all } Performs the test run and displays the results prior to actually implementing the actions as if the operation was performed, however, no changes are actually made. If you specify -n with the -d option, it displays which backups would be deleted, but does not actually delete the backups. If you specify -n with the -c option, it displays the keep or nokeep action, but does not actually change the backup status. -dry-run If you specify -n alone with no other options or if you specify -n with only the -s option, it displays which active backups would be marked as obsolete, but does not actually change the backup status. In addition, no compression is performed on uncompressed, archived WAL files even if WAL compression is enabled for the database server.

Example

The following code sample performs a dry run for the specified database server displaying which active backups are evaluated as obsolete according to the retention policy, but does not actually change the backup status:

```
-bash-4.2$ bart MANAGE -s acctg -n
INFO: processing server 'acctg', backup '1482770807519'
INFO: processing server 'acctg', backup '1482770803000'
INFO: marking backup '1482770803000' as obsolete
```

```
INFO: 1 WAL file(s) marked obsolete
INFO: processing server 'acctg', backup '1482770735155'
INFO: marking backup '1482770735155' as obsolete
INFO: 2 incremental(s) of backup '1482770735155' will be marked obsolete
INFO: marking incremental backup '1482770780423' as obsolete
INFO: marking incremental backup '1482770763227' as obsolete
INFO: 3 WAL file(s) marked obsolete
INFO: 1 Unused WAL file(s) present
INFO: 2 Unused file(s) (WALs included) present, use 'MANAGE -l' for the
list
```

The following code sample marks active backups as obsolete according to the retention policy for the specified database server:

```
-bash-4.2$ bart MANAGE -s acctg
INFO: processing server 'acctg', backup '1482770807519'
INFO: processing server 'acctg', backup '1482770803000'
INFO: marking backup '1482770803000' as obsolete
INFO: 1 WAL file(s) marked obsolete
INFO: processing server 'acctg', backup '1482770735155'
INFO: marking backup '1482770735155' as obsolete
INFO: 2 incremental(s) of backup '1482770735155' will be marked obsolete
INFO: marking incremental backup '1482770780423' as obsolete
INFO: marking incremental backup '1482770763227' as obsolete
INFO: 3 WAL file(s) marked obsolete
INFO: 1 Unused WAL file(s) present
INFO: 2 Unused file(s) (WALs included) present, use 'MANAGE -l' for the list
```

The following code sample lists backups marked as obsolete for the specified database server:

```
-bash-4.2$ bart MANAGE -s acctg -l
SERVER NAME: acctg
BACKUP ID: 1482770803000
BACKUP STATUS: obsolete
BACKUP TIME: 2016-12-26 11:46:43 EST
BACKUP SIZE: 59.52 MB
WAL FILE(s): 1
WAL FILE: 00000001000000100000055
SERVER NAME: acctg
BACKUP ID: 1482770735155
BACKUP STATUS: obsolete
BACKUP TIME: 2016-12-26 11:45:35 EST
BACKUP SIZE: 59.52 MB
INCREMENTAL BACKUP(s): 2
BACKUP ID: 1482770780423
BACKUP PARENT: 1482770735155
BACKUP STATUS: obsolete
BACKUP TIME: 2016-12-26 11:45:35 EST
BACKUP SIZE: 59.52 MB
BACKUP ID: 1482770763227
BACKUP PARENT: 1482770735155
BACKUP STATUS: obsolete
BACKUP TIME: 2016-12-26 11:45:35 EST
BACKUP SIZE: 59.52 MB
```

The following code sample deletes the obsolete backups for the specified database server:

```
-bash-4.2$ bart MANAGE -s acctg -d
INFO: removing all obsolete backups of server 'acctg'
INFO: removing obsolete backup '1482770803000'
INFO: 1 WAL file(s) will be removed
INFO: removing WAL file '000000010000000100000055'
INFO: removing obsolete backup '1482770735155'
INFO: 3 WAL file(s) will be removed
INFO: 2 incremental(s) of backup '1482770735155' will be removed
INFO: removing obsolete incremental backup '1482770780423'
INFO: removing obsolete incremental backup '1482770763227'
INFO: removing WAL file '00000001000000100000054'
INFO: removing WAL file '000000010000000100000053'
INFO: removing WAL file '00000001000000100000052'
INFO: 8 Unused file(s) will be removed
INFO: removing (unused) file '00000001000000100000056.00000028.backup'
INFO: removing (unused) file '000000010000000100000056'
INFO: removing (unused) file '00000001000000100000055.00000028.backup'
INFO: removing (unused) file '000000010000000100000054.000000028.backup'
INFO: removing (unused) file '00000001000000100000053.00000028.backup'
INFO: removing (unused) file '00000001000000100000052.00000028.backup'
INFO: removing (unused) file '000000010000000100000051'
INFO: removing (unused) file
'0000000100000001510000280000000152000000.mbm'
```

The following code sample changes the specified backup to keep status to retain it indefinitely:

```
-bash-4.2$ bart MANAGE -s acctg -c keep -i 1482770807519
INFO: changing status of backup '1482770807519' of server 'acctg' from
'active' to 'keep'
INFO: 1 WAL file(s) changed
-bash-4.2$ bart SHOW-BACKUPS -s acctg -i 1482770807519 -t
SERVER NAME : acctg
BACKUP ID: 1482770807519
BACKUP NAME: none
BACKUP PARENT : none
BACKUP STATUS: keep
BACKUP TIME: 2016-12-26 11:46:47 EST
BACKUP SIZE: 59.52 MB
WAL(S) SIZE : 16.00 MB
NO. OF WALS: 1
FIRST WAL FILE : 00000001000000100000057
CREATION TIME: 2016-12-26 11:52:47 EST
LAST WAL FILE : 00000001000000100000057
CREATION TIME : 2016-12-26 11:52:47 EST
```

The following code sample resets the specified backup to active status:

```
-bash-4.2$ bart MANAGE -s acctg -c nokeep -i 1482770807519
INFO: changing status of backup '1482770807519' of server 'acctg' from
'keep' to 'active'
INFO: 1 WAL file(s) changed
-bash-4.2$ bart SHOW-BACKUPS -s acctg -i 1482770807519 -t
SERVER NAME : acctg
BACKUP ID: 1482770807519
BACKUP NAME: none
BACKUP PARENT: none
BACKUP STATUS: active
BACKUP TIME : 2016-12-26 11:46:47 EST
BACKUP SIZE: 59.52 MB
WAL(S) SIZE: 16.00 MB
NO. OF WALS: 1
FIRST WAL FILE : 00000001000000100000057
CREATION TIME : 2016-12-26 11:52:47 EST
LAST WAL FILE: 00000001000000100000057
CREATION TIME : 2016-12-26 11:52:47 EST
```

The following code sample uses the enabled wal_compression parameter in the BART configuration file as shown by the following:

```
host = 127.0.0.1
port = 5445
user = enterprisedb
cluster_owner = enterprisedb
allow_incremental_backups = disabled
wal_compression = enabled
description = "Accounting"
```

When the MANAGE subcommand is invoked, the following message is displayed indicating that WAL file compression is performed:

```
-bash-4.2$ bart MANAGE -s acctg
INFO: 4 WAL file(s) compressed
WARNING: 'retention_policy' is not set for server 'acctg'
```

The following code sample shows the archived WAL files in compressed format:

```
-bash-4.2$ pwd
/opt/backup/acctg
-bash-4.2$ ls -l archived_wals
total 160
-rw----- 1 enterprisedb enterprisedb 27089 Dec 26 12:16
00000001000000010000005B.gz
-rw----- 1 enterprisedb enterprisedb 305 Dec 26 12:17
00000001000000010000005C.00000028.backup
-rw----- 1 enterprisedb enterprisedb 27112 Dec 26 12:17
00000001000000010000005C.gz
-rw----- 1 enterprisedb enterprisedb 65995 Dec 26 12:18
000000010000000100000005D.gz
```

```
-rw----- 1 enterprisedb enterprisedb 305 Dec 26 12:18
00000001000000010000005E.00000028.backup
-rw----- 1 enterprisedb enterprisedb 27117 Dec 26 12:18
000000010000000100000005E.gz
```

4.1.6 RESTORE

The **RESTORE** subcommand restores a backup and its archived WAL files for the designated database server to the specified directory location.

Syntax for Restore:

To view detailed information about the RESTORE subcommand, see the EDB Backup and Recovery User Guide available at the EDB website.

If the backup is restored to a different database cluster directory than where the original database cluster resided, then some operations dependent upon the database cluster location may fail. This happens if the supporting service scripts are not updated to reflect the new directory location of restored backup.

For information about the use and modification of service scripts, see the EDB Advanced Server Installation Guide available at the EDB website.

The following table describes the RESTORE options:

Options	Description
<pre>-s <server_name>server <server_name></server_name></server_name></pre>	<pre><server_name> is the name of the database server to be restored.</server_name></pre>
<pre>-prestore-path <restore_path>restore-path <restore_path></restore_path></restore_path></pre>	<pre><restore_path> is the directory path where the backup of the database server is to be restored. The directory must be empty and have the proper ownership and privileges assigned to it.</restore_path></pre>
<pre>-i { <backup_id> <backup_name>}</backup_name></backup_id></pre>	backup_id is the backup identifier of the backup to be used for the restoration and <backup_name> is the user-defined</backup_name>
<pre>backupid { <backup_id> <backup_name>}</backup_name></backup_id></pre>	alphanumeric name for the backup. If the option is omitted, the latest backup is restored by default.

Options	Description
<pre>-r <remote_user@remote_host_address>remote-host <remote_user@remote_host_address></remote_user@remote_host_address></remote_user@remote_host_address></pre>	<pre><remote_user> is the user account on the remote database server host that accepts a passwordless SSH/SCP login connection and is the owner of the directory where the backup is to be restored. <remote_host_address> is the IP address of the remote host to which the backup is to be restored. This option must be specified if the remote_host parameter for this database server is not set in the BART configuration file. For information about the remote_host parameter, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.</remote_host_address></remote_user></pre>
<pre>-w <number_of_workers>workers <number_of_workers></number_of_workers></number_of_workers></pre>	<pre><number_of_workers> is the number of worker processes to run in parallel to stream the modified blocks of an incremental backup to the restore location. If the -w option is omitted, the default is 1 worker process. For example, if four worker processes are specified, four receiver processes on the restore host and four streamer processes on the BART host are used. The output of each streamer process is connected to the input of a receiver process. When the receiver gets to the point where it needs a modified block file, it obtains those modified blocks from its input. With this method, the modified block files are never written to the restore host disk.</number_of_workers></pre>
<pre>-t <timeline_id>target-tli <timeline_id></timeline_id></timeline_id></pre>	<pre><timeline_id> is the integer identifier of the timeline to be used for replaying the archived WAL files for point-in-time recovery.</timeline_id></pre>
<pre>-x <target_xid>target-xid <target_xid></target_xid></target_xid></pre>	<target_xid> is the integer identifier of the transaction ID that determines the transaction up to and including, which point-in-time recovery encompasses.</target_xid>
<pre>-g <target_timestamp>target-timestamp <target_timestamp></target_timestamp></target_timestamp></pre>	<target_timestamp> is the timestamp that determines the point in time up to and including, which point-in-time recovery encompasses.</target_timestamp>
-c copy-wals	Specify this option to copy archived WAL files from the BART backup catalog to <restore_path>/archived_wals directory. The restore_command retrieves the WAL files from <restore_path>/archived_wals for the database server archive recovery. If the -c option is omitted and the copy_wals_during_restore parameter in the BART configuration file is not enabled in a manner applicable to this database server, then the restore_command in the postgresql.conf retrieves the archived WAL files directly from the BART backup catalog. For information about the copy_wals_during_restore parameter, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.</restore_path></restore_path>
disable-checksum	While restoring a backup, specify this option to skip verifying the MD5 or SHA256 checksum files. If you set thechecksum-algorithm=NONE option with the BART scanner or while taking a backup, you also need to specify thedisable checksum option while restoring an incremental backup.

Examples

The following code sample restores a database server(named mktg) to the /opt/restore directory up to timestamp 2015-12-15 10:47:00:

```
-bash-4.1$ bart RESTORE -s mktg -i 1450194208824 -p /opt/restore -t 1 -g '2015-12-15 10:47:00'
INFO: restoring backup '1450194208824' of server 'mktg'
INFO: restoring backup to enterprisedb@192.168.2.24:/opt/restore
INFO: base backup restored
INFO: WAL file(s) will be streamed from the BART host
INFO: writing recovery settings to postgresql.auto.conf file
INFO: archiving is disabled
INFO: tablespace(s) restored
```

The following parameters are set in the postgresql.auto.conf file:

```
restore_command = 'scp -o BatchMode=yes -o PasswordAuthentication=no
enterprisedb@192.168.2.22:/opt/backup/mktg/archived_wals/%f %p'
recovery_target_time = '2015-12-15 10:47:00'
recovery_target_timeline = 1
```

The following is a list of the restored files and subdirectories:

```
[root@localhost restore]# pwd
/opt/restore
[root@localhost restore]# ls -l
total 108
-rw----- 1 enterprisedb enterprisedb 208 Dec 15 10:43 backup_label
drwx----- 6 enterprisedb enterprisedb 4096 Dec 2 10:38 base
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 10:42 dbms_pipe
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 11:00 global
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_clog\
-rw----- 1 enterprisedb enterprisedb 4438 Dec 2 10:38 pg_hba.conf
-rw----- 1 enterprisedb enterprisedb 1636 Nov 10 15:38 pg_ident.conf
drwxr-xr-x 2 enterprisedb enterprisedb 4096 Dec 15 10:42 pg_log
drwx----- 4 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_multixact
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 10:42 pg_notify
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_serial
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_snapshots
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 10:42 pg_stat
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 10:43 pg_stat_tmp
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_subtrans
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 11:00 pg tblspc
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_twophase
-rw----- 1 enterprisedb enterprisedb 4 Nov 10 15:38 PG_VERSION
drwx----- 2 enterprisedb enterprisedb 4096 Dec 15 11:00 pg_xlog
-rw----- 1 enterprisedb enterprisedb 23906 Dec 15 11:00
postgresql.conf
-rw-r--r-- 1 enterprisedb enterprisedb 217 Dec 15 11:00
postgresql.auto.conf
```

Example

The following code sample performs a RESTORE operation with the copy_wals_during_restore parameter enabled to copy the archived WAL files to the local <restore_path>/archived_wals directory:

```
-bash-4.1$ bart RESTORE -s hr -i hr_2017-03-29T13:50 -p
/opt/restore_pg96 -t 1 -g '2017-03-29 14:01:00'
INFO: restoring backup 'hr_2017-03-29T13:50' of server 'hr'
INFO: base backup restored
INFO: copying WAL file(s) to
postgres@192.168.2.24:/opt/restore_pg96/archived_wals
INFO: writing recovery settings to postgresql.auto.conf file
INFO: archiving is disabled
INFO: permissions set on $PGDATA
INFO: restore completed successfully
```

The following parameters are set in the postgresql.auto.conf file:

```
restore_command = 'cp archived_wals/%f %p'
recovery_target_time = '2017-03-29 14:01:00'
recovery_target_timeline = 1
```

The following is a list of the restored files and subdirectories:

```
-bash-4.1$ pwd
/opt/restore_pg96
-bash-4.1$ ls -l
total 128
drwxr-xr-x 2 postgres postgres 4096 Mar 29 14:27 archived_wals
-rw----- 1 postgres postgres 206 Mar 29 13:50 backup_label
drwx----- 5 postgres postgres 4096 Mar 29 12:25 base
drwx----- 2 postgres postgres 4096 Mar 29 14:27 global
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_clog
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_commit_ts
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_dynshmem
-rw---- 1 postgres postgres 4212 Mar 29 13:18 pg_hba.conf
-rw----- 1 postgres postgres 1636 Mar 29 12:25 pg_ident.conf
drwxr-xr-x 2 postgres postgres 4096 Mar 29 13:45 pg_log
drwx----- 4 postgres postgres 4096 Mar 29 12:25 pg_logical
drwx----- 4 postgres postgres 4096 Mar 29 12:25 pg_multixact
drwx----- 2 postgres postgres 4096 Mar 29 13:43 pg_notify
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_replslot
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_serial
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_snapshots
drwx----- 2 postgres postgres 4096 Mar 29 13:43 pg_stat
drwx----- 2 postgres postgres 4096 Mar 29 13:50 pg_stat_tmp
drwx---- 2 postgres postgres 4096 Mar 29 12:25 pg_subtrans
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_tblspc
drwx---- 2 postgres postgres 4096 Mar 29 12:25 pg_twophase
-rw----- 1 postgres postgres 4 Mar 29 12:25 PG_VERSION
drwx----- 3 postgres postgres 4096 Mar 29 14:27 pg_xlog
-rw----- 1 postgres postgres 169 Mar 29 13:24 postgresql.auto.conf
-rw-r--r- 1 postgres postgres 21458 Mar 29 14:27 postgresql.conf
-rw-r--r 1 postgres postgres 118 Mar 29 14:27 postgresql.auto.conf
```

The following code sample displays restoring an incremental backup taken using --checksum-algorithm=NONE option. To restore this incremental backup, you must specify the --disable-checksum option to skip verifying MD5 or SHA256 checksum files.

```
[edb@localhost bin]$ ./bart restore -s ppas11 -i 1593507779811 -p
```

```
/home/edb/RESTORE/ --disable-checksum
INFO: restoring incremental backup '1593507779811' of server 'ppas11'
INFO: base backup restored
INFO: writing recovery.conf file
INFO: WAL file(s) will be streamed from the BART host
INFO: archiving is disabled
INFO: permissions set on $PGDATA
INFO: incremental restore completed successfully
```

4.1.7 SHOW-SERVERS

The SHOW-SERVERS subcommand displays information for the managed database servers listed in the BART configuration file.

Syntax:

```
bart SHOW-SERVERS [ -s { <server_name> | all } ]
```

The following table describes the SHOW-SERVERS option:

Option

Description

Example

The following code sample shows all the database servers managed by BART as returned by the SHOW-SERVERS subcommand:

```
-bash-4.2$ bart SHOW-SERVERS
SERVER NAME : acctg
BACKUP FRIENDLY NAME: acctg_%year-%month-%dayT%hour:%minute
HOST NAME : 127.0.0.1
USER NAME: enterprisedb
PORT: 5444
REMOTE HOST:
RETENTION POLICY: 6 Backups
DISK UTILIZATION: 0.00 bytes
NUMBER OF ARCHIVES: 0
ARCHIVE PATH : /opt/backup/acctg/archived_wals
ARCHIVE COMMAND: (disabled)
XLOG METHOD: fetch
WAL COMPRESSION: disabled
TABLESPACE PATH(s):
INCREMENTAL BACKUP : DISABLED
DESCRIPTION: "Accounting"
SERVER NAME : hr
```

```
BACKUP FRIENDLY NAME: hr_%year-%month-%dayT%hour:%minute
HOST NAME: 192.168.2.24
USER NAME : postgres
PORT: 5432
REMOTE HOST: postgres@192.168.2.24
RETENTION POLICY: 6 Backups
DISK UTILIZATION: 0.00 bytes
NUMBER OF ARCHIVES: 0
ARCHIVE PATH : /opt/backup/hr/archived_wals
ARCHIVE COMMAND : (disabled)
XLOG METHOD: fetch
WAL COMPRESSION: disabled
TABLESPACE PATH(s):
INCREMENTAL BACKUP : DISABLED
DESCRIPTION: "Human Resources"
SERVER NAME : mktg
BACKUP FRIENDLY NAME: mktg_%year-%month-%dayT%hour:%minute
HOST NAME: 192.168.2.24
USER NAME: repuser
PORT: 5444
REMOTE HOST: enterprisedb@192.168.2.24
RETENTION POLICY: 6 Backups
DISK UTILIZATION: 0.00 bytes
NUMBER OF ARCHIVES: 0
ARCHIVE PATH : /opt/backup/mktg/archived_wals
ARCHIVE COMMAND: (disabled)
XLOG METHOD: fetch
WAL COMPRESSION: disabled
TABLESPACE PATH(s):
INCREMENTAL BACKUP : DISABLED\
DESCRIPTION: "Marketing"
```

4.1.8 SHOW-BACKUPS

The SHOW-BACKUPS subcommand displays the backup information for the managed database servers.

Syntax:

```
bart SHOW-BACKUPS [ -s { <server_name> | all } ]
[ -i { <backup_id> | <backup_name> | all } ]
[ -t ]
```

The following table describes the SHOW-BACKUPS options:

Options

Description

```
-s { <server_name> | all }

<server_name> is the name of the database server whose backup information is to be displayed.

--server { <server_name> | If all is specified or if the option is omitted, the backup information for all database servers is displayed.
```

Options

Description

Example

The following code sample shows the backup from database server dev:

```
-bash-4.2$ bart SHOW-BACKUPS -s dev
                         BACKUP ID
SERVER NAME
                                          BACKUP NAME
                                                                     BACKUP PARENT
BACKUP TIME
                         BACKUP SIZE
                                          WAL(s) SIZE
                                                                     WAL FILES STATUS
dev
                         1477579596637
                                          dev_2016-10-27T10:46:36
                                                                     none
2016-10-27 10:46:37 EDT 54.50 MB
                                          96.00 MB
                                                                     6
                                                                               active
```

The following code sample shows detailed information using the -t option:

-bash-4.2\$ bart SHOW-BACKUPS -s dev -i 1477579596637 -t

The following code sample shows a listing of an incremental backup along with its parent backup:

```
-bash-4.2$ bart SHOW-BACKUPS
SERVER NAME
                         BACKUP ID
                                            BACKUP NAME
                                                              BACKUP PARENT
BACKUP TIME
                         BACKUP SIZE
                                            WAL(s) SIZE
                                                              WAL FILES
                                                                               STATUS
acctg
                         1477580293193
                                            acctg_2016-10-27 none
2016-10-27 10:58:13 EDT 16.45 MB
                                            16.00 MB
                                                                               active
acctg 1477580111358 acctg_2016-10-27 none 2016-10-27 10:55:11 EDT 59.71
MB 16.00 MB 1 active
```

The following code sample shows the complete, detailed information of the incremental backup and the parent backup:

```
-bash-4.2$ bart SHOW-BACKUPS -t
SERVER NAME : acctg
BACKUP ID : 1477580293193
BACKUP NAME : none
BACKUP PARENT : acctg_2016-10-27
```

BACKUP STATUS : active

```
BACKUP TIME: 2016-10-27 10:58:13 EDT
BACKUP SIZE: 16.45 MB
WAL(S) SIZE : 16.00 MB
NO. OF WALS: 1
FIRST WAL FILE: 000000010000002000000009
CREATION TIME : 2016-10-27 10:58:13 EDT
LAST WAL FILE: 0000000100000002000000009
CREATION TIME : 2016-10-27 10:58:13 EDT
SERVER NAME : acctg
BACKUP ID: 1477580111358
BACKUP NAME: acctg_2016-10-27
BACKUP PARENT: none
BACKUP STATUS: active
BACKUP TIME: 2016-10-27 10:55:11 EDT
BACKUP SIZE: 59.71 MB
WAL(S) SIZE: 16.00 MB
NO. OF WALS: 1
FIRST WAL FILE: 00000001000000200000000
CREATION TIME: 2016-10-27 10:55:12 EDT
LAST WAL FILE: 00000001000000200000000
CREATION TIME: 2016-10-27 10:55:12 EDT
```

4.1.9 VERIFY-CHKSUM

The VERIFY-CHKSUM subcommand verifies the MD5 checksums of the full backups and any user-defined tablespaces for the specified database server or for all database servers. The checksum is verified by comparing the current checksum of the backup against the checksum when the backup was taken.

!!! Note The VERIFY-CHKSUM subcommand is only used for tar format backups.

Syntax:

```
bart VERIFY-CHKSUM
[ -s { <server_name> | all } ]
[ -i { <backup_id> | <backup_name> | all } ]
```

The following table describes the VERIFY-CHKSUM options:

Options

Description

```
-s { <server_name>
| all }
--server {
<server_name> |
all }
```

<server_name> is the name of the database server whose tar backup checksums are to be verified.

If all is specified or if the -s option is omitted, the checksums of all tar backups are verified for all database servers.

Options

Description

<backup_id> is the backup identifier of a tar format full backup whose checksum is to
be verified along with any user-defined tablespaces. <backup_name> is the userdefined alphanumeric name for the full backup.

If all is specified or if the -i option is omitted, the checksums of all tar backups for the relevant database server are verified.

Example

The following code sample verifies the checksum of all tar format backups of the specified database server:

```
-bash-4.1$ bart VERIFY-CHKSUM -s acctg -i all
SERVER NAME BACKUP ID VERIFY
acctg 1430239348243 OK
acctg 1430232284202 OK
acctg 1430232016284 OK
acctg 1430231949065 OK
acctg 1429821844271 OK
```

4.1.10 Running the BART WAL Scanner

The BART WAL scanner is used to process each WAL file to find and record modified blocks in a corresponding MBM file. As a BART account user, use the BART WAL scanner to invoke the bart-scanner program located in the BART_HOME / bin directory.

For detailed information about the WAL scanner and its usage, see the EDB Backup and Recovery User Guide available at the EDB website.

Syntax:

```
bart-scanner
[ -d ]
[ -c <config_file_path> ]
{ -h |
    -v |
    --daemon |
    -p <mbm_file_path> |
    <wal_file_path> |
    RELOAD |
    STOP
    --checksum-algorithm }
```

When the bart-scanner program is invoked, it forks a separate process for each database server enabled with the allow_incremental_backups parameter.

The WAL scanner processes can run in either the foreground or background depending upon usage of the --daemon

option:

- If the ——daemon option is specified, the WAL scanner process runs in the background. All output messages can be viewed in the BART log file.
- If the --daemon option is omitted, the WAL scanner process runs in the foreground. All output messages can be viewed from the terminal running the program as well as in the BART log file.

The following table describes the VERIFY-CHKSUM options.

Options	Description
-hhelp	Displays general syntax and information on WAL scanner usage.
-vversion	Displays the WAL scanner version information.
-ddebug	Displays debugging output while executing the WAL scanner with any of its options.
<pre>-c <config_file_path>config-path <config_file_path></config_file_path></config_file_path></pre>	Specifies <config_file_path> as the full directory path to a BART configuration file. Use this option if you do not want to use the default BART configuration file <bart_home>/etc/bart.cfg</bart_home></config_file_path>
daemon	Runs the WAL scanner as a background process.
<pre>-p <mbm_file_path> print <mbm_file_path></mbm_file_path></mbm_file_path></pre>	Specifies the full directory path to an MBM file whose content is to be printed. The archived_wals directory as specified in the the archive_path parameter in the bart.cfg file contains the MBM files.
<wal_file_path></wal_file_path>	Specifies the full directory path to a WAL file to be scanned. The archive path directory contains the WAL files. Use it if a WAL file in the archive path is missing its MBM file. This option is to be used for assisting the EnterpriseDB support team for debugging problems that may have been encountered.
RELOAD	Reloads the BART configuration file. The keyword RELOAD is case-insensitive. The RELOAD option is useful if you make changes to the configuration file after the WAL scanner has been started. It will reload the configuration file and adjust the WAL scanners accordingly. For example, if a server section allowing incremental backups is removed from the BART configuration file, then the process attached to that server will stop. Similarly, if a server allowing incremental backups is added, a new WAL scanner process will be launched to scan the WAL files of that server.
STOP	Stops the WAL scanner. The keyword STOP is not case-sensitive.
checksum- algorithm	While invoking the WAL scanner, you can specify one of the following values with thechecksum-algorithm option:checksum-algorithm=MD5 (default) to generate MD5 checksum fileschecksum-algorithm=SHA256 to generate SHA256 checksum fileschecksum-algorithm=NONE to skip generating checksum files.

Example

The following code sample demonstrates starting the WAL scanner to run interactively. The WAL scanner begins scanning existing WAL files in the archive path that have not yet been scanned (that is, there is no corresponding MBM file for the WAL file):

```
INFO: WAL file to be processed: 000000010000000000000001
```

The following code sample is the content of the archive path showing the MBM files created for the WAL files. (The user name and group name of the files have been removed from the example to list the WAL files and MBM files in a more readable manner):

```
[root@localhost archived_wals]# pwd
/opt/backup/acctg/archived_wals
[root@localhost archived_wals]# ls -l
total 81944
-rw----- 1 ... 16777216 Dec 20 09:10 0000000100000000000000ED
-rw----- 1 ... 16777216 Dec 20 09:06 000000010000000000000EE
-rw----- 1 ... 16777216 Dec 20 09:11 000000010000000000000EF
-rw----- 1 ... 16777216 Dec 20 09:16 000000010000000000000F1
                          Dec 20 09:16 000000010000000000000000000001.00000028.backup
-rw----- 1 ... 305
-rw-rw-r-- 1 ... 161
                          Dec 20 09:18
0000000100000000ED0000280000000EE000000.mbm
-rw-rw-r-- 1 ... 161
                          Dec 20 09:18
0000000100000000EE0000280000000EF000000.mbm
-rw-rw-r-- 1 ... 161
                          Dec 20 09:18
0000000100000000EF0000280000000F0000000.mbm
-rw-rw-r-- 1 ... 161
                          Dec 20 09:18
000000010000000F00000280000000F1000000.mbm
-rw-rw-r-- 1 ... 161
                          Dec 20 09:18
\tt 0000000100000000F10000280000000F2000000.mbm
```

To stop the interactively running WAL scanner, either enter ctrl-C at the terminal running the WAL scanner or invoke the bart-scanner program from another terminal with the STOP option:

```
-bash-4.2$ bart-scanner STOP
-bash-4.2$
```

The terminal on which the WAL scanner was running interactively appears as follows after it has been stopped:

The following code sample demonstrates invoking the WAL scanner to run as a background process with the --daemon option:

```
-bash-4.2$ bart-scanner --daemon
-bash-4.2$
```

The WAL scanner runs as a background process. There is also a separate background process for each database server that has been enabled for WAL scanning with the allow_incremental_backups parameter in the BART configuration file:

```
-bash-4.2$ ps -ef | grep bart
enterpr+ 4340 1 0 09:48 ? 00:00:00 bart-scanner --daemon
enterpr+ 4341 4340 0 09:48 ? 00:00:00 bart-scanner --daemon
enterpr+ 4415 3673 0 09:50 pts/0 00:00:00 grep --color=auto bart
```

To stop the WAL scanner processes, invoke the WAL scanner with the STOP option:

```
-bash-4.2$ bart-scanner STOP
-bash-4.2$
```

The following command demonstrates scanning an individual WAL file:

```
-bash-4.2$ bart-scanner /opt/backup/acctg/archived_wals/0000000100000000000000FF -bash-4.2$
```

To print the content of an MBM file for assisting the EnterpriseDB support team for debugging problems that may have been encountered, use the -p option to specify the file as shown in the following code sample:

```
-bash-4.2$ bart-scanner -p
opt/backup/acctg/archived_wals/000000010000000FF0000280000000100000000.mbm
Header:
Version: 1.0:90500:1.2.0
Scan Start: 2016-12-20 10:02:11 EST, Scan End: 2016-12-20 10:02:11 EST, Diff: 0
Start LSN: ff000028, End LSN: 100000000, TLI: 1
flags: 0, Check Sum: f9cfe66ae2569894d6746b61503a767d
Path: base/14845/16384
NodeTag: BLOCK_CHANGE
Relation: relPath base/14845/16384, isTSNode 0, Blocks
*.....
First modified block: 0
Total modified blocks: 1
Path: base/14845/16391
NodeTag: BLOCK_CHANGE
Relation: relPath base/14845/16391, isTSNode 0, Blocks
*.....
First modified block: 0
Total modified blocks: 1
```

4.2 Additional Examples

This section lists examples of the following BART operations.

- Restoring a database cluster with tablespaces.
- Restoring an incremental backup.
- · Managing backups.
- Managing incremental backups.

Restoring a Database Cluster with Tablespaces

The following code sample illustrates taking a backup and restoring a database cluster on a remote host containing tablespaces. For detailed information regarding using tablespaces, see the EDB Backup and Recovery User Guide available at the EDB website.

On an Advanced Server database running on a remote host, the following tablespaces are created for use by two tables:

```
edb=# CREATE TABLESPACE tblspc_1 LOCATION '/mnt/tablespace_1';
edb=# CREATE TABLESPACE tblspc_2 LOCATION '/mnt/tablespace_2';
CREATE TABLESPACE
edb=# \db
           List of tablespaces
Name
         Owner | Location
pg_default | enterprisedb
pg_global | enterprisedb
tblspc_1 | enterprisedb | /mnt/tablespace_1
tblspc_2
         | enterprisedb | /mnt/tablespace_2
(4 rows)
edb=# CREATE TABLE tbl_tblspc_1 (c1 TEXT) TABLESPACE tblspc_1;
CREATE TABLE
edb=# CREATE TABLE tbl tblspc 2 (c1 TEXT) TABLESPACE tblspc 2;
CREATE TABLE
edb=# \d tbl_tblspc_1
Table "enterprisedb.tbl_tblspc_1"
Column | Type | Modifiers
-----
c1 | text |
Tablespace: "tblspc_1"
edb=# \d tbl_tblspc_2
Table "enterprisedb.tbl_tblspc_2"
Column | Type | Modifiers
_____
c1
   | text |
Tablespace: "tblspc_2"
```

The following code sample shows the OIDs assigned to the tablespaces and the symbolic links to the tablespace directories:

```
-bash-4.1$ pwd
/opt/PostgresPlus/9.6AS/data/pg_tblspc
-bash-4.1$ ls -l
total 0
lrwxrwxrwx 1 enterprisedb enterprisedb 17 Nov 16 16:17 16587 ->/mnt/tablespace_1
lrwxrwxrwx 1 enterprisedb enterprisedb 17 Nov 16 16:17 16588 ->/mnt/tablespace_2
```

The BART configuration file contains the following settings. Note that the tablespace_path parameter does not have to be set at this point.

```
[BART]
bart_host= enterprisedb@192.168.2.22
```

```
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log

[ACCTG]
host = 192.168.2.24
port = 5444
user = repuser
cluster_owner = enterprisedb
remote_host = enterprisedb@192.168.2.24
tablespace_path =
description = "Accounting"
```

After the necessary configuration steps are performed to ensure BART manages the remote database server, a full backup is taken as shown in the following code sample:

```
-bash-4.1$ bart BACKUP -s acctg
INFO: creating backup for server 'acctg'
INFO: backup identifier: '1447709811516'
54521/54521 kB (100%), 3/3 tablespaces
INFO: backup completed successfully
INFO: backup checksum: 594f69fe7d26af991d4173d3823e174f of 16587.tar
INFO: backup checksum: 7a5507567729a21c98a15c948ff6c015 of base.tar
INFO: backup checksum: ae8c62604c409635c9d9e82b29cc0399 of 16588.tar
INFO:
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1447709811516
BACKUP NAME: none
BACKUP LOCATION: /opt/backup/acctg/1447709811516
BACKUP SIZE: 53.25 MB
BACKUP FORMAT: tar
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 3
ChkSum File
594f69fe7d26af991d4173d3823e174f 16587.tar
7a5507567729a21c98a15c948ff6c015 base.tar
ae8c62604c409635c9d9e82b29cc0399 16588.tar
TABLESPACE(s): 2
Oid Name Location
16587 tblspc_1 /mnt/tablespace_1
16588 tblspc_2 /mnt/tablespace_2
START WAL LOCATION: 00000001000000000000000000000
BACKUP METHOD: streamed
BACKUP FROM: master
START TIME: 2015-11-16 16:36:51 EST
STOP TIME: 2015-11-16 16:36:52 EST
TOTAL DURATION: 1 sec(s)
```

Note that in the output from the preceding example, checksums are generated for the tablespaces as well as the full backup.

Within the backup subdirectory 1447709811516 of the BART backup catalog, the tablespace data is stored with file names 16587.tar.gz and 16588.tar.gz as shown below:

```
-bash-4.1$ pwd
/opt/backup/acctg
-bash-4.1$ ls -l
total 8
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 16:36 1447709811516
drwx---- 2 enterprisedb enterprisedb 4096 Nov 16 16:43 archived_wals
-bash-4.1$ ls -l 1447709811516
total 54536
-rw-rw-r-- 1 enterprisedb enterprisedb 19968 Nov 16 16:36 16587.tar
-rw-rw-r-- 1 enterprisedb enterprisedb 19968 Nov 16 16:36 16588.tar
-rw-rw-r-- 1 enterprisedb enterprisedb 949 Nov 16 17:05 backupinfo
-rw-rw-r-- 1 enterprisedb enterprisedb 55792640 Nov 16 16:36 base.tar
```

When you are ready to restore the backup, in addition to creating the directory to which the main database cluster is to be restored, you must prepare the directories to which the tablespaces are to be restored.

On the remote host, directories <code>/opt/restore_tblspc_1</code> and <code>/opt/restore_tblspc_2</code> are created and assigned the proper ownership and permissions as shown by the following example. The main database cluster is to be restored to <code>/opt/restore</code>.

```
[root@localhost opt]# mkdir restore_tblspc_1
[root@localhost opt]# chown enterprisedb restore_tblspc_1
[root@localhost opt]# chgrp enterprisedb restore_tblspc_1
[root@localhost opt]# chmod 700 restore_tblspc_1
[root@localhost opt]# mkdir restore_tblspc_2
[root@localhost opt]# chown enterprisedb restore_tblspc_2
[root@localhost opt]# chgrp enterprisedb restore_tblspc_2
[root@localhost opt]# chmod 700 restore_tblspc_2
[root@localhost opt]# ls -l
total 20
drwxr-xr-x 3 root daemon 4096 Nov 10 15:38 PostgresPlus
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:40 restore
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:40
restore_tblspc_1
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:41
restore_tblspc_2
drwxr-xr-x. 2 root root 4096 Nov 22 2013 rh
```

Set the tablespace_path parameter in the BART configuration file to specify the tablespace directories. The remote host user and IP address are specified by the remote_host configuration parameter.

```
host = 192.168.2.24
port = 5444
user = repuser
cluster_owner = enterprisedb
remote_host = enterprisedb@192.168.2.24
tablespace_path =
16587=/opt/restore_tblspc_1;16588=/opt/restore_tblspc_2
```

```
description = "Accounting"
```

The following code sample demonstrates invoking the RESTORE subcommand:

```
-bash-4.1$ bart RESTORE -s acctg -i 1447709811516 -p /opt/restore
INFO: restoring backup '1447709811516' of server 'acctg'
INFO: restoring backup to enterprisedb@192.168.2.24:/opt/restore
INFO: base backup restored
INFO: archiving is disabled
INFO: tablespace(s) restored
```

The following code sample shows the restored full backup (including the restored tablespaces):

```
bash-4.1$ pwd
/opt
-bash-4.1$ ls -l restore
total 104
-rw----- 1 enterprisedb enterprisedb 206 Nov 16 16:36 backup_label.old
drwx----- 6 enterprisedb enterprisedb 4096 Nov 10 15:38 base
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:46 global
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_clog
-rw----- 1 enterprisedb enterprisedb 4438 Nov 10 16:23 pg_hba.conf
-rw----- 1 enterprisedb enterprisedb 1636 Nov 10 15:38 pg_ident.conf
drwxr-xr-x 2 enterprisedb enterprisedb 4096 Nov 16 17:45 pg_log
drwx----- 4 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_multixact
drwx---- 2 enterprisedb enterprisedb 4096 Nov 16 17:45 pg_notify
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_serial
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_snapshots
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:47 pg_stat
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:47 pg_stat_tmp
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_subtrans
drwx----- 2 enterprisedb enterprisedb 4096 Nov 16 17:42 pg_tblspc
drwx----- 2 enterprisedb enterprisedb 4096 Nov 10 15:38 pg_twophase
-rw----- 1 enterprisedb enterprisedb 4 Nov 10 15:38 PG_VERSION
drwx----- 3 enterprisedb enterprisedb 4096 Nov 16 17:47 pg_xlog
-rw----- 1 enterprisedb enterprisedb 23906 Nov 16 17:42 postgresql.conf
-rw----- 1 enterprisedb enterprisedb 61 Nov 16 17:45 postmaster.opts
-bash-4.1$
-bash-4.1$ ls -l restore_tblspc_1
drwx----- 3 enterprisedb enterprisedb 4096 Nov 16 16:18
PG_9.6_201306121
-bash-4.1$ ls -l restore_tblspc_2
total 4
drwx----- 3 enterprisedb enterprisedb 4096 Nov 16 16:18
PG_9.6_201306121
```

The symbolic links in the pg_tblspc subdirectory point to the restored directory location:

```
bash-4.1$ pwd
/opt/restore/pg_tblspc
-bash-4.1$ ls -l
total 0
lrwxrwxrwx 1 enterprisedb enterprisedb 21 Nov 16 17:42 16587 ->
/opt/restore_tblspc_1
```

```
lrwxrwxrwx 1 enterprisedb enterprisedb 21 Nov 16 17:42 16588 ->
/opt/restore_tblspc_2
```

psql queries also show the restored tablespaces:

Restoring an Incremental Backup

Restoring an incremental backup may require additional setup steps depending upon the host on which the incremental backup is to be restored. For more information, see the EDB Backup and Recovery User Guide available at the EDB website.

This section provides an example of creating backup chains and then restoring an incremental backup.

Creating a Backup Chain

A backup chain is the set of backups consisting of a full backup and all of its successive incremental backups. Tracing back on the parent backups of all incremental backups in the chain eventually leads back to that single, full backup.

In the following example, the allow_incremental_backups parameter is set to enabled in the BART configuration file to permit incremental backups on the listed database server:

```
[BART]
bart_host= enterprisedb@192.168.2.27
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log

[ACCTG]
host = 127.0.0.1
port = 5445
user = enterprisedb
cluster_owner = enterprisedb
allow_incremental_backups = enabled
description = "Accounting"
```

After the database server has been started with WAL archiving enabled to the BART backup catalog, the WAL scanner is started:

```
-bash-4.2$ bart-scanner --daemon
```

First, a full backup is taken.

```
-bash-4.2$ bart BACKUP -s acctg --backup-name full_1
INFO: creating backup for server 'acctg'
INFO: backup identifier: '1490649204327'\
63364/63364 kB (100%), 1/1 tablespace
INFO: backup completed successfully
INFO: backup checksum: aae27d4a7c09dffc82f423221154db7e of base.tar
INFO:
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1490649204327
BACKUP NAME: full_1
BACKUP PARENT: none
BACKUP LOCATION: /opt/backup/acctg/1490649204327
BACKUP SIZE: 61.88 MB
BACKUP FORMAT: tar
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 1
ChkSum File
aae27d4a7c09dffc82f423221154db7e base.tar
TABLESPACE(s): 0
BACKUP METHOD: streamed
BACKUP FROM: master
START TIME: 2017-03-27 17:13:24 EDT
STOP TIME: 2017-03-27 17:13:25 EDT
TOTAL DURATION: 1 sec(s)
```

A series of incremental backups are taken. The first incremental backup specifies the full backup as the parent. Each successive incremental backup then uses the preceding incremental backup as its parent.

```
-bash-4.2$ bart BACKUP -s acctg -F p --parent full_1 --backup-name
incr_1-a
INFO: creating incremental backup for server 'acctg'
INFO: checking mbm files /opt/backup/acctg/archived wals
INFO: new backup identifier generated 1490649255649
INFO: reading directory /opt/backup/acctg/archived_wals
INFO: all files processed
NOTICE: pg_stop_backup complete, all required WAL segments have been
INFO: incremental backup completed successfully
INFO:
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1490649255649
BACKUP NAME: incr_1-a
BACKUP PARENT: 1490649204327
BACKUP LOCATION: /opt/backup/acctg/1490649255649
BACKUP SIZE: 16.56 MB
BACKUP FORMAT: plain
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 0
TABLESPACE(s): 0
START WAL LOCATION: 0000000100000000000000010
```

```
STOP WAL LOCATION: 0000000100000000000000010
BACKUP METHOD: pg_start_backup
BACKUP FROM: master
START TIME: 2017-03-27 17:14:15 EDT
STOP TIME: 2017-03-27 17:14:16 EDT
TOTAL DURATION: 1 sec(s)
-bash-4.2$ bart BACKUP -s acctg -F p --parent incr_1-a --backup-name
incr_1-b
INFO: creating incremental backup for server 'acctg'
INFO: checking mbm files /opt/backup/acctg/archived wals
INFO: new backup identifier generated 1490649336845
INFO: reading directory /opt/backup/acctg/archived_wals
INFO: all files processed
NOTICE: pg_stop_backup complete, all required WAL segments have been
archived
INFO: incremental backup completed successfully
-bash-4.2$ bart BACKUP -s acctg -F p --parent incr_1-b --backup-name
incr_1-c
INFO: creating incremental backup for server 'acctg'
INFO: checking mbm files /opt/backup/acctg/archived_wals
INFO: new backup identifier generated 1490649414316
INFO: reading directory /opt/backup/acctg/archived_wals
INFO: all files processed
NOTICE: pg_stop_backup complete, all required WAL segments have been
archived
INFO: incremental backup completed successfully
```

The following output of the SHOW-BACKUPS subcommand lists the backup chain, which are backups full_1, incr_1-a, incr_1-b, and incr_1-c.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                           BACKUP NAME
                                         BACKUP PARENT BACKUP TIME ...
          1490649414316 incr_1-c
                                         incr_1-b
                                                       2017-03-27 17:16:55 ...
acctg
           1490649336845 incr_1-b
                                         incr_1-a
                                                       2017-03-27 17:15:37 ...
acctg
           1490649255649 incr_1-a
                                         full_1
                                                       2017-03-27 17:14:16 ...
acctg
acctg
           1490649204327 full_1
                                         none
                                                       2017-03-27 17:13:25 ...
```

For the full backup full_1, the BACKUP PARENT field contains none. For each incremental backup, the BACKUP PARENT field contains the backup identifier or name of its parent backup.

A second backup chain is created in the same manner with the BACKUP subcommand. The following example shows the addition of the resulting, second backup chain consisting of full backup full_2 and incremental backups incr_2-a and incr_2-b.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                                      BACKUP PARENT BACKUP TIME ...
                         BACKUP NAME
           1490649605607 incr_2-b
                                      incr_2-a
                                                      2017-03-27 17:20:06 ...
acctg
           1490649587702 incr_2-a
                                       full_2
                                                      2017-03-27 17:19:48 ...
acctg
           1490649528633 full_2
                                                      2017-03-27 17:18:49 ...
acctg
                                      none
```

acctg	1490649414316 incr_1-c	incr_1-b	2017-03-27 17:16:55
acctg	1490649336845 incr_1-b	incr_1-a	2017-03-27 17:15:37
acctg	1490649255649 incr_1-a	full_1	2017-03-27 17:14:16
acctg	1490649204327 full_1	none	2017-03-27 17:13:25

The following additional incremental backups starting with $incr_1-b-1$, which designates $incr_1-b$ as the parent, results in the forking from that backup into a second line of backups in the chain consisting of full_1, $incr_1-a$, $incr_1-b-1$, $incr_1-b-2$, and $incr_1-b-3$ as shown in the following list:

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                           BACKUP NAME
                                          BACKUP PARENT
                                                           BACKUP TIME
                                          incr_1-b-2
             1490649791430 incr_1-b-3
                                                           2017-03-27 17:23:12 ...
acctg
                                          incr_1-b-1
                                                           2017-03-27 17:22:44 ...
             1490649763929 incr_1-b-2
acctg
acctg
             1490649731672 incr_1-b-1
                                          incr_1-b
                                                           2017-03-27 17:22:12 ...
             1490649605607 incr_2-b
                                          incr_2-a
                                                           2017-03-27 17:20:06 ...
acctg
acctg
             1490649587702 incr_2-a
                                          full_2
                                                           2017-03-27 17:19:48 ...
                                                           2017-03-27 17:18:49 ...
             1490649528633 full_2
acctg
                                          none
                                                           2017-03-27 17:16:55 ...
             1490649414316 incr_1-c
                                          incr_1-b
acctg
acctg
             1490649336845 incr_1-b
                                          incr_1-a
                                                           2017-03-27 17:15:37 ...
                                                           2017-03-27 17:14:16 ...
acctg
             1490649255649 incr_1-a
                                          full_1
             1490649204327 full_1
                                                           2017-03-27 17:13:25 ...
acctg
                                          none
```

Restoring an Incremental Backup

Restoring an incremental backup is done with the **RESTORE** subcommand in the same manner as for restoring a full backup. Specify the backup identifier or backup name of the incremental backup to be restored as shown in the following example.

```
-bash-4.2$ bart RESTORE -s acctg -p /opt/restore -i incr_1-b
INFO: restoring incremental backup 'incr_1-b' of server 'acctg'
INFO: base backup restored
INFO: archiving is disabled
INFO: permissions set on $PGDATA
INFO: incremental restore completed successfully
```

Restoring incremental backup $incr_1-b$ as shown by the preceding example results in the restoration of full backup $full_1$, then incremental backups $incr_1-a$ and finally, $incr_1-b$.

Managing Backups

This section illustrates evaluating, marking, and deleting backups using the MANAGE subcommand using a redundancy retention policy and a recovery window retention policy. For detailed information about the MANAGE subcommand, see the EDB Backup and Recovery User Guide available at the EDB website.

Using a Redundancy Retention Policy

The following code sample uses a redundancy retention policy to evaluate, mark, and delete backups as shown by the following server configuration:

```
[ACCTG]
host = 127.0.0.1
port = 5444
```

```
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 BACKUPS
description = "Accounting"
```

The following list is the set of backups. Note that the last backup in the list has been marked as keep.

-bash-4.1\$ b	art SHOW-BACKUPS	-s acctg		
SERVER NAME	BACKUP ID	BACKUP TIME	BACKUP SIZE	WAL(s) SIZE
WAL FILES	STATUS			
acctg	1428768344061	2015-04-11 12:05:46 E	DT 5.72 MB	48.00 MB
3	active			
acctg	1428684537299	2015-04-10 12:49:00 E	DT 5.72 MB	272.00 MB
17	active			
acctg	1428589759899	2015-04-09 10:29:27 E	DT 5.65 MB	96.00 MB
6	active			
acctg	1428502049836	2015-04-08 10:07:30 E	DT 55.25 MB	96.00 MB
6	active			
acctg	1428422324880	2015-04-07 11:58:45 E	DT 54.53 MB	32.00 MB
2	active			
acctg	1428355371389	2015-04-06 17:22:53 E	DT 5.71 MB	16.00 MB
1	keep			

Invoke the MANAGE subcommand with the -n option to perform a dry run to observe which active backups would be changed to obsolete according to the retention policy as shown in the following code sample:

```
-bash-4.1$ bart MANAGE -s acctg -n

INFO: processing server 'acctg', backup '1428768344061'

INFO: processing server 'acctg', backup '1428684537299'

INFO: processing server 'acctg', backup '1428589759899'

INFO: processing server 'acctg', backup '1428502049836'

INFO: marking backup '1428502049836' as obsolete

INFO: 6 WAL file(s) marked obsolete

INFO: processing server 'acctg', backup '1428422324880'

INFO: marking backup '1428422324880' as obsolete

INFO: 2 WAL file(s) marked obsolete

INFO: processing server 'acctg', backup '1428355371389'
```

The dry run shows that backups 1428502049836 and 1428422324880 would be marked as obsolete.

!!! Note A dry run does not change the backup status. The two backups that would be considered obsolete are still marked as active:

-bash-4.1\$ b	art SHOW-BACKUPS	s -s acctg		
SERVER NAME	BACKUP ID	BACKUP TIME	BACKUP SIZE	WAL(s) SIZE
WAL FILES	STATUS			
acctg	1428768344061	2015-04-11 12:05:46 EDT	5.72 MB	48.00 MB
3	active			
acctg	1428684537299	2015-04-10 12:49:00 EDT	5.72 MB	272.00 MB
17	active			
acctg	1428589759899	2015-04-09 10:29:27 EDT	5.65 MB	96.00 MB
6	active			
acctg	1428502049836	2015-04-08 10:07:30 EDT	55.25 MB	96.00 MB
6	active			
acctg	1428422324880	2015-04-07 11:58:45 EDT	54.53 MB	32.00 MB

```
2 active
acctg 1428355371389 2015-04-06 17:22:53 EDT 5.71 MB 16.00 MB
1 keep
```

Invoke the MANAGE subcommand omitting the -n option to change and mark the status of the backups as obsolete:

```
-bash-4.1$ bart MANAGE -s acctg
INFO: processing server 'acctg', backup '1428768344061'
INFO: processing server 'acctg', backup '1428684537299'
INFO: processing server 'acctg', backup '1428589759899'
INFO: processing server 'acctg', backup '1428502049836'
INFO: marking backup '1428502049836' as obsolete
INFO: 6 WAL file(s) marked obsolete
INFO: processing server 'acctg', backup '1428422324880'
INFO: marking backup '1428422324880' as obsolete
INFO: 2 WAL file(s) marked obsolete
INFO: processing server 'acctg', backup '1428355371389'
```

The obsolete backups can be observed in a number of ways. Use the MANAGE subcommand with the -l option to list the obsolete backups:

```
-bash-4.1$ bart MANAGE -s acctg -l
INFO: 6 WAL file(s) will be removed
SERVER NAME: acctg
BACKUP ID: 1428502049836
BACKUP STATUS: obsolete
BACKUP TIME: 2015-04-08 10:07:30 EDT
BACKUP SIZE: 55.25 MB
WAL FILE(s): 6
WAL FILE: 00000001000000100000003
WAL FILE: 000000010000001000000002
WAL FILE: 00000001000000100000001
WAL FILE: 00000001000000100000000
WAL FILE: 000000010000000000000000
WAL FILE: 0000000100000000000000000000
INFO: 2 WAL file(s) will be removed
SERVER NAME: acctg
BACKUP ID: 1428422324880
BACKUP STATUS: obsolete
BACKUP TIME: 2015-04-07 11:58:45 EDT
BACKUP SIZE: 54.53 MB
WAL FILE(s): 2
WAL FILE: 000000010000000000000001
```

The STATUS field of the SHOW-BACKUPS subcommand displays the current status:

```
-bash-4.1$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                            BACKUP TIME
                                                     BACKUP SIZE
                                                                   WAL(s) SIZE
WAL FILES
             STATUS
             1428768344061 2015-04-11 12:05:46 EDT 5.72 MB
acctg
                                                                   48.00 MB
3
             active
             1428684537299 2015-04-10 12:49:00 EDT
                                                                    272.00 MB
acctg
                                                     5.72 MB
17
             active
             1428589759899 2015-04-09 10:29:27 EDT 5.65 MB
acctg
                                                                   96.00 MB
```

6	active					
acctg	1428502049836	2015-04-08	10:07:30	EDT	55.25 MB	96.00 MB
6	obsolete					
acctg	1428422324880	2015-04-07	11:58:45	EDT	54.53 MB	32.00 MB
2	obsolete					
acctg	1428355371389	2015-04-06	17:22:53	EDT	5.71 MB	16.00 MB
1	keep					

The details of an individual backup can be displayed using the SHOW-BACKUPS subcommand with the -t option. Note the status in the BACKUP STATUS field.

```
-bash-4.1$ bart SHOW-BACKUPS -s acctg -i 1428502049836 -t
SERVER NAME : acctg
BACKUP ID: 1428502049836
```

BACKUP NAME : none

BACKUP STATUS : obsolete

BACKUP TIME: 2015-04-08 10:07:30 EDT

BACKUP SIZE: 55.25 MB WAL(S) SIZE: 96.00 MB

NO. OF WALS: 6

CREATION TIME: 2015-04-08 10:07:30 EDT LAST WAL FILE: 00000001000000100000003 CREATION TIME: 2015-04-09 10:25:46 EDT

Use the MANAGE subcommand with the -d option to physically delete the obsolete backups including the unneeded WAL files.

```
-bash-4.1$ bart MANAGE -s acctg -d
INFO: removing all obsolete backups of server 'acctg'
INFO: removing obsolete backup '1428502049836'
INFO: 6 WAL file(s) will be removed
INFO: removing WAL file '000000010000000100000003'
INFO: removing WAL file '00000001000000100000002'
INFO: removing WAL file '000000010000000100000001'
INFO: removing WAL file '000000010000001000000000'
INFO: removing WAL file '00000001000000000000000001'
INFO: removing obsolete backup '1428422324880'
INFO: 2 WAL file(s) will be removed
INFO: removing WAL file '0000000100000000000000001'
```

The SHOW-BACKUPS subcommand now displays the remaining backups marked as active or keep:

-bash-4.1\$ bar	t SHOW-BACKUPS -s	acctg		
SERVER NAME	BACKUP ID	BACKUP TIME	BACKUP SIZE	WAL(s)
SIZE				
WAL FILES	STATUS			
acctg	1428768344061	2015-04-11 12:05:46 EDT	5.72 MB	48.00 MB
3	active			
acctg	1428684537299	2015-04-10 12:49:00 EDT	5.72 MB	272.00 MB
17	active			
acctg	1428589759899	2015-04-09 10:29:27 EDT	5.65 MB	96.00 MB
6	active			

acctg	1428355371389	2015-04-06 17:22:53 EDT	5.71 MB	16.00 MB
1	keep			

Using a Recovery Window Retention Policy

This section illustrates the evaluation, marking, and deletion of backup using a recovery window retention policy. To use the recovery window retention policy, set the retention_policy parameter to the desired length of time for the recovery window.

This section provides examples of the following:

- How to view the calculated recovery window.
- How to evaluate, mark, and delete backup using a recovery window retention policy.

Viewing the Recovery Window

You can view the actual, calculated recovery window by invoking any of the following subcommands:

- MANAGE subcommand in debug mode (along with the -n option).
- SHOW-SERVERS subcommand.

Viewing the Recovery Window Using the Manage Subcommand

When invoking BART in debug mode with the MANAGE subcommand and the -n option, the length of the recovery window is calculated based on the retention_policy setting and the current date/time.

For example, using the following retention_policy settings:

```
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 DAYS
backup-name = acctg_%year-%month-%dayT%hour:%minute:%second
description = "Accounting"
[DEV]
host = 127.0.0.1
port = 5445
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 WEEKS
description = "Development"
[HR]
host = 127.0.0.1
port = 5432
user = postgres
```

```
retention_policy = 3 MONTHS
description = "Human Resources"
```

If the MANAGE subcommand is invoked in debug mode along with the -n option on 2015-04-17, the following results are displayed:

```
-bash-4.1$ bart -d MANAGE -n

DEBUG: Server: acctg, Now: 2015-04-17 16:34:03 EDT, RetentionWindow:

259200 (secs) ==> 72 hour(s)

DEBUG: Server: dev, Now: 2015-04-17 16:34:03 EDT, RetentionWindow:

1814400 (secs) ==> 504 hour(s)

DEBUG: Server: hr, Now: 2015-04-17 16:34:03 EDT, RetentionWindow:

7776000 (secs) ==> 2160 hour(s)
```

For server acctg, 72 hours translates to a recovery window of 3 days.

For server dev, 504 hours translates to a recovery window of 21 days (3 weeks).

For server hr, 2160 hours translates to a recovery window of 90 days (3 months).

For a setting of <max_number> MONTHS, the calculated total number of days for the recovery window is dependent upon the actual number of days in the preceding months from the current date/time. Thus, <max_number> MONTHS is not always exactly equivalent to <max_number> x 30 DAYS. For example, if the current date/time is in the month of March, a 1-month recovery window would be equivalent to only 28 days because the preceding month is February. Thus, for a current date of March 31, a 1-month recovery window would start on March 3. However, the typical result is that the day of the month of the starting recovery window boundary will be the same day of the month of when the MANAGE subcommand is invoked.

Viewing the Recovery Window Using the Show-Servers Subcommand

This section provides an example of viewing the recovery window using the SHOW-SERVERS subcommand; the RETENTION POLICY field displays the start of the recovery window.

In the following code sample, the recovery window retention policy setting considers the backups taken within a 3-day recovery window as the active backups.

```
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 DAYS
description = "Accounting"
```

The start of the 3-day recovery window displayed in the RETENTION POLICY field is 2015-04-07 14:57:36 EDT when the SHOW-SERVERS subcommand is invoked on 2015-04-10.

At this current point in time, backups taken on or after 2015-04-07 14:57:36 EDT would be considered active. Backups taken prior to 2015-04-07 14:57:36 EDT would be considered obsolete except for backups marked as keep.

```
-bash-4.1$ date
Fri Apr 10 14:57:33 EDT 2015
-bash-4.1$
-bash-4.1$ bart SHOW-SERVERS -s acctg
```

SERVER NAME : acctg
HOST NAME : 127.0.0.1
USER NAME : enterprisedb

PORT : 5444

REMOTE HOST :

RETENTION POLICY : 2015-04-07 14:57:36 EDT

DISK UTILIZATION : 824.77 MB

NUMBER OF ARCHIVES: 37

ARCHIVE PATH : /opt/backup/acctg/archived_wals

ARCHIVE COMMAND : cp %p /opt/backup/acctg/archived_wals/%f

XLOG METHOD : fetch
WAL COMPRESSION : disabled

TABLESPACE PATH(s):

DESCRIPTION: "Accounting"

In the following code sample, the recovery window retention policy setting considers the backups taken within a 3-week recovery window as the active backups.

```
[DEV]
host = 127.0.0.1
port = 5445
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 WEEKS
description = "Development"
```

The start of the 3-week recovery window displayed in the RETENTION POLICY field is 2015-03-20 14:59:42 EDT when the SHOW-SERVERS subcommand is invoked on 2015-04-10.

At this current point in time, backups taken on or after 2015-03-20 14:59:42 EDT would be considered active. Backups taken prior to 2015-03-20 14:59:42 EDT would be considered obsolete except for backups marked as keep.

```
-bash-4.1$ date
Fri Apr 10 14:59:39 EDT 2015
-bash-4.1$
-bash-4.1$ bart SHOW-SERVERS -s dev
SERVER NAME : dev
HOST NAME: 127.0.0.1
USER NAME: enterprisedb
PORT: 5445
REMOTE HOST:
RETENTION POLICY: 2015-03-20 14:59:42 EDT
DISK UTILIZATION: 434.53 MB
NUMBER OF ARCHIVES: 22
ARCHIVE PATH: /opt/backup/dev/archived_wals
ARCHIVE COMMAND: cp %p /opt/backup/dev/archived_wals/%f
XLOG METHOD: fetch
WAL COMPRESSION: disabled
TABLESPACE PATH(s):
DESCRIPTION: "Development"
```

In the following code sample, the recovery window retention policy setting considers the backups taken within a 3-month recovery window as the active backups.

```
[HR]
host = 127.0.0.1
port = 5432
user = postgres
retention_policy = 3 MONTHS
description = "Human Resources"
```

The start of the 3-month recovery window displayed in the RETENTION POLICY field is 2015-01-10 14:04:23 EST when the SHOW-SERVERS subcommand is invoked on 2015-04-10.

At this current point in time, backups taken on or after 2015-01-10 14:04:23 EST would be considered active. Backups taken prior to 2015-01-10 14:04:23 EST would be considered obsolete, except for backups marked as keep.

```
-bash-4.1$ date
Fri Apr 10 15:04:19 EDT 2015
-bash-4.1$
-bash-4.1$ bart SHOW-SERVERS -s hr
SERVER NAME : hr
HOST NAME: 127.0.0.1
USER NAME: postgres
PORT: 5432
REMOTE HOST:
RETENTION POLICY: 2015-01-10 14:04:23 EST
DISK UTILIZATION: 480.76 MB
NUMBER OF ARCHIVES: 26
ARCHIVE PATH : /opt/backup/hr/archived_wals
ARCHIVE COMMAND: scp %p
enterprisedb@192.168.2.22:/opt/backup/hr/archived_wals/%f
XLOG METHOD: fetch
WAL COMPRESSION: disabled
TABLESPACE PATH(s):
DESCRIPTION: "Human Resources"
```

Evaluating, Marking, and Deleting Backup Using a Recovery Window Retention Policy

The following code sample uses a recovery window retention policy to evaluate, mark, and delete backups as shown by the following server configuration:

```
[DEV]
host = 127.0.0.1
port = 5445
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 DAYS
description = "Development"
```

The following is the current set of backups. Note that the last backup in the list has been marked as keep.

```
-bash-4.1$ bart SHOW-BACKUPS -s dev

SERVER NAME BACKUP ID BACKUP TIME BACKUP SIZE WAL(s) SIZE

WAL FILES STATUS

dev 1428933278236 2015-04-13 09:54:40 EDT 5.65 MB 16.00 MB

1 active
```

dev 2	1428862187757 active	2015-04-12	14:09:50	EDT	5.65 MB	32.00 MB
dev 2	1428768351638 active	2015-04-11	12:05:54	EDT	5.65 MB	32.00 MB
dev 14	1428684544008 active	2015-04-10	12:49:06	EDT	5.65 MB	224.00 MB
dev 3	1428590536488 active	2015-04-09	10:42:18	EDT	5.65 MB	48.00 MB
dev 5	1428502171990 keep	2015-04-08	10:09:34	EDT	5.65 MB	80.00 MB

The current date and time is 2015-04-13 16:46:35 EDT as shown below:

```
-bash-4.1$ date
Mon Apr 13 16:46:35 EDT 2015
```

Thus, a 3-day recovery window would evaluate backups prior to 2015-04-10 16:46:35 EDT as obsolete except for those marked as keep.

Invoke the MANAGE subcommand with the -n option to perform a dry run to observe which active backups would be changed to obsolete according to the retention policy.

```
-bash-4.1$ bart MANAGE -s dev -n
INFO: processing server 'dev', backup '1428933278236'
INFO: processing server 'dev', backup '1428862187757'
INFO: processing server 'dev', backup '1428768351638'
INFO: processing server 'dev', backup '1428684544008'
INFO: marking backup '1428684544008' as obsolete
INFO: 14 WAL file(s) marked obsolete
INFO: 1 Unused WAL file(s) present
INFO: processing server 'dev', backup '1428590536488'
INFO: marking backup '1428590536488' as obsolete
INFO: 3 WAL file(s) marked obsolete
INFO: 1 Unused WAL file(s) present
INFO: processing server 'dev', backup '1428502171990'
```

The dry run shows that backups 1428684544008 and 1428590536488 would be marked as obsolete.

Also note that a dry run does not change the backup status. The two backups that would be considered obsolete are still marked as active:

-bash-4.1\$ bar	t SHOW-BACKUPS	-s dev\		
SERVER NAME	BACKUP ID	BACKUP TIME	BACKUP SIZE	WAL(s) SIZE
WAL FILES	STATUS			
dev	1428933278236	2015-04-13 09:54:40 EDT	5.65 MB	16.00 MB
1	active			
dev	1428862187757	2015-04-12 14:09:50 EDT	5.65 MB	32.00 MB
2	active			
dev	1428768351638	2015-04-11 12:05:54 EDT	5.65 MB	32.00 MB
2	active			
dev	1428684544008	2015-04-10 12:49:06 EDT	5.65 MB	224.00 MB
14	active			
dev	1428590536488	2015-04-09 10:42:18 EDT	5.65 MB	48.00 MB
3	active			
dev	1428502171990	2015-04-08 10:09:34 EDT	5.65 MB	80.00 MB

```
5 keep
```

Invoke the MANAGE subcommand omitting the -n option to change and mark the status of the backups as obsolete:

```
-bash-4.1$ bart MANAGE -s dev
INFO: processing server 'dev', backup '1428933278236'
INFO: processing server 'dev', backup '1428862187757'
INFO: processing server 'dev', backup '1428768351638'
INFO: processing server 'dev', backup '1428684544008'
INFO: marking backup '1428684544008' as obsolete
INFO: 14 WAL file(s) marked obsolete
INFO: 1 Unused WAL file(s) present
INFO: processing server 'dev', backup '1428590536488'
INFO: marking backup '1428590536488' as obsolete
INFO: 3 WAL file(s) marked obsolete
INFO: 1 Unused WAL file(s) present
INFO: processing server 'dev', backup '1428502171990'
```

The obsolete backups can be observed in a number of ways. Use the MANAGE subcommand with the -l option to list the obsolete backups:

```
-bash-4.1$ bart MANAGE -s dev -l
INFO: 14 WAL file(s) will be removed
INFO: 1 Unused WAL file(s) will be removed
SERVER NAME: dev
BACKUP ID: 1428684544008
BACKUP STATUS: obsolete
BACKUP TIME: 2015-04-10 12:49:06 EDT
BACKUP SIZE: 5.65 MB
WAL FILE(s): 14
UNUSED WAL FILE(s): 1
WAL FILE: 000000010000000000000002E
WAL FILE: 000000010000000000000002D
WAL FILE: 000000010000000000000002B
WAL FILE: 000000010000000000000002A
WAL FILE: 000000010000000000000028
WAL FILE: 0000000100000000000000027
WAL FILE: 0000000100000000000000025
WAL FILE: 000000010000000000000024
WAL FILE: 000000010000000000000022
WAL FILE: 0000000100000000000000001
UNUSED WAL FILE: 000000010000000000000F.00000028
INFO: 3 WAL file(s) will be removed
INFO: 1 Unused WAL file(s) will be removed
SERVER NAME: dev
BACKUP ID: 1428590536488
BACKUP STATUS: obsolete
BACKUP TIME: 2015-04-09 10:42:18 EDT\
BACKUP SIZE: 5.65 MB
WAL FILE(s): 3
UNUSED WAL FILE(s): 1
```

UNUSED WAL FILE: 0000000100000000000000F.00000028

The STATUS field of the SHOW-BACKUPS subcommand displays the current status:

-bash-4.1\$ bart SHOW-BACKUPS -s dev -i 1428684544008 -t

```
-bash-4.1$ bart SHOW-BACKUPS -s dev
SERVER NAME BACKUP ID
                              BACKUP TIME
                                                         BACKUP SIZE WAL(s) SIZE
WAL FILES
             STATUS
dev
             1428933278236
                              2015-04-13 09:54:40 EDT
                                                         5.65 MB
                                                                     16.00 MB
1
             active
dev
             1428862187757
                              2015-04-12 14:09:50 EDT
                                                         5.65 MB
                                                                     32.00 MB
2
             active
dev
             1428768351638
                              2015-04-11 12:05:54 EDT
                                                         5.65 MB
                                                                     32.00 MB
2
             active
             1428684544008
                              2015-04-10 12:49:06 EDT
                                                                     224.00 MB
dev
                                                         5.65 MB
14
             obsolete
             1428590536488
                              2015-04-09 10:42:18 EDT
                                                                     48.00 MB
dev
                                                         5.65 MB
3
             obsolete
dev
             1428502171990
                              2015-04-08 10:09:34 EDT
                                                         5.65 MB
                                                                     80.00 MB
5
             keep
```

The details of an individual backup can be displayed using the SHOW-BACKUPS subcommand with the -t option. Note the status in the BACKUP STATUS field.

```
SERVER NAME
              : dev
BACKUP ID
               : 1428684544008
BACKUP NAME
               : none
BACKUP STATUS : obsolete
BACKUP TIME
               : 2015-04-10 12:49:06 EDT
BACKUP SIZE
              : 5.65 MB
WAL(S) SIZE
              : 224.00 MB
NO. OF WALS
              : 14
FIRST WAL FILE : 000000010000000000000021
CREATION TIME : 2015-04-10 12:49:06 EDT
LAST WAL FILE : 00000001000000000000002E
```

Use the MANAGE subcommand with the -d option to physically delete the obsolete backups including the unneeded WAL files.

CREATION TIME : 2015-04-11 12:02:15 EDT

The SHOW-BACKUPS subcommand now displays the remaining backups marked as active or keep:

-bash-4.1\$ b	art SHOW-BACKUP	S -s dev				
SERVER NAME	BACKUP ID	BACKUP TIME			BACKUP SIZE	WAL(s) SIZE
WAL FILES	STATUS					
dev	1428933278236	2015-04-13	09:54:40	EDT	5.65 MB	16.00 MB
1	active					
dev	1428862187757	2015-04-12	14:09:50	EDT	5.65 MB	32.00 MB
2	active					
dev	1428768351638	2015-04-11	12:05:54	EDT	5.65 MB	32.00 MB
2	active					
dev	1428502171990	2015-04-08	10:09:34	EDT	5.65 MB	80.00 MB
5	keep					

Managing Incremental Backups

This section illustrates evaluating, marking, and deleting incremental backups using the MANAGE and DELETE subcommands utilizing redundancy retention policy and recovery window retention policy. For detailed information about the MANAGE and DELETE subcommands, as well as the redundancy retention and recovery window retention policy, see the EDB Backup and Recovery User Guide available at the EDB website.

- Using a Redundancy Retention Policy provides an example of using the MANAGE and DELETE subcommands when a 3 backup redundancy retention policy is in effect.
- Using a Recovery Window Retention Policy provides an example of using the MANAGE and DELETE subcommands when a 1-day recovery window retention policy is in effect.

Using a Redundancy Retention Policy

The following code sample uses the MANAGE and DELETE subcommands to evaluate, mark, and delete incremental backups when a 3 backup redundancy retention policy is in effect. The example uses the following server configuration:

```
[ACCTG]
host = 192.168.2.24
port = 5445
user = enterprisedb
cluster_owner = enterprisedb
remote_host = enterprisedb@192.168.2.24
```

```
allow_incremental_backups = enabled
retention_policy = 3 BACKUPS
description = "Accounting"
```

The example uses the following set of backups. In these code samples, some columns have been omitted from the SHOW-BACKUPS output to display the relevant information in a more observable manner.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                             ... BACKUP PARENT
                                                  BACKUP TIME
                                                                          ... STATUS
                                                  2016-12-14 16:08:17 EST ... active
acctg
             1481749696905
                             ... 1481749673603
             1481749673603
acctg
                             ... 1481749651927
                                                  2016-12-14 16:07:53 EST ... active
             1481749651927
                             ... 1481749619582
                                                  2016-12-14 16:07:32 EST ... active
acctg
acctg
             1481749619582
                                 none
                                                  2016-12-14 16:07:00 EST ... active
```

There is one backup chain. The first backup is the initial full backup.

Backup chain: 1481749619582 => 1481749651927 => 1481749673603 => 1481749696905

The MANAGE subcommand is invoked as shown by the following:

```
-bash-4.2$ bart MANAGE -s acctg
INFO: processing server 'acctg', backup '1481749619582'
INFO: 2 Unused WAL file(s) present
INFO: 4 Unused file(s) (WALs included) present, use 'MANAGE -l' for the list
```

The following code sample shows the resulting status of the backups:

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME
                                                                          ... STATUS
               BACKUP ID
                             ... BACKUP PARENT
                                                  BACKUP TIME
acctg
               1481749696905 ... 1481749673603
                                                  2016-12-14 16:08:17 EST ... active
acctg
               1481749673603 ... 1481749651927
                                                  2016-12-14 16:07:53 EST ... active
               1481749651927 ... 1481749619582
                                                  2016-12-14 16:07:32 EST ... active
acctg
               1481749619582 ... none
                                                  2016-12-14 16:07:00 EST ... active
acctg
```

The status remains active for all backups. Even though the total number of backups exceeds the 3 backup redundancy retention policy, it is only the total number of full backups that is used to determine if the redundancy retention policy has been exceeded. Additional full backups are added including a second backup chain. The following example shows the resulting list of backups:

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                           ... BACKUP PARENT
                                                  BACKUP TIME
                                                                          ... STATUS
             1481750365397 ... none
                                                 2016-12-14 16:19:26 EST ... active
acctg
             1481750098924 ... 1481749997807
                                                 2016-12-14 16:14:59 EST ... active
acctg
             1481749997807 ... none
                                                 2016-12-14 16:13:18 EST ... active
acctg
acctg
             1481749992003 ... none
                                                 2016-12-14 16:13:12 EST ... active
acctg
             1481749696905 ... 1481749673603
                                                 2016-12-14 16:08:17 EST ... active
             1481749673603 ... 1481749651927
                                                 2016-12-14 16:07:53 EST ... active
acctg
             1481749651927 ... 1481749619582
                                                 2016-12-14 16:07:32 EST ... active
acctg
                                                 2016-12-14 16:07:00 EST ... active
acctg
             1481749619582 ... none
```

Second backup chain: 1481749997807 => 1481750098924

The MANAGE subcommand is invoked, but now with a total of four active full backups.

```
-bash-4.2$ bart MANAGE -s acctg
```

```
INFO: processing server 'acctg', backup '1481750365397'
INFO: processing server 'acctg', backup '1481749992003'
INFO: processing server 'acctg', backup '1481749619582'
INFO: processing server 'acctg', backup '1481749619582'
INFO: marking backup '1481749619582' as obsolete
INFO: 3 incremental(s) of backup '1481749619582' will be marked obsolete
INFO: marking incremental backup '1481749696905' as obsolete
INFO: marking incremental backup '1481749673603' as obsolete
INFO: marking incremental backup '1481749651927' as obsolete
INFO: 4 WAL file(s) marked obsolete
INFO: 2 Unused WAL file(s) present
INFO: 4 Unused file(s) (WALs included) present, use 'MANAGE -l' for the
list
```

The oldest full backup and its chain of incremental backups are now marked as obsolete.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                              ... BACKUP PARENT
                                                   BACKUP TIME
STATUS
             1481750365397
                                                   2016-12-14 16:19:26 EST ...
acctg
                              ... none
active
             1481750098924
                              ... 1481749997807
                                                   2016-12-14 16:14:59 EST ...
acctg
active
acctg
             1481749997807
                                                   2016-12-14 16:13:18 EST ...
                              ... none
active
                                                   2016-12-14 16:13:12 EST ...
acctg
             1481749992003
                              ... none
active
                                                   2016-12-14 16:08:17 EST ...
acctg
             1481749696905
                              ... 1481749673603
obsolete
acctg
             1481749673603
                              ... 1481749651927
                                                   2016-12-14 16:07:53 EST ...
obsolete
                                                   2016-12-14 16:07:32 EST ...
             1481749651927
                              ... 1481749619582
acctg
obsolete
             1481749619582
                                                   2016-12-14 16:07:00 EST ...
acctg
                              ... none
obsolete
```

Invoking the MANAGE subcommand with the -d option deletes the entire obsolete backup chain.

```
INFO: removing (unused) file
'000000010000000FB0000280000000FC000000.mbm'
```

The following code sample shows the remaining full backups and the second backup chain.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
                            ... BACKUP PARENT
                                                                         ... STATUS
SERVER NAME BACKUP ID
                                                 BACKUP TIME
acctg
            1481750365397
                             ... none
                                                 2016-12-14 16:19:26 EST ... active
                             ... 1481749997807
acctg
            1481750098924
                                                 2016-12-14 16:14:59 EST ... active
                                                 2016-12-14 16:13:18 EST ... active
            1481749997807
acctg
                             ... none
                                                 2016-12-14 16:13:12 EST ... active
acctg
            1481749992003
                             ... none
```

Using a Recovery Window Retention Policy

The following example demonstrates using the MANAGE and DELETE subcommands to evaluate, mark, and delete incremental backups when a 1-day recovery window retention policy is in effect. The example uses the following server configuration:

```
host = 192.168.2.24
port = 5445
user = enterprisedb
cluster_owner = enterprisedb
remote_host = enterprisedb@192.168.2.24
allow_incremental_backups = enabled
retention_policy = 1 DAYS
description = "Accounting"
```

The example uses the following set of backups. In the samples, some columns have been omitted from the SHOW-BACKUPS output to display the relevant information in a more observable manner.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg

SERVER NAME BACKUP ID ... BACKUP PARENT BACKUP TIME ... STATUS

acctg 1481559303348 ... 1481554203288 2016-12-12 11:15:03 EST ... active

acctg 1481559014359 ... 1481554802918 2016-12-12 11:10:14 EST ... active

acctg 1481554802918 ... 1481553914533 2016-12-12 10:00:03 EST ... active

acctg 1481554203288 ... 1481553651165 2016-12-12 09:50:03 EST ... active

acctg 1481553914533 ... 1481553088053 2016-12-12 09:45:14 EST ... active

acctg 1481553651165 ... none 2016-12-12 09:40:51 EST ... active

acctg 1481553088053 ... 1481552078404 2016-12-12 09:31:28 EST ... active

acctg 1481552078404 ... none 2016-12-12 09:14:39 EST ... active
```

There are two backup chains. In each of the following chains, the first backup is the initial full backup.

```
First backup chain: 1481552078404 => 1481553088053 => 1481553914533 => 1481554802918 => 1481559014359
```

Second backup chain: 1481553651165 => 1481554203288 => 1481559303348

The MANAGE subcommand is invoked when the first full backup 1481552078404 falls out of the recovery window. When the MANAGE subcommand is invoked, it is 2016-12-13 09:20:03 EST, thus making the start of the 1-day recovery window at 2016-12-12 09:20:03 EST exactly one day earlier. This backup was taken at 2016-12-12 09:14:39 EST, which is about 5 ½ minutes before the start of the recovery window, thus making the backup

obsolete.

```
-bash-4.2$ date

Tue Dec 13 09:20:03 EST 2016
-bash-4.2$ bart MANAGE -s acctg

INFO: processing server 'acctg', backup '1481553651165'

INFO: processing server 'acctg', backup '1481552078404'

INFO: marking backup '1481552078404' as obsolete

INFO: 4 incremental(s) of backup '1481552078404' will be marked obsolete

INFO: marking incremental backup '1481559014359' as obsolete

INFO: marking incremental backup '1481554802918' as obsolete

INFO: marking incremental backup '1481553914533' as obsolete

INFO: marking incremental backup '1481553088053' as obsolete

INFO: 7 WAL file(s) marked obsolete

INFO: 1 Unused WAL file(s) present

INFO: 2 Unused file(s) (WALs included) present, use 'MANAGE -l' for the list
```

The incremental backup date and time are within the recovery window since they were taken after the start of the recovery window of 2016-12-12 09:20:03 EST, but all backups in the chain are marked as obsolete.

-bash-4.2\$ bart	SHOW-BACKUPS -s acctg\	
SERVER NAME	BACKUP ID BACKUP PARENT	BACKUP TIME
STATUS		
acctg	1481559303348 1481554203288	2016-12-12 11:15:03 EST
active		
acctg	1481559014359 1481554802918	2016-12-12 11:10:14 EST
obsolete		
acctg	1481554802918 1481553914533	2016-12-12 10:00:03 EST
obsolete		
acctg	1481554203288 1481553651165	2016-12-12 09:50:03 EST
active		
acctg	1481553914533 1481553088053	2016-12-12 09:45:14 EST
obsolete		
acctg	1481553651165 none	2016-12-12 09:40:51 EST
active		
acctg	1481553088053 1481552078404	2016-12-12 09:31:28 EST
obsolete		
acctg	1481552078404 none	2016-12-12 09:14:39 EST
obsolete		

The following code sample shows how the entire backup chain is changed back to active status by invoking the MANAGE subcommand with the -c nokeep option on the full backup of the chain.

```
-bash-4.2$ bart MANAGE -s acctg -c nokeep -i 1481552078404
INFO: changing status of backup '1481552078404' of server 'acctg' from 'obsolete' to 'active'
INFO: status of 4 incremental(s) of backup '1481552078404' will be changed
INFO: changing status of incremental backup '1481559014359' of server 'acctg' from 'obsolete' to 'active'
INFO: changing status of incremental backup '1481554802918' of server 'acctg' from 'obsolete' to 'active'
INFO: changing status of incremental backup '1481553914533' of server 'acctg' from 'obsolete' to 'active'
INFO: changing status of incremental backup '1481553088053' of server
```

```
'acctg' from 'obsolete' to 'active'
INFO: 7 WAL file(s) changed
```

The backup chain has now been reset to active status.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
                                ... BACKUP PARENT
SERVER NAME
               BACKUP ID
                                                     BACKUP TIME
STATUS
                                                     2016-12-12 11:15:03 EST ...
acctg
               1481559303348
                                ... 1481554203288
active
                                 ... 1481554802918
acctg
               1481559014359
                                                     2016-12-12 11:10:14 EST ...
active
acctg
               1481554802918
                                ... 1481553914533
                                                     2016-12-12 10:00:03 EST ...
active
                                 ... 1481553651165
                                                     2016-12-12 09:50:03 EST ...
acctg
               1481554203288
active
               1481553914533
                                 ... 1481553088053
                                                     2016-12-12 09:45:14 EST ...
acctg
active
acctg
               1481553651165
                                ... none
                                                     2016-12-12 09:40:51 EST ...
active
               1481553088053
                                                     2016-12-12 09:31:28 EST ...
acctg
                                 ... 1481552078404
active
acctg
               1481552078404
                                ... none
                                                     2016-12-12 09:14:39 EST ...
active
```

The following code sample shows usage of the DELETE subcommand on an incremental backup. The specified incremental backup 1481554802918 in the first backup chain as well as its successive incremental backup 1481559014359 are deleted.

The results show that the backups 1481554802918 and 1481559014359 are no longer listed by the SHOW-BACKUPS subcommand.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID
                             ... BACKUP PARENT
                                                 BACKUP TIME
STATUS
acctg
             1481559303348
                             ... 1481554203288
                                                 2016-12-12 11:15:03 EST
active
             1481554203288
                                                 2016-12-12 09:50:03 EST
acctg
                             ... 1481553651165
active
acctg
             1481553914533
                             ... 1481553088053
                                                 2016-12-12 09:45:14 EST
active
acctg
             1481553651165
                             ... none
                                                 2016-12-12 09:40:51 EST
```

```
active
acctg 1481553088053 ... 1481552078404 2016-12-12 09:31:28 EST ...
active
acctg 1481552078404 ... none 2016-12-12 09:14:39 EST ...
active
```

The MANAGE subcommand is invoked again. This time both backup chains are marked obsolete since the full backups of both chains fall out of the start of the recovery window, which is now 2016-12-12 09:55:03 EST.

```
-bash-4.2$ date

Tue Dec 13 09:55:03 EST 2016
-bash-4.2$ bart MANAGE -s acctg

INFO: processing server 'acctg', backup '1481553651165'

INFO: marking backup '1481553651165' as obsolete

INFO: 2 incremental(s) of backup '1481553651165' will be marked obsolete

INFO: marking incremental backup '1481559303348' as obsolete

INFO: marking incremental backup '1481554203288' as obsolete

INFO: 38 WAL file(s) marked obsolete

INFO: processing server 'acctg', backup '1481552078404'

INFO: marking backup '1481552078404' as obsolete

INFO: 2 incremental(s) of backup '1481552078404' will be marked obsolete

INFO: marking incremental backup '1481553914533' as obsolete

INFO: marking incremental backup '1481553088053' as obsolete

INFO: 7 WAL file(s) marked obsolete
```

The following code sample shows both backup chains marked as obsolete.

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
                                                    BACKUP TIME
SERVER NAME
               BACKUP ID
                               ... BACKUP PARENT
... STATUS
               1481559303348
                               ... 1481554203288
                                                    2016-12-12 11:15:03 EST
acctg
... obsolete
acctg
               1481554203288
                               ... 1481553651165
                                                    2016-12-12 09:50:03 EST
... obsolete
                               ... 1481553088053
               1481553914533
                                                    2016-12-12 09:45:14 EST
acctg
... obsolete
                               ... none
acctg
               1481553651165
                                                    2016-12-12 09:40:51 EST
... obsolete
               1481553088053
                                ... 1481552078404
                                                    2016-12-12 09:31:28 EST
acctg
... obsolete
                                                    2016-12-12 09:14:39 EST
acctg
               1481552078404
                                ... none
... obsolete
```

The following code sample demonstrates using the MANAGE subcommand with the -c keep option to keep a backup chain indefinitely. The MANAGE subcommand with the -c keep option must specify the backup identifier or backup name of the full backup of the chain, and not any incremental backup.

```
-bash-4.2$ bart MANAGE -s acctg -c keep -i 1481553651165
INFO: changing status of backup '1481553651165' of server 'acctg' from 'obsolete' to 'keep'
INFO: status of 2 incremental(s) of backup '1481553651165' will be changed
INFO: changing status of incremental backup '1481559303348' of server 'acctg' from 'obsolete' to 'keep'
INFO: changing status of incremental backup '1481554203288' of server
```

```
'acctg' from 'obsolete' to 'keep'
INFO: 38 WAL file(s) changed
```

The full backup 1481553651165 and its successive incremental backups 1481554203288 and 1481559303348 have been changed to keep status:

```
-bash-4.2$ bart SHOW-BACKUPS -s acctg
SERVER NAME BACKUP ID ... BACKUP PARENT
                                              BACKUP TIME
... STATUS
           1481559303348 ... 1481554203288
                                              2016-12-12 11:15:03 EST
acctg
... keep
                                             2016-12-12 09:50:03 EST
acctg
           1481554203288 ... 1481553651165
... keep
         1481553914533 ... 1481553088053 2016-12-12 09:45:14 EST
acctg
... obsolete
                                              2016-12-12 09:40:51 EST
acctg
       1481553651165
                           ... none
... keep
           1481553088053 ... 1481552078404
                                              2016-12-12 09:31:28 EST
acctg
... obsolete
            1481552078404
                                              2016-12-12 09:14:39 EST
acctg
                           ... none
... obsolete
```

Finally, the MANAGE subcommand with the -d option is used to delete the obsolete backup chain:

```
-bash-4.2$ bart MANAGE -s acctg -d
INFO: removing all obsolete backups of server 'acctg'
INFO: removing obsolete backup '1481552078404'
INFO: 7 WAL file(s) will be removed
INFO: 2 incremental(s) of backup '1481552078404' will be removed
INFO: removing obsolete incremental backup '1481553914533'
INFO: removing obsolete incremental backup '1481553088053'
INFO: removing WAL file '000000010000000000000001'
INFO: removing WAL file '0000000100000000000000BF'
INFO: removing WAL file '0000000100000000000000BE'
INFO: removing WAL file '0000000100000000000000BD'
INFO: removing WAL file '0000000100000000000000BC'
INFO: removing WAL file '0000000100000000000000BB'
INFO: 48 Unused file(s) will be removed
INFO: removing (unused) file '00000001000000000000000000FA'
INFO: removing (unused) file '000000010000000000000BB.00000028.backup'
```

Only the backup chain with the keep status remains as shown below:

... keep

4.3 Sample BART System with Local and Remote Database Servers

This section describes a sample BART managed backup and recovery system consisting of both local and remote database servers. The complete steps to configure and operate the system are provided.

For detailed information about configuring a BART system, see the *EDB Backup and Recovery Installation and Upgrade Guide*. For detailed information about the operational procedures and BART subcommands, see the *EDB Backup and Recovery User Guide*. These guides are available at the *EDB website*.

The environment for this sample system is as follows:

- BART on host 192.168.2.22 running with BART user account enterprisedb
- Local Advanced Server on host 192.168.2.22 running with user account enterprisedb
- Remote Advanced Server on host 192.168.2.24 running with user account enterprisedb
- Remote PostgreSQL server on host 192.168.2.24 running with user account postgres

Passwordless SSH/SCP connections are required between the following:

- BART on host 192.168.2.22 and the local Advanced Server on the same host 192.168.2.22
- BART on host 192.168.2.22 and the remote Advanced Server on host 192.168.2.24
- BART on host 192.168.2.22 and the remote PostgreSQL server on host 192.168.2.24

The following sections demonstrate configuring and taking full backups only. To support incremental backups as well, enable the allow_incremental_backups parameter for the desired database servers and use the WAL scanner program.

- The BART Configuration File shows the settings used in the BART configuration file.
- Establishing SSH/SCP Passwordless Connections provides an example of how to establish an SSH/SCP passwordless connection.
- Configuring a Replication Database User provides an example of how to configure the replication database user.
- WAL Archiving Configuration Parameters provides an example of how to configure WAL archiving.
- Creating the BART Backup Catalog provides information about creating a BART Backup Catalog.
- Starting the Database Servers with WAL Archiving provides example of starting the database servers with WAL archiving.
- Taking a Full Backup illustrates taking the first full backup of the database servers.
- Using Point-In-Time Recovery demonstrates the point-in-time recovery operation on the remote PostgreSQL database server.

The BART Configuration File

The following code sample shows the settings used in the BART configuration file for the examples that follow:

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
retention_policy = 6 BACKUPS
```

```
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
cluster_owner = enterprisedb
backup_name = acctg_%year-%month-%dayT%hour:%minute
archive_command = 'cp %p %a/%f'
description = "Accounting"
[MKTG]
host = 192.168.2.24
port = 5444
user = repuser
cluster_owner = enterprisedb
backup_name = mktg_%year-%month-%dayT%hour:%minute
remote_host = enterprisedb@192.168.2.24
description = "Marketing"
[HR]
host = 192.168.2.24
port = 5432
user = postgres
cluster_owner = postgres
backup_name = hr_%year-%month-%dayT%hour:%minute
remote_host = postgres@192.168.2.24
copy_wals_during_restore = enabled
description = "Human Resources"
```

Establishing SSH/SCP Passwordless Connections

This section demonstrates how passwordless SSH/SCP connections are established with the authorized public keys files.

Generating a Public Key File for the BART User Account

The BART user account is enterprised with a home directory of /opt/PostgresPlus/9.6AS.

To generate the public key file, as a root user, first create the .ssh subdirectory in the BART user's home directory and assign ownership of this directory to the enterprised user, ensuring there are no groups or other users that can access the .ssh directory.

```
[root@localhost 9.6AS]# pwd
/opt/PostgresPlus/9.6AS
[root@localhost 9.6AS]# mkdir .ssh
[root@localhost 9.6AS]# chown enterprisedb .ssh
[root@localhost 9.6AS]# chgrp enterprisedb .ssh
[root@localhost 9.6AS]# chmod 700 .ssh
[root@localhost 9.6AS]# ls -la | grep ssh
```

```
drwx----- 2 enterprisedb enterprisedb 4096 Apr 23 13:02 .ssh
```

Generate the public key file:

```
[user@localhost ~]$ su - enterprisedb
Password:
-bash-4.1$ pwd
/opt/PostgresPlus/9.6AS
-bash-4.1$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key
(/opt/PostgresPlus/9.6AS/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in
/opt/PostgresPlus/9.6AS/.ssh/id_rsa.
Your public key has been saved in
/opt/PostgresPlus/9.6AS/.ssh/id_rsa.pub.
The key fingerprint is:
de:65:34:d6:b1:d2:32:3c:b0:43:c6:a3:c0:9f:f4:64
enterprisedb@localhost.localdomain
The key's randomart image is:
+----[ RSA 2048]----+
            .+ . |
         o .oE+ o o |
          + * o.X + |
           + .+ *
           S
             0
          . . 0
```

The following are the resulting files. id_rsa.pub is the public key file of BART user account enterprisedb.

```
-bash-4.1$ ls -l .ssh
total 8
-rw----- 1 enterprisedb enterprisedb 1675 Apr 23 13:04 id_rsa
-rw-r--r-- 1 enterprisedb enterprisedb 416 Apr 23 13:04 id_rsa.pub
```

Configuring Access between Local Advanced Server and the BART Host

Even when the Advanced Server database is on the same host as the BART user account, and the Advanced Server database cluster owner is also the BART user account (enterprised) is this case), a passwordless SSH/SCP connection must be established from the same user account to itself.

On the BART host where the public key file was just generated (as shown in Generating a Public Key File for the BART User Account), create the authorized keys file by appending the public key file to any existing authorized keys file.

Log into the BART host as the BART user account and append the public key file, id_rsa.pub onto the authorized_keys file in the same .ssh directory.

```
[user@localhost ~]$ su - enterprisedb
Password:
Last login: Thu Mar 23 10:27:35 EDT 2017 on pts/0
-bash-4.2$ pwd
/opt/PostgresPlus/9.6AS
-bash-4.2$ ls -l .ssh
total 12
-rw----- 1 enterprisedb enterprisedb 1675 Mar 23 09:54 id_rsa
-rw-r--r 1 enterprisedb enterprisedb 416 Mar 23 09:54 id_rsa.pub
-rw-r--r 1 enterprisedb enterprisedb 345 Mar 23 10:05 known hosts
-bash-4.2$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
-bash-4.2$ ls -l .ssh
total 16
-rw-rw-r-- 1 enterprisedb enterprisedb 416 Mar 23 10:33 authorized keys
-rw----- 1 enterprisedb enterprisedb 1675 Mar 23 09:54 id rsa
-rw-r--r 1 enterprisedb enterprisedb 416 Mar 23 09:54 id_rsa.pub
-rw-r--r 1 enterprisedb enterprisedb 345 Mar 23 10:05 known_hosts
```

The authorized_keys file must have file permission 600 as set by the following chmod 600 command, or the passwordless connection will fail:

```
-bash-4.2$ chmod 600 ~/.ssh/authorized_keys
-bash-4.2$ ls -l .ssh
total 16
-rw----- 1 enterprisedb enterprisedb 416 Mar 23 10:33 authorized_keys
-rw----- 1 enterprisedb enterprisedb 1675 Mar 23 09:54 id_rsa
-rw-r--r- 1 enterprisedb enterprisedb 416 Mar 23 09:54 id_rsa.pub
-rw-r--r- 1 enterprisedb enterprisedb 345 Mar 23 10:05 known_hosts
```

Test the passwordless connection. Use the ssh command to verify that you can access the same user account as you are currently logged in as (enterprisedb) without being prompted for a password:

```
-bash-4.2$ ssh enterprisedb@127.0.0.1
Last login: Thu Mar 23 10:27:50 2017
-bash-4.2$ exit
logout
Connection to 127.0.0.1 closed.
```

Configuring Access from Remote Advanced Server to BART Host

On the remote host 192.168.2.24, create the public key file for the remote database server user account, enterprisedb, for access to the BART user account, enterprisedb, on the BART host 192.168.2.22.

Create the .ssh directory for user account enterprisedb on the remote host:

```
[root@localhost 9.6AS]# pwd
/opt/PostgresPlus/9.6AS
[root@localhost 9.6AS]# mkdir .ssh
[root@localhost 9.6AS]# chown enterprisedb .ssh
[root@localhost 9.6AS]# chgrp enterprisedb .ssh
[root@localhost 9.6AS]# chmod 700 .ssh
[root@localhost 9.6AS]# ls -la | grep ssh
drwx----- 2 enterprisedb enterprisedb 4096 Apr 23 13:08 .ssh
```

Generate the public key file on the remote host for user account enterprisedb:

```
[user@localhost ~]$ su - enterprisedb
Password:
-bash-4.1$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key
(/opt/PostgresPlus/9.6AS/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in
/opt/PostgresPlus/9.6AS/.ssh/id_rsa.
Your public key has been saved in
/opt/PostgresPlus/9.6AS/.ssh/id_rsa.pub.
The key fingerprint is:
15:27:1e:1e:61:4b:48:66:67:0b:b2:be:fc:ea:ea:e6
enterprisedb@localhost.localdomain
The key's randomart image is:
+--[ RSA 2048]---+
     ..=.@..
      =.0 0
      . *
     . S
     0
 +Eoo..
```

Copy the generated public key file, id_rsa.pub, to the BART user account, enterprisedb, on the BART host, 192.168.2.22:

```
-bash-4.1$ scp ~/.ssh/id_rsa.pub enterprisedb@192.168.2.22:/tmp/tmp.pub
The authenticity of host '192.168.2.22 (192.168.2.22)' can't be
established.
RSA key fingerprint is b8:a9:97:31:79:16:b8:2b:b0:60:5a:91:38:d7:68:22.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.2.22' (RSA) to the list of known hosts.
enterprisedb@192.168.2.22's password:
id_rsa.pub
```

Log into the BART host as the BART user account and append the temporary public key file, /tmp/tmp.pub onto the authorized_keys file owned by the BART user account.

```
-bash-4.1$ ssh enterprisedb@192.168.2.22
enterprisedb@192.168.2.22's password:
Last login: Tue Apr 21 17:03:24 2015 from 192.168.2.22
-bash-4.1$ pwd
/opt/PostgresPlus/9.6AS
-bash-4.1$ cat /tmp/tmp.pub >> ~/.ssh/authorized_keys
-bash-4.1$ ls -l .ssh
total 12
-rw-rw-r-- 1 enterprisedb enterprisedb 416 Apr 23 13:15 authorized_keys
-rw------ 1 enterprisedb enterprisedb 1675 Apr 23 13:04 id_rsa
```

```
-rw-r--r-- 1 enterprisedb enterprisedb 416 Apr 23 13:04 id_rsa.pub
```

The authorized_keys file must have file permission 600 as set by the following chmod 600 command, otherwise the passwordless connection fails:

```
-bash-4.1$ chmod 600 ~/.ssh/authorized_keys
-bash-4.1$ ls -l .ssh
total 12
-rw----- 1 enterprisedb enterprisedb 416 Apr 23 13:15 authorized_keys
-rw----- 1 enterprisedb enterprisedb 1675 Apr 23 13:04 id_rsa
-rw-r--r- 1 enterprisedb enterprisedb 416 Apr 23 13:04 id_rsa.pub
-bash-4.1$ rm /tmp/tmp.pub
-bash-4.1$ exit
logout
Connection to 192.168.2.22 closed.
```

Test the passwordless connection. From the remote host, verify that you can log into the BART host with the BART user account without being prompted for a password:

```
-bash-4.1$ ssh enterprisedb@192.168.2.22
Last login: Thu Apr 23 13:14:48 2015 from 192.168.2.24
-bash-4.1$ exit
logout
Connection to 192.168.2.22 closed.
```

Configuring Access from the BART Host to a Remote Advanced Server

On the BART host 192.168.2.22, copy the public key file for the BART user account, enterprisedb, for access to the remote database server user account, enterprisedb, on the remote host 192.168.2.24.

The following lists the current SSH keys files in the BART user's .ssh directory on the BART host:

```
[user@localhost ~]$ su - enterprisedb
Password:
-bash-4.1$ pwd
/opt/PostgresPlus/9.6AS
-bash-4.1$ ls -l .ssh
total 12
-rw----- 1 enterprisedb enterprisedb 416 Apr 23 13:15 authorized_keys
-rw----- 1 enterprisedb enterprisedb 1675 Apr 23 13:04 id_rsa
-rw-r--r-- 1 enterprisedb enterprisedb 416 Apr 23 13:04 id_rsa.pub
```

The public key file, id_rsa.pub, for BART user account enterprisedb on the BART host that was earlier generated in Generating a Public Key File for the BART User Account, is now copied to the remote Advanced Server host on 192.168.2.24:

```
-bash-4.1$ scp ~/.ssh/id_rsa.pub enterprisedb@192.168.2.24:/tmp/tmp.pub
The authenticity of host '192.168.2.24 (192.168.2.24)' can't be
established.
RSA key fingerprint is 59:41:fb:0c:ae:64:3d:3f:a2:d9:90:95:cf:2c:99:f2.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.2.24' (RSA) to the list of known
hosts.
```

```
enterprisedb@192.168.2.24's password:
id_rsa.pub
```

Log into the enterprisedb user account on the remote host and copy the public key file onto the authorized_keys file of the remote enterprisedb user account under its .ssh directory:

```
-bash-4.1$ ssh enterprisedb@192.168.2.24
enterprisedb@192.168.2.24's password:
Last login: Tue Apr 21 09:53:18 2015 from 192.168.2.22
-bash-4.1$ pwd
/opt/PostgresPlus/9.6AS
-bash-4.1$ ls -l .ssh
total 12
-rw----- 1 enterprisedb enterprisedb 1675 Apr 23 13:11 id_rsa
-rw-r--r- 1 enterprisedb enterprisedb 416 Apr 23 13:11 id_rsa.pub
-rw-r--r- 1 enterprisedb enterprisedb 394 Apr 23 13:12 known_hosts
-bash-4.1$ cat /tmp/tmp.pub >> ~/.ssh/authorized_keys
```

Adjust the file permission on authorized_keys:

```
-bash-4.1$ chmod 600 ~/.ssh/authorized_keys
-bash-4.1$ ls -l .ssh

total 16
-rw------ 1 enterprisedb enterprisedb 416 Apr 23 13:26 authorized_keys
-rw------ 1 enterprisedb enterprisedb 1675 Apr 23 13:11 id_rsa
-rw-r--r-- 1 enterprisedb enterprisedb 416 Apr 23 13:11 id_rsa.pub
-rw-r--r-- 1 enterprisedb enterprisedb 394 Apr 23 13:12 known_hosts
-bash-4.1$ rm /tmp/tmp.pub
-bash-4.1$ exit
logout
Connection to 192.168.2.24 closed.
```

While logged into the BART host, test the passwordless connection from the BART host to the remote Advanced Server host:

```
-bash-4.1$ ssh enterprisedb@192.168.2.24
Last login: Thu Apr 23 13:25:53 2015 from 192.168.2.22
-bash-4.1$ exit
logout
Connection to 192.168.2.24 closed.
```

Configuring Access from a Remote PostgreSQL Server to a BART Host

On the remote host (192.168.2.24), create a public key file owned by the database server user account (postgres), allowing access to the BART user account (enterprisedb) on the BART host (192.168.2.22).

Create the .ssh directory for the postgres user account on the remote host:

```
[root@localhost 9.6]# cd /opt/PostgreSQL/9.6
[root@localhost 9.6]# mkdir .ssh
[root@localhost 9.6]# chown postgres .ssh
[root@localhost 9.6]# chgrp postgres .ssh
[root@localhost 9.6]# chmod 700 .ssh
```

```
[root@localhost 9.6]# ls -la | grep ssh
drwx----- 2 postgres postgres 4096 Apr 23 13:32 .ssh
```

Create and copy the generated public key file, id_rsa.pub, to the BART user account (enterprisedb), on the BART host (192.168.2.22):

```
[user@localhost ~]$ su - postgres
Password:
-bash-4.1$ pwd
/opt/PostgreSQL/9.6
-bash-4.1$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/opt/PostgreSQL/9.6/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /opt/PostgreSQL/9.6/.ssh/id_rsa.
Your public key has been saved in /opt/PostgreSQL/9.6/.ssh/id_rsa.pub.
The key fingerprint is:
1f:f8:76:d6:fc:a5:1a:c5:5a:66:66:01:d0:a0:ca:ba
postgres@localhost.localdomain
The key's randomart image is:
+--[ RSA 2048]----+
           0+.
          . . .
      o S . 0 |
        o . @
          + = 0 .
          . 0 . 0.
     Ε
           . . . .
-bash-4.1$ ls -l .ssh
total 8
-rw----- 1 postgres postgres 1671 Apr 23 13:36 id_rsa
-rw-r--r- 1 postgres postgres 412 Apr 23 13:36 id_rsa.pub
-bash-4.1$ scp ~/.ssh/id_rsa.pub enterprisedb@192.168.2.22:/tmp/tmp.pub
The authenticity of host '192.168.2.22 (192.168.2.22)' can't be
established.
RSA key fingerprint is b8:a9:97:31:79:16:b8:2b:b0:60:5a:91:38:d7:68:22.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.2.22' (RSA) to the list of known
hosts.
enterprisedb@192.168.2.22's password:
id_rsa.pub
```

Log into the BART host as the BART user account and append the temporary public key file, /tmp/tmp.pub, onto the authorized_keys file owned by the BART user account.

```
-bash-4.1$ ssh enterprisedb@192.168.2.22
enterprisedb@192.168.2.22's password:
Last login: Thu Apr 23 13:19:25 2015 from 192.168.2.24
-bash-4.1$ pwd
/opt/PostgresPlus/9.6AS
-bash-4.1$ cat /tmp/tmp.pub >> ~/.ssh/authorized_keys
```

```
-bash-4.1$ ls -l .ssh

total 16

-rw----- 1 enterprisedb enterprisedb 828 Apr 23 13:40 authorized_keys

-rw----- 1 enterprisedb enterprisedb 1675 Apr 23 13:04 id_rsa

-rw-r--r- 1 enterprisedb enterprisedb 416 Apr 23 13:04 id_rsa.pub

-rw-r--r- 1 enterprisedb enterprisedb 394 Apr 23 13:24 known_hosts

-bash-4.1$ rm /tmp/tmp.pub

-bash-4.1$ exit

logout

Connection to 192.168.2.22 closed.
```

Make sure the authorized_keys file has file permission 600 as shown, or the passwordless connection will fail. Test the passwordless connection; from the remote host, while logged in as user account postgres, verify that you can log into the BART host with the BART user account without being prompted for a password:

```
-bash-4.1$ pwd
/opt/PostgreSQL/9.6
-bash-4.1$ ssh enterprisedb@192.168.2.22
Last login: Thu Apr 23 13:40:10 2015 from 192.168.2.24
-bash-4.1$ exit
logout
Connection to 192.168.2.22 closed.
```

Configuring Access from the BART Host to Remote PostgreSQL

Copy the public key file on the BART host that is owned by the BART user account (enterprisedb) to the remote database server user account (postgres), on the remote host (192.168.2.24).

The following lists the current SSH keys files in the BART user's .ssh directory on the BART host:

```
[user@localhost ~]$ su - enterprisedb
Password:
-bash-4.1$ ls -l .ssh
total 16
-rw----- 1 enterprisedb enterprisedb 828 Apr 23 13:40 authorized_keys
-rw----- 1 enterprisedb enterprisedb 1675 Apr 23 13:04 id_rsa
-rw-r--r- 1 enterprisedb enterprisedb 416 Apr 23 13:04 id_rsa.pub
-rw-r--r- 1 enterprisedb enterprisedb 394 Apr 23 13:24 known_hosts
```

The public key file, id_rsa.pub, for BART user account enterprisedb on the BART host that was earlier generated in Generating a Public Key File for the BART User Account, now resides on the remote PostgreSQL host:

```
-bash-4.1$ scp ~/.ssh/id_rsa.pub postgres@192.168.2.24:/tmp/tmp.pub postgres@192.168.2.24's password: id_rsa.pub
```

Log into the postgres user account on the remote host and copy the public key file onto the authorized_keys file of postgres under its .ssh directory:

```
-bash-4.1$ ssh postgres@192.168.2.24
postgres@192.168.2.24's password:
Last login: Mon Jan 26 18:08:36 2015 from 192.168.2.19
-bash-4.1$ pwd
```

```
/opt/PostgreSQL/9.6
-bash-4.1$ cat /tmp/tmp.pub >> ~/.ssh/authorized_keys
```

Adjust the file permissions on authorized_keys:

```
-bash-4.1$ ls -l .ssh
total 16
-rw-rw-r-- 1 postgres postgres 416 Apr 23 13:52 authorized_keys
-rw----- 1 postgres postgres 1671 Apr 23 13:36 id_rsa
-rw-r--r 1 postgres postgres 412 Apr 23 13:36 id_rsa.pub
-rw-r--r 1 postgres postgres 394 Apr 23 13:36 known_hosts
-bash-4.1$ chmod 600 ~/.ssh/authorized_keys
-bash-4.1$ ls -l .ssh
total 16
-rw----- 1 postgres postgres 416 Apr 23 13:52 authorized_keys
-rw----- 1 postgres postgres 1671 Apr 23 13:36 id_rsa
-rw-r--r- 1 postgres postgres 412 Apr 23 13:36 id_rsa.pub
-rw-r--r 1 postgres postgres 394 Apr 23 13:36 known_hosts
-bash-4.1$ rm /tmp/tmp.pub
-bash-4.1$ exit
logout
Connection to 192.168.2.24 closed.
```

Test the passwordless connection from the BART host to the remote PostgreSQL host:

```
[user@localhost ~]$ su - enterprisedb
Password:
-bash-4.1$ ssh postgres@192.168.2.24
Last login: Thu Apr 23 13:52:25 2015 from 192.168.2.22
-bash-4.1$ exit
logout
Connection to 192.168.2.24 closed.
```

Configuring a Replication Database User

This section demonstrates how a replication database user is established.

All database servers must use a superuser as the replication database user.

The replication database user for each database server is specified by the user parameter in the BART configuration file as shown by the following:

```
host = 127.0.0.1
port = 5444
user = enterprisedb <=== Replication Database User
cluster_owner = enterprisedb
backup_name = acctg_%year-%month-%dayT%hour:%minute
archive_command = 'cp %p %a/%f'
description = "Accounting"

[MKTG]
host = 192.168.2.24</pre>
```

```
port = 5444
user = repuser <=== Replication Database User
cluster_owner = enterprisedb
backup_name = mktg_%year-%month-%dayT%hour:%minute
remote_host = enterprisedb@192.168.2.24
description = "Marketing"

[HR]
host = 192.168.2.24
port = 5432
user = postgres <=== Replication Database User
cluster_owner = enterprisedb
backup_name = hr_%year-%month-%dayT%hour:%minute
remote_host = postgres@192.168.2.24
copy_wals_during_restore = enabled
description = "Human Resources"</pre>
```

Add entries to the .pgpass file on each server to allow the BART user account to initiate a backup without being prompted for credentials. The .pgpass file is located in /opt/PostgresPlus/9.6AS/.pgpass:

```
127.0.0.1:5444:*:enterprisedb:password
192.168.2.24:5444:*:repuser:password
192.168.2.24:5432:*:postgres:password
```

For more information about using a .pgpass file, please see the PostgreSQL documentation.

While connected to MKTG on 192.168.2.24, execute the following CREATE ROLE command to create the replication database superuser:

```
CREATE ROLE repuser WITH LOGIN SUPERUSER PASSWORD 'password';
```

Access is granted in the pg_hba.conf file for the local Advanced Server:

```
### TYPE
             DATABASE
                                  USER
                                                                METHOD
### "local" is for Unix domain socket connections only
             all
                                                                md5
local
                                 all
### IPv4
             local connections:
                              enterprisedb 127.0.0.1/32
                                                                md5
host
           template1
           edb
                              enterprisedb 127.0.0.1/32
                                                                md5
host
#host
           all
                              all
                                             127.0.0.1/32
                                                                md5
### IPv6 local connections:
host
           all.
                              all
                                             ::1/128
                                                                md5
### Allow replication connections from localhost, by a user with the
### replication privilege.
#local
           replication
                              enterprisedb
                                                                md5
host
           replication
                              enterprisedb
                                             127.0.0.1/32
                                                                md5
```

Similarly, access is granted in the pg_hba.conf file for the remote Advanced Server installation:

```
### TYPE DATABASE USER ADDRESS METHOD
### "local" is for Unix domain socket connections only
local all all md5
### IPv4 local connections:
host template1 repuser 192.168.2.22/32 md5
```

```
host
           all
                              enterprisedb 127.0.0.1/32
                                                                   md5
           all
                                            127.0.0.1/32
                                                                  md5
#host
### IPv6 local connections:
           all
                              all
                                                                  md5
host
                                             ::1/128
### Allow replication connections from localhost, by a user with the
### replication privilege.
         replication
#local
                              enterprisedb
                                                                  md5
host
         replication
                              repuser
                                             192.168.2.22/32
                                                                  md5
```

Access is also granted in the pg_hba.conf file for the remote PostgreSQL server:

```
### TYPE
             DATABASE
                                 USER
                                              ADDRESS
                                                                     METHOD
### "local" is for Unix domain socket connections only
local
           all
                               all
                                                                     md5
### IPv4 local connections:
host
           template1
                               postgres
                                           192.168.2.22/32
                                                                     md5
           all
                                           127.0.0.1/32
                                                                     md5
host
                              all
### IPv6 local connections:
host
           all
                               all
                                           ::1/128
                                                                     md5
### Allow replication connections from localhost, by a user with the
q# replication privilege.
           replication
#local
                                                                     md5
                               postgres
host
           replication
                                           192.168.2.22/32
                                                                     md5
                               postgres
```

WAL Archiving Configuration Parameters

Use the following parameters in the postgresql.conf file to enable WAL archiving. The postgresql.conf file for the local Advanced Server database (ACCTG) is set as follows:

When the INIT subcommand is invoked, the Postgres archive_command configuration parameter in the postgresql.auto.conf file will be set based on the BART archive_command parameter located in the BART configuration file.

!!! Note If the Postgres archive_command is already set, invoke the INIT subcommand with the -- no-configure option to prevent the archive_command from being reset. For details, see INIT.

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
retention_policy = 6 BACKUPS
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log
```

```
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
cluster_owner = enterprisedb
backup_name = acctg_%year-%month-%dayT%hour:%minute
archive_command = 'cp %p %a/%f'
description = "Accounting"
```

When the INIT subcommand is invoked, the postgresql.auto.conf file contains the following:

```
### Do not edit this file manually!
### It will be overwritten by ALTER SYSTEM command.
archive_command = 'cp %p /opt/backup/acctg/archived_wals/%f'
```

The archive_command uses the cp command instead of scp since the BART backup catalog is local to this database cluster and the BART user account (the account that owns the backup catalog, enterprisedb), is the same user account running Advanced Server. The result is that there is no directory permission conflict during the archive operation.

The postgresql.conf file for the remote Advanced Server, MKTG is set as follows:

When the INIT subcommand is invoked, the Postgres archive_command configuration parameter in the postgresql.auto.conf file will be set by the default BART format of the BART archive_command parameter (since it is not explicitly set for this database server in the BART configuration file).

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
retention_policy = 6 BACKUPS
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log
.
.
.
[MKTG]
host = 192.168.2.24
port = 5444
user = repuser
cluster_owner = enterprisedb
backup_name = mktg_%year-%month-%dayT%hour:%minute
remote_host = enterprisedb@192.168.2.24
```

```
description = "Marketing"
```

The default BART archive_command format is:

```
archive_command = 'scp %p %h:%a/%f'
```

The postgresql.auto.conf file contains the following after the INIT subcommand is invoked:

```
### Do not edit this file manually!
### It will be overwritten by ALTER SYSTEM command.
archive_command = 'scp %p
enterprisedb@192.168.2.22:/opt/backup/hr/archived_wals/%f'
```

The archive_command uses the scp command since the BART backup catalog is remote relative to this database cluster. The BART user account, enterprisedb, is specified on the scp command since this is the user account owning the BART backup catalog where the archived WAL files are to be copied. The result is that there is no directory permission conflict during the archive operation.

The postgresql.conf file for the remote PostgreSQL server (HR) is set as follows:

When the INIT subcommand is invoked, the Postgres archive_command configuration parameter in the postgresql.auto.conf file will be set by the default BART format of the BART archive_command parameter (since it is not explicitly set for this database server in the BART configuration file):

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
retention_policy = 6 BACKUPS
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log
[HR]
host = 192.168.2.24
port = 5432
user = postgres
cluster_owner = postgres
backup_name = hr_%year-%month-%dayT%hour:%minute
remote_host = postgres@192.168.2.24
copy_wals_during_restore = enabled
description = "Human Resources"
```

The default BART archive_command format is:

```
archive_command = 'scp %p %h:%a/%f'
```

The postgresql.auto.conf file contains the following after the INIT subcommand is invoked:

```
### Do not edit this file manually!
### It will be overwritten by ALTER SYSTEM command.
archive_command = 'scp %p
enterprisedb@192.168.2.22:/opt/backup/hr/archived_wals/%f'
```

The archive_command uses the scp command since the BART backup catalog is remote relative to this database cluster. The BART user account, enterprisedb, is specified on the scp command since this is the user account owning the BART backup catalog where the archived WAL files are to be copied. The result is that there is no directory permission conflict during the archive operation.

Creating the BART Backup Catalog (backup_path)

Create the directory specified by the backup_path configuration parameter.

```
[BART]
bart_host= enterprisedb@192.168.2.22
backup_path = /opt/backup
pg_basebackup_path = /usr/edb/as11/bin/pg_basebackup
retention_policy = 6 BACKUPS
logfile = /tmp/bart.log
scanner_logfile = /tmp/bart_scanner.log
```

Ensure that the directory is owned by the BART user account:

```
[root@localhost opt]# pwd
/opt
[root@localhost opt]# mkdir backup
[root@localhost opt]# chown enterprisedb backup
[root@localhost opt]# chgrp enterprisedb backup
[root@localhost opt]# chmod 700 backup
[root@localhost opt]# ls -l | grep backup
drwx----- 2 enterprisedb enterprisedb 4096 Apr 23 15:36 backup
```

Use the BART INIT subcommand to complete the directory structure and set the Postgres archive_command configuration parameter.

Before invoking any BART subcommands, set up a profile under the BART user account's home directory to set the LD_LIBRARY_PATH and PATH environment variables. For more information regarding setting this variable, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.

The -o option is specified with the INIT subcommand to force the setting of the Postgres archive_command configuration parameter when archive_mode is off or if the Postgres archive_command parameter is already set and needs to be overridden.

```
[user@localhost ~]$ su - enterprisedb
Password:
-bash-4.1$ bart INIT -o
```

```
INFO: setting archive_command for server 'acctg'
WARNING: archive_command is set. server restart is required
INFO: setting archive_command for server 'hr'
WARNING: archive_command is set. server restart is required
INFO: setting archive_command for server 'mktg'
WARNING: archive_command is set. server restart is required
```

The BART SHOW-SERVERS subcommand displays the following:

```
-bash-4.1$ bart SHOW-SERVERS
SERVER NAME :
                            acctg
BACKUP FRIENDLY NAME:
                            acctg_%year-%month-%dayT%hour:%minute
HOST NAME:
                            127.0.0.1
USER NAME:
                            enterprisedb
PORT:
                            5444
REMOTE HOST:
RETENTION POLICY:
                            6 Backups
DISK UTILIZATION:
                            0.00 bytes
NUMBER OF ARCHIVES:
ARCHIVE PATH:
                            /opt/backup/acctg/archived_wals
ARCHIVE COMMAND:
                           (disabled)
XLOG METHOD:
                           fetch
WAL COMPRESSION:
                           disabled
TABLESPACE PATH(s):
INCREMENTAL BACKUP:
                           DISABLED
                           "Accounting"
DESCRIPTION:
SERVER NAME:
BACKUP FRIENDLY NAME:
                           hr_%year-%month-%dayT%hour:%minute
HOST NAME:
                           192.168.2.24
USER NAME:
                           postgres
PORT:
                           5432
REMOTE HOST:
                           postgres@192.168.2.24
RETENTION POLICY:
                           6 Backups
DISK UTILIZATION:
                           0.00 bytes
NUMBER OF ARCHIVES:
ARCHIVE PATH:
                           /opt/backup/hr/archived_wals
ARCHIVE COMMAND:
                           (disabled)
XLOG METHOD:
                           fetch
                           disabled
WAL COMPRESSION:
TABLESPACE PATH(s):
INCREMENTAL BACKUP :
                           DISABLED
                           "Human Resources"
DESCRIPTION:
SERVER NAME:
                           mktg
                           mktg_%year-%month-%dayT%hour:%minute
BACKUP FRIENDLY NAME:
HOST NAME:
                           192.168.2.24
USER NAME:
                           repuser
PORT:
                           5444
REMOTE HOST:
                           enterprisedb@192.168.2.24
RETENTION POLICY:
                           6 Backups
DISK UTILIZATION:
                           0.00 bytes
NUMBER OF ARCHIVES:
ARCHIVE PATH:
                           /opt/backup/mktg/archived_wals
ARCHIVE COMMAND:
                           (disabled)
XLOG METHOD:
                           fetch
WAL COMPRESSION:
                           disabled
```

```
TABLESPACE PATH(s):
INCREMENTAL BACKUP:
                           DISABLED
DESCRIPTION:
                           "Marketing"
-bash-4.1$ cd /opt/backup
-bash-4.1$ pwd
/opt/backup
-bash-4.1$ ls -l
total 12
drwxrwxr-x 3 enterprisedb enterprisedb 4096 Mar 29 13:16 acctg
drwxrwxr-x 3 enterprisedb enterprisedb 4096 Mar 29 13:16 hr
drwxrwxr-x 3 enterprisedb enterprisedb 4096 Mar 29 13:16 mktg
-bash-4.1$ ls -l acctg
total 4
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Mar 29 13:16 archived wals
-bash-4.1$ ls -l hr
total 4
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Mar 29 13:16 archived_wals
-bash-4.1$ ls -l mktg
total 4
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Mar 29 13:16 archived_wals
```

The ARCHIVE PATH field displays the full directory path to where the WAL files are copied. This directory path must match the directory path specified in the Postgres archive_command parameter of the postgresql.conf file or the postgresql.auto.conf file of each database server.

Starting the Database Servers with WAL Archiving

After the BART backup catalog directory structure has been configured, start the archiving of WAL files from the database servers by restarting each database server.

On BART host 192.168.2.22:

```
[root@localhost data]# service ppas-9.6 restart
```

On remote host 192.168.2.24:

```
[root@localhost data]# service ppas-9.6 restart
[root@localhost data]# service postgresql-9.6 restart
```

In the BART backup catalog, verify that the WAL files are archiving.

Archived WAL files may not appear very frequently depending upon how often WAL archiving is set to switch to a new segment file with the archive_timeout parameter in your database server configuration settings.

Verify that there are no archiving-related errors in the database server log files.

Taking a Full Backup

The following code sample shows the first full backup of the database servers.

```
-bash-4.1$ bart BACKUP -s acctg -z
```

```
INFO: creating backup for server 'acctg'
INFO: backup identifier: '1490809695281'
60776/60776 kB (100%), 1/1 tablespace
INFO: backup completed successfully
INFO: backup checksum: 37f3defb98ca88dcf05079815555dfc2 of base.tar.gz
INFO:
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1490809695281
BACKUP NAME: acctg_2017-03-29T13:48
BACKUP PARENT: none
BACKUP LOCATION: /opt/backup/acctg/1490809695281
BACKUP SIZE: 6.10 MB
BACKUP FORMAT: tar.gz
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 1
ChkSum File
37f3defb98ca88dcf05079815555dfc2 base.tar.gz
TABLESPACE(s): 0
START WAL LOCATION: 0000000100000000000000004
STOP WAL LOCATION: 0000000100000000000000004
BACKUP METHOD: streamed
BACKUP FROM: primary
START TIME: 2017-03-29 13:48:15 EDT
STOP TIME: 2017-03-29 13:48:17 EDT
TOTAL DURATION: 2 sec(s)
-bash-4.1$ bart BACKUP -s mktg -z
INFO: creating backup for server 'mktg'
INFO: backup identifier: '1490809751193'
61016/61016 kB (100%), 1/1 tablespace
INFO: backup completed successfully
INFO: backup checksum: 8b010e130a105e76d01346bb56dfcf14 of base.tar.gz
INFO:
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1490809751193
BACKUP NAME: mktg_2017-03-29T13:49
BACKUP PARENT: none
BACKUP LOCATION: /opt/backup/mktg/1490809751193
BACKUP SIZE: 6.13 MB
BACKUP FORMAT: tar.gz
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 1
ChkSum File
8b010e130a105e76d01346bb56dfcf14 base.tar.gz
TABLESPACE(s): 0
START WAL LOCATION: 000000010000000100000085
BACKUP METHOD: streamed
```

```
BACKUP FROM: primary
START TIME: 2017-03-29 13:49:11 EDT
STOP TIME: 2017-03-29 13:49:14 EDT
TOTAL DURATION: 3 sec(s)
-bash-4.1$ bart BACKUP -s hr -z
INFO: creating backup for server 'hr'
INFO: backup identifier: '1490809824946'
38991/38991 kB (100%), 1/1 tablespace
INFO: backup completed successfully
INFO: backup checksum: 277e8a1a80ba3474f541eb316a417c9a of base.tar.gz
INFO:
BACKUP DETAILS:
BACKUP STATUS: active
BACKUP IDENTIFIER: 1490809824946
BACKUP NAME: hr_2017-03-29T13:50
BACKUP PARENT: none
BACKUP LOCATION: /opt/backup/hr/1490809824946
BACKUP SIZE: 2.59 MB
BACKUP FORMAT: tar.gz
BACKUP TIMEZONE: US/Eastern
XLOG METHOD: fetch
BACKUP CHECKSUM(s): 1
ChkSum File
277e8a1a80ba3474f541eb316a417c9a base.tar.gz
TABLESPACE(s): 0
BACKUP METHOD: streamed
BACKUP FROM: primary
START TIME: 2017-03-29 13:50:25 EDT
STOP TIME: 2017-03-29 13:50:26 EDT
TOTAL DURATION: 1 sec(s)
```

The following code sample shows the backup directories created for each backup of each database server. The backup ID is used as the backup directory name.

```
-bash-4.1$ cd /opt/backup
-bash-4.1$ ls -l
total 12
drwxrwxr-x 4 enterprisedb enterprisedb 4096 Mar 29 13:48 acctg
drwxrwxr-x 4 enterprisedb enterprisedb 4096 Mar 29 13:50 hr
drwxrwxr-x 4 enterprisedb enterprisedb 4096 Mar 29 13:49 mktg
-bash-4.1$ ls -l acctg
drwx----- 2 enterprisedb enterprisedb 4096 Mar 29 13:48 1490809695281
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Mar 29 13:48 archived_wals
-bash-4.1$ ls -l hr
total 8
drwx----- 2 enterprisedb enterprisedb 4096 Mar 29 13:50 1490809824946
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Mar 29 13:50 archived_wals
-bash-4.1$ ls -l mktg
total 8
drwx----- 2 enterprisedb enterprisedb 4096 Mar 29 13:49 1490809751193
drwxrwxr-x 2 enterprisedb enterprisedb 4096 Mar 29 13:49 archived_wals
```

Using Point-In-Time Recovery

This section demonstrates using the point-in-time recovery operation on the remote PostgreSQL database server.

The following tables were created about two minutes apart with WAL archiving enabled:

```
postgres=# \dt
          List of relations
Schema
            Name
                 | Type
                             | Owner
-----
public | hr_rmt_t1_1356 | table
                             postgres
public | hr_rmt_t1_1358 | table
                             postgres
public | hr_rmt_t1_1400 | table
                             postgres
public | hr_rmt_t1_1402 | table
                             postgres
public | hr_rmt_t1_1404 | table
                             postgres
public | hr_rmt_t1_1406 | table
                             postgres
(6 rows)
```

In the table name <a href="hr_rmt_t<n>_<hhmi>", "n" represents the active timeline. <hhmi>" is the approximate time the table was created. For example, hr_rmt_t1_1356 was created at approximately 1:56 PM while timeline #1 is active.

The PostgreSQL database server was then stopped. WAL files that have been created, but not yet archived must be identified, and then saved. The following archived WAL files are in the BART backup catalog:

The following sample lists the current PostgreSQL server WAL files. The unarchived WAL files are marked with two stars (**).

Copies of the unarchived WAL files are saved to a temporary location:

On the remote host, a directory is created to which the PostgreSQL database cluster is to be restored. This restore path is named /opt/restore_pg96 and is owned by user account postgres.

```
[user@localhost ~]$ su root
Password:
[root@localhost user]# cd /opt
[root@localhost opt]# mkdir restore_pg96
[root@localhost opt]# chown postgres restore_pg96
[root@localhost opt]# chgrp postgres restore_pg96
[root@localhost opt]# chmod 700 restore_pg96
[root@localhost opt]# ls -l
total 16
drwxr-xr-x 4 root daemon 4096 Mar 29 12:10 PostgresPlus
drwxr-xr-x 3 root daemon 4096 Mar 29 12:25 PostgreSQL
drwx----- 2 postgres postgres 4096 Mar 29 14:15 restore_pg96
drwxr-xr-x. 2 root root 4096 Nov 22 2013 rh
```

In the BART configuration file, the remote user and remote host IP address, postgres@192.168.2.24, have been set with the remote_host parameter. If not given in the BART configuration file, this information must then be specified by the --remote-host option when giving the RESTORE subcommand (for example, bart RESTORE --remote-host postgres@192.168.2.24 ...).

```
[HR]
host = 192.168.2.24
port = 5432
user = postgres
cluster_owner = postgres
backup_name = hr_%year-%month-%dayT%hour:%minute
remote_host = postgres@192.168.2.24
copy_wals_during_restore = enabled
description = "Human Resources"
```

Use the SHOW-BACKUPS subcommand to identify the backup to use with the RESTORE subcommand.

SERVER NAME	BACKUP ID	BACKUP NAME	BACKUP PARENT			
BACKUP TIME						
BACKUP SIZE	WAL(s) SIZE	WAL FILES	STATUS			
acctg	1490809695281	acctg_2017-03-29T13:48	none			
2017-03-29 13:48:17	EDT					
6.10 MB	32.00 MB	2	active			
hr	1490809824946	hr_2017-03-29T13:50	none			
2017-03-29 13:50:26 EDT						

```
2.59 MB 32.00 MB 2 active mktg 1490809751193 mktg_2017-03-29T13:49 none 2017-03-29 13:49:14 EDT 6.13 MB 64.00 MB 4 active
```

The -t option with the SHOW-BACKUPS subcommand displays additional backup information:

```
-bash-4.1$ bart SHOW-BACKUPS -s hr -i 1490809824946 -t
SERVER NAME
            : hr
BACKUP ID
             : 1490809824946
BACKUP NAME
             : hr_2017-03-29T13:50
BACKUP PARENT : none
BACKUP STATUS : active
BACKUP TIME
            : 2017-03-29 13:50:26 EDT
BACKUP SIZE
             : 2.59 MB
WAL(S) SIZE
            : 32.00 MB
NO. OF WALS
             : 2
CREATION TIME : 2017-03-29 13:50:31 EDT
LAST WAL FILE : 000000010000000000000000
CREATION TIME : 2017-03-29 14:07:35 EDT
```

A recovery is made using timeline 1 to 2017-03-29 14:01:00.

```
-bash-4.1$ bart RESTORE -s hr -i hr_2017-03-29T13:50 -p
/opt/restore_pg96 -t 1 -g '2017-03-29 14:01:00'
INFO: restoring backup 'hr_2017-03-29T13:50' of server 'hr'
INFO: base backup restored
INFO: copying WAL file(s) to
postgres@192.168.2.24:/opt/restore_pg96/archived_wals
INFO: writing recovery settings to postgresql.auto.conf file
INFO: archiving is disabled
INFO: permissions set on $PGDATA
INFO: restore completed successfully
```

The following example shows the restored backup files in the restore path directory, /opt/restore_pg96:

```
-bash-4.1$ pwd
/opt/restore_pg96
-bash-4.1$ ls -l
total 128
drwxr-xr-x 2 postgres postgres 4096 Mar 29 14:27 archived_wals
-rw----- 1 postgres postgres 206 Mar 29 13:50 backup_label
drwx----- 5 postgres postgres 4096 Mar 29 12:25 base
drwx----- 2 postgres postgres 4096 Mar 29 14:27 global
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_clog
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_commit_ts
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg dynshmem
-rw---- 1 postgres postgres 4212 Mar 29 13:18 pg_hba.conf
-rw----- 1 postgres postgres 1636 Mar 29 12:25 pg_ident.conf
drwxr-xr-x 2 postgres postgres 4096 Mar 29 13:45 pg_log
drwx----- 4 postgres postgres 4096 Mar 29 12:25 pg_logical
drwx---- 4 postgres postgres 4096 Mar 29 12:25 pg_multixact
drwx----- 2 postgres postgres 4096 Mar 29 13:43 pg_notify
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_replslot
```

```
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_serial
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_snapshots
drwx----- 2 postgres postgres 4096 Mar 29 13:43 pg_stat
drwx----- 2 postgres postgres 4096 Mar 29 13:50 pg_stat_tmp
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_subtrans
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_tblspc
drwx----- 2 postgres postgres 4096 Mar 29 12:25 pg_twophase
-rw----- 1 postgres postgres 4 Mar 29 12:25 PG_VERSION
drwx----- 3 postgres postgres 4096 Mar 29 14:27 pg_xlog
-rw----- 1 postgres postgres 169 Mar 29 13:24 postgresql.auto.conf
-rw-r---- 1 postgres postgres 21458 Mar 29 14:27 postgresql.conf
-rw-r---- 1 postgres postgres 118 Mar 29 14:27 postgresql.auto.conf
```

Copy the saved, unarchived WAL files to the restore path pg_xlog subdirectory (/opt/restore_pg96/pg_xlog):

```
-bash-4.1$ pwd
/opt/restore_pg96/pg_xlog
-bash-4.1$ ls -l
total 16388
-rw----- 1 postgres postgres 16777216 Mar 29 13:50
0000000100000000000000000
drwx----- 2 postgres postgres 4096 Mar 29 14:27 archive_status
-bash-4.1$ ls -l /tmp/unarchived_pg96_wals
total 32768
-rw----- 1 postgres postgres 16777216 Mar 29 14:07
0000000100000000000000004
-rw----- 1 postgres postgres 16777216 Mar 29 13:50
000000010000000000000005
-bash-4.1$ cp -p /tmp/unarchived_pg96_wals/* .
-bash-4.1$ ls -l
total 49156
-rw----- 1 postgres postgres 16777216 Mar 29 13:50
0000000100000000000000000
-rw----- 1 postgres postgres 16777216 Mar 29 14:07
000000100000000000000004
-rw----- 1 postgres postgres 16777216 Mar 29 13:50
0000000100000000000000005
drwx----- 2 postgres postgres 4096 Mar 29 14:27 archive_status
```

Inspect the /opt/restore_pg96/postgresql.auto.conf file to verify that it contains the correct recovery settings:

```
restore_command = 'cp archived_wals/%f %p'
recovery_target_time = '2017-03-29 14:01:00'
recovery_target_timeline = 1
```

Note that the command restores from the archived_wals subdirectory of /opt/restore_pg96 since the copy_wals_during_restore parameter in the BART configuration file is set to enabled for database server hr.

Start the database server to initiate the point-in-time recovery operation:

```
[user@localhost ~]$ su postgres
Password:
bash-4.1$ cd /opt/restore_pg96
bash-4.1$ /opt/PostgreSQL/9.6/bin/pg_ctl start -D /opt/restore_pg96 -l
```

```
/opt/restore_pg96/pg_log/logfile
server starting
```

Inspect the database server log file to ensure the operation did not result in any errors:

```
2017-03-29 14:33:23 EDT LOG: database system was interrupted; last known
up at 2017-03-29 13:50:25 EDT
2017-03-29 14:33:23 EDT LOG: starting point-in-time recovery to
2017-03-29 14:01:00-04
2017-03-29 14:33:23 EDT LOG: restored log file
"000000010000000000000000002" from archive
2017-03-29 14:33:23 EDT LOG: redo starts at 0/2000098
2017-03-29 14:33:23 EDT LOG: consistent recovery state reached at
0/20000C0
2017-03-29 14:33:23 EDT LOG: restored log file
"0000000100000000000000003" from archive
2017-03-29 14:33:23 EDT LOG: recovery stopping before commit of
transaction 1762, time 2017-03-29 14:02:28.100072-04
2017-03-29 14:33:23 EDT LOG: redo done at 0/303F390
2017-03-29 14:33:23 EDT LOG: last completed transaction was at log time
2017-03-29 14:00:43.351333-04
cp: cannot stat `archived wals/00000002.history': No such file or
directory
2017-03-29 14:33:23 EDT LOG: selected new timeline ID: 2
cp: cannot stat `archived_wals/00000001.history': No such file or
directory
2017-03-29 14:33:23 EDT LOG: archive recovery complete
2017-03-29 14:33:23 EDT LOG: MultiXact member wraparound protections are
now enabled
2017-03-29 14:33:23 EDT LOG: database system is ready to accept
connections
2017-03-29 14:33:23 EDT LOG: autovacuum launcher started
```

The tables that exist in the recovered database cluster are:

Since recovery was up to and including 2017-03-29 14:01:00, the following tables created after 14:01 are not present:

```
public | hr_rmt_t1_1402 | table | postgres
public | hr_rmt_t1_1404 | table | postgres
public | hr_rmt_t1_1406 | table | postgres
```

The BART RESTORE operation stops WAL archiving by adding an archive_mode = off parameter at the very end of the postgresql.conf file. This last parameter in the file overrides any other previous setting of the same parameter in the file. Delete the last setting and restart the database server to start WAL archiving.

```
### Add settings for extensions here
```

5 EDB Postgres Backup and Recovery User Guide

5.1 Introduction

The EDB Backup and Recovery Tool (BART) is an administrative utility that provides simplified backup and recovery management for multiple local or remote EDB Advanced Server and PostgreSQL database servers.

BART provides the following features:

- Support for complete, hot, physical backups of multiple Advanced Servers and PostgreSQL database servers
- Support for two types of backups full base backups and block-level incremental backups
- Backup and recovery management of database servers on local or remote hosts
- A single, centralized catalog for backup data
- Retention policy support for defining and managing how long backups should be kept
- The capability to store the backup data in compressed format
- · Verified backup data with checksums
- Backup information displayed in an easy-to-read format
- A simplified point-in-time recovery process

This guide provides the following information about using BART:

- an overview of the BART components and concepts.
- backup and recovery management process.
- using tablespaces.

For information about installing BART, see the *EDB Backup and Recovery Installation and Upgrade Guide*, for examples of BART operations and subcommand usage, see the *EDB Backup and Recovery Reference Guide*. These guides are available at the *EDB website*.

Conventions Used in this Guide

The following is a list of conventions used throughout this document.

- Much of the information in this document applies interchangeably to the PostgreSQL and EDB Advanced Server database systems. The term *Advanced Server* is used to refer to EDB Advanced Server. The term *Postgres* is used to generically refer to both PostgreSQL and Advanced Server. When a distinction needs to be made between these two database systems, the specific names, PostgreSQL or Advanced Server are used.
- The installation directory of the PostgreSQL or Advanced Server products is referred to as POSTGRES_INSTALL_HOME:
 - For PostgreSQL Linux installations, this defaults to /opt/PostgreSQL/<x.x> for version 10 and earlier. For later versions, the installation directory is /var/lib/pgsql/<x>.

For Advanced Server Linux installations performed using the interactive installer for version 10 and earlier, this defaults to /opt/PostgresPlus/<x.x>AS or /opt/edb/as<x.x>. For Advanced Server Linux installations performed with an RPM package, this defaults to /usr/ppas-<x.x> or /usr/edb/as<x.x>.
 For Advanced Server Linux installations performed with an RPM package for version 11 or later, this defaults to /usr/edb/as<xx>.

Restrictions on pg_basebackup

BART takes full backups using the pg_basebackup utility program under the following conditions:

- The backup is taken on a standby server.
- The --with-pg_basebackup option is specified with the BACKUP subcommand (see Backup).
- The number of thread count in effect is 1, and the with-pg_basebackup option is not specified with the BACKUP subcommand.
- Database servers can only be backed up using pg_basebackup utility program of the same or later version than the database server version.

In the global section of the BART configuration file, the pg_basebackup_path parameter specifies the complete directory path to the pg_basebackup program. For information about the pg_basebackup_path parameter and the thread_count, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.

For information about pg_basebackup, see the PostgreSQL Core Documentation.

5.2 Overview

BART provides a simplified interface for the continuous archiving and point-in-time recovery method provided with Postgres database servers. This consists of the following processes:

- Capturing a complete image of a database cluster as a full base backup or referred to simply as a full backup.
- Capturing a modified image of a database cluster called a *block-level incremental backup* or referred as *incremental backup*, which is similar to a full backup, but contains the modified blocks of the relation files that have been changed since a previous backup.
- Archiving the Write-Ahead Log segments (WAL files), which continuously record changes to be made to the database files.
- Performing *Point-In-Time Recovery* (PITR) to a specified transaction ID or timestamp with respect to a timeline using a full backup along with successive, block-level incremental backups that reside in the same backup chain, and the WAL files.

Detailed information regarding WAL files and point-in-time recovery is documented in the PostgreSQL Core Documentation.

The general term *backup* refers to both full backups and incremental backups.

When taking a full backup of a standby server, BART uses the PostgreSQL pg_basebackup utility program. However, it must be noted that for standby servers, you can only take a full backup, but cannot take an incremental or parallel backups. For information about standby servers, see the PostgreSQL Documentation.

BART uses a centralized backup catalog, a single configuration file, and a command line interface controlling the necessary operations to simplify the management process. Reasonable defaults are automatically used for various

backup and restore options. BART also performs the necessary recovery file configuration required for point-in-time recovery using its command line interface.

BART also provides the following features to enhance backup management:

- Automation of the WAL archiving command configuration.
- Usage of retention policies to evaluate, categorize, and delete obsolete backups.
- Compression of WAL files to conserve disk space.
- Customizable naming of backups to ease their usage.
- Easy access to comprehensive information about each backup.

The key components of a BART installation are:

- BART Host. The host system on which BART is installed. BART operations are invoked from this host system. The database server backups and archived WAL files are stored on this host as well.
- BART User Account. Linux operating system user account you choose to run BART. The BART user account owns the BART backup catalog directory.
- BART Configuration File. File in editable text format containing the configuration information that BART uses.
- BART Backup Catalog. File system directory structure containing all of the backups for the database servers that BART manages. It is also the default archive_path to store archived WAL files.
- BART Backupinfo File. File in text format containing information for a BART backup. A backupinfo file resides in each backup subdirectory within the BART backup catalog.
- BART Command Line Utility Program. Single, executable file named bart, which is used to manage all BART operations.
- BART WAL Scanner Program. Single, executable file named bart-scanner, which is used to scan WAL files to locate and record the modified blocks for incremental backups.

Other concepts and terms referred to in this document include the following:

Postgres Database Cluster. Also commonly called the data directory, this is the file system directory where all of the data files related to a particular Postgres database server instance are stored. (Each specific running instance is identified by its host and port number when connecting to a database.) The database cluster is identified by the -D option when it is created, started, stopped, etc. by the Postgres initdb and pg_ctl commands. A full backup is a copy of a database cluster.

The terms database cluster and database server are used somewhat interchangeably throughout this document, though a single database server can run multiple database clusters.

- Postgres User Account. Linux operating system user account that runs the Advanced Server or PostgreSQL database server and owns the database cluster directory.
 - By default, the database user account is enterprised when Advanced Server is installed to support compatibility with Oracle databases.
 - By default, the database user account is **postgres** when Advanced Server is installed in PostgreSQL compatible mode. For a PostgreSQL database server, the default database user account is also **postgres**.
 - The BART configuration parameter cluster_owner must be set to the database user account for each database server.
- Replication Database User. For each database server that BART manages, a database superuser must be selected to
 act as the replication database user. This database user is used to connect to the database server when backups are
 taken. The database superusers created with an initial Postgres database server installation (enterprisedb or
 postgres) may be used for this purpose. The BART configuration parameter user must be set to this replication
 database user for each database server.
- Secure Shell (SSH)/Secure Copy (SCP). Linux utility programs used to log into hosts (SSH) and copy files (SCP)

between hosts. A valid user account must be specified that exists on the target host and in fact is the user account under which the SSH or SCP operations occur.

For information on how all of these components are configured and used with BART, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.

Supported BART Operations

The following tables are not a conclusive list of the scenarios supported by BART, but instead provides an overview of some of the most common scenarios in both pg_basebackup (thread count=1) as well as parallel backup mode (thread count greater than 1).

	-Fp-xlog- method=fetch	-Fp-xlog- method=stream	-Ft-xlog- method=fetch	-Ft-xlog- method=stream
Primary Database Server/Full backup	Supported	Supported	Supported	Supported
Primary Database Server/Incremental backup	Supported	Supported	Not Supported	Not Supported
Standby Database Server/Full backup	Supported	Supported	Supported	Supported
Standby Database Server/Incremental backup	Not Supported	Not Supported	Not Supported	Not Supported

Backup

	Wal compression by BART	WAL scanner
Primary Database Server/Full backup	Supported	N/A
Primary Database Server/Incremental backup	Not Supported	N/A
Standby Database Server/Full backup	Supported	N/A
Standby Database Server/Incremental backup	Not Supported	N/A

Wal Archiving

	Wal compression = enabled	Wal compression = disabled
Restore	Supported	Supported
Parallel restore	Supported	Supported

Restore

5.2.1 Block-Level Incremental Backup

This section describes the basic concepts of a block-level incremental backup (referred to as an incremental backup). An incremental backup is a unique functionality of BART.

An incremental backup provides a number of advantages when compared to using a full backup:

- The amount of time required to produce an incremental backup is generally less than a full backup, as modified relation blocks are saved instead of all full relation files of the database cluster.
- An incremental backup uses less disk space than a full backup.

Generally, all BART features (such as retention policy management) apply to incremental backups and full backups. See Managing Incremental Backups for more information.

5.2.1.1 Incremental Backup Limitations and Requirements

The following limitations apply to incremental backup:

- If you have restored a full or incremental backup, you must take a full backup before enabling incremental backup.
- If a standby node has been promoted to the role of a primary node, you must take a full backup before enabling incremental backup on the cluster.
- On a standby database server, you cannot take an incremental backup.

You must meet the following requirements before implementing incremental backup:

- You must create or select an operating system account to be used as the BART user account.
- You must create or select the replication database user, which must be a superuser.
- In the BART configuration file:
 - You must set the cluster_owner parameter to the Linux operating system user account that owns the database cluster directory from which incremental backups are to be taken.
 - You must enable the allow_incremental_backups parameter.
- A passwordless SSH/SCP connection must be established between the BART user account on the BART host and the cluster_owner user account on the database server.

It must be noted that a passwordless SSH/SCP connection must be established even if BART and the database server are running on the same host and the BART user account and the cluster_owner user account are the same account.

- In addition to the BART host (where the BART backup catalog resides), the BART package must also be installed on every remote database server on which incremental backups are to be restored. To restore an incremental backup, the bart program must be executable on the remote host by the remote user (the remote user is specified by the remote_host parameter in the BART configuration file or by the -r option when using the RESTORE subcommand to restore the incremental backup).
- When restoring incremental backups on a remote database server, a passwordless SSH/SCP connection must be established from the BART user account on the BART host to the remote user on the remote host (the remote user is specified by the remote_host parameter in the BART configuration file or by the -r option when using the RESTORE subcommand to restore the incremental backup).
- Compression of archived WAL files in the BART backup catalog is not permitted for database servers on which
 incremental backups are to be taken. The wal_compression setting in the BART configuration file must be
 disabled for those database servers.
- The incremental backup must be on the same timeline as the parent backup. The timeline changes after each recovery operation so an incremental backup cannot use a parent backup from an earlier timeline.

For information about configuring these requirements, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.

The following section provides an overview of the basic incremental backup concepts.

5.2.1.2 Concept Overview

Using incremental backups involves the following sequence of steps:

1. Configure BART, and enable and initiate archiving of WAL files to the archive_path in the same manner as done for full backups.

The default archive_path is the BART backup catalog (<backup_path>/<server_name>/archived_wals). Using the <archive_path> parameter in the server section of the BART configuration file, you can specify the location where WAL files will be archived.

For more information about the archive_path parameter and configuring BART, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.

- 2. Take an initial full backup with the BACKUP subcommand. This full backup establishes the parent of the first incremental backup.
- 3. Scan all WAL files produced by database servers on which incremental backups are to be taken. These WAL files are scanned once they have been archived to the archive_path.
 - Each scanned WAL file results in a modified block map (MBM) file that records the location of modified blocks obtained from the corresponding WAL file. The BART WAL scanner program bart-scanner performs this process.
- 4. Take incremental backups using the BACKUP subcommand with the --parent option to specify the backup identifier or name of a previous, full backup or an incremental backup. Any previous backup may be chosen as the parent as long as all backups belong to the same timeline.
- 5. The incremental backup process identifies which WAL files may contain changes from when the parent backup was taken to the starting point of the incremental backup. The corresponding MBM files are used to locate and copy the modified blocks to the incremental backup directory along with other database cluster directories and files. Instead of backing up all, full relation files, only the modified blocks are copied and saved. In addition, the relevant MBM files are condensed into one consolidated block map (CBM) file that is stored with the incremental backup.
 - Multiple block copier threads can be used to copy the modified blocks to the incremental backup directory. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for information about setting the thread_count parameter in the BART configuration file. See Backup for information about using the thread_count option with the BACKUP subcommand.
- 6. Invoke the restore process for an incremental backup using the RESTORE subcommand in the same manner as restoring a full backup. The -i option specifies the backup identifier or name of the incremental backup to restore. The restore process begins by going back through the chain of past, parent incremental backups until the initial full backup starting the chain is identified. This full backup provides the initial set of directories and files to be restored to the location specified with the -p option. Each subsequent incremental backup in the chain is then restored. Restoration of an incremental backup uses its CBM file to restore the modified blocks from the incremental backup.

The following sections provide some additional information on these procedures.

5.2.1.3 WAL Scanning – Preparation for an Incremental Backup

The WAL scanner program (bart-scanner) scans the WAL files created from the time of the parent backup up to the start of the incremental backup to determine which blocks have modified since the parent backup, and records the information in a file called the *modified block map (MBM) file*. One MBM file is created for each WAL file.

The MBM file is stored in the directory where archived_wals will be stored, as specified in the archive_path parameter in the bart.cfg file. If the archive_path is not specified, the default archived_wals directory is:

```
<backup_path>/<server_name>/<archived_wals>
```

Where:

<backup_path> is the BART backup catalog parent directory specified in the global section of the BART configuration
file.

<server_name> is the lowercase conversion of the database server name specified in the server section of the BART configuration file.

The following code snippet is the content of the archive path showing the MBM files created for the WAL files. (The user name and group name of the files have been removed from the example to list the WAL files and MBM files in a more comparable manner):

```
[root@localhost archived_wals]# pwd
/opt/backup/acctg/archived_wals
[root@localhost archived_wals]# ls -l
total 131104
-rw----- 1 ... 16777216 Oct 12 09:38 000000010000000100000078
-rw----- 1 ... 16777216 Oct 12 09:38 000000010000000100000079
-rw----- 1 ... 16777216 Oct 12 09:38 00000001000000010000007A
-rw----- 1 ... 16777216 Oct 12 09:35 00000001000000010000007B
-rw----- 1 ... 16777216 Oct 12 09:38 00000001000000010000007C
-rw----- 1 ... 16777216 Oct 12 09:39 00000001000000010000007D
-rw----- 1 ... 16777216 Oct 12 09:42 0000000100000010000007E
-rw----- 1 ... 16777216 Oct 12 09:47 0000000100000010000007F
-rw-rw-r-- 1 ... ... 161 Oct 12 09:49 0000000100000001780000280000000179000000.mbm
-rw-rw-r-- 1 ... ... 684 Oct 12 09:49 000000010000000179000028000000017A000000.mbm
-rw-rw-r-- 1 ... 161 Oct 12 09:49 00000001000000017A000028000000017B000000.mbm
-rw-rw-r-- 1 ... 161 Oct 12 09:49 00000001000000017B000028000000017C000000.mbm
-rw-rw-r-- 1 ... 1524 Oct 12 09:49 00000001000000017C000028000000017D000000.mbm
-rw-rw-r-- 1 ... ... 161 Oct 12 09:49 00000001000000017D000028000000017E000000.mbm
-rw-rw-r-- 1 ... 161 Oct 12 09:49 00000001000000017E000028000000017F000000.mbm
-rw-rw-r-- 1 ... 161 Oct 12 09:49 00000001000000017F0000280000000180000000.mbm
```

MBM files have the suffix, .mbm.

In preparation for any incremental backup, the WAL files should be scanned as soon as they are copied to the archive_path. Thus, the WAL scanner should be running as soon as the WAL files from the database cluster are archived to the archive_path. If the archive_path contains WAL files that have not yet been scanned, starting the WAL scanner begins scanning these files. If WAL file fails to be scanned (resulting in a missing MBM file), you can use the WAL scanner to specify an individual WAL file.

Under certain conditions such as when the Network File System (NFS) is used to copy WAL files to the archive_path,

the WAL files may have been missed by the WAL scanner program for scanning and creation of MBM files. Use the scan_interval parameter in the BART configuration file to initiate force scanning of WAL files in the archive_path to ensure MBM files are generated. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for more information about the scan_interval parameter.

See Running the BART WAL Scanner for information about using the WAL scanner.

5.2.1.4 Performing an Incremental Backup

The WAL files produced at the time of the parent backup up to the start of the incremental backup contain information about which blocks were modified during that time interval. That information is consolidated into an MBM file for each WAL file by the WAL scanner.

The MBM files for the relevant WAL files are read, and the information is used to copy the modified blocks from the database cluster to the archived_wals directory as specified in the archive_path parameter in the bart.cfg file. When compared to a full backup, the number and sizes of relation files can be significantly less for an incremental backup.

For comparison, the following is an abbreviated list of the files copied to the archived base subdirectory of a full backup for one database:

<pre>[root@localhost 14845]# pwd /opt/backup/acctg/1476301238969/base/base/14845 [root@localhost 14845]# ls</pre>								
112	13182_vm	14740	16467	16615	2608_vm	2655	2699	2995
113	13184	14742	16471	174	2609	2656	2701	2995_vm
1247	13186	14745	16473	175	2609_fsm	2657	2702	2996
 1247_fsm	13187	14747	16474	2187	2609_vm	2658	2703	2998
 1247_vm	13187_fsm	14748	16476	2328	2610	2659	2704	2998_vm
1249	13187_vm	14749	16477	2328_fsm	2610_fsm	2660	2753	2999
1249_fsm	13189	14752	16479	2328_vm	2610_vm	2661	2753_fsm	2999_vm
 1249_∨m	13191	14754	16488	2336	2611	2662	2753_vm	3079
1255	13192	14755	16490	2336_vm	2611_vm	2663	2754	3079_fsm
•••				•				
13182_fsm	14739	16465	16603	2608_fsm	2654	2696	2893_vm	3501_vm
• • •								

In contrast, the following is the content of the archived base subdirectory of the same database from a subsequent incremental backup:

```
[root@localhost 14845]# pwd
/opt/backup/acctg/1476301835391/base/base/14845
[root@localhost 14845]# ls
          1249
                    1259
                              16384 17006 2608
1247
                                                      2610
                                                                 2658
                                                                       2663
                                                                             2678
. . .
1247_fsm 1249_fsm 1259_fsm
                                     17009
                                            2608_fsm
                                                      2610_fsm
                              16387
                                                                2659
                                                                       2673
                                                                             2679
1247_vm
          1249_vm
                    1259_vm
                              16389
                                     17011
                                            2608_vm
                                                      2610_vm
                                                                 2662
                                                                       2674 2703
```

The information from the MBM files are consolidated into one file called a *consolidated block map* (CBM) file. During the restore operation for the incremental backup, the CBM file is used to identify the modified blocks to be restored for that backup. In addition, the incremental backup also stores other required subdirectories and files from the database cluster as is done for full backups.

Before taking an incremental backup, an initial full backup must be taken with the BACKUP subcommand. This full backup establishes the parent of the first incremental backup.

The syntax for taking a full backup is:

```
bart BACKUP -s { <server_name> | all } [ -F { p | t } ]
  [ -z ] [ -c <compression_level> ]
  [ --backup-name <backup_name> ]
  [ --thread-count <number_of_threads> ]
  [ { --with-pg_basebackup | --no-pg_basebackup } ]
  [-checksum-algorithm ]
```

The syntax for taking an incremental backup is:

```
bart BACKUP -s { <server_name> | all } [ -F p]
[ --parent { <backup_id> | <backup_name> } ]
[ --backup-name <backup_name> ]
[ --thread-count <number_of_threads> ]
[ --check ]
[ --checksum-algorithm ]
```

You must specify the following before taking an incremental backup:

- -Fp option for plain text format as incremental backup can only be taken in the plain text format.
- --check option to verify if the required MBM files are present in the archived_wals directory. The --parent option must be specified when the --check option is used.

See BACKUP for more information about using the BACKUP subcommand.

5.2.1.5 Restoring an Incremental Backup

Restoring an incremental backup may require additional steps depending upon the host on which the incremental backup is to be restored:

• Restoring an Incremental Backup on a BART Host - This section outlines restoring an incremental backup onto the same host where BART has been installed.

 Restoring an Incremental Backup on a Remote Host - This section outlines restoring an incremental backup onto a remote host where BART has not been installed.

Restoring an Incremental Backup on a BART Host

Specify a backup identifier or name, and include the -i option when invoking the RESTORE subcommand to restore an incremental backup. All RESTORE options may be used in the same manner as when restoring a full backup. See RESTORE command for more details.

First, all files from the full backup from the beginning of the backup chain are restored. For each incremental backup, the CBM file is used to identify and restore blocks from the incremental backup. If there are new relations or databases identified in the CBM file, then relevant relation files are copied. If consolidated block map information is found indicating the drop of a relation or a database, then the relevant files are removed from the restore directory. Similarly, if there is any indication of a table truncation, then the related files are truncated.

Also note that you can use the —w option of the RESTORE subcommand to specify a multiple number of parallel worker processes to stream the modified blocks to the restore host.

!!! Note If you set the BART scanner or backup with the --checksum-algorithm=NONE option, then you must specify the --disable checksum option while restoring an incremental backup.

Restoring an Incremental Backup on a Remote Host

Ensure the bart program is available on the remote host when restoring an incremental backup on a remote host; the invocation of the RESTORE subcommand for an incremental backup results in the execution of the bart program on the remote host to restore the modified blocks to their proper location.

To restore an incremental backup onto a remote host where BART has not been installed, perform the following steps:

Step 1: Install BART on the remote host to which an incremental backup is to be restored.

No editing is needed in the bart.cfg file installed on the remote host.

Step 2: Determine the Linux operating system user account on the remote host to be used as the remote user. This user is specified by the remote_host parameter in the BART configuration file or by the -r option when using the RESTORE subcommand to restore the incremental backup. The remote user must be the owner of the directory where the incremental backup is to be restored on the remote host. By default, the user account is enterprised for Advanced Server or postgres for PostgreSQL.

Step 3: Ensure a passwordless SSH/SCP connection is established from the BART user on the BART host to the remote user on the remote host. For information about creating a passwordless SSH/SCP connection, see the *EDB Backup and Recovery Installation and Upgrade Guide*, available at the EDB website.

When restoring an incremental backup, specify the backup identifier or name of the incremental backup that will be restored. See the RESTORE documentation for more details. To view an example of restoring an incremental backup, see the EDB Backup and Recovery Reference Guide available at the EDB website.

!!! Note If you set the BART scanner or backup with the --checksum-algorithm=NONE option, then you must specify the --disable checksum option while restoring an incremental backup.

5.2.2 Creating a Backup Chain

A backup chain is the set of backups consisting of a full backup and all of its successive incremental backups. Tracing back on the parent backups of all incremental backups in the chain eventually leads back to that single, full backup.

It is possible to have a *multi-forked* backup chain, which is two or more successive lines of incremental backups, all of which begin with the same, full backup. Thus, within the chain there is a backup that serves as the parent of more than one incremental backup.

Since restoration of an incremental backup is dependent upon first restoring the full backup, then all successive incremental backups up to, and including the incremental backup specified by the to note the following:

RESTORE

- Deletion or corruption of the full backup destroys the entire backup chain. It is not possible to restore any of the incremental backups that were part of that chain.
- Deletion or corruption of an incremental backup within the chain results in the inability to restore any incremental backup that was added to that successive line of backups following the deleted or corrupted backup. The full backup and incremental backups prior to the deleted or corrupted backup can still be restored.

The actions of retention policy management are applied to the full backup and all of its successive incremental backups within the chain in an identical manner as if they were one backup. Thus, use of retention policy management does not result in the breakup of a backup chain.

See the *EDB Backup and Recovery Reference Guide*, available at the *EDB website* for examples of creating a backup chain and restoring an incremental backup.

5.3 Using BART

After installing and configuring the BART host and the database servers, you can start using BART.

This section describes how to perform backup and recovery management operations using BART. Review the sections that follow before proceeding with any BART operation.

5.3.1 BART Management Overview

After configuring BART, you can begin the backup and recovery management process. The following steps will help you get started:

- 1. Run the CHECK-CONFIG subcommand without the -s option. When the CHECK-CONFIG subcommand is used without specifying the -s option, it checks the parameters in the global section of the BART configuration file.
- 2. Run the INIT subcommand (if you have not already done so) to finish creation of the BART backup catalog, which results in the complete directory structure to which backups and WAL files are saved. This step must be done before restarting the database servers with enabled WAL archiving, otherwise the copy operation in the archive_command parameter of the postgresql.conf file or the postgresql.auto.conf file fails due

- to the absence of the target archive directory. When the directory structure is complete, the archived_wals subdirectory should exist for each database server.
- 3. Start the Postgres database servers with archiving enabled. Verify that the WAL files are appearing in the archive_path. The archiving frequency is dependent upon other postgresql.conf configuration parameters. Check the Postgres database server log files to ensure there are no archiving errors. Archiving should be operational before taking a backup in order to ensure that the WAL files that may be created during the backup process are archived.
- 4. Start the WAL scanner if you intend to take incremental backups. Since the WAL scanner processes the WAL files copied to the archive path, it is advantageous to commence the WAL scanning as soon as the WAL files begin to appear in the archive_path in order to keep the scanning in pace with the WAL archiving.
- 5. Run the BART CHECK-CONFIG subcommand for each database server with the -s option specifying the server name. This ensures the database server is properly configured for taking backups.
- 6. Create a full backup for each database server. The full backup establishes the starting point of when point-in-time recovery can begin and also establishes the initial parent backup for any incremental backups to be taken.

There are now a number of other BART management processes you may perform:

- Execute the BACKUP subcommand to create additional full backups or incremental backups.
- Use the VERIFY-CHKSUM subcommand to verify the checksum of the full backups.
- Display database server information with the SHOW-SERVERS subcommand.
- Display backup information with the SHOW-BACKUPS subcommand.
- Compress the archived WAL files in the archive_path by enabling WAL compression in the BART configuration file and then invoking the MANAGE subcommand.
- Determine and set the retention policy for backups in the BART configuration file.
- Establish the procedure for using the MANAGE subcommand to enforce the retention policy for backups. This may include using cron jobs to schedule the MANAGE subcommand.

5.3.1.1 Performing a Restore Operation

The following steps describe the process of restoring a backup:

Step 1: Use your system-specific command to shut down the database server.

Step 2: Inspect the pg_wal subdirectory (inspect the pg_xlog subdirectory if you are using server 9.6 version) of the data directory and save any WAL files that have not yet been archived to the archive_path. If there are files that have not been archived, save them to a temporary location.

Step 3: If you want to restore to current data directory, it is recommended to make a copy of the current data directory and then delete all files and subdirectories under the data directory if you have enough extra space. If there is not enough space, then make a copy of pg_wal directory (or pg_xlog if you are using server 9.6 version) until the server is successfully restored.

If you want to restore to a new, empty directory, create the directory on which you want to restore the backed up database cluster. Ensure the data directory can be written to by the BART user account or by the user account specified by the remote_host configuration parameter, or by the --remote-host option of the RESTORE subcommand (if

these are to be used).

Step 4: Perform the same process for tablespaces as described in Step 3. The tablespace_path parameter in the BART configuration file must contain the tablespace directory paths to which the tablespace data files are to be restored. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for more information about this parameter.

Step 5: Identify the backup to use for the restore operation and obtain the backup ID or backup name.

To use the latest backup, omit the -i option; the RESTORE subcommand uses that backup by default. The backups can be listed with the SHOW-BACKUPS subcommand.

Step 6: Run the BART RESTORE subcommand.

- Minimal recovery settings will be saved in the postgresql.auto.conf file and archive recovery will proceed
 only until consistency is reached, with no restoration of files from the archive. See Restore for detailed information
 about Restore subcommand.
- If the -c option is specified or if the copy_wals_during_restore BART configuration parameter is enabled for this database server, then the following actions occur:
 - In addition to restoring the database cluster to the directory specified by the -p restore_path option, the archived WAL files of the backup are copied from the BART backup catalog to the subdirectory restore_path/archived_wals.
 - If recovery settings are saved in the postgresql.auto.conf file, the command string set in the
 restore_command parameter retrieves the WAL files from this archived_wals subdirectory relative to the
 restore_path parent directory as: restore_command = cp archived_wals/%f %p

You must ensure that valid options are specified when using the **RESTORE** subcommand. BART will not generate an error message if invalid option values or invalid option combinations are provided. BART will accept the invalid options and pass them to the **postgresql.auto.conf** file, which will then be processed by the database server when it is restarted.

Step 7: Copy any saved WAL files from Step 2 to the restore_path/pg_xlog subdirectory.

Step 8: Inspect the restored directories and data files of the restored database cluster in directory restore_path.

All files and directories must be owned by the user account that you intend to use to start the database server.

Recursively change the user and group ownership of the restore_path directory, its files, and its subdirectories if necessary. There must only be directory access privileges for the user account that will start the database server. No other groups or users can have access to the directory.

Step 9: The postgresql.auto.conf file should be configured to recover only until the cluster reaches consistency. In either case, the settings may be modified as desired.

Step 10: Disable WAL archiving at this point. The BART RESTORE subcommand adds archive_mode = off to the end of the postgresql.conf file.

- If you want to restart the database server with WAL archiving enabled, ensure that this additional parameter is deleted.
- The original archive_mode parameter still resides in the postgresql.conf file in its initial location with its last setting.

Step 11: Start the database server to initiate recovery. After completion, check the database server log file to ensure the recovery was successful.

If the backup is restored to a different location than where the original database cluster resided, operations dependent

upon the database cluster location may fail if supporting service scripts are not updated to reflect the location where the backup has been restored. For information about the use and modification of service scripts, see the EDB Advanced Server Installation Guide available at the EDB website.

See Restore for more information about using the BART Restore subcommand.

An example of a restore operation is documented in the EDB Backup and Recovery Reference Guide available at the EDB website.

!!! Note If you set the backup --checksum-algorithm=NONE option, then you must specify the --disable checksum option while restoring a backup.

5.3.1.2 Point-In-Time Recovery Operation

The following steps outline how to perform a point-in-time recovery operation for a database cluster:

- 1. Use your system-specific command to shut down the database server.
- 2. If you want to: a. restore the database cluster and tablespace files to new, empty directories, create the new directories with the appropriate directory ownership and permissions. b. reuse the existing database cluster directories, delete all the files and subdirectories in the existing directories. We strongly recommend that you make a copy of this data before deleting it. Be sure to save any recent WAL files in the pg_wal subdirectory (pg_xlog subdirectory if you are using server 9.6 version) that have not been archived to archive_path.
- 3. Run the BART SHOW-BACKUPS -s <server_name> subcommand to list the backup IDs and backup names of the backups for the database server. You will need to provide the appropriate backup ID or backup name with the BART RESTORE subcommand, unless you intend to restore the latest backup in which case the -i option of the RESTORE subcommand for specifying the backup ID or backup name may be omitted.
- 4. Run the BART RESTORE subcommand with the appropriate options.
 - The backup is restored to the directory specified by the -p restore_path option.
 - In addition, if the RESTORE subcommand -c option is specified or if the enabled setting of the copy_wals_during_restore BART configuration parameter is applicable to the database server, then the required archived WAL files from the archive_path are copied to the restore_path/archived_wals subdirectory.
 - Ensure the restore_path directory and all subdirectories and files in the restore_path are owned by the proper Postgres user account (for example, enterprisedb or postgres). Also ensure that only the Postgres user account has access permission to the restore_path directory.

Use the chown command to make the appropriate adjustments to file permissions; for example, the following command changes the ownership of restore_path to enterprisedb:

```
chown -R enterprisedb:enterprisedb restore_path
```

The following command restricts access to restore_path:

chmod 700 restore_path

- 5. Copy any saved WAL files from Step 2 that were not archived to the BART backup catalog to the restore_path/pg_wal subdirectory (pg_xlog subdirectory if you are using server 9.6 version).
- 6. Identify the timeline ID you wish to use to perform the restore operation.

The available timeline IDs can be identified by the first non-zero digit of the WAL file names reading from left to right.

7. Verify that the postgresql.auto.conf file created in the directory specified with the RESTORE subcommand's -p <restore_path> option was generated with the correct recovery parameter settings.

If the RESTORE subcommand -c option is specified or if the enabled setting of the copy_wals_during_restore BART configuration parameter is applicable to the database server, then the restore_command parameter retrieves the archived WAL files from the <restore_path>/archived_wals subdirectory that was created by the RESTORE subcommand, otherwise the restore_command retrieves the archived WAL files from the BART backup catalog.

- 8. The BART RESTORE subcommand disables WAL archiving in the restored database cluster. If you want to immediately enable WAL archiving, modify the postgresql.conf file by deleting the archive_mode = off parameter that BART appends to the end of the file.
- 9. Start the database server, which will then perform the point-in-time recovery operation if recovery settings are saved in the postgresql.auto.conf file.

For a detailed description of the RESTORE subcommand, see Basic BART Subcommand Usage. An example of a Point-in-Time Recovery operation is documented in the EDB Backup and Recovery Reference Guide available at the EDB website. See Restore for more information about using the Restore subcommand.

5.3.2 Managing Backups Using a Retention Policy

Over the course of time when using BART, the number of backups can grow significantly. This ultimately leads to a large consumption of disk space unless an administrator periodically performs the process of deleting the oldest backups that are no longer needed. This process of determining when a backup is old enough to be deleted and then actually deleting such backups can be done and automated with the following basic steps:

- 1. Determine and set a retention policy in the BART configuration file. A *retention policy* is a rule that determines when a backup is considered obsolete. The retention policy can be applied globally to all servers, but each server can override the global retention policy with its own.
- 2. Use the MANAGE subcommand to categorize and manage backups according to the retention policy.
- 3. Create a cron job to periodically run the MANAGE subcommand to evaluate the backups and then list and/or delete the obsolete backups.

Retention policy management applies differently to incremental backups than to full backups. See Managing Incremental Backups for information about how retention policy management is applied to each backup type.

The following sections describe how retention policy management generally applies to backups, and its specific usage and effect on full backups.

5.3.2.1 Overview - Managing Backups Using a Retention Policy

The BART retention policy results in the categorization of each backup in one of three statuses – *active*, *obsolete*, and *keep*.

- Active. The backup satisfies the retention policy applicable to its server. Such backups would be considered necessary to ensure the recovery safety for the server and thus should be retained.
- Obsolete. The backup does not satisfy the retention policy applicable to its server. The backup is no longer considered necessary for the recovery safety of the server and thus can be deleted.
- Keep. The backup is to be retained regardless of the retention policy applicable to its server. The backup is considered vital to the recovery safety for the server and thus should not be deleted for an indefinite period of time.

There are two types of retention policies - redundancy retention policy and recovery window retention policy.

- Redundancy Retention Policy The redundancy retention policy relies on a specified, maximum number of most recent backups to retain for a given server. When the number of backups exceeds that maximum number, the oldest backups are considered obsolete (except for backups marked as keep).
- Recovery Window Retention Policy The recovery window retention policy relies on a time frame (the recovery window) for when a backup should be considered active. The boundaries defining the recovery window are the current date/time (the ending boundary of the recovery window) and the date/time going back in the past for a specified length of time (the starting boundary of the recovery window).
 - If the date/time the backup was taken is within the recovery window (that is, the backup date/time is on or after the starting date/time of the recovery window), then the backup is considered active, otherwise it is considered obsolete (except for backups marked as keep).
 - Thus, for the recovery window retention policy, the recovery window time frame dynamically shifts, so the end of
 the recovery window is always the current date/time when the MANAGE subcommand is run. As you run the
 MANAGE subcommand at future points in time, the starting boundary of the recovery window moves forward in
 time. At some future point, the date/time of when a backup was taken will be earlier than the starting boundary of
 the recovery window. This is when an active backup's status will be considered obsolete.
 - You can see the starting boundary of the recovery window at any point in time by running the SHOW-SERVERS subcommand. The RETENTION POLICY field of the SHOW-SERVERS subcommand displays the starting boundary of the recovery window.

5.3.2.2 Marking the Backup Status

When a backup is initially created with the BACKUP subcommand, it is always recorded with active status. Use the MANAGE subcommand to evaluate if the backup status should be changed to obsolete in accordance with the retention policy. You can review the current status of your backups with the SHOW-BACKUPS subcommand.

Active backups are evaluated and also marked (that is, internally recorded by BART) as obsolete only when the MANAGE subcommand is invoked either with no options or with only the -s option.

Once a backup has been marked as obsolete, you cannot change it back to active unless you perform the following steps:

- Use the MANAGE subcommand with the -c option along with the backup identifier or name with the -i option. To keep this particular backup indefinitely, use -c keep, otherwise use -c nokeep.
- If you use the -c nokeep option, the backup status is changed back to active. When the MANAGE subcommand is used the next time, the backup is re-evaluated to determine if its status needs to be changed back to obsolete based on the current retention policy in the BART configuration file.

After setting the retention_policy parameter and running the MANAGE subcommand if you change the retention_policy parameter, the current, marked status of the backups are probably inconsistent with the new retention_policy setting. To modify the backup status to be consistent with the new retention_policy setting, you need to run the MANAGE subcommand with:

- the -c nokeep option to change the obsolete status to active status if there are backups currently marked as obsolete that would no longer be considered obsolete under a new retention policy. You can also specify the all option to change all backups back to active status, including those currently marked as keep.
- no options or with only the -s option to reset the marked status based on the new retention_policy setting in the BART configuration file.

See MANAGE for usage information for the MANAGE subcommand.

5.3.2.3 Setting the Retention Policy

The retention policy is determined by the retention_policy parameter in the BART configuration file. It can be applied globally to all servers, but each server can override the global retention policy with its own. For information about creating a global retention policy and an individual database server retention policy, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website.

There are two types of retention policies - redundancy retention policy and the recovery window retention policy as described in the following sections.

Redundancy Retention Policy

To use the redundancy retention policy, set retention_policy = max_number BACKUPS where max_number is a positive integer designating the maximum number of most recent backups.

Additional Restrictions:

- The keyword BACKUPS must always be specified in plural form (for example, 1 BACKUPS).
- BART will accept a maximum integer value of 2,147,483,647 for max_number; however, you should use a realistic, practical value based on your system environment.

The redundancy retention policy is the default type of retention policy if all keywords BACKUPS, DAYS, WEEKS, and MONTHS following the max_number integer are omitted as shown by the following example:

```
retention_policy = 3
```

In the following example, the redundancy retention policy setting considers the three most recent backups as the active backups. Any older backups, except those marked as keep, are considered obsolete:

```
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
archive_command = 'cp %p %a/%f'
retention_policy = 3 BACKUPS
description = "Accounting"
```

The SHOW-SERVERS subcommand displays the 3 Backups redundancy retention policy in the RETENTION POLICY field:

-bash-4.1\$ bart SHOW-SERVERS -s acctg

SERVER NAME : acctg
HOST NAME : 127.0.0.1
USER NAME : enterprisedb

PORT : 5444

REMOTE HOST

RETENTION POLICY : 3 Backups
DISK UTILIZATION : 627.04 MB

NUMBER OF ARCHIVES : 25

ARCHIVE PATH : /opt/backup/acctg/archived_wals

ARCHIVE COMMAND : cp %p /opt/backup/acctg/archived_wals/%f

XLOG METHOD : fetch
WAL COMPRESSION : disabled

TABLESPACE PATH(s) :

DESCRIPTION : "Accounting"

Recovery Window Retention Policy

To use the recovery window retention policy, set the retention_policy parameter to the desired length of time for the recovery window in one of the following ways:

- Set to max_number DAYS to define the start date/time recovery window boundary as the number of days specified by max_number going back in time from the current date/time.
- Set to max_number WEEKS to define the start date/time recovery window boundary as the number of weeks specified by max_number going back in time from the current date/time.
- Set to max_number MONTHS to define the start date/time recovery window boundary as the number of months specified by max_number going back in time from the current date/time.

Additional Restrictions:

- The keywords DAYS, WEEKS, and MONTHS must always be specified in plural form (for example, 1 DAYS, 1 WEEKS, or 1 MONTHS).
- BART will accept a maximum integer value of 2,147,483,647 for max_number, however, a realistic, practical value based on your system environment must always be used.

A backup is considered active if the date/time of the backup is equal to or greater than the start of the recovery window date/time.

You can view the actual, calculated recovery window by:

- Invoking the MANAGE subcommand in debug mode, along with the -n option.
- Using the SHOW-SERVERS subcommand.

5.3.2.4 Managing the Backups Based on the Retention Policy

The MANAGE subcommand is used to evaluate and categorize backups according to the retention policy set in the BART

configuration file. When a backup is first created with the BACKUP subcommand, it is active. You can use the MANAGE subcommand to change the status of an active backup to obsolete. Obsolete backups can then be deleted.

This section covers following aspects of backup management:

- The rules for deleting backups depending upon the backup status and the subcommand used.
- The process to retain a backup indefinitely by marking it as keep. This section also provides information about resetting backups status (that are marked as obsolete and keep) back to active status.
- The general process for evaluating, marking, and then deleting obsolete backups.

Deletions Permitted Under a Retention Policy

This section describes how and under what conditions backups may be deleted under a retention policy.

You must use the MANAGE subcommand to delete obsolete backups. Use the DELETE subcommand only for special administrative purposes.

The deletion behavior of the MANAGE subcommand and the DELETE subcommand are based on different aspects of the retention policy.

- The MANAGE subcommand deletion relies solely upon how a backup status is currently marked (that is, internally recorded by BART). The current setting of the retention_policy parameter in the BART configuration file is ignored.
- The DELETE subcommand relies solely upon the current setting of the retention_policy parameter in the BART configuration file. The current active, obsolete, or keep status of a backup is ignored.

The specific deletion rules for the MANAGE and DELETE subcommands are as follows:

- MANAGE subcommand: The MANAGE subcommand with the -d option can only delete backups marked as obsolete. This deletion occurs regardless of the current retention_policy setting in the BART configuration file. The deletion of backups relies on the last occasion when the backups have been marked.
- DELETE subcommand:
 - Under a redundancy retention policy currently set with the retention_policy parameter in the BART configuration file, any backup regardless of its marked status, can be deleted with the DELETE subcommand when the backup identifier or name is specified with the -i option and if the current total number of backups for the specified database server is greater than the maximum number of redundancy backups currently specified with the retention_policy parameter.

If the total number of backups is less than or equal to the specified, maximum number of redundancy backups, then no additional backups can be deleted using **DELETE** with the -i backup option.

- o Under a recovery window retention policy currently set with the retention_policy parameter in the BART configuration file, any backup regardless of its marked status, can be deleted with the DELETE subcommand when the backup identifier or name is specified with the -i option, and if the backup date/time is not within the recovery window currently specified with the retention_policy parameter. If the backup date/time is within the recovery window, then it cannot be deleted using DELETE with the -i backup option.
- Invoking the DELETE subcommand with the -i all option results in the deletion of all backups regardless of the retention policy and regardless of whether the status is marked as active, obsolete, or keep.

The following table summarizes the deletion rules of backups according to their marked status. An entry of Yes indicates the backup may be deleted under the specified circumstances. An entry of No indicates that the backup may

not be deleted.

Operation	Redundancy Retention Policy			Recovery Window Retention Policy		
	Active	Obsolete	Кеер	Active	Obsolete	Keep
MANAGE - d	No	Yes	No	No	Yes	No
DELETE -i *backup*	Yes (see Note 1)	Yes (see Note 1)	Yes (see Note 1_)	Yes (see Note 2	Yes (see Note 2	Yes (see Note 2)
DELETE -i all	Yes	Yes	Yes	Yes	Yes	Yes

!!! Note Redundancy Retention Policy (Note 1): Deletion occurs only if the total number of backups for the specified database server is greater than the specified, maximum number of redundancy backups currently set with the redundancy_policy parameter in the BART configuration file.

!!! Note Recovery Window Retention Policy (Note 2): Deletion occurs only if the backup is not within the recovery window currently set with the redundancy_policy parameter in the BART configuration file.

Marking Backups for Indefinite Keep Status

There may be certain backups that you wish to keep for an indefinite period of time and do not wish to delete based upon the retention policy applied to the database server. Such backups can be marked as keep to exclude them from being marked as obsolete. Use the MANAGE subcommand with the -c keep option to retain such backups indefinitely.

Evaluating, Marking, and Deleting Obsolete Backups

When the MANAGE subcommand is invoked, BART evaluates active backups:

- If you include the -s option when invoking the MANAGE subcommand, BART evaluates backups for the database server.
- If you include the -s all option when invoking the MANAGE subcommand, BART evaluates backups for all database servers.
- If the -s option is omitted, the command evaluates the current number of backups for the database server based on the redundancy retention policy or the current date/time for a recovery window retention policy.

!!! Note The status of backups currently marked as obsolete or keep is not changed. To re-evaluate such backups and then classify them, their status must first be reset to active with the MANAGE -c nokeep option. See Marking the Backup Status for more information.

See the *EDB Backup and Recovery Reference Guide* available at the *EDB website* to review examples of how to evaluate, mark, and delete backups using a redundancy retention policy and recovery window retention policy, as well as examples of MANAGE subcommand.

5.3.2.5 Managing Incremental Backups

The following section summarizes how retention policy management affects incremental backups.

- The retention policy rules are applied to full backups.
 - A redundancy retention policy uses the number of full backups to determine if a backup is obsolete. Incremental
 backups are excluded from the comparison count against the
 retention_policy setting for the maximum
 number of backups.
 - A recovery window retention policy uses the backup date/time of any full backups to determine if a backup is obsolete. The backup date/time of any successive incremental backups in the chain are ignored when comparing with the recovery window.
- The retention status of all incremental backups in a chain is set to the same status applied to the full backup of the
- The actions applied by the MANAGE and DELETE subcommands on a full backup are applied to all incremental backups in the chain in the same manner.
- Thus, a backup chain (that is, the full backup and all its successive incremental backups) are treated by retention policy management as if they are all one, single backup.
 - The status setting applied to the full backup is also applied to all incremental backups in its chain.
 - If a full backup is marked as obsolete and then deleted according to the retention policy, all incremental backups in the chain are also marked obsolete and then deleted as well.

The following are some specific points regarding the MANAGE and DELETE subcommands on incremental backups.

- MANAGE subcommand:
 - When the MANAGE subcommand is invoked, the status applied to the full backup is also applied to all successive incremental backups in the chain.
 - The MANAGE subcommand with the -c { keep | nokeep} option cannot specify the backup identifier or backup name of an incremental backup with -i backup option. The -i backup option can only specify the backup identifier or backup name of a full backup.
 - You can also use the -i all option to take a backup of all backups. When the MANAGE subcommand with the -c
 { keep | nokeep } option is applied to a full backup, the same status change is made to all incremental backups in the chain.
- DELETE subcommand:
 - The DELETE subcommand with the -s server -i backup option specifies the backup identifier or backup name of an incremental backup in which case that incremental backup along with all its successive incremental backups are deleted, thus shortening that backup chain.

Using a Redundancy Retention Policy with Incremental Backups

When a redundancy retention policy is used and the MANAGE subcommand is invoked, the status of the oldest active full backup is changed to obsolete if the number of full backups exceeds the maximum number specified by the retention_policy parameter in the BART configuration file.

!!! Note When a full backup is changed from active to obsolete, all successive incremental backups in the chain of the full backup are also changed from active to obsolete.

When determining the number of backups that exceeds the number specified by the retention_policy parameter, only full backups are counted for the comparison. Incremental backups are not included in the count for the comparison against the retention_policy parameter setting.

See the *EDB Backup and Recovery Reference Guide* available at the *EDB website* for examples demonstrating use of the MANAGE and DELETE subcommands when a redundancy retention policy is in effect.

Using a Recovery Window Retention Policy with Incremental Backups

If the MANAGE command is invoked when BART is configured to use a recovery window retention policy, the status of active full backups are changed to obsolete if the date/time of the full backup is outside of the recovery window.

!!! Note If a full backup is changed from active to obsolete, all successive incremental backups in the chain of the full backup are also changed from active to obsolete.

The status of an incremental backup is changed to obsolete regardless of whether or not the date/time of when the incremental backup was taken still lies within the recovery window.

See the *EDB Backup and Recovery Reference Guide* available at the *EDB website* for examples demonstrating use of the MANAGE and DELETE subcommands when a recovery window retention policy is in effect.

5.3.3 Basic BART Subcommand Usage

This section briefly describes the BART subcommands and options. You can invoke the bart program (located in the BART_HOME /bin directory) with the desired options and subcommands to manage your BART installation.

To view examples of BART subcommands, see the *EDB Backup and Recovery Reference Guide* available at the *EDB website*.

Syntax for invoking BART:

```
bart [ general_option ]... [ subcommand ] [subcommand_option ]...
```

- When invoking a subcommand, the subcommand name is not case-sensitive (that is, the subcommand can be specified in uppercase, lowercase, or mixed case).
- Each subcommand has a number of its own applicable options that are specified following the subcommand. All options are available in both single-character and multi-character forms.
- Keywords are case-sensitive; options are generally specified in lowercase unless specified otherwise in this section.
- When invoking BART, the current user must be the BART user account (operating system user account used to run the BART command line program). For example, enterprised or postgres can be selected as the BART user account when the managed database servers are Advanced Server or PostgreSQL respectively.
- The chosen operating system user account must own the BART backup catalog directory, be able to run the bart program and the bart scanner program, and have a passwordless SSH/SCP connection established between database servers managed by BART.

You can specify one or more of the following general options:

Options	Description
-hhelp	Displays general syntax and information on BART usage. All subcommands support a help option $(-h,help)$. If the help option is specified, information is displayed regarding that particular subcommand. The subcommand, itself, is not executed.
-vversion	Displays the BART version information.

Options	Description
-ddebug	Displays debugging output while executing BART subcommands.
<pre>-c <config_file_path>config-path <config_file_path></config_file_path></config_file_path></pre>	Specifies config_file_path as the full directory path to a BART configuration file. Use this option if you do not want to use the default BART configuration file <bart_home>/etc/bart.cfg.</bart_home>

Troubleshooting: Setting Path Environment Variable

If execution of BART subcommands fails with the following error message, then you need to set the LD_LIBRARY_PATH to include the libpq library directory:

```
./bart: symbol lookup error: ./bart: undefined symbol: PQping
```

Workaround: Set the LD_LIBRARY_PATH environment variable for the BART user account to include the directory containing the libpq library. This directory is POSTGRES_INSTALL_HOME/lib.

It is suggested that the PATH and the LD_LIBRARY_PATH environment variable settings be placed in the BART user account's profile. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for details.

In the following sections, the help option is omitted from the syntax diagrams for the purpose of providing readability for the subcommand options.

5.3.3.1 CHECK-CONFIG

The CHECK-CONFIG subcommand checks the parameter settings in the BART configuration file as well as the database server configuration for which the -s option is specified.

Syntax:

```
bart CHECK-CONFIG [ -s server_name ]
```

The following table describes the option.

Options Description -s <server_name> or server_name is the name of the database server to be checked for proper configuration. If the option is omitted, the settings of the global section of the BART configuration file are checked.

- When the -s option is omitted, the global section [BART] parameters including bart_host, backup_path, and pg_basebackup_path are checked.
- When the -s option is specified, the server section parameters are checked. In addition, certain database server postgresql.conf parameters are also checked, which include the following:
 - The cluster_owner parameter must be set to the user account owning the database cluster directory.
 - A passwordless SSH/SCP connection must be set between the BART user and the user account specified by the cluster_owner parameter.
 - A database superuser must be specified by the BART user parameter.

- The pg_hba.conf file must contain a replication entry for the database superuser specified by the BART user parameter.
- The archive_mode parameter in the postgresql.conf file must be enabled.
- The archive_command parameter in the postgresql.auto.conf or the postgresql.conf file must he set
- The allow_incremental_backups parameter in the BART configuration file must be enabled for database servers for which incremental backups are to be taken.
- Archiving of WAL files to the archive_path must be in process.
- The WAL scanner program must be running.

The CHECK-CONFIG subcommand displays an error message if the required configuration is not properly set.

5.3.3.2 INIT

The INIT subcommand is used to create the BART backup catalog directory, rebuild the BART backupinfo file, and set the archive_command in the PostgreSQL server based on the archive_command setting in the bart.cfg file.

!!! Note If the archive_mode configuration parameter is set to off, then the -o option must be used to set the Postgres archive_command using the BART archive_command parameter in the BART configuration file even if the archive_command is not currently set in postgresql.conf nor in postgresql.auto.conf file.

Syntax:

```
bart INIT [ -s { <server_name> | all } ] [ -o ]
  [ -r [ -i { <backup_id> | <backup_name> | all } ] ]
  [--no-configure]
```

All subcommand options are generally specified in lowercase. The following table describes the command options:

Options	Description
-s { <server_name> all }server {<server_name> all }</server_name></server_name>	server_name is the name of the database server to which the INIT actions are to be applied. If all is specified or if the option is omitted, the actions are applied to all servers.
-o override	Overrides the existing, active Postgres archive_command configuration parameter setting in the postgresql.conf file or the postgresql.auto.conf file using the BART archive_command parameter in the BART configuration file. The INIT generated archive command string is written to the postgresql.auto.conf file.
-r rebuild	Rebuilds the backupinfo file (a text file named backupinfo) located in each backup subdirectory. This option is only intended for recovering from a situation where the backupinfo file has become corrupt. If the backup was initially created with a user-defined backup name, and then the INIT -r option is invoked to rebuild that backupinfo file, the user-defined backup name is no longer available. Thus, future references to the backup must use the backup identifier.

Description Options -i { <backup_id> | <backup_name> <backup_id> is an integer, backup identifier and <backup_name> is the user-defined | all } alphanumeric name for the backup. If all is specified or if the option is omitted, the backupinfo files of all backups for the database servers specified by the -s option are --backupid { recreated. The -i option can only be used with the -r option. <backup_id> <backup_name> | all } --no-Prevents the archive_command from being set in the PostgreSQL server. configure

Archive Command Setting

After the archive_command is set, you need to either restart the PostgreSQL server or reload the configuration file in the PostgreSQL server based on the following conditions.

- If the archive_mode is set to off and archive_command is not set in the PostgreSQL server, the archive_command is set based on the archive_command setting in the bart.cfg and also sets the archive_mode to on. In this case, you need to restart the PostgreSQL server using pg_ctl restart
- If the archive_mode is set to on and archive_command is not set in the PostgreSQL server, the archive_command is set based on the archive_command setting in the bart.cfg. In this case, you need to reload the configuration in the PostgreSQL server using pg_reload_conf() or pg_ctl reload.
- If the archive_mode is set to off and archive_command is already set in the PostgreSQL server, the archive_mode is set to on. In this case, you need to restart the PostgreSQL server using pg_ctl restart
- If the archive_mode is set to on and archive_command is already set in the PostgreSQL server, then the archive_command is not set unless -o option is specified.

5.3.3.3 BACKUP

The BACKUP subcommand is used to create a full backup or an incremental backup.

Syntax for full backup:

```
bart BACKUP -s { <server_name> | all } [ -F { p | t } ]
  [ -z ] [ -c <compression_level> ]
  [ --backup-name <backup_name> ]
  [ --thread-count <number_of_threads> ]
  [ { --with-pg_basebackup | --no-pg_basebackup } ]
```

!!! Note While taking a backup, if a file (for example, database server log file) exceeding 1 GB size is stored in the \$PGDATA directory, the backup will fail. To avoid such backup failure, you need to store large files (exceeding 1 GB) outside the \$PGDATA directory.

Syntax for incremental Backup:

```
bart BACKUP -s { <server_name> | all } [ -F p]
[ --parent { <backup_id> | <backup_name> } ]
```

```
[ --backup-name <backup_name> ]
[ --thread-count <number_of_threads> ]
[ --check ]
[--checksum-algorithm ]
```

!!! Note To take an incremental backup, you must take a full backup first followed by incremental backup.

Please Note:

- While a BACKUP subcommand is in progress, no other subcommands must be invoked. Any subcommands invoked while a backup is in progress will skip and ignore the backups.
- For full backup, the target default format is a tar file, whereas for incremental backup, only plain format must be specified.
- The backup is saved in the <backup_path>/<server_name>/<backup_id> directory, where <backup_path> is the value assigned to the <backup_path> parameter in the BART configuration file, <server_name> is the lowercase name of the database server as listed in the configuration file, and <backup_id> is a backup identifier assigned by BART to the particular backup.
- MD5 checksums of the full backup and any user-defined tablespaces are saved as well for tar backups.
- Before performing the backup, BART checks to ensure if there is enough disk space to completely store the backup in the BART backup catalog.
- In the postgresql.conf file, ensure the wal_keep_segments configuration parameter is set to a sufficiently large value. A low setting of the wal_keep_segments configuration parameter may result in the deletion of some WAL files before the BART BACKUP subcommand saves them to the archive_path. For information about the wal_keep_segments parameter, see the PostgreSQL Core Documentation.
- In the BART configuration file, setting xlog_method=stream will instruct the server to stream the transaction log in parallel with creation of the backup for a specific database server; otherwise the transaction log files are collected upon completion of the backup. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for details about database server setting.
 - !!! Note If the transaction log streaming method is used, the -Fp option for a plain text backup format must be specified with the BACKUP subcommand.
- When you use BART to take a backup of PostgreSQL server, multiple backups can be taken simultaneously and if a backup is interrupted, the backup mode is terminated automatically without the need to run pg_stop_backup() command manually to terminate the backup.

Options

Along with the BACKUP subcommand, specify the following option:

Options

Description

```
-s {
server_name
| all }
--server {
server_name
| all }
```

server_name is the database server name to be backed up as specified in the BART configuration file. If all is specified, all servers are backed up. This option is mandatory. If all is specified, and a connection to a database server listed in the BART configuration file cannot be opened, the backup for that database server is skipped, but the backup operation continues for the other database servers.

Specify the following options as required. If you do not specify any of the following options, the backup is created using

default settings.

Options	Description
-F { p t }format { p t }	Specify this option to provide the backup file format. Use p for plain text or t for tar. If the option is omitted, the default is tar format. For taking incremental backups, the option -Fp must be specified.
-z gzip	This is applicable only for full backup. Specify this option to use gzip compression on the tar file output using the default compression level. This option is applicable only for the tar format.
<pre>-c <compression_level>compress-level <compression_level></compression_level></compression_level></pre>	This is applicable only for full backup. Specify this option to use the gzip compression level on the tar file output. <pre>compression_level</pre> is a digit from 1 through 9, with 9 being the best compression. This option is applicable only for the tar format.
<pre>parent { backup_id backup_name }</pre>	Specify this option to take an incremental backup. <backup_id> is the backup identifier of a parent backup. <backup_name> is the user-defined alphanumeric name of a parent backup. The parent is a backup taken prior to the incremental backup. The parent backup can be either a full backup or an incremental backup. The option _Fp must be specified since an incremental backup can only be taken in plain text format. An incremental backup cannot be taken on a standby database server. See Block-Level Incremental Backup for additional information on incremental backups.</backup_name></backup_id>
backup-name <backup_name></backup_name>	Specify this option to assign a user-defined, alphanumeric friendly name to the backup. The maximum permitted length of backup name is 49 characters. The backup name may include the following variables to be substituted by the timestamp values when the backup is taken: 1) <code>%year - 4-digit year</code> , 2) <code>%month - 2-digit month</code> , 3) <code>%day - 2-digit day</code> , 4) <code>%hour 2-digit hour</code> , 5) <code>%minute - 2-digit minute</code> , and 6) <code>%second - 2-digit second</code> . To include the percent sign (%) as a character in the backup name, specify <code>%%</code> in the alphanumeric string. If the backup name contains space characters (i.e. more than one word) or when referenced with the option <code>-i</code> by other subcommands (such as <code>restore</code>), enclose the string in single quotes or double quotes. See <code>backup name examples</code> . If the <code>backup-name</code> option is not specified, and the <code>backup_name</code> parameter is not set for this database server in the BART configuration file, then the backup can only be referenced in other BART subcommands by the BART assigned backup identifier.
thread-count <number_of_threads></number_of_threads>	Use this option to use the number of worker threads to run in parallel to copy blocks for a backup. If the optionthread-count is omitted, then the thread_count parameter in the BART configuration file applicable to this database server is used. If the optionthread-count is not enabled for this database server, then the thread_count setting in the global section of the BART configuration file is used. If the optionthread-count is not set in the global section as well, the default number of threads is 1. If parallel backup is run with N number of worker threads, then it will initiate N+ 1 concurrent connections with the server. Thread count will not be effective if backup is taken on a standby server. For more information about thethread-count parameter, see the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website

Options	Description
with- pg_basebackup	This is applicable only for full backup. Specify this option to use pg_basebackup to take a full backup. The number of thread counts in effect is ignored as given by the thread_count parameter in the BART configuration file. When taking a full backup, if the thread count in effect is greater than 1, then the pg_basebackup utility is not used to take the full backup (parallel worker threads are used) unless the optionwith-pg_basebackup is specified with the BACKUP subcommand.
	This is applicable only for full backup. Specify this option if you do not want pg_basebackup to be used to take a full backup.
no-pg_basebackup	When taking a full backup, if the thread count in effect is only 1, then the
	pg_basebackup utility is used to take the full backup unless the optionno-pg_basebackup is specified with the BACKUP subcommand.
check	This is applicable only for incremental backup. Specify this option to verify if the required MBM files are present in the archive_path parameter in the bart.cfg file before taking an incremental backup. The option

Backup Name Examples

The following examples demonstrate using the --backup-name clause:

```
./bart backup -s ppas12 -Ft --backup-name "YEAR = %year
MONTH = %month DAY = %day"
./bart backup -s ppas12 -Ft --backup-name "YEAR = %year
MONTH = %month DAY = %day %%"
./bart show-backups -s ppas12 -i "test backup"
```

Error messages

The following table lists the error messages that may be encountered when using the BACKUP subcommand:

error message	Cause
edb@localhost bin]\$./bart backup -s mktg -Ft	
WARNING: xlog_method is empty, defaulting to global policy	
ERROR: backup failed for server 'mktg'	Insufficient free disk space.
free disk space is not enough to backup the server 'mktg'	
space available 13.35 GB, approximately required 14.65 GB	

Cause error message ERROR: backup failed for server 'mktg' command failed with exit code 1 The wal_keep_segments configuration parameter is not pg_basebackup: could not get transaction set to a sufficiently large value in the postgresql.conf log end position from server: ERROR: file. requested WAL segment 00000001000000D50000006B has already been removed ERROR: backup failed for server 'mktg' connection to the server failed: could not A connection to a database server listed in the BART connect to server: Connection refused configuration file fails. As a result the backup for that database server is skipped, but the backup operation Is the server running on host continues for other database servers "172.16.114.132" and accepting TCP/IP connections on port 5444?

5.3.3.4 SHOW-SERVERS

The SHOW-SERVERS subcommand displays the information for the managed database servers listed in the BART configuration file.

Syntax:

```
bart SHOW-SERVERS [ -s { <server_name> | all } ]
```

The following table describes the command options.

```
Options

-s {

<server_name>
| all }

--server {

<server_name>
| all }

| all }

| all }

| all }
```

5.3.3.5 SHOW-BACKUPS

The SHOW-BACKUPS subcommand displays the backup information for the managed database servers.

Syntax:

```
bart SHOW-BACKUPS [ -s { <server_name> | all } ]
  [ -i { <backup_id> | <backup_name> | all } ]
  [ -t ]
```

The following table describes the command options:

Options Description <server_name> is the name of the database server whose backup information is to be displayed. -s { <server_name> If all is specified or if the option is omitted, the backup information for all database all } servers is displayed with the exception as described by the following note: --server { If SHOW-BACKUPS is invoked while the BART BACKUP subcommand is in progress, <server_name> | all } backups affected by the backup process are shown in progress status in the displayed backup information. -i { <backup_id> | <backup_id> is a backup identifier and <backup_name> is the user-defined <backup_name> | all } alphanumeric name for the backup. --backupid { If all is specified or if the option is omitted, all backup information for the relevant <backup_id> | database server is displayed. <backup_name> | all } -t Displays more backup information in a list format. If the option is omitted, the default is a tabular format. --toggle

5.3.3.6 VERIFY-CHKSUM

The VERIFY-CHKSUM subcommand verifies the MD5 checksums of the full backups and any user-defined tablespaces for the specified database server or for all database servers. The checksum is verified by comparing the current checksum of the backup against the checksum when the backup was taken.

!!! Note The VERIFY-CHKSUM subcommand is only used for tar format backups. It is not applicable to plain format backups.

Syntax:

```
bart VERIFY-CHKSUM
  [ -s { <server_name> | all } ]
  [ -i { <backup_id> | <backup_name> | all } ]
```

The following table describes the command options:

Options

Description

```
-s { <server_name> |
all }
--server {
<server_name> | all
}
```

<server_name> is the name of the database server whose tar backup checksums
are to be verified. If all is specified or if the -s option is omitted, the checksums are
verified for all database servers.

Options

Description

<backup_id> is the backup identifier of a tar format full backup whose checksum is
to be verified along with any user-defined tablespaces.

<backup_name> is the user-defined alphanumeric name for the full backup.
If all is specified or if the -i option is omitted, the checksums of all tar backups for
the relevant database server are verified.

5.3.3.7 MANAGE

The MANAGE subcommand can be invoked to:

- Evaluate backups, mark their status, and delete obsolete backups based on the retention_policy parameter in the BART configuration file (See Managing Backups Using a Retention Policy for information about retention policy management).
- Compress the archived WAL files based on the wal_compression parameter in the BART configuration file (See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for information about setting this parameter).

Syntax:

```
bart MANAGE [ -s { <server_name> | all} ]
  [ -l ] [ -d ]
  [ -c { keep | nokeep }
    -i { <backup_id> | <backup_name> | all } ]
  [ -n ]
```

The following summarizes the actions performed when the MANAGE subcommand is invoked:

- When the MANAGE subcommand is invoked with no options or with only the -s <server_name> or -s all option, the following actions are performed:
 - For the server specified by the -s option, or for all servers (if -s all is specified or the -s option is omitted), active backups are marked as obsolete in accordance with the retention policy.
 - All backups that were marked obsolete or keep prior to invoking the MANAGE subcommand remain marked with the same prior status.
 - If WAL compression is enabled for the database server, then any uncompressed, archived WAL files in the BART backup catalog of the database server are compressed with gzip.
- When the MANAGE subcommand is invoked with any other option besides the -s option, the following actions are performed:
 - ∘ For the server specified by the ¬s option, or for all servers, the action performed is determined by the other specified options (that is, ¬l to list obsolete backups, ¬d to delete obsolete backups, ¬c to keep or to return backups to active status, or ¬n to perform a dry run of any action).
 - No marking of active backups to obsolete status is performed regardless of the retention policy.
 - o All backups that were marked obsolete or keep prior to invoking the MANAGE subcommand remain marked with the same prior status unless the −c option (without the −n option) is specified to change the backup status of the particular backup or all backups referenced with the −i option.

• No compression is applied to any uncompressed, archived WAL file in the BART backup catalog regardless of whether or not WAL compression is enabled.

The following are additional considerations when using WAL compression:

- Compression of archived WAL files is not permitted for database servers on which incremental backups are to be taken.
- The gzip compression program must be installed on the BART host and be accessible in the PATH of the BART user account.
- When the RESTORE subcommand is invoked, if the -c option is specified or if the copy_wals_during_restore BART configuration parameter is enabled for the database server, then the following actions occur:
 - If compressed, archived WAL files are stored in the BART backup catalog and the location to which the WAL files are to be restored is on a remote host relative to the BART host:
 - the archived WAL files are transmitted across the network to the remote host in compressed format only if the gzip compression program is accessible in the PATH of the remote user account that is used to log into the remote host when performing the RESTORE operation.
 - This remote user is specified with either the remote_host parameter in the BART configuration file or the RESTORE -r option (see RESTORE).
 - Transmission of compressed WAL files results in less network traffic. After the compressed WAL files are transmitted across the network, the RESTORE subcommand uncompresses the files for the point-in-time recovery operation.
 - If the gzip program is not accessible on the remote host in the manner described in the previous bullet point, then the compressed, archived WAL files are first uncompressed while on the BART host, then transmitted across the network to the remote host in uncompressed format.
- When the RESTORE subcommand is invoked without the -c option and the copy_wals_during_restore BART configuration parameter is disabled for the database server, then any compressed, archived WAL files needed for the RESTORE operation are uncompressed in the BART backup catalog. The uncompressed WAL files can then be saved to the remote host by the restore_command in the postgresql.auto.conf file when the database server archive recovery begins.

The following table describes the command options:

Options	Description
s {	
<pre><server_name> all }server {</server_name></pre>	<pre><server_name> is the name of the database server to which the actions are to be applied. If all is specified or if the -s option is omitted, the actions are applied to all database servers.</server_name></pre>
<pre><server_name> all }</server_name></pre>	
-l list- obsolete	Lists the backups marked as obsolete.
-d delete- obsolete	Delete the backups marked as obsolete. This action physically deletes the backup along with its archived WAL files and any MBM files for incremental backups.

Options Description -c { keep | Specify keep to change the status of a backup to keep to retain it indefinitely. nokeep } Specify nokeep to change the status of any backup back to active status. The backup can then be re-evaluated and possibly be marked to obsolete according to the retention policy by --changesubsequent usage of the MANAGE subcommand. status { keep The -i option must be included when using the -c option. | nokeep } -i { <backup_id> | <backup_name> <backup_id> is a backup identifier and <backup_name> is the user-defined alphanumeric | all } name for the backup. If all is specified, then actions are applied to all backups. --backupid { The -c option must be included when using the -i option. <backup_id> | <backup_name> | all } Performs the test run and displays the results prior to actually implementing the actions as if the operation was performed, however, no changes are actually made. If -n is specified with the -d option, it displays which backups would be deleted, but does not actually delete the backups. If -n is specified with the -c option, it displays the keep or nokeep action, but does not actually change the backup from its current status. -dry-run If -n is specified alone with no other options, or with only the -s option, it displays which active backups would be marked as obsolete, but does not actually change the backup status. In addition, no compression is performed on uncompressed, archived WAL files even if WAL compression is enabled for the database server.

5.3.3.8 RESTORE

The RESTORE subcommand restores a backup and its archived WAL files for the designated database server to the specified directory location. If the appropriate RESTORE options are specified, all recovery settings will be saved in the postgresql.auto.conf file.

Syntax:

For information about using a continuous archive backup for recovery, see the PostgreSQL Core Documentation. This reference material provides detailed information about the underlying point-in-time recovery process and the meaning and usage of the restore options that are generated into the postgresql.auto.conf file by BART.

Please note:

- For special requirements when restoring an incremental backup to a remote database server, see Restoring an Incremental Backup on a Remote Host.
- Check to ensure that the host where the backup is to be restored contains enough disk space for the backup and its archived WAL files. The RESTORE subcommand may result in an error while copying files if there is not enough disk space available.
- See Performing a Restore Operation to view steps on how to perform a restore operation and see Point-In-Time Recovery Operation to view steps on how to perform a point-in-time recovery operation.
- If the backup is restored to a different database cluster directory than where the original database cluster resided, certain operations dependent upon the database cluster location may fail. This happens if their supporting service scripts are not updated to reflect the new directory location of restored backup. For information about the usage and modification of service scripts, see the EDB Advanced Server Installation Guide available at the EDB website.

The following table describes the command options:

Options	Description
<pre>-s <server_name>server <server_name></server_name></server_name></pre>	<pre><server_name> is the name of the database server to be restored.</server_name></pre>
<pre>-p <restore_path>restore-path <restore_path></restore_path></restore_path></pre>	<pre><restore_path> is the directory path where the backup of the database server is to be restored. The directory must be empty and have the proper ownership and privileges assigned to it.</restore_path></pre>
<pre>-i {<backup_id> <backup_name>} backupid {<backup_id> <backup_name>}</backup_name></backup_id></backup_name></backup_id></pre>	<pre></pre>
-r orremote-host <remote_user @remote_host_address></remote_user 	<pre><remote_user> is the user account on the remote database server host that accepts a passwordless SSH/SCP login connection and is the owner of the directory where the backup is to be restored and <remote_host_address> is the IP address of the remote host to which the backup is to be restored. This option must be specified if the <remote_host> parameter for this database server is not set in the BART configuration file. If the BART user account is not the same as the operating system account owning the <restore_path> directory given with the -p option, use the <remote_host> BART configuration parameter or the RESTORE subcommand -r option to specify the <restore_path> directory owner even when restoring to a directory on the same host as the BART host. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for information about the <remote_host> parameter.</remote_host></restore_path></remote_host></restore_path></remote_host></remote_host_address></remote_user></pre>
<pre>-w <number_of_workers>workers <number_of_workers></number_of_workers></number_of_workers></pre>	<number_of_workers> is the specification of the number of worker processes to run in parallel to stream the modified blocks of an incremental backup to the restore location. For example, if 4 worker processes are specified, 4 receiver processes on the restore host and 4 streamer processes on the BART host are used. The output of each streamer process is connected to the input of a receiver process. When the receiver gets to the point where it needs a modified block file, it obtains those modified blocks from its input. With this method, the modified block files are never written to the restore host disk. If the _w option is omitted, the default is 1 worker process.</number_of_workers>
<pre>-t <timeline_id>target-tli <timeline_id></timeline_id></timeline_id></pre>	<timeline_id> is the integer identifier of the timeline to be used for replaying the archived WAL files for point-in-time recovery.</timeline_id>

Options	Description
<pre>-x <target_xid>target-xid <target_xid></target_xid></target_xid></pre>	<target_xid> is the integer identifier of the transaction ID that determines the transaction up to and including, which point-in-time recovery encompasses. Include either the -x <target_xid> or thetarget-xid <target_xid> option if point-in-time recovery is desired.</target_xid></target_xid></target_xid>
<pre>-g <target_timestamp>target-timestamp <target_timestamp></target_timestamp></target_timestamp></pre>	<pre><target_timestamp> is the timestamp that determines the point in time up to and including, which point-in-time recovery encompasses. Include either the target-timestamp <target_timestamp> or the -g <target_timestamp> option if point-in-time recovery is desired.</target_timestamp></target_timestamp></target_timestamp></pre>
-c copy-wals	Specify this option to copy archived WAL files from the BART backup catalog to <restore_path>/archived_wals directory. If recovery settings are saved in the postgresql.auto.conf file for point-in-time recovery, the restore_command retrieves the WAL files from <restore_path>/archived_wals for the database server archive recovery. If the -c option is omitted and the copy_wals_during_restore parameter in the BART configuration file is not enabled in a manner applicable to this database server, the restore_command in the postgresql.auto.conf file is generated by default to retrieve the archived WAL files directly from the BART backup catalog. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for information about the copy_wals_during_restore parameter.</restore_path></restore_path>
	While restoring a backup, specify this option to skip verifying the MD5 or SHA256 checksum files.
disable-checksum	If you set thechecksum-algorithm=NONE option with the BART scanner or while taking a backup, you must specify thedisable checksum option while restoring an incremental backup.

5.3.3.9 **DELETE**

The DELETE subcommand removes the subdirectory and data files from the BART backup catalog for the specified backups along with its archived WAL files.

Syntax:

```
bart DELETE -s <server_name>
  -i { all |
      [']{ <backup_id> | <backup_name> },... }[']
    }
[ -n ]
```

!!! Note While invoking the DELETE subcommand, you must specify a specific database server.

For database servers under a retention policy, there are conditions where certain backups may not be deleted. See Deletions Permitted Under a Retention Policy for information about permitted backup deletions.

The following table describes the command options:

Options Description

Options Description -s <server_name> <server_name> is the name of the database server whose backups are to be deleted. --server <server_name> -i { all | [']{ <backup_id> | <backup_id> is the backup identifier of the backup to be deleted and <backup_name> },... } <backup name> is the user-defined alphanumeric name for the backup. Multiple backup identifiers and backup names may be specified in a comma-['] } separated list. The list must be enclosed within single quotes if there is any white --backupid { all | ['] space appearing before or after each comma. If all is specified, all of the backups and their archived WAL files for the specified { <backup_id> | database server are deleted. <backup_name> },... } ['] } Displays the results as if the deletions were done, however, no physical removal of the files are actually performed. In other words, a test run is performed so that you can see the potential results prior to actually initiating the action. After the deletion, the BART backup catalog for the database server no longer --dry-run contains the corresponding directory for the deleted backup ID. The archived_wals subdirectory no longer contains the WAL files of the backup.

5.3.4 Running the BART WAL Scanner

Use the BART WAL scanner to invoke the bart-scanner program located in the BART_HOME/bin directory. When invoking the WAL scanner, the current user must be the BART user account.

Syntax:

```
bart-scanner
  [ -d ]
  [ -c <config_file_path> ]
  { -h |
    -v |
    -daemon |
    -p <mbm_file_path> |
    <wal_file_path> |
    RELOAD |
    STOP
    --checksum-algorithm }
```

!!! Note For clarity, the syntax diagram shows only the single-character option form (for example, -d), but the multi-character option form (for example, --debug) is supported as well.

The WAL scanner processes each WAL file to find and record modified blocks in a corresponding modified block map (MBM) file. The default approach is that the WAL scanner gets notified whenever a new WAL file is added to the archived_wals directory specified in the archive_path parameter of the configuration file. It then scans the WAL file and produces the MBM file.

The default approach does not work in some cases; for example when the WAL files are shipped to the archive_path using the Network File System (NFS) and also in case of some specific platforms. This results in the WAL files being

copied to the <a href="archived_wals" archived_wals" archived_wals" archived_wals directory, but the WAL scanner does not scan them (as WAL scanner is not aware of WAL file) and produce the MBM files. This results in the failure of an incremental backup. This can be avoided by using the timer-based WAL scanning approach, which is done by using the scan_interval parameter in the BART configuration file. The value for scan_interval is the number of seconds after which the WAL scanner will initiate force scanning of the new WAL files. See the EDB Backup and Recovery Installation and Upgrade Guide available at the EDB website for more information about scan_interval parameter.

!!! Note After upgrading to BART 2.6, users who have set this parameter to a non-default value may see increased CPU consumption on the part of bart-scanner. If this is an issue, consider increasing the configured value of scan_interval parameter, or removing the setting if it is not required.

When the bart-scanner program is invoked, it forks a separate process for each database server enabled with the allow_incremental_backups parameter.

The WAL scanner processes can run in either the foreground or background depending upon usage of the --daemon option. Use the --daemon option to run the WAL scanner process in the background so that all output messages can be viewed in the BART log file. If the --daemon option is omitted, the WAL scanner process runs in the foreground and all output messages can be viewed from the terminal running the program as well as in the BART log file.

See the *EDB Backup and Recovery Installation and Upgrade Guide* available at the *EDB website* for additional information about WAL scanning, allow_incremental_backups, and logfile parameters.

!!! Note The BART user account's LD_LIBRARY_PATH environment variable may need to be set to include the directory containing the libpq library if invocation of the WAL scanner program fails. See Basic BART Subcommand Usage for information about setting the LD_LIBRARY_PATH environment variable.

The following table describes the scanner options:

Options	Description
-hhelp	Displays general syntax and information on WAL scanner usage.
-vversion	Displays the WAL scanner version information.
-ddebug	Displays debugging output while executing the WAL scanner with any of its options.
<pre>-c <config_file_path>config-path <config_file_path></config_file_path></config_file_path></pre>	Use this option to specify the <pre>config_file_path</pre> of a BART configuration file if you do not want to use the default BART configuration file path <pre>BART_HOME/etc/bart.cfg</pre> .
daemon	Runs the WAL scanner as a background process.
<pre>-p <mbm_file_path> print <mbm_file_path></mbm_file_path></mbm_file_path></pre>	Use this option to specify the full directory path to an MBM file whose content is to be printed. The directory specified in the archive_path parameter in the bart.cfg file contains the MBM files.
<wal_file_path></wal_file_path>	Specify the full directory path to a WAL file to be scanned. The directory specified in the archive_path parameter in the bart.cfg file contains the WAL files. Use this option if a WAL file in the archive path is missing its MBM file. This option is to be used for assisting the EnterpriseDB support team for debugging problems that may have been encountered.
RELOAD	Reloads the BART configuration file. The keyword RELOAD is not case-sensitive. The RELOAD option is useful if you make changes to the configuration file after the WAL scanner has been started. It will reload the configuration file and adjust the WAL scanners accordingly. For example, if a server section allowing incremental backups is removed from the BART configuration file, then the process attached to that server will stop. Similarly, if a server allowing incremental backups is added, a new WAL scanner process will be launched to scan the WAL files of that server.

Options	Description
STOP	Stops the WAL scanner. The keyword STOP is not case-sensitive.
	While invoking the WAL scanner, you can specify one of the following values with the
checksum-	checksum-algorithm option:checksum-algorithm=MD5 (default) to generate
algorithm	MD5 checksum fileschecksum-algorithm=SHA256 to generate SHA256 checksum fileschecksum-algorithm=NONE to skip generating checksum files.

5.4 Using Tablespaces

If the database cluster contains user-defined tablespaces (that is, tablespaces created with the CREATE TABLESPACE command):

- You can take full backups with the BACKUP subcommand in either tar (-Ft) or plain text (-Fp) backup file format.
- You must take incremental backups in the plain text (-Fp) backup file format.
- You can take full backups using the transaction log streaming method (xlog_method = stream in the BART configuration file) --with-pg_basebackup and the BACKUP subcommand in either tar (-Ft) or plain text (-Fp) backup file format.

!!! Note If the particular database cluster you plan to back up contains tablespaces created by the CREATE TABLESPACE command, then you must set the tablespace_path parameter in the BART configuration file before you perform a BART RESTORE operation.

The tablespace_path parameter specifies the directory paths to which you want the tablespaces to be restored. It takes the following format:

```
OID_1=tablespace_path_1;OID_2=tablespace_path_2 ...
```

Where OID_1, OID_2, ... are the Object Identifiers of the tablespaces. You can find the OIDs of the tablespaces and their corresponding soft links to the directories by listing the contents of the POSTGRES_INSTALL_HOME/data/pg_tblspc subdirectory as shown in the following example:

```
[root@localhost pg_tblspc ]# pwd
/opt/PostgresPlus/9.6AS/data/pg_tblspc
[root@localhost pg_tblspc]# ls -l
total 0
lrwxrwxrwx 1 enterprisedb enterprisedb 17 Aug 22 16:38 16644 -> /mnt/tablespace_1
lrwxrwxrwx 1 enterprisedb enterprisedb 17 Aug 22 16:38 16645 -> /mnt/tablespace_2
```

The OIDs are 16644 and 16645 to directories /mnt/tablespace_1 and /mnt/tablespace_2, respectively.

If you later wish to restore the tablespaces to the same locations as indicated in the preceding example, the BART configuration file must contain the following entry:

```
[ACCTG]
host = 127.0.0.1
port = 5444
user = enterprisedb
cluster_owner = enterprisedb
tablespace_path = 16644=/mnt/tablespace_1;16645=/mnt/tablespace_2
description = "Accounting"
```

If you later wish to restore the tablespaces to different locations, specify the new directory locations in the tablespace_path parameter.

In either case, the directories specified in the tablespace_path parameter must exist and be empty at the time you perform the BART RESTORE operation.

If the database server is running on a remote host (in other words you are also using the remote_host configuration parameter or will specify the --remote-host option with the RESTORE subcommand), the specified tablespace directories must exist on the specified remote host.

To view example of backing up and restoring a database cluster on a remote host containing tablespaces, see the *EDB Backup and Recovery Reference Guide* available at the *EDB website*.

The directories must be owned by the user account with which you intend to start the database server (typically the Postgres user account) with no access by other users or groups as is required for the directory path to which the main full backup is to be restored.

To view a sample BART managed backup and recovery system consisting of both local and remote database servers, see the *EDB Backup and Recovery Reference Guide* available at the *EDB website*.