



EDB Postgres Advanced Server

Version 12

| | | |
|---------|--|----|
| 1 | Advanced Server Installation Guide for Linux | 12 |
| 1.1 | Supported Platforms | 12 |
| 1.2 | Using a Package Manager to Install Advanced Server | 12 |
| 1.3 | Installation Troubleshooting | 29 |
| 1.4 | Managing an Advanced Server Installation | 30 |
| 1.5 | Uninstalling Advanced Server | 35 |
| 2 | Advanced Server Installation Guide for Windows | 36 |
| 2.1 | Requirements Overview | 36 |
| 2.2 | Installing Advanced Server with the Interactive Installer | 37 |
| 2.2.1 | Performing a Graphical Installation on Windows | 38 |
| 2.2.2 | Invoking the Graphical Installer from the Command Line | 46 |
| 2.2.2.1 | Performing an Unattended Installation | 46 |
| 2.2.2.2 | Performing an Installation with Limited Privileges | 47 |
| 2.2.2.3 | Reference - Command Line Options | 51 |
| 2.2.3 | Using StackBuilder Plus | 55 |
| 2.2.4 | Using the Update Monitor | 57 |
| 2.2.5 | Installation Troubleshooting | 58 |
| 2.3 | Managing an Advanced Server Installation | 59 |
| 2.3.1 | Using the Windows Services Applet | 59 |
| 2.3.2 | Using pg_ctl to Control Advanced Server | 60 |
| 2.3.3 | Controlling Server Startup Behavior on Windows | 60 |
| 2.4 | Configuring Advanced Server | 61 |
| 2.4.1 | Modifying the postgresql.conf File | 62 |
| 2.4.2 | Modifying the pg_hba.conf File | 62 |
| 2.4.3 | Setting Advanced Server Environment Variables | 63 |
| 2.4.4 | Connecting to Advanced Server with psql | 64 |
| 2.4.5 | Connecting to Advanced Server with the pgAdmin 4 Client | 64 |
| 2.5 | Uninstalling Advanced Server | 66 |
| 3 | Release Notes | 67 |
| 4 | Advanced Server Upgrade Guide | 79 |
| 4.1 | Supported Platforms | 79 |
| 4.2 | Limitations | 79 |
| 4.3 | Upgrading an Installation With pg_upgrade | 79 |
| 4.3.1 | Performing an Upgrade | 80 |
| 4.3.1.1 | Linking versus Copying | 81 |
| 4.3.2 | Invoking pg_upgrade | 81 |
| 4.3.2.1 | Command Line Options - Reference | 82 |
| 4.3.3 | Upgrading to Advanced Server 13 | 84 |
| 4.3.4 | Upgrading a pgAgent Installation | 91 |
| 4.3.5 | pg_upgrade Troubleshooting | 91 |
| 4.3.6 | Reverting to the Old Cluster | 92 |
| 4.4 | Performing a Minor Version Update of an RPM Installation | 93 |
| 4.5 | Using StackBuilder Plus to Perform a Minor Version Update | 93 |
| 5 | Database Compatibility for Oracle Developers Reference Guide | 96 |
| 5.1 | Introduction | 96 |
| 5.2 | The SQL Language | 97 |
| 5.2.1 | SQL Syntax | 97 |
| 5.2.1.1 | Lexical Structure | 97 |

| | | |
|----------|--|-----|
| 5.2.1.2 | Identifiers and Key Words | 98 |
| 5.2.1.3 | Constants | 99 |
| 5.2.1.4 | Comments | 101 |
| 5.2.2 | Data Types | 101 |
| 5.2.2.1 | Numeric Types | 102 |
| 5.2.2.2 | Character Types | 104 |
| 5.2.2.3 | Binary Data | 105 |
| 5.2.2.4 | Date/Time Types | 105 |
| 5.2.2.5 | Boolean Types | 109 |
| 5.2.2.6 | XML Type | 109 |
| 5.2.3 | Functions and Operators | 110 |
| 5.2.3.1 | Logical Operators | 110 |
| 5.2.3.2 | Comparison Operators | 111 |
| 5.2.3.3 | Mathematical Functions and Operators | 112 |
| 5.2.3.4 | String Functions and Operators | 114 |
| 5.2.3.5 | Pattern Matching String Functions | 119 |
| 5.2.3.6 | Pattern Matching Using the LIKE Operator | 123 |
| 5.2.3.7 | Data Type Formatting Functions | 124 |
| 5.2.3.8 | Date/Time Functions and Operators | 130 |
| 5.2.3.9 | Sequence Manipulation Functions | 145 |
| 5.2.3.10 | Conditional Expressions | 146 |
| 5.2.3.11 | Aggregate Functions | 149 |
| 5.2.3.12 | Subquery Expressions | 155 |
| 5.3 | System Catalog Tables | 157 |
| 6 | Database Compatibility for Oracle | 161 |
| 6.1 | Introduction | 161 |
| 6.2 | Packages | 161 |
| 6.2.1 | Package Components | 162 |
| 6.2.2 | Creating Packages | 168 |
| 6.2.3 | Referencing a Package | 171 |
| 6.2.4 | Using Packages With User Defined Types | 171 |
| 6.2.5 | Dropping a Package | 174 |
| 6.3 | Built-In Packages | 175 |
| 6.3.1 | DBMS_ALERT | 175 |
| 6.3.2 | DBMS_AQ | 182 |
| 6.3.2.1 | ENQUEUE | 184 |
| 6.3.2.2 | DEQUEUE | 186 |
| 6.3.2.3 | REGISTER | 189 |
| 6.3.2.4 | UNREGISTER | 190 |
| 6.3.3 | DBMS_AQADM | 192 |
| 6.3.3.1 | ALTER_QUEUE | 193 |
| 6.3.3.2 | ALTER_QUEUE_TABLE | 194 |
| 6.3.3.3 | CREATE_QUEUE | 195 |
| 6.3.3.4 | CREATE_QUEUE_TABLE | 196 |
| 6.3.3.5 | DROP_QUEUE | 198 |
| 6.3.3.6 | DROP_QUEUE_TABLE | 198 |
| 6.3.3.7 | PURGE_QUEUE_TABLE | 199 |
| 6.3.3.8 | START_QUEUE | 200 |

| | | |
|----------|---------------------|-----|
| 6.3.3.9 | STOP_QUEUE | 201 |
| 6.3.4 | DBMS_CRYPTO | 201 |
| 6.3.4.1 | DECRYPT | 202 |
| 6.3.4.2 | ENCRYPT | 204 |
| 6.3.4.3 | HASH | 206 |
| 6.3.4.4 | MAC | 207 |
| 6.3.4.5 | RANDOMBYTES | 207 |
| 6.3.4.6 | RANDOMINTEGER | 208 |
| 6.3.4.7 | RANDOMNUMBER | 208 |
| 6.3.5 | DBMS_JOB | 209 |
| 6.3.5.1 | BROKEN | 210 |
| 6.3.5.2 | CHANGE | 211 |
| 6.3.5.3 | INTERVAL | 211 |
| 6.3.5.4 | NEXT_DATE | 212 |
| 6.3.5.5 | REMOVE | 212 |
| 6.3.5.6 | RUN | 213 |
| 6.3.5.7 | SUBMIT | 213 |
| 6.3.5.8 | WHAT | 214 |
| 6.3.6 | DBMS_LOB | 215 |
| 6.3.6.1 | APPEND | 216 |
| 6.3.6.2 | COMPARE | 217 |
| 6.3.6.3 | CONVERTTOBLOB | 217 |
| 6.3.6.4 | CONVERTTOCLOB | 218 |
| 6.3.6.5 | COPY | 219 |
| 6.3.6.6 | ERASE | 220 |
| 6.3.6.7 | GET_STORAGE_LIMIT | 221 |
| 6.3.6.8 | GETLENGTH | 221 |
| 6.3.6.9 | INSTR | 221 |
| 6.3.6.10 | READ | 222 |
| 6.3.6.11 | SUBSTR | 223 |
| 6.3.6.12 | TRIM | 223 |
| 6.3.6.13 | WRITE | 224 |
| 6.3.6.14 | WRITEAPPEND | 224 |
| 6.3.7 | DBMS_LOCK | 225 |
| 6.3.8 | DBMS_MVIEW | 225 |
| 6.3.8.1 | GET_MV_DEPENDENCIES | 226 |
| 6.3.8.2 | REFRESH | 226 |
| 6.3.8.3 | REFRESH_ALL_MVIEWS | 228 |
| 6.3.8.4 | REFRESH_DEPENDENT | 229 |
| 6.3.9 | DBMS_OUTPUT | 230 |
| 6.3.10 | DBMS_PIPE | 238 |
| 6.3.10.1 | CREATE_PIPE | 239 |
| 6.3.10.2 | NEXT_ITEM_TYPE | 240 |
| 6.3.10.3 | PACK_MESSAGE | 242 |
| 6.3.10.4 | PURGE | 242 |
| 6.3.10.5 | RECEIVE_MESSAGE | 244 |
| 6.3.10.6 | REMOVE_PIPE | 244 |
| 6.3.10.7 | RESET_BUFFER | 246 |

| | | |
|-----------|---|-----|
| 6.3.10.8 | SEND_MESSAGE | 247 |
| 6.3.10.9 | UNIQUE_SESSION_NAME | 247 |
| 6.3.10.10 | UNPACK_MESSAGE | 248 |
| 6.3.10.11 | Comprehensive Example | 248 |
| 6.3.11 | DBMS_PROFILER | 251 |
| 6.3.12 | DBMS_RANDOM | 270 |
| 6.3.13 | DBMS_REDACT | 274 |
| 6.3.14 | DBMS_RLS | 296 |
| 6.3.15 | DBMS_SCHEDULER | 306 |
| 6.3.15.1 | Using Calendar Syntax to Specify a Repeating Interval | 308 |
| 6.3.15.2 | CREATE_JOB | 309 |
| 6.3.15.3 | CREATE_PROGRAM | 311 |
| 6.3.15.4 | CREATE_SCHEDULE | 312 |
| 6.3.15.5 | DEFINE_PROGRAM_ARGUMENT | 313 |
| 6.3.15.6 | DISABLE | 315 |
| 6.3.15.7 | DROP_JOB | 315 |
| 6.3.15.8 | DROP_PROGRAM | 316 |
| 6.3.15.9 | DROP_PROGRAM_ARGUMENT | 317 |
| 6.3.15.10 | DROP_SCHEDULE | 318 |
| 6.3.15.11 | ENABLE | 318 |
| 6.3.15.12 | EVALUATE_CALENDAR_STRING | 319 |
| 6.3.15.13 | RUN_JOB | 320 |
| 6.3.15.14 | SET_JOB_ARGUMENT_VALUE | 320 |
| 6.3.16 | DBMS_SESSION | 321 |
| 6.3.17 | DBMS_SQL | 322 |
| 6.3.17.1 | BIND_VARIABLE | 323 |
| 6.3.17.2 | BIND_VARIABLE_CHAR | 325 |
| 6.3.17.3 | BIND VARIABLE RAW | 325 |
| 6.3.17.4 | CLOSE_CURSOR | 326 |
| 6.3.17.5 | COLUMN_VALUE | 326 |
| 6.3.17.6 | COLUMN_VALUE_CHAR | 327 |
| 6.3.17.7 | COLUMN_VALUE_RAW | 328 |
| 6.3.17.8 | DEFINE_COLUMN | 329 |
| 6.3.17.9 | DEFINE_COLUMN_CHAR | 330 |
| 6.3.17.10 | DEFINE_COLUMN_RAW | 331 |
| 6.3.17.11 | DESCRIBE_COLUMNS | 331 |
| 6.3.17.12 | EXECUTE | 332 |
| 6.3.17.13 | EXECUTE_AND_FETCH | 333 |
| 6.3.17.14 | FETCH_ROWS | 334 |
| 6.3.17.15 | IS_OPEN | 336 |
| 6.3.17.16 | LAST_ROW_COUNT | 336 |
| 6.3.17.17 | OPEN_CURSOR | 338 |
| 6.3.17.18 | PARSE | 338 |
| 6.3.18 | DBMS.Utility | 340 |
| 6.3.19 | UTL_ENCODE | 353 |
| 6.3.19.1 | BASE64_DECODE | 354 |
| 6.3.19.2 | BASE64_ENCODE | 355 |
| 6.3.19.3 | MIMEHEADER_DECODE | 356 |

| | | |
|-----------|---|-----|
| 6.3.19.4 | MIMEHEADER_ENCODE | 356 |
| 6.3.19.5 | QUOTED_PRINTABLE_DECODE | 357 |
| 6.3.19.6 | QUOTED_PRINTABLE_ENCODE | 358 |
| 6.3.19.7 | TEXT_DECODE | 359 |
| 6.3.19.8 | TEXT_ENCODE | 360 |
| 6.3.19.9 | UUDECODE | 361 |
| 6.3.19.10 | UUENCODE | 362 |
| 6.3.20 | UTL_FILE | 363 |
| 6.3.21 | UTL_HTTP | 378 |
| 6.3.22 | UTL_MAIL | 396 |
| 6.3.23 | UTL_RAW | 400 |
| 6.3.24 | UTL_SMTP | 404 |
| 6.3.25 | UTL_URL | 412 |
| 7 | Database Compatibility for Oracle Developers | 415 |
| 7.1 | Introduction | 415 |
| 7.1.1 | Configuration Parameters Compatible with Oracle Databases | 416 |
| 7.1.1.1 | edb_redwood_date | 417 |
| 7.1.1.2 | edb_redwood_raw_names | 417 |
| 7.1.1.3 | edb_redwood_strings | 418 |
| 7.1.1.4 | edb_stmt_level_tx | 419 |
| 7.1.1.5 | oracle_home | 421 |
| 7.1.2 | About the Examples Used in this Guide | 421 |
| 7.2 | SQL Tutorial | 422 |
| 7.2.1 | Sample Database | 422 |
| 7.2.1.1 | Sample Database Installation | 422 |
| 7.2.1.2 | Sample Database Description | 423 |
| 7.2.2 | Creating a New Table | 434 |
| 7.2.3 | Populating a Table With Rows | 434 |
| 7.2.4 | Querying a Table | 435 |
| 7.2.5 | Joins Between Tables | 437 |
| 7.2.6 | Aggregate Functions | 440 |
| 7.2.7 | Updates | 441 |
| 7.2.8 | Deletions | 442 |
| 7.2.9 | The SQL Language | 443 |
| 7.3 | Advanced Concepts | 443 |
| 7.3.1 | Views | 443 |
| 7.3.2 | Foreign Keys | 444 |
| 7.3.3 | The ROWNUM Pseudo-Column | 445 |
| 7.3.4 | Synonyms | 446 |
| 7.3.5 | Hierarchical Queries | 449 |
| 7.3.5.1 | Defining the Parent/Child Relationship | 450 |
| 7.3.5.2 | Selecting the Root Nodes | 450 |
| 7.3.5.3 | Organization Tree in the Sample Application | 451 |
| 7.3.5.4 | Node Level | 452 |
| 7.3.5.5 | Ordering the Siblings | 452 |
| 7.3.5.6 | Retrieving the Root Node with CONNECT_BY_ROOT | 453 |
| 7.3.5.7 | Retrieving a Path with SYS_CONNECT_BY_PATH | 457 |
| 7.3.6 | Multidimensional Analysis | 458 |

| | | |
|---------|--|-----|
| 7.3.6.1 | ROLLUP Extension | 460 |
| 7.3.6.2 | CUBE Extension | 462 |
| 7.3.6.3 | GROUPING SETS Extension | 466 |
| 7.3.6.4 | GROUPING Function | 471 |
| 7.3.6.5 | GROUPING_ID Function | 474 |
| 7.4 | Profile Management | 476 |
| 7.4.1 | Creating a New Profile | 477 |
| 7.4.1.1 | Creating a Password Function | 479 |
| 7.4.2 | Altering a Profile | 481 |
| 7.4.3 | Dropping a Profile | 482 |
| 7.4.4 | Associating a Profile with an Existing Role | 482 |
| 7.4.5 | Unlocking a Locked Account | 484 |
| 7.4.6 | Creating a New Role Associated with a Profile | 485 |
| 7.4.7 | Backing up Profile Management Functions | 486 |
| 7.5 | Optimizer Hints | 487 |
| 7.5.1 | Default Optimization Modes | 488 |
| 7.5.2 | Access Method Hints | 490 |
| 7.5.3 | Specifying a Join Order | 494 |
| 7.5.4 | Joining Relations Hints | 494 |
| 7.5.5 | Global Hints | 497 |
| 7.5.6 | Using the APPEND Optimizer Hint | 499 |
| 7.5.7 | Parallelism Hints | 500 |
| 7.5.8 | Conflicting Hints | 504 |
| 7.6 | dblink_ora | 504 |
| 7.6.1 | dblink_ora Functions and Procedures | 505 |
| 7.6.1.1 | dblink_ora_connect() | 505 |
| 7.6.1.2 | dblink_ora_status() | 506 |
| 7.6.1.3 | dblink_ora_disconnect() | 506 |
| 7.6.1.4 | dblink_ora_record() | 507 |
| 7.6.1.5 | dblink_ora_call() | 507 |
| 7.6.1.6 | dblink_ora_exec() | 507 |
| 7.6.1.7 | dblink_ora_copy() | 508 |
| 7.6.2 | Calling dblink_ora Functions | 508 |
| 7.7 | Open Client Library | 509 |
| 7.8 | Oracle Catalog Views | 510 |
| 7.9 | Tools and Utilities | 510 |
| 7.10 | ECPGPlus | 510 |
| 7.11 | System Catalog Tables | 511 |
| 8 | Database Compatibility for Oracle | 511 |
| 8.1 | Introduction | 511 |
| 8.2 | EDB Loader | 512 |
| 8.3 | EDB Wrap | 541 |
| 8.4 | Dynamic Runtime Instrumentation Tools Architecture (DRITA) | 545 |
| 9 | ECPGPlus Guide | 573 |
| 9.1 | Introduction | 573 |
| 9.2 | ECPGPlus - Overview | 573 |
| 9.3 | Using Embedded SQL | 579 |
| 9.4 | Using Descriptors | 585 |

| | | |
|---------------|---|-----|
| 9.5 | Building and Executing Dynamic SQL Statements | 595 |
| 9.6 | Error Handling | 610 |
| 9.7 | Reference | 614 |
| 10 | Language Pack Guide | 643 |
| 10.1 | Supported Database Server Versions | 644 |
| 10.2 | Installing Language Pack | 645 |
| 10.3 | Using the Procedural Languages | 649 |
| 11 | Advanced Server User Guide | 651 |
| 11.1 | Introduction | 663 |
| 11.1.1 | What's New | 664 |
| 11.1.2 | Conventions Used in this Guide | 665 |
| 11.1.3 | About the Examples Used in this Guide | 665 |
| 11.2 | Enhanced Compatibility Features | 675 |
| 11.3 | Database Administration | 680 |
| 11.3.1 | Configuration Parameters | 680 |
| 11.3.1.1 | Setting Configuration Parameters | 681 |
| 11.3.1.2 | Summary of Configuration Parameters | 682 |
| 11.3.1.3 | Configuration Parameters by Functionality | 687 |
| 11.3.1.3.1 | Top Performance Related Parameters | 687 |
| 11.3.1.3.1.1 | shared_buffers | 687 |
| 11.3.1.3.1.2 | work_mem | 688 |
| 11.3.1.3.1.3 | maintenance_work_mem | 688 |
| 11.3.1.3.1.4 | wal_buffers | 689 |
| 11.3.1.3.1.5 | checkpoint_segments | 690 |
| 11.3.1.3.1.6 | checkpoint_completion_target | 690 |
| 11.3.1.3.1.7 | checkpoint_timeout | 690 |
| 11.3.1.3.1.8 | max_wal_size | 691 |
| 11.3.1.3.1.9 | min_wal_size | 691 |
| 11.3.1.3.1.10 | bgwriter_delay | 691 |
| 11.3.1.3.1.11 | seq_page_cost | 692 |
| 11.3.1.3.1.12 | random_page_cost | 692 |
| 11.3.1.3.1.13 | effective_cache_size | 693 |
| 11.3.1.3.1.14 | synchronous_commit | 694 |
| 11.3.1.3.1.15 | edb_max_spins_per_delay | 694 |
| 11.3.1.3.1.16 | pg_prewarm.autoprewarm | 695 |
| 11.3.1.3.1.17 | pg_prewarm.autoprewarm_interval | 695 |
| 11.3.1.3.2 | Resource Usage / Memory | 696 |
| 11.3.1.3.3 | Resource Usage / EDB Resource Manager | 697 |
| 11.3.1.3.4 | Query Tuning | 698 |
| 11.3.1.3.5 | Query Tuning / Planner Method Configuration | 698 |
| 11.3.1.3.6 | Reporting and Logging / What to Log | 699 |
| 11.3.1.3.7 | Auditing Settings | 700 |
| 11.3.1.3.7.1 | edb_audit | 700 |
| 11.3.1.3.7.2 | edb_audit_directory | 701 |
| 11.3.1.3.7.3 | edb_audit_filename | 701 |
| 11.3.1.3.7.4 | edb_audit_rotation_day | 702 |
| 11.3.1.3.7.5 | edb_audit_rotation_size | 702 |
| 11.3.1.3.7.6 | edb_audit_rotation_seconds | 702 |

| | | |
|---------------|--|-----|
| 11.3.1.3.7.7 | edb_audit_connect | 703 |
| 11.3.1.3.7.8 | edb_audit_disconnect | 703 |
| 11.3.1.3.7.9 | edb_audit_statement | 703 |
| 11.3.1.3.7.10 | edb_audit_tag | 704 |
| 11.3.1.3.7.11 | edb_audit_destination | 704 |
| 11.3.1.3.7.12 | edb_log_every_bulk_value | 705 |
| 11.3.1.3.8 | Client Connection Defaults / Locale and Formatting | 705 |
| 11.3.1.3.9 | Client Connection Defaults / Statement Behavior | 706 |
| 11.3.1.3.10 | Client Connection Defaults / Other Defaults | 707 |
| 11.3.1.3.11 | Compatibility Options | 708 |
| 11.3.1.3.12 | Customized Options | 715 |
| 11.3.1.3.13 | Ungrouped | 722 |
| 11.3.2 | Index Advisor | 724 |
| 11.3.2.1 | Index Advisor Components | 725 |
| 11.3.2.2 | Index Advisor Configuration | 726 |
| 11.3.2.3 | Using Index Advisor | 727 |
| 11.3.2.4 | Reviewing the Index Advisor Recommendations | 730 |
| 11.3.2.5 | Limitations | 734 |
| 11.3.3 | SQL Profiler | 734 |
| 11.3.4 | pgsnmpd | 736 |
| 11.3.5 | EDB Audit Logging | 738 |
| 11.3.5.1 | Audit Logging Configuration Parameters | 738 |
| 11.3.5.2 | Selecting SQL Statements to Audit | 740 |
| 11.3.5.3 | Enabling Audit Logging | 748 |
| 11.3.5.4 | Audit Log File | 753 |
| 11.3.5.5 | Using Error Codes to Filter Audit Logs | 758 |
| 11.3.5.6 | Using Command Tags to Filter Audit Logs | 758 |
| 11.3.5.7 | Redacting Passwords from Audit Logs | 759 |
| 11.3.6 | Unicode Collation Algorithm | 760 |
| 11.4 | Configuration Parameters | 767 |
| 11.4.1 | Setting Configuration Parameters | 767 |
| 11.4.2 | Summary of Configuration Parameters | 769 |
| 11.4.3 | Configuration Parameters by Functionality | 773 |
| 11.4.3.1 | Top Performance Related Parameters | 774 |
| 11.4.3.1.1 | shared_buffers | 774 |
| 11.4.3.1.2 | work_mem | 774 |
| 11.4.3.1.3 | maintenance_work_mem | 775 |
| 11.4.3.1.4 | wal_buffers | 775 |
| 11.4.3.1.5 | checkpoint_segments | 776 |
| 11.4.3.1.6 | checkpoint_completion_target | 776 |
| 11.4.3.1.7 | checkpoint_timeout | 777 |
| 11.4.3.1.8 | max_wal_size | 777 |
| 11.4.3.1.9 | min_wal_size | 777 |
| 11.4.3.1.10 | bgwriter_delay | 778 |
| 11.4.3.1.11 | seq_page_cost | 778 |
| 11.4.3.1.12 | random_page_cost | 779 |
| 11.4.3.1.13 | effective_cache_size | 779 |
| 11.4.3.1.14 | synchronous_commit | 780 |

| | | |
|-------------|--|-----|
| 11.4.3.1.15 | edb_max_spins_per_delay | 781 |
| 11.4.3.1.16 | pg_prewarm.autoprewarm | 781 |
| 11.4.3.1.17 | pg_prewarm.autoprewarm_interval | 782 |
| 11.4.3.2 | Resource Usage / Memory | 782 |
| 11.4.3.3 | Resource Usage / EDB Resource Manager | 783 |
| 11.4.3.4 | Query Tuning | 784 |
| 11.4.3.5 | Query Tuning / Planner Method Configuration | 785 |
| 11.4.3.6 | Reporting and Logging / What to Log | 785 |
| 11.4.3.7 | Auditing Settings | 787 |
| 11.4.3.7.1 | edb_audit | 787 |
| 11.4.3.7.2 | edb_audit_directory | 787 |
| 11.4.3.7.3 | edb_audit_filename | 788 |
| 11.4.3.7.4 | edb_audit_rotation_day | 788 |
| 11.4.3.7.5 | edb_audit_rotation_size | 788 |
| 11.4.3.7.6 | edb_audit_rotation_seconds | 789 |
| 11.4.3.7.7 | edb_audit_connect | 789 |
| 11.4.3.7.8 | edb_audit_disconnect | 789 |
| 11.4.3.7.9 | edb_audit_statement | 790 |
| 11.4.3.7.10 | edb_audit_tag | 790 |
| 11.4.3.7.11 | edb_audit_destination | 791 |
| 11.4.3.7.12 | edb_log_every_bulk_value | 791 |
| 11.4.3.8 | Client Connection Defaults / Locale and Formatting | 791 |
| 11.4.3.9 | Client Connection Defaults / Statement Behavior | 792 |
| 11.4.3.10 | Client Connection Defaults / Other Defaults | 793 |
| 11.4.3.11 | Compatibility Options | 794 |
| 11.4.3.12 | Customized Options | 801 |
| 11.4.3.13 | Ungrouped | 809 |
| 11.5 | Index Advisor | 811 |
| 11.5.1 | Index Advisor Components | 812 |
| 11.5.2 | Index Advisor Configuration | 813 |
| 11.5.3 | Using Index Advisor | 814 |
| 11.5.4 | Reviewing the Index Advisor Recommendations | 817 |
| 11.5.5 | Limitations | 820 |
| 11.6 | SQL Profiler | 821 |
| 11.7 | pgsnmpd | 822 |
| 11.8 | EDB Audit Logging | 825 |
| 11.8.1 | Audit Logging Configuration Parameters | 825 |
| 11.8.2 | Selecting SQL Statements to Audit | 826 |
| 11.8.3 | Enabling Audit Logging | 835 |
| 11.8.4 | Audit Log File | 839 |
| 11.8.5 | Using Error Codes to Filter Audit Logs | 844 |
| 11.8.6 | Using Command Tags to Filter Audit Logs | 845 |
| 11.8.7 | Redacting Passwords from Audit Logs | 846 |
| 11.9 | Unicode Collation Algorithm | 847 |
| 11.10 | EDB Resource Manager | 854 |
| 11.11 | libpq C Library | 873 |
| 11.12 | Debugger | 882 |
| 11.13 | Performance Analysis and Tuning | 897 |

| | | |
|-------|---|-----|
| 11.14 | EDB Clone Schema | 906 |
| 11.15 | Enhanced SQL and Other Miscellaneous Features | 927 |
| 11.16 | System Catalog Tables | 931 |
| 11.17 | Advanced Server Keywords | 935 |
| 12 | EDB Plus User Guide | 936 |
| 12.1 | What's New | 936 |
| 12.2 | EDB Plus | 936 |
| 12.3 | Installing EDB Plus | 937 |
| 12.4 | Using EDB Plus | 941 |
| 12.5 | Using a Secure Sockets Layer (SSL) Connection | 943 |
| 12.6 | Command Summary | 950 |

1 Advanced Server Installation Guide for Linux

The *EDB Postgres Advanced Server Installation Guide* is a comprehensive guide to installing EDB Postgres Advanced Server (Advanced Server). In this guide you will find detailed information about:

- Software prerequisites for performing an Advanced Server 12 installation on a Linux host.
- Using a package manager to install and update Advanced Server and its supporting components or utilities on a Linux host.
- Managing an Advanced Server installation.
- Configuring an Advanced Server package installation.
- Uninstalling Advanced Server and its components.

1.1 Supported Platforms

For information about the platforms and versions supported by Advanced Server, visit the EnterpriseDB website at:

<https://www.enterprisedb.com/blog/edb-supported-products-and-platforms>

Limitations

The following limitations apply to EDB Postgres Advanced Server:

- The `data` directory of a production database should not be stored on an NFS file system.
- The LLVM JIT package is supported on RHEL 7.x only. LLVM JIT is not supported on RHEL 6.x and PPC-LE 64 (running RHEL or CentOS 7.x).

1.2 Using a Package Manager to Install Advanced Server

You can use the yum package manager to install Advanced Server or Advanced Server supporting components. yum will attempt to satisfy package dependencies as it installs a package, but requires access to the Advanced Server repositories. If your system does not have access to a repository via the Internet, you can use RPM to install an individual package or create a local repository, but you may be required to manually satisfy package dependencies.

You can list the dependencies of a package by running the following command:

- On Fedora/RHEL/CentOS: `repoquery --requires --resolve <package_name>`
- On Debian/Ubuntu: `apt-cache depends <package_name>`
- On SLES: `zypper search --requires --verbose <package_name>`

Where, `<package_name>` is the name of the package that you want to install.

Installing the server package creates a database superuser named `enterprisedb`. The user is assigned a user ID (UID) and a group ID (GID) of `26`. The user has no default password; use the `passwd` command to assign a password for the user. The default shell for the user is `bash`, and the user's home directory is `/var/lib/edb/as12`.

By default, Advanced Server logging is configured to write files to the `log` subdirectory of the `data` directory, rotating the files each day and retaining one week of log entries. You can customize the logging behavior of the server by [modifying the postgresql.conf File](#).

The RPM installers place Advanced Server components in the directories listed in the table below:

| Component | Location |
|-----------------------------------|---|
| Executables | <code>/usr/edb/as12/bin</code> |
| Libraries | <code>/usr/edb/as12/lib</code> |
| Cluster configuration files | <code>/etc/edb/as12</code> |
| Documentation | <code>/usr/edb/as12/share/doc</code> |
| Contrib | <code>/usr/edb/as12/share/contrib</code> |
| Data | <code>/var/lib/edb/as12/data</code> |
| Logs | <code>/var/log/as12</code> |
| Lock files | <code>/var/lock/as12</code> |
| Log rotation file | <code>/etc/logrotate.d/as12</code> |
| Sudo configuration file | <code>/etc/sudoers.d/as12</code> |
| Binary to access VIP without sudo | <code>/usr/edb/as12/bin/secure</code> |
| Backup area | <code>/var/lib/edb/as12/backups</code> |
| Templates | <code>/usr/edb/as12/share</code> |
| Procedural Languages | <code>/usr/edb/as12/lib or /usr/edb/as12/lib64</code> |
| Development Headers | <code>/usr/edb/as12/include</code> |
| Shared data | <code>/usr/edb/as12/share</code> |
| Regression tests | <code>/usr/edb/as12/lib/pgxs/src/test/regress</code> |
| SGML Documentation | <code>/usr/edb/as12/share/doc</code> |

Installing Advanced Server on a Linux Host

Before using an RPM package to install Advanced Server on a Linux host, you must:

Install Linux-specific Software

You must install `xterm`, `konsole`, or `gnome-terminal` before executing any console-based program installed by EnterpriseDB installers.

Install Migration Toolkit or EDB*Plus Installation Prerequisites (Optional)

Before using an RPM to install Migration Toolkit or EDB*Plus, you must first install Java version 1.7 or later for Migration Toolkit and Java version 1.8 or later for EDB*Plus. On a Linux system, you can use the yum package manager to install Java. Open a terminal window, assume superuser privileges, and enter:

```
# yum install java
```

Follow the onscreen instructions to complete the installation.

Request Credentials to the EnterpriseDB Repository

You must have credentials that allow access to the EnterpriseDB repository. For information about requesting credentials, visit:

<https://info.enterprisedb.com/rs/069-ALB-339/images/Repository%20Access%2004-09-2019.pdf>.

After receiving your repository credentials you can:

1. Create the repository configuration file.
2. Modify the file, providing your user name and password.
3. Install the repository keys and additional prerequisite software.
4. Install Advanced Server and supporting components.

Creating a Repository Configuration File and Installing Advanced Server

To create the repository configuration file, assume superuser privileges and invoke the following command:

```
yum -y install https://yum.enterprisedb.com/edb-repo-rpms/edb-repo-latest.noarch.rpm
```

The repository configuration file is named `edb.repo`. The file resides in `/etc/yum.repos.d`.

After creating the `edb.repo` file, use your choice of editor to set the value of the `enabled` parameter to `1`, and replace the `username` and `password` placeholders in the `baseurl` specification with the name and password of a registered EnterpriseDB user.

```
[edb]
name=EnterpriseDB RPMs $releasever - $basearch
baseurl=https://<*username>:<*password>@\ yum.enterprisedb.com/edb/redhat/rhel-$releasever-$basearch
enabled=1
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/ENTERPRISEDB-GPG-KEY
```

After saving your changes to the configuration file, you must download and install the repository keys:

1. Use the following command to download the repository key:

```
curl -o /etc/pki/rpm-gpg/ENTERPRISEDB-GPG-KEY https://yum.enterprisedb.com/ENTERPRISEDB-GPG-KEY
```

2. Use the following command to install the key:

```
rpm --import /etc/pki/rpm-gpg/ENTERPRISEDB-GPG-KEY
```

3. Use the following commands to install the prerequisite software:

```
# yum -y install epel-release
```

```
# yum makecache
```

Then, you can use `yum install` command to install Advanced Server. For example, to install the server and its core components, use the command:

```
yum install edb-as12-server
```

When you install an RPM package that is signed by a source that is not recognized by your system, `yum` may ask for your permission to import the key to your local server. If prompted, and you are satisfied that the packages come from a trustworthy source, enter a `y`, and press `Return` to continue.

After installing Advanced Server, you must configure the installation. For more information, see [Configuring a Package Installation](#).

During the installation, yum may encounter a dependency that it cannot resolve. If it does, it will provide a list of the required dependencies that you must manually resolve.

Advanced Server RPM Packages

The tables that follow list the RPM packages that are available from EnterpriseDB. You can also use the yum search command to access a list of the packages that are currently available from your configured repository. Open a command line, assume superuser privileges, and enter:

```
yum search <package>
```

Where `package` is the search term that specifies the name (or partial name) of a package.

Please note: The available package list is subject to change.

| Package Name | Package Installs |
|--|--|
| edb-as12-server | This package installs core components of the Advanced Server database server. |
| edb-as12-server-client | Client programs and utilities that you can use to access and manage Advanced Server. |
| edb-as12-server-contrib | <p>Installs contributed tools and utilities that are distributed with Advanced Server. Files for these modules are installed in:</p> <p>Documentation: <code>/usr/edb/as12/share/doc</code></p> <p>Loadable modules: <code>/usr/edb/as12/lib</code></p> <p>Binaries: <code>/usr/edb/as12/bin</code></p> |
| edb-as12-server-core | Includes the programs needed to create the core functionality behind the Advanced Server database. |
| edb-as12-server-devel | Installs the header files and libraries needed to compile C or C++ applications that directly interact with an Advanced Server server and the ecpg or ecpgPlus C preprocessor. |
| edb-as12-server-docs | Installs the readme file. |
| edb-as12-server-edb-modules | Installs supporting modules for Advanced Server |
| edb-as12-server-indexadvisor | Installs Advanced Server's Index Advisor feature. The Index Advisor utility helps determine which columns you should index to improve performance in a given workload. |
| edb-as12-server-libs | Provides the essential shared libraries for any Advanced Server client program or interface. |
| edb-as12-server-llvmjit | This package contains support for Just in Time (JIT) compiling parts of EDBAS queries. |
| edb-as12-server-pldebugger | This package implements an API for debugging PL/pgSQL functions on Advanced Server. |
| edb-as12-server-plperl | Installs the PL/Perl procedural language for Advanced Server. Please note that the <code>edb-as12-server-plperl</code> package is dependent on the platform-supplied version of Perl. |
| edb-as12-server-plpython edb-as12-server-plpython3 | Installs the PL/Python procedural language for Advanced Server. Please note that the <code>edb-as12-server-plpython</code> package is dependent on the platform-supplied version of Python. By default, the <code>edb-as12-server-plpython3</code> package is not installed; you need to explicitly install the package using yum. |
| edb-as12-server-ptcl | Installs the PL/Tcl procedural language for Advanced Server. Please note that the <code>edb-as12-ptcl</code> package is dependent on the platform-supplied version of TCL. |
| edb-as12-server-sqlprofiler | This package installs Advanced Server's SQL Profiler feature. SQL Profiler helps identify and optimize SQL code. |

| Package Name | Package Installs |
|----------------------------------|---|
| edb-as12-server-sqlprotect | This package installs Advanced Server's SQL Protect feature. SQL Protect provides protection against SQL injection attacks. |
| edb-as12-server-sslutils | This package installs functionality that provides SSL support. |
| edb-as12-server-cloneschema | This package installs the EDB Clone Schema extension. For more information about EDB Clone Schema, see the EDB Postgres Advanced Server Guide. |
| edb-as12-server-parallel-clone | This package installs functionality that supports the EDB Clone Schema extension. |
| edb-as12-pgagent | Installs pgAgent; pgAgent is a job scheduler for Advanced Server. Before installing this package, you must install EPEL; for detailed information about installing EPEL, see Installing Advanced Server on a Linux Host . |
| edb-as12-edbplus | The edb-edbplus package contains the files required to install the EDB Plus command line client. EDB Plus commands are compatible with Oracle's SQL*Plus. |
| edb-as12-pgsnmpd | SNMP (Simple Network Management Protocol) is a protocol that allows you to supervise an apparatus connected to the network. |
| edb-as12-pljava | Installs PL/Java, providing access to Java stored procedures, triggers and functions via the JDBC interface. |
| edb-as12-pgpool40-extensions | This package creates pgPool extensions required by the server for use with pgpool. |
| edb-as12-postgis-2.5 | Installs POSTGIS meta RPMs. |
| edb-as12-postgis-core | This package provides support for geographic objects to the PostgreSQL object-relational database. In effect, PostGIS "spatially enables" the PostgreSQL server, allowing it to be used as a backend spatial database for geographic information systems (GIS), much like ESRI's SDE or Oracle's Spatial extension. |
| edb-as12-postgis-docs | This package installs pdf documentation of PostGIS. |
| edb-as12-postgis-jdbc | This package installs the essential jdbc driver for PostGIS. |
| edb-as12-postgis-utils | This package installs the utilities for PostGIS. |
| edb-as12-slony-replication | Installs the meta RPM for Slony-I. |
| edb-as12-slony-replication-core | Slony-I builds a master-slave system that includes all features and capabilities needed to replicate large databases to a reasonably limited number of slave systems. |
| edb-as12-slony-replication-docs | This package contains the Slony project documentation (in pdf form). |
| edb-as12-slony-replication-tools | This package contains the Slony altperl tools and utilities that are useful when deploying Slony replication environments. Before installing this package, you must install EPEL; for detailed information about installing EPEL, see Installing Advanced Server on a Linux Host . |
| edb-as12-libicu | These packages contain supporting library files. |

The following table lists the packages for Advanced Server 12 supporting components.

| Package Name | Package Installs |
|--------------|------------------|
|--------------|------------------|

| Package Name | Package Installs |
|-----------------------------|---|
| edb-pgpool40 | <p>This package contains the pgPool-II installer. The pgpool-II utility package acts as a middleman between client applications and Server database servers. pgpool-II functionality is transparent to client applications; client applications connect to pgpool-II instead of directly to Advanced Server, and pgpool-II manages the connection. EnterpriseDB supports the following pgpool-II features:</p> <ul style="list-style-type: none"> • Load balancing • Connection pooling • High availability • Connection limits <p>pgpool-II runs as a service on Linux systems, and is not supported on Windows systems.</p> |
| edb-jdbc | The edb-jdbc package includes the .jar files needed for Java programs to access an Advanced Server database. |
| edb-migrationtoolkit | The edb-migrationtoolkit package installs Migration Toolkit, facilitating migration to an Advanced Server database from Oracle, PostgreSQL, MySQL, Sybase and SQL Server. |
| edb-oci | The edb-oci package installs the EnterpriseDB Open Client library, allowing applications that use the Oracle Call Interface API to connect to an Advanced Server database. |
| edb-oci-devel | This package installs the OCI include files; install this package if you are developing C/C++ applications that require these files. |
| edb-odbc | This package installs the driver needed for applications to access an Advanced Server system via ODBC. |
| edb-odbc-devel | This package installs the ODBC include files; install this package if you are developing C/C++ applications that require these files. |
| edb-pgbouncer112 | This package contains PgBouncer (a lightweight connection pooler). This package requires the libevent package. |
| ppas-xdb | This package contains the xDB installer; xDB provides asynchronous cross-database replication. For more information, visit http://www.enterprisedb.com/faq-xdb-multi-master |
| ppas-xdb-console | This package provides support for xDB. |
| ppas-xdb-libs | This package provides support for xDB. |
| ppas-xdb-publisher | This package provides support for xDB. |
| ppas-xdb-subscriber | This package provides support for xDB. |
| edb-pem | The edb-pem package installs Management Tool that efficiently manages, monitor, and tune large Postgres deployments from a single remote GUI console. |
| edb-pem-agent | This package is an agent component of Postgres Enterprise Manager. |
| edb-pem-docs | This package contains documentation for various languages, which are in HTML format. |
| edb-pem-server | This package contains server components of Postgres Enterprise Manager. |
| edb-pgadmin4 | This package is a management tool for PostgreSQL capable of hosting the Python application and presenting it to the user as a desktop application. |
| edb-pgadmin4-desktop-common | This package installs the desktop components of pgAdmin4 for all window managers. |
| edb-pgadmin4-desktop-gnome | This package installs the gnome desktop components of pgAdmin4 |
| edb-pgadmin4-docs | This package contains documentation of pgAdmin4. |

| Package Name | Package Installs |
|----------------------|---|
| edb-pgadmin4-web | This package contains the required files to run pgAdmin4 as a web application. |
| edb-efm37 | This package installs EDB Failover Manager that adds fault tolerance to database clusters to minimize downtime when a master database fails by keeping data online in high availability configurations. |
| edb-rs | This package is a java-based replication framework that provides asynchronous replication across Postgres and EPAS database servers. It supports master-slave, master-master, and hybrid configurations. |
| edb-rs-client | This package is a java-based command-line tool that is used to configure and operate a replication network via different commands by interacting with the EPRS server. |
| edb-rs-datavalidator | This package is a java-based command-line tool that provides row and column level data comparison of a source and target database table. The supported RDBMS servers include PostgreSQL, EPAS, Oracle, and MS SQL Server. |
| edb-rs-libs | This package contains certain libraries that are commonly used by ERPS Server, EPRS Client, and Monitoring modules. |
| edb-rs-monitor | This package is a java-based application that provides monitoring capabilities to ensure a smooth functioning of the EPRS replication cluster. |
| edb-rs-server | This package is a java-based replication framework that provides asynchronous replication across Postgres and EPAS database servers. It supports master-slave, master-master, and hybrid configurations. |
| edb-bart | This package installs the Backup and Recovery Tool (BART) to support online backup and recovery across local and remote PostgreSQL and EDB Advanced Servers. |
| libevent-edb | This package contains supporting library files. |
| libiconv-edb | This package contains supporting library files. |
| libevent-edb-devel | This package contains supporting library files. |

Updating an RPM Installation

If you have an existing Advanced Server RPM installation, you can use yum to upgrade your repository configuration file and update to a more recent product version. To update the `edb.repo` file, assume superuser privileges and enter:

```
yum upgrade edb-repo
```

yum will update the `edb.repo` file to enable access to the current EDB repository, configured to connect with the credentials specified in your `edb.repo` file. Then, you can use yum to upgrade all packages whose names include the expression `edb`:

```
yum upgrade edb*
```

Please note that the `yum upgrade` command will only perform an update between minor releases; to update between major releases, you must use `pg_upgrade`.

For more information about using yum commands and options, enter `yum --help` on your command line, or visit:

https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Deployment_Guide/ch-yum.html

Installing Advanced Server on a Debian or Ubuntu Host

To install Advanced Server on a Debian or Ubuntu host, you must have credentials that allow access to the EnterpriseDB repository. To request credentials for the repository, visit:

<https://www.enterprisedb.com/repository-access-request>

The following steps will walk you through using the EnterpriseDB apt repository to install a debian package. When using the commands, replace the `username` and `password` with the credentials provided by EnterpriseDB.

1. Assume superuser privileges:

```
sudo su -
```

2. Configure the EnterpriseDB repository:

```
sh -c 'echo "deb https://<username>:<password>@apt.enterprisedb.com/$(lsb_release -cs)-edb/ $(lsb_release -cs) main" > /etc/apt/sources.list.d/edb-$(lsb_release -cs).list'
```

3. Add support to your system for secure APT repositories:

```
apt-get install apt-transport-https
```

4. Add the EBD signing key:

```
wget -q -O - https://<username>:<password>@apt.enterprisedb.com/edb-deb.gpg.key | apt-key add -
```

5. Update the repository metadata:

```
apt-get update
```

6. Install Debian package:

```
apt-get install edb-as12
```

Note

Some Advanced Server components require a Java installation. Before using a native package to add Migration Toolkit to your system, please ensure that Java version 7 is installed on your Advanced Server host. Before using a native package to install EDB*Plus, please ensure that Java version 8 or later is installed.

If you are installing an Advanced Server supporting component that requires Java (such as MTK or EDB*Plus), make sure that you install Java version 8 before installing the supporting components.

The Debian package manager places Advanced Server and supporting components in the directories listed in the following table:

| Component | Location |
|----------------------------------|--|
| Server | /usr/lib/edb-as/12/ |
| Data and Configuration Directory | /var/lib/edb-as/12/main /etc/edb-as/12/main/ |
| pgAgent | /usr/lib/edb-as/12 |
| Pgpool | /usr/edb/pgpool4.0/ |
| Postgis | /usr/lib/edb-as/12/ |
| PGSNMPD | /usr/lib/edb-as/12 |
| Slony Replication | /usr/lib/edb-as/12 |
| pgBouncer | /usr/edb/pgbouncer1.12/ |
| pgBouncer Configuration Files | /etc/edb/pgbouncer1.12/edb-pgbouncer-1.12.ini /etc/edb/pgbouncer1.12/userlist.txt |

| Component | Location |
|--------------|----------------------------|
| SQL-Profiler | /usr/lib/edb-as/12/lib |
| SQL-Protect | /usr/lib/edb-as/12/lib |
| SSLUTILS | /usr/lib/edb-as/12/lib |
| PL-PERL | /usr/lib/edb-as/12/lib |
| PL-PYTHON | /usr/lib/edb-as/12/lib |
| PLTCL | /usr/lib/edb-as/12/lib |
| EFM | /usr/edb/efm-3.7 |
| JDBC | /usr/edb/jdbc |
| MTK | /usr/edb/migrationtoolkit/ |

Advanced Server Debian Packages

The table that follows lists some of the Debian packages that are available from EnterpriseDB. You can also use the `apt list` command to access a list of the packages that are currently available from your configured repository. Open a command line, assume superuser privileges, and enter:

```
| apt list edb*
```

Please note: The available package list is subject to change.

| Package Name | Package Installs |
|------------------------------|--|
| edb-as12-server | Installs core components of the Advanced Server database server. |
| edb-as12-server-client | Includes client programs and utilities that you can use to access and manage Advanced Server. |
| edb-as12-server-core | Includes the programs needed to create the core functionality behind the Advanced Server database. |
| edb-as12-server-dev | The <code>edb-as12-server-dev</code> package contains the header files and libraries needed to compile C or C++ applications that directly interact with an Advanced Server server and the ecpg or ecpgPlus C preprocessor. |
| edb-as12-server-doc | Installs the readme file. |
| edb-as12-server-edb-modules | Installs supporting modules for Advanced Server |
| edb-as12-server-indexadvisor | Installs Advanced Server's Index Advisor feature. The Index Advisor utility helps determine which columns you should index to improve performance in a given workload. |
| edb-as12-server-pldebugger | This package implements an API for debugging PL/pgSQL functions on Advanced Server. |
| edb-as12-server-plpython | Installs the PL/Python procedural language for Advanced Server. Please note that the <code>edb-as12-server-plpython</code> package is dependent on the platform-supplied version of Python. By default, the <code>edb-as12-server-plpython3</code> package is not installed; you need to explicitly install the package using apt. |
| edb-as12-server-pltcl | Installs the PL/Tcl procedural language for Advanced Server. Please note that the <code>edb-as12-pltcl</code> package is dependent on the platform-supplied version of TCL. |

| Package Name | Package Installs |
|---------------------------------|--|
| edb-as12-server-sqlprofiler | This package installs Advanced Server's SQL Profiler feature. SQL Profiler helps identify and optimize SQL code. |
| edb-as12-server-sqlprotect | This package installs Advanced Server's SQL Protect feature. SQL Protect provides protection against SQL injection attacks. |
| edb-as12-server-sslutils | This package installs functionality that provides SSL support. |
| edb-as12-server-cloneschema | This package installs the EDB Clone Schema extension. For more information about EDB Clone Schema, see the EDB Postgres Advanced Server Guide. |
| edb-as12-server-parallel-clone | This package installs functionality that supports the EDB Clone Schema extension. |
| edb-as12-edbplus | The <code>edb-edbplus</code> package contains the files required to install the EDB Plus command line client. EDB Plus commands are compatible with Oracle's SQL*Plus. |
| edb-as12-pgsnmpd | SNMP (Simple Network Management Protocol) is a protocol that allows you to supervise an apparatus connected to the network. |
| edb-as12-pgadmin4 | pgAdmin 4 provides a graphical management interface for Advanced Server and PostgreSQL databases. |
| edb-as12-pgadmin-apache | Apache support module for pgAdmin 4. |
| edb-as12-pgadmin4-common | pgAdmin 4 supporting files. |
| edb-as12-pgadmin4-doc | pgAdmin 4 documentation module. |
| edb-as12-pgpool40-extensions | This package creates pgPool extensions required by the server. |
| edb-as12-postgis-2.5 | This package installs POSTGIS support for geospatial data. |
| edb-as12-postgis-2.5-scripts | This package installs POSTGIS support for geospatial data. |
| edb-as12-postgis-doc-2.5 | This package provides support for POSTGIS. |
| edb-as12-postgis-gui-2.5 | This package provides support for POSTGIS. |
| edb-as12-postgis-jdbc | This package provides support for POSTGIS. |
| edb-as12-postgis-scripts | This package provides support for POSTGIS. |
| edb-as12-pgagent | This package installs pgAgent; pgAgent is a job scheduler for Advanced Server. Before installing this package, you must install EPEL; for detailed information about installing EPEL, see Installing Advanced Server on a Linux Host . |
| edb-as12-slony-replication | This package installs the meta RPM for Slony-I. |
| edb-as12-slony-replication-core | This package contains core portions of Slony-I to build a master-slave system that includes all features and capabilities needed to replicate large databases to a reasonably limited number of slave systems. |

| Package Name | Package Installs |
|--------------|------------------|
|--------------|------------------|

| | |
|----------------------------------|---|
| edb-as12-slony-replication-docs | This package contains the Slony project documentation (in pdf form). |
| edb-as12-slony-replication-tools | This package contains the Slony altperl tools and utilities that are useful when deploying Slony replication environments. Before installing this package, you must install EPEL; for detailed information about installing EPEL, see Installing Advanced Server on a Linux Host . |
| edb-as12-hdfs-fdw | The Hadoop Data Adapter allows you to query and join data from Hadoop environments with your Postgres or Advanced Server instances. It is YARN Ready certified with HortonWorks, and provides optimizations for performance with predicate pushdown support. |
| edb-as12-hdfs-fdw-doc | Documentation for the Hadoop Data Adapter. |
| edb-as12-mongo-fdw | This EnterpriseDB Advanced Server extension implements a Foreign Data Wrapper for MongoDB. |
| edb-as12-mongo-fdw-doc | Documentation for the Foreign Data Wrapper for MongoDB. |
| edb-as12-mysql-fdw | This EnterpriseDB Advanced Server extension implements a Foreign Data Wrapper for MySQL. |
| edb-pgpool40 | <p>This package contains the pgPool-II installer. The pgpool-II utility package acts as a middleman between client applications and Server database servers. pgpool-II functionality is transparent to client applications; client applications connect to pgpool-II instead of directly to Advanced Server, and pgpool-II manages the connection. EnterpriseDB supports the following pgpool-II features:</p> <ul style="list-style-type: none"> • Load balancing • Connection pooling • High availability • Connection limits <p>pgpool-II runs as a service on Linux systems, and is not supported on Windows systems.</p> |
| edb-jdbc | The edb-jdbc package includes the .jar files needed for Java programs to access an Advanced Server database. |
| edb-migrationtoolkit | The edb-migrationtoolkit package installs Migration Toolkit, facilitating migration to an Advanced Server database from Oracle, PostgreSQL, MySQL, Sybase and SQL Server. |
| edb-pgbouncer19 | This package contains PgBouncer (a lightweight connection pooler). This package requires the libevent package. |
| edb-efm37 | This package installs EDB Failover Manager that adds fault tolerance to database clusters to minimize downtime when a master database fails by keeping data online in high availability configurations. |

Configuring a Package Installation

The packages that install the database server component create a service configuration file (on version 6.x hosts) or unit file (on version 7.x hosts), and service startup scripts.

Creating a Database Cluster and Starting the Service

The PostgreSQL `initdb` command creates a database cluster; when installing Advanced Server with an RPM package, the `initdb` executable is in `/usr/edb/asx.x/bin`. After installing Advanced Server, you must manually configure the service and invoke `initdb` to create your cluster. When invoking `initdb`, you can:

- Specify environment options on the command line.
- Include the `service` command on RHEL or CentOS 6.x and use a service configuration file to configure the environment.
- Include the `systemd` service manager on RHEL or CentOS 7.x and use a service configuration file to configure the environment.

To review the `initdb` documentation, visit:

<https://www.postgresql.org/docs/12/static/app-initdb.html>

After specifying any options in the service configuration file, you can create the database cluster and start the service; these steps are platform specific.

On RHEL or CentOS 6.x

To create a database cluster in the `PGDATA` directory that listens on the port specified by the `PGPORT` environment variable specified in the service configuration file (described in [Using a Service Configuration File on CentOS or Redhat 6.x](#)), assume `root` privileges, and invoke the `service` script:

```
| service edb-as-12 initdb
```

You can also assign a locale to the cluster when invoking `initdb`. By default, `initdb` will use the value specified by the `$LANG` operating system variable, but if you append a preferred locale when invoking the script, the cluster will use the alternate value. For example, to create a database cluster that uses simplified Chinese, invoke the command:

```
| service edb-as-12 initdb zh_CH.UTF-8
```

After creating a database cluster, start the database server with the command:

```
| service edb-as-12 start
```

On RHEL or CentOS 7.x

To invoke `initdb` on a RHEL or CentOS 7.x system, with the options specified in the service configuration file, assume the identity of the operating system superuser:

```
| su - root
```

To initialize a cluster with the non-default values, you can use the `PGSETUP_INITDB_OPTIONS` environment variable by invoking the `edb-as-12-setup` cluster initialization script that resides under `EPAS_Home/bin`.

To invoke `initdb` export the `PGSETUP_INITDB_OPTIONS` environment variable with the following command:

```
| PGSETUP_INITDB_OPTIONS="-E UTF-8" /usr/edb/as12/bin/edb-as-12-setup initdb
```

After creating the cluster, use `systemctl` to start, stop, or restart the service:

```
| systemctl { start \| stop \| restart } edb-as-12
```

On Debian 9x or Ubuntu 18.04

You can initialize multiple clusters using the bundled scripts. To create a new cluster, assume `root` privileges, and invoke the bundled script:

```
| /usr/bin/epas_createcluster 12 main2
```

To start a new cluster, use the following command:

```
/usr/bin/epas_ctlcluster 12 main2 start
```

To list all the available clusters, use the following command:

```
/usr/bin/epas_lsclusters
```

Note

The data directory is created under `/var/lib/edb-as/12/main2` and configuration directory is created under `/etc/edb-as/12/main/`.

Using a Service Configuration File on CentOS or Redhat 6.x

On a CentOS or RedHat version 6.x host, the RPM installer creates a service configuration file named `edb-as-12.sysconfig` in `/etc/sysconfig/edb/as12` (see Figure 4.1). Please note that options specified in the service configuration file are only enforced if `initdb` is invoked via the service command; if you manually invoke `initdb` (at the command line), you must specify the other options (such as the location of the `data` directory and installation mode) on the command line.

```
centos@tusharcentos6:/etc/sysconfig/edb/as12
File Edit View Search Terminal Help
PGENGINE=/usr/edb/as12/bin
PGPORT=5444
PGDATA=/var/lib/edb/as12/data
PGLOG=/var/lib/edb/as12/pgstartup.log
#INITDBOPTS="--no-redwood-compat"
```

The file contains the following environment variables:

- `PGENGINE` specifies the location of the engine and utility executable files.
- `PGPORT` specifies the listener port for the database server.
- `PGDATA` specifies the path to the data directory.
- `PGLOG` specifies the location of the log file to which the server writes startup information.
- Use `INITDBOPTS` to specify any `initdb` option or options that you wish to apply to the new cluster. For more information, see [Specifying Cluster Options with INITDBOPTS](#).

You can modify the `edb-as-12.sysconfig` file before using the service command to invoke the `startup` script to change the listener port, data directory location, startup log location or installation mode. If you plan to create more than one instance on the same system, you may wish to copy the `edb-as-12.sysconfig` file (and the associated `edb-as-12` startup script) and modify the file contents for each additional instance that resides on the same host.

Specifying Cluster Options with INITDBOPTS

You can use the `INITDBOPTS` variable to specify your cluster configuration preferences. By default, the `INITDBOPTS` variable is commented out in the service configuration file; unless modified, when you run the service startup script, the new cluster will be created in a mode compatible with Oracle databases. Clusters created in this mode will contain a database named `edb`, and have a database superuser named `enterprisedb`.

Clusters created in PostgreSQL mode do not include compatibility features. To create a new cluster in PostgreSQL mode, remove the pound sign (#) in front of the `INITDBOPTS` variable, enabling the `--no-redwood-compat` option. Clusters created in PostgreSQL mode will contain a database named `postgres`, and have a database superuser named `postgres`.

If you initialize the database using Oracle compatibility mode, the installation includes:

- Data dictionary views compatible with Oracle databases.
- Oracle data type conversions.
- Date values displayed in a format compatible with Oracle syntax.
- Support for Oracle-styled concatenation rules (if you concatenate a string value with a `NULL` value, the returned value is the value of the string).
- Support for the following Oracle built-in packages:

| Package | Functionality compatible with Oracle Databases |
|-----------------------------|--|
| <code>dbms_alert</code> | Provides the capability to register for, send, and receive alerts. |
| <code>dbms_job</code> | Provides the capability for the creation, scheduling, and managing of jobs. |
| <code>dbms_lob</code> | Provides the capability to manage on large objects. |
| <code>dbms_output</code> | Provides the capability to send messages to a message buffer, or get messages from the message buffer. |
| <code>dbms_pipe</code> | Provides the capability to send messages through a pipe within or between sessions connected to the same database cluster. |
| <code>dbms_rls</code> | Enables the implementation of Virtual Private Database on certain Advanced Server database objects. |
| <code>dbms_sql</code> | Provides an application interface to the EnterpriseDB dynamic SQL functionality. |
| <code>dbms_utility</code> | Provides various utility programs. |
| <code>dbms_aqadm</code> | Provides supporting procedures for Advanced Queueing functionality. |
| <code>dbms_aq</code> | Provides message queueing and processing for Advanced Server. |
| <code>dbms_profiler</code> | Collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a performance profiling session. |
| <code>dbms_random</code> | Provides a number of methods to generate random values. |
| <code>dbms_redact</code> | Enables the redacting or masking of data that is returned by a query. |
| <code>dbms_lock</code> | Provides support for the DBMS_LOCK.SLEEP procedure. |
| <code>dbms_scheduler</code> | Provides a way to create and manage jobs, programs, and job schedules. |
| <code>dbms_crypto</code> | Provides functions and procedures to encrypt or decrypt RAW, BLOB or CLOB data. You can also use DBMS_CRYPTO functions to generate cryptographically strong random values. |
| <code>dbms_mview</code> | Provides a way to manage and refresh materialized views and their dependencies. |
| <code>dbms_session</code> | Provides support for the DBMS_SESSION.SET_ROLE procedure. |
| <code>utl_encode</code> | Provides a way to encode and decode data. |
| <code>utl_http</code> | Provides a way to use the HTTP or HTTPS protocol to retrieve information found at an URL. |
| <code>utl_file</code> | Provides the capability to read from, and write to files on the operating system's file system. |
| <code>utl_smtp</code> | Provides the capability to send e-mails over the Simple Mail Transfer Protocol (SMTP). |
| <code>utl_mail</code> | Provides the capability to manage e-mail. |
| <code>utl_url</code> | Provides a way to escape illegal and reserved characters within an URL. |
| <code>utl_raw</code> | Provides a way to manipulate or retrieve the length of raw data types. |

You may also specify multiple `initdb` options. For example, the following statement:

```
| INITDBOPTS="--no-redwood-compat -U alice --locale=en_US.UTF-8"
```

Creates a database cluster (without compatibility features for Oracle) that contains a database named `postgres` that is owned by a user named `alice`; the cluster uses `UTF-8` encoding.

In addition to the cluster configuration options documented in the PostgreSQL core documentation, Advanced Server supports the following `initdb` options:

```
| --no-redwood-compat
```

Include the `--no-redwood-compat` keywords to instruct the server to create the cluster in PostgreSQL mode.

When the cluster is created in PostgreSQL mode, the name of the database superuser will be `postgres`, the name of the default database will be `postgres`, and Advanced Server's features compatible with Oracle databases will not be available to the cluster.

--redwood-like

Include the `--redwood-like` keywords to instruct the server to use an escape character (an empty string ("")) following the `LIKE` (or PostgreSQL-compatible `ILIKE`) operator in a SQL statement that is compatible with Oracle syntax.

--icu-short-form

Include the `--icu-short-form` keywords to create a cluster that uses a default ICU (International Components for Unicode) collation for all databases in the cluster. For more information about Unicode collations, please refer to the *EDB Postgres Advanced Server Guide* available at:

<https://www.enterprisedb.com/edb-docs>

For more information about using `initdb`, and the available cluster configuration options, see the PostgreSQL Core Documentation available at:

<https://www.postgresql.org/docs/12/static/app-initdb.html>

You can also view online help for `initdb` by assuming superuser privileges and entering:

`/<path_to_initdb_installation_directory>/initdb --help`

Where `path_to_initdb_installation_directory` specifies the location of the `initdb` binary file.

Modifying the Data Directory Location on CentOS or Redhat 7.x

On a CentOS or RedHat version 7.x host, the unit file is named `edb-as-12.service` and resides in `/usr/lib/systemd/system`. The unit file contains references to the location of the Advanced Server `data` directory. You should avoid making any modifications directly to the unit file because it may be overwritten during package upgrades.

By default, data files reside under `/var/lib/edb/as13/data` directory. To use a data directory that resides in a non-default location, perform the following steps:

1. create a copy of the unit file under the `/etc` directory:

```
cp /usr/lib/systemd/system/edb-as-13.service /etc/systemd/system/
```

2. After copying the unit file to the new location, create the service file (`/etc/systemd/system/edb-as-13.service`).
3. Update the following values with new location of data directory in the `/lib/systemd/system/edb-as-13.service` file:

```
Environment=PGDATA=/var/lib/edb/as13/data
```

```
PIDFile=/var/lib/edb/as13/data/postmaster.pid
```

1. Delete the entire content of `/etc/systemd/system/edb-as-13.service` file, except the following line:

```
.include /lib/systemd/system/edb-as-13.service
```

1. Run the following command to initialize the cluster at the new location:

```
PGSETUP_INITDB_OPTIONS="-E UTF-8" /usr/edb/as13/bin/edb-as-13-setup initdb
```

1. Use the following command to reload `systemd`, updating the modified service scripts:

```
systemctl daemon-reload
```

1. Start the Advanced Server service with the following command:

```
systemctl start edb-as-13
```

For information about setting an environment variable, see the [Database Compatibility for Oracle Developers Reference Guide](#).

Starting Multiple Postmasters with Different Clusters

You can configure Advanced Server to use multiple postmasters, each with its own database cluster. The steps required are version specific to the Linux host.

On RHEL or CentOS 6.x

The `edb-as12-server-core` RPM contains a script that starts the Advanced Server instance. The script can be copied, allowing you to run multiple services, with unique `data` directories and that monitor different ports. You must have `root` access to invoke or modify the script.

The example that follows creates a second instance on an Advanced Server host; the secondary instance is named `secondary`:

1. Create a hard link in `/etc/rc.d/init.d` (or equivalent location) to the `edb-as-12` service (named `secondary-edb-as-12`):

```
In edb-as-12 secondary-edb-as-12
```

Be sure to pick a name that is not already used in `/etc/rc.d/init.d`.

1. Create a file in `/etc/sysconfig/edb/as12/` named `secondary-edb-as-12`. This file is where you would typically define `PGDATA` and `PGOPTS`. Since `$PGDATA/postgresql.conf` will override many of these settings (except `PGDATA`) you might notice unexpected results on startup.
2. Create the target `PGDATA` directory.
3. Assume the identity of the Advanced Server database superuser (`enterprisedb`) and invoke `initdb` on the target `PGDATA`. For information about using `initdb`, please see the `PostgreSQL` Core Documentation available at:
`https://www.postgresql.org/docs/12/static/app-initdb.html`.
4. Edit the `postgresql.conf` file to specify the port, address, TCP/IP settings, etc. for the `secondary` instance.
5. Start the postmaster with the following command:

```
service secondary-edb-as-12 start
```

On RHEL or CentOS 7.x

The `edb-as12-server-core` RPM for version 7.x contains a unit file that starts the Advanced Server instance. The file allows you to start multiple services, with unique `data` directories and that monitor different ports. You must have `root` access to invoke or modify the script.

The example that follows creates an Advanced Server installation with two instances; the secondary instance is named `secondary`:

1. Make a copy of the default file with the new name. As noted at the top of the file, all modifications must reside

under `/etc`. You must pick a name that is not already used `in /etc/systemd/system`.

```
cp /usr/lib/systemd/system/edb-as-12.service /etc/systemd/system/secondary-edb-as-12.service
```

2. Edit the file, changing `PGDATA` to point to the new `data` directory that you will create the cluster against.
3. Create the target `PGDATA` with user `enterprisedb`.
4. Run `initdb`, specifying the setup script:

```
/usr/edb/as12/bin/edb-as-12-setup initdb secondary-edb-as-12
```

1. Edit the `postgresql.conf` file for the new instance, specifying the port, the IP address, TCP/IP settings, etc.
2. Make sure that new cluster runs after a reboot:

```
systemctl enable secondary-edb-as-12
```

3. Start the second cluster with the following command:

```
systemctl start secondary-edb-as-12
```

Creating an Advanced Server Repository on an Isolated Network

You can create a local repository to act as a host for the Advanced Server RPM packages if the server on which you wish to install Advanced Server (or supporting components) cannot directly access the EnterpriseDB repository. Please note that this is a high-level listing of the steps required; you will need to modify the process for your individual network.

To create and use a local repository, you must:

1. Use yum to install the `epel-release`, `yum-utils`, and `createrepo` packages:

```
yum install epel-release
yum install yum-utils
yum install createrepo
```

1. Create a directory in which to store the repository:

```
mkdir /srv/repos
```

2. Copy the RPM installation packages to your local repository. You can download the individual packages or use a tarball to populate the repository. The packages are available from the EnterpriseDB repository at:

```
https://yum.enterprisedb.com/
```

3. Sync the RPM packages and create the repository.

```
reposync -r edbas12 -p /srv/repos createrepo /srv/repos
```

4. Install your preferred webserver on the host that will act as your local repository, and ensure that the repository directory is accessible to the other servers on your network.
5. On each isolated database server, configure yum to pull updates from the mirrored repository on your local network. For example, you might create a repository configuration file called `/etc/yum.repos.d/edb-repo` with connection information that specifies:

```
[edbas12]
name=EnterpriseDB Advanced Server 12
```

```
baseurl=https: yum.your_domain.com/edbas12
enabled=1
gpgcheck=0
```

After specifying the location and connection information for your local repository, you can use yum commands to install Advanced Server and its supporting components on the isolated servers. For example:

`yum install edb-as12-server`

For more information about creating a local yum repository, visit:

<https://wiki.centos.org/HowTos/CreateLocalRepos>

1.3 Installation Troubleshooting

Difficulty Displaying Java-based Applications

If you encounter difficulty displaying Java-based server features (controls or text not being displayed correctly, or blank windows), upgrading to the latest `libxcb-xlib` libraries should correct the problem on most distributions. Please visit the following link for other possible work-arounds:

http://bugs.sun.com/bugdatabase/view_bug.do?bug_id=6532373

The Installation Fails to Complete Due to Existing data Directory Contents

If an installation fails to complete due to an existing content in the data directory, the server will write an error message to the server logs:

`A data directory is neither empty, or a recognisable data directory.`

If you encounter a similar message, you should confirm that the data directory is empty; the presence of files (including the system-generated `lost+found` folder) will prevent the installation from completing. Either remove the files from the data directory, or specify a different location for the data directory before re-invoking the installer to complete the installation.

Difficulty Installing the EPEL Release Package

If you encounter difficulty when installing the `EPEL` release package, you can use the following command to install the `epel-release` package:

```
yum -y install https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm
```

Please note that you may need to enable the `[extras]` repository definition in the `CentOS-Base.repo` file (located in `/etc/yum.repos.d`). If `yum` cannot access a repository that contains `epel-release`, you will get an error message:

`No package epel available. Error: Nothing to do`

If you receive this error, you can download the `EPEL` rpm package, and install it manually; download the rpm package, assume superuser privileges, navigate into the directory that contains the package, and install `EPEL` with the command:

`yum install epel-release`

1.4 Managing an Advanced Server Installation

Unless otherwise noted, the commands and paths noted in the following section assume that you have performed an installation using the native packages.

Starting and Stopping Advanced Server and Supporting Components

A service is a program that runs in the background and requires no user interaction (in fact, a service provides no user interface); a service can be configured to start at boot time, or manually on demand. Services are best controlled using the platform-specific operating system service control utility. Many of the Advanced Server supporting components are services.

The following table lists the names of the services that control Advanced Server and services that control Advanced Server supporting components:

| Advanced Server Component Name | Linux Service Name | Debian Service Name |
|--------------------------------|--------------------------|----------------------------|
| Advanced Server | edb-as-12 | edb-as12-main |
| pgAgent | edb-pgagent-12 | edb-as12-pgagent |
| PgBouncer | edb-pgbouncer-112 | edb-pgbouncer112 |
| pgPool-II | edb-pgpool-37 | edb-pgpool37 |
| Slony | edb-slony-replication-12 | edb-as12-slony-replication |
| EFM | efm-3.7 | efm-3.7 |

You can use the Linux command line to control Advanced Server's database server and the services of Advanced Server's supporting components. The commands that control the Advanced Server service on a Linux platform are host specific.

Controlling a Service on CentOS or RHEL 7.x

If your installation of Advanced Server resides on version 7.x of RHEL and CentOS, you must use the `systemctl` command to control the Advanced Server service and supporting components.

The `systemctl` command must be in your search path and must be invoked with superuser privileges. To use the command, open a command line, and enter:

```
systemctl action service_name
```

Where:

`service_name` specifies the name of the service.

`action` specifies the action taken by the service command. Specify:

- `start` to start the service.
- `stop` to stop the service.
- `restart` to stop and then start the service.
- `status` to discover the current status of the service.

Controlling a Service on CentOS or RHEL 6.x

On version 6.x of RHEL or CentOS Linux, you can control a service at the command line with the `service` command. The `service` command can be used to manage an Advanced Server cluster, as well as the services of component software installed with Advanced Server.

Using the `service` command to change the status of a service allows the Linux service controller to keep track of the server status (the `pg_ctl` command does not alert the service controller to changes in the status of a server). The command must be in your search path and must be invoked with superuser privileges. Open a command line, and issue the command:

```
| service service_name action
```

The Linux `service` command invokes a script (with the same name as the service) that resides in `/etc/init.d`. If your Linux distribution does not support the `service` command, you can call the script directly by entering:

```
| /etc/init.d/service_name action
```

Where:

`service_name` specifies the name of the service.

`action` specifies the action taken by the service command. Specify:

- `start` to start the service.
- `stop` to stop the service.
- `condstop` to stop the service without displaying a notice if the server is already stopped.
- `restart` to stop and then start the service.
- `condrestart` to restart the service without displaying a notice if the server is already stopped.
- `try-restart` to restart the service without displaying a notice if the server is already stopped.
- `status` to discover the current status of the service.

Controlling a Service on Debian 9x or Ubuntu 18.04

If your installation of Advanced Server resides on version 9x of Debian or 18.04 of Ubuntu, assume superuser privileges and invoke the following commands (using bundled scripts) to manage the service. Use the following commands to:

- Discover the current status of a service:

```
| /usr/bin/epas_ctlcluster 12 main status
```

- Stop a service:

```
| /usr/bin/epas_ctlcluster 12 main stop
```

- Restart a service:

```
| /usr/bin/epas_ctlcluster 12 main restart
```

- Reload a service:

```
| /usr/bin/epas_ctlcluster 12 main reload
```

- Control the component services:

```
| systemctl restart edb-as@12-main
```

Using pg_ctl to Control Advanced Server

You can use the `pg_ctl` utility to control an Advanced Server service from the command line on any platform. `pg_ctl` allows you to start, stop, or restart the Advanced Server database server, reload the configuration parameters, or display the status of a running server. To invoke the utility, assume the identity of the cluster owner, navigate into the home directory of Advanced Server, and issue the command:

```
./bin/pg_ctl -D <data_directory action>
```

`data_directory` is the location of the data controlled by the Advanced Server cluster.

`action` specifies the action taken by the `pg_ctl` utility. Specify:

- `start` to start the service.
- `stop` to stop the service.
- `restart` to stop and then start the service.
- `reload` sends the server a `SIGHUP` signal, reloading configuration parameters
- `status` to discover the current status of the service.

For more information about using the `pg_ctl` utility, or the command line options available, please see the official PostgreSQL Core Documentation available at:

```
https://www.postgresql.org/docs/12/static/app-pg-ctl.html
```

Choosing Between `pg_ctl` and the Service Command

You can use the `pg_ctl` utility to manage the status of an Advanced Server cluster, but it is important to note that `pg_ctl` does not alert the operating system service controller to changes in the status of a server, so it is beneficial to use the service command whenever possible.

Configuring Component Services to AutoStart at System Reboot

After installing, configuring, and starting the services of Advanced Server supporting components on a Linux system, you must manually configure your system to autostart the service when your system reboots. To configure a service to autostart on a Linux system, open a command line, assume superuser privileges, and enter the following command.

On a Redhat-compatible Linux system, enter:

```
/sbin/chkconfig *service_name* on
```

Where `service_name` specifies the name of the service.

Modifying the `postgresql.conf` File

Configuration parameters in the `postgresql.conf` file specify server behavior with regards to auditing, authentication, encryption, and other behaviors. On a RHEL or CentOS host, the `postgresql.conf` file resides in the `data` directory under your Advanced Server installation. On a Debian or Ubuntu host, server configuration files are located in the `/etc/edb-as/12/main` directory.

```

# PostgreSQL Client Authentication Configuration File
# =====
#
# Refer to the "Client Authentication" section in the PostgreSQL
# documentation for a complete description of this file. A short
# synopsis follows.
#
# This file controls: which hosts are allowed to connect, how clients
# are authenticated, which PostgreSQL user names they can use, which
# databases they can access. Records take one of these forms:
#
# local      DATABASE  USER  METHOD  [OPTIONS]
# host       DATABASE  USER  ADDRESS  METHOD  [OPTIONS]
# hostssl    DATABASE  USER  ADDRESS  METHOD  [OPTIONS]
# hostnossal DATABASE  USER  ADDRESS  METHOD  [OPTIONS]
#
# (The uppercase items must be replaced by actual values.)
#
# The first field is the connection type: "local" is a Unix-domain
# socket, "host" is either a plain or SSL-encrypted TCP/IP socket,
# "hostssl" is an SSL-encrypted TCP/IP socket, and "hostnossal" is a
# plain TCP/IP socket.
#
# DATABASE can be "all", "sameuser", "samerole", "replication", a
# database name, or a comma-separated list thereof. The "all"
# keyword does not match "replication". Access to replication
# must be enabled in a separate record (see example below).
#
# USER can be "all", a user name, a group name prefixed with "+", or a
# comma-separated list thereof. In both the DATABASE and USER fields
# you can also write a file name prefixed with "@" to include names
# from a separate file.
#
# ADDRESS specifies the set of hosts the record matches. It can be a
# host name, or it is made up of an IP address and a CIDR mask that is
# an integer (between 0 and 32 (IPv4) or 128 (IPv6) inclusive) that
# specifies the number of significant bits in the mask. A host name
# that starts with a dot (.) matches a suffix of the actual host name.
# Alternatively, you can write an IP address and netmask in separate
# columns to specify the set of hosts. Instead of a CIDR-address, you
# can write "samehost" to match any of the server's own IP addresses,
# or "samenet" to match any address in any subnet that the server is
# directly connected to.
#
# METHOD can be "trust", "reject", "md5", "password", "scram-sha-256",
# "gss", "sspi", "ident", "peer", "pam", "ldap", "radius" or "cert".
# Note that "password" sends passwords in clear text; "md5" or
# "scram-sha-256" are preferred since they send encrypted passwords.
#
# OPTIONS are a set of options for the authentication in the format

```

Parameters that are preceded by a pound sign (#) are set to their default value (as shown in the parameter setting). To change a parameter value, remove the pound sign and enter a new value. After setting or changing a parameter, you must either **reload** or **restart** the server for the new parameter value to take effect.

Within the **postgresql.conf** file, some parameters contain comments that indicate **change requires restart**. To view

a list of the parameters that require a server restart, execute the following query at the psql command line:

```
| SELECT name FROM pg_settings WHERE context = 'postmaster';
```

Modifying the pg_hba.conf File

Appropriate authentication methods provide protection and security. Entries in the `pg_hba.conf` file specify the authentication method or methods that the server will use when authenticating connecting clients. Before connecting to the server, you may be required to modify the authentication properties specified in the `pg_hba.conf` file.

When you invoke the `initdb` utility to create a cluster, `initdb` creates a `pg_hba.conf` file for that cluster that specifies the type of authentication required from connecting clients.

The default authentication configuration specified in the `pg_hba.conf` file is:

| # | TYPE | DATABASE | USER | ADDRESS | METHOD |
|--|------|----------|------|--------------|--------|
| # "local" is for Unix domain socket connections only | | | | | |
| local all all | | | | | peer |
| # IPv4 local connections: | | | | | |
| host all all | | | | 127.0.0.1/32 | ident |
| # IPv6 local connections: | | | | | |
| host all all | | | | ::1/128 | ident |
| # Allow replication connections from localhost, by a user with the | | | | | |
| # replication privilege. | | | | | |
| local replication all | | | | | peer |
| host replication all | | | | 127.0.0.1/32 | ident |
| host replication all | | | | ::1/128 | ident |

To modify the `pg_hba.conf` file, open the file with your choice of editor. After modifying the authentication settings in the `pg_hba.conf` file, use the Linux command line to restart the server and apply the changes.

For more information about authentication, and modifying the `pg_hba.conf` file, see the PostgreSQL Core Documentation at:

<https://www.postgresql.org/docs/12/static/auth-pg-hba-conf.html>

Managing Authentication on a Debian or Ubuntu Host

By default, the server is running with the peer or md5 permission on a Debian or Ubuntu host. You can change the authentication method by modifying the `pg_hba.conf` file, located in:

```
| /etc/edb-as/12/main/pg_hba.conf
```

For more information about modifying the `pg_hba.conf` file, please review the PostgreSQL core documentation.

Connecting to Advanced Server with psql

`psql` is a command line client application that allows you to execute SQL commands and view the results. To open the `psql` client, the client must be in your search path. The executable resides in the `bin` directory, under your Advanced Server installation.

Use the following command and options to start the `psql` client:

```
| psql -d edb -U enterpriseedb
```

Where:

`-d` specifies the database to which `psql` will connect;

`-U` specifies the identity of the database user that will be used for the session.

For more information about using the command line client, please refer to the PostgreSQL Core Documentation at:

<https://www.postgresql.org/docs/12/static/app-psql.html>

1.5 Uninstalling Advanced Server

Note that after uninstalling Advanced Server, the cluster data files remain intact and the service user persists. You may manually remove the cluster `data` and service user from the system.

Uninstalling an RPM Package

You can use variations of the rpm or yum command to remove installed packages. Note that removing a package does not damage the Advanced Server `data` directory.

Include the `-e` option when invoking the rpm command to remove an installed package; the command syntax is:

```
| rpm -e <package_name>
```

Where `package_name` is the name of the package that you would like to remove.

You can use the `yum remove` command to remove a package installed by yum. To remove a package, open a terminal window, assume superuser privileges, and enter the command:

```
| yum remove <package_name>
```

Where `package_name` is the name of the package that you would like to remove.

Note

yum and RPM will not remove a package that is required by another package. If you attempt to remove a package that satisfies a package dependency, yum or RPM will provide a warning.

To uninstall Advanced Server and its dependent packages; use the following command:

```
| yum remove <edb-as12-server*>
```

Uninstalling Advanced Server Components on a Debian or Ubuntu Host

1. To uninstall Advanced Server, invoke the following command:

```
| apt-get remove <edb-as12-server*>
```

Please note: The configuration files and data directory remains intact.

2. To uninstall Advanced Server, configuration files, and data directory, invoke the following command:

```
apt-get purge <edb-as12-server*>
```

2 Advanced Server Installation Guide for Windows

The EDB Postgres Advanced Server Installation Guide is a comprehensive guide to installing EDB Postgres Advanced Server (Advanced Server). In this guide you will find detailed information about:

- Software prerequisites for Advanced Server 12 installation on Windows.
- Graphical installation options available through the interactive setup wizard on Windows.
- Managing an Advanced Server installation.
- Configuring an Advanced Server package installation.
- Uninstalling Advanced Server and its components.

2.1 Requirements Overview

For information about the platforms and versions supported by Advanced Server, visit the EnterpriseDB website at:

<https://www.enterprisedb.com/blog/edb-supported-products-and-platforms>

Limitations

The following limitations apply to EDB Postgres Advanced Server:

- The `data` directory of a production database should not be stored on an NFS file system.

Windows Installation Prerequisites

User Privileges

To perform an Advanced Server installation on a Windows system, you must have administrator privileges. If you are installing Advanced Server on a Windows system that is configured with `User Account Control` enabled, you can assume sufficient privileges to invoke the graphical installer by right clicking on the name of the installer and selecting `Run as administrator` from the context menu.

Windows-specific Software Requirements

You should apply Windows operating system updates before invoking the Advanced Server installer. If (during the installation process) the installer encounters errors, exit the installation, and ensure that your version of Windows is up-to-date before restarting the installer.

Migration Toolkit or EDB*Plus Installation Pre-requisites

Before using a StackBuilder Plus to install Migration Toolkit or EDB*Plus, you must first install Java (version 1.8 or later). If you are using Windows, Java installers and instructions are available online at:

<http://www.java.com/en/download/manual.jsp>

2.2 Installing Advanced Server with the Interactive Installer

You can use the Advanced Server interactive installer to install Advanced Server on Windows. The interactive installer is available from the EnterpriseDB website at:

<https://www.enterprisedb.com/advanced-downloads>

You can invoke the graphical installer in different installation modes to perform an Advanced Server installation:

- For information about using the graphical installer, see [Performing a Graphical Installation on Windows](#).
- For information about performing an unattended installation, see [Performing an Unattended Installation](#).
- For information about performing an installation with limited privileges, see [Performing an Installation with Limited Privileges](#).
- For information about the command line options you can include when invoking the installer, see [Reference - Command Line Options](#).

During the installation, the graphical installer copies a number of temporary files to the location specified by the `TEMP` environment variable. You can optionally specify an alternate location for the temporary files by modifying the `TEMP` environment variable.

If invoking the installer from the command line, you can set the value of the variable on the command line. Use the command:

`SET TEMP=temp_file_location`

Where `temp_file_location` specifies the alternate location for the temporary files and must match the permissions with the `TEMP` environment variable.

Note

If you are invoking the installer to perform a system upgrade, the installer will preserve the configuration options specified during the previous installation.

Setting Cluster Preferences during a Graphical Installation

During an installation, the graphical installer invokes the PostgreSQL `initdb` utility to initialize a cluster. If you are using the graphical installer, you can use the `INITDBOPTS` environment variable to specify your `initdb` preferences. Before invoking the graphical installer, set the value of `INITDBOPTS` at the command line, specifying one or more cluster options. For example:

`SET INITDBOPTS= -k -E=UTF-8`

If you specify values in `INITDBOPTS` that are also provided by the installer (such as the `-D` option, which specifies the installation directory), the value specified in the graphical installer will supersede the value if specified in `INITDBOPTS`.

For more information about using `initdb` cluster configuration options, see the PostgreSQL Core Documentation available at:

<https://www.postgresql.org/docs/12/static/app-initdb.html>

In addition to the cluster configuration options documented in the PostgreSQL core documentation, Advanced Server supports the following `initdb` options:

--no-redwood-compat

Include the `--no-redwood-compat` keywords to instruct the server to create the cluster in PostgreSQL mode. When the cluster is created in PostgreSQL mode, the name of the database superuser will be `postgres`, the name of the default database will be `postgres`, and Advanced Server's features compatible with Oracle databases will not be available to the cluster.

--redwood-like

Include the `--redwood-like` keywords to instruct the server to use an escape character (an empty string (")) following the `LIKE` (or PostgreSQL compatible `ILIKE`) operator in a SQL statement that is compatible with Oracle syntax.

--icu-short-form

Include the `--icu-short-form` keywords to create a cluster that uses a default ICU (International Components for Unicode) collation for all databases in the cluster. For more information about Unicode collations, please refer to the *EDB Postgres Advanced Server Guide* available at:

<https://www.enterprisedb.com/edb-docs>

2.2.1 Performing a Graphical Installation on Windows

A graphical installation is a quick and easy way to install Advanced Server 12 on a Windows system. Use the wizard's dialogs to specify information about your system and system usage; when you have completed the dialogs, the installer performs an installation based on the selections made during the setup process.

To invoke the wizard, you must have administrator privileges. Assume administrator privileges, and double-click the `edb-as12-server-12.x.x-x-windows-x64` executable file.

Note

To install Advanced Server on some versions of Windows, you may be required to right click on the file icon and select `Run as Administrator` from the context menu to invoke the installer with `Administrator` privileges.

When the `Language Selection` popup opens, select an installation language and click `OK` to continue to the `Setup` window.




Click `Next` to continue.

The EnterpriseDB License Agreement opens.



Carefully review the license agreement before highlighting the appropriate radio button; click **Next** to continue.

The Installation Directory window opens.




By default, the Advanced Server installation directory is:

C:\Program Files\edb\as12

You can accept the default installation location, and click **Next** to continue, or optionally click the **File Browser** icon to open the **Browse For Folder** dialog to choose an alternate installation directory.

Note

The **data** directory of a production database should not be stored on an NFS file system.



The **Select Components** window contains a list of optional components that you can install with the Advanced Server **Setup** wizard. You can omit a module from the Advanced Server installation by deselecting the box next to

the components name.

The **Setup** wizard can install the following components while installing Advanced Server 12:

EDB Postgres Advanced Server

Select the **EDB Postgres Advanced Server** option to install Advanced Server 12.

pgAdmin 4

Select the **pgAdmin 4** option to install the pgAdmin 4 client. pgAdmin 4 provides a powerful graphical interface for database management and monitoring.

StackBuilder Plus

The **StackBuilder Plus** utility is a graphical tool that can update installed products, or download and add supporting modules (and the resulting dependencies) after your Advanced Server setup and installation completes. See [Using StackBuilder Plus](#) for more information about StackBuilder Plus.

Command Line Tools

The **Command Line Tools** option installs command line tools and supporting client libraries including:

- libpq
- psql
- EDB*Loader
- ecpgPlus
- pg_basebackup, pg_dump, and pg_restore
- pg_bench
- and more.

Note

The **Command Line Tools** are required if you are installing Advanced Server or pgAdmin 4.

After selecting the components you wish to install, click **Next** to open the **Additional Directories** window.



By default, the Advanced Server **data** files are saved to:


C:\Program Files\edb\as12\data

The default location of the Advanced Server **Write-Ahead Log (WAL) Directory** is:

C:\Program Files\edb\as12\data\pg_wal

Advanced Server uses write-ahead logs to promote transaction safety and speed transaction processing; when you make a change to a table, the change is stored in shared memory and a record of the change is written to the write-ahead log. When you perform a **COMMIT**, Advanced Server writes contents of the write-ahead log to disk.

Accept the default file locations, or use the **File Browser** icon to select an alternate location; click **Next** to continue to the **Advanced Server Dialect** window.



Use the drop-down listbox on the **Advanced Server Dialect** window to choose a server dialect. The server dialect specifies the compatibility features supported by Advanced Server.

By default, Advanced Server installs in **Compatible with Oracle** mode; you can choose between **Compatible with Oracle** and **Compatible with PostgreSQL** installation modes.

Compatible with Oracle

If you select **Compatible with Oracle**, the installation will include the following features:

- Data dictionary views that is compatible with Oracle databases.
- Oracle data type conversions.
- Date values displayed in a format compatible with Oracle syntax.
- Support for Oracle-styled concatenation rules (if you concatenate a string value with a **NULL** value, the returned value is the value of the string).
- Schemas (**dbo** and **sys**) compatible with Oracle databases added to the **SEARCH_PATH**.
- Support for the following Oracle built-in packages:

| Package | Functionality compatible with Oracle Databases |
|-----------------------|---|
| dbms_alert | Provides the capability to register for, send, and receive alerts. |
| dbms_job | Provides the capability for the creation, scheduling, and managing of jobs. |
| dbms_lob | Provides the capability to manage on large objects. |
| dbms_output | Provides the capability to send messages to a message buffer, or get messages from the message buffer. |
| dbms_pipe | Provides the capability to send messages through a pipe within or between sessions connected to the same database cluster. |
| dbms_rls | Enables the implementation of Virtual Private Database on certain Advanced Server database objects. |
| dbms_sql | Provides an application interface to the EnterpriseDB dynamic SQL functionality. |
| dbms_utility | Provides various utility programs. |
| dbms_aqadm | Provides supporting procedures for Advanced Queueing functionality. |
| dbms_aq | Provides message queueing and processing for Advanced Server. |
| dbms_profiler | Collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a performance profiling session. |
| dbms_random | Provides a number of methods to generate random values. |
| dbms_redact | Enables the redacting or masking of data that is returned by a query. |
| dbms_lock | Provides support for the DBMS_LOCK.SLEEP procedure. |
| dbms_scheduler | Provides a way to create and manage jobs, programs, and job schedules. |
| dbms_crypto | Provides functions and procedures to encrypt or decrypt RAW, BLOB or CLOB data. You can also use DBMS_CRYPTO functions to generate cryptographically strong random values. |

| Package | Functionality compatible with Oracle Databases |
|--------------|---|
| dbms_mview | Provides a way to manage and refresh materialized views and their dependencies. |
| dbms_session | Provides support for the <code>DBMS_SESSION.SET_ROLE</code> procedure. |
| utl_encode | Provides a way to encode and decode data. |
| utl_http | Provides a way to use the HTTP or HTTPS protocol to retrieve information found at an URL. |
| utl_file | Provides the capability to read from, and write to files on the operating system's file system. |
| utl_smtp | Provides the capability to send e-mails over the Simple Mail Transfer Protocol (SMTP). |
| utl_mail | Provides the capability to manage e-mail. |
| utl_url | Provides a way to escape illegal and reserved characters within an URL. |
| utl_raw | Provides a way to manipulate or retrieve the length of raw data types. |

This is not a comprehensive list of the compatibility features for Oracle included when Advanced Server is installed in `Compatible with Oracle` mode; for more information, see the *Database Compatibility for Oracle Developer's Guide* available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

If you choose to install in `Compatible with Oracle` mode, the Advanced Server superuser name is `enterprisedb`.


Compatible with PostgreSQL

If you select `Compatible with PostgreSQL`, Advanced Server will exhibit compatibility with PostgreSQL version 12. If you choose to install in `Compatible with PostgreSQL` mode, the default Advanced Server superuser name is `postgres`.

For detailed information about PostgreSQL functionality, visit the official PostgreSQL website at:

<http://www.postgresql.org>


After specifying a configuration mode, click `Next` to continue to the `Password` window.



Advanced Server uses the password specified on the `Password` window for the database superuser. The specified password must conform to any security policies existing on the Advanced Server host.

After you enter a password in the `Password` field, confirm the password in the `Retype Password` field, and click `Next` to continue.

The `Additional Configuration` window opens.




Use the fields on the **Additional Configuration** window to specify installation details:

- Use the **Port** field to specify the port number that Advanced Server should listen to for connection requests from client applications. The default is **5444**.
- If the **Locale** field is set to **[Default locale]**, Advanced Server uses the system locale as the working locale. Use the drop-down listbox next to **Locale** to specify an alternate locale for Advanced Server.
- By default, the **Setup** wizard installs corresponding sample data for the server dialect specified by the compatibility mode (**Oracle** or **PostgreSQL**). Clear the check box next to **Install sample tables and procedures** if you do not wish to have sample data installed.

After verifying the information on the **Additional Configuration** window, click **Next** to open the **Dynatune Dynamic Tuning: Server Utilization** window.

The graphical **Setup** wizard facilitates performance tuning via the Dynatune Dynamic Tuning feature. Dynatune functionality allows Advanced Server to make optimal usage of the system resources available on the host machine on which it is installed.




The **edb_dynatune** configuration parameter determines how Advanced Server allocates system resources. Use the radio buttons on the **Server Utilization** window to set the initial value of the **edb_dynatune** configuration parameter:

- Select **Development** to set the value of **edb_dynatune** to **33**. A low value dedicates the least amount of the host machine's resources to the database server. This is a good choice for a development machine.
- Select **General Purpose** to set the value of **edb_dynatune** to **66**. A mid-range value dedicates a moderate amount of system resources to the database server. This would be a good setting for an application server with a fixed number of applications running on the same host as Advanced Server.
- Select **Dedicated** to set the value of **edb_dynatune** to **100**. A high value dedicates most of the system resources to the database server. This is a good choice for a dedicated server host.

After the installation is complete, you can adjust the value of **edb_dynatune** by editing the **postgresql.conf** file, located in the **data** directory of your Advanced Server installation. After editing the **postgresql.conf** file, you must restart the server for your changes to take effect.

Select the appropriate setting for your system, and click **Next** to continue to the **Dynatune Dynamic Tuning**:

Workload Profile window.



Use the radio buttons on the **Workload Profile** window to specify the initial value of the `edb_dynatune_profile` configuration parameter. The `edb_dynatune_profile` parameter controls performance-tuning aspects based on the type of work that the server performs.


- Select **Transaction Processing (OLTP systems)** to specify an `edb_dynatune_profile` value of `oltp`. Recommended when Advanced Server is supporting heavy online transaction processing.
- Select **General Purpose (OLTP and reporting workloads)** to specify an `edb_dynatune_profile` value of `mixed`. Recommended for servers that provide a mix of transaction processing and data reporting.
- Select **Reporting (Complex queries or OLAP workloads)** to specify an `edb_dynatune_profile` value of `reporting`. Recommended for database servers used for heavy data reporting.

After the installation is complete, you can adjust the value of `edb_dynatune_profile` by editing the `postgresql.conf` file, located in the `data` directory of your Advanced Server installation. After editing the `postgresql.conf` file, you must restart the server for your changes to take effect.

For more information about `edb_dynatune` and other performance-related topics, see the *EDB Postgres Advanced Server Guide* available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>


Click **Next** to continue. The **Update Notification Service** window opens.



When enabled, the update notification service notifies you of any new updates and security patches available for your installation of Advanced Server.



By default, Advanced Server is configured to start the service when the system boots; clear the `Install Update Notification Service` check box, or accept the default, and click **Next** to continue.

The **Pre Installation Summary** opens.




The **Pre Installation Summary** provides an overview of the options specified during the **Setup** process. Review the options before clicking **Next**; click **Back** to navigate back through the dialogs and update any options.

The **Ready to Install** window confirms that the installer has the information it needs about your configuration preferences to install Advanced Server. Click **Next** to continue.



As each supporting module is unpacked and installed, the module's installation is confirmed with a progress bar.

Before the **Setup** wizard completes the Advanced Server installation, it offers to **Launch StackBuilder Plus at exit?**



You can clear the **StackBuilder Plus** check box and click **Finish** to complete the Advanced Server installation, or accept the default and proceed to StackBuilder Plus.

EDB Postgres StackBuilder Plus is included with the installation of Advanced Server and its core supporting components. StackBuilder Plus is a graphical tool that can update installed products, or download and add supporting modules (and the resulting dependencies) after your Advanced Server setup and installation completes. See [Using StackBuilder Plus](#) for more information about StackBuilder Plus.

2.2.2 Invoking the Graphical Installer from the Command Line

The command line options of the Advanced Server installer offer functionality for Windows systems that reside in situations where a graphical installation may not work because of limited resources or privileges. You can:

- Include the **--mode unattended** option when invoking the installer to perform an installation without user input.
- Invoke the installer with the **--extract-only** option to perform a minimal installation when you don't hold the privileges required to perform a complete installation.

Not all command line options are suitable for all situations. For a complete reference guide to the command line options, see [Reference - Command Line Options](#).

Note

If you are invoking the installer from the command line to perform a system upgrade, the installer will ignore command line options, and preserve the configuration of the previous installation.

2.2.2.1 Performing an Unattended Installation

To specify that the installer should run without user interaction, include the **--mode unattended** command line option. In unattended mode, the installer uses one of the following sources for configuration parameters:

- command line options (specified when invoking the installer)
- parameters specified in an option file
- Advanced Server installation defaults

You can embed the non-interactive Advanced Server installer within another application installer; during the installation process, a progress bar allows the user to view the progression of the installation.

You must have administrative privileges to install Advanced Server using the `--mode unattended` option. If you are using the `--mode unattended` option to install Advanced Server with a client, the calling client must be invoked with superuser or administrative privileges.

To start the installer in unattended mode, navigate to the directory that contains the executable file, and enter:

```
edb-as12-server-12.x.x-x-windows-x64.exe --mode unattended --superpassword database_superuser_password  
--servicepassword system_password
```

When invoking the installer, include the `--servicepassword` option to specify an operating system password for the user installing Advanced Server.

Use the `--superpassword` option to specify a password that conforms to the password security policies defined on the host; enforced password policies on your system may not accept the default password (`enterprisedb`).

2.2.2.2 Performing an Installation with Limited Privileges

To perform an abbreviated installation of Advanced Server without access to administrative privileges, invoke the installer from the command line and include the `--extract-only` option. The `--extract-only` option extracts the binary files in an unaltered form, allowing you to experiment with a minimal installation of Advanced Server.

If you invoke the installer with the `--extract-only` options, you can either manually create a cluster and start the service, or run the installation script. To manually create the cluster, you must:

- Use `initdb` to initialize the cluster
- Configure the cluster
- Use `pg_ctl` to start the service

For more information about the `initdb` and `pg_ctl` commands, please see the PostgreSQL Core Documentation at:

<https://www.postgresql.org/docs/12/static/app-initdb.html>

<https://www.postgresql.org/docs/12/static/app-pg-ctl.html>

If you include the `--extract-only` option, the installer steps through a shortened form of the `Setup` wizard. During the brief installation process, the installer generates an installation script that can be later used to complete a more complete installation. You must have administrative privileges to invoke the installation script.

The installation script:

- Initializes the database cluster if the cluster is empty.
- Configures the server to start at boot-time.
- Establishes initial values for Dynature (dynamic tuning) variables.

The scripted Advanced Server installation does not create menu shortcuts or provide access to EDB Postgres StackBuilder Plus, and no modifications are made to registry files. The Advanced Server Update Monitor will not detect components installed by the scripted installation, and will not issue alerts for available updates to those components.


To perform a limited installation and generate an installation script, download and unpack the Advanced Server installer. Navigate into the directory that contains the installer, and invoke the installer with the command:

```
edb-as12-server-12.x.x-x-windows.exe --extract-only yes
```

A dialog opens, prompting you to choose an installation language. Select a language for the installation from the drop-down listbox, and click `OK` to continue. The `Setup Wizard` opens.




Click [Next](#) to continue.




On Windows, the default Advanced Server installation directory is:

[C:\Program Files\edb\as12](#)

You can accept the default installation location and click [Next](#) to continue to the [Ready to Install](#) window, or optionally click the [File Browser](#) icon to choose an alternate installation directory.



Click [Next](#) to proceed with the Advanced Server installation. During the installation, progress bars and popups mark the installation progress. The installer notifies you when the installation is complete.



After completing the minimal installation, you can execute a script to initialize a cluster and start the service. The script is (by default) located in:

`C:\Program Files\edb`

To execute the installation script, open a command line and assume administrative privileges. Navigate to the directory that contains the script, and execute the command:

`cscript runAsAdmin.vbs`

The installation script executes at the command line, prompting you for installation configuration information. The default configuration value is displayed in square braces immediately before each prompt; update the default value or press `Enter` to accept the default value and continue.

Performing a Scripted Installation

The following dialog is an example of a scripted installation. The actual installation dialog may vary, and will reflect the options specified during the installation.

Specify the installation directory is the directory where Advanced Server is installed:

Please enter the installation directory [C:\Program Files\edb\as12] :

Specify the directory in which Advanced Server data will be stored:

Please enter the data directory path: [C:\Program Files\edb\as12\data] :

Specify the WAL directory (where the write-ahead log will be written):

Please enter the Write-Ahead Log (WAL) directory path:

[C:\Program Files\edb\as12\data\pg_wal] :

The database mode specifies the database dialect with which the Advanced Server installation is compatible. The optional values are `oracle` or `postgresql`.

Please enter database mode: [oracle] :

Compatible with Oracle Mode

Specify `oracle` mode to include the following functionality:

- Data dictionary views and data type conversions compatible with Oracle databases.
- Date values displayed in a format compatible with Oracle syntax.
- Oracle-styled concatenation rules (if you concatenate a string value with a `NULL` value, the returned value is the value of the string).
- Schemas (`dbo` and `sys`) compatible with Oracle databases added to the `SEARCH_PATH`.
- Support for the Oracle built-in packages.

If you choose to install in **Compatible with Oracle** mode, the Advanced Server superuser name is **enterprisedb**.

Compatible with PostgreSQL Mode

Specify **postgresql** to install Advanced Server with complete compatibility with Postgres version 12.

For more information about PostgreSQL functionality, see the PostgreSQL Core Documentation available at:

<https://www.enterprisedb.com/edb-docs>

If you choose to install in **Compatible with PostgreSQL** mode, the Advanced Server superuser name is **postgres**.

Specify a port number for the Advanced Server listener to listen on:

NOTE: We will not be able to examine, if port is currently used by other application.

Please enter port: [5444] :

Specify a locale for the Advanced Server installation. If you accept the **DEFAULT** value, the locale defaults to the locale of the host system.

Please enter the locale: [**DEFAULT**] :

You can optionally install sample tables and procedures. Press **Return**, or enter **Y** to accept the default and install the sample tables and procedures; enter an **N** and press **Return** to skip this step.

Install sample tables and procedures? (Y/n): [**Y**] :

Specify a password for the database superuser. By default, the database superuser is named **enterprisedb**.

Please enter the password for the SuperUser(**enterprisedb**): [] :

Specify a password for the service account user.

Please enter the password for the ServiceAccount(**enterprisedb**): [] :

The server utilization value is used as an initial value for the **edb_dynatune** configuration parameter. **edb_dynatune** determines how Advanced Server allocates system resources.

- A low value dedicates the least amount of the host machine's resources to the database server; a low value is a good choice for a development machine.
- A mid-range value dedicates a moderate amount of system resources to the database server. A mid-range value is a good setting for an application server with a fixed number of applications running on the same host as Advanced Server.
- A high value dedicates most of the system resources to the database server. This is a good choice for a dedicated server host.

Specify a value between **1** and **100**:

Please enter the Server Utilization: [**66**] :

After the installation is complete, you can adjust the value of **edb_dynatune** by editing the **postgresql.conf** file, located in the **data** directory of your Advanced Server installation. After editing the **postgresql.conf** file, you must restart the server for the changes to take effect.

The workload profile value is used as an initial value for the **edb_dynatune_profile** configuration parameter. **edb_dynatune_profile** controls performance-tuning based on the type of work that the server performs.

- Specify **oltp** if the server will be supporting heavy online transaction workloads.
- Specify **mixed** if the server will provide a mix of transaction processing and data reporting.
- Specify **reporting** if the database server will be used for heavy data reporting.

Specify a value between 1 and 100:

Please enter the Workload Profile: [oltp] :

After the installation is complete, you can adjust the value of `edb_dynatune_profile` by editing the `postgresql.conf` file, located in the `data` directory of your Advanced Server installation, and restarting the server.

After processing, the Advanced Server installation is complete.

2.2.2.3 Reference - Command Line Options

You can optionally include the following parameters for an Advanced Server installation on the command line, or in a configuration file when invoking the Advanced Server installer.

`--create_samples { yes | no }`

Use the `--create_samples` option to specify whether the installer should create the sample tables and procedures for the database dialect specified with the `--databasemode` parameter. The default is `yes`.

`--databasemode { oracle | postgresql }`

Use the `--databasemode` parameter to specify a database dialect. The default is `oracle`.

`--datadir data_directory`

Use the `--datadir` parameter to specify a location for the cluster's data directory. `data_directory` is the name of the directory; include the complete path to the desired directory.

`--debuglevel { 0 | 1 | 2 | 3 | 4 }`

Use the `--debuglevel` parameter to set the level of detail written to the `debug_log` file (see `--debugtrace`). Higher values produce more detail in a longer trace file. The default is `2`.

`--debugtrace debug_log`

Use the `--debugtrace` parameter to troubleshoot installation problems. `debug_log` is the name of the file that contains troubleshooting details.

`--disable-components component_list`

Use the `--disable-components` parameter to specify a list of Advanced Server components to exclude from the installation. By default, `component_list` contains "" (the empty string). `component_list` is a comma-separated list containing one or more of the following components:

`dbserver`

EDB Postgres Advanced Server 12.

`pgadmin4`

The EDB Postgres pgAdmin 4 provides a powerful graphical interface for database management and monitoring.

`--enable_acledit { 1 | 0 }`

The `--enable_acledit 1` option instructs the installer to grant permission to the user specified by the `--`

`serviceaccount` option to access the Advanced Server binaries and `data` directory. By default, this option is disabled if `--enable_acledit 0` is specified or if the `--enable_acledit` option is completely omitted.

Note

Specification of this option is valid only when installing on Windows. The `--enable_acledit 1` option should be specified when a `discretionary access control list` (DACL) needs to be set for allowing access to objects on which Advanced Server is to be installed. See the following for information on a DACL:

[https://msdn.microsoft.com/en-us/library/windows/desktop/aa446597\(v=vs.85\).aspx](https://msdn.microsoft.com/en-us/library/windows/desktop/aa446597(v=vs.85).aspx)

In order to perform future operations such as upgrading Advanced Server, access to the `data` directory must exist for the service account user specified by the `--serviceaccount` option. By specifying the `--enable_acledit 1` option, access to the `data` directory by the service account user is provided.

`--enable-components component_list`

Although this option is listed when you run the installer with the `--help` option, the `--enable-components` parameter has absolutely no effect on which components are installed. All components will be installed regardless of what is specified in `component_list`. In order to install only specific selected components, you must use the `--disable-components` parameter previously described in this section to list the components you do not want to install.

`--extract-only { yes | no }`

Include the `--extract-only` parameter to indicate that the installer should extract the Advanced Server binaries without performing a complete installation. Superuser privileges are not required for the `--extract-only` option. The default value is `no`.

`--help`

Include the `--help` parameter to view a list of the optional parameters.

`--installer-language { en | ja | zh_CN | zh_TW | ko }`

Use the `--installer-language` parameter to specify an installation language for Advanced Server. The default is `en`.

`en` specifies English.

`ja` specifies Japanese

`zh_CN` specifies Chinese Simplified.

`zh_TW` specifies Traditional Chinese.

`ko` specifies Korean.

`--install_runtimes { yes | no } (Windows only.)`

Include `--install_runtimes` to specify whether the installer should install the Microsoft Visual C++ runtime libraries. Default is `yes`.

`--locale locale`

Specifies the locale for the Advanced Server cluster. By default, the installer will use the locale detected by `initdb`.

`--mode { unattended }`

Use the `--mode unattended` parameter to specify that the installer should perform an installation that requires no user input during the installation process.

--optionfile config_file

Use the `--optionfile` parameter to specify the name of a file that contains the installation configuration parameters. `config_file` must specify the complete path to the configuration parameter file.

--prefix installation_dir/as12.x

Use the `--prefix` parameter to specify an installation directory for Advanced Server. The installer will append a version-specific sub-directory (that is, `as12`) to the specified directory. The default installation directory is:

`C:\Program Files\edb\as12`

--serverport port_number

Use the `--serverport` parameter to specify a listener port number for Advanced Server.

If you are installing Advanced Server in unattended mode, and do not specify a value using the `--serverport` parameter, the installer will use port `5444`, or the first available port after port `5444` as the default listener port.

--server_utilization {33 | 66 | 100}

Use the `--server_utilization` parameter to specify a value for the `edb_dynatune` configuration parameter. The `edb_dynatune` configuration parameter determines how Advanced Server allocates system resources.

- A value of `33` is appropriate for a system used for development. A low value dedicates the least amount of the host machine's resources to the database server.
- A value of `66` is appropriate for an application server with a fixed number of applications. A mid-range value dedicates a moderate amount of system resources to the database server. The default value is `66`.
- A value of `100` is appropriate for a host machine that is dedicated to running Advanced Server. A high value dedicates most of the system resources to the database server.

When the installation is complete, you can adjust the value of `edb_dynatune` by editing the `postgresql.conf` file, located in the `data` directory of your Advanced Server installation. After editing the `postgresql.conf` file, you must restart the server for the changes to take effect.

--serviceaccount user_account_name

Use the `--serviceaccount` parameter to specify the name of the user account that owns the server process.

- If `--databasemode` is set to `oracle` (the default), the default value of `--serviceaccount` is `enterprisedb`.
- If `--databasemode` is `postgresql`, the default value of `--serviceaccount` is set to `postgres`.

Please note that for security reasons, the `--serviceaccount` parameter must specify the name of an account that does not hold administrator privileges.

If you specify both the `--serviceaccount` option and the `--enable_acledit 1` option when invoking the installer, the database service and pgAgent will use the same service account, thereby having the required permissions to access the Advanced Server binaries and `data` directory.

Note

Specification of the `--enable_acledit` option is permitted only when installing on Windows.

Note

If you do not include the `--serviceaccount` option when invoking the installer, the `NetworkService` account will own the database service, and the pgAgent service will be owned by either `enterprisedb` or `postgres` (depending on the installation mode).

--servicename service_name

Use the `--servicename` parameter to specify the name of the Advanced Server service. The default is `edb-as-12`.

`--servicepassword user_password`

Use `--servicepassword` to specify the OS system password. If unspecified, the value of `--servicepassword` defaults to the value of `--superpassword`.

`--superaccount super_user_name`

Use the `--superaccount` parameter to specify the user name of the database superuser.

- If `--databasemode` is set to `oracle` (the default), the default value of `--superaccount` is `enterprisedb`.
- If `--databasemode` is set to `postgresql`, the default value of `--superaccount` is set to `postgres`.

`--superpassword superuser_password`

Use `--superpassword` to specify the database superuser password. If you are installing in non-interactive mode, `--superpassword` defaults to `enterprisedb`.

`--unattendedmodeui { none | minimal | minimalWithDialogs }`

Use the `--unattendedmodeui` parameter to specify installer behavior during an unattended installation.

Include `--unattendedmodeui none` to specify that the installer should not display progress bars during the Advanced Server installation.

Include `--unattendedmodeui minimal` to specify that the installer should display progress bars during the installation process. This is the default behavior.

Include `--unattendedmodeui minimalWithDialogs` to specify that the installer should display progress bars and report any errors encountered during the installation process (in additional dialogs).

`--version`

Include the `--version` parameter to retrieve version information about the installer:

EDB Postgres Advanced Server 12 --- Built on 2018-03-15 00:04:00 IB: 15.12.1-201511121057

`--workload_profile {oltp | mixed | reporting}`

Use the `--workload_profile` parameter to specify an initial value for the `edb_dynatune_profile` configuration parameter. `edb_dynatune_profile` controls aspects of performance-tuning based on the type of work that the server performs.

- Specify `oltp` if the Advanced Server installation will be used to support heavy online transaction processing workloads.
- The default value is `oltp`.
- Specify `mixed` if Advanced Server will provide a mix of transaction processing and data reporting.
- Specify `reporting` if Advanced Server will be used for heavy data reporting.

After the installation is complete, you can adjust the value of `edb_dynatune_profile` by editing the `postgresql.conf` file, located in the `data` directory of your Advanced Server installation. After editing the `postgresql.conf` file, you must restart the server for the changes to take effect.

For more information about `edb_dynatune` and other performance-related topics, see the *EDB Postgres Advanced Server Guide* available at:

<https://www.enterprisedb.com/edb-docs>

2.2.3 Using StackBuilder Plus

Please note: StackBuilder Plus is supported only on Windows systems.

The StackBuilder Plus utility provides a graphical interface that simplifies the process of updating, downloading, and installing modules that complement your Advanced Server installation. When you install a module with StackBuilder Plus, StackBuilder Plus automatically resolves any software dependencies.


You can invoke StackBuilder Plus at any time after the installation has completed by selecting the **StackBuilder Plus** menu option from the **Apps** menu. Enter your system password (if prompted), and the StackBuilder Plus welcome window opens.



Use the drop-down listbox on the welcome window to select your Advanced Server installation.

StackBuilder Plus requires Internet access; if your installation of Advanced Server resides behind a firewall (with restricted Internet access), StackBuilder Plus can download program installers through a proxy server. The module provider determines if the module can be accessed through an HTTP proxy or an FTP proxy; currently, all updates are transferred via an HTTP proxy and the FTP proxy information is not used.

If the selected Advanced Server installation has restricted Internet access, use the **Proxy Servers** on the **Welcome** window to open the **Proxy servers** dialog.



Enter the IP address and port number of the proxy server in the **HTTP proxy** on the **Proxy Servers** dialog. Currently, all StackBuilder Plus modules are distributed via HTTP proxy (FTP proxy information is ignored). Click **OK** to continue.



The tree control on the StackBuilder Plus module selection window displays a node for each module category.

Expand a module, and highlight a component name in the tree control to review a detailed description of the component. To add one or more components to the installation or to upgrade a component, check the box to the left of a module name and click **Next**.

StackBuilder Plus confirms the packages selected.



Use the browse icon (...) to the right of the **Download directory** field to open a file selector, and choose an alternate location to store the downloaded installers. Click **Next** to connect to the server and download the required installation files.

When the download completes, a window opens that confirms the installation files have been downloaded and are ready for installation.




You can check the box next to **Skip Installation**, and select **Next** to exit StackBuilder Plus without installing the downloaded files, or leave the box unchecked and click **Next** to start the installation process.

Each downloaded installer has different requirements. As the installers execute, they may prompt you to confirm acceptance of license agreements, to enter passwords, and provide configuration information.

During the installation process, you may be prompted by one (or more) of the installers to restart your system. Select **No** or **Restart Later** until all installations are completed. When the last installation has completed, reboot the system to apply all of the updates.

You may occasionally encounter packages that don't install successfully. If a package fails to install, StackBuilder Plus will alert you to the installation error with a popup dialog, and write a message to the log file at %TEMP%.



When the installation is complete, StackBuilder Plus will alert you to the success or failure of the installations of the requested packages. If you were prompted by an installer to restart your computer, reboot now.

2.2.4 Using the Update Monitor

The Update Monitor utility polls the EnterpriseDB website and alerts you to security updates and enhancements for Windows hosts as they become available. Update Monitor is automatically installed by the graphical installer.


When Update Monitor is actively monitoring, the Postgres elephant icon is displayed in the system tray.



If you have installed more than one version of Advanced Server, Update Monitor watches for updates and alerts for all installed versions. When Update Monitor finds an update or alert, it displays an alert symbol to let you know that an update or alert is available for one of the Advanced Server installations.




Right click on the symbol to access the Update Monitor context menu.



Click [Install components](#) to open StackBuilder Plus, and check for component updates. See [Using StackBuilder Plus](#) for detailed information about the update process.

If one or more alerts are available for your Advanced Server installation, the context menu displays the alert count. Select [View alerts](#) to display the [EnterpriseDB Advanced Server Alerts](#) window.



The technical alert will provide information about updates and issues that might impact your installation.

- Click **Run StackBuilder Plus** to open StackBuilder Plus and install relevant software updates.
- Click **Next** to access each successive alert.

When you have reviewed all of the alerts, the icon no longer displays the alert symbol; use the X icon in the upper-right corner of the dialog to close Update Monitor.

2.2.5 Installation Troubleshooting

Difficulty Displaying Java-based Applications

If you encounter difficulty displaying Java-based server features (controls or text not being displayed correctly, or blank windows), upgrading to the latest **libxcb-xlib** libraries should correct the problem on most distributions. Please visit the following link for other possible work-arounds:

http://bugs.sun.com/bugdatabase/view_bug.do?bug_id=6532373

--mode unattended Authentication Errors

Authentication errors from component modules during unattended installations may indicate that the specified values of **--servicepassword** or **--superpassword** may be incorrect.

Errors During an Advanced Server Installation

If you encounter an error during the installation process, exit the installation, and ensure that your version of Windows is up-to-date. After applying any outstanding operating system updates, re-invoke the Advanced Server installer.

The Installation Fails to Complete Due to Existing Data Directory Contents

If an installation fails to complete due to existing content in the data directory, the server will write an error message to the server logs:

A data directory is neither empty, or a recognisable data directory .

If you encounter a similar message, you should confirm that the data directory is empty; the presence of files (including the system-generated **lost+found** folder) will prevent the installation from completing. Either remove the files from the data directory, or specify a different location for the data directory before re-invoking the installer to complete the installation.

2.3 Managing an Advanced Server Installation

Unless otherwise noted, the commands and paths noted in the following section assume that you have performed an installation with the interactive installer.

Starting and Stopping Advanced Server and Supporting Components

A service is a program that runs in the background and requires no user interaction (in fact, a service provides no user interface); a service can be configured to start at boot time, or manually on demand. Services are best controlled using the platform-specific operating system service control utility. Many of the Advanced Server supporting components are services.


The following table lists the names of the services that control Advanced Server and services that control Advanced Server supporting components:

| Advanced Server Component Name | Windows Service Name |
|--------------------------------|--|
| Advanced Server | edb-as-12 |
| pgAgent | EDB Postgres Advanced Server 12 Scheduling Agent |
| PgBouncer | edb-pgbouncer-1.12 |
| Slony | edb-slony-replication-12 |

You can use the command line or the Windows Services applet to control Advanced Server's database server and the services of Advanced Server's supporting components on a Windows host.

2.3.1 Using the Windows Services Applet

The Windows operating system includes a graphical service controller that offers control of Advanced Server and the services associated with Advanced Server components. The **Services** utility can be accessed through the **Administrative Tools** section of the Windows **Control Panel**.



The **Services** window displays an alphabetized list of services; the **edb-as-12** service controls Advanced Server.

- Use the `Stop the service` option to stop the instance of Advanced Server. Please note that any user (or client application) connected to the Advanced Server instance will be abruptly disconnected if you stop the service.
 - Use the `Start the service` option to start the Advanced Server service.
 - Use the `Pause the service` option to tell Advanced Server to reload the server configuration parameters without disrupting user sessions for many of the configuration parameters. See [Configuring Advanced Server](#) for more information about the parameters that can be updated with a server reload.
 - Use the `Restart the service` option to stop and then start the Advanced Server. Please note that any user sessions will be terminated when you stop the service. This option is useful to reset server parameters that only take effect on server start.
-

2.3.2 Using pg_ctl to Control Advanced Server

You can use the `pg_ctl` utility to control an Advanced Server service from the command line on any platform. `pg_ctl` allows you to start, stop, or restart the Advanced Server database server, reload the configuration parameters, or display the status of a running server. To invoke the utility, assume the identity of the cluster owner, navigate into the home directory of Advanced Server, and issue the command:

```
./bin/pg_ctl -D data_directory action
data_directory
data_directory is the location of the data controlled by the Advanced Server cluster.
action
```

`action` specifies the action taken by the `pg_ctl` utility. Specify:

- `start` to start the service.
- `stop` to stop the service.
- `restart` to stop and then start the service.
- `reload` sends the server a `SIGHUP` signal, reloading configuration parameters
- `status` to discover the current status of the service.


For more information about using the `pg_ctl` utility, or the command line options available, please see the official PostgreSQL Core Documentation available at:

<https://www.postgresql.org/docs/12/static/app-pg-ctl.html>

2.3.3 Controlling Server Startup Behavior on Windows

You can use the Windows Services utility to control the startup behavior of the server. Right click on the name of the service you wish to update, and select `Properties` from the context menu to open the `Properties` dialog.

Use the drop-down listbox in the `Startup type` field to specify how the Advanced Server service will behave when the host starts.




- Specify **Automatic (Delayed Start)** to instruct the service controller to start after boot.
- Specify **Automatic** to instruct the service controller to start and stop the server whenever the system starts or stops.
- Specify **Manual** to instruct the service controller that the server must be started manually.
- Specify **Disabled** to instruct the service controller to disable the service; after disabling the service, you must stop the service or restart the server to make the change take effect. Once disabled, the server's status cannot be changed until **Startup type** is reset to **Automatic (Delayed Start)**, **Automatic**, or **Manual**.

2.4 Configuring Advanced Server

You can easily update parameters that determine the behavior of Advanced Server and supporting components by modifying the following configuration files:


- The **postgresql.conf** file determines the initial values of Advanced Server configuration parameters.
- The **pg_hba.conf** file specifies your preferences for network authentication and authorization.
- The **pg_ident.conf** file maps operating system identities (user names) to Advanced Server identities (roles) when using **ident**-based authentication.

You can use your editor of choice to open a configuration file, or on Windows navigate through the **EDB Postgres** menu to open a file.



2.4.1 Modifying the postgresql.conf File

Configuration parameters in the `postgresql.conf` file specify server behavior with regards to auditing, authentication, encryption, and other behaviors. The `postgresql.conf` file resides in the `data` directory under your Advanced Server installation.



```
# -----
# PostgreSQL configuration file
# -----
#
# This file consists of lines of the form:
#
#   name = value
#
# (The "=" is optional.) Whitespace may be used. Comments are introduced with
# "#" anywhere on a line. The complete list of parameter names and allowed
# values can be found in the PostgreSQL documentation.
#
# The commented-out settings shown in this file represent the default values.
# Re-commenting a setting is NOT sufficient to revert it to the default value;
# you need to reload the server.
#
# This file is read on server startup and when the server receives a SIGHUP
# signal. If you edit the file on a running system, you have to SIGHUP the
# server for the changes to take effect, run "pg_ctl reload", or execute
# "SELECT pg_reload_conf()". Some parameters, which are marked below,
# require a server shutdown and restart to take effect.
#
# Any parameter can also be given as a command-line option to the server, e.g.,
# "postgres -c log_connections=on". Some parameters can be changed at run time
# with the "SET" SQL command.
#
# Memory units: kB = Kilobytes      Time units: ms = milliseconds
#                 MB = megabytes        s = seconds
#                 GB = gigabytes         min = minutes
```

Parameters that are preceded by a pound sign (#) are set to their default value (as shown in the parameter setting). To change a parameter value, remove the pound sign and enter a new value. After setting or changing a parameter, you must either *reload* or *restart* the server for the new parameter value to take effect.

Within the `postgresql.conf` file, some parameters contain comments that indicate `change requires restart`. To view a list of the parameters that require a server restart, execute the following query at the `psql` command line:

```
| SELECT name FROM pg_settings WHERE context = 'postmaster';
```

2.4.2 Modifying the pg_hba.conf File

Appropriate authentication methods provide protection and security. Entries in the `pg_hba.conf` file specify the authentication method or methods that the server will use when authenticating connecting clients. Before connecting to the server, you may be required to modify the authentication properties specified in the `pg_hba.conf` file.

When you invoke the `initdb` utility to create a cluster, `initdb` creates a `pg_hba.conf` file for that cluster that specifies the type of authentication required from connecting clients. The following figure displays the default authentication configuration specified in the `pg_hba.conf` file.

```

# TYPE DATABASE USER ADDRESS METHOD
# "local" is for Unix domain socket connections only
local all all scram-sha-256
# IPv4 local connections:
host all all 127.0.0.1/32 scram-sha-256
# IPv6 local connections:
host all all ::1/128 scram-sha-256
# Allow replication connections from localhost, by a user with the
# replication privilege.
local replication all scram-sha-256
host replication all 127.0.0.1/32 scram-sha-256
host replication all ::1/128 scram-sha-256

```

To modify the `pg_hba.conf` file, open the file with your choice of editor. After modifying the authentication settings in the `pg_hba.conf` file, use the Windows services utility to restart the server and apply the changes.

For more information about authentication, and modifying the `pg_hba.conf` file, see the PostgreSQL Core Documentation at:

<https://www.postgresql.org/docs/12/static/auth-pg-hba-conf.html>

2.4.3 Setting Advanced Server Environment Variables

The graphical installer provides a script that simplifies the task of setting environment variables for Windows users. The script sets the environment variables for your current shell session; when your shell session ends, the environment variables are destroyed. You may wish to invoke `pgplus_env.bat` from your system-wide shell startup script, so that environment variables are automatically defined for each shell session.

The `pgplus_env` script is created during the Advanced Server installation process and reflects the choices made during installation. To invoke the script, open a command line and enter:

```
C:\Program Files\edb\as12\pgplus_env.bat
```

As the `pgplus_env.bat` script executes, it sets the following environment variables:

```

PATH="C:\Program Files\edb\as12\bin";%PATH%
EDBHOME=C:\Program Files\edb\as12
PGDATA=C:\Program Files\edb\as12\data
PGDATABASE=edb
REM @SET PGUSER=enterprisedb
PGPORT=5444
PGLOCALEDIR=C:\Program Files\edb\as12\share\locale

```

If you have used an installer created by EnterpriseDB to install PostgreSQL, the `pg_env` script performs the same function:

```
C:\Progra~1\PostgreSQL\12\pg_env.bat
```

As the `pg_env.bat` script executes on PostgreSQL, it sets the following environment variables:

```

PATH="C:\Program Files\PostgreSQL\12\bin";%PATH%
PGDATA=C:\Program Files\PostgreSQL\12\data
PGDATABASE=postgres
PGUSER=postgres

```

```
PGPORT=5432
PGLOCALEDIR=C:\Program Files\PostgreSQL\12\share\locale
```

2.4.4 Connecting to Advanced Server with psql

psql is a command line client application that allows you to execute SQL commands and view the results. To open the psql client, the client must be in your search path. The executable resides in the **bin** directory, under your Advanced Server installation.

Use the following command and options to start the psql client:

```
| psql -d edb -U enterpriseedb
```

```
File Edit Format View Help
C:\Program Files\edb\as12\bin>psql.exe -U enterpriseedb -d edb
Password for user enterpriseedb:
psql (12.0.1, server 12.0.1)
WARNING: Console code page (437) differs from windows code page (1252)
         8-bit characters might not work correctly. See psql reference
         page "Notes for windows users" for details.
Type "help" for help.

edb=#
```

Where:

-d specifies the database to which psql will connect;

-U specifies the identity of the database user that will be used for the session.

If you have performed an installation with the interactive installer, you can access the psql client by selecting **EDB-PSQL** from the **EDB Postgres** menu. When the client opens, provide connection information for your session.

For more information about using the command line client, please refer to the PostgreSQL Core Documentation at:

```
| https://www.postgresql.org/docs/12/static/app-psql.html
```

2.4.5 Connecting to Advanced Server with the pgAdmin 4 Client

pgAdmin 4 provides an interactive graphical interface that you can use to manage your database and database objects. Easy-to-use dialogs and online help simplify tasks such as object creation, role management, and granting or revoking privileges. The tabbed browser panel provides quick access to information about the object currently selected in the pgAdmin tree control.


The client is distributed with the graphical installer. To open pgAdmin, select pgAdmin4 from the **EDB Postgres** menu. The client opens in your default browser.



To connect to the Advanced Server database server, expand the **Servers** node of the **Browser** tree control, and right click on the **EDB Postgres Advanced Server** node. When the context menu opens, select **Connect Server**. The **Connect to Server** dialog opens.



Provide the password associated with the database superuser in the **Password** field, and click **OK** to connect.



When the client connects, you can use the **Browser** tree control to retrieve information about existing database objects, or to create new objects. For more information about using the pgAdmin client, use the **Help** drop-down menu to access the online help files.

2.5 Uninstalling Advanced Server


Note that after uninstalling Advanced Server, the cluster data files remain intact and the service user persists. You may manually remove the cluster data and service user from the system.

Using Advanced Server Uninstallers at the Command Line


The Advanced Server interactive installer creates an uninstaller that you can use to remove Advanced Server or components that reside on a Windows host. The uninstaller is created in `C:\Program Files\edb\as12`. To open the uninstaller, assume superuser privileges, navigate into the directory that contains the uninstaller, and enter:

`uninstall-edb-as12-server.exe`

The uninstaller opens.




You can remove the **Entire application** (the default), or select the radio button next to **Individual components** to select components for removal; if you select **Individual components**, a dialog will open, prompting you to select the components you wish to remove. After making your selection, click **Next**.



If you have elected to remove components that are dependent on Advanced Server, those components will be removed first; click **Yes** to acknowledge that you wish to continue.

Progress bars are displayed as the software is removed. When the uninstallation is complete, an **Info** dialog opens to confirm that Advanced Server (and/or its components) has been removed.



3 Release Notes

With the release of EDB Postgres Advanced Server 12.0, EnterpriseDB continues to lead as the only worldwide company to deliver innovative and low cost open-source-derived database solutions with commercial quality, ease of use, compatibility, scalability, and performance for small or large-scale enterprises.

EDB Postgres Advanced Server 12.0 is built on open-source PostgreSQL 12.0, which introduces a myriad of enhancements that enable databases to scale up and scale out in more efficient ways.

EDB Postgres Advanced Server 12.0 adds a number of new outstanding features, including:

- Interval Partition
- Logical decoding on standby
- COMPOUND TRIGGER
- MEDIAN and LISTAGG
- CAST_MULTISET

For more information about EnterpriseDB software, visit:

<https://www.enterprisedb.com/enterprise-postgres/edb-postgres-platform>

Installers and Documentation

EDB Postgres Advanced Server v12 is packaged and delivered as interactive installers for Windows; visit the EnterpriseDB website:

<https://www.enterprisedb.com/advanced-downloads>

RPM Packages are available for Linux from:

<https://yum.enterprisedb.com/>

Debian/Ubuntu Packages are available for download from:

<https://apt.enterprisedb.com/>

If you need to request repository access, visit:

<https://www.enterprisedb.com/repository-access-request>

Please note: the EnterpriseDB repository structure has changed; for detailed installation instructions, please see the EDB Postgres Advanced Server Linux Installation Guide.

Documentation is provided on the EnterpriseDB website. Visit:

<https://www.enterprisedb.com/edb-docs/p/edb-postgres-advanced-server>

Supported Platforms

EDB Postgres Advanced Server v12 supports 64 bit Linux and Windows server platforms. The Advanced Server 12 RPM packages are supported on the following 64-bit Linux platforms:

- Red Hat Enterprise Linux (x86_64) 6.x and 7.x
- CentOS (x86_64) 6.x and 7.x

- OEL (x86_64) 6.x and 7.x
- PPC-LE 8 running RHEL 7.x

The Advanced Server 12 native packages are supported on the following 64-bit Linux platforms:

- Debian 9
- Ubuntu 18.04 LTS

Graphical installers are supported on the following 64-bit Windows platforms:

- Windows 2019
- Windows Server 2016
- Windows Server 2012 R2 Server

Additional information about supported platforms is available on the EnterpriseDB website:

<https://www.enterprisedb.com/services-support/edb-supported-products-and-platforms>

Component Certification

The following components are included in the EBD Postgres Advanced Server v12 release:

- Procedural Language Packs – PL/Perl 5.26, PL/Python 3.7, PL/TCL 8.6
- CloneSchema 1.10
- Parallel Clone 1.5
- pgAgent 4.15
- Slony 2.2.8
- Connectors JDBC 42.2.8, ODBC 12.00.0000 .NET 4.0.10.1, OCI 11.0.3.1
- pgAdmin 4.15
- pgBouncer 1.12.0
- pgPool-II & pgPool-IIExtensions 4.0.6
- MTK 53.0.0
- EDBPlus 38.0.0
- PostGIS-JDBC 2.2.1

EDB Postgres Advanced Server v12 Features

The major highlights of this release are :

- Advanced Server now supports interval partitioning. Interval partitioning is an extension to range partitioning where the system automatically creates the partition if a new tuple doesn't fit into existing partitions.
- For better Redwood compatibility, Advanced Server now treats the SELECT UNIQUE clause synonymous to SELECT DISTINCT clause (i.e. they can be used interchangeably in all places).
- Advanced Server now supports COMPOUND TRIGGERS. This provides a way to define a single trigger on a table that enables the user to specify actions for different DML event points. It allows defining BEFORE STATEMENT, AFTER STATEMENT, BEFORE EACH ROW, AFTER EACH ROW timing events in a single trigger.
- Advanced Server allows logical decoding on a standby server. With this feature, the user can now create a logical replication slot on standby server.
- Advanced Server now provides CAST(MULTISET..). CAST(MULTISET) is an extension to the redwood style casting function CAST(.. AS ..) where the subquery result is like a table output CAST to a collection type.
- Advanced Server now supports two aggregate functions: LISTAGG and MEDIAN
- The LISTAGG function supports string aggregation that concatenates data from multiple rows into a single row in an ordered fashion.
- The MEDIAN function calculates a median value from the set of provided values. The aggregate function can be used with numeric, date/time and interval types and much like other aggregate functions, it can be used with window function as well.
- In addition to supporting %type and %rowtype variables for functions and procedures, Advanced Server now

supports the same inside a package as well.

- Advanced Server now supports a Redwood compatible to_timestamp implementation, which facilitates stricter input string parsing for timestamp columns for EDB Loader.
- Advanced Server has re-implemented ROWIDs using identity data type after PG v12 PostgreSQL community removed support for OIDs (to which ROWIDs were mapped).
- Advanced Server has added the SYS_GUID function to generate and return a globally unique identifier in the form of 16-bytes of RAW data.
- Advanced Server now offers new view that provide information which is compatible with the Oracle data dictionary views:

For information about Advanced Server features that are compatible with Oracle databases, see the following guides:

- *Database Compatibility for Oracle Developer's Guide*
- *Database Compatibility for Oracle Developer's Reference Guide*
- *Database Compatibility for Oracle Developer's Tools and Utilities Guide*
- *Database Compatibility for Oracle Developer's Built-in Package Guide*

Community PostgreSQL 12 Updates

Advanced Server 12 integrates all of the community PostgreSQL 12 features. To review a complete list of changes to the community PostgreSQL project and the contributors names, see the PostgreSQL 12 Release Notes at:

<https://www.postgresql.org/docs/12/release-12.html>

The following updates are available in PostgreSQL 12:

Partitioning Updates

Improve performance of many operations on partitioned tables.

Allow tables with thousands of child partitions to be processed efficiently by operations that only affect a small number of partitions.

Allow foreign keys to reference partitioned tables.

Improve speed of COPY into partitioned tables.

Allow partition bounds to be any expression.

Such expressions are evaluated at partitioned-table creation time. Previously, only simple constants were allowed as partition bounds.

Allow CREATE TABLE's tablespace specification for a partitioned table to affect the tablespace of its children.

Avoid sorting when partitions are already being scanned in the necessary order.

ALTER TABLE ATTACH PARTITION is now performed with reduced locking requirements.

Add partition introspection functions.

The new function pg_partition_root() returns the top-most parent of a partition tree, pg_partition_ancestors() reports all ancestors of a partition, and pg_partition_tree() displays information about partitions.

Include partitioned indexes in the system view pg_indexes .

Add psql command \dP to list partitioned tables and indexes.

Improve psql \d and \z display of partitioned tables.

Fix bugs that could cause ALTER TABLE DETACH PARTITION to leave behind incorrect dependency state, allowing subsequent operations to misbehave, for example by not dropping a former partition child index when its table is dropped.

Indexes

Improve performance and space utilization of btree indexes with many duplicates.

Previously, duplicate index entries were stored unordered within their duplicate groups. This caused overhead during index inserts, wasted space due to excessive page splits, and it reduced VACUUM's ability to recycle entire pages. Duplicate index entries are now sorted in heap-storage order.

Indexes pg_upgrade'd from previous releases will not have these benefits.

Allow multi-column btree indexes to be smaller.

Internal pages and min/max leaf page indicators now only store index keys until the change key, rather than all indexed keys. This also improves the locality of index access.

Indexes pg_upgrade'd from previous releases will not have these benefits.

Improve speed of btree index insertions by reducing locking overhead.

Support INCLUDE columns in GiST indexes.

Add support for nearest-neighbor (KNN) searches of SP-GiST indexes.

Reduce the WAL write overhead of GiST, GIN, and SP-GiST index creation.

Allow index-only scans to be more efficient on indexes with many columns.

Improve the performance of vacuum scans of GiST indexes.

Delete empty leaf pages during GiST VACUUM.

Reduce locking requirements for index renaming.

Optimizer

Allow CREATE STATISTICS to create most-common-value statistics for multiple columns.

This improves optimization for queries that test several columns, requiring an estimate of the combined effect of several WHERE clauses. If the columns are correlated and have non-uniform distributions then multi-column statistics will allow much better estimates.

Allow common table expressions (CTEs) to be inlined into the outer query.

Specifically, CTEs are automatically inlined if they have no side-effects, are not recursive, and are referenced only once in the query. Inlining can be prevented by specifying MATERIALIZED, or forced for multiply-referenced CTEs by specifying NOT MATERIALIZED. Previously, CTEs were never inlined and were always evaluated before the rest of the query.

Allow control over when generic plans are used for prepared statements.

This is controlled by the plan_cache_mode server parameter.

Improve optimization of partition and UNION ALL queries that have only a single child.

Improve processing of domains that have no check constraints.

Domains that are being used purely as type aliases no longer cause optimization difficulties.

Pre-evaluate calls of LEAST and GREATEST when their arguments are constants.

Improve optimizer's ability to verify that partial indexes with IS NOT NULL conditions are usable in queries.

Usability can now be recognized in more cases where the calling query involves casts or large x IN (array) clauses.

Compute ANALYZE statistics using the collation defined for each column.

Previously, the database's default collation was used for all statistics. This potentially gives better optimizer behavior for columns with non-default collations.

Improve selectivity estimates for inequality comparisons on ctid columns.

Improve optimization of joins on columns of type tid.

These changes primarily improve the efficiency of self-joins on ctid columns.

Fix the leakproofness designations of some btree comparison operators and support functions.

This allows some optimizations that previously would not have been applied in the presence of security barrier views or row-level security.

General Performance

Enable Just-in-Time (JIT) compilation by default, if the server has been built with support for it.

Note that this support is not built by default, but has to be selected explicitly while configuring the build.

Speed up keyword lookup.

Improve search performance for multi-byte characters in position() and related functions.

Allow toasted values to be minimally decompressed.

This is useful for routines that only need to examine the initial portion of a toasted field.

Allow ALTER TABLE ... SET NOT NULL to avoid unnecessary table scans.

This can be optimized when the table's column constraints can be recognized as disallowing nulls.

Allow ALTER TABLE ... SET DATA TYPE changing between timestamp and timestamptz to avoid a table rewrite when the session time zone is UTC.

In the UTC time zone, these two data types are binary compatible.

Improve speed in converting strings to int2 or int4 integers.

Allow parallelized queries when in SERIALIZABLE isolation mode.

Previously, parallelism was disabled when in this mode.

Use pread() and pwrite() for random I/O.

This reduces the number of system calls required for I/O.

Improve the speed of setting the process title on FreeBSD.

Monitoring

Allow logging of statements from only a percentage of transactions.

The parameter log_transaction_sample_rate controls this.

Add progress reporting to CREATE INDEX and REINDEX operations.

Progress is reported in the pg_stat_progress_create_index system view.

Add progress reporting to CLUSTER and VACUUM FULL.

Progress is reported in the pg_stat_progress_cluster system view.

Add progress reporting to pg_checksums.

This is enabled with the option --progress.

Add counter of checksum failures to pg_stat_database.

Add tracking of global objects in system view pg_stat_database.

Global objects are shown with a pg_stat_database.datid value of zero.

Add the ability to list the contents of the archive directory.

The function is pg_ls_archive_statusdir().

Add the ability to list the contents of temporary directories.

The function, pg_ls_tmpdir(), optionally allows specification of a tablespace.

Add information about the client certificate to the system view pg_stat_ssl.

The new columns are client_serial and issuer_dn. Column clientdn has been renamed to client_dn for clarity.

Restrict visibility of rows in pg_stat_ssl for unprivileged users.

At server start, emit a log message including the server version number.

Prevent logging “incomplete startup packet” if a new connection is immediately closed.

This avoids log spam from certain forms of monitoring.

Include the application_name, if set, in log_connections log messages.

Make the walreceiver set its application name to the cluster name, if set.

Add the timestamp of the last received standby message to pg_stat_replication.

Add a wait event for fsync of WAL segments.

Authentication

Add GSSAPI encryption support.

This feature allows TCP/IP connections to be encrypted when using GSSAPI authentication, without having to set up a separate encryption facility such as SSL. In support of this, add hostgssenc and hostnogssenc record types in pg_hba.conf for selecting connections that do or do not use GSSAPI encryption, corresponding to the existing hostssl and hostnossal record types. There is also a new gssencmode libpq option, and a pg_stat_gssapi system view.

Allow the clientcert pg_hba.conf option to check that the database user name matches the client certificate's common name.

This new check is enabled with clientcert=verify-full.

Allow discovery of an LDAP server using DNS SRV records.

This avoids the requirement of specifying ldapserver. It is only supported if PostgreSQL is compiled with OpenLDAP.

Server Configuration

Add ability to enable/disable cluster checksums using pg_checksums.

The cluster must be shut down for these operations.

Reduce the default value of autovacuum_vacuum_cost_delay to 2ms.

This allows autovacuum operations to proceed faster by default.

Allow vacuum_cost_delay to specify sub-millisecond delays, by accepting fractional values.

Allow time-based server parameters to use units of microseconds (us).

Allow fractional input for integer server parameters.

For example, SET work_mem = '30.1GB' is now allowed, even though work_mem is an integer parameter. The value will be rounded to an integer after any required units conversion.

Allow units to be defined for floating-point server parameters.

Add wal_recycle and wal_init_zero server parameters to control WAL file recycling.

Avoiding file recycling can be beneficial on copy-on-write file systems like ZFS.

Add server parameter tcp_user_timeout to control the server's TCP timeout.

Allow control of the minimum and maximum SSL protocol versions.

The server parameters are ssl_min_protocol_version and ssl_max_protocol_version.

Add server parameter ssl_library to report the SSL library version used by the server.

Add server parameter shared_memory_type to control the type of shared memory to use.

This allows selection of System V shared memory, if desired.

Streaming Replication And Recovery

Allow some recovery parameters to be changed with reload.

These parameters are archive_cleanup_command, promote_trigger_file, recovery_end_command, and recovery_min_apply_delay.

Allow the streaming replication timeout (wal_sender_timeout) to be set per connection.

Previously, this could only be set cluster-wide.

Add function pg_promote() to promote standbys to primaries.

Previously, this operation was only possible by using pg_ctl or creating a trigger file.

Allow replication slots to be copied.

The functions for this are pg_copy_physical_replication_slot() and pg_copy_logical_replication_slot().

Make max_wal_senders not count as part of max_connections.

Add an explicit value of current for recovery_target_timeline.

Make recovery fail if a two-phase transaction status file is corrupt.

Previously, a warning was logged and recovery continued, allowing the transaction to be lost.

Utility Commands

Add REINDEX CONCURRENTLY option to allow reindexing without locking out writes.

This is also controlled by the reindexdb application's --concurrently option.

Add support for generated columns.

The content of generated columns are computed from expressions (including references to other columns in the same table) rather than being specified by INSERT or UPDATE commands.

Add a WHERE clause to COPY FROM to control which rows are accepted.

This provides a simple way to filter incoming data.

Allow enumerated values to be added more flexibly.

Previously, ALTER TYPE ... ADD VALUE could not be called in a transaction block, unless it was part of the same transaction that created the enumerated type. Now it can be called in a later transaction, so long as the new enumerated value is not referenced until after it is committed.

Add commands to end a transaction and start a new one.

The commands are COMMIT AND CHAIN and ROLLBACK AND CHAIN.

Add VACUUM and CREATE TABLE options to prevent VACUUM from truncating trailing empty pages.

These options are vacuum_truncate and toast.vacuum_truncate. Use of these options reduces VACUUM's locking requirements, but prevents returning disk space to the operating system.

Allow VACUUM to skip index cleanup.

This change adds a VACUUM command option INDEX_CLEANUP as well as a table storage option vacuum_index_cleanup. Use of this option reduces the ability to reclaim space and can lead to index bloat, but it is helpful when the main goal is to freeze old tuples.

Add the ability to skip VACUUM and ANALYZE operations on tables that cannot be locked immediately.

This option is called SKIP_LOCKED.

Allow VACUUM and ANALYZE to take optional Boolean argument specifications.

Prevent TRUNCATE, VACUUM and ANALYZE from requesting a lock on tables for which the user lacks permission.

This prevents unauthorized locking, which could interfere with user queries.

Add EXPLAIN option SETTINGS to output non-default optimizer settings.

This output can also be obtained when using auto_explain by setting auto_explain.log_settings.

Add OR REPLACE option to CREATE AGGREGATE.

Allow modifications of system catalogs' options using ALTER TABLE.

Modifications of catalogs' reloptions and autovacuum settings are now supported. (Setting allow_system_table_mods is still required.)

Use all key columns' names when selecting default constraint names for foreign keys.

Previously, only the first column name was included in the constraint name, resulting in ambiguity for multi-column foreign keys.

Data Types

Update assorted knowledge about Unicode to match Unicode 12.1.0.

This fixes, for example, cases where psql would misformat output involving combining characters.

Update Snowball stemmer dictionaries with support for new languages.

This adds word stemming support for Arabic, Indonesian, Irish, Lithuanian, Nepali, and Tamil to full text search.

Allow creation of collations that report string equality for strings that are not bit-wise equal.

This feature supports “nondeterministic” collations that can define case- and accent-agnostic equality comparisons. Thus, for example, a case-insensitive uniqueness constraint on a text column can be made more easily than before. This is only supported for ICU collations.

Add support for ICU collation attributes on older ICU versions.

This allows customization of the collation rules in a consistent way across all ICU versions.

Allow data type name to more seamlessly be compared to other text types.

Type name now behaves much like a domain over type text that has default collation “C”. This allows cross-type comparisons to be processed more efficiently.

Functions

Add support for the SQL/JSON path language.

This allows execution of complex queries on JSON values using an SQL-standard language.

Add support for hyperbolic functions.

Also add log10() as an alias for log(), for standards compliance.

Improve the accuracy of statistical aggregates like variance() by using more precise algorithms.

Allow date_trunc() to have an additional argument to control the time zone.

This is faster and simpler than using the AT TIME ZONE clause.

Adjust to_timestamp()/to_date() functions to be more forgiving of template mismatches.

This new behavior more closely matches the Oracle functions of the same name.

Fix assorted bugs in XML functions.

Specifically, in XMLTABLE, xpath(), and xmlexists(), fix some cases where nothing was output for a node, or an unexpected error was thrown, or necessary escaping of XML special characters was omitted.

Allow the BY VALUE clause in XMLEXISTS and XMLTABLE.

This SQL-standard clause has no effect in PostgreSQL's implementation, but it was unnecessarily being rejected.

Prevent current_schema() and current_schemas() from being run by parallel workers, as they are not parallel-safe.

Allow RECORD and RECORD[] to be used as column types in a query's column definition list for a table function that is declared to return RECORD.

PL/PgSQL

Allow SQL commands and variables with the same names as those commands to be used in the same PL/pgSQL function.

For example, allow a variable called comment to exist in a function that calls the COMMENT SQL command. Previously this combination caused a parse error.

Add new optional warning and error checks to PL/pgSQL.

The new checks allow for run-time validation of INTO column counts and single-row results.

Client Interfaces

Add connection parameter `tcp_user_timeout` to control libpq's TCP timeout.

Allow libpq (and thus psql) to report only the SQLSTATE value in error messages.

Add libpq function `PQresultMemorySize()` to report the memory used by a query result.

Remove the no-display/debug flag from libpq's options connection parameter.

This allows this parameter to be set by `postgres_fdw`.

Allow ecpg to create variables of data type bytea.

This allows ECPG clients to interact with bytea data directly, rather than using an encoded form.

Add PREPARE AS support to ECPG.

Client Applications

Allow vacuumdb to select tables for vacuum based on their wraparound horizon.

The options are `--min-xid-age` and `--min-mxid-age`.

Allow vacuumdb to disable waiting for locks or skipping all-visible pages.

The options are `--skip-locked` and `--disable-page-skipping`.

Add colorization to the output of command-line utilities.

This is enabled by setting the environment variable `PG_COLOR` to always or auto. The specific colors used can be adjusted by setting the environment variable `PG_COLORS`, using ANSI escape codes for colors. For example, the default behavior is equivalent to `PG_COLORS="error=01;31:warning=01;35:locus=01"`.

psql

Add CSV table output mode in psql.

This is controlled by `\pset format csv` or the command-line `--csv` option.

Show the manual page URL in psql's `\help` output for a SQL command.

Display the IP address in psql's `\conninfo` (Fabien Coelho)

Improve tab completion of CREATE TABLE, CREATE TRIGGER, CREATE EVENT TRIGGER, ANALYZE, EXPLAIN, VACUUM, ALTER TABLE, ALTER INDEX, ALTER DATABASE, and ALTER INDEX ALTER COLUMN.

pgbench

Allow values produced by queries to be assigned to pgbench variables.

The command for this is `\gset`.

Improve precision of pgbench's --rate option.

Improve pgbench's error reporting with clearer messages and return codes.

Server Applications

Allow control of log file rotation via pg_ctl.

Previously, this was only possible via an SQL function or a process signal.

Properly detach the new server process during pg_ctl start.

This prevents the server from being shut down if the shell script that invoked pg_ctl is interrupted later.

Allow pg_upgrade to use the file system's cloning feature, if there is one.

The --clone option has the advantages of --link, while preventing the old cluster from being changed after the new cluster has started.

Allow specification of the socket directory to use in pg_upgrade.

This is controlled by --socketdir; the default is the current directory.

Allow pg_checksums to disable fsync operations.

This is controlled by the --no-sync option.

Allow pg_rewind to disable fsync operations.

Fix pg_test_fsync to report accurate open_datasync durations on Windows.

pg_dump, pg_dumpall, pg_restore

When pg_dump emits data with INSERT commands rather than COPY, allow more than one data row to be included in each INSERT.

The option controlling this is --rows-per-insert.

Allow pg_dump to emit INSERT... ON CONFLICT DO NOTHING.

This avoids conflict failures during restore. The option is --on-conflict-do-nothing.

Decouple the order of operations in a parallel pg_dump from the order used by a subsequent parallel pg_restore.

This allows pg_restore to perform more-fully-parallelized parallel restores, especially in cases where the original dump was not done in parallel. Scheduling of a parallel pg_dump is also somewhat improved.

Allow the extra_float_digits setting to be specified for pg_dump and pg_dumpall.

This is primarily useful for making dumps that are exactly comparable across different source server versions. It is not recommended for normal use, as it may result in loss of precision when the dump is restored.

Add --exclude-database option to pg_dumpall.

Source Code

Add CREATE ACCESS METHOD command to create new table types.

This enables the development of new table access methods, which can optimize storage for different use cases. The existing heap access method remains the default.

Add planner support function interfaces to improve optimizer estimates, inlining, and indexing for functions.

This allows extensions to create planner support functions that can provide function-specific selectivity, cost, and row-count estimates that can depend on the function's arguments. Support functions can also supply simplified representations and index conditions, greatly expanding optimization possibilities.

Simplify renumbering manually-assigned OIDs, and establish a new project policy for management of such OIDs.

Patches that manually assign OIDs for new built-in objects (such as new functions) should now randomly choose OIDs in the range 8000—9999. At the end of a development cycle, the OIDs used by committed patches will be renumbered down to lower numbers, currently somewhere in the 4xxx range, using the new `renumber_oids.pl` script. This approach should greatly reduce the odds of OID collisions between different in-process patches.

While there is no specific policy reserving any OIDs for external use, it is recommended that forks and other projects needing private manually-assigned OIDs use numbers in the high 7xxx range. This will avoid conflicts with recently-merged patches, and it should be a long time before the core project reaches that range.

Build Cygwin binaries using dynamic instead of static libraries.

Remove configure switch `--disable-strong-random`.

A strong random-number source is now required.

`printf`-family functions, as well as `strerror` and `strerror_r`, now behave uniformly across platforms within Postgres code.

Notably, `printf` understands `%m` everywhere; on Windows, `strerror` copes with Winsock error codes (it used to do so in backend but not frontend code); and `strerror_r` always follows the GNU return convention.

Require a C99-compliant compiler, and MSVC 2013 or later on Windows.

Use pandoc, not lynx, for generating plain-text documentation output files.

This affects only the `INSTALL` file generated during `make dist` and the seldom-used plain-text `postgres.txt` output file. Pandoc produces better output than lynx and avoids some locale/encoding issues. Pandoc version 1.13 or later is required.

Support use of images in the PostgreSQL documentation.

Additional Modules

Allow `ORDER BY` sorts and `LIMIT` clauses to be pushed to `postgres_fdw` foreign servers in more cases.

Improve optimizer cost accounting for `postgres_fdw` queries.

Properly honor `WITH CHECK OPTION` on views that reference `postgres_fdw` tables.

While `CHECK OPTIONS` on `postgres_fdw` tables are ignored (because the reference is foreign), views on such tables are considered local, so this change enforces `CHECK OPTIONS` on them. Previously, only `INSERTs` and `UPDATEs` with `RETURNING` clauses that returned `CHECK OPTION` values were validated.

How to Report Problems

To report any issues you are having please contact EnterpriseDB's technical support staff:

- Email: support@enterprisedb.com
- Phone: +1-732-331-1320 or 1-800-235-5891 (US Only)

4 Advanced Server Upgrade Guide

The EDB Postgres Advanced Server Upgrade Guide is a comprehensive guide about upgrading EDB Postgres Advanced Server (Advanced Server). In this guide you will find detailed information about using:

- `pg_upgrade` to upgrade from an earlier version of Advanced Server to Advanced Server 13.
- `yum` to perform a minor version upgrade on a Linux host.
- `StackBuilder Plus` to perform a minor version upgrade on a Windows host.

4.1 Supported Platforms

For information about the platforms and versions supported by Advanced Server, visit the EnterpriseDB website at:

<https://www.enterprisedb.com/services-support/edb-supported-products-and-platforms#epas>

4.2 Limitations

- The `data` directory of a production database should not be stored on an NFS file system.
- The `pg_upgrade` utility cannot upgrade a partitioned table if a foreign key refers to the partitioned table.
- If you are upgrading from the version 9.4 server or a lower version of Advanced Server, and you use partitioned tables that include a `SUBPARTITION BY` clause, you must use `pg_dump` and `pg_restore` to upgrade an existing Advanced Server installation to a later version of Advanced Server. To upgrade, you must:
 1. Use `pg_dump` to preserve the content of the subpartitioned table.
 2. Drop the table from the Advanced Server 9.4 database or a lower version of Advanced Server database.
 3. Use `pg_upgrade` to upgrade the rest of the Advanced Server database to a more recent version.
 4. Use `pg_restore` to restore the subpartitioned table to the latest upgraded Advanced Server database.
- If you perform an upgrade of the Advanced Server installation, you must rebuild any hash-partitioned table on the upgraded server.

4.3 Upgrading an Installation With `pg_upgrade`

While minor upgrades between versions are fairly simple and require only the installation of new executables, past major version upgrades has been both expensive and time consuming. `pg_upgrade` facilitates migration between any version of Advanced Server (version 9.0 or later), and any subsequent release of Advanced Server that is supported on the same platform.

Without `pg_upgrade`, to migrate from an earlier version of Advanced Server to Advanced Server 13, you must

export all of your data using `pg_dump`, install the new release, run `initdb` to create a new cluster, and then import your old data.

`pg_upgrade` can reduce both the amount of time required and the disk space required for many major-version upgrades.

The `pg_upgrade` utility performs an in-place transfer of existing data between Advanced Server and any subsequent version.

Several factors determine if an in-place upgrade is practical:

- The on-disk representation of user-defined tables must not change between the original version and the upgraded version.
- The on-disk representation of data types must not change between the original version and the upgraded version.
- To upgrade between major versions of Advanced Server with `pg_upgrade`, both versions must share a common binary representation for each data type. Therefore, you cannot use `pg_upgrade` to migrate from a 32-bit to a 64-bit Linux platform.

Before performing a version upgrade, `pg_upgrade` will verify that the two clusters (the old cluster and the new cluster) are compatible.

If the upgrade involves a change in the on-disk representation of database objects or data, or involves a change in the binary representation of data types, `pg_upgrade` will be unable to perform the upgrade; to upgrade, you will have to `pg_dump` the old data and then import that data into the new cluster.

The `pg_upgrade` executable is distributed with Advanced Server 13, and is installed as part of the `Database Server` component; no additional installation or configuration steps are required.

4.3.1 Performing an Upgrade

To upgrade an earlier version of Advanced Server to the current version, you must:

- Install the current version of Advanced Server. The new installation must contain the same supporting server components as the old installation.
- Empty the target database or create a new target cluster with `initdb`.
- Place the `pg_hba.conf` file for both databases in `trust` authentication mode (to avoid authentication conflicts).
- Shut down the old and new Advanced Server services.
- Invoke the `pg_upgrade` utility.

When `pg_upgrade` starts, it performs a compatibility check to ensure that all required executables are present and contain the expected version numbers. The verification process also checks the old and new `$PGDATA` directories to ensure that the expected files and subdirectories are in place. If the verification process succeeds, `pg_upgrade` starts the old `postmaster` and runs `pg_dumpall --schema-only` to capture the metadata contained in the old cluster. The script produced by `pg_dumpall` is used in a later step to recreate all user-defined objects in the new cluster.

Note that the script produced by `pg_dumpall` recreates only user-defined objects and not system-defined objects. The new cluster will *already* contain the system-defined objects created by the latest version of Advanced Server.

After extracting the metadata from the old cluster, `pg_upgrade` performs the bookkeeping tasks required to sync the new cluster with the existing data.

`pg_upgrade` runs the `pg_dumpall` script against the new cluster to create (empty) database objects of the same shape and type as those found in the old cluster. Then, `pg_upgrade` links or copies each table and index from the old cluster to the new cluster.

If you are upgrading from a version of Advanced Server prior to 9.5 to Advanced Server 13 and have installed the `edb_dblink_oci` or `edb_dblink_libpq` extension, you must drop the extension before performing an upgrade. To drop the extension, connect to the server with the `psql` or `PEM` client, and invoke the commands:

```
DROP EXTENSION edb_dblink_oci;
DROP EXTENSION edb_dblink_libpq;
```

When you have completed upgrading, you can use the `CREATE EXTENSION` command to add the current versions of the extensions to your installation.

4.3.1.1 Linking versus Copying

When invoking `pg_upgrade`, you can use a command-line option to specify whether `pg_upgrade` should *copy* or *link* each table and index in the old cluster to the new cluster.

Linking is much faster because `pg_upgrade` simply creates a second name (a hard link) for each file in the cluster; linking also requires no extra workspace because `pg_upgrade` does not make a copy of the original data. When linking the old cluster and the new cluster, the old and new clusters share the data; note that after starting the new cluster, your data can no longer be used with the previous version of Advanced Server.

If you choose to copy data from the old cluster to the new cluster, `pg_upgrade` will still reduce the amount of time required to perform an upgrade compared to the traditional `dump/restore` procedure. `pg_upgrade` uses a file-at-a-time mechanism to copy data files from the old cluster to the new cluster (versus the row-by-row mechanism used by `dump/restore`). When you use `pg_upgrade`, you avoid building indexes in the new cluster; each index is simply copied from the old cluster to the new cluster. Finally, using a `dump/restore` procedure to upgrade requires a great deal of workspace to hold the intermediate text-based dump of all of your data, while `pg_upgrade` requires very little extra workspace.

Data that is stored in user-defined tablespaces is not copied to the new cluster; it stays in the same location in the file system, but is copied into a subdirectory whose name reflects the version number of the new cluster. To manually relocate files that are stored in a tablespace after upgrading, move the files to the new location and update the symbolic links (located in the `pg_tblspc` directory under your cluster's `data` directory) to point to the files.

4.3.2 Invoking pg_upgrade

When invoking `pg_upgrade`, you must specify the location of the old and new cluster's `PGDATA` and executable (`/bin`) directories, as well as the name of the Advanced Server superuser, and the ports on which the installations are listening. A typical call to invoke `pg_upgrade` to migrate from Advanced Server 12 to Advanced Server 13 takes the form:

```
pg_upgrade
--old-datadir <path_to_12_data_directory>
--new-datadir <path_to_13_data_directory>
--user <superuser_name>
--old-bindir <path_to_12_bin_directory>
--new-bindir <path_to_13_bin_directory>
--old-port <12_port> --new-port <13_port>
```

Where:

`--old-datadir <path_to_12_data_directory>`

Use the `--old-datadir` option to specify the complete path to the `data` directory within the Advanced Server 12 installation.

`--new-datadir <path_to_13_data_directory>`

Use the `--new-datadir` option to specify the complete path to the `data` directory within the Advanced Server 13 installation.

`--username <superuser_name>`

Include the `--username` option to specify the name of the Advanced Server superuser. The superuser name should be the same in both versions of Advanced Server. By default, when Advanced Server is installed in Oracle mode, the superuser is named `enterprisedb`. If installed in PostgreSQL mode, the superuser is named `postgres`.

If the Advanced Server superuser name is not the same in both clusters, the clusters will not pass the `pg_upgrade` consistency check.

`--old-bindir <path_to_12_bin_directory>`

Use the `--old-bindir` option to specify the complete path to the `bin` directory in the Advanced Server 12 installation.

`--new-bindir <path_to_13_bin_directory>`

Use the `--new-bindir` option to specify the complete path to the `bin` directory in the Advanced Server 13 installation.

`--old-port <12_port>`

Include the `--old-port` option to specify the port on which Advanced Server 12 listens for connections.

`--new-port <13_port>`

Include the `--new-port` option to specify the port on which Advanced Server 13 listens for connections.

4.3.2.1 Command Line Options - Reference

`pg_upgrade` accepts the following command line options; each option is available in a long form or a short form:

`-b <path_to_old_bin_directory>`

`--old-bindir <path_to_old_bin_directory>`

Use the `-b` or `--old-bindir` keyword to specify the location of the old cluster's executable directory.

`-B <path_to_new_bin_directory>`

`--new-bindir <path_to_new_bin_directory>`

Use the `-B` or `--new-bindir` keyword to specify the location of the new cluster's executable directory.

-c**--check**

Include the **-c** or **--check** keyword to specify that **pg_upgrade** should perform a consistency check on the old and new cluster without performing a version upgrade.

-d <path_to_old_data_directory>**--old-datadir <path_to_old_data_directory>**

Use the **-d** or **--old-datadir** keyword to specify the location of the old cluster's **data** directory.

-D <path_to_new_data_directory>**--new-datadir <path_to_new_data_directory>**

Use the **-D** or **--new-datadir** keyword to specify the location of the new cluster's **data** directory.

Data that is stored in user-defined tablespaces is not copied to the new cluster; it stays in the same location in the file system, but is copied into a subdirectory whose name reflects the version number of the new cluster. To manually relocate files that are stored in a tablespace after upgrading, you must move the files to the new location and update the symbolic links (located in the **pg_tblspc** directory under your cluster's **data** directory) to point to the files.

-j**--jobs**

Include the **-j** or **--jobs** keyword to specify the number of simultaneous processes or threads to use during the upgrade.

-k**--link**

Include the **-k** or **--link** keyword to create a hard link from the new cluster to the old cluster. See [Linking versus Copying](#) for more information about using a symbolic link.

-o <options>**--old-options <options>**

Use the **-o** or **--old-options** keyword to specify options that will be passed to the old **postgres** command. Enclose options in single or double quotes to ensure that they are passed as a group.

-O <options>**--new-options <options>**

Use the **-O** or **--new-options** keyword to specify options to be passed to the new **postgres** command. Enclose options in single or double quotes to ensure that they are passed as a group.

-p <old_port_number>**--old-port <old_port_number>**

Include the **-p** or **--old-port** keyword to specify the port number of the Advanced Server installation that you are upgrading.

-P <new_port_number>

--new-port <new_port_number>

Include the `-P` or `--new-port` keyword to specify the port number of the new Advanced Server installation.

Note

If the original Advanced Server installation is using port number `5444` when you invoke the Advanced Server 13 installer, the installer will recommend using listener port `5445` for the new installation of Advanced Server.

-r**--retain**

During the upgrade process, `pg_upgrade` creates four append-only log files; when the upgrade is completed, `pg_upgrade` deletes these files. Include the `-r` or `--retain` option to specify that the server should retain the `pg_upgrade` log files.

-U <user_name>**--username <user_name>**

Include the `-U` or `--username` keyword to specify the name of the Advanced Server database superuser. The same superuser must exist in both clusters.

-v**--verbose**

Include the `-v` or `--verbose` keyword to enable verbose output during the upgrade process.

-V**--version**

Use the `-V` or `--version` keyword to display version information for `pg_upgrade`.

-?**-h****--help**

Use `-?, -h` or `--help` options to display `pg_upgrade` help information.

4.3.3 Upgrading to Advanced Server 13

You can use `pg_upgrade` to upgrade from an existing installation of Advanced Server into the cluster built by the Advanced Server 13 installer or into an alternate cluster created using the `initdb` command. In this section, we will provide the details of upgrading into the cluster provided by the installer.

The basic steps to perform an upgrade into an empty cluster created with the `initdb` command are the same as the steps to upgrade into the cluster created by the Advanced Server 13 installer, but you can omit Step 2 (*Empty the `edb` database*), and substitute the location of the alternate cluster when specifying a target cluster for the upgrade.

If a problem occurs during the upgrade process, you can revert to the previous version. See [Reverting to the old](#)

[cluster](#) Section for detailed information about this process.

You must be an operating system superuser or Windows Administrator to perform an Advanced Server upgrade.

Step 1 - Install the New Server

Install Advanced Server 13, specifying the same non-server components that were installed during the previous Advanced Server installation.

The new cluster and the old cluster must reside in different directories.

Step 2 - Empty the target database

The target cluster must not contain any data; you can create an empty cluster using the `initdb` command, or you can empty a database that was created during the installation of Advanced Server 13. If you have installed Advanced Server in PostgreSQL mode, the installer creates a single database named `postgres`; if you have installed Advanced Server in Oracle mode, it creates a database named `postgres` and a database named `edb`.

The easiest way to empty the target database is to drop the database and then create a new database. Before invoking the `DROP DATABASE` command, you must disconnect any users and halt any services that are currently using the database.

On Windows, navigate through the `Control Panel` to the `Services` manager; highlight each service in the `Services` list, and select `Stop`.

On Linux, open a terminal window, assume superuser privileges, and manually stop each service; for example, if you are on Linux 6.x, invoke the following command to stop the pgAgent service:

```
service edb-pgagent-13 stop
```

After stopping any services that are currently connected to Advanced Server, you can use the EDB-PSQL command line client to drop and create a database. When the client opens, connect to the `template1` database as the database superuser; if prompted, provide authentication information. Then, use the following command to drop your database:

```
DROP DATABASE <database_name>;
```

Where `<database_name>` is the name of the database.


Then, create an empty database based on the contents of the `template1` database.

```
CREATE DATABASE <database_name>;
```

Step 3 - Set both servers in trust mode

During the upgrade process, `pg_upgrade` will connect to the old and new servers several times; to make the connection process easier, you can edit the `pg_hba.conf` file, setting the authentication mode to `trust`. To modify the `pg_hba.conf` file, navigate through the `Start` menu to the `EDB Postgres` menu; to the `Advanced Server` menu, and open the `Expert Configuration` menu; select the `Edit pg_hba.conf` menu option to open the `pg_hba.conf` file.

You must allow trust authentication for the previous Advanced Server installation, and Advanced Server 13 servers. Edit the `pg_hba.conf` file for both installations of Advanced Server as shown in the following figure.



After editing each file, save the file and exit the editor.

If the system is required to maintain `md5` authentication mode during the upgrade process, you can specify user passwords for the database superuser in a password file (`pgpass.conf` on Windows, `.pgpass` on Linux). For more information about configuring a password file, see the [PostgreSQL Core Documentation](#).

Step 4 - Stop All Component Services and Servers

Before you invoke `pg_upgrade`, you must stop any services that belong to the original Advanced Server installation, Advanced Server 13, or the supporting components. This ensures that a service will not attempt to access either cluster during the upgrade process.

The services that are most likely to be running in your installation are:

| Service: | On Linux: | On Windows: |
|---|----------------|--|
| Postgres Plus Advanced Server 9.0 | ppas-9.0 | ppas-9.0 |
| Postgres Plus Advanced Server 9.1 | ppas-9.1 | ppas-9.1 |
| Postgres Plus Advanced Server 9.2 | ppas-9.2 | ppas-9.2 |
| Postgres Plus Advanced Server 9.3 | ppas-9.3 | ppas-9.3 |
| Postgres Plus Advanced Server 9.4 | ppas-9.4 | ppas-9.4 |
| Postgres Plus Advanced Server 9.5 | ppas-9.5 | ppas-9.5 |
| EnterpriseDB Postgres Advanced Server 9.6 | edb-as-9.6 | edb-as-9.6 |
| EnterpriseDB Postgres Advanced Server 10 | edb-as-10 | edb-as-10 |
| EnterpriseDB Postgres Advanced Server 11 | edb-as-11 | edb-as-11 |
| EnterpriseDB Postgres Advanced Server 12 | edb-as-13 | edb-as-13 |
| Advanced Server 9.0 Scheduling Agent | ppasAgent-90 | Postgres Plus Advanced Server 90 Scheduling Agent |
| Advanced Server 9.1 Scheduling Agent | ppasAgent-91 | Postgres Plus Advanced Server 91 Scheduling Agent |
| Advanced Server 9.2 Scheduling Agent | ppas-agent-9.2 | Postgres Plus Advanced Server 9.2 Scheduling Agent |
| Advanced Server 9.3 Scheduling Agent | ppas-agent-9.3 | Postgres Plus Advanced Server 9.3 Scheduling Agent |
| Advanced Server 9.4 Scheduling Agent | ppas-agent-9.4 | Postgres Plus Advanced Server 9.4 Scheduling Agent |
| Advanced Server 9.5 Scheduling Agent | ppas-agent-9.5 | Postgres Plus Advanced Server 9.5 Scheduling Agent |

| Service: | On Linux: | On Windows: |
|--|--|--|
| Advanced Server 9.6 Scheduling Agent (pgAgent) | edb-pgagent-9.6 | EnterpriseDB Postgres Advanced Server 9.6 Scheduling Agent |
| Infinite Cache 9.2 | ppas-infinitecache-9.2 | N/A |
| Infinite Cache 9.3 | ppas-infinitecache-9.3 | N/A |
| Infinite Cache 9.4 | ppas-infinitecache | N/A |
| Infinite Cache 9.5 | ppas-infinitecache | N/A |
| Infinite Cache 9.6 | edb-icache | N/A |
| Infinite Cache 10 | edb-icache | N/A |
| PgBouncer 9.0 | pgbouncer-90 | pgbouncer-90 |
| PgBouncer 9.1 | pgbouncer-91 | pgbouncer-91 |
| PgBouncer 9.2 | pgbouncer-9.2 | pgbouncer-9.2 |
| PgBouncer 9.3 | pgbouncer-9.3 | pgbouncer-9.3 |
| PgBouncer | Pgbouncer | Pgbouncer |
| PgBouncer 1.6 | ppas-pgbouncer-1.6 or ppas-pgbouncer16 | ppas-pgbouncer-1.6 |
| PgBouncer 1.7 | edb-pgbouncer-1.7 | edb-pgbouncer-1.7 |
| PgPool 9.2 | ppas-pgpool-9.2 | N/A |
| PgPool 9.3 | ppas-pgpool-9.3 | N/A |
| PgPool | ppas-pgpool | N/A |
| PgPool 3.4 | ppas-pgpool-3.4 or ppas-pgpool34 or | N/A |
| pgPool-II | edb-pgpool-3.5 | N/A |
| Slony 9.2 | ppas-replication-9.2 | ppas-replication-9.2 |
| Slony 9.3 | ppas-replication-9.3 | ppas-replication-9.3 |
| Slony 9.4 | ppas-replication-9.4 | ppas-replication-9.4 |
| Slony 9.5 | ppas-replication-9.5 | ppas-replication-9.5 |
| Slony 9.6 | edb-slony-replication-9.6 | edb-slony-replication-9.6 |
| xDB Publication Server 9.0 | edb-xdbpubserver-90 | Publication Service 90 |
| xDB Publication Server 9.1 | edb-xdbpubserver-91 | Publication Service 91 |
| xDB Subscription Server | edb-xdbsubserver-90 | Subscription Service 90 |
| xDB Subscription Server | edb-xdbsubserver-91 | Subscription Service 91 |
| EDB Replication Server v6.x | edb-xdbpubserver | Publication Service for xDB Replication Server |
| EDB Subscription Server v6.x | edb-xdbsubserver | Subscription Service for xDB Replication Server |

To stop a service on Windows:

Open the **Services** applet; highlight each Advanced Server or supporting component service displayed in the list, and select **Stop**.

To stop a service on Linux:

Open a terminal window and manually stop each service at the command line.

Step 5 for Linux only - Assume the identity of the cluster owner

If you are using Linux, assume the identity of the Advanced Server cluster owner. (The following example assumes Advanced Server was installed in the default, compatibility with Oracle database mode, thus assigning **enterprisedb** as the cluster owner. If installed in compatibility with PostgreSQL database mode, **postgres** is the cluster owner.)

```
| su - enterpriseedb
```

Enter the Advanced Server cluster owner password if prompted. Then, set the path to include the location of the `pg_upgrade` executable:

```
| export PATH=$PATH:/usr/edb/as13/bin
```

During the upgrade process, `pg_upgrade` writes a file to the current working directory of the `enterpriseedb` user; you must invoke `pg_upgrade` from a directory where the `enterpriseedb` user has `write` privileges. After performing the above commands, navigate to a directory in which the `enterpriseedb` user has sufficient privileges to write a file.

```
| cd /tmp
```

Proceed to Step 6.

Step 5 for Windows only - Assume the identity of the cluster owner

If you are using Windows, open a terminal window, assume the identity of the Advanced Server cluster owner and set the path to the `pg_upgrade` executable.

If the `--serviceaccount <service_account_user>` parameter was specified during the initial installation of Advanced Server, then `<service_account_user>` is the Advanced Server cluster owner and is the user to be given with the `RUNAS` command.

```
| RUNAS /USER:<service account user> "CMD.EXE"
SET PATH=%PATH%;C:\Program Files\edb\as13\bin
```

During the upgrade process, `pg_upgrade` writes a file to the current working directory of the service account user; you must invoke `pg_upgrade` from a directory where the service account user has `write` privileges. After performing the above commands, navigate to a directory in which the service account user has sufficient privileges to write a file.

```
| cd %TEMP%
```

Proceed to Step 6.

If the `--serviceaccount` parameter was omitted during the initial installation of Advanced Server, then the default owner of the Advanced Server service and the database cluster is `NT AUTHORITY\NetworkService`.

When `NT AUTHORITY\NetworkService` is the service account user, the `RUNAS` command may not be usable as it prompts for a password and the `NT AUTHORITY\NetworkService` account is not assigned a password. Thus, there is typically a failure with an error message such as, “Unable to acquire user password”.

Under this circumstance a Windows utility program named `PsExec` must be used to run `CMD.EXE` as the service account `NT AUTHORITY\NetworkService`.

The `PsExec` program must be obtained by downloading `PsTools`, which is available at the following site:

<https://technet.microsoft.com/en-us/sysinternals/bb897553.aspx>.

You can then use the following command to run `CMD.EXE` as `NT AUTHORITY\NetworkService`, and then set the path to the `pg_upgrade` executable.

```
| psexec.exe -u "NT AUTHORITY\NetworkService" CMD.EXE
SET PATH=%PATH%;C:\Program Files\edb\as13\bin
```

During the upgrade process, `pg_upgrade` writes a file to the current working directory of the service account user; you must invoke `pg_upgrade` from a directory where the service account user has `write` privileges. After performing the above commands, navigate to a directory in which the service account user has sufficient privileges to write a file.

```
| cd %TEMP%
```

Proceed with Step 6.

Step 6 - Perform a consistency check

Before attempting an upgrade, perform a consistency check to assure that the old and new clusters are compatible and properly configured. Include the `--check` option to instruct `pg_upgrade` to perform the consistency check.

The following example demonstrates invoking `pg_upgrade` to perform a consistency check on Linux:

```
pg_upgrade -d /var/lib/edb/as12/data
-D /var/lib/edb/as13/data -U enterprise
-b /usr/edb/as12/bin -B /usr/edb/as13/bin -p 5444 -P 5445 --check
```

If the command is successful, it will return `*Clusters are compatible*`.

If you are using Windows, you must quote any directory names that contain a space:

```
pg_upgrade.exe
-d "C:\Program Files\ PostgresPlus\12AS \data"
-D "C:\Program Files\edb\as13\data" -U enterprise
-b "C:\Program Files\PostgresPlus\12AS\bin"
-B "C:\Program Files\edb\as13\bin" -p 5444 -P 5445 --check
```

During the consistency checking process, `pg_upgrade` will log any discrepancies that it finds to a file located in the directory from which `pg_upgrade` was invoked. When the consistency check completes, review the file to identify any missing components or upgrade conflicts. You must resolve any conflicts before invoking `pg_upgrade` to perform a version upgrade.

If `pg_upgrade` alerts you to a missing component, you can use StackBuilder Plus to add the component that contains the component. Before using StackBuilder Plus, you must restart the Advanced Server 13 service. After restarting the service, open StackBuilder Plus by navigating through the `Start` menu to the `Advanced Server 13` menu, and selecting `StackBuilder Plus`. Follow the onscreen advice of the StackBuilder Plus wizard to download and install the missing components.

When `pg_upgrade` has confirmed that the clusters are compatible, you can perform a version upgrade.

Step 7 - Run pg_upgrade

After confirming that the clusters are compatible, you can invoke `pg_upgrade` to upgrade the old cluster to the new version of Advanced Server.

On Linux:

```
pg_upgrade -d /var/lib/edb/as12/data
-D /var/lib/edb/as13/data -U enterprise
-b /usr/edb/as12/bin -B /usr/edb/as13/bin -p 5444 -P 5445
```

On Windows:

```
pg_upgrade.exe -d "C:\Program Files\PostgresPlus\12AS\data"
-D "C:\Program Files\edb\as13\data" -U enterprise
-b "C:\Program Files\PostgresPlus\12AS\bin"
```

```
-B "C:\Program Files\edb\as13\bin" -p 5444 -P 5445
```

`pg_upgrade` will display the progress of the upgrade onscreen:

```
$ pg_upgrade -d /var/lib/edb/as12/data -D /var/lib/edb/as13/data -U
enterprisedb -b /usr/edb/as12/bin -B /usr/edb/as13/bin -p 5444 -P 5445
Performing Consistency Checks
```

| | |
|---|----|
| Checking current, bin, and data directories | ok |
| Checking cluster versions | ok |
| Checking database user is a superuser | ok |
| Checking for prepared transactions | ok |
| Checking for reg* system OID user data types | ok |
| Checking for contrib/isn with bigint-passing mismatch | ok |
| Creating catalog dump | ok |
| Checking for presence of required libraries | ok |
| Checking database user is a superuser | ok |
| Checking for prepared transactions | ok |

If `pg_upgrade` fails after this point, you must re-initdb the new cluster before continuing.

Performing Upgrade

| | |
|---|----|
| Analyzing all rows in the new cluster | ok |
| Freezing all rows on the new cluster | ok |
| Deleting files from new pg_clog | ok |
| Copying old pg_clog to new server | ok |
| Setting next transaction ID for new cluster | ok |
| Resetting WAL archives | ok |
| Setting frozenid counters in new cluster | ok |
| Creating databases in the new cluster | ok |
| Adding support functions to new cluster | ok |
| Restoring database schema to new cluster | ok |
| Removing support functions from new cluster | ok |
| Copying user relation files | ok |
| Setting next OID for new cluster | ok |
| Creating script to analyze new cluster | ok |
| Creating script to delete old cluster | ok |

Upgrade Complete

Optimizer statistics are not transferred by `pg_upgrade` so, once you start the new server, consider running:
`analyze_new_cluster.sh`

Running this script will delete the old cluster's data files:
`delete_old_cluster.sh`

While `pg_upgrade` runs, it may generate SQL scripts that handle special circumstances that it has encountered during your upgrade. For example, if the old cluster contains large objects, you may need to invoke a script that defines the default permissions for the objects in the new cluster. When performing the pre-upgrade consistency check `pg_upgrade` will alert you to any script that you may be required to run manually.

You must invoke the scripts after `pg_upgrade` completes. To invoke the scripts, connect to the new cluster as a

database superuser with the EDB-PSQL command line client, and invoke each script using the `\i` option:

```
\i <complete_path_to_script/script.sql>
```

It is generally unsafe to access tables referenced in rebuild scripts until the rebuild scripts have completed; accessing the tables could yield incorrect results or poor performance. Tables not referenced in rebuild scripts can be accessed immediately.

If `pg_upgrade` fails to complete the upgrade process, the old cluster will be unchanged, except that `$PGDATA/global/pg_control` is renamed to `pg_control.old` and each tablespace is renamed to `tablespace.old`. To revert to the pre-invocation state:

1. Delete any tablespace directories created by the new cluster.
2. Rename `$PGDATA/global/pg_control`, removing the `.old` suffix.
3. Rename the old cluster tablespace directory names, removing the `.old` suffix.
4. Remove any database objects (from the new cluster) that may have been moved before the upgrade failed.

After performing these steps, resolve any upgrade conflicts encountered before attempting the upgrade again.

When the upgrade is complete, `pg_upgrade` may also recommend vacuuming the new cluster, and will provide a script that allows you to delete the old cluster.

Note

Before removing the old cluster, ensure that the cluster has been upgraded as expected, and that you have preserved a backup of the cluster in case you need to revert to a previous version.

Step 8 - Restore the authentication settings in the pg_hba.conf file

If you modified the `pg_hba.conf` file to permit `trust` authentication, update the contents of the `pg_hba.conf` file to reflect your preferred authentication settings.

Step 9 - Move and identify user-defined tablespaces (Optional)

If you have data stored in a user-defined tablespace, you must manually relocate tablespace files after upgrading; move the files to the new location and update the symbolic links (located in the `pg_tblspc` directory under your cluster's `data` directory) to point to the files.

4.3.4 Upgrading a pgAgent Installation

If your existing Advanced Server installation uses pgAgent, you can use a script provided with the Advanced Server 13 installer to update pgAgent. The script is named `dbms_job.upgrade.script.sql`, and is located in the `/share/contrib/` directory under your Advanced Server installation.

If you are using `pg_upgrade` to upgrade your installation, you should:

1. Perform the upgrade.
2. Invoke the `dbms_job.upgrade.script.sql` script to update the catalog files. If your existing pgAgent installation was performed with a script, the update will convert the installation to an extension.

4.3.5 pg_upgrade Troubleshooting

The troubleshooting tips in this section address problems you may encounter when using `pg_upgrade`.

Upgrade Error - There seems to be a postmaster servicing the cluster

If `pg_upgrade` reports that a postmaster is servicing the cluster, please stop all Advanced Server services and try the upgrade again.

Upgrade Error - fe_sendauth: no password supplied

If `pg_upgrade` reports an authentication error that references a missing password, please modify the `pg_hba.conf` files in the old and new cluster to enable `trust` authentication, or configure the system to use a `pgpass.conf` file.

Upgrade Error - New cluster is not empty; exiting

If `pg_upgrade` reports that the new cluster is not empty, please empty the new cluster. The target cluster may not contain any user-defined databases.

Upgrade Error - Failed to load library

If the original Advanced Server cluster included libraries that are not included in the Advanced Server 13 cluster, `pg_upgrade` will alert you to the missing component during the consistency check by writing an entry to the `loadable_libraries.txt` file in the directory from which you invoked `pg_upgrade`. Generally, for missing libraries that are not part of a major component upgrade, perform the following steps:

1. Restart the Advanced Server service.

Use StackBuilder Plus to download and install the missing module. Then:

2. Stop the Advanced Server service.
3. Resume the upgrade process: invoke `pg_upgrade` to perform consistency checking.
4. When you have resolved any remaining problems noted in the consistency checks, invoke `pg_upgrade` to perform the data migration from the old cluster to the new cluster.

4.3.6 Reverting to the Old Cluster

The method used to revert to a previous cluster varies with the options specified when invoking `pg_upgrade`:

- If you specified the `--check` option when invoking `pg_upgrade`, an upgrade has not been performed, and no modifications have been made to the old cluster; you can re-use the old cluster at any time.
- If you included the `--link` option when invoking `pg_upgrade`, the data files are shared between the old and new cluster after the upgrade completes. If you have started the server that is servicing the new cluster, the new server has written to those shared files and it is unsafe to use the old cluster.
- If you ran `pg_upgrade` without the `--link` specification or have not started the new server, the old cluster is unchanged, except that the `.old` suffix has been appended to the `$PGDATA/global/pg_control` and tablespace directories.

- To reuse the old cluster, delete the tablespace directories created by the new cluster and remove the `.old` suffix from `$PGDATA/global/pg_control` and the old cluster tablespace directory names and restart the server that services the old cluster.
-

4.4 Performing a Minor Version Update of an RPM Installation

If you used an RPM package to install Advanced Server or its supporting components, you can use `yum` to perform a minor version upgrade to a more recent version. To review a list of the package updates that are available for your system, open a command line, assume root privileges, and enter the command:

```
yum check-update <package_name>
```

Where `<package_name>` is the search term for which you wish to search for updates. Please note that you can include wild-card values in the search term. To use `yum update` to install an updated package, use the command:

```
yum update <package_name>
```

Where `<package_name>` is the name of the package you wish to update. Include wild-card values in the update command to update multiple related packages with a single command. For example, use the following command to update all packages whose names include the expression `edb`:

```
yum update edb*
```

Note

The `yum update` command will only perform an update between minor releases; to update between major releases, you must use `pg_upgrade`.

For more information about using yum commands and options, enter `yum --help` on your command line, or visit:

https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Deployment_Guide/ch-yum.html


4.5 Using StackBuilder Plus to Perform a Minor Version Update

The StackBuilder Plus utility provides a graphical interface that simplifies the process of updating, downloading, and installing modules that complement your Advanced Server installation. When you install a module with StackBuilder Plus, StackBuilder Plus automatically resolves any software dependencies.

Note

StackBuilder Plus is supported only on Windows systems.


You can invoke StackBuilder Plus at any time after the installation has completed by selecting the `StackBuilder Plus` menu option from the `Apps` menu. Enter your system password (if prompted), and the StackBuilder Plus welcome window opens (shown in the following figure).




Use the drop-down listbox on the welcome window to select your Advanced Server installation.

StackBuilder Plus requires Internet access; if your installation of Advanced Server resides behind a firewall (with restricted Internet access), StackBuilder Plus can download program installers through a proxy server. The module provider determines if the module can be accessed through an HTTP proxy or an FTP proxy; currently, all updates are transferred via an HTTP proxy and the FTP proxy information is not used.

If the selected Advanced Server installation has restricted Internet access, use the [Proxy Servers](#) on the [Welcome to StackBuilder Plus Window!](#) to open the [Proxy servers](#) dialog (shown in the following figure).





Enter the IP address and port number of the proxy server in the [HTTP proxy](#) on the [Proxy servers](#) dialog. Currently, all StackBuilder Plus modules are distributed via HTTP proxy (FTP proxy information is ignored). Click [OK](#) to continue.




The tree control on the StackBuilder Plus module selection window (shown in the following figure) displays a node for each module category. To perform an Advanced Server update, expand the [Database Server](#) module in the tree control and check the box to the left of the Advanced Server upgraded version. Then, click [Next](#).

If prompted, enter your email address and password on the StackBuilder Plus registration window (shown in the following figure).




StackBuilder Plus confirms the packages selected. The **Selected packages** dialog will display the name and version of the installer; click **Next** to continue.

When the download completes, a window opens that confirms the installation files have been downloaded and are ready for installation.



You can check the box next to **Skip Installation**, and select **Next** to exit StackBuilder Plus without installing the downloaded files, or leave the box unchecked and click **Next** to start the installation process.



When the upgrade is complete, StackBuilder Plus will alert you to the success or failure of the installation of the requested package. If you were prompted by an installer to restart your computer, reboot now.

Note

If the update fails to install, StackBuilder Plus will alert you to the installation error with a popup dialog and write a message to the log file at **%TEMP%**.

5 Database Compatibility for Oracle Developers Reference Guide

5.1 Introduction

Database Compatibility for Oracle means that an application runs in an Oracle environment as well as in the EDB Postgres Advanced Server (Advanced Server) environment with minimal or no changes to the application code.

This guide provides reference material about the compatible data types supported by Advanced Server. Reference information about:

- Compatible SQL Language syntax is provided in the *Database Compatibility for Oracle Developers SQL Guide*.
- Compatible Catalog Views is provided in the *Database Compatibility for Oracle Developers Catalog View Guide*.

Developing an application that is compatible with Oracle databases in the Advanced Server requires special attention to which features are used in the construction of the application. For example, developing a compatible application means selecting:

- Data types to define the application's database tables that are compatible with Oracle databases
- SQL statements that are compatible with Oracle SQL
- System and built-in functions for use in SQL statements and procedural logic that are compatible with Oracle databases
- Stored Procedure Language (SPL) to create database server-side application logic for stored procedures, functions, triggers, and packages
- System catalog views that are compatible with Oracle's data dictionary

For detailed information about Advanced Server's compatibility features and extended functionality, see the complete library of Advanced Server documentation, available at:

<https://www.enterprisedb.com/edb-docs>

5.2 The SQL Language

The following sections describe the subset of the Advanced Server SQL language compatible with Oracle databases. The following SQL syntax, data types, and functions work in both EDB Postgres Advanced Server and Oracle.

The Advanced Server documentation set includes syntax and commands for extended functionality (functionality that does not provide database compatibility for Oracle or support Oracle-styled applications) that is not included in this guide.

This section is organized into the following sections:

- General discussion of Advanced Server SQL syntax and language elements
 - Data types
 - Built-in functions
-

5.2.1 SQL Syntax

This section describes the general syntax of SQL. It forms the foundation for understanding the following chapters that include detail about how the SQL commands are applied to define and modify data.

5.2.1.1 Lexical Structure

SQL input consists of a sequence of commands. A **command** is composed of a sequence of **tokens**, terminated by a semicolon (`;`). The end of the input stream also terminates a command. Which tokens are valid depends on the syntax of the particular command.

A token can be a **key word**, an **identifier**, a **quoted identifier**, a **literal** (or **constant**), or a special character symbol. Tokens are normally separated by **whitespace** (space, tab, new line), but need not be if there is no ambiguity (which is generally only the case if a special character is adjacent to some other token type).

Additionally, **comments** can occur in SQL input. They are not tokens -they are effectively equivalent to whitespace.

For example, the following is (syntactically) valid SQL input:

```
SELECT * FROM MY_TABLE;
UPDATE MY_TABLE SET A = 5;
INSERT INTO MY_TABLE VALUES (3, 'hi there');
```

This is a sequence of three commands, one per line (although this is not required; more than one command can be on a line, and commands can usually be split across lines).

The SQL syntax is not very consistent regarding what tokens identify commands and which are operands or parameters. The first few tokens are generally the command name, so in the above example we would usually speak of a **SELECT**, an **UPDATE**, and an **INSERT** command. But for instance the **UPDATE** command always requires a **SET** token to appear in a certain position, and this particular variation of **INSERT** also requires a **VALUES** token in order to be complete. The precise syntax rules for each command are described in *Database Compatibility for Oracle Developers SQL Guide*.

5.2.1.2 Identifiers and Key Words

Tokens such as **SELECT**, **UPDATE**, or **VALUES** in the example above are examples of **key words**, that is, words that have a fixed meaning in the SQL language. The tokens **MY_TABLE** and **A** are examples of **identifiers**. They identify names of tables, columns, or other database objects, depending on the command they are used in. Therefore they are sometimes simply called, **names**. Key words and identifiers have the same **lexical structure**, meaning that one cannot know whether a token is an identifier or a key word without knowing the language.

SQL identifiers and key words must begin with a letter (**a-z** or **A-Z**). Subsequent characters in an identifier or key word can be letters, underscores, digits (**0-9**), dollar signs (\$), or number signs (#).

Identifier and key word names are case insensitive. Therefore

```
| UPDATE MY_TABLE SET A = 5;
```

can equivalently be written as:

```
| uPDAte my_TabLE SeT a = 5;
```

A convention often used is to write key words in upper case and names in lower case, e.g.,

```
| UPDATE my_table SET a = 5;
```

There is a second kind of identifier: the **delimited identifier** or **quoted identifier**. It is formed by enclosing an arbitrary sequence of characters in double-quotes (""). A delimited identifier is always an identifier, never a key word. So **"select"** could be used to refer to a column or table named **"select"**, whereas an unquoted select would be taken as a key word and would therefore provoke a parse error when used where a table or column name is expected. The example can be written with quoted identifiers like this:

```
| UPDATE "my_table" SET "a" = 5;
```

Quoted identifiers can contain any character, except the character with the numeric code zero.

To include a double quote, use two double quotes. This allows you to construct table or column names that would otherwise not be possible (such as ones containing spaces or ampersands). The length limitation still applies.

Quoting an identifier also makes it case-sensitive, whereas unquoted names are always folded to lower case. For example, the identifiers `FOO`, `foo`, and `"foo"` are considered the same by Advanced Server, but `"Foo"` and `"FOO"` are different from these three and each other. The folding of unquoted names to lower case is not compatible with Oracle databases. In Oracle syntax, unquoted names are folded to upper case: for example, `foo` is equivalent to `"FOO"` not `"foo"`. If you want to write portable applications you are advised to always quote a particular name or never quote it.

5.2.1.3 Constants

The kinds of implicitly-typed constants in Advanced Server are `strings` and `numbers`. Constants can also be specified with explicit types, which can enable more accurate representation and more efficient handling by the system. These alternatives are discussed in the following subsections.

String Constants

A `string constant` in SQL is an arbitrary sequence of characters bounded by single quotes `('')`, for example `'This is a string'`. To include a single-quote character within a string constant, write two adjacent single quotes, e.g. `'Dianne's horse'`. Note that this is not the same as a double-quote character `("")`.

Numeric Constants

Numeric constants are accepted in these general forms:

```
digits
digits.[digits][e[+-]digits]
[digits].digits[e[+-]digits]
digitse[+-]digits
```

where `digits` is one or more decimal digits (0 through 9). At least one digit must be before or after the decimal point, if one is used. At least one digit must follow the exponent marker `(e)`, if one is present. There may not be any spaces or other characters embedded in the constant. Note that any leading plus or minus sign is not actually considered part of the constant; it is an operator applied to the constant.

These are some examples of valid numeric constants:

```
42
3.5
4.
.001
5e2
1.925e-3
```

A numeric constant that contains neither a decimal point nor an exponent is initially presumed to be type

`INTEGER` if its value fits in type `INTEGER` (32 bits); otherwise it is presumed to be type `BIGINT` if its value fits in type `BIGINT` (64 bits); otherwise it is taken to be type `NUMBER`. Constants that contain decimal points and/or exponents are always initially presumed to be type `NUMBER`.

The initially assigned data type of a numeric constant is just a starting point for the type resolution algorithms. In most cases the constant will be automatically coerced to the most appropriate type depending on context. When necessary, you can force a numeric value to be interpreted as a specific data type by casting it as described in the following section.

Constants of Other Types

CAST

A constant of an arbitrary type can be entered using the following notation:

`CAST('string' AS type)`

The string constant's text is passed to the input conversion routine for the type called `type`. The result is a constant of the indicated type. The explicit type cast may be omitted if there is no ambiguity as to the type the constant must be (for example, when it is assigned directly to a table column), in which case it is automatically coerced.

`CAST` can also be used to specify runtime type conversions of arbitrary expressions.

CAST (MULTISET)

`MULTISET` is an extension to `CAST` that converts subquery results into a nested table type. The synopsis is:

`CAST (MULTISET (< subquery >) AS < datatype >)`

Where `subquery` is a query returning one or more rows and `datatype` is a nested table type.

`CAST(MULTISET)` is used to store a collection of data in a table.

Example

The following example demonstrates using `MULTISET`:

```
edb=# CREATE OR REPLACE TYPE project_table_t AS TABLE OF VARCHAR2(25);
CREATE TYPE
edb=# CREATE TABLE projects (person_id NUMBER(10), project_name VARCHAR2(20));
CREATE TABLE
edb=# CREATE TABLE pers_short (person_id NUMBER(10), last_name VARCHAR2(25));
CREATE TABLE

edb=# INSERT INTO projects VALUES (1, 'Teach');
INSERT 0 1
edb=# INSERT INTO projects VALUES (1, 'Code');
INSERT 0 1
edb=# INSERT INTO projects VALUES (2, 'Code');
INSERT 0 1
edb=# INSERT INTO pers_short VALUES (1, 'Morgan');
INSERT 0 1
edb=# INSERT INTO pers_short VALUES (2, 'Kolk');
INSERT 0 1
edb=# INSERT INTO pers_short VALUES (3, 'Scott');
INSERT 0 1
```

```

edb=# COMMIT;
COMMIT

edb=# SELECT e.last_name, CAST(MULTISET(
edb(#  SELECT p.project_name
edb(#  FROM projects p
edb(# WHERE p.person_id = e.person_id
edb(# ORDER BY p.project_name) AS project_table_t)
edb# FROM pers_short e;
last_name | project_table_t
-----+-----
Morgan   | {Code,Teach}
Kolk     | {Code}
Scott    | {}
(3 rows)

```

5.2.1.4 Comments

A comment is an arbitrary sequence of characters beginning with double dashes and extending to the end of the line, e.g.:

-- This is a standard SQL comment

Alternatively, C-style block comments can be used:

```

/* multiline comment
 * block
 */

```

where the comment begins with /* and extends to the matching occurrence of */.

5.2.2 Data Types

The following table shows the built-in general-purpose data types:

| Name | Alias | Description |
|-----------------------|----------------------------------|--|
| BLOB | LONG RAW, RAW(<i>n</i>), BYTEA | Binary data |
| BOOLEAN | | Logical Boolean (true/false) |
| CHAR [(<i>n</i>)] | CHARACTER [(<i>n</i>)] | Fixed-length character string of <i>n</i> characters |
| CLOB | LONG, LONG VARCHAR | Long character string |
| DATE | TIMESTAMP | Date and time to the second |

| | | |
|--|---|--|
| | FLOAT, | |
| DOUBLE PRECISION | FLOAT(25) – FLOAT(53) | Double precision floating-point number |
| INTEGER | INT, BINARY_INTEGER, PLS_INTEGER | Signed four-byte integer |
| NUMBER | DEC, DECIMAL, NUMERIC | Exact numeric with optional decimal places |
| NUMBER(p [, s]) | DEC(p [, s]), DECIMAL(p [, s]), NUMERIC(p [, s]) | Exact numeric of maximum precision, p, and optional scale, s |
| REAL TIMESTAMP [(p)] TIMESTAMP [(p)] WITH TIME ZONE | FLOAT(1) – FLOAT(24) | Single precision floating-point number Date and time with optional, fractional second precision, p Date and time with optional, fractional second precision, p, and with time zone |
| VARCHAR2(n) XMLTYPE | CHAR VARYING(n), CHARACTER VARYING(n), VARCHAR(n) | Variable-length character string with a maximum length of n characters XML data |

5.2.2.1 Numeric Types

Numeric types consist of four-byte integers, four-byte and eight-byte floating-point numbers, and fixed-precision decimals. The following table lists the available types:

| Name | Storage Size | Description | Range |
|------------------|--------------|--|---|
| BINARY_INTEGER | 4 bytes | Signed integer, Alias for INTEGER | -2,147,483,648 to +2,147,483,647 |
| DOUBLE PRECISION | 8 bytes | Variable-precision, inexact | 15 decimal digits precision |
| INTEGER | 4 bytes | Usual choice for integer | -2,147,483,648 to +2,147,483,647 |
| NUMBER | Variable | User-specified precision, exact | Up to 1000 digits of precision |
| NUMBER(p [, s]) | Variable | Exact numeric of maximum precision, p, and optional scale, s | Up to 1000 digits of precision |
| PLS_INTEGER | 4 bytes | Signed integer, Alias for INTEGER | -2,147,483,648 to +2,147,483,647 |
| REAL | 4 bytes | Variable-precision, inexact | 6 decimal digits precision |
| ROWID | 8 bytes | Signed 8 bit integer. | -9223372036854775808 to 9223372036854775807 |

The following sections describe the types in detail.

Integer Types

The `BINARY_INTEGER`, `INTEGER`, `PLS_INTEGER`, and `ROWID` types store whole numbers (without fractional components) as specified in table [Numeric Types](#). Attempts to store values outside of the allowed range will result in an error.

Arbitrary Precision Numbers

The type, `NUMBER`, can store practically an unlimited number of digits of precision and perform calculations exactly. It is especially recommended for storing monetary amounts and other quantities where exactness is required. However, the `NUMBER` type is very slow compared to the floating-point types described in the next section.

In what follows we use these terms: The `scale` of a `NUMBER` is the count of decimal digits in the fractional part, to the right of the decimal point. The `precision` of a `NUMBER` is the total count of significant digits in the whole number, that is, the number of digits to both sides of the decimal point. So the number 23.5141 has a precision of 6 and a scale of 4. Integers can be considered to have a scale of zero.

Both the precision and the scale of the `NUMBER` type can be configured. To declare a column of type `NUMBER` use the syntax

| `NUMBER(precision, scale)`

The precision must be positive, the scale zero or positive. Alternatively,

| `NUMBER(precision)`

selects a scale of 0. Specifying `NUMBER` without any precision or scale creates a column in which numeric values of any precision and scale can be stored, up to the implementation limit on precision. A column of this kind will not coerce input values to any particular scale, whereas `NUMBER` columns with a declared scale will coerce input values to that scale. (The SQL standard requires a default scale of 0, i.e., coercion to integer precision. For maximum portability, it is best to specify the precision and scale explicitly.)

If the precision or scale of a value is greater than the declared precision or scale of a column, the system will attempt to round the value. If the value cannot be rounded so as to satisfy the declared limits, an error is raised.

Floating-Point Types

The data types `REAL` and `DOUBLE PRECISION` are `inexact`, variable-precision numeric types. In practice, these types are usually implementations of IEEE Standard 754 for Binary Floating-Point Arithmetic (single and double precision, respectively), to the extent that the underlying processor, operating system, and compiler support it.

Inexact means that some values cannot be converted exactly to the internal format and are stored as approximations, so that storing and printing back out a value may show slight discrepancies. Managing these errors and how they propagate through calculations is the subject of an entire branch of mathematics and computer science and will not be discussed further here, except for the following points:

If you require exact storage and calculations (such as for monetary amounts), use the `NUMBER` type instead.

If you want to do complicated calculations with these types for anything important, especially if you rely on certain behavior in boundary cases (infinity, underflow), you should evaluate the implementation carefully.

Comparing two floating-point values for equality may or may not work as expected.

On most platforms, the `REAL` type has a range of at least `1E-37` to `1E+37` with a precision of at least 6 decimal digits. The `DOUBLE PRECISION` type typically has a range of around `1E-307` to `1E+308` with a precision of at

least 15 digits. Values that are too large or too small will cause an error. Rounding may take place if the precision of an input number is too high. Numbers too close to zero that are not representable as distinct from zero will cause an underflow error.

Advanced Server also supports the SQL standard notations `FLOAT` and `FLOAT(p)` for specifying inexact numeric types. Here, `p` specifies the minimum acceptable precision in binary digits. Advanced Server accepts `FLOAT(1)` to `FLOAT(24)` as selecting the `REAL` type, while `FLOAT(25)` to `FLOAT(53)` as selecting `DOUBLE PRECISION`. Values of `p` outside the allowed range draw an error. `FLOAT` with no precision specified is taken to mean `DOUBLE PRECISION`.

5.2.2.2 Character Types

The following table lists the general-purpose character types available in Advanced Server:

| Name | Description |
|---------------------------|---|
| <code>CHAR[(n)]</code> | Fixed-length character string, blank-padded to the size specified by <code>n</code> |
| <code>CLOB</code> | Large variable-length up to 1 GB |
| <code>LONG</code> | Variable unlimited length. |
| <code>NVARCHAR(n)</code> | Variable-length national character string, with limit. |
| <code>NVARCHAR2(n)</code> | Variable-length national character string, with limit. |
| <code>STRING</code> | Alias for <code>VARCHAR2</code> . |
| <code>VARCHAR(n)</code> | Variable-length character string, with limit (considered deprecated, but supported for compatibility) |
| <code>VARCHAR2(n)</code> | Variable-length character string, with limit |

Where `n` is a positive integer; these types can store strings up to `n` characters in length. An attempt to assign a value that exceeds the length of `n` will result in an error, unless the excess characters are all spaces, in which case the string will be truncated to the maximum length.

The storage requirement for data of these types is the actual string plus 1 byte if the string is less than 127 bytes, or 4 bytes if the string is 127 bytes or greater. In the case of `CHAR`, the padding also requires storage. Long strings are compressed by the system automatically, so the physical requirement on disk may be less. Long values are stored in background tables so they do not interfere with rapid access to the shorter column values.

The database character set determines the character set used to store textual values.

CHAR

If you do not specify a value for `n`, `n` will default to 1. If the string to be assigned is shorter than `n`, values of type `CHAR` will be space-padded to the specified width (`n`), and will be stored and displayed that way.

Padding spaces are treated as semantically insignificant. That is, trailing spaces are disregarded when comparing two values of type `CHAR`, and they will be removed when converting a `CHAR` value to one of the other string types.

If you explicitly cast an over-length value to a `CHAR(n)` type, the value will be truncated to `n` characters without raising an error (as specified by the SQL standard).

VARCHAR, VARCHAR2, NVARCHAR and NVARCHAR2

If the string to be assigned is shorter than `n`, values of type `VARCHAR`, `VARCHAR2`, `NVARCHAR` and `NVARCHAR2` will store the shorter string without padding.

Note that trailing spaces **are** semantically significant in **VARCHAR** values.

If you explicitly cast a value to a **VARCHAR** type, an over-length value will be truncated to **n** characters without raising an error (as specified by the SQL standard).

CLOB

You can store a large character string in a **CLOB** type. **CLOB** is semantically equivalent to **VARCHAR2** except no length limit is specified. Generally, you should use a **CLOB** type if the maximum string length is not known.

The longest possible character string that can be stored in a **CLOB** type is about 1 GB.

Note

The **CLOB** data type is actually a **DOMAIN** based on the PostgreSQL **TEXT** data type. For information on a **DOMAIN**, see the PostgreSQL core documentation at <https://www.postgresql.org/docs/current/static/sql-createdomain.html>

Thus, usage of the **CLOB** type is limited by what can be done for **TEXT** such as a maximum size of approximately 1 GB.

For usage of larger amounts of data, instead of using the **CLOB** data type, use the PostgreSQL **Large Objects** feature that relies on the **pg_largeobject** system catalog. For information on large objects, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/largeobjects.html>

5.2.2.3 Binary Data

The following table shows data types that allow the storage of binary strings:

| Name | Storage Size | Description |
|-----------|--|--|
| BINARY | The length of the binary string. | Fixed-length binary string, with a length between 1 and 8300. |
| BLOB | The actual binary string plus 1 byte if the binary string is less than 127 bytes, or 4 bytes if the binary string is 127 bytes or greater. | Variable-length binary string, with a maximum size of 1 GB. |
| VARBINARY | The length of the binary string | Variable-length binary string, with a length between 1 and 8300. |

A binary string is a sequence of octets (or bytes). Binary strings are distinguished from character strings by two characteristics: First, binary strings specifically allow storing octets of value zero and other "non-printable" octets (defined as octets outside the range 32 to 126). Second, operations on binary strings process the actual bytes, whereas the encoding and processing of character strings depends on locale settings.

5.2.2.4 Date/Time Types

The following discussion of the date/time types assumes that the configuration parameter, `edb_redwood_date`, has been set to `TRUE` whenever a table is created or altered.

Advanced Server supports the date/time types shown in the following table:

| Name | Storage Size | Description | Low Value | High Value | Resolution |
|--------------------------------|--------------|------------------------------|------------------|-----------------|---------------------------|
| DATE | 8 bytes | Date and time | 4713 BC | 5874897 AD | 1 second |
| INTERVAL DAY TO SECOND [(p)] | 12 bytes | Period of time | -178000000 years | 178000000 years | 1 microsecond / 14 digits |
| INTERVAL YEAR TO MONTH | 12 bytes | Period of time | -178000000 years | 178000000 years | 1 microsecond / 14 digits |
| TIMESTAMP [(p)] | 8 bytes | Date and time | 4713 BC | 5874897 AD | 1 microsecond |
| TIMESTAMP [(p)] WITH TIME ZONE | 8 bytes | Date and time with time zone | 4713 BC | 5874897 AD | 1 microsecond |

When `DATE` appears as the data type of a column in the data definition language (DDL) commands, `CREATE TABLE` or `ALTER TABLE`, it is translated to `TIMESTAMP` at the time the table definition is stored in the database. Thus, a time component will also be stored in the column along with the date.

When `DATE` appears as a data type of a variable in an SPL declaration section, or the data type of a formal parameter in an SPL procedure or an SPL function, or the return type of an SPL function, it is always translated to `TIMESTAMP` and thus can handle a time component if present.

`TIMESTAMP` accepts an optional precision value `p` which specifies the number of fractional digits retained in the seconds field. The allowed range of `p` is from 0 to 6 with the default being 6.

When `TIMESTAMP` values are stored as double precision floating-point numbers (currently the default), the effective limit of precision may be less than 6. `TIMESTAMP` values are stored as seconds before or after midnight 2000-01-01. Microsecond precision is achieved for dates within a few years of 2000-01-01, but the precision degrades for dates further away. When `TIMESTAMP` values are stored as eight-byte integers (a compile-time option), microsecond precision is available over the full range of values. However eight-byte integer timestamps have a more limited range of dates than shown above: from 4713 BC up to 294276 AD.

`TIMESTAMP (p) WITH TIME ZONE` is similar to `TIMESTAMP (p)`, but includes the time zone as well.

INTERVAL Types

`INTERVAL` values specify a period of time. Values of `INTERVAL` type are composed of fields that describe the value of the data. The following table lists the fields allowed in an `INTERVAL` type:

| Field Name | INTERVAL Values Allowed |
|------------|---|
| YEAR | Integer value (positive or negative) |
| MONTH | 0 through 11 |
| DAY | Integer value (positive or negative) |
| HOUR | 0 through 23 |
| MINUTE | 0 through 59 |
| SECOND | 0 through 59.9(p) where 9(p) is the precision of fractional seconds |

The fields must be presented in descending order – from `YEARS` to `MONTHS`, and from `DAYS` to `HOURS`,

MINUTES and then **SECONDS**.

Advanced Server supports two **INTERVAL** types compatible with Oracle databases.

The first variation supported by Advanced Server is **INTERVAL DAY TO SECOND [(p)]**. **INTERVAL DAY TO SECOND [(p)]** stores a time interval in days, hours, minutes and seconds.

| **p** specifies the precision of the **second** field.

Advanced Server interprets the value:

| **INTERVAL '1 2:34:5.678' DAY TO SECOND(3)**

as 1 day, 2 hours, 34 minutes, 5 seconds and 678 thousandths of a second.

Advanced Server interprets the value:

| **INTERVAL '1 23' DAY TO HOUR**

as 1 day and 23 hours.

Advanced Server interprets the value:

| **INTERVAL '2:34' HOUR TO MINUTE**

as 2 hours and 34 minutes.

Advanced Server interprets the value:

| **INTERVAL '2:34:56.129' HOUR TO SECOND(2)**

as 2 hours, 34 minutes, 56 seconds and 13 thousandths of a second. Note that the fractional second is rounded up to 13 because of the specified precision.

The second variation supported by Advanced Server that is compatible with Oracle databases is **INTERVAL YEAR TO MONTH**. This variation stores a time interval in years and months.

Advanced Server interprets the value:

| **INTERVAL '12-3' YEAR TO MONTH**

as 12 years and 3 months.

Advanced Server interprets the value:

| **INTERVAL '456' YEAR(2)**

as 12 years and 3 months.

Advanced Server interprets the value:

| **INTERVAL '300' MONTH**

as 25 years.

Date/Time Input

Date and time input is accepted in ISO 8601 SQL-compatible format, the Oracle default **dd-MON-yy** format, as well as a number of other formats provided that there is no ambiguity as to which component is the year, month, and day. However, use of the **TO_DATE** function is strongly recommended to avoid ambiguities.

Any date or time literal input needs to be enclosed in single quotes, like text strings. The following SQL standard syntax is also accepted:

```
| type 'value'
```

`type` is either `DATE` or `TIMESTAMP`.

`value` is a date/time text string.

Dates

The following block shows some possible input formats for dates, all of which equate to January 8, 1999.

Example

```
January 8, 1999
1999-01-08
1999-Jan-08
Jan-08-1999
08-Jan-1999
08-Jan-99
Jan-08-99
19990108
990108
```

The date values can be assigned to a `DATE` or `TIMESTAMP` column or variable. The hour, minute, and seconds fields will be set to zero if the date value is not appended with a time value.

Times

Some examples of the time component of a date or time stamp are shown in the following table:

| Example | Description |
|--------------|---|
| 04:05:06.789 | ISO 8601 |
| 04:05:06 | ISO 8601 |
| 04:05 | ISO 8601 |
| 040506 | ISO 8601 |
| 04:05 AM | Same as 04:05; AM does not affect value |
| 04:05 PM | Same as 16:05; input hour must be <= 12 |

Time Stamps

Valid input for time stamps consists of a concatenation of a date and a time. The date portion of the time stamp can be formatted according to any of the examples shown in the block under section [Dates](#). The time portion of the time stamp can be formatted according to any of examples shown in table under section [Times](#).

The following is an example of a time stamp which follows the Oracle default format.

```
| 08-JAN-99 04:05:06
```

The following is an example of a time stamp which follows the ISO 8601 standard.

```
| 1999-01-08 04:05:06
```

Date/Time Output

The default output format of the date/time types will be either (`dd-MON-yy`) referred to as the **Redwood date style**, compatible with Oracle databases, or (`yyy-mm-dd`) referred to as the ISO 8601 format, depending upon the application interface to the database. Applications that use JDBC such as SQL Interactive always present the date in ISO 8601 form. Other applications such as PSQL present the date in Redwood form.

The following table shows examples of the output formats for the two styles, Redwood and ISO 8601:

| Description | Example |
|-----------------------|---------------------|
| Redwood style | 31-DEC-05 07:37:16 |
| ISO 8601/SQL standard | 1997-12-17 07:37:16 |

Internals

Advanced Server uses Julian dates for all date/time calculations. Julian dates correctly predict or calculate any date after 4713 BC based on the assumption that the length of the year is 365.2425 days.

5.2.2.5 Boolean Types

Advanced Server provides the standard SQL type **BOOLEAN**. **BOOLEAN** can have one of only two states: **TRUE** or **FALSE**. A third state, **UNKNOWN**, is represented by the SQL **NULL** value.

| Name | Storage Size | Description |
|---------|--------------|------------------------------|
| BOOLEAN | 1 byte | Logical Boolean (true/false) |

The valid literal value for representing the true state is **TRUE**. The valid literal for representing the false state is **FALSE**.

5.2.2.6 XML Type

The **XMLTYPE** data type is used to store XML data. Its advantage over storing XML data in a character field is that it checks the input values for well-formedness, and there are support functions to perform type-safe operations on it.

The XML type can store well-formed “documents”, as defined by the XML standard, as well as “content” fragments, which are defined by the production **XMLDecl? content** in the XML standard. Roughly, this means that content fragments can have more than one top-level element or character node.

Note

Oracle does not support the storage of content fragments in **XMLTYPE** columns.

The following example shows the creation and insertion of a row into a table with an **XMLTYPE** column.

```

CREATE TABLE books (
    content      XMLTYPE
);

INSERT INTO books VALUES (XMLPARSE (DOCUMENT '<?xml
version="1.0"?><book><title>Manual</title><chapter>...</chapter></book>'));
SELECT * FROM books;

content
-----
<book><title>Manual</title><chapter>...</chapter></book>
(1 row)

```

5.2.3 Functions and Operators

Advanced Server provides a large number of functions and operators for the built-in data types.

5.2.3.1 Logical Operators

The usual logical operators are available: `AND`, `OR`, `NOT`

SQL uses a three-valued Boolean logic where the null value represents "unknown". Observe the following truth tables:

AND/OR Truth Table

| a | b | a AND b | a OR b |
|-------|-------|---------|--------|
| True | True | True | True |
| True | False | False | True |
| True | Null | Null | True |
| False | False | False | False |
| False | Null | False | Null |
| Null | Null | Null | Null |

NOT Truth Table

| a | NOT a |
|-------|-------|
| True | False |
| False | True |
| Null | Null |

The operators `AND` and `OR` are commutative, that is, you can switch the left and right operand without affecting

the result.

5.2.3.2 Comparison Operators

The usual comparison operators are shown in the following table:

| Operator | Description |
|----------|--------------------------|
| < | Less than |
| > | Greater than |
| <= | Less than or equal to |
| >= | Greater than or equal to |
| = | Equal |
| <> | Not equal |
| != | Not equal |

Comparison operators are available for all data types where this makes sense. All comparison operators are binary operators that return values of type `BOOLEAN`; expressions like `1 < 2 < 3` are not valid (because there is no `<` operator to compare a Boolean value with `3`).

In addition to the comparison operators, the special `BETWEEN`` construct is available.

`a BETWEEN x AND y`

is equivalent to

`a >= x AND a <= y`

Similarly,

`a NOT BETWEEN x AND y`

is equivalent to

`a < x OR a > y`

There is no difference between the two respective forms apart from the CPU cycles required to rewrite the first one into the second one internally.

To check whether a value is or is not null, use the constructs

`expression IS NULL`

`expression IS NOT NULL`

Do not write `expression = NULL` because `NULL` is not "equal to" `NULL`. (The null value represents an unknown value, and it is not known whether two unknown values are equal.) This behavior conforms to the SQL standard.

Some applications may expect that `expression = NULL` returns true if `expression` evaluates to the null value. It is highly recommended that these applications be modified to comply with the SQL standard.

5.2.3.3 Mathematical Functions and Operators

Mathematical operators are provided for many Advanced Server types. For types without common mathematical conventions for all possible permutations (e.g., date/time types) the actual behavior is described in subsequent sections.

The following table shows the available mathematical operators:

| Operator | Description | Example | Result |
|----------|------------------------------------|---------|--------|
| + | Addition | 2 + 3 | 5 |
| - | Subtraction | 2 - 3 | -1 |
| * | Multiplication | 2 * 3 | 6 |
| / | Division (See the following note.) | 4 / 2 | 2 |
| ** | Exponentiation Operator | 2 ** 3 | 8 |

Note

If the `db_dialect` configuration parameter in the `postgresql.conf` file is set to `redwood`, then division of a pair of `INTEGER` data types does not result in a truncated value. Any fractional result is retained as shown by the following example:

```
edb=# SET db_dialect TO redwood;
SET
edb=# SHOW db_dialect;
db_dialect
-----
redwood
(1 row)

edb=# SELECT CAST('10' AS INTEGER) / CAST('3' AS INTEGER) FROM dual;
?column?
-----
3.333333333333333
(1 row)
```

This behavior is compatible with Oracle databases where there is no native `INTEGER` data type, and any `INTEGER` data type specification is internally converted to `NUMBER(38)`, which results in retaining any fractional result.

If the `db_dialect` configuration parameter is set to `postgres`, then division of a pair of `INTEGER` data types results in a truncated value as shown by the following example:

```
edb=# SET db_dialect TO postgres;
SET
edb=# SHOW db_dialect;
db_dialect
-----
postgres
(1 row)

edb=# SELECT CAST('10' AS INTEGER) / CAST('3' AS INTEGER) FROM dual;
?column?
-----
3
(1 row)
```

This behavior is compatible with PostgreSQL databases where division involving any pair of `INTEGER`, `SMALLINT`, or `BIGINT` data types results in truncation of the result. The same truncated result is returned by Advanced Server when `db_dialect` is set to `postgres` as shown in the previous example.

Note however, that even when `db_dialect` is set to `redwood`, only division with a pair of `INTEGER` data types results in no truncation of the result. Division that includes only `SMALLINT` or `BIGINT` data types, with or without an `INTEGER` data type, does result in truncation in the PostgreSQL fashion without retaining the fractional portion as shown by the following where `INTEGER` and `SMALLINT` are involved in the division:

```
edb=# SHOW db_dialect;
db_dialect
-----
redwood
(1 row)

edb=# SELECT CAST('10' AS INTEGER) / CAST('3' AS SMALLINT) FROM dual;
?column?
-----
3
(1 row)
```

The following table shows the available mathematical functions. Many of these functions are provided in multiple forms with different argument types. Except where noted, any given form of a function returns the same data type as its argument. The functions working with **DOUBLE PRECISION** data are mostly implemented on top of the host system's C library; accuracy and behavior in boundary cases may therefore vary depending on the host system.

| Function | Return Type | Description | Example | Result |
|---|--|---|-----------------|--------------------|
| ABS(x) | Same as x | Absolute value | ABS(-17.4) | 17.4 |
| CEIL(DOUBLE PRECISION or NUMBER) | Same as input | Smallest integer not less than argument | CEIL(-42.8) | -42 |
| EXP(DOUBLE PRECISION or NUMBER) | Same as input | Exponential | EXP(1.0) | 2.7182818284590452 |
| FLOOR(DOUBLE PRECISION or NUMBER) | Same as input | Largest integer not greater than argument | FLOOR(-42.8) | 43 |
| LN(DOUBLE PRECISION or NUMBER) | Same as input | Natural logarithm | LN(2.0) | 0.6931471805599453 |
| LOG(b NUMBER, x NUMBER) | NUMBER | Logarithm to base b | LOG(2.0, 64.0) | 6.000000000000000 |
| MOD(y, x) | Same as argument types | Remainder of y/x | MOD(9, 4) | 1 |
| NVL(x, y) | Same as argument types; where both arguments are of the same data type | If x is null, then NVL returns y | NVL(9, 0) | 9 |
| POWER(a DOUBLE PRECISION, b DOUBLE PRECISION) | DOUBLE PRECISION | a raised to the power of b | POWER(9.0, 3.0) | 729.0000000000000 |

| | | | | |
|--|---------------|---|-------------------------------------|----------------------|
| POWER(a NUMBER, b NUMBER) | NUMBER | a raised to the power of b | POWER(9.0, 3.0) | 729.0000000000000000 |
| ROUND(DOUBLE PRECISION or NUMBER) | Same as input | Round to nearest integer | ROUND(42.4) | 42 |
| ROUND(v NUMBER, s INTEGER) | NUMBER | Round to s decimal places | ROUND(42.4382, 2) | 42.44 |
| SIGN(DOUBLE PRECISION or NUMBER) | Same as input | Sign of the argument (-1, 0, +1) | SIGN(-8.4) | -1 |
| SQRT(DOUBLE PRECISION or NUMBER) | Same as input | Square root | SQRT(2.0) | 1.414213562373095 |
| TRUNC(DOUBLE PRECISION or NUMBER) | Same as input | Truncate toward zero | TRUNC(42.8) | 42 |
| TRUNC(v NUMBER, s INTEGER) | NUMBER | Truncate to s decimal places | TRUNC(42.4382, 2) | 42.43 |
| WIDTH_BUCKET(op NUMBER, b1 NUMBER, b2 NUMBER, count INTEGER) | INTEGER | Return the bucket to which op would be assigned in an equidepth histogram with count buckets, in the range b1 to b2 | WIDTH_BUCKET(5.35, 0.024, 10.06, 5) | 3 |

The following table shows the available trigonometric functions. All trigonometric functions take arguments and return values of type **DOUBLE PRECISION**.

| Function | Description |
|-------------|------------------------|
| ACOS(x) | Inverse cosine |
| ASIN(x) | Inverse sine |
| ATAN(x) | Inverse tangent |
| ATAN2(x, y) | Inverse tangent of x/y |
| COS(x) | Cosine |
| SIN(x) | Sine |
| TAN(x) | Tangent |

5.2.3.4 String Functions and Operators

This section describes functions and operators for examining and manipulating string values. Strings in this context include values of the types **CHAR**, **VARCHAR2**, and **CLOB**. Unless otherwise noted, all of the functions listed below work on all of these types, but be wary of potential effects of automatic padding when using the **CHAR** type. Generally, the functions described here also work on data of non-string types by converting that data to a string representation first.

| Function | Return Type | Description | Example | Result |
|------------------|-------------|----------------------|----------------------|--------------|
| string string | CLOB | String concatenation | 'Enterprise' 'DB' | EnterpriseDB |

| | | | | |
|---|----------|--|--|----------|
| CONCAT(string, string) | CLOB | String concatenation | 'a' 'b' | ab |
| HEXTORAW(varchar2) | RAW | Converts a VARCHAR2 value to a RAW value | HEXTORAW('303132') | '012' |
| RAWTOHEX(raw) | VARCHAR2 | Converts a RAW value to a HEXADECIMAL value | RAWTOHEX('012') | '303132' |
| INSTR(string, set, [start [, occurrence]]) | INTEGER | Finds the location of a set of characters in a string, starting at position start in the string, string, and looking for the first, second, third and so on occurrences of the set. Returns 0 if the set is not found. | INSTR('PETER PIPER PIKED a PECK of PICKLED PEPPERS','PI',1,3) | 30 |
| INSTRB(string, set) | INTEGER | Returns the position of the set within the string. Returns 0 if set is not found. | INSTRB('PETER PIPER PIKED a PECK of PICKLED PEPPERS', 'PICK') | 13 |
| INSTRB(string, set, start) | INTEGER | Returns the position of the set within the string, beginning at start. Returns 0 if set is not found. | INSTRB('PETER PIPER PIKED a PECK of PICKLED PEPPERS','PICK', 14) | 30 |
| INSTRB(string, set, start, occurrence) | INTEGER | Returns the position of the specified occurrence of set within the string, beginning at start. Returns 0 if set is not found. | INSTRB('PETER PIPER PIKED a PECK of PICKLED PEPPERS','PICK', 1, 2) | 30 |
| LOWER(string) | CLOB | Convert string to lower case | LOWER('TOM') | tom |
| SUBSTR(string, start [, count]) | CLOB | Extract substring starting from start and going for count characters. If count is not specified, the string is clipped from the start till the end. | SUBSTR('This is a test',6,2) | is |
| SUBSTRB(string, start [, count]) | CLOB | Same as SUBSTR except start and count are in number of bytes. | SUBSTRB('abc',3) (assuming a double- byte character set) | c |
| SUBSTR2(string, start [, count]) | CLOB | Alias for SUBSTR. | SUBSTR2('This is a test',6,2) | is |
| SUBSTR2(string, start [, count]) | CLOB | Alias for SUBSTRB. | SUBSTR2('abc',3) (assuming a double- byte character set) | c |
| SUBSTR4(string, start [, count]) | CLOB | Alias for SUBSTR. | SUBSTR4('This is a test',6,2) | is |
| SUBSTR4(string, start [, count]) | CLOB | Alias for SUBSTRB. | SUBSTR4('abc',3) (assuming a double- byte character set) | c |
| SUBSTRC(string, start [, count]) | CLOB | Alias for SUBSTR. | SUBSTRC('This is a test',6,2) | is |
| SUBSTRC(string, start [, count]) | CLOB | Alias for SUBSTRB. | SUBSTRC('abc',3) (assuming a double- byte character set) | c |
| TRIM([LEADING TRAILING BOTH] [characters] FROM string) | CLOB | Remove the longest string containing only the characters (a space by default) from the start/end/both ends of the string. | TRIM(BOTH 'x' FROM 'xTomxx') | Tom |

| | | | | |
|-----------------------|------|--|---------------------------|--------|
| LTRIM(string [, set]) | CLOB | Removes all the characters specified in set from the left of a given string. If set is not specified, a blank space is used as default. | LTRIM('abcdefghi', 'abc') | defghi |
| RTRIM(string [, set]) | CLOB | Removes all the characters specified in set from the right of a given string. If set is not specified, a blank space is used as default. | RTRIM('abcdefghi', 'ghi') | abcdef |
| UPPER(string) | CLOB | Convert string to upper case | UPPER('tom') | TOM |

Additional string manipulation functions are available and are listed in the following table. Some of them are used internally to implement the SQL-standard string functions listed in the above table.

| Function | Return Type | Description | Example | Result |
|--|---|---|---|-----------|
| ASCII(string) | INTEGER | ASCII code of the first byte of the argument | ASCII('x') | 120 |
| CHR(INTEGER) | CLOB | Character with the given ASCII code | CHR(65) | A |
| DECODE(expr, expr1a, expr1b [, expr2a, expr2b]... [, default]) | Same as argument types of expr1b, expr2b,..., default | Finds first match of expr with expr1a, expr2a, etc. When match found, returns corresponding parameter pair, expr1b, expr2b, etc. If no match found, returns default. If no match found and default not specified, returns null. | DECODE(3, 1,'One', 2,'Two', 3,'Three', 'Not found') | Three |
| INITCAP(string) | CLOB | Convert the first letter of each word to uppercase and the rest to lowercase. Words are sequences of alphanumeric characters separated by non-alphanumeric characters. | INITCAP('hi THOMAS') | Hi Thomas |
| LENGTH | INTEGER | Returns the number of characters in a string value. | LENGTH('Côte d"azur') | 11 |
| LENGTHC | INTEGER | This function is identical in functionality to LENGTH; the function name is supported for compatibility. | LENGTHC('Côte d"azur') | 11 |
| LENGTH2 | INTEGER | This function is identical in functionality to LENGTH; the function name is supported for compatibility. | LENGTH2('Côte d"azur') | 11 |
| LENGTH4 | INTEGER | This function is identical in functionality to LENGTH; the function name is supported for compatibility. | LENGTH4('Côte d"azur') | 11 |
| LENGTHB | INTEGER | Returns the number of bytes required to hold the given value. | LENGTHB('Côte d"azur') | 12 |
| LPAD(string, length INTEGER [, fill]) | CLOB | Fill up string to size, length by prepending the characters, fill (a space by default). If string is already longer than length then it is truncated (on the right). | LPAD('hi', 5, 'xy') | xyhi |
| REPLACE(string, search_string [, replace_string]) | CLOB | Replaces one value in a string with another. If you do not specify a value for replace_string, the search_string value when found, is removed. | REPLACE('GEORGE', 'GE', 'EG') | EGOREG |
| RPAD(string, length INTEGER [, fill]) | CLOB | Fill up string to size, length by appending the characters, fill (a space by default). If string is already longer than length then it is truncated. | RPAD('hi', 5, 'xy') | hixyx |

| | | | | |
|-----------------------------|------|--|--------------------------------|-------|
| TRANSLATE(string, from, to) | CLOB | Any character in string that matches a character in the from set is replaced by the corresponding character in the to set. | TRANSLATE('12345', '14', 'ax') | a23x5 |
|-----------------------------|------|--|--------------------------------|-------|

Truncation of String Text Resulting from Concatenation with NULL

Note

This section describes a functionality that is not compatible with Oracle databases, which may lead to some inconsistency when converting data from Oracle to Advanced Server.

For Advanced Server, when a column value is `NULL`, the concatenation of the column with a text string may result in either of the following:

- Return of the text string
- Disappearance of the text string (that is, a null result)

The result is dependent upon the data type of the `NULL` column and the way in which the concatenation is done.

If one uses the string concatenation operator `'||'`, then the types that have implicit coercion to text as listed in Table Data Types with Implicit Coercion to Text will not truncate the string if one of the input parameters is `NULL`, whereas for other types it will truncate the string unless the explicit type cast is used (that is, `::text`). Also, to see the consistent behavior in the presence of nulls, one can use the `CONCAT` function.

The following query lists the data types that have implicit coercion to text:

```
SELECT castsource::regtype, casttarget::regtype, castfunc::regproc,
CASE castcontext
    WHEN 'e' THEN 'explicit'
    WHEN 'a' THEN 'implicit in assignment'
    WHEN 'i' THEN 'implicit in expressions'
END as castcontext,
CASE castmethod
    WHEN 'f' THEN 'function'
    WHEN 'i' THEN 'input/output function'
    WHEN 'b' THEN 'binary-coercible'
END as castmethod
FROM pg_cast
WHERE casttarget::regtype::text = 'text'
AND castcontext='i';
```

The result of the query is listed in the following table:

| castsource | casttarget | castfunc | castcontext | castmethod |
|-------------------|------------|-----------------|-------------------------|-----------------------|
| character | text | pg_catalog.text | implicit in expressions | function |
| character varying | text | - | implicit in expressions | binary-coercible |
| "char" | text | pg_catalog.text | implicit in expressions | function |
| name | text | pg_catalog.text | implicit in expressions | function |
| pg_node_tree | text | - | implicit in expressions | binary-coercible |
| pg_ndistinct | text | - | implicit in expressions | input/output function |
| pg_dependencies | text | - | implicit in expressions | input/output function |
| integer | text | - | implicit in expressions | input/output function |
| smallint | text | - | implicit in expressions | input/output function |

| | | | | |
|-----------------------------|------|-----------------|-------------------------|-----------------------|
| oid | text | - | implicit in expressions | input/output function |
| date | text | - | implicit in expressions | input/output function |
| double precision | text | - | implicit in expressions | input/output function |
| real | text | - | implicit in expressions | input/output function |
| time with time zone | text | - | implicit in expressions | input/output function |
| time without time zone | text | - | implicit in expressions | input/output function |
| timestamp with time zone | text | - | implicit in expressions | input/output function |
| interval | text | - | implicit in expressions | input/output function |
| bigrnt | text | - | implicit in expressions | input/output function |
| numeric | text | - | implicit in expressions | input/output function |
| timestamp without time zone | text | - | implicit in expressions | input/output function |
| record | text | - | implicit in expressions | input/output function |
| boolean | text | pg_catalog.text | implicit in expressions | function |
| bytea | text | - | implicit in expressions | input/output function |

For information on the column output, see the `pg_cast` system catalog in the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/catalog-pg-cast.html>

So for example, data type `UUID` is not in this list and therefore does not have the implicit coercion to text. As a result, certain concatenation attempts with a `NULL UUID` column results in a truncated text result.

The following table is created for this example with a single row with all `NULL` column values.

```
CREATE TABLE null_concat_types (
    boolean_type BOOLEAN,
    uuid_type UUID,
    char_type CHARACTER
);
```

```
INSERT INTO null_concat_types VALUES (NULL, NULL, NULL);
```

Columns `boolean_type` and `char_type` have the implicit coercion to text while column `uuid_type` does not.

Thus, string concatenation with the concatenation operator `'||'` against columns `boolean_type` or `char_type` results in the following:

```
SELECT 'x=' || boolean_type || 'y' FROM null_concat_types;
```

```
?column?
```

```
-----  
x=y  
(1 row)
```

```
SELECT 'x=' || char_type || 'y' FROM null_concat_types;
```

```
?column?
```

```
-----  
x=y  
(1 row)
```

But concatenation with column `uuid_type` results in the loss of the `x=` string:

```
SELECT 'x=' || uuid_type || 'y' FROM null_concat_types;
```

```
?column?
```

```
-----
y
(1 row)
```

However, using explicit casting with `::text` prevents the loss of the `x=` string:

```
SELECT 'x=' || uuid_type::text || 'y' FROM null_concat_types;
```

```
?column?
```

```
-----
x=y
(1 row)
```

Using the `CONCAT` function also preserves the `x=` string:

```
SELECT CONCAT('x=',uuid_type) || 'y' FROM null_concat_types;
```

```
?column?
```

```
-----
x=y
(1 row)
```

Thus, depending upon the data type of a `NULL` column, explicit casting or the `CONCAT` function should be used to avoid loss of some text string.

SYS_GUID

The `SYS_GUID` function generates and returns a globally unique identifier; the identifier takes the form of 16 bytes of `RAW` data. The `SYS_GUID` function is based on the `uuid-ossp` module to generate universally unique identifiers. The synopsis is:

`SYS_GUID()`

Example

The following example adds a column to the table `EMP`, inserts a unique identifier, and returns a 16-byte `RAW` value:

```
edb=# CREATE TABLE EMP(C1 RAW (16) DEFAULT SYS_GUID() PRIMARY KEY, C2 INT);
CREATE TABLE
edb=# INSERT INTO EMP(C2) VALUES (1);
INSERT 0 1
edb=# SELECT * FROM EMP;
   c1      | c2
-----+-----
\xb944970d3a1b42a7a2119265c49cbb7f | 1
(1 row)
```

5.2.3.5 Pattern Matching String Functions

Advanced Server offers support for the `REGEXP_COUNT`, `REGEXP_INSTR` and `REGEXP_SUBSTR` functions. These functions search a string for a pattern specified by a regular expression, and return information about occurrences of the pattern within the string. The pattern should be a POSIX-style regular expression; for more information about forming a POSIX-style regular expression, please refer to the core documentation at:

<https://www.postgresql.org/docs/current/static/functions-matching.html>

REGEXP_COUNT

`REGEXP_COUNT` searches a string for a regular expression, and returns a count of the times that the regular expression occurs. The signature is:

```
INTEGER REGEXP_COUNT
(
    srcstr TEXT,
    pattern TEXT,
    position DEFAULT 1
    modifier DEFAULT NULL
)
```

Parameters

`srcstr`

`srcstr` specifies the string to search.

`pattern`

`pattern` specifies the regular expression for which `REGEXP_COUNT` will search.

`position`

`position` is an integer value that indicates the position in the source string at which `REGEXP_COUNT` will begin searching. The default value is `1`.

`modifier`

`modifier` specifies values that control the pattern matching behavior. The default value is `NULL`. For a complete list of the modifiers supported by Advanced Server, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/functions-matching.html>

Example

In the following simple example, `REGEXP_COUNT` returns a count of the number of times the letter `i` is used in the character string `'reinitializing'`:

```
edb=# SELECT REGEXP_COUNT('reinitializing', 'i', 1) FROM DUAL;
regexp_count
-----
      5
(1 row)
```

In the first example, the command instructs `REGEXP_COUNT` begins counting in the first position; if we modify the command to start the count on the 6th position:

```
edb=# SELECT REGEXP_COUNT('reinitializing', 'i', 6) FROM DUAL;
```

```
regexp_count
-----
```

```
3
(1 row)
```

`REGEXP_COUNT` returns 3; the count now excludes any occurrences of the letter i that occur before the 6th position.

REGEXP_INSTR

`REGEXP_INSTR` searches a string for a POSIX-style regular expression. This function returns the position within the string where the match was located. The signature is:

```
INTEGER REGEXP_INSTR
(
srcstr      TEXT,
pattern     TEXT,
position    INT DEFAULT 1,
occurrence  INT DEFAULT 1,
returnparam INT DEFAULT 0,
modifier    TEXT DEFAULT NULL,
subexpression INT DEFAULT 0,
)
```

Parameters:

`srcstr`

`srcstr` specifies the string to search.

`pattern`

`pattern` specifies the regular expression for which `REGEXP_INSTR` will search.

`position`

`position` specifies an integer value that indicates the start position in a source string. The default value is 1.

`occurrence`

`occurrence` specifies which match is returned if more than one occurrence of the pattern occurs in the string that is searched. The default value is 1.

`returnparam`

`returnparam` is an integer value that specifies the location within the string that `REGEXP_INSTR` should return. The default value is 0. Specify:

0 to return the location within the string of the first character that matches `pattern`.

A value greater than 0 to return the position of the first character following the end of the `pattern`.

`modifier`

`modifier` specifies values that control the pattern matching behavior. The default value is `NULL`. For a complete list of the modifiers supported by Advanced Server, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/functions-matching.html>

subexpression

`subexpression` is an integer value that identifies the portion of the `pattern` that will be returned by `REGEXP_INSTR`. The default value of `subexpression` is 0.

If you specify a value for `subexpression`, you must include one (or more) set of parentheses in the `pattern` that isolate a portion of the value being searched for. The value specified by `subexpression` indicates which set of parentheses should be returned; for example, if `subexpression` is 2, `REGEXP_INSTR` will return the position of the second set of parentheses.

Example

In the following simple example, `REGEXP_INSTR` searches a string that contains a phone number for the first occurrence of a pattern that contains three consecutive digits:

```
edb=# SELECT REGEXP_INSTR('800-555-1212', '[0-9][0-9][0-9]', 1, 1) FROM DUAL;
regexp_instr
-----
      1
(1 row)
```

The command instructs `REGEXP_INSTR` to return the position of the first occurrence. If we modify the command to return the start of the second occurrence of three consecutive digits:

```
edb=# SELECT REGEXP_INSTR('800-555-1212', '[0-9][0-9][0-9]', 1, 2) FROM DUAL;
regexp_instr
-----
      5
(1 row)
```

`REGEXP_INSTR` returns 5; the second occurrence of three consecutive digits begins in the 5th position.

REGEXP_SUBSTR

The `REGEXP_SUBSTR` function searches a string for a pattern specified by a POSIX compliant regular expression. `REGEXP_SUBSTR` returns the string that matches the pattern specified in the call to the function. The signature of the function is:

```
TEXT REGEXP_SUBSTR
(
    srcstr    TEXT,
    pattern   TEXT,
    position   INT DEFAULT 1,
    occurrence INT DEFAULT 1,
    modifier   TEXT DEFAULT NULL,
    subexpression INT DEFAULT 0
)
```

Parameters:

srcstr

`srcstr` specifies the string to search.

pattern

`pattern` specifies the regular expression for which `REGEXP_SUBSTR` will search.

`position`

`position` specifies an integer value that indicates the start position in a source string. The default value is `1`.

`occurrence`

`occurrence` specifies which match is returned if more than one occurrence of the pattern occurs in the string that is searched. The default value is `1`.

`modifier`

`modifier` specifies values that control the pattern matching behavior. The default value is `NULL`. For a complete list of the modifiers supported by Advanced Server, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/functions-matching.html>

`subexpression`

`subexpression` is an integer value that identifies the portion of the `pattern` that will be returned by `REGEXP_SUBSTR`. The default value of `subexpression` is `0`.

If you specify a value for `subexpression`, you must include one (or more) set of parentheses in the `pattern` that isolate a portion of the value being searched for. The value specified by `subexpression` indicates which set of parentheses should be returned; for example, if `subexpression` is `2`, `REGEXP_SUBSTR` will return the value contained within the second set of parentheses.

Example

In the following simple example, `REGEXP_SUBSTR` searches a string that contains a phone number for the first set of three consecutive digits:

```
edb=# SELECT REGEXP_SUBSTR('800-555-1212', '[0-9][0-9][0-9]', 1, 1) FROM
DUAL;
regexp_substr
-----
800
(1 row)
```

It locates the first occurrence of three digits and returns the string `(800)`; if we modify the command to check for the second occurrence of three consecutive digits:

```
edb=# SELECT REGEXP_SUBSTR('800-555-1212', '[0-9][0-9][0-9]', 1, 2) FROM
DUAL;
regexp_substr
-----
555
(1 row)
```

`REGEXP_SUBSTR` returns `555`, the contents of the second substring.

5.2.3.6 Pattern Matching Using the LIKE Operator

Advanced Server provides pattern matching using the traditional SQL `LIKE` operator. The syntax for the `LIKE` operator is as follows.

```
string LIKE pattern [ ESCAPE escape-character ]
string NOT LIKE pattern [ ESCAPE escape-character ]
```

Every `pattern` defines a set of strings. The `LIKE` expression returns `TRUE` if `string` is contained in the set of strings represented by `pattern`. As expected, the `NOT LIKE` expression returns `FALSE` if `LIKE` returns `TRUE`, and vice versa. An equivalent expression is `NOT (string LIKE pattern)`.

If `pattern` does not contain percent signs or underscore, then the pattern only represents the string itself; in that case `LIKE` acts like the equals operator. An underscore (`_`) in `pattern` stands for (matches) any single character; a percent sign (`%`) matches any string of zero or more characters.

Some examples:

```
'abc' LIKE 'abc'    true
'abc' LIKE 'a%'    true
'abc' LIKE '_b_'    true
'abc' LIKE 'c'     false
```

`LIKE` pattern matches always cover the entire string. To match a pattern anywhere within a string, the pattern must therefore start and end with a percent sign.`

To match a literal underscore or percent sign without matching other characters, the respective character in `pattern` must be preceded by the escape character. The default escape character is the backslash but a different one may be selected by using the `ESCAPE` clause. To match the escape character itself, write two escape characters.

Note that the backslash already has a special meaning in string literals, so to write a pattern constant that contains a backslash you must write two backslashes in an SQL statement. Thus, writing a pattern that actually matches a literal backslash means writing four backslashes in the statement. You can avoid this by selecting a different escape character with `ESCAPE`; then a backslash is not special to `LIKE` anymore. (But it is still special to the string literal parser, so you still need two of them.)

It's also possible to select no escape character by writing `ESCAPE ''`. This effectively disables the escape mechanism, which makes it impossible to turn off the special meaning of underscore and percent signs in the pattern.

5.2.3.7 Data Type Formatting Functions

The Advanced Server formatting functions described in the following table provide a powerful set of tools for converting various data types (date/time, integer, floating point, numeric) to formatted strings and for converting from formatted strings to specific data types. These functions all follow a common calling convention: the first argument is the value to be formatted and the second argument is a string template that defines the output or input format.

| Function | Return Type | Description | Example | Result |
|--|-----------------------|--|---|------------------------------|
| <code>TO_CHAR(DATE [, format])</code> | <code>VARCHAR2</code> | Convert a date/time to a string with output, format. If omitted default format is DD-MON-YY. | <code>TO_CHAR(SYSDATE, 'MM/DD/YYYY HH12:MI:SS AM')</code> | 07/25/2007 09:43:02 AM |

| | | | | |
|-----------------------------------|-------------|--|---|---------------------------|
| TO_CHAR(TIMESTAMP [, format]) | VARCHAR2 | Convert a timestamp to a string with output, format. If omitted default format is DD-MON-YY. | TO_CHAR(CURRENT_TIMESTAMP, 'MM/DD/YYYY HH12:MI:SS AM') | 08/13/2015 08:55:22 PM |
| TO_CHAR(INTEGER [, format]) | VARCHAR2 | Convert an integer to a string with output, format | TO_CHAR(2412, '999,999S') | 2,412+ |
| TO_CHAR(NUMBER [, format]) | VARCHAR2 | Convert a decimal number to a string with output, format | TO_CHAR(10125.35, '999,999.99') | 10,125.35 |
| TO_CHAR(DOUBLE PRECISION, format) | VARCHAR2 | Convert a floating-point number to a string with output, format | TO_CHAR(CAST(123.5282 AS REAL), '999.99') | 123.53 |
| TO_DATE(string [, format]) | DATE | Convert a date formatted string to a DATE data type | TO_DATE('2007-07-04 13:39:10', 'YYYY-MM-DD HH24:MI:SS') | 04-JUL-07 |
| TO_NUMBER(string [, format]) | NUMBER | Convert a number formatted string to a NUMBER data type | TO_NUMBER('2,412-', '999,999S') | -2412 |
| TO_TIMESTAMP(string, format) | TIMESTAMPTZ | Convert a timestamp formatted string to a TIMESTAMPTZ data type | TO_TIMESTAMP('05 Dec 2000 08:30:25 pm', 'DD Mon YYYY hh12:mi:ss pm') | 05-DEC-00 20:30:25 +05:30 |
| TO_TIMESTAMP_TZ(string, format) | TIMESTAMPTZ | Convert a timestamp formatted string to a TIMESTAMPTZ data type | TO_TIMESTAMP_TZ('2003/12/13 10:13:18 -8:00', 'YYYY/MM/DD HH:MI:SS TZH:TZM') | 13-DEC-03 23:43:18 +05:30 |

TO_CHAR, TO_DATE, TO_TIMESTAMP, and TO_TIMESTAMP_TZ

In an output template string (for `TO_CHAR`), there are certain patterns that are recognized and replaced with appropriately-formatted data from the value to be formatted. Any text that is not a template pattern is simply copied verbatim. Similarly, in an input template string (for anything but `TO_CHAR`), template patterns identify the parts of the input data string to be looked at and the values to be found there.

If you do not specify a date, month, or year when calling `TO_TIMESTAMP`, `TO_TIMESTAMP_TZ` or `TO_DATE`, then by default the output format considers the first date of a current month or current year respectively. In the following example, date, month, and year is not specified in the input string; `TO_TIMESTAMP`, `TO_TIMESTAMP_TZ`, and `TO_DATE` returns a default value of the first date of a current month and current year.

```
edb=# select to_timestamp('12', 'HH');
      to_timestamp
-----
```

```
01-MAY-20 12:00:00 +05:30
(1 row)
```

```
edb=# select to_timestamp_tz('12', 'HH');
      to_timestamp_tz
-----
```

```
01-MAY-20 12:00:00 +05:30
(1 row)
```

```
edb=# select to_date('12', 'HH');
      to_date
```

01-MAY-20 12:00:00

(1 row)

The following table shows the template patterns available for formatting date values using the `TO_CHAR`, `TO_DATE`, `TO_TIMESTAMP`, and `TO_TIMESTAMP_TZ` functions.

| Pattern | Description |
|-----------------------------------|---|
| HH | Hour of day (01-12) |
| HH12 | Hour of day (01-12) |
| HH24 | Hour of day (00-23) |
| MI | Minute (00-59) |
| SS | Second (00-59) |
| SSSS | Seconds past midnight (0-86399) |
| FFn | Fractional seconds where n is an optional integer from 1 to 9 for the number of digits to return. If omitted, the default is 6. |
| AM or A.M. or PM or P.M. | Meridian indicator (uppercase) |
| am or a.m. or pm or p.m. | Meridian indicator (lowercase) |
| Y,YYY | Year (4 and more digits) with comma |
| YEAR | Year (spelled out) |
| SYEAR | Year (spelled out) (BC dates prefixed by a minus sign) |
| YYYY | Year (4 and more digits) |
| YYYYY | Year (4 and more digits) (BC dates prefixed by a minus sign) |
| YYY | Last 3 digits of year |
| YY | Last 2 digits of year |
| Y | Last digit of year |
| IYYY | ISO year (4 and more digits) |
| IYY | Last 3 digits of ISO year |
| IY | Last 2 digits of ISO year |
| I | Last 1 digit of ISO year |
| BC or B.C. or AD or A.D. | Era indicator (uppercase) |
| bc or b.c. or ad or a.d. | Era indicator (lowercase) |
| MONTH | Full uppercase month name |
| Month | Full mixed-case month name |
| month | Full lowercase month name |
| MON | Abbreviated uppercase month name (3 chars in English, localized lengths vary) |
| Mon | Abbreviated mixed-case month name (3 chars in English, localized lengths vary) |
| mon | Abbreviated lowercase month name (3 chars in English, localized lengths vary) |
| MM | Month number (01-12) |
| DAY | Full uppercase day name |

| | |
|-----|---|
| Day | Full mixed-case day name |
| day | Full lowercase day name |
| DY | Abbreviated uppercase day name (3 chars in English, localized lengths vary) |
| Dy | Abbreviated mixed-case day name (3 chars in English, localized lengths vary) |
| dy | Abbreviated lowercase day name (3 chars in English, localized lengths vary) |
| DDD | Day of year (001-366) |
| DD | Day of month (01-31) |
| D | Day of week (1-7; Sunday is 1) |
| W | Week of month (1-5) (The first week starts on the first day of the month) |
| WW | Week number of year (1-53) (The first week starts on the first day of the year) |
| IW | ISO week number of year; the first Thursday of the new year is in week 1 |
| CC | Century (2 digits); the 21st century starts on 2001-01-01 |
| SCC | Same as CC except BC dates are prefixed by a minus sign |
| J | Julian Day (days since January 1, 4712 BC) |
| Q | Quarter |
| RM | Month in Roman numerals (I-XII; I=January) (uppercase) |
| rm | Month in Roman numerals (i-xii; i=January) (lowercase) |

First 2 digits of the year when given only the last 2 digits of the year. Result is based upon an algorithm using the current year and the given 2-digit year. The first 2 digits of the given 2-digit year will be the same as the first 2 digits of the current year with the following exceptions:

| | |
|------|---|
| RR | If the given 2-digit year is < 50 and the last 2 digits of the current year is >= 50, then the first 2 digits for the given year is 1 greater than the first 2 digits of the current year. If the given 2-digit year is >= 50 and the last 2 digits of the current year is < 50, then the first 2 digits for the given year is 1 less than the first 2 digits of the current year. |
| RRRR | Only affects TO_DATE function. Allows specification of 2-digit or 4-digit year. If 2-digit year given, then returns first 2 digits of year like RR format. If 4-digit year given, returns the given 4-digit year. |
| TZH | Time-zone hours |
| TZM | Time-zone minutes |

Date and Time Modifiers

Certain modifiers may be applied to any template pattern to alter its behavior. For example, `FMMonth` is the `Month` pattern with the `FM` modifier. The following table shows the modifier patterns for date/time formatting.

| Modifier | Description | Example |
|-----------|---|-----------------|
| FM prefix | Fill mode (suppress padding blanks and zeros) | FMMonth |
| TH suffix | Uppercase ordinal number suffix | DDTH |
| th suffix | Lowercase ordinal number suffix | DDth |
| FX prefix | Fixed format global option (see usage notes) | FX Month DD Day |
| SP suffix | Spell mode | DDSP |

Usage notes for date/time formatting:

- `FM` suppresses leading zeroes and trailing blanks that would otherwise be added to make the output of a pattern fixed-width.
- `TO_TIMESTAMP`, `TO_TIMESTAMP_TZ`, and `TO_DATE` skip multiple blank spaces in the input string if the `FX` option is not used. `FX` must be specified as the first item in the template. For example:

TO_TIMESTAMP('2000 - JUN', 'YYYY-MON') is correct, but
 TO_TIMESTAMP('2000 JUN', 'FXXXXXX MON') and
 TO_TIMESTAMP_TZ('2000 JUN', 'FXXXXXX MON') returns an error
 because TO_TIMESTAMP and TO_TIMESTAMP_TZ expects one space only.

- Ordinary text is allowed in `TO_CHAR` templates and will be output literally.
- In conversions from string to `timestamp`, `timestamptz`, or `date`, the `CC` field is ignored if there is a `YY`, `YYYY` or `Y,YYYY` field. If `CC` is used with `YY` or `Y` then the year is computed as $(CC-1)*100+YY$.

The following table shows some examples of the use of the `TO_CHAR` and `TO_DATE` functions:

| Expression | Result |
|---|-------------------------|
| <code>TO_CHAR(CURRENT_TIMESTAMP, 'Day, DD HH12:MI:SS')</code> | 'Tuesday , 06 05:39:18' |
| <code>TO_CHAR(CURRENT_TIMESTAMP, 'FMDay, FMDD HH12:MI:SS')</code> | 'Tuesday, 6 05:39:18' |
| <code>TO_CHAR(-0.1, '99.99')</code> | ' -.10' |
| <code>TO_CHAR(-0.1, 'FM9.99')</code> | '-.1' |
| <code>TO_CHAR(0.1, '0.9')</code> | ' 0.1' |
| <code>TO_CHAR(12, '9990999.9')</code> | ' 0012.0' |
| <code>TO_CHAR(12, 'FM9990999.9')</code> | '0012.' |
| <code>TO_CHAR(485, '999')</code> | ' 485' |
| <code>TO_CHAR(-485, '999')</code> | '-485' |
| <code>TO_CHAR(1485, '9,999')</code> | ' 1,485' |
| <code>TO_CHAR(1485, '9G999')</code> | ' 1,485' |
| <code>TO_CHAR(148.5, '999.999')</code> | ' 148.500' |
| <code>TO_CHAR(148.5, 'FM999.999')</code> | '148.5' |
| <code>TO_CHAR(148.5, 'FM999.990')</code> | '148.500' |
| <code>TO_CHAR(148.5, '999D999')</code> | ' 148.500' |
| <code>TO_CHAR(3148.5, '9G999D999')</code> | ' 3,148.500' |
| <code>TO_CHAR(-485, '999S')</code> | '485-' |
| <code>TO_CHAR(-485, '999MI')</code> | '485-' |
| <code>TO_CHAR(485, '999MI')</code> | '485 ' |
| <code>TO_CHAR(485, 'FM999MI')</code> | '485' |
| <code>TO_CHAR(-485, '999PR')</code> | '<485>' |
| <code>TO_CHAR(485, 'L999')</code> | '\$ 485' |
| <code>TO_CHAR(485, 'RN')</code> | ' CDLXXXV' |
| <code>TO_CHAR(485, 'FMRN')</code> | 'CDLXXXV' |
| <code>TO_CHAR(5.2, 'FMRN')</code> | 'V' |
| <code>TO_CHAR(12, '99V999')</code> | ' 12000' |
| <code>TO_CHAR(12.4, '99V999')</code> | ' 12400' |
| <code>TO_CHAR(12.45, '99V9')</code> | ' 125' |

The following table shows some examples of the use of the `TO_TIMESTAMP_TZ` function:

| Expression | Result |
|--|-----------------------------|
| <code>TO_TIMESTAMP_TZ('12-JAN-2010', 'DD-MONTH-YYYY')</code> | '12-JAN-10 00:00:00 +05:30' |
| <code>TO_TIMESTAMP_TZ('03-APR-07 09:12:21 P.M.', 'DD-MON-YY HH12:MI:SS A.M')</code> | '03-APR-07 09:12:21 +05:30' |
| <code>TO_TIMESTAMP_TZ('2003/12/13 10:13:18 -8:00', 'YYYY/MM/DD HH:MI:SS TZH:TZM')</code> | '13-DEC-03 23:43:18 +05:30' |

| | |
|---|---------------------------------|
| TO_TIMESTAMP_TZ('20-MAR-20 04:30:00 +08:00', 'DD-MON-YY HH:MI:SS TZH:TZM'); | '20-MAR-20 02:00:00 +05:30' |
| TO_TIMESTAMP_TZ('10-Sep-02 14:10:10.123000', 'DD-MON-RR HH24:MI:SS.FF'); | '10-SEP-02 14:10:10.123 +05:30' |

IMMUTABLE TO_CHAR(TIMESTAMP, format) Function

There are certain cases of the `TO_CHAR` function that can result in usage of an `IMMUTABLE` form of the function. Basically, a function is `IMMUTABLE` if the function does not modify the database, and the function returns the same, consistent value dependent upon only its input parameters. That is, the settings of configuration parameters, the locale, the content of the database, etc. do not affect the results returned by the function.

For more information about function volatility categories `VOLATILE`, `STABLE`, and `IMMUTABLE`, see the PostgreSQL Core documentation at:

<https://www.postgresql.org/docs/current/static/xfunc-volatility.html>

A particular advantage of an `IMMUTABLE` function is that it can be used in the `CREATE INDEX` command to create an index based on that function.

In order for the `TO_CHAR` function to use the `IMMUTABLE` form the following conditions must be satisfied:

- The first parameter of the `TO_CHAR` function must be of data type `TIMESTAMP`.
- The format specified in the second parameter of the `TO_CHAR` function must not affect the return value of the function based on factors such as language, locale, etc. For example a format of '`YYYY-MM-DD HH24:MI:SS`' can be used for an `IMMUTABLE` form of the function since, regardless of locale settings, the result of the function is the date and time expressed solely in numeric form. However, a format of '`DD-MON-YYYY`' cannot be used for an `IMMUTABLE` form of the function because the 3-character abbreviation of the month may return different results depending upon the locale setting.

Format patterns that result in a non-immutable function include any variations of spelled out or abbreviated months (`MONTH`, `MON`), days (`DAY`, `DY`), median indicators (`AM`, `PM`), or era indicators (`BC`, `AD`).

For the following example, a table with a `TIMESTAMP` column is created.

```
CREATE TABLE ts_tbl (ts_col TIMESTAMP);
```

The following shows the successful creation of an index with the `IMMUTABLE` form of the `TO_CHAR` function.

```
edb=# CREATE INDEX ts_idx ON ts_tbl (TO_CHAR(ts_col,'YYYY-MM-DD HH24:MI:SS'));
CREATE INDEX
edb=# \dS ts_idx
          Index "public.ts_idx"
 Column |      Type       |           Definition
-----+-----+-----+
 to_char | character varying | to_char(ts_col, 'YYYY-MM-DD
 HH24:MI:SS'::character varying)
 btree, for table "public.ts_tbl"
```

The following results in an error because the format specified in the `TO_CHAR` function prevents the use of the `IMMUTABLE` form since the 3-character month abbreviation, `MON`, may result in different return values based on the locale setting.

```
edb=# CREATE INDEX ts_idx_2 ON ts_tbl (TO_CHAR(ts_col, 'DD-MON-YYYY'));
ERROR: functions in index expression must be marked IMMUTABLE
```

TO_NUMBER

The following table lists the template patterns available for formatting numeric values:

| Pattern | Description |
|------------|---|
| 9 | Value with the specified number of digits |
| 0 | Value with leading zeroes |
| . (period) | Decimal point |
| , (comma) | Group (thousand) separator |
| \$ | Dollar sign |
| S | Sign anchored to number (uses locale). |
| L | Currency symbol (uses locale) |

Note that 9 results in a value with the same number of digits as there are 9s. If a digit is not available, the server ignores the corresponding 9s. The S pattern does not support + and \$ pattern does not support decimal points in the expression.

The following table shows some examples of the use of the TO_NUMBER function:

| Expression | Result |
|--|--------------|
| TO_NUMBER('-65', 'S99') | ' -65' |
| TO_NUMBER('\$65', 'L99') | ' 65' |
| TO_NUMBER('9678584', '9999999') | ' 9678584' |
| TO_NUMBER('123,456,789', '999,999,999') | ' 123456789' |
| TO_NUMBER('1210.73', '9999.99') | ' 1210.73' |
| TO_NUMBER('1210.73') | ' 1210.73' |
| TO_NUMBER('0101.010','FM99999999.99999') | ' 101.010' |

Numeric Modifiers

The following table shows the modifier pattern for numeric formatting:

| Pattern | Description | Example |
|-----------|---|---------|
| FM prefix | fill mode (suppress trailing zeroes and padding blanks) | FM99.99 |

5.2.3.8 Date/Time Functions and Operators

The Date/Time Functions table shows the available functions for date/time value processing, with details appearing in the following subsections. The following table illustrates the behaviors of the basic arithmetic operators (+, -). For formatting functions, refer to [IMMUTABLE TO_CHAR\(TIMESTAMP, format\) Function](#). You should be familiar with the background information on date/time data types, see [Date/Time Types](#).

| Operator | Example | Result |
|----------|-----------------------|--------------------|
| plus (+) | DATE '2001-09-28' + 7 | 05-OCT-01 00:00:00 |

| | | |
|-----------|---|--------------------|
| plus (+) | TIMESTAMP '2001-09-28 13:30:00' + 3 | 01-OCT-01 13:30:00 |
| minus (-) | DATE '2001-10-01' - 7 | 24-SEP-01 00:00:00 |
| minus (-) | TIMESTAMP '2001-09-28 13:30:00' - 3 | 25-SEP-01 13:30:00 |
| minus (-) | TIMESTAMP '2001-09-29 03:00:00' - TIMESTAMP '2001-09-27 12:00:00' | @ 1 day 15 hours |

In the date/time functions of the following table the use of the **DATE** and **TIMESTAMP** data types are interchangeable.

| Function | Return Type | Description | Example | Result |
|-----------------------------------|------------------|--|---|------------------------|
| ADD_MONTHS (DATE, NUMBER) | DATE | Add months to a date | ADD_MONTHS ('28-FEB-97', .3.8) | 31-MAY-97 00:00:00 |
| CURRENT_DATE | DATE | Current date | CURRENT_DATE | 04-JUL-07 |
| CURRENT_TIMESTAMP | TIMESTAMP | Returns the current date and time | CURRENT_TIMESTAMP | 04-JUL-07 15:33:23.484 |
| EXTRACT(field FROM TIMESTAMP) | DOUBLE PRECISION | Get subfield | EXTRACT(hour FROM TIMESTAMP '2001-02-16 20:38:40') | 20 |
| LAST_DAY(DATE) | DATE | Returns the last day of the month represented by the given date. If the given date contains a time portion, it is carried forward to the result unchanged. | LAST_DAY('14-APR-98') | 30-APR-98 00:00:00 |
| LOCALTIMESTAMP [(precision)] | TIMESTAMP | Current date and time (start of current transaction) | LOCALTIMESTAMP | 04-JUL-07 15:33:23.484 |
| MONTHS_BETWEEN (DATE, DATE) | NUMBER | Number of months between two dates | MONTHS_BETWEEN ('28-FEB-07', '30-NOV-06') | 3 |
| NEXT_DAY(DATE, dayofweek) | DATE | Date falling on dayofweek following specified date | NEXT_DAY('16 (-APR-07', 'FRI') | 20-APR-07 00:00:00 |
| NEW_TIME(DATE, VARCHAR, VARCHAR) | DATE | Converts a date and time to an alternate time zone | NEW_TIME (TO_DATE '2005/05/29 01:45', 'AST', 'PST') | 2005/05/29 21:45:00 |
| NUMTODSINTERVAL(NUMBER, INTERVAL) | INTERVAL | Converts a number to a specified day or second interval | SELECT numtodsinterval(100, 'hour'); | 4 days 04:00:00 |

| | | | | |
|---|-----------|--|--|-----------------------------------|
| NUMTOYMINTERVAL(NUMBER, INTERVAL) | INTERVAL | Converts a number to a specified year or month interval. | SELECT numtoyminterval(100, 'month'); | 8 years 4 mons |
| ROUND(DATE [, format]) | DATE | Date rounded according to format | ROUND(TO_DATE('29-MAY-05'), 'MON') | 01-JUN-05 00:00:00 |
| SYS_EXTRACT_UTCTIMESTAMP WITH TIME ZONE | TIMESTAMP | Returns Coordinated Universal Time | SYS_EXTRACT_UTCTIMESTAMP(CAST ('24-MAR-11 12:30:00PM -04:00' AS TIMESTAMP WITH TIME ZONE)) | 24-MAR-11 16:30:00 |
| SYSDATE | DATE | Returns current date and time | SYSDATE | 01-AUG-12 11:12:34 |
| SYS_TIMESTAMP() | TIMESTAMP | Returns current date and time | SYSTIMESTAMP | 01-AUG-12 11:11:23. 665229 -07:00 |
| TRUNC(DATE [format]) | DATE | Truncate according to format | TRUNC(TO_DATE ('29-MAY-05'), 'MON') | 01-MAY-05 00:00:00 |

ADD_MONTHS

The **ADD_MONTHS** functions adds (or subtracts if the second parameter is negative) the specified number of months to the given date. The resulting day of the month is the same as the day of the month of the given date except when the day is the last day of the month in which case the resulting date always falls on the last day of the month.

Any fractional portion of the number of months parameter is truncated before performing the calculation.

If the given date contains a time portion, it is carried forward to the result unchanged.

The following are examples of the **ADD_MONTHS** function.

```
SELECT ADD_MONTHS('13-JUN-07',4) FROM DUAL;
```

```
add_months
-----
13-OCT-07 00:00:00
(1 row)
```

```
SELECT ADD_MONTHS('31-DEC-06',2) FROM DUAL;
```

```
add_months
-----
28-FEB-07 00:00:00
(1 row)
```

```
SELECT ADD_MONTHS('31-MAY-04',-3) FROM DUAL;
```

```
add_months
-----
```

```
29-FEB-04 00:00:00
(1 row)
```

CURRENT DATE/TIME

Advanced Server provides a number of functions that return values related to the current date and time. These functions all return values based on the start time of the current transaction.

- CURRENT_DATE
- CURRENT_TIMESTAMP
- LOCALTIMESTAMP
- LOCALTIMESTAMP(precision)

`CURRENT_DATE` returns the current date and time based on the start time of the current transaction. The value of `CURRENT_DATE` will not change if called multiple times within a transaction.

```
SELECT CURRENT_DATE FROM DUAL;
```

```
date
```

```
-----
```

```
06-AUG-07
```

`CURRENT_TIMESTAMP` returns the current date and time. When called from a single SQL statement, it will return the same value for each occurrence within the statement. If called from multiple statements within a transaction, may return different values for each occurrence. If called from a function, may return a different value than the value returned by `current_timestamp` in the caller.

```
SELECT CURRENT_TIMESTAMP, CURRENT_TIMESTAMP FROM DUAL;
```

```
current_timestamp | current_timestamp
```

```
-----+-----
```

```
02-SEP-13 17:52:29.261473 +05:00 | 02-SEP-13 17:52:29.261474 +05:00
```

`LOCALTIMESTAMP` can optionally be given a precision parameter which causes the result to be rounded to that many fractional digits in the seconds field. Without a precision parameter, the result is given to the full available precision.

```
SELECT LOCALTIMESTAMP FROM DUAL;
```

```
timestamp
```

```
-----
```

```
06-AUG-07 16:11:35.973
```

```
(1 row)
```

```
SELECT LOCALTIMESTAMP(2) FROM DUAL;
```

```
timestamp
```

```
-----
```

```
06-AUG-07 16:11:44.58
```

```
(1 row)
```

Since these functions return the start time of the current transaction, their values do not change during the transaction. This is considered a feature: the intent is to allow a single transaction to have a consistent notion of the “current” time, so that multiple modifications within the same transaction bear the same time stamp. Other database systems may advance these values more frequently.

EXTRACT

The `EXTRACT` function retrieves subfields such as year or hour from date/time values. The `EXTRACT` function returns values of type `DOUBLE PRECISION`. The following are valid field names:

YEAR

The year field

```
SELECT EXTRACT(YEAR FROM TIMESTAMP '2001-02-16 20:38:40') FROM DUAL;
```

```
date_part
```

```
-----  
2001  
(1 row)
```

MONTH

The number of the month within the year (1 - 12)

```
SELECT EXTRACT(MONTH FROM TIMESTAMP '2001-02-16 20:38:40') FROM DUAL;
```

```
date_part
```

```
-----  
2  
(1 row)
```

DAY

The day (of the month) field (1 - 31)

```
SELECT EXTRACT(DAY FROM TIMESTAMP '2001-02-16 20:38:40') FROM DUAL;
```

```
date_part
```

```
-----  
16  
(1 row)
```

HOUR

The hour field (0 - 23)

```
SELECT EXTRACT(HOUR FROM TIMESTAMP '2001-02-16 20:38:40') FROM DUAL;
```

```
date_part
```

```
-----  
20  
(1 row)
```

MINUTE

The minutes field (0 - 59)

```
SELECT EXTRACT(MINUTE FROM TIMESTAMP '2001-02-16 20:38:40') FROM DUAL;
```

```
date_part
```

```
38
(1 row)
```

SECOND

The seconds field, including fractional parts (0 - 59)

```
SELECT EXTRACT(SECOND FROM TIMESTAMP '2001-02-16 20:38:40') FROM DUAL;
```

```
date_part
-----
40
(1 row)
```

MONTHS_BETWEEN

The **MONTHS_BETWEEN** function returns the number of months between two dates. The result is a numeric value which is positive if the first date is greater than the second date or negative if the first date is less than the second date.

The result is always a whole number of months if the day of the month of both date parameters is the same, or both date parameters fall on the last day of their respective months.

The following are some examples of the **MONTHS_BETWEEN** function.

```
SELECT MONTHS_BETWEEN('15-DEC-06','15-OCT-06') FROM DUAL;

months_between
-----
2
(1 row)
```

```
SELECT MONTHS_BETWEEN('15-OCT-06','15-DEC-06') FROM DUAL;

months_between
-----
-2
(1 row)
```

```
SELECT MONTHS_BETWEEN('31-JUL-00','01-JUL-00') FROM DUAL;

months_between
-----
0.967741935
(1 row)
```

```
SELECT MONTHS_BETWEEN('01-JAN-07','01-JAN-06') FROM DUAL;

months_between
-----
12
(1 row)
```

NEXT_DAY

The `NEXT_DAY` function returns the first occurrence of the given weekday strictly greater than the given date. At least the first three letters of the weekday must be specified - e.g., `SAT`. If the given date contains a time portion, it is carried forward to the result unchanged.

The following are examples of the `NEXT_DAY` function.

```
SELECT NEXT_DAY(TO_DATE('13-AUG-07','DD-MON-YY'),'SUNDAY') FROM DUAL;
```

```
next_day
-----
19-AUG-07 00:00:00
(1 row)
```

```
SELECT NEXT_DAY(TO_DATE('13-AUG-07','DD-MON-YY'),'MON') FROM DUAL;
```

```
next_day
-----
20-AUG-07 00:00:00
(1 row)
```

NEW_TIME

The `NEW_TIME` function converts a date and time from one time zone to another. `NEW_TIME` returns a value of type `DATE`. The syntax is:

```
NEW_TIME(DATE, time_zone1, time_zone2)
```

`time_zone1` and `time_zone2` must be string values from the Time Zone column of the following table:

| Time Zone | Offset from UTC | Description |
|-----------|-----------------|-----------------------------|
| AST | UTC+4 | Atlantic Standard Time |
| ADT | UTC+3 | Atlantic Daylight Time |
| BST | UTC+11 | Bering Standard Time |
| BDT | UTC+10 | Bering Daylight Time |
| CST | UTC+6 | Central Standard Time |
| CDT | UTC+5 | Central Daylight Time |
| EST | UTC+5 | Eastern Standard Time |
| EDT | UTC+4 | Eastern Daylight Time |
| GMT | UTC | Greenwich Mean Time |
| HST | UTC+10 | Alaska-Hawaii Standard Time |
| HDT | UTC+9 | Alaska-Hawaii Daylight Time |
| MST | UTC+7 | Mountain Standard Time |
| MDT | UTC+6 | Mountain Daylight Time |
| NST | UTC+3:30 | Newfoundland Standard Time |
| PST | UTC+8 | Pacific Standard Time |
| PDT | UTC+7 | Pacific Daylight Time |
| YST | UTC+9 | Yukon Standard Time |
| YDT | UTC+8 | Yukon Daylight Time |

Following is an example of the `NEW_TIME` function:

```
SELECT NEW_TIME(TO_DATE('08-13-07 10:35:15','MM-DD-YY HH24:MI:SS'),'AST',
'PST') "Pacific Standard Time" FROM DUAL;
```

Pacific Standard Time

```
-----  
13-AUG-07 06:35:15  
(1 row)
```

NUMTODSINTERVAL

The `NUMTODSINTERVAL` function converts a numeric value to a time interval that includes day through second interval units. When calling the function, specify the smallest fractional interval type to be included in the result set. The valid interval types are `DAY`, `HOUR`, `MINUTE`, and `SECOND`.

The following example converts a numeric value to a time interval that includes days and hours:

```
SELECT numtodsinterval(100, 'hour');  
numtodsinterval  
-----  
4 days 04:00:00  
(1 row)
```

The following example converts a numeric value to a time interval that includes minutes and seconds:

```
SELECT numtodsinterval(100, 'second');  
numtodsinterval  
-----  
1 min 40 secs  
(1 row)
```

NUMTOYMINTEGRAL

The `NUMTOYMINTEGRAL` function converts a numeric value to a time interval that includes year through month interval units. When calling the function, specify the smallest fractional interval type to be included in the result set. The valid interval types are `YEAR` and `MONTH`.

The following example converts a numeric value to a time interval that includes years and months:

```
SELECT numtoyminterval(100, 'month');  
numtoyminterval  
-----  
8 years 4 mons  
(1 row)
```

The following example converts a numeric value to a time interval that includes years only:

```
SELECT numtoyminterval(100, 'year');  
numtoyminterval  
-----  
100 years  
(1 row)
```

ROUND

The **ROUND** function returns a date rounded according to a specified template pattern. If the template pattern is omitted, the date is rounded to the nearest day. The following table shows the template patterns for the **ROUND** function.

| Pattern | Description |
|---|--|
| CC, SCC | Returns January 1, cc01 where cc is first 2 digits of the given year if last 2 digits <= 50, or 1 greater than the first 2 digits of the given year if last 2 digits > 50; (for AD years) |
| YYYY, YYYY, YEAR, SYEAR, YYY, YY, Y | Returns January 1, yyyy where yyyy is rounded to the nearest year; rounds down on June 30, rounds up on July 1 |
| IYYY, IYY, IY, I | Rounds to the beginning of the ISO year which is determined by rounding down if the month and day is on or before June 30th, or by rounding up if the month and day is July 1st or later |
| Q | Returns the first day of the quarter determined by rounding down if the month and day is on or before the 15th of the second month of the quarter, or by rounding up if the month and day is on the 16th of the second month or later of the quarter |
| MONTH, MON, MM, RM | Returns the first day of the specified month if the day of the month is on or prior to the 15th; returns the first day of the following month if the day of the month is on the 16th or later |
| WW | Round to the nearest date that corresponds to the same day of the week as the first day of the year |
| IW | Round to the nearest date that corresponds to the same day of the week as the first day of the ISO year |
| W | Round to the nearest date that corresponds to the same day of the week as the first day of the month |
| DDD, DD, J | Rounds to the start of the nearest day; 11:59:59 AM or earlier rounds to the start of the same day; 12:00:00 PM or later rounds to the start of the next day |
| DAY, DY, D | Rounds to the nearest Sunday |
| HH, HH12, HH24 | Round to the nearest hour |
| MI | Round to the nearest minute |

Following are examples of usage of the **ROUND** function.

The following examples round to the nearest hundred years.

```
SELECT TO_CHAR(ROUND(TO_DATE('1950','YYYY'),'CC'),'DD-MON-YYYY') "Century"
FROM DUAL;
```

```
Century
-----
01-JAN-1901
(1 row)
```

```
SELECT TO_CHAR(ROUND(TO_DATE('1951','YYYY'),'CC'),'DD-MON-YYYY') "Century"
FROM DUAL;
```

```
Century
-----
01-JAN-2001
(1 row)
```

The following examples round to the nearest year.

```
SELECT TO_CHAR(ROUND(TO_DATE('30-JUN-1999','DD-MON-YYYY'),'Y'),'DD-MON-
```

```
YYYY') "Year" FROM DUAL;
```

Year

```
01-JAN-1999
(1 row)
```

```
SELECT TO_CHAR(ROUND(TO_DATE('01-JUL-1999','DD-MON-YYYY'),'Y'),'DD-MON-
YYYY') "Year" FROM DUAL;
```

Year

```
01-JAN-2000
(1 row)
```

The following examples round to the nearest ISO year. The first example rounds to 2004 and the ISO year for 2004 begins on December 29th of 2003. The second example rounds to 2005 and the ISO year for 2005 begins on January 3rd of that same year.

(An ISO year begins on the first Monday from which a 7 day span, Monday thru Sunday, contains at least 4 days of the new year. Thus, it is possible for the beginning of an ISO year to start in December of the prior year.)

```
SELECT TO_CHAR(ROUND(TO_DATE('30-JUN-2004','DD-MON-YYYY'),'IYYY'),'DD-MON-
YYYY') "ISO Year" FROM DUAL;
```

ISO Year

```
29-DEC-2003
(1 row)
```

```
SELECT TO_CHAR(ROUND(TO_DATE('01-JUL-2004','DD-MON-YYYY'),'IYYY'),'DD-MON-
YYYY') "ISO Year" FROM DUAL;
```

ISO Year

```
03-JAN-2005
(1 row)
```

The following example round to the nearest quarter:

```
SELECT ROUND(TO_DATE('15-FEB-07','DD-MON-YY'),'Q') "Quarter" FROM DUAL;
```

Quarter

```
01-JAN-07 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('16-FEB-07','DD-MON-YY'),'Q') "Quarter" FROM DUAL;
```

Quarter

```
01-APR-07 00:00:00
(1 row)
```

The following example round to the nearest month:

```
SELECT ROUND(TO_DATE('15-DEC-07','DD-MON-YY'),'MONTH') "Month" FROM DUAL;
```

Month

```
01-DEC-07 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('16-DEC-07','DD-MON-YY'),'MONTH') "Month" FROM DUAL;
```

Month

```
01-JAN-08 00:00:00
(1 row)
```

The following examples round to the nearest week. The first day of 2007 lands on a Monday so in the first example, January 18th is closest to the Monday that lands on January 15th. In the second example, January 19th is closer to the Monday that falls on January 22nd.

```
SELECT ROUND(TO_DATE('18-JAN-07','DD-MON-YY'),'WW') "Week" FROM DUAL;
```

Week

```
15-JAN-07 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('19-JAN-07','DD-MON-YY'),'WW') "Week" FROM DUAL;
```

Week

```
22-JAN-07 00:00:00
(1 row)
```

The following examples round to the nearest ISO week. An ISO week begins on a Monday. In the first example, January 1, 2004 is closest to the Monday that lands on December 29, 2003. In the second example, January 2, 2004 is closer to the Monday that lands on January 5, 2004.

```
SELECT ROUND(TO_DATE('01-JAN-04','DD-MON-YY'),'IW') "ISO Week" FROM DUAL;
```

ISO Week

```
29-DEC-03 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('02-JAN-04','DD-MON-YY'),'IW') "ISO Week" FROM DUAL;
```

ISO Week

```
05-JAN-04 00:00:00
(1 row)
```

The following examples round to the nearest week where a week is considered to start on the same day as the first day of the month.

```
SELECT ROUND(TO_DATE('05-MAR-07','DD-MON-YY'),'W') "Week" FROM DUAL;
```

Week

```
-----
08-MAR-07 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('04-MAR-07','DD-MON-YY'),'W') "Week" FROM DUAL;
```

Week

```
-----
01-MAR-07 00:00:00
(1 row)
```

The following examples round to the nearest day.

```
SELECT ROUND(TO_DATE('04-AUG-07 11:59:59 AM','DD-MON-YY HH:MI:SS AM'),'J')
"Day" FROM DUAL;
```

Day

```
-----
04-AUG-07 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('04-AUG-07 12:00:00 PM','DD-MON-YY HH:MI:SS AM'),'J')
"Day" FROM DUAL;
```

Day

```
-----
05-AUG-07 00:00:00
(1 row)
```

The following examples round to the start of the nearest day of the week (Sunday).

```
SELECT ROUND(TO_DATE('08-AUG-07','DD-MON-YY'),'DAY') "Day of Week" FROM
DUAL;
```

Day of Week

```
-----
05-AUG-07 00:00:00
(1 row)
```

```
SELECT ROUND(TO_DATE('09-AUG-07','DD-MON-YY'),'DAY') "Day of Week" FROM
DUAL;
```

Day of Week

```
-----
12-AUG-07 00:00:00
(1 row)
```

The following examples round to the nearest hour.

```
SELECT TO_CHAR(ROUND(TO_DATE('09-AUG-07 08:29','DD-MON-YY HH:MI'),'HH'),'DD-
MON-YY HH24:MI:SS') "Hour" FROM DUAL;
```

Hour

```
-----
09-AUG-07 08:00:00
(1 row)
```

```
SELECT TO_CHAR(ROUND(TO_DATE('09-AUG-07 08:30','DD-MON-YY HH:MI'),'HH'),'DD-MON-YY HH24:MI:SS') "Hour" FROM DUAL;
```

Hour

09-AUG-07 09:00:00

(1 row)

The following examples round to the nearest minute.

```
SELECT TO_CHAR(ROUND(TO_DATE('09-AUG-07 08:30:29','DD-MON-YY HH:MI:SS'),'MI'),'DD-MON-YY HH24:MI:SS') "Minute" FROM DUAL;
```

Minute

09-AUG-07 08:30:00

(1 row)

```
SELECT TO_CHAR(ROUND(TO_DATE('09-AUG-07 08:30:30','DD-MON-YY HH:MI:SS'),'MI'),'DD-MON-YY HH24:MI:SS') "Minute" FROM DUAL;
```

Minute

09-AUG-07 08:31:00

(1 row)

SYSDATE

The **SYSDATE** function returns the current date and time (timestamp without timezone) of the operating system on which the database server resides. The function is **STABLE** and requires no arguments.

When called from a single SQL statement, it will return the same value for each occurrence within the statement. If called from multiple statements within a transaction, may return different values for each occurrence. If called from a function, may return a different value than the value returned by **SYSDATE** in the caller.

The following example demonstrates a call to **SYSDATE**:

```
SELECT SYSDATE, SYSDATE FROM DUAL;
```

| sysdate | | sysdate |
|---------|--|---------|
|---------|--|---------|

28-APR-20 16:45:28 | 28-APR-20 16:45:28

(1 row)

TRUNC

The **TRUNC** function returns a date truncated according to a specified template pattern. If the template pattern is omitted, the date is truncated to the nearest day. The following table shows the template patterns for the **TRUNC** function.

| Pattern | Description |
|---------|---|
| CC, SCC | Returns January 1, cc 01 where cc is first 2 digits of the given year |

| | |
|---------------------------------------|--|
| YYYY, YYYY, YEAR, SYEAR, YY, YY, Y | Returns January 1, yyyy where yyyy is the given year |
| IYYY, IYY, IY, I | Returns the start date of the ISO year containing the given date |
| Q | Returns the first day of the quarter containing the given date |
| MONTH, MON, MM, RM | Returns the first day of the specified month |
| WW | Returns the largest date just prior to, or the same as the given date that corresponds to the same day of the week as the first day of the year |
| IW | Returns the start of the ISO week containing the given date |
| W | Returns the largest date just prior to, or the same as the given date that corresponds to the same day of the week as the first day of the month |
| DDD, DD, J | Returns the start of the day for the given date |
| DAY, DY, D | Returns the start of the week (Sunday) containing the given date |
| HH, HH12, HH24 | Returns the start of the hour |
| MI | Returns the start of the minute |

Following are examples of usage of the `TRUNC` function.

The following example truncates down to the hundred years unit.

```
SELECT TO_CHAR(TRUNC(TO_DATE('1951','YYYY'),'CC'),'DD-MON-YYYY') "Century"
FROM DUAL;
```

Century

```
-----
01-JAN-1901
(1 row)
```

The following example truncates down to the year.

```
SELECT TO_CHAR(TRUNC(TO_DATE('01-JUL-1999','DD-MON-YYYY'),'Y'),'DD-MON-
YYYY') "Year" FROM DUAL;
```

Year

```
-----
01-JAN-1999
(1 row)
```

The following example truncates down to the beginning of the ISO year.

```
SELECT TO_CHAR(TRUNC(TO_DATE('01-JUL-2004','DD-MON-YYYY'),'IYYY'),'DD-MON-
YYYY') "ISO Year" FROM DUAL;
```

ISO Year

```
-----
29-DEC-2003
(1 row)
```

The following example truncates down to the start date of the quarter.

```
SELECT TRUNC(TO_DATE('16-FEB-07','DD-MON-YY'),'Q') "Quarter" FROM DUAL;
```

Quarter

```
-----
01-JAN-07 00:00:00
(1 row)
```

The following example truncates to the start of the month.

```
SELECT TRUNC(TO_DATE('16-DEC-07','DD-MON-YY'),'MONTH') "Month" FROM DUAL;
```

Month

```
01-DEC-07 00:00:00
(1 row)
```

The following example truncates down to the start of the week determined by the first day of the year. The first day of 2007 lands on a Monday so the Monday just prior to January 19th is January 15th.

```
SELECT TRUNC(TO_DATE('19-JAN-07','DD-MON-YY'),'WW') "Week" FROM DUAL;
```

Week

```
15-JAN-07 00:00:00
(1 row)
```

The following example truncates to the start of an ISO week. An ISO week begins on a Monday. January 2, 2004 falls in the ISO week that starts on Monday, December 29, 2003.

```
SELECT TRUNC(TO_DATE('02-JAN-04','DD-MON-YY'),'IW') "ISO Week" FROM DUAL;
```

ISO Week

```
29-DEC-03 00:00:00
(1 row)
```

The following example truncates to the start of the week where a week is considered to start on the same day as the first day of the month.

```
SELECT TRUNC(TO_DATE('21-MAR-07','DD-MON-YY'),'W') "Week" FROM DUAL;
```

Week

```
15-MAR-07 00:00:00
(1 row)
```

The following example truncates to the start of the day.

```
SELECT TRUNC(TO_DATE('04-AUG-07 12:00:00 PM','DD-MON-YY HH:MI:SS AM'),'J')
"Day" FROM DUAL;
```

Day

```
04-AUG-07 00:00:00
(1 row)
```

The following example truncates to the start of the week (Sunday).

```
SELECT TRUNC(TO_DATE('09-AUG-07','DD-MON-YY'),'DAY') "Day of Week" FROM
DUAL;
```

Day of Week

```
05-AUG-07 00:00:00
```

```
(1 row)
```

The following example truncates to the start of the hour.

```
SELECT TO_CHAR(TRUNC(TO_DATE('09-AUG-07 08:30','DD-MON-YY HH:MI'),'HH'),'DD-MON-YY HH24:MI:SS') "Hour" FROM DUAL;
```

```
Hour
```

```
-----
```

```
09-AUG-07 08:00:00
```

```
(1 row)
```

The following example truncates to the minute.

```
SELECT TO_CHAR(TRUNC(TO_DATE('09-AUG-07 08:30:30','DD-MON-YY HH:MI:SS'),'MI'),'DD-MON-YY HH24:MI:SS') "Minute" FROM DUAL;
```

```
Minute
```

```
-----
```

```
09-AUG-07 08:30:00
```

```
(1 row)
```

5.2.3.9 Sequence Manipulation Functions

This section describes Advanced Server's functions for operating on sequence objects. Sequence objects (also called sequence generators or just sequences) are special single-row tables created with the [CREATE SEQUENCE](#) command. A sequence object is usually used to generate unique identifiers for rows of a table. The sequence functions, listed below, provide simple, multiuser-safe methods for obtaining successive sequence values from sequence objects.

```
sequence.NEXTVAL
```

```
sequence.CURRVAL
```

sequence is the identifier assigned to the sequence in the [CREATE SEQUENCE](#) command. The following describes the usage of these functions.

NEXTVAL

Advance the sequence object to its next value and return that value. This is done atomically: even if multiple sessions execute **NEXTVAL** concurrently, each will safely receive a distinct sequence value.

CURRVAL

Return the value most recently obtained by **NEXTVAL** for this sequence in the current session. (An error is reported if **NEXTVAL** has never been called for this sequence in this session.) Notice that because this is returning a session-local value, it gives a predictable answer whether or not other sessions have executed **NEXTVAL** since the current session did.

If a sequence object has been created with default parameters, **NEXTVAL** calls on it will return successive values beginning with 1. Other behaviors can be obtained by using special parameters in the [CREATE SEQUENCE](#) command.

Important: To avoid blocking of concurrent transactions that obtain numbers from the same sequence, a `NEXTVAL` operation is never rolled back; that is, once a value has been fetched it is considered used, even if the transaction that did the `NEXTVAL` later aborts. This means that aborted transactions may leave unused "holes" in the sequence of assigned values.

5.2.3.10 Conditional Expressions

The following section describes the SQL-compliant conditional expressions available in Advanced Server.

CASE

The SQL `CASE` expression is a generic conditional expression, similar to if/else statements in other languages:

```
CASE WHEN condition THEN result
      [ WHEN ... ]
      [ ELSE result ]
END
```

`CASE` clauses can be used wherever an expression is valid. `condition` is an expression that returns a `BOOLEAN` result. If the result is `TRUE` then the value of the `CASE` expression is the `result` that follows the condition. If the result is `FALSE` any subsequent `WHEN` clauses are searched in the same manner. If `no WHEN condition` is `TRUE` then the value of the `CASE` expression is the `result` in the `ELSE` clause. If the `ELSE` clause is omitted and no condition matches, the result is `null`.

An example:

```
SELECT * FROM test;
```

```
a
---
1
2
3
(3 rows)
```

```
SELECT a,
CASE WHEN a=1 THEN 'one'
      WHEN a=2 THEN 'two'
      ELSE 'other'
END
FROM test;
```

```
a | case
---+-----
1 | one
2 | two
3 | other
(3 rows)
```

The data types of all the `result` expressions must be convertible to a single output type.

The following “simple” `CASE` expression is a specialized variant of the general form above:

```
CASE expression
  WHEN value THEN result
  [ WHEN ... ]
  [ ELSE result ]
END
```

The `expression` is computed and compared to all the `value` specifications in the `WHEN` clauses until one is found that is equal. If no match is found, the `result` in the `ELSE` clause (or a null value) is returned.

The example above can be written using the simple `CASE` syntax:

```
SELECT a,
CASE a WHEN 1 THEN 'one'
      WHEN 2 THEN 'two'
      ELSE 'other'
END
```

`FROM test;`

```
a | case
---+-----
1 | one
2 | two
3 | other
(3 rows)
```

A `CASE` expression does not evaluate any subexpressions that are not needed to determine the result. For example, this is a possible way of avoiding a division-by-zero failure:

```
| SELECT ... WHERE CASE WHEN x <> 0 THEN y/x > 1.5 ELSE false END;
```

COALESCE

The `COALESCE` function returns the first of its arguments that is not null. Null is returned only if all arguments are null.

```
COALESCE(value [, value2] ... )
```

It is often used to substitute a default value for null values when data is retrieved for display or further computation. For example:

```
| SELECT COALESCE(description, short_description, '(none)') ...
```

Like a `CASE` expression, `COALESCE` will not evaluate arguments that are not needed to determine the result; that is, arguments to the right of the first non-null argument are not evaluated. This SQL-standard function provides capabilities similar to `NVL` and `IFNULL`, which are used in some other database systems.

NULLIF

The `NULLIF` function returns a null value if `value1` and `value2` are equal; otherwise it returns `value1`.

```
NULLIF(value1, value2)
```

This can be used to perform the inverse operation of the `COALESCE` example given above:

```
| SELECT NULLIF(value1, '(none)') ...
```

If `value1` is (none), return a null, otherwise return `value1`.

NVL

The `NVL` function returns the first of its arguments that is not null. `NVL` evaluates the first expression; if that expression evaluates to `NULL`, `NVL` returns the second expression.

```
NVL(expr1, expr2)
```

The return type is the same as the argument types; all arguments must have the same data type (or be coercible to a common type). `NVL` returns `NULL` if all arguments are `NULL`.

The following example computes a bonus for non-commissioned employees. If an employee is a commissioned employee, this expression returns the employees commission; if the employee is not a commissioned employee (that is, his commission is `NULL`), this expression returns a bonus that is 10% of his salary.

```
bonus = NVL(emp.commission, emp.salary * .10)
```

NVL2

`NVL2` evaluates an expression, and returns either the second or third expression, depending on the value of the first expression. If the first expression is not `NULL`, `NVL2` returns the value in `expr2`; if the first expression is `NULL`, `NVL2` returns the value in `expr3`.

```
NVL2(expr1, expr2, expr3)
```

The return type is the same as the argument types; all arguments must have the same data type (or be coercible to a common type).

The following example computes a bonus for commissioned employees - if a given employee is a commissioned employee, this expression returns an amount equal to 110% of his commission; if the employee is not a commissioned employee (that is, his commission is `NULL`), this expression returns `0`.

```
bonus = NVL2(emp.commission, emp.commission * 1.1, 0)
```

GREATEST and LEAST

The `GREATEST` and `LEAST` functions select the largest or smallest value from a list of any number of expressions.

```
GREATEST(value [, value2] ...) 
LEAST(value [, value2] ...) 
```

The expressions must all be convertible to a common data type, which will be the type of the result. Null values in the list are ignored. The result will be null only if all the expressions evaluate to null.

Note

The `GREATEST` and `LEAST` are not in the SQL standard, but are a common extension.

5.2.3.11 Aggregate Functions

Aggregate functions compute a single result value from a set of input values. The built-in aggregate functions are listed in the following tables.

| Function | Argument Type | Return Type | Description |
|------------------------|---|--|--|
| AVG (expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | NUMBER for any integer type, DOUBLE PRECISION for a floating-point argument, otherwise the same as the argument data type | The average (arithmetic mean) of all input values |
| COUNT(*) | | BIGINT | Number of input rows |
| COUNT (expression) | Any | BIGINT | Number of input rows for which the value of expression is not null |
| MAX (expression) | Any numeric, string, date/time, or bytea type | Same as argument type | Maximum value of expression across all input values |
| MEDIAN (expression) | BIGINT, DOUBLE PRECISION, INTEGER, INTERVAL, NUMERIC, REAL, SMALLINT, TIMESTAMP, TIMESTAMPZ | NUMBER for any integer type, DOUBLE PRECISION for a floating-point argument, otherwise the same as the argument data type | Identifies the middle value of an expression |
| MIN (expression) | Any numeric, string, date/time, or bytea type | Same as argument type | Minimum value of expression across all input values |
| SUM (expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | BIGINT for SMALLINT or INTEGER arguments, NUMBER for BIGINT arguments, DOUBLE PRECISION for floating-point arguments, otherwise the same as the argument data type | Sum of expression across all input values |

It should be noted that except for `COUNT`, these functions return a null value when no rows are selected. In particular, `SUM` of no rows returns null, not zero as one might expect. The `COALESCE` function may be used to substitute zero for null when necessary.

The following table shows the aggregate functions typically used in statistical analysis. (These are separated out merely to avoid cluttering the listing of more-commonly-used aggregates.) Where the description mentions `N`, it means the number of input rows for which all the input expressions are non-null. In all cases, null is returned if the computation is meaningless, for example when `N` is zero.

| Function | Argument Type | Return Type | Description |
|-------------------------------|------------------|------------------|---|
| <code>CORR(Y, X)</code> | DOUBLE PRECISION | DOUBLE PRECISION | Correlation coefficient |
| <code>COVAR_POP(Y, X)</code> | DOUBLE PRECISION | DOUBLE PRECISION | Population covariance |
| <code>COVAR_SAMP(Y, X)</code> | DOUBLE PRECISION | DOUBLE PRECISION | Sample covariance |
| <code>REGR_AVGX(Y, X)</code> | DOUBLE PRECISION | DOUBLE PRECISION | Average of the independent variable ($\text{sum}(X) / N$) |

| | | | |
|-------------------------|---|---|---|
| REGR_AVGY(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | Average of the dependent variable ($\text{sum}(Y) / N$) |
| REGR_COUNT(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | Number of input rows in which both expressions are nonnull |
| REGR_INTERCEPT(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | y-intercept of the least-squares-fit linear equation determined by the (X, Y) pairs |
| REGR_R2(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | Square of the correlation coefficient |
| REGR_SLOPE(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | Slope of the least-squares-fit linear equation determined by the (X, Y) pairs |
| REGR_SXX(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | $\text{Sum } (X^2) - \text{sum } (X)^2 / N$ ("sum of squares" of the independent variable) |
| REGR_SXY(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | $\text{Sum } (X*Y) - \text{sum } (X) * \text{sum } (Y) / N$ ("sum of products" of independent times dependent variable) |
| REGR_SYY(Y, X) | DOUBLE PRECISION | DOUBLE PRECISION | $\text{Sum } (Y^2) - \text{sum } (Y)^2 / N$ ("sum of squares" of the dependent variable) |
| STDDEV(expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | DOUBLE PRECISION for floating-point arguments, otherwise NUMBER | Historic alias for STDDEV_SAMP |
| STDDEV_POP(expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | DOUBLE PRECISION for floating-point arguments, otherwise NUMBER | Population standard deviation of the input values |
| STDDEV_SAMP(expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | DOUBLE PRECISION for floating-point arguments, otherwise NUMBER | Sample standard deviation of the input values |
| VARIANCE(expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | DOUBLE PRECISION for floating-point arguments, otherwise NUMBER | Historical alias for VAR_SAMP |
| VAR_POP(expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | DOUBLE PRECISION for floating-point arguments, otherwise NUMBER | Population variance of the input values (square of the population standard deviation) |
| VAR_SAMP(expression) | INTEGER, REAL, DOUBLE PRECISION, NUMBER | DOUBLE PRECISION for floating-point arguments, otherwise NUMBER | Sample variance of the input values (square of the sample standard deviation) |

LISTAGG

Advanced Server has added the `LISTAGG` function to support string aggregation. `LISTAGG` is an aggregate function that concatenates data from multiple rows into a single row in an ordered manner. You can optionally include a custom delimiter for your data.

The `LISTAGG` function mandates the use of an `ORDER BY` clause under a `WITHIN GROUP` clause to concatenate values of the measure column, and then generate the ordered aggregated data.

Objective

- `LISTAGG` can be used without any grouping. In this case, the `LISTAGG` function operates on all rows in a table and returns a single row.
- `LISTAGG` can be used with the `GROUP BY` clause. In this case, the `LISTAGG` function operates on each group and returns an aggregated output for each group.
- `LISTAGG` can be used with the `OVER` clause. In this case, the `LISTAGG` function partitions a query result set into groups based on the expression in the `query_partition_by_clause` and then aggregates data in each group.

Synopsis

```
LISTAGG( <measure_expr> [, <delimiter> ] ) WITHIN GROUP( <order_by_clause> )
[ OVER <query_partition_by_clause> ]
```

Parameters

`measure_expr`

`measure_expr` (mandatory) specifies the column or expression that assigns a value to aggregate. `NULL` values are ignored.

`delimiter`

`delimiter` (optional) specifies a string that separates the concatenated values in the result row. The `delimiter` can be a `NULL` value, string, character literal, column name, or constant expression. If ignored, the `LISTAGG` function uses a `NULL` value by default.

`order_by_clause`

`order_by_clause` (mandatory) determines the sort order in which the concatenated values are returned.

`query_partition_by_clause`

`query_partition_by_clause` (optional) allows `LISTAGG` function to be used as an analytic function and sets the range of records for each group in the `OVER` clause.

Return Type

The `LISTAGG` function returns a string value.

Examples

The following example concatenates the values in the `EMP` table and lists all the employees separated by a `delimiter` comma.

First, create a table named `EMP` and then insert records into the `EMP` table.

```
edb=# CREATE TABLE EMP
edb-#   (EMPNO NUMBER(4) NOT NULL,
edb-#     ENAME VARCHAR2(10),
edb-#     JOB VARCHAR2(9),
edb-#     MGR NUMBER(4),
edb-#     HIREDATE DATE,
edb-#     SAL NUMBER(7, 2),
edb-#     COMM NUMBER(7, 2),
edb-#     DEPTNO NUMBER(2));
CREATE TABLE

edb=# INSERT INTO EMP VALUES
```

```

edb=#      (7499, 'ALLEN', 'SALESMAN', 7698,
edb(#      TO_DATE('20-FEB-1981', 'DD-MON-YYYY'), 1600, 300, 30);
INSERT 0 1
edb=# INSERT INTO EMP VALUES
edb-#      (7521, 'WARD', 'SALESMAN', 7698,
edb(#      TO_DATE('22-FEB-1981', 'DD-MON-YYYY'), 1250, 500, 30);
INSERT 0 1
edb=# INSERT INTO EMP VALUES
edb-#      (7566, 'JONES', 'MANAGER', 7839,
edb(#      TO_DATE('2-APR-1981', 'DD-MON-YYYY'), 2975, NULL, 20);
INSERT 0 1
edb=# INSERT INTO EMP VALUES
edb-#      (7654, 'MARTIN', 'SALESMAN', 7698,
edb(#      TO_DATE('28-SEP-1981', 'DD-MON-YYYY'), 1250, 1400, 30);
INSERT 0 1
edb=# INSERT INTO EMP VALUES
edb-#      (7698, 'BLAKE', 'MANAGER', 7839,
edb(#      TO_DATE('1-MAY-1981', 'DD-MON-YYYY'), 2850, NULL, 30);
INSERT 0 1

edb=# SELECT LISTAGG(ENAME, ',') WITHIN GROUP (ORDER BY ENAME) FROM EMP;
listagg
-----
ALLEN,BLAKE,JONES,MARTIN,WARD
(1 row)

```

The following example uses `PARTITION BY` clause with `LISTAGG` in `EMP` table and generates output based on a partition by `DEPTNO` that applies to each partition and not on the entire table.

```

edb=# SELECT DISTINCT DEPTNO, LISTAGG(ENAME, ',') WITHIN GROUP (ORDER BY
ENAME) OVER(PARTITION BY DEPTNO) FROM EMP;
deptno |      listagg
-----+
 30 | ALLEN,BLAKE,MARTIN,WARD
 20 | JONES
(2 rows)

```

The following example is identical to the previous example, except it includes the `GROUP BY` clause.

```

edb=# SELECT DEPTNO, LISTAGG(ENAME, ',') WITHIN GROUP (ORDER BY ENAME) FROM
EMP GROUP BY DEPTNO;
deptno |      listagg
-----+
 20 | JONES
 30 | ALLEN,BLAKE,MARTIN,WARD
(2 rows)

```

MEDIAN

The `MEDIAN` function calculates the middle value of an expression from a given range of values; `NULL` values are ignored. The `MEDIAN` function returns an error if a query does not reference the user-defined table.

Objective

- `MEDIAN` can be used without any grouping. In this case, the `MEDIAN` function operates on all rows in a

table and returns a single row.

- **MEDIAN** can be used with the **OVER** clause. In this case, the **MEDIAN** function partitions a query result set into groups based on the **expression** specified in the **PARTITION BY** clause and then aggregates data in each group.

Synopsis

```
MEDIAN( <median_expression> ) [ OVER ( [ PARTITION BY... ] ) ]
```

Parameters

median_expression

median_expression (mandatory) is a target column or expression that the **MEDIAN** function operates on and returns a median value. It can be a numeric, datetime, or interval data type.

PARTITION BY

PARTITION BY clause (optional) allows a **MEDIAN** function to be used as an analytic function and sets the range of records for each group in the **OVER** clause.

Return Types

The return type is determined by the input data type of **expression**. The following table illustrates the return type for each input type.

| Input Type | Return Type |
|-------------------------|------------------|
| BIGINT | NUMERIC |
| FLOAT, DOUBLE PRECISION | DOUBLE PRECISION |
| INTEGER | NUMERIC |
| INTERVAL | INTERVAL |
| NUMERIC | NUMERIC |
| REAL | REAL |
| SMALLINT | NUMERIC |
| TIMESTAMP | TIMESTAMP |
| TIMESTAMPTZ | TIMESTAMPTZ |

Examples

In the following example, a query returns the median salary for each department in the **EMP** table:

```
edb=# SELECT * FROM EMP;
empno|ename|job|mgr|hiredate|sal|comm|deptno
-----+-----+-----+-----+-----+-----+-----+
7369 |SMITH|CLERK|7902|17-DEC-80 00:00:00|800.00||20
7499 |ALLEN|SALESMAN|7698|20-FEB-81 00:00:00|1600.00|300.0|30
7521 |WARD|SALESMAN|7698|22-FEB-81 00:00:00|1250.00|500.00|30
7566 |JONES|MANAGER|7839|02-APR-81 00:00:00|2975.00||20
7654 |MARTIN|SALESMAN|7698|28-SEP-81 00:00:00|1250.00|1400.00|30
(5 rows)
```

```
edb=# SELECT MEDIAN(SAL) FROM EMP;
```

```
median
```

```
-----+
1250
(1 row)
```

The following example uses `PARTITION BY` clause with `MEDIAN` in `EMP` table and returns the median salary based on a partition by `DEPTNO`:

```
edb=# SELECT EMPNO, ENAME, DEPTNO, MEDIAN (SAL) OVER (PARTITION BY DEPTNO)
FROM EMP;
empno | ename | deptno | median
-----+-----+-----+
7369 | SMITH | 20 | 1887.5
7566 | JONES | 20 | 1887.5
7499 | ALLEN | 30 | 1250
7521 | WARD | 30 | 1250
7654 | MARTIN | 30 | 1250
(5 rows)
```

The `MEDIAN` function can be compared with `PERCENTILE_CONT`. In the following example, `MEDIAN` generates the same result as `PERCENTILE_CONT`:

```
edb=# SELECT MEDIAN (SAL), PERCENTILE_CONT(0.5) WITHIN GROUP(ORDER BY SAL)
FROM EMP;
median | percentile_cont
-----+-----
1250 |      1250
(1 row)
```

STATS_MODE

The `STATS_MODE` function takes a set of values as an argument and returns the value that occurs with the highest frequency. If multiple values are appearing with the same frequency, the `STATS_MODE` function arbitrarily chooses the first value and returns only that one value.

Objective

- `STATS_MODE` function can be used without any grouping. In this case, the `STATS_MODE` function operates on all the rows in a table and returns a single value.
- `STATS_MODE` can be used as an ordered-set aggregate function using the `WITHIN GROUP` clause. In this case, the `STATS_MODE` function operates on the ordered data set.
- `STATS_MODE` can be used with the `GROUP BY` clause. In this case, the `STATS_MODE` function operates on each group and returns the most frequent and aggregated output for each group.

Synopsis

`STATS_MODE(<expr>)`

OR

`STATS_MODE() WITHIN GROUP (ORDER BY sort_expression)`

Parameters

`expr`

An expression or value to assign to the column.

Return Type

The `STATS_MODE` function returns a value that appears frequently. However, if all the values of a column are `NULL`, the `STATS_MODE` returns `NULL`.

Examples

The following example returns the mode of salary in the `EMP` table:

```
edb=# SELECT * FROM EMP;
empno|ename|job|mgr|hiredate|sal|comm|deptno
-----+-----+-----+-----+-----+-----+
7369 |SMITH|CLERK|7902|17-DEC-80 00:00:00|800.00|| 20
7499 |ALLEN|SALESMAN|7698|20-FEB-81 00:00:00|1600.00|300.0| 30
7521 |WARD|SALESMAN|7698|22-FEB-81 00:00:00|1250.00|500.00| 30
7566 |JONES|MANAGER|7839|02-APR-81 00:00:00|2975.00|| 20
7654 |MARTIN|SALESMAN|7698|28-SEP-81 00:00:00|1250.00|1400.00| 30
(5 rows)
```

```
edb=# SELECT STATS_MODE(SAL) FROM EMP;
stats_mode
-----
1250.00
(1 row)
```

The following example uses `GROUP BY` and `ORDER BY` clause with `STATS_MODE` in `EMP` table and returns the salary based on a partition by `DEPTNO`:

```
edb=# SELECT STATS_MODE(SAL) FROM EMP GROUP BY DEPTNO ORDER BY DEPTNO;
stats_mode
-----
800.00
1250.00
(2 rows)
```

The following example uses the `WITHIN GROUP` clause with the `STATS_MODE` function to perform aggregation on the ordered data set.

```
SELECT STATS_MODE() WITHIN GROUP(ORDER BY SAL) FROM EMP;
stats_mode
-----
1250.00
(1 row)
```

5.2.3.12 Subquery Expressions

This section describes the SQL-compliant subquery expressions available in Advanced Server. All of the expression forms documented in this section return Boolean (`TRUE/FALSE`) results.

EXISTS

The argument of `EXISTS` is an arbitrary `SELECT` statement, or subquery. The subquery is evaluated to determine whether it returns any rows. If it returns at least one row, the result of `EXISTS` is `TRUE`; if the subquery returns no rows, the result of `EXISTS` is `FALSE`.

EXISTS(subquery)

The subquery can refer to variables from the surrounding query, which will act as constants during any one evaluation of the subquery.

The subquery will generally only be executed far enough to determine whether at least one row is returned, not all the way to completion. It is unwise to write a subquery that has any side effects (such as calling sequence functions); whether the side effects occur or not may be difficult to predict.

Since the result depends only on whether any rows are returned, and not on the contents of those rows, the output list of the subquery is normally uninteresting. A common coding convention is to write all `EXISTS` tests in the form `EXISTS(SELECT 1 WHERE ...)`. There are exceptions to this rule however, such as subqueries that use `INTERSECT`.

This simple example is like an inner join on `deptno`, but it produces at most one output row for each `dept` row, even though there are multiple matching `emp` rows:

```
SELECT dname FROM dept WHERE EXISTS (SELECT 1 FROM emp WHERE emp.deptno =
dept.deptno);
```

`dname`

ACCOUNTING
RESEARCH
SALES
(3 rows)

IN

The right-hand side is a parenthesized subquery, which must return exactly one column. The left-hand expression is evaluated and compared to each row of the subquery result. The result of `IN` is `TRUE` if any equal subquery row is found. The result is `FALSE` if no equal row is found (including the special case where the subquery returns no rows).

expression IN (subquery)

Note that if the left-hand expression yields `NULL`, or if there are no equal right-hand values and at least one right-hand row yields `NULL`, the result of the `IN` construct will be `NULL`, not `FALSE`. This is in accordance with SQL's normal rules for Boolean combinations of null values.

As with `EXISTS`, it's unwise to assume that the subquery will be evaluated completely.

NOT IN

The right-hand side is a parenthesized subquery, which must return exactly one column. The left-hand expression is evaluated and compared to each row of the subquery result. The result of `NOT IN` is `TRUE` if only unequal subquery rows are found (including the special case where the subquery returns no rows). The result is `FALSE` if any equal row is found.

expression NOT IN (subquery)

Note that if the left-hand expression yields `NULL`, or if there are no equal right-hand values and at least one right-hand row yields `NULL`, the result of the `NOT IN` construct will be `NULL`, not `TRUE`. This is in accordance with SQL's normal rules for Boolean combinations of null values.

As with `EXISTS`, it's unwise to assume that the subquery will be evaluated completely.

ANY/SOME

The right-hand side is a parenthesized subquery, which must return exactly one column. The left-hand expression is evaluated and compared to each row of the subquery result using the given operator, which must yield a Boolean result. The result of `ANY` is `TRUE` if any true result is obtained. The result is `FALSE` if no true result is found (including the special case where the subquery returns no rows).

expression operator ANY (subquery)

expression operator SOME (subquery)

`SOME` is a synonym for `ANY`. `IN` is equivalent to `= ANY`.

Note that if there are no successes and at least one right-hand row yields `NULL` for the operator's result, the result of the `ANY` construct will be `NULL`, not `FALSE`. This is in accordance with SQL's normal rules for Boolean combinations of null values.

As with `EXISTS`, it's unwise to assume that the subquery will be evaluated completely.

ALL

The right-hand side is a parenthesized subquery, which must return exactly one column. The left-hand expression is evaluated and compared to each row of the subquery result using the given operator, which must yield a Boolean result. The result of `ALL` is `TRUE` if all rows yield true (including the special case where the subquery returns no rows). The result is `FALSE` if any false result is found. The result is `NULL` if the comparison does not return `FALSE` for any row, and it returns `NULL` for at least one row.

expression operator ALL (subquery)

`NOT IN` is equivalent to `<> ALL`. As with `EXISTS`, it's unwise to assume that the subquery will be evaluated completely.

5.3 System Catalog Tables

The following system catalog tables contain definitions of database objects. The layout of the system tables is subject to change; if you are writing an application that depends on information stored in the system tables, it would be prudent to use an existing catalog view, or create a catalog view to isolate the application from changes to the system table.

dual

`dual` is a single-row, single-column table that is provided for compatibility with Oracle databases only.

| Column | Type | Modifiers | Description |
|--------|-------------|-----------|----------------------------------|
| dummy | VARCHAR2(1) | | Provided for compatibility only. |

edb_dir

The `edb_dir` table contains one row for each alias that points to a directory created with the `CREATE DIRECTORY` command. A directory is an alias for a pathname that allows a user limited access to the host file system.

You can use a directory to fence a user into a specific directory tree within the file system. For example, the `UTL_FILE` package offers functions that permit a user to read and write files and directories in the host file system, but only allows access to paths that the database administrator has granted access to via a `CREATE DIRECTORY` command.

| Column | Type | Modifiers | Description |
|-------------------------------|-----------------------|-----------|---|
| dirname | "name" | not null | The name of the alias. |
| dirowner dirpath diracl | oid text aclitem[] | not null | The OID of the user that owns the alias. The directory name to which access is granted. The access control list that determines which users may access the alias. |

edb_password_history

The `edb_password_history` table contains one row for each password change. The table is shared across all databases within a cluster.

| Column | Type | References | Description |
|---|-------------------------|---------------|---|
| passhistroleid passhistpassword passhistpasswordsetat | oid text timestamptz | pg_authid.oid | The ID of a role. Role password in md5 encrypted form. The time the password was set. |

edb_policy

The `edb_policy` table contains one row for each policy.

| Column | Type | Modifiers | Description |
|--------------|---------|-----------|---|
| policyname | name | not null | The policy name. |
| policygroup | oid | not null | Currently unused. |
| policyobject | oid | not null | The OID of the table secured by this policy (the object_schema plus the object_name). |
| policykind | char | not null | The kind of object secured by this policy: 'r' for a table 'v' for a view = for a synonym Currently always 'r'. |
| policyproc | oid | not null | The OID of the policy function (function_schema plus policy_function). |
| policyinsert | boolean | not null | True if the policy is enforced by INSERT statements. |

| Column | Type | Modifiers | Description |
|-------------------|------------|-----------|--|
| policyselect | boolean | not null | True if the policy is enforced by SELECT statements. |
| policydelete | boolean | not null | True if the policy is enforced by DELETE statements. |
| policyupdate | boolean | not null | True if the policy is enforced by UPDATE statements. |
| policyindex | boolean | not null | Currently unused. |
| policyenabled | boolean | not null | True if the policy is enabled. |
| policyupdatecheck | boolean | not null | True if rows updated by an UPDATE statement must satisfy the policy. |
| policystatic | boolean | not null | Currently unused. |
| policytype | integer | not null | Currently unused. |
| policyopts | integer | not null | Currently unused. |
| policyseccols | int2vector | not null | The column numbers for columns listed in sec_relevant_cols. |

edb_profile

The `edb_profile` table stores information about the available profiles. `edb_profiles` is shared across all databases within a cluster.

| Column | Type | References | Description |
|--|--|-----------------|---|
| oid prfname prffailedloginattempts prfpasswordlocktime prfpasswordlifetime prfpasswordgracetime prfpasswordreusetime prfpasswordreusemax prfpasswordallowhashed | oid name integer integer integer integer integer integer integer | | Row identifier (hidden attribute; must be explicitly selected). The name of the profile. The number of failed login attempts allowed by the profile. -1 indicates that the value from the default profile should be used. -2 indicates no limit on failed login attempts. The password lock time associated with the profile (in seconds). -1 indicates that the value from the default profile should be used. -2 indicates that the account should be locked permanently. The password life time associated with the profile (in seconds). -1 indicates that the value from the default profile should be used. -2 indicates that the password never expires. The password grace time associated with the profile (in seconds). -1 indicates that the value from the default profile should be used. -2 indicates that the password never expires. The number of seconds that a user must wait before reusing a password. -1 indicates that the value from the default profile should be used. -2 indicates that the old passwords can never be reused. The number of password changes that have to occur before a password can be reused. -1 indicates that the value from the default profile should be used. -2 indicates that the old passwords can never be reused. The password allow hashed parameter specifies whether an encrypted password to be allowed for use or not. The possible values can be true/on/yes/1, false/off/no/0, and DEFAULT. |
| prfpasswordverifyfuncdb | oid | pg_database.oid | The OID of the database in which the password verify function exists. |
| prfpasswordverifyfunc | oid | pg_proc.oid | The OID of the password verify function associated with the profile. |

edb_variable

The `edb_variable` table contains one row for each package level variable (each variable declared within a package).

| Column | Type | Modifiers | Description |
|------------|----------|-----------|---|
| varname | "name" | not null | The name of the variable. |
| varpackage | oid | not null | The OID of the pg_namespace row that stores the package. |
| vartype | oid | not null | The OID of the pg_type row that defines the type of the variable. |
| | | | + if the variable is visible outside of the package. - if the variable is only visible within the package. |
| varaccess | "char" | not null | Note: Public variables are declared within the package header; private variables are declared within the package body. Contains the source of the variable declaration, including any default value expressions for the variable. |
| varsrvc | text | | |
| varseq | smallint | not null | The order in which the variable was declared in the package. |

pg_synonym

The `pg_synonym` table contains one row for each synonym created with the `CREATE SYNONYM` command or `CREATE PUBLIC SYNONYM` command.

| Column | Type | Modifiers | Description |
|--------------|--------|-----------|--|
| synname | "name" | not null | The name of the synonym. |
| synnamespace | oid | not null | Replaces synowner. Contains the OID of the pg_namespace row where the synonym is stored |
| synowner | oid | not null | The OID of the user that owns the synonym. |
| synobjschema | "name" | not null | The schema in which the referenced object is defined. |
| synobjname | "name" | not null | The name of the referenced object. The (optional) name of the database link in which the referenced object is defined. |
| synlink | text | | |

product_component_version

The `product_component_version` table contains information about feature compatibility; an application can query this table at installation or run time to verify that features used by the application are available with this deployment.

| Column | Type | Description |
|---------|------------------------|--------------------------|
| product | character varying (74) | The name of the product. |

| Column | Type | Description |
|---------|------------------------|------------------------------------|
| version | character varying (74) | The version number of the product. |
| status | character varying (74) | The status of the release. |

6 Database Compatibility for Oracle

6.1 Introduction

Database Compatibility for Oracle means that an application runs in an Oracle environment as well as in the EDB Postgres Advanced Server (Advanced Server) environment with minimal or no changes to the application code. This guide focuses solely on the features that are related to the package support provided by Advanced Server.

For more information about using other compatibility features offered by Advanced Server, please see the complete set of Advanced Server guides, available at:

<https://www.enterprisedb.com/edb-docs/>

6.2 Packages

This chapter discusses the concept of packages in Advanced Server. A *package* is a named collection of functions, procedures, variables, cursors, user-defined record types, and records that are referenced using a common qualifier – the package identifier. Packages have the following characteristics:

- Packages provide a convenient means of organizing the functions and procedures that perform a related purpose. Permission to use the package functions and procedures is dependent upon one privilege granted to the entire package. All of the package programs must be referenced with a common name.
- Certain functions, procedures, variables, types, etc. in the package can be declared as *public*. Public entities are visible and can be referenced by other programs that are given `EXECUTE` privilege on the package. For public functions and procedures, only their signatures are visible - the program names, parameters if any, and return types of functions. The SPL code of these functions and procedures is not accessible to others, therefore applications that utilize a package are dependent only upon the information available in the signature – not in the procedural logic itself.
- Other functions, procedures, variables, types, etc. in the package can be declared as *private*. Private entities can be referenced and used by function and procedures within the package, but not by other external applications. Private entities are for use only by programs within the package.
- Function and procedure names can be overloaded within a package. One or more functions/procedures can be defined with the same name, but with different signatures. This provides the capability to create identically named programs that perform the same job, but on different types of input.

6.2.1 Package Components

Packages consist of two main components:

- The *package specification*: This is the public interface, (these are the elements which can be referenced outside the package). We declare all database objects that are to be a part of our package within the specification.
- The *package body*: This contains the actual implementation of all the database objects declared within the package specification.

The package body implements the specifications in the package specification. It contains implementation details and private declarations which are invisible to the application. You can debug, enhance or replace a package body without changing the specifications. Similarly, you can change the body without recompiling the calling programs because the implementation details are invisible to the application.

Package Specification Syntax

The package specification defines the user interface for a package (the API). The specification lists the functions, procedures, types, exceptions and cursors that are visible to a user of the package.

The syntax used to define the interface for a package is:

```
CREATE [ OR REPLACE ] PACKAGE <package_name>
[ <authorization_clause> ]
{ IS | AS }
[ <declaration>; ] ...
[ <procedure_or_function_declaration> ] ...
END [ <package_name> ] ;
```

Where `authorization_clause` :=

```
{ AUTHID DEFINER } | { AUTHID CURRENT_USER }
```

Where `procedure_or_function_declaration` :=

```
procedure_declaration | function_declaration
```

Where `procedure_declaration` :=

```
PROCEDURE proc_name [ argument_list ];
[ restrictionPragma; ]
```

Where `function_declaration` :=

```
FUNCTION func_name [ argument_list ]
  RETURN rettype [ DETERMINISTIC ];
[ restrictionPragma; ]
```

Where `argument_list` :=

```
( argument_declaration [, ...] )
```

Where `argument_declaration` :=

`argname [IN | IN OUT | OUT] argtype [DEFAULT value]`

Where `restrictionPragma` :=

`PRAGMA RESTRICT_REFERENCES(name, restrictions)`

Where `restrictions` :=

`restriction [, ...]`

Parameters

`package_name`

`package_name` is an identifier assigned to the package - each package must have a name unique within the schema.

`AUTHID DEFINER`

If you omit the `AUTHID` clause or specify `AUTHID DEFINER`, the privileges of the package owner are used to determine access privileges to database objects.

`AUTHID CURRENT_USER`

If you specify `AUTHID CURRENT_USER`, the privileges of the current user executing a program in the package are used to determine access privileges.

`declaration`

`declaration` is an identifier of a public variable. A public variable can be accessed from outside of the package using the syntax `package_name.variable`. There can be zero, one, or more public variables. Public variable definitions must come before procedure or function declarations.

`declaration` can be any of the following:

- Variable Declaration
- Record Declaration
- Collection Declaration
- `REF CURSOR` and Cursor Variable Declaration
- `TYPE` Definitions for Records, Collections, and `REF CURSORS`
- Exception
- Object Variable Declaration

`proc_name`

The name of a public procedure.

`argname`

The name of an argument. The argument is referenced by this name within the function or procedure body.

`IN | IN OUT | OUT`

The argument mode. `IN` declares the argument for input only. This is the default. `IN OUT` allows the argument to receive a value as well as return a value. `OUT` specifies the argument is for output only.

`argtype`

The data type(s) of an argument. An argument type may be a base data type, a copy of the type of an existing column using `%TYPE`, or a user-defined type such as a nested table or an object type. A length must not be specified for any base type - for example, specify `VARCHAR2`, not `VARCHAR2(10)`.

The type of a column is referenced by writing `tablename.columnname %TYPE`; using this can sometimes help make a procedure independent from changes to the definition of a table.

DEFAULT value

The `DEFAULT` clause supplies a default value for an input argument if one is not supplied in the invocation. `DEFAULT` may not be specified for arguments with modes `IN OUT` or `OUT`.

func_name

The name of a public function.

rettype

The return data type.

DETERMINISTIC

`DETERMINISTIC` is a synonym for `IMMUTABLE`. A `DETERMINISTIC` function cannot modify the database and always reaches the same result when given the same argument values; it does not do database lookups or otherwise use information not directly present in its argument list. If you include this clause, any call of the function with all-constant arguments can be immediately replaced with the function value.

restriction

The following keywords are accepted for compatibility and ignored:

RNDS

RNPS

TRUST

WNDS

WNPS

Package Body Syntax

Package implementation details reside in the package body; the package body may contain objects that are not visible to the package user. Advanced Server supports the following syntax for the package body:

```
CREATE [ OR REPLACE ] PACKAGE BODY <package_name>
{ IS | AS }
[ <private_declaration>; ] ...
[ <procedure_or_function_definition> ] ...
[ <package_initializer> ]
END [ <package_name> ] ;
```

Where `procedure_or_function_definition` :=

`procedure_definition` | `function_definition`

Where `procedure_definition` :=

```
PROCEDURE proc_name[ argument_list ]
[ options_list ]
{ IS | AS }
```

```
procedure_body
END [ proc_name ] ;
```

Where `procedure_body` :=

```
[ PRAGMA AUTONOMOUS_TRANSACTION; ]
[ declaration; ] [, ...]
BEGIN
  statement; [...]
[ EXCEPTION
  { WHEN exception [OR exception] [...] THEN statement; }
  [...]
]
```

Where `function_definition` :=

```
FUNCTION func_name [ argument_list ]
  RETURN rettype [ DETERMINISTIC ]
  [ options_list ]
  { IS | AS }
    function_body
  END [ func_name ] ;
```

Where `function_body` :=

```
[ PRAGMA AUTONOMOUS_TRANSACTION; ]
[ declaration; ] [, ...]
BEGIN
  statement; [...]
[ EXCEPTION
  { WHEN exception [ OR exception ] [...] THEN statement; }
  [...]
]
```

Where `argument_list` :=

```
( argument_declaration [, ...] )
```

Where `argument_declaration` :=

```
argname [ IN | IN OUT | OUT ] argtype [ DEFAULT value ]
```

Where `options_list` :=

```
option [ ... ]
```

Where `option` :=

```
STRICT
LEAKPROOF
PARALLEL { UNSAFE | RESTRICTED | SAFE }
COST execution_cost
ROWS result_rows
SET config_param { TO value | = value | FROM CURRENT }
```

Where `package_initializer` :=

```
BEGIN
statement; [...]
END;
```

Parameters

package_name

`package_name` is the name of the package for which this is the package body. There must be an existing package specification with this name.

private_declaration

`private_declaration` is an identifier of a private variable that can be accessed by any procedure or function within the package. There can be zero, one, or more private variables. `private_declaration` can be any of the following:

- Variable Declaration
- Record Declaration
- Collection Declaration
- `REF CURSOR` and Cursor Variable Declaration
- `TYPE` Definitions for Records, Collections, and `REF CURSORS`
- Exception
- Object Variable Declaration

proc_name

The name of the procedure being created.

PRAGMA AUTONOMOUS_TRANSACTION

`PRAGMA AUTONOMOUS_TRANSACTION` is the directive that sets the procedure as an autonomous transaction.

declaration

A variable, type, `REF CURSOR`, or subprogram declaration. If subprogram declarations are included, they must be declared after all other variable, type, and `REF CURSOR` declarations.

statement

An SPL program statement. Note that a `DECLARE - BEGIN - END` block is considered an SPL statement unto itself. Thus, the function body may contain nested blocks.

exception

An exception condition name such as `NO_DATA_FOUND`, `OTHERS`, etc.

func_name

The name of the function being created.

rettype

The return data type, which may be any of the types listed for `argtype`. As for `argtype`, a length must not be specified for `rettype`.

DETERMINISTIC

Include `DETERMINISTIC` to specify that the function will always return the same result when given the same argument values. A `DETERMINISTIC` function must not modify the database.

Note

The **DETERMINISTIC** keyword is equivalent to the PostgreSQL **IMMUTABLE** option.

Note

If **DETERMINISTIC** is specified for a public function in the package body, it must also be specified for the function declaration in the package specification. (For private functions, there is no function declaration in the package specification.)

PRAGMA AUTONOMOUS_TRANSACTION

PRAGMA AUTONOMOUS_TRANSACTION is the directive that sets the function as an autonomous transaction.

argname

The name of a formal argument. The argument is referenced by this name within the procedure body.

IN | IN OUT | OUT

The argument mode. **IN** declares the argument for input only. This is the default. **IN OUT** allows the argument to receive a value as well as return a value. **OUT** specifies the argument is for output only.

argtype

The data type(s) of an argument. An argument type may be a base data type, a copy of the type of an existing column using **%TYPE**, or a user-defined type such as a nested table or an object type. A length must not be specified for any base type - for example, specify **VARCHAR2**, not **VARCHAR2(10)**.

The type of a column is referenced by writing **tablename.columnname%TYPE**; using this can sometimes help make a procedure independent from changes to the definition of a table.

DEFAULT value

The **DEFAULT** clause supplies a default value for an input argument if one is not supplied in the procedure call. **DEFAULT** may not be specified for arguments with modes **IN OUT** or **OUT**.

Please note: The following options are not compatible with Oracle databases; they are extensions to Oracle package syntax provided by Advanced Server only.

STRICT

The **STRICT** keyword specifies that the function will not be executed if called with a **NULL** argument; instead the function will return **NULL**.

LEAKPROOF

The **LEAKPROOF** keyword specifies that the function will not reveal any information about arguments, other than through a return value.

PARALLEL { UNSAFE | RESTRICTED | SAFE }

The **PARALLEL** clause enables the use of parallel sequential scans (parallel mode). A parallel sequential scan uses multiple workers to scan a relation in parallel during a query in contrast to a serial sequential scan.

When set to **UNSAFE**, the procedure or function cannot be executed in parallel mode. The presence of such a procedure or function forces a serial execution plan. This is the default setting if the **PARALLEL** clause is omitted.

When set to **RESTRICTED**, the procedure or function can be executed in parallel mode, but the execution is restricted to the parallel group leader. If the qualification for any particular relation has anything that is parallel restricted, that relation won't be chosen for parallelism.

When set to `SAFE`, the procedure or function can be executed in parallel mode with no restriction.

`execution_cost`

`execution_cost` specifies a positive number giving the estimated execution cost for the function, in units of `cpu_operator_cost`. If the function returns a set, this is the cost per returned row. The default is `0.0025`.

`result_rows`

`result_rows` is the estimated number of rows that the query planner should expect the function to return. The default is `1000`.

`SET`

Use the `SET` clause to specify a parameter value for the duration of the function:

`config_param` specifies the parameter name.

`value` specifies the parameter value.

`FROM CURRENT` guarantees that the parameter value is restored when the function ends.

`package_initializer`

The statements in the `package_initializer` are executed once per user's session when the package is first referenced.

Note

The `STRICT`, `LEAKPROOF`, `PARALLEL`, `COST`, `ROWS` and `SET` keywords provide extended functionality for Advanced Server and are not supported by Oracle.

6.2.2 Creating Packages

A package is not an executable piece of code; rather it is a repository of code. When you use a package, you actually execute or make reference to an element within a package.

Creating the Package Specification

The package specification contains the definition of all the elements in the package that can be referenced from outside of the package. These are called the public elements of the package, and they act as the package interface. The following code sample is a package specification:

```
--  
-- Package specification for the 'emp_admin' package.  
--  
CREATE OR REPLACE PACKAGE emp_admin  
IS  
    FUNCTION get_dept_name (  
        p_deptno NUMBER DEFAULT 10  
    )  
        RETURN VARCHAR2;
```

```

FUNCTION update_emp_sal (
    p_empho NUMBER,
    p_raise NUMBER
)
RETURN NUMBER;
PROCEDURE hire_emp (
    p_empho      NUMBER,
    p_ename      VARCHAR2,
    p_job        VARCHAR2,
    p_sal        NUMBER,
    p_hiredate   DATE DEFAULT sysdate,
    p_comm       NUMBER DEFAULT 0,
    p_mgr        NUMBER,
    p_deptno     NUMBER DEFAULT 10
);
PROCEDURE fire_emp (
    p_empho NUMBER
);
END emp_admin;

```

This code sample creates the `emp_admin` package specification. This package specification consists of two functions and two stored procedures. We can also add the `OR REPLACE` clause to the `CREATE PACKAGE` statement for convenience.

Creating the Package Body

The body of the package contains the actual implementation behind the package specification. For the above `emp_admin` package specification, we shall now create a package body which will implement the specifications. The body will contain the implementation of the functions and stored procedures in the specification.

```

-- 
-- Package body for the 'emp_admin' package.
-- 
CREATE OR REPLACE PACKAGE BODY emp_admin
IS
    --
    -- Function that queries the 'dept' table based on the department
    -- number and returns the corresponding department name.
    --
    FUNCTION get_dept_name (
        p_deptno    IN NUMBER DEFAULT 10
    )
    RETURN VARCHAR2
    IS
        v_dname      VARCHAR2(14);
    BEGIN
        SELECT dname INTO v_dname FROM dept WHERE deptno = p_deptno;
        RETURN v_dname;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            DBMS_OUTPUT.PUT_LINE('Invalid department number ' || p_deptno);
            RETURN '';
    END;
    --

```

```

-- Function that updates an employee's salary based on the
-- employee number and salary increment/decrement passed
-- as IN parameters. Upon successful completion the function
-- returns the new updated salary.
--
FUNCTION update_emp_sal (
    p_empno      IN NUMBER,
    p_raise      IN NUMBER
)
RETURN NUMBER
IS
    v_sal      NUMBER := 0;
BEGIN
    SELECT sal INTO v_sal FROM emp WHERE empno = p_empno;
    v_sal := v_sal + p_raise;
    UPDATE emp SET sal = v_sal WHERE empno = p_empno;
    RETURN v_sal;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('Employee ' || p_empno || ' not found');
        RETURN -1;
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
        DBMS_OUTPUT.PUT_LINE(SQLERRM);
        DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
        DBMS_OUTPUT.PUT_LINE(SQLCODE);
        RETURN -1;
END;
--
-- Procedure that inserts a new employee record into the 'emp' table.
--
PROCEDURE hire_emp (
    p_empno      NUMBER,
    p_ename      VARCHAR2,
    p_job       VARCHAR2,
    p_sal       NUMBER,
    p_hiredate   DATE  DEFAULT sysdate,
    p_comm      NUMBER DEFAULT 0,
    p_mgr       NUMBER,
    p_deptno    NUMBER DEFAULT 10
)
AS
BEGIN
    INSERT INTO emp(empno, ename, job, sal, hiredate, comm, mgr, deptno)
        VALUES(p_empno, p_ename, p_job, p_sal,
               p_hiredate, p_comm, p_mgr, p_deptno);
END;
--
-- Procedure that deletes an employee record from the 'emp' table based
-- on the employee number.
--
PROCEDURE fire_emp (
    p_empno      NUMBER
)
AS

```

```

BEGIN
  DELETE FROM emp WHERE empno = p_empno;
END;
END;

```

6.2.3 Referencing a Package

To reference the types, items and subprograms that are declared within a package specification, we use the dot notation. For example:

```

package_name.type_name
package_name.item_name
package_name.subprogram_name

```

To invoke a function from the `emp_admin` package specification, we will execute the following SQL command.

```
SELECT emp_admin.get_dept_name(10) FROM DUAL;
```

Here we are invoking the `get_dept_name` function declared within the package `emp_admin`. We are passing the department number as an argument to the function, which will return the name of the department. Here the value returned should be `ACCOUNTING`, which corresponds to department number `10`.

6.2.4 Using Packages With User Defined Types

The following example incorporates the various user-defined types discussed in earlier chapters within the context of a package.

The package specification of `emp_rpt` shows the declaration of a record type, `emprec_typ`, and a weakly-typed `REF CURSOR`, `emp_refcur`, as publicly accessible along with two functions and two procedures. Function, `open_emp_by_dept`, returns the `REF CURSOR` type, `EMP_REFCUR`. Procedures, `fetch_emp` and `close_refcur`, both declare a weakly-typed `REF CURSOR` as a formal parameter.

```

CREATE OR REPLACE PACKAGE emp_rpt
IS
  TYPE emprec_typ IS RECORD (
    empno    NUMBER(4),
    ename    VARCHAR(10)
  );
  TYPE emp_refcur IS REF CURSOR;

  FUNCTION get_dept_name (
    p_deptno  IN NUMBER
  ) RETURN VARCHAR2;
  FUNCTION open_emp_by_dept (
    p_deptno  IN emp.deptno%TYPE
  
```

```
) RETURN EMP_REFCUR;
PROCEDURE fetch_emp (
    p_refcur  IN OUT SYS_REFCURSOR
);
PROCEDURE close_refcur (
    p_refcur  IN OUT SYS_REFCURSOR
);
END emp_rpt;
```

The package body shows the declaration of several private variables - a static cursor, `dept_cur`, a table type, `depttab_typ`, a table variable, `t_dept`, an integer variable, `t_dept_max`, and a record variable, `r_emp`.

```
CREATE OR REPLACE PACKAGE BODY emp_rpt
IS
    CURSOR dept_cur IS SELECT * FROM dept;
    TYPE depttab_typ IS TABLE OF dept%ROWTYPE
        INDEX BY BINARY_INTEGER;
    t_dept      DEPTTAB_TYP;
    t_dept_max  INTEGER := 1;
    r_emp       EMPREC_TYP;

    FUNCTION get_dept_name (
        p_deptno  IN NUMBER
    ) RETURN VARCHAR2
    IS
    BEGIN
        FOR i IN 1..t_dept_max LOOP
            IF p_deptno = t_dept(i).deptno THEN
                RETURN t_dept(i).dname;
            END IF;
        END LOOP;
        RETURN 'Unknown';
    END;

    FUNCTION open_emp_by_dept(
        p_deptno      IN emp.deptno%TYPE
    ) RETURN EMP_REFCUR
    IS
        emp_by_dept EMP_REFCUR;
    BEGIN
        OPEN emp_by_dept FOR SELECT empno, ename FROM emp
            WHERE deptno = p_deptno;
        RETURN emp_by_dept;
    END;

    PROCEDURE fetch_emp (
        p_refcur  IN OUT SYS_REFCURSOR
    )
    IS
    BEGIN
        DBMS_OUTPUT.PUT_LINE('EMPNO ENAME');
        DBMS_OUTPUT.PUT_LINE('----- -----');
        LOOP
            FETCH p_refcur INTO r_emp;
            EXIT WHEN p_refcur%NOTFOUND;
            DBMS_OUTPUT.PUT_LINE(r_emp.empno || ' ' || r_emp.ename);
        END LOOP;
    END;
END emp_rpt;
```

```

    END LOOP;
END;

PROCEDURE close_refcur (
    p_refcur    IN OUT SYS_REFCURSOR
)
IS
BEGIN
    CLOSE p_refcur;
END;
BEGIN
    OPEN dept_cur;
LOOP
    FETCH dept_cur INTO t_dept(t_dept_max);
    EXIT WHEN dept_cur%NOTFOUND;
    t_dept_max := t_dept_max + 1;
END LOOP;
CLOSE dept_cur;
t_dept_max := t_dept_max - 1;
END emp_rpt;

```

This package contains an initialization section that loads the private table variable, `t_dept`, using the private static cursor, `dept_cur.t_dept` serves as a department name lookup table in function, `get_dept_name`.

Function, `open_emp_by_dept` returns a `REF CURSOR` variable for a result set of employee numbers and names for a given department. This `REF CURSOR` variable can then be passed to procedure, `fetch_emp`, to retrieve and list the individual rows of the result set. Finally, procedure, `close_refcur`, can be used to close the `REF CURSOR` variable associated with this result set.

The following anonymous block runs the package function and procedures. In the anonymous block's declaration section, note the declaration of cursor variable, `v_emp_cur`, using the package's public `REF CURSOR` type, `EMP_REFCUR`. `v_emp_cur` contains the pointer to the result set that is passed between the package function and procedures.

```

DECLARE
    v_deptno dept.deptno%TYPE DEFAULT 30;
    v_emp_cur emp_rpt.EMP_REFCUR;
BEGIN
    v_emp_cur := emp_rpt.open_emp_by_dept(v_deptno);
    DBMS_OUTPUT.PUT_LINE('EMPLOYEES IN DEPT #' || v_deptno ||
        ':' || emp_rpt.get_dept_name(v_deptno));
    emp_rpt.fetch_emp(v_emp_cur);
    DBMS_OUTPUT.PUT_LINE('*****');
    DBMS_OUTPUT.PUT_LINE(v_emp_cur%ROWCOUNT || ' rows were retrieved');
    emp_rpt.close_refcur(v_emp_cur);
END;

```

The following is the result of this anonymous block.

```

EMPLOYEES IN DEPT #30: SALES
EMPNO  ENAME
-----
7499  ALLEN
7521  WARD
7654  MARTIN
7698  BLAKE
7844  TURNER

```

```
7900 JAMES
*****
6 rows were retrieved
```

The following anonymous block illustrates another means of achieving the same result. Instead of using the package procedures, `fetch_emp` and `close_refcur`, the logic of these programs is coded directly into the anonymous block. In the anonymous block's declaration section, note the addition of record variable, `r_emp`, declared using the package's public record type, `EMPREC_TYP`.

```
DECLARE
  v_deptno  dept.deptno%TYPE DEFAULT 30;
  v_emp_cur  emp_rpt.EMP_REFCUR;
  r_emp      emp_rpt.EMPREC_TYP;
BEGIN
  v_emp_cur := emp_rpt.open_emp_by_dept(v_deptno);
  DBMS_OUTPUT.PUT_LINE('EMPLOYEES IN DEPT #' || v_deptno ||
    ': ' || emp_rpt.get_dept_name(v_deptno));
  DBMS_OUTPUT.PUT_LINE('EMPNO ENAME');
  DBMS_OUTPUT.PUT_LINE('----- -----');
  LOOP
    FETCH v_emp_cur INTO r_emp;
    EXIT WHEN v_emp_cur%NOTFOUND;
    DBMS_OUTPUT.PUT_LINE(r_emp.empno || ' ' ||
      r_emp.ename);
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('*****');
  DBMS_OUTPUT.PUT_LINE(v_emp_cur%ROWCOUNT || ' rows were retrieved');
  CLOSE v_emp_cur;
END;
```

The following is the result of this anonymous block.

```
EMPLOYEES IN DEPT #30: SALES
EMPNO  ENAME
-----
7499  ALLEN
7521  WARD
7654  MARTIN
7698  BLAKE
7844  TURNER
7900  JAMES
*****
6 rows were retrieved
```

6.2.5 Dropping a Package

The syntax for deleting an entire package or just the package body is as follows:

```
DROP PACKAGE [ BODY ] package_name;
```

If the keyword, `BODY`, is omitted, both the package specification and the package body are deleted - i.e., the

entire package is dropped. If the keyword, `BODY`, is specified, then only the package body is dropped. The package specification remains intact. `package_name` is the identifier of the package to be dropped.

Following statement will destroy only the package body of `emp_admin`:

```
DROP PACKAGE BODY emp_admin;
```

The following statement will drop the entire `emp_admin` package:

```
DROP PACKAGE emp_admin;
```

6.3 Built-In Packages

This chapter describes the built-in packages that are provided with Advanced Server. For certain packages, non-superusers must be explicitly granted the `EXECUTE` privilege on the package before using any of the package's functions or procedures. For most of the built-in packages, `EXECUTE` privilege has been granted to `PUBLIC` by default.

For information about using the `GRANT` command to provide access to a package, see the *Database Compatibility for Oracle Developers SQL Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

All built-in packages are owned by the special `sys` user which must be specified when granting or revoking privileges on built-in packages:

```
GRANT EXECUTE ON PACKAGE SYS.UTL_FILE TO john;
```

6.3.1 DBMS_ALERT

The `DBMS_ALERT` package provides the capability to register for, send, and receive alerts. The following table lists the supported procedures:

| Function/Procedure | Return Type | Description |
|--|-------------|--|
| <code>REGISTER(name)</code> | n/a | Register to be able to receive alerts named, <code>name</code> |
| <code>REMOVE(name)</code> | n/a | Remove registration for the alert named, <code>name</code> |
| <code>REMOVEALL</code> | n/a | Remove registration for all alerts. |
| <code>SIGNAL(name, message)</code> | n/a | Signals the alert named, <code>name</code> , with <code>message</code> |
| <code>WAITANY(name OUT, message OUT, status OUT, timeout)</code> | n/a | Wait for any registered alert to occur. |
| <code>WAITONE(name, message OUT, status OUT, timeout)</code> | n/a | Wait for the specified alert, <code>name</code> , to occur. |

Advanced Server's implementation of `DBMS_ALERT` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Advanced Server allows a maximum of `500` concurrent alerts. You can use the `dbms_alert.max_alerts` GUC variable (located in the `postgresql.conf` file) to specify the maximum number of concurrent alerts allowed on a system.

To set a value for the `dbms_alert.max_alerts` variable, open the `postgresql.conf` file (located by default in `/opt/PostgresPlus/13AS/data`) with your choice of editor, and edit the `dbms_alert.max_alerts` parameter as shown:

```
dbms_alert.max_alerts = alert_count
alert_count
```

`alert_count` specifies the maximum number of concurrent alerts. By default, the value of `dbms_alert.max_alerts` is `100`. To disable this feature, set `dbms_alert.max_alerts` to `0`.

For the `dbms_alert.max_alerts` GUC to function correctly, the `custom_variable_classes` parameter must contain `dbms_alerts`:

```
custom_variable_classes = 'dbms_alert, ...'
```

After editing the `postgresql.conf` file parameters, you must restart the server for the changes to take effect.

REGISTER

The `REGISTER` procedure enables the current session to be notified of the specified alert.

```
REGISTER(<name> VARCHAR2)
```

Parameters

`name`

Name of the alert to be registered.

Examples

The following anonymous block registers for an alert named, `alert_test`, then waits for the signal.

```
DECLARE
  v_name      VARCHAR2(30) := 'alert_test';
  v_msg       VARCHAR2(80);
  v_status    INTEGER;
  v_timeout   NUMBER(3) := 120;
BEGIN
  DBMS_ALERT.REGISTER(v_name);
  DBMS_OUTPUT.PUT_LINE('Registered for alert ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
  DBMS_ALERT.WAITONE(v_name,v_msg,v_status,v_timeout);
  DBMS_OUTPUT.PUT_LINE('Alert name : ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Alert msg  : ' || v_msg);
  DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
  DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
  DBMS_ALERT.REMOVE(v_name);
END;
```

```
Registered for alert alert_test
Waiting for signal...
```

REMOVE

The **REMOVE** procedure unregisters the session for the named alert.

```
REMOVE(<name> VARCHAR2)
```

Parameters

name

Name of the alert to be unregistered.

REMOVEALL

The **REMOVEALL** procedure unregisters the session for all alerts.

```
REMOVEALL
```

SIGNAL

The **SIGNAL** procedure signals the occurrence of the named alert.

```
SIGNAL(<name> VARCHAR2, <message> VARCHAR2)
```

Parameters

name

Name of the alert.

message

Information to pass with this alert.

Examples

The following anonymous block signals an alert for **alert_test**.

```
DECLARE
  v_name  VARCHAR2(30) := 'alert_test';
BEGIN
  DBMS_ALERT.SIGNAL(v_name,'This is the message from ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Issued alert for ' || v_name);
END;
Issued alert for alert_test
```

WAITANY

The `WAITANY` procedure waits for any of the registered alerts to occur.

```
WAITANY(<name> OUT VARCHAR2, <message> OUT VARCHAR2,
<status> OUT INTEGER, <timeout> NUMBER)
```

Parameters

`name`

Variable receiving the name of the alert.

`message`

Variable receiving the message sent by the `SIGNAL` procedure.

`status`

Status code returned by the operation. Possible values are: 0 – alert occurred; 1 – timeout occurred.

`timeout`

Time to wait for an alert in seconds.

Examples

The following anonymous block uses the `WAITANY` procedure to receive an alert named, `alert_test` or `any_alert`:

```
DECLARE
  v_name      VARCHAR2(30);
  v_msg       VARCHAR2(80);
  v_status    INTEGER;
  v_timeout   NUMBER(3) := 120;
BEGIN
  DBMS_ALERT.REGISTER('alert_test');
  DBMS_ALERT.REGISTER('any_alert');
  DBMS_OUTPUT.PUT_LINE('Registered for alert alert_test and any_alert');
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
  DBMS_ALERT.WAITANY(v_name,v_msg,v_status,v_timeout);
  DBMS_OUTPUT.PUT_LINE('Alert name : ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Alert msg  : ' || v_msg);
  DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
  DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
  DBMS_ALERT.REMOVEALL;
END;
```

Registered for alert alert_test and any_alert

Waiting for signal...

An anonymous block in a second session issues a signal for `any_alert`:

```
DECLARE
  v_name  VARCHAR2(30) := 'any_alert';
BEGIN
  DBMS_ALERT.SIGNAL(v_name,'This is the message from ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Issued alert for ' || v_name);
END;
```

Issued alert for any_alert

Control returns to the first anonymous block and the remainder of the code is executed:

```
Registered for alert alert_test and any_alert
Waiting for signal...
Alert name : any_alert
Alert msg : This is the message from any_alert
Alert status : 0
Alert timeout: 120 seconds
```

WAITONE

The **WAITONE** procedure waits for the specified registered alert to occur.

```
WAITONE(<name> VARCHAR2, <message> OUT VARCHAR2,
<status> OUT INTEGER, <timeout> NUMBER)
```

Parameters

name

Name of the alert.

message

Variable receiving the message sent by the **SIGNAL** procedure.

status

Status code returned by the operation. Possible values are: 0 – alert occurred; 1 – timeout occurred.

timeout

Time to wait for an alert in seconds.

Examples

The following anonymous block is similar to the one used in the **WAITANY** example except the **WAITONE** procedure is used to receive the alert named, **alert_test**.

```
DECLARE
  v_name      VARCHAR2(30) := 'alert_test';
  v_msg       VARCHAR2(80);
  v_status    INTEGER;
  v_timeout   NUMBER(3) := 120;
BEGIN
  DBMS_ALERT.REGISTER(v_name);
  DBMS_OUTPUT.PUT_ DBMS_ALERT.REGISTER(v_name);
  DBMS_OUTPUT.PUT_LINE('Registered for alert ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
  DBMS_ALERT.WAITONE(v_name,v_msg,v_status,v_timeout);
  DBMS_OUTPUT.PUT_LINE('Alert name : ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Alert msg : ' || v_msg);
  DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
  DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
  DBMS_ALERT.REMOVE(v_name);LINE('Registered for alert ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
```

```

DBMS_ALERT.WAITONE(v_name,v_msg,v_status,v_timeout);
DBMS_OUTPUT.PUT_LINE('Alert name : ' || v_name);
DBMS_OUTPUT.PUT_LINE('Alert msg : ' || v_msg);
DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
DBMS_OUTPUT.PUT_LINE('Alert timeout: ' || v_timeout || ' seconds');
DBMS_ALERT.REMOVE(v_name);
END;

```

Registered for alert alert_test
Waiting for signal...

Signal sent for `alert_test` sent by an anonymous block in a second session:

```

DECLARE
  v_name  VARCHAR2(30) := 'alert_test';
BEGIN
  DBMS_ALERT.SIGNAL(v_name,'This is the message from ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Issued alert for ' || v_name);
END;

```

Issued alert for alert_test

First session is alerted, control returns to the anonymous block, and the remainder of the code is executed:

```

Registered for alert alert_test
Waiting for signal...
Alert name : alert_test
Alert msg : This is the message from alert_test
Alert status : 0
Alert timeout: 120 seconds

```

Comprehensive Example

The following example uses two triggers to send alerts when the `dept` table or the `emp` table is changed. An anonymous block listens for these alerts and displays messages when an alert is received.

The following are the triggers on the `dept` and `emp` tables:

```

CREATE OR REPLACE TRIGGER dept_alert_trig
  AFTER INSERT OR UPDATE OR DELETE ON dept
DECLARE
  v_action      VARCHAR2(25);
BEGIN
  IF INSERTING THEN
    v_action := ' added department(s)';
  ELSIF UPDATING THEN
    v_action := ' updated department(s)';
  ELSIF DELETING THEN
    v_action := ' deleted department(s)';
  END IF;
  DBMS_ALERT.SIGNAL('dept_alert',USER || v_action || 'on ' ||
    SYSDATE);
END;

```

```

CREATE OR REPLACE TRIGGER emp_alert_trig
  AFTER INSERT OR UPDATE OR DELETE ON emp
DECLARE
  v_action      VARCHAR2(25);
BEGIN
  IF INSERTING THEN
    v_action := ' added employee(s)';
  ELSIF UPDATING THEN
    v_action := ' updated employee(s)';
  ELSIF DELETING THEN
    v_action := ' deleted employee(s)';
  END IF;
  DBMS_ALERT.SIGNAL('emp_alert',USER || v_action || 'on' ||
                     SYSDATE);
END;

```

The following anonymous block is executed in a session while updates to the `dept` and `emp` tables occur in other sessions:

```

DECLARE
  v_dept_alert  VARCHAR2(30) := 'dept_alert';
  v_emp_alert   VARCHAR2(30) := 'emp_alert';
  v_name        VARCHAR2(30);
  v_msg         VARCHAR2(80);
  v_status      INTEGER;
  v_timeout     NUMBER(3) := 60;
BEGIN
  DBMS_ALERT.REGISTER(v_dept_alert);
  DBMS_ALERT.REGISTER(v_emp_alert);
  DBMS_OUTPUT.PUT_LINE('Registered for alerts dept_alert and emp_alert');
  DBMS_OUTPUT.PUT_LINE('Waiting for signal...');
LOOP
  DBMS_ALERT.WAITANY(v_name,v_msg,v_status,v_timeout);
  EXIT WHEN v_status != 0;
  DBMS_OUTPUT.PUT_LINE('Alert name : ' || v_name);
  DBMS_OUTPUT.PUT_LINE('Alert msg  : ' || v_msg);
  DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
  DBMS_OUTPUT.PUT_LINE('-----' ||
                       '-----');
END LOOP;
DBMS_OUTPUT.PUT_LINE('Alert status : ' || v_status);
DBMS_ALERT.REMOVEALL;
END;

```

Registered for alerts dept_alert and emp_alert
Waiting for signal...

The following changes are made by user, mary:

```

INSERT INTO dept VALUES (50,'FINANCE','CHICAGO');
INSERT INTO emp (empno,ename,deptno) VALUES (9001,'JONES',50);
INSERT INTO emp (empno,ename,deptno) VALUES (9002,'ALICE',50);

```

The following change is made by user, john:

```

INSERT INTO dept VALUES (60,'HR','LOS ANGELES');

```

The following is the output displayed by the anonymous block receiving the signals from the triggers:

```
Registered for alerts dept_alert and emp_alert
Waiting for signal...
Alert name : dept_alert
Alert msg   : mary added department(s) on 25-OCT-07 16:41:01
Alert status : 0
-----
Alert name : emp_alert
Alert msg   : mary added employee(s) on 25-OCT-07 16:41:02
Alert status : 0
-----
Alert name : dept_alert
Alert msg   : john added department(s) on 25-OCT-07 16:41:22
Alert status : 0
-----
Alert status : 1
```

6.3.2 DBMS_AQ

EDB Postgres Advanced Server Advanced Queueing provides message queueing and message processing for the Advanced Server database. User-defined messages are stored in a queue; a collection of queues is stored in a queue table. Procedures in the DBMS_AQADM package create and manage message queues and queue tables. Use the DBMS_AQ package to add messages to a queue or remove messages from a queue, or register or unregister a PL/SQL callback procedure.

Advanced Server also provides extended (non-compatible) functionality for the DBMS_AQ package with SQL commands, see the *Database Compatibility for Oracle Developers SQL Guide* for detailed information about the following SQL commands:

- [ALTER QUEUE](#)
- [ALTER QUEUE TABLE](#)
- [CREATE QUEUE](#)
- [CREATE QUEUE TABLE](#)
- [DROP QUEUE](#)
- [DROP QUEUE TABLE](#)

The DBMS_AQ package provides procedures that allow you to enqueue a message, dequeue a message, and manage callback procedures. The supported procedures are:

| Function/Procedure | Return Type | Description |
|----------------------------|-------------|--|
| ENQUEUE | n/a | Post a message to a queue. |
| DEQUEUE | n/a | Retrieve a message from a queue if or when a message is available. |
| REGISTER | n/a | Register a callback procedure. |
| UNREGISTER | n/a | Unregister a callback procedure. |

Advanced Server's implementation of DBMS_AQ is a partial implementation when compared to Oracle's version. Only those procedures listed in the table above are supported.

Advanced Server supports use of the constants listed below:

| Constant | Description | For Parameters |
|---------------------------------|---|--|
| DBMS_AQ.BROWSE (0) | Read the message without locking. | dequeue_options_t.dequeue_mode |
| DBMS_AQ.LOCKED (1) | This constant is defined, but will return an error if used. | dequeue_options_t.dequeue_mode |
| DBMS_AQ.REMOVE (2) | Delete the message after reading; the default. | dequeue_options_t.dequeue_mode |
| DBMS_AQ.REMOVE_NODATA (3) | This constant is defined, but will return an error if used. | dequeue_options_t.dequeue_mode |
| DBMS_AQ.FIRST_MESSAGE (0) | Return the first available message that matches the search terms. | dequeue_options_t.navigation |
| DBMS_AQ.NEXT_MESSAGE (1) | Return the next available message that matches the search terms. | dequeue_options_t.navigation |
| DBMS_AQ.NEXT_TRANSACTION (2) | This constant is defined, but will return an error if used. | dequeue_options_t.navigation |
| DBMS_AQ.FOREVER (-1) | Wait forever if a message that matches the search term is not found, the default. | dequeue_options_t.wait |
| DBMS_AQ.NO_WAIT (0) | Do not wait if a message that matches the search term is not found. | dequeue_options_t.wait |
| DBMS_AQ.ON_COMMIT (0) | The dequeue is part of the current transaction. | enqueue_options_t.visibility, dequeue_options_t.visibility |
| DBMS_AQ.IMMEDIATE (1) | This constant is defined, but will return an error if used. | enqueue_options_t.visibility, dequeue_options_t.visibility |
| DBMS_AQ.PERSISTENT (0) | The message should be stored in a table. | enqueue_options_t.delivery_mode |
| DBMS_AQ.BUFFERED (1) | This constant is defined, but will return an error if used. | enqueue_options_t.delivery_mode |
| DBMS_AQ.READY (0) | Specifies that the message is ready to process. | message_properties_t.state |
| DBMS_AQ.WAITING (1) | Specifies that the message is waiting to be processed. | message_properties_t.state |
| DBMS_AQ.PROCESSED (2) | Specifies that the message has been processed. | message_properties_t.state |
| DBMS_AQ.EXPIRED (3) | Specifies that the message is in the exception queue. | message_properties_t.state |
| DBMS_AQ.NO_DELAY (0) | This constant is defined, but will return an error if used | message_properties_t.delay |
| DBMS_AQ.NEVER (NULL) | This constant is defined, but will return an error if used | message_properties_t.expiration |
| DBMS_AQ.NAMESPACE_AQ (0) | Accept notifications from DBMS_AQ queues. | sys.aq\$_reg_info.namespace |
| DBMS_AQ.NAMESPACE_ANONYMOUS (1) | This constant is defined, but will return an error if used | sys.aq\$_reg_info.namespace |

The DBMS_AQ configuration parameters listed in the following table can be defined in the [postgresql.conf](#) file. After the configuration parameters are defined, you can invoke the DBMS_AQ package to use and manage messages held in queues and queue tables.

| Parameter | Description |
|-----------|-------------|
| | |

| | |
|--|--|
| <code>dbms_aq.max_workers</code> | The maximum number of workers to run. |
| <code>dbms_aq.max_idle_time</code> | The idle time a worker must wait before exiting. |
| <code>dbms_aq.min_work_time</code> | The minimum time a worker can run before exiting. |
| <code>dbms_aq.launch_delay</code> | The minimum time between creating workers. |
| <code>dbms_aq.batch_size</code> | The maximum number of messages to process in a single transaction. The default batch size is 10. |
| <code>dbms_aq.max_databases</code> | The size of DBMS_AQ's hash table of databases. The default value is 1024. |
| <code>dbms_aq.max_pending_retries</code> | The size of DBMS_AQ's hash table of pending retries. The default value is 1024. |

6.3.2.1 ENQUEUE

The `ENQUEUE` procedure adds an entry to a queue. The signature is:

```
ENQUEUE(
<queue_name> IN VARCHAR2,
<enqueue_options> IN DBMS_AQ.ENQUEUE_OPTIONS_T,
<message_properties> IN DBMS_AQ.MESSAGE_PROPERTIES_T,
<payload> IN <type_name>,
<msgid> OUT RAW)
```

Parameters

`queue_name`

The name (optionally schema-qualified) of an existing queue. If you omit the schema name, the server will use the schema specified in the `SEARCH_PATH`. Please note that unlike Oracle, unquoted identifiers are converted to lower case before storing. To include special characters or use a case-sensitive name, enclose the name in double quotes.

For detailed information about creating a queue, see `DBMS_AQADM.CREATE_QUEUE`.

`enqueue_options`

`enqueue_options` is a value of the type, `enqueue_options_t`:

```
DBMS_AQ.ENQUEUE_OPTIONS_T IS RECORD(
visibility BINARY_INTEGER DEFAULT ON_COMMIT,
relative_msgid RAW(16) DEFAULT NULL,
sequence_deviation BINARY_INTEGER DEFAULT NULL,
transformation VARCHAR2(61) DEFAULT NULL,
delivery_mode PLS_INTEGER NOT NULL DEFAULT PERSISTENT);
```

Currently, the only supported parameter values for `enqueue_options_t` are:

| | |
|---------------------------------|--------------------------|
| <code>visibility</code> | <code>ON_COMMIT</code> . |
| <code>delivery_mode</code> | <code>PERSISTENT</code> |
| <code>sequence_deviation</code> | <code>NULL</code> |
| <code>transformation</code> | <code>NULL</code> |

| | |
|----------------|------------|
| visibility | ON_COMMIT. |
| relative_msgid | NULL |

message_properties

`message_properties` is a value of the type, `message_properties_t`:

```
message_properties_t IS RECORD(
    priority INTEGER,
    delay INTEGER,
    expiration INTEGER,
    correlation CHARACTER VARYING(128) COLLATE pg_catalog."C",
    attempts INTEGER,
    recipient_list "AQ$_RECIPIENT_LIST_T",
    exception_queue CHARACTER VARYING(61) COLLATE pg_catalog."C",
    enqueue_time TIMESTAMP WITHOUT TIME ZONE,
    state INTEGER,
    original_msgid BYTEA,
    transaction_group CHARACTER VARYING(30) COLLATE pg_catalog."C",
    delivery_mode INTEGER
    DBMS_AQ.PERSISTENT);
```

The supported values for `message_properties_t` are:

| | |
|------------------------------|---|
| <code>priority</code> | If the queue table definition includes a <code>sort_list</code> that references <code>priority</code> , this parameter affects the order that messages are dequeued. A lower value indicates a higher dequeue priority. |
| <code>delay</code> | Specify the number of seconds that will pass before a message is available for dequeuing or <code>NO_DELAY</code> . |
| <code>expiration</code> | Use the <code>expiration</code> parameter to specify the number of seconds until a message expires. |
| <code>correlation</code> | Use <code>correlation</code> to specify a message that will be associated with the entry; the default is <code>NULL</code> . |
| <code>attempts</code> | This is a system-maintained value that specifies the number of attempts to dequeue the message. |
| <code>recipient_list</code> | This parameter is not supported. |
| <code>exception_queue</code> | Use the <code>exception_queue</code> parameter to specify the name of an exception queue to which a message will be moved if it expires or is dequeued by a transaction that rolls back too many times. |
| <code>enqueue_time</code> | <code>enqueue_time</code> is the time the record was added to the queue; this value is provided by the system. |

This parameter is maintained by DBMS_AQ; state can be:

| | |
|--------------------------------|---|
| <code>DBMS_AQ.WAITING</code> | – the delay has not been reached. |
| <code>DBMS_AQ.READY</code> | – the queue entry is ready for processing. |
| <code>DBMS_AQ.PROCESSED</code> | – the queue entry has been processed. |
| <code>DBMS_AQ.EXPIRED</code> | – the queue entry has been moved to the exception queue. |
| <code>original_msgid</code> | This parameter is accepted for compatibility and ignored. |
| <code>transaction_group</code> | This parameter is accepted for compatibility and ignored. |
| <code>delivery_mode</code> | This parameter is not supported; specify a value of <code>DBMS_AQ.PERSISTENT</code> . |

payload

Use the `payload` parameter to provide the data that will be associated with the queue entry. The payload type must match the type specified when creating the corresponding queue table (see `DBMS_AQADM.CREATE_QUEUE_TABLE`).

msgid

Use the `msgid` parameter to retrieve a unique (system-generated) message identifier.

Example

The following anonymous block calls `DBMS_AQ.ENQUEUE`, adding a message to a queue named `work_order`:

```
DECLARE
```

```
enqueue_options  DBMS_AQ.ENQUEUE_OPTIONS_T;
message_properties DBMS_AQ.MESSAGE_PROPERTIES_T;
message_handle    raw(16);
payload          work_order;
```

```
BEGIN
```

```
payload := work_order('Smith', 'system upgrade');
```

```
DBMS_AQ.ENQUEUE(
queue_name      => 'work_order',
enqueue_options  => enqueue_options,
message_properties => message_properties,
payload        => payload,
msgid         => message_handle
);
END;
```

6.3.2.2 DEQUEUE

The `DEQUEUE` procedure dequeues a message. The signature is:

```
DEQUEUE(
<queue_name> IN VARCHAR2,
<dequeue_options> IN DBMS_AQ.DEQUEUE_OPTIONS_T,
<message_properties> OUT DBMS_AQ.MESSAGE_PROPERTIES_T,
<payload> OUT <type_name>,
<msgid> OUT RAW)
```

Parameters

queue_name

The name (optionally schema-qualified) of an existing queue. If you omit the schema name, the server will use the schema specified in the `SEARCH_PATH`. Please note that unlike Oracle, unquoted identifiers are converted to lower case before storing. To include special characters or use a case-sensitive name, enclose the name in double quotes.

For detailed information about creating a queue, see `DBMS_AQADM.CREATE_QUEUE`.

`dequeue_options`

`dequeue_options` is a value of the type, `dequeue_options_t`:

```
DEQUEUE_OPTIONS_T IS RECORD(
    consumer_name CHARACTER VARYING(30),
    dequeue_mode INTEGER,
    navigation INTEGER,
    visibility INTEGER,
    wait INTEGER,
    msgid BYTEA,
    correlation CHARACTER VARYING(128),
    deq_condition CHARACTER VARYING(4000),
    transformation CHARACTER VARYING(61),
    delivery_mode INTEGER);
```

Currently, the supported parameter values for `dequeue_options_t` are:

`consumer_name` Must be `NULL`.

The locking behavior of the dequeue operation. Must be either:

`DBMS_AQ.BROWSE` – Read the message without obtaining a lock.

`dequeue_mode` `DBMS_AQ.LOCKED` – Read the message after acquiring a lock.

`DBMS_AQ.REMOVE` – Read the message before deleting the message.

`DBMS_AQ.REMOVE_NODATA` – Read the message, but do not delete the message.

Identifies the message that will be retrieved. Must be either:

`navigation` `FIRST_MESSAGE` – The first message within the queue that matches the search term.

`NEXT_MESSAGE` – The next message that is available that matches the first term.

`visibility` Must be `ON_COMMIT` – if you roll back the current transaction the dequeued item will remain in the queue.

Must be a number larger than 0, or:

`wait` `DBMS_AQ.FOREVER` – Wait indefinitely.

`DBMS_AQ.NO_WAIT` – Do not wait.

`msgid` The message ID of the message that will be dequeued.

`correlation` Accepted for compatibility, and ignored.

`deq_condition` A `VARCHAR2` expression that evaluates to a `BOOLEAN` value indicating if the message should be dequeued

`transformation` Accepted for compatibility, and ignored.

`delivery_mode` Must be `PERSISTENT`; buffered messages are not supported at this time.

`message_properties`

`message_properties` is a value of the type, `message_properties_t`:

```
message_properties_t IS RECORD(
    priority INTEGER,
    delay INTEGER,
    expiration INTEGER,
    correlation CHARACTER VARYING(128) COLLATE pg_catalog."C",
    attempts INTEGER,
    recipient_list "AQ$_RECIPIENT_LIST_T",
    exception_queue CHARACTER VARYING(61) COLLATE pg_catalog."C",
    enqueue_time TIMESTAMP WITHOUT TIME ZONE,
    state INTEGER,
    original_msgid BYTEA,
    transaction_group CHARACTER VARYING(30) COLLATE pg_catalog."C",
    delivery_mode INTEGER
DBMS_AQ.PERSISTENT);
```

The supported values for `message_properties_t` are:

| | |
|--------------------------------|--|
| <code>priority</code> | If the queue table definition includes a <code>sort_list</code> that references priority, this parameter affects the order that messages are dequeued. A lower value indicates a higher dequeue priority. |
| <code>delay</code> | Specify the number of seconds that will pass before a message is available for dequeuing or <code>NO_DELAY</code> . |
| <code>expiration</code> | Use the expiration parameter to specify the number of seconds until a message expires. |
| <code>correlation</code> | Use correlation to specify a message that will be associated with the entry; the default is <code>NULL</code> . |
| <code>attempts</code> | This is a system-maintained value that specifies the number of attempts to dequeue the message. |
| <code>recipient_list</code> | This parameter is not supported. |
| <code>exception_queue</code> | Use the <code>exception_queue</code> parameter to specify the name of an exception queue to which a message will be moved if it expires or is dequeued by a transaction that rolls back too many times. |
| <code>enqueue_time</code> | <code>enqueue_time</code> is the time the record was added to the queue; this value is provided by the system. |
| <code>state</code> | This parameter is maintained by <code>DBMS_AQ</code> ; state can be: <code>DBMS_AQ.WAITING</code> – the delay has not been reached. <code>DBMS_AQ.READY</code> – the queue entry is ready for processing. <code>DBMS_AQ.PROCESSED</code> – the queue entry has been processed. <code>DBMS_AQ.EXPIRED</code> – the queue entry has been moved to the exception queue. |
| <code>original_msgid</code> | This parameter is accepted for compatibility and ignored. |
| <code>transaction_group</code> | This parameter is accepted for compatibility and ignored. |
| <code>delivery_mode</code> | This parameter is not supported; specify a value of <code>DBMS_AQ.PERSISTENT</code> . |

payload

Use the `payload` parameter to retrieve the payload of a message with a dequeue operation. The `payload` type must match the type specified when creating the queue table.

msgid

Use the `msgid` parameter to retrieve a unique message identifier.

Example

The following anonymous block calls `DBMS_AQ.DEQUEUE`, retrieving a message from the queue and a payload:

```

DECLARE
    dequeue_options  DBMS_AQ.DEQUEUE_OPTIONS_T;
    message_properties DBMS_AQ.MESSAGE_PROPERTIES_T;
    message_handle    raw(16);
    payload          work_order;

BEGIN
    dequeue_options.dequeue_mode := DBMS_AQ.BROWSE;

    DBMS_AQ.DEQUEUE(
        queue_name      => 'work_queue',
        dequeue_options => dequeue_options,
        message_properties => message_properties,
        payload         => payload,
        msgid          => message_handle
    );
    DBMS_OUTPUT.PUT_LINE(
        'The next work order is [' || payload.subject || ']'
    );
END;

```

The payload is displayed by `DBMS_OUTPUT.PUT_LINE`.

6.3.2.3 REGISTER

Use the `REGISTER` procedure to register an email address, procedure or URL that will be notified when an item is enqueued or dequeued. The signature is:

```
REGISTER(
    <reg_list> IN SYS.AQ$_REG_INFO_LIST,
    <count> IN NUMBER)
```

Parameters

`reg_list`

`reg_list` is a list of type `AQ$_REG_INFO_LIST`; that provides information about each subscription that you would like to register. Each entry within the list is of the type `AQ$_REG_INFO`, and may contain:

`Y{0..2}`

| Attribute | Type | Description |
|------------------------|----------------|---|
| <code>name</code> | VARCHAR2 (128) | The (optionally schema-qualified) name of the subscription. |
| <code>namespace</code> | NUMERIC | The only supported value is <code>DBMS_AQ.NAMESPACE_AQ (0)</code> |

Describes the action that will be performed upon notification. Currently, only calls to PL/SQL procedures are supported. The call should take the form:

`plsql://schema.procedure`

`callback` `VARCHAR2(4000)`

Where:

`schema` specifies the schema in which the procedure resides.

`procedure` specifies the name of the procedure that will be notified.

`context` `RAW (16)` Any user-defined value required by the callback procedure.

`count`

`count` is the number of entries in `reg_list`.

Example

The following anonymous block calls `DBMS_AQ.REGISTER`, registering procedures that will be notified when an item is added to or removed from a queue. A set of attributes (of `sys.aq$_reg_info` type) is provided for each subscription identified in the `DECLARE` section:

```

DECLARE
    subscription1 sys.aq$_reg_info;
    subscription2 sys.aq$_reg_info;
    subscription3 sys.aq$_reg_info;
    subscriptionlist sys.aq$_reg_info_list;
BEGIN
    subscription1 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://assign_worker?PR=0',HEXTORAW('FFFF'));
    subscription2 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://add_to_history?PR=1',HEXTORAW('FFFF'));
    subscription3 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
    'plsql://reserve_parts?PR=2',HEXTORAW('FFFF'));

    subscriptionlist := sys.aq$_reg_info_list(subscription1,
    subscription2, subscription3);
    dbms_aq.register(subscriptionlist, 3);
    commit;

END;
/

```

The `subscriptionlist` is of type `sys.aq$_reg_info_list`, and contains the previously described `sys.aq$_reg_info` objects. The list name and an object count are passed to `dbms_aq.register`.

6.3.2.4 UNREGISTER

Use the `UNREGISTER` procedure to turn off notifications related to enqueueing and dequeuing. The signature is:

```
UNREGISTER(
<reg_list> IN SYS.AQ$_REG_INFO_LIST,
<count> IN NUMBER)
```

Parameter

reg_list

`reg_list` is a list of type `AQ$_REG_INFO_LIST`; that provides information about each subscription that you would like to register. Each entry within the list is of the type `AQ$_REG_INFO`, and may contain:

`Y{0..2}`

| Attribute | Type | Description |
|------------------------|----------------------------|---|
| <code>name</code> | <code>VARCHAR2(128)</code> | The (optionally schema-qualified) name of the subscription. |
| <code>namespace</code> | <code>NUMERIC</code> | The only supported value is <code>DBMS_AQ.NAMESPACE_AQ (0)</code> |

Describes the action that will be performed upon notification. Currently, only calls to PL/SQL procedures are supported. The call should take the form:

`plsql://schema.procedure`

| | | |
|-----------------------|-----------------------------|--------|
| <code>callback</code> | <code>VARCHAR2(4000)</code> | Where: |
|-----------------------|-----------------------------|--------|

schema specifies the schema in which the procedure resides.

procedure specifies the name of the procedure that will be notified.

| | | |
|----------------------|-----------------------|---|
| <code>context</code> | <code>RAW (16)</code> | Any user-defined value required by the procedure. |
|----------------------|-----------------------|---|

count

`count` is the number of entries in `reg_list`.

Example

The following anonymous block calls `DBMS_AQ.UNREGISTER`, disabling the notifications specified in the example for `DBMS_AQ.REGISTER`:

```
DECLARE
subscription1 sys.aq$_reg_info;
subscription2 sys.aq$_reg_info;
subscription3 sys.aq$_reg_info;
subscriptionlist sys.aq$_reg_info_list;
BEGIN
subscription1 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
'plsql://assign_worker?PR=0',HEXTORAW('FFFF'));
subscription2 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
'plsql://add_to_history?PR=1',HEXTORAW('FFFF'));
subscription3 := sys.aq$_reg_info('q', DBMS_AQ.NAMESPACE_AQ,
'plsql://reserve_parts?PR=2',HEXTORAW('FFFF'));

subscriptionlist := sys.aq$_reg_info_list(subscription1,
subscription2, subscription3);
dbms_aq.unregister(subscriptionlist, 3);
commit;
```

```
END;
/
```

The `subscriptionlist` is of type `sys.aq$_reg_info_list`, and contains the previously described `sys.aq$_reg_info` objects. The list name and an object count are passed to `dbms_aq.unregister`.

6.3.3 DBMS_AQADM

EDB Postgres Advanced Server Advanced Queueing provides message queueing and message processing for the Advanced Server database. User-defined messages are stored in a queue; a collection of queues is stored in a queue table. Procedures in the DBMS_AQADM package create and manage message queues and queue tables. Use the DBMS_AQ package to add messages to a queue or remove messages from a queue, or register or unregister a PL/SQL callback procedure.

Advanced Server also provides extended (non-compatible) functionality for the DBMS_AQ package with SQL commands, see the *Database Compatibility for Oracle Developers SQL Guide* for detailed information about the following SQL commands:

- `ALTER QUEUE`
- `ALTER QUEUE TABLE`
- `CREATE QUEUE`
- `CREATE QUEUE TABLE`
- `DROP QUEUE`
- `DROP QUEUE TABLE`

The DBMS_AQADM package provides procedures that allow you to create and manage queues and queue tables.

| Function/Procedure | Return Type | Description |
|---------------------------------|-------------|---|
| <code>ALTER_QUEUE</code> | n/a | Modify an existing queue. |
| <code>ALTER_QUEUE_TABLE</code> | n/a | Modify an existing queue table. |
| <code>CREATE_QUEUE</code> | n/a | Create a queue. |
| <code>CREATE_QUEUE_TABLE</code> | n/a | Create a queue table. |
| <code>DROP_QUEUE</code> | n/a | Drop an existing queue. |
| <code>DROP_QUEUE_TABLE</code> | n/a | Drop an existing queue table. |
| <code>PURGE_QUEUE_TABLE</code> | n/a | Remove one or more messages from a queue table. |
| <code>START_QUEUE</code> | n/a | Make a queue available for enqueueing and dequeueing procedures. |
| <code>STOP_QUEUE</code> | n/a | Make a queue unavailable for enqueueing and dequeueing procedures |

Advanced Server's implementation of `DBMS_AQADM` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Advanced Server supports use of the arguments listed below:

| Constant | Description | For Parameters |
|--|---|-------------------------------|
| <code>DBMS_AQADM.TRANSACTIONAL(1)</code> | This constant is defined, but will return an error if used. | <code>message_grouping</code> |

| Constant | Description | For Parameters |
|---------------------------------------|---|---------------------------------|
| DBMS_AQADM.NONE(0) | Use to specify message grouping for a queue table. | message_grouping |
| DBMS_AQADM.NORMAL_QUEUE(0) | Use with <code>create queue</code> to specify <code>queue_type</code> . | queue_type |
| DBMS_AQADM.EXCEPTION_QUEUE (1) | Use with <code>create queue</code> to specify <code>queue_type</code> . | queue_type |
| DBMS_AQADM.INFINITE(-1) | Use with <code>create queue</code> to specify <code>retention_time</code> . | retention_time |
| DBMS_AQADM.PERSISTENT (0) | The message should be stored in a table. | enqueue_options_t.delivery_mode |
| DBMS_AQADM.BUFFERED (1) | This constant is defined, but will return an error if used. | enqueue_options_t.delivery_mode |
| DBMS_AQADM.PERSISTENT_OR_BUFFERED (2) | This constant is defined, but will return an error if used. | enqueue_options_t.delivery_mode |

6.3.3.1 ALTER_QUEUE

Use the `ALTER_QUEUE` procedure to modify an existing queue. The signature is:

```
ALTER_QUEUE(
  <max_retries> IN NUMBER DEFAULT NULL,
  <retry_delay> IN NUMBER DEFAULT 0
  <retention_time> IN NUMBER DEFAULT 0,
  <auto_commit> IN BOOLEAN DEFAULT TRUE)
  <comment> IN VARCHAR2 DEFAULT NULL,
```

Parameters

`queue_name`

The name of the new queue.

`max_retries`

`max_retries` specifies the maximum number of attempts to remove a message with a `dequeue` statement. The value of `max_retries` is incremented with each `ROLLBACK` statement. When the number of failed attempts reaches the value specified by `max_retries`, the message is moved to the exception queue. Specify `0` to indicate that no retries are allowed.

`retry_delay`

`retry_delay` specifies the number of seconds until a message is scheduled for re-processing after a `ROLLBACK`. Specify `0` to indicate that the message should be retried immediately (the default).

`retention_time`

`retention_time` specifies the length of time (in seconds) that a message will be stored after being dequeued. You can also specify `0` (the default) to indicate the message should not be retained after dequeuing, or

| `INFINITE` to retain the message forever.

`auto_commit`

| This parameter is accepted for compatibility and ignored.

`comment`

| `comment` specifies a comment associated with the queue.

Example

The following command alters a queue named `work_order`, setting the `retry_delay` parameter to 5 seconds:

```
EXEC DBMS_AQADM.ALTER_QUEUE(queue_name => 'work_order', retry_delay => 5);
```

6.3.3.2 ALTER_QUEUE_TABLE

Use the `ALTER_QUEUE_TABLE` procedure to modify an existing queue table. The signature is:

```
ALTER_QUEUE_TABLE (
  <queue_table> IN VARCHAR2,
  <comment> IN VARCHAR2 DEFAULT NULL,
  <primary_instance> IN BINARY_INTEGER DEFAULT 0,
  <secondary_instance> IN BINARY_INTEGER DEFAULT 0,
```

Parameters

`queue_table`

| The (optionally schema-qualified) name of the queue table.

`comment`

| Use the `comment` parameter to provide a comment about the queue table.

`primary_instance`

| `primary_instance` is accepted for compatibility and stored, but is ignored.

`secondary_instance`

| `secondary_instance` is accepted for compatibility, but is ignored.

Example

The following command modifies a queue table named `work_order_table`:

```
EXEC DBMS_AQADM.ALTER_QUEUE_TABLE
  (queue_table => 'work_order_table', comment => 'This queue table
contains work orders for the shipping department.');
```

The queue table is named `work_order_table`; the command adds a comment to the definition of the queue table.

6.3.3.3 CREATE_QUEUE

Use the `CREATE_QUEUE` procedure to create a queue in an existing queue table. The signature is:

```
CREATE_QUEUE(
  <queue_name> IN VARCHAR2
  <queue_table> IN VARCHAR2,
  <queue_type> IN BINARY_INTEGER DEFAULT NORMAL_QUEUE,
  <max_retries> IN NUMBER DEFAULT 5,
  <retry_delay> IN NUMBER DEFAULT 0
  <retention_time> IN NUMBER DEFAULT 0,
  <dependency_tracking> IN BOOLEAN DEFAULT FALSE,
  <comment> IN VARCHAR2 DEFAULT NULL,
  <auto_commit> IN BOOLEAN DEFAULT TRUE)
```

Parameters

`queue_name`

The name of the new queue.

`queue_table`

The name of the table in which the new queue will reside.

`queue_type`

The type of the new queue. The valid values for `queue_type` are:

`DBMS_AQADM.NORMAL_QUEUE` – This value specifies a normal queue (the default).

`DBMS_AQADM.EXCEPTION_QUEUE` – This value specifies that the new queue is an exception queue. An exception queue will support only dequeue operations.

`max_retries`

`max_retries` specifies the maximum number of attempts to remove a message with a dequeue statement. The value of `max_retries` is incremented with each `ROLLBACK` statement. When the number of failed attempts reaches the value specified by `max_retries`, the message is moved to the exception queue. The default value for a system table is `0`; the default value for a user created table is `5`.

`retry_delay`

`retry_delay` specifies the number of seconds until a message is scheduled for re-processing after a `ROLLBACK`. Specify `0` to indicate that the message should be retried immediately (the default).

`retention_time`

`retention_time` specifies the length of time (in seconds) that a message will be stored after being dequeued. You can also specify `0` (the default) to indicate the message should not be retained after dequeuing, or `INFINITE` to retain the message forever.

`dependency_tracking`

This parameter is accepted for compatibility and ignored.

`comment`

`comment` specifies a comment associated with the queue.

`auto_commit`

This parameter is accepted for compatibility and ignored.

Example

The following anonymous block creates a queue named `work_order` in the `work_order_table` table:

```
BEGIN
DBMS_AQADM.CREATE_QUEUE ( queue_name => 'work_order', queue_table =>
'work_order_table', comment => 'This queue contains pending work orders.');
END;
```

6.3.3.4 CREATE_QUEUE_TABLE

Use the `CREATE_QUEUE_TABLE` procedure to create a queue table. The signature is:

```
CREATE_QUEUE_TABLE (
<queue_table> IN VARCHAR2,
<queue_payload_type> IN VARCHAR2,
<storage_clause> IN VARCHAR2 DEFAULT NULL,
<sort_list> IN VARCHAR2 DEFAULT NULL,
<multiple_consumers> IN BOOLEAN DEFAULT FALSE,
<message_grouping> IN BINARY_INTEGER DEFAULT NONE,
<comment> IN VARCHAR2 DEFAULT NULL,
<auto_commit> IN BOOLEAN DEFAULT TRUE,
<primary_instance> IN BINARY_INTEGER DEFAULT 0,
<secondary_instance> IN BINARY_INTEGER DEFAULT 0,
<compatible> IN VARCHAR2 DEFAULT NULL,
<secure> IN BOOLEAN DEFAULT FALSE)
```

Parameters

`queue_table`

The (optionally schema-qualified) name of the queue table.

`queue_payload_type`

The user-defined type of the data that will be stored in the queue table. Please note that to specify a `RAW` data type, you must create a user-defined type that identifies a `RAW` type.

`storage_clause`

Use the `storage_clause` parameter to specify attributes for the queue table. Please note that only the `TABLESPACE` option is enforced; all others are accepted for compatibility and ignored. Use the `TABLESPACE` clause to specify the name of a tablespace in which the table will be created.

`storage_clause` may be one or more of the following:

`TABLESPACE tablespace_name, PCTFREE integer, PCTUSED integer, INITTRANS`

integer, MAXTRANS integer or STORAGE storage_option.

`storage_option` may be one or more of the following:

MINEXTENTS integer, MAXEXTENTS integer, PCTINCREASE integer, INITIAL size_clause, NEXT, FREELISTS integer, OPTIMAL size_clause, BUFFER_POOL {KEEP|RECYCLE|DEFAULT}.

sort_list

`sort_list` controls the dequeuing order of the queue; specify the names of the column(s) that will be used to sort the queue (in ascending order). The currently accepted values are the following combinations of `enq_time` and `priority`:

`enq_time, priority`

`priority, enq_time`

`priority`

`enq_time`

multiple_consumers

`multiple_consumers` queue tables is not supported.

message_grouping

If specified, `message_grouping` must be `NONE`.

comment

Use the `comment` parameter to provide a comment about the queue table.

auto_commit

`auto_commit` is accepted for compatibility, but is ignored.

primary_instance

`primary_instance` is accepted for compatibility and stored, but is ignored.

secondary_instance

`secondary_instance` is accepted for compatibility, but is ignored.

compatible

`compatible` is accepted for compatibility, but is ignored.

secure

`secure` is accepted for compatibility, but is ignored.

Example

The following anonymous block first creates a type (`work_order`) with attributes that hold a name (a `VARCHAR2`), and a project description (a `TEXT`). The block then uses that type to create a queue table:

```
BEGIN
```

```

CREATE TYPE work_order AS (name VARCHAR2, project TEXT, completed BOOLEAN);

EXEC DBMS_AQADM.CREATE_QUEUE_TABLE
  (queue_table => 'work_order_table',
   queue_payload_type => 'work_order',
   comment => 'Work order message queue table');

END;

```

The queue table is named `work_order_table`, and contains a payload of a type `work_order`. A comment notes that this is the `Work order message queue table`.

6.3.3.5 DROP_QUEUE

Use the `DROP_QUEUE` procedure to delete a queue. The signature is:

```

DROP_QUEUE(
  <queue_name> IN VARCHAR2,
  <auto_commit> IN BOOLEAN DEFAULT TRUE)

```

Parameters

`queue_name`

| The name of the queue that you wish to drop.

`auto_commit`

| `auto_commit` is accepted for compatibility, but is ignored.

Example

The following anonymous block drops the queue named `work_order`:

```

BEGIN
DBMS_AQADM.DROP_QUEUE(queue_name => 'work_order');
END;

```

6.3.3.6 DROP_QUEUE_TABLE

Use the `DROP_QUEUE_TABLE` procedure to delete a queue table. The signature is:

```

DROP_QUEUE_TABLE(
  <queue_table> IN VARCHAR2,
  <force> IN BOOLEAN default FALSE,
  <auto_commit> IN BOOLEAN default TRUE)

```

Parameters

queue_table

The (optionally schema-qualified) name of the queue table.

force

The `force` keyword determines the behavior of the `DROP_QUEUE_TABLE` command when dropping a table that contain entries:

If the target table contains entries and force is `FALSE`, the command will fail, and the server will issue an error.

If the target table contains entries and force is `TRUE`, the command will drop the table and any dependent objects.

auto_commit

`auto_commit` is accepted for compatibility, but is ignored.

Example

The following anonymous block drops a table named `work_order_table`:

```
BEGIN
  DBMS_AQADM.DROP_QUEUE_TABLE ('work_order_table', force => TRUE);
END;
```

6.3.3.7 PURGE_QUEUE_TABLE

Use the `PURGE_QUEUE_TABLE` procedure to delete messages from a queue table. The signature is:

```
PURGE_QUEUE_TABLE(
  <queue_table> IN VARCHAR2,
  <purge_condition> IN VARCHAR2,
  <purge_options> IN aq$_purge_options_t)
```

Parameters

queue_table

`queue_table` specifies the name of the queue table from which you are deleting a message.

purge_condition

Use `purge_condition` to specify a condition (a SQL `WHERE` clause) that the server will evaluate when deciding which messages to purge.

purge_options

`purge_options` is an object of the type `aq$_purge_options_t`. An `aq$_purge_options_t` object contains:

| Attribute | Type | Description |
|-----------|------|-------------|
|-----------|------|-------------|

| Attribute | Type | Description |
|----------------------------|---------|--|
| Block | Boolean | Specify <code>TRUE</code> if an exclusive lock should be held on all queues within the table; the default is <code>FALSE</code> . |
| <code>delivery_mode</code> | INTEGER | <code>delivery_mode</code> specifies the type of message that will be purged. The only accepted value is <code>DBMS_AQ.PERSISTENT</code> . |

Example

The following anonymous block removes any messages from the `work_order_table` with a value in the `completed` column of `YES`:

```
DECLARE
    purge_options dbms_aqadm.aq$_purge_options_t;
BEGIN
    dbms_aqadm.purge_queue_table('work_order_table', 'completed = YES',
    purge_options);
END;
```

6.3.3.8 START_QUEUE

Use the `START_QUEUE` procedure to make a queue available for enqueueing and dequeuing. The signature is:

```
START_QUEUE(
    <queue_name> IN VARCHAR2,
    <enqueue> IN BOOLEAN DEFAULT TRUE,
    <dequeue> IN BOOLEAN DEFAULT TRUE)
```

Parameters

`queue_name`

`queue_name` specifies the name of the queue that you are starting.

`enqueue`

Specify `TRUE` to enable enqueueing (the default), or `FALSE` to leave the current setting unchanged.

`dequeue`

Specify `TRUE` to enable dequeuing (the default), or `FALSE` to leave the current setting unchanged.

Example

The following anonymous block makes a queue named `work_order` available for enqueueing:

```
BEGIN
DBMS_AQADM.START_QUEUE
(queue_name => 'work_order');
END;
```

6.3.3.9 STOP_QUEUE

Use the `STOP_QUEUE` procedure to disable enqueueing or dequeuing on a specified queue. The signature is:

```
STOP_QUEUE(
  <queue_name> IN VARCHAR2,
  <enqueue> IN BOOLEAN DEFAULT TRUE,
  <dequeue> IN BOOLEAN DEFAULT TRUE,
  <wait> IN BOOLEAN DEFAULT TRUE)
```

Parameters

`queue_name`

`queue_name` specifies the name of the queue that you are stopping.

`enqueue`

Specify `TRUE` to disable enqueueing (the default), or `FALSE` to leave the current setting unchanged.

`dequeue`

Specify `TRUE` to disable dequeuing (the default), or `FALSE` to leave the current setting unchanged.

`wait`

Specify `TRUE` to instruct the server to wait for any uncompleted transactions to complete before applying the specified changes; while waiting to stop the queue, no transactions are allowed to enqueue or dequeue from the specified queue. Specify `FALSE` to stop the queue immediately.

Example

The following anonymous block disables enqueueing and dequeuing from the queue named `work_order`:

```
BEGIN
DBMS_AQADM.STOP_QUEUE(queue_name =>'work_order', enqueue=>TRUE,
dequeue=>TRUE, wait=>TRUE);
END;
```

Enqueueing and dequeuing will stop after any outstanding transactions complete.

6.3.4 DBMS_CRYPTO

The `DBMS_CRYPTO` package provides functions and procedures that allow you to encrypt or decrypt `RAW`, `BLOB` or `CLOB` data. You can also use `DBMS_CRYPTO` functions to generate cryptographically strong random values.

The following table lists the `DBMS_CRYPTO` Functions and Procedures.

| Function/Procedure | Return Type | Description |
|---------------------------------------|-------------|---|
| DECRYPT(src, typ, key, iv) | RAW | Decrypts RAW data. |
| DECRYPT(dst INOUT, src, typ, key, iv) | N/A | Decrypts BLOB data. |
| DECRYPT(dst INOUT, src, typ, key, iv) | N/A | Decrypts CLOB data. |
| ENCRYPT(src, typ, key, iv) | RAW | Encrypts RAW data. |
| ENCRYPT(dst INOUT, src, typ, key, iv) | N/A | Encrypts BLOB data. |
| ENCRYPT(dst INOUT, src, typ, key, iv) | N/A | Encrypts CLOB data. |
| HASH(src, typ) | RAW | Applies a hash algorithm to RAW data. |
| HASH(src) | RAW | Applies a hash algorithm to CLOB data. |
| MAC(src, typ, key) | RAW | Returns the hashed MAC value of the given RAW data using the specified hash algorithm and key. |
| MAC(src, typ, key) | RAW | Returns the hashed MAC value of the given CLOB data using the specified hash algorithm and key. |
| RANDOMBYTES(number_bytes) | RAW | Returns a specified number of cryptographically strong random bytes. |
| RANDOMINTEGER() | INTEGER | Returns a random INTEGER. |
| RANDOMNUMBER() | NUMBER | Returns a random NUMBER. |

DBMS_CRYPTO functions and procedures support the following error messages:

ORA-28239 - DBMS_CRYPTO.KeyNull

ORA-28829 - DBMS_CRYPTO.CipherSuiteNull

ORA-28827 - DBMS_CRYPTO.CipherSuiteInvalid

Unlike Oracle, Advanced Server will not return error ORA-28233 if you re-encrypt previously encrypted information.

Please note that RAW and BLOB are synonyms for the PostgreSQL BYTEA data type, and CLOB is a synonym for TEXT.

6.3.4.1 DECRYPT

The DECRYPT function or procedure decrypts data using a user-specified cipher algorithm, key and optional initialization vector. The signature of the DECRYPT function is:

```
DECRYPT
(<src> IN RAW, <typ> IN INTEGER, <key> IN RAW, <iv> IN RAW
DEFAULT NULL) RETURN RAW
```

The signature of the DECRYPT procedure is:

```
DECRYPT
(<dst> INOUT BLOB, <src> IN BLOB, <typ> IN INTEGER, <key> IN RAW,
```

<iv> IN RAW DEFAULT NULL)

or

DECRYPT

(<dst> INOUT CLOB, <src> IN CLOB, <typ> IN INTEGER, <key> IN RAW,
<iv> IN RAW DEFAULT NULL)

When invoked as a procedure, **DECRYPT** returns **BLOB** or **CLOB** data to a user-specified **BLOB**.

Parameters

dst

dst specifies the name of a **BLOB** to which the output of the **DECRYPT** procedure will be written. The **DECRYPT** procedure will overwrite any existing data currently in **dst**.

src

src specifies the source data that will be decrypted. If you are invoking **DECRYPT** as a function, specify **RAW** data; if invoking **DECRYPT** as a procedure, specify **BLOB** or **CLOB** data.

typ

typ specifies the block cipher type and any modifiers. This should match the type specified when the **src** was encrypted. Advanced Server supports the following block cipher algorithms, modifiers and cipher suites:

Block Cipher Algorithms

| | |
|----------------|------------------------|
| ENCRYPT_DES | CONSTANT INTEGER := 1; |
| ENCRYPT_3DES | CONSTANT INTEGER := 3; |
| ENCRYPT_AES | CONSTANT INTEGER := 4; |
| ENCRYPT_AES128 | CONSTANT INTEGER := 6; |

Block Cipher Modifiers

| | |
|-----------|--------------------------|
| CHAIN_CBC | CONSTANT INTEGER := 256; |
| CHAIN_ECB | CONSTANT INTEGER := 768; |

Block Cipher Padding Modifiers

| | |
|-----------|---------------------------|
| PAD_PKCS5 | CONSTANT INTEGER := 4096; |
| PAD_NONE | CONSTANT INTEGER := 8192; |

Block Cipher Suites

| | |
|----------------|--|
| DES_CBC_PKCS5 | CONSTANT INTEGER := ENCRYPT_DES + CHAIN_CBC + PAD_PKCS5; |
| DES3_CBC_PKCS5 | CONSTANT INTEGER := ENCRYPT_3DES + CHAIN_CBC + PAD_PKCS5; |
| AES_CBC_PKCS5 | CONSTANT INTEGER := ENCRYPT_AES + CHAIN_CBC + PAD_PKCS5; |

key

key specifies the user-defined decryption key. This should match the key specified when the **src** was encrypted.

iv

iv (optional) specifies an initialization vector. If an initialization vector was specified when the **src** was encrypted, you must specify an initialization vector when decrypting the **src**. The default is **NULL**.

Examples

The following example uses the `DBMS_CRYPTO.DECRYPT` function to decrypt an encrypted password retrieved from the `passwords` table:

```
CREATE TABLE passwords
(
    principal VARCHAR2(90) PRIMARY KEY, -- username
    ciphertext RAW(9) -- encrypted password
);

CREATE FUNCTION get_password(username VARCHAR2) RETURN RAW AS
typ    INTEGER := DBMS_CRYPTO.DES_CBC_PKCS5;
key    RAW(128) := 'my secret key';
iv     RAW(100) := 'my initialization vector';
password RAW(2048);
BEGIN

    SELECT ciphertext INTO password FROM passwords WHERE principal = username;

    RETURN dbms_crypto.decrypt(password, typ, key, iv);
END;
```

Note that when calling `DECRYPT`, you must pass the same cipher type, key value and initialization vector that was used when `ENCRYPTING` the target.

6.3.4.2 ENCRYPT

The `ENCRYPT` function or procedure uses a user-specified algorithm, key, and optional initialization vector to encrypt `RAW`, `BLOB` or `CLOB` data. The signature of the `ENCRYPT` function is:

```
ENCRYPT
(<src> IN RAW, <typ> IN INTEGER, <key> IN RAW,
<iv> IN RAW DEFAULT NULL) RETURN RAW
```

The signature of the `ENCRYPT` procedure is:

```
ENCRYPT
(<dst> INOUT BLOB, <src> IN BLOB, <typ> IN INTEGER, <key> IN RAW,
<iv> IN RAW DEFAULT NULL)
```

or

```
ENCRYPT
(<dst> INOUT BLOB, <src> IN CLOB, <typ> IN INTEGER, <key> IN RAW,
<iv> IN RAW DEFAULT NULL)
```

When invoked as a procedure, `ENCRYPT` returns `BLOB` or `CLOB` data to a user-specified `BLOB`.

Parameters

`dst`

`dst` specifies the name of a `BLOB` to which the output of the `ENCRYPT` procedure will be written. The

ENCRYPT procedure will overwrite any existing data currently in `dst`.

src

`src` specifies the source data that will be encrypted. If you are invoking **ENCRYPT** as a function, specify `RAW` data; if invoking **ENCRYPT** as a procedure, specify `BLOB` or `CLOB` data.

typ

`typ` specifies the block cipher type that will be used by **ENCRYPT**, and any modifiers. Advanced Server supports the block cipher algorithms, modifiers and cipher suites listed below:

Block Cipher Algorithms

| | |
|-----------------------------|-------------------------------------|
| <code>ENCRYPT_DES</code> | <code>CONSTANT INTEGER := 1;</code> |
| <code>ENCRYPT_3DES</code> | <code>CONSTANT INTEGER := 3;</code> |
| <code>ENCRYPT_AES</code> | <code>CONSTANT INTEGER := 4;</code> |
| <code>ENCRYPT_AES128</code> | <code>CONSTANT INTEGER := 6;</code> |

Block Cipher Modifiers

| | |
|------------------------|---------------------------------------|
| <code>CHAIN_CBC</code> | <code>CONSTANT INTEGER := 256;</code> |
| <code>CHAIN_ECB</code> | <code>CONSTANT INTEGER := 768;</code> |

Block Cipher Padding**Modifiers**

| | |
|------------------------|--|
| <code>PAD_PKCS5</code> | <code>CONSTANT INTEGER := 4096;</code> |
| <code>PAD_NONE</code> | <code>CONSTANT INTEGER := 8192;</code> |

Block Cipher Suites

| | |
|-----------------------------|--|
| <code>DES_CBC_PKCS5</code> | <code>CONSTANT INTEGER := ENCRYPT_DES + CHAIN_CBC + PAD_PKCS5;</code> |
| <code>DES3_CBC_PKCS5</code> | <code>CONSTANT INTEGER := ENCRYPT_3DES + CHAIN_CBC + PAD_PKCS5;</code> |
| <code>AES_CBC_PKCS5</code> | <code>CONSTANT INTEGER := ENCRYPT_AES + CHAIN_CBC + PAD_PKCS5;</code> |

key

`key` specifies the encryption key.

iv

`iv` (optional) specifies an initialization vector. By default, `iv` is `NULL`.

Examples

The following example uses the `DBMS_CRYPTO.DES_CBC_PKCS5` Block Cipher Suite (a pre-defined set of algorithms and modifiers) to encrypt a value retrieved from the `passwords` table:

```
CREATE TABLE passwords
(
    principal VARCHAR2(90) PRIMARY KEY, -- username
    ciphertext RAW(9) -- encrypted password
);
CREATE PROCEDURE set_password(username VARCHAR2, cleartext RAW) AS
typ      INTEGER := DBMS_CRYPTO.DES_CBC_PKCS5;
key      RAW(128) := 'my secret key';
iv       RAW(100) := 'my initialization vector';
encrypted RAW(2048);
BEGIN
    encrypted := dbms_crypto.encrypt(cleartext, typ, key, iv);
```

```
UPDATE passwords SET ciphertext = encrypted WHERE principal = username;
END;
```

ENCRYPT uses a key value of **my secret key** and an initialization vector of **my initialization vector** when encrypting the **password**; specify the same key and initialization vector when decrypting the **password**.

6.3.4.3 HASH

The **HASH** function uses a user-specified algorithm to return the hash value of a **RAW** or **CLOB** value. The **HASH** function is available in three forms:

```
HASH
(<src> IN RAW, <typ> IN INTEGER) RETURN RAW
```

```
HASH
(<src> IN CLOB, <typ> IN INTEGER) RETURN RAW
```

Parameters

src

src specifies the value for which the hash value will be generated. You can specify a **RAW**, a **BLOB**, or a **CLOB** value.

typ

typ specifies the **HASH** function type. Advanced Server supports the **HASH** function types listed below:

HASH Functions

| | |
|-----------------|------------------------|
| HASH_MD4 | CONSTANT INTEGER := 1; |
| HASH_MD5 | CONSTANT INTEGER := 2; |
| HASH_SH1 | CONSTANT INTEGER := 3; |

Examples

The following example uses **DBMS_CRYPTO.HASH** to find the **md5** hash value of the string, **cleartext source**:

```
DECLARE
typ INTEGER := DBMS_CRYPTO.HASH_MD5;
hash_value RAW(100);
BEGIN
hash_value := DBMS_CRYPTO.HASH('cleartext source', typ);
END;
```

6.3.4.4 MAC

The `MAC` function uses a user-specified `MAC` function to return the hashed `MAC` value of a `RAW` or `CLOB` value. The `MAC` function is available in three forms:

```
MAC
(<src> IN RAW, <typ> IN INTEGER, <key> IN RAW) RETURN RAW
```

```
MAC
(<src> IN CLOB, <typ> IN INTEGER, <key> IN RAW) RETURN RAW
```

Parameters

`src`

`src` specifies the value for which the `MAC` value will be generated. Specify a `RAW`, `BLOB`, or `CLOB` value.

`typ`

`typ` specifies the `MAC` function used. Advanced Server supports the `MAC` functions listed below.

MAC Functions

| | |
|-----------------------|-------------------------------------|
| <code>HMAC_MD5</code> | <code>CONSTANT INTEGER := 1;</code> |
| <code>HMAC_SH1</code> | <code>CONSTANT INTEGER := 2;</code> |

`key`

`key` specifies the key that will be used to calculate the hashed `MAC` value.

Examples

The following example finds the hashed `MAC` value of the string `cleartext source`:

```
DECLARE
typ INTEGER := DBMS_CRYPTO.HMAC_MD5;
key RAW(100) := 'my secret key';
mac_value RAW(100);
BEGIN
mac_value := DBMS_CRYPTO.MAC('cleartext source', typ, key);
END;
```

`DBMS_CRYPTO.MAC` uses a key value of `my secret` key when calculating the `MAC` value of `cleartext source`.

6.3.4.5 RANDOMBYTES

The `RANDOMBYTES` function returns a `RAW` value of the specified length, containing cryptographically random bytes. The signature is:

RANDOMBYTES
(<number_bytes> IN INTEGER) RETURNS RAW

Parameter

number_bytes

| number_bytes specifies the number of random bytes to be returned

Examples

The following example uses RANDOMBYTES to return a value that is 1024 bytes long:

```
DECLARE
  result RAW(1024);
BEGIN
  result := DBMS_CRYPTO.RANDOMBYTES(1024);
END;
```

6.3.4.6 RANDOMINTEGER

The RANDOMINTEGER() function returns a random INTEGER between 0 and 268,435,455. The signature is:

RANDOMINTEGER() RETURNS INTEGER

Examples

The following example uses the RANDOMINTEGER function to return a cryptographically strong random INTEGER value:

```
DECLARE
  result INTEGER;
BEGIN
  result := DBMS_CRYPTO.RANDOMINTEGER();
  DBMS_OUTPUT.PUT_LINE(result);
END;
```

6.3.4.7 RANDOMNUMBER

The RANDOMNUMBER() function returns a random NUMBER between 0 and 268,435,455. The signature is:

RANDOMNUMBER() RETURNS NUMBER

Examples

The following example uses the RANDOMNUMBER function to return a cryptographically strong random number:

```

DECLARE
  result NUMBER;
BEGIN
  result := DBMS_CRYPTO.RANDOMNUMBER();
  DBMS_OUTPUT.PUT_LINE(result);
END;

```

6.3.5 DBMS_JOB

The `DBMS_JOB` package provides for the creation, scheduling, and managing of jobs. A job runs a stored procedure which has been previously stored in the database. The `SUBMIT` procedure is used to create and store a job definition. A job identifier is assigned to a job along with its associated stored procedure and the attributes describing when and how often the job is to be run.

This package relies on the `pgAgent` scheduler. By default, the Advanced Server installer installs `pgAgent`, but you must start the `pgAgent` service manually prior to using `DBMS_JOB`. If you attempt to use this package to schedule a job after un-installing `pgAgent`, `DBMS_JOB` will throw an error. `DBMS_JOB` verifies that `pgAgent` is installed, but does not verify that the service is running.

The following table lists the supported `DBMS_JOB` procedures:

| Function/Procedure | Return Type | Description |
|---|-------------|---|
| <code>BROKEN(job, broken [, next_date])</code> | n/a | Specify that a given job is either broken or not broken. |
| <code>CHANGE(job, what, next_date, interval, instance, force>)</code> | n/a | Change the job's parameters. |
| <code>INTERVAL(job, interval)</code> | n/a | Set the execution frequency by means of a date function that is recalculated each time the job is run. This value becomes the next date/time for execution. |
| <code>NEXT_DATE(job, next_date)</code> | n/a | Set the next date/time the job is to be run. |
| <code>REMOVE(job)</code> | n/a | Delete the job definition from the database. |
| <code>RUN(job)</code> | n/a | Forces execution of a job even if it is marked broken. |
| <code>SUBMIT(job OUT, what [, next_date [, interval [, no_parse]]])</code> | n/a | Creates a job and stores its definition in the database. |
| <code>WHAT(job, what)</code> | n/a | Change the stored procedure run by a job. |

Advanced Server's implementation of `DBMS_JOB` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Before using `DBMS_JOB`, a database superuser must create the `pgAgent` extension. Use the `psql` client to connect to a database and invoke the command:

```
| CREATE EXTENSION pgagent;
```

When and how often a job is run is dependent upon two interacting parameters – `next_date` and `interval`. The `next_date` parameter is a date/time value that specifies the next date/time when the job is to be executed. The `interval` parameter is a string that contains a date function that evaluates to a date/time value.

Just prior to any execution of the job, the expression in the `interval` parameter is evaluated. The resulting value replaces the `next_date` value stored with the job. The job is then executed. In this manner, the expression in `interval` is repeatedly re-evaluated prior to each job execution, supplying the `next_date` date/time for the next execution.

The following examples use the following stored procedure, `job_proc`, which simply inserts a timestamp into table, `jobrun`, containing a single `VARCHAR2` column.

```
CREATE TABLE jobrun (
    runtime VARCHAR2(40)
);

CREATE OR REPLACE PROCEDURE job_proc
IS
BEGIN
    INSERT INTO jobrun VALUES ('job_proc run at ' || TO_CHAR(SYSDATE,
        'yyyy-mm-dd hh24:mi:ss'));
END;
```

6.3.5.1 BROKEN

The `BROKEN` procedure sets the state of a job to either broken or not broken. A broken job cannot be executed except by using the `RUN` procedure.

`BROKEN(<job> BINARY_INTEGER, <broken> BOOLEAN [, <next_date> DATE])`

Parameters

`job`

Identifier of the job to be set as broken or not broken.

`broken`

If set to `TRUE` the job's state is set to broken. If set to `FALSE` the job's state is set to not broken. Broken jobs cannot be run except by using the `RUN` procedure.

`next_date`

Date/time when the job is to be run. The default is `SYSDATE`.

Examples

Set the state of a job with job identifier 104 to broken:

```
BEGIN
    DBMS_JOB.BROKEN(104,true);
END;
```

Change the state back to not broken:

```
BEGIN
    DBMS_JOB.BROKEN(104,false);
```

```
END;
```

6.3.5.2 CHANGE

The **CHANGE** procedure modifies certain job attributes including the stored procedure to be run, the next date/time the job is to be run, and how often it is to be run.

```
CHANGE(<job> BINARY_INTEGER <what> VARCHAR2, <next_date> DATE,
      <interval> VARCHAR2, <instance> BINARY_INTEGER, <force> BOOLEAN)
```

Parameters

job

Identifier of the job to modify.

what

Stored procedure name. Set this parameter to null if the existing value is to remain unchanged.

next_date

Date/time when the job is to be run next. Set this parameter to null if the existing value is to remain unchanged.

interval

Date function that when evaluated, provides the next date/time the job is to run. Set this parameter to null if the existing value is to remain unchanged.

instance

This argument is ignored, but is included for compatibility.

force

This argument is ignored, but is included for compatibility.

Examples

Change the job to run next on December 13, 2007. Leave other parameters unchanged.

```
BEGIN
  DBMS_JOB.CHANGE(104,NULL,TO_DATE('13-DEC-07','DD-MON-YY'),NULL, NULL,
  NULL);
END;
```

6.3.5.3 INTERVAL

The **INTERVAL** procedure sets the frequency of how often a job is to be run.

```
INTERVAL(<job> BINARY_INTEGER, <interval> VARCHAR2)
```

Parameters

job

Identifier of the job to modify.

interval

Date function that when evaluated, provides the next date/time the job is to be run. If **interval** is **NULL** and the job is complete, the job is removed from the queue.

Examples

Change the job to run once a week:

```
BEGIN
  DBMS_JOB.INTERVAL(104,'SYSDATE + 7');
END;
```

6.3.5.4 NEXT_DATE

The **NEXT_DATE** procedure sets the date/time of when the job is to be run next.

```
NEXT_DATE(<job> BINARY_INTEGER, <next_date> DATE)
```

Parameters

job

Identifier of the job whose next run date is to be set.

next_date

Date/time when the job is to be run next.

Examples

Change the job to run next on December 14, 2007:

```
BEGIN
  DBMS_JOB.NEXT_DATE(104, TO_DATE('14-DEC-07','DD-MON-YY'));
END;
```

6.3.5.5 REMOVE

The **REMOVE** procedure deletes the specified job from the database. The job must be resubmitted using the **SUBMIT** procedure in order to have it executed again. Note that the stored procedure that was associated with the job is not deleted.

```
REMOVE(<job> BINARY_INTEGER)
```

Parameter

job

Identifier of the job that is to be removed from the database.

Examples

Remove a job from the database:

```
BEGIN
  DBMS_JOB.REMOVE(104);
END;
```

6.3.5.6 RUN

The **RUN** procedure forces the job to be run, even if its state is broken.

```
RUN(<job> BINARY_INTEGER)
```

Parameter

job

Identifier of the job to be run.

Examples

Force a job to be run.

```
BEGIN
  DBMS_JOB.RUN(104);
END;
```

6.3.5.7 SUBMIT

The **SUBMIT** procedure creates a job definition and stores it in the database. A job consists of a job identifier, the stored procedure to be executed, when the job is to be first run, and a date function that calculates the next date/time the job is to be run.

```
SUBMIT(<job> OUT BINARY_INTEGER, <what> VARCHAR2
[, <next_date> DATE [, <interval> VARCHAR2 [, <no_parse> BOOLEAN ]]])
```

Parameters

`job`

Identifier assigned to the job.

`what`

Name of the stored procedure to be executed by the job.

`next_date`

Date/time when the job is to be run next. The default is `SYSDATE`.

`interval`

Date function that when evaluated, provides the next date/time the job is to run. If `interval` is set to null, then the job is run only once. Null is the default.

`no_parse`

If set to `TRUE`, do not syntax-check the stored procedure upon job creation – check only when the job first executes. If set to `FALSE`, check the procedure upon job creation. The default is `FALSE`.

Note

The `no_parse` option is not supported in this implementation of `SUBMIT()`. It is included for compatibility only.

Examples

The following example creates a job using stored procedure, `job proc`. The job will execute immediately and run once a day thereafter as set by the `interval` parameter, `SYSDATE + 1`.

```
DECLARE
    jobid INTEGER;
BEGIN
    DBMS_JOB.SUBMIT(jobid,'job_proc',SYSDATE,
        'SYSDATE + 1');
    DBMS_OUTPUT.PUT_LINE('jobid: ' || jobid);
END;
```

jobid: 104

The job immediately executes procedure, `job_proc`, populating table, `jobrun`, with a row:

```
SELECT * FROM jobrun;
```

runtime

```
-----  
job_proc run at 2007-12-11 11:43:25  
(1 row)
```

6.3.5.8 WHAT

The **WHAT** procedure changes the stored procedure that the job will execute.

```
WHAT(<job> BINARY_INTEGER, <what> VARCHAR2)
```

Parameters

job

Identifier of the job for which the stored procedure is to be changed.

what

Name of the stored procedure to be executed.

Examples

Change the job to run the `list_emp` procedure:

```
BEGIN
  DBMS_JOB.WHAT(104,'list_emp;');
END;
```

6.3.6 DBMS_LOB

The **DBMS_LOB** package provides the capability to operate on large objects. The following table lists the supported functions and procedures:

| Function/Procedure | Return Type | Description |
|---|---------------|---|
| APPEND(dest_lob IN OUT, src_lob) | n/a | Appends one large object to another. |
| COMPARE(lob_1, lob_2 [, amount [, offset_1 [, offset_2]]]) | INTEGER | Compares two large objects. |
| CONVERTOBLOB(dest_lob IN OUT, src_clob, amount, dest_offset IN OUT, src_offset IN OUT, blob_csid, lang_context IN OUT, <warning> OUT) | n/a | Converts character data to binary. |
| CONVERTTOCLOB(dest_lob IN OUT, src_blob, amount, dest_offset IN OUT, src_offset IN OUT, blob_csid, lang_context IN OUT, warning OUT) | n/a | Converts binary data to character. |
| COPY(dest_lob IN OUT, src_lob, amount [, dest_offset [, src_offset]]) | n/a | Copies one large object to another. |
| ERASE(lob_loc IN OUT, amount IN OUT [, offset]) | n/a | Erase a large object. |
| GET_STORAGE_LIMIT(lob_loc) | INTEGER | Get the storage limit for large objects. |
| GETLENGTH(lob_loc) | INTEGER | Get the length of the large object. |
| INSTR(lob_loc, pattern [, offset [, nth]]) | INTEGER | Get the position of the nth occurrence of a pattern in the large object starting at <code>offset</code> |
| READ(lob_loc, amount IN OUT, offset, buffer OUT) | n/a | Read a large object. |
| SUBSTR(lob_loc [, amount [, offset]]) | RAW, VARCHAR2 | Get part of a large object. |

| | | |
|---|-----|--|
| TRIM(lob_loc IN OUT, newlen) | n/a | Trim a large object to the specified length. |
| WRITE(lob_loc IN OUT, amount, offset, buffer) | n/a | Write data to a large object. |
| WRITEAPPEND(lob_loc IN OUT, amount, buffer) | n/a | Write data from the buffer to the end of a large object. |

Advanced Server's implementation of `DBMS_LOB` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variables available in the package.

| Public Variables | Data Type | Value |
|-------------------------|-----------|------------|
| compress off | INTEGER | 0 |
| compress_on | INTEGER | 1 |
| deduplicate_off | INTEGER | 0 |
| deduplicate_on | INTEGER | 4 |
| default_csid | INTEGER | 0 |
| default_lang_ctx | INTEGER | 0 |
| encrypt_off | INTEGER | 0 |
| encrypt_on | INTEGER | 1 |
| file_READONLY | INTEGER | 0 |
| lobmaxsize | INTEGER | 1073741823 |
| lob_READONLY | INTEGER | 0 |
| lob_readwrite | INTEGER | 1 |
| no_warning | INTEGER | 0 |
| opt_compress | INTEGER | 1 |
| opt_deduplicate | INTEGER | 4 |
| opt_encrypt | INTEGER | 2 |
| warn_inconvertible_char | INTEGER | 1 |

In the following sections, lengths and offsets are measured in bytes if the large objects are `BLOBs`. Lengths and offsets are measured in characters if the large objects are `CLOBs`.

6.3.6.1 APPEND

The `APPEND` procedure provides the capability to append one large object to another. Both large objects must be of the same type.

```
APPEND(<dest_lob> IN OUT { BLOB | CLOB }, <src_lob> { BLOB | CLOB })
```

Parameters

`dest_lob`

Large object locator for the destination object. Must be the same data type as `src_lob`.

`src_lob`

Large object locator for the source object. Must be the same data type as `dest_lob`.

6.3.6.2 COMPARE

The `COMPARE` procedure performs an exact byte-by-byte comparison of two large objects for a given length at given offsets. The large objects being compared must be the same data type.

```
<status> INTEGER COMPARE(<lob_1> { BLOB | CLOB },
<lob_2> { BLOB | CLOB }
[, <amount> INTEGER [, <offset_1> INTEGER [, <offset_2> INTEGER ]]])
```

Parameters

`lob_1`

Large object locator of the first large object to be compared. Must be the same data type as `lob_2`.

`lob_2`

Large object locator of the second large object to be compared. Must be the same data type as `lob_1`.

`amount`

If the data type of the large objects is `BLOB`, then the comparison is made for `amount` bytes. If the data type of the large objects is `CLOB`, then the comparison is made for `amount` characters. The default is the maximum size of a large object.

`offset_1`

Position within the first large object to begin the comparison. The first byte/character is offset 1. The default is 1.

`offset_2`

Position within the second large object to begin the comparison. The first byte/character is offset 1. The default is 1.

`status`

Zero if both large objects are exactly the same for the specified length for the specified offsets. Non-zero, if the objects are not the same. `NULL` if `amount`, `offset_1`, or `offset_2` are less than zero.

6.3.6.3 CONVERTTOBLOB

The `CONVERTTOBLOB` procedure provides the capability to convert character data to binary.

```
CONVERTTOBLOB(<dest_lob> IN OUT BLOB, <src_clob> CLOB,
<amount> INTEGER, <dest_offset> IN OUT INTEGER,
<src_offset> IN OUT INTEGER, <blob_csid> NUMBER,
```

<lang_context> IN OUT INTEGER, <warning> OUT INTEGER)

Parameters

`dest_lob`

| `BLOB` large object locator to which the character data is to be converted.

`src_clob`

| `CLOB` large object locator of the character data to be converted.

`amount`

| Number of characters of `src_clob` to be converted.

`dest_offset IN`

| Position in bytes in the destination `BLOB` where writing of the source `CLOB` should begin. The first byte is offset 1.

`dest_offset OUT`

| Position in bytes in the destination `BLOB` after the write operation completes. The first byte is offset 1.

`src_offset IN`

| Position in characters in the source `CLOB` where conversion to the destination `BLOB` should begin. The first character is offset 1.

`src_offset OUT`

| Position in characters in the source `CLOB` after the conversion operation completes. The first character is offset 1.

`blob_csid`

| Character set ID of the converted, destination `BLOB`.

`lang_context IN`

| Language context for the conversion. The default value of 0 is typically used for this setting.

`lang_context OUT`

| Language context after the conversion completes.

`warning`

| 0 if the conversion was successful, 1 if an invertible character was encountered.

6.3.6.4 CONVERTTOCLOB

The `CONVERTTOCLOB` procedure provides the capability to convert binary data to character.

`CONVERTTOCLOB(<dest_lob> IN OUT CLOB, <src_blob> BLOB,`

```
<amount> INTEGER, <dest_offset> IN OUT INTEGER,
<src_offset> IN OUT INTEGER, <blob_csid> NUMBER,
<lang_context> IN OUT INTEGER, <warning> OUT INTEGER)
```

Parameters

`dest_lob`

| `CLOB` large object locator to which the binary data is to be converted.

`src_blob`

| `BLOB` large object locator of the binary data to be converted.

`amount`

| Number of bytes of `src_blob` to be converted.

`dest_offset IN`

| Position in characters in the destination `CLOB` where writing of the source `BLOB` should begin. The first character is offset 1.

`dest_offset OUT`

| Position in characters in the destination `CLOB` after the write operation completes. The first character is offset 1.

`src_offset IN`

| Position in bytes in the source `BLOB` where conversion to the destination `CLOB` should begin. The first byte is offset 1.

`src_offset OUT`

| Position in bytes in the source `BLOB` after the conversion operation completes. The first byte is offset 1.

`blob_csid`

| Character set ID of the converted, destination `CLOB`.

`lang_context IN`

| Language context for the conversion. The default value of 0 is typically used for this setting.

`lang_context OUT`

| Language context after the conversion completes.

`warning`

| 0 if the conversion was successful, 1 if an incomparable character was encountered.

6.3.6.5 COPY

The `COPY` procedure provides the capability to copy one large object to another. The source and destination

large objects must be the same data type.

```
COPY(<dest_lob> IN OUT { BLOB | CLOB }, <src_lob>
{ BLOB | CLOB },
<amount> INTEGER
[, <dest_offset> INTEGER [, <src_offset> INTEGER ]])
```

Parameters

`dest_lob`

Large object locator of the large object to which `src_lob` is to be copied. Must be the same data type as `src_lob`.

`src_lob`

Large object locator of the large object to be copied to `dest_lob`. Must be the same data type as `dest_lob`.

`amount`

Number of bytes/characters of `src_lob` to be copied.

`dest_offset`

Position in the destination large object where writing of the source large object should begin. The first position is offset 1. The default is 1.

`src_offset`

Position in the source large object where copying to the destination large object should begin. The first position is offset 1. The default is 1.

6.3.6.6 ERASE

The `ERASE` procedure provides the capability to erase a portion of a large object. To erase a large object means to replace the specified portion with zero-byte fillers for `BLOBS` or with spaces for `CLOBS`. The actual size of the large object is not altered.

```
ERASE(<lob_loc> IN OUT { BLOB | CLOB }, <amount> IN OUT INTEGER
[, <offset> INTEGER ])
```

Parameters

`lob_loc`

Large object locator of the large object to be erased.

`amount IN`

Number of bytes/characters to be erased.

`amount OUT`

Number of bytes/characters actually erased. This value can be smaller than the input value if the end of the large object is reached before `amount` bytes/characters have been erased.

offset

Position in the large object where erasing is to begin. The first byte/character is position 1. The default is 1.

6.3.6.7 GET_STORAGE_LIMIT

The `GET_STORAGE_LIMIT` function returns the limit on the largest allowable large object.

```
<size> INTEGER GET_STORAGE_LIMIT(<lob_loc> BLOB)
```

```
<size> INTEGER GET_STORAGE_LIMIT(<lob_loc> CLOB)
```

Parameters

size

Maximum allowable size of a large object in this database.

lob_loc

This parameter is ignored, but is included for compatibility.

6.3.6.8 GETLENGTH

The `GETLENGTH` function returns the length of a large object.

```
<amount> INTEGER GETLENGTH(<lob_loc> BLOB)
```

```
<amount> INTEGER GETLENGTH(<lob_loc> CLOB)
```

Parameters

lob_loc

Large object locator of the large object whose length is to be obtained.

amount

Length of the large object in bytes for `BLOBS` or characters for `CLOBS`.

6.3.6.9 INSTR

The `INSTR` function returns the location of the nth occurrence of a given pattern within a large object.

```
<position> INTEGER INSTR(<lob_loc> { BLOB | CLOB },
<pattern> { RAW | VARCHAR2 } [, <offset> INTEGER [, <nth> INTEGER ]])
```

Parameters

`lob_loc`

Large object locator of the large object in which to search for pattern.

`pattern`

Pattern of bytes or characters to match against the large object, `lob. pattern` must be `RAW` if `lob_loc` is a `BLOB`. `pattern` must be `VARCHAR2` if `lob_loc` is a `CLOB`.

`offset`

Position within `lob_loc` to start search for `pattern`. The first byte/character is position 1. The default is 1.

`nth`

Search for `pattern`, `nth` number of times starting at the position given by `offset`. The default is 1.

`position`

Position within the large object where `pattern` appears the nth time specified by `nth` starting from the position given by `offset`.

6.3.6.10 READ

The `READ` procedure provides the capability to read a portion of a large object into a buffer.

```
READ(<lob_loc> { BLOB | CLOB }, <amount> IN OUT BINARY_INTEGER,
<offset> INTEGER, <buffer> OUT { RAW | VARCHAR2 })
```

Parameters

`lob_loc`

Large object locator of the large object to be read.

`amount IN`

Number of bytes/characters to read.

`amount OUT`

Number of bytes/characters actually read. If there is no more data to be read, then `amount` returns 0 and a `DATA_NOT_FOUND` exception is thrown.

`offset`

Position to begin reading. The first byte/character is position 1.

buffer

Variable to receive the large object. If `lob_loc` is a `BLOB`, then `buffer` must be `RAW`. If `lob_loc` is a `CLOB`, then `buffer` must be `VARCHAR2`.

6.3.6.11 SUBSTR

The `SUBSTR` function provides the capability to return a portion of a large object.

```
<data> { RAW | VARCHAR2 } SUBSTR(<lob_loc> { BLOB | CLOB }
[, <amount> INTEGER [, <offset> INTEGER ]])
```

Parameters**lob_loc**

Large object locator of the large object to be read.

amount

Number of bytes/characters to be returned. Default is 32,767.

offset

Position within the large object to begin returning data. The first byte/character is position 1. The default is 1.

data

Returned portion of the large object to be read. If `lob_loc` is a `BLOB`, the return data type is `RAW`. If `lob_loc` is a `CLOB`, the return data type is `VARCHAR2`.

6.3.6.12 TRIM

The `TRIM` procedure provides the capability to truncate a large object to the specified length.

```
TRIM(<lob_loc> IN OUT { BLOB | CLOB }, <newlen> INTEGER)
```

Parameters**lob_loc**

Large object locator of the large object to be trimmed.

newlen

Number of bytes/characters to which the large object is to be trimmed.

6.3.6.13 WRITE

The `WRITE` procedure provides the capability to write data into a large object. Any existing data in the large object at the specified offset for the given length is overwritten by data given in the buffer.

```
WRITE(<lob_loc> IN OUT { BLOB | CLOB },
<amount> BINARY_INTEGER,
<offset> INTEGER, <buffer> { RAW | VARCHAR2 })
```

Parameters

`lob_loc`

Large object locator of the large object to be written.

`amount`

The number of bytes/characters in `buffer` to be written to the large object.

`offset`

The offset in bytes/characters from the beginning of the large object (origin is 1) for the write operation to begin.

`buffer`

Contains data to be written to the large object. If `lob_loc` is a `BLOB`, then `buffer` must be `RAW`. If `lob_loc` is a `CLOB`, then `buffer` must be `VARCHAR2`.

6.3.6.14 WRITEAPPEND

The `WRITEAPPEND` procedure provides the capability to add data to the end of a large object.

```
WRITEAPPEND(<lob_loc> IN OUT { BLOB | CLOB },
<amount> BINARY_INTEGER, <buffer> { RAW | VARCHAR2 })
```

Parameters

`lob_loc`

Large object locator of the large object to which data is to be appended.

`amount`

Number of bytes/characters from `buffer` to be appended to the large object.

`buffer`

Data to be appended to the large object. If `lob_loc` is a `BLOB`, then `buffer` must be `RAW`. If `lob_loc` is a `CLOB`, then `buffer` must be `VARCHAR2`.

6.3.7 DBMS_LOCK

Advanced Server provides support for the `DBMS_LOCK.SLEEP` procedure.

| Function/Procedure | Return Type | Description |
|-----------------------------|-------------|---|
| <code>SLEEP(seconds)</code> | n/a | Suspends a session for the specified number of <code>seconds</code> . |

Advanced Server's implementation of `DBMS_LOCK` is a partial implementation when compared to Oracle's version. Only `DBMS_LOCK.SLEEP` is supported.

SLEEP

The `SLEEP` procedure suspends the current session for the specified number of seconds.

`SLEEP(<seconds> NUMBER)`

Parameters

`seconds`

`seconds` specifies the number of seconds for which you wish to suspend the session. `seconds` can be a fractional value; for example, enter `1.75` to specify one and three-fourths of a second.

6.3.8 DBMS_MVIEW

Use procedures in the `DBMS_MVIEW` package to manage and refresh materialized views and their dependencies. Advanced Server provides support for the following `DBMS_MVIEW` procedures:

| Procedure | Return Type | Description |
|--|-------------|---|
| <code>GET_MV_DEPENDENCIES(list VARCHAR2, deplist VARCHAR2);</code> | n/a | The <code>GET_MV_DEPENDENCIES</code> procedure returns a list of dependencies for a specified view. |
| <code>REFRESH(list VARCHAR2, method VARCHAR2, rollback seq VARCHAR2, push deferred rpc BOOLEAN, refresh after errors BOOLEAN, purge option NUMBER, parallelism NUMBER, heap size NUMBER, atomic_refresh BOOLEAN);</code> | n/a | This variation of the <code>REFRESH</code> procedure refreshes all views named in a comma-separated list of view names. |
| <code>REFRESH(tab dbms_utility.uncl_array, method VARCHAR2, rollback seq VARCHAR2, push deferred rpc BOOLEAN, refresh after errors BOOLEAN, purge option NUMBER, parallelism NUMBER, heap size NUMBER, atomic_refresh BOOLEAN, nested BOOLEAN);</code> | n/a | This variation of the <code>REFRESH</code> procedure refreshes all views named in a table of <code>dbms_utility.uncl_array</code> values. |
| <code>REFRESH_ALL_MVIEWS(number of failures BINARY_INTEGER, method VARCHAR2, rollback seq VARCHAR2, refresh_after_errors BOOLEAN, atomic_refresh BOOLEAN);</code> | n/a | The <code>REFRESH_ALL_MVIEWS</code> procedure refreshes all materialized views. |

| | | |
|---|-----|---|
| REFRESH_DEPENDENT(number_of_failures BINARY_INTEGER, list VARCHAR2, method VARCHAR2, rollback_seq VARCHAR2, refresh_after_errors BOOLEAN, atomic_refresh BOOLEAN, nested BOOLEAN); | n/a | This variation of the REFRESH_DEPENDENT procedure refreshes all views that are dependent on the views listed in a comma-separated list. |
| REFRESH_DEPENDENT(number_of_failures BINARY_INTEGER, tab dbms_utility.uncl_array, method VARCHAR2, rollback_seq VARCHAR2, refresh_after_errors BOOLEAN, atomic_refresh BOOLEAN, nested BOOLEAN); | n/a | This variation of the REFRESH_DEPENDENT procedure refreshes all views that are dependent on the views listed in a table of dbms_utility.uncl_array values. |

Advanced Server's implementation of **DBMS_MVIEW** is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

6.3.8.1 GET_MV_DEPENDENCIES

When given the name of a materialized view, **GET_MV_DEPENDENCIES** returns a list of items that depend on the specified view. The signature is:

```
GET_MV_DEPENDENCIES(  
<list> IN VARCHAR2,  
<deplist> OUT VARCHAR2);
```

Parameters

list

| **list** specifies the name of a materialized view, or a comma-separated list of materialized view names.

deplist

| **deplist** is a comma-separated list of schema-qualified dependencies. **deplist** is a **VARCHAR2** value.

Examples

The following example:

```
DECLARE  
    deplist VARCHAR2(1000);  
BEGIN  
    DBMS_MVIEW.GET_MV_DEPENDENCIES('public.emp_view', deplist);  
    DBMS_OUTPUT.PUT_LINE(deplist);  
END;
```

Displays a list of the dependencies on a materialized view named **public.emp_view**.

6.3.8.2 REFRESH

Use the `REFRESH` procedure to refresh all views specified in either a comma-separated list of view names, or a table of `DBMS.Utility.UNCL_ARRAY` values. The procedure has two signatures; use the first form when specifying a comma-separated list of view names:

```
REFRESH(
<list> IN VARCHAR2,
<method> IN VARCHAR2 DEFAULT NULL,
<rollback_seg> IN VARCHAR2 DEFAULT NULL,
<push_deferred_rpc> IN BOOLEAN DEFAULT TRUE,
<refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
<purge_option> IN NUMBER DEFAULT 1,
<parallelism> IN NUMBER DEFAULT 0,
<heap_size> IN NUMBER DEFAULT 0,
<atomic_refresh> IN BOOLEAN DEFAULT TRUE,
<nested> IN BOOLEAN DEFAULT FALSE);
```

Use the second form to specify view names in a table of `DBMS.Utility.UNCL_ARRAY` values:

```
REFRESH(
<tab> IN OUT DBMS.Utility.UNCL_ARRAY,
<method> IN VARCHAR2 DEFAULT NULL,
<rollback_seg> IN VARCHAR2 DEFAULT NULL,
<push_deferred_rpc> IN BOOLEAN DEFAULT TRUE,
<refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
<purge_option> IN NUMBER DEFAULT 1,
<parallelism> IN NUMBER DEFAULT 0,
<heap_size> IN NUMBER DEFAULT 0,
<atomic_refresh> IN BOOLEAN DEFAULT TRUE,
<nested> IN BOOLEAN DEFAULT FALSE);
```

Parameters

list

`list` is a `VARCHAR2` value that specifies the name of a materialized view, or a comma-separated list of materialized view names. The names may be schema-qualified.

tab

`tab` is a table of `DBMS.Utility.UNCL_ARRAY` values that specify the name (or names) of a materialized view.

method

`method` is a `VARCHAR2` value that specifies the refresh method that will be applied to the specified view (or views). The only supported method is `C`; this performs a complete refresh of the view.

rollback_seg

`rollback_seg` is accepted for compatibility and ignored. The default is `NULL`.

push_deferred_rpc

`push_deferred_rpc` is accepted for compatibility and ignored. The default is `TRUE`.

refresh_after_errors

`refresh_after_errors` is accepted for compatibility and ignored. The default is `FALSE`.

purge_option

purge_option is accepted for compatibility and ignored. The default is **1**.

parallelism

parallelism is accepted for compatibility and ignored. The default is **0**.

heap_size IN NUMBER DEFAULT 0,

heap_size is accepted for compatibility and ignored. The default is **0**.

atomic_refresh

atomic_refresh is accepted for compatibility and ignored. The default is **TRUE**.

nested

nested is accepted for compatibility and ignored. The default is **FALSE**.

Examples

The following example uses **DBMS_MVIEW.REFRESH** to perform a **COMPLETE** refresh on the **public.emp_view** materialized view:

```
EXEC DBMS_MVIEW.REFRESH(list => 'public.emp_view', method => 'C');
```

6.3.8.3 REFRESH_ALL_MVIEWS

Use the **REFRESH_ALL_MVIEWS** procedure to refresh any materialized views that have not been refreshed since the table or view on which the view depends has been modified. The signature is:

```
REFRESH_ALL_MVIEWS(
  <number_of_failures> OUT BINARY_INTEGER,
  <method> IN VARCHAR2 DEFAULT NULL,
  <rollback_seg> IN VARCHAR2 DEFAULT NULL,
  <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
  <atomic_refresh> IN BOOLEAN DEFAULT TRUE);
```

Parameters**number_of_failures**

number_of_failures is a **BINARY_INTEGER** that specifies the number of failures that occurred during the refresh operation.

method

method is a **VARCHAR2** value that specifies the refresh method that will be applied to the specified view (or views). The only supported method is **C**; this performs a complete refresh of the view.

rollback_seg

rollback_seg is accepted for compatibility and ignored. The default is **NULL**.

refresh_after_errors

`refresh_after_errors` is accepted for compatibility and ignored. The default is `FALSE`.

atomic_refresh

`atomic_refresh` is accepted for compatibility and ignored. The default is `TRUE`.

Examples

The following example performs a `COMPLETE` refresh on all materialized views:

```
DECLARE
  errors INTEGER;
BEGIN
  DBMS_MVIEW.REFRESH_ALL_MVIEWS(errors, method => 'C');
END;
```

Upon completion, `errors` contains the number of failures.

6.3.8.4 REFRESH_DEPENDENT

Use the `REFRESH_DEPENDENT` procedure to refresh all material views that are dependent on the views specified in the call to the procedure. You can specify a comma-separated list or provide the view names in a table of `DBMS.Utility.UNCL_ARRAY` values.

Use the first form of the procedure to refresh all material views that are dependent on the views specified in a comma-separated list:

```
REFRESH_DEPENDENT(
  <number_of_failures> OUT BINARY_INTEGER,
  <list> IN VARCHAR2,
  <method> IN VARCHAR2 DEFAULT NULL,
  <rollback_seg> IN VARCHAR2 DEFAULT NULL,
  <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
  <atomic_refresh> IN BOOLEAN DEFAULT TRUE,
  <nested> IN BOOLEAN DEFAULT FALSE);
```

Use the second form of the procedure to refresh all material views that are dependent on the views specified in a table of `DBMS.Utility.UNCL_ARRAY` values:

```
REFRESH_DEPENDENT(
  <number_of_failures> OUT BINARY_INTEGER,
  <tab> IN DBMS.Utility.UNCL_ARRAY,
  <method> IN VARCHAR2 DEFAULT NULL,
  <rollback_seg> IN VARCHAR2 DEFAULT NULL,
  <refresh_after_errors> IN BOOLEAN DEFAULT FALSE,
  <atomic_refresh> IN BOOLEAN DEFAULT TRUE,
  <nested> IN BOOLEAN DEFAULT FALSE);
```

Parameters

number_of_failures

`number_of_failures` is a `BINARY_INTEGER` that contains the number of failures that occurred during the refresh operation.

list

`list` is a `VARCHAR2` value that specifies the name of a materialized view, or a comma-separated list of materialized view names. The names may be schema-qualified.

tab

`tab` is a table of `DBMS.Utility.UNCL_ARRAY` values that specify the name (or names) of a materialized view.

method

`method` is a `VARCHAR2` value that specifies the refresh method that will be applied to the specified view (or views). The only supported method is `C`; this performs a complete refresh of the view.

rollback_seg

`rollback_seg` is accepted for compatibility and ignored. The default is `NULL`.

refresh_after_errors

`refresh_after_errors` is accepted for compatibility and ignored. The default is `FALSE`.

atomic_refresh

`atomic_refresh` is accepted for compatibility and ignored. The default is `TRUE`.

nested

`nested` is accepted for compatibility and ignored. The default is `FALSE`.

Examples

The following example performs a `COMPLETE` refresh on all materialized views dependent on a materialized view named `emp_view` that resides in the `public` schema:

```
DECLARE
  errors INTEGER;
BEGIN
  DBMS_MVIEW.REFRESH_DEPENDENT(errors, list => 'public.emp_view', method =>
'C');
END;
```

Upon completion, `errors` contains the number of failures.

6.3.9 DBMS_OUTPUT

The `DBMS_OUTPUT` package provides the capability to send messages (lines of text) to a message buffer, or get messages from the message buffer. A message buffer is local to a single session. Use the `DBMS_PIPE` package to send messages between sessions.

The procedures and functions available in the `DBMS_OUTPUT` package are listed in the following table.

| Function/Procedure | Return Type | Description |
|---------------------------------------|-------------|--|
| DISABLE | n/a | Disable the capability to send and receive messages. |
| ENABLE(buffer_size) | n/a | Enable the capability to send and receive messages. |
| GET_LINE(line OUT, status OUT) | n/a | Get a line from the message buffer. |
| GET_LINES(lines OUT, numlines IN OUT) | n/a | Get multiple lines from the message buffer. |
| NEW_LINE | n/a | Puts an end-of-line character sequence. |
| PUT(item) | n/a | Puts a partial line without an end-of-line character sequence. |
| PUT_LINE(item) | n/a | Puts a complete line with an end-of-line character sequence. |
| SERVEROUTPUT(stdout) | n/a | Direct messages from PUT, PUT_LINE, or NEW_LINE to either standard output or the message buffer. |

The following table lists the public variables available in the DBMS_OUTPUT package.

| Public Variables | Data Type | Value | Description |
|------------------|-----------|-------|--------------------|
| chararr | TABLE | | For message lines. |

CHARARR

The CHARARR is for storing multiple message lines.

```
TYPE chararr IS TABLE OF VARCHAR2(32767) INDEX BY BINARY_INTEGER;
```

DISABLE

The DISABLE procedure clears out the message buffer. Any messages in the buffer at the time the DISABLE procedure is executed will no longer be accessible. Any messages subsequently sent with the PUT, PUT_LINE, or NEW_LINE procedures are discarded. No error is returned to the sender when the PUT, PUT_LINE, or NEW_LINE procedures are executed and messages have been disabled.

Use the ENABLE procedure or SERVEROUTPUT(TRUE) procedure to re-enable the sending and receiving of messages.

```
DISABLE
```

Examples

This anonymous block disables the sending and receiving messages in the current session.

```
BEGIN
  DBMS_OUTPUT.DISABLE;
END;
```

ENABLE

The ENABLE procedure enables the capability to send messages to the message buffer or retrieve messages from the message buffer. Running SERVEROUTPUT(TRUE) also implicitly performs the ENABLE procedure.

The destination of a message sent with `PUT`, `PUT_LINE`, or `NEW_LINE` depends upon the state of `SERVEROUTPUT`.

- If the last state of `SERVEROUTPUT` is `TRUE`, the message goes to standard output of the command line.
- If the last state of `SERVEROUTPUT` is `FALSE`, the message goes to the message buffer.

```
ENABLE [ (<buffer_size> INTEGER) ]
```

Parameter

`buffer_size`

Maximum length of the message buffer in bytes. If a `buffer_size` of less than 2000 is specified, the buffer size is set to 2000.

Examples

The following anonymous block enables messages. Setting `SERVEROUTPUT(TRUE)` forces them to standard output.

```
BEGIN
  DBMS_OUTPUT.ENABLE;
  DBMS_OUTPUT.SERVEROUTPUT(TRUE);
  DBMS_OUTPUT.PUT_LINE('Messages enabled');
END;
```

Messages enabled

The same effect could have been achieved by simply using `SERVEROUTPUT(TRUE)`.

```
BEGIN
  DBMS_OUTPUT.SERVEROUTPUT(TRUE);
  DBMS_OUTPUT.PUT_LINE('Messages enabled');
END;
```

Messages enabled

The following anonymous block enables messages, but setting `SERVEROUTPUT(FALSE)` directs messages to the message buffer.

```
BEGIN
  DBMS_OUTPUT.ENABLE;
  DBMS_OUTPUT.SERVEROUTPUT(FALSE);
  DBMS_OUTPUT.PUT_LINE('Message sent to buffer');
END;
```

GET_LINE

The `GET_LINE` procedure provides the capability to retrieve a line of text from the message buffer. Only text that has been terminated by an end-of-line character sequence is retrieved – that is complete lines generated using `PUT_LINE`, or by a series of `PUT` calls followed by a `NEW_LINE` call.

```
GET_LINE(<line> OUT VARCHAR2, <status> OUT INTEGER)
```

Parameters

line

Variable receiving the line of text from the message buffer.

status

0 if a line was returned from the message buffer, 1 if there was no line to return.

Examples

The following anonymous block writes the `emp` table out to the message buffer as a comma-delimited string for each row.

```
EXEC DBMS_OUTPUT.SERVEROUTPUT(FALSE);

DECLARE
    v_emprec      VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    DBMS_OUTPUT.ENABLE;
    FOR i IN emp_cur LOOP
        v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
                    NVL(LTRIM(TO_CHAR(i.mgr,'9999')),'') || ',' || i.hiredate ||
                    ',' || i.sal || ',' ||
                    NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),'') || ',' || i.deptno;
        DBMS_OUTPUT.PUT_LINE(v_emprec);
    END LOOP;
END;
```

The following anonymous block reads the message buffer and inserts the messages written by the prior example into a table named `messages`. The rows in `messages` are then displayed.

```
CREATE TABLE messages (
    status      INTEGER,
    msg         VARCHAR2(100)
);

DECLARE
    v_line      VARCHAR2(100);
    v_status    INTEGER := 0;
BEGIN
    DBMS_OUTPUT.GET_LINE(v_line,v_status);
    WHILE v_status = 0 LOOP
        INSERT INTO messages VALUES(v_status, v_line);
        DBMS_OUTPUT.GET_LINE(v_line,v_status);
    END LOOP;
END;

SELECT msg FROM messages;
```

msg

```
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
```

```

7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
(14 rows)

```

GET_LINES

The `GET_LINES` procedure provides the capability to retrieve one or more lines of text from the message buffer into a collection. Only text that has been terminated by an end-of-line character sequence is retrieved – that is complete lines generated using `PUT_LINE`, or by a series of `PUT` calls followed by a `NEW_LINE` call.

```
GET_LINES(<lines> OUT CHARARR, <numlines> IN OUT INTEGER)
```

Parameters

`lines`

Table receiving the lines of text from the message buffer. See `CHARARR` for a description of `lines`.

`numlines IN`

Number of lines to be retrieved from the message buffer.

`numlines OUT`

Actual number of lines retrieved from the message buffer. If the output value of `numlines` is less than the input value, then there are no more lines left in the message buffer.

Examples

The following example uses the `GET_LINES` procedure to store all rows from the `emp` table that were placed on the message buffer, into an array.

```

EXEC DBMS_OUTPUT.SERVEROUTPUT(FALSE);

DECLARE
    v_emprec      VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    DBMS_OUTPUT.ENABLE;
    FOR i IN emp_cur LOOP
        v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
                    NVL(LTRIM(TO_CHAR(i.mgr,'9999')),"") || ',' || i.hiredate ||
                    ',' || i.sal || ',' ||
                    NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),"") || ',' || i.deptno;
        DBMS_OUTPUT.PUT_LINE(v_emprec);
    END LOOP;
END;

```

DECLARE

```

v_lines      DBMS_OUTPUT.CHARARR;
v_numlines   INTEGER := 14;
v_status     INTEGER := 0;
BEGIN
  DBMS_OUTPUT.GET_LINES(v_lines,v_numlines);
  FOR i IN 1..v_numlines LOOP
    INSERT INTO messages VALUES(v_numlines, v_lines(i));
  END LOOP;
END;

```

SELECT msg FROM messages;

| msg |
|---|
| 7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20 |
| 7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30 |
| 7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30 |
| 7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20 |
| 7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30 |
| 7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30 |
| 7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10 |
| 7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20 |
| 7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10 |
| 7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30 |
| 7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20 |
| 7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30 |
| 7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20 |
| 7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10 |

(14 rows)

NEW_LINE

The `NEW_LINE` procedure writes an end-of-line character sequence in the message buffer.

NEW_LINE

Parameter

The `NEW_LINE` procedure expects no parameters.

PUT

The `PUT` procedure writes a string to the message buffer. No end-of-line character sequence is written at the end of the string. Use the `NEW_LINE` procedure to add an end-of-line character sequence.

PUT(<item> VARCHAR2)

Parameter

item

Text written to the message buffer.

Examples

The following example uses the `PUT` procedure to display a comma-delimited list of employees from the `emp` table.

```

DECLARE
  CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
  FOR i IN emp_cur LOOP
    DBMS_OUTPUT.PUT(i.empno);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.ename);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.job);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.mgr);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.hiredate);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.sal);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.comm);
    DBMS_OUTPUT.PUT(',');
    DBMS_OUTPUT.PUT(i.deptno);
    DBMS_OUTPUT.NEW_LINE;
  END LOOP;
END;

```

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
 7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
 7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
 7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
 7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
 7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
 7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
 7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
 7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
 7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
 7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
 7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
 7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
 7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10

PUT_LINE

The `PUT_LINE` procedure writes a single line to the message buffer including an end-of-line character sequence.

```
PUT_LINE(<item> VARCHAR2)
```

Parameter

`item`

Text to be written to the message buffer.

Examples

The following example uses the `PUT_LINE` procedure to display a comma-delimited list of employees from the `emp` table.

```

DECLARE
    v_emprec      VARCHAR2(120);
    CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
    FOR i IN emp_cur LOOP
        v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
                    NVL(LTRIM(TO_CHAR(i.mgr,'9999')),"") || ',' || i.hiredate ||
                    ',' || i.sal || ',' ||
                    NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),"") || ',' || i.deptno;
        DBMS_OUTPUT.PUT_LINE(v_emprec);
    END LOOP;
END;

```

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
 7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
 7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
 7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
 7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
 7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
 7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
 7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
 7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
 7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
 7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
 7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
 7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
 7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10

SERVEROUTPUT

The `SERVEROUTPUT` procedure provides the capability to direct messages to standard output of the command line or to the message buffer. Setting `SERVEROUTPUT(TRUE)` also performs an implicit execution of `ENABLE`.

The default setting of `SERVEROUTPUT` is implementation dependent. For example, in Oracle SQL*Plus, `SERVEROUTPUT(FALSE)` is the default. In PSQL, `SERVEROUTPUT(TRUE)` is the default. Also note that in Oracle SQL*Plus, this setting is controlled using the SQL*Plus `SET` command, not by a stored procedure as implemented in Advanced Server.

```
SERVEROUTPUT(<stdout> BOOLEAN)
```

To get an Oracle-style display output, you can set the `dbms_output.serveroutput` to `FALSE` in the `postgresql.conf` file; this disables the message output. The default is `TRUE`, which enables the message output.

Parameter

`stdout`

Set to `TRUE` if subsequent `PUT`, `PUT_LINE`, or `NEW_LINE` commands are to send text directly to standard output of the command line. Set to `FALSE` if text is to be sent to the message buffer.

Examples

The following anonymous block sends the first message to the command line and the second message to the message buffer.

```
BEGIN
  DBMS_OUTPUT.SERVEROUTPUT(TRUE);
  DBMS_OUTPUT.PUT_LINE('This message goes to the command line');
  DBMS_OUTPUT.SERVEROUTPUT(FALSE);
  DBMS_OUTPUT.PUT_LINE('This message goes to the message buffer');
END;
```

This message goes to the command line

If within the same session, the following anonymous block is executed, the message stored in the message buffer from the prior example is flushed and displayed on the command line as well as the new message.

```
BEGIN
  DBMS_OUTPUT.SERVEROUTPUT(TRUE);
  DBMS_OUTPUT.PUT_LINE('Flush messages from the buffer');
END;
```

This message goes to the message buffer
Flush messages from the buffer

6.3.10 DBMS_PIPE

The `DBMS_PIPE` package provides the capability to send messages through a pipe within or between sessions connected to the same database cluster.

The procedures and functions available in the `DBMS_PIPE` package are listed in the following table:

| Function/Procedure | Return Type | Description |
|---|-----------------------|---|
| <code>CREATE PIPE(pipename [, maxpipesize] [, private])</code> | <code>INTEGER</code> | Explicitly create a private pipe if <code>private</code> is “true” (the default) or a public pipe if <code>private</code> is “false”. |
| <code>NEXT_ITEM_TYPE</code> | <code>INTEGER</code> | Determine the data type of the next item in a received message. |
| <code>PACK_MESSAGE(item)</code> | n/a | Place <code>item</code> in the session’s local message buffer. |
| <code>PURGE(pipename)</code> | n/a | Remove unreceived messages from the specified pipe. |
| <code>RECEIVE_MESSAGE(pipename [, timeout])</code> | <code>INTEGER</code> | Get a message from a specified pipe. |
| <code>REMOVE_PIPE(pipename)</code> | <code>INTEGER</code> | Delete an explicitly created pipe. |
| <code>RESET_BUFFER</code> | n/a | Reset the local message buffer. |
| <code>SEND_MESSAGE(pipename [, timeout] [, maxpipesize])</code> | <code>INTEGER</code> | Send a message on a pipe. |
| <code>UNIQUE_SESSION_NAME</code> | <code>VARCHAR2</code> | Obtain a unique session name. |
| <code>UNPACK_MESSAGE(item OUT)</code> | n/a | Retrieve the next data item from a message into a type-compatible variable, <code>item</code> . |

Pipes are categorized as implicit or explicit. An *implicit pipe* is created if a reference is made to a pipe name that was not previously created by the `CREATE_PIPE` function. For example, if the `SEND_MESSAGE` function is

executed using a non-existent pipe name, a new implicit pipe is created with that name. An *explicit pipe* is created using the `CREATE_PIPE` function whereby the first parameter specifies the pipe name for the new pipe.

Pipes are also categorized as private or public. A *private pipe* can only be accessed by the user who created the pipe. Even a superuser cannot access a private pipe that was created by another user. A *public pipe* can be accessed by any user who has access to the `DBMS_PIPE` package.

A public pipe can only be created by using the `CREATE_PIPE` function with the third parameter set to `FALSE`. The `CREATE_PIPE` function can be used to create a private pipe by setting the third parameter to `TRUE` or by omitting the third parameter. All implicit pipes are private.

The individual data items or “lines” of a message are first built-in a *local message buffer*, unique to the current session. The `PACK_MESSAGE` procedure builds the message in the session’s local message buffer. The `SEND_MESSAGE` function is then used to send the message through the pipe.

Receipt of a message involves the reverse operation. The `RECEIVE_MESSAGE` function is used to get a message from the specified pipe. The message is written to the session’s local message buffer. The `UNPACK_MESSAGE` procedure is then used to transfer the message data items from the message buffer to program variables. If a pipe contains multiple messages, `RECEIVE_MESSAGE` gets the messages in *FIFO* (first-in-first-out) order.

Each session maintains separate message buffers for messages created with the `PACK_MESSAGE` procedure and messages retrieved by the `RECEIVE_MESSAGE` function. Thus messages can be both built and received in the same session. However, if consecutive `RECEIVE_MESSAGE` calls are made, only the message from the last `RECEIVE_MESSAGE` call will be preserved in the local message buffer.

6.3.10.1 CREATE_PIPE

The `CREATE_PIPE` function creates an explicit public pipe or an explicit private pipe with a specified name.

```
<status> INTEGER CREATE_PIPE(<pipename> VARCHAR2
[, <maxpipesize> INTEGER ] [, <private> BOOLEAN ])
```

Parameters

`pipename`

Name of the pipe.

`maxpipesize`

Maximum capacity of the pipe in bytes. Default is 8192 bytes.

`private`

Create a public pipe if set to `FALSE`. Create a private pipe if set to `TRUE`. This is the default.

`status`

Status code returned by the operation. 0 indicates successful creation.

Examples

The following example creates a private pipe named `messages`:

```
DECLARE
```

```

v_status      INTEGER;
BEGIN
  v_status := DBMS_PIPE.CREATE_PIPE('messages');
  DBMS_OUTPUT.PUT_LINE('CREATE_PIPE status: ' || v_status);
END;
CREATE_PIPE status: 0

```

The following example creates a public pipe named `mailbox`:

```

DECLARE
  v_status      INTEGER;
BEGIN
  v_status := DBMS_PIPE.CREATE_PIPE('mailbox',8192, FALSE);
  DBMS_OUTPUT.PUT_LINE('CREATE_PIPE status: ' || v_status);
END;
CREATE_PIPE status: 0

```

6.3.10.2 NEXT_ITEM_TYPE

The `NEXT_ITEM_TYPE` function returns an integer code identifying the data type of the next data item in a message that has been retrieved into the session's local message buffer. As each item is moved off of the local message buffer with the `UNPACK_MESSAGE` procedure, the `NEXT_ITEM_TYPE` function will return the data type code for the next available item. A code of 0 is returned when there are no more items left in the message.

`<typecode>` INTEGER `NEXT_ITEM_TYPE`

Parameters

`typecode`

Code identifying the data type of the next data item as shown in the following table.

| Type Code | Data Type |
|-----------|--------------------|
| 0 | No more data items |
| 9 | NUMBER |
| 11 | VARCHAR2 |
| 13 | DATE |
| 23 | RAW |

Note

The type codes list in the table are not compatible with Oracle databases. Oracle assigns a different numbering sequence to the data types.

Examples

The following example shows a pipe packed with a `NUMBER` item, a `VARCHAR2` item, a `DATE` item, and a `RAW` item. A second anonymous block then uses the `NEXT_ITEM_TYPE` function to display the type code of each item.

`DECLARE`

```

v_number      NUMBER := 123;
v_varchar     VARCHAR2(20) := 'Character data';
v_date        DATE := SYSDATE;
v_raw         RAW(4) := '21222324';
v_status      INTEGER;

BEGIN
  DBMS_PIPE.PACK_MESSAGE(v_number);
  DBMS_PIPE.PACK_MESSAGE(v_varchar);
  DBMS_PIPE.PACK_MESSAGE(v_date);
  DBMS_PIPE.PACK_MESSAGE(v_raw);
  v_status := DBMS_PIPE.SEND_MESSAGE('datatypes');
  DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
EXCEPTION
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
    DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

```

SEND_MESSAGE status: 0

```

DECLARE
  v_number      NUMBER;
  v_varchar     VARCHAR2(20);
  v_date        DATE;
  v_timestamp   TIMESTAMP;
  v_raw         RAW(4);
  v_status      INTEGER;

BEGIN
  v_status := DBMS_PIPE.RECEIVE_MESSAGE('datatypes');
  DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_number);
  DBMS_OUTPUT.PUT_LINE('NUMBER Item : ' || v_number);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_varchar);
  DBMS_OUTPUT.PUT_LINE('VARCHAR2 Item : ' || v_varchar);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_date);
  DBMS_OUTPUT.PUT_LINE('DATE Item : ' || v_date);
  DBMS_OUTPUT.PUT_LINE('-----');

  v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
  DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_raw);
  DBMS_OUTPUT.PUT_LINE('RAW Item : ' || v_raw);
  DBMS_OUTPUT.PUT_LINE('-----');

```

```

v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
DBMS_OUTPUT.PUT_LINE('NEXT_ITEM_TYPE: ' || v_status);
DBMS_OUTPUT.PUT_LINE('-----');
EXCEPTION
WHEN OTHERS THEN
DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

```

RECEIVE_MESSAGE status: 0

NEXT_ITEM_TYPE: 9

NUMBER Item : 123

NEXT_ITEM_TYPE: 11

VARCHAR2 Item : Character data

NEXT_ITEM_TYPE: 13

DATE Item : 02-OCT-07 11:11:43

NEXT_ITEM_TYPE: 23

RAW Item : 21222324

NEXT_ITEM_TYPE: 0

6.3.10.3 PACK_MESSAGE

The `PACK_MESSAGE` procedure places an item of data in the session's local message buffer. `PACK_MESSAGE` must be executed at least once before issuing a `SEND_MESSAGE` call.

`PACK_MESSAGE(<item> { DATE | NUMBER | VARCHAR2 | RAW })`

Use the `UNPACK_MESSAGE` procedure to obtain data items once the message is retrieved using a `RECEIVE_MESSAGE` call.

Parameters

`item`

An expression evaluating to any of the acceptable parameter data types. The value is added to the session's local message buffer.

6.3.10.4 PURGE

The `PURGE` procedure removes the unreceived messages from a specified implicit pipe.

```
PURGE(<pipename> VARCHAR2)
```

Use the [REMOVE_PIPE](#) function to delete an explicit pipe.

Parameters

pipename

Name of the pipe.

Examples

Two messages are sent on a pipe:

```
DECLARE
    v_status      INTEGER;
BEGIN
    DBMS_PIPE.PACK_MESSAGE('Message #1');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);

    DBMS_PIPE.PACK_MESSAGE('Message #2');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
END;
```

```
SEND_MESSAGE status: 0
SEND_MESSAGE status: 0
```

Receive the first message and unpack it:

```
DECLARE
    v_item      VARCHAR2(80);
    v_status    INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
    DBMS_PIPE.UNPACK_MESSAGE(v_item);
    DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
END;
```

```
RECEIVE_MESSAGE status: 0
Item: Message #1
```

Purge the pipe:

```
EXEC DBMS_PIPE.PURGE('pipe');
```

Try to retrieve the next message. The [RECEIVE_MESSAGE](#) call returns status code 1 indicating it timed out because no message was available.

```
DECLARE
    v_item      VARCHAR2(80);
    v_status    INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
```

```
END;
```

```
RECEIVE_MESSAGE status: 1
```

6.3.10.5 RECEIVE_MESSAGE

The `RECEIVE_MESSAGE` function obtains a message from a specified pipe.

```
<status> INTEGER RECEIVE_MESSAGE(<pipename> VARCHAR2
[, <timeout> INTEGER ])
```

Parameters

`pipename`

Name of the pipe.

`timeout`

Wait time (seconds). Default is 86400000 (1000 days).

`status`

Status code returned by the operation.

The possible status codes are:

| Status Code | Description |
|-------------|-----------------------------------|
| 0 | Success |
| 1 | Time out |
| 2 | Message too large .for the buffer |

6.3.10.6 REMOVE_PIPE

The `REMOVE_PIPE` function deletes an explicit private or explicit public pipe.

```
<status> INTEGER REMOVE_PIPE(<pipename> VARCHAR2)
```

Use the `REMOVE_PIPE` function to delete explicitly created pipes – i.e., pipes created with the `CREATE_PIPE` function.

Parameters

`pipename`

Name of the pipe.

status

Status code returned by the operation. A status code of 0 is returned even if the named pipe is non-existent.

Examples

Two messages are sent on a pipe:

```
DECLARE
    v_status      INTEGER;
BEGIN
    v_status := DBMS_PIPE.CREATE_PIPE('pipe');
    DBMS_OUTPUT.PUT_LINE('CREATE_PIPE status : ' || v_status);

    DBMS_PIPE.PACK_MESSAGE('Message #1');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);

    DBMS_PIPE.PACK_MESSAGE('Message #2');
    v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
    DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
END;

CREATE_PIPE status : 0
SEND_MESSAGE status: 0
SEND_MESSAGE status: 0
```

Receive the first message and unpack it:

```
DECLARE
    v_item      VARCHAR2(80);
    v_status    INTEGER;
BEGIN
    v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
    DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
    DBMS_PIPE.UNPACK_MESSAGE(v_item);
    DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
END;

RECEIVE_MESSAGE status: 0
Item: Message #1
```

Remove the pipe:

```
SELECT DBMS_PIPE.REMOVE_PIPE('pipe') FROM DUAL;

remove_pipe
-----
0
(1 row)
```

Try to retrieve the next message. The `RECEIVE_MESSAGE` call returns status code 1 indicating it timed out because the pipe had been deleted.

```
DECLARE
    v_item      VARCHAR2(80);
    v_status    INTEGER;
```

```
BEGIN
  v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
  DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
END;
```

RECEIVE_MESSAGE status: 1

6.3.10.7 RESET_BUFFER

The `RESET_BUFFER` procedure resets a “pointer” to the session’s local message buffer back to the beginning of the buffer. This has the effect of causing subsequent `PACK_MESSAGE` calls to overwrite any data items that existed in the message buffer prior to the `RESET_BUFFER` call.

`RESET_BUFFER`

Examples

A message to John is written to the local message buffer. It is replaced by a message to Bob by calling `RESET_BUFFER`. The message is sent on the pipe.

```
DECLARE
  v_status      INTEGER;
BEGIN
  DBMS_PIPE.PACK_MESSAGE('Hi, John');
  DBMS_PIPE.PACK_MESSAGE('Can you attend a meeting at 3:00, today?');
  DBMS_PIPE.PACK_MESSAGE('If not, is tomorrow at 8:30 ok with you?');
  DBMS_PIPE.RESET_BUFFER;
  DBMS_PIPE.PACK_MESSAGE('Hi, Bob');
  DBMS_PIPE.PACK_MESSAGE('Can you attend a meeting at 9:30, tomorrow?');
  v_status := DBMS_PIPE.SEND_MESSAGE('pipe');
  DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE status: ' || v_status);
END;
```

SEND_MESSAGE status: 0

The message to Bob is in the received message.

```
DECLARE
  v_item      VARCHAR2(80);
  v_status    INTEGER;
BEGIN
  v_status := DBMS_PIPE.RECEIVE_MESSAGE('pipe',1);
  DBMS_OUTPUT.PUT_LINE('RECEIVE_MESSAGE status: ' || v_status);
  DBMS_PIPE.UNPACK_MESSAGE(v_item);
  DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
  DBMS_PIPE.UNPACK_MESSAGE(v_item);
  DBMS_OUTPUT.PUT_LINE('Item: ' || v_item);
END;
```

RECEIVE_MESSAGE status: 0

Item: Hi, Bob

Item: Can you attend a meeting at 9:30, tomorrow?

6.3.10.8 SEND_MESSAGE

The `SEND_MESSAGE` function sends a message from the session's local message buffer to the specified pipe.

```
<status> SEND_MESSAGE(<pipename> VARCHAR2 [, <timeout> INTEGER ]
[, <maxpipesize> INTEGER ])
```

Parameters

`pipename`

 | Name of the pipe.

`timeout`

 | Wait time (seconds). Default is 86400000 (1000 days).

`maxpipesize`

 | Maximum capacity of the pipe in bytes. Default is 8192 bytes.

`status`

 | Status code returned by the operation.

The possible status codes are:

| Status Code | Description |
|-------------|----------------------|
| 0 | Success |
| 1 | Time out |
| 3 | Function interrupted |

6.3.10.9 UNIQUE_SESSION_NAME

The `UNIQUE_SESSION_NAME` function returns a name, unique to the current session.

```
<name> VARCHAR2 UNIQUE_SESSION_NAME
```

Parameters

`name`

 | Unique session name.

Examples

The following anonymous block retrieves and displays a unique session name.

```
DECLARE
  v_session  VARCHAR2(30);
BEGIN
  v_session := DBMS_PIPE.UNIQUE_SESSION_NAME;
  DBMS_OUTPUT.PUT_LINE('Session Name: ' || v_session);
END;
```

```
Session Name: PG$PIPE$5$2752
```

6.3.10.10 UNPACK_MESSAGE

The **UNPACK_MESSAGE** procedure copies the data items of a message from the local message buffer to a specified program variable. The message must be placed in the local message buffer with the **RECEIVE_MESSAGE** function before using **UNPACK_MESSAGE**.

```
UNPACK_MESSAGE(<item> OUT { DATE | NUMBER | VARCHAR2 | RAW })
```

Parameters

item

Type-compatible variable that receives a data item from the local message buffer.

6.3.10.11 Comprehensive Example

The following example uses a pipe as a “mailbox”. The procedures to create the mailbox, add a multi-item message to the mailbox (up to three items), and display the full contents of the mailbox are enclosed in a package named, **mailbox**.

```
CREATE OR REPLACE PACKAGE mailbox
IS
  PROCEDURE create_mailbox;
  PROCEDURE add_message (
    p_mailbox  VARCHAR2,
    p_item_1   VARCHAR2,
    p_item_2   VARCHAR2 DEFAULT 'END',
    p_item_3   VARCHAR2 DEFAULT 'END'
  );
  PROCEDURE empty_mailbox (
    p_mailbox  VARCHAR2,
    p_waittime INTEGER DEFAULT 10
  );
END mailbox;
```

```

CREATE OR REPLACE PACKAGE BODY mailbox
IS
  PROCEDURE create_mailbox
  IS
    v_mailbox  VARCHAR2(30);
    v_status   INTEGER;
  BEGIN
    v_mailbox := DBMS_PIPE.UNIQUE_SESSION_NAME;
    v_status := DBMS_PIPE.CREATE_PIPE(v_mailbox,1000, FALSE);
    IF v_status = 0 THEN
      DBMS_OUTPUT.PUT_LINE('Created mailbox: ' || v_mailbox);
    ELSE
      DBMS_OUTPUT.PUT_LINE('CREATE_PIPE failed - status: ' ||
                           v_status);
    END IF;
  END create_mailbox;

  PROCEDURE add_message (
    p_mailbox  VARCHAR2,
    p_item_1   VARCHAR2,
    p_item_2   VARCHAR2 DEFAULT 'END',
    p_item_3   VARCHAR2 DEFAULT 'END'
  )
  IS
    v_item_cnt  INTEGER := 0;
    v_status    INTEGER;
  BEGIN
    DBMS_PIPE.PACK_MESSAGE(p_item_1);
    v_item_cnt := 1;
    IF p_item_2 != 'END' THEN
      DBMS_PIPE.PACK_MESSAGE(p_item_2);
      v_item_cnt := v_item_cnt + 1;
    END IF;
    IF p_item_3 != 'END' THEN
      DBMS_PIPE.PACK_MESSAGE(p_item_3);
      v_item_cnt := v_item_cnt + 1;
    END IF;
    v_status := DBMS_PIPE.SEND_MESSAGE(p_mailbox);
    IF v_status = 0 THEN
      DBMS_OUTPUT.PUT_LINE('Added message with ' || v_item_cnt ||
                           ' item(s) to mailbox ' || p_mailbox);
    ELSE
      DBMS_OUTPUT.PUT_LINE('SEND_MESSAGE in add_message failed - ' ||
                           'status: ' || v_status);
    END IF;
  END add_message;

  PROCEDURE empty_mailbox (
    p_mailbox  VARCHAR2,
    p_waittime INTEGER DEFAULT 10
  )
  IS
    v_msgno   INTEGER DEFAULT 0;
    v_itemno  INTEGER DEFAULT 0;
    v_item    VARCHAR2(100);

```

```

v_status  INTEGER;
BEGIN
  v_status := DBMS_PIPE.RECEIVE_MESSAGE(p_mailbox,p_waittime);
  WHILE v_status = 0 LOOP
    v_msgno := v_msgno + 1;
    DBMS_OUTPUT.PUT_LINE('***** Start message #' || v_msgno ||
      ' *****');
    BEGIN
      LOOP
        v_status := DBMS_PIPE.NEXT_ITEM_TYPE;
        EXIT WHEN v_status = 0;
        DBMS_PIPE.UNPACK_MESSAGE(v_item);
        v_itemno := v_itemno + 1;
        DBMS_OUTPUT.PUT_LINE('Item #' || v_itemno || ':' ||
          v_item);
      END LOOP;
      DBMS_OUTPUT.PUT_LINE('***** End message #' || v_msgno ||
        ' *****');
      DBMS_OUTPUT.PUT_LINE('*');
      v_itemno := 0;
      v_status := DBMS_PIPE.RECEIVE_MESSAGE(p_mailbox,1);
    END;
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('Number of messages received: ' || v_msgno);
  v_status := DBMS_PIPE.REMOVE_PIPE(p_mailbox);
  IF v_status = 0 THEN
    DBMS_OUTPUT.PUT_LINE('Deleted mailbox ' || p_mailbox);
  ELSE
    DBMS_OUTPUT.PUT_LINE('Could not delete mailbox - status: ' ||
      || v_status);
  END IF;
  END empty_mailbox;
END mailbox;

```

The following demonstrates the execution of the procedures in `mailbox`. The first procedure creates a public pipe using a name generated by the `UNIQUE_SESSION_NAME` function.

```
EXEC mailbox.create_mailbox;
```

Created mailbox: PG\$PIPE\$13\$3940

Using the mailbox name, any user in the same database with access to the `mailbox` package and `DBMS_PIPE` package can add messages:

```
EXEC mailbox.add_message('PG$PIPE$13$3940','Hi, John','Can you attend a
meeting at 3:00, today?','-- Mary');
```

Added message with 3 item(s) to mailbox PG\$PIPE\$13\$3940

```
EXEC mailbox.add_message('PG$PIPE$13$3940','Don''t forget to submit your
report','Thanks','-- Joe');
```

Added message with 3 item(s) to mailbox PG\$PIPE\$13\$3940

Finally, the contents of the mailbox can be emptied:

```
EXEC mailbox.empty_mailbox('PG$PIPE$13$3940');
```

***** Start message #1 *****

Item #1: Hi, John

Item #2: Can you attend a meeting at 3:00, today?

Item #3: -- Mary

***** End message #1 *****

*

***** Start message #2 *****

Item #1: Don't forget to submit your report

Item #2: Thanks,

Item #3: Joe

***** End message #2 *****

*

Number of messages received: 2

Deleted mailbox PG\$PIPE\$13\$3940

6.3.11 DBMS_PROFILER

The `DBMS_PROFILER` package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a performance profiling session; use the functions and procedures listed below to control the profiling tool.

| Function/Procedure | Return Type | Description |
|--|--------------------------|--|
| <code>FLUSH_DATA</code> | Status Code or Exception | Flushes performance data collected in the current session without terminating the session (profiling continues). |
| <code>GET_VERSION(major OUT, minor OUT)</code> | n/a | Returns the version number of this package. |
| <code>INTERNAL_VERSION_CHECK</code> | Status Code | Confirms that the current version of the profiler will work with the current database. |
| <code>PAUSE_PROFILER</code> | Status Code or Exception | Pause data collection. |
| <code>RESUME_PROFILER</code> | Status Code or Exception | Resume data collection. |
| <code>START_PROFILER(run_comment, run_comment1 [, run_number OUT])</code> | Status Code or Exception | Start data collection. |
| <code>STOP_PROFILER</code> | Status Code or Exception | Stop data collection and flush performance data to the <code>PLSQL_PROFILER_RAWDATA</code> table. |

The functions within the `DBMS_PROFILER` package return a status code to indicate success or failure; the `DBMS_PROFILER` procedures raise an exception only if they encounter a failure. The status codes and messages returned by the functions, and the exceptions raised by the procedures are listed in the table below.

| Status Code | Message | Exception | Description |
|-------------|----------------------------|-------------------------------|---|
| -1 | <code>error_version</code> | <code>version_mismatch</code> | The profiler version and the database are incompatible. |
| 0 | <code>success</code> | n/a | The operation completed successfully. |
| 1 | <code>error_param</code> | <code>profiler_error</code> | The operation received an incorrect parameter. |

| Status Code | Message | Exception | Description |
|-------------|----------|----------------|--------------------------------------|
| 2 | error_io | profiler_error | The data flush operation has failed. |

FLUSH_DATA

The `FLUSH_DATA` function/procedure flushes the data collected in the current session without terminating the profiler session. The data is flushed to the tables described in the Advanced Server Performance Features Guide. The function and procedure signatures are:

```
<status> INTEGER FLUSH_DATA
```

```
FLUSH_DATA
```

Parameters

`status`

| Status code returned by the operation.

GET_VERSION

The `GET_VERSION` procedure returns the version of `DBMS_PROFILER`. The procedure signature is:

```
GET_VERSION(<major> OUT INTEGER, <minor> OUT INTEGER)
```

Parameters

`major`

| The major version number of `DBMS_PROFILER`.

`minor`

| The minor version number of `DBMS_PROFILER`.

INTERNAL_VERSION_CHECK

The `INTERNAL_VERSION_CHECK` function confirms that the current version of `DBMS_PROFILER` will work with the current database. The function signature is:

```
<status> INTEGER INTERNAL_VERSION_CHECK
```

Parameters

`status`

| Status code returned by the operation.

PAUSE_PROFILER

The `PAUSE_PROFILER` function/procedure pauses a profiling session. The function and procedure signatures

are:

```
<status> INTEGER PAUSE_PROFILER  
PAUSE_PROFILER
```

Parameters

status

| Status code returned by the operation.

RESUME_PROFILER

The **RESUME_PROFILER** function/procedure pauses a profiling session. The function and procedure signatures are:

```
<status> INTEGER RESUME_PROFILER  
RESUME_PROFILER
```

Parameters

status

| Status code returned by the operation.

START_PROFILER

The **START_PROFILER** function/procedure starts a data collection session. The function and procedure signatures are:

```
<status> INTEGER START_PROFILER(<run_comment> TEXT := SYSDATE,  
<run_comment1> TEXT := " [, <run_number> OUT INTEGER ])  
  
START_PROFILER(<run_comment> TEXT := SYSDATE,  
<run_comment1> TEXT := " [, <run_number> OUT INTEGER ])
```

Parameters

run_comment

| A user-defined comment for the profiler session. The default value is **SYSDATE**.

run_comment1

| An additional user-defined comment for the profiler session. The default value is ".

run_number

| The session number of the profiler session.

status

| Status code returned by the operation.

STOP_PROFILER

The `STOP_PROFILER` function/procedure stops a profiling session and flushes the performance information to the `DBMS_PROFILER` tables and view. The function and procedure signatures are:

```
<status> INTEGER STOP_PROFILER
```

```
STOP_PROFILER
```

Parameters

`status`

Status code returned by the operation.

Using DBMS_PROFILER

The `DBMS_PROFILER` package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a profiling session; you can review the performance information in the tables and views provided by the profiler.

`DBMS_PROFILER` works by recording a set of performance-related counters and timers for each line of PL/pgSQL or SPL statement that executes within a profiling session. The counters and timers are stored in a table named `SYS.PLSQL_PROFILER_DATA`. When you complete a profiling session, `DBMS_PROFILER` will write a row to the performance statistics table for each line of PL/pgSQL or SPL code that executed within the session. For example, if you execute the following function:

```
1 - CREATE OR REPLACE FUNCTION getBalance(acctNumber INTEGER)
2 - RETURNS NUMERIC AS $$

3 - DECLARE
4 - result NUMERIC;
5 - BEGIN
6 - SELECT INTO result balance FROM acct WHERE id = acctNumber;

7 -
8 - IF (result IS NULL) THEN
9 -   RAISE INFO 'Balance is null';

10- END IF;

11-
12- RETURN result;
13- END;

14- $$ LANGUAGE 'plpgsql';
```

`DBMS_PROFILER` adds one `PLSQL_PROFILER_DATA` entry for each line of code within the `getBalance()` function (including blank lines and comments). The entry corresponding to the `SELECT` statement executed exactly one time; and required a very small amount of time to execute. On the other hand, the entry corresponding to the `RAISE INFO` statement executed once or not at all (depending on the value for the `balance` column).

Some of the lines in this function contain no executable code so the performance statistics for those lines will always contain zero values.

To start a profiling session, invoke the `DBMS_PROFILER.START_PROFILER` function (or procedure). Once you've invoked `START_PROFILER`, Advanced Server will profile every PL/pgSQL or SPL function, procedure, trigger, or anonymous block that your session executes until you either stop or pause the profiler (by calling `STOP_PROFILER` or `PAUSE_PROFILER`).

It is important to note that when you start (or resume) the profiler, the profiler will only gather performance statistics for functions/procedures/triggers that start after the call to `START_PROFILER` (or `RESUME_PROFILER`).

While the profiler is active, Advanced Server records a large set of timers and counters in memory; when you invoke the `STOP_PROFILER` (or `FLUSH_DATA`) function/procedure, `DBMS_PROFILER` writes those timers and counters to a set of three tables:

- `SYS.PLSQL_PROFILER_RAWDATA`

Contains the performance counters and timers for each statement executed within the session.

- `SYS.PLSQL_PROFILER_RUNS`

Contains a summary of each run (aggregating the information found in `PLSQL_PROFILER_RAWDATA`).

- `SYS.PLSQL_PROFILER_UNITS`

Contains a summary of each code unit (function, procedure, trigger, or anonymous block) executed within a session.

In addition, `DBMS_PROFILER` defines a view, `SYS.PLSQL_PROFILER_DATA`, which contains a subset of the `PLSQL_PROFILER_RAWDATA` table.

Please note that a non-superuser may gather profiling information, but may not view that profiling information unless a superuser grants specific privileges on the profiling tables (stored in the `SYS` schema). This permits a non-privileged user to gather performance statistics without exposing information that the administrator may want to keep secret.

Querying the DBMS_PROFILER Tables and View

The following step-by-step example uses `DBMS_PROFILER` to retrieve performance information for procedures, functions, and triggers included in the sample data distributed with Advanced Server.

1. Open the EDB-PSQL command line, and establish a connection to the Advanced Server database. Use an `EXEC` statement to start the profiling session:

```
acctg=# EXEC dbms_profiler.start_profiler('profile list_emp');
```

EDB-SPL Procedure successfully completed

Note

(The call to `start_profiler()` includes a comment that `DBMS_PROFILER` associates with the profiler session).

1. Then call the `list_emp` function:

```
acctg=# SELECT list_emp();
INFO: EMPNO  ENAME
INFO: ----- -----
INFO: 7369   SMITH
INFO: 7499   ALLEN
```

```

INFO: 7521  WARD
INFO: 7566  JONES
INFO: 7654  MARTIN
INFO: 7698  BLAKE
INFO: 7782  CLARK
INFO: 7788  SCOTT
INFO: 7839  KING
INFO: 7844  TURNER
INFO: 7876  ADAMS
INFO: 7900  JAMES
INFO: 7902  FORD
INFO: 7934  MILLER
list_emp
-----

```

(1 row)

1. Stop the profiling session with a call to `dbms_profiler.stop_profiler`:

```
acctg=# EXEC dbms_profiler.stop_profiler;
```

EDB-SPL Procedure successfully completed

1. Start a new session with the `dbms_profiler.start_profiler` function (followed by a new comment):

```
acctg=# EXEC dbms_profiler.start_profiler('profile get_dept_name and
emp_sal_trig');
```

EDB-SPL Procedure successfully completed

1. Invoke the `get_dept_name` function:

```
acctg=# SELECT get_dept_name(10);
get_dept_name
-----
```

ACCOUNTING
(1 row)

1. Execute an `UPDATE` statement that causes a trigger to execute:

```
acctg=# UPDATE memp SET sal = 500 WHERE empno = 7902;
INFO: Updating employee 7902
INFO: ..Old salary: 3000.00
INFO: ..New salary: 500.00
INFO: ..Raise: -2500.00
INFO: User enterprisedb updated employee(s) on 04-FEB-14
UPDATE 1
```

1. Terminate the profiling session and flush the performance information to the profiling tables:

```
acctg=# EXEC dbms_profiler.stop_profiler;
```

EDB-SPL Procedure successfully completed

1. Now, query the `plsql_profiler_runs` table to view a list of the profiling sessions, arranged by `runid`:

```
acctg=# SELECT * FROM plsql_profiler_runs;
```

```

runid | related_run | run_owner |      run_date
|       run_comment        | run_total_time | run_system_info
| run_comment1 | spare1
-----+-----+-----+-----+
-----+-----+-----+-----+
+----+
 1 |      | enterprisedb | 04-FEB-14 09:32:48.874315 | profile
list_emp           |      4154 |
  |      |
 2 |      | enterprisedb | 04-FEB-14 09:41:30.546503 | profile
get_dept_name and emp_sal_trig |      2088 |
  |      |
(2 rows)

```

1. Query the `plsql_profiler_units` table to view the amount of time consumed by each unit (each function, procedure, or trigger):

```

acctg=# SELECT * FROM plsql_profiler_units;
runid | unit_number | unit_type | unit_owner |
unit_name      | unit_timestamp | total_time | spare1 | spare2
-----+-----+-----+-----+
-----+-----+-----+-----+
 1 | 16999 | FUNCTION | enterprisedb |
list_emp()           |      |      4 |      |
 2 | 17002 | FUNCTION | enterprisedb |
user_audit_trig()    |      |      1 |      |
 2 | 17000 | FUNCTION | enterprisedb | get_dept_name(p_deptno
numeric) |      |      1 |      |
 2 | 17004 | FUNCTION | enterprisedb |
emp_sal_trig()      |      |      1 |      |
(4 rows)

```

1. Query the `plsql_profiler_rawdata` table to view a list of the wait event counters and wait event times:

```

acctg=# SELECT runid, sourcecode, func_oid, line_number, exec_count,
tuples_returned, time_total FROM plsql_profiler_rawdata;

```

```

runid |          sourcecode          |
func_oid | line_number | exec_count | tuples_returned | time_total
-----+-----+-----+-----+
-----+-----+-----+-----+
 1 | DECLARE
  | 16999 |      1 |      0 |      0 |      0
  | v_empno   NUMERIC(4);
  | 16999 |      2 |      0 |      0 |      0
  | v_ename   VARCHAR(10);
  | 16999 |      3 |      0 |      0 |      0
  | emp_cur CURSOR FOR
  | 16999 |      4 |      0 |      0 |      0
  |      SELECT empno, ename FROM memp ORDER BY empno;
  | 16999 |      5 |      0 |      0 |      0
  | BEGIN
  | 16999 |      6 |      0 |      0 |      0
  | OPEN emp_cur;
  | 16999 |      7 |      0 |      0 |      0
  | RAISE INFO 'EMPNO  ENAME';

```

```

| 16999 |     8 |     1 |          0 |  0.001621
1 | RAISE INFO '----';
| 16999 |     9 |     1 |          0 |  0.000301
1 | LOOP
| 16999 |    10 |     1 |          0 |  4.6e-05
1 |     FETCH emp_cur INTO v_empno, v_ename;
| 16999 |    11 |     1 |          0 |  0.001114
1 |     EXIT WHEN NOT FOUND;
| 16999 |    12 |    15 |          0 |  0.000206
1 |     RAISE INFO '% %', v_empno, v_ename;
| 16999 |    13 |    15 |          0 |  8.3e-05
1 | END LOOP;
| 16999 |    14 |    14 |          0 |  0.000773
1 | CLOSE emp_cur;
| 16999 |    15 |     0 |          0 |      0
1 | RETURN;
| 16999 |    16 |     1 |          0 |  1e-05
1 | END;
| 16999 |    17 |     1 |          0 |      0
1 |
| 16999 |    18 |     0 |          0 |      0
2 | DECLARE
| 17002 |     1 |     0 |          0 |      0
2 | v_action      VARCHAR(24);
| 17002 |     2 |     0 |          0 |      0
2 | v_text        TEXT;
| 17002 |     3 |     0 |          0 |      0
2 | BEGIN
| 17002 |     4 |     0 |          0 |      0
2 | IF TG_OP = 'INSERT' THEN
| 17002 |     5 |     0 |          0 |      0
2 |   v_action := ' added employee(s) on ';
| 17002 |     6 |     1 |          0 |  0.000143
2 | ELSIF TG_OP = 'UPDATE' THEN
| 17002 |     7 |     0 |          0 |      0
2 |   v_action := ' updated employee(s) on ';
| 17002 |     8 |     0 |          0 |      0
2 | ELSIF TG_OP = 'DELETE' THEN
| 17002 |     9 |     1 |          0 |  3.2e-05
2 |   v_action := ' deleted employee(s) on ';
| 17002 |    10 |     0 |          0 |      0
2 | END IF;
| 17002 |    11 |     0 |          0 |      0
2 | v_text := 'User ' || USER || v_action || CURRENT_DATE;
| 17002 |    12 |     0 |          0 |      0
2 | RAISE INFO '%', v_text;
| 17002 |    13 |     1 |          0 |  0.000383
2 | RETURN NULL;
| 17002 |    14 |     1 |          0 |  6.3e-05
2 | END;
| 17002 |    15 |     1 |          0 |  3.6e-05
2 |
| 17002 |    16 |     0 |          0 |      0
2 | DECLARE
| 17000 |     1 |     0 |          0 |      0

```

```

2 |   v_dname      VARCHAR(14);
| 17000 |     2 |     0 |     0 |     0
2 | BEGIN
| 17000 |     3 |     0 |     0 |     0
2 |   SELECT INTO v_dname dname FROM dept WHERE deptno = p_deptno;
| 17000 |     4 |     0 |     0 |     0
2 |   RETURN v_dname;
| 17000 |     5 |     1 |     0 | 0.000647
2 | IF NOT FOUND THEN
| 17000 |     6 |     1 |     0 | 2.6e-05
2 |   RAISE INFO 'Invalid department number %', p_deptno;
| 17000 |     7 |     0 |     0 |     0
2 |   RETURN "";
| 17000 |     8 |     0 |     0 |     0
2 | END IF;
| 17000 |     9 |     0 |     0 |     0
2 | END;
| 17000 |    10 |     0 |     0 |     0
2 |
| 17000 |    11 |     0 |     0 |     0
2 | DECLARE
| 17004 |     1 |     0 |     0 |     0
2 |   sal_diff  NUMERIC(7,2);
| 17004 |     2 |     0 |     0 |     0
2 | BEGIN
| 17004 |     3 |     0 |     0 |     0
2 | IF TG_OP = 'INSERT' THEN
| 17004 |     4 |     0 |     0 |     0
2 |   RAISE INFO 'Inserting employee %', NEW.empno;
| 17004 |     5 |     1 |     0 | 8.4e-05
2 |   RAISE INFO '..New salary: %', NEW.sal;
| 17004 |     6 |     0 |     0 |     0
2 |   RETURN NEW;
| 17004 |     7 |     0 |     0 |     0
2 | END IF;
| 17004 |     8 |     0 |     0 |     0
2 | IF TG_OP = 'UPDATE' THEN
| 17004 |     9 |     0 |     0 |     0
2 |   sal_diff := NEW.sal - OLD.sal;
| 17004 |    10 |     1 |     0 | 0.000355
2 |   RAISE INFO 'Updating employee %', OLD.empno;
| 17004 |    11 |     1 |     0 | 0.000177
2 |   RAISE INFO '..Old salary: %', OLD.sal;
| 17004 |    12 |     1 |     0 | 5.5e-05
2 |   RAISE INFO '..New salary: %', NEW.sal;
| 17004 |    13 |     1 |     0 | 3.1e-05
2 |   RAISE INFO '..Raise : %', sal_diff;
| 17004 |    14 |     1 |     0 | 2.8e-05
2 |   RETURN NEW;
| 17004 |    15 |     1 |     0 | 2.7e-05
2 | END IF;
| 17004 |    16 |     1 |     0 | 1e-06
2 | IF TG_OP = 'DELETE' THEN
| 17004 |    17 |     0 |     0 |     0
2 |   RAISE INFO 'Deleting employee %', OLD.empno;

```

```

| 17004 |    18 |    0 |    0 |    0
2|      RAISE INFO '..Old salary: %', OLD.sal;
| 17004 |    19 |    0 |    0 |    0
2|      RETURN OLD;
| 17004 |    20 |    0 |    0 |    0
2|      END IF;
| 17004 |    21 |    0 |    0 |    0
2| END;
| 17004 |    22 |    0 |    0 |    0
2|
| 17004 |    23 |    0 |    0 |    0
(68 rows)

```

1. Query the `plsql_profiler_data` view to review a subset of the information found in `plsql_profiler_rawdata` table:

```

acctg=# SELECT * FROM plsql_profiler_data;
runid | unit_number | line# | total_occur | total_time | min_time | max_time
| spare1 | spare2 | spare3 | spare4
-----+-----+-----+-----+-----+-----+
-----+-----+-----+
1 | 16999 | 1 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 2 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 3 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 4 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 5 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 6 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 7 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 8 | 1 | 0.001621 | 0.001621 |
0.001621 | | | |
1 | 16999 | 9 | 1 | 0.000301 | 0.000301 |
0.000301 | | | |
1 | 16999 | 10 | 1 | 4.6e-05 | 4.6e-05 | 4.6e-
05 | | | |
1 | 16999 | 11 | 1 | 0.001114 | 0.001114 |
0.001114 | | | |
1 | 16999 | 12 | 15 | 0.000206 | 5e-06 | 7.8e-
05 | | | |
1 | 16999 | 13 | 15 | 8.3e-05 | 2e-06 | 4.7e-
05 | | | |
1 | 16999 | 14 | 14 | 0.000773 | 4.7e-05 |
0.000116 | | | |
1 | 16999 | 15 | 0 | 0 | 0 |
0 | | | |
1 | 16999 | 16 | 1 | 1e-05 | 1e-05 | 1e-
05 | | | |
1 | 16999 | 17 | 1 | 0 | 0 |
0 | | | |
1 | 16999 | 18 | 0 | 0 | 0 |

```

| | | | | | | |
|----------|-------|----|---|----------|----------|-------|
| 0 | | | | | | |
| 2 | 17002 | 1 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 2 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 3 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 4 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 5 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 6 | 1 | 0.000143 | 0.000143 | |
| 0.000143 | | | | | | |
| 2 | 17002 | 7 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 8 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 9 | 1 | 3.2e-05 | 3.2e-05 | 3.2e- |
| 05 | | | | | | |
| 2 | 17002 | 10 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 11 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 12 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17002 | 13 | 1 | 0.000383 | 0.000383 | |
| 0.000383 | | | | | | |
| 2 | 17002 | 14 | 1 | 6.3e-05 | 6.3e-05 | 6.3e- |
| 05 | | | | | | |
| 2 | 17002 | 15 | 1 | 3.6e-05 | 3.6e-05 | 3.6e- |
| 05 | | | | | | |
| 2 | 17002 | 16 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 1 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 2 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 3 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 4 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 5 | 1 | 0.000647 | 0.000647 | |
| 0.000647 | | | | | | |
| 2 | 17000 | 6 | 1 | 2.6e-05 | 2.6e-05 | 2.6e- |
| 05 | | | | | | |
| 2 | 17000 | 7 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 8 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 9 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 10 | 0 | 0 | 0 | |
| 0 | | | | | | |
| 2 | 17000 | 11 | 0 | 0 | 0 | |
| 0 | | | | | | |

| | | | | | |
|----|----------|----|---|----------|----------|
| 2 | 17004 | 1 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 2 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 3 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 4 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 5 | 1 | 8.4e-05 | 8.4e-05 |
| 05 | | | | 8.4e- | |
| 2 | 17004 | 6 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 7 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 8 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 9 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 10 | 1 | 0.000355 | 0.000355 |
| | 0.000355 | | | | |
| 2 | 17004 | 11 | 1 | 0.000177 | 0.000177 |
| | 0.000177 | | | | |
| 2 | 17004 | 12 | 1 | 5.5e-05 | 5.5e-05 |
| 05 | | | | 5.5e- | |
| 2 | 17004 | 13 | 1 | 3.1e-05 | 3.1e-05 |
| 05 | | | | 3.1e- | |
| 2 | 17004 | 14 | 1 | 2.8e-05 | 2.8e-05 |
| 05 | | | | 2.8e- | |
| 2 | 17004 | 15 | 1 | 2.7e-05 | 2.7e-05 |
| 05 | | | | 2.7e- | |
| 2 | 17004 | 16 | 1 | 1e-06 | 1e-06 |
| 06 | | | | 1e- | |
| 2 | 17004 | 17 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 18 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 19 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 20 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 21 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 22 | 0 | 0 | 0 |
| 0 | | | | | |
| 2 | 17004 | 23 | 0 | 0 | 0 |
| 0 | | | | | |

(68 rows)

DBMS_PROFILER - Reference

The Advanced Server installer creates the following tables and views that you can query to review PL/SQL performance profile information:

| Table Name | Description |
|------------------------|---|
| PLSQL_PROFILER_RUNS | Table containing information about all profiler runs, organized by <code>runid</code> . |
| PLSQL_PROFILER_UNITS | Table containing information about all profiler runs, organized by unit. |
| PLSQL_PROFILER_DATA | View containing performance statistics. |
| PLSQL_PROFILER_RAWDATA | Table containing the performance statistics <code>and</code> the extended performance statistics for DRITA counters and timers. |

PLSQL_PROFILER_RUNS

The `PLSQL_PROFILER_RUNS` table contains the following columns:

| Column | Data Type | Description |
|------------------------------|-----------------------------|---|
| <code>runid</code> | INTEGER (NOT NULL) | Unique identifier (<code>plsql_profiler_runnumber</code>) |
| <code>related_run</code> | INTEGER | The <code>runid</code> of a related run. |
| <code>run_owner</code> | TEXT | The role that recorded the profiling session. |
| <code>run_date</code> | TIMESTAMP WITHOUT TIME ZONE | The profiling session start time. |
| <code>run_comment</code> | TEXT | User comments relevant to this run |
| <code>run_total_time</code> | BIGINT | Run time (in microseconds) |
| <code>run_system_info</code> | TEXT | Currently Unused |
| <code>run_comment1</code> | TEXT | Additional user comments |
| <code>spare1</code> | TEXT | Currently Unused |

PLSQL_PROFILER_UNITS

The `PLSQL_PROFILER_UNITS` table contains the following columns:

| Column | Data Type | Description |
|-----------------------------|-----------------------------|---|
| <code>runid</code> | INTEGER | Unique identifier (<code>plsql_profiler_runnumber</code>) |
| <code>unit_number</code> | OID | Corresponds to the OID of the row in the <code>pg_proc</code> table that identifies the unit. |
| <code>unit_type</code> | TEXT | PL/SQL function, procedure, trigger or anonymous block |
| <code>unit_owner</code> | TEXT | The identity of the role that owns the unit. |
| <code>unit_name</code> | TEXT | The complete signature of the unit. |
| <code>unit_timestamp</code> | TIMESTAMP WITHOUT TIME ZONE | Creation date of the unit (currently NULL). |
| <code>total_time</code> | BIGINT | Time spent within the unit (in milliseconds) |
| <code>spare1</code> | BIGINT | Currently Unused |
| <code>spare2</code> | BIGINT | Currently Unused |

PLSQL_PROFILER_DATA

The `PLSQL_PROFILER_DATA` view contains the following columns:

| Column | Data Type | Description |
|--------------------------|-----------|---|
| <code>runid</code> | INTEGER | Unique identifier (<code>plsql_profiler_runnumber</code>) |
| <code>unit_number</code> | OID | Object ID of the unit that contains the current line. |
| <code>line#</code> | INTEGER | Current line number of the profiled workload. |

| Column | Data Type | Description |
|-------------|------------------|--|
| total_occur | BIGINT | The number of times that the line was executed. |
| total_time | DOUBLE PRECISION | The amount of time spent executing the line (in seconds) |
| min_time | DOUBLE PRECISION | The minimum execution time for the line. |
| max_time | DOUBLE PRECISION | The maximum execution time for the line. |
| spare1 | NUMBER | Currently Unused |
| spare2 | NUMBER | Currently Unused |
| spare3 | NUMBER | Currently Unused |
| spare4 | NUMBER | Currently Unused |

PLSQL_PROFILER_RAWDATA

The `PLSQL_PROFILER_RAWDATA` table contains the statistical and wait events information that is found in the `PLSQL_PROFILER_DATA` view, as well as the performance statistics returned by the DRITA counters and timers.

| Column | Data Type | Description |
|-----------------|------------------|--|
| runid | INTEGER | The run identifier (plsql_profiler_runnumber). |
| sourcecode | TEXT | The individual line of profiled code. |
| func_oid | OID | Object ID of the unit that contains the current line. |
| line_number | INTEGER | Current line number of the profiled workload. |
| exec_count | BIGINT | The number of times that the line was executed. |
| tuples_returned | BIGINT | Currently Unused |
| time_total | DOUBLE PRECISION | The amount of time spent executing the line (in seconds) |
| time_shortest | DOUBLE PRECISION | The minimum execution time for the line. |
| time_longest | DOUBLE PRECISION | The maximum execution time for the line. |
| num_scans | BIGINT | Currently Unused |
| tuples_fetched | BIGINT | Currently Unused |
| tuples_inserted | BIGINT | Currently Unused |
| tuples_updated | BIGINT | Currently Unused |
| tuples_deleted | BIGINT | Currently Unused |
| blocks_fetched | BIGINT | Currently Unused |
| blocks_hit | BIGINT | Currently Unused |
| wal_write | BIGINT | A server has waited for a write to the write-ahead log buffer (expect this value to be high). |
| wal_flush | BIGINT | A server has waited for the write-ahead log to flush to disk. |
| wal_file_sync | BIGINT | A server has waited for the write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to open_sync). |
| db_file_read | BIGINT | A server has waited for the completion of a read (from disk). |
| db_file_write | BIGINT | A server has waited for the completion of a write (to disk). |
| db_file_sync | BIGINT | A server has waited for the operating system to flush all changes to disk. |

| Column | Data Type | Description |
|-----------------------------|-----------|--|
| db_file_extend | BIGINT | A server has waited for the operating system while adding a new page to the end of a file. |
| sql_parse | BIGINT | Currently Unused. |
| query_plan | BIGINT | A server has generated a query plan. |
| other_lwlock_acquire | BIGINT | A server has waited for other light-weight lock to protect data. |
| shared_plan_cache_collision | BIGINT | A server has waited for the completion of the shared_plan_cache_collision event. |
| shared_plan_cache_insert | BIGINT | A server has waited for the completion of the shared_plan_cache_insert event. |
| shared_plan_cache_hit | BIGINT | A server has waited for the completion of the shared_plan_cache_hit event. |
| shared_plan_cache_miss | BIGINT | A server has waited for the completion of the shared_plan_cache_miss event. |
| shared_plan_cache_lock | BIGINT | A server has waited for the completion of the shared_plan_cache_lock event. |
| shared_plan_cache_busy | BIGINT | A server has waited for the completion of the shared_plan_cache_busy event. |
| shmemindexlock | BIGINT | A server has waited to find or allocate space in the shared memory. |
| oidgenlock | BIGINT | A server has waited to allocate or assign an OID. |
| xidgenlock | BIGINT | A server has waited to allocate or assign a transaction ID. |
| procarraylock | BIGINT | A server has waited to get a snapshot or clearing a transaction ID at transaction end. |
| sinvalreadlock | BIGINT | A server has waited to retrieve or remove messages from shared invalidation queue. |
| sinvalwritelock | BIGINT | A server has waited to add a message to the shared invalidation queue. |
| walbufmappinglock | BIGINT | A server has waited to replace a page in WAL buffers. |
| walwritelock | BIGINT | A server has waited for WAL buffers to be written to disk. |
| controlfilelock | BIGINT | A server has waited to read or update the control file or creation of a new WAL file. |
| checkpointlock | BIGINT | A server has waited to perform a checkpoint. |
| clogcontrollock | BIGINT | A server has waited to read or update the transaction status. |
| subtranscontrollock | BIGINT | A server has waited to read or update the sub-transaction information. |
| multixactgenlock | BIGINT | A server has waited to read or update the shared multixact state. |
| multixactoffsetcontrollock | BIGINT | A server has waited to read or update multixact offset mappings. |
| multixactmembercontrollock | BIGINT | A server has waited to read or update multixact member mappings. |
| relcacheinitlock | BIGINT | A server has waited to read or write the relation cache initialization file. |
| checkpointercommlock | BIGINT | A server has waited to manage the fsync requests. |
| twophasesstatelock | BIGINT | A server has waited to read or update the state of prepared transactions. |
| tablespacecreatelock | BIGINT | A server has waited to create or drop the tablespace. |

| Column | Data Type | Description |
|-------------------------------------|-----------|--|
| btreevacuumlock | BIGINT | A server has waited to read or update the vacuum related information for a B-tree index. |
| addinshmeminitlock | BIGINT | A server has waited to manage space allocation in shared memory. |
| autovacuumlock | BIGINT | The autovacuum launcher waiting to read or update the current state of autovacuum workers. |
| autovacuumschedulelock | BIGINT | A server has waited to ensure that the table selected for a vacuum still needs vacuuming. |
| syncscanlock | BIGINT | A server has waited to get the start location of a scan on a table for synchronized scans. |
| relationmappinglock | BIGINT | A server has waited to update the relation map file used to store catalog to file node mapping. |
| asyncctllock | BIGINT | A server has waited to read or update shared notification state. |
| asyncqueuelock | BIGINT | A server has waited to read or update the notification messages. |
| serializableexacthashlock | BIGINT | A server has waited to retrieve or store information about serializable transactions. |
| serializablefinishedlistlock | BIGINT | A server has waited to access the list of finished serializable transactions. |
| serializablepredicateunlocklistlock | BIGINT | A server has waited to perform an operation on a list of locks held by serializable transactions. |
| oldserxidlock | BIGINT | A server has waited to read or record the conflicting serializable transactions. |
| syncrepllock | BIGINT | A server has waited to read or update information about synchronous replicas. |
| backgroundworkerlock | BIGINT | A server has waited to read or update the background worker state. |
| dynamicsharedmemorycontrollock | BIGINT | A server has waited to read or update the dynamic shared memory state. |
| autofilelock | BIGINT | A server has waited to update the <code>postgresql.auto.conf</code> file. |
| replicationslotallocationlock | BIGINT | A server has waited to allocate or free a replication slot. |
| replicationslotcontrollock | BIGINT | A server has waited to read or update replication slot state. |
| commitscontrollock | BIGINT | A server has waited to read or update transaction commit timestamps. |
| commitslock | BIGINT | A server has waited to read or update the last value set for the transaction timestamp. |
| replicationoriginlock | BIGINT | A server has waited to set up, drop, or use replication origin. |
| multixacttruncationlock | BIGINT | A server has waited to read or truncate multixact information. |
| oldsnapshottimemaplock | BIGINT | A server has waited to read or update old snapshot control information. |
| backendrandomlock | BIGINT | A server has waited to generate a random number. |
| logicalrepworkerlock | BIGINT | A server has waited for the action on logical replication worker to finish. |
| clogtruncationlock | BIGINT | A server has waited to truncate the write-ahead log or waiting for write-ahead log truncation to finish. |
| bulkloadlock | BIGINT | A server has waited for the <code>bulkloadlock</code> to bulk upload the data. |

| Column | Data Type | Description |
|----------------------------------|-----------|--|
| edbresourcemanagerlock | BIGINT | The <code>edbresourcemanagerlock</code> provides detail aboutedb resource manager lock module. |
| wal_write_time | BIGINT | The amount of time that the server has waited for a <code>wal_write</code> wait event to write to the write-ahead log buffer (expect this value to be high). |
| wal_flush_time | BIGINT | The amount of time that the server has waited for a <code>wal_flush</code> wait event to write-ahead log to flush to disk. |
| wal_file_sync_time | BIGINT | The amount of time that the server has waited for a <code>wal_file_sync</code> wait event to write-ahead log to sync to disk (related to the <code>wal_sync_method</code> parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to <code>open_sync</code>). |
| db_file_read_time | BIGINT | The amount of time that the server has waited for the <code>db_file_read</code> wait event for completion of a read (from disk). |
| db_file_write_time | BIGINT | The amount of time that the server has waited for the <code>db_file_write</code> wait event for completion of a write (to disk). |
| db_file_sync_time | BIGINT | The amount of time that the server has waited for the <code>db_file_sync</code> wait event to sync all changes to disk. |
| db_file_extend_time | BIGINT | The amount of time that the server has waited for the <code>db_file_extend</code> wait event while adding a new page to the end of a file. |
| sql_parse_time | BIGINT | The amount of time that the server has waited for the <code>sql_parse</code> wait event to parse a SQL statement. |
| query_plan_time | BIGINT | The amount of time that the server has waited for the <code>query_plan</code> wait event to compute the execution plan for a SQL statement. |
| other_lwlock_acquire_time | BIGINT | The amount of time that the server has waited for the <code>other_lwlock_acquire</code> wait event to protect data. |
| shared_plan_cache_collision_time | BIGINT | The amount of time that the server has waited for the <code>shared_plan_cache_collision</code> wait event. |
| shared_plan_cache_insert_time | BIGINT | The amount of time that the server has waited for the <code>shared_plan_cache_insert</code> wait event. |
| shared_plan_cache_hit_time | BIGINT | The amount of time that the server has waited for the <code>shared_plan_cache_hit</code> wait event. |
| shared_plan_cache_miss_time | BIGINT | The amount of time that the server has waited for the <code>shared_plan_cache_miss</code> wait event. |
| shared_plan_cache_lock_time | BIGINT | The amount of time that the server has waited for the <code>shared_plan_cache_lock</code> wait event. |
| shared_plan_cache_busy_time | BIGINT | The amount of time that the server has waited for the <code>shared_plan_cache_busy</code> wait event. |
| shmemindexlock_time | BIGINT | The amount of time that the server has waited for the <code>shmemindexlock</code> wait event to find or allocate space in the shared memory. |
| oidgenlock_time | BIGINT | The amount of time that the server has waited for the <code>oidgenlock</code> wait event to allocate or assign an OID. |
| xidgenlock_time | BIGINT | The amount of time that the server has waited for the <code>xidgenlock</code> wait event to allocate or assign a transaction ID. |
| procarraylock_time | BIGINT | The amount of time that the server has waited for a <code>procarraylock</code> wait event to clear a transaction ID at transaction end. |

| Column | Data Type | Description |
|---------------------------------|-----------|--|
| sinvalreadlock_time | BIGINT | The amount of time that the server has waited for a <code>sinvalreadlock</code> wait event to retrieve or remove messages from shared invalidation queue. |
| sinvalwritelock_time | BIGINT | The amount of time that the server has waited for a <code>sinvalwritelock</code> wait event to add a message to the shared invalidation queue. |
| walbufmappinglock_time | BIGINT | The amount of time that the server has waited for a <code>walbufmappinglock</code> wait event to replace a page in WAL buffers. |
| walwritelock_time | BIGINT | The amount of time that the server has waited for a <code>walwritelock</code> wait event to write the WAL buffers to disk. |
| controlfilelock_time | BIGINT | The amount of time that the server has waited for a <code>controlfilelock</code> wait event to read or update the control file or to create a new WAL file. |
| checkpointlock_time | BIGINT | The amount of time that the server has waited for a <code>checkpointlock</code> wait event to perform a checkpoint. |
| clogcontrollock_time | BIGINT | The amount of time that the server has waited for a <code>clogcontrollock</code> wait event to read or update the transaction status. |
| subtranscontrollock_time | BIGINT | The amount of time that the server has waited for the <code>subtranscontrollock</code> wait event to read or update the sub-transaction information. |
| multixactgenlock_time | BIGINT | The amount of time that the server has waited for the <code>multixactgenlock</code> wait event to read or update the shared multixact state. |
| multixactoffsetcontrollock_time | BIGINT | The amount of time that the server has waited for the <code>multixactoffsetcontrollock</code> wait event to read or update multixact offset mappings. |
| multixactmembercontrollock_time | BIGINT | The amount of time that the server has waited for the <code>multixactmembercontrollock</code> wait event to read or update multixact member mappings. |
| relcacheinitlock_time | BIGINT | The amount of time that the server has waited for the <code>relcacheinitlock</code> wait event to read or write the relation cache initialization file. |
| checkpointercommlock_time | BIGINT | The amount of time that the server has waited for the <code>checkpointercommlock</code> wait event to manage the fsync requests. |
| twophasestatelock_time | BIGINT | The amount of time that the server has waited for the <code>twophasestatelock</code> wait event to read or update the state of prepared transactions. |
| tablespacecreatelock_time | BIGINT | The amount of time that the server has waited for the <code>tablespacecreatelock</code> wait event to create or drop the tablespace. |
| btreevacuumlock_time | BIGINT | The amount of time that the server has waited for the <code>btreevacuumlock</code> wait event to read or update the vacuum related information for a B-tree index. |
| addinshmeminitlock_time | BIGINT | The amount of time that the server has waited for the <code>addinshmeminitlock</code> wait event to manage space allocation in shared memory. |
| autovacuumlock_time | BIGINT | The amount of time that the server has waited for the <code>autovacuumlock</code> wait event to read or update the current state of autovacuum workers. |
| autovacuumschedulelock_time | BIGINT | The amount of time that the server has waited for the <code>autovacuumschedulelock</code> wait event to ensure that the table selected for a vacuum still needs vacuuming. |

| Column | Data Type | Description |
|--|-----------|---|
| syncscanlock_time | BIGINT | The amount of time that the server has waited for the <code>syncscanlock</code> wait event to get the start location of a scan on a table for synchronized scans. |
| relationmappinglock_time | BIGINT | The amount of time that the server has waited for the <code>relationmappinglock</code> wait event to update the relation map file used to store catalog to file node mapping. |
| asyncctllock_time | BIGINT | The amount of time that the server has waited for the <code>asyncctllock</code> wait event to read or update shared notification state. |
| asyncqueuelock_time | BIGINT | The amount of time that the server has waited for the <code>asyncqueuelock</code> wait event to read or update the notification messages. |
| serializablexacthashlock_time | BIGINT | The amount of time that the server has waited for the <code>serializablexacthashlock</code> wait event to retrieve or store information about serializable transactions. |
| serializablefinishedlistlock_time | BIGINT | The amount of time that the server has waited for the <code>serializablefinishedlistlock</code> wait event to access the list of finished serializable transactions. |
| serializablepredicateunlocklistlock_time | BIGINT | The amount of time that the server has waited for the <code>serializablepredicateunlocklistlock</code> wait event to perform an operation on a list of locks held by serializable transactions. |
| oldserxidlock_time | BIGINT | The amount of time that the server has waited for the <code>oldserxidlock</code> wait event to read or record the conflicting serializable transactions. |
| syncreplock_time | BIGINT | The amount of time that the server has waited for the <code>syncreplock</code> wait event to read or update information about synchronous replicas. |
| backgroundworkerlock_time | BIGINT | The amount of time that the server has waited for the <code>backgroundworkerlock</code> wait event to read or update the background worker state. |
| dynamicsharedmemorycontrollock_time | BIGINT | The amount of time that the server has waited for the <code>dynamicsharedmemorycontrollock</code> wait event to read or update the dynamic shared memory state. |
| autofilelock_time | BIGINT | The amount of time that the server has waited for the <code>autofilelock</code> wait event to update the <code>postgresql.auto.conf</code> file. |
| replicationslotallocationlock_time | BIGINT | The amount of time that the server has waited for the <code>replicationslotallocationlock</code> wait event to allocate or free a replication slot. |
| replicationslotcontrollock_time | BIGINT | The amount of time that the server has waited for the <code>replicationslotcontrollock</code> wait event to read or update replication slot state. |
| commitscontrollock_time | BIGINT | The amount of time that the server has waited for the <code>commitscontrollock</code> wait event to read or update transaction commit timestamps. |
| commitslock_time | BIGINT | The amount of time that the server has waited for the <code>commitslock</code> wait event to read or update the last value set for the transaction timestamp. |
| replicationoriginlock_time | BIGINT | The amount of time that the server has waited for the <code>replicationoriginlock</code> wait event to set up, drop, or use replication origin. |
| multixacttruncationlock_time | BIGINT | The amount of time that the server has waited for the <code>multixacttruncationlock</code> wait event to read or truncate multixact information. |

| Column | Data Type | Description |
|-----------------------------|-----------|---|
| oldsnapshotmemlock_time | BIGINT | The amount of time that the server has waited for the <code>oldsnapshotmemlock</code> wait event to read or update old snapshot control information. |
| backendrandomlock_time | BIGINT | The amount of time that the server has waited for the <code>backendrandomlock</code> wait event to generate a random number. |
| logicalrepworkerlock_time | BIGINT | The amount of time that the server has waited for the <code>logicalrepworkerlock</code> wait event for an action on logical replication worker to finish. |
| clogtruncationlock_time | BIGINT | The amount of time that the server has waited for the <code>clogtruncationlock</code> wait event to truncate the write-ahead log or waiting for write-ahead log truncation to finish. |
| bulkloadlock_time | BIGINT | The amount of time that the server has waited for the <code>bulkloadlock</code> wait event to bulk upload the data. |
| edbresourcemanagerlock_time | BIGINT | The amount of time that the server has waited for the <code>edbresourcemanagerlock</code> wait event. |
| totalwaits | BIGINT | The total number of event waits. |
| totalwaittime | BIGINT | The total time spent waiting for an event. |

6.3.12 DBMS_RANDOM

The `DBMS_RANDOM` package provides a number of methods to generate random values. The procedures and functions available in the `DBMS_RANDOM` package are listed in the following table.

| Function/Procedure | Return Type | Description |
|-------------------------------|-------------|--|
| <code>INITIALIZE(val)</code> | n/a | Initializes the <code>DBMS_RANDOM</code> package with the specified seed <code>value</code> . Deprecated, but supported for backward compatibility. |
| <code>NORMAL()</code> | NUMBER | Returns a random <code>NUMBER</code> . |
| <code>RANDOM</code> | INTEGER | Returns a random <code>INTEGER</code> with a value greater than or equal to -2^{31} and less than 2^{31} . Deprecated, but supported for backward compatibility. |
| <code>SEED(val)</code> | n/a | Resets the seed with the specified <code>value</code> . |
| <code>SEED(val)</code> | n/a | Resets the seed with the specified <code>value</code> . |
| <code>STRING(opt, len)</code> | VARCHAR2 | Returns a random string. |
| <code>TERMINATE</code> | n/a | <code>TERMINATE</code> has no effect. Deprecated, but supported for backward compatibility. |
| <code>VALUE</code> | NUMBER | Returns a random number with a value greater than or equal to <code>0</code> and less than <code>1</code> , with 38 digit precision. |
| <code>VALUE(low, high)</code> | NUMBER | Returns a random number with a value greater than or equal to <code>low</code> and less than <code>high</code> . |

INITIALIZE

The `INITIALIZE` procedure initializes the `DBMS_RANDOM` package with a seed value. The signature is:

```
INITIALIZE(<val> IN INTEGER)
```

This procedure should be considered deprecated; it is included for backward compatibility only.

Parameters

`val`

| `val` is the seed value used by the `DBMS_RANDOM` package algorithm.

Example

The following code snippet demonstrates a call to the `INITIALIZE` procedure that initializes the `DBMS_RANDOM` package with the seed value, `6475`.

```
DBMS_RANDOM.INITIALIZE(6475);
```

NORMAL

The `NORMAL` function returns a random number of type `NUMBER`. The signature is:

```
<result> NUMBER NORMAL()
```

Parameters

`result`

| `result` is a random value of type `NUMBER`.

Example

The following code snippet demonstrates a call to the `NORMAL` function:

```
x:= DBMS_RANDOM.NORMAL();
```

RANDOM

The `RANDOM` function returns a random `INTEGER` value that is greater than or equal to -2^{31} and less than 2^{31} . The signature is:

```
<result> INTEGER RANDOM()
```

This function should be considered deprecated; it is included for backward compatibility only.

Parameters

`result`

| `result` is a random value of type `INTEGER`.

Example

The following code snippet demonstrates a call to the `RANDOM` function. The call returns a random number:

```
x := DBMS_RANDOM.RANDOM();
```

SEED

The first form of the `SEED` procedure resets the seed value for the `DBMS_RANDOM` package with an `INTEGER` value. The `SEED` procedure is available in two forms; the signature of the first form is:

```
SEED(<val> IN INTEGER)
```

Parameters

`val`

`val` is the seed value used by the `DBMS_RANDOM` package algorithm.

Example

The following code snippet demonstrates a call to the `SEED` procedure; the call sets the seed value at `8495`.

```
DBMS_RANDOM.SEED(8495);
```

SEED

The second form of the `SEED` procedure resets the seed value for the `DBMS_RANDOM` package with a string value. The `SEED` procedure is available in two forms; the signature of the second form is:

```
SEED(<val> IN VARCHAR2)
```

Parameters

`val`

`val` is the seed value used by the `DBMS_RANDOM` package algorithm.

Example

The following code snippet demonstrates a call to the `SEED` procedure; the call sets the seed value to `abc123`.

```
DBMS_RANDOM.SEED('abc123');
```

STRING

The `STRING` function returns a random `VARCHAR2` string in a user-specified format. The signature of the `STRING` function is:

```
<result> VARCHAR2 STRING(<opt> IN CHAR, <len> IN NUMBER)
```

Parameters

`opt`

Formatting option for the returned string. `option` may be:

| Option | Specifies Formatting Option |
|----------------------------------|-----------------------------|
| <code>U</code> or <code>U</code> | Uppercase alpha string |
| <code>L</code> or <code>L</code> | Lowercase alpha string |

| Option | Specifies Formatting Option |
|--------|--------------------------------|
| a or A | Mixed case string |
| x or X | Uppercase alpha-numeric string |
| p or P | Any printable characters |

len

The length of the returned string.

result

result is a random value of type VARCHAR2.

Example

The following code snippet demonstrates a call to the STRING function; the call returns a random alpha-numeric character string that is 10 characters long.

```
x := DBMS_RANDOM.STRING('X', 10);
```

TERMINATE

The TERMINATE procedure has no effect. The signature is:

TERMINATE

The TERMINATE procedure should be considered deprecated; the procedure is supported for compatibility only.

VALUE

The VALUE function returns a random NUMBER that is greater than or equal to 0, and less than 1, with 38 digit precision. The VALUE function has two forms; the signature of the first form is:

```
<result> NUMBER VALUE()
```

Parameters

result

result is a random value of type NUMBER.

Example

The following code snippet demonstrates a call to the VALUE function. The call returns a random NUMBER:

```
x := DBMS_RANDOM.VALUE();
```

VALUE

The VALUE function returns a random NUMBER with a value that is between user-specified boundaries. The VALUE function has two forms; the signature of the second form is:

<result> NUMBER VALUE(<low> IN NUMBER, <high> IN NUMBER)

Parameters

`low`

| `low` specifies the lower boundary for the random value. The random value may be equal to `low`.

`high`

| `high` specifies the upper boundary for the random value; the random value will be less than `high`.

`result`

| `result` is a random value of type `NUMBER`.

Example

The following code snippet demonstrates a call to the `VALUE` function. The call returns a random `NUMBER` with a value that is greater than or equal to 1 and less than 100:

```
x := DBMS_RANDOM.VALUE(1, 100);
```

6.3.13 DBMS_REDACT

The `DBMS REDACT` package enables the redacting or masking of data returned by a query. The `DBMS_REDACT` package provides a procedure to create policies, alter policies, enable policies, disable policies, and drop policies. The procedures available in the `DBMS_REDACT` package are listed in the following table.

| Function/Procedure | Function or Procedure | Return Type | Description |
|--|-----------------------|-------------|--|
| <code>ADD_POLICY(object schema, object name, policy name, policy description, column name, column description, function_type, function parameters, expression, enable, regexp pattern, regexp replace string, regexp position, regexp occurrence, regexp_match_parameter, custom_function_expression)</code> | Procedure | n/a | Adds a data redaction policy. |
| <code>ALTER_POLICY(object schema, object name, policy name, action, column name, function type, function parameters, expression, regexp pattern, regexp replace string, regexp position, regexp occurrence, regexp match parameter, policy_description, column_description, custom_function_expression)</code> | Procedure | n/a | Alters the existing data redaction policy. |
| <code>DISABLE_POLICY(object_schema, object_name, policy_name)</code> | Procedure | n/a | Disables the existing data redaction policy. |

| | | | |
|---|-----------|-----|---|
| <code>ENABLE_POLICY(object_schema, object_name, policy_name)</code> | Procedure | n/a | Enables a previously disabled data redaction policy. |
| <code>DROP_POLICY(object_schema, object_name, policy_name)</code> | Procedure | n/a | Drops a data redaction policy. |
| <code>UPDATE FULL REDACTION VALUES(number_val, binfloat_val, bindouble_val, char_val, varchar_val, nchar_val, nvarchar_val, datecol_val, ts_val, tswtz_val, blob_val, clob_val, nclob_val)</code> | Procedure | n/a | Updates the full redaction default values for the specified datatype. |

The data redaction feature uses the `DBMS_REDACT` package to define policies or conditions to redact data in a column based on the table column type and redaction type.

Note that you must be the owner of the table to create or change the data redaction policies. The users are exempted from all the column redaction policies, which the table owner or super-user is by default.

Using DBMS_REDACT Constants and Function Parameters

The `DBMS_REDACT` package uses the constants and redacts the column data by using any one of the data redaction types. The redaction type can be decided based on the `function_type` parameter of `dbms_redact.add_policy` and `dbms_redact.alter_policy` procedure. The below table highlights the values for `function_type` parameters of `dbms_redact.add_policy` and `dbms_redact.alter_policy`.

| Constant | Type | Value | Description |
|----------------------|----------------------|-------|---|
| <code>NONE</code> | <code>INTEGER</code> | 0 | No redaction, zero effect on the result of a query against table. |
| <code>FULL</code> | <code>INTEGER</code> | 1 | Full redaction, redacts full values of the column data. |
| <code>PARTIAL</code> | <code>INTEGER</code> | 2 | Partial redaction, redacts a portion of the column data. |
| <code>RANDOM</code> | <code>INTEGER</code> | 4 | Random redaction, each query results in a different random value depending on the datatype of the column. |
| <code>REGEXP</code> | <code>INTEGER</code> | 5 | Regular Expression based redaction, searches for the pattern of data to redact. |
| <code>CUSTOM</code> | <code>INTEGER</code> | 99 | Custom redaction type. |

The following table shows the values for the `action` parameter of `dbms_redact.alter_policy`.

| Constant | Type | Value | Description |
|-------------------------------------|----------------------|-------|--|
| <code>ADD_COLUMN</code> | <code>INTEGER</code> | 1 | Adds a column to the redaction policy. |
| <code>DROP_COLUMN</code> | <code>INTEGER</code> | 2 | Drops a column from the redaction policy. |
| <code>MODIFY_EXPRESSION</code> | <code>INTEGER</code> | 3 | Modifies the expression of a redaction policy. The redaction is applied when the expression evaluates to the <code>BOOLEAN</code> value to <code>TRUE</code> . |
| <code>MODIFY_COLUMN</code> | <code>INTEGER</code> | 4 | Modifies a column in the redaction policy to change the redaction function type or function parameter. |
| <code>SET_POLICY_DESCRIPTION</code> | <code>INTEGER</code> | 5 | Sets the redaction policy description. |

| SET_COLUMN_DESCRIPTION | INTEGER | 6 | Sets a description for the redaction performed on the column. |
|---|---|---|---|
| The partial data redaction enables you to redact only a portion of the column data. To use partial redaction, you must set the <code>dbms_redact.add_policy</code> procedure <code>function_type</code> parameter to <code>dbms_redact.partial</code> and use the <code>function_parameters</code> parameter to specify the partial redaction behavior. | | | |
| The data redaction feature provides a predefined format to configure policies that use the following datatype: | | | |
| <ul style="list-style-type: none"> • Character • Number • Datetime | | | |
| The following table highlights the format descriptor for partial redaction with respect to datatype. The example described below shows how to perform a redaction for a string datatype (in this scenario, a Social Security Number (SSN)), a <code>Number</code> datatype, and a <code>DATE</code> datatype. | | | |
| Datatype | Format Descriptor | Description | Examples |
| Character | <code>REDACT_PARTIAL_INPUT_FORMAT</code> | Specifies the input format. Enter <code>V</code> for each character from the input string to be possibly redacted. Enter <code>F</code> for each character from the input string that can be considered as a separator such as blank spaces or hyphens. | Consider ' <code>VVVFVVFVVVV,VVV-VVV,X,1,5</code> ' for masking first 5 digits of SSN strings such as <code>123-45-6789</code> , adding hyphen to format it and thereby resulting in strings such as <code>XXX-6789</code> . The field value <code>VVVFVVFVVVV</code> for matching SSN strings such as <code>123-45-6789</code> . |
| | <code>REDACT_PARTIAL_OUTPUT_FORMAT</code> | Specifies the output format. Enter <code>V</code> for each character from the input string to be possibly redacted. Replace each <code>F</code> character from the input format with a character such as a hyphen or any other separator. | The field value <code>VVV-VV-VVVV</code> can be used to redact SSN strings into <code>XXX-XX-6789</code> where <code>X</code> comes from <code>REDACT_PARTIAL_MASKCHAR</code> field. |
| | <code>REDACT_PARTIAL_MASKCHAR</code> | Specifies the character to be used for redaction. | The value <code>X</code> for redacting SSN strings into <code>XXX-XX-6789</code> . |
| | <code>REDACT_PARTIAL_MASKFROM</code> | Specifies which <code>V</code> within the input format from which to start the redaction. | The value <code>1</code> for redacting SSN strings starting at the first <code>V</code> of the input format of <code>VVVFVVFVVVV</code> into strings such as <code>XXX-XX-6789</code> . |
| | <code>REDACT_PARTIAL_MASKTO</code> | Specifies which <code>V</code> within the input format at which to end the redaction. | The value <code>5</code> for redacting SSN strings up to and including the fifth <code>V</code> within the input format of <code>VVVFVVFVVVV</code> into strings such as <code>XXX-XX-6789</code> . |

| | | | |
|----------|---|---|--|
| Number | <code>REDACT_PARTIAL_MASKCHAR</code> | Specifies the character to be displayed in the range between 0 and 9. | '9, 1, 5' for redacting the first five digits of the Social Security Number <code>123456789</code> into <code>999996789</code> . |
| | <code>REDACT_PARTIAL_MASKFROM</code> | Specifies the start digit position for redaction. | |
| | <code>REDACT_PARTIAL_MASKTO</code> | Specifies the end digit position for redaction. | |
| Datetime | <code>REDACT_PARTIAL_DATE_MONTH</code> | 'm' redacts the month. To mask a specific month, specify 'm#' where # indicates the month specified by its number between 1 and 12. | <code>m3</code> displays as March. |
| | <code>REDACT_PARTIAL_DATE_DAY</code> | 'd' redacts the day of the month. To mask with a day of the month, append 1-31 to a lowercase d. | <code>d3</code> displays as 03. |
| | <code>REDACT_PARTIAL_DATE_YEAR</code> | 'y' redacts the year. To mask with a year, append 1-9999 to a lowercase y. | <code>y1960</code> displays as 60. |
| | <code>REDACT_PARTIAL_DATE_HOUR</code> | 'h' redacts the hour. To mask with an hour, append 0-23 to a lowercase h. | <code>h18</code> displays as 18. |
| | <code>REDACT_PARTIAL_DATE_MINUTE</code> | 'm' redacts the minute. To mask with a minute, append 0-59 to a lowercase m. | <code>m20</code> displays as 20. |
| | <code>REDACT_PARTIAL_DATE_SECOND</code> | 's' redacts the second. To mask with a second, append 0-59 to a lowercase s. | <code>s40</code> displays as 40. |

The following table represents `function_parameters` values that can be used in partial redaction.

| Function Parameter | Data Type | Value | Description |
|--------------------|-----------|-------|-------------|
| | | | |

| | | |
|--------------------------|--|---|
| | | Redacts the first numbers of SS |
| REDACT_US_SSN_F5 | VARCHAR2 'VVVFVVFVVVV,VVV-VV-VVVV,X,1,5' | Example: The number 123-456789 becomes XX-6789. |
| | | Redacts the last numbers of SS |
| REDACT_US_SSN_L4 | VARCHAR2 'VVVFVVFVVVV,VVV-VV-VVVV,X,6,9' | Example: The number 123-456789 becomes 45-XXXX. |
| | | Redacts the entire SSN. |
| REDACT_US_SSN_ENTIRE | VARCHAR2 'VVVFVVFVVVV,VVV-VV-VVVV,X,1,9' | Example: The number 123-456789 becomes XX-XXXX. |
| | | Redacts the first numbers of SS when the column datatype is integer |
| REDACT_NUM_US_SSN_F5 | VARCHAR2 '9,1,5' | Example: The number 12345 becomes 999996789. |
| | | Redacts the last numbers of SS when the column datatype is integer |
| REDACT_NUM_US_SSN_L4 | VARCHAR2 '9,6,9' | Example: The number 12345 becomes 123459999. |
| | | Redacts the entire SSN when the column is a numeric datatype. |
| REDACT_NUM_US_SSN_ENTIRE | VARCHAR2 '9,1,9' | Example: The number 12345 becomes 999999999. |

| | | | |
|---------------------------|----------|--|---|
| REDACT_ZIP_CODE | VARCHAR2 | 'VVVVV,VVVVV,X,1,5' | Redacts a 5 digit zip code. Example: 123 becomes XXX. |
| REDACT_NUM_ZIP_CODE | VARCHAR2 | '9,1,5' | Redacts a 5 digit zip code when the column is a numeric datatype. Example: 123 becomes 9999. |
| REDACT_CCN16_F12 | VARCHAR2 | 'VVVVFVVVVFVVVVFVVVV,VVVV-VVVV-VVVV,*1,12' | Redacts a 16 character credit card number and displays only 12 digits. Example: 123 5678 9000 235 becomes ****-2358. |
| REDACT_DATE_MILLENNIUM | VARCHAR2 | 'm1d1y2000' | Redacts a date that is in the DD-MMM-YYYY format. Example: Redacts all date to 01-JAN-2000. |
| REDACT_DATE_EPOCH | VARCHAR2 | 'm1d1y1970' | Redacts all dates to 01-JAN-70. |
| REDACT_AMEX_CCN_FORMATTED | VARCHAR2 | 'VVVVFVVVVVVFVVVV,VVVV-VVVV-VVVV,*1,10' | Redacts the American Express credit card number and replaces the first 4 digits with * except the last 5 digits. Example: The card number 12345678901234500 becomes ***1234500. |

| | | |
|------------------------|---|---|
| REDACT_AMEX_CCN_NUMBER | VARCHAR2 '0,1,10' | Redacts the American Express credit card number and replaces the first digit with 0 except the last 5 digits |
| REDACT_SIN_FORMATTED | VARCHAR2 'VVVFVVVFVVV,VVV-VVV-VVV,*1,6' | Example: The card number 1567890 34500 becomes 0000 000000 34500 |
| REDACT_SIN_NUMBER | VARCHAR2 '9,1,6' | Redacts the Social Insurance Number by replacing the first 6 digits by *. Example: 123456789 becomes ***-789. |
| REDACT_SIN_UNFORMATTED | VARCHAR2 'VVVVVVVVV,VVVVVVVVVV,*1,6' | Redacts the Social Insurance Number by replacing the first 6 digits by *. Example: 123456789 becomes *****. |
| REDACT_CCN_FORMATTED | VARCHAR2 'VVVVFVVVVFVVVVFVVVV,VVVV-VVVV-VVVV,*1,12' | Redacts a credit card number by * and displays only 4 digits. Example: The card number 15678-9000-467 becomes ****-****-4671. |

| | | |
|-----------------------------|--|--|
| REDACT_CCN_NUMBER | VARCHAR2 '9,1,12' | Redacts a credit card number by redacting the last 4 digits. |
| REDACT_NA_PHONE_FORMATTED | VARCHAR2 'VVVFVVVFVVVV,VVV-VVV-VVVV,X,4,10' | Example: The card number 123456789000 becomes 000000000000. |
| REDACT_NA_PHONE_NUMBER | VARCHAR2 '0,4,10' | Redacts the North American phone number by redacting the area code. |
| REDACT_NA_PHONE_UNFORMATTED | VARCHAR2 'VVVVVVVV,VVVVVVVV,X,4,10' | Example: 1234567890 becomes 1230000000. |
| REDACT_UK_NIN_FORMATTED | VARCHAR2 'VVFVVFVVFVVFV,VV VV VV VV VV,X,3,8' | Redacts the UK National Insurance Number by redacting leaving the alphabetic characters. |

| | | | |
|---|----------|---------------------------|---|
| REDACT_UK_NIN_UNFORMATTED | VARCHAR2 | 'VVVVVVVV,VVVVVVVV,X,3,8' | Redacts the UI National Insurance Number by leaving the alphabetic characters. |
| Example: NY220134D becomes NYXXXXXXD. | | | |

A regular expression-based redaction searches for patterns of data to redact. The `regexp_pattern` search the values in order for the `regexp_replace_string` to change the value. The following table illustrates the `regexp_pattern` values that you can use during `REGEXP` based redaction.

| Function Parameter and Description | Data Type | Value |
|--|-----------|---|
| RE_PATTERN_CC_L6_T4: Searches for the middle digits of a credit card number that includes 6 leading digits and 4 trailing digits. The <code>regexp_replace_string</code> setting to use with the format is <code>RE_REDACT_CC_MIDDLE_DIGITS</code> that replaces the identified pattern with the characters specified by the <code>RE_REDACT_CC_MIDDLE_DIGITS</code> parameter. | VARCHAR2 | '(\d\d\d\d\d)(\d\d\d*)(\d\d\d\d)' |
| RE_PATTERN_ANY_DIGIT: Searches for any digit and replaces the identified pattern with the characters specified by the following values of the <code>regexp_replace_string</code> parameter. <code>regexp_replace_string=></code> <code>RE_REDACT_WITH_SINGLE_X</code> (replaces any matched digit with the X character). <code>regexp_replace_string=></code> <code>RE_REDACT_WITH_SINGLE_1</code> (replaces any matched digit with the 1 character). | VARCHAR2 | '\d' |
| RE_PATTERN_US_PHONE: Searches for the U.S phone number and replaces the identified pattern with the characters specified by the <code>regexp_replace_string</code> parameter. <code>regexp_replace_string=></code> <code>RE_REDACT_US_PHONE_L7</code> (searches the phone number and then replaces the last 7 digits). | VARCHAR2 | '((\d\d\d) \d\d\d)-(\d\d\d)-(\d\d\d\d)' |

| Function Parameter and Description | Data Type | Value |
|---|-----------|--|
| RE_PATTERN_EMAIL_ADDRESS: Searches for the email address and replaces the identified pattern with the characters specified by the following values of the regexp_replace_string parameter. regexp_replace_string=> RE_REDACT_EMAIL_NAME (finds the email address and redacts the email username). regexp_replace_string=> RE_REDACT_EMAIL_DOMAIN (finds the email address and redacts the email domain). regexp_replace_string=> RE_REDACT_EMAIL_ENTIRE (finds the email address and redacts the entire email address). | VARCHAR2 | '([A-Za-z0-9._%+-]+)@([A-Za-z0-9.-]+.[A-Za-z]{2,4})' |
| RE_PATTERN_IP_ADDRESS: Searches for an IP address and replaces the identified pattern with the characters specified by the regexp_replace_string parameter. The regexp_replace_string parameter to be used is RE_REDACT_IP_L3 that replaces the last section of an IP address with 999 and indicates it is redacted. | VARCHAR2 | '(\d{1,3}\.\d{1,3}\.\d{1,3})\.\d{1,3}' |
| RE_PATTERN_AMEX_CCN: Searches for the American Express credit card number. The regexp_replace_string parameter to be used is RE_REDACT_AMEX_CCN that redacts all of the digits except the last 5. | VARCHAR2 | '.*(\d\d\d\d)\$' |
| RE_PATTERN_CCN: Searches for the credit card number other than American Express credit cards. The regexp_replace_string parameter to be used is RE_REDACT_CCN that redacts all of the digits except the last 4. | VARCHAR2 | '.*(\d\d\d\d)\$' |

| Function Parameter and Description | Data Type | Value |
|---|-----------|------------------------------------|
| RE_PATTERN_US_SSN: Searches the SSN number and replaces the identified pattern with the characters specified by the regexp_replace_string parameter. \1-XXX-XXXX or 'XXX-XXX-\3' will return 123-XXX-XXXX or XXX-XXX-6789 for the value '123-45-6789' respectively. | VARCHAR2 | '(\d{2}\d)-(\d{2})-(\d{2}\d{2}\d)' |

The below table illustrates the `regexp_replace_string` values that you can use during `REGEXP` based redaction.

| Function Parameter | Data Type | Value | Description |
|----------------------------|-----------|-------------|--|
| RE_REDACT_CC_MIDDLE_DIGITS | VARCHAR2 | '1XXXXXX\3' | Redacts the middle digits of a credit card number according to the <code>regexp_pattern</code> parameter with the <code>RE_PATTERN_CC_L6_T4</code> format and replaces each redacted character with an X. Example: The credit card number 1234 5678 9000 2490 becomes 1234 56XX XXXX 2490. |
| RE_REDACT_WITH_SINGLE_X | VARCHAR2 | 'X' | Replaces the data with a single X character for each matching pattern as specified by setting the <code>regexp_pattern</code> parameter with the <code>RE_PATTERN_ANY_DIGIT</code> format. Example: The credit card number 1234 5678 9000 2490 becomes XXXX XXXX XXXX XXXX. |
| RE_REDACT_WITH_SINGLE_1 | VARCHAR2 | '1' | Replaces the data with a single 1 digit for each of the data digits as specified by setting the <code>regexp_pattern</code> parameter with the <code>RE_PATTERN_ANY_DIGIT</code> format. Example: The credit card number 1234 5678 9000 2490 becomes 1111 1111 1111 1111. |

| | | |
|------------------------|---------------------------|---|
| RE_REDACT_US_PHONE_L7 | VARCHAR2 '\1-XXX-XXXX' | Redacts the last 7 digits of U.S phone number according to the regexp_pattern parameter with the RE_PATTERN_US_PHONE format and replaces each redacted character with an X. |
| RE_REDACT_EMAIL_NAME | VARCHAR2 'xxxx@\12' | Example: The phone number 123-444-5900 becomes 123-XXX-XXXX. |
| RE_REDACT_EMAIL_DOMAIN | VARCHAR2 '\1@xxxxx.com' | Redacts the email name according to the regexp_pattern parameter with the RE_PATTERN_EMAIL_ADDRESS format and replaces the email username with the four x characters. |
| RE_REDACT_EMAIL_ENTIRE | VARCHAR2 'xxxx@xxxxx.com' | Example: The email address sjohn@example.com becomes xxxx@example.com. |
| RE_REDACT_IP_L3 | VARCHAR2 '\1.999' | Redacts the entire email address according to the regexp_pattern parameter with the RE_PATTERN_EMAIL_ADDRESS format and replaces the email address with the x characters. |
| | | Example: The email address sjohn@example.com becomes xxxx@xxxxx.com. |
| | | Redacts the last 3 digits of an IP address according to the regexp_pattern parameter with the RE_PATTERN_IP_ADDRESS format. |
| | | Example: The IP address 172.0.1.258 becomes 172.0.1.999, which is an invalid IP address. |

| | | |
|--------------------|--------------------|---|
| RE_REDACT_AMEX_CCN | VARCHAR2 '*****\1' | Redacts the first 10 digits of an American Express credit card number according to the regexp_pattern parameter with the RE_PATTERN_AMEX_CCN format. Example: 123456789062816 becomes *****62816. |
| RE_REDACT_CCN | VARCHAR2 '*****\1' | Redacts the first 12 digits of a credit card number as specified by the regexp_pattern parameter with the RE_PATTERN_CCN format. Example: 8749012678345671 becomes *****5671. |

The following tables show the `regexp_position` value and `regexp_occurrence` values that you can use during `REGEXP` based redaction.

| Function Parameter | Data Type | Value | Description |
|--------------------|-----------|-------|---|
| RE_BEGINNING | INTEGER | 1 | Specifies the position of a character where search must begin. By default, the value is 1 that indicates the search begins at the first character of source_char. |

| Function Parameter | Data Type | Value | Description |
|--------------------|-----------|-------|--|
| RE_ALL | INTEGER | 0 | Specifies the replacement occurrence of a substring. If the value is 0, then the replacement of each matching substring occurs. |
| RE_FIRST | INTEGER | 1 | Specifies the replacement occurrence of a substring. If the value is 1, then the replacement of the first matching substring occurs. |

The following table shows the `regexp_match_parameter` values that you can use during `REGEXP` based redaction which lets you change the default matching behavior of a function.

| Function Parameter | Data Type | Value | Description |
|----------------------|-----------|-------|---|
| RE_CASE_SENSITIVE | VARCHAR2 | 'c' | Specifies the case-sensitive matching. |
| RE_CASE_INSENSITIVE | VARCHAR2 | 'i' | Specifies the case-insensitive matching. |
| RE_MULTIPLE_LINES | VARCHAR2 | 'm' | Treats the source string as multiple lines but if you omit this parameter, then it indicates as a single line. |
| RE_NEWLINE_WILDCARD | VARCHAR2 | 'n' | Specifies the period (.), but if you omit this parameter, then the period does not match the newline character. |
| RE_IGNORE_WHITESPACE | VARCHAR2 | 'x' | Ignores the whitespace characters. |

Note

If you create a redaction policy based on a numeric type column, then make sure that the result after redaction is a number and accordingly set the replacement string to avoid runtime errors.

Note

If you create a redaction policy based on a character type column, then make sure that a length of the result after redaction is compatible with the column type and accordingly set the replacement string to avoid runtime errors.

ADD_POLICY

The `add_policy` procedure creates a new data redaction policy for a table.

```
PROCEDURE add_policy (
    <object_schema>          IN VARCHAR2 DEFAULT NULL,
    <object_name>            IN VARCHAR2,
    <policy_name>             IN VARCHAR2,
    <policy_description>      IN VARCHAR2 DEFAULT NULL,
    <column_name>             IN VARCHAR2 DEFAULT NULL,
    <column_description>      IN VARCHAR2 DEFAULT NULL,
    <function_type>           IN INTEGER DEFAULT DBMS_REDACT.FULL,
    <function_parameters>     IN VARCHAR2 DEFAULT NULL,
    <expression>              IN VARCHAR2,
    <enable>                 IN BOOLEAN DEFAULT TRUE,
    <regexp_pattern>          IN VARCHAR2 DEFAULT NULL,
    <regexp_replace_string>   IN VARCHAR2 DEFAULT NULL,
    <regexp_position>         IN INTEGER DEFAULT DBMS_REDACT.RE_BEGINNING,
    <regexp_occurrence>       IN INTEGER DEFAULT DBMS_REDACT.RE_ALL,
    <regexp_match_parameter>  IN VARCHAR2 DEFAULT NULL,
    <custom_function_expression> IN VARCHAR2 DEFAULT NULL
)
```

Parameters

`object_schema`

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

`object_name`

Name of the table on which the data redaction policy is created.

`policy_name`

Name of the policy to be added. Ensure that the `policy_name` is unique for the table on which the policy is created.

`policy_description`

Specify the description of a redaction policy.

`column_name`

Name of the column to which the redaction policy applies. To redact more than one column, use the `alter_policy` procedure to add additional columns.

`column_description`

Description of the column to be redacted. The `column_description` is not supported, but if you specify the description for a column then, you will get a warning message.

`function_type`

The type of redaction function to be used. The possible values are `NONE`, `FULL`, `PARTIAL`, `RANDOM`, `REGEXP`, and `CUSTOM`.

`function_parameters`

Specifies the function parameters for the partition redaction and is applicable only for partial redaction.

expression

Specifies the Boolean expression for the table and determines how the policy is to be applied. The redaction occurs if this policy expression is evaluated to `TRUE`.

enable

When set to `TRUE`, the policy is enabled upon creation. The default is set as `TRUE`. When set to `FALSE`, the policy is disabled but the policy can be enabled by calling the `enable_policy` procedure.

regexp_pattern

Specifies the regular expression pattern to redact data. If the `regexp_pattern` does not match, then the `NULL` value is returned.

regexp_replace_string

Specifies the replacement string value.

regexp_position

Specifies the position of a character where search must begin. By default, the function parameter is `RE_BEGINNING`.

regexp_occurrence

Specifies the replacement occurrence of a substring. If the constant is `RE_ALL`, then the replacement of each matching substring occurs. If the constant is `RE_FIRST`, then the replacement of the first matching substring occurs.

regexp_match_parameter

Changes the default matching behavior of a function. The possible `regexp_match_parameter` constants can be '`RE_CASE_SENSITIVE`', '`RE_CASE_INSENSITIVE`', '`RE_MULTIPLE_LINES`', '`RE_NEWLINE_WILDCARD`', '`RE_IGNORE_WHITESPACE`'.

Note

For more information on `constants`, `function_parameters`, or `regexp` (regular expressions) see, Using `DBMS_REDACT` Constants and Function Parameters.

custom_function_expression

The `custom_function_expression` is applicable only for the `CUSTOM` redaction type. The `custom_function_expression` is a function expression that is, schema-qualified function with a parameter such as `schema_name.function_name(argument1, ...)` that allows a user to use their redaction logic to redact the column data.

Example

The following example illustrates how to create a policy and use full redaction for values in the `payment_details_tab` table `customer_id` column.

```
edb=# CREATE TABLE payment_details_tab (
customer_id NUMBER      NOT NULL,
card_string VARCHAR2(19) NOT NULL);
CREATE TABLE

edb=# BEGIN
INSERT INTO payment_details_tab VALUES (4000, '1234-1234-1234-1234');
```

```
INSERT INTO payment_details_tab VALUES (4001, '2345-2345-2345-2345');
END;
```

EDB-SPL Procedure successfully completed

```
edb=# CREATE USER redact_user;
CREATE ROLE
edb=# GRANT SELECT ON payment_details_tab TO redact_user;
GRANT
```

```
\c edb base_user
```

```
BEGIN
DBMS_REDACT.add_policy(
    object_schema      => 'public',
    object_name        => 'payment_details_tab',
    policy_name        => 'redactPolicy_001',
    policy_description => 'redactPolicy_001 for payment_details_tab
table',
    column_name        => 'customer_id',
    function_type       => DBMS_REDACT.full,
    expression          => '1=1',
    enable              => TRUE);
END;
```

Redacted Result:

```
edb=# \c edb redact_user
You are now connected to database "edb" as user "redact_user".
edb=> select customer_id from payment_details_tab order by 1;
customer_id
-----
0
0
(2 rows)
```

ALTER_POLICY

The `alter_policy` procedure alters or modifies an existing data redaction policy for a table.

```
PROCEDURE alter_policy (
<object_schema>           IN VARCHAR2 DEFAULT NULL,
<object_name>             IN VARCHAR2,
<policy_name>             IN VARCHAR2,
<action>                  IN INTEGER DEFAULT DBMS_REDACT.ADD_COLUMN,
<column_name>             IN VARCHAR2 DEFAULT NULL,
<function_type>            IN INTEGER DEFAULT DBMS_REDACT.FULL,
<function_parameters>     IN VARCHAR2 DEFAULT NULL,
<expression>               IN VARCHAR2 DEFAULT NULL,
<regexp_pattern>          IN VARCHAR2 DEFAULT NULL,
<regexp_replace_string>   IN VARCHAR2 DEFAULT NULL,
<regexp_position>         IN INTEGER DEFAULT DBMS_REDACT.RE_BEGINNING,
<regexp_occurrence>       IN INTEGER DEFAULT DBMS_REDACT.RE_ALL,
```

```
<regexp_match_parameter>    IN VARCHAR2 DEFAULT NULL,
<policy_description>       IN VARCHAR2 DEFAULT NULL,
<column_description>       IN VARCHAR2 DEFAULT NULL,
<custom_function_expression> IN VARCHAR2 DEFAULT NULL
)
```

Parameters

object_schema

Specifies the name of the schema in which the object resides and on which the data redaction policy will be altered. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

object_name

Name of the table to which to alter a data redaction policy.

policy_name

Name of the policy to be altered.

action

The action to perform. For more information about action parameters see, [DBMS_REDACT Constants and Function Parameters](#).

column_name

Name of the column to which the redaction policy applies.

function_type

The type of redaction function to be used. The possible values are `NONE`, `FULL`, `PARTIAL`, `RANDOM`, `REGEXP`, and `CUSTOM`.

function_parameters

Specifies the function parameters for the redaction function.

expression

Specifies the Boolean expression for the table and determines how the policy is to be applied. The redaction occurs if this policy expression is evaluated to `TRUE`.

regexp_pattern

Enables the use of regular expressions to redact data. If the `regexp_pattern` does not match the data, then the `NULL` value is returned.

regexp_replace_string

Specifies the replacement string value.

regexp_position

Specifies the position of a character where search must begin. By default, the function parameter is `RE_BEGINNING`.

regexp_occurrence

Specifies the replacement occurrence of a substring. If the constant is `RE_ALL`, then the replacement of each matching substring occurs. If the constant is `RE_FIRST`, then the replacement of the first matching

substring occurs.

regexp_match_parameter

Changes the default matching behavior of a function. The possible regexp_match_parameter constants can be 'RE_CASE_SENSITIVE', 'RE_CASE_INSENSITIVE', 'RE_MULTIPLE_LINES', 'RE_NEWLINE_WILDCARD', 'RE_IGNORE_WHITESPACE'.

Note

For more information on constants, function parameters, or regexp (regular expressions) see, [Using DBMS_REDACT Constants and Function Parameters](#).

policy_description

Specify the description of a redaction policy.

column_description

Description of the column to be redacted. The column_description is not supported, but if you specify the description for a column then, you will get a warning message.

custom_function_expression

The custom function expression is applicable only for the CUSTOM redaction type. The custom_function_expression is a function expression that is, schema-qualified function with a parameter such as schema_name.function_name(argument1, ...) that allows a user to use their redaction logic to redact the column data.

Example

The following example illustrates to alter a policy partial redaction for values in the payment_details_tab table card_string (usually a credit card number) column.

```
\c edb base_user

BEGIN
DBMS_REDACT.alter_policy (
object_schema      => 'public',
object_name        => 'payment_details_tab',
policy_name        => 'redactPolicy_001',
action             => DBMS_REDACT.ADD_COLUMN,
column_name        => 'card_string',
function_type      => DBMS_REDACT.partial,
function_parameters => DBMS_REDACT.REDACT_CCN16_F12);
END;
```

Redacted Result:

```
edb=# \c - redact_user
You are now connected to database "edb" as user "redact_user".
edb=> SELECT * FROM payment_details_tab;
customer_id | card_string
-----+-----
 0 | ****-****-****-1234
 0 | ****-****-****-2345
(2 rows)
```

DISABLE_POLICY

The `disable_policy` procedure disables an existing data redaction policy.

```
PROCEDURE disable_policy (
    <object_schema>      IN VARCHAR2 DEFAULT NULL,
    <object_name>        IN VARCHAR2,
    <policy_name>        IN VARCHAR2
)
```

Parameters

`object_schema`

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

`object_name`

Name of the table for which to disable a data redaction policy.

`policy_name`

Name of the policy to be disabled.

Example

The following example illustrates how to disable a policy.

```
\c edb base_user

BEGIN
    DBMS_REDACT.disable_policy(
        object_schema => 'public',
        object_name => 'payment_details_tab',
        policy_name => 'redactPolicy_001');
END;
```

Redacted Result: Data is no longer redacted after disabling a policy.

ENABLE_POLICY

The `enable_policy` procedure enables the previously disabled data redaction policy.

```
PROCEDURE enable_policy (
    <object_schema>      IN VARCHAR2 DEFAULT NULL,
    <object_name>        IN VARCHAR2,
    <policy_name>        IN VARCHAR2
)
```

Parameters

`object_schema`

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify `NULL` then the given object is searched by the order specified by `search_path` setting.

object_name

Name of the table to which to enable a data redaction policy.

policy_name

Name of the policy to be enabled.

Example

The following example illustrates how to enable a policy.

```
\c edb base_user
```

```
BEGIN
```

```
DBMS_REDACT.enable_policy(
    object_schema => 'public',
    object_name => 'payment_details_tab',
    policy_name => 'redactPolicy_001');
```

```
END;
```

Redacted Result: Data is redacted after enabling a policy.

DROP_POLICY

The **drop_policy** procedure drops a data redaction policy by removing the masking policy from a table.

```
PROCEDURE drop_policy (
    <object_schema>    IN VARCHAR2 DEFAULT NULL,
    <object_name>      IN VARCHAR2,
    <policy_name>      IN VARCHAR2
)
```

Parameters**object_schema**

Specifies the name of the schema in which the object resides and on which the data redaction policy will be applied. If you specify **NULL** then the given object is searched by the order specified by **search_path** setting.

object_name

Name of the table from which to drop a data redaction policy.

policy_name

Name of the policy to be dropped.

Example

The following example illustrates how to drop a policy.

```
\c edb base_user
```

```
BEGIN
```

```
DBMS_REDACT.drop_policy(
    object_schema => 'public',
```

```

object_name => 'payment_details_tab',
policy_name => 'redactPolicy_001');
END

```

Redacted Result: The server drops the specified policy.

UPDATE_FULL_REDACTION_VALUES

The `update_full_redaction_values` procedure updates the default displayed values for a data redaction policy and these default values can be viewed using the `redaction_values_for_type_full` view that use the full redaction type.

```

PROCEDURE update_full_redaction_values (
<number_val>      IN NUMBER          DEFAULT NULL,
<bfloat_val>      IN FLOAT4         DEFAULT NULL,
<bindouble_val>   IN FLOAT8         DEFAULT NULL,
<char_val>        IN CHAR           DEFAULT NULL,
<varchar_val>     IN VARCHAR2       DEFAULT NULL,
<nchar_val>       IN NCHAR          DEFAULT NULL,
<nvarchar_val>    IN NVARCHAR2      DEFAULT NULL,
<datecol_val>     IN DATE           DEFAULT NULL,
<ts_val>          IN TIMESTAMP      DEFAULT NULL,
<tswtz_val>       IN TIMESTAMPTZ    DEFAULT NULL,
<blob_val>        IN BLOB           DEFAULT NULL,
<clob_val>        IN CLOB           DEFAULT NULL,
<nclob_val>       IN CLOB           DEFAULT NULL
)

```

Parameters

`number_val`

Updates the default value for columns of the `NUMBER` datatype.

`bfloat_val`

The `FLOAT4` datatype is a random value. The binary float datatype is not supported.

`bindouble_val`

The `FLOAT8` datatype is a random value. The binary double datatype is not supported.

`char_val`

Updates the default value for columns of the `CHAR` datatype.

`varchar_val`

Updates the default value for columns of the `VARCHAR2` datatype.

`nchar_val`

The `nchar_val` is mapped to `CHAR` datatype and returns the `CHAR` value.

`nvarchar_val`

The `nvarchar_val` is mapped to `VARCHAR2` datatype and returns the `VARCHAR` value.

datecol_val

Updates the default value for columns of the **DATE** datatype.

ts_val

Updates the default value for columns of the **TIMESTAMP** datatype.

tswtz_val

Updates the default value for columns of the **TIMESTAMPTZ** datatype.

blob_val

Updates the default value for columns of the **BLOB** datatype.

clob_val

Updates the default value for columns of the **CLOB** datatype.

nclob_val

The **nclob_val** is mapped to **CLOB** datatype and returns the **CLOB** value.

Example

The following example illustrates how to update the full redaction values but before updating the values, you can:

1. View the default values using **redaction_values_for_type_full** view as shown below:

```
edb=# \x
Expanded display is on.
edb=# SELECT number_value, char_value, varchar_value, date_value,
    timestamp_value, timestamp_with_time_zone_value, blob_value,
    clob_value
FROM redaction_values_for_type_full;
-[ RECORD 1 ]-----+
number_value          | 0
char_value            |
varchar_value         |
date_value            | 01-JAN-01 00:00:00
timestamp_value       | 01-JAN-01 01:00:00
timestamp_with_time_zone_value | 31-DEC-00 20:00:00 -05:00
blob_value            | \x5b72656461637465645d
clob_value            | [redacted]
(1 row)
```

1. Now, update the default values for full redaction type. The **NULL** values will be ignored.

```
\c edb base_user
```

```
edb=# BEGIN
```

```
DBMS_REDACT.update_full_redaction_values (
    number_val => 9999999,
    char_val => 'Z',
    varchar_val => 'V',
    datecol_val => to_date('17/10/2018', 'DD/MM/YYYY'),
    ts_val => to_timestamp('17/10/2018 11:12:13', 'DD/MM/YYYY HH24:MI:SS'),
    tswtz_val => NULL,
```

```

blob_val => 'NEW REDACTED VALUE',
clob_val => 'NEW REDACTED VALUE');
END;

```

1. You can now see the updated values using `redaction_values_for_type_full` view.

EDB-SPL Procedure successfully completed

```

edb=# SELECT number_value, char_value, varchar_value, date_value,
    timestamp_value, timestamp_with_time_zone_value, blob_value,
    clob_value
FROM redaction_values_for_type_full;
-[ RECORD 1 ]-----+
number_value      | 9999999
char_value        | Z
varchar_value     | V
date_value        | 17-OCT-18 00:00:00
timestamp_value   | 17-OCT-18 11:12:13
timestamp_with_time_zone_value | 31-DEC-00 20:00:00 -05:00
blob_value        | \x4e45572052454441435445442056414c5545
clob_value        | NEW REDACTED VALUE
(1 row)

```

Redacted Result:

```

edb=# \c edb redact_user
You are now connected to database "edb" as user "redact_user".

edb=> select * from payment_details_tab order by 1;
customer_id | card_string
-----+
 9999999 | V
 9999999 | V
(2 rows)

```

6.3.14 DBMS_RLS

The `DBMS_RLS` package enables the implementation of Virtual Private Database on certain Advanced Server database objects.

| Function/Procedure | Function or Procedure | Return Type | Description |
|--|-----------------------|-------------|---|
| <code>ADD_POLICY(object_schema, object_name, policy_name, function_schema, policy_function [, statement_types [, update_check [, enable [, static_policy [, policy_type [, long_predicate [, sec_relevant_cols [, sec_relevant_cols_opt]]]]]]]])</code> | Procedure | n/a | Add a security policy to a database object. |

| | | | |
|--|-----------|-----|--|
| DROP_POLICY(object_schema, object_name, policy_name) | Procedure | n/a | Remove a security policy from a database object. |
| ENABLE_POLICY(object_schema, object_name, policy_name, enable) | Procedure | n/a | Enable or disable a security policy. |

Advanced Server's implementation of `DBMS_RLS` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Virtual Private Database is a type of fine-grained access control using security policies. *Fine-grained access control* in Virtual Private Database means that access to data can be controlled down to specific rows as defined by the security policy.

The rules that encode a security policy are defined in a *policy function*, which is an SPL function with certain input parameters and return value. The *security policy* is the named association of the policy function to a particular database object, typically a table.

Note

In Advanced Server, the policy function can be written in any language supported by Advanced Server such as SQL, PL/pgSQL and SPL.

Note

The database objects currently supported by Advanced Server Virtual Private Database are tables. Policies cannot be applied to views or synonyms.

The advantages of using Virtual Private Database are the following:

- Provides a fine-grained level of security. Database object level privileges given by the `GRANT` command determine access privileges to the entire instance of a database object, while Virtual Private Database provides access control for the individual rows of a database object instance.
- A different security policy can be applied depending upon the type of SQL command (`INSERT`, `UPDATE`, `DELETE`, or `SELECT`).
- The security policy can vary dynamically for each applicable SQL command affecting the database object depending upon factors such as the session user of the application accessing the database object.
- Invocation of the security policy is transparent to all applications that access the database object and thus, individual applications do not have to be modified to apply the security policy.
- Once a security policy is enabled, it is not possible for any application (including new applications) to circumvent the security policy except by the system privilege noted by the following.
- Even superusers cannot circumvent the security policy except by the system privilege noted by the following.

Note

The only way security policies can be circumvented is if the `EXEMPT ACCESS POLICY` system privilege has been granted to a user. The `EXEMPT ACCESS POLICY` privilege should be granted with extreme care as a user with this privilege is exempted from all policies in the database.

The `DBMS_RLS` package provides procedures to create policies, remove policies, enable policies, and disable policies.

The process for implementing Virtual Private Database is as follows:

- Create a policy function. The function must have two input parameters of type `VARCHAR2`. The first input parameter is for the schema containing the database object to which the policy is to apply and the second input parameter is for the name of that database object. The function must have a `VARCHAR2` return type. The function must return a string in the form of a `WHERE` clause predicate. This predicate is dynamically

appended as an `AND` condition to the SQL command that acts upon the database object. Thus, rows that do not satisfy the policy function predicate are filtered out from the SQL command result set.

- Use the `ADD_POLICY` procedure to define a new policy, which is the association of a policy function with a database object. With the `ADD_POLICY` procedure, you can also specify the types of SQL commands (`INSERT`, `UPDATE`, `DELETE`, or `SELECT`) to which the policy is to apply, whether or not to enable the policy at the time of its creation, and if the policy should apply to newly inserted rows or the modified image of updated rows.
- Use the `ENABLE POLICY` procedure to disable or enable an existing policy.
- Use the `DROP_POLICY` procedure to remove an existing policy. The `DROP_POLICY` procedure does not drop the policy function or the associated database object.

Once policies are created, they can be viewed in the catalog views, compatible with Oracle databases: `ALL_POLICIES`, `DBA_POLICIES`, or `USER_POLICIES`. The supported compatible views are listed in the *Database Compatibility for Oracle Developers Catalog Views Guide*, available at the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs/>

The `SYS_CONTEXT` function is often used with `DBMS_RLS`. The signature is:

`SYS_CONTEXT(<namespace>, <attribute>)`

Where:

`namespace` is a `VARCHAR2`; the only accepted value is `USERENV`. Any other value will return `NULL`.

`attribute` is a `VARCHAR2`. `attribute` may be:

| attribute Value | Equivalent Value |
|-----------------------------|--|
| <code>SESSION_USER</code> | <code>pg_catalog.session_user</code> |
| <code>CURRENT_USER</code> | <code>pg_catalog.current_user</code> |
| <code>CURRENT_SCHEMA</code> | <code>pg_catalog.current_schema</code> |
| <code>HOST</code> | <code>pg_catalog.inet_host</code> |
| <code>IP_ADDRESS</code> | <code>pg_catalog.inet_client_addr</code> |
| <code>SERVER_HOST</code> | <code>pg_catalog.inet_server_addr</code> |

Note

The examples used to illustrate the `DBMS_RLS` package are based on a modified copy of the sample `emp` table provided with Advanced Server along with a role named `salesmgr` that is granted all privileges on the table. You can create the modified copy of the `emp` table named `vpemp` and the `salesmgr` role as shown by the following:

```
CREATE TABLE public.vpemp AS SELECT empno, ename, job, sal, comm, deptno
FROM emp;
ALTER TABLE vpemp ADD authid VARCHAR2(12);
UPDATE vpemp SET authid = 'researchmgr' WHERE deptno = 20;
UPDATE vpemp SET authid = 'salesmgr' WHERE deptno = 30;
SELECT * FROM vpemp;
```

| empno | ename | job | sal | comm | deptno | authid |
|-------|--------|-----------|---------|------|--------|-------------|
| 7782 | CLARK | MANAGER | 2450.00 | | 10 | |
| 7839 | KING | PRESIDENT | 5000.00 | | 10 | |
| 7934 | MILLER | CLERK | 1300.00 | | 10 | |
| 7369 | SMITH | CLERK | 800.00 | | 20 | researchmgr |
| 7566 | JONES | MANAGER | 2975.00 | | 20 | researchmgr |
| 7788 | SCOTT | ANALYST | 3000.00 | | 20 | researchmgr |
| 7876 | ADAMS | CLERK | 1100.00 | | 20 | researchmgr |

```

7902 | FORD | ANALYST | 3000.00 |      | 20 | researchmgr
7499 | ALLEN | SALESMAN | 1600.00 | 300.00 | 30 | salesmgr
7521 | WARD | SALESMAN | 1250.00 | 500.00 | 30 | salesmgr
7654 | MARTIN | SALESMAN | 1250.00 | 1400.00 | 30 | salesmgr
7698 | BLAKE | MANAGER | 2850.00 |      | 30 | salesmgr
7844 | TURNER | SALESMAN | 1500.00 | 0.00 | 30 | salesmgr
7900 | JAMES | CLERK | 950.00 |      | 30 | salesmgr
(14 rows)

```

```

CREATE ROLE salesmgr WITH LOGIN PASSWORD 'password';
GRANT ALL ON vpemp TO salesmgr;

```

ADD_POLICY

The `ADD_POLICY` procedure creates a new policy by associating a policy function with a database object.

You must be a superuser to execute this procedure.

```

ADD_POLICY(<object_schema> VARCHAR2, <object_name> VARCHAR2,
<policy_name> VARCHAR2, <function_schema> VARCHAR2,
<policy_function> VARCHAR2
[, <statement_types> VARCHAR2
[, <update_check> BOOLEAN
[, <enable> BOOLEAN
[, <static_policy> BOOLEAN
[, <policy_type> INTEGER
[, <long_predicate> BOOLEAN
[, <sec_relevant_cols> VARCHAR2
[, <sec_relevant_cols_opt> INTEGER ]]]]]])

```

Parameters

`object_schema`

Name of the schema containing the database object to which the policy is to be applied.

`object_name`

Name of the database object to which the policy is to be applied. A given database object may have more than one policy applied to it.

`policy_name`

Name assigned to the policy. The combination of database object (identified by `object_schema` and `object_name`) and policy name must be unique within the database.

`function_schema`

Name of the schema containing the policy function.

Note

The policy function may belong to a package in which case `function_schema` must contain the name of the schema in which the package is defined.

`policy_function`

Name of the SPL function that defines the rules of the security policy. The same function may be specified in more than one policy.

Note

The policy function may belong to a package in which case `policy_function` must also contain the package name in dot notation (that is, `package_name.function_name`).

`statement_types`

Comma-separated list of SQL commands to which the policy applies. Valid SQL commands are `INSERT`, `UPDATE`, `DELETE`, and `SELECT`. The default is `INSERT,UPDATE,DELETE,SELECT`.

Note

Advanced Server accepts `INDEX` as a statement type, but it is ignored. Policies are not applied to index operations in Advanced Server.

`update_check`

Applies to `INSERT` and `UPDATE` SQL commands only.

When set to `TRUE`, the policy is applied to newly inserted rows and to the modified image of updated rows. If any of the new or modified rows do not qualify according to the policy function predicate, then the `INSERT` or `UPDATE` command throws an exception and no rows are inserted or modified by the `INSERT` or `UPDATE` command.

When set to `FALSE`, the policy is not applied to newly inserted rows or the modified image of updated rows. Thus, a newly inserted row may not appear in the result set of a subsequent SQL command that invokes the same policy. Similarly, rows which qualified according to the policy prior to an `UPDATE` command may not appear in the result set of a subsequent SQL command that invokes the same policy.

The default is `FALSE`.

`enable`

When set to `TRUE`, the policy is enabled and applied to the SQL commands given by the `statement_types` parameter. When set to `FALSE` the policy is disabled and not applied to any SQL commands. The policy can be enabled using the `ENABLE_POLICY` procedure. The default is `TRUE`.

`static_policy`

In Oracle, when set to `TRUE`, the policy is *static*, which means the policy function is evaluated once per database object the first time it is invoked by a policy on that database object. The resulting policy function predicate string is saved in memory and reused for all invocations of that policy on that database object while the database server instance is running.

When set to `FALSE`, the policy is *dynamic*, which means the policy function is re-evaluated and the policy function predicate string regenerated for all invocations of the policy.

The default is `FALSE`.

Note

In Oracle 10g, the `policy_type` parameter was introduced, which is intended to replace the `static_policy` parameter. In Oracle, if the `policy_type` parameter is not set to its default value of `NULL`, the `policy_type` parameter setting overrides the `static_policy` setting.

Note

The setting of `static_policy` is ignored by Advanced Server. Advanced Server implements only the dynamic policy, regardless of the setting of the `static_policy` parameter.

policy_type

In Oracle, determines when the policy function is re-evaluated, and hence, if and when the predicate string returned by the policy function changes. The default is `NULL`.

Note

The setting of this parameter is ignored by Advanced Server. Advanced Server always assumes a dynamic policy.

long_predicate

In Oracle, allows predicates up to 32K bytes if set to `TRUE`, otherwise predicates are limited to 4000 bytes. The default is `FALSE`.

Note

The setting of this parameter is ignored by Advanced Server. An Advanced Server policy function can return a predicate of unlimited length for all practical purposes.

sec_relevant_cols

Comma-separated list of columns of `object_name`. Provides *column-level Virtual Private Database* for the listed columns. The policy is enforced if any of the listed columns are referenced in a SQL command of a type listed in `statement_types`. The policy is not enforced if no such columns are referenced.

The default is `NULL`, which has the same effect as if all of the database object's columns were included in `sec_relevant_cols`.

sec_relevant_cols_opt

In Oracle, if `sec_relevant_cols_opt` is set to `DBMS_RLS.ALL_ROWS` (INTEGER constant of value 1), then the columns listed in `sec_relevant_cols` return `NULL` on all rows where the applied policy predicate is false. (If `sec_relevant_cols_opt` is not set to `DBMS_RLS.ALL_ROWS`, these rows would not be returned at all in the result set.) The default is `NULL`.

Note

Advanced Server does not support the `DBMS_RLS.ALL_ROWS` functionality. Advanced Server throws an error if `sec_relevant_cols_opt` is set to `DBMS_RLS.ALL_ROWS` (INTEGER value of 1).

Examples

This example uses the following policy function:

```
CREATE OR REPLACE FUNCTION verify_session_user (
    p_schema      VARCHAR2,
    p_object      VARCHAR2
)
RETURN VARCHAR2
IS
BEGIN
    RETURN 'authid = SYS_CONTEXT("USERENV", "SESSION_USER")';
END;
```

This function generates the predicate `authid = SYS_CONTEXT('USERENV', 'SESSION_USER')`, which is added to the `WHERE` clause of any SQL command of the type specified in the `ADD_POLICY` procedure.

This limits the effect of the SQL command to those rows where the content of the `authid` column is the same as the session user.

Note

This example uses the `SYS_CONTEXT` function to return the login user name. In Oracle the `SYS_CONTEXT` function is used to return attributes of an *application context*. The first parameter of the `SYS_CONTEXT` function is the name of an application context while the second parameter is the name of an attribute set within the application context. `USERENV` is a special built-in namespace that describes the current session. Advanced Server does not support application contexts, but only this specific usage of the `SYS_CONTEXT` function.

The following anonymous block calls the `ADD_POLICY` procedure to create a policy named `secure_update` to be applied to the `vpemp` table using function `verify_session_user` whenever an `INSERT`, `UPDATE`, or `DELETE` SQL command is given referencing the `vpemp` table.

```

DECLARE
    v_object_schema      VARCHAR2(30) := 'public';
    v_object_name        VARCHAR2(30) := 'vpemp';
    v_policy_name        VARCHAR2(30) := 'secure_update';
    v_function_schema    VARCHAR2(30) := 'enterprisedb';
    v_policy_function    VARCHAR2(30) := 'verify_session_user';
    v_statement_types    VARCHAR2(30) := 'INSERT,UPDATE,DELETE';
    v_update_check       BOOLEAN    := TRUE;
    v_enable             BOOLEAN    := TRUE;
BEGIN
    DBMS_RLS.ADD_POLICY(
        v_object_schema,
        v_object_name,
        v_policy_name,
        v_function_schema,
        v_policy_function,
        v_statement_types,
        v_update_check,
        v_enable
    );
END;

```

After successful creation of the policy, a terminal session is started by user `salesmgr`. The following query shows the content of the `vpemp` table:

```

edb=# \c edb salesmgr
Password for user salesmgr:
You are now connected to database "edb" as user "salesmgr".
edb=> SELECT * FROM vpemp;
empno | ename | job | sal | comm | deptno | authid
-----+-----+-----+-----+-----+
7782 | CLARK | MANAGER | 2450.00 |   | 10 |
7839 | KING | PRESIDENT | 5000.00 |   | 10 |
7934 | MILLER | CLERK | 1300.00 |   | 10 |
7369 | SMITH | CLERK | 800.00 |   | 20 | researchmgr
7566 | JONES | MANAGER | 2975.00 |   | 20 | researchmgr
7788 | SCOTT | ANALYST | 3000.00 |   | 20 | researchmgr
7876 | ADAMS | CLERK | 1100.00 |   | 20 | researchmgr
7902 | FORD | ANALYST | 3000.00 |   | 20 | researchmgr
7499 | ALLEN | SALESMAN | 1600.00 | 300.00 | 30 | salesmgr
7521 | WARD | SALESMAN | 1250.00 | 500.00 | 30 | salesmgr
7654 | MARTIN | SALESMAN | 1250.00 | 1400.00 | 30 | salesmgr
7698 | BLAKE | MANAGER | 2850.00 |   | 30 | salesmgr
7844 | TURNER | SALESMAN | 1500.00 | 0.00 | 30 | salesmgr
7900 | JAMES | CLERK | 950.00 |   | 30 | salesmgr

```

(14 rows)

An unqualified `UPDATE` command (no `WHERE` clause) is issued by the `salesmgr` user:

```
edb=> UPDATE vpemp SET comm = sal * .75;
UPDATE 6
```

Instead of updating all rows in the table, the policy restricts the effect of the update to only those rows where the `authid` column contains the value `salesmgr` as specified by the policy function predicate `authid = SYS_CONTEXT('USERENV', 'SESSION_USER')`.

The following query shows that the `comm` column has been changed only for those rows where `authid` contains `salesmgr`. All other rows are unchanged.

```
edb=> SELECT * FROM vpemp;
empno | ename | job | sal | comm | deptno | authid
-----+-----+-----+-----+-----+
7782 | CLARK | MANAGER | 2450.00 |    | 10 |
7839 | KING  | PRESIDENT | 5000.00 |    | 10 |
7934 | MILLER | CLERK | 1300.00 |    | 10 |
7369 | SMITH | CLERK | 800.00 |    | 20 | researchmgr
7566 | JONES | MANAGER | 2975.00 |    | 20 | researchmgr
7788 | SCOTT | ANALYST | 3000.00 |    | 20 | researchmgr
7876 | ADAMS | CLERK | 1100.00 |    | 20 | researchmgr
7902 | FORD  | ANALYST | 3000.00 |    | 20 | researchmgr
7499 | ALLEN | SALESMAN | 1600.00 | 1200.00 | 30 | salesmgr
7521 | WARD  | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
7654 | MARTIN | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
7698 | BLAKE  | MANAGER | 2850.00 | 2137.50 | 30 | salesmgr
7844 | TURNER | SALESMAN | 1500.00 | 1125.00 | 30 | salesmgr
7900 | JAMES  | CLERK | 950.00 | 712.50 | 30 | salesmgr
```

(14 rows)

Furthermore, since the `update_check` parameter was set to `TRUE` in the `ADD_POLICY` procedure, the following `INSERT` command throws an exception since the value given for the `authid` column, `researchmgr`, does not match the session user, which is `salesmgr`, and hence, fails the policy.

```
edb=> INSERT INTO vpemp VALUES (9001,'SMITH','ANALYST',3200.00,NULL,20,
'researchmgr');
ERROR: policy with check option violation
DETAIL: Policy predicate was evaluated to FALSE with the updated values
```

If `update_check` was set to `FALSE`, the preceding `INSERT` command would have succeeded.

The following example illustrates the use of the `sec_relevant_cols` parameter to apply a policy only when certain columns are referenced in the SQL command. The following policy function is used for this example, which selects rows where the employee salary is less than `2000`.

```
CREATE OR REPLACE FUNCTION sal_lt_2000 (
  p_schema      VARCHAR2,
  p_object      VARCHAR2
)
RETURN VARCHAR2
IS
BEGIN
  RETURN 'sal < 2000';
END
```

The policy is created so that it is enforced only if a `SELECT` command includes columns `sal` or `comm`:

```

DECLARE
    v_object_schema      VARCHAR2(30) := 'public';
    v_object_name        VARCHAR2(30) := 'vpemp';
    v_policy_name        VARCHAR2(30) := 'secure_salary';
    v_function_schema    VARCHAR2(30) := 'enterprisedb';
    v_policy_function   VARCHAR2(30) := 'sal_lt_2000';
    v_statement_types    VARCHAR2(30) := 'SELECT';
    v_sec_relevant_cols VARCHAR2(30) := 'sal,comm';

BEGIN
    DBMS_RLS.ADD_POLICY(
        v_object_schema,
        v_object_name,
        v_policy_name,
        v_function_schema,
        v_policy_function,
        v_statement_types,
        sec_relevant_cols => v_sec_relevant_cols
    );
END;

```

If a query does not reference columns `sal` or `comm`, then the policy is not applied. The following query returns all 14 rows of table `vpemp`:

```

edb=# SELECT empno, ename, job, deptno, authid FROM vpemp;
empno | ename | job | deptno | authid
-----+-----+-----+
7782 | CLARK | MANAGER | 10 |
7839 | KING | PRESIDENT | 10 |
7934 | MILLER | CLERK | 10 |
7369 | SMITH | CLERK | 20 | researchmgr
7566 | JONES | MANAGER | 20 | researchmgr
7788 | SCOTT | ANALYST | 20 | researchmgr
7876 | ADAMS | CLERK | 20 | researchmgr
7902 | FORD | ANALYST | 20 | researchmgr
7499 | ALLEN | SALESMAN | 30 | salesmgr
7521 | WARD | SALESMAN | 30 | salesmgr
7654 | MARTIN | SALESMAN | 30 | salesmgr
7698 | BLAKE | MANAGER | 30 | salesmgr
7844 | TURNER | SALESMAN | 30 | salesmgr
7900 | JAMES | CLERK | 30 | salesmgr
(14 rows)

```

If the query references the `sal` or `comm` columns, then the policy is applied to the query eliminating any rows where `sal` is greater than or equal to `2000` as shown by the following:

```

edb=# SELECT empno, ename, job, sal, comm, deptno, authid FROM vpemp;
empno | ename | job | sal | comm | deptno | authid
-----+-----+-----+-----+
7934 | MILLER | CLERK | 1300.00 | | 10 |
7369 | SMITH | CLERK | 800.00 | | 20 | researchmgr
7876 | ADAMS | CLERK | 1100.00 | | 20 | researchmgr
7499 | ALLEN | SALESMAN | 1600.00 | 1200.00 | 30 | salesmgr
7521 | WARD | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr
7654 | MARTIN | SALESMAN | 1250.00 | 937.50 | 30 | salesmgr

```

```
7844 | TURNER | SALESMAN | 1500.00 | 1125.00 |    30 | salesmgr
7900 | JAMES  | CLERK   | 950.00 | 712.50 |    30 | salesmgr
(8 rows)
```

DROP_POLICY

The `DROP_POLICY` procedure deletes an existing policy. The policy function and database object associated with the policy are not deleted by the `DROP_POLICY` procedure.

You must be a superuser to execute this procedure.

```
DROP_POLICY(<object_schema> VARCHAR2, <object_name> VARCHAR2,
<policy_name> VARCHAR2)
```

Parameters

`object_schema`

Name of the schema containing the database object to which the policy applies.

`object_name`

Name of the database object to which the policy applies.

`policy_name`

Name of the policy to be deleted.

Examples

The following example deletes policy `secure_update` on table `public.vpemp`:

```
DECLARE
  v_object_schema      VARCHAR2(30) := 'public';
  v_object_name        VARCHAR2(30) := 'vpemp';
  v_policy_name        VARCHAR2(30) := 'secure_update';
BEGIN
  DBMS_RLS.DROP_POLICY(
    v_object_schema,
    v_object_name,
    v_policy_name
  );
END;
```

ENABLE_POLICY

The `ENABLE_POLICY` procedure enables or disables an existing policy on the specified database object.

You must be a superuser to execute this procedure.

```
ENABLE_POLICY(<object_schema> VARCHAR2, <object_name> VARCHAR2,
<policy_name> VARCHAR2, <enable> BOOLEAN)
```

Parameters

object_schema

Name of the schema containing the database object to which the policy applies.

object_name

Name of the database object to which the policy applies.

policy_name

Name of the policy to be enabled or disabled.

enable

When set to `TRUE`, the policy is enabled. When set to `FALSE`, the policy is disabled.

Examples

The following example disables policy `secure_update` on table `public.vpemp`:

```
DECLARE
  v_object_schema      VARCHAR2(30) := 'public';
  v_object_name        VARCHAR2(30) := 'vpemp';
  v_policy_name        VARCHAR2(30) := 'secure_update';
  v_enable              BOOLEAN := FALSE;
BEGIN
  DBMS_RLS.ENABLE_POLICY(
    v_object_schema,
    v_object_name,
    v_policy_name,
    v_enable
  );
END;
```

6.3.15 DBMS_SCHEDULER

The `DBMS_SCHEDULER` package provides a way to create and manage Oracle-styled jobs, programs and job schedules. The `DBMS_SCHEDULER` package implements the following functions and procedures:

| Function/Procedure | Return Type | Description |
|--|-------------|--|
| <code>CREATE_JOB(job_name, job_type, job_action, number_of_arguments, start_date, repeat_interval, end_date, job_class, enabled, auto_drop, comments)</code> | n/a | Use the first form of the <code>CREATE_JOB</code> procedure to create a job, specifying program and schedule details by means of parameters. |
| <code>CREATE_JOB(job_name, program_name, schedule_name, job_class, enabled, auto_drop, comments)</code> | n/a | Use the second form of <code>CREATE_JOB</code> to create a job that uses a named program and named schedule. |
| <code>CREATE_PROGRAM(program_name, program_type, program_action, number_of_arguments, enabled, comments)</code> | n/a | Use <code>CREATE_PROGRAM</code> to create a program. |

| | | |
|--|-----|--|
| CREATE_SCHEDULE(schedule_name, start_date, repeat_interval, end_date, comments) | n/a | Use the CREATE_SCHEDULE procedure to create a schedule. |
| DEFINE_PROGRAM_ARGUMENT(program_name, argument_position, argument_name, argument_type, default_value, out_argument) | n/a | Use the first form of the DEFINE_PROGRAM_ARGUMENT procedure to define a program argument that has a default value. |
| DEFINE_PROGRAM_ARGUMENT(program_name, argument_position, argument_name, argument_type, out_argument) | n/a | Use the first form of the DEFINE_PROGRAM_ARGUMENT procedure to define a program argument that does not have a default value. |
| DISABLE(name, force, commit_semantics) | n/a | Use the DISABLE procedure to disable a job or program. |
| DROP_JOB(job_name, force, defer, commit_semantics) | n/a | Use the DROP_JOB procedure to drop a job. |
| DROP_PROGRAM(program_name, force) | n/a | Use the DROP_PROGRAM procedure to drop a program. |
| DROP_PROGRAM_ARGUMENT(program_name, argument_position) | n/a | Use the first form of DROP_PROGRAM_ARGUMENT to drop a program argument by specifying the argument position. |
| DROP_PROGRAM_ARGUMENT(program_name, argument_name) | n/a | Use the second form of DROP_PROGRAM_ARGUMENT to drop a program argument by specifying the argument name. |
| DROP_SCHEDULE(schedule_name, force) | n/a | Use the DROP_SCHEDULE procedure to drop a schedule. |
| ENABLE(name, commit_semantics) | n/a | Use the ENABLE command to enable a program or job. |
| EVALUATE_CALENDAR_STRING(calendar_string, start_date, return_date_after, next_run_date) | n/a | Use EVALUATE_CALENDAR_STRING to review the execution date described by a user-defined calendar schedule. |
| RUN_JOB(job_name, use_current_session, manually) | n/a | Use the RUN_JOB procedure to execute a job immediately. |
| SET_JOB_ARGUMENT_VALUE(job_name, argument_position, argument_value) | n/a | Use the first form of SET_JOB_ARGUMENT value to set the value of a job argument described by the argument's position. |
| SET_JOB_ARGUMENT_VALUE(job_name, argument_name, argument_value) | n/a | Use the second form of SET_JOB_ARGUMENT value to set the value of a job argument described by the argument's name. |

Advanced Server's implementation of `DBMS_SCHEDULER` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The `DBMS_SCHEDULER` package is dependent on the pgAgent service; you must have a pgAgent service installed and running on your server before using `DBMS_SCHEDULER`.

Before using `DBMS_SCHEDULER`, a database superuser must create the catalog tables in which the `DBMS_SCHEDULER` programs, schedules and jobs are stored. Use the `psql` client to connect to the database, and invoke the command:

```
| CREATE EXTENSION dbms_scheduler;
```

By default, the `dbms_scheduler` extension resides in the `contrib/dbms_scheduler_ext` subdirectory (under the Advanced Server installation).

Note that after creating the `DBMS_SCHEDULER` tables, only a superuser will be able to perform a dump or reload of the database.

6.3.15.1 Using Calendar Syntax to Specify a Repeating Interval

The `CREATE_JOB` and `CREATE_SCHEDULE` procedures use Oracle-styled calendar syntax to define the interval with which a job or schedule is repeated. You should provide the scheduling information in the `repeat_interval` parameter of each procedure.

`repeat_interval` is a value (or series of values) that define the interval between the executions of the scheduled job. Each value is composed of a token, followed by an equal sign, followed by the unit (or units) on which the schedule will execute. Multiple token values must be separated by a semi-colon (:).

For example, the following value:

```
| FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;BYMINUTE=45
```

Defines a schedule that is executed each weeknight at 5:45.

The token types and syntax described in the table below are supported by Advanced Server:

| Token type | Syntax | Valid Values |
|------------|----------------------------|--|
| FREQ | FREQ=predefined_interval | Where predefined_interval is one of the following: YEARLY, MONTHLY, WEEKLY, DAILY, HOURLY, MINUTELY. The SECONDLY keyword is not supported. |
| BYMONTH | BYMONTH=month(, month)... | Where month is the three-letter abbreviation of the month name: JAN FEB MAR APR MAY JUN JUL AUG SEP OCT NOV DEC |
| BYMONTH | BYMONTH=month (, month)... | Where month is the numeric value representing the month: 1 2 3 4 5 6 7 8 9 10 11 12 |
| BYMONTHDAY | BYMONTHDAY=day_of_month | Where day_of_month is a value from 1 through 31 |
| BYDAY | BYDAY=weekday | Where weekday is a three-letter abbreviation or single-digit value representing the day of the week. |
| | | Monday MON 1 Tuesday TUE 2 Wednesday WED 3 Thursday THU 4 Friday FRI 5 Saturday SAT 6 Sunday SUN 7 |

| Token type | Syntax | Valid Values |
|------------|-------------------------|--|
| BYDATE | BYDATE=date (, date)... | Where date is YYYYMMDD. YYYY is a four-digit year representation of the year, MM is a two-digit representation of the month, and DD is a two-digit day representation of the day. |
| BYDATE | BYDATE=date (, date)... | Where date is MMDD. MM is a two-digit representation of the month, and DD is a two-digit day representation of the day |
| BYHOUR | BYHOUR=hour | Where hour is a value from 0 through 23. |
| BYMINUTE | BYMINUTE=minute | Where minute is a value from 0 through 59. |

6.3.15.2 CREATE_JOB

Use the `CREATE_JOB` procedure to create a job. The procedure comes in two forms; the first form of the procedure specifies a schedule within the job definition, as well as a job action that will be invoked when the job executes:

```
CREATE_JOB(
<job_name> IN VARCHAR2,
<job_type> IN VARCHAR2,
<job_action> IN VARCHAR2,
<number_of_arguments> IN PLS_INTEGER DEFAULT 0,
<start_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
<repeat_interval> IN VARCHAR2 DEFAULT NULL,
<end_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
<job_class> IN VARCHAR2 DEFAULT 'DEFAULT_JOB_CLASS',
<enabled> IN BOOLEAN DEFAULT FALSE,
<auto_drop> IN BOOLEAN DEFAULT TRUE,
<comments> IN VARCHAR2 DEFAULT NULL)
```

The second form uses a job schedule to specify the schedule on which the job will execute, and specifies the name of a program that will execute when the job runs:

```
CREATE_JOB(
<job_name> IN VARCHAR2,
<program_name> IN VARCHAR2,
<schedule_name> IN VARCHAR2,
<job_class> IN VARCHAR2 DEFAULT 'DEFAULT_JOB_CLASS',
<enabled> IN BOOLEAN DEFAULT FALSE,
```

```
<auto_drop> IN BOOLEAN DEFAULT TRUE,
<comments> IN VARCHAR2 DEFAULT NULL)
```

Parameters

job_name

`job_name` specifies the optionally schema-qualified name of the job being created.

job_type

`job_type` specifies the type of job. The current implementation of `CREATE_JOB` supports a job type of `PLSQL_BLOCK` or `STORED_PROCEDURE`.

job_action

If `job_type` is `PLSQL_BLOCK`, `job_action` specifies the content of the PL/SQL block that will be invoked when the job executes. The block must be terminated with a semi-colon (`;`).

If `job_type` is `STORED_PROCEDURE`, `job_action` specifies the optionally schema-qualified name of the procedure.

number_of_arguments

`number_of_arguments` is an `INTEGER` value that specifies the number of arguments expected by the job. The default is `0`.

start_date

`start_date` is a `TIMESTAMP WITH TIME ZONE` value that specifies the first time that the job is scheduled to execute. The default value is `NULL`, indicating that the job should be scheduled to execute when the job is enabled.

repeat_interval

`repeat_interval` is a `VARCHAR2` value that specifies how often the job will repeat. If a `repeat_interval` is not specified, the job will execute only once. The default value is `NULL`.

end_date

`end_date` is a `TIMESTAMP WITH TIME ZONE` value that specifies a time after which the job will no longer execute. If a date is specified, the `end_date` must be after `start_date`. The default value is `NULL`.

Please note that if an `end_date` is not specified and a `repeat_interval` is specified, the job will repeat indefinitely until it is disabled.

program_name

`program_name` is the name of a program that will be executed by the job.

schedule_name

`schedule_name` is the name of the schedule associated with the job.

job_class

`job_class` is accepted for compatibility and ignored.

enabled

`enabled` is a `BOOLEAN` value that specifies if the job is enabled when created. By default, a job is created in a disabled state, with `enabled` set to `FALSE`. To enable a job, specify a value of `TRUE` when creating the

job, or enable the job with the `DBMS_SCHEDULER.ENABLE` procedure.

auto_drop

The `auto_drop` parameter is accepted for compatibility and is ignored. By default, a job's status will be changed to `DISABLED` after the time specified in `end_date`.

comments

Use the `comments` parameter to specify a comment about the job.

Example

The following example demonstrates a call to the `CREATE_JOB` procedure:

```
EXEC
DBMS_SCHEDULER.CREATE_JOB (
    job_name      => 'update_log',
    job_type      => 'PLSQL_BLOCK',
    job_action    => 'BEGIN INSERT INTO my_log VALUES(current_timestamp);
                      END;',
    start_date    => '01-JUN-15 09:00:00.000000',
    repeat_interval => 'FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;',
    end_date      => NULL,
    enabled        => TRUE,
    comments       => 'This job adds a row to the my_log table.');
```

The code fragment creates a job named `update_log` that executes each weeknight at 5:00. The job executes a PL/SQL block that inserts the current timestamp into a logfile (`my_log`). Since no `end_date` is specified, the job will execute until it is disabled by the `DBMS_SCHEDULER.DISABLE` procedure.

6.3.15.3 CREATE_PROGRAM

Use the `CREATE_PROGRAM` procedure to create a `DBMS_SCHEDULER` program. The signature is:

```
CREATE_PROGRAM(
    <program_name> IN VARCHAR2,
    <program_type> IN VARCHAR2,
    <program_action> IN VARCHAR2,
    <number_of_arguments> IN PLS_INTEGER DEFAULT 0,
    <enabled> IN BOOLEAN DEFAULT FALSE,
    <comments> IN VARCHAR2 DEFAULT NULL)
```

Parameters

program_name

`program_name` specifies the name of the program that is being created.

program_type

`program_type` specifies the type of program. The current implementation of `CREATE_PROGRAM` supports a `program_type` of `PLSQL_BLOCK` or `PROCEDURE`.

program_action

If `program_type` is `PLSQL_BLOCK`, `program_action` contains the PL/SQL block that will execute when the program is invoked. The PL/SQL block must be terminated with a semi-colon (`;`).

If `program_type` is `PROCEDURE`, `program_action` contains the name of the stored procedure.

number_of_arguments

If `program_type` is `PLSQL_BLOCK`, this argument is ignored.

If `program_type` is `PROCEDURE`, `number_of_arguments` specifies the number of arguments required by the procedure. The default value is `0`.

enabled

`enabled` specifies if the program is created enabled or disabled:

- If `enabled` is `TRUE`, the program is created enabled.
- If `enabled` is `FALSE`, the program is created disabled; use the `DBMS_SCHEDULER.ENABLE` program to enable a disabled program.

The default value is `FALSE`.

comments

Use the `comments` parameter to specify a comment about the program; by default, this parameter is `NULL`.

Example

The following call to the `CREATE_PROGRAM` procedure creates a program named `update_log`:

```
EXEC
DBMS_SCHEDULER.CREATE_PROGRAM (
    program_name    => 'update_log',
    program_type    => 'PLSQL_BLOCK',
    program_action  => 'BEGIN INSERT INTO my_log VALUES(current_timestamp);
                           END;',
    enabled         => TRUE,
    comment         => 'This program adds a row to the my_log table.');
```

`update_log` is a PL/SQL block that adds a row containing the current date and time to the `my_log` table. The program will be enabled when the `CREATE_PROGRAM` procedure executes.

6.3.15.4 CREATE_SCHEDULE

Use the `CREATE_SCHEDULE` procedure to create a job schedule. The signature of the `CREATE_SCHEDULE` procedure is:

```
CREATE_SCHEDULE(
    <schedule_name> IN VARCHAR2,
    <start_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
    <repeat_interval> IN VARCHAR2,
    <end_date> IN TIMESTAMP WITH TIME ZONE DEFAULT NULL,
```

<comments> IN VARCHAR2 DEFAULT NULL)

Parameters

schedule_name

`schedule_name` specifies the name of the schedule.

start_date

`start_date` is a `TIMESTAMP WITH TIME ZONE` value that specifies the date and time that the schedule is eligible to execute. If a `start_date` is not specified, the date that the job is enabled is used as the `start_date`. By default, `start_date` is `NULL`.

repeat_interval

`repeat_interval` is a `VARCHAR2` value that specifies how often the job will repeat. If a `repeat_interval` is not specified, the job will execute only once, on the date specified by `start_date`.

Note

You must provide a value for either `start_date` or `repeat_interval`; if both `start_date` and `repeat_interval` are `NULL`, the server will return an error.

end_date

`end_date` is a `TIMESTAMP WITH TIME ZONE` value that specifies a time after which the schedule will no longer execute. If a date is specified, the `end_date` must be after the `start_date`. The default value is `NULL`.

Note

If a `repeat_interval` is specified and an `end_date` is not specified, the schedule will repeat indefinitely until it is disabled.

comments

Use the `comments` parameter to specify a comment about the schedule; by default, this parameter is `NULL`.

Example

The following code fragment calls `CREATE_SCHEDULE` to create a schedule named `weeknights_at_5`:

```
EXEC
DBMS_SCHEDULER.CREATE_SCHEDULE (
    schedule_name  => 'weeknights_at_5',
    start_date     => '01-JUN-13 09:00:00.000000'
    repeat_interval => 'FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17';
    comments       => 'This schedule executes each weeknight at 5:00');
```

The schedule executes each weeknight, at 5:00 pm, effective after June 1, 2013. Since no `end_date` is specified, the schedule will execute indefinitely until it is disabled with `DBMS_SCHEDULER.DISABLE`.

6.3.15.5 DEFINE_PROGRAM_ARGUMENT

Use the `DEFINE_PROGRAM_ARGUMENT` procedure to define a program argument. The

`DEFINE_PROGRAM_ARGUMENT` procedure comes in two forms; the first form defines an argument with a default value:

```
DEFINE_PROGRAM_ARGUMENT(
<program_name> IN VARCHAR2,
<argument_position> IN PLS_INTEGER,
<argument_name> IN VARCHAR2 DEFAULT NULL,
<argument_type> IN VARCHAR2,
<default_value> IN VARCHAR2,
<out_argument> IN BOOLEAN DEFAULT FALSE)
```

The second form defines an argument without a default value:

```
DEFINE_PROGRAM_ARGUMENT(
<program_name> IN VARCHAR2,
<argument_position> IN PLS_INTEGER,
<argument_name> IN VARCHAR2 DEFAULT NULL,
<argument_type> IN VARCHAR2,
<out_argument> IN BOOLEAN DEFAULT FALSE)
```

Parameters

`program_name`

`program_name` is the name of the program to which the arguments belong.

`argument_position`

`argument_position` specifies the position of the argument as it is passed to the program.

`argument_name`

`argument_name` specifies the optional name of the argument. By default, `argument_name` is `NULL`.

`argument_type` IN `VARCHAR2`

`argument_type` specifies the data type of the argument.

`default_value`

`default_value` specifies the default value assigned to the argument. `default_value` will be overridden by a value specified by the job when the job executes.

`out_argument` IN `BOOLEAN` `DEFAULT FALSE`

`out_argument` is not currently used; if specified, the value must be `FALSE`.

Example

The following code fragment uses the `DEFINE_PROGRAM_ARGUMENT` procedure to define the first and second arguments in a program named `add_emp`:

```
EXEC
DBMS_SCHEDULER.DEFINE_PROGRAM_ARGUMENT(
program_name      => 'add_emp',
argument_position => 1,
argument_name     => 'dept_no',
argument_type     => 'INTEGER',
default_value     => '20');
```

```
EXEC
DBMS_SCHEDULER.DEFINE_PROGRAM_ARGUMENT(
    program_name      => 'add_emp',
    argument_position => 2,
    argument_name     => 'emp_name',
    argument_type     => 'VARCHAR2');
```

The first argument is an `INTEGER` value named `dept_no` that has a default value of `20`. The second argument is a `VARCHAR2` value named `emp_name`; the second argument does not have a default value.

6.3.15.6 DISABLE

Use the `DISABLE` procedure to disable a program or a job. The signature of the `DISABLE` procedure is:

```
DISABLE(
    <name> IN VARCHAR2,
    <force> IN BOOLEAN DEFAULT FALSE,
    <commit_semantics> IN VARCHAR2 DEFAULT 'STOP_ON_FIRST_ERROR')
```

Parameters

`name`

`name` specifies the name of the program or job that is being disabled.

`force`

`force` is accepted for compatibility, and ignored.

`commit_semantics`

`commit_semantics` instructs the server how to handle an error encountered while disabling a program or job. By default, `commit_semantics` is set to `STOP_ON_FIRST_ERROR`, instructing the server to stop when it encounters an error. Any programs or jobs that were successfully disabled prior to the error will be committed to disk.

The `TRANSACTIONAL` and `ABSORB_ERRORS` keywords are accepted for compatibility, and ignored.

Example

The following call to the `DISABLE` procedure disables a program named `update_emp`:

```
DBMS_SCHEDULER.DISABLE('update_emp');
```

6.3.15.7 DROP_JOB

Use the `DROP_JOB` procedure to `DROP` a job, `DROP` any arguments that belong to the job, and eliminate any future job executions. The signature of the procedure is:

```
DROP_JOB(
<job_name> IN VARCHAR2,
<force> IN BOOLEAN DEFAULT FALSE,
<defer> IN BOOLEAN DEFAULT FALSE,
<commit_semantics> IN VARCHAR2 DEFAULT 'STOP_ON_FIRST_ERROR')
```

Parameters

`job_name`

`job_name` specifies the name of the job that is being dropped.

`force`

`force` is accepted for compatibility, and ignored.

`defer`

`defer` is accepted for compatibility, and ignored.

`commit_semantics`

`commit_semantics` instructs the server how to handle an error encountered while dropping a program or job. By default, `commit_semantics` is set to `STOP_ON_FIRST_ERROR`, instructing the server to stop when it encounters an error.

The `TRANSACTIONAL` and `ABSORB_ERRORS` keywords are accepted for compatibility, and ignored.

Example

The following call to `DROP_JOB` drops a job named `update_log`:

```
DBMS_SCHEDULER.DROP_JOB('update_log');
```

6.3.15.8 `DROP_PROGRAM`

The `DROP_PROGRAM` procedure to drop a program. The signature of the `DROP_PROGRAM` procedure is:

```
DROP_PROGRAM(
<program_name> IN VARCHAR2,
<force> IN BOOLEAN DEFAULT FALSE)
```

Parameters

`program_name`

`program_name` specifies the name of the program that is being dropped.

`force`

`force` is a `BOOLEAN` value that instructs the server how to handle programs with dependent jobs.

Specify `FALSE` to instruct the server to return an error if the program is referenced by a job.

Specify `TRUE` to instruct the server to disable any jobs that reference the program before dropping the program.

The default value is `FALSE`.

Example

The following call to `DROP_PROGRAM` drops a job named `update_emp`:

```
DBMS_SCHEDULER.DROP_PROGRAM('update_emp');
```

6.3.15.9 DROP_PROGRAM_ARGUMENT

Use the `DROP_PROGRAM_ARGUMENT` procedure to drop a program argument. The `DROP_PROGRAM_ARGUMENT` procedure comes in two forms; the first form uses an argument position to specify which argument to drop:

```
DROP_PROGRAM_ARGUMENT(
  <program_name> IN VARCHAR2,
  <argument_position> IN PLS_INTEGER)
```

The second form takes the argument name:

```
DROP_PROGRAM_ARGUMENT(
  <program_name> IN VARCHAR2,
  <argument_name> IN VARCHAR2)
```

Parameters

`program_name`

`program_name` specifies the name of the program that is being modified.

`argument_position`

`argument_position` specifies the position of the argument that is being dropped.

`argument_name`

`argument_name` specifies the name of the argument that is being dropped.

Examples

The following call to `DROP_PROGRAM_ARGUMENT` drops the first argument in the `update_emp` program:

```
DBMS_SCHEDULER.DROP_PROGRAM_ARGUMENT('update_emp', 1);
```

The following call to `DROP_PROGRAM_ARGUMENT` drops an argument named `emp_name`:

```
DBMS_SCHEDULER.DROP_PROGRAM_ARGUMENT('update_emp', 'emp_name');
```

6.3.15.10 DROP_SCHEDULE

Use the `DROP_SCHEDULE` procedure to drop a schedule. The signature is:

```
DROP_SCHEDULE(
  <schedule_name> IN VARCHAR2,
  <force> IN BOOLEAN DEFAULT FALSE)
```

Parameters

`schedule_name`

`schedule_name` specifies the name of the schedule that is being dropped.

`force`

`force` specifies the behavior of the server if the specified schedule is referenced by any job:

- Specify `FALSE` to instruct the server to return an error if the specified schedule is referenced by a job. This is the default behavior.
- Specify `TRUE` to instruct the server to disable to any jobs that use the specified schedule before dropping the schedule. Any running jobs will be allowed to complete before the schedule is dropped.

Example

The following call to `DROP_SCHEDULE` drops a schedule named `weeknights_at_5`:

```
DBMS_SCHEDULER.DROP_SCHEDULE('weeknights_at_5', TRUE);
```

The server will disable any jobs that use the schedule before dropping the schedule.

6.3.15.11 ENABLE

Use the `ENABLE` procedure to enable a disabled program or job.

The signature of the `ENABLE` procedure is:

```
ENABLE(
  <name> IN VARCHAR2,
  <commit_semantics> IN VARCHAR2 DEFAULT 'STOP_ON_FIRST_ERROR')
```

Parameters

`name`

`name` specifies the name of the program or job that is being enabled.

`commit_semantics`

`commit_semantics` instructs the server how to handle an error encountered while enabling a program or job. By default, `commit_semantics` is set to `STOP_ON_FIRST_ERROR`, instructing the server to stop when it encounters an error.

The `TRANSACTIONAL` and `ABSORB_ERRORS` keywords are accepted for compatibility, and ignored.

Example

The following call to `DBMS_SCHEDULER.ENABLE` enables the `update_emp` program:

```
DBMS_SCHEDULER.ENABLE('update_emp');
```

6.3.15.12 EVALUATE CALENDAR STRING

Use the `EVALUATE CALENDAR STRING` procedure to evaluate the `repeat_interval` value specified when creating a schedule with the `CREATE_SCHEDULE` procedure. The `EVALUATE CALENDAR STRING` procedure will return the date and time that a specified schedule will execute without actually scheduling the job.

The signature of the `EVALUATE CALENDAR STRING` procedure is:

```
EVALUATE CALENDAR STRING(
<calendar_string> IN VARCHAR2,
<start_date> IN TIMESTAMP WITH TIME ZONE,
<return_date_after> IN TIMESTAMP WITH TIME ZONE,
<next_run_date> OUT TIMESTAMP WITH TIME ZONE)
```

Parameters

`calendar_string`

`calendar_string` is the calendar string that describes a `repeat_interval` that is being evaluated.

`start_date` IN `TIMESTAMP WITH TIME ZONE`

`start_date` is the date and time after which the `repeat_interval` will become valid.

`return_date_after`

Use the `return_date_after` parameter to specify the date and time that `EVALUATE CALENDAR STRING` should use as a starting date when evaluating the `repeat_interval`.

For example, if you specify a `return_date_after` value of `01-APR-13 09.00.00.000000`, `EVALUATE CALENDAR STRING` will return the date and time of the first iteration of the schedule after April 1st, 2013.

`next_run_date` OUT `TIMESTAMP WITH TIME ZONE`

`next_run_date` is an `OUT` parameter that will contain the first occurrence of the schedule after the date specified by the `return_date_after` parameter.

Example

The following example evaluates a calendar string and returns the first date and time that the schedule will be executed after June 15, 2013:

```
DECLARE
result  TIMESTAMP;
BEGIN
```

```

DBMS_SCHEDULER.EVALUATE_CALENDAR_STRING
(
  'FREQ=DAILY;BYDAY=MON,TUE,WED,THU,FRI;BYHOUR=17;',
  '15-JUN-2013', NULL, result
);

  DBMS_OUTPUT.PUT_LINE('next_run_date: ' || result);
END;
/

```

next_run_date: 17-JUN-13 05.00.00.000000 PM

June 15, 2013 is a Saturday; the schedule will not execute until Monday, June 17, 2013 at 5:00 pm.

6.3.15.13 RUN_JOB

Use the `RUN_JOB` procedure to execute a job immediately. The signature of the `RUN_JOB` procedure is:

```

RUN_JOB(
  <job_name> IN VARCHAR2,
  <use_current_session> IN BOOLEAN DEFAULT TRUE

```

Parameters

`job_name`

`job_name` specifies the name of the job that will execute.

`use_current_session`

By default, the job will execute in the current session. If specified, `use_current_session` must be set to `TRUE`; if `use_current_session` is set to `FALSE`, Advanced Server will return an error.

Example

The following call to `RUN_JOB` executes a job named `update_log`:

```
DBMS_SCHEDULER.RUN_JOB('update_log', TRUE);
```

Passing a value of `TRUE` as the second argument instructs the server to invoke the job in the current session.

6.3.15.14 SET_JOB_ARGUMENT_VALUE

Use the `SET_JOB_ARGUMENT_VALUE` procedure to specify a value for an argument. The `SET_JOB_ARGUMENT_VALUE` procedure comes in two forms; the first form specifies which argument should be modified by position:

```
SET_JOB_ARGUMENT_VALUE(
<job_name> IN VARCHAR2,
<argument_position> IN PLS_INTEGER,
<argument_value> IN VARCHAR2)
```

The second form uses an argument name to specify which argument to modify:

```
SET_JOB_ARGUMENT_VALUE(
<job_name> IN VARCHAR2,
<argument_name> IN VARCHAR2,
<argument_value> IN VARCHAR2)
```

Argument values set by the `SET_JOB_ARGUMENT_VALUE` procedure override any values set by default.

Parameters

`job_name`

`job_name` specifies the name of the job to which the modified argument belongs.

`argument_position`

Use `argument_position` to specify the argument position for which the value will be set.

`argument_name`

Use `argument_name` to specify the argument by name for which the value will be set.

`argument_value`

`argument_value` specifies the new value of the argument.

Examples

The following example assigns a value of `30` to the first argument in the `update_emp` job:

```
DBMS_SCHEDULER.SET_JOB_ARGUMENT_VALUE('update_emp', 1, '30');
```

The following example sets the `emp_name` argument to `SMITH`:

```
DBMS_SCHEDULER.SET_JOB_ARGUMENT_VALUE('update_emp', 'emp_name', 'SMITH');
```

6.3.16 DBMS_SESSION

Advanced Server provides support for the following `DBMS_SESSION.SET_ROLE` procedure:

| Function/Procedure | Return Type | Description |
|---------------------------------|-------------|---|
| <code>SET_ROLE(role_cmd)</code> | n/a | Executes a SET ROLE statement followed by the string value specified in <code>role_cmd</code> . |

Advanced Server's implementation of `DBMS_SESSION` is a partial implementation when compared to Oracle's

version. Only `DBMS_SESSION.SET_ROLE` is supported.

SET_ROLE

The `SET ROLE` procedure sets the current session user to the role specified in `role_cmd`. After invoking the `SET_ROLE` procedure, the current session will use the permissions assigned to the specified role. The signature of the procedure is:

```
SET_ROLE(<role_cmd>)
```

The `SET_ROLE` procedure appends the value specified for `role_cmd` to the `SET ROLE` statement, and then invokes the statement.

Parameters

`role_cmd`

`role_cmd` specifies a role name in the form of a string value.

Example

The following call to the `SET_ROLE` procedure invokes the `SET ROLE` command to set the identity of the current session user to manager:

```
edb=# exec DBMS_SESSION.SET_ROLE('manager');
```

6.3.17 DBMS_SQL

The `DBMS_SQL` package provides an application interface compatible with Oracle databases to the EnterpriseDB dynamic SQL functionality. With `DBMS_SQL` you can construct queries and other commands at run time (rather than when you write the application). EnterpriseDB Advanced Server offers native support for dynamic SQL; `DBMS_SQL` provides a way to use dynamic SQL in a fashion compatible with Oracle databases without modifying your application.

`DBMS_SQL` assumes the privileges of the current user when executing dynamic SQL statements.

| Function/Procedure | Function or Procedure | Return Type | Description |
|--|-----------------------|-------------|---|
| <code>BIND_VARIABLE(c, name, value [, out_value_size])</code> | Procedure | n/a | Bind a value to a variable. |
| <code>BIND_VARIABLE_CHAR(c, name, value [, out_value_size])</code> | Procedure | n/a | Bind a CHAR value to a variable. |
| <code>BIND_VARIABLE_RAW(c, name, value [, out_value_size])</code> | Procedure | n/a | Bind a RAW value to a variable. |
| <code>CLOSE_CURSOR(c IN OUT)</code> | Procedure | n/a | Close a cursor. |
| <code>COLUMN_VALUE(c, position, value OUT [, column_error OUT [, actual_length OUT]])</code> | Procedure | n/a | Return a column value into a variable. |
| <code>COLUMN_VALUE_CHAR(c, position, value OUT [, column_error OUT [, actual_length OUT]])</code> | Procedure | n/a | Return a CHAR column value into a variable. |

| | | | |
|--|-----------|---------|--|
| COLUMN_VALUE_RAW(c, position, value OUT [, column_error OUT [, actual_length OUT]]) | Procedure | n/a | Return a RAW column value into a variable. |
| DEFINE_COLUMN(c, position, column [, column_size]) | Procedure | n/a | Define a column in the SELECT list. |
| DEFINE_COLUMN_CHAR(c, position, column, column_size) | Procedure | n/a | Define a CHAR column in the SELECT list. |
| DEFINE_COLUMN_RAW(c, position, column, column_size) | Procedure | n/a | Define a RAW column in the SELECT list. |
| DESCRIBE_COLUMNS | Procedure | n/a | Defines columns to hold a cursor result set. |
| EXECUTE(c) | Function | INTEGER | Execute a cursor. |
| EXECUTE_AND_FETCH(c [, exact]) | Function | INTEGER | Execute a cursor and fetch a single row. |
| FETCH_ROWS(c) | Function | INTEGER | Fetch rows from the cursor. |
| IS_OPEN(c) | Function | BOOLEAN | Check if a cursor is open. |
| LAST_ROW_COUNT | Function | INTEGER | Return cumulative number of rows fetched. |
| OPEN_CURSOR | Function | INTEGER | Open a cursor. |
| PARSE(c, statement, language_flag) | Procedure | n/a | Parse a statement. |

Advanced Server's implementation of `DBMS_SQL` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variable available in the `DBMS_SQL` package.

| Public Variables | Data Type | Value | Description |
|------------------|-----------|-------|--|
| native | INTEGER | 1 | Provided for compatibility with Oracle syntax. See <code>DBMS_SQL.PARSE</code> for more information. |
| V6 | INTEGER | 2 | Provided for compatibility with Oracle syntax. See <code>DBMS_SQL.PARSE</code> for more information. |
| V7 | INTEGER | 3 | Provided for compatibility with Oracle syntax. See <code>DBMS_SQL.PARSE</code> for more information |

6.3.17.1 BIND_VARIABLE

The `BIND_VARIABLE` procedure provides the capability to associate a value with an `IN` or `IN OUT` bind variable in a SQL command.

```
BIND_VARIABLE(<c> INTEGER, <name> VARCHAR2,
<value> { BLOB | CLOB | DATE | FLOAT | INTEGER | NUMBER |
TIMESTAMP | VARCHAR2 }
[, <out_value_size> INTEGER ])
```

Parameters

C

Cursor ID of the cursor for the SQL command with bind variables.

name

Name of the bind variable in the SQL command.

value

Value to be assigned.

out_value_size

If **name** is an **IN OUT** variable, defines the maximum length of the output value. If not specified, the length of **value** is assumed.

Examples

The following anonymous block uses bind variables to insert a row into the **emp** stable.

```

DECLARE
    curid      INTEGER;
    v_sql       VARCHAR2(150) := 'INSERT INTO emp VALUES ' ||
                                '(:p_empno, :p_ename, :p_job, :p_mgr, ' ||
                                ':p_hiredate, :p_sal, :p_comm, :p_deptno)';
    v_empno    emp.empno%TYPE;
    v_ename    emp.ename%TYPE;
    v_job      emp.job%TYPE;
    v_mgr      emp.mgr%TYPE;
    v_hiredate emp.hiredate%TYPE;
    v_sal      emp.sal%TYPE;
    v_comm     emp.comm%TYPE;
    v_deptno   emp.deptno%TYPE;
    v_status   INTEGER;

BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQLPARSE(curid,v_sql,DBMS_SQL.native);
    v_empno  := 9001;
    v_ename   := 'JONES';
    v_job     := 'SALESMAN';
    v_mgr     := 7369;
    v_hiredate := TO_DATE('13-DEC-07','DD-MON-YY');
    v_sal     := 8500.00;
    v_comm    := 1500.00;
    v_deptno  := 40;
    DBMS_SQL.BIND_VARIABLE(curid,:p_empno,v_empno);
    DBMS_SQL.BIND_VARIABLE(curid,:p_ename,v_ename);
    DBMS_SQL.BIND_VARIABLE(curid,:p_job,v_job);
    DBMS_SQL.BIND_VARIABLE(curid,:p_mgr,v_mgr);
    DBMS_SQL.BIND_VARIABLE(curid,:p_hiredate,v_hiredate);
    DBMS_SQL.BIND_VARIABLE(curid,:p_sal,v_sal);
    DBMS_SQL.BIND_VARIABLE(curid,:p_comm,v_comm);
    DBMS_SQL.BIND_VARIABLE(curid,:p_deptno,v_deptno);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    DBMS_SQL CLOSE_CURSOR(curid);
END;
```

Number of rows processed: 1

6.3.17.2 BIND_VARIABLE_CHAR

The `BIND_VARIABLE_CHAR` procedure provides the capability to associate a `CHAR` value with an `IN` or `IN OUT` bind variable in a SQL command.

```
BIND_VARIABLE_CHAR(<c> INTEGER, <name> VARCHAR2, <value> CHAR
[, <out_value_size> INTEGER ])
```

Parameters

`c`

Cursor ID of the cursor for the SQL command with bind variables.

`name`

Name of the bind variable in the SQL command.

`value`

Value of type `CHAR` to be assigned.

`out_value_size`

If `name` is an `IN OUT` variable, defines the maximum length of the output value. If not specified, the length of `value` is assumed.

6.3.17.3 BIND VARIABLE RAW

The `BIND_VARIABLE_RAW` procedure provides the capability to associate a `RAW` value with an `IN` or `IN OUT` bind variable in a SQL command.

```
BIND_VARIABLE_RAW(<c> INTEGER, <name> VARCHAR2, <value> RAW
[, <out_value_size> INTEGER ])
```

Parameters

`c`

Cursor ID of the cursor for the SQL command with bind variables.

`name`

Name of the bind variable in the SQL command.

`value`

| Value of type **RAW** to be assigned.

out_value_size

| If **name** is an **IN OUT** variable, defines the maximum length of the output value. If not specified, the length of **value** is assumed.

6.3.17.4 CLOSE_CURSOR

The **CLOSE_CURSOR** procedure closes an open cursor. The resources allocated to the cursor are released and it can no longer be used.

```
CLOSE_CURSOR(<c> IN OUT INTEGER)
```

Parameters

c

| Cursor ID of the cursor to be closed.

Examples

The following example closes a previously opened cursor:

```
DECLARE
    curid      INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    .
    .
    .
    DBMS_SQL.CLOSE_CURSOR(curid);
END;
```

6.3.17.5 COLUMN_VALUE

The **COLUMN_VALUE** procedure defines a variable to receive a value from a cursor.

```
COLUMN_VALUE(<c> INTEGER, <position> INTEGER, <value> OUT { BLOB |
CLOB | DATE | FLOAT | INTEGER | NUMBER | TIMESTAMP | VARCHAR2 }
[, <column_error> OUT NUMBER [, <actual_length> OUT INTEGER ]])
```

Parameters

c

| Cursor id of the cursor returning data to the variable being defined.

position

Position within the cursor of the returned data. The first value in the cursor is position 1.

value

Variable receiving the data returned in the cursor by a prior fetch call.

column_error

Error number associated with the column, if any.

actual_length

Actual length of the data prior to any truncation.

Examples

The following example shows the portion of an anonymous block that receives the values from a cursor using the **COLUMN_VALUE** procedure.

```

DECLARE
    curid      INTEGER;
    v_empno    NUMBER(4);
    v_ename    VARCHAR2(10);
    v_hiredate DATE;
    v_sal      NUMBER(7,2);
    v_comm     NUMBER(7,2);
    v_sql      VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, '
                           'comm FROM emp';
    v_status   INTEGER;
BEGIN
    .
    .
    .
    LOOP
        v_status := DBMS_SQL.FETCH_ROWS(curid);
        EXIT WHEN v_status = 0;
        DBMS_SQL.COLUMN_VALUE(curid,1,v_empno);
        DBMS_SQL.COLUMN_VALUE(curid,2,v_ename);
        DBMS_SQL.COLUMN_VALUE(curid,3,v_hiredate);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,5,v_comm);
        DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || RPAD(v_ename,10) || ' ' ||
                             TO_CHAR(v_hiredate,'yyyy-mm-dd') || ' ' ||
                             TO_CHAR(v_sal,'9,999.99') || ' ' ||
                             TO_CHAR(NVL(v_comm,0),'9,999.99'));
    END LOOP;
    DBMS_SQL.CLOSE_CURSOR(curid);
END;

```

6.3.17.6 COLUMN_VALUE_CHAR

The `COLUMN_VALUE_CHAR` procedure defines a variable to receive a `CHAR` value from a cursor.

```
COLUMN_VALUE_CHAR(<c> INTEGER, <position> INTEGER, <value> OUT CHAR
[, <column_error> OUT NUMBER [, <actual_length> OUT INTEGER ]])
```

Parameters

`c`

Cursor id of the cursor returning data to the variable being defined.

`position`

Position within the cursor of the returned data. The first value in the cursor is position 1.

`value`

Variable of data type `CHAR` receiving the data returned in the cursor by a prior fetch call.

`column_error`

Error number associated with the column, if any.

`actual_length`

Actual length of the data prior to any truncation.

6.3.17.7 `COLUMN_VALUE_RAW`

The `COLUMN_VALUE_RAW` procedure defines a variable to receive a `RAW` value from a cursor.

```
COLUMN_VALUE_RAW(<c> INTEGER, <position> INTEGER, <value> OUT RAW
[, <column_error> OUT NUMBER [, <actual_length> OUT INTEGER ]])
```

Parameters

`c`

Cursor id of the cursor returning data to the variable being defined.

`position`

Position within the cursor of the returned data. The first value in the cursor is position 1.

`value`

Variable of data type `RAW` receiving the data returned in the cursor by a prior fetch call.

`column_error`

Error number associated with the column, if any.

`actual_length`

Actual length of the data prior to any truncation.

6.3.17.8 DEFINE_COLUMN

The `DEFINE_COLUMN` procedure defines a column or expression in the `SELECT` list that is to be returned and retrieved in a cursor.

```
DEFINE_COLUMN(<c> INTEGER, <position> INTEGER, <column> { BLOB |
CLOB | DATE | FLOAT | INTEGER | NUMBER | TIMESTAMP | VARCHAR2 }
[, <column_size> INTEGER ])
```

Parameters

`c`

Cursor id of the cursor associated with the `SELECT` command.

`position`

Position of the column or expression in the `SELECT` list that is being defined.

`column`

A variable that is of the same data type as the column or expression in position `position` of the `SELECT` list.

`column_size`

The maximum length of the returned data. `column_size` must be specified only if `column` is `VARCHAR2`.

Returned data exceeding `column_size` is truncated to `column_size` characters.

Examples

The following shows how the `empno`, `ename`, `hiredate`, `sal`, and `comm` columns of the `emp` table are defined with the `DEFINE_COLUMN` procedure.

```
DECLARE
    curid      INTEGER;
    v_empno    NUMBER(4);
    v_ename    VARCHAR2(10);
    v_hiredate DATE;
    v_sal      NUMBER(7,2);
    v_comm     NUMBER(7,2);
    v_sql      VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status   INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQLPARSE(curid,v_sql,DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid,1,v_empno);
    DBMS_SQL.DEFINE_COLUMN(curid,2,v_ename,10);
    DBMS_SQL.DEFINE_COLUMN(curid,3,v_hiredate);
    DBMS_SQL.DEFINE_COLUMN(curid,4,v_sal);
    DBMS_SQL.DEFINE_COLUMN(curid,5,v_comm);

    .
    .

END;
```

The following shows an alternative to the prior example that produces the exact same results. Note that the lengths of the data types are irrelevant – the `empno`, `sal`, and `comm` columns will still return data equivalent to `NUMBER(4)` and `NUMBER(7,2)`, respectively, even though `v_num` is defined as `NUMBER(1)` (assuming the declarations in the `COLUMN_VALUE` procedure are of the appropriate maximum sizes). The `ename` column will return data up to ten characters in length as defined by the `length` parameter in the `DEFINE_COLUMN` call, not by the data type declaration, `VARCHAR2(1)` declared for `v_varchar`. The actual size of the returned data is dictated by the `COLUMN_VALUE` procedure.

```

DECLARE
    curid      INTEGER;
    v_num      NUMBER(1);
    v_varchar   VARCHAR2(1);
    v_date     DATE;
    v_sql      VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status   INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQLPARSE(curid,v_sql,DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid,1,v_num);
    DBMS_SQL.DEFINE_COLUMN(curid,2,v_varchar,10);
    DBMS_SQL.DEFINE_COLUMN(curid,3,v_date);
    DBMS_SQL.DEFINE_COLUMN(curid,4,v_num);
    DBMS_SQL.DEFINE_COLUMN(curid,5,v_num);

    .
    .

END;

```

6.3.17.9 DEFINE_COLUMN_CHAR

The `DEFINE_COLUMN_CHAR` procedure defines a `CHAR` column or expression in the `SELECT` list that is to be returned and retrieved in a cursor.

```
DEFINE_COLUMN_CHAR(<c> INTEGER, <position> INTEGER, <column>
CHAR, <column_size> INTEGER)
```

Parameters

`c`

Cursor id of the cursor associated with the `SELECT` command.

`position`

Position of the column or expression in the `SELECT` list that is being defined.

`column`

A `CHAR` variable.

`column_size`

The maximum length of the returned data. Returned data exceeding `column_size` is truncated to `column_size` characters.

6.3.17.10 DEFINE_COLUMN_RAW

The `DEFINE_COLUMN_RAW` procedure defines a `RAW` column or expression in the `SELECT` list that is to be returned and retrieved in a cursor.

```
DEFINE_COLUMN_RAW(<c> INTEGER, <position> INTEGER, <column> RAW,
<column_size> INTEGER)
```

Parameters

`c`

Cursor id of the cursor associated with the `SELECT` command.

`position`

Position of the column or expression in the `SELECT` list that is being defined.

`column`

A `RAW` variable.

`column_size`

The maximum length of the returned data. Returned data exceeding `column_size` is truncated to `column_size` characters.

6.3.17.11 DESCRIBE COLUMNS

The `DESCRIBE_COLUMNS` procedure describes the columns returned by a cursor.

```
DESCRIBE_COLUMNS(<c> INTEGER, <col_cnt> OUT INTEGER, <desc_t> OUT
DESC_TAB);
```

Parameters

`c`

The cursor ID of the cursor.

`col_cnt`

The number of columns in cursor result set.

`desc_tab`

The table that contains a description of each column returned by the cursor. The descriptions are of type `DESC_REC`, and contain the following values:

| Column Name | Type |
|----------------------------------|---------------|
| <code>col_type</code> | INTEGER |
| <code>col_max_len</code> | INTEGER |
| <code>col_name</code> | VARCHAR2(128) |
| <code>col_name_len</code> | INTEGER |
| <code>col_schema_name</code> | VARCHAR2(128) |
| <code>col_schema_name_len</code> | INTEGER |
| <code>col_precision</code> | INTEGER |
| <code>col_scale</code> | INTEGER |
| <code>col_charsetid</code> | INTEGER |
| <code>col_charsetform</code> | INTEGER |
| <code>col_null_ok</code> | BOOLEAN |

6.3.17.12 EXECUTE

The `EXECUTE` function executes a parsed SQL command or SPL block.

```
<status> INTEGER EXECUTE(<c> INTEGER)
```

Parameters

c

Cursor ID of the parsed SQL command or SPL block to be executed.

status

Number of rows processed if the SQL command was `DELETE`, `INSERT`, or `UPDATE`. `status` is meaningless for all other commands.

Examples

The following anonymous block inserts a row into the `dept` table.

```
DECLARE
    curid      INTEGER;
    v_sql       VARCHAR2(50);
    v_status    INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    v_sql := 'INSERT INTO dept VALUES (50, "HR", "LOS ANGELES")';
    DBMS_SQLPARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    DBMS_SQLCLOSE_CURSOR(curid);
END;
```

6.3.17.13 EXECUTE_AND_FETCH

Function `EXECUTE_AND_FETCH` executes a parsed `SELECT` command and fetches one row.

```
<status> INTEGER EXECUTE_AND_FETCH(<c> INTEGER
[, <exact> BOOLEAN ])
```

Parameters

`c`

Cursor id of the cursor for the `SELECT` command to be executed.

`exact`

If set to `TRUE`, an exception is thrown if the number of rows in the result set is not exactly equal to 1. If set to `FALSE`, no exception is thrown. The default is `FALSE`. A `NO DATA FOUND` exception is thrown if `exact` is `TRUE` and there are no rows in the result set. A `TOO_MANY_ROWS` exception is thrown if `exact` is `TRUE` and there is more than one row in the result set.

`status`

Returns 1 if a row was successfully fetched, 0 if no rows to fetch. If an exception is thrown, no value is returned.

Examples

The following stored procedure uses the `EXECUTE_AND_FETCH` function to retrieve one employee using the employee's name. An exception will be thrown if the employee is not found, or there is more than one employee with the same name.

```
CREATE OR REPLACE PROCEDURE select_by_name(
    p_ename      emp.ename%TYPE
)
IS
    curid      INTEGER;
    v_empno     emp.empno%TYPE;
    v_hiredate  emp.hiredate%TYPE;
    v_sal       emp.sal%TYPE;
    v_comm      emp.comm%TYPE;
    v_dname     dept.dname%TYPE;
    v_disp_date VARCHAR2(10);
    v_sql       VARCHAR2(120) := 'SELECT empno, hiredate, sal, '
                                || 'NVL(comm, 0), dname '
                                || 'FROM emp e, dept d '
                                || 'WHERE ename = :p_ename '
                                || 'AND e.deptno = d.deptno';
    v_status    INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQLPARSE(curid,v_sql,DBMS_SQL.native);
    DBMS_SQLBIND_VARIABLE(curid,':p_ename',UPPER(p_ename));
    DBMS_SQLDEFINE_COLUMN(curid,1,v_empno);
    DBMS_SQLDEFINE_COLUMN(curid,2,v_hiredate);
    DBMS_SQLDEFINE_COLUMN(curid,3,v_sal);
```

```

DBMS_SQL.DEFINE_COLUMN(curid,4,v_comm);
DBMS_SQL.DEFINE_COLUMN(curid,5,v_dname,14);
v_status := DBMS_SQL.EXECUTE_AND_FETCH(curid,TRUE);
DBMS_SQL.COLUMN_VALUE(curid,1,v_empno);
DBMS_SQL.COLUMN_VALUE(curid,2,v_hiredate);
DBMS_SQL.COLUMN_VALUE(curid,3,v_sal);
DBMS_SQL.COLUMN_VALUE(curid,4,v_comm);
DBMS_SQL.COLUMN_VALUE(curid,5,v_dname);
v_disp_date := TO_CHAR(v_hiredate, 'MM/DD/YYYY');
DBMS_OUTPUT.PUT_LINE('Number : ' || v_empno);
DBMS_OUTPUT.PUT_LINE('Name : ' || UPPER(p_ename));
DBMS_OUTPUT.PUT_LINE('Hire Date : ' || v_disp_date);
DBMS_OUTPUT.PUT_LINE('Salary : ' || v_sal);
DBMS_OUTPUT.PUT_LINE('Commission: ' || v_comm);
DBMS_OUTPUT.PUT_LINE('Department: ' || v_dname);
DBMS_SQL.CLOSE_CURSOR(curid);
EXCEPTION
WHEN NO_DATA_FOUND THEN
  DBMS_OUTPUT.PUT_LINE('Employee ' || p_ename || ' not found');
  DBMS_SQL.CLOSE_CURSOR(curid);
WHEN TOO_MANY_ROWS THEN
  DBMS_OUTPUT.PUT_LINE('Too many employees named, ' ||
    p_ename || ', found');
  DBMS_SQL.CLOSE_CURSOR(curid);
WHEN OTHERS THEN
  DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
  DBMS_OUTPUT.PUT_LINE(SQLERRM);
  DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
  DBMS_OUTPUT.PUT_LINE(SQLCODE);
  DBMS_SQL.CLOSE_CURSOR(curid);
END;

EXEC select_by_name('MARTIN')

Number : 7654
Name : MARTIN
Hire Date : 09/28/1981
Salary : 1250
Commission: 1400
Department: SALES

```

6.3.17.14 **FETCH_ROWS**

The **FETCH_ROWS** function retrieves a row from a cursor.

```
<status> INTEGER FETCH_ROWS(<c> INTEGER)
```

Parameters

c

Cursor ID of the cursor from which to fetch a row.

status

Returns 1 if a row was successfully fetched, 0 if no more rows to fetch.

Examples

The following examples fetches the rows from the `emp` table and displays the results.

```

DECLARE
    curid      INTEGER;
    v_empno    NUMBER(4);
    v_ename    VARCHAR2(10);
    v_hiredate DATE;
    v_sal      NUMBER(7,2);
    v_comm     NUMBER(7,2);
    v_sql      VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, ' ||
                                'comm FROM emp';
    v_status   INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQLPARSE(curid,v_sql,DBMS_SQL.native);
    DBMS_SQL.DEFINE_COLUMN(curid,1,v_empno);
    DBMS_SQL.DEFINE_COLUMN(curid,2,v_ename,10);
    DBMS_SQL.DEFINE_COLUMN(curid,3,v_hiredate);
    DBMS_SQL.DEFINE_COLUMN(curid,4,v_sal);
    DBMS_SQL.DEFINE_COLUMN(curid,5,v_comm);

    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('EMPNO  ENAME      HIREDATE  SAL      COMM');
    DBMS_OUTPUT.PUT_LINE('----- ----- ----- ----- -----');
    LOOP
        v_status := DBMS_SQL.FETCH_ROWS(curid);
        EXIT WHEN v_status = 0;
        DBMS_SQL.COLUMN_VALUE(curid,1,v_empno);
        DBMS_SQL.COLUMN_VALUE(curid,2,v_ename);
        DBMS_SQL.COLUMN_VALUE(curid,3,v_hiredate);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
        DBMS_SQL.COLUMN_VALUE(curid,5,v_comm);
        DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || RPAD(v_ename,10) || ' ' ||
                             TO_CHAR(v_hiredate,'yyyy-mm-dd') || ' ' ||
                             TO_CHAR(v_sal,'9,999.99') || ' ' ||
                             TO_CHAR(NVL(v_comm,0),'9,999.99'));
    END LOOP;
    DBMS_SQL CLOSE_CURSOR(curid);
END;

EMPNO  ENAME      HIREDATE  SAL      COMM
-----
7369  SMITH      1980-12-17  800.00   .00
7499  ALLEN      1981-02-20  1,600.00  300.00
7521  WARD       1981-02-22  1,250.00  500.00
7566  JONES      1981-04-02  2,975.00  .00
7654  MARTIN     1981-09-28  1,250.00  1,400.00

```

| | | | | |
|------|--------|------------|----------|-----|
| 7698 | BLAKE | 1981-05-01 | 2,850.00 | .00 |
| 7782 | CLARK | 1981-06-09 | 2,450.00 | .00 |
| 7788 | SCOTT | 1987-04-19 | 3,000.00 | .00 |
| 7839 | KING | 1981-11-17 | 5,000.00 | .00 |
| 7844 | TURNER | 1981-09-08 | 1,500.00 | .00 |
| 7876 | ADAMS | 1987-05-23 | 1,100.00 | .00 |
| 7900 | JAMES | 1981-12-03 | 950.00 | .00 |
| 7902 | FORD | 1981-12-03 | 3,000.00 | .00 |
| 7934 | MILLER | 1982-01-23 | 1,300.00 | .00 |

6.3.17.15 IS_OPEN

The `IS_OPEN` function provides the capability to test if the given cursor is open.

```
<status> BOOLEAN IS_OPEN(<c> INTEGER)
```

Parameters

c

Cursor ID of the cursor to be tested.

status

Set to `TRUE` if the cursor is open, set to `FALSE` if the cursor is not open.

6.3.17.16 LAST_ROW_COUNT

The `LAST_ROW_COUNT` function returns the number of rows that have been currently fetched.

```
<rowcnt> INTEGER LAST_ROW_COUNT
```

Parameters

rowcnt

Number of row fetched thus far.

Examples

The following example uses the `LAST_ROW_COUNT` function to display the total number of rows fetched in the query.

```
DECLARE
    curid      sINTEGER;
    v_empno    NUMBER(4);
    v_ename    VARCHAR2(10);
```

```

v_hiredate DATE;
v_sal      NUMBER(7,2);
v_comm     NUMBER(7,2);
v_sql      VARCHAR2(50) := 'SELECT empno, ename, hiredate, sal, '' ||
                           'comm FROM emp';
v_status   INTEGER;
BEGIN
  curid := DBMS_SQL.OPEN_CURSOR;
  DBMS_SQLPARSE(curid,v_sql,DBMS_SQL.native);
  DBMS_SQL.DEFINE_COLUMN(curid,1,v_empno);
  DBMS_SQL.DEFINE_COLUMN(curid,2,v_ename,10);
  DBMS_SQL.DEFINE_COLUMN(curid,3,v_hiredate);
  DBMS_SQL.DEFINE_COLUMN(curid,4,v_sal);
  DBMS_SQL.DEFINE_COLUMN(curid,5,v_comm);

  v_status := DBMS_SQL.EXECUTE(curid);
  DBMS_OUTPUT.PUT_LINE('EMPNO  ENAME    HIREDATE  SAL    COMM');
  DBMS_OUTPUT.PUT_LINE('----- ----- ----- ----- ----- ' ||
  '-----');
LOOP
  v_status := DBMS_SQL.FETCH_ROWS(curid);
  EXIT WHEN v_status = 0;
  DBMS_SQL.COLUMN_VALUE(curid,1,v_empno);
  DBMS_SQL.COLUMN_VALUE(curid,2,v_ename);
  DBMS_SQL.COLUMN_VALUE(curid,3,v_hiredate);
  DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
  DBMS_SQL.COLUMN_VALUE(curid,4,v_sal);
  DBMS_SQL.COLUMN_VALUE(curid,5,v_comm);
  DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || RPAD(v_ename,10) || ' ' ||
    TO_CHAR(v_hiredate,'yyyy-mm-dd') || ' ' ||
    TO_CHAR(v_sal,'9,999.99') || ' ' ||
    TO_CHAR(NVL(v_comm,0),'9,999.99'));
END LOOP;
DBMS_OUTPUT.PUT_LINE('Number of rows: ' || DBMS_SQL.LAST_ROW_COUNT);
DBMS_SQL CLOSE_CURSOR(curid);
END;

```

| EMPNO | ENAME | HIREDATE | SAL | COMM |
|-------|--------|------------|----------|----------|
| 7369 | SMITH | 1980-12-17 | 800.00 | .00 |
| 7499 | ALLEN | 1981-02-20 | 1,600.00 | 300.00 |
| 7521 | WARD | 1981-02-22 | 1,250.00 | 500.00 |
| 7566 | JONES | 1981-04-02 | 2,975.00 | .00 |
| 7654 | MARTIN | 1981-09-28 | 1,250.00 | 1,400.00 |
| 7698 | BLAKE | 1981-05-01 | 2,850.00 | .00 |
| 7782 | CLARK | 1981-06-09 | 2,450.00 | .00 |
| 7788 | SCOTT | 1987-04-19 | 3,000.00 | .00 |
| 7839 | KING | 1981-11-17 | 5,000.00 | .00 |
| 7844 | TURNER | 1981-09-08 | 1,500.00 | .00 |
| 7876 | ADAMS | 1987-05-23 | 1,100.00 | .00 |
| 7900 | JAMES | 1981-12-03 | 950.00 | .00 |
| 7902 | FORD | 1981-12-03 | 3,000.00 | .00 |
| 7934 | MILLER | 1982-01-23 | 1,300.00 | .00 |

Number of rows: 14

6.3.17.17 OPEN_CURSOR

The `OPEN_CURSOR` function creates a new cursor. A cursor must be used to parse and execute any dynamic SQL statement. Once a cursor has been opened, it can be re-used with the same or different SQL statements. The cursor does not have to be closed and re-opened in order to be re-used.

```
<c> INTEGER OPEN_CURSOR
```

Parameters

c

Cursor ID number associated with the newly created cursor.

Examples

The following example creates a new cursor:

```
DECLARE
    curid      INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    .
    .
    .
END;
```

6.3.17.18 PARSE

The `PARSE` procedure parses a SQL command or SPL block. If the SQL command is a DDL command, it is immediately executed and does not require running the `EXECUTE` function.

```
PARSE(<c> INTEGER, <statement> VARCHAR2, <language_flag> INTEGER)
```

Parameters

c

Cursor ID of an open cursor.

`statement`

SQL command or SPL block to be parsed. A SQL command must not end with the semi-colon terminator, however an SPL block does require the semi-colon terminator.

`language_flag`

Language flag provided for compatibility with Oracle syntax. Use `DBMS_SQL.V6`, `DBMS_SQL.V7` or `DBMS_SQL.native`. This flag is ignored, and all syntax is assumed to be in EnterpriseDB Advanced Server form.

Examples

The following anonymous block creates a table named, `job`. Note that DDL statements are executed immediately by the `PARSE` procedure and do not require a separate `EXECUTE` step.

```
DECLARE
    curid      INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQL.PARSE(curid, 'CREATE TABLE job (jobno NUMBER(3), ' ||
        'jname VARCHAR2(9))',DBMS_SQL.native);
    DBMS_SQL CLOSE_CURSOR(curid);
END;
```

The following inserts two rows into the `job` table.

```
DECLARE
    curid      INTEGER;
    v_sql      VARCHAR2(50);
    v_status   INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    v_sql := 'INSERT INTO job VALUES (100, "ANALYST")';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    v_sql := 'INSERT INTO job VALUES (200, "CLERK")';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_OUTPUT.PUT_LINE('Number of rows processed: ' || v_status);
    DBMS_SQL CLOSE_CURSOR(curid);
END;
```

Number of rows processed: 1

Number of rows processed: 1

The following anonymous block uses the `DBMS_SQL` package to execute a block containing two `INSERT` statements. Note that the end of the block contains a terminating semi-colon, while in the prior example, each individual `INSERT` statement does not have a terminating semi-colon.

```
DECLARE
    curid      INTEGER;
    v_sql      VARCHAR2(100);
    v_status   INTEGER;
BEGIN
    curid := DBMS_SQL.OPEN_CURSOR;
    v_sql := 'BEGIN ' ||
        'INSERT INTO job VALUES (300, "MANAGER"); ' ||
        'INSERT INTO job VALUES (400, "SALESMAN"); ' ||
        'END;';
    DBMS_SQL.PARSE(curid, v_sql, DBMS_SQL.native);
    v_status := DBMS_SQL.EXECUTE(curid);
    DBMS_SQL CLOSE_CURSOR(curid);
END;
```

6.3.18 DBMS.Utility

The `DBMS.Utility` package provides support for the following various utility programs:

| Function/Procedure | Function or Procedure | Return Type | Description |
|---|-----------------------|----------------|--|
| <code>ANALYZE_DATABASE(method [, estimate_rows [, estimate_percent [, method_opt]]])</code> | Procedure | n/a | Analyze database tables. |
| <code>ANALYZE_PART_OBJECT(schema, object_name [, object_type [, command_type [, command_opt [, sample_clause]]]])</code> | Procedure | n/a | Analyze a partitioned table. |
| <code>ANALYZE_SCHEMA(schema, method [, estimate_rows [, estimate_percent [, method_opt]]])</code> | Procedure | n/a | Analyze schema tables. |
| <code>CANONICALIZE(name, canon_name OUT, canon_len)</code> | Procedure | n/a | Canonicalizes a string – e.g., strips off white space. |
| <code>COMMA_TO_TABLE(list, tablen OUT, tab OUT)</code> | Procedure | n/a | Convert a comma-delimited list of names to a table of names. |
| <code>DB_VERSION(version OUT, compatibility OUT)</code> | Procedure | n/a | Get the database version. |
| <code>EXEC_DDL_STATEMENT(parse_string)</code> | Procedure | n/a | Execute a DDL statement. |
| <code>FORMAT_CALL_STACK</code> | Function | TEXT | Formats the current call stack. |
| <code>GET_CPU_TIME</code> | Function | NUMBER | Get the current CPU time. |
| <code>GET_DEPENDENCY(type, schema, name)</code> | Procedure | n/a | Get objects that are dependent upon the given object.. |
| <code>GET_HASH_VALUE(name, base, hash_size)</code> | Function | NUMBER | Compute a hash value. |
| <code>GET_PARAMETER_VALUE(parnam, intval OUT, strval OUT)</code> | Procedure | BINARY_INTEGER | Get database initialization parameter settings. |
| <code>GET_TIME</code> | Function | NUMBER | Get the current time. |
| <code>NAME_TOKENIZE(name, a OUT, b OUT, c OUT, dblink OUT, nextpos OUT)</code> | Procedure | n/a | Parse the given name into its component parts. |
| <code>TABLE_TO_COMMA(tab, tablen OUT, list OUT)</code> | Procedure | n/a | Convert a table of names to a comma-delimited list. |

Advanced Server's implementation of `DBMS.Utility` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variables available in the `DBMS.Utility` package.

| Public Variables | Data Type | Value | Description |
|---------------------------|-------------|-------|-----------------------------------|
| inv_error_on_restrictions | PLS_INTEGER | 1 | Used by the INVALIDATE procedure. |
| Iname_array | TABLE | | For lists of long names. |
| uncl_array | TABLE | | For lists of users and names. |

LNAME_ARRAY

The `LNAME_ARRAY` is for storing lists of long names including fully-qualified names.

```
TYPE Iname_array IS `TABLE OF VARCHAR2(4000) INDEX BY BINARY_INTEGER;
```

UNCL_ARRAY

The `UNCL_ARRAY` is for storing lists of users and names.

```
TYPE uncl_array IS TABLE OF VARCHAR2(227) INDEX BY BINARY_INTEGER;
```

ANALYZE_DATABASE, ANALYZE_SCHEMA and ANALYZE_PART_OBJECT

The `ANALYZE_DATABASE()`, `ANALYZE_SCHEMA()` and `ANALYZE_PART_OBJECT()` procedures provide the capability to gather statistics on tables in the database. When you execute the `ANALYZE` statement, Postgres samples the data in a table and records distribution statistics in the `pg_statistics` system table.

`ANALYZE_DATABASE`, `ANALYZE_SCHEMA`, and `ANALYZE_PART_OBJECT` differ primarily in the number of tables that are processed:

- `ANALYZE DATABASE` analyzes all tables in all schemas within the current database.
- `ANALYZE SCHEMA` analyzes all tables in a given schema (within the current database).
- `ANALYZE_PART_OBJECT` analyzes a single table.

The syntax for the `ANALYZE` commands are:

```
ANALYZE_DATABASE(<method> VARCHAR2 [, <estimate_rows> NUMBER  
[, <estimate_percent> NUMBER [, <method_opt> VARCHAR2 ]]])
```

```
ANALYZE_SCHEMA(<schema> VARCHAR2, <method> VARCHAR2  
[, <estimate_rows> NUMBER [, <estimate_percent> NUMBER  
[, <method_opt> VARCHAR2 ]]])
```

```
ANALYZE_PART_OBJECT(<schema> VARCHAR2, <object_name> VARCHAR2  
[, <object_type> CHAR [, <command_type> CHAR  
[, <command_opt> VARCHAR2 [, <sample_clause> ]]]])
```

Parameters - `ANALYZE_DATABASE` and `ANALYZE_SCHEMA`

method

method determines whether the `ANALYZE` procedure populates the `pg_statistics` table or removes entries from the `pg_statistics` table. If you specify a method of `DELETE`, the `ANALYZE` procedure removes the relevant rows from `pg_statistics`. If you specify a method of `COMPUTE` or `ESTIMATE`, the `ANALYZE`

procedure analyzes a table (or multiple tables) and records the distribution information in `pg_statistics`. There is no difference between `COMPUTE` and `ESTIMATE`; both methods execute the Postgres `ANALYZE` statement. All other parameters are validated and then ignored.

`estimate_rows`

Number of rows upon which to base estimated statistics. One of `estimate_rows` or `estimate_percent` must be specified if method is `ESTIMATE`.

This argument is ignored, but is included for compatibility.

`estimate_percent`

Percentage of rows upon which to base estimated statistics. One of `estimate_rows` or `estimate_percent` must be specified if method is `ESTIMATE`.

This argument is ignored, but is included for compatibility.

`method_opt`

Object types to be analyzed. Any combination of the following:

[FOR TABLE]

[FOR ALL [INDEXED] COLUMNS] [SIZE n]

[FOR ALL INDEXES]

This argument is ignored, but is included for compatibility.

Parameters - `ANALYZE_PART_OBJECT`

`schema`

Name of the schema whose objects are to be analyzed.

`object_name`

Name of the partitioned object to be analyzed.

`object_type`

Type of object to be analyzed. Valid values are: `T` – table, `I` – index.

This argument is ignored, but is included for compatibility.

`command_type`

Type of analyze functionality to perform. Valid values are: `E` - gather estimated statistics based upon a specified number of rows or a percentage of rows in the `sample_clause` clause; `C` - compute exact statistics; or `V` – validate the structure and integrity of the partitions.

This argument is ignored, but is included for compatibility.

`command_opt`

For `command_type` `C` or `E`, can be any combination of:

[FOR TABLE]

[FOR ALL COLUMNS]

[FOR ALL LOCAL INDEXES]

For `command_type` V, can be `CASCADE` if `object_type` is T.

This argument is ignored, but is included for compatibility.

sample_clause

If `command_type` is E, contains the following clause to specify the number of rows or percentage or rows on which to base the estimate.

`SAMPLE n { ROWS | PERCENT }`

This argument is ignored, but is included for compatibility.

CANONICALIZE

The `CANONICALIZE` procedure performs the following operations on an input string:

- If the string is not double-quoted, verifies that it uses the characters of a legal identifier. If not, an exception is thrown. If the string is double-quoted, all characters are allowed.
- If the string is not double-quoted and does not contain periods, uppercases all alphabetic characters and eliminates leading and trailing spaces.
- If the string is double-quoted and does not contain periods, strips off the double quotes.
- If the string contains periods and no portion of the string is double-quoted, uppercases each portion of the string and encloses each portion in double quotes.
- If the string contains periods and portions of the string are double-quoted, returns the double-quoted portions unchanged including the double quotes and returns the non-double-quoted portions upercased and enclosed in double quotes.

```
CANONICALIZE(<name> VARCHAR2, <canon_name> OUT VARCHAR2,
<canon_len> BINARY_INTEGER)
```

Parameters

name

String to be canonicalized.

canon_name

The canonicalized string.

canon_len

Number of bytes in `name` to canonicalize starting from the first character.

Examples

The following procedure applies the `CANONICALIZE` procedure on its input parameter and displays the results.

```
CREATE OR REPLACE PROCEDURE canonicalize (
    p_name      VARCHAR2,
    p_length    BINARY_INTEGER DEFAULT 30
)
IS
    v_canon    VARCHAR2(100);
BEGIN
```

```

DBMS_UTILITY.CANONICALIZE(p_name,v_canon,p_length);
DBMS_OUTPUT.PUT_LINE('Canonicalized name ==>' || v_canon || '<==');
DBMS_OUTPUT.PUT_LINE('Length: ' || LENGTH(v_canon));
EXCEPTION
WHEN OTHERS THEN
DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

EXEC canonicalize('Identifier')
Canonicalized name ==>IDENTIFIER<==
Length: 10

EXEC canonicalize("Identifier")
Canonicalized name ==>Identifier<==
Length: 10

EXEC canonicalize("_+142%")
Canonicalized name ==>_+142%<==
Length: 6

EXEC canonicalize('abc.def.ghi')
Canonicalized name ==>"ABC"."DEF"."GHI"<==
Length: 17

EXEC canonicalize("abc.def.ghi")
Canonicalized name ==>abc.def.ghi<==
Length: 11

EXEC canonicalize("abc".def."ghi")
Canonicalized name ==>"abc"."DEF"."ghi"<==
Length: 17

EXEC canonicalize("abc.def".ghi)
Canonicalized name ==>"abc.def"."GHI"<==
Length: 15

```

COMMA_TO_TABLE

The `COMMA_TO_TABLE` procedure converts a comma-delimited list of names into a table of names. Each entry in the list becomes a table entry. The names must be formatted as valid identifiers.

```
COMMA_TO_TABLE(<list> VARCHAR2, <tablen> OUT BINARY_INTEGER,
<tab> OUT { LNAME_ARRAY | UNCL_ARRAY })
```

Parameters

`list`

Comma-delimited list of names.

`tablen`

Number of entries in `tab`.

tab

Table containing the individual names in [list](#).

LNAME_ARRAY

A [DBMS.Utility LNAME_ARRAY](#) (as described in the [LNAME_ARRAY](#) section).

UNCL_ARRAY

A [DBMS.Utility UNCL_ARRAY](#) (as described in the [UNCL_ARRAY](#) section).

Examples

The following procedure uses the [COMMA_TO_TABLE](#) procedure to convert a list of names to a table. The table entries are then displayed.

```
CREATE OR REPLACE PROCEDURE comma_to_table (
    p_list      VARCHAR2
)
IS
    r_lname      DBMS_UTILITY.LNAME_ARRAY;
    v_length     BINARY_INTEGER;
BEGIN
    DBMS_UTILITY.COMMA_TO_TABLE(p_list,v_length,r_lname);
    FOR i IN 1..v_length LOOP
        DBMS_OUTPUT.PUT_LINE(r_lname(i));
    END LOOP;
END;
```

EXEC comma_to_table('edb.dept, edb.emp, edb.jobhist')

edb.dept
edb.emp
edb.jobhist

DB_VERSION

The [DB_VERSION](#) procedure returns the version number of the database.

DB_VERSION(<version> OUT VARCHAR2, <compatibility> OUT VARCHAR2)

Parameters**version**

Database version number.

compatibility

Compatibility setting of the database. (To be implementation-defined as to its meaning.)

Examples

The following anonymous block displays the database version information.

DECLARE

```

v_version      VARCHAR2(150);
v_compat       VARCHAR2(150);
BEGIN
  DBMS_UTILITY.DB_VERSION(v_version,v_compat);
  DBMS_OUTPUT.PUT_LINE('Version: ' || v_version);
  DBMS_OUTPUT.PUT_LINE('Compatibility: ' || v_compat);
END;

```

Version: EnterpriseDB 10.0.0 on i686-pc-linux-gnu, compiled by GCC gcc (GCC) 4.1.2 20080704 (Red Hat 4.1.2-48), 32-bit
 Compatibility: EnterpriseDB 10.0.0 on i686-pc-linux-gnu, compiled by GCC gcc (GCC) 4.1.220080704 (Red Hat 4.1.2-48), 32-bit

EXEC_DDL_STATEMENT

The `EXEC_DDL_STATEMENT` provides the capability to execute a `DDL` command.

```
EXEC_DDL_STATEMENT(<parse_string>) VARCHAR2
```

Parameters

`parse_string`

The DDL command to be executed.

Examples

The following anonymous block creates the `job` table.

```

BEGIN
  DBMS_UTILITY.EXEC_DDL_STATEMENT(
    'CREATE TABLE job (' ||
    'jobno NUMBER(3,' ||
    'jname VARCHAR2(9))'
  );
END;

```

If the `parse_string` does not include a valid DDL statement, Advanced Server returns the following error:

```
edb=# exec dbms_utility.exec_ddl_statement('select rownum from dual');
ERROR: EDB-20001: 'parse_string' must be a valid DDL statement
```

In this case, Advanced Server's behavior differs from Oracle's; Oracle accepts the invalid `parse_string` without complaint.

FORMAT_CALL_STACK

The `FORMAT_CALL_STACK` function returns the formatted contents of the current call stack.

```
DBMS_UTILITY.FORMAT_CALL_STACK
return VARCHAR2
```

This function can be used in a stored procedure, function or package to return the current call stack in a readable format. This function is useful for debugging purposes.

GET_CPU_TIME

The `GET_CPU_TIME` function returns the CPU time in hundredths of a second from some arbitrary point in time.

```
<cpuTime> NUMBER GET_CPU_TIME
```

Parameters

`cpuTime`

Number of hundredths of a second of CPU time.

Examples

The following `SELECT` command retrieves the current CPU time, which is 603 hundredths of a second or .0603 seconds.

```
SELECT DBMS_UTILITY.GET_CPU_TIME FROM DUAL;
```

`get_cpu_time`

```
-----  
603
```

GET_DEPENDENCY

The `GET_DEPENDENCY` procedure provides the capability to list the objects that are dependent upon the specified object. `GET_DEPENDENCY` does not show dependencies for functions or procedures.

```
GET_DEPENDENCY(<type> VARCHAR2, <schema> VARCHAR2,  
<name> VARCHAR2)
```

Parameters

`type`

Name of the object type of `name`. Valid values are `INDEX`, `PACKAGE`, `PACKAGE BODY`, `SEQUENCE`, `TABLE`, `TRIGGER`, `TYPE` and `VIEW`.

`schema`

Name of the schema in which `name` exists.

`name`

Name of the object for which dependencies are to be obtained.

Examples

The following anonymous block finds dependencies on the `EMP` table.

```
BEGIN  
  DBMS_UTILITY.GET_DEPENDENCY('TABLE','public','EMP');  
END;
```

`DEPENDENCIES ON public.EMP`

```
-----  
*TABLE public.EMP()
```

```
* CONSTRAINT c public.emp()
* CONSTRAINT f public.emp()
* CONSTRAINT p public.emp()
* TYPE public.emp()
* CONSTRAINT c public.emp()
* CONSTRAINT f public.jobhist()
* VIEW .empname_view()
```

GET_HASH_VALUE

The `GET_HASH_VALUE` function provides the capability to compute a hash value for a given string.

```
<hash> NUMBER GET_HASH_VALUE(<name> VARCHAR2, <base> NUMBER,
<hash_size> NUMBER)
```

Parameters

`name`

The string for which a hash value is to be computed.

`base`

Starting value at which hash values are to be generated.

`hash_size`

The number of hash values for the desired hash table.

`hash`

The generated hash value.

Examples

The following anonymous block creates a table of hash values using the `ename` column of the `emp` table and then displays the key along with the hash value. The hash values start at 100 with a maximum of 1024 distinct values.

```
DECLARE
    v_hash      NUMBER;
    TYPE hash_tab IS TABLE OF NUMBER INDEX BY VARCHAR2(10);
    r_hash      HASH_TAB;
    CURSOR emp_cur IS SELECT ename FROM emp;
BEGIN
    FOR r_emp IN emp_cur LOOP
        r_hash(r_emp.ename) :=
            DBMS.Utility.GET_HASH_VALUE(r_emp.ename,100,1024);
    END LOOP;
    FOR r_emp IN emp_cur LOOP
        DBMS_OUTPUT.PUT_LINE(RPAD(r_emp.ename,10) || ' ' ||
            r_hash(r_emp.ename));
    END LOOP;
END;
```

SMITH 377

| | |
|--------|------|
| ALLEN | 740 |
| WARD | 718 |
| JONES | 131 |
| MARTIN | 176 |
| BLAKE | 568 |
| CLARK | 621 |
| SCOTT | 1097 |
| KING | 235 |
| TURNER | 850 |
| ADAMS | 156 |
| JAMES | 942 |
| FORD | 775 |
| MILLER | 148 |

GET_PARAMETER_VALUE

The `GET_PARAMETER_VALUE` procedure provides the capability to retrieve database initialization parameter settings.

```
<status> BINARY_INTEGER GET_PARAMETER_VALUE(<parnam> VARCHAR2,
<intval> OUT INTEGER, <strval> OUT VARCHAR2)``
```

Parameters

`parnam`

Name of the parameter whose value is to be returned. The parameters are listed in the `pg_settings` system view.

`intval`

Value of an integer parameter or the length of `strval`.

`strval`

Value of a string parameter.

`status`

Returns 0 if the parameter value is `INTEGER` or `BOOLEAN`. Returns 1 if the parameter value is a string.

Examples

The following anonymous block shows the values of two initialization parameters.

```
DECLARE
  v_intval  INTEGER;
  v_strval  VARCHAR2(80);
BEGIN
  DBMS_OUTPUT.PUT_LINE('max_fsm_pages' || ':' || v_intval);
  DBMS_OUTPUT.PUT_LINE('client_encoding' || ':' || v_strval);
END;
```

max_fsm_pages: 72625

```
client_encoding: SQL_ASCII
```

GET_TIME

The `GET_TIME` function provides the capability to return the current time in hundredths of a second.

```
<time> NUMBER GET_TIME
```

Parameters

`time`

Number of hundredths of a second from the time in which the program is started.

Examples

The following example shows calls to the `GET_TIME` function.

```
SELECT DBMS_UTILITY.GET_TIME FROM DUAL;
```

`get_time`

1555860

```
SELECT DBMS_UTILITY.GET_TIME FROM DUAL;
```

`get_time`

1556037

NAME_TOKENIZE

The `NAME_TOKENIZE` procedure parses a name into its component parts. Names without double quotes are uppercased. The double quotes are stripped from names with double quotes.

```
NAME_TOKENIZE(<name> VARCHAR2, <a> OUT VARCHAR2,
<b> OUT VARCHAR2, <c> OUT VARCHAR2, <dblink> OUT VARCHAR2,
<nextpos> OUT BINARY_INTEGER)
```

Parameters

`name`

String containing a name in the following format:

`a[.b[.c]][@dblink]`

`a`

Returns the leftmost component.

`b`

Returns the second component, if any.

C

Returns the third component, if any.

`dblink`

Returns the database link name.

`nextpos`

Position of the last character parsed in name.

Examples

The following stored procedure is used to display the returned parameter values of the `NAME_TOKENIZE` procedure for various names.

```
CREATE OR REPLACE PROCEDURE name_tokenize (
    p_name      VARCHAR2
)
IS
    v_a      VARCHAR2(30);
    v_b      VARCHAR2(30);
    v_c      VARCHAR2(30);
    v_dblink  VARCHAR2(30);
    v_nextpos BINARY_INTEGER;
BEGIN
    DBMS_OUTPUT.PUT_LINE('name : ' || p_name);
    DBMS_OUTPUT.PUT_LINE('a   : ' || v_a);
    DBMS_OUTPUT.PUT_LINE('b   : ' || v_b);
    DBMS_OUTPUT.PUT_LINE('c   : ' || v_c);
    DBMS_OUTPUT.PUT_LINE('dblink : ' || v_dblink);
    DBMS_OUTPUT.PUT_LINE('nextpos: ' || v_nextpos);
END;
```

Tokenize the name, `emp`:

```
BEGIN
    name_tokenize('emp');
END;
```

```
name : emp
a   : EMP
b   :
c   :
dblink :
nextpos: 3
```

Tokenize the name, `edb.list_emp`:

```
BEGIN
    name_tokenize('edb.list_emp');
END;
```

```
name : edb.list_emp
a   : EDB
```

```
b    : LIST_EMP
c    :
dblink :
nextpos: 12
```

Tokenize the name, "edb"."Emp_Admin".update_emp_sal :

```
BEGIN
  name_tokenize("edb"."Emp_Admin".update_emp_sal);
END;

name  :"edb"."Emp_Admin".update_emp_sal
a    :edb
b    :Emp_Admin
c    :UPDATE_EMP_SAL
dblink :
nextpos: 32
```

Tokenize the name `edb.emp@edb_dblink`:

```
BEGIN
  name_tokenize('edb.emp@edb_dblink');
END;

name  :edb.emp@edb_dblink
a    :EDB
b    :EMP
c    :
dblink :EDB_DBLINK
nextpos: 18
```

TABLE_TO_COMMA

The `TABLE_TO_COMMA` procedure converts table of names into a comma-delimited list of names. Each table entry becomes a list entry. The names must be formatted as valid identifiers.

```
TABLE_TO_COMMA(<tab> { LNAME_ARRAY | UNCL_ARRAY },
<tablen> OUT BINARY_INTEGER, <list> OUT VARCHAR2)
```

Parameters

`tab`

Table containing names.

`LNAME_ARRAY`

A `DBMS.Utility LNAME_ARRAY` (as described in the [LNAME ARRAY](#) section).

`UNCL_ARRAY`

A `DBMS.Utility UNCL_ARRAY` (as described the [UNCL_ARRAY](#) section).

`tablen`

Number of entries in `list`.

list

Comma-delimited list of names from `tab`.

Examples

The following example first uses the `COMMA_TO_TABLE` procedure to convert a comma-delimited list to a table. The `TABLE_TO_COMMA` procedure then converts the table back to a comma-delimited list that is displayed.

```
CREATE OR REPLACE PROCEDURE table_to_comma (
    p_list      VARCHAR2
)
IS
    r_lname     DBMS_UTILITY.LNAME_ARRAY;
    v_length    BINARY_INTEGER;
    v_listlen   BINARY_INTEGER;
    v_list      VARCHAR2(80);
BEGIN
    DBMS_UTILITY.COMMA_TO_TABLE(p_list,v_length,r_lname);
    DBMS_OUTPUT.PUT_LINE('Table Entries');
    DBMS_OUTPUT.PUT_LINE('-----');
    FOR i IN 1..v_length LOOP
        DBMS_OUTPUT.PUT_LINE(r_lname(i));
    END LOOP;
    DBMS_OUTPUT.PUT_LINE('-----');
    DBMS_UTILITY.TABLE_TO_COMMA(r_lname,v_listlen,v_list);
    DBMS_OUTPUT.PUT_LINE('Comma-Delimited List: ' || v_list);
END;

EXEC table_to_comma('edb.dept,edb.emp,edb.jobhist')
```

Table Entries

edb.dept
edb.emp
edb.jobhist

Comma-Delimited List: edb.dept, edb.emp, edb.jobhist

6.3.19 UTL_ENCODE

The `UTL_ENCODE` package provides a way to encode and decode data. Advanced Server supports the following functions and procedures:

| Function/Procedure | Return Type | Description |
|--------------------|-------------|--|
| BASE64_DECODE(r) | RAW | Use the BASE64_DECODE function to translate a Base64 encoded string to the original RAW value. |
| BASE64_ENCODE(r) | RAW | Use the BASE64_ENCODE function to translate a RAW string to an encoded Base64 value. |

| | | |
|--|----------|---|
| BASE64_ENCODE(loid) | TEXT | Use the BASE64_ENCODE function to translate a TEXT string to an encoded Base64 value. |
| MIMEHEADER_DECODE(buf) | VARCHAR2 | Use the MIMEHEADER_DECODE function to translate an encoded MIMEHEADER formatted string to its original value. |
| MIMEHEADER_ENCODE(buf, encode_charset, encoding) | VARCHAR2 | Use the MIMEHEADER_ENCODE function to convert and encode a string in MIMEHEADER format. |
| QUOTED_PRINTABLE_DECODE(r) | RAW | Use the QUOTED_PRINTABLE_DECODE function to translate an encoded string to a RAW value. |
| QUOTED_PRINTABLE_ENCODE(r) | RAW | Use the QUOTED_PRINTABLE_ENCODE function to translate an input string to a quoted-printable formatted RAW value. |
| TEXT_DECODE(buf, encode_charset, encoding) | VARCHAR2 | Use the TEXT_DECODE function to decode a string encoded by TEXT_ENCODE. |
| TEXT_ENCODE(buf, encode_charset, encoding) | VARCHAR2 | Use the TEXT_ENCODE function to translate a string to a user-specified character set, and then encode the string. |
| UUDECODE(r) | RAW | Use the UUDECODE function to translate a uuencode encoded string to a RAW value. |
| UUENCODE(r, type, filename, permission) | RAW | Use the UUENCODE function to translate a RAW string to an encoded uuencode value. |

6.3.19.1 BASE64_DECODE

Use the `BASE64 DECODE` function to translate a Base64 encoded string to the original value originally encoded by `BASE64_ENCODE`. The signature is:

```
BASE64_DECODE (<r> IN RAW)
```

This function returns a `RAW` value.

Parameters

r

| r is the string that contains the Base64 encoded data that will be translated to `RAW` form.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display `BYTEA` or `RAW` values onscreen in readable form. For more information, refer to the Postgres Core Documentation available at:

| <https://www.postgresql.org/docs/current/static/datatype-binary.html>

The following example first encodes (using `BASE64_ENCODE`), and then decodes (using `BASE64_DECODE`) a string that contains the text abc:

```
edb=# SELECT UTL_ENCODE.BASE64_ENCODE(CAST ('abc' AS RAW));
base64_encode
-----
```

```
YWJj
(1 row)
```

```
edb=# SELECT UTL_ENCODE.BASE64_DECODE(CAST ('YWJj' AS RAW));
base64_decode
-----
abc
(1 row)
```

6.3.19.2 BASE64_ENCODE

Use the `BASE64_ENCODE` function to translate and encode a string in Base64 format (as described in RFC 4648). This function can be useful when composing `MIME` email that you intend to send using the `UTL_SMTP` package. The `BASE64_ENCODE` function has two signatures:

```
BASE64_ENCODE(<r> IN RAW)
```

and

```
BASE64_ENCODE(<loid> IN OID)
```

This function returns a `RAW` value or an `OID`.

Parameters

`r`

`r` specifies the `RAW` string that will be translated to Base64.

`loid`

`loid` specifies the object ID of a large object that will be translated to Base64.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display `BYTEA` or `RAW` values onscreen in readable form. For more information, refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/current/static/datatype-binary.html>

The following example first encodes (using `BASE64_ENCODE`), and then decodes (using `BASE64_DECODE`) a string that contains the text `abc`:

```
edb=# SELECT UTL_ENCODE.BASE64_ENCODE(CAST ('abc' AS RAW));
base64_encode
-----
YWJj
(1 row)
```

```
edb=# SELECT UTL_ENCODE.BASE64_DECODE(CAST ('YWJj' AS RAW));
base64_decode
-----
abc
(1 row)
```

6.3.19.3 MIMEHEADER_DECODE

Use the `MIMEHEADER_DECODE` function to decode values that are encoded by the `MIMEHEADER_ENCODE` function. The signature is:

```
MIMEHEADER_DECODE(<buf> IN VARCHAR2)
```

This function returns a `VARCHAR2` value.

Parameters

`buf`

`buf` contains the value (encoded by `MIMEHEADER_ENCODE`) that will be decoded.

Examples

The following examples use the `MIMEHEADER_ENCODE` and `MIMEHEADER_DECODE` functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.MIMEHEADER_ENCODE('What is the date?') FROM DUAL;
mimeheader_encode
-----
```

```
=?UTF8?Q?What is the date??=
```

```
(1 row)
```

```
edb=# SELECT UTL_ENCODE.MIMEHEADER_DECODE('=?UTF8?Q?What is the date??=') FROM DUAL;
mimeheader_decode
-----
```

```
What is the date?
```

```
(1 row)
```

6.3.19.4 MIMEHEADER_ENCODE

Use the `MIMEHEADER_ENCODE` function to convert a string into mime header format, and then encode the string. The signature is:

```
MIMEHEADER_ENCODE(<buf> IN VARCHAR2, <encode_charset> IN VARCHAR2
DEFAULT NULL, <encoding> IN INTEGER DEFAULT NULL)
```

This function returns a `VARCHAR2` value.

Parameters

`buf`

`buf` contains the string that will be formatted and encoded. The string is a `VARCHAR2` value.

`encode_charset`

`encode_charset` specifies the character set to which the string will be converted before being formatted and encoded. The default value is `NULL`.

`encoding`

`encoding` specifies the encoding type used when encoding the string. You can specify:

- `Q` to enable quoted-printable encoding. If you do not specify a value, `MIMEHEADER_ENCODE` will use quoted-printable encoding.
- `B` to enable base-64 encoding.

Examples

The following examples use the `MIMEHEADER_ENCODE` and `MIMEHEADER_DECODE` functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.MIMEHEADER_ENCODE('What is the date?') FROM DUAL;
      mimeheader_encode
```

```
-----  
=?UTF8?Q?What is the date??=  
(1 row)
```

```
edb=# SELECT UTL_ENCODE.MIMEHEADER_DECODE('=?UTF8?Q?What is the date??=')  
FROM DUAL;  
      mimeheader_decode
```

```
-----  
What is the date?  
(1 row)
```

6.3.19.5 QUOTED_PRINTABLE_DECODE

Use the `QUOTED_PRINTABLE_DECODE` function to translate an encoded quoted-printable string into a decoded `RAW` string.

The signature is:

```
QUOTED_PRINTABLE_DECODE(<r> IN RAW)
```

This function returns a `RAW` value.

Parameters

`r`

`r` contains the encoded string that will be decoded. The string is a `RAW` value, encoded by `QUOTED_PRINTABLE_ENCODE`.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display `BYTEA` or `RAW` values onscreen in readable form. For more information, refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/current/static/datatype-binary.html>

The following example first encodes and then decodes a string:

```
edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_ENCODE('E=mc2') FROM DUAL;
quoted_printable_encode
-----
```

```
E=3Dmc2
(1 row)
```

```
edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_DECODE('E=3Dmc2') FROM DUAL;
quoted_printable_decode
-----
```

```
E=mc2
(1 row)
```

6.3.19.6 QUOTED_PRINTABLE_ENCODE

Use the `QUOTED_PRINTABLE_ENCODE` function to translate and encode a string in quoted-printable format. The signature is:

```
QUOTED_PRINTABLE_ENCODE(<r> IN RAW)
```

This function returns a `RAW` value.

Parameters

`r`

`r` contains the string (a `RAW` value) that will be encoded in a quoted-printable format.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display `BYTEA` or `RAW` values onscreen in readable form. For more information, refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/current/static/datatype-binary.html>

The following example first encodes and then decodes a string:

```
edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_ENCODE('E=mc2') FROM DUAL;
quoted_printable_encode
-----
E=3Dmc2
(1 row)

edb=# SELECT UTL_ENCODE.QUOTED_PRINTABLE_DECODE('E=3Dmc2') FROM DUAL;
quoted_printable_decode
-----
E=mc2
(1 row)
```

6.3.19.7 TEXT_DECODE

Use the `TEXT_DECODE` function to translate and decode an encoded string to the `VARCHAR2` value that was originally encoded by the `TEXT_ENCODE` function. The signature is:

```
TEXT_DECODE(<buf> IN VARCHAR2, <encode_charset> IN VARCHAR2 DEFAULT
NULL, <encoding> IN PLS_INTEGER DEFAULT NULL)
```

This function returns a `VARCHAR2` value.

Parameters

`buf`

`buf` contains the encoded string that will be translated to the original value encoded by `TEXT_ENCODE`.

`encode_charset`

`encode_charset` specifies the character set to which the string will be translated before encoding. The default value is `NULL`.

`encoding`

`encoding` specifies the encoding type used by `TEXT_DECODE`. Specify:

- `UTL_ENCODE.BASE64` to specify base-64 encoding.
- `UTL_ENCODE.QUOTED_PRINTABLE` to specify quoted printable encoding. This is the default.

Examples

The following example uses the `TEXT_ENCODE` and `TEXT_DECODE` functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.TEXT_ENCODE('What is the date?', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_encode
-----
V2hhCBpcyB0aGUgZGF0ZT8=
(1 row)
```

```
edb=# SELECT UTL_ENCODE.TEXT_DECODE('V2hhCBpcyB0aGUgZGF0ZT8=', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_decode
-----
```

What is the date?
(1 row)

6.3.19.8 TEXT_ENCODE

Use the `TEXT_ENCODE` function to translate a string to a user-specified character set, and then encode the string. The signature is:

```
TEXT_DECODE(<buf> IN VARCHAR2, <encode_charset> IN VARCHAR2 DEFAULT
NULL, <encoding> IN PLS_INTEGER DEFAULT NULL)
```

This function returns a `VARCHAR2` value.

Parameters

`buf`

`buf` contains the encoded string that will be translated to the specified character set and encoded by `TEXT_ENCODE`.

`encode_charset`

`encode_charset` specifies the character set to which the value will be translated before encoding. The default value is `NULL`.

`encoding`

`encoding` specifies the encoding type used by `TEXT_ENCODE`. Specify:

- `UTL_ENCODE.BASE64` to specify base-64 encoding.
- `UTL_ENCODE.QUOTED_PRINTABLE` to specify quoted printable encoding. This is the default.

Examples

The following example uses the `TEXT_ENCODE` and `TEXT_DECODE` functions to first encode, and then decode a string:

```
edb=# SELECT UTL_ENCODE.TEXT_ENCODE('What is the date?', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_encode
-----
```

V2hhCBpcyB0aGUgZGF0ZT8=
(1 row)

```
edb=# SELECT UTL_ENCODE.TEXT_DECODE('V2hhCBpcyB0aGUgZGF0ZT8=', 'BIG5',
UTL_ENCODE.BASE64) FROM DUAL;
text_decode
-----
```

What is the date?

(1 row)

6.3.19.9 UUDECODE

Use the `UUDECODE` function to translate and decode a uuencode encoded string to the `RAW` value that was originally encoded by the `UUENCODE` function. The signature is:

`UUDECODE(<r> IN RAW)`

This function returns a `RAW` value.

If you are using the Advanced Server `UUDECODE` function to decode uuencoded data that was created by the Oracle implementation of the `UTL_ENCODE.UUENCODE` function, then you must first set the Advanced Server configuration parameter `utl_encode.uudecode_redwood` to `TRUE` before invoking the Advanced Server `UUDECODE` function on the Oracle-created data. (For example, this situation may occur if you migrated Oracle tables containing uuencoded data to an Advanced Server database.)

The uuencoded data created by the Oracle version of the `UUENCODE` function results in a format that differs from the uuencoded data created by the Advanced Server `UUENCODE` function. As a result, attempting to use the Advanced Server `UUDECODE` function on the Oracle uuencoded data results in an error unless the configuration parameter `utl_encode.uudecode_redwood` is set to `TRUE`.

However, if you are using the Advanced Server `UUDECODE` function on uuencoded data created by the Advanced Server `UUENCODE` function, then `utl_encode.uudecode_redwood` must be set to `FALSE`, which is the default setting.

Parameters

`r`

`r` contains the uuencoded string that will be translated to `RAW`.

Examples

Before executing the following example, invoke the command:

`SET bytea_output = escape;`

This command instructs the server to escape any non-printable characters, and to display `BYTEA` or `RAW` values onscreen in readable form. For more information, refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/current/static/datatype-binary.html>

The following example uses `UUENCODE` and `UUDECODE` to first encode and then decode a string:

```
edb=# SET bytea_output = escape;
SET
edb=# SELECT UTL_ENCODE.UUENCODE('What is the date?') FROM DUAL;
          uuencode
-----
begin 0 uuencode.txt\01215VAA="!I<RIT:&4@9&%T93\\`012`012end\012
(1 row)
```

```
edb=# SELECT UTL_ENCODE.UUDECODE
edb# ('begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\`012`\012end\012')
edb# FROM DUAL;
uudecode
-----
What is the date?
(1 row)
```

6.3.19.10 UUENCODE

Use the `UUENCODE` function to translate `RAW` data into a uuencode formatted encoded string. The signature is:

```
UUENCODE(<r> IN RAW, <type> IN INTEGER DEFAULT 1, <filename> IN
VARCHAR2 DEFAULT NULL, <permission> IN VARCHAR2 DEFAULT NULL)
```

This function returns a `RAW` value.

Parameters

`r`

`r` contains the `RAW` string that will be translated to uuencode format.

`type`

`type` is an `INTEGER` value or constant that specifies the type of uuencoded string that will be returned; the default value is `1`. The possible values are:

| Value | Constant |
|-------|--------------|
| 1 | complete |
| 2 | header_piece |
| 3 | middle_piece |
| 4 | end_piece |

`filename`

`filename` is a `VARCHAR2` value that specifies the file name that you want to embed in the encoded form; if you do not specify a file name, `UUENCODE` will include a filename of `uuencode.txt` in the encoded form.

`permission`

`permission` is a `VARCHAR2` that specifies the permission mode; the default value is `NULL`.

Examples

Before executing the following example, invoke the command:

```
SET bytea_output = escape;
```

This command instructs the server to escape any non-printable characters, and to display `BYTEA` or `RAW` values onscreen in readable form. For more information, refer to the Postgres Core Documentation available at:

<https://www.postgresql.org/docs/current/static/datatype-binary.html>

The following example uses `UUENCODE` and `UUECODE` to first encode and then decode a string:

```
edb=# SET bytea_output = escape;
SET
edb=# SELECT UTL_ENCODE.UUENCODE('What is the date?') FROM DUAL;
uuencode
-----
begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\`012`\012end\012
(1 row)

edb=# SELECT UTL_ENCODE.UUDECODE
edb-# ('begin 0 uuencode.txt\01215VAA="!I<R!T:&4@9&%T93\\`012`\012end\012')
edb-# FROM DUAL;
uudecode
-----
What is the date?
(1 row)
```

6.3.20 UTL_FILE

The `UTL_FILE` package provides the capability to read from, and write to files on the operating system's file system. Non-superusers must be granted `EXECUTE` privilege on the `UTL_FILE` package by a superuser before using any of the functions or procedures in the package. For example the following command grants the privilege to user `mary`:

```
GRANT EXECUTE ON PACKAGE SYS.UTL_FILE TO mary;
```

Also, the operating system username, `enterprisedb`, must have the appropriate read and/or write permissions on the directories and files to be accessed using the `UTL_FILE` functions and procedures. If the required file permissions are not in place, an exception is thrown in the `UTL_FILE` function or procedure.

A handle to the file to be written to, or read from is used to reference the file. The *file handle* is defined by a public variable in the `UTL_FILE` package named, `UTL_FILE.FILE_TYPE`. A variable of type `FILE_TYPE` must be declared to receive the file handle returned by calling the `FOPEN` function. The file handle is then used for all subsequent operations on the file.

References to directories on the file system are done using the directory name or alias that is assigned to the directory using the `CREATE DIRECTORY` command. The procedures and functions available in the `UTL_FILE` package are listed in the following table:

| Function/Procedure | Return Type | Description |
|--|-------------|--|
| FCLOSE(file IN OUT) | n/a | Closes the specified file identified by file. |
| FCLOSE_ALL | n/a | Closes all open files. |
| FCOPY(location, filename, dest_dir, dest_file [, start_line [, end_line]]) | n/a | Copies filename in the directory identified by location to file, dest_file, in directory, dest_dir, starting from line, start_line, to line, end_line. |
| FFLUSH(file) | n/a | Forces data in the buffer to be written to disk in the file identified by file. |

| | | |
|--|-----------|---|
| FOPEN(location, filename, open_mode [, max_linesize]) | FILE_TYPE | Opens file, filename, in the directory identified by location. |
| FREMOVE(location, filename) | n/a | Removes the specified file from the file system. |
| FRENAME(location, filename, dest_dir, dest_file [, overwrite]) | n/a | Renames the specified file. |
| GET_LINE(file, buffer OUT) | n/a | Reads a line of text into variable, buffer, from the file identified by file. |
| IS_OPEN(file) | BOOLEAN | Determines whether or not the given file is open. |
| NEW_LINE(file [, lines]) | n/a | Writes an end-of-line character sequence into the file. |
| PUT(file, buffer) | n/a | Writes buffer to the given file. PUT does not write an end-of-line character sequence. |
| PUT_LINE(file, buffer) | n/a | Writes buffer to the given file. An end-of-line character sequence is added by the PUT_LINE procedure. |
| PUTF(file, format [, arg1] [, ...]) | n/a | Writes a formatted string to the given file. Up to five substitution parameters, arg1,...arg5 may be specified for replacement in format. |

Advanced Server's implementation of `UTL_FILE` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

UTL_FILE Exception Codes

If a call to a `UTL_FILE` procedure or function raises an exception, you can use the condition name to catch the exception. The `UTL_FILE` package reports the following exception codes compatible with Oracle databases:

| Exception Code | Condition name |
|----------------|---------------------|
| -29283 | invalid_operation |
| -29285 | write_error |
| -29284 | read_error |
| -29282 | invalid_filehandle |
| -29287 | invalid_maxlinesize |
| -29281 | invalid_mode |
| -29280 | invalid_path |

Setting File Permissions with `utl_file.umask`

When a `UTL_FILE` function or procedure creates a file, there are default file permissions as shown by the following.

```
-rw----- 1 enterprisebd enterprisebd 21 Jul 24 16:08 utlfile
```

Note that all permissions are denied on users belonging to the `enterprisebd` group as well as all other users. Only the `enterprisebd` user has read and write permissions on the created file.

If you wish to have a different set of file permissions on files created by the `UTL_FILE` functions and procedures, you can accomplish this by setting the `utl_file.umask` configuration parameter.

The `utl_file.umask` parameter sets the *file mode creation mask* or simply, the *mask*, in a manner similar to the Linux `umask` command. This is for usage only within the Advanced Server `UTL_FILE` package.

Note

The `utl_file.umask` parameter is not supported on Windows systems.

The value specified for `utl_file.umask` is a 3 or 4-character octal string that would be valid for the Linux `umask` command. The setting determines the permissions on files created by the `UTL_FILE` functions and procedures. (Refer to any information source regarding Linux or Unix systems for information on file permissions and the usage of the `umask` command.)

The following is an example of setting the file permissions with `utl_file.umask`.

First, set up the directory in the file system to be used by the `UTL_FILE` package. Be sure the operating system account, `enterprisedb` or `postgres`, whichever is applicable, can read and write in the directory.

```
mkdir /tmp/utldir
chmod 777 /tmp/utldir
```

The `CREATE DIRECTORY` command is issued in `psql` to create the directory database object using the file system directory created in the preceding step.

```
CREATE DIRECTORY utldir AS '/tmp/utldir';
```

Set the `utl_file.umask` configuration parameter. The following setting allows the file owner any permission. Group users and other users are permitted any permission except for the execute permission.

```
SET utl_file.umask TO '0011';
```

In the same session during which the `utl_file.umask` parameter is set to the desired value, run the `UTL_FILE` functions and procedures.

```
DECLARE
  v_utlfile    UTL_FILE.FILE_TYPE;
  v_directory   VARCHAR2(50) := 'utldir';
  v_filename    VARCHAR2(20) := 'utlfile';
BEGIN
  v_utlfile := UTL_FILE.FOPEN(v_directory, v_filename, 'w');
  UTL_FILE.PUT_LINE(v_utlfile, 'Simple one-line file');
  DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
  UTL_FILE.FCLOSE(v_utlfile);
END;
```

The permission settings on the resulting file show that group users and other users have read and write permissions on the file as well as the file owner.

```
$ pwd
/tmp/utldir
$ ls -l
total 4
-rw-rw-rw- 1 enterprisedb enterprisedb 21 Jul 24 16:04 utlfile
```

This parameter can also be set on a per role basis with the `ALTER ROLE` command, on a per database basis with the `ALTER DATABASE` command, or for the entire database server instance by setting it in the `postgresql.conf` file.

FCLOSE

The `FCLOSE` procedure closes an open file.

```
FCLOSE(<file> IN OUT FILE_TYPE)
```

Parameters

file

Variable of type `FILE_TYPE` containing a file handle of the file to be closed.

FCLOSE_ALL

The `FCLOSE_ALL` procedures closes all open files. The procedure executes successfully even if there are no open files to close.

FCLOSE_ALL

FCOPY

The `FCOPY` procedure copies text from one file to another.

```
FCOPY(<location> VARCHAR2, <filename> VARCHAR2,
      <dest_dir> VARCHAR2, <dest_file> VARCHAR2
      [, <start_line> PLS_INTEGER [, <end_line> PLS_INTEGER ] ])
```

Parameters

location

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be copied.

filename

Name of the source file to be copied.

dest_dir

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory to which the file is to be copied.

dest_file

Name of the destination file.

start_line

Line number in the source file from which copying will begin. The default is 1.

end_line

Line number of the last line in the source file to be copied. If omitted or null, copying will go to the last line of the file.

Examples

The following makes a copy of a file, `C:\TEMP\EMPDIR\empfile.csv`, containing a comma-delimited list of employees from the `emp` table. The copy, `empcopy.csv`, is then listed.

```
CREATE DIRECTORY empdir AS 'C:/TEMP/EMPDIR';
```

```
DECLARE
```

```

v_empfile    UTL_FILE.FILE_TYPE;
v_src_dir    VARCHAR2(50) := 'empdir';
v_src_file   VARCHAR2(20) := 'empfile.csv';
v_dest_dir   VARCHAR2(50) := 'empdir';
v_dest_file  VARCHAR2(20) := 'empcopy.csv';
v_emprec     VARCHAR2(120);
v_count      INTEGER := 0;
BEGIN
  UTL_FILE.FCOPY(v_src_dir,v_src_file,v_dest_dir,v_dest_file);
  v_empfile := UTL_FILE.FOPEN(v_dest_dir,v_dest_file,'r');
  DBMS_OUTPUT.PUT_LINE('The following is the destination file, "' ||
    v_dest_file || "'");
LOOP
  UTL_FILE.GET_LINE(v_empfile,v_emprec);
  DBMS_OUTPUT.PUT_LINE(v_emprec);
  v_count := v_count + 1;
END LOOP;
EXCEPTION
  WHEN NO_DATA_FOUND THEN
    UTL_FILE.FCLOSE(v_empfile);
    DBMS_OUTPUT.PUT_LINE(v_count || ' records retrieved');
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
    DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

```

The following is the destination file, 'empcopy.csv'

7369,SMITH,CLERK,7902,17-DEC-80,800,,20
 7499,ALLEN,SALESMAN,7698,20-FEB-81,1600,300,30
 7521,WARD,SALESMAN,7698,22-FEB-81,1250,500,30
 7566,JONES,MANAGER,7839,02-APR-81,2975,,20
 7654,MARTIN,SALESMAN,7698,28-SEP-81,1250,1400,30
 7698,BLAKE,MANAGER,7839,01-MAY-81,2850,,30
 7782,CLARK,MANAGER,7839,09-JUN-81,2450,,10
 7788,SCOTT,ANALYST,7566,19-APR-87,3000,,20
 7839,KING,PRESIDENT,,17-NOV-81,5000,,10
 7844,TURNER,SALESMAN,7698,08-SEP-81,1500,0,30
 7876,ADAMS,CLERK,7788,23-MAY-87,1100,,20
 7900,JAMES,CLERK,7698,03-DEC-81,950,,30
 7902,FORD,ANALYST,7566,03-DEC-81,3000,,20
 7934,MILLER,CLERK,7782,23-JAN-82,1300,,10
 14 records retrieved

FFLUSH

The **FFLUSH** procedure flushes unwritten data from the write buffer to the file.

FFLUSH(<file> FILE_TYPE)

Parameters

file

Variable of type **FILE_TYPE** containing a file handle.

Examples

Each line is flushed after the `NEW_LINE` procedure is called.

```

DECLARE
  v_empfile    UTL_FILE.FILE_TYPE;
  v_directory   VARCHAR2(50) := 'empdir';
  v_filename    VARCHAR2(20) := 'empfile.csv';
  CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
  v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
  FOR i IN emp_cur LOOP
    UTL_FILE.PUT(v_empfile,i.empno);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.ename);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.job);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.mgr);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.hiredate);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.sal);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.comm);
    UTL_FILE.PUT(v_empfile,';');
    UTL_FILE.PUT(v_empfile,i.deptno);
    UTL_FILE.NEW_LINE(v_empfile);
    UTL_FILE.FFLUSH(v_empfile);
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
  UTL_FILE.FCLOSE(v_empfile);
END;

```

FOPEN

The `FOPEN` function opens a file for I/O.

```

filetype FILE_TYPE FOPEN(<location> VARCHAR2,
<filename> VARCHAR2,<open_mode> VARCHAR2
[, <max_linesize> BINARY_INTEGER ])

```

Parameters

location

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be opened.

filename

Name of the file to be opened.

open_mode

Mode in which the file will be opened. Modes are: `a` - append to file; `r` - read from file; `w` - write to file.

max_linesize

Maximum size of a line in characters. In read mode, an exception is thrown if an attempt is made to read a line exceeding `max_linesize`. In write and append modes, an exception is thrown if an attempt is made to write a line exceeding `max_linesize`. The end-of-line character(s) are not included in determining if the maximum line size is exceeded. This behavior is not compatible with Oracle databases; Oracle does count the end-of-line character(s).

filetype

Variable of type `FILE_TYPE` containing the file handle of the opened file.

FREMOVE

The `FREMOVE` procedure removes a file from the system.

```
FREMOVE(<location> VARCHAR2, <filename> VARCHAR2)
```

An exception is thrown if the file to be removed does not exist.

Parameters

location

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be removed.

filename

Name of the file to be removed.

Examples

The following removes file `empfile.csv`.

```
DECLARE
    v_directory    VARCHAR2(50) := 'empdir';
    v_filename     VARCHAR2(20) := 'empfile.csv';
BEGIN
    UTL_FILE.FREMOVE(v_directory,v_filename);
    DBMS_OUTPUT.PUT_LINE('Removed file: ' || v_filename);
EXCEPTION
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
        DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;
```

Removed file: empfile.csv

FRENAME

The `FRENAME` procedure renames a given file. This effectively moves a file from one location to another.

```
FRENAME(<location> VARCHAR2, <filename> VARCHAR2,
<dest_dir> VARCHAR2, <dest_file> VARCHAR2,
[ <overwrite> BOOLEAN ])
```

Parameters

location

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory containing the file to be renamed.

filename

Name of the source file to be renamed.

dest_dir

Directory name, as stored in `pg_catalog.edb_dir.dirname`, of the directory to which the renamed file is to exist.

dest_file

New name of the original file.

overwrite

Replaces any existing file named `dest_file` in `dest_dir` if set to `TRUE`, otherwise an exception is thrown if set to `FALSE`. This is the default.

Examples

The following renames a file, `C:\TEMP\EMPDIR\empfile.csv`, containing a comma-delimited list of employees from the `emp` table. The renamed file, `C:\TEMP\NEWDIR\newemp.csv`, is then listed.

```
CREATE DIRECTORY "newdir" AS 'C:/TEMP/NEWDIR';

DECLARE
    v_empfile    UTL_FILE.FILE_TYPE;
    v_src_dir    VARCHAR2(50) := 'empdir';
    v_src_file   VARCHAR2(20) := 'empfile.csv';
    v_dest_dir   VARCHAR2(50) := 'newdir';
    v_dest_file  VARCHAR2(50) := 'newemp.csv';
    v_replace    BOOLEAN := FALSE;
    v_emprec     VARCHAR2(120);
    v_count      INTEGER := 0;
BEGIN
    UTL_FILE.FRENAME(v_src_dir,v_src_file,v_dest_dir,
                      v_dest_file,v_replace);
    v_empfile := UTL_FILE.FOPEN(v_dest_dir,v_dest_file,'r');
    DBMS_OUTPUT.PUT_LINE('The following is the renamed file, "' ||
                         v_dest_file || "'");
    LOOP
        UTL_FILE.GET_LINE(v_empfile,v_emprec);
        DBMS_OUTPUT.PUT_LINE(v_emprec);
        v_count := v_count + 1;
    END LOOP;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            UTL_FILE.FCLOSE(v_empfile);
            DBMS_OUTPUT.PUT_LINE(v_count || ' records retrieved');
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
            DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
    END;
```

```
The following is the renamed file, 'newemp.csv'
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
14 records retrieved
```

GET_LINE

The `GET_LINE` procedure reads a line of text from a given file up to, but not including the end-of-line terminator. A `NO_DATA_FOUND` exception is thrown when there are no more lines to read.

```
GET_LINE(<file> FILE_TYPE, <buffer> OUT VARCHAR2)
```

Parameters

`file`

Variable of type `FILE_TYPE` containing the file handle of the opened file.

`buffer`

Variable to receive a line from the file.

Examples

The following anonymous block reads through and displays the records in file `empfile.csv`.

```
DECLARE
  v_empfile  UTL_FILE.FILE_TYPE;
  v_directory  VARCHAR2(50) := 'empdir';
  v_filename   VARCHAR2(20) := 'empfile.csv';
  v_emprec     VARCHAR2(120);
  v_count      INTEGER := 0;
BEGIN
  v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'r');
  LOOP
    UTL_FILE.GET_LINE(v_empfile,v_emprec);
    DBMS_OUTPUT.PUT_LINE(v_emprec);
    v_count := v_count + 1;
  END LOOP;
EXCEPTION
  WHEN NO_DATA_FOUND THEN
    UTL_FILE.FCLOSE(v_empfile);
```

```

DBMS_OUTPUT.PUT_LINE('End of file ' || v_filename || ' - ' ||
    v_count || ' records retrieved');
WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
    DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
End of file empfile.csv - 14 records retrieved

```

IS_OPEN

The **IS_OPEN** function determines whether or not the given file is open.

```
<status> BOOLEAN IS_OPEN(<file> FILE_TYPE)
```

Parameters

file

Variable of type **FILE_TYPE** containing the file handle of the file to be tested.

status

TRUE if the given file is open, **FALSE** otherwise.

NEW_LINE

The **NEW_LINE** procedure writes an end-of-line character sequence in the file.

```
NEW_LINE(<file> FILE_TYPE [, <lines> INTEGER ])
```

Parameters

file

Variable of type **FILE_TYPE** containing the file handle of the file to which end-of-line character sequences are to be written.

lines

Number of end-of-line character sequences to be written. The default is one.

Examples

A file containing a double-spaced list of employee records is written.

```

DECLARE
  v_empfile    UTL_FILE.FILE_TYPE;
  v_directory  VARCHAR2(50) := 'empdir';
  v_filename   VARCHAR2(20) := 'empfile.csv';
  CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
  v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
  FOR i IN emp_cur LOOP
    UTL_FILE.PUT(v_empfile,i.empno);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.ename);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.job);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.mgr);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.hiredate);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.sal);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.comm);
    UTL_FILE.PUT(v_empfile,',');
    UTL_FILE.PUT(v_empfile,i.deptno);
    UTL_FILE.NEW_LINE(v_empfile,2);
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
  UTL_FILE.FCLOSE(v_empfile);
END;

```

Created file: empfile.csv

This file is then displayed:

```

C:\TEMP\EMPDIR>TYPE empfile.csv

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20

7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30

7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30

7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20

7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30

7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30

7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10

7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20

```

```
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
```

PUT

The `PUT` procedure writes a string to the given file. No end-of-line character sequence is written at the end of the string. Use the `NEW_LINE` procedure to add an end-of-line character sequence.

```
PUT(<file> FILE_TYPE, <buffer> { DATE | NUMBER | TIMESTAMP |  
VARCHAR2 })
```

Parameters

`file`

Variable of type `FILE_TYPE` containing the file handle of the file to which the given string is to be written.

`buffer`

Text to be written to the specified file.

Examples

The following example uses the `PUT` procedure to create a comma-delimited file of employees from the `emp` table.

```
DECLARE  
v_empfile    UTL_FILE.FILE_TYPE;  
v_directory   VARCHAR2(50) := 'empdir';  
v_filename    VARCHAR2(20) := 'empfile.csv';  
CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;  
BEGIN  
v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');  
FOR i IN emp_cur LOOP  
    UTL_FILE.PUT(v_empfile,i.empno);  
    UTL_FILE.PUT(v_empfile,'>');  
    UTL_FILE.PUT(v_empfile,i.ename);  
    UTL_FILE.PUT(v_empfile,'>');  
    UTL_FILE.PUT(v_empfile,i.job);  
    UTL_FILE.PUT(v_empfile,'>');  
    UTL_FILE.PUT(v_empfile,i.mgr);  
    UTL_FILE.PUT(v_empfile,'>');  
    UTL_FILE.PUT(v_empfile,i.hiredate);  
    UTL_FILE.PUT(v_empfile,'>');  
    UTL_FILE.PUT(v_empfile,i.sal);
```

```

UTL_FILE.PUT(v_empfile,'');
UTL_FILE.PUT(v_empfile,i.comm);
UTL_FILE.PUT(v_empfile,' ');
UTL_FILE.PUT(v_empfile,i.deptno);
UTL_FILE.NEW_LINE(v_empfile);
END LOOP;
DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
UTL_FILE.FCLOSE(v_empfile);
END;

```

Created file: empfile.csv

The following is the contents of `empfile.csv` created above:

```
C:\TEMP\EMPDIR>TYPE empfile.csv

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10

```

PUT_LINE

The `PUT_LINE` procedure writes a single line to the given file including an end-of-line character sequence.

```
PUT_LINE(<file> FILE_TYPE,
        <buffer> {DATE|NUMBER|TIMESTAMP|VARCHAR2})
```

Parameters

`file`

Variable of type `FILE_TYPE` containing the file handle of the file to which the given line is to be written.

`buffer`

Text to be written to the specified file.

Examples

The following example uses the `PUT_LINE` procedure to create a comma-delimited file of employees from the `emp` table.

```
DECLARE
    v_empfile    UTL_FILE.FILE_TYPE;
```

```

v_directory  VARCHAR2(50) := 'empdir';
v_filename   VARCHAR2(20) := 'empfile.csv';
v_emprec     VARCHAR2(120);
CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
  v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
  FOR i IN emp_cur LOOP
    v_emprec := i.empno || ',' || i.ename || ',' || i.job || ',' ||
      NVL(LTRIM(TO_CHAR(i.mgr,'9999')),"") || ',' || i.hiredate ||
      ',' || i.sal || ',' ||
      NVL(LTRIM(TO_CHAR(i.comm,'9990.99')),"") || ',' || i.deptno;
    UTL_FILE.PUT_LINE(v_empfile,v_emprec);
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
  UTL_FILE.FCLOSE(v_empfile);
END;

```

The following is the contents of `empfile.csv` created above:

```

C:\TEMP\EMPDIR>TYPE empfile.csv

7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10

```

PUTF

The `PUTF` procedure writes a formatted string to the given file.

```

PUTF(<file> FILE_TYPE, <format> VARCHAR2 [, <arg1> VARCHAR2]
 [, ...])

```

Parameters

`file`

Variable of type `FILE_TYPE` containing the file handle of the file to which the formatted line is to be written.

`format`

String to format the text written to the file. The special character sequence, `%s`, is substituted by the value of arg. The special character sequence, `\n`, indicates a new line. Note, however, in Advanced Server, a new line character must be specified with two consecutive backslashes instead of one - `\n`. This characteristic is not compatible with Oracle databases.

arg1

Up to five arguments, `arg1`, ..., `arg5`, to be substituted in the format string for each occurrence of `%s`. The first arg is substituted for the first occurrence of `%s`, the second arg is substituted for the second occurrence of `%s`, etc.

Examples

The following anonymous block produces formatted output containing data from the `emp` table. Note the use of the E literal syntax and double backslashes for the new line character sequence in the format string which are not compatible with Oracle databases.

```
DECLARE
  v_empfile  UTL_FILE.FILE_TYPE;
  v_directory VARCHAR2(50) := 'empdir';
  v_filename  VARCHAR2(20) := 'empfile.csv';
  v_format    VARCHAR2(200);
  CURSOR emp_cur IS SELECT * FROM emp ORDER BY empno;
BEGIN
  v_format := E'%s %s, %s\\nSalary: $%s Commission: $%s\\n\\n';
  v_empfile := UTL_FILE.FOPEN(v_directory,v_filename,'w');
  FOR i IN emp_cur LOOP
    UTL_FILE.PUTF(v_empfile,v_format,i.empno,i.ename,i.job,i.sal,
      NVL(i.comm,0));
  END LOOP;
  DBMS_OUTPUT.PUT_LINE('Created file: ' || v_filename);
  UTL_FILE.FCLOSE(v_empfile);
EXCEPTION
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('SQLERRM: ' || SQLERRM);
    DBMS_OUTPUT.PUT_LINE('SQLCODE: ' || SQLCODE);
END;
```

Created file: empfile.csv

The following is the contents of `empfile.csv` created above:

```
C:\TEMP\EMPDIR>TYPE empfile.csv
7369 SMITH, CLERK
Salary: $800.00 Commission: $0
7499 ALLEN, SALESMAN
Salary: $1600.00 Commission: $300.00
7521 WARD, SALESMAN
Salary: $1250.00 Commission: $500.00
7566 JONES, MANAGER
Salary: $2975.00 Commission: $0
7654 MARTIN, SALESMAN
Salary: $1250.00 Commission: $1400.00
7698 BLAKE, MANAGER
Salary: $2850.00 Commission: $0
7782 CLARK, MANAGER
Salary: $2450.00 Commission: $0
7788 SCOTT, ANALYST
Salary: $3000.00 Commission: $0
7839 KING, PRESIDENT
Salary: $5000.00 Commission: $0
7844 TURNER, SALESMAN
```

```

Salary: $1500.00 Commission: $0.00
7876 ADAMS, CLERK
Salary: $1100.00 Commission: $0
7900 JAMES, CLERK
Salary: $950.00 Commission: $0
7902 FORD, ANALYST
Salary: $3000.00 Commission: $0
7934 MILLER, CLERK
Salary: $1300.00 Commission: $0

```

6.3.21 UTL_HTTP

The `UTL_HTTP` package provides a way to use the HTTP or HTTPS protocol to retrieve information found at an URL. Advanced Server supports the following functions and procedures:

| Function/Procedure | Return Type | Description |
|--|----------------------|---|
| BEGIN_REQUEST(url, method, http_version) | UTL_HTTP.REQ | Initiates a new HTTP request. |
| END_REQUEST(r IN OUT) | n/a | Ends an HTTP request before allowing it to complete. |
| END_RESPONSE(r IN OUT) | n/a | Ends the HTTP response. |
| END_OF_BODY(r IN OUT) | n/a | Ends package body |
| GET_BODY_CHARSET | VARCHAR2 | Returns the default character set of the body of future HTTP requests. |
| GET_BODY_CHARSET(charset OUT) | n/a | Returns the default character set of the body of future HTTP requests. |
| GET_FOLLOW_REDIRECT(max_redirects OUT) | n/a | Current setting for the maximum number of redirections allowed. |
| GET_HEADER(r IN OUT, n, name OUT, value OUT) | n/a | Returns the nth header of the HTTP response. |
| GET_HEADER_BY_NAME(r IN OUT, name, value OUT, n) | n/a | Returns the HTTP response header for the specified name. |
| GET_HEADER_COUNT(r IN OUT) | INTEGER | Returns the number of HTTP response headers. |
| GET_RESPONSE(r IN OUT) | UTL_HTTP.RESP | Returns the HTTP response. |
| GET_RESPONSE_ERROR_CHECK(enable OUT) | n/a | Returns whether or not response error check is set. |
| GET_TRANSFER_TIMEOUT(timeout OUT) | n/a | Returns the transfer timeout setting for HTTP requests. |
| READ_LINE(r IN OUT, data OUT, remove_crlf) | n/a | Returns the HTTP response body in text form until the end of line. |
| READ_RAW(r IN OUT, data OUT, len) | n/a | Returns the HTTP response body in binary form for a specified number of bytes. |
| READ_TEXT(r IN OUT, data OUT, len) | n/a | Returns the HTTP response body in text form for a specified number of characters. |
| REQUEST(url) | VARCHAR2 | Returns the content of a web page. |
| REQUEST_PIECES(url, max_pieces) | UTL_HTTP.HTML_PIECES | Returns a table of 2000-byte segments retrieved from an URL. |

| | | |
|--|-----|--|
| SET_BODY_CHARSET(charset) | n/a | Sets the default character set of the body of future HTTP requests. |
| SET_FOLLOW_REDIRECT(max_redirects) | n/a | Sets the maximum number of times to follow the redirect instruction. |
| SET_FOLLOW_REDIRECT(r IN OUT, max_redirects) | n/a | Sets the maximum number of times to follow the redirect instruction for an individual request. |
| SET_HEADER(r IN OUT, name, value) | n/a | Sets the HTTP request header. |
| SET_RESPONSE_ERROR_CHECK(enable) | n/a | Determines whether or not HTTP 4xx and 5xx status codes are to be treated as errors. |
| SET_TRANSFER_TIMEOUT(timeout) | n/a | Sets the default, transfer timeout value for HTTP requests. |
| SET_TRANSFER_TIMEOUT(r IN OUT, timeout) | n/a | Sets the transfer timeout value for an individual HTTP request. |
| WRITE_LINE(r IN OUT, data) | n/a | Writes CRLF terminated data to the HTTP request body in TEXT form. |
| WRITE_RAW(r IN OUT, data) | n/a | Writes data to the HTTP request body in BINARY form. |
| WRITE_TEXT(r IN OUT, data) | n/a | Writes data to the HTTP request body in TEXT form. |

Advanced Server's implementation of `UTL_HTTP` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

Note

In Advanced Server, an `HTTP 4xx` or `HTTP 5xx` response produces a database error; in Oracle, this is configurable but `FALSE` by default.

In Advanced Server, the `UTL_HTTP` text interfaces expect the downloaded data to be in the database encoding. All currently-available interfaces are text interfaces. In Oracle, the encoding is detected from HTTP headers; in the absence of the header, the default is configurable and defaults to `ISO-8859-1`.

Advanced Server ignores all cookies it receives.

The `UTL_HTTP` exceptions that can be raised in Oracle are not recognized by Advanced Server. In addition, the error codes returned by Advanced Server are not the same as those returned by Oracle.

UTL_HTTP Exception Codes

If a call to a `UTL_HTTP` procedure or function raises an exception, you can use the condition name to catch the exception. The `UTL_HTTP` package reports the following exception codes compatible with Oracle databases:

| Exception Code | Condition Name | Description | Raised Where |
|----------------|----------------|--|--|
| -29266 | END_OF_BODY | The end of HTTP response body is reached | READ_LINE, READ_RAW, and READ_TEXT functions |

To use the `UTL_HTTP.END_OF_BODY` exception, first you need to run the `utl_http_public.sql` file from the `contrib/utl_http` directory of your installation directory.

There are various public constants available with `UTL_HTTP`. These are listed in the following tables.

The following table contains `UTL_HTTP` public constants defining HTTP versions and port assignments.

HTTP VERSIONS

HTTP VERSIONS

| | |
|------------------|--------------------------------------|
| HTTP_VERSION_1_0 | CONSTANT VARCHAR2(64) := 'HTTP/1.0'; |
|------------------|--------------------------------------|

| | |
|------------------|--------------------------------------|
| HTTP_VERSION_1_1 | CONSTANT VARCHAR2(64) := 'HTTP/1.1'; |
|------------------|--------------------------------------|

STANDARD PORT ASSIGNMENTS

| | |
|-------------------|-------------------------|
| DEFAULT_HTTP_PORT | CONSTANT INTEGER := 80; |
|-------------------|-------------------------|

| | |
|--------------------|--------------------------|
| DEFAULT_HTTPS_PORT | CONSTANT INTEGER := 443; |
|--------------------|--------------------------|

The following table contains **UTL_HTTP** public status code constants.

1XX INFORMATIONAL

| | |
|---------------|--------------------------|
| HTTP_CONTINUE | CONSTANT INTEGER := 100; |
|---------------|--------------------------|

| | |
|--------------------------|--------------------------|
| HTTP_SWITCHING_PROTOCOLS | CONSTANT INTEGER := 101; |
|--------------------------|--------------------------|

| | |
|-----------------|--------------------------|
| HTTP_PROCESSING | CONSTANT INTEGER := 102; |
|-----------------|--------------------------|

2XX SUCCESS

| | |
|---------|--------------------------|
| HTTP_OK | CONSTANT INTEGER := 200; |
|---------|--------------------------|

| | |
|--------------|--------------------------|
| HTTP_CREATED | CONSTANT INTEGER := 201; |
|--------------|--------------------------|

| | |
|---------------|--------------------------|
| HTTP_ACCEPTED | CONSTANT INTEGER := 202; |
|---------------|--------------------------|

| | |
|-----------------------------|--------------------------|
| HTTP_NON_AUTHORITATIVE_INFO | CONSTANT INTEGER := 203; |
|-----------------------------|--------------------------|

| | |
|-----------------|--------------------------|
| HTTP_NO_CONTENT | CONSTANT INTEGER := 204; |
|-----------------|--------------------------|

| | |
|--------------------|--------------------------|
| HTTP_RESET_CONTENT | CONSTANT INTEGER := 205; |
|--------------------|--------------------------|

| | |
|----------------------|--------------------------|
| HTTP_PARTIAL_CONTENT | CONSTANT INTEGER := 206; |
|----------------------|--------------------------|

| | |
|-------------------|--------------------------|
| HTTP_MULTI_STATUS | CONSTANT INTEGER := 207; |
|-------------------|--------------------------|

| | |
|-----------------------|--------------------------|
| HTTP_ALREADY_REPORTED | CONSTANT INTEGER := 208; |
|-----------------------|--------------------------|

| | |
|--------------|--------------------------|
| HTTP_IM_USED | CONSTANT INTEGER := 226; |
|--------------|--------------------------|

3XX REDIRECTION

| | |
|-----------------------|--------------------------|
| HTTP_MULTIPLE_CHOICES | CONSTANT INTEGER := 300; |
|-----------------------|--------------------------|

| | |
|------------------------|--------------------------|
| HTTP_MOVED_PERMANENTLY | CONSTANT INTEGER := 301; |
|------------------------|--------------------------|

| | |
|------------|--------------------------|
| HTTP_FOUND | CONSTANT INTEGER := 302; |
|------------|--------------------------|

| | |
|----------------|--------------------------|
| HTTP_SEE_OTHER | CONSTANT INTEGER := 303; |
|----------------|--------------------------|

| | |
|-------------------|--------------------------|
| HTTP_NOT_MODIFIED | CONSTANT INTEGER := 304; |
|-------------------|--------------------------|

| | |
|----------------|--------------------------|
| HTTP_USE_PROXY | CONSTANT INTEGER := 305; |
|----------------|--------------------------|

| | |
|-------------------|--------------------------|
| HTTP_SWITCH_PROXY | CONSTANT INTEGER := 306; |
|-------------------|--------------------------|

| | |
|-------------------------|--------------------------|
| HTTP_TEMPORARY_REDIRECT | CONSTANT INTEGER := 307; |
|-------------------------|--------------------------|

| | |
|-------------------------|--------------------------|
| HTTP_PERMANENT_REDIRECT | CONSTANT INTEGER := 308; |
|-------------------------|--------------------------|

4XX CLIENT ERROR

| | |
|------------------|--------------------------|
| HTTP_BAD_REQUEST | CONSTANT INTEGER := 400; |
|------------------|--------------------------|

| | |
|-------------------|--------------------------|
| HTTP_UNAUTHORIZED | CONSTANT INTEGER := 401; |
|-------------------|--------------------------|

| | |
|-----------------------|--------------------------|
| HTTP_PAYMENT_REQUIRED | CONSTANT INTEGER := 402; |
|-----------------------|--------------------------|

| | |
|----------------|--------------------------|
| HTTP_FORBIDDEN | CONSTANT INTEGER := 403; |
|----------------|--------------------------|

| | |
|----------------|--------------------------|
| HTTP_NOT_FOUND | CONSTANT INTEGER := 404; |
|----------------|--------------------------|

| | |
|-------------------------|--------------------------|
| HTTP_METHOD_NOT_ALLOWED | CONSTANT INTEGER := 405; |
|-------------------------|--------------------------|

| | |
|---------------------|--------------------------|
| HTTP_NOT_ACCEPTABLE | CONSTANT INTEGER := 406; |
|---------------------|--------------------------|

| | |
|--------------------------|--------------------------|
| HTTP_PROXY_AUTH_REQUIRED | CONSTANT INTEGER := 407; |
|--------------------------|--------------------------|

| | |
|-----------------------|--------------------------|
| HTTP_REQUEST_TIME_OUT | CONSTANT INTEGER := 408; |
|-----------------------|--------------------------|

| | |
|---------------|--------------------------|
| HTTP_CONFLICT | CONSTANT INTEGER := 409; |
|---------------|--------------------------|

| | |
|-----------|--------------------------|
| HTTP_GONE | CONSTANT INTEGER := 410; |
|-----------|--------------------------|

4XX CLIENT ERROR

| | |
|---|--------------------------|
| HTTP_LENGTH_REQUIRED | CONSTANT INTEGER := 411; |
| HTTP_PRECONDITION_FAILED | CONSTANT INTEGER := 412; |
| HTTP_REQUEST_ENTITY_TOO_LARGE | CONSTANT INTEGER := 413; |
| HTTP_REQUEST_URI_TOO_LARGE | CONSTANT INTEGER := 414; |
| HTTP_UNSUPPORTED_MEDIA_TYPE | CONSTANT INTEGER := 415; |
| HTTP_REQ_RANGE_NOT_SATISFIABLE | CONSTANT INTEGER := 416; |
| HTTP_EXPECTATION_FAILED | CONSTANT INTEGER := 417; |
| HTTP_I_AM_A_TEAPOT | CONSTANT INTEGER := 418; |
| HTTP_AUTHENTICATION_TIME_OUT | CONSTANT INTEGER := 419; |
| HTTP_ENHANCE_YOUR_CALM | CONSTANT INTEGER := 420; |
| HTTP_UNPROCESSABLE_ENTITY | CONSTANT INTEGER := 422; |
| HTTP_LOCKED | CONSTANT INTEGER := 423; |
| HTTP_FAILED_DEPENDENCY | CONSTANT INTEGER := 424; |
| HTTP_UNORDERED_COLLECTION | CONSTANT INTEGER := 425; |
| HTTP_UPGRADE_REQUIRED | CONSTANT INTEGER := 426; |
| HTTP_PRECONDITION_REQUIRED | CONSTANT INTEGER := 428; |
| HTTP_TOO_MANY_REQUESTS | CONSTANT INTEGER := 429; |
| HTTP_REQUEST_HEADER_FIELDS_TOO_LARGE | CONSTANT INTEGER := 431; |
| HTTP_NO_RESPONSE | CONSTANT INTEGER := 444; |
| HTTP_RETRY_WITH | CONSTANT INTEGER := 449; |
| HTTP_BLOCKED_BY_WINDOWS_PARENTAL_CONTROLS | CONSTANT INTEGER := 450; |
| HTTP_REDIRECT | CONSTANT INTEGER := 451; |
| HTTP_REQUEST_HEADER_TOO_LARGE | CONSTANT INTEGER := 494; |
| HTTP_CERT_ERROR | CONSTANT INTEGER := 495; |
| HTTP_NO_CERT | CONSTANT INTEGER := 496; |
| HTTP_HTTP_TO_HTTPS | CONSTANT INTEGER := 497; |
| HTTP_CLIENT_CLOSED_REQUEST | CONSTANT INTEGER := 499; |

5XX SERVER ERROR

| | |
|--------------------------------------|--------------------------|
| HTTP_INTERNAL_SERVER_ERROR | CONSTANT INTEGER := 500; |
| HTTP_NOT_IMPLEMENTED | CONSTANT INTEGER := 501; |
| HTTP_BAD_GATEWAY | CONSTANT INTEGER := 502; |
| HTTP_SERVICE_UNAVAILABLE | CONSTANT INTEGER := 503; |
| HTTP_GATEWAY_TIME_OUT | CONSTANT INTEGER := 504; |
| HTTP_VERSION_NOT_SUPPORTED | CONSTANT INTEGER := 505; |
| HTTP_VARIANT_ALSO_NEGOTIATES | CONSTANT INTEGER := 506; |
| HTTP_INSUFFICIENT_STORAGE | CONSTANT INTEGER := 507; |
| HTTP_LOOP_DETECTED | CONSTANT INTEGER := 508; |
| HTTP_BANDWIDTH_LIMIT_EXCEEDED | CONSTANT INTEGER := 509; |
| HTTP_NOT_EXTENDED | CONSTANT INTEGER := 510; |
| HTTP_NETWORK_AUTHENTICATION_REQUIRED | CONSTANT INTEGER := 511; |
| HTTP_NETWORK_READ_TIME_OUT_ERROR | CONSTANT INTEGER := 598; |
| HTTP_NETWORK_CONNECT_TIME_OUT_ERROR | CONSTANT INTEGER := 599; |

HTML_PIECES

The `UTL_HTTP` package declares a type named `HTML_PIECES`, which is a table of type `VARCHAR2(2000)` indexed by `BINARY_INTEGER`. A value of this type is returned by the `REQUEST_PIECES` function.

```
TYPE html_pieces IS TABLE OF VARCHAR2(2000) INDEX BY BINARY_INTEGER;
```

REQ

The `REQ` record type holds information about each HTTP request.

```
TYPE req IS RECORD (
    url      VARCHAR2(32767),    -- URL to be accessed
    method   VARCHAR2(64),      -- HTTP method
    http_version VARCHAR2(64),  -- HTTP version
    private_hdl INTEGER        -- Holds handle for this request
);
```

RESP

The `RESP` record type holds information about the response from each HTTP request.

```
TYPE resp IS RECORD (
    status_code  INTEGER,       -- HTTP status code
    reason_phrase VARCHAR2(256), -- HTTP response reason phrase
    http_version VARCHAR2(64),  -- HTTP version
    private_hdl INTEGER        -- Holds handle for this response
);
```

BEGIN_REQUEST

The `BEGIN_REQUEST` function initiates a new HTTP request. A network connection is established to the web server with the specified URL. The signature is:

```
BEGIN_REQUEST(<url> IN VARCHAR2, <method> IN VARCHAR2 DEFAULT
'GET ', <http_version> IN VARCHAR2 DEFAULT NULL) RETURN
UTL_HTTP.REQ
```

The `BEGIN_REQUEST` function returns a record of type `UTL_HTTP.REQ`.

Parameters

`url`

`url` is the Uniform Resource Locator from which `UTL_HTTP` will return content.

`method`

`method` is the HTTP method to be used. The default is `GET`.

`http_version`

`http_version` is the HTTP protocol version sending the request. The specified values should be either `HTTP/1.0` or `HTTP/1.1`. The default is null in which case the latest HTTP protocol version supported by the

| UTL_HTTP package is used which is 1.1.

END_REQUEST

The END_REQUEST procedure terminates an HTTP request. Use the END_REQUEST procedure to terminate an HTTP request without completing it and waiting for the response. The normal process is to begin the request, get the response, then close the response. The signature is:

```
END_REQUEST(<r> IN OUT UTL_HTTP.REQ)
```

Parameters

r

| r is the HTTP request record.

END_RESPONSE

The END_RESPONSE procedure terminates the HTTP response. The END_RESPONSE procedure completes the HTTP request and response. This is the normal method to end the request and response process. The signature is:

```
END_RESPONSE(<r> IN OUT UTL_HTTP.RESP)
```

Parameters

r

| r is the HTTP response record.

GET_BODY_CHARSET

The GET BODY CHARSET program is available in the form of both a procedure and a function. A call to GET_BODY_CHARSET returns the default character set of the body of future HTTP requests.

The procedure signature is:

```
GET_BODY_CHARSET(<charset> OUT VARCHAR2)
```

The function signature is:

```
GET_BODY_CHARSET() RETURN VARCHAR2
```

This function returns a VARCHAR2 value.

Parameters

charset

| charset is the character set of the body.

Examples

The following is an example of the GET_BODY_CHARSET function.

```
edb=# SELECT UTL_HTTP.GET_BODY_CHARSET() FROM DUAL;
get_body_charset
-----
ISO-8859-1
(1 row)
```

GET_FOLLOW_REDIRECT

The `GET_FOLLOW_REDIRECT` procedure returns the current setting for the maximum number of redirections allowed. The signature is:

```
GET_FOLLOW_REDIRECT(<max_redirects> OUT INTEGER)
```

Parameters

`max_redirects`

`max_redirects` is maximum number of redirections allowed.

GET_HEADER

The `GET_HEADER` procedure returns the `nth` header of the HTTP response. The signature is:

```
GET_HEADER(<r> IN OUT UTL_HTTP.RESP, <n> INTEGER, <name> OUT
VARCHAR2, <value> OUT VARCHAR2)
```

Parameters

`r`

`r` is the HTTP response record.

`n`

`n` is the `nth` header of the HTTP response record to retrieve.

`name`

`name` is the name of the response header.

`value`

`value` is the value of the response header.

Examples

The following example retrieves the header count, then the headers.

```
DECLARE
  v_req      UTL_HTTP.REQ;
  v_resp     UTL_HTTP.RESP;
  v_name     VARCHAR2(30);
  v_value    VARCHAR2(200);
  v_header_cnt  INTEGER;
BEGIN
```

```
-- Initiate request and get response
v_req := UTL_HTTP.BEGIN_REQUEST('www.enterprisedb.com');
v_resp := UTL_HTTP.GET_RESPONSE(v_req);

-- Get header count
v_header_cnt := UTL_HTTP.GET_HEADER_COUNT(v_resp);
DBMS_OUTPUT.PUT_LINE('Header Count: ' || v_header_cnt);

-- Get all headers
FOR i IN 1 .. v_header_cnt LOOP
    UTL_HTTP.GET_HEADER(v_resp, i, v_name, v_value);
    DBMS_OUTPUT.PUT_LINE(v_name || ': ' || v_value);
END LOOP;

-- Terminate request
UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output from the example.

```
Header Count: 23
Age: 570
Cache-Control: must-revalidate
Content-Type: text/html; charset=utf-8
Date: Wed, 30 Apr 2015 14:57:52 GMT
ETag: "aab02f2bd2d696eed817ca89ef411dda"
Expires: Sun, 19 Nov 1978 05:00:00 GMT
Last-Modified: Wed, 30 Apr 2015 14:15:49 GMT
RTSS: 1-1307-3
Server: Apache/2.2.3 (Red Hat)
Set-Cookie:
SESS2771d0952de2a1a84d322a262e0c173c=jn1u1j1etmdi5gg4lh8hakvs01;
expires=Fri, 23-May-2015 18:21:43 GMT; path=/; domain=.enterprisedb.com
Vary: Accept-Encoding
Via: 1.1 varnish
X-EDB-Backend: ec
X-EDB-Cache: HIT
X-EDB-Cache-Address: 10.31.162.212
X-EDB-Cache-Server: ip-10-31-162-212
X-EDB-Cache-TTL: 600.000
X-EDB-Cacheable: MAYBE: The user has a cookie of some sort. Maybe it's
double choc-chip!
X-EDB-Do-GZIP: false
X-Powered-By: PHP/5.2.17
X-Varnish: 484508634 484506789
transfer-encoding: chunked
Connection: keep-alive
```

GET_HEADER_BY_NAME

The **GET_HEADER_BY_NAME** procedure returns the header of the HTTP response according to the specified name. The signature is:

```
GET_HEADER_BY_NAME(<r> IN OUT UTL_HTTP.RESP, <name> VARCHAR2,
```

```
<value> OUT VARCHAR2, <n> INTEGER DEFAULT 1)
```

Parameters

r

| r is the HTTP response record.

name

| name is the name of the response header to retrieve.

value

| value is the value of the response header.

n

| n is the nth header of the HTTP response record to retrieve according to the values specified by name. The default is 1.

Examples

The following example retrieves the header for Content-Type.

```
DECLARE
  v_req      UTL_HTTP.REQ;
  v_resp     UTL_HTTP.RESP;
  v_name     VARCHAR2(30) := 'Content-Type';
  v_value    VARCHAR2(200);
BEGIN
  v_req := UTL_HTTP.BEGIN_REQUEST('www.enterprisedb.com');
  v_resp := UTL_HTTP.GET_RESPONSE(v_req);
  UTL_HTTP.GET_HEADER_BY_NAME(v_resp, v_name, v_value);
  DBMS_OUTPUT.PUT_LINE(v_name || ':' || v_value);
  UTL_HTTP.END_RESPONSE(v_resp);
END;
```

Content-Type: text/html; charset=utf-8

GET_HEADER_COUNT

The GET_HEADER_COUNT function returns the number of HTTP response headers. The signature is:

```
GET_HEADER_COUNT(<r> IN OUT UTL_HTTP.RESP) RETURN INTEGER
```

This function returns an INTEGER value.

Parameters

r

| r is the HTTP response record.

GET_RESPONSE

The `GET_RESPONSE` function sends the network request and returns any HTTP response. The signature is:

```
GET_RESPONSE(<r> IN OUT UTL_HTTP.REQ) RETURN UTL_HTTP.RESP
```

This function returns a `UTL_HTTP.RESP` record.

Parameters

`r`

`| r` is the HTTP request record.

GET_RESPONSE_ERROR_CHECK

The `GET_RESPONSE_ERROR_CHECK` procedure returns whether or not response error check is set. The signature is:

```
GET_RESPONSE_ERROR_CHECK(<enable> OUT BOOLEAN)
```

Parameters

`enable`

`| enable` returns `TRUE` if response error check is set, otherwise it returns `FALSE`.

GET_TRANSFER_TIMEOUT

The `GET_TRANSFER_TIMEOUT` procedure returns the current, default transfer timeout setting for HTTP requests. The signature is:

```
GET_TRANSFER_TIMEOUT(<timeout> OUT INTEGER)
```

Parameters

`timeout`

`| timeout` is the transfer timeout setting in seconds.

READ_LINE

The `READ_LINE` procedure returns the data from the HTTP response body in text form until the end of line is reached. A `CR` character, a `LF` character, a `CR LF` sequence, or the end of the response body constitutes the end of line. The signature is:

```
READ_LINE(<r> IN OUT UTL_HTTP.RESP, <data> OUT VARCHAR2,
<remove_crlf> BOOLEAN DEFAULT FALSE)
```

Parameters

`r`

`| r` is the HTTP response record.

`data`

| `data` is the response body in text form.

remove_crlf

| Set `remove_crlf` to `TRUE` to remove new line characters, otherwise set to `FALSE`. The default is `FALSE`.

Examples

The following example retrieves and displays the body of the specified website.

```
DECLARE
    v_req      UTL_HTTP.REQ;
    v_resp     UTL_HTTP.RESP;
    v_value    VARCHAR2(1024);
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.enterprisedb.com');
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    LOOP
        UTL_HTTP.READ_LINE(v_resp, v_value, TRUE);
        DBMS_OUTPUT.PUT_LINE(v_value);
    END LOOP;
    EXCEPTION
        WHEN OTHERS THEN
            UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output.

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" xml:lang="en" lang="en" dir="ltr">

<!-- _____ HEAD _____ -->

<head>
<meta http-equiv="Content-Type" content="text/html; charset=utf-8" />

<title>EnterpriseDB | The Postgres Database Company</title>

<meta http-equiv="Content-Type" content="text/html; charset=utf-8" />
<meta name="keywords" content="postgres, postgresql, postgresql installer,
mysql migration, open source database, training, replication" />
<meta name="description" content="The leader in open source database
products, services, support, training and expertise based on PostgreSQL.
Free downloads, documentation, and tutorials." />
<meta name="abstract" content="The Enterprise PostgreSQL Company" />
<link rel="EditURI" type="application/rsd+xml" title="RSD" href="http://
www.enterprisedb.com/blogapi/rsd" />
<link rel="alternate" type="application/rss+xml" title="EnterpriseDB RSS"
href="http://www.enterprisedb.com/rss.xml" />
<link rel="shortcut icon" href="/sites/all/themes/edb_pixelcrayons/
favicon.ico" type="image/x-icon" />
<link type="text/css" rel="stylesheet" media="all" href="/sites/default/
files/css/css_db11adabae0aed6b79a2c3c52def4754.css" />
<!--[if IE 6]>
<link type="text/css" rel="stylesheet" media="all" href="/sites/all/themes/
```

```

oho_basic/css/ie6.css?g" />
<![endif]-->
<!--[if IE 7]>
<link type="text/css" rel="stylesheet" media="all" href="/sites/all/themes/
oho_basic/css/ie7.css?g" />
<![endif]-->
<script type="text/javascript" src="/sites/default/files/js/
js_74d97b1176812e2fd6e43d62503a5204.js"></script>
<script type="text/javascript">
<!--//--><![CDATA[//><!--

```

READ_RAW

The `READ_RAW` procedure returns the data from the HTTP response body in binary form. The number of bytes returned is specified by the `len` parameter. The signature is:

```
READ_RAW(<r> IN OUT UTL_HTTP.RESP, <data> OUT RAW, <len> INTEGER)
```

Parameters

`r`

`r` is the HTTP response record.

`data`

`data` is the response body in binary form.

`len`

Set `len` to the number of bytes of data to be returned.

Examples

The following example retrieves and displays the first 150 bytes in binary form.

```

DECLARE
  v_req      UTL_HTTP.REQ;
  v_resp     UTL_HTTP.RESP;
  v_data     RAW;
BEGIN
  v_req := UTL_HTTP.BEGIN_REQUEST('http://www.enterprisedb.com');
  v_resp := UTL_HTTP.GET_RESPONSE(v_req);
  UTL_HTTP.READ_RAW(v_resp, v_data, 150);
  DBMS_OUTPUT.PUT_LINE(v_data);
  UTL_HTTP.END_RESPONSE(v_resp);
END;

```

The following is the output from the example.

```

\x3c21444f43545950452068746d6c205055424c494320222d2f2f5733432f2f4454442058485
44d4c20312e30205374726963742f2f454e220d0a202022687474703a2f2f7777772e77332e6f
72672f54522f7868746d6c312f4454442f7868746d6c312d7374726963742e647464223e0d0a3
c68746d6c20786d6c6e733d22687474703a2f2f7777772e77332e6f72672f313939392f

```

READ_TEXT

The `READ_TEXT` procedure returns the data from the HTTP response body in text form. The maximum number of characters returned is specified by the `len` parameter. The signature is:

```
READ_TEXT(<r> IN OUT UTL_HTTP.RESP, <data> OUT VARCHAR2, <len> INTEGER)
```

Parameters

`r`

`r` is the HTTP response record.

`data`

`data` is the response body in text form.

`len`

Set `len` to the maximum number of characters to be returned.

Examples

The following example retrieves the first 150 characters.

```
DECLARE
  v_req      UTL_HTTP.REQ;
  v_resp     UTL_HTTP.RESP;
  v_data     VARCHAR2(150);
BEGIN
  v_req := UTL_HTTP.BEGIN_REQUEST('http://www.enterprisedb.com');
  v_resp := UTL_HTTP.GET_RESPONSE(v_req);
  UTL_HTTP.READ_TEXT(v_resp, v_data, 150);
  DBMS_OUTPUT.PUT_LINE(v_data);
  UTL_HTTP.END_RESPONSE(v_resp);
END;
```

The following is the output.

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/
```

REQUEST

The `REQUEST` function returns the first 2000 bytes retrieved from a user-specified URL. The signature is:

```
REQUEST(<url> IN VARCHAR2) RETURN VARCHAR2
```

If the data found at the given URL is longer than 2000 bytes, the remainder will be discarded. If the data found at the given URL is shorter than 2000 bytes, the result will be shorter than 2000 bytes.

Parameters

`url`

`url` is the Uniform Resource Locator from which `UTL_HTTP` will return content.

Example

The following command returns the first 2000 bytes retrieved from the EnterpriseDB website:

```
SELECT UTL_HTTP.REQUEST('http://www.enterprisedb.com/') FROM DUAL;
```

REQUEST_PIECES

The `REQUEST_PIECES` function returns a table of 2000-byte segments retrieved from an URL. The signature is:

```
REQUEST_PIECES(<url> IN VARCHAR2, <max_pieces> NUMBER IN
DEFAULT 32767) RETURN UTL_HTTP.HTML_PIECES
```

Parameters

`url`

`url` is the Uniform Resource Locator from which `UTL_HTTP` will return content.

`max_pieces`

`max_pieces` specifies the maximum number of 2000-byte segments that the `REQUEST_PIECES` function will return. If `max_pieces` specifies more units than are available at the specified `url`, the final unit will contain fewer bytes.

Example

The following example returns the first four 2000 byte segments retrieved from the EnterpriseDB website:

```
DECLARE
    result UTL_HTTP.HTML_PIECES;
BEGIN
result := UTL_HTTP.REQUEST_PIECES('http://www.enterprisedb.com/', 4);
END
```

SET_BODY_CHARSET

The `SET_BODY_CHARSET` procedure sets the default character set of the body of future HTTP requests. The signature is:

```
SET_BODY_CHARSET(<charset> VARCHAR2 DEFAULT NULL)
```

Parameters

`charset`

`charset` is the character set of the body of future requests. The default is null in which case the database character set is assumed.

SET_FOLLOW_REDIRECT

The `SET_FOLLOW_REDIRECT` procedure sets the maximum number of times the HTTP redirect instruction is to be followed in the response to this request or future requests. This procedures has two signatures:

```
SET_FOLLOW_REDIRECT(<max_redirects> IN INTEGER DEFAULT 3)
```

and

```
SET_FOLLOW_REDIRECT(<r> IN OUT UTL_HTTP.REQ, <max_redirects>
IN INTEGER DEFAULT 3)``
```

Use the second form to change the maximum number of redirections for an individual request that a request inherits from the session default settings.

Parameters

`r`

`r` is the HTTP request record.

`max_redirects`

`max_redirects` is maximum number of redirections allowed. Set to 0 to disable redirections. The default is 3.

SET_HEADER

The `SET_HEADER` procedure sets the HTTP request header. The signature is:

```
SET_HEADER(<r> IN OUT UTL_HTTP.REQ, <name> IN VARCHAR2, <value>
IN VARCHAR2 DEFAULT NULL)
```

Parameters

`r`

`r` is the HTTP request record.

`name`

`name` is the name of the request header.

`value`

`value` is the value of the request header. The default is null.

SET_RESPONSE_ERROR_CHECK

The `SET_RESPONSE_ERROR_CHECK` procedure determines whether or not HTTP 4xx and 5xx status codes returned by the `GET_RESPONSE` function should be interpreted as errors. The signature is:

```
SET_RESPONSE_ERROR_CHECK(<enable> IN BOOLEAN DEFAULT FALSE)
```

Parameters

`enable`

Set `enable` to `TRUE` if HTTP 4xx and 5xx status codes are to be treated as errors, otherwise set to `FALSE`.

The default is FALSE.

SET_TRANSFER_TIMEOUT

The `SET_TRANSFER_TIMEOUT` procedure sets the default, transfer timeout setting for waiting for a response from an HTTP request. This procedure has two signatures:

```
SET_TRANSFER_TIMEOUT(<timeout> IN INTEGER DEFAULT 60)
```

and

```
SET_TRANSFER_TIMEOUT(<r> IN OUT UTL_HTTP.REQ, <timeout> IN  
INTEGER DEFAULT 60)
```

Use the second form to change the transfer timeout setting for an individual request that a request inherits from the session default settings.

Parameters

`r`

`r` is the HTTP request record.

`timeout`

`timeout` is the transfer timeout setting in seconds for HTTP requests. The default is 60 seconds.

WRITE_LINE

The `WRITE_LINE` procedure writes data to the HTTP request body in text form; the text is terminated with a CRLF character pair. The signature is:

```
WRITE_LINE(<r> IN OUT UTL_HTTP.REQ, <data> IN VARCHAR2)
```

Parameters

`r`

`r` is the HTTP request record.

`data`

`data` is the request body in `TEXT` form.

Example

The following example writes data (`Account balance $500.00`) in text form to the request body to be sent using the HTTP `POST` method. The data is sent to a hypothetical web application (`post.php`) that accepts and processes data.

```
DECLARE  
  v_req      UTL_HTTP.REQ;  
  v_resp     UTL_HTTP.RESP;  
BEGIN  
  v_req := UTL_HTTP.BEGIN_REQUEST('http://www.example.com/post.php',  
    'POST');
```

```

UTL_HTTP.SET_HEADER(v_req, 'Content-Length', '23');
UTL_HTTP.WRITE_LINE(v_req, 'Account balance $500.00');
v_resp := UTL_HTTP.GET_RESPONSE(v_req);
DBMS_OUTPUT.PUT_LINE('Status Code: ' || v_resp.status_code);
DBMS_OUTPUT.PUT_LINE('Reason Phrase: ' || v_resp.reason_phrase);
UTL_HTTP.END_RESPONSE(v_resp);
END;

```

Assuming the web application successfully processed the `POST` method, the following output would be displayed:

```

Status Code: 200
Reason Phrase: OK

```

WRITE_RAW

The `WRITE_RAW` procedure writes data to the HTTP request body in binary form. The signature is:

```
WRITE_RAW(<r> IN OUT UTL_HTTP.REQ, <data> IN RAW)
```

Parameters

`r`

| `r` is the HTTP request record.

`data`

| `data` is the request body in binary form.

Example

The following example writes data in binary form to the request body to be sent using the HTTP `POST` method to a hypothetical web application that accepts and processes such data.

```

DECLARE
    v_req      UTL_HTTP.REQ;
    v_resp     UTL_HTTP.RESP;
BEGIN
    v_req := UTL_HTTP.BEGIN_REQUEST('http://www.example.com/post.php',
        'POST');
    UTL_HTTP.SET_HEADER(v_req, 'Content-Length', '23');
    UTL_HTTP.WRITE_RAW(v_req, HEXTORAW
        ('54657374696e6720504f5354206d6574686f6420696e20485454502072657175657374'));
    v_resp := UTL_HTTP.GET_RESPONSE(v_req);
    DBMS_OUTPUT.PUT_LINE('Status Code: ' || v_resp.status_code);
    DBMS_OUTPUT.PUT_LINE('Reason Phrase: ' || v_resp.reason_phrase);
    UTL_HTTP.END_RESPONSE(v_resp);
END;

```

The text string shown in the `HEXTORAW` function is the hexadecimal translation of the text `Testing POST method in HTTP request`.

Assuming the web application successfully processed the `POST` method, the following output would be displayed:

Status Code: 200
Reason Phrase: OK

WRITE_TEXT

The `WRITE_TEXT` procedure writes data to the HTTP request body in text form. The signature is:

```
WRITE_TEXT(<r> IN OUT UTL_HTTP.REQ, <data> IN VARCHAR2)
```

Parameters

`r`

`r` is the HTTP request record.

`data`

`data` is the request body in text form.

Example

The following example writes data (`Account balance $500.00`) in text form to the request body to be sent using the HTTP `POST` method. The data is sent to a hypothetical web application (`post.php`) that accepts and processes data.

```
DECLARE
  v_req      UTL_HTTP.REQ;
  v_resp     UTL_HTTP.RESP;
BEGIN
  v_req := UTL_HTTP.BEGIN_REQUEST('http://www.example.com/post.php',
  'POST');
  UTL_HTTP.SET_HEADER(v_req, 'Content-Length', '23');
  UTL_HTTP.WRITE_TEXT(v_req, 'Account balance $500.00');
  v_resp := UTL_HTTP.GET_RESPONSE(v_req);
  DBMS_OUTPUT.PUT_LINE('Status Code: ' || v_resp.status_code);
  DBMS_OUTPUT.PUT_LINE('Reason Phrase: ' || v_resp.reason_phrase);
  UTL_HTTP.END_RESPONSE(v_resp);
END;
```

Assuming the web application successfully processed the `POST` method, the following output would be displayed:

Status Code: 200
Reason Phrase: OK

END_OF_BODY

The `END_OF_BODY` exception will be raised when it reaches the end of the HTTP response body.

Example

The following example handles the exception and writes `Exception caught` in text form to the request body to be sent using the HTTP `POST` method. The data is sent to a hypothetical web application (`post.php`) that accepts and processes data.

```

DECLARE
    req UTL_HTTP.REQ;
    resp UTL_HTTP.RESP;
    value VARCHAR2(32768);
BEGIN
    req := UTL_HTTP.BEGIN_REQUEST('https://www.google.com/');
    resp := UTL_HTTP.GET_RESPONSE(req);
LOOP
    UTL_HTTP.READ_LINE(resp, value, TRUE);
END LOOP;
    UTL_HTTP.END_RESPONSE(resp);
EXCEPTION
    WHEN UTL_HTTP.END_OF_BODY THEN
        DBMS_OUTPUT.PUT_LINE('Exception caught');
        UTL_HTTP.END_RESPONSE(resp);
END;

```

Assuming the web application successfully processed the **POST** method, the following output will be displayed:

Output is:

Exception caught

6.3.22 UTL_MAIL

The **UTL_MAIL** package provides the capability to manage e-mail. Advanced Server supports the following procedures:

| Function/Procedure | Return Type | Description |
|---|-------------|---|
| SEND(sender, recipients, cc, bcc, subject, message [, mime_type [, priority]]) | n/a | Packages and sends an e-mail to an SMTP server. |
| SEND_ATTACH_RAW(sender, recipients, cc, bcc, subject, message, mime_type, priority, attachment [, att_inline [, att_mime_type [, att_filename]]]) | n/a | Same as the SEND procedure, but with BYTEA or large object attachments. |
| SEND_ATTACH_VARCHAR2(sender, recipients, cc, bcc, subject, message, mime_type, priority, attachment [, att_inline [, att_mime_type [, att_filename]]]) | n/a | Same as the SEND procedure, but with VARCHAR2 attachments. |

Note

An administrator must grant execute privileges to each user or group before they can use this package.

SEND

The **SEND** procedure provides the capability to send an e-mail to an SMTP server.

```

SEND(<sender> VARCHAR2, <recipients> VARCHAR2, <cc> VARCHAR2,
<bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2
[, <mime_type> VARCHAR2 [, <priority> PLS_INTEGER ]])

```

Parameters

sender

E-mail address of the sender.

recipients

Comma-separated e-mail addresses of the recipients.

cc

Comma-separated e-mail addresses of copy recipients.

bcc

Comma-separated e-mail addresses of blind copy recipients.

subject

Subject line of the e-mail.

message

Body of the e-mail.

mime_type

Mime type of the message. The default is `text/plain; charset=us-ascii`.

priority

Priority of the e-mail. The default is 3.

Examples

The following anonymous block sends a simple e-mail message.

```
DECLARE
  v_sender      VARCHAR2(30);
  v_recipients  VARCHAR2(60);
  v_subj        VARCHAR2(20);
  v_msg         VARCHAR2(200);
BEGIN
  v_sender := 'jsmith@enterprisedb.com';
  v_recipients := 'ajones@enterprisedb.com,rrogers@enterprisedb.com';
  v_subj := 'Holiday Party';
  v_msg := 'This year''s party is scheduled for Friday, Dec. 21 at ' ||
           '6:00 PM. Please RSVP by Dec. 15th.';
  UTL_MAIL.SEND(v_sender,v_recipients,NULL,NULL,v_subj,v_msg);
END;
```

SEND_ATTACH_RAW

The `SEND_ATTACH_RAW` procedure provides the capability to send an e-mail to an SMTP server with an attachment containing either `BYTEA` data or a large object (identified by the large object's `OID`). The call to `SEND_ATTACH_RAW` can be written in two ways:

```
SEND_ATTACH_RAW(<sender> VARCHAR2, <recipients> VARCHAR2,
<cc> VARCHAR2, <bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2,
<mime_type> VARCHAR2, <priority> PLS_INTEGER,
<attachment> BYTEA[, <att_inline> BOOLEAN
[, <att_mime_type> VARCHAR2[, <att_filename> VARCHAR2 ]]])
```

or

```
SEND_ATTACH_RAW(<sender> VARCHAR2, <recipients> VARCHAR2,
<cc> VARCHAR2, <bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2,
<mime_type> VARCHAR2, <priority> PLS_INTEGER, <attachment> OID
[, <att_inline> BOOLEAN [, <att_mime_type> VARCHAR2
[, <att_filename> VARCHAR2 ]]])
```

Parameters

sender

E-mail address of the sender.

recipients

Comma-separated e-mail addresses of the recipients.

cc

Comma-separated e-mail addresses of copy recipients.

bcc

Comma-separated e-mail addresses of blind copy recipients.

subject

Subject line of the e-mail.

message

Body of the e-mail.

mime_type

Mime type of the message. The default is `text/plain; charset=us-ascii`.

priority

Priority of the e-mail. The default is `3`.

attachment

The attachment.

att_inline

If set to `TRUE`, then the attachment is viewable inline, `FALSE` otherwise. The default is `TRUE`.

att_mime_type

Mime type of the attachment. The default is `application/octet`.

att_filename

The file name containing the attachment. The default is `NULL`.

SEND_ATTACH_VARCHAR2

The `SEND_ATTACH_VARCHAR2` procedure provides the capability to send an e-mail to an SMTP server with a text attachment.

```
SEND_ATTACH_VARCHAR2(<sender> VARCHAR2, <recipients> VARCHAR2, <cc>
VARCHAR2, <bcc> VARCHAR2, <subject> VARCHAR2, <message> VARCHAR2,
<mime_type> VARCHAR2, <priority> PLS_INTEGER, <attachment> VARCHAR2 [,,
<att_inline> BOOLEAN [, <att_mime_type> VARCHAR2 [, <att_filename> VARCHAR2
]]])
```

Parameters

`sender`

E-mail address of the sender.

`recipients`

Comma-separated e-mail addresses of the recipients.

`cc`

Comma-separated e-mail addresses of copy recipients.

`bcc`

Comma-separated e-mail addresses of blind copy recipients.

`subject`

Subject line of the e-mail.

`message`

Body of the e-mail.

`mime_type`

Mime type of the message. The default is `text/plain; charset=us-ascii`.

`priority`

Priority of the e-mail. The default is `3`.

`attachment`

The `VARCHAR2` attachment.

`att_inline`

If set to `TRUE`, then the attachment is viewable inline, `FALSE` otherwise. The default is `TRUE`.

`att_mime_type`

Mime type of the attachment. The default is `text/plain; charset=us-ascii`.

att_filename

The file name containing the attachment. The default is `NULL`.

6.3.23 UTL_RAW

The `UTL_RAW` package allows you to manipulate or retrieve the length of raw data types.

Note

An administrator must grant execute privileges to each user or group before they can use this package.

| Function/Procedure | Function or Procedure | Return Type | Description |
|--|-----------------------|-------------|--|
| <code>CAST_TO_RAW(c IN VARCHAR2)</code> | Function | RAW | Converts a VARCHAR2 string to a RAW value. |
| <code>CAST_TO_VARCHAR2(r IN RAW)</code> | Function | VARCHAR2 | Converts a RAW value to a VARCHAR2 string. |
| <code>CONCAT(r1 IN RAW, r2 IN RAW, r3 IN RAW,...)</code> | Function | RAW | Concatenate multiple RAW values into a single RAW value. |
| <code>CONVERT(r IN RAW, to_charset IN VARCHAR2, from_charset IN VARCHAR2)</code> | Function | RAW | Converts encoded data from one encoding to another, and returns the result as a RAW value. |
| <code>LENGTH(r IN RAW)</code> | Function | NUMBER | Returns the length of a RAW value. |
| <code>SUBSTR(r IN RAW, pos IN INTEGER, len IN INTEGER)</code> | Function | RAW | Returns a portion of a RAW value. |

Advanced Server's implementation of `UTL_RAW` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

CAST_TO_RAW

The `CAST_TO_RAW` function converts a `VARCHAR2` string to a `RAW` value. The signature is:

```
CAST_TO_RAW(<c> VARCHAR2)
```

The function returns a `RAW` value if you pass a non-`NULL` value; if you pass a `NULL` value, the function will return `NULL`.

Parameters

c

The `VARCHAR2` value that will be converted to `RAW`.

Example

The following example uses the `CAST_TO_RAW` function to convert a `VARCHAR2` string to a `RAW` value:

```

DECLARE
v VARCHAR2;
r RAW;
BEGIN
v := 'Accounts';
dbms_output.put_line(v);
r := UTL_RAW.CAST_TO_RAW(v);
dbms_output.put_line(r);
END;

```

The result set includes the content of the original string and the converted `RAW` value:

```

Accounts
\x4163636f756e7473

```

CAST_TO_VARCHAR2

The `CAST_TO_VARCHAR2` function converts `RAW` data to `VARCHAR2` data. The signature is:

```
CAST_TO_VARCHAR2(<r> RAW)
```

The function returns a `VARCHAR2` value if you pass a non-`NULL` value; if you pass a `NULL` value, the function will return `NULL`.

Parameters

`r`

The `RAW` value that will be converted to a `VARCHAR2` value.

Example

The following example uses the `CAST_TO_VARCHAR2` function to convert a `RAW` value to a `VARCHAR2` string:

```

DECLARE
r RAW;
v VARCHAR2;
BEGIN
r := '\x4163636f756e7473';
dbms_output.put_line(v);
v := UTL_RAW.CAST_TO_VARCHAR2(r);
dbms_output.put_line(r);
END;

```

The result set includes the content of the original string and the converted `RAW` value:

```

\x4163636f756e7473
Accounts

```

CONCAT

The `CONCAT` function concatenates multiple `RAW` values into a single `RAW` value. The signature is:

`CONCAT(<r1> RAW, <r2> RAW, <r3> RAW,...)`

The function returns a `RAW` value. Unlike the Oracle implementation, the Advanced Server implementation is a variadic function, and does not place a restriction on the number of values that can be concatenated.

Parameters

`r1, r2, r3,...`

| The `RAW` values that `CONCAT` will concatenate.

Example

The following example uses the `CONCAT` function to concatenate multiple `RAW` values into a single `RAW` value:

```
SELECT UTL_RAW.CAST_TO_VARCHAR2(UTL_RAW.CONCAT('lx61', 'lx62', 'lx63')) FROM
DUAL;
concat
-----
abc
(1 row)
```

The result (the concatenated values) is then converted to `VARCHAR2` format by the `CAST_TO_VARCHAR2` function.

CONVERT

The `CONVERT` function converts a string from one encoding to another encoding and returns the result as a `RAW` value. The signature is:

`CONVERT(<r> RAW, <to_charset> VARCHAR2, <from_charset> VARCHAR2)`

The function returns a `RAW` value.

Parameters

`r`

| The `RAW` value that will be converted.

`to_charset`

| The name of the encoding to which `r` will be converted.

`from_charset`

| The name of the encoding from which `r` will be converted.

Example

The following example uses the `UTL_RAW.CAST_TO_RAW` function to convert a `VARCHAR2` string (`Accounts`) to a raw value, and then convert the value from `UTF8` to `LATIN7`, and then from `LATIN7` to `UTF8`:

```
DECLARE
r RAW;
v VARCHAR2;
BEGIN
```

```
v:= 'Accounts';
dbms_output.put_line(v);
r:= UTL_RAW.CAST_TO_RAW(v);
dbms_output.put_line(r);
r:= UTL_RAW.CONVERT(r, 'UTF8', 'LATIN7');
dbms_output.put_line(r);
r:= UTL_RAW.CONVERT(r, 'LATIN7', 'UTF8');
dbms_output.put_line(r);
```

The example returns the `VARCHAR2` value, the `RAW` value, and the converted values:

```
Accounts
\x4163636f756e7473
\x4163636f756e7473
\x4163636f756e7473
```

LENGTH

The `LENGTH` function returns the length of a `RAW` value. The signature is:

```
LENGTH(<r> RAW)
```

The function returns a `RAW` value.

Parameters

`r`

The `RAW` value that `LENGTH` will evaluate.

Example

The following example uses the `LENGTH` function to return the length of a `RAW` value:

```
SELECT UTL_RAW.LENGTH(UTL_RAW.CAST_TO_RAW('Accounts')) FROM DUAL;
length
-----
8
(1 row)
```

The following example uses the `LENGTH` function to return the length of a `RAW` value that includes multi-byte characters:

```
SELECT UTL_RAW.LENGTH(UTL_RAW.CAST_TO_RAW('hello'));
length
-----
12
(1 row)
```

SUBSTR

The `SUBSTR` function returns a substring of a `RAW` value. The signature is:

```
SUBSTR (<r> RAW, <pos> INTEGER, <len> INTEGER)
```

This function returns a `RAW` value.

Parameters

`r`

The `RAW` value from which the substring will be returned.

`pos`

The position within the `RAW` value of the first byte of the returned substring.

- If `pos` is `0` or `1`, the substring begins at the first byte of the `RAW` value.
- If `pos` is greater than one, the substring begins at the first byte specified by `pos`. For example, if `pos` is `3`, the substring begins at the third byte of the value.
- If `pos` is negative, the substring begins at `pos` bytes from the end of the source value. For example, if `pos` is `-3`, the substring begins at the third byte from the end of the value.

`len`

The maximum number of bytes that will be returned.

Example

The following example uses the `SUBSTR` function to select a substring that begins `3` bytes from the start of a `RAW` value:

```
SELECT UTL_RAW.SUBSTR(UTL_RAW.CAST_TO_RAW('Accounts'), 3, 5) FROM DUAL;
substr
-----
count
(1 row)
```

The following example uses the `SUBSTR` function to select a substring that starts `5` bytes from the end of a `RAW` value:

```
SELECT UTL_RAW.SUBSTR(UTL_RAW.CAST_TO_RAW('Accounts'), -5 , 3) FROM DUAL;
substr
-----
oun
(1 row)
```

6.3.24 UTL_SMTP

The `UTL_SMTP` package provides the capability to send e-mails over the Simple Mail Transfer Protocol (SMTP).

Note

An administrator must grant execute privileges to each user or group before they can use this package.

| Function/Procedure | Function or Procedure | Return Type | Description |
|-----------------------------------|-----------------------|-------------|-------------------------|
| <code>CLOSE_DATA(c IN OUT)</code> | Procedure | n/a | Ends an e-mail message. |

| | | | |
|--|-----------|------------|--|
| COMMAND(c IN OUT, cmd [, arg]) | Both | REPLY | Execute an SMTP command. |
| COMMAND_REPLIES(c IN OUT, cmd [, arg]) | Function | REPLIES | Execute an SMTP command where multiple reply lines are expected. |
| DATA(c IN OUT, body VARCHAR2) | Procedure | n/a | Specify the body of an e-mail message. |
| EHLO(c IN OUT, domain) | Procedure | n/a | Perform initial handshaking with an SMTP server and return extended information. |
| HELO(c IN OUT, domain) | Procedure | n/a | Perform initial handshaking with an SMTP server |
| HELP(c IN OUT [, command]) | Function | REPLIES | Send the HELP command. |
| MAIL(c IN OUT, sender [, parameters]) | Procedure | n/a | Start a mail transaction. |
| NOOP(c IN OUT) | Both | REPLY | Send the null command. |
| OPEN_CONNECTION(host [, port [, tx_timeout]]) | Function | CONNECTION | Open a connection. |
| OPEN_DATA(c IN OUT) | Both | REPLY | Send the DATA command. |
| QUIT(c IN OUT) | Procedure | n/a | Terminate the SMTP session and disconnect. |
| RCPT(c IN OUT, recipient [, parameters]) | Procedure | n/a | Specify the recipient of an e-mail message. |
| RSET(c IN OUT) | Procedure | n/a | Terminate the current mail transaction. |
| VRFY(c IN OUT, recipient) | Function | REPLY | Validate an e-mail address. |
| WRITE_DATA(c IN OUT, data) | Procedure | n/a | Write a portion of the e-mail message. |

Advanced Server's implementation of `UTL_SMTP` is a partial implementation when compared to Oracle's version. Only those functions and procedures listed in the table above are supported.

The following table lists the public variables available in the `UTL_SMTP` package.

| Public Variables | Data Type | Value | Description |
|------------------|-----------|-------|------------------------------------|
| connection | RECORD | | Description of an SMTP connection. |
| reply | RECORD | | SMTP reply line. |

CONNECTION

The `CONNECTION` record type provides a description of an SMTP connection.

```
TYPE connection IS RECORD (
    host      VARCHAR2(255),
    port      PLS_INTEGER,
    tx_timeout PLS_INTEGER
);
```

REPLY/REPLIES

The `REPLY` record type provides a description of an SMTP reply line. `REPLIES` is a table of multiple SMTP reply lines.

```
TYPE reply IS RECORD (
    code      INTEGER,
```

```
text      VARCHAR2(508)
);
TYPE replies IS TABLE OF reply INDEX BY BINARY_INTEGER;
```

CLOSE_DATA

The `CLOSE_DATA` procedure terminates an e-mail message by sending the following sequence:

```
| <CR><LF>.<CR><LF>
```

This is a single period at the beginning of a line.

```
CLOSE_DATA(<c> IN OUT CONNECTION)
```

Parameters

`c`

The SMTP connection to be closed.

COMMAND

The `COMMAND` procedure provides the capability to execute an SMTP command. If you are expecting multiple reply lines, use `COMMAND_REPLIES`.

```
<reply> REPLY COMMAND(<c> IN OUT CONNECTION, <cmd> VARCHAR2
[, <arg> VARCHAR2 ])
```

```
COMMAND(<c> IN OUT CONNECTION, <cmd> VARCHAR2 [, <arg> VARCHAR2 ])
```

Parameters

`c`

The SMTP connection to which the command is to be sent.

`cmd`

The SMTP command to be processed.

`arg`

An argument to the SMTP command. The default is null.

`reply`

SMTP reply to the command. If SMTP returns multiple replies, only the last one is returned in `reply`.

See [Reply/Replies](#) for a description of `REPLY` and `REPLIES`.

COMMAND_REPLIES

The `COMMAND_REPLIES` function processes an SMTP command that returns multiple reply lines. Use `COMMAND` if only a single reply line is expected.

```
<replies> REPLIES COMMAND(<c> IN OUT CONNECTION, <cmd> VARCHAR2
[, <arg> VARCHAR2 ])
```

Parameters

c

The SMTP connection to which the command is to be sent.

cmd

The SMTP command to be processed.

arg

An argument to the SMTP command. The default is null.

replies

SMTP reply lines to the command. See [Reply/Replies](#) for a description of REPLY and REPLIES.

DATA

The DATA procedure provides the capability to specify the body of the e-mail message. The message is terminated with a <CR><LF>.<CR><LF> sequence.

```
DATA(<c> IN OUT CONNECTION, <body> VARCHAR2)
```

Parameters

c

The SMTP connection to which the command is to be sent.

body

Body of the e-mail message to be sent.

EHLO

The EHLO procedure performs initial handshaking with the SMTP server after establishing the connection. The EHLO procedure allows the client to identify itself to the SMTP server according to RFC 821. RFC 1869 specifies the format of the information returned in the server's reply. The HELO procedure performs the equivalent functionality, but returns less information about the server.

```
EHLO(<c> IN OUT CONNECTION, <domain> VARCHAR2)
```

Parameters

c

The connection to the SMTP server over which to perform handshaking.

domain

Domain name of the sending host.

HELO

The `HELO` procedure performs initial handshaking with the SMTP server after establishing the connection. The `HELO` procedure allows the client to identify itself to the SMTP server according to RFC 821. The `EHLO` procedure performs the equivalent functionality, but returns more information about the server.

```
HELO(<c> IN OUT, <domain*> VARCHAR2)
```

Parameters

`c`

The connection to the SMTP server over which to perform handshaking.

`domain`

Domain name of the sending host.

HELP

The `HELP` function provides the capability to send the `HELP` command to the SMTP server.

```
<replies> REPLIES HELP(<c> IN OUT CONNECTION [, <command> VARCHAR2 ])
```

Parameters

`c`

The SMTP connection to which the command is to be sent.

`command`

Command on which help is requested.

`replies`

SMTP reply lines to the command. See [Reply/Replies](#) for a description of `REPLY` and `REPLIES`.

MAIL

The `MAIL` procedure initiates a mail transaction.

```
MAIL(<c> IN OUT CONNECTION, <sender> VARCHAR2  
[, <parameters> VARCHAR2 ])
```

Parameters

`c`

Connection to SMTP server on which to start a mail transaction.

`sender`

The sender's e-mail address.

`parameters`

Mail command parameters in the format, `key=value` as defined in RFC 1869.

NOOP

The `NOOP` function/procedure sends the null command to the SMTP server. The `NOOP` has no effect upon the server except to obtain a successful response.

```
<reply> REPLY NOOP(<c> IN OUT CONNECTION)
```

```
NOOP(<c> IN OUT CONNECTION)
```

Parameters

`c`

The SMTP connection on which to send the command.

`reply`

SMTP reply to the command. If SMTP returns multiple replies, only the last one is returned in `reply`. See [Reply/Replies](#) for a description of `REPLY` and `REPLIES`.

OPEN_CONNECTION

The `OPEN_CONNECTION` functions open a connection to an SMTP server.

```
<c> CONNECTION OPEN_CONNECTION(<host> VARCHAR2 [, <port>
PLS_INTEGER [, <tx_timeout> PLS_INTEGER DEFAULT NULL]])
```

Parameters

`host`

Name of the SMTP server.

`port`

Port number on which the SMTP server is listening. The default is 25.

`tx_timeout`

Time out value in seconds. Do not wait is indicated by specifying 0. Wait indefinitely is indicated by setting timeout to null. The default is null.

`c`

Connection handle returned by the SMTP server.

OPEN_DATA

The `OPEN_DATA` procedure sends the `DATA` command to the SMTP server.

```
OPEN_DATA(<c> IN OUT CONNECTION)
```

Parameters

c

| SMTP connection on which to send the command.

QUIT

The **QUIT** procedure closes the session with an SMTP server.

QUIT(<c> IN OUT CONNECTION)

Parameters

c

| SMTP connection to be terminated.

RCPT

The **RCPT** procedure provides the e-mail address of the recipient. To schedule multiple recipients, invoke **RCPT** multiple times.

**RCPT(<c> IN OUT CONNECTION, <recipient> VARCHAR2
[, <parameters> VARCHAR2])**

Parameters

c

| Connection to SMTP server on which to add a recipient.

recipient

| The recipient's e-mail address.

parameters

| Mail command parameters in the format, key=value as defined in RFC 1869.

RSET

The **RSET** procedure provides the capability to terminate the current mail transaction.

RSET(<c> IN OUT CONNECTION)

Parameters

c

| SMTP connection on which to cancel the mail transaction.

VRFY

The `VRFY` function provides the capability to validate and verify the recipient's e-mail address. If valid, the recipient's full name and fully qualified mailbox is returned.

```
<reply> REPLY VRFY(<c> IN OUT CONNECTION, <recipient> VARCHAR2)
```

Parameters

`c`

The SMTP connection on which to verify the e-mail address.

`recipient`

The recipient's e-mail address to be verified.

`reply`

SMTP reply to the command. If SMTP returns multiple replies, only the last one is returned in `reply`. See [Reply/Replies](#) for a description of `REPLY` and `REPLIES`.

WRITE_DATA

The `WRITE_DATA` procedure provides the capability to add `VARCHAR2` data to an e-mail message. The `WRITE_DATA` procedure may be repetitively called to add data.

```
WRITE_DATA(<c> IN OUT CONNECTION, <data> VARCHAR2)
```

Parameters

`c`

The SMTP connection on which to add data.

`data`

Data to be added to the e-mail message. The data must conform to the RFC 822 specification.

Comprehensive Example

The following procedure constructs and sends a text e-mail message using the `UTL_SMTP` package.

```
CREATE OR REPLACE PROCEDURE send_mail (
    p_sender      VARCHAR2,
    p_recipient   VARCHAR2,
    p_subj        VARCHAR2,
    p_msg         VARCHAR2,
    p_mailhost    VARCHAR2
)
IS
    v_conn        UTL_SMTP.CONNECTION;
    v_crlf        CONSTANT VARCHAR2(2) := CHR(13) || CHR(10);
    v_port        CONSTANT PLS_INTEGER := 25;
BEGIN
    v_conn := UTL_SMTP.OPEN_CONNECTION(p_mailhost,v_port);
    UTL_SMTP.HELO(v_conn,p_mailhost);
```

```

UTL_SMTP.MAIL(v_conn,p_sender);
UTL_SMTP.RCPT(v_conn,p_recipient);
UTL_SMTP.DATA(v_conn, SUBSTR(
  'Date: ' || TO_CHAR(SYSDATE,
  'Dy, DD Mon YYYY HH24:MI:SS') || v_crlf
  || 'From: ' || p_sender || v_crlf
  || 'To: ' || p_recipient || v_crlf
  || 'Subject: ' || p_subj || v_crlf
  || p_msg
  , 1, 32767));
UTL_SMTP.QUIT(v_conn);
END;

EXEC send_mail('asmith@enterprisedb.com','pjones@enterprisedb.com','Holiday
Party','Are you planning to attend?','smtp.enterprisedb.com');
```

The following example uses the `OPEN_DATA`, `WRITE_DATA`, and `CLOSE_DATA` procedures instead of the `DATA` procedure.

```

CREATE OR REPLACE PROCEDURE send_mail_2 (
  p_sender      VARCHAR2,
  p_recipient   VARCHAR2,
  p_subj        VARCHAR2,
  p_msg         VARCHAR2,
  p_mailhost    VARCHAR2
)
IS
  v_conn      UTL_SMTP.CONNECTION;
  v_crlf     CONSTANT VARCHAR2(2) := CHR(13) || CHR(10);
  v_port      CONSTANT PLS_INTEGER := 25;
BEGIN
  v_conn := UTL_SMTP.OPEN_CONNECTION(p_mailhost,v_port);
  UTL_SMTP.HELO(v_conn,p_mailhost);
  UTL_SMTP.MAIL(v_conn,p_sender);
  UTL_SMTP.RCPT(v_conn,p_recipient);
  UTL_SMTP.OPEN_DATA(v_conn);
  UTL_SMTP.WRITE_DATA(v_conn,'From: ' || p_sender || v_crlf);
  UTL_SMTP.WRITE_DATA(v_conn,'To: ' || p_recipient || v_crlf);
  UTL_SMTP.WRITE_DATA(v_conn,'Subject: ' || p_subj || v_crlf);
  UTL_SMTP.WRITE_DATA(v_conn,v_crlf || p_msg);
  UTL_SMTP.CLOSE_DATA(v_conn);
  UTL_SMTP.QUIT(v_conn);
END;
```

```

EXEC send_mail_2('asmith@enterprisedb.com','pjones@enterprisedb.com','Holiday
Party','Are you planning to attend?','smtp.enterprisedb.com');
```

6.3.25 UTL_URL

The `UTL_URL` package provides a way to escape illegal and reserved characters within an URL.

| Function/Procedure | Return Type | Description |
|---|-------------|---|
| ESCAPE(url, escape_reserved_chars, url_charset) | VARCHAR2 | Use the ESCAPE function to escape any illegal and reserved characters in a URL. |
| UNESCAPE(url, url_charset) | VARCHAR2 | The UNESCAPE function to convert an URL to its original form. |

The `UTL_URL` package will return the `BAD_URL` exception if the call to a function includes an incorrectly-formed URL.

ESCAPE

Use the `ESCAPE` function to escape illegal and reserved characters within an URL. The signature is:

```
ESCAPE(<url> VARCHAR2, <escape_reserved_chars> BOOLEAN,
<url_charset> VARCHAR2)
```

Reserved characters are replaced with a percent sign, followed by the two-digit hex code of the ascii value for the escaped character.

Parameters

`url`

`url` specifies the Uniform Resource Locator that `UTL_URL` will escape.

`escape_reserved_chars`

`escape_reserved_chars` is a `BOOLEAN` value that instructs the `ESCAPE` function to escape reserved characters as well as illegal characters:

- If `escaped_reserved_chars` is `FALSE`, `ESCAPE` will escape only the illegal characters in the specified URL.
- If `escape_reserved_chars` is `TRUE`, `ESCAPE` will escape both the illegal characters and the reserved characters in the specified URL.

By default, `escape_reserved_chars` is `FALSE`.

Within an URL, legal characters are:

| Uppercase A through Z | Lowercase a through z | 0 through 9 |
|-----------------------|-----------------------|----------------------|
| asterisk (*) | exclamation point (!) | hyphen (-) |
| left parenthesis () | period (.) | right parenthesis () |
| single-quote (') | tilde (~) | underscore (_) |

Some characters are legal in some parts of an URL, while illegal in others; to review comprehensive rules about illegal characters, refer to RFC 2396. Some examples of characters that are considered illegal in any part of an URL are:

| Illegal Character | Escape Sequence |
|-----------------------|-----------------|
| a blank space () | %20 |
| curly braces ({ or }) | %7b and %7d |
| hash mark (#) | %23 |

The `ESCAPE` function considers the following characters to be reserved, and will escape them if `escape_reserved_chars` is set to `TRUE`:

| Reserved Character | Escape Sequence |
|--------------------|-----------------|
| ampersand (&) | %5C |
| at sign (@) | %25 |
| colon (:) | %3a |
| comma (,) | %2c |
| dollar sign (\$) | %24 |
| equal sign (=) | %3d |
| plus sign (+) | %2b |
| question mark (?) | %3f |
| semi-colon (;) | %3b |
| slash (/) | %2f |

url_charset

`url_charset` specifies a character set to which a given character will be converted before it is escaped. If `url_charset` is `NULL`, the character will not be converted. The default value of `url_charset` is `ISO-8859-1`.

Examples

The following anonymous block uses the `ESCAPE` function to escape the blank spaces in the URL:

```
DECLARE
  result varchar2(400);
BEGIN
  result := UTL_URL.ESCAPE('http://www.example.com/Using the ESCAPE
function.html');
  DBMS_OUTPUT.PUT_LINE(result);
END;
```

The resulting (escaped) URL is:

```
http://www.example.com/Using%20the%20ESCAPE%20function.html
```

If you include a value of `TRUE` for the `escape_reserved_chars` parameter when invoking the function:

```
DECLARE
  result varchar2(400);
BEGIN
  result := UTL_URL.ESCAPE('http://www.example.com/Using the ESCAPE
function.html', TRUE);
  DBMS_OUTPUT.PUT_LINE(result);
END;
```

The `ESCAPE` function escapes the reserved characters as well as the illegal characters in the URL:

```
http%3A%2F%2Fwww.example.com%2FUsing%20the%20ESCAPE%20function.html
```

UNESCAPE

The `UNESCAPE` function removes escape characters added to an URL by the `ESCAPE` function, converting the

URL to its original form.

The signature is:

```
UNESCAPE(<url> VARCHAR2, <url_charset> VARCHAR2)
```

Parameters

`url`

`url` specifies the Uniform Resource Locator that `UTL_URL` will unescape.

`url_charset`

After unescaping a character, the character is assumed to be in `url_charset` encoding, and will be converted from that encoding to database encoding before being returned. If `url_charset` is `NULL`, the character will not be converted. The default value of `url_charset` is `ISO-8859-1`.

Examples

The following anonymous block uses the `ESCAPE` function to escape the blank spaces in the URL:

```
DECLARE
  result varchar2(400);
BEGIN
  result := UTL_URL.UNESCAPE('http://www.example.com
Using%20the%20UNESCAPE%20function.html');
  DBMS_OUTPUT.PUT_LINE(result);
END;
```

The resulting (unespaced) URL is:

```
http://www.example.com/Using the UNESCAPE function.html
```

7 Database Compatibility for Oracle Developers

7.1 Introduction

Database Compatibility for Oracle means that an application runs in an Oracle environment as well as in the EDB Postgres Advanced Server (Advanced Server) environment with minimal or no changes to the application code. Developing an application that is compatible with Oracle databases in the Advanced Server requires special attention to which features are used in the construction of the application. For example, developing a compatible application means choosing compatible:

- System and built-in functions for use in SQL statements and procedural logic.
- Stored Procedure Language (SPL) when creating database server-side application logic for stored procedures, functions, triggers, and packages.
- Data types that are compatible with Oracle databases
- SQL statements that are compatible with Oracle SQL

- System catalog views that are compatible with Oracle's data dictionary

For detailed information about the compatible SQL syntax, data types, and views, see the *Database Compatibility for Oracle Developers SQL Guide*.

The compatibility offered by the procedures and functions that are part of the Built-in packages is documented in the *Database Compatibility for Oracle Developers Built-in Packages Guide*.

For information about using the compatible tools and utilities (EDB*Plus, EDB*Loader, DRITA, and EDB Wrap) that are included with an Advanced Server installation, see the *Database Compatibility for Oracle Developers Tools and Utilities Guide*.

For applications written using the Oracle Call Interface (OCI), EnterpriseDB's Open Client Library (OCL) provides interoperability with these applications. For detailed information about using the Open Client Library, see the *EDB Postgres Advanced Server OCL Connector Guide*.

Advanced Server contains a rich set of features that enables development of database applications for either PostgreSQL or Oracle. For more information about all of the features of Advanced Server, see the user documentation available at the EnterpriseDB website.

Advanced Server documentation is available at:

<https://www.enterprisedb.com/edb-docs>

7.1.1 Configuration Parameters Compatible with Oracle Databases

EDB Postgres Advanced Server supports the development and execution of applications compatible with PostgreSQL and Oracle. Some system behaviors can be altered to act in a more PostgreSQL or in a more Oracle compliant manner; these behaviors are controlled by configuration parameters. Modifying the parameters in the `postgresql.conf` file changes the behavior for all databases in the cluster, while a user or group can `SET` the parameter value on the command line, effecting only their session. These parameters are:

- `edb_redwood_date` – Controls whether or not a time component is stored in `DATE` columns. For behavior compatible with Oracle databases, set `edb_redwood_date` to `TRUE`. See [edb_redwood_date](#).
 - `edb_redwood_raw_names` – Controls whether database object names appear in uppercase or lowercase letters when viewed from Oracle system catalogs. For behavior compatible with Oracle databases, `edb_redwood_raw_names` is set to its default value of `FALSE`. To view database object names as they are actually stored in the PostgreSQL system catalogs, set `edb_redwood_raw_names` to `TRUE`. See [edb_redwood_raw_names](#).
 - `edb_redwood_strings` – Equates `NULL` to an empty string for purposes of string concatenation operations. For behavior compatible with Oracle databases, set `edb_redwood_strings` to `TRUE`. See [edb_redwood_strings](#).
 - `edb_stmt_level_tx` – Isolates automatic rollback of an aborted SQL command to statement level rollback only – the entire, current transaction is not automatically rolled back as is the case for default PostgreSQL behavior. For behavior compatible with Oracle databases, set `edb_stmt_level_tx` to `TRUE`; however, use only when absolutely necessary. See [edb_stmt_level_tx](#).
 - `oracle_home` – Point Advanced Server to the correct Oracle installation directory. See [oracle_home](#).
-

7.1.1.1 `edb_redwood_date`

When `DATE` appears as the data type of a column in the commands, it is translated to `TIMESTAMP` at the time the table definition is stored in the data base if the configuration parameter `edb_redwood_date` is set to `TRUE`. Thus, a time component will also be stored in the column along with the date. This is consistent with Oracle's `DATE` data type.

If `edb_redwood_date` is set to `FALSE` the column's data type in a `CREATE TABLE` or `ALTER TABLE` command remains as a native PostgreSQL `DATE` data type and is stored as such in the database. The PostgreSQL `DATE` data type stores only the date without a time component in the column.

Regardless of the setting of `edb_redwood_date`, when `DATE` appears as a data type in any other context such as the data type of a variable in an SPL declaration section, or the data type of a formal parameter in an SPL procedure or SPL function, or the return type of an SPL function, it is always internally translated to a `TIMESTAMP` and thus, can handle a time component if present.

See the *Database Compatibility for Oracle Developers Reference Guide* for more information about date/time data types.

7.1.1.2 `edb_redwood_raw_names`

When `edb_redwood_raw_names` is set to its default value of `FALSE`, database object names such as table names, column names, trigger names, program names, user names, etc. appear in uppercase letters when viewed from Oracle catalogs (for a complete list of supported catalog views, see the *Database Compatibility for Oracle Catalog Views Guide*). In addition, quotation marks enclose names that were created with enclosing quotation marks.

When `edb_redwood_raw_names` is set to `TRUE`, the database object names are displayed exactly as they are stored in the PostgreSQL system catalogs when viewed from the Oracle catalogs. Thus, names created without enclosing quotation marks appear in lowercase as expected in PostgreSQL. Names created with enclosing quotation marks appear exactly as they were created, but without the quotation marks.

For example, the following user name is created, and then a session is started with that user.

```
CREATE USER reduser IDENTIFIED BY password;
edb=# \c - reduser
Password for user reduser:
You are now connected to database "edb" as user "reduser".
```

When connected to the database as `reduser`, the following tables are created.

```
CREATE TABLE all_lower (col INTEGER);
CREATE TABLE ALL_UPPER (COL INTEGER);
CREATE TABLE "Mixed_Case" ("Col" INTEGER);
```

When viewed from the Oracle catalog, `USER_TABLES`, with `edb_redwood_raw_names` set to the default value `FALSE`, the names appear in uppercase except for the `Mixed_Case` name, which appears as created and also with enclosing quotation marks.

```
edb=> SELECT * FROM USER_TABLES;
schema_name | table_name | tablespace_name | status | temporary
-----+-----+-----+-----+
REDUSER   | ALL_LOWER |           | VALID  | N
```

```
REDUSER | ALL_UPPER |          | VALID | N
REDUSER | "Mixed_Case" |       | VALID | N
(3 rows)
```

When viewed with `edb_redwood_raw_names` set to `TRUE`, the names appear in lowercase except for the `Mixed_Case` name, which appears as created, but now without the enclosing quotation marks.

```
edb=> SET edb_redwood_raw_names TO true;
SET
edb=> SELECT * FROM USER_TABLES;
schema_name | table_name | tablespace_name | status | temporary
-----+-----+-----+-----+
reduser | all_lower |          | VALID | N
reduser | all_upper |          | VALID | N
reduser | Mixed_Case |         | VALID | N
(3 rows)
```

These names now match the case when viewed from the PostgreSQL `pg_tables` catalog.

```
edb=> SELECT schemaname, tablename, tableowner FROM pg_tables WHERE
tableowner = 'reduser';
schemaname | tablename | tableowner
-----+-----+
reduser | all_lower | reduser
reduser | all_upper | reduser
reduser | Mixed_Case | reduser
(3 rows)
```

7.1.1.3 `edb_redwood_strings`

In Oracle, when a string is concatenated with a null variable or null column, the result is the original string; however, in PostgreSQL concatenation of a string with a null variable or null column gives a null result. If the `edb_redwood_strings` parameter is set to `TRUE`, the aforementioned concatenation operation results in the original string as done by Oracle. If `edb_redwood_strings` is set to `FALSE`, the native PostgreSQL behavior is maintained.

The following example illustrates the difference.

The sample application introduced in the next section contains a table of employees. This table has a column named `comm` that is null for most employees. The following query is run with `edb_redwood_string` set to `FALSE`. The concatenation of a null column with non-empty strings produces a final result of null, so only employees that have a commission appear in the query result. The output line for all other employees is null.

```
SET edb_redwood_strings TO off;

SELECT RPAD(ename,10) || ' ' || TO_CHAR(sal,'99,999.99') || ' ' ||
TO_CHAR(comm,'99,999.99') "EMPLOYEE COMPENSATION" FROM emp;

EMPLOYEE COMPENSATION
-----
ALLEN    1,600.00  300.00
```

| | | |
|--------|----------|----------|
| WARD | 1,250.00 | 500.00 |
| MARTIN | 1,250.00 | 1,400.00 |
| | | |
| TURNER | 1,500.00 | .00 |

(14 rows)

The following is the same query executed when `edb_redwood_strings` is set to `TRUE`. Here, the value of a null column is treated as an empty string. The concatenation of an empty string with a non-empty string produces the non-empty string. This result is consistent with the results produced by Oracle for the same query.

```
SET edb_redwood_strings TO on;

SELECT RPAD(ename,10) || ' ' || TO_CHAR(sal,'99,999.99') || ' ' ||
TO_CHAR(comm,'99,999.99') "EMPLOYEE COMPENSATION" FROM emp;

EMPLOYEE COMPENSATION
-----
SMITH      800.00
ALLEN     1,600.00   300.00
WARD      1,250.00   500.00
JONES     2,975.00
MARTIN    1,250.00  1,400.00
BLAKE      2,850.00
CLARK      2,450.00
SCOTT     3,000.00
KING      5,000.00
TURNER    1,500.00   .00
ADAMS     1,100.00
JAMES      950.00
FORD      3,000.00
MILLER    1,300.00
(14 rows)
```

7.1.1.4 `edb_stmt_level_tx`

In Oracle, when a runtime error occurs in a SQL command, all the updates on the database caused by that single command are rolled back. This is called *statement level transaction isolation*. For example, if a single `UPDATE` command successfully updates five rows, but an attempt to update a sixth row results in an exception, the updates to all six rows made by this `UPDATE` command are rolled back. The effects of prior SQL commands that have not yet been committed or rolled back are pending until a `COMMIT` or `ROLLBACK` command is executed.

In PostgreSQL, if an exception occurs while executing a SQL command, all the updates on the database since the start of the transaction are rolled back. In addition, the transaction is left in an aborted state and either a `COMMIT`

or `ROLLBACK` command must be issued before another transaction can be started.

If `edb_stmt_level_tx` is set to `TRUE`, then an exception will not automatically roll back prior uncommitted database updates, emulating the Oracle behavior. If `edb_stmt_level_tx` is set to `FALSE`, then an exception will roll back uncommitted database updates.

Note

Use `edb_stmt_level_tx` set to `TRUE` only when absolutely necessary, as this may cause a negative performance impact.

The following example run in PSQL shows that when `edb_stmt_level_tx` is `FALSE`, the abort of the second `INSERT` command also rolls back the first `INSERT` command. Note that in PSQL, the command `\set AUTOCOMMIT off` must be issued, otherwise every statement commits automatically defeating the purpose of this demonstration of the effect of `edb_stmt_level_tx`.

```
\set AUTOCOMMIT off
SET edb_stmt_level_tx TO off;

INSERT INTO emp (empno,ename,deptno) VALUES (9001, 'JONES', 40);
INSERT INTO emp (empno,ename,deptno) VALUES (9002, 'JONES', 00);
ERROR: insert or update on table "emp" violates foreign key constraint
"emp_ref_dept_fk"
DETAIL: Key (deptno)=(0) is not present in table "dept".

COMMIT;
SELECT empno, ename, deptno FROM emp WHERE empno > 9000;

empno | ename | deptno
-----+-----+
(0 rows)
```

In the following example, with `edb_stmt_level_tx` set to `TRUE`, the first `INSERT` command has not been rolled back after the error on the second `INSERT` command. At this point, the first `INSERT` command can either be committed or rolled back.

```
\set AUTOCOMMIT off
SET edb_stmt_level_tx TO on;

INSERT INTO emp (empno,ename,deptno) VALUES (9001, 'JONES', 40);
INSERT INTO emp (empno,ename,deptno) VALUES (9002, 'JONES', 00);
ERROR: insert or update on table "emp" violates foreign key constraint
"emp_ref_dept_fk"
DETAIL: Key (deptno)=(0) is not present in table "dept".

SELECT empno, ename, deptno FROM emp WHERE empno > 9000;

empno | ename | deptno
-----+-----+
9001 | JONES | 40
(1 row)

COMMIT;
```

A `ROLLBACK` command could have been issued instead of the `COMMIT` command in which case the insert of employee number `9001` would have been rolled back as well.

7.1.1.5 oracle_home

Before creating a link to an Oracle server, you must direct Advanced Server to the correct Oracle home directory. Set the `LD_LIBRARY_PATH` environment variable on Linux (or `PATH` on Windows) to the `lib` directory of the Oracle client installation directory.

For Windows only, you can instead set the value of the `oracle_home` configuration parameter in the `postgresql.conf` file. The value specified in the `oracle_home` configuration parameter will override the Windows `PATH` environment variable.

The `LD_LIBRARY_PATH` environment variable on Linux (`PATH` environment variable or `oracle_home` configuration parameter on Windows) must be set properly each time you start Advanced Server.

When using a Linux service script to start Advanced Server, be sure `LD_LIBRARY_PATH` has been set within the service script so it is in effect when the script invokes the `pg_ctl` utility to start Advanced Server.

For Windows only: To set the `oracle_home` configuration parameter in the `postgresql.conf` file, edit the file, adding the following line:

```
oracle_home = 'lib_directory'
```

Substitute the name of the Windows directory that contains `oci.dll` for `lib_directory`.

After setting the `oracle_home` configuration parameter, you must restart the server for the changes to take effect. Restart the server from the Windows Services console.

7.1.2 About the Examples Used in this Guide

The examples shown in this guide are illustrated using the PSQL program. The prompt that normally appears when using PSQL is omitted in these examples to provide extra clarity for the point being demonstrated.

Examples and output from examples are shown in fixed-width, grey font on a light background.

Also note the following points:

- During installation of the EDB Postgres Advanced Server the selection for configuration and defaults compatible with Oracle databases must be chosen in order to reproduce the same results as the examples shown in this guide. A default compatible configuration can be verified by issuing the following commands in PSQL and obtaining the same results as shown below.

```
SHOW edb_redwood_date;
```

```
edb_redwood_date
```

```
on
```

```
SHOW datestyle;
```

```
DateStyle
```

Redwood, DMY

```
SHOW edb_redwood_strings;
edb_redwood_strings
-----
on
```

- The examples use the sample tables, `dept`, `emp`, and `jobhist`, created and loaded when Advanced Server is installed. The `emp` table is installed with triggers that must be disabled in order to reproduce the same results as shown in this guide. Log onto Advanced Server as the `enterprisedb` superuser and disable the triggers by issuing the following command.

```
ALTER TABLE emp DISABLE TRIGGER USER;
```

The triggers on the `emp` table can later be re-activated with the following command.

```
ALTER TABLE emp ENABLE TRIGGER USER;
```

7.2 SQL Tutorial

Advanced Server is a *relational database management system* (RDBMS). That means it is a system for managing data stored in *relations*. A relation is essentially a mathematical term for a *table*. The notion of storing data in tables is so commonplace today that it might seem inherently obvious, but there are a number of other ways of organizing databases. Files and directories on Unix-like operating systems form an example of a hierarchical database. A more modern development is the object-oriented database.

Each table is a named collection of *rows*. Each row of a given table has the same set of named *columns*, and each column is of a specific *data type*. Whereas columns have a fixed order in each row, it is important to remember that SQL does not guarantee the order of the rows within the table in any way (although they can be explicitly sorted for display).

Tables are grouped into *databases*, and a collection of databases managed by a single Advanced Server instance constitutes a database *cluster*.

7.2.1 Sample Database

Throughout this documentation we will be working with a sample database to help explain some basic to advanced level database concepts.

7.2.1.1 Sample Database Installation

When Advanced Server is installed a sample database named, `edb`, is automatically created. This sample database contains the tables and programs used throughout this document by executing the script, `edb-sample.sql`, located in the `/usr/edb/as13/share` directory.

This script does the following:

- Creates the sample tables and programs in the currently connected database
- Grants all permissions on the tables to the `PUBLIC` group

The tables and programs will be created in the first schema of the search path in which the current user has permission to create tables and procedures. You can display the search path by issuing the command:

```
SHOW SEARCH_PATH;
```

Altering the search path can be done using commands in PSQL.

7.2.1.2 Sample Database Description

The sample database represents employees in an organization.


It contains three types of records: employees, departments, and historical records of employees.

Each employee has an identification number, name, hire date, salary, and manager. Some employees earn a commission in addition to their salary. All employee-related information is stored in the `emp` table.

The sample company is regionally diverse, so the database keeps track of the location of the departments. Each company employee is assigned to a department. Each department is identified by a unique department number and a short name. Each department is associated with one location. All department-related information is stored in the `dept` table.

The company also tracks information about jobs held by the employees. Some employees have been with the company for a long time and have held different positions, received raises, switched departments, etc. When a change in employee status occurs, the company records the end date of the former position. A new job record is added with the start date and the new job title, department, salary, and the reason for the status change. All employee history is maintained in the `jobhist` table.

The following is an entity relationship diagram of the sample database tables.



The following is the `edb-sample.sql` script.

```
--  
-- Script that creates the 'sample' tables, views, procedures,
```

```

-- functions, triggers, etc.
--
-- Start new transaction - commit all or nothing
--
BEGIN;
/
--

-- Create and load tables used in the documentation examples.
--
-- Create the 'dept' table
--

CREATE TABLE dept (
    deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc         VARCHAR2(13)
);
--

-- Create the 'emp' table
--

CREATE TABLE emp (
    empno      NUMBER(4) NOT NULL CONSTRAINT emp_pk PRIMARY KEY,
    ename       VARCHAR2(10),
    job         VARCHAR2(9),
    mgr         NUMBER(4),
    hiredate    DATE,
    sal         NUMBER(7,2) CONSTRAINT emp_sal_ck CHECK (sal > 0),
    comm        NUMBER(7,2),
    deptno     NUMBER(2) CONSTRAINT emp_ref_dept_fk
                REFERENCES dept(deptno)
);
--

-- Create the 'jobhist' table
--

CREATE TABLE jobhist (
    empno      NUMBER(4) NOT NULL,
    startdate   DATE NOT NULL,
    enddate     DATE,
    job         VARCHAR2(9),
    sal         NUMBER(7,2),
    comm        NUMBER(7,2),
    deptno     NUMBER(2),
    chgdesc    VARCHAR2(80),
    CONSTRAINT jobhist_pk PRIMARY KEY (empno, startdate),
    CONSTRAINT jobhist_ref_emp_fk FOREIGN KEY (empno)
        REFERENCES emp(empno) ON DELETE CASCADE,
    CONSTRAINT jobhist_ref_dept_fk FOREIGN KEY (deptno)
        REFERENCES dept (deptno) ON DELETE SET NULL,
    CONSTRAINT jobhist_date_chk CHECK (startdate <= enddate)
);
--

-- Create the 'salesemp' view
--

CREATE OR REPLACE VIEW salesemp AS
    SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job = 'SALESMAN';
--
```

```
-- Sequence to generate values for function 'new_empno'.
--
CREATE SEQUENCE next_empno START WITH 8000 INCREMENT BY 1;
--
-- Issue PUBLIC grants
--
GRANT ALL ON emp TO PUBLIC;
GRANT ALL ON dept TO PUBLIC;
GRANT ALL ON jobhist TO PUBLIC;
GRANT ALL ON salesemp TO PUBLIC;
GRANT ALL ON next_empno TO PUBLIC;
--
-- Load the 'dept' table
--
INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT INTO dept VALUES (20,'RESEARCH','DALLAS');
INSERT INTO dept VALUES (30,'SALES','CHICAGO');
INSERT INTO dept VALUES (40,'OPERATIONS','BOSTON');
--
-- Load the 'emp' table
--
INSERT INTO emp VALUES (7369,'SMITH','CLERK',7902,'17-DEC-80',800,NULL,20);
INSERT INTO emp VALUES (7499,'ALLEN','SALESMAN',7698,'20-FEB-81',1600,300,30);
INSERT INTO emp VALUES (7521,'WARD','SALESMAN',7698,'22-FEB-81',1250,500,30);
INSERT INTO emp VALUES (7566,'JONES','MANAGER',7839,'02-APR-81',2975,NULL,20);
INSERT INTO emp VALUES (7654,'MARTIN','SALESMAN',7698,'28-SEP-81',1250,1400,30);
INSERT INTO emp VALUES (7698,'BLAKE','MANAGER',7839,'01-MAY-81',2850,NULL,30);
INSERT INTO emp VALUES (7782,'CLARK','MANAGER',7839,'09-JUN-81',2450,NULL,10);
INSERT INTO emp VALUES (7788,'SCOTT','ANALYST',7566,'19-APR-87',3000,NULL,20);
INSERT INTO emp VALUES (7839,'KING','PRESIDENT',NULL,'17-NOV-81',5000,NULL,10);
INSERT INTO emp VALUES (7844,'TURNER','SALESMAN',7698,'08-SEP-81',1500,0,30);
INSERT INTO emp VALUES (7876,'ADAMS','CLERK',7788,'23-MAY-87',1100,NULL,20);
INSERT INTO emp VALUES (7900,'JAMES','CLERK',7698,'03-DEC-81',950,NULL,30);
INSERT INTO emp VALUES (7902,'FORD','ANALYST',7566,'03-DEC-81',3000,NULL,20);
INSERT INTO emp VALUES (7934,'MILLER','CLERK',7782,'23-JAN-82',1300,NULL,10);
--
-- Load the 'jobhist' table
--
INSERT INTO jobhist VALUES (7369,'17-DEC-80',NULL,'CLERK',800,NULL,20,'New Hire');
INSERT INTO jobhist VALUES (7499,'20-FEB-81',NULL,'SALESMAN',1600,300,30,'New Hire');
INSERT INTO jobhist VALUES (7521,'22-FEB-81',NULL,'SALESMAN',1250,500,30,'New Hire');
INSERT INTO jobhist VALUES (7566,'02-APR-81',NULL,'MANAGER',2975,NULL,20,'New Hire');
INSERT INTO jobhist VALUES (7654,'28-SEP-81',NULL,'SALESMAN',1250,1400,30,'New Hire');
```

```

INSERT INTO jobhist VALUES (7698,'01-MAY-81',NULL,'MANAGER',2850,NULL,30,'New Hire');
INSERT INTO jobhist VALUES (7782,'09-JUN-81',NULL,'MANAGER',2450,NULL,10,'New Hire');
INSERT INTO jobhist VALUES (7788,'19-APR-87','12-APR-88','CLERK',1000,NULL,20,'New Hire');
INSERT INTO jobhist VALUES (7788,'13-APR-88','04-MAY-89','CLERK',1040,NULL,20,'Raise');
INSERT INTO jobhist VALUES (7788,'05-MAY-90',NULL,'ANALYST',3000,NULL,20,'Promoted to Analyst');
INSERT INTO jobhist VALUES (7839,'17-NOV-81',NULL,'PRESIDENT',5000,NULL,10,'New Hire');
INSERT INTO jobhist VALUES (7844,'08-SEP-81',NULL,'SALESMAN',1500,0,30,'New Hire');
INSERT INTO jobhist VALUES (7876,'23-MAY-87',NULL,'CLERK',1100,NULL,20,'New Hire');
INSERT INTO jobhist VALUES (7900,'03-DEC-81','14-JAN-83','CLERK',950,NULL,10,'New Hire');
INSERT INTO jobhist VALUES (7900,'15-JAN-83',NULL,'CLERK',950,NULL,30,'Changed to Dept 30');
INSERT INTO jobhist VALUES (7902,'03-DEC-81',NULL,'ANALYST',3000,NULL,20,'New Hire');
INSERT INTO jobhist VALUES (7934,'23-JAN-82',NULL,'CLERK',1300,NULL,10,'New Hire');

-- 
-- Populate statistics table and view (pg_statistic/pg_stats)
-- 

ANALYZE dept;
ANALYZE emp;
ANALYZE jobhist;

-- 
-- Procedure that lists all employees' numbers and names
-- from the 'emp' table using a cursor.
-- 

CREATE OR REPLACE PROCEDURE list_emp
IS
    v_empno      NUMBER(4);
    v_ename      VARCHAR2(10);
    CURSOR emp_cur IS
        SELECT empno, ename FROM emp ORDER BY empno;
BEGIN
    OPEN emp_cur;
    DBMS_OUTPUT.PUT_LINE('EMPNO  ENAME');
    DBMS_OUTPUT.PUT_LINE('-----  -----');
    LOOP
        FETCH emp_cur INTO v_empno, v_ename;
        EXIT WHEN emp_cur%NOTFOUND;
        DBMS_OUTPUT.PUT_LINE(v_empno || '  ' || v_ename);
    END LOOP;
    CLOSE emp_cur;
END;
/
-- 
-- Procedure that selects an employee row given the employee
-- number and displays certain columns.

```

```

--  

CREATE OR REPLACE PROCEDURE select_emp (
    p_empno      IN NUMBER
)
IS
    v_ename      emp.ename%TYPE;
    v_hiredate   emp.hiredate%TYPE;
    v_sal        emp.sal%TYPE;
    v_comm       emp.comm%TYPE;
    v_dname      dept.dname%TYPE;
    v_disp_date  VARCHAR2(10);
BEGIN
    SELECT ename, hiredate, sal, NVL(comm, 0), dname
        INTO v_ename, v_hiredate, v_sal, v_comm, v_dname
        FROM emp e, dept d
        WHERE empno = p_empno
        AND e.deptno = d.deptno;
    v_disp_date := TO_CHAR(v_hiredate, 'MM/DD/YYYY');
    DBMS_OUTPUT.PUT_LINE('Number : ' || p_empno);
    DBMS_OUTPUT.PUT_LINE('Name : ' || v_ename);
    DBMS_OUTPUT.PUT_LINE('Hire Date : ' || v_disp_date);
    DBMS_OUTPUT.PUT_LINE('Salary : ' || v_sal);
    DBMS_OUTPUT.PUT_LINE('Commission: ' || v_comm);
    DBMS_OUTPUT.PUT_LINE('Department: ' || v_dname);
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('Employee ' || p_empno || ' not found');
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
        DBMS_OUTPUT.PUT_LINE(SQLERRM);
        DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
        DBMS_OUTPUT.PUT_LINE(SQLCODE);
END;
/  

--  

-- Procedure that queries the 'emp' table based on  

-- department number and employee number or name. Returns  

-- employee number and name as IN OUT parameters and job,  

-- hire date, and salary as OUT parameters.  

--  

CREATE OR REPLACE PROCEDURE emp_query (
    p_deptno     IN NUMBER,
    p_empno      IN OUT NUMBER,
    p_ename      IN OUT VARCHAR2,
    p_job        OUT VARCHAR2,
    p_hiredate   OUT DATE,
    p_sal        OUT NUMBER
)
IS
BEGIN
    SELECT empno, ename, job, hiredate, sal
        INTO p_empno, p_ename, p_job, p_hiredate, p_sal
        FROM emp
        WHERE deptno = p_deptno
        AND (empno = p_empno)

```

```

    OR ename = UPPER(p_ename));
END;
/
-- 
-- Procedure to call 'emp_query_caller' with IN and IN OUT
-- parameters. Displays the results received from IN OUT and
-- OUT parameters.
-- 
CREATE OR REPLACE PROCEDURE emp_query_caller
IS
    v_deptno      NUMBER(2);
    v_empno       NUMBER(4);
    v_ename       VARCHAR2(10);
    v_job         VARCHAR2(9);
    v_hiredate    DATE;
    v_sal         NUMBER;
BEGIN
    v_deptno := 30;
    v_empno := 0;
    v_ename := 'Martin';
    emp_query(v_deptno, v_empno, v_ename, v_job, v_hiredate, v_sal);
    DBMS_OUTPUT.PUT_LINE('Department : ' || v_deptno);
    DBMS_OUTPUT.PUT_LINE('Employee No: ' || v_empno);
    DBMS_OUTPUT.PUT_LINE('Name     : ' || v_ename);
    DBMS_OUTPUT.PUT_LINE('Job      : ' || v_job);
    DBMS_OUTPUT.PUT_LINE('Hire Date : ' || v_hiredate);
    DBMS_OUTPUT.PUT_LINE('Salary   : ' || v_sal);
EXCEPTION
    WHEN TOO_MANY_ROWS THEN
        DBMS_OUTPUT.PUT_LINE('More than one employee was selected');
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('No employees were selected');
END;
/
-- 
-- Function to compute yearly compensation based on semimonthly
-- salary.
-- 
CREATE OR REPLACE FUNCTION emp_comp (
    p_sal      NUMBER,
    p_comm     NUMBER
) RETURN NUMBER
IS
BEGIN
    RETURN (p_sal + NVL(p_comm, 0)) * 24;
END;
/
-- 
-- Function that gets the next number from sequence, 'next_empno',
-- and ensures it is not already in use as an employee number.
-- 
CREATE OR REPLACE FUNCTION new_empno RETURN NUMBER
IS
    v_cnt      INTEGER := 1;
    v_new_empno NUMBER;

```

```

BEGIN
  WHILE v_cnt > 0 LOOP
    SELECT next_empno.nextval INTO v_new_empno FROM dual;
    SELECT COUNT(*) INTO v_cnt FROM emp WHERE empno = v_new_empno;
  END LOOP;
  RETURN v_new_empno;
END;
/
-- 
-- EDB-SPL function that adds a new clerk to table 'emp'. This function
-- uses package 'emp_admin'.
-- 

CREATE OR REPLACE FUNCTION hire_clerk (
  p_ename      VARCHAR2,
  p_deptno     NUMBER
) RETURN NUMBER
IS
  v_empno      NUMBER(4);
  v_ename      VARCHAR2(10);
  v_job        VARCHAR2(9);
  v_mgr        NUMBER(4);
  v_hiredate   DATE;
  v_sal        NUMBER(7,2);
  v_comm        NUMBER(7,2);
  v_deptno     NUMBER(2);
BEGIN
  v_empno := new_empno;
  INSERT INTO emp VALUES (v_empno, p_ename, 'CLERK', 7782,
    TRUNC(SYSDATE), 950.00, NULL, p_deptno);
  SELECT empno, ename, job, mgr, hiredate, sal, comm, deptno INTO
    v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno
    FROM emp WHERE empno = v_empno;
  DBMS_OUTPUT.PUT_LINE('Department : ' || v_deptno);
  DBMS_OUTPUT.PUT_LINE('Employee No: ' || v_empno);
  DBMS_OUTPUT.PUT_LINE('Name      : ' || v_ename);
  DBMS_OUTPUT.PUT_LINE('Job       : ' || v_job);
  DBMS_OUTPUT.PUT_LINE('Manager   : ' || v_mgr);
  DBMS_OUTPUT.PUT_LINE('Hire Date : ' || v_hiredate);
  DBMS_OUTPUT.PUT_LINE('Salary    : ' || v_sal);
  DBMS_OUTPUT.PUT_LINE('Commission : ' || v_comm);
  RETURN v_empno;
EXCEPTION
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
    DBMS_OUTPUT.PUT_LINE(SQLERRM);
    DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
    DBMS_OUTPUT.PUT_LINE(SQLCODE);
    RETURN -1;
END;
/
-- 
-- PostgreSQL PL/pgSQL function that adds a new salesman
-- to table 'emp'.
-- 

CREATE OR REPLACE FUNCTION hire_salesman (

```

```

p_ename      VARCHAR,
p_sal        NUMERIC,
p_comm       NUMERIC
) RETURNS NUMERIC
AS $$

DECLARE
    v_empno     NUMERIC(4);
    v_ename     VARCHAR(10);
    v_job       VARCHAR(9);
    v_mgr       NUMERIC(4);
    v_hiredate  DATE;
    v_sal       NUMERIC(7,2);
    v_comm      NUMERIC(7,2);
    v_deptno    NUMERIC(2);

BEGIN
    v_empno := new_empno();
    INSERT INTO emp VALUES (v_empno, p_ename, 'SALESMAN', 7698,
                           CURRENT_DATE, p_sal, p_comm, 30);
    SELECT INTO
        v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno
        empno, ename, job, mgr, hiredate, sal, comm, deptno
    FROM emp WHERE empno = v_empno;
    RAISE INFO 'Department : %', v_deptno;
    RAISE INFO 'Employee No: %', v_empno;
    RAISE INFO 'Name    : %', v_ename;
    RAISE INFO 'Job     : %', v_job;
    RAISE INFO 'Manager : %', v_mgr;
    RAISE INFO 'Hire Date : %', v_hiredate;
    RAISE INFO 'Salary   : %', v_sal;
    RAISE INFO 'Commission : %', v_comm;
    RETURN v_empno;
EXCEPTION
    WHEN OTHERS THEN
        RAISE INFO 'The following is SQLERRM:';
        RAISE INFO '%', SQLERRM;
        RAISE INFO 'The following is SQLSTATE:';
        RAISE INFO '%', SQLSTATE;
        RETURN -1;
END;
$$ LANGUAGE 'plpgsql';
/
-- 
-- Rule to INSERT into view 'salesemp'
-- 

CREATE OR REPLACE RULE salesemp_i AS ON INSERT TO salesemp
DO INSTEAD
    INSERT INTO emp VALUES (NEW.empno, NEW.ename, 'SALESMAN', 7698,
                           NEW.hiredate, NEW.sal, NEW.comm, 30);
-- 
-- Rule to UPDATE view 'salesemp'
-- 

CREATE OR REPLACE RULE salesemp_u AS ON UPDATE TO salesemp
DO INSTEAD
    UPDATE emp SET empno = NEW.empno,
                  ename = NEW.ename,

```

```

hiredate = NEW.hiredate,
sal    = NEW.sal,
comm   = NEW.comm
WHERE empno = OLD.empno;
-- 
-- Rule to DELETE from view 'salesemp'
-- 
CREATE OR REPLACE RULE salesemp_d AS ON DELETE TO salesemp
DO INSTEAD
DELETE FROM emp WHERE empno = OLD.empno;
-- 
-- After statement-level trigger that displays a message after
-- an insert, update, or deletion to the 'emp' table. One message
-- per SQL command is displayed.
-- 
CREATE OR REPLACE TRIGGER user_audit_trig
AFTER INSERT OR UPDATE OR DELETE ON emp
DECLARE
v_action      VARCHAR2(24);
BEGIN
IF INSERTING THEN
  v_action := ' added employee(s) on ';
ELSIF UPDATING THEN
  v_action := ' updated employee(s) on ';
ELSIF DELETING THEN
  v_action := ' deleted employee(s) on ';
END IF;
DBMS_OUTPUT.PUT_LINE('User ' || USER || v_action ||
TO_CHAR(SYSDATE,'YYYY-MM-DD'));
END;
/
-- 
-- Before row-level trigger that displays employee number and
-- salary of an employee that is about to be added, updated,
-- or deleted in the 'emp' table.
-- 
CREATE OR REPLACE TRIGGER emp_sal_trig
BEFORE DELETE OR INSERT OR UPDATE ON emp
FOR EACH ROW
DECLARE
sal_diff      NUMBER;
BEGIN
IF INSERTING THEN
  DBMS_OUTPUT.PUT_LINE('Inserting employee ' || :NEW.empno);
  DBMS_OUTPUT.PUT_LINE('..New salary: ' || :NEW.sal);
END IF;
IF UPDATING THEN
  sal_diff := :NEW.sal - :OLD.sal;
  DBMS_OUTPUT.PUT_LINE('Updating employee ' || :OLD.empno);
  DBMS_OUTPUT.PUT_LINE('..Old salary: ' || :OLD.sal);
  DBMS_OUTPUT.PUT_LINE('..New salary: ' || :NEW.sal);
  DBMS_OUTPUT.PUT_LINE('..Raise   : ' || sal_diff);
END IF;
IF DELETING THEN
  DBMS_OUTPUT.PUT_LINE('Deleting employee ' || :OLD.empno);

```

```

DBMS_OUTPUT.PUT_LINE(..Old salary: ' || :OLD.sal);
END IF;
END;
/
-- 
-- Package specification for the 'emp_admin' package.
-- 
CREATE OR REPLACE PACKAGE emp_admin
IS
    FUNCTION get_dept_name (
        p_deptno      NUMBER
    ) RETURN VARCHAR2;
    FUNCTION update_emp_sal (
        p_empno      NUMBER,
        p_raise      NUMBER
    ) RETURN NUMBER;
    PROCEDURE hire_emp (
        p_empno      NUMBER,
        p_ename      VARCHAR2,
        p_job       VARCHAR2,
        p_sal       NUMBER,
        p_hiredate   DATE,
        p_comm      NUMBER,
        p_mgr       NUMBER,
        p_deptno     NUMBER
    );
    PROCEDURE fire_emp (
        p_empno      NUMBER
    );
END emp_admin;
/
-- 
-- Package body for the 'emp_admin' package.
-- 
CREATE OR REPLACE PACKAGE BODY emp_admin
IS
    --
    -- Function that queries the 'dept' table based on the department
    -- number and returns the corresponding department name.
    --
    FUNCTION get_dept_name (
        p_deptno      IN NUMBER
    ) RETURN VARCHAR2
    IS
        v_dname      VARCHAR2(14);
    BEGIN
        SELECT dname INTO v_dname FROM dept WHERE deptno = p_deptno;
        RETURN v_dname;
    EXCEPTION
        WHEN NO_DATA_FOUND THEN
            DBMS_OUTPUT.PUT_LINE('Invalid department number ' || p_deptno);
            RETURN '';
    END;
    --
    -- Function that updates an employee's salary based on the

```

```

-- employee number and salary increment/decrement passed
-- as IN parameters. Upon successful completion the function
-- returns the new updated salary.
--
FUNCTION update_emp_sal (
    p_empno      IN NUMBER,
    p_raise      IN NUMBER
) RETURN NUMBER
IS
    v_sal      NUMBER := 0;
BEGIN
    SELECT sal INTO v_sal FROM emp WHERE empno = p_empno;
    v_sal := v_sal + p_raise;
    UPDATE emp SET sal = v_sal WHERE empno = p_empno;
    RETURN v_sal;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('Employee ' || p_empno || ' not found');
        RETURN -1;
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('The following is SQLERRM:');
        DBMS_OUTPUT.PUT_LINE(SQLERRM);
        DBMS_OUTPUT.PUT_LINE('The following is SQLCODE:');
        DBMS_OUTPUT.PUT_LINE(SQLCODE);
        RETURN -1;
END;

-- Procedure that inserts a new employee record into the 'emp' table.
--
PROCEDURE hire_emp (
    p_empno      NUMBER,
    p_ename      VARCHAR2,
    p_job       VARCHAR2,
    p_sal       NUMBER,
    p_hiredate   DATE,
    p_comm      NUMBER,
    p_mgr       NUMBER,
    p_deptno    NUMBER
)
AS
BEGIN
    INSERT INTO emp(empno, ename, job, sal, hiredate, comm, mgr, deptno)
        VALUES(p_empno, p_ename, p_job, p_sal,
               p_hiredate, p_comm, p_mgr, p_deptno);
END;

-- Procedure that deletes an employee record from the 'emp' table based
-- on the employee number.
--
PROCEDURE fire_emp (
    p_empno      NUMBER
)
AS
BEGIN
    DELETE FROM emp WHERE empno = p_empno;

```

```
END;
END;
/
COMMIT;
```

7.2.2 Creating a New Table

A new table is created by specifying the table name, along with all column names and their types. The following is a simplified version of the `emp` sample table with just the minimal information needed to define a table.

```
CREATE TABLE emp (
    empno      NUMBER(4),
    ename      VARCHAR2(10),
    job        VARCHAR2(9),
    mgr        NUMBER(4),
    hiredate   DATE,
    sal        NUMBER(7,2),
    comm       NUMBER(7,2),
    deptno    NUMBER(2)
);
```

You can enter this into PSQL with line breaks. PSQL will recognize that the command is not terminated until the semicolon.

White space (i.e., spaces, tabs, and newlines) may be used freely in SQL commands. That means you can type the command aligned differently than the above, or even all on one line. Two dashes ("--) introduce comments. Whatever follows them is ignored up to the end of the line. SQL is case insensitive about key words and identifiers, except when identifiers are double-quoted to preserve the case (not done above).

`VARCHAR(10)` specifies a data type that can store arbitrary character strings up to 10 characters in length. `NUMBER(7,2)` is a fixed point number with precision 7 and scale 2. `NUMBER(4)` is an integer number with precision 4 and scale 0.

Advanced Server supports the usual SQL data types `INTEGER`, `SMALLINT`, `NUMBER`, `REAL`, `DOUBLE PRECISION`, `CHAR`, `VARCHAR2`, `DATE`, and `TIMESTAMP` as well as various synonyms for these types.

If you don't need a table any longer or want to recreate it differently you can remove it using the following command:

```
DROP TABLE tablename;
```

7.2.3 Populating a Table With Rows

The `INSERT` statement is used to populate a table with rows:

```
INSERT INTO emp VALUES (7369,'SMITH','CLERK',7902,'17-DEC-80',800,NULL,20);
```

Note that all data types use rather obvious input formats. Constants that are not simple numeric values usually

must be surrounded by single quotes ('), as in the example. The **DATE** type is actually quite flexible in what it accepts, but for this tutorial we will stick to the unambiguous format shown here.

The syntax used so far requires you to remember the order of the columns. An alternative syntax allows you to list the columns explicitly:

```
INSERT INTO emp(empno,ename,job,mgr,hiredate,sal,comm,deptno)
VALUES (7499,'ALLEN','SALESMAN',7698,'20-FEB-81',1600,300,30);
```

You can list the columns in a different order if you wish or even omit some columns, e.g., if the commission is unknown:

```
INSERT INTO emp(empno,ename,job,mgr,hiredate,sal,deptno)
VALUES (7369,'SMITH','CLERK',7902,'17-DEC-80',800,20);
```

Many developers consider explicitly listing the columns better style than relying on the order implicitly.

7.2.4 Querying a Table

To retrieve data from a table, the table is *queried*. An SQL **SELECT** statement is used to do this. The statement is divided into a select list (the part that lists the columns to be returned), a table list (the part that lists the tables from which to retrieve the data), and an optional qualification (the part that specifies any restrictions). The following query lists all columns of all employees in the table in no particular order.

```
SELECT * FROM emp;
```

Here, “*” in the select list means all columns. The following is the output from this query.

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|--------|-----------|------|--------------------|---------|---------|--------|
| 7369 | SMITH | CLERK | 7902 | 17-DEC-80 00:00:00 | 800.00 | | 20 |
| 7499 | ALLEN | SALESMAN | 7698 | 20-FEB-81 00:00:00 | 1600.00 | 300.00 | 30 |
| 7521 | WARD | SALESMAN | 7698 | 22-FEB-81 00:00:00 | 1250.00 | 500.00 | 30 |
| 7566 | JONES | MANAGER | 7839 | 02-APR-81 00:00:00 | 2975.00 | | 20 |
| 7654 | MARTIN | SALESMAN | 7698 | 28-SEP-81 00:00:00 | 1250.00 | 1400.00 | 30 |
| 7698 | BLAKE | MANAGER | 7839 | 01-MAY-81 00:00:00 | 2850.00 | | 30 |
| 7782 | CLARK | MANAGER | 7839 | 09-JUN-81 00:00:00 | 2450.00 | | 10 |
| 7788 | SCOTT | ANALYST | 7566 | 19-APR-87 00:00:00 | 3000.00 | | 20 |
| 7839 | KING | PRESIDENT | | 17-NOV-81 00:00:00 | 5000.00 | | 10 |
| 7844 | TURNER | SALESMAN | 7698 | 08-SEP-81 00:00:00 | 1500.00 | 0.00 | 30 |
| 7876 | ADAMS | CLERK | 7788 | 23-MAY-87 00:00:00 | 1100.00 | | 20 |
| 7900 | JAMES | CLERK | 7698 | 03-DEC-81 00:00:00 | 950.00 | | 30 |
| 7902 | FORD | ANALYST | 7566 | 03-DEC-81 00:00:00 | 3000.00 | | 20 |
| 7934 | MILLER | CLERK | 7782 | 23-JAN-82 00:00:00 | 1300.00 | | 10 |

(14 rows)

You may specify any arbitrary expression in the select list. For example, you can do:

```
SELECT ename, sal, sal * 24 AS yearly_salary, deptno FROM emp;
```

| ename | sal | yearly_salary | deptno |
|-------|-----|---------------|--------|
|-------|-----|---------------|--------|

| | | | |
|--------|---------|-----------|----|
| SMITH | 800.00 | 19200.00 | 20 |
| ALLEN | 1600.00 | 38400.00 | 30 |
| WARD | 1250.00 | 30000.00 | 30 |
| JONES | 2975.00 | 71400.00 | 20 |
| MARTIN | 1250.00 | 30000.00 | 30 |
| BLAKE | 2850.00 | 68400.00 | 30 |
| CLARK | 2450.00 | 58800.00 | 10 |
| SCOTT | 3000.00 | 72000.00 | 20 |
| KING | 5000.00 | 120000.00 | 10 |
| TURNER | 1500.00 | 36000.00 | 30 |
| ADAMS | 1100.00 | 26400.00 | 20 |
| JAMES | 950.00 | 22800.00 | 30 |
| FORD | 3000.00 | 72000.00 | 20 |
| MILLER | 1300.00 | 31200.00 | 10 |

(14 rows)

Notice how the `AS` clause is used to re-label the output column. (The `AS` clause is optional.)

A query can be qualified by adding a `WHERE` clause that specifies which rows are wanted. The `WHERE` clause contains a Boolean (truth value) expression, and only rows for which the Boolean expression is true are returned. The usual Boolean operators (`AND`, `OR`, and `NOT`) are allowed in the qualification. For example, the following retrieves the employees in department 20 with salaries over \$1000.00:

```
SELECT ename, sal, deptno FROM emp WHERE deptno = 20 AND sal > 1000;
```

| ename | sal | deptno |
|-------|---------|--------|
| JONES | 2975.00 | 20 |
| SCOTT | 3000.00 | 20 |
| ADAMS | 1100.00 | 20 |
| FORD | 3000.00 | 20 |

(4 rows)

You can request that the results of a query be returned in sorted order:

```
SELECT ename, sal, deptno FROM emp ORDER BY ename;
```

| ename | sal | deptno |
|--------|---------|--------|
| ADAMS | 1100.00 | 20 |
| ALLEN | 1600.00 | 30 |
| BLAKE | 2850.00 | 30 |
| CLARK | 2450.00 | 10 |
| FORD | 3000.00 | 20 |
| JAMES | 950.00 | 30 |
| JONES | 2975.00 | 20 |
| KING | 5000.00 | 10 |
| MARTIN | 1250.00 | 30 |
| MILLER | 1300.00 | 10 |
| SCOTT | 3000.00 | 20 |
| SMITH | 800.00 | 20 |
| TURNER | 1500.00 | 30 |
| WARD | 1250.00 | 30 |

(14 rows)

You can request that duplicate rows be removed from the result of a query:

```
SELECT DISTINCT job FROM emp;
```

```
job
```

```
-----
ANALYST
CLERK
MANAGER
PRESIDENT
SALESMAN
(5 rows)
```

The following section shows how to obtain rows from more than one table in a single query.

7.2.5 Joins Between Tables

Thus far, our queries have only accessed one table at a time. Queries can access multiple tables at once, or access the same table in such a way that multiple rows of the table are being processed at the same time. A query that accesses multiple rows of the same or different tables at one time is called a *join* query. For example, say you wish to list all the employee records together with the name and location of the associated department. To do that, we need to compare the `deptno` column of each row of the `emp` table with the `deptno` column of all rows in the `dept` table, and select the pairs of rows where these values match. This would be accomplished by the following query:

```
SELECT emp.ename, emp.sal, dept.deptno, dept.dname, dept.loc FROM emp, dept
WHERE emp.deptno = dept.deptno;
```

| ename | sal | deptno | dname | loc |
|--------|---------|--------|------------|----------|
| MILLER | 1300.00 | 10 | ACCOUNTING | NEW YORK |
| CLARK | 2450.00 | 10 | ACCOUNTING | NEW YORK |
| KING | 5000.00 | 10 | ACCOUNTING | NEW YORK |
| SCOTT | 3000.00 | 20 | RESEARCH | DALLAS |
| JONES | 2975.00 | 20 | RESEARCH | DALLAS |
| SMITH | 800.00 | 20 | RESEARCH | DALLAS |
| ADAMS | 1100.00 | 20 | RESEARCH | DALLAS |
| FORD | 3000.00 | 20 | RESEARCH | DALLAS |
| WARD | 1250.00 | 30 | SALES | CHICAGO |
| TURNER | 1500.00 | 30 | SALES | CHICAGO |
| ALLEN | 1600.00 | 30 | SALES | CHICAGO |
| BLAKE | 2850.00 | 30 | SALES | CHICAGO |
| MARTIN | 1250.00 | 30 | SALES | CHICAGO |
| JAMES | 950.00 | 30 | SALES | CHICAGO |

(14 rows)

Observe two things about the result set:

- There is no result row for department 40. This is because there is no matching entry in the `emp` table for department 40, so the join ignores the unmatched rows in the `dept` table. Shortly we will see how this can be fixed.
- It is more desirable to list the output columns qualified by table name rather than using * or leaving out the qualification as follows:

```
SELECT ename, sal, dept.deptno, dname, loc FROM emp, dept WHERE emp.deptno =
dept.deptno;
```

Since all the columns had different names (except for `deptno` which therefore must be qualified), the parser automatically found out which table they belong to, but it is good style to fully qualify column names in join queries:

Join queries of the kind seen thus far can also be written in this alternative form:

```
SELECT emp.ename, emp.sal, dept.deptno, dept.dname, dept.loc FROM emp INNER
JOIN dept ON emp.deptno = dept.deptno;
```

This syntax is not as commonly used as the one above, but we show it here to help you understand the following topics.

You will notice that in all the above results for joins no employees were returned that belonged to department 40 and as a consequence, the record for department 40 never appears. Now we will figure out how we can get the department 40 record in the results despite the fact that there are no matching employees. What we want the query to do is to scan the `dept` table and for each row to find the matching `emp` row. If no matching row is found we want some “empty” values to be substituted for the `emp` table’s columns. This kind of query is called an *outer join*. (The joins we have seen so far are *inner joins*.) The command looks like this:

```
SELECT emp.ename, emp.sal, dept.deptno, dept.dname, dept.loc FROM dept LEFT
OUTER JOIN emp ON emp.deptno = dept.deptno;
```

| ename | sal | deptno | dname | loc |
|--------|---------|--------|------------|----------|
| MILLER | 1300.00 | 10 | ACCOUNTING | NEW YORK |
| CLARK | 2450.00 | 10 | ACCOUNTING | NEW YORK |
| KING | 5000.00 | 10 | ACCOUNTING | NEW YORK |
| SCOTT | 3000.00 | 20 | RESEARCH | DALLAS |
| JONES | 2975.00 | 20 | RESEARCH | DALLAS |
| SMITH | 800.00 | 20 | RESEARCH | DALLAS |
| ADAMS | 1100.00 | 20 | RESEARCH | DALLAS |
| FORD | 3000.00 | 20 | RESEARCH | DALLAS |
| WARD | 1250.00 | 30 | SALES | CHICAGO |
| TURNER | 1500.00 | 30 | SALES | CHICAGO |
| ALLEN | 1600.00 | 30 | SALES | CHICAGO |
| BLAKE | 2850.00 | 30 | SALES | CHICAGO |
| MARTIN | 1250.00 | 30 | SALES | CHICAGO |
| JAMES | 950.00 | 30 | SALES | CHICAGO |
| | | 40 | OPERATIONS | BOSTON |

(15 rows)

This query is called a *left outer join* because the table mentioned on the left of the join operator will have each of its rows in the output at least once, whereas the table on the right will only have those rows output that match some row of the left table. When a left-table row is selected for which there is no right-table match, empty (`NULL`) values are substituted for the right-table columns.

An alternative syntax for an outer join is to use the outer join operator, “`(+)`”, in the join condition within the `WHERE` clause. The outer join operator is placed after the column name of the table for which null values should be substituted for unmatched rows. So for all the rows in the `dept` table that have no matching rows in the `emp` table, Advanced Server returns null for any select list expressions containing columns of `emp`. Hence the above example could be rewritten as:

```
SELECT emp.ename, emp.sal, dept.deptno, dept.dname, dept.loc FROM dept, emp
WHERE emp.deptno(+) = dept.deptno;
```

| ename | sal | deptno | dname | loc |
|--------|---------|--------|------------|----------|
| MILLER | 1300.00 | 10 | ACCOUNTING | NEW YORK |
| CLARK | 2450.00 | 10 | ACCOUNTING | NEW YORK |
| KING | 5000.00 | 10 | ACCOUNTING | NEW YORK |
| SCOTT | 3000.00 | 20 | RESEARCH | DALLAS |
| JONES | 2975.00 | 20 | RESEARCH | DALLAS |
| SMITH | 800.00 | 20 | RESEARCH | DALLAS |
| ADAMS | 1100.00 | 20 | RESEARCH | DALLAS |
| FORD | 3000.00 | 20 | RESEARCH | DALLAS |
| WARD | 1250.00 | 30 | SALES | CHICAGO |
| TURNER | 1500.00 | 30 | SALES | CHICAGO |
| ALLEN | 1600.00 | 30 | SALES | CHICAGO |
| BLAKE | 2850.00 | 30 | SALES | CHICAGO |
| MARTIN | 1250.00 | 30 | SALES | CHICAGO |
| JAMES | 950.00 | 30 | SALES | CHICAGO |
| | | 40 | OPERATIONS | BOSTON |

(15 rows)

| MILLER | 1300.00 | 10 | ACCOUNTING | NEW YORK |
|--------|---------|----|------------|----------|
| CLARK | 2450.00 | 10 | ACCOUNTING | NEW YORK |
| KING | 5000.00 | 10 | ACCOUNTING | NEW YORK |
| SCOTT | 3000.00 | 20 | RESEARCH | DALLAS |
| JONES | 2975.00 | 20 | RESEARCH | DALLAS |
| SMITH | 800.00 | 20 | RESEARCH | DALLAS |
| ADAMS | 1100.00 | 20 | RESEARCH | DALLAS |
| FORD | 3000.00 | 20 | RESEARCH | DALLAS |
| WARD | 1250.00 | 30 | SALES | CHICAGO |
| TURNER | 1500.00 | 30 | SALES | CHICAGO |
| ALLEN | 1600.00 | 30 | SALES | CHICAGO |
| BLAKE | 2850.00 | 30 | SALES | CHICAGO |
| MARTIN | 1250.00 | 30 | SALES | CHICAGO |
| JAMES | 950.00 | 30 | SALES | CHICAGO |
| | | 40 | OPERATIONS | BOSTON |

(15 rows)

We can also join a table against itself. This is called a *self join*. As an example, suppose we wish to find the name of each employee along with the name of that employee's manager. So we need to compare the `mgr` column of each `emp` row to the `empno` column of all other `emp` rows.

```
SELECT e1.ename || ' works for ' || e2.ename AS "Employees and their Managers"
FROM emp e1, emp e2 WHERE e1.mgr = e2.empno;
```

Employees and their Managers

| FORD works for JONES |
|------------------------|
| SCOTT works for JONES |
| WARD works for BLAKE |
| TURNER works for BLAKE |
| MARTIN works for BLAKE |
| JAMES works for BLAKE |
| ALLEN works for BLAKE |
| MILLER works for CLARK |
| ADAMS works for SCOTT |
| CLARK works for KING |
| BLAKE works for KING |
| JONES works for KING |
| SMITH works for FORD |

(13 rows)

Here, the `emp` table has been re-labeled as `e1` to represent the employee row in the select list and in the join condition, and also as `e2` to represent the matching employee row acting as manager in the select list and in the join condition. These kinds of aliases can be used in other queries to save some typing, for example:

```
SELECT e.ename, e.mgr, d.deptno, d.dname, d.loc FROM emp e, dept d WHERE
e.deptno = d.deptno;
```

| ename | mgr | deptno | dname | loc |
|--------|------|--------|------------|----------|
| MILLER | 7782 | 10 | ACCOUNTING | NEW YORK |
| CLARK | 7839 | 10 | ACCOUNTING | NEW YORK |
| KING | | 10 | ACCOUNTING | NEW YORK |
| SCOTT | 7566 | 20 | RESEARCH | DALLAS |
| JONES | 7839 | 20 | RESEARCH | DALLAS |

```

SMITH | 7902 | 20 | RESEARCH | DALLAS
ADAMS | 7788 | 20 | RESEARCH | DALLAS
FORD | 7566 | 20 | RESEARCH | DALLAS
WARD | 7698 | 30 | SALES | CHICAGO
TURNER | 7698 | 30 | SALES | CHICAGO
ALLEN | 7698 | 30 | SALES | CHICAGO
BLAKE | 7839 | 30 | SALES | CHICAGO
MARTIN | 7698 | 30 | SALES | CHICAGO
JAMES | 7698 | 30 | SALES | CHICAGO
(14 rows)

```

This style of abbreviating will be encountered quite frequently.

7.2.6 Aggregate Functions

Like most other relational database products, Advanced Server supports aggregate functions. An aggregate function computes a single result from multiple input rows. For example, there are aggregates to compute the `COUNT`, `SUM`, `AVG` (average), `MAX` (maximum), and `MIN` (minimum) over a set of rows.

As an example, the highest and lowest salaries can be found with the following query:

```
SELECT MAX(sal) highest_salary, MIN(sal) lowest_salary FROM emp;
```

```

highest_salary | lowest_salary
-----+-----
      5000.00 |     800.00
(1 row)

```

If we wanted to find the employee with the largest salary, we may be tempted to try:

```
SELECT ename FROM emp WHERE sal = MAX(sal);
```

```
ERROR: aggregates not allowed in WHERE clause
```

This does not work because the aggregate function, `MAX`, cannot be used in the `WHERE` clause. This restriction exists because the `WHERE` clause determines the rows that will go into the aggregation stage so it has to be evaluated before aggregate functions are computed. However, the query can be restated to accomplish the intended result by using a *subquery*:

```
SELECT ename FROM emp WHERE sal = (SELECT MAX(sal) FROM emp);
```

```

ename
-----
KING
(1 row)

```

The subquery is an independent computation that obtains its own result separately from the outer query.

Aggregates are also very useful in combination with the `GROUP BY` clause. For example, the following query gets the highest salary in each department.

```
SELECT deptno, MAX(sal) FROM emp GROUP BY deptno;
```

```
deptno | max
-----+-----
 10 | 5000.00
 20 | 3000.00
 30 | 2850.00
(3 rows)
```

This query produces one output row per department. Each aggregate result is computed over the rows matching that department. These grouped rows can be filtered using the `HAVING` clause.

```
SELECT deptno, MAX(sal) FROM emp GROUP BY deptno HAVING AVG(sal) > 2000;
```

```
deptno | max
-----+-----
 10 | 5000.00
 20 | 3000.00
(2 rows)
```

This query gives the same results for only those departments that have an average salary greater than 2000.

Finally, the following query takes into account only the highest paid employees who are analysts in each department.

```
SELECT deptno, MAX(sal) FROM emp WHERE job = 'ANALYST' GROUP BY deptno
HAVING AVG(sal) > 2000;
```

```
deptno | max
-----+-----
 20 | 3000.00
(1 row)
```

There is a subtle distinction between the `WHERE` and `HAVING` clauses. The `WHERE` clause filters out rows before grouping occurs and aggregate functions are applied. The `HAVING` clause applies filters on the results after rows have been grouped and aggregate functions have been computed for each group.

So in the previous example, only employees who are analysts are considered. From this subset, the employees are grouped by department and only those groups where the average salary of analysts in the group is greater than 2000 are in the final result. This is true of only the group for department 20 and the maximum analyst salary in department 20 is 3000.00.

7.2.7 Updates

The column values of existing rows can be changed using the `UPDATE` command. For example, the following sequence of commands shows the before and after results of giving everyone who is a manager a 10% raise:

```
SELECT ename, sal FROM emp WHERE job = 'MANAGER';
```

```
ename | sal
-----+-----
JONES | 2975.00
BLAKE | 2850.00
```

```
CLARK | 2450.00
(3 rows)
```

```
UPDATE emp SET sal = sal * 1.1 WHERE job = 'MANAGER';
```

```
SELECT ename, sal FROM emp WHERE job = 'MANAGER';
```

```
ename | sal
-----+
JONES | 3272.50
BLAKE | 3135.00
CLARK | 2695.00
(3 rows)
```

7.2.8 Deletions

Rows can be removed from a table using the **DELETE** command. For example, the following sequence of commands shows the before and after results of deleting all employees in department 20.

```
SELECT ename, deptno FROM emp;
```

```
ename | deptno
-----+
SMITH | 20
ALLEN | 30
WARD | 30
JONES | 20
MARTIN | 30
BLAKE | 30
CLARK | 10
SCOTT | 20
KING | 10
TURNER | 30
ADAMS | 20
JAMES | 30
FORD | 20
MILLER | 10
(14 rows)
```

```
DELETE FROM emp WHERE deptno = 20;
```

```
SELECT ename, deptno FROM emp;
ename | deptno
```

```
-----+
ALLEN | 30
WARD | 30
MARTIN | 30
BLAKE | 30
CLARK | 10
KING | 10
```

```
TURNER | 30
JAMES | 30
MILLER | 10
(9 rows)
```

Be extremely careful of giving a `DELETE` command without a `WHERE` clause such as the following:

```
| DELETE FROM tablename;
```

This statement will remove all rows from the given table, leaving it completely empty. The system will not request confirmation before doing this.

7.2.9 The SQL Language

Advanced Server supports SQL language that is compatible with Oracle syntax as well as syntax and commands for extended functionality (functionality that does not provide database compatibility for Oracle or support Oracle-styled applications).

The *Database Compatibility for Oracle Developer's SQL Guide* provides detailed information about:

- Compatible SQL syntax and language elements
- Data types
- Supported SQL command syntax

To review a copy of the guide, visit the Advanced Server website at:

```
| https://www.enterprisedb.com/edb-docs
```

7.3 Advanced Concepts

The previous section discussed the basics of using SQL to store and access your data in Advanced Server. This section discusses more advanced SQL features that may simplify management and prevent loss or corruption of your data.

7.3.1 Views

Consider the following `SELECT` command.

```
SELECT ename, sal, sal * 24 AS yearly_salary, deptno FROM emp;
```

| ename | sal | yearly_salary | deptno |
|-------|-------|---------------|--------|
| ----- | ----- | ----- | ----- |

| | | | |
|--------|---------|-----------|----|
| SMITH | 800.00 | 19200.00 | 20 |
| ALLEN | 1600.00 | 38400.00 | 30 |
| WARD | 1250.00 | 30000.00 | 30 |
| JONES | 2975.00 | 71400.00 | 20 |
| MARTIN | 1250.00 | 30000.00 | 30 |
| BLAKE | 2850.00 | 68400.00 | 30 |
| CLARK | 2450.00 | 58800.00 | 10 |
| SCOTT | 3000.00 | 72000.00 | 20 |
| KING | 5000.00 | 120000.00 | 10 |
| TURNER | 1500.00 | 36000.00 | 30 |
| ADAMS | 1100.00 | 26400.00 | 20 |
| JAMES | 950.00 | 22800.00 | 30 |
| FORD | 3000.00 | 72000.00 | 20 |
| MILLER | 1300.00 | 31200.00 | 10 |

(14 rows)

If this is a query that is used repeatedly, a shorthand method of reusing this query without re-typing the entire `SELECT` command each time is to create a *view* as shown below.

```
CREATE VIEW employee_pay AS SELECT ename, sal, sal * 24 AS yearly_salary,
deptno FROM emp;
```

The view name, `employee_pay`, can now be used like an ordinary table name to perform the query.

```
SELECT * FROM employee_pay;
```

| ename | sal | yearly_salary | deptno |
|--------|---------|---------------|--------|
| SMITH | 800.00 | 19200.00 | 20 |
| ALLEN | 1600.00 | 38400.00 | 30 |
| WARD | 1250.00 | 30000.00 | 30 |
| JONES | 2975.00 | 71400.00 | 20 |
| MARTIN | 1250.00 | 30000.00 | 30 |
| BLAKE | 2850.00 | 68400.00 | 30 |
| CLARK | 2450.00 | 58800.00 | 10 |
| SCOTT | 3000.00 | 72000.00 | 20 |
| KING | 5000.00 | 120000.00 | 10 |
| TURNER | 1500.00 | 36000.00 | 30 |
| ADAMS | 1100.00 | 26400.00 | 20 |
| JAMES | 950.00 | 22800.00 | 30 |
| FORD | 3000.00 | 72000.00 | 20 |
| MILLER | 1300.00 | 31200.00 | 10 |

(14 rows)

Making liberal use of views is a key aspect of good SQL database design. Views provide a consistent interface that encapsulate details of the structure of your tables which may change as your application evolves.

Views can be used in almost any place a real table can be used. Building views upon other views is not uncommon.

7.3.2 Foreign Keys

Suppose you want to make sure all employees belong to a valid department. This is called maintaining the *referential integrity* of your data. In simplistic database systems this would be implemented (if at all) by first looking at the `dept` table to check if a matching record exists, and then inserting or rejecting the new employee record. This approach has a number of problems and is very inconvenient. Advanced Server can make it easier for you.

A modified version of the `emp` table presented in [Creating a New Table](#) is shown in this section with the addition of a foreign key constraint. The modified `emp` table looks like the following:

```
CREATE TABLE emp (
    empno      NUMBER(4) NOT NULL CONSTRAINT emp_pk PRIMARY KEY,
    ename      VARCHAR2(10),
    job        VARCHAR2(9),
    mgr        NUMBER(4),
    hiredate   DATE,
    sal         NUMBER(7,2),
    comm       NUMBER(7,2),
    deptno    NUMBER(2) CONSTRAINT emp_ref_dept_fk
               REFERENCES dept(deptno)
);
```

If an attempt is made to issue the following `INSERT` command in the sample `emp` table, the foreign key constraint, `emp_ref_dept_fk`, ensures that department `50` exists in the `dept` table. Since it does not, the command is rejected.

```
INSERT INTO emp VALUES (8000,'JONES','CLERK',7902,'17-AUG-07',1200,NULL,50);
```

ERROR: insert or update on table "emp" violates foreign key constraint

"emp_ref_dept_fk"

DETAIL: Key (deptno)=(50) is not present in table "dept".

The behavior of foreign keys can be finely tuned to your application. Making correct use of foreign keys will definitely improve the quality of your database applications, so you are strongly encouraged to learn more about them.

7.3.3 The ROWNUM Pseudo-Column

`ROWNUM` is a pseudo-column that is assigned an incremental, unique integer value for each row based on the order the rows were retrieved from a query. Therefore, the first row retrieved will have `ROWNUM` of `1`; the second row will have `ROWNUM` of `2` and so on.

This feature can be used to limit the number of rows retrieved by a query. This is demonstrated in the following example:

```
SELECT empno, ename, job FROM emp WHERE ROWNUM < 5;
```

| empno | ename | job |
|-------|-------|----------|
| 7369 | SMITH | CLERK |
| 7499 | ALLEN | SALESMAN |
| 7521 | WARD | SALESMAN |
| 7566 | JONES | MANAGER |

(4 rows)

The `ROWNUM` value is assigned to each row before any sorting of the result set takes place. Thus, the result set is returned in the order given by the `ORDER BY` clause, but the `ROWNUM` values may not necessarily be in ascending order as shown in the following example:

```
SELECT ROWNUM, empno, ename, job FROM emp WHERE ROWNUM < 5 ORDER BY ename;
```

| rownum | empno | ename | job |
|--------|-------|-------|----------|
| 2 | 7499 | ALLEN | SALESMAN |
| 4 | 7566 | JONES | MANAGER |
| 1 | 7369 | SMITH | CLERK |
| 3 | 7521 | WARD | SALESMAN |

(4 rows)

The following example shows how a sequence number can be added to every row in the `jobhist` table. First a new column named, `seqno`, is added to the table and then `seqno` is set to `ROWNUM` in the `UPDATE` command.

```
ALTER TABLE jobhist ADD seqno NUMBER(3);
UPDATE jobhist SET seqno = ROWNUM;
```

The following `SELECT` command shows the new `seqno` values.

```
SELECT seqno, empno, TO_CHAR(startdate,'DD-MON-YY') AS start, job FROM
jobhist;
```

| seqno | empno | start | job |
|-------|-------|-----------|-----------|
| 1 | 7369 | 17-DEC-80 | CLERK |
| 2 | 7499 | 20-FEB-81 | SALESMAN |
| 3 | 7521 | 22-FEB-81 | SALESMAN |
| 4 | 7566 | 02-APR-81 | MANAGER |
| 5 | 7654 | 28-SEP-81 | SALESMAN |
| 6 | 7698 | 01-MAY-81 | MANAGER |
| 7 | 7782 | 09-JUN-81 | MANAGER |
| 8 | 7788 | 19-APR-87 | CLERK |
| 9 | 7788 | 13-APR-88 | CLERK |
| 10 | 7788 | 05-MAY-90 | ANALYST |
| 11 | 7839 | 17-NOV-81 | PRESIDENT |
| 12 | 7844 | 08-SEP-81 | SALESMAN |
| 13 | 7876 | 23-MAY-87 | CLERK |
| 14 | 7900 | 03-DEC-81 | CLERK |
| 15 | 7900 | 15-JAN-83 | CLERK |
| 16 | 7902 | 03-DEC-81 | ANALYST |
| 17 | 7934 | 23-JAN-82 | CLERK |

(17 rows)

7.3.4 Synonyms

A *synonym* is an identifier that can be used to reference another database object in a SQL statement. A synonym is useful in cases where a database object would normally require full qualification by schema name to be properly referenced in a SQL statement. A synonym defined for that object simplifies the reference to a single, unqualified

name.

Advanced Server supports synonyms for:

- tables
- views
- materialized views
- sequences
- procedures
- functions
- types
- objects that are accessible through a database link
- other synonyms

Neither the referenced schema or referenced object must exist at the time that you create the synonym; a synonym may refer to a non-existent object or schema. A synonym will become invalid if you drop the referenced object or schema. You must explicitly drop a synonym to remove it.

As with any other schema object, Advanced Server uses the search path to resolve unqualified synonym names. If you have two synonyms with the same name, an unqualified reference to a synonym will resolve to the first synonym with the given name in the search path. If `public` is in your search path, you can refer to a synonym in that schema without qualifying that name.

When Advanced Server executes an SQL command, the privileges of the current user are checked against the synonym's underlying database object; if the user does not have the proper permissions for that object, the SQL command will fail.

Creating a Synonym

Use the `CREATE SYNONYM` command to create a synonym. The syntax is:

```
CREATE [OR REPLACE] [PUBLIC] SYNONYM [<schema>.]<syn_name>
    FOR <object_schema>.<object_name>[<@dblink_name>];
```

Parameters:

`syn_name`

`syn_name` is the name of the synonym. A synonym name must be unique within a schema.

`schema`

`schema` specifies the name of the schema that the synonym resides in. If you do not specify a schema name, the synonym is created in the first existing schema in your search path.

`object_name`

`object_name` specifies the name of the object.

`object_schema`

`object_schema` specifies the name of the schema that the object resides in.

`dblink_name`

`dblink_name` specifies the name of the database link through which a target object may be accessed.

Include the `REPLACE` clause to replace an existing synonym definition with a new synonym definition.

Include the `PUBLIC` clause to create the synonym in the `public` schema. Compatible with Oracle databases, the `CREATE PUBLIC SYNONYM` command creates a synonym that resides in the `public` schema:

```
CREATE [OR REPLACE] PUBLIC SYNONYM <syn_name> FOR
<object_schema>.<object_name>;
```

This just a shorthand way to write:

```
CREATE [OR REPLACE] SYNONYM public.<syn_name> FOR
<object_schema>.<object_name>;
```

The following example creates a synonym named `personnel` that refers to the `enterprisedb.emp` table.

```
CREATE SYNONYM personnel FOR enterprisedb.emp;
```

Unless the synonym is schema qualified in the `CREATE SYNONYM` command, it will be created in the first existing schema in your search path. You can view your search path by executing the following command:

```
SHOW SEARCH_PATH;
```

```
search_path
-----
development,accounting
(1 row)
```

In our example, if a schema named `development` does not exist, the synonym will be created in the schema named `accounting`.

Now, the `emp` table in the `enterprisedb` schema can be referenced in any SQL statement (DDL or DML), by using the synonym, `personnel`:

```
INSERT INTO personnel VALUES (8142,'ANDERSON','CLERK',7902,'17-DEC-06',1300,NULL,20);
```

```
SELECT * FROM personnel;
```

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|----------|-----------|------|--------------------|---------|---------|--------|
| 7369 | SMITH | CLERK | 7902 | 17-DEC-80 00:00:00 | 800.00 | | 20 |
| 7499 | ALLEN | SALESMAN | 7698 | 20-FEB-81 00:00:00 | 1600.00 | 300.00 | 30 |
| 7521 | WARD | SALESMAN | 7698 | 22-FEB-81 00:00:00 | 1250.00 | 500.00 | 30 |
| 7566 | JONES | MANAGER | 7839 | 02-APR-81 00:00:00 | 2975.00 | | 20 |
| 7654 | MARTIN | SALESMAN | 7698 | 28-SEP-81 00:00:00 | 1250.00 | 1400.00 | 30 |
| 7698 | BLAKE | MANAGER | 7839 | 01-MAY-81 00:00:00 | 2850.00 | | 30 |
| 7782 | CLARK | MANAGER | 7839 | 09-JUN-81 00:00:00 | 2450.00 | | 10 |
| 7788 | SCOTT | ANALYST | 7566 | 19-APR-87 00:00:00 | 3000.00 | | 20 |
| 7839 | KING | PRESIDENT | | 17-NOV-81 00:00:00 | 5000.00 | | 10 |
| 7844 | TURNER | SALESMAN | 7698 | 08-SEP-81 00:00:00 | 1500.00 | 0.00 | 30 |
| 7876 | ADAMS | CLERK | 7788 | 23-MAY-87 00:00:00 | 1100.00 | | 20 |
| 7900 | JAMES | CLERK | 7698 | 03-DEC-81 00:00:00 | 950.00 | | 30 |
| 7902 | FORD | ANALYST | 7566 | 03-DEC-81 00:00:00 | 3000.00 | | 20 |
| 7934 | MILLER | CLERK | 7782 | 23-JAN-82 00:00:00 | 1300.00 | | 10 |
| 8142 | ANDERSON | CLERK | 7902 | 17-DEC-06 00:00:00 | 1300.00 | | 20 |

(15 rows)

Deleting a Synonym

To delete a synonym, use the command, `DROP SYNONYM`. The syntax is:

```
DROP [PUBLIC] SYNONYM [<schema>.] <syn_name>
```

Parameters:`syn_name`

`syn_name` is the name of the synonym. A synonym name must be unique within a schema.

`schema`

`schema` specifies the name of the schema in which the synonym resides.

Like any other object that can be schema-qualified, you may have two synonyms with the same name in your search path. To disambiguate the name of the synonym that you are dropping, include a schema name. Unless a synonym is schema qualified in the `DROP SYNONYM` command, Advanced Server deletes the first instance of the synonym it finds in your search path.

You can optionally include the `PUBLIC` clause to drop a synonym that resides in the `public` schema. Compatible with Oracle databases, the `DROP PUBLIC SYNONYM` command drops a synonym that resides in the `public` schema:

`DROP PUBLIC SYNONYM <syn_name>;`

The following example drops the synonym, `personnel`:

```
DROP SYNONYM personnel;
```

7.3.5 Hierarchical Queries

A *hierarchical query* is a type of query that returns the rows of the result set in a hierarchical order based upon data forming a parent-child relationship. A hierarchy is typically represented by an inverted tree structure. The tree is comprised of interconnected *nodes*. Each node may be connected to none, one, or multiple *child* nodes. Each node is connected to one *parent* node except for the top node which has no parent. This node is the *root* node. Each tree has exactly one root node. Nodes that don't have any children are called *leaf* nodes. A tree always has at least one leaf node - e.g., the trivial case where the tree is comprised of a single node. In this case it is both the root and the leaf.

In a hierarchical query the rows of the result set represent the nodes of one or more trees.

Note

It is possible that a single, given row may appear in more than one tree and thus appear more than once in the result set.

The hierarchical relationship in a query is described by the `CONNECT BY` clause which forms the basis of the order in which rows are returned in the result set. The context of where the `CONNECT BY` clause and its associated optional clauses appear in the `SELECT` command is shown below.

```
SELECT <select_list> FROM <table_expression> [ WHERE ... ]
[ START WITH <start_expression> ]
CONNECT BY { PRIOR <parent_expr> = <child_expr> |
<child_expr> = PRIOR <parent_expr> }
[ ORDER SIBLINGS BY <column1> [ ASC | DESC ]
[, <column2> [ ASC | DESC ] ] ...
[ GROUP BY ...]
[ HAVING ...]
[ <other> ...]
```

`select_list` is one or more expressions that comprise the fields of the result set. `table_expression` is one or more tables or views from which the rows of the result set originate. `other` is any additional legal `SELECT` command clauses. The clauses pertinent to hierarchical queries, `START WITH`, `CONNECT BY`, and `ORDER SIBLINGS BY` are described in the following sections.

Note

At this time, Advanced Server does not support the use of `AND` (or other operators) in the `CONNECT BY` clause.

7.3.5.1 Defining the Parent/Child Relationship

For any given row, its parent and its children are determined by the `CONNECT BY` clause. The `CONNECT BY` clause must consist of two expressions compared with the equals (=) operator. In addition, one of these two expressions must be preceded by the keyword, `PRIOR`.

For any given row, to determine its children:

1. Evaluate `parent_expr` on the given row
2. Evaluate `child_expr` on any other row resulting from the evaluation of `table_expression`
3. If `parent_expr = child_expr`, then this row is a child node of the given parent row
4. Repeat the process for all remaining rows in `table_expression`. All rows that satisfy the equation in step 3 are the children nodes of the given parent row.

Note

The evaluation process to determine if a row is a child node occurs on every row returned by `table_expression` before the `WHERE` clause is applied to `table_expression`.

By iteratively repeating this process treating each child node found in the prior steps as a parent, an inverted tree of nodes is constructed. The process is complete when the final set of child nodes has no children of their own - these are the leaf nodes.


A `SELECT` command that includes a `CONNECT BY` clause typically includes the `START WITH` clause. The `START WITH` clause determines the rows that are to be the root nodes - i.e., the rows that are the initial parent nodes upon which the algorithm described previously is to be applied. This is further explained in the following section.

7.3.5.2 Selecting the Root Nodes

The `START WITH` clause is used to determine the row(s) selected by `table_expression` that are to be used as the root nodes. All rows selected by `table_expression` where `start_expression` evaluates to true become a root node of a tree. Thus, the number of potential trees in the result set is equal to the number of root nodes. As a consequence, if the `START WITH` clause is omitted, then every row returned by `table_expression` is a root of its own tree.

7.3.5.3 Organization Tree in the Sample Application

Consider the `emp` table of the sample application. The rows of the `emp` table form a hierarchy based upon the `mgr` column which contains the employee number of the employee's manager. Each employee has at most, one manager. `KING` is the president of the company so he has no manager, therefore `KING`'s `mgr` column is null. Also, it is possible for an employee to act as a manager for more than one employee. This relationship forms a typical, tree-structured, hierarchical organization chart as illustrated below.



To form a hierarchical query based upon this relationship, the `SELECT` command includes the clause, `CONNECT BY PRIOR empno = mgr`. For example, given the company president, `KING`, with employee number 7839, any employee whose `mgr` column is 7839 reports directly to `KING` which is true for `JONES`, `BLAKE`, and `CLARK` (these are the child nodes of `KING`). Similarly, for employee, `JONES`, any other employee with `mgr` column equal to 7566 is a child node of `JONES` - these are `SCOTT` and `FORD` in this example.

The top of the organization chart is `KING` so there is one root node in this tree. The `START WITH mgr IS NULL` clause selects only `KING` as the initial root node.

The complete `SELECT` command is shown below.

```

SELECT ename, empno, mgr
FROM emp
START WITH mgr IS NULL
CONNECT BY PRIOR empno = mgr;
  
```

The rows in the query output traverse each branch from the root to leaf moving in a top-to-bottom, left-to-right order. Below is the output from this query.

| ename | | empno | | mgr |
|--------|--|-------|--|------|
| KING | | 7839 | | |
| JONES | | 7566 | | 7839 |
| SCOTT | | 7788 | | 7566 |
| ADAMS | | 7876 | | 7788 |
| FORD | | 7902 | | 7566 |
| SMITH | | 7369 | | 7902 |
| BLAKE | | 7698 | | 7839 |
| ALLEN | | 7499 | | 7698 |
| WARD | | 7521 | | 7698 |
| MARTIN | | 7654 | | 7698 |
| TURNER | | 7844 | | 7698 |
| JAMES | | 7900 | | 7698 |
| CLARK | | 7782 | | 7839 |
| MILLER | | 7934 | | 7782 |

(14 rows)

7.3.5.4 Node Level

`LEVEL` is a pseudo-column that can be used wherever a column can appear in the `SELECT` command. For each row in the result set, `LEVEL` returns a non-zero integer value designating the depth in the hierarchy of the node represented by this row. The `LEVEL` for root nodes is 1. The `LEVEL` for direct children of root nodes is 2, and so on.

The following query is a modification of the previous query with the addition of the `LEVEL` pseudo-column. In addition, using the `LEVEL` value, the employee names are indented to further emphasize the depth in the hierarchy of each row.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr
FROM emp START WITH mgr IS NULL
CONNECT BY PRIOR empno = mgr;
```

The output from this query follows.

| level | employee | empno | mgr |
|-------|----------|-------|------|
| 1 | KING | 7839 | |
| 2 | JONES | 7566 | 7839 |
| 3 | SCOTT | 7788 | 7566 |
| 4 | ADAMS | 7876 | 7788 |
| 3 | FORD | 7902 | 7566 |
| 4 | SMITH | 7369 | 7902 |
| 2 | BLAKE | 7698 | 7839 |
| 3 | ALLEN | 7499 | 7698 |
| 3 | WARD | 7521 | 7698 |
| 3 | MARTIN | 7654 | 7698 |
| 3 | TURNER | 7844 | 7698 |
| 3 | JAMES | 7900 | 7698 |
| 2 | CLARK | 7782 | 7839 |
| 3 | MILLER | 7934 | 7782 |

(14 rows)

Nodes that share a common parent and are at the same level are called *siblings*. For example in the above output, employees `ALLEN`, `WARD`, `MARTIN`, `TURNER`, and `JAMES` are siblings since they are all at level three with parent, `BLAKE`. `JONES`, `BLAKE`, and `CLARK` are siblings since they are at level two and `KING` is their common parent.

7.3.5.5 Ordering the Siblings

The result set can be ordered so the siblings appear in ascending or descending order by selected column value(s) using the `ORDER SIBLINGS BY` clause. This is a special case of the `ORDER BY` clause that can be used only with hierarchical queries.

The previous query is further modified with the addition of `ORDER SIBLINGS BY ename ASC`.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr
```

```
FROM emp START WITH mgr IS NULL
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

The output from the prior query is now modified so the siblings appear in ascending order by name. Siblings **BLAKE**, **CLARK**, and **JONES** are now alphabetically arranged under **KING**. Siblings **ALLEN**, **JAMES**, **MARTIN**, **TURNER**, and **WARD** are alphabetically arranged under **BLAKE**, and so on.

| level | employee | empno | mgr |
|-------|----------|-------|------|
| 1 | KING | 7839 | |
| 2 | BLAKE | 7698 | 7839 |
| 3 | ALLEN | 7499 | 7698 |
| 3 | JAMES | 7900 | 7698 |
| 3 | MARTIN | 7654 | 7698 |
| 3 | TURNER | 7844 | 7698 |
| 3 | WARD | 7521 | 7698 |
| 2 | CLARK | 7782 | 7839 |
| 3 | MILLER | 7934 | 7782 |
| 2 | JONES | 7566 | 7839 |
| 3 | FORD | 7902 | 7566 |
| 4 | SMITH | 7369 | 7902 |
| 3 | SCOTT | 7788 | 7566 |
| 4 | ADAMS | 7876 | 7788 |

(14 rows)

This final example adds the **WHERE** clause and starts with three root nodes. After the node tree is constructed, the **WHERE** clause filters out rows in the tree to form the result set.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr
FROM emp WHERE mgr IN (7839, 7782, 7902, 7788)
START WITH ename IN ('BLAKE','CLARK','JONES')
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

The output from the query shows three root nodes (level one) - **BLAKE**, **CLARK**, and **JONES**. In addition, rows that do not satisfy the **WHERE** clause have been eliminated from the output.

| level | employee | empno | mgr |
|-------|----------|-------|------|
| 1 | BLAKE | 7698 | 7839 |
| 1 | CLARK | 7782 | 7839 |
| 2 | MILLER | 7934 | 7782 |
| 1 | JONES | 7566 | 7839 |
| 3 | SMITH | 7369 | 7902 |
| 3 | ADAMS | 7876 | 7788 |

(6 rows)

7.3.5.6 Retrieving the Root Node with CONNECT_BY_ROOT

CONNECT_BY_ROOT is a unary operator that can be used to qualify a column in order to return the column's

value of the row considered to be the root node in relation to the current row.

Note

A *unary operator* operates on a single operand, which in the case of `CONNECT_BY_ROOT`, is the column name following the `CONNECT_BY_ROOT` keyword.

In the context of the `SELECT` list, the `CONNECT_BY_ROOT` operator is shown by the following.

```
SELECT [...] CONNECT_BY_ROOT <column> [, ...]
  FROM <table_expression> ...
```

The following are some points to note about the `CONNECT_BY_ROOT` operator.

- The `CONNECT_BY_ROOT` operator can be used in the `SELECT` list, the `WHERE` clause, the `GROUP BY` clause, the `HAVING` clause, the `ORDER BY` clause, and the `ORDER SIBLINGS BY` clause as long as the `SELECT` command is for a hierarchical query.
- The `CONNECT_BY_ROOT` operator cannot be used in the `CONNECT BY` clause or the `START WITH` clause of the hierarchical query.
- It is possible to apply `CONNECT_BY_ROOT` to an expression involving a column, but to do so, the expression must be enclosed within parentheses.

The following query shows the use of the `CONNECT_BY_ROOT` operator to return the employee number and employee name of the root node for each employee listed in the result set based on trees starting with employees `BLAKE`, `CLARK`, and `JONES`.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr,
CONNECT_BY_ROOT empno "mgr empno",
CONNECT_BY_ROOT ename "mgr ename"
FROM emp
START WITH ename IN ('BLAKE','CLARK','JONES')
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

Note that the output from the query shows that all of the root nodes in columns `mgr empno` and `mgr ename` are one of the employees, `BLAKE`, `CLARK`, or `JONES`, listed in the `START WITH` clause.

| level | employee | empno | mgr | mgr empno | mgr ename |
|-------|----------|-------|------|-----------|-----------|
| 1 | BLAKE | 7698 | 7839 | 7698 | BLAKE |
| 2 | ALLEN | 7499 | 7698 | 7698 | BLAKE |
| 2 | JAMES | 7900 | 7698 | 7698 | BLAKE |
| 2 | MARTIN | 7654 | 7698 | 7698 | BLAKE |
| 2 | TURNER | 7844 | 7698 | 7698 | BLAKE |
| 2 | WARD | 7521 | 7698 | 7698 | BLAKE |
| 1 | CLARK | 7782 | 7839 | 7782 | CLARK |
| 2 | MILLER | 7934 | 7782 | 7782 | CLARK |
| 1 | JONES | 7566 | 7839 | 7566 | JONES |
| 2 | FORD | 7902 | 7566 | 7566 | JONES |
| 3 | SMITH | 7369 | 7902 | 7566 | JONES |
| 2 | SCOTT | 7788 | 7566 | 7566 | JONES |
| 3 | ADAMS | 7876 | 7788 | 7566 | JONES |

(13 rows)

The following is a similar query, but producing only one tree starting with the single, top-level, employee where the `mgr` column is null.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr,
```

```
CONNECT_BY_ROOT empno "mgr empno",
CONNECT_BY_ROOT ename "mgr ename"
FROM emp START WITH mgr IS NULL
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

In the following output, all of the root nodes in columns `mgr empno` and `mgr ename` indicate `KING` as the root for this particular query.

| level | employee | empno | mgr | mgr empno | mgr ename |
|-------|----------|-------|------|-----------|-----------|
| 1 | KING | 7839 | | 7839 | KING |
| 2 | BLAKE | 7698 | 7839 | 7839 | KING |
| 3 | ALLEN | 7499 | 7698 | 7839 | KING |
| 3 | JAMES | 7900 | 7698 | 7839 | KING |
| 3 | MARTIN | 7654 | 7698 | 7839 | KING |
| 3 | TURNER | 7844 | 7698 | 7839 | KING |
| 3 | WARD | 7521 | 7698 | 7839 | KING |
| 2 | CLARK | 7782 | 7839 | 7839 | KING |
| 3 | MILLER | 7934 | 7782 | 7839 | KING |
| 2 | JONES | 7566 | 7839 | 7839 | KING |
| 3 | FORD | 7902 | 7566 | 7839 | KING |
| 4 | SMITH | 7369 | 7902 | 7839 | KING |
| 3 | SCOTT | 7788 | 7566 | 7839 | KING |
| 4 | ADAMS | 7876 | 7788 | 7839 | KING |

(14 rows)

By contrast, the following example omits the `START WITH` clause thereby resulting in fourteen trees.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr,
CONNECT_BY_ROOT empno "mgr empno",
CONNECT_BY_ROOT ename "mgr ename"
FROM emp
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

The following is the output from the query. Each node appears at least once as a root node under the `mgr empno` and `mgr ename` columns since even the leaf nodes form the top of their own trees.

| level | employee | empno | mgr | mgr empno | mgr ename |
|-------|----------|-------|------|-----------|-----------|
| 1 | ADAMS | 7876 | 7788 | 7876 | ADAMS |
| 1 | ALLEN | 7499 | 7698 | 7499 | ALLEN |
| 1 | BLAKE | 7698 | 7839 | 7698 | BLAKE |
| 2 | ALLEN | 7499 | 7698 | 7698 | BLAKE |
| 2 | JAMES | 7900 | 7698 | 7698 | BLAKE |
| 2 | MARTIN | 7654 | 7698 | 7698 | BLAKE |
| 2 | TURNER | 7844 | 7698 | 7698 | BLAKE |
| 2 | WARD | 7521 | 7698 | 7698 | BLAKE |
| 1 | CLARK | 7782 | 7839 | 7782 | CLARK |
| 2 | MILLER | 7934 | 7782 | 7782 | CLARK |
| 1 | FORD | 7902 | 7566 | 7902 | FORD |
| 2 | SMITH | 7369 | 7902 | 7902 | FORD |
| 1 | JAMES | 7900 | 7698 | 7900 | JAMES |
| 1 | JONES | 7566 | 7839 | 7566 | JONES |
| 2 | FORD | 7902 | 7566 | 7566 | JONES |

| | | | | | |
|---|--------|------|------|------|--------|
| 3 | SMITH | 7369 | 7902 | 7566 | JONES |
| 2 | SCOTT | 7788 | 7566 | 7566 | JONES |
| 3 | ADAMS | 7876 | 7788 | 7566 | JONES |
| 1 | KING | 7839 | 7839 | 7839 | KING |
| 2 | BLAKE | 7698 | 7839 | 7839 | KING |
| 3 | ALLEN | 7499 | 7698 | 7839 | KING |
| 3 | JAMES | 7900 | 7698 | 7839 | KING |
| 3 | MARTIN | 7654 | 7698 | 7839 | KING |
| 3 | TURNER | 7844 | 7698 | 7839 | KING |
| 3 | WARD | 7521 | 7698 | 7839 | KING |
| 2 | CLARK | 7782 | 7839 | 7839 | KING |
| 3 | MILLER | 7934 | 7782 | 7839 | KING |
| 2 | JONES | 7566 | 7839 | 7839 | KING |
| 3 | FORD | 7902 | 7566 | 7839 | KING |
| 4 | SMITH | 7369 | 7902 | 7839 | KING |
| 3 | SCOTT | 7788 | 7566 | 7839 | KING |
| 4 | ADAMS | 7876 | 7788 | 7839 | KING |
| 1 | MARTIN | 7654 | 7698 | 7654 | MARTIN |
| 1 | MILLER | 7934 | 7782 | 7934 | MILLER |
| 1 | SCOTT | 7788 | 7566 | 7788 | SCOTT |
| 2 | ADAMS | 7876 | 7788 | 7788 | SCOTT |
| 1 | SMITH | 7369 | 7902 | 7369 | SMITH |
| 1 | TURNER | 7844 | 7698 | 7844 | TURNER |
| 1 | WARD | 7521 | 7698 | 7521 | WARD |

(39 rows)

The following illustrates the unary operator effect of `CONNECT_BY_ROOT`. As shown in this example, when applied to an expression that is not enclosed in parentheses, the `CONNECT BY ROOT` operator affects only the term, `ename`, immediately following it. The subsequent concatenation of `|| 'manages' || ename` is not part of the `CONNECT_BY_ROOT` operation, hence the second occurrence of `ename` results in the value of the currently processed row while the first occurrence of `ename` results in the value from the root node.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr,
CONNECT_BY_ROOT ename || ' manages ' || ename "top mgr/employee"
FROM emp
START WITH ename IN ('BLAKE','CLARK','JONES')
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

The following is the output from the query. Note the values produced under the `top mgr/employee` column.

| level | employee | empno | mgr | top mgr/employee |
|-------|----------|-------|------|----------------------|
| 1 | BLAKE | 7698 | 7839 | BLAKE manages BLAKE |
| 2 | ALLEN | 7499 | 7698 | BLAKE manages ALLEN |
| 2 | JAMES | 7900 | 7698 | BLAKE manages JAMES |
| 2 | MARTIN | 7654 | 7698 | BLAKE manages MARTIN |
| 2 | TURNER | 7844 | 7698 | BLAKE manages TURNER |
| 2 | WARD | 7521 | 7698 | BLAKE manages WARD |
| 1 | CLARK | 7782 | 7839 | CLARK manages CLARK |
| 2 | MILLER | 7934 | 7782 | CLARK manages MILLER |
| 1 | JONES | 7566 | 7839 | JONES manages JONES |
| 2 | FORD | 7902 | 7566 | JONES manages FORD |
| 3 | SMITH | 7369 | 7902 | JONES manages SMITH |
| 2 | SCOTT | 7788 | 7566 | JONES manages SCOTT |
| 3 | ADAMS | 7876 | 7788 | JONES manages ADAMS |

(13 rows)

The following example uses the `CONNECT_BY_ROOT` operator on an expression enclosed in parentheses.

```
SELECT LEVEL, LPAD (' ', 2 * (LEVEL - 1)) || ename "employee", empno, mgr,
CONNECT_BY_ROOT ('Manager ' || ename || ' is emp # ' || empno)
"top mgr/empno"
FROM emp
START WITH ename IN ('BLAKE','CLARK','JONES')
CONNECT BY PRIOR empno = mgr
ORDER SIBLINGS BY ename ASC;
```

The following is the output of the query. Note that the values of both `ename` and `empno` are affected by the `CONNECT_BY_ROOT` operator and as a result, return the values from the root node as shown under the `top mgr/empno` column.

| level | employee | empno | mgr | top mgr/empno |
|-------|----------|-------|------|-----------------------------|
| 1 | BLAKE | 7698 | 7839 | Manager BLAKE is emp # 7698 |
| 2 | ALLEN | 7499 | 7698 | Manager BLAKE is emp # 7698 |
| 2 | JAMES | 7900 | 7698 | Manager BLAKE is emp # 7698 |
| 2 | MARTIN | 7654 | 7698 | Manager BLAKE is emp # 7698 |
| 2 | TURNER | 7844 | 7698 | Manager BLAKE is emp # 7698 |
| 2 | WARD | 7521 | 7698 | Manager BLAKE is emp # 7698 |
| 1 | CLARK | 7782 | 7839 | Manager CLARK is emp # 7782 |
| 2 | MILLER | 7934 | 7782 | Manager CLARK is emp # 7782 |
| 1 | JONES | 7566 | 7839 | Manager JONES is emp # 7566 |
| 2 | FORD | 7902 | 7566 | Manager JONES is emp # 7566 |
| 3 | SMITH | 7369 | 7902 | Manager JONES is emp # 7566 |
| 2 | SCOTT | 7788 | 7566 | Manager JONES is emp # 7566 |
| 3 | ADAMS | 7876 | 7788 | Manager JONES is emp # 7566 |

(13 rows)

7.3.5.7 Retrieving a Path with `SYS_CONNECT_BY_PATH`

`SYS_CONNECT_BY_PATH` is a function that works within a hierarchical query to retrieve the column values of a specified column that occur between the current node and the root node. The signature of the function is:

`SYS_CONNECT_BY_PATH (<column>, <delimiter>)`

The function takes two arguments:

`column` is the name of a column that resides within a table specified in the hierarchical query that is calling the function.

`delimiter` is the `varchar` value that separates each entry in the specified column.

The following example returns a list of employee names, and their managers; if the manager has a manager, that name is appended to the result:

```
edb=# SELECT level, ename , SYS_CONNECT_BY_PATH(ename, '/') managers
      FROM emp
```

```
CONNECT BY PRIOR empno = mgr
START WITH mgr IS NULL
ORDER BY level, ename, managers;
```

| level | ename | managers |
|-------|--------|-------------------------|
| 1 | KING | /KING |
| 2 | BLAKE | /KING/BLAKE |
| 2 | CLARK | /KING/CLARK |
| 2 | JONES | /KING/JONES |
| 3 | ALLEN | /KING/BLAKE/ALLEN |
| 3 | FORD | /KING/JONES/FORD |
| 3 | JAMES | /KING/BLAKE/JAMES |
| 3 | MARTIN | /KING/BLAKE/MARTIN |
| 3 | MILLER | /KING/CLARK/MILLER |
| 3 | SCOTT | /KING/JONES/SCOTT |
| 3 | TURNER | /KING/BLAKE/TURNER |
| 3 | WARD | /KING/BLAKE/WARD |
| 4 | ADAMS | /KING/JONES/SCOTT/ADAMS |
| 4 | SMITH | /KING/JONES/FORD/SMITH |

(14 rows)

Within the result set:

- The `level` column displays the number of levels that the query returned.
- The `ename` column displays the employee name.
- The `managers` column contains the hierarchical list of managers.

The Advanced Server implementation of `SYS_CONNECT_BY_PATH` does not support use of:

- `SYS_CONNECT_BY_PATH` inside `CONNECT_BY_PATH`
- `SYS_CONNECT_BY_PATH` inside `SYS_CONNECT_BY_PATH`

7.3.6 Multidimensional Analysis

Multidimensional analysis refers to the process commonly used in data warehousing applications of examining data using various combinations of dimensions. *Dimensions* are categories used to classify data such as time, geography, a company's departments, product lines, and so forth. The results associated with a particular set of dimensions are called *facts*. Facts are typically figures associated with product sales, profits, volumes, counts, etc.

In order to obtain these facts according to a set of dimensions in a relational database system, SQL aggregation is typically used. *SQL aggregation* basically means data is grouped according to certain criteria (dimensions) and the result set consists of aggregates of facts such as counts, sums, and averages of the data in each group.

The `GROUP BY` clause of the SQL `SELECT` command supports the following extensions that simplify the process of producing aggregate results.

- `ROLLUP` extension
- `CUBE` extension
- `GROUPING SETS` extension

In addition, the `GROUPING` function and the `GROUPING_ID` function can be used in the `SELECT` list or the `HAVING` clause to aid with the interpretation of the results when these extensions are used.

Note

The sample `dept` and `emp` tables are used extensively in this discussion to provide usage examples. The following changes were applied to these tables to provide more informative results.

```
UPDATE dept SET loc = 'BOSTON' WHERE deptno = 20;
INSERT INTO emp (empno,ename,job,deptno) VALUES (9001,'SMITH','CLERK',40);
INSERT INTO emp (empno,ename,job,deptno) VALUES (9002,'JONES','ANALYST',40);
INSERT INTO emp (empno,ename,job,deptno) VALUES (9003,'ROGERS','MANAGER',40);
```

The following rows from a join of the `emp` and `dept` tables are used:

```
SELECT loc, dname, job, empno FROM emp e, dept d
WHERE e.deptno = d.deptno
ORDER BY 1, 2, 3, 4;
```

| loc | dname | job | empno |
|----------|------------|-----------|-------|
| BOSTON | OPERATIONS | ANALYST | 9002 |
| BOSTON | OPERATIONS | CLERK | 9001 |
| BOSTON | OPERATIONS | MANAGER | 9003 |
| BOSTON | RESEARCH | ANALYST | 7788 |
| BOSTON | RESEARCH | ANALYST | 7902 |
| BOSTON | RESEARCH | CLERK | 7369 |
| BOSTON | RESEARCH | CLERK | 7876 |
| BOSTON | RESEARCH | MANAGER | 7566 |
| CHICAGO | SALES | CLERK | 7900 |
| CHICAGO | SALES | MANAGER | 7698 |
| CHICAGO | SALES | SALESMAN | 7499 |
| CHICAGO | SALES | SALESMAN | 7521 |
| CHICAGO | SALES | SALESMAN | 7654 |
| CHICAGO | SALES | SALESMAN | 7844 |
| NEW YORK | ACCOUNTING | CLERK | 7934 |
| NEW YORK | ACCOUNTING | MANAGER | 7782 |
| NEW YORK | ACCOUNTING | PRESIDENT | 7839 |

(17 rows)

The `loc`, `dname`, and `job` columns are used for the dimensions of the SQL aggregations used in the examples. The resulting facts of the aggregations are the number of employees obtained by using the `COUNT(*)` function.

A basic query grouping the `loc`, `dname`, and `job` columns is given by the following.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY loc, dname, job
ORDER BY 1, 2, 3;
```

The rows of this result set using the basic `GROUP BY` clause without extensions are referred to as the *base aggregate* rows.

| loc | dname | job | employees |
|--------|------------|---------|-----------|
| BOSTON | OPERATIONS | ANALYST | 1 |
| BOSTON | OPERATIONS | CLERK | 1 |
| BOSTON | OPERATIONS | MANAGER | 1 |
| BOSTON | RESEARCH | ANALYST | 2 |
| BOSTON | RESEARCH | CLERK | 2 |

```
BOSTON | RESEARCH | MANAGER |    1
CHICAGO | SALES    | CLERK   |    1
CHICAGO | SALES    | MANAGER |    1
CHICAGO | SALES    | SALESMAN |    4
NEW YORK | ACCOUNTING | CLERK |    1
NEW YORK | ACCOUNTING | MANAGER |    1
NEW YORK | ACCOUNTING | PRESIDENT |    1
(12 rows)
```

The `ROLLUP` and `CUBE` extensions add to the base aggregate rows by providing additional levels of subtotals to the result set.

The `GROUPING SETS` extension provides the ability to combine different types of groupings into a single result set.

The `GROUPING` and `GROUPING_ID` functions aid in the interpretation of the result set.

The additions provided by these extensions are discussed in more detail in the subsequent sections.

7.3.6.1 ROLLUP Extension

The `ROLLUP` extension produces a hierarchical set of groups with subtotals for each hierarchical group as well as a grand total. The order of the hierarchy is determined by the order of the expressions given in the `ROLLUP` expression list. The top of the hierarchy is the leftmost item in the list. Each successive item proceeding to the right moves down the hierarchy with the rightmost item being the lowest level.

The syntax for a single `ROLLUP` is as follows:

```
ROLLUP ( { <expr_1> | ( <expr_1a> [, <expr_1b>] ... ) }
          [, <expr_2> | ( <expr_2a> [, <expr_2b>] ... )] ... )
```

Each `expr` is an expression that determines the grouping of the result set. If enclosed within parenthesis as `(expr_1a, expr_1b, ...)` then the combination of values returned by `expr_1a` and `expr_1b` defines a single grouping level of the hierarchy.

The base level of aggregates returned in the result set is for each unique combination of values returned by the expression list.

In addition, a subtotal is returned for the first item in the list (`expr_1` or the combination of `(expr_1a, expr_1b, ...)`, whichever is specified) for each unique value. A subtotal is returned for the second item in the list (`expr_2` or the combination of `(expr_2a, expr_2b, ...)`, whichever is specified) for each unique value, within each grouping of the first item and so on. Finally a grand total is returned for the entire result set.

For the subtotal rows, null is returned for the items across which the subtotal is taken.

The `ROLLUP` extension specified within the context of the `GROUP BY` clause is shown by the following:

```
SELECT <select_list> FROM ...
GROUP BY [...] ROLLUP ( <expression_list> ) [, ...]
```

The items specified in `select_list` must also appear in the `ROLLUP expression_list`; or they must be aggregate functions such as `COUNT`, `SUM`, `AVG`, `MIN`, or `MAX`; or they must be constants or functions whose return values are independent of the individual rows in the group (for example, the `SYSDATE` function).

The `GROUP BY` clause may specify multiple `ROLLUP` extensions as well as multiple occurrences of other `GROUP BY` extensions and individual expressions.

The `ORDER BY` clause should be used if you want the output to display in a hierarchical or other meaningful structure. There is no guarantee on the order of the result set if no `ORDER BY` clause is specified.

The number of grouping levels or totals is $n + 1$ where n represents the number of items in the `ROLLUP` expression list. A parenthesized list counts as one item.

The following query produces a rollup based on a hierarchy of columns `loc, dname`, then `job`.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP (loc, dname, job)
ORDER BY 1, 2, 3;
```

The following is the result of the query. There is a count of the number of employees for each unique combination of `loc, dname`, and `job`, as well as subtotals for each unique combination of `loc` and `dname`, for each unique value of `loc`, and a grand total displayed on the last line.

| loc | dname | job | employees |
|-----------|------------|-----------|-----------|
| BOSTON | OPERATIONS | ANALYST | 1 |
| BOSTON | OPERATIONS | CLERK | 1 |
| BOSTON | OPERATIONS | MANAGER | 1 |
| BOSTON | OPERATIONS | | 3 |
| BOSTON | RESEARCH | ANALYST | 2 |
| BOSTON | RESEARCH | CLERK | 2 |
| BOSTON | RESEARCH | MANAGER | 1 |
| BOSTON | RESEARCH | | 5 |
| BOSTON | | | 8 |
| CHICAGO | SALES | CLERK | 1 |
| CHICAGO | SALES | MANAGER | 1 |
| CHICAGO | SALES | SALESMAN | 4 |
| CHICAGO | SALES | | 6 |
| CHICAGO | | | 6 |
| NEW YORK | ACCOUNTING | CLERK | 1 |
| NEW YORK | ACCOUNTING | MANAGER | 1 |
| NEW YORK | ACCOUNTING | PRESIDENT | 1 |
| NEW YORK | ACCOUNTING | | 3 |
| NEW YORK | | | 3 |
| | | | 17 |
| (20 rows) | | | |

The following query shows the effect of combining items in the `ROLLUP` list within parenthesis.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP (loc, (dname, job))
ORDER BY 1, 2, 3;
```

In the output, note that there are no subtotals for `loc` and `dname` combinations as in the prior example.

| loc | dname | job | employees |
|--------|------------|---------|-----------|
| BOSTON | OPERATIONS | ANALYST | 1 |
| BOSTON | OPERATIONS | CLERK | 1 |

| loc | dname | job | | employees |
|----------|------------|-----------|--|-----------|
| BOSTON | OPERATIONS | MANAGER | | 1 |
| BOSTON | RESEARCH | ANALYST | | 2 |
| BOSTON | RESEARCH | CLERK | | 2 |
| BOSTON | RESEARCH | MANAGER | | 1 |
| BOSTON | | | | 8 |
| CHICAGO | SALES | CLERK | | 1 |
| CHICAGO | SALES | MANAGER | | 1 |
| CHICAGO | SALES | SALESMAN | | 4 |
| CHICAGO | | | | 6 |
| NEW YORK | ACCOUNTING | CLERK | | 1 |
| NEW YORK | ACCOUNTING | MANAGER | | 1 |
| NEW YORK | ACCOUNTING | PRESIDENT | | 1 |
| NEW YORK | | | | 3 |
| | | | | 17 |

(16 rows)

If the first two columns in the `ROLLUP` list are enclosed in parenthesis, the subtotal levels differ as well.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP ((loc, dname), job)
ORDER BY 1, 2, 3;
```

Now there is a subtotal for each unique `loc` and `dname` combination, but none for unique values of `loc`.

| loc | dname | job | | employees |
|----------|------------|-----------|--|-----------|
| BOSTON | OPERATIONS | ANALYST | | 1 |
| BOSTON | OPERATIONS | CLERK | | 1 |
| BOSTON | OPERATIONS | MANAGER | | 1 |
| BOSTON | OPERATIONS | | | 3 |
| BOSTON | RESEARCH | ANALYST | | 2 |
| BOSTON | RESEARCH | CLERK | | 2 |
| BOSTON | RESEARCH | MANAGER | | 1 |
| BOSTON | RESEARCH | | | 5 |
| CHICAGO | SALES | CLERK | | 1 |
| CHICAGO | SALES | MANAGER | | 1 |
| CHICAGO | SALES | SALESMAN | | 4 |
| CHICAGO | SALES | | | 6 |
| NEW YORK | ACCOUNTING | CLERK | | 1 |
| NEW YORK | ACCOUNTING | MANAGER | | 1 |
| NEW YORK | ACCOUNTING | PRESIDENT | | 1 |
| NEW YORK | ACCOUNTING | | | 3 |
| | | | | 17 |

(17 rows)

7.3.6.2 CUBE Extension

The `CUBE` extension is similar to the `ROLLUP` extension. However, unlike `ROLLUP`, which produces groupings and results in a hierarchy based on a left to right listing of items in the `ROLLUP` expression list, a `CUBE` produces groupings and subtotals based on every permutation of all items in the `CUBE` expression list. Thus, the

result set contains more rows than a `ROLLUP` performed on the same expression list.

The syntax for a single `CUBE` is as follows:

```
CUBE ( { <expr_1> | ( <expr_1a> [, <expr_1b> ] ... ) }
      [, <expr_2> | ( <expr_2a> [, <expr_2b> ] ... ) ] ... )
```

Each `expr` is an expression that determines the grouping of the result set. If enclosed within parenthesis as `(expr_1a, expr_1b, ...)` then the combination of values returned by `expr_1a` and `expr_1b` defines a single group.

The base level of aggregates returned in the result set is for each unique combination of values returned by the expression list.

In addition, a subtotal is returned for the first item in the list (`expr_1` or the combination of `(expr_1a, expr_1b, ...)`, whichever is specified) for each unique value. A subtotal is returned for the second item in the list (`expr_2` or the combination of `(expr_2a, expr_2b, ...)`, whichever is specified) for each unique value. A subtotal is also returned for each unique combination of the first item and the second item. Similarly, if there is a third item, a subtotal is returned for each unique value of the third item, each unique value of the third item and first item combination, each unique value of the third item and second item combination, and each unique value of the third item, second item, and first item combination. Finally a grand total is returned for the entire result set.

For the subtotal rows, null is returned for the items across which the subtotal is taken.

The `CUBE` extension specified within the context of the `GROUP BY` clause is shown by the following:

```
SELECT <select_list> FROM ...
GROUP BY [... ,] CUBE ( <expression_list> ) [, ...]
```

The items specified in `select_list` must also appear in the `CUBE expression_list`; or they must be aggregate functions such as `COUNT`, `SUM`, `AVG`, `MIN`, or `MAX`; or they must be constants or functions whose return values are independent of the individual rows in the group (for example, the `SYSDATE` function).

The `GROUP BY` clause may specify multiple `CUBE` extensions as well as multiple occurrences of other `GROUP BY` extensions and individual expressions.

The `ORDER BY` clause should be used if you want the output to display in a meaningful structure. There is no guarantee on the order of the result set if no `ORDER BY` clause is specified.

The number of grouping levels or totals is `2` raised to the power of `n` where `n` represents the number of items in the `CUBE` expression list. A parenthesized list counts as one item.

The following query produces a cube based on permutations of columns `loc`, `dname`, and `job`.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY CUBE (loc, dname, job)
ORDER BY 1, 2, 3;
```

The following is the result of the query. There is a count of the number of employees for each combination of `loc`, `dname`, and `job`, as well as subtotals for each combination of `loc` and `dname`, for each combination of `loc` and `job`, for each combination of `dname` and `job`, for each unique value of `loc`, for each unique value of `dname`, for each unique value of `job`, and a grand total displayed on the last line.

| loc | dname | job | employees |
|--------|------------|---------|-----------|
| BOSTON | OPERATIONS | ANALYST | 1 |
| BOSTON | OPERATIONS | CLERK | 1 |
| BOSTON | OPERATIONS | MANAGER | 1 |
| BOSTON | OPERATIONS | | 3 |
| BOSTON | RESEARCH | ANALYST | 2 |

| | | | |
|----------|------------|-----------|----|
| BOSTON | RESEARCH | CLERK | 2 |
| BOSTON | RESEARCH | MANAGER | 1 |
| BOSTON | RESEARCH | | 5 |
| BOSTON | | ANALYST | 3 |
| BOSTON | | CLERK | 3 |
| BOSTON | | MANAGER | 2 |
| BOSTON | | | 8 |
| CHICAGO | SALES | CLERK | 1 |
| CHICAGO | SALES | MANAGER | 1 |
| CHICAGO | SALES | SALESMAN | 4 |
| CHICAGO | SALES | | 6 |
| CHICAGO | | CLERK | 1 |
| CHICAGO | | MANAGER | 1 |
| CHICAGO | | SALESMAN | 4 |
| CHICAGO | | | 6 |
| NEW YORK | ACCOUNTING | CLERK | 1 |
| NEW YORK | ACCOUNTING | MANAGER | 1 |
| NEW YORK | ACCOUNTING | PRESIDENT | 1 |
| NEW YORK | ACCOUNTING | | 3 |
| NEW YORK | | CLERK | 1 |
| NEW YORK | | MANAGER | 1 |
| NEW YORK | | PRESIDENT | 1 |
| NEW YORK | | | 3 |
| | ACCOUNTING | CLERK | 1 |
| | ACCOUNTING | MANAGER | 1 |
| | ACCOUNTING | PRESIDENT | 1 |
| | ACCOUNTING | | 3 |
| | OPERATIONS | ANALYST | 1 |
| | OPERATIONS | CLERK | 1 |
| | OPERATIONS | MANAGER | 1 |
| | OPERATIONS | | 3 |
| | RESEARCH | ANALYST | 2 |
| | RESEARCH | CLERK | 2 |
| | RESEARCH | MANAGER | 1 |
| | RESEARCH | | 5 |
| | SALES | CLERK | 1 |
| | SALES | MANAGER | 1 |
| | SALES | SALESMAN | 4 |
| | SALES | | 6 |
| | | ANALYST | 3 |
| | | CLERK | 5 |
| | | MANAGER | 4 |
| | | PRESIDENT | 1 |
| | | SALESMAN | 4 |
| | | | 17 |

(50 rows)

The following query shows the effect of combining items in the `CUBE` list within parenthesis.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY CUBE (loc, (dname, job))
ORDER BY 1, 2, 3;
```

In the output note that there are no subtotals for permutations involving `loc` and `dname` combinations, `loc` and `job` combinations, or for `dname` by itself, or for `job` by itself.

| loc | dname | job | employees |
|----------|------------|-----------|-----------|
| BOSTON | OPERATIONS | ANALYST | 1 |
| BOSTON | OPERATIONS | CLERK | 1 |
| BOSTON | OPERATIONS | MANAGER | 1 |
| BOSTON | RESEARCH | ANALYST | 2 |
| BOSTON | RESEARCH | CLERK | 2 |
| BOSTON | RESEARCH | MANAGER | 1 |
| BOSTON | | | 8 |
| CHICAGO | SALES | CLERK | 1 |
| CHICAGO | SALES | MANAGER | 1 |
| CHICAGO | SALES | SALESMAN | 4 |
| CHICAGO | | | 6 |
| NEW YORK | ACCOUNTING | CLERK | 1 |
| NEW YORK | ACCOUNTING | MANAGER | 1 |
| NEW YORK | ACCOUNTING | PRESIDENT | 1 |
| NEW YORK | | | 3 |
| | ACCOUNTING | CLERK | 1 |
| | ACCOUNTING | MANAGER | 1 |
| | ACCOUNTING | PRESIDENT | 1 |
| | OPERATIONS | ANALYST | 1 |
| | OPERATIONS | CLERK | 1 |
| | OPERATIONS | MANAGER | 1 |
| | RESEARCH | ANALYST | 2 |
| | RESEARCH | CLERK | 2 |
| | RESEARCH | MANAGER | 1 |
| | SALES | CLERK | 1 |
| | SALES | MANAGER | 1 |
| | SALES | SALESMAN | 4 |
| | | | 17 |

(28 rows)

The following query shows another variation whereby the first expression is specified outside of the `CUBE` extension.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY loc, CUBE (dname, job)
ORDER BY 1, 2, 3;
```

In this output, the permutations are performed for `dname` and `job` within each grouping of `loc`.

| loc | dname | job | employees |
|--------|------------|---------|-----------|
| BOSTON | OPERATIONS | ANALYST | 1 |
| BOSTON | OPERATIONS | CLERK | 1 |
| BOSTON | OPERATIONS | MANAGER | 1 |
| BOSTON | OPERATIONS | | 3 |
| BOSTON | RESEARCH | ANALYST | 2 |
| BOSTON | RESEARCH | CLERK | 2 |
| BOSTON | RESEARCH | MANAGER | 1 |
| BOSTON | RESEARCH | | 5 |
| BOSTON | | ANALYST | 3 |
| BOSTON | | CLERK | 3 |
| BOSTON | | MANAGER | 2 |
| BOSTON | | | 8 |

```

CHICAGO | SALES    | CLERK      |      1
CHICAGO | SALES    | MANAGER    |      1
CHICAGO | SALES    | SALESMAN   |      4
CHICAGO | SALES    |           |      6
CHICAGO |           | CLERK      |      1
CHICAGO |           | MANAGER    |      1
CHICAGO |           | SALESMAN   |      4
CHICAGO |           |           |      6
NEW YORK | ACCOUNTING | CLERK      |      1
NEW YORK | ACCOUNTING | MANAGER    |      1
NEW YORK | ACCOUNTING | PRESIDENT |      1
NEW YORK | ACCOUNTING |           |      3
NEW YORK |           | CLERK      |      1
NEW YORK |           | MANAGER    |      1
NEW YORK |           | PRESIDENT |      1
NEW YORK |           |           |      3
(28 rows)

```

7.3.6.3 GROUPING SETS Extension

The use of the **GROUPING SETS** extension within the **GROUP BY** clause provides a means to produce one result set that is actually the concatenation of multiple results sets based upon different groupings. In other words, a **UNION ALL** operation is performed combining the result sets of multiple groupings into one result set.

Note that a **UNION ALL** operation, and therefore the **GROUPING SETS** extension, do not eliminate duplicate rows from the result sets that are being combined together.

The syntax for a single **GROUPING SETS** extension is as follows:

```

GROUPING SETS (
{ <expr_1> | ( <expr_1a> [, <expr_1b> ] ... ) |
  ROLLUP ( <expr_list> ) | CUBE ( <expr_list> )
} [, ...])

```

A **GROUPING SETS** extension can contain any combination of one or more comma-separated expressions, lists of expressions enclosed within parenthesis, **ROLLUP** extensions, and **CUBE** extensions.

The **GROUPING SETS** extension is specified within the context of the **GROUP BY** clause as shown by the following:

```

SELECT <select_list> FROM ...
GROUP BY [...] GROUPING SETS ( <expression_list> ) [, ...]

```

The items specified in **select_list** must also appear in the **GROUPING SETS expression_list**; or they must be aggregate functions such as **COUNT**, **SUM**, **AVG**, **MIN**, or **MAX**; or they must be constants or functions whose return values are independent of the individual rows in the group (for example, the **SYSDATE** function).

The **GROUP BY** clause may specify multiple **GROUPING SETS** extensions as well as multiple occurrences of other **GROUP BY** extensions and individual expressions.

The **ORDER BY** clause should be used if you want the output to display in a meaningful structure. There is no guarantee on the order of the result set if no **ORDER BY** clause is specified.

The following query produces a union of groups given by columns `loc`, `dname`, and `job`.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY GROUPING SETS (loc, dname, job)
ORDER BY 1, 2, 3;
```

The result is as follows:

| loc | dname | job | employees |
|----------|------------|-----------|-----------|
| BOSTON | | | 8 |
| CHICAGO | | | 6 |
| NEW YORK | | | 3 |
| | ACCOUNTING | | 3 |
| | OPERATIONS | | 3 |
| | RESEARCH | | 5 |
| | SALES | | 6 |
| | | ANALYST | 3 |
| | | CLERK | 5 |
| | | MANAGER | 4 |
| | | PRESIDENT | 1 |
| | | SALESMAN | 4 |

(12 rows)

This is equivalent to the following query, which employs the use of the `UNION ALL` operator.

```
SELECT loc AS "loc", NULL AS "dname", NULL AS "job", COUNT(*) AS "employees"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY loc
UNION ALL
SELECT NULL, dname, NULL, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY dname
UNION ALL
SELECT NULL, NULL, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY job
ORDER BY 1, 2, 3;
```

The output from the `UNION ALL` query is the same as the `GROUPING SETS` output.

| loc | dname | job | employees |
|----------|------------|-----------|-----------|
| BOSTON | | | 8 |
| CHICAGO | | | 6 |
| NEW YORK | | | 3 |
| | ACCOUNTING | | 3 |
| | OPERATIONS | | 3 |
| | RESEARCH | | 5 |
| | SALES | | 6 |
| | | ANALYST | 3 |
| | | CLERK | 5 |
| | | MANAGER | 4 |
| | | PRESIDENT | 1 |

| | | |
|--|----------|---|
| | SALESMAN | 4 |
|--|----------|---|

(12 rows)

The following example shows how various types of `GROUP BY` extensions can be used together within a `GROUPING SETS` expression list.

```
SELECT loc, dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY GROUPING SETS (loc, ROLLUP (dname, job), CUBE (job, loc))
ORDER BY 1, 2, 3;
```

The following is the output from this query.

| loc | dname | job | employees |
|----------|------------|-----------|-----------|
| BOSTON | | ANALYST | 3 |
| BOSTON | | CLERK | 3 |
| BOSTON | | MANAGER | 2 |
| BOSTON | | | 8 |
| BOSTON | | | 8 |
| CHICAGO | | CLERK | 1 |
| CHICAGO | | MANAGER | 1 |
| CHICAGO | | SALESMAN | 4 |
| CHICAGO | | | 6 |
| CHICAGO | | | 6 |
| NEW YORK | | CLERK | 1 |
| NEW YORK | | MANAGER | 1 |
| NEW YORK | | PRESIDENT | 1 |
| NEW YORK | | | 3 |
| NEW YORK | | | 3 |
| | ACCOUNTING | CLERK | 1 |
| | ACCOUNTING | MANAGER | 1 |
| | ACCOUNTING | PRESIDENT | 1 |
| | ACCOUNTING | | 3 |
| | OPERATIONS | ANALYST | 1 |
| | OPERATIONS | CLERK | 1 |
| | OPERATIONS | MANAGER | 1 |
| | OPERATIONS | | 3 |
| | RESEARCH | ANALYST | 2 |
| | RESEARCH | CLERK | 2 |
| | RESEARCH | MANAGER | 1 |
| | RESEARCH | | 5 |
| | SALES | CLERK | 1 |
| | SALES | MANAGER | 1 |
| | SALES | SALESMAN | 4 |
| | SALES | | 6 |
| | | ANALYST | 3 |
| | | CLERK | 5 |
| | | MANAGER | 4 |
| | | PRESIDENT | 1 |
| | | SALESMAN | 4 |
| | | | 17 |
| | | | 17 |

(38 rows)

The output is basically a concatenation of the result sets that would be produced individually from `GROUP BY loc,`

GROUP BY ROLLUP (dname, job), and **GROUP BY CUBE (job, loc)**. These individual queries are shown by the following.

```
SELECT loc, NULL AS "dname", NULL AS "job", COUNT(*) AS "employees"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY loc
ORDER BY 1;
```

The following is the result set from the **GROUP BY loc** clause.

| loc | dname | job | employees |
|----------|-------|-----|-----------|
| BOSTON | | | 8 |
| CHICAGO | | | 6 |
| NEW YORK | | | 3 |

(3 rows)

The following query uses the **GROUP BY ROLLUP (dname, job)** clause.

```
SELECT NULL AS "loc", dname, job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP (dname, job)
ORDER BY 2, 3;
```

The following is the result set from the **GROUP BY ROLLUP (dname, job)** clause.

| loc | dname | job | employees |
|-----|------------|-----------|-----------|
| | ACCOUNTING | CLERK | 1 |
| | ACCOUNTING | MANAGER | 1 |
| | ACCOUNTING | PRESIDENT | 1 |
| | ACCOUNTING | | 3 |
| | OPERATIONS | ANALYST | 1 |
| | OPERATIONS | CLERK | 1 |
| | OPERATIONS | MANAGER | 1 |
| | OPERATIONS | | 3 |
| | RESEARCH | ANALYST | 2 |
| | RESEARCH | CLERK | 2 |
| | RESEARCH | MANAGER | 1 |
| | RESEARCH | | 5 |
| | SALES | CLERK | 1 |
| | SALES | MANAGER | 1 |
| | SALES | SALESMAN | 4 |
| | SALES | | 6 |
| | | | 17 |

(17 rows)

The following query uses the **GROUP BY CUBE (job, loc)** clause.

```
SELECT loc, NULL AS "dname", job, COUNT(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY CUBE (job, loc)
ORDER BY 1, 3;
```

The following is the result set from the **GROUP BY CUBE (job, loc)** clause.

| loc | dname | job | employees |
|----------|-------|-----------|-----------|
| BOSTON | | ANALYST | 3 |
| BOSTON | | CLERK | 3 |
| BOSTON | | MANAGER | 2 |
| BOSTON | | | 8 |
| CHICAGO | | CLERK | 1 |
| CHICAGO | | MANAGER | 1 |
| CHICAGO | | SALESMAN | 4 |
| CHICAGO | | | 6 |
| NEW YORK | | CLERK | 1 |
| NEW YORK | | MANAGER | 1 |
| NEW YORK | | PRESIDENT | 1 |
| NEW YORK | | | 3 |
| | | ANALYST | 3 |
| | | CLERK | 5 |
| | | MANAGER | 4 |
| | | PRESIDENT | 1 |
| | | SALESMAN | 4 |
| | | | 17 |

(18 rows)

If the previous three queries are combined with the `UNION ALL` operator, a concatenation of the three results sets is produced.

```

SELECT loc AS "loc", NULL AS "dname", NULL AS "job", COUNT(*) AS "employees"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY loc
UNION ALL
SELECT NULL, dname, job, count(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP (dname, job)
UNION ALL
SELECT loc, NULL, job, count(*) AS "employees" FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY CUBE (job, loc)
ORDER BY 1, 2, 3;

```

The following is the output, which is the same as when the `GROUP BY GROUPING SETS (loc, ROLLUP (dname, job), CUBE (job, loc))` clause is used.

| loc | dname | job | employees |
|----------|-------|----------|-----------|
| BOSTON | | ANALYST | 3 |
| BOSTON | | CLERK | 3 |
| BOSTON | | MANAGER | 2 |
| BOSTON | | | 8 |
| BOSTON | | | 8 |
| CHICAGO | | CLERK | 1 |
| CHICAGO | | MANAGER | 1 |
| CHICAGO | | SALESMAN | 4 |
| CHICAGO | | | 6 |
| CHICAGO | | | 6 |
| NEW YORK | | CLERK | 1 |
| NEW YORK | | MANAGER | 1 |

| | | |
|------------------------|-----------|---|
| NEW YORK | PRESIDENT | 1 |
| NEW YORK | 3 | |
| NEW YORK | 3 | |
| ACCOUNTING CLERK | 1 | |
| ACCOUNTING MANAGER | 1 | |
| ACCOUNTING PRESIDENT | 1 | |
| ACCOUNTING 3 | | |
| OPERATIONS ANALYST | 1 | |
| OPERATIONS CLERK | 1 | |
| OPERATIONS MANAGER | 1 | |
| OPERATIONS 3 | | |
| RESEARCH ANALYST | 2 | |
| RESEARCH CLERK | 2 | |
| RESEARCH MANAGER | 1 | |
| RESEARCH 5 | | |
| SALES CLERK | 1 | |
| SALES MANAGER | 1 | |
| SALES SALESMAN | 4 | |
| SALES 6 | | |
| ANALYST | 3 | |
| CLERK | 5 | |
| MANAGER | 4 | |
| PRESIDENT | 1 | |
| SALESMAN | 4 | |
| 17 | | |
| 17 | | |

(38 rows)

7.3.6.4 GROUPING Function

When using the `ROLLUP`, `CUBE`, or `GROUPING SETS` extensions to the `GROUP BY` clause, it may sometimes be difficult to differentiate between the various levels of subtotals generated by the extensions as well as the base aggregate rows in the result set. The `GROUPING` function provides a means of making this distinction.

The general syntax for use of the `GROUPING` function is shown by the following.

```
SELECT [ <expr> ... ] GROUPING( <col_expr> ) [, <expr> ] ...
FROM ...
GROUP BY [ ... ]
{ ROLLUP | CUBE | GROUPING SETS } ( [ ... ] <col_expr>
[ , ... ] ) [ , ... ]
```

The `GROUPING` function takes a single parameter that must be an expression of a dimension column specified in the expression list of a `ROLLUP`, `CUBE`, or `GROUPING SETS` extension of the `GROUP BY` clause.

The return value of the `GROUPING` function is either a 0 or 1. In the result set of a query, if the column expression specified in the `GROUPING` function is null because the row represents a subtotal over multiple values of that column then the `GROUPING` function returns a value of 1. If the row returns results based on a particular value of the column specified in the `GROUPING` function, then the `GROUPING` function returns a value of 0. In the latter case, the column can be null as well as non-null, but in any case, it is for a particular value of that column, not a subtotal across multiple values.

The following query shows how the return values of the `GROUPING` function correspond to the subtotal lines.

```
SELECT loc, dname, job, COUNT(*) AS "employees",
GROUPING(loc) AS "gf_loc",
GROUPING(dname) AS "gf_dname",
GROUPING(job) AS "gf_job"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP (loc, dname, job)
ORDER BY 1, 2, 3;
```

In the three right-most columns displaying the output of the `GROUPING` functions, a value of 1 appears on a subtotal line wherever a subtotal is taken across values of the corresponding columns.

| loc | dname | job | employees | gf_loc | gf_dname | gf_job |
|----------|------------|-----------|-----------|--------|----------|--------|
| BOSTON | OPERATIONS | ANALYST | 1 | 0 | 0 | 0 |
| BOSTON | OPERATIONS | CLERK | 1 | 0 | 0 | 0 |
| BOSTON | OPERATIONS | MANAGER | 1 | 0 | 0 | 0 |
| BOSTON | OPERATIONS | | 3 | 0 | 0 | 1 |
| BOSTON | RESEARCH | ANALYST | 2 | 0 | 0 | 0 |
| BOSTON | RESEARCH | CLERK | 2 | 0 | 0 | 0 |
| BOSTON | RESEARCH | MANAGER | 1 | 0 | 0 | 0 |
| BOSTON | RESEARCH | | 5 | 0 | 0 | 1 |
| BOSTON | | | 8 | 0 | 1 | 1 |
| CHICAGO | SALES | CLERK | 1 | 0 | 0 | 0 |
| CHICAGO | SALES | MANAGER | 1 | 0 | 0 | 0 |
| CHICAGO | SALES | SALESMAN | 4 | 0 | 0 | 0 |
| CHICAGO | SALES | | 6 | 0 | 0 | 1 |
| CHICAGO | | | 6 | 0 | 1 | 1 |
| NEW YORK | ACCOUNTING | CLERK | 1 | 0 | 0 | 0 |
| NEW YORK | ACCOUNTING | MANAGER | 1 | 0 | 0 | 0 |
| NEW YORK | ACCOUNTING | PRESIDENT | 1 | 0 | 0 | 0 |
| NEW YORK | ACCOUNTING | | 3 | 0 | 0 | 1 |
| NEW YORK | | | 3 | 0 | 1 | 1 |
| | | | 17 | 1 | 1 | 1 |

(20 rows)

These indicators can be used as screening criteria for particular subtotals. For example, using the previous query, you can display only those subtotals for `loc` and `dname` combinations by using the `GROUPING` function in a `HAVING` clause.

```
SELECT loc, dname, job, COUNT(*) AS "employees",
GROUPING(loc) AS "gf_loc",
GROUPING(dname) AS "gf_dname",
GROUPING(job) AS "gf_job"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY ROLLUP (loc, dname, job)
HAVING GROUPING(loc) = 0
    AND GROUPING(dname) = 0
    AND GROUPING(job) = 1
ORDER BY 1, 2;
```

This query produces the following result:

| loc | | dname | | job | | employees | | gf_loc | | gf_dname | | gf_job |
|----------|--|------------|--|-----|--|-----------|--|--------|--|----------|--|--------|
| BOSTON | | OPERATIONS | | | | 3 | | 0 | | 0 | | 1 |
| BOSTON | | RESEARCH | | | | 5 | | 0 | | 0 | | 1 |
| CHICAGO | | SALES | | | | 6 | | 0 | | 0 | | 1 |
| NEW YORK | | ACCOUNTING | | | | 3 | | 0 | | 0 | | 1 |
| (4 rows) | | | | | | | | | | | | |

The **GROUPING** function can be used to distinguish a subtotal row from a base aggregate row or from certain subtotal rows where one of the items in the expression list returns null as a result of the column on which the expression is based being null for one or more rows in the table, as opposed to representing a subtotal over the column.

To illustrate this point, the following row is added to the **emp** table. This provides a row with a null value for the **job** column.

```
INSERT INTO emp (empno,ename,deptno) VALUES (9004,'PETERS',40);
```

The following query is issued using a reduced number of rows for clarity.

```
SELECT loc, job, COUNT(*) AS "employees",
       GROUPING(loc) AS "gf_loc",
       GROUPING(job) AS "gf_job"
  FROM emp e, dept d
 WHERE e.deptno = d.deptno AND loc = 'BOSTON'
 GROUP BY CUBE (loc, job)
 ORDER BY 1, 2;
```

Note that the output contains two rows containing **BOSTON** in the **loc** column and spaces in the **job** column (fourth and fifth entries in the table).

| loc | | job | | employees | | gf_loc | | gf_job |
|-----------|--|---------|--|-----------|--|--------|--|--------|
| BOSTON | | ANALYST | | 3 | | 0 | | 0 |
| BOSTON | | CLERK | | 3 | | 0 | | 0 |
| BOSTON | | MANAGER | | 2 | | 0 | | 0 |
| BOSTON | | | | 1 | | 0 | | 0 |
| BOSTON | | | | 9 | | 0 | | 1 |
| | | ANALYST | | 3 | | 1 | | 0 |
| | | CLERK | | 3 | | 1 | | 0 |
| | | MANAGER | | 2 | | 1 | | 0 |
| | | | | 1 | | 1 | | 0 |
| | | | | 9 | | 1 | | 1 |
| (10 rows) | | | | | | | | |

The fifth row where the **GROUPING** function on the **job** column (af job) returns 1 indicates this is a subtotal over all jobs. Note that the row contains a subtotal value of 9 in the **employees** column.

The fourth row where the **GROUPING** function on the **job** column as well as on the **loc** column returns 0 indicates this is a base aggregate of all rows where **loc** is **BOSTON** and **job** is null, which is the row inserted for this example. The **employees** column contains 1, which is the count of the single such row inserted.

Also note that in the ninth row (next to last) the **GROUPING** function on the **job** column returns 0 while the **GROUPING** function on the **loc** column returns 1 indicating this is a subtotal over all locations where the **job** column is null, which again, is a count of the single row inserted for this example.

7.3.6.5 GROUPING_ID Function

The `GROUPING_ID` function provides a simplification of the `GROUPING` function in order to determine the subtotal level of a row in the result set from a `ROLLBACK`, `CUBE`, or `GROUPING SETS` extension.

The `GROUPING` function takes only one column expression and returns an indication of whether or not a row is a subtotal over all values of the given column. Thus, multiple `GROUPING` functions may be required to interpret the level of subtotals for queries with multiple grouping columns.

The `GROUPING_ID` function accepts one or more column expressions that have been used in the `ROLLBACK`, `CUBE`, or `GROUPING SETS` extensions and returns a single integer that can be used to determine over which of these columns a subtotal has been aggregated.

The general syntax for use of the `GROUPING_ID` function is shown by the following.

```
SELECT [ <expr> ... ]
  GROUPING_ID( <col_expr_1> [, <col_expr_2> ] ... )
  [, <expr> ] ...
FROM ...
GROUP BY [...,]
{ ROLLUP | CUBE | GROUPING SETS }([...] <col_expr_1>
[, <col_expr_2>] [, ...]) [, ...]
```

The `GROUPING_ID` function takes one or more parameters that must be expressions of dimension columns specified in the expression list of a `ROLLUP`, `CUBE`, or `GROUPING SETS` extension of the `GROUP BY` clause.

The `GROUPING_ID` function returns an integer value. This value corresponds to the base-10 interpretation of a bit vector consisting of the concatenated 1's and 0's that would be returned by a series of `GROUPING` functions specified in the same left-to-right order as the ordering of the parameters specified in the `GROUPING_ID` function.

The following query shows how the returned values of the `GROUPING_ID` function represented in column `gid` correspond to the values returned by two `GROUPING` functions on columns `loc` and `dname`.

```
SELECT loc, dname, COUNT(*) AS "employees",
  GROUPING(loc) AS "gf_loc", GROUPING(dname) AS "gf_dname",
  GROUPING_ID(loc, dname) AS "gid"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY CUBE (loc, dname)
ORDER BY 6, 1, 2;
```

In the following output, note the relationship between a bit vector consisting of the `gf_loc` value and `gf_dname` value compared to the integer given in `gid`.

| loc | dname | employees | gf_loc | gf_dname | gid |
|----------|------------|-----------|--------|----------|-----|
| BOSTON | OPERATIONS | 3 | 0 | 0 | 0 |
| BOSTON | RESEARCH | 5 | 0 | 0 | 0 |
| CHICAGO | SALES | 6 | 0 | 0 | 0 |
| NEW YORK | ACCOUNTING | 3 | 0 | 0 | 0 |
| BOSTON | | 8 | 0 | 1 | 1 |
| CHICAGO | | 6 | 0 | 1 | 1 |
| NEW YORK | | 3 | 0 | 1 | 1 |
| | ACCOUNTING | 3 | 1 | 0 | 2 |
| | OPERATIONS | 3 | 1 | 0 | 2 |
| | RESEARCH | 5 | 1 | 0 | 2 |

| | | | | | |
|-----------|-------|----|---|---|---|
| | SALES | 6 | 1 | 0 | 2 |
| | | 17 | 1 | 1 | 3 |
| (12 rows) | | | | | |

The following table provides specific examples of the `GROUPING_ID` function calculations based on the `GROUPING` function return values for four rows of the output.

| loc | dname | Bit Vector gf_loc gf_dname | GROUPING_ID gid |
|--------|------------|-------------------------------|--------------------|
| BOSTON | OPERATIONS | $0 * 2^1 + 0 * 2^0$ | 0 |
| BOSTON | null | $0 * 2^1 + 1 * 2^0$ | 1 |
| null | ACCOUNTING | $1 * 2^1 + 0 * 2^0$ | 2 |
| null | null | $1 * 2^1 + 1 * 2^0$ | 3 |

The following table summarizes how the `GROUPING_ID` function return values correspond to the grouping columns over which aggregation occurs.

| Aggregation by Column | Bit Vector gf_loc gf_dname | GROUPING_ID gid |
|-----------------------|-------------------------------|--------------------|
| loc, dname | 0 0 | 0 |
| loc | 0 1 | 1 |
| dname | 1 0 | 2 |
| Grand Total | 1 1 | 3 |

So to display only those subtotals by `dname`, the following simplified query can be used with a `HAVING` clause based on the `GROUPING_ID` function.

```
SELECT loc, dname, COUNT(*) AS "employees",
GROUPING(loc) AS "gf_loc", GROUPING(dname) AS "gf_dname",
GROUPING_ID(loc, dname) AS "gid"
FROM emp e, dept d
WHERE e.deptno = d.deptno
GROUP BY CUBE (loc, dname)
HAVING GROUPING_ID(loc, dname) = 2
ORDER BY 6, 1, 2;
```

The following is the result of the query.

| loc | dname | employees | gf_loc | gf_dname | gid |
|------------|-------|-----------|--------|----------|-----|
| ACCOUNTING | 3 | 1 | 0 | 2 | |
| OPERATIONS | 3 | 1 | 0 | 2 | |
| RESEARCH | 5 | 1 | 0 | 2 | |
| SALES | 6 | 1 | 0 | 2 | |

(4 rows)

7.4 Profile Management

Advanced Server allows a database superuser to create named *profiles*. Each profile defines rules for password management that augment `password` and `md5` authentication. The rules in a profile can:

- count failed login attempts
- lock an account due to excessive failed login attempts
- mark a password for expiration
- define a grace period after a password expiration
- define rules for password complexity
- define rules that limit password re-use

A profile is a named set of password attributes that allow you to easily manage a group of roles that share comparable authentication requirements. If the password requirements change, you can modify the profile to have the new requirements applied to each user that is associated with that profile.

After creating the profile, you can associate the profile with one or more users. When a user connects to the server, the server enforces the profile that is associated with their login role. Profiles are shared by all databases within a cluster, but each cluster may have multiple profiles. A single user with access to multiple databases will use the same profile when connecting to each database within the cluster.

Advanced Server creates a profile named `default` that is associated with a new role when the role is created unless an alternate profile is specified. If you upgrade to Advanced Server from a previous server version, existing roles will automatically be assigned to the `default` profile. You cannot delete the `default` profile.

The `default` profile specifies the following attributes:

| | |
|--------------------------|-----------|
| FAILED_LOGIN_ATTEMPTS | UNLIMITED |
| PASSWORD_LOCK_TIME | UNLIMITED |
| PASSWORD_LIFE_TIME | UNLIMITED |
| PASSWORD_GRACE_TIME | UNLIMITED |
| PASSWORD_REUSE_TIME | UNLIMITED |
| PASSWORD_REUSE_MAX | UNLIMITED |
| PASSWORD_VERIFY_FUNCTION | NULL |
| PASSWORD_ALLOW_HASHED | TRUE |

A database superuser can use the `ALTER PROFILE` command to modify the values specified by the `default` profile. For more information about modifying a profile, see [Altering a Profile](#).

7.4.1 Creating a New Profile

Use the `CREATE PROFILE` command to create a new profile. The syntax is:

```
CREATE PROFILE <profile_name>
[LIMIT {<parameter value>} ...];
```

Include the `LIMIT` clause and one or more space-delimited `parameter/value` pairs to specify the rules enforced by Advanced Server.

Parameters:

`profile_name` specifies the name of the profile.

`parameter` specifies the attribute limited by the profile.

`value` specifies the parameter limit.

Advanced Server supports the `value` shown below for each `parameter`:

`FAILED_LOGIN_ATTEMPTS` specifies the number of failed login attempts that a user may make before the server locks the user out of their account for the length of time specified by `PASSWORD_LOCK_TIME`. Supported values are:

- An `INTEGER` value greater than `0`.
- `DEFAULT` - the value of `FAILED_LOGIN_ATTEMPTS` specified in the `DEFAULT` profile.
- `UNLIMITED` – the connecting user may make an unlimited number of failed login attempts.

`PASSWORD_LOCK_TIME` specifies the length of time that must pass before the server unlocks an account that has been locked because of `FAILED_LOGIN_ATTEMPTS`. Supported values are:

- A `NUMERIC` value greater than or equal to `0`. To specify a fractional portion of a day, specify a decimal value. For example, use the value `4.5` to specify `4` days, `12` hours.
- `DEFAULT` - the value of `PASSWORD_LOCK_TIME` specified in the `DEFAULT` profile.
- `UNLIMITED` – the account is locked until it is manually unlocked by a database superuser.

`PASSWORD_LIFE_TIME` specifies the number of days that the current password may be used before the user is prompted to provide a new password. Include the `PASSWORD_GRACE_TIME` clause when using the `PASSWORD_LIFE_TIME` clause to specify the number of days that will pass after the password expires before connections by the role are rejected. If `PASSWORD_GRACE_TIME` is not specified, the password will expire on the day specified by the default value of `PASSWORD_GRACE_TIME`, and the user will not be allowed to execute any command until a new password is provided. Supported values are:

- A `NUMERIC` value greater than or equal to `0`. To specify a fractional portion of a day, specify a decimal value. For example, use the value `4.5` to specify `4` days, `12` hours.
- `DEFAULT` - the value of `PASSWORD_LIFE_TIME` specified in the `DEFAULT` profile.
- `UNLIMITED` – The password does not have an expiration date.

`PASSWORD_GRACE_TIME` specifies the length of the grace period after a password expires until the user is forced to change their password. When the grace period expires, a user will be allowed to connect, but will not be allowed to execute any command until they update their expired password. Supported values are:

- A `NUMERIC` value greater than or equal to `0`. To specify a fractional portion of a day, specify a decimal value. For example, use the value `4.5` to specify `4` days, `12` hours.
- `DEFAULT` - the value of `PASSWORD_GRACE_TIME` specified in the `DEFAULT` profile.
- `UNLIMITED` – The grace period is infinite.

`PASSWORD_REUSE_TIME` specifies the number of days a user must wait before re-using a password. The `PASSWORD_REUSE_TIME` and `PASSWORD_REUSE_MAX` parameters are intended to be used together. If you specify a finite value for one of these parameters while the other is `UNLIMITED`, old passwords can

never be reused. If both parameters are set to **UNLIMITED** there are no restrictions on password reuse. Supported values are:

- A **NUMERIC** value greater than or equal to **0**. To specify a fractional portion of a day, specify a decimal value. For example, use the value **4.5** to specify **4** days, **12** hours.
- **DEFAULT** - the value of **PASSWORD_REUSE_TIME** specified in the **DEFAULT** profile.
- **UNLIMITED** – The password can be re-used without restrictions.

PASSWORD_REUSE_MAX specifies the number of password changes that must occur before a password can be reused. The **PASSWORD_REUSE_TIME** and **PASSWORD_REUSE_MAX** parameters are intended to be used together. If you specify a finite value for one of these parameters while the other is **UNLIMITED**, old passwords can never be reused. If both parameters are set to **UNLIMITED** there are no restrictions on password reuse. Supported values are:

- An **INTEGER** value greater than or equal to **0**.
- **DEFAULT** - the value of **PASSWORD_REUSE_MAX** specified in the **DEFAULT** profile.
- **UNLIMITED** – The password can be re-used without restrictions.

PASSWORD_VERIFY_FUNCTION specifies password complexity. Supported values are:

- The name of a PL/SQL function.
- **DEFAULT** - the value of **PASSWORD_VERIFY_FUNCTION** specified in the **DEFAULT** profile.
- **NULL**

PASSWORD_ALLOW_HASHED specifies whether an encrypted password to be allowed for use or not. If you specify the value as **TRUE**, the system allows a user to change the password by specifying a hash computed encrypted password on the client side. However, if you specify the value as **FALSE**, then a password must be specified in a plain-text form in order to be validated effectively, else an error will be thrown if a server receives an encrypted password. Supported values are:

- A **BOOLEAN** value **TRUE/ON/YES/1** or **FALSE/OFF/NO/0**.
- **DEFAULT** - the value of **PASSWORD_ALLOW_HASHED** specified in the **DEFAULT** profile.

Note

The **PASSWORD_ALLOW_HASHED** is not an Oracle-compatible parameter.

Notes

Use **DROP PROFILE** command to remove the profile.

Examples

The following command creates a profile named **acctg**. The profile specifies that if a user has not authenticated with the correct password in five attempts, the account will be locked for one day:

```
CREATE PROFILE acctg LIMIT
  FAILED_LOGIN_ATTEMPTS 5
  PASSWORD_LOCK_TIME 1;
```

The following command creates a profile named **sales**. The profile specifies that a user must change their password every 90 days:

```
CREATE PROFILE sales LIMIT
  PASSWORD_LIFE_TIME 90
  PASSWORD_GRACE_TIME 3;
```

If the user has not changed their password before the 90 days specified in the profile has passed, they will be issued a warning at login. After a grace period of 3 days, their account will not be allowed to invoke any commands until they change their password.

The following command creates a profile named `accts`. The profile specifies that a user cannot re-use a password within 180 days of the last use of the password, and must change their password at least 5 times before re-using the password:

```
CREATE PROFILE accts LIMIT
    PASSWORD_REUSE_TIME 180
    PASSWORD_REUSE_MAX 5;
```

The following command creates a profile named `resources`; the profile calls a user-defined function named `password_rules` that will verify that the password provided meets their standards for complexity:

```
CREATE PROFILE resources LIMIT
    PASSWORD_VERIFY_FUNCTION password_rules;
```

7.4.1.1 Creating a Password Function

When specifying `PASSWORD_VERIFY_FUNCTION`, you can provide a customized function that specifies the security rules that will be applied when your users change their password. For example, you can specify rules that stipulate that the new password must be at least n characters long, and may not contain a specific value.

The password function has the following signature:

```
<function_name> (<user_name> VARCHAR2,
    <new_password> VARCHAR2,
    <old_password> VARCHAR2) RETURN boolean
```

Where:

`user_name` is the name of the user.

`new_password` is the new password.

`old_password` is the user's previous password. If you reference this parameter within your function:

When a database superuser changes their password, the third parameter will always be `NULL`.

When a user with the `CREATEROLE` attribute changes their password, the parameter will pass the previous password if the statement includes the `REPLACE` clause. Note that the `REPLACE` clause is optional syntax for a user with the `CREATEROLE` privilege.

When a user that is not a database superuser and does not have the `CREATEROLE` attribute changes their password, the third parameter will contain the previous password for the role.

The function returns a Boolean value. If the function returns true and does not raise an exception, the password is accepted; if the function returns false or raises an exception, the password is rejected. If the function raises an exception, the specified error message is displayed to the user. If the function does not raise an exception, but returns false, the following error message is displayed:

`ERROR: password verification for the specified password failed`

The function must be owned by a database superuser, and reside in the `sys` schema.

Example:

The following example creates a profile and a custom function; then, the function is associated with the profile. The following `CREATE PROFILE` command creates a profile named `acctg_pwd_profile`:

```
CREATE PROFILE acctg_pwd_profile;
```

The following commands create a (schema-qualified) function named `verify_password`:

```
CREATE OR REPLACE FUNCTION sys.verify_password(user_name varchar2,
new_password varchar2, old_password varchar2)
RETURN boolean IMMUTABLE
IS
BEGIN
  IF (length(new_password) < 5)
  THEN
    raise_application_error(-20001, 'too short');
  END IF;

  IF substring(new_password FROM old_password) IS NOT NULL
  THEN
    raise_application_error(-20002, 'includes old password');
  END IF;

  RETURN true;
END;
```

The function first ensures that the password is at least 5 characters long, and then compares the new password to the old password. If the new password contains fewer than 5 characters, or contains the old password, the function raises an error.

The following statement sets the ownership of the `verify_password` function to the `enterprisedb` database superuser:

```
ALTER FUNCTION verify_password(varchar2, varchar2, varchar2) OWNER TO
enterprisedb;
```

Then, the `verify_password` function is associated with the profile:

```
ALTER PROFILE acctg_pwd_profile LIMIT PASSWORD_VERIFY_FUNCTION
verify_password;
```

The following statements confirm that the function is working by first creating a test user `(alice)`, and then attempting to associate invalid and valid passwords with her role:

```
CREATE ROLE alice WITH LOGIN PASSWORD 'temp_password' PROFILE
acctg_pwd_profile;
```

Then, when `alice` connects to the database and attempts to change her password, she must adhere to the rules established by the profile function. A non-superuser without `CREATEROLE` must include the `REPLACE` clause when changing a password:

```
edb=> ALTER ROLE alice PASSWORD 'hey';
ERROR: missing REPLACE clause
```

The new password must be at least 5 characters long:

```
edb=> ALTER USER alice PASSWORD 'hey' REPLACE 'temp_password';
ERROR: EDB-20001: too short
```

CONTEXT: edb-spl function verify_password(character varying,character varying,character varying) line 5 at procedure/function invocation statement

If the new password is acceptable, the command completes without error:

```
edb=> ALTER USER alice PASSWORD 'hello' REPLACE 'temp_password';
ALTER ROLE
```

If `alice` decides to change her password, the new password must not contain the old password:

```
edb=> ALTER USER alice PASSWORD 'helloworld' REPLACE 'hello';
ERROR: EDB-20002: includes old password
CONTEXT: edb-spl function verify_password(character varying,character varying,character varying) line 10 at procedure/function invocation statement
```

To remove the verify function, set `password_verify_function` to `NULL`:

```
ALTER PROFILE acctg_pwd_profile LIMIT password_verify_function NULL;
```

Then, all password constraints will be lifted:

```
edb=# ALTER ROLE alice PASSWORD 'hey';
ALTER ROLE
```

7.4.2 Altering a Profile

Use the `ALTER PROFILE` command to modify a user-defined profile; Advanced Server supports two forms of the command:

```
ALTER PROFILE <profile_name> RENAME TO <new_name>;
ALTER PROFILE <profile_name>
    LIMIT {<parameter value>}[...];
```

Include the `LIMIT` clause and one or more space-delimited `parameter/value` pairs to specify the rules enforced by Advanced Server, or use `ALTER PROFILE...RENAME TO` to change the name of a profile.

Parameters:

`profile_name` specifies the name of the profile.

`new_name` specifies the new name of the profile.

`parameter` specifies the attribute limited by the profile.

`value` specifies the parameter limit.

See the table in [Creating a New Profile](#), for a complete list of accepted parameter/value pairs.

Examples

The following example modifies a profile named `acctg_profile`:

```
ALTER PROFILE acctg_profile
    LIMIT FAILED_LOGIN_ATTEMPTS 3 PASSWORD_LOCK_TIME 1;
```

`acctg_profile` will count failed connection attempts when a login role attempts to connect to the server. The profile specifies that if a user has not authenticated with the correct password in three attempts, the account will be locked for one day.

The following example changes the name of `acctg_profile` to `payables_profile`:

```
ALTER PROFILE acctg_profile RENAME TO payables_profile;
```

7.4.3 Dropping a Profile

Use the `DROP PROFILE` command to drop a profile. The syntax is:

```
> DROP PROFILE [IF EXISTS] <profile_name> [CASCADE|RESTRICT];
```

Include the `IF EXISTS` clause to instruct the server to not throw an error if the specified profile does not exist. The server will issue a notice if the profile does not exist.

Include the optional `CASCADE` clause to reassign any users that are currently associated with the profile to the `default` profile, and then drop the profile. Include the optional `RESTRICT` clause to instruct the server to not drop any profile that is associated with a role. This is the default behavior.

Parameters

`profile_name`

The name of the profile being dropped.

Examples

The following example drops a profile named `acctg_profile`:

```
DROP PROFILE acctg_profile CASCADE;
```

The command first re-associates any roles associated with the `acctg_profile` profile with the `default` profile, and then drops the `acctg_profile` profile.

The following example drops a profile named `acctg_profile`:

```
DROP PROFILE acctg_profile RESTRICT;
```

The `RESTRICT` clause in the command instructs the server to not drop `acctg_profile` if there are any roles associated with the profile.

7.4.4 Associating a Profile with an Existing Role

After creating a profile, you can use the `ALTER USER... PROFILE` or `ALTER ROLE... PROFILE` command to associate the profile with a role. The command syntax related to profile management functionality is:

> `ALTER USER|ROLE <name> [[WITH] option[...]]`

where `option` can be the following compatible clauses:

```
PROFILE <profile_name>
| ACCOUNT {LOCK|UNLOCK}
| PASSWORD EXPIRE [AT '<timestamp>']
```

or `option` can be the following non-compatible clauses:

```
| PASSWORD SET AT '<timestamp>'
| LOCK TIME '<timestamp>'
| STORE PRIOR PASSWORD {'<password>' '<timestamp>} [, ...]
```

For information about the administrative clauses of the `ALTER USER` or `ALTER ROLE` command that are supported by Advanced Server, please see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-commands.html>

Only a database superuser can use the `ALTER USER|ROLE` clauses that enforce profile management. The clauses enforce the following behaviors:

Include the `PROFILE` clause and a `profile_name` to associate a pre-defined profile with a role, or to change which pre-defined profile is associated with a user.

Include the `ACCOUNT` clause and the `LOCK` or `UNLOCK` keyword to specify that the user account should be placed in a locked or unlocked state.

Include the `LOCK TIME 'timestamp'` clause and a date/time value to lock the role at the specified time, and unlock the role at the time indicated by the `PASSWORD_LOCK_TIME` parameter of the profile assigned to this role. If `LOCK TIME` is used with the `ACCOUNT LOCK` clause, the role can only be unlocked by a database superuser with the `ACCOUNT UNLOCK` clause.

Include the `PASSWORD EXPIRE` clause with the `AT 'timestamp'` keywords to specify a date/time when the password associated with the role will expire. If you omit the `AT 'timestamp'` keywords, the password will expire immediately.

Include the `PASSWORD SET AT 'timestamp'` keywords to set the password modification date to the time specified.

Include the `STORE PRIOR PASSWORD {'password' 'timestamp'} [, ...]` clause to modify the password history, adding the new password and the time the password was set.

Each login role may only have one profile. To discover the profile that is currently associated with a login role, query the `profile` column of the `DBA_USERS` view.

Parameters

`name`

The name of the role with which the specified profile will be associated.

`password`

The password associated with the role.

`profile_name`

The name of the profile that will be associated with the role.

timestamp

The date and time at which the clause will be enforced. When specifying a value for `timestamp`, enclose the value in single-quotes.

Examples

The following command uses the `ALTER USER... PROFILE` command to associate a profile named `acctg` with a user named `john`:

```
ALTER USER john PROFILE acctg_profile;
```

The following command uses the `ALTER ROLE... PROFILE` command to associate a profile named `acctg` with a user named `john`:

```
ALTER ROLE john PROFILE acctg_profile;
```

7.4.5 Unlocking a Locked Account

A database superuser can use clauses of the `ALTER USER|ROLE...` command to lock or unlock a role. The syntax is:

```
ALTER USER|ROLE <name>
  ACCOUNT {LOCK|UNLOCK}
  LOCK TIME '<timestamp>'
```

Include the `ACCOUNT LOCK` clause to lock a role immediately; when locked, a role's `LOGIN` functionality is disabled. When you specify the `ACCOUNT LOCK` clause without the `LOCK TIME` clause, the state of the role will not change until a superuser uses the `ACCOUNT UNLOCK` clause to unlock the role.

Use the `ACCOUNT UNLOCK` clause to unlock a role.

Use the `LOCK TIME 'timestamp'` clause to instruct the server to lock the account at the time specified by the given timestamp for the length of time specified by the `PASSWORD_LOCK_TIME` parameter of the profile associated with this role.

Combine the `LOCK TIME 'timestamp'` clause and the `ACCOUNT LOCK` clause to lock an account at a specified time until the account is unlocked by a superuser invoking the `ACCOUNT UNLOCK` clause.

Parameters

name

The name of the role that is being locked or unlocked.

timestamp

The date and time at which the role will be locked. When specifying a value for `timestamp`, enclose the value in single-quotes.

Note

This command (available only in Advanced Server) is implemented to support Oracle-styled profile management.

Examples

The following example uses the `ACCOUNT LOCK` clause to lock the role named `john`. The account will remain locked until the account is unlocked with the `ACCOUNT UNLOCK` clause:

```
ALTER ROLE john ACCOUNT LOCK;
```

The following example uses the `ACCOUNT UNLOCK` clause to unlock the role named `john`:

```
ALTER USER john ACCOUNT UNLOCK;
```

The following example uses the `LOCK TIME 'timestamp'` clause to lock the role named `john` on September 4, 2015:

```
ALTER ROLE john LOCK TIME 'September 4 12:00:00 2015';
```

The role will remain locked for the length of time specified by the `PASSWORD_LOCK_TIME` parameter.

The following example combines the `LOCK TIME 'timestamp'` clause and the `ACCOUNT LOCK` clause to lock the role named `john` on September 4, 2015:

```
ALTER ROLE john LOCK TIME 'September 4 12:00:00 2015' ACCOUNT LOCK;
```

The role will remain locked until a database superuser uses the `ACCOUNT UNLOCK` command to unlock the role.

7.4.6 Creating a New Role Associated with a Profile

A database superuser can use clauses of the `CREATE USER|ROLE` command to assign a named profile to a role when creating the role, or to specify profile management details for a role. The command syntax related to profile management functionality is:

```
> CREATE USER|ROLE <name> [[WITH] <option> [...]]
```

where `option` can be the following compatible clauses:

```
PROFILE <profile_name>
| ACCOUNT {LOCK|UNLOCK}
| PASSWORD EXPIRE [AT '<timestamp>']
```

or `option` can be the following non-compatible clauses:

```
| LOCK TIME '<timestamp>'
```

For information about the administrative clauses of the `CREATE USER` or `CREATE ROLE` command that are supported by Advanced Server, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-commands.html>

`CREATE ROLE|USER... PROFILE` adds a new role with an associated profile to an Advanced Server database cluster.

Roles created with the `CREATE USER` command are (by default) login roles. Roles created with the `CREATE ROLE` command are (by default) not login roles. To create a login account with the `CREATE ROLE` command,

you must include the `LOGIN` keyword.

Only a database superuser can use the `CREATE USER|ROLE` clauses that enforce profile management; these clauses enforce the following behaviors:

Include the `PROFILE` clause and a `profile_name` to associate a pre-defined profile with a role, or to change which pre-defined profile is associated with a user.

Include the `ACCOUNT` clause and the `LOCK` or `UNLOCK` keyword to specify that the user account should be placed in a locked or unlocked state.

Include the `LOCK TIME 'timestamp'` clause and a date/time value to lock the role at the specified time, and unlock the role at the time indicated by the `PASSWORD LOCK_TIME` parameter of the profile assigned to this role. If `LOCK TIME` is used with the `ACCOUNT LOCK` clause, the role can only be unlocked by a database superuser with the `ACCOUNT UNLOCK` clause.

Include the `PASSWORD EXPIRE` clause with the optional `AT 'timestamp'` keywords to specify a date/time when the password associated with the role will expire. If you omit the `AT 'timestamp'` keywords, the password will expire immediately.

Each login role may only have one profile. To discover the profile that is currently associated with a login role, query the `profile` column of the `DBA_USERS` view.

Parameters

`name`

The name of the role.

`profile_name`

The name of the profile associated with the role.

`timestamp`

The date and time at which the clause will be enforced. When specifying a value for `timestamp`, enclose the value in single-quotes.

Examples

The following example uses `CREATE USER` to create a login role named `john` who is associated with the `acctg_profile` profile:

```
CREATE USER john PROFILE acctg_profile IDENTIFIED BY "1safepwd";
```

`john` can log in to the server, using the password `1safepwd`.

The following example uses `CREATE ROLE` to create a login role named `john` who is associated with the `acctg_profile` profile:

```
CREATE ROLE john PROFILE acctg_profile LOGIN PASSWORD "1safepwd";
```

`john` can log in to the server, using the password `1safepwd`.

7.4.7 Backing up Profile Management Functions

A profile may include a `PASSWORD_VERIFY_FUNCTION` clause that refers to a user-defined function that specifies the behavior enforced by Advanced Server. Profiles are global objects; they are shared by all of the databases within a cluster. While profiles are global objects, user-defined functions are database objects.

Invoking `pg_dumpall` with the `-g` or `-r` option will create a script that recreates the definition of any existing profiles, but that does not recreate the user-defined functions that are referred to by the `PASSWORD_VERIFY_FUNCTION` clause. You should use the `pg_dump` utility to explicitly dump (and later restore) the database in which those functions reside.

The script created by `pg_dump` will contain a command that includes the clause and function name:

```
ALTER PROFILE... LIMIT PASSWORD_VERIFY_FUNCTION <function_name>
```

to associate the restored function with the profile with which it was previously associated.

If the `PASSWORD_VERIFY_FUNCTION` clause is set to `DEFAULT` or `NULL`, the behavior will be replicated by the script generated by the `pg_dumpall -g` or `pg_dumpall -r` command.

7.5 Optimizer Hints

When you invoke a `DELETE`, `INSERT`, `SELECT` or `UPDATE` command, the server generates a set of execution plans; after analyzing those execution plans, the server selects a plan that will (generally) return the result set in the least amount of time. The server's choice of plan is dependent upon several factors:

- The estimated execution cost of data handling operations.
- Parameter values assigned to parameters in the `Query Tuning` section of the `postgresql.conf` file.
- Column statistics that have been gathered by the `ANALYZE` command.

As a rule, the query planner will select the least expensive plan. You can use an *optimizer hint* to influence the server as it selects a query plan. An optimizer hint is a directive (or multiple directives) embedded in a comment-like syntax that immediately follows a `DELETE`, `INSERT`, `SELECT` or `UPDATE` command. Keywords in the comment instruct the server to employ or avoid a specific plan when producing the result set.

Synopsis

```
{ DELETE | INSERT | SELECT | UPDATE } /*+ { <hint> [ <comment> ] } [...] */
<statement_body>
```

```
{ DELETE | INSERT | SELECT | UPDATE } --+ { <hint> [ <comment> ] } [...]
<statement_body>
```

Optimizer hints may be included in either of the forms shown above. Note that in both forms, a plus sign (+) must immediately follow the `/*` or `--` opening comment symbols, with no intervening space, or the server will not interpret the following tokens as hints.

If you are using the first form, the hint and optional comment may span multiple lines. The second form requires all hints and comments to occupy a single line; the remainder of the statement must start on a new line.

Description

Please Note:

- The database server will always try to use the specified hints if at all possible.
- If a planner method parameter is set so as to disable a certain plan type, then this plan will not be used even if it is specified in a hint, unless there are no other possible options for the planner. Examples of planner method

parameters are `enable_indexscan`, `enable_seqscan`, `enable_hashjoin`, `enable_mergejoin`, and `enable_nestloop`. These are all Boolean parameters.

- Remember that the hint is embedded within a comment. As a consequence, if the hint is misspelled or if any parameter to a hint such as view, table, or column name is misspelled, or non-existent in the SQL command, there will be no indication that any sort of error has occurred. No syntax error will be given and the entire hint is simply ignored.
- If an alias is used for a table or view name in the SQL command, then the alias name, not the original object name, must be used in the hint. For example, in the command, `SELECT /*+ FULL(acct) */ * FROM accounts acct ..., acct`, the alias for `accounts`, must be specified in the `FULL` hint, not the table name, `accounts`.

Use the `EXPLAIN` command to ensure that the hint is correctly formed and the planner is using the hint. See the Advanced Server documentation set for information on the `EXPLAIN` command.

In general, optimizer hints should not be used in production applications (where table data changes throughout the life of the application). By ensuring that dynamic columns are `ANALYZED` frequently, the column statistics will be updated to reflect value changes, and the planner will use such information to produce the least cost plan for any given command execution. Use of optimizer hints defeats the purpose of this process and will result in the same plan regardless of how the table data changes.

Parameters

hint

An optimizer hint directive.

comment

A string with additional information. Note that there are restrictions as to what characters may be included in the comment. Generally, `comment` may only consist of alphabetic, numeric, the underscore, dollar sign, number sign and space characters. These must also conform to the syntax of an identifier. Any subsequent hint will be ignored if the comment is not in this form.

statement_body

The remainder of the `DELETE`, `INSERT`, `SELECT`, or `UPDATE` command.

The following sections describe the optimizer hint directives in more detail.

7.5.1 Default Optimization Modes

There are a number of optimization modes that can be chosen as the default setting for an Advanced Server database cluster. This setting can also be changed on a per session basis by using the `ALTER SESSION` command as well as in individual `DELETE`, `SELECT`, and `UPDATE` commands within an optimizer hint. The configuration parameter that controls these default modes is named `OPTIMIZER_MODE`.

The following table shows the possible values.

| Hint | Description |
|-----------------------|--|
| <code>ALL_ROWS</code> | Optimizes for retrieval of all rows of the result set. |

| | |
|-----------------|---|
| CHOOSE | Does no default optimization based on assumed number of rows to be retrieved from the result set. This is the default. |
| FIRST_ROWS | Optimizes for retrieval of only the first row of the result set. |
| FIRST_ROWS_10 | Optimizes for retrieval of the first 10 rows of the results set. |
| FIRST_ROWS_100 | Optimizes for retrieval of the first 100 rows of the result set. |
| FIRST_ROWS_1000 | Optimizes for retrieval of the first 1000 rows of the result set. |
| FIRST_ROWS(n) | Optimizes for retrieval of the first <i>n</i> rows of the result set. This form may not be used as the object of the <code>ALTER SESSION SET OPTIMIZER_MODE</code> command. It may only be used in the form of a hint in a SQL command. |

These optimization modes are based upon the assumption that the client submitting the SQL command is interested in viewing only the first “n” rows of the result set and will then abandon the remainder of the result set. Resources allocated to the query are adjusted as such.

Examples

Alter the current session to optimize for retrieval of the first 10 rows of the result set.

```
ALTER SESSION SET OPTIMIZER_MODE = FIRST_ROWS_10;
```

The current value of the `OPTIMIZER_MODE` parameter can be shown by using the `SHOW` command. Note that this command is a utility dependent command. In PSQL, the `SHOW` command is used as follows:

```
SHOW OPTIMIZER_MODE;
```

```
optimizer_mode
```

```
-----  
first_rows_10  
(1 row)
```

The `SHOW` command, compatible with Oracle databases, has the following syntax:

```
SHOW PARAMETER OPTIMIZER_MODE;
```

```
NAME
```

```
-----  
VALUE
```

```
optimizer_mode  
first_rows_10
```

The following example shows an optimization mode used in a `SELECT` command as a hint:

```
SELECT /*+ FIRST_ROWS(7) */ * FROM emp;
```

```
empno|ename|job|mgr|hiredate|sal|comm|deptno
```

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | JOB_TITLE | HIRE_DATE | SALARY | MGR_ID | COMMISSION_PCT |
|-------------|------------|-----------|-----------|--------------------|---------|---------|----------------|
| 7369 | SMITH | CLERK | CLERK | 17-DEC-80 00:00:00 | 800.00 | | 20 |
| 7499 | ALLEN | SALESMAN | SALESMAN | 20-FEB-81 00:00:00 | 1600.00 | 30 | 300.00 |
| 7521 | WARD | SALESMAN | SALESMAN | 22-FEB-81 00:00:00 | 1250.00 | 500.00 | 30 |
| 7566 | JONES | MANAGER | MANAGER | 02-APR-81 00:00:00 | 2975.00 | | 20 |
| 7654 | MARTIN | SALESMAN | SALESMAN | 28-SEP-81 00:00:00 | 1250.00 | 1400.00 | 30 |
| 7698 | BLAKE | MANAGER | MANAGER | 01-MAY-81 00:00:00 | 2850.00 | | 30 |
| 7782 | CLARK | MANAGER | MANAGER | 09-JUN-81 00:00:00 | 2450.00 | | 10 |
| 7788 | SCOTT | ANALYST | ANALYST | 19-APR-87 00:00:00 | 3000.00 | | 20 |
| 7839 | KING | PRESIDENT | PRESIDENT | 17-NOV-81 00:00:00 | 5000.00 | | 10 |
| 7844 | TURNER | SALESMAN | SALESMAN | 08-SEP-81 00:00:00 | 1500.00 | 0.00 | 30 |
| 7876 | ADAMS | CLERK | CLERK | 23-MAY-87 00:00:00 | 1100.00 | | 20 |
| 7900 | JAMES | CLERK | CLERK | 03-DEC-81 00:00:00 | 950.00 | | 30 |
| 7902 | FORD | ANALYST | ANALYST | 03-DEC-81 00:00:00 | 3000.00 | | 20 |
| 7934 | MILLER | CLERK | CLERK | 23-JAN-82 00:00:00 | 1300.00 | | 10 |

7.5.2 Access Method Hints

The following hints influence how the optimizer accesses relations to create the result set.

| Hint | Description |
|---------------------------------|---|
| FULL(table) | Perform a full sequential scan on <code>table</code> . |
| INDEX(table [index] [...]) | Use <code>index</code> on <code>table</code> to access the relation. |
| NO INDEX(table [index] [...]) | Do not use <code>index</code> on <code>table</code> to access the relation. |

In addition, the `ALL ROWS`, `FIRST ROWS`, and `FIRST ROWS(n)` hints can be used.

Examples

The sample application does not have sufficient data to illustrate the effects of optimizer hints so the remainder of the examples in this section will use a banking database created by the `pgbench` application located in the Advanced Server `bin` subdirectory.

The following steps create a database named, `bank`, populated by the tables, `pgbench_accounts`, `pgbench_branches`, `pgbench_tellers`, and `pgbench_history`. The `-s 20` option specifies a scaling factor of twenty, which results in the creation of twenty branches, each with 100,000 accounts, resulting in a total of 2,000,000 rows in the `pgbench_accounts` table and twenty rows in the `pgbench_branches` table. Ten tellers are assigned to each branch resulting in a total of 200 rows in the `pgbench_tellers` table.

The following initializes the `pabench` application in the `bank` database.

```
createdb -U enterpriseDB bank  
CREATE DATABASE
```

```
ngbench -i -s 20 -U enterpriseedb bank
```

```
NOTICE: table "pgbench_history" does not exist, skipping  
NOTICE: table "pgbench_tellers" does not exist, skipping  
NOTICE: table "pgbench_accounts" does not exist, skipping  
NOTICE: table "pgbench_branches" does not exist, skipping
```

creating tables...

```
100000 of 2000000 tuples (5%) done (elapsed 0.11 s, remaining 2.10 s)
200000 of 2000000 tuples (10%) done (elapsed 0.22 s, remaining 1.98 s)
300000 of 2000000 tuples (15%) done (elapsed 0.33 s, remaining 1.84 s)
400000 of 2000000 tuples (20%) done (elapsed 0.42 s, remaining 1.67 s)
500000 of 2000000 tuples (25%) done (elapsed 0.52 s, remaining 1.57 s)
600000 of 2000000 tuples (30%) done (elapsed 0.62 s, remaining 1.45 s)
700000 of 2000000 tuples (35%) done (elapsed 0.73 s, remaining 1.35 s)
800000 of 2000000 tuples (40%) done (elapsed 0.87 s, remaining 1.31 s)
900000 of 2000000 tuples (45%) done (elapsed 0.98 s, remaining 1.19 s)
1000000 of 2000000 tuples (50%) done (elapsed 1.09 s, remaining 1.09 s)
1100000 of 2000000 tuples (55%) done (elapsed 1.22 s, remaining 1.00 s)
1200000 of 2000000 tuples (60%) done (elapsed 1.36 s, remaining 0.91 s)
1300000 of 2000000 tuples (65%) done (elapsed 1.51 s, remaining 0.82 s)
1400000 of 2000000 tuples (70%) done (elapsed 1.65 s, remaining 0.71 s)
1500000 of 2000000 tuples (75%) done (elapsed 1.78 s, remaining 0.59 s)
1600000 of 2000000 tuples (80%) done (elapsed 1.93 s, remaining 0.48 s)
1700000 of 2000000 tuples (85%) done (elapsed 2.10 s, remaining 0.37 s)
1800000 of 2000000 tuples (90%) done (elapsed 2.23 s, remaining 0.25 s)
1900000 of 2000000 tuples (95%) done (elapsed 2.37 s, remaining 0.12 s)
2000000 of 2000000 tuples (100%) done (elapsed 2.48 s, remaining 0.00 s)
```

vacuum...

set primary keys...

done.

A total of 500,00 transactions are then processed. This will populate the `pgbench_history` table with 500,000 rows.

```
pgbench -U enterpriseedb -t 500000 bank
```

starting vacuum...end.

transaction type: <builtin: TPC-B (sort of)>

scaling factor: 20

query mode: simple

number of clients: 1

number of threads: 1

number of transactions per client: 500000

number of transactions actually processed: 500000/500000

latency average: 0.000 ms

tps = 1464.338375 (including connections establishing)

tps = 1464.350357 (excluding connections establishing)

The table definitions are shown below:

```
\d pgbench_accounts
```

Table "public.pgbench_accounts"

| Column | Type | Modifiers |
|--------|------|-----------|
|--------|------|-----------|

```
+-----+-----+
```

| | | |
|-----|---------|----------|
| aid | integer | not null |
|-----|---------|----------|

| | | |
|-----|---------|--|
| bid | integer | |
|-----|---------|--|

| | | |
|----------|---------|--|
| abalance | integer | |
|----------|---------|--|

| | | |
|--------|---------------|--|
| filler | character(84) | |
|--------|---------------|--|

Indexes:

"pgbench_accounts_pkey" PRIMARY KEY, btree (aid)

```
\d pgbench_branches
```

Table "public.pgbench_branches"

| Column | Type | Modifiers |
|----------|---------------|-----------|
| bid | integer | not null |
| bbalance | integer | |
| filler | character(88) | |

Indexes:

- "pgbench_branches_pkey" PRIMARY KEY, btree (bid)

```
\d pgbench_tellers
```

Table "public.pgbench_tellers"

| Column | Type | Modifiers |
|----------|---------------|-----------|
| tid | integer | not null |
| bid | integer | |
| tbalance | integer | |
| filler | character(84) | |

Indexes:

- "pgbench_tellers_pkey" PRIMARY KEY, btree (tid)

```
\d pgbench_history
```

Table "public.pgbench_history"

| Column | Type | Modifiers |
|--------|-----------------------------|-----------|
| tid | integer | |
| bid | integer | |
| aid | integer | |
| delta | integer | |
| mtime | timestamp without time zone | |
| filler | character(22) | |

The **EXPLAIN** command shows the plan selected by the query planner. In the following example, **aid** is the primary key column, so an indexed search is used on index, **pgbench_accounts_pkey**.

```
EXPLAIN SELECT * FROM pgbench_accounts WHERE aid = 100;
```

QUERY PLAN

Index Scan using pgbench_accounts_pkey on pgbench_accounts (cost=0.43..8.45
rows=1 width=97)

 Index Cond: (aid = 100)
(2 rows)

The **FULL** hint is used to force a full sequential scan instead of using the index as shown below:

```
EXPLAIN SELECT /*+ FULL(pgbench_accounts) */ * FROM pgbench_accounts WHERE  
aid = 100;
```

QUERY PLAN

Seq Scan on pgbench_accounts (cost=0.00..58781.69 rows=1 width=97)

Filter: (aid = 100)
(2 rows)

The `NO_INDEX` hint forces a parallel sequential scan instead of use of the index as shown below:

```
EXPLAIN SELECT /*+ NO_INDEX(pgbench_accounts pgbench_accounts_pkey) */ *
FROM pgbench_accounts WHERE aid = 100;
```

QUERY PLAN

```
-----
Gather (cost=1000.00..45094.80 rows=1 width=97)
Workers Planned: 2
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..44094.70 rows=1
width=97)
      Filter: (aid = 100)
(4 rows)
```

In addition to using the `EXPLAIN` command as shown in the prior examples, more detailed information regarding whether or not a hint was used by the planner can be obtained by setting the `trace_hints` configuration parameter as follows:

```
SET trace_hints TO on;
```

The `SELECT` command with the `NO_INDEX` hint is repeated below to illustrate the additional information produced when the `trace_hints` configuration parameters is set.

```
EXPLAIN SELECT /*+ NO_INDEX(pgbench_accounts pgbench_accounts_pkey) */ *
FROM pgbench_accounts WHERE aid = 100;
```

INFO: [HINTS] Index Scan of [pgbench_accounts].[pgbench_accounts_pkey]
rejected due to NO_INDEX hint.

QUERY PLAN

```
-----
Gather (cost=1000.00..45094.80 rows=1 width=97)
Workers Planned: 2
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..44094.70 rows=1
width=97)
      Filter: (aid = 100)
(4 rows)
```

Note that if a hint is ignored, the `INFO: [HINTS]` line will not appear. This may be an indication that there was a syntax error or some other misspelling in the hint as shown in the following example where the index name is misspelled.

```
EXPLAIN SELECT /*+ NO_INDEX(pgbench_accounts pgbench_accounts_xxx) */ * FROM
pgbench_accounts WHERE aid = 100;
```

QUERY PLAN

```
-----
Index Scan using pgbench_accounts_pkey on pgbench_accounts (cost=0.43..8.45
rows=1 width=97)
Index Cond: (aid = 100)
(2 rows)
```

7.5.3 Specifying a Join Order

Include the **ORDERED** directive to instruct the query optimizer to join tables in the order in which they are listed in the **FROM** clause. If you do not include the **ORDERED** keyword, the query optimizer will choose the order in which to join the tables.

For example, the following command allows the optimizer to choose the order in which to join the tables listed in the **FROM** clause:

```
SELECT e.ename, d.dname, h.startdate
  FROM emp e, dept d, jobhist h
 WHERE d.deptno = e.deptno
   AND h.empno = e.empno;
```

The following command instructs the optimizer to join the tables in the ordered specified:

```
SELECT /*+ ORDERED */ e.ename, d.dname, h.startdate
  FROM emp e, dept d, jobhist h
 WHERE d.deptno = e.deptno
   AND h.empno = e.empno;
```

In the **ORDERED** version of the command, Advanced Server will first join **emp e** with **dept d** before joining the results with **jobhist h**. Without the **ORDERED** directive, the join order is selected by the query optimizer.

Note

The **ORDERED** directive does not work for Oracle-style outer joins (those joins that contain a '+' sign).

7.5.4 Joining Relations Hints

When two tables are to be joined, there are three possible plans that may be used to perform the join.

- *Nested Loop Join* – A table is scanned once for every row in the other joined table.
- *Merge Sort Join* – Each table is sorted on the join attributes before the join starts. The two tables are then scanned in parallel and the matching rows are combined to form the join rows.
- *Hash Join* – A table is scanned and its join attributes are loaded into a hash table using its join attributes as hash keys. The other joined table is then scanned and its join attributes are used as hash keys to locate the matching rows from the first table.

The following table lists the optimizer hints that can be used to influence the planner to use one type of join plan over another.

| Hint | Description |
|------------------------------------|--|
| <code>USE_HASH(table [...])</code> | Use a hash join for <code>table</code> . |

`NO_USE_HASH(table [...])`

| Do not use a hash join for `table`.

`USE_MERGE(table [...])`

| Use a merge sort join for `table`.

`NO_USE_MERGE(table [...])`

| Do not use a merge sort join for `table`.

`USE_NL(table [...])`

| Use a nested loop join for `table`.

`NO_USE_NL(table [...])`

| Do not use a nested loop join for `table`.

Examples

In the following example, the `USE_HASH` hint is used for a join on the `pgbench_branches` and `pgbench_accounts` tables. The query plan shows that a hash join is used by creating a hash table from the join attribute of the `pgbench_branches` table.

```
EXPLAIN SELECT /*+ USE_HASH(b) */ b.bid, a.aid, abalance FROM
pgbench_branches b, pgbench_accounts a WHERE b.bid = a.bid;
```

QUERY PLAN

```
-----
Hash Join (cost=21.45..81463.06 rows=2014215 width=12)
  Hash Cond: (a.bid = b.bid)
    -> Seq Scan on pgbench_accounts a (cost=0.00..53746.15 rows=2014215
width=12)
        -> Hash (cost=21.20..21.20 rows=20 width=4)
            -> Seq Scan on pgbench_branches b (cost=0.00..21.20 rows=20
width=4)
(5 rows)
```

Next, the `NO_USE_HASH(a b)` hint forces the planner to use an approach other than hash tables. The result is a merge join.

```
EXPLAIN SELECT /*+ NO_USE_HASH(a b) */ b.bid, a.aid, abalance FROM
pgbench_branches b, pgbench_accounts a WHERE b.bid = a.bid;
```

QUERY PLAN

```
-----
Merge Join (cost=333526.08..368774.94 rows=2014215 width=12)
  Merge Cond: (b.bid = a.bid)
    -> Sort (cost=21.63..21.68 rows=20 width=4)
        Sort Key: b.bid
        -> Seq Scan on pgbench_branches b (cost=0.00..21.20 rows=20
width=4)
    -> Materialize (cost=333504.45..343575.53 rows=2014215 width=12)
        -> Sort (cost=333504.45..338539.99 rows=2014215 width=12)
            Sort Key: a.bid
            -> Seq Scan on pgbench_accounts a (cost=0.00..53746.15
rows=2014215 width=12)
```

(9 rows)

Finally, the `USE_MERGE` hint forces the planner to use a merge join.

```
EXPLAIN SELECT /*+ USE_MERGE(a) */ b.bid, a.aid, abalance FROM
pgbench_branches b, pgbench_accounts a WHERE b.bid = a.bid;
```

QUERY PLAN

```
-----
Merge Join (cost=333526.08..368774.94 rows=2014215 width=12)
  Merge Cond: (b.bid = a.bid)
    -> Sort (cost=21.63..21.68 rows=20 width=4)
      Sort Key: b.bid
      -> Seq Scan on pgbench_branches b  (cost=0.00..21.20 rows=20
width=4)
    -> Materialize (cost=333504.45..343575.53 rows=2014215 width=12)
      -> Sort (cost=333504.45..338539.99 rows=2014215 width=12)
        Sort Key: a.bid
        -> Seq Scan on pgbench_accounts a  (cost=0.00..53746.15
rows=2014215 width=12)
(9 rows)
```

In this three-table join example, the planner first performs a hash join on the `pgbench_branches` and `pgbench_history` tables, then finally performs a hash join of the result with the `pgbench_accounts` table.

```
EXPLAIN SELECT h.mtime, h.delta, b.bid, a.aid FROM pgbench_history h,
pgbench_branches b, pgbench_accounts a WHERE h.bid = b.bid AND h.aid = a.aid;
```

QUERY PLAN

```
-----
Hash Join (cost=86814.29..123103.29 rows=500000 width=20)
  Hash Cond: (h.aid = a.aid)
    -> Hash Join (cost=21.45..15081.45 rows=500000 width=20)
      Hash Cond: (h.bid = b.bid)
      -> Seq Scan on pgbench_history h  (cost=0.00..8185.00
rows=500000 width=20)
      -> Hash (cost=21.20..21.20 rows=20 width=4)
        -> Seq Scan on pgbench_branches b  (cost=0.00..21.20 rows=20
width=4)
    -> Hash (cost=53746.15..53746.15 rows=2014215 width=4)
      -> Seq Scan on pgbench_accounts a  (cost=0.00..53746.15 rows=2014215
width=4)
(9 rows)
```

This plan is altered by using hints to force a combination of a merge sort join and a hash join.

```
EXPLAIN SELECT /*+ USE_MERGE(h b) USE_HASH(a) */ h.mtime, h.delta, b.bid,
a.aid FROM pgbench_history h, pgbench_branches b, pgbench_accounts a WHERE
h.bid = b.bid AND h.aid = a.aid;
```

QUERY PLAN

```
-----
Hash Join (cost=152583.39..182562.49 rows=500000 width=20)
```

```

Hash Cond: (h.aid = a.aid)
-> Merge Join (cost=65790.55..74540.65 rows=500000 width=20)
  Merge Cond: (b.bid = h.bid)
  -> Sort (cost=21.63..21.68 rows=20 width=4)
    Sort Key: b.bid
    -> Seq Scan on pgbench_branches b (cost=0.00..21.20 rows=20
width=4)
-> Materialize (cost=65768.92..68268.92 rows=500000 width=20)
  -> Sort (cost=65768.92..67018.92 rows=500000 width=20)
    Sort Key: h.bid
    -> Seq Scan on pgbench_history h (cost=0.00..8185.00
rows=500000 width=20)
-> Hash (cost=53746.15..53746.15 rows=2014215 width=4)
  -> Seq Scan on pgbench_accounts a (cost=0.00..53746.15
rows=2014215 width=4)
(13 rows)

```

7.5.5 Global Hints

Thus far, hints have been applied directly to tables that are referenced in the SQL command. It is also possible to apply hints to tables that appear in a view when the view is referenced in the SQL command. The hint does not appear in the view, itself, but rather in the SQL command that references the view.

When specifying a hint that is to apply to a table within a view, the view and table names are given in dot notation within the hint argument list.

Synopsis

`<hint>(<view>.<table>)`

Parameters

`hint`

Any of the hints in table [Access Method Hints, Joining Relations Hints](#).

`view`

The name of the view containing `table`.

`table`

The table on which the hint is to be applied.

Examples

A view named, `tx`, is created from the three-table join of `pgbench_history`, `pgbench_branches`, and `pgbench_accounts` shown in the final example of [Joining Relations Hints](#).

```
CREATE VIEW tx AS SELECT h.mtime, h.delta, b.bid, a.aid FROM pgbench_history
h, pgbench_branches b, pgbench_accounts a WHERE h.bid = b.bid AND h.aid =
a.aid;
```

The query plan produced by selecting from this view is shown below:

```
EXPLAIN SELECT * FROM tx;
```

QUERY PLAN

```
-----
Hash Join (cost=86814.29..123103.29 rows=500000 width=20)
  Hash Cond: (h.aid = a.aid)
    -> Hash Join (cost=21.45..15081.45 rows=500000 width=20)
      Hash Cond: (h.bid = b.bid)
        -> Seq Scan on pgbench_history h (cost=0.00..8185.00 rows=500000
width=20)
          -> Hash (cost=21.20..21.20 rows=20 width=4)
            -> Seq Scan on pgbench_branches b (cost=0.00..21.20 rows=20
width=4)
    -> Hash (cost=53746.15..53746.15 rows=2014215 width=4)
      -> Seq Scan on pgbench_accounts a (cost=0.00..53746.15
rows=2014215 width=4)
(9 rows)
```

The same hints that were applied to this join at the end of [Joining Relations Hints](#) can be applied to the view as follows:

```
EXPLAIN SELECT /*+ USE_MERGE(tx.h tx.b) USE_HASH(tx.a) */ * FROM tx;
```

QUERY PLAN

```
-----
Hash Join (cost=152583.39..182562.49 rows=500000 width=20)
  Hash Cond: (h.aid = a.aid)
    -> Merge Join (cost=65790.55..74540.65 rows=500000 width=20)
      Merge Cond: (b.bid = h.bid)
        -> Sort (cost=21.63..21.68 rows=20 width=4)
          Sort Key: b.bid
            -> Seq Scan on pgbench_branches b (cost=0.00..21.20 rows=20
width=4)
        -> Materialize (cost=65768.92..68268.92 rows=500000 width=20)
          -> Sort (cost=65768.92..67018.92 rows=500000 width=20)
            Sort Key: h.bid
              -> Seq Scan on pgbench_history h (cost=0.00..8185.00
rows=500000 width=20)
    -> Hash (cost=53746.15..53746.15 rows=2014215 width=4)
      -> Seq Scan on pgbench_accounts a (cost=0.00..53746.15
rows=2014215 width=4)
(13 rows)
```

In addition to applying hints to tables within stored views, hints can be applied to tables within subqueries as illustrated by the following example. In this query on the sample application `emp` table, employees and their managers are listed by joining the `emp` table with a subquery of the `emp` table identified by the alias, `b`.

```
SELECT a.empno, a.ename, b.empno "mgr empno", b.ename "mgr ename" FROM emp
a, (SELECT * FROM emp) b WHERE a.mgr = b.empno;
```

| empno | ename | mgr empno | mgr ename |
|-------|-------|-----------|-----------|
| 7369 | SMITH | 7902 | FORD |
| 7499 | ALLEN | 7698 | BLAKE |

```

7521 | WARD |    7698 | BLAKE
7566 | JONES |   7839 | KING
7654 | MARTIN | 7698 | BLAKE
7698 | BLAKE |   7839 | KING
7782 | CLARK |   7839 | KING
7788 | SCOTT |   7566 | JONES
7844 | TURNER | 7698 | BLAKE
7876 | ADAMS |   7788 | SCOTT
7900 | JAMES |   7698 | BLAKE
7902 | FORD |    7566 | JONES
7934 | MILLER | 7782 | CLARK
(13 rows)

```

The plan chosen by the query planner is shown below:

```
EXPLAIN SELECT a.empno, a.ename, b.empno "mgr empno", b.ename "mgr ename"
FROM emp a, (SELECT * FROM emp) b WHERE a.mgr = b.empno;
```

QUERY PLAN

```

Hash Join (cost=1.32..2.64 rows=13 width=22)
  Hash Cond: (a.mgr = emp.empno)
    -> Seq Scan on emp a (cost=0.00..1.14 rows=14 width=16)
    -> Hash (cost=1.14..1.14 rows=14 width=11)
      -> Seq Scan on emp (cost=0.00..1.14 rows=14 width=11)
(5 rows)

```

A hint can be applied to the `emp` table within the subquery to perform an index scan on index, `emp_pk`, instead of a table scan. Note the difference in the query plans.

```
EXPLAIN SELECT /*+ INDEX(b.emp emp_pk) */ a.empno, a.ename, b.empno "mgr
empno", b.ename "mgr ename" FROM emp a, (SELECT * FROM emp) b WHERE a.mgr =
b.empno;
```

QUERY PLAN

```

Merge Join (cost=4.17..13.11 rows=13 width=22)
  Merge Cond: (a.mgr = emp.empno)
    -> Sort (cost=1.41..1.44 rows=14 width=16)
      Sort Key: a.mgr
        -> Seq Scan on emp a (cost=0.00..1.14 rows=14 width=16)
    -> Index Scan using emp_pk on emp (cost=0.14..12.35 rows=14 width=11)
(6 rows)

```

7.5.6 Using the APPEND Optimizer Hint

By default, Advanced Server will add new data into the first available free-space in a table (vacated by vacuumed records). Include the `APPEND` directive after an `INSERT` or `SELECT` command to instruct the server to bypass mid-table free space, and affix new rows to the end of the table. This optimizer hint can be particularly useful when bulk loading data.

The syntax is:

```
/*+APPEND*/
```

For example, the following command, compatible with Oracle databases, instructs the server to append the data in the `INSERT` statement to the end of the `sales` table:

```
INSERT /*+APPEND*/ INTO sales VALUES  
(10, 10, '01-Mar-2011', 10, 'OR');
```

Note that Advanced Server supports the `APPEND` hint when adding multiple rows in a single `INSERT` statement:

```
INSERT /*+APPEND*/ INTO sales VALUES  
(20, 20, '01-Aug-2011', 20, 'NY'),  
(30, 30, '01-Feb-2011', 30, 'FL'),  
(40, 40, '01-Nov-2011', 40, 'TX');
```

The `APPEND` hint can also be included in the `SELECT` clause of an `INSERT INTO` statement:

```
INSERT INTO sales_history SELECT /*+APPEND*/ FROM sales;
```

7.5.7 Parallelism Hints

The `PARALLEL` optimizer hint is used to force parallel scanning.

The `NO_PARALLEL` optimizer hint prevents usage of a parallel scan.

Synopsis

```
PARALLEL (<table> [ <parallel_degree> | DEFAULT ])
```

```
NO_PARALLEL (<table>)
```

Description

Parallel scanning is the usage of multiple background workers to simultaneously perform a scan of a table (that is, in parallel) for a given query. This process provides performance improvement over other methods such as the sequential scan.

Parameters

`table`

The table to which the parallel hint is to be applied.

`parallel_degree | DEFAULT`

`parallel_degree` is a positive integer that specifies the desired number of workers to use for a parallel scan. If specified, the lesser of `parallel_degree` and configuration parameter `max_parallel_workers_per_gather` is used as the planned number of workers. For information on the `max_parallel_workers_per_gather` parameter, see [Asynchronous Behavior](#) located under [Resource Consumption](#) in the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/runtime-config-resource.html>

If `DEFAULT` is specified, then the maximum possible parallel degree is used.

If both `parallel_degree` and `DEFAULT` are omitted, then the query optimizer determines the parallel degree. In this case, if `table` has been set with the `parallel_workers` storage parameter, then this value is used as the parallel degree, otherwise the optimizer uses the maximum possible parallel degree as if `DEFAULT` was specified. For information on the `parallel_workers` storage parameter, see the [Storage Parameters](#) located under `CREATE TABLE` in the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-createtable.html>

Regardless of the circumstance, the parallel degree never exceeds the setting of configuration parameter `max_parallel_workers_per_gather`.

Examples

The following configuration parameter settings are in effect:

```
SHOW max_worker_processes;
```

```
max_worker_processes
```

```
8
```

```
(1 row)
```

```
SHOW max_parallel_workers_per_gather;
```

```
max_parallel_workers_per_gather
```

```
2
```

```
(1 row)
```

The following example shows the default scan on table `pgbench_accounts`. Note that a sequential scan is shown in the query plan.

```
SET trace_hints TO on;
```

```
EXPLAIN SELECT * FROM pgbench_accounts;
```

```
QUERY PLAN
```

```
Seq Scan on pgbench_accounts (cost=0.00..53746.15 rows=2014215 width=97)
```

```
(1 row)
```

The following example uses the `PARALLEL` hint. In the query plan, the Gather node, which launches the background workers, indicates that two workers are planned to be used.

Note

If `trace_hints` is set to `on`, the `INFO: [HINTS]` lines appear stating that `PARALLEL` has been accepted for `pgbench_accounts` as well as other hint information. For the remaining examples, these lines will not be displayed as they generally show the same output (that is, `trace_hints` has been reset to `off`).

```
EXPLAIN SELECT /*+ PARALLEL(pgbench_accounts) */ * FROM pgbench_accounts;
```

```
INFO: [HINTS] SeqScan of [pgbench_accounts] rejected due to PARALLEL hint.
```

```
INFO: [HINTS] PARALLEL on [pgbench_accounts] accepted.
```

```
INFO: [HINTS] Index Scan of [pgbench_accounts].[pgbench_accounts_pkey]  
rejected due to PARALLEL hint.
```

QUERY PLAN

```
Gather (cost=1000.00..244418.06 rows=2014215 width=97)
Workers Planned: 2
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..41996.56
rows=839256 width=97)
(3 rows)
```

Now, the `max_parallel_workers_per_gather` setting is increased:

```
SET max_parallel_workers_per_gather TO 6;
```

```
SHOW max_parallel_workers_per_gather;
```

```
max_parallel_workers_per_gather
```

```
6
(1 row)
```

The same query on `pgbench_accounts` is issued again with no parallel degree specification in the `PARALLEL` hint. Note that the number of planned workers has increased to `4` as determined by the optimizer.

```
EXPLAIN SELECT /*+ PARALLEL(pgbench_accounts) */ * FROM pgbench_accounts;
```

QUERY PLAN

```
Gather (cost=1000.00..241061.04 rows=2014215 width=97)
Workers Planned: 4
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..38639.54
rows=503554 width=97)
(3 rows)
```

Now, a value of `6` is specified for the parallel degree parameter of the `PARALLEL` hint. The planned number of workers is now returned as this specified value:

```
EXPLAIN SELECT /*+ PARALLEL(pgbench_accounts 6) */ * FROM pgbench_accounts;
```

QUERY PLAN

```
Gather (cost=1000.00..239382.52 rows=2014215 width=97)
Workers Planned: 6
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..36961.03
rows=335702 width=97)
(3 rows)
```

The same query is now issued with the `DEFAULT` setting for the parallel degree. The results indicate that the maximum allowable number of workers is planned.

```
EXPLAIN SELECT /*+ PARALLEL(pgbench_accounts DEFAULT) */ * FROM
pgbench_accounts;
```

QUERY PLAN

```
-----
Gather (cost=1000.00..239382.52 rows=2014215 width=97)
Workers Planned: 6
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..36961.03
rows=335702 width=97)
(3 rows)
```

Table `pgbench_accounts` is now altered so that the `parallel_workers` storage parameter is set to `3`.

Note

This format of the `ALTER TABLE` command to set the `parallel_workers` parameter is not compatible with Oracle databases.

The `parallel_workers` setting is shown by the PSQL `\d+` command.

```
ALTER TABLE pgbench_accounts SET (parallel_workers=3);
```

```
\d+ pgbench_accounts
      Table "public.pgbench_accounts"
Column | Type   | Modifiers | Storage | Stats target | Description
-----+-----+-----+-----+-----+
aid    | integer | not null | plain   |          |
bid    | integer |           | plain   |          |
abalance| integer |           | plain   |          |
filler | character(84) |           | extended |          |
Indexes:
  "pgbench_accounts_pkey" PRIMARY KEY, btree (aid)
Options: fillfactor=100, parallel_workers=3
```

Now, when the `PARALLEL` hint is given with no parallel degree, the resulting number of planned workers is the value from the `parallel_workers` parameter:

```
EXPLAIN SELECT /*+ PARALLEL(pgbench_accounts) */ * FROM pgbench_accounts;
```

QUERY PLAN

```
-----
Gather (cost=1000.00..242522.97 rows=2014215 width=97)
Workers Planned: 3
-> Parallel Seq Scan on pgbench_accounts (cost=0.00..40101.47
rows=649747 width=97)
(3 rows)
```

Specifying a parallel degree value or `DEFAULT` in the `PARALLEL` hint overrides the `parallel_workers` setting.

The following example shows the `NO_PARALLEL` hint. Note that with `trace hints` set to `on`, the `INFO: [HINTS]` message states that the parallel scan was rejected due to the `NO_PARALLEL` hint.

```
EXPLAIN SELECT /*+ NO_PARALLEL(pgbench_accounts) */ * FROM pgbench_accounts;
INFO: [HINTS] Parallel SeqScan of [pgbench_accounts] rejected due to
NO_PARALLEL hint.
```

QUERY PLAN

```
-----
Seq Scan on pgbench_accounts (cost=0.00..53746.15 rows=2014215 width=97)
(1 row)
```

7.5.8 Conflicting Hints

If a command includes two or more conflicting hints, the server will ignore the contradictory hints. The following table lists hints that are contradictory to each other.

| Hint | Conflicting Hint |
|----------------------------|----------------------------|
| ALL_ROWS | FIRST_ROWS - all formats |
| FULL(table) | INDEX(table [index]) |
| | PARALLEL(table [degree]) |
| INDEX(table) | FULL(table) |
| | NO_INDEX(table) |
| | PARALLEL(table [degree]) |
| INDEX(table index) | FULL(table) |
| | NO_INDEX(table index) |
| | PARALLEL(table [degree]) |
| PARALLEL(table [degree]) | FULL(table) |
| | INDEX(table) |
| | NO_PARALLEL(table) |
| USE_HASH(table) | NO_USE_HASH(table) |
| USE_MERGE(table) | NO_USE_MERGE(table) |
| USE_NL(table) | NO_USE_NL(table) |

7.6 dblink_ora

dblink_ora provides an OCI-based database link that allows you to `SELECT`, `INSERT`, `UPDATE` or `DELETE` data stored on an Oracle system from within Advanced Server.

Connecting to an Oracle Database

To enable Oracle connectivity, download Oracle's freely available OCI drivers from their website, presently at:

<http://www.oracle.com/technetwork/database/database-technologies/instant-client/overview/index.html>

For Linux, if the Oracle instant client that you've downloaded does not include the `libclntsh.so` library, you must create a symbolic link named `libclntsh.so` that points to the downloaded version. Navigate to the instant client directory and execute the following command:

```
| ln -s libclntsh.so.<version> libclntsh.so
```

Where `version` is the version number of the `libclntsh.so` library. For example:

```
| ln -s libclntsh.so.12.1 libclntsh.so
```

Before creating a link to an Oracle server, you must tell Advanced Server where to find the OCI driver.

Set the `LD_LIBRARY_PATH` environment variable on Linux (or `PATH` on Windows) to the `lib` directory of the Oracle client installation directory.

For Windows only, you can instead set the value of the `oracle_home` configuration parameter in the `postgresql.conf` file. The value specified in the `oracle_home` configuration parameter will override the Windows `PATH` environment variable.

The `LD_LIBRARY_PATH` environment variable on Linux (`PATH` environment variable or `oracle_home` configuration parameter on Windows) must be set properly each time you start Advanced Server.

When using a Linux service script to start Advanced Server, be sure `LD_LIBRARY_PATH` has been set within the service script so it is in effect when the script invokes the `pg_ctl` utility to start Advanced Server.

For Windows only: To set the `oracle_home` configuration parameter in the `postgresql.conf` file, edit the file, adding the following line:

```
| oracle_home = 'lib_directory'
```

Substitute the name of the Windows directory that contains `oci.dll` for `lib_directory`.

After setting the `oracle_home` configuration parameter, you must restart the server for the changes to take effect. Restart the server from the Windows Services console.

7.6.1 dblink_ora Functions and Procedures

dblink_ora supports the following functions and procedures.

7.6.1.1 dblink_ora_connect()

The `dblink_ora_connect()` function establishes a connection to an Oracle database with user-specified connection information. The function comes in two forms; the signature of the first form is:

```
dblink_ora_connect(<conn_name>, <server_name>, <service_name>, <user_name>,
<password>, <port>, <asDBA>)
```

Where:

`conn_name` specifies the name of the link.

`server_name` specifies the name of the host.

`service_name` specifies the name of the service.

`user_name` specifies the name used to connect to the server.

`password` specifies the password associated with the user name.

`port` specifies the port number.

`asDBA` is `True` if you wish to request `SYSDBA` privileges on the Oracle server. This parameter is optional; if omitted, the default value is `FALSE`.

The first form of `dblink_ora_connect()` returns a `TEXT` value.

The signature of the second form of the `dblink_ora_connect()` function is:

```
dblink_ora_connect(<foreign_server_name>, <asDBA>)
```

Where:

`foreign_server_name` specifies the name of a foreign server.

`asDBA` is `True` if you wish to request `SYSDBA` privileges on the Oracle server. This parameter is optional; if omitted, the default value is `FALSE`.

The second form of the `dblink_ora_connect()` function allows you to use the connection properties of a pre-defined foreign server when establishing a connection to the server.

Before invoking the second form of the `dblink_ora_connect()` function, use the `CREATE SERVER` command to store the connection properties for the link to a system table. When you call the `dblink_ora_connect()` function, substitute the server name specified in the `CREATE SERVER` command for the name of the link.

The second form of `dblink_ora_connect()` returns a `TEXT` value.

7.6.1.2 `dblink_ora_status()`

The `dblink_ora_status()` function returns the database connection status. The signature is:

```
dblink_ora_status(<conn_name>)
```

Where:

`conn_name` specifies the name of the link.

If the specified connection is active, the function returns a `TEXT` value of `OK`.

7.6.1.3 `dblink_ora_disconnect()`

The `dblink_ora_disconnect()` function closes a database connection. The signature is:

```
dblink_ora_disconnect(<conn_name>)
```

Where:

`conn_name` specifies the name of the link.

The function returns a `TEXT` value.

7.6.1.4 `dblink_ora_record()`

The `dblink_ora_record()` function retrieves information from a database. The signature is:

```
dblink_ora_record(<conn_name>, <query_text>)
```

Where:

`conn_name` specifies the name of the link.

`query_text` specifies the text of the SQL `SELECT` statement that will be invoked on the Oracle server.

The function returns a `SETOF` record.

7.6.1.5 `dblink_ora_call()`

The `dblink_ora_call()` function executes a non-`SELECT` statement on an Oracle database and returns a result set. The signature is:

```
dblink_ora_call(<conn_name>, <command>, <iterations>)
```

Where:

`conn_name` specifies the name of the link.

`command` specifies the text of the SQL statement that will be invoked on the Oracle server.

`iterations` specifies the number of times the statement is executed.

The function returns a `SETOF` record.

7.6.1.6 `dblink_ora_exec()`

The `dblink_ora_exec()` procedure executes a DML or DDL statement in the remote database. The signature is:

```
dblink_ora_exec(<conn_name>, <command>)
```

Where:

`conn_name` specifies the name of the link.

`command` specifies the text of the `INSERT`, `UPDATE`, or `DELETE` SQL statement that will be invoked on the Oracle server.

The function returns a `VOID`.

7.6.1.7 dblink_ora_copy()

The `dblink ora copy()` function copies an Oracle table to an EnterpriseDB table. The `dblink_ora_copy()` function returns a `BIGINT` value that represents the number of rows copied. The signature is:

```
dblink_ora_copy(<conn_name>, <command>, <schema_name>, <table_name>,
<truncate>, <count>)
```

Where:

`conn_name` specifies the name of the link.

`command` specifies the text of the SQL `SELECT` statement that will be invoked on the Oracle server.

`schema_name` specifies the name of the target schema.

`table_name` specifies the name of the target table.

`truncate` specifies if the server should `TRUNCATE` the table prior to copying; specify `TRUE` to indicate that the server should `TRUNCATE` the table. `truncate` is optional; if omitted, the value is `FALSE`.

`count` instructs the server to report status information every `n` record, where `n` is the number specified. During the execution of the function, Advanced Server raises a notice of severity `INFO` with each iteration of the count. For example, if FeedbackCount is `10`, `dblink_ora_copy()` raises a notice every `10` records. `count` is optional; if omitted, the value is `0`.

7.6.2 Calling dblink_ora Functions

The following command establishes a connection using the `dblink_ora_connect()` function:

```
SELECT dblink_ora_connect('acctg', 'localhost', 'xe', 'hr', 'pwd', 1521);
```

The example connects to a service named `xe` running on port `1521` (on the `localhost`) with a user name of `hr` and a password of `pwd`. You can use the connection name `acctg` to refer to this connection when calling other `dblink_ora` functions.

The following command uses the `dblink_ora_copy()` function over a connection named `edb_conn` to copy the `empid` and `deptno` columns from a table (on an Oracle server) named `ora_acctg` to a table located in the `public` schema on an instance of Advanced Server named `as_acctg`. The `TRUNCATE` option is enforced, and a feedback count of `3` is specified:

```
edb=# SELECT dblink_ora_copy('edb_conn','select empid, deptno FROM ora_acctg', 'public', 'as_acctg', true, 3);
```

```
INFO: Row: 0
INFO: Row: 3
INFO: Row: 6
INFO: Row: 9
INFO: Row: 12
```

```
dblink_ora_copy
-----
12
(1 row)
```

The following `SELECT` statement uses `dblink_ora_record()` function and the `acctg` connection to retrieve information from the Oracle server:


```
SELECT * FROM dblink_ora_record( 'acctg', 'SELECT first_name from employees') AS t1(id VARCHAR);
```

The command retrieves a list that includes all of the entries in the `first_name` column of the `employees` table.

7.7 Open Client Library

The Open Client Library provides application interoperability with the Oracle Call Interface – an application that was formerly “locked in” can now work with either an EDB Postgres Advanced Server or an Oracle database with minimal to no changes to the application code. The EnterpriseDB implementation of the Open Client Library is written in C.

The following diagram compares the Open Client Library and Oracle Call Interface application stacks.



For detailed usage information about the Open Client Library and the supported functions, see the *EDB Postgres Advanced Server OCL Connector Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Note

EnterpriseDB does not support use of the Open Client Library with Oracle Real Application Clusters (RAC) and Oracle Exadata; the aforementioned Oracle products have not been evaluated nor certified with this EnterpriseDB product.

7.8 Oracle Catalog Views

The Oracle Catalog Views provide information about database objects in a manner compatible with the Oracle data dictionary views. Information about the supported views is now available in the *Database Compatibility for Oracle Developer's Catalog Views Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

7.9 Tools and Utilities

Compatible tools and utility programs can allow a developer to work with Advanced Server in a familiar environment. The tools supported by Advanced Server include:

- EDB*Plus
- EDB*Loader
- EDB Wrap
- The Dynamic Runtime Instrumentation Tools Architecture (DRITA)

For detailed information about the functionality supported by Advanced Server, see the *Database Compatibility for Oracle Developer's Tools and Utilities Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

7.10 ECPGPlus

EnterpriseDB has enhanced ECPG (the PostgreSQL pre-compiler) to create ECPGPlus. ECPGPlus allows you to include embedded SQL commands in C applications; when you use ECPGPlus to compile an application that contains embedded SQL commands, the SQL code is syntax-checked and translated into C.

ECPGPlus supports Pro*C compatible syntax in C programs when connected to an Advanced Server database. ECPGPlus supports:

- Oracle Dynamic SQL – Method 4 (ODS-M4).

- Pro*C compatible anonymous blocks.
- A `CALL` statement compatible with Oracle databases.

As part of ECPGPlus' Pro*C compatibility, you do not need to include the `BEGIN DECLARE SECTION` and `END DECLARE SECTION` directives.

For more information about using ECPGPlus, see the *EDB Postgres Advanced Server ECPG Connector Guide* available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

7.11 System Catalog Tables

The system catalog tables contain definitions of database objects that are available to Advanced Server; the layout of the system tables is subject to change. If you are writing an application that depends on information stored in the system tables, it would be prudent to use an existing catalog view, or create a catalog view to isolate the application from changes to the system table.

For detailed information about the system catalog tables, see the *Database Compatibility for Oracle Developer's Catalog Views Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

8 Database Compatibility for Oracle

8.1 Introduction

The tools and utilities documented in this guide allow a developer that is accustomed to working with Oracle utilities to work with Advanced Server in a familiar environment.

The sections in this guide describe compatible tools and utilities that are supported by Advanced Server. These include:

- EDB*Loader
- EDB Wrap
- Dynamic Runtime Instrumentation

The EDB*Plus command line client provides a user interface to Advanced Server that supports SQL*Plus commands; EDB*Plus allows you to:

- Query database objects
- Execute stored procedures
- Format output from SQL commands
- Execute batch scripts
- Execute OS commands

- Record output

For detailed installation and usage information about EDB*Plus, see the *EDB Plus User's Guide*, available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs/p/edbplus>

For detailed information about the features supported by Advanced Server, consult the complete library of Advanced Server guides available at:

<https://www.enterprisedb.com/edb-docs>

8.2 EDB Loader

EDB*Loader is a high-performance bulk data loader that provides an interface compatible with Oracle databases for Advanced Server. The EDB*Loader command line utility loads data from an input source, typically a file, into one or more tables using a subset of the parameters offered by Oracle SQL*Loader.

EDB*Loader features include:

- Support for the Oracle SQL*Loader data loading methods - conventional path load, direct path load, and parallel direct path load
- Syntax for control file directives compatible with Oracle SQL*Loader
- Input data with delimiter-separated or fixed-width fields
- Bad file for collecting rejected records
- Loading of multiple target tables
- Discard file for collecting records that do not meet the selection criteria of any target table
- Log file for recording the EDB*Loader session and any error messages
- Data loading from standard input and remote loading, particularly useful for large data sources on remote hosts

These features are explained in detail in the following sections.

Note

The following are important version compatibility restrictions between the EDB*Loader client and the database server.

- When you invoke the EDB*Loader program (called `edbldr`), you pass in parameters and directive information to the database server. **We strongly recommend that the version 13 EDB*Loader client (the `edbldr` program supplied with Advanced Server 13) be used to load data only into version 13 of the database server. In general, the EDB*Loader client and database server should be the same version.**
- Use of a version 13, 12, 11, 10, 9.6, 9.5, 9.4 or 9.3 EDB*Loader client is not supported for database servers version 9.2 or earlier.

Data Loading Methods

As with Oracle SQL*Loader, EDB*Loader supports three data loading methods:

- Conventional path load
- Direct path load
- Parallel direct path load

Conventional path load is the default method used by EDB*Loader. Basic insert processing is used to add rows to

the table.

The advantage of a conventional path load over the other methods is that table constraints and database objects defined on the table such as primary keys, not null constraints, check constraints, unique indexes, foreign key constraints, and triggers are enforced during a conventional path load.

One exception is that the Advanced Server *rules* defined on the table are not enforced. EDB*Loader can load tables on which rules are defined, but the rules are not executed. As a consequence, partitioned tables implemented using rules cannot be loaded using EDB*Loader.

Note

Advanced Server rules are created by the `CREATE RULE` command. Advanced Server rules are not the same database objects as rules and rule sets used in Oracle.

EDB*Loader also supports direct path loads. A direct path load is faster than a conventional path load, but requires the removal of most types of constraints and triggers from the table. For more information, see [Direct Path Load](#).

Finally, EDB*Loader supports parallel direct path loads. A parallel direct path load provides even greater performance improvement by permitting multiple EDB*Loader sessions to run simultaneously to load a single table. For more information, see [Parallel Direct Path Load](#).

General Usage

EDB*Loader can load data files with either delimiter-separated or fixed-width fields, in single-byte or multi-byte character sets. The delimiter can be a string consisting of one or more single-byte or multi-byte characters. Data file encoding and the database encoding may be different. Character set conversion of the data file to the database encoding is supported.

Each EDB*Loader session runs as a single, independent transaction. If an error should occur during the EDB*Loader session that aborts the transaction, all changes made during the session are rolled back.

Generally, formatting errors in the data file do not result in an aborted transaction. Instead, the badly formatted records are written to a text file called the *bad file*. The reason for the error is recorded in the *log file*.

Records causing database integrity errors do result in an aborted transaction and rollback. As with formatting errors, the record causing the error is written to the bad file and the reason is recorded in the log file.

Note

EDB*Loader differs from Oracle SQL*Loader in that a database integrity error results in a rollback in EDB*Loader. In OracleSQL*Loader, only the record causing the error is rejected. Records that were previously inserted into the table are retained and loading continues after the rejected record.

The following are examples of types of formatting errors that do not abort the transaction:

- Attempt to load non-numeric value into a numeric column
- Numeric value is too large for a numeric column
- Character value is too long for the maximum length of a character column
- Attempt to load improperly formatted date value into a date column

The following are examples of types of database errors that abort the transaction and result in the rollback of all changes made in the EDB*Loader session:

- Violation of a unique constraint such as a primary key or unique index
- Violation of a referential integrity constraint
- Violation of a check constraint
- Error thrown by a trigger fired as a result of inserting rows

Building the EDB Loader Control File

When you invoke EDB*Loader, the list of arguments provided must include the name of a control file. The control file includes the instructions that EDB*Loader uses to load the table (or tables) from the input data file. The control file includes information such as:

- The name of the input data file containing the data to be loaded.
- The name of the table or tables to be loaded from the data file.
- Names of the columns within the table or tables and their corresponding field placement in the data file.
- Specification of whether the data file uses a delimiter string to separate the fields, or if the fields occupy fixed column positions.
- Optional selection criteria to choose which records from the data file to load into a given table.
- The name of the file that will collect illegally formatted records.
- The name of the discard file that will collect records that do not meet the selection criteria of any table.

The syntax for the EDB*Loader control file is as follows:

```
[ OPTIONS (<param=value> [, <param=value>] ...) ]
LOAD DATA
[ CHARACTERSET <charset> ]
[ INFILE '{ <data_file> | <stdin> }' ]
[ BADFILE '<bad_file>' ]
[ DISCARDFILE '<discard_file>' ]
[ { DISCARDMAX | DISCARDS } <max_discard_recs> ]
[ INSERT | APPEND | REPLACE | TRUNCATE ]
[ PRESERVE BLANKS ]
{ INTO TABLE <target_table>
  [ WHEN <field_condition> [ AND <field_condition> ] ...]
  [ FIELDS TERMINATED BY '<termstring>' ]
  [ OPTIONALLY ENCLOSED BY '<enclstring>' ] ]
[ RECORDS DELIMITED BY '<delimstring>' ]
[.TRAILING NULLCOLS ]
  (<field_def> [, <field_def>] ...)
} ...
```

where `field_def` defines a field in the specified `data_file` that describes the location, data format, or value of the data to be inserted into `column_name` of the `target_table`. The syntax of `field_def` is the following:

```
<column_name> {
  CONSTANT <val> |
  FILLER [ POSITION (<start:end>) ] [ <fieldtype> ] |
  BOUNDFILLER [ POSITION (<start:end>) ] [ <fieldtype> ] |
  [ POSITION (<start:end>) ] [ <fieldtype> ]
  [ NULLIF <field_condition> [ AND <field_condition> ] ...]
  [ PRESERVE BLANKS ] [ "<expr>" ]
}
```

where `fieldtype` is one of:

```
CHAR [(<length>)] | DATE [(<length>)] [ "<datemask>" ] |
INTEGER EXTERNAL [(<length>)] |
FLOAT EXTERNAL [(<length>)] | DECIMAL EXTERNAL [(length)] |
ZONED EXTERNAL [(<length>)] | ZONED [(<precision> ,<scale>)]
```

Description

The specification of `data_file`, `bad_file`, and `discard_file` may include the full directory path or a relative directory

path to the file name. If the file name is specified alone or with a relative directory path, the file is then assumed to exist (in the case of `data_file`), or is created (in the case of `bad_file` or `discard_file`), relative to the current working directory from which `edbldr` is invoked.

You can include references to environment variables within the EDB*Loader control file when referring to a directory path and/or file name. Environment variable references are formatted differently on Windows systems than on Linux systems:

- On Linux, the format is `$ENV VARIABLE` or `${ENV_VARIABLE}`
- On Windows, the format is `%ENV_VARIABLE%`

Where `ENV_VARIABLE` is the environment variable that is set to the directory path and/or file name.

The `EDBLDR_ENV_STYLE` environment variable instructs Advanced Server to interpret environment variable references as Windows-styled references or Linux-styled references irregardless of the operating system on which EDB*Loader resides. You can use this environment variable to create portable control files for EDB*Loader.

- On a Windows system, set `EDBLDR_ENV_STYLE` to `linux` or `unix` to instruct Advanced Server to recognize Linux-style references within the control file.
- On a Linux system, set `EDBLDR_ENV_STYLE` to `windows` to instruct Advanced Server to recognize Windows-style references within the control file.

The operating system account `enterprisedb` must have read permission on the directory and file specified by `data_file`.

The operating system account `enterprisedb` must have write permission on the directories where `bad_file` and `discard_file` are to be written.

Note

The file names for `data_file`, `bad_file`, and `discard_file` should include extensions of `.dat`, `.bad`, and `.dsc`, respectively. If the provided file name does not contain an extension, EDB*Loader assumes the actual file name includes the appropriate aforementioned extension.

If an EDB*Loader session results in data format errors and the `BADFILE` clause is not specified, nor is the `BAD` parameter given on the command line when `edbldr` is invoked, a bad file is created with the name `control_file_base.bad` in the current working directory from which `edbldr` is invoked. `control_file_base` is the base name of the control file (that is, the file name without any extension) used in the `edbldr` session.

If all of the following conditions are true, the discard file is not created even if the EDB*Loader session results in discarded records:

- The `DISCARDFILE` clause for specifying the discard file is not included in the control file.
- The `DISCARD` parameter for specifying the discard file is not included on the command line.
- The `DISCARDMAX` clause for specifying the maximum number of discarded records is not included in the control file.
- The `DISCARDS` clause for specifying the maximum number of discarded records is not included in the control file.
- The `DISCARDMAX` parameter for specifying the maximum number of discarded records is not included on the command line.

If neither the `DISCARDFILE` clause nor the `DISCARD` parameter for explicitly specifying the discard file name are specified, but `DISCARDMAX` or `DISCARDS` is specified, then the EDB*Loader session creates a discard file using the data file name with an extension of `.dsc`.

Note

There is a distinction between keywords `DISCARD` and `DISCARDS`. `DISCARD` is an EDB*Loader command line parameter used to specify the discard file name (see [General Usage](#)). `DISCARDS` is a clause of the `LOAD DATA` directive that may only appear in the control file. Keywords `DISCARDS` and `DISCARDMAX` provide the same functionality of specifying the maximum number of discarded records allowed before terminating the EDB*Loader session. Records loaded into the database before termination of the EDB*Loader session due to exceeding the

`DISCARDS` or `DISCARDMAX` settings are kept in the database and are not rolled back.

If one of `INSERT`, `APPEND`, `REPLACE`, or `TRUNCATE` is specified, it establishes the default action of how rows are to be added to target tables. If omitted, the default action is as if `INSERT` had been specified.

If the `FIELDS TERMINATED BY` clause is specified, then the `POSITION (start:end)` clause may not be specified for any `field_def`. Alternatively if the `FIELDS TERMINATED BY` clause is not specified, then every `field_def` must contain either the `POSITION (start:end)` clause, the `fieldtype(length)` clause, or the `CONSTANT` clause.

Parameters

`OPTIONS param=value`

Use the `OPTIONS` clause to specify `param=value` pairs that represent an EDB*Loader directive. If a parameter is specified in both the `OPTIONS` clause and on the command line when `edbldr` is invoked, the command line setting is used.

Specify one or more of the following parameter/value pairs:

> `DIRECT= { FALSE | TRUE }`

If `DIRECT` is set to `TRUE` EDB*Loader performs a direct path load instead of a conventional path load. The default value of `DIRECT` is `FALSE`.

You should not set `DIRECT=true` when loading the data into a replicated table. If you are using EDB*Loader to load data into a replicated table and set `DIRECT=true`, indexes may omit rows that are in a table or may potentially contain references to rows that have been deleted. EnterpriseDB does not support direct inserts to load data into replicated tables.

For information on direct path loads see, [Direct Path Load](#).

`ERRORS=error_count`

`error_count` specifies the number of errors permitted before aborting the EDB*Loader session. The default is `50`.

`FREEZE= { FALSE | TRUE }`

Set `FREEZE` to `TRUE` to indicate that the data should be copied with the rows `frozen`. A tuple guaranteed to be visible to all current and future transactions is marked as frozen to prevent transaction ID wrap-around. For more information about frozen tuples, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/routine-vacuuming.html>

You must specify a data-loading type of `TRUNCATE` in the control file when using the `FREEZE` option. `FREEZE` is not supported for direct loading.

By default, `FREEZE` is `FALSE`.

`PARALLEL= { FALSE | TRUE }`

Set `PARALLEL` to `TRUE` to indicate that this EDB*Loader session is one of a number of concurrent EDB*Loader sessions participating in a parallel direct path load. The default value of `PARALLEL` is `FALSE`.

When `PARALLEL` is `TRUE`, the `DIRECT` parameter must also be set to `TRUE`. For more information about parallel direct path loads, see [Parallel Direct Path Load](#).

`ROWS=n`

`n` specifies the number of rows that EDB*Loader will commit before loading the next set of `n` rows.

If EDB*Loader encounters an invalid row during a load (in which the `ROWS` parameter is specified), those rows committed prior to encountering the error will remain in the destination table.

`SKIP=skip_count`

`skip_count` specifies the number of records at the beginning of the input data file that should be skipped before loading begins. The default is `0`.

`SKIP_INDEX_MAINTENANCE={ FALSE | TRUE }`

If `SKIP_INDEX_MAINTENANCE` is `TRUE`, index maintenance is not performed as part of a direct path load, and indexes on the loaded table are marked as invalid. The default value of `SKIP_INDEX_MAINTENANCE` is `FALSE`.

Please note: During a parallel direct path load, target table indexes are not updated, and are marked as invalid after the load is complete.

You can use the `REINDEX` command to rebuild an index. For more information about the `REINDEX` command, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-reindex.html>

`charset`

Use the `CHARACTERSET` clause to identify the character set encoding of `data_file` where `charset` is the character set name. This clause is required if the data file encoding differs from the control file encoding. (The control file encoding must always be in the encoding of the client where `edbldr` is invoked.)

Examples of `charset` settings are `UTF8`, `SQL_ASCII`, and `SJIS`.

For more information about client to database character set conversion, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/multibyte.html>

`data_file`

File containing the data to be loaded into `target_table`. Each record in the data file corresponds to a row to be inserted into `target_table`.

If an extension is not provided in the file name, EDB*Loader assumes the file has an extension of `.dat`, for example, `mydatafile.dat`.

Note

If the `DATA` parameter is specified on the command line when `edbldr` is invoked, the file given by the command line `DATA` parameter is used instead.

If the `INFILE` clause is omitted as well as the command line `DATA` parameter, then the data file name is assumed to be identical to the control file name, but with an extension of `.dat`.

`stdin`

Specify `stdin` (all lowercase letters) if you want to use standard input to pipe the data to be loaded directly to EDB*Loader. This is useful for data sources generating a large number of records to be loaded.

`bad_file`

A file that receives `data_file` records that cannot be loaded due to errors. The bad file is generated for collecting rejected or bad records.

From Advanced Server version 12 and onwards, a bad file will be generated only if there are any bad or

rejected records. However, if there is an existing bad file with identical name and location, and no bad records are generated after invoking a new version of `edbldr`, the existing bad file remains untouched.

If an extension is not provided in the file name, EDB*Loader assumes the file has an extension of `.bad`, for example, `mybadfile.bad`.

Note

If the `BAD` parameter is specified on the command line when `edbldr` is invoked, the file given by the command line `BAD` parameter is used instead.

`discard_file`

File that receives input data records that are not loaded into any table because none of the selection criteria are met for tables with the `WHEN` clause, and there are no tables without a `WHEN` clause. (All records meet the selection criteria of a table without a `WHEN` clause.)

If an extension is not provided in the file name, EDB*Loader assumes the file has an extension of `.dsc`, for example, `mydiscardfile.dsc`.

Note

If the `DISCARD` parameter is specified on the command line when `edbldr` is invoked, the file given by the command line `DISCARD` parameter is used instead.

{ `DISCARDMAX` | `DISCARDS` } `max_discard_recs`

Maximum number of discarded records that may be encountered from the input data records before terminating the EDB*Loader session. (A discarded record is described in the preceding description of the `discard_file` parameter.) Either keyword `DISCARDMAX` or `DISCARDS` may be used preceding the integer value specified by `max_discard_recs`.

For example, if `max_discard_recs` is `0`, then the EDB*Loader session is terminated if and when a first discarded record is encountered. If `max_discard_recs` is `1`, then the EDB*Loader session is terminated if and when a second discarded record is encountered.

When the EDB*Loader session is terminated due to exceeding `max_discard_recs`, prior input data records that have been loaded into the database are retained. They are not rolled back.

`INSERT` | `APPEND` | `REPLACE` | `TRUNCATE`

Specifies how data is to be loaded into the target tables. If one of `INSERT`, `APPEND`, `REPLACE`, or `TRUNCATE` is specified, it establishes the default action for all tables, overriding the default of `INSERT`.

`INSERT`

Data is to be loaded into an empty table. EDB*Loader throws an exception and does not load any data if the table is not initially empty.

Note

If the table contains rows, the `TRUNCATE` command must be used to empty the table prior to invoking EDB*Loader. EDB*Loader throws an exception if the `DELETE` command is used to empty the table instead of the `TRUNCATE` command. Oracle SQL*Loader allows the table to be emptied by using either the `DELETE` or `TRUNCATE` command.

`APPEND`

Data is to be added to any existing rows in the table. The table may be initially empty as well.

`REPLACE`

The `REPLACE` keyword and `TRUNCATE` keywords are functionally identical. The table is truncated by EDB*Loader prior to loading the new data.

Note

Delete triggers on the table are not fired as a result of the `REPLACE` operation.

`TRUNCATE`

The table is truncated by EDB*Loader prior to loading the new data. Delete triggers on the table are not fired as a result of the `TRUNCATE` operation.

`PRESERVE BLANKS`

For all target tables, retains leading white space when the optional enclosure delimiters are not present and leaves trailing white space intact when fields are specified with a predetermined size. When omitted, the default behavior is to trim leading and trailing white space.

`target_table`

Name of the table into which data is to be loaded. The table name may be schema-qualified (for example, `enterprisedb.emp`). The specified target must not be a view.

`field_condition`

Conditional clause taking the following form:

`[() { (start:end) | column_name } { = | != | <> } 'val' []]`

This conditional clause is used for the `WHEN` clause, which is part of the `INTO TABLE target_table` clause, and the `NULLIF` clause, which is part of the field definition denoted as `field_def` in the syntax diagram.

`start` and `end` are positive integers specifying the column positions in `data_file` that mark the beginning and end of a field that is to be compared with the constant `val`. The first character in each record begins with a `start` value of `1`.

`column_name` specifies the name assigned to a field definition of the data file as defined by `field_def` in the syntax diagram.

Use of either `(start:end)` or `column_name` defines the portion of the record in `data_file` that is to be compared with the value specified by 'val' to evaluate as either true or false.

All characters used in the `field_condition` text (particularly in the `val` string) must be valid in the database encoding. (For performing data conversion, EDB*Loader first converts the characters in `val` string to the database encoding and then to the data file encoding.)

In the `WHEN field_condition [AND field_condition]` clause, if all such conditions evaluate to `TRUE` for a given record, then EDB*Loader attempts to insert that record into `target_table`. If the insert operation fails, the record is written to `bad_file`.

If for a given record, none of the `WHEN` clauses evaluate to `TRUE` for all `INTO TABLE` clauses, the record is written to `discard_file`, if a discard file was specified for the EDB*Loader session.

See the description of the `NULLIF` clause in this Parameters list for the effect of `field_condition` on this clause.

`termstring`

String of one or more characters that separates each field in `data_file`. The characters may be single-byte or multi-byte as long as they are valid in the database encoding. Two consecutive appearances of `termstring` with no intervening character results in the corresponding column set to null.

enclstring

String of one or more characters used to enclose a field value in `data_file`. The characters may be single-byte or multi-byte as long as they are valid in the database encoding. Use `enclstring` on fields where `termstring` appears as part of the data.

delimstring

String of one or more characters that separates each record in `data_file`. The characters may be single-byte or multi-byte as long as they are valid in the database encoding. Two consecutive appearances of `delimstring` with no intervening character results in no corresponding row loaded into the table. The last record (in other words, the end of the data file) must also be terminated by the `delimstring` characters, otherwise the final record is not loaded into the table.

Note

The `RECORDS DELIMITED BY 'delimstring'` clause is not compatible with Oracle databases.

TRAILING NULLCOLS

If `TRAILING NULLCOLS` is specified, then the columns in the column list for which there is no data in `data_file` for a given record, are set to null when the row is inserted. This applies only to one or more consecutive columns at the end of the column list.

If fields are omitted at the end of a record and `TRAILING NULLCOLS` is not specified, EDB*Loader assumes the record contains formatting errors and writes it to the bad file.

column_name

Name of a column in `target table` into which a field value defined by `field def` is to be inserted. If the field definition includes the `FILLER` or `BOUNDFILLER` clause, then `column_name` is not required to be the name of a column in the table. It can be any identifier name since the `FILLER` and `BOUNDFILLER` clauses prevent the loading of the field data into a table column.

CONSTANT val

Specifies a constant that is type-compatible with the column data type to which it is assigned in a field definition. Single or double quotes may enclose `val`. If `val` contains white space, then enclosing quotation marks must be used.

The use of the `CONSTANT` clause completely determines the value to be assigned to a column in each inserted row. No other clause may appear in the same field definition.

If the `TERMINATED BY` clause is used to delimit the fields in `data_file`, there must be no delimited field in `data_file` corresponding to any field definition with a `CONSTANT` clause. In other words, EDB*Loader assumes there is no field in `data_file` for any field definition with a `CONSTANT` clause.

FILLER

Specifies that the data in the field defined by the field definition is not to be loaded into the associated column if the identifier of the field definition is an actual column name in the table. In such case, the column is set to null. Use of the `FILLER` or `BOUNDFILLER` clause is the only circumstance in which the field definition does not have to be identified by an actual column name.

Unlike the `BOUNDFILLER` clause, an identifier defined with the `FILLER` clause must not be referenced in a `SQL` expression. See the discussion of the `expr` parameter.

BOUNDFILLER

Specifies that the data in the field defined by the field definition is not to be loaded into the associated column if the identifier of the field definition is an actual column name in the table. In such case, the column is set to null. Use of the `FILLER` or `BOUNDFILLER` clause is the only circumstance in which the field definition does

not have to be identified by an actual column name.

Unlike the `FILLER` clause, an identifier defined with the `BOUNDFILLER` clause may be referenced in a SQL expression. See the discussion of the `expr` parameter.

`POSITION (start:end)`

Defines the location of the field in a record in a fixed-width field data file. `start` and `end` are positive integers. The first character in the record has a start value of `1`.

`CHAR [(<length>)] | DATE [(<length>)] ["<datemask>"] |`

`INTEGER EXTERNAL [<length>] |`

`FLOAT EXTERNAL [<length>] | DECIMAL EXTERNAL [<length>] |`

`ZONED EXTERNAL [<length>] | ZONED [<precision>[,<scale>]]]`

Field type that describes the format of the data field in `data_file`.

Note

Specification of a field type is optional (for descriptive purposes only) and has no effect on whether or not EDB*Loader successfully inserts the data in the field into the table column. Successful loading depends upon the compatibility of the column data type and the field value. For example, a column with data type `NUMBER(7,2)` successfully accepts a field containing `2600`, but if the field contains a value such as `26XX`, the insertion fails and the record is written to `bad_file`.

Please note that `ZONED` data is not human-readable; `ZONED` data is stored in an internal format where each digit is encoded in a separate nibble/nybble/4-bit field. In each `ZONED` value, the last byte contains a single digit (in the high-order 4 bits) and the sign (in the low-order 4 bits).

`length`

Specifies the length of the value to be loaded into the associated column.

If the `POSITION (start:end)` clause is specified along with a `fieldtype(length)` clause, then the ending position of the field is overridden by the specified `length` value. That is, the length of the value to be loaded into the column is determined by the `length` value beginning at the `start` position, and not by the `end` position of the `POSITION (start:end)` clause. Thus, the value to be loaded into the column may be shorter than the field defined by `POSITION (start:end)`, or it may go beyond the `end` position depending upon the specified `length` size.

If the `FIELDS TERMINATED BY 'termstring'` clause is specified as part of the `INTO TABLE` clause, and a field definition contains the `fieldtype(length)` clause, then a record is accepted as long as the specified `length` values are greater than or equal to the field lengths as determined by the `termstring` characters enclosing all such fields of the record. If the specified `length` value is less than a field length as determined by the enclosing `termstring` characters for any such field, then the record is rejected.

If the `FIELDS TERMINATED BY 'termstring'` clause is not specified, and the `POSITION (start:end)` clause is not included with a field containing the `fieldtype(length)` clause, then the starting position of this field begins with the next character following the ending position of the preceding field. The ending position of the preceding field is either the end of its `length` value if the preceding field contains the `fieldtype(length)` clause, or by its `end` parameter if the field contains the `POSITION (start:end)` clause without the `fieldtype(length)` clause.

`precision`

Use `precision` to specify the length of the `ZONED` value.

If the `precision` value specified for `ZONED` conflicts with the length calculated by the server based on information provided with the `POSITION` clause, EDB*Loader will use the value specified for `precision`.

scale

scale specifies the number of digits to the right of the decimal point in a **ZONED** value.

datemask

Specifies the ordering and abbreviation of the day, month, and year components of a date field.

Note

If the **DATE** field type is specified along with a SQL expression for the column, then **datemask** must be specified after **DATE** and before the SQL expression. See the following discussion of the **expr** parameter.

NULLIF field_condition [AND field_condition] ...

See the description of **field_condition** previously listed in this Parameters section for the syntax of **field_condition**.

If all field conditions evaluate to **TRUE**, then the column identified by **column_name** in the field definition is set to null. If any field condition evaluates to **FALSE**, then the column is set to the appropriate value as would normally occur according to the field definition.

PRESERVE BLANKS

For the column on which this option appears, retains leading white space when the optional enclosure delimiters are not present and leaves trailing white space intact when fields are specified with a predetermined size. When omitted, the default behavior is to trim leading and trailing white space.

expr

A SQL expression returning a scalar value that is type-compatible with the column data type to which it is assigned in a field definition. Double quotes must enclose **expr**. **expr** may contain a reference to any column in the field list (except for fields with the **FILLER** clause) by prefixing the column name by a colon character **(:)**.

expr may also consist of a SQL **SELECT** statement. If a **SELECT** statement is used then the following rules must apply:

- The **SELECT** statement must be enclosed within parentheses **(SELECT ...)**.
- The select list must consist of exactly one expression following the **SELECT** keyword.
- The result set must not return more than one row. If no rows are returned, then the returned value of the resulting expression is null.

The following is the syntax for use of the **SELECT** statement:

"(SELECT expr [FROM table_list [WHERE condition]])"

Note

Omitting the **FROM table list** clause is not compatible with Oracle databases. If no tables need to be specified, use of the **FROM DUAL** clause is compatible with Oracle databases.

EDB Loader Control File Examples

The following are some examples of control files and their corresponding data files.

Delimiter-Separated Field Data File

The following control file uses a delimiter-separated data file that appends rows to the **emp** table:

```

LOAD DATA
INFILE 'emp.dat'
BADFILE 'emp.bad'
APPEND
INTO TABLE emp
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY """
TRAILING NULLCOLS
(
    empno,
    ename,
    job,
    mgr,
    hiredate,
    sal,
    deptno,
    comm
)

```

In the preceding control file, the `APPEND` clause is used to allow the insertion of additional rows into the `emp` table.

The following is the corresponding delimiter-separated data file:

```

9101,ROGERS,CLERK,7902,17-DEC-10,1980.00,20
9102,PETERSON,SALESMAN,7698,20-DEC-10,2600.00,30,2300.00
9103,WARREN,SALESMAN,7698,22-DEC-10,5250.00,30,2500.00
9104,"JONES, JR.",MANAGER,7839,02-APR-09,7975.00,20

```

The use of the `TRAILING NULLCOLS` clause allows the last field supplying the `comm` column to be omitted from the first and last records. The `comm` column is set to null for the rows inserted from these records.

The double quotation mark enclosure character surrounds the value `JONES, JR.` in the last record since the comma delimiter character is part of the field value.

The following query displays the rows added to the table after the EDB*Loader session:

```

SELECT * FROM emp WHERE empno > 9100;

empno|ename |job |mgr| hiredate | sal |comm |deptno
-----+-----+-----+-----+-----+-----+
9101| ROGERS | CLERK |7902| 17-DEC-10 00:00:00|1980.00|      | 20
9102| PETERSON | SALESMAN|7698| 20-DEC-10 00:00:00|2600.00| 2300.00| 30
9103| WARREN | SALESMAN|7698| 22-DEC-10 00:00:00|5250.00| 2500.00| 30
9104| JONES, JR.| MANAGER |7839| 02-APR-09 00:00:00|7975.00|      | 20
(4 rows)

```

Fixed-Width Field Data File

The following example is a control file that loads the same rows into the `emp` table, but uses a data file containing fixed-width fields:

```

LOAD DATA
INFILE 'emp_fixed.dat'
BADFILE 'emp_fixed.bad'
APPEND
INTO TABLE emp
TRAILING NULLCOLS

```

```
(  
    empno POSITION (1:4),  
    ename POSITION (5:14),  
    job POSITION (15:23),  
    mgr POSITION (24:27),  
    hiredate POSITION (28:38),  
    sal POSITION (39:46),  
    deptno POSITION (47:48),  
    comm POSITION (49:56)  
)
```

In the preceding control file, the `FIELDS TERMINATED BY` and `OPTIONALLY ENCLOSED BY` clauses are absent. Instead, each field now includes the `POSITION` clause.

The following is the corresponding data file containing fixed-width fields:

```
9101ROGERS CLERK 790217-DEC-10 1980.0020  
9102PETERSON SALESMAN 769820-DEC-10 2600.0030 2300.00  
9103WARREN SALESMAN 769822-DEC-10 5250.0030 2500.00  
9104JONES, JR.MANAGER 783902-APR-09 7975.0020
```

Single Physical Record Data File – RECORDS DELIMITED BY Clause

The following example is a control file that loads the same rows into the `emp` table, but uses a data file with one physical record. Each individual record that is to be loaded as a row in the table is terminated by the semicolon character (`;`) specified by the `RECORDS DELIMITED BY` clause.

```
LOAD DATA  
INFILE 'emp_recdelim.dat'  
BADFILE 'emp_recdelim.bad'  
APPEND  
INTO TABLE emp  
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY ""  
RECORDS DELIMITED BY ';'  
TRAILING NULLCOLS  
(  
    empno,  
    ename,  
    job,  
    mgr,  
    hiredate,  
    sal,  
    deptno,  
    comm  
)
```

The following is the corresponding data file. The content is a single, physical record in the data file. The record delimiter character is included following the last record (that is, at the end of the file).

```
9101,ROGERS,CLERK,7902,17-DEC-10,1980.00,20,;9102,PETERSON,SALESMAN,7698,20-  
DEC-10,2600.00,30,2300.00;9103,WARREN,SALESMAN,7698,22-DEC-  
10,5250.00,30,2500.00;9104,"JONES, JR.",MANAGER,7839,02-APR-09,7975.00,20,;
```

FILLER Clause

The following control file illustrates the use of the `FILLER` clause in the data fields for the `sal` and `comm` columns. EDB*Loader ignores the values in these fields and sets the corresponding columns to null.

```

LOAD DATA
INFILE  'emp_fixed.dat'
BADFILE 'emp_fixed.bad'
APPEND
INTO TABLE emp
TRAILING NULLCOLS
(
    empno      POSITION (1:4),
    ename      POSITION (5:14),
    job        POSITION (15:23),
    mgr        POSITION (24:27),
    hiredate   POSITION (28:38),
    sal        FILLER POSITION (39:46),
    deptno    POSITION (47:48),
    comm      FILLER POSITION (49:56)
)

```

Using the same fixed-width data file as in the prior fixed-width field example, the resulting rows in the table appear as follows:

```
SELECT * FROM emp WHERE empno > 9100;
```

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|------------|----------|------|--------------------|-----|------|--------|
| 9101 | ROGERS | CLERK | 7902 | 17-DEC-10 00:00:00 | | | 20 |
| 9102 | PETERSON | SALESMAN | 7698 | 20-DEC-10 00:00:00 | | | 30 |
| 9103 | WARREN | SALESMAN | 7698 | 22-DEC-10 00:00:00 | | | 30 |
| 9104 | JONES, JR. | MANAGER | 7839 | 02-APR-09 00:00:00 | | | 20 |

(4 rows)

BOUNDFILLER Clause

The following control file illustrates the use of the **BOUNDFILLER** clause in the data fields for the **job** and **mgr** columns. EDB*Loader ignores the values in these fields and sets the corresponding columns to null in the same manner as the **FILLER** clause. However, unlike columns with the **FILLER** clause, columns with the **BOUNDFILLER** clause are permitted to be used in an expression as shown for column **jobdesc**.

```

LOAD DATA
INFILE  'emp.dat'
BADFILE 'emp.bad'
APPEND
INTO TABLE empjob
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY ""
TRAILING NULLCOLS
(
    empno,
    ename,
    job    BOUNDFILLER,
    mgr    BOUNDFILLER,
    hiredate  FILLER,
    sal    FILLER,
    deptno  FILLER,
    comm    FILLER,
    jobdesc  ":job || ' for manager ' || :mgr"
)

```

The following is the delimiter-separated data file used in this example.

```
9101,ROGERS,CLERK,7902,17-DEC-10,1980.00,20
9102,PETERSON,SALESMAN,7698,20-DEC-10,2600.00,30,2300.00
9103,WARREN,SALESMAN,7698,22-DEC-10,5250.00,30,2500.00
9104,"JONES, JR.",MANAGER,7839,02-APR-09,7975.00,20
```

The following table is loaded using the preceding control file and data file.

```
CREATE TABLE empjob (
    empno      NUMBER(4) NOT NULL CONSTRAINT empjob_pk PRIMARY KEY,
    ename      VARCHAR2(10),
    job        VARCHAR2(9),
    mgr        NUMBER(4),
    jobdesc    VARCHAR2(25)
);
```

The resulting rows in the table appear as follows:

```
SELECT * FROM empjob;

empno | ename | job | mgr |      jobdesc
-----+-----+-----+-----+
9101 | ROGERS |   |   | CLERK for manager 7902
9102 | PETERSON |   |   | SALESMAN for manager 7698
9103 | WARREN |   |   | SALESMAN for manager 7698
9104 | JONES, JR. |   |   | MANAGER for manager 7839
(4 rows)
```

Field Types with Length Specification

The following example is a control file that contains the field type clauses with the length specification:

```
LOAD DATA
INFILE 'emp_fixed.dat'
BADFILE 'emp_fixed.bad'
APPEND
INTO TABLE emp
TRAILING NULLCOLS
(
    empno    CHAR(4),
    ename    CHAR(10),
    job      POSITION (15:23) CHAR(9),
    mgr      INTEGER EXTERNAL(4),
    hiredate  DATE(11) "DD-MON-YY",
    sal      DECIMAL EXTERNAL(8),
    deptno   POSITION (47:48),
    comm     POSITION (49:56) DECIMAL EXTERNAL(8)
)
```

Note

The **POSITION** clause and the **fieldtype(length)** clause can be used individually or in combination as long as each field definition contains at least one of the two clauses.

The following is the corresponding data file containing fixed-width fields:

```
9101ROGERS CLERK 790217-DEC-10 1980.0020
9102PETERSON SALESMAN 769820-DEC-10 2600.0030 2300.00
9103WARREN SALESMAN 769822-DEC-10 5250.0030 2500.00
9104JONES, JR. MANAGER 783902-APR-09 7975.0020
```

The resulting rows in the table appear as follows:

```
SELECT * FROM emp WHERE empno > 9100;
```

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|------------|----------|------|--------------------|---------|---------|--------|
| 9101 | ROGERS | CLERK | 7902 | 17-DEC-10 00:00:00 | 1980.00 | | 20 |
| 9102 | PETERSON | SALESMAN | 7698 | 20-DEC-10 00:00:00 | 2600.00 | 2300.00 | 30 |
| 9103 | WARREN | SALESMAN | 7698 | 22-DEC-10 00:00:00 | 5250.00 | 2500.00 | 30 |
| 9104 | JONES, JR. | MANAGER | 7839 | 02-APR-09 00:00:00 | 7975.00 | | 20 |

(4 rows)

NULLIF Clause

The following example uses the **NULLIF** clause on the sal column to set it to null for employees of job **MANAGER** as well as on the comm column to set it to null if the employee is not a **SALESMAN** and is not in department **30**. In other words, a comm value is accepted if the employee is a **SALESMAN** or is a member of department **30**.

The following is the control file:

```
LOAD DATA
INFILE 'emp_fixed_2.dat'
BADFILE 'emp_fixed_2.bad'
APPEND
INTO TABLE emp
TRAILING NULLCOLS
(
    empno      POSITION (1:4),
    ename      POSITION (5:14),
    job        POSITION (15:23),
    mgr        POSITION (24:27),
    hiredate   POSITION (28:38),
    sal        POSITION (39:46) NULLIF job = 'MANAGER',
    deptno    POSITION (47:48),
    comm      POSITION (49:56) NULLIF job <> 'SALESMAN' AND deptno <> '30'
)
```

The following is the corresponding data file:

```
9101ROGERS CLERK 790217-DEC-10 1980.0020
9102PETERSON SALESMAN 769820-DEC-10 2600.0030 2300.00
9103WARREN SALESMAN 769822-DEC-10 5250.0030 2500.00
9104JONES, JR. MANAGER 783902-APR-09 7975.0020
9105ARNOLDS CLERK 778213-SEP-10 3750.0030 800.00
9106JACKSON ANALYST 756603-JAN-11 4500.0040 2000.00
9107MAXWELL SALESMAN 769820-DEC-10 2600.0010 1600.00
```

The resulting rows in the table appear as follows:

```
SELECT empno, ename, job, NVL(TO_CHAR(sal),'--null--) "sal",
NVL(TO_CHAR(comm),'--null--) "comm", deptno FROM emp WHERE empno > 9100;
```

| empno | ename | job | sal | comm | deptno |
|-------|------------|----------|----------|----------|--------|
| 9101 | ROGERS | CLERK | 1980.00 | --null-- | 20 |
| 9102 | PETERSON | SALESMAN | 2600.00 | 2300.00 | 30 |
| 9103 | WARREN | SALESMAN | 5250.00 | 2500.00 | 30 |
| 9104 | JONES, JR. | MANAGER | --null-- | --null-- | 20 |
| 9105 | ARNOLDS | CLERK | 3750.00 | 800.00 | 30 |
| 9106 | JACKSON | ANALYST | 4500.00 | --null-- | 40 |
| 9107 | MAXWELL | SALESMAN | 2600.00 | 1600.00 | 10 |

(7 rows)

Note

The `sal` column for employee `JONES, JR.` is null since the job is `MANAGER`.

The `comm` values from the data file for employees `PETERSON`, `WARREN`, `ARNOLDS`, and `MAXWELL` are all loaded into the `comm` column of the `emp` table since these employees are either `SALESMAN` or members of department `30`.

The `comm` value of `2000.00` in the data file for employee `JACKSON` is ignored and the `comm` column of the `emp` table set to null since this employee is neither a `SALESMAN` nor is a member of department `30`.

SELECT Statement in a Field Expression

The following example uses a `SELECT` statement in the expression of the field definition to return the value to be loaded into the column.

```
LOAD DATA
INFILE  'emp_fixed.dat'
BADFILE 'emp_fixed.bad'
APPEND
INTO TABLE emp
TRAILING NULLCOLS
(
    empno      POSITION (1:4),
    ename      POSITION (5:14),
    job        POSITION (15:23) "(SELECT dname FROM dept WHERE deptno =
:deptno)",
    mgr        POSITION (24:27),
    hiredate   POSITION (28:38),
    sal        POSITION (39:46),
    deptno    POSITION (47:48),
    comm      POSITION (49:56)
)
```

The content of the `dept` table used in the `SELECT` statement is the following:

```
SELECT * FROM dept;
```

| deptno | dname | loc |
|--------|------------|----------|
| 10 | ACCOUNTING | NEW YORK |
| 20 | RESEARCH | DALLAS |
| 30 | SALES | CHICAGO |
| 40 | OPERATIONS | BOSTON |

(4 rows)

The following is the corresponding data file:

```
9101ROGERS CLERK 790217-DEC-10 1980.0020
9102PETERSON SALESMAN 769820-DEC-10 2600.0030 2300.00
9103WARREN SALESMAN 769822-DEC-10 5250.0030 2500.00
9104JONES, JR. MANAGER 783902-APR-09 7975.0020
```

The resulting rows in the table appear as follows:

```
SELECT * FROM emp WHERE empno > 9100;
```

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|------------|----------|------|--------------------|---------|---------|--------|
| 9101 | ROGERS | RESEARCH | 7902 | 17-DEC-10 00:00:00 | 1980.00 | | 20 |
| 9102 | PETERSON | SALES | 7698 | 20-DEC-10 00:00:00 | 2600.00 | 2300.00 | 30 |
| 9103 | WARREN | SALES | 7698 | 22-DEC-10 00:00:00 | 5250.00 | 2500.00 | 30 |
| 9104 | JONES, JR. | RESEARCH | 7839 | 02-APR-09 00:00:00 | 7975.00 | | 20 |

(4 rows)

Note

The `job` column contains the value from the `dname` column of the `dept` table returned by the `SELECT` statement instead of the job name from the data file.

Multiple INTO TABLE Clauses

The following example illustrates the use of multiple `INTO TABLE` clauses. For this example, two empty tables are created with the same data definition as the `emp` table. The following `CREATE TABLE` commands create these two empty tables, while inserting no rows from the original `emp` table:

```
CREATE TABLE emp_research AS SELECT * FROM emp WHERE deptno = 99;
CREATE TABLE emp_sales AS SELECT * FROM emp WHERE deptno = 99;
```

The following control file contains two `INTO TABLE` clauses. Also note that there is no `APPEND` clause so the default operation of `INSERT` is used, which requires that tables `emp_research` and `emp_sales` be empty.

```
LOAD DATA
INFILE      'emp_multtbl.dat'
BADFILE     'emp_multtbl.bad'
DISCARDFILE 'emp_multtbl.dsc'
INTO TABLE emp_research
WHEN (47:48) = '20'
TRAILING NULLCOLS
(
    empno    POSITION (1:4),
    ename    POSITION (5:14),
    job      POSITION (15:23),
    mgr      POSITION (24:27),
    hiredate  POSITION (28:38),
    sal      POSITION (39:46),
    deptno   CONSTANT '20',
    comm     POSITION (49:56)
)
INTO TABLE emp_sales
WHEN (47:48) = '30'
TRAILING NULLCOLS
(
```

```

empno  POSITION (1:4),
ename   POSITION (5:14),
job     POSITION (15:23),
mgr     POSITION (24:27),
hiredate POSITION (28:38),
sal     POSITION (39:46),
deptno  CONSTANT '30',
comm    POSITION (49:56) "ROUND(:comm + (:sal * .25), 0)"
)

```

The **WHEN** clauses specify that when the field designated by columns 47 thru 48 contains **20**, the record is inserted into the **emp_research** table and when that same field contains **30**, the record is inserted into the **emp_sales** table. If neither condition is true, the record is written to the discard file named **emp_multitbl.dsc**.

The **CONSTANT** clause is given for column **deptno** so the specified constant value is inserted into **deptno** for each record. When the **CONSTANT** clause is used, it must be the only clause in the field definition other than the column name to which the constant value is assigned.

Finally, column **comm** of the **emp_sales** table is assigned a SQL expression. Column names may be referenced in the expression by prefixing the column name with a colon character (**:**).

The following is the corresponding data file:

```

9101ROGERS CLERK 790217-DEC-10 1980.0020
9102PETERSON SALESMAN 769820-DEC-10 2600.0030 2300.00
9103WARREN SALESMAN 769822-DEC-10 5250.0030 2500.00
9104JONES, JR. MANAGER 783902-APR-09 7975.0020
9105ARNOLDS CLERK 778213-SEP-10 3750.0010
9106JACKSON ANALYST 756603-JAN-11 4500.0040

```

Since the records for employees **ARNOLDS** and **JACKSON** contain **10** and **40** in columns 47 thru 48, which do not satisfy any of the **WHEN** clauses, EDB*Loader writes these two records to the discard file, **emp_multitbl.dsc**, whose content is shown by the following:

```

9105ARNOLDS CLERK 778213-SEP-10 3750.0010
9106JACKSON ANALYST 756603-JAN-11 4500.0040

```

The following are the rows loaded into the **emp_research** and **emp_sales** tables:

```
SELECT * FROM emp_research;
```

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|------------|---------|------|--------------------|---------|-------|--------|
| 9101 | ROGERS | CLERK | 7902 | 17-DEC-10 00:00:00 | 1980.00 | 20.00 | |
| 9104 | JONES, JR. | MANAGER | 7839 | 02-APR-09 00:00:00 | 7975.00 | 20.00 | |

(2 rows)

```
SELECT * FROM emp_sales;
```

| empno | ename | job | mgr | hiredate | sal | comm | deptno |
|-------|----------|----------|------|--------------------|---------|---------|--------|
| 9102 | PETERSON | SALESMAN | 7698 | 20-DEC-10 00:00:00 | 2600.00 | 2950.00 | 30.00 |
| 9103 | WARREN | SALESMAN | 7698 | 22-DEC-10 00:00:00 | 5250.00 | 3813.00 | 30.00 |

(2 rows)

Invoking EDB Loader

You must have superuser privileges to run EDB*Loader. Use the following command to invoke EDB*Loader from the command line:

```
edbldr [ -d <dbname> ] [ -p <port> ] [ -h <host> ]
[ USERID={ <username/password> | <username>/ | <username> | / } ]
[ { -c | connstr= } <CONNECTION_STRING> ]
  CONTROL=<control_file>
[ DATA=<data_file> ]
[ BAD=<bad_file>]
[ DISCARD=<discard_file> ]
[ DISCARDMAX=<max_discard_recs> ]
[ HANDLE_CONFLICTS={ FALSE | TRUE } ]
[ LOG=<log_file> ]
[ PARFILE=<param_file> ]
[ DIRECT={ FALSE | TRUE } ]
[ FREEZE={ FALSE | TRUE } ]
[ ERRORS=<error_count> ]
[ PARALLEL={ FALSE | TRUE } ]
[ ROWS=<n> ]
[ SKIP=<skip_count> ]
[ SKIP_INDEX_MAINTENANCE={ FALSE | TRUE } ]
[edb_resource_group=<group_name> ]
```

Description

If the `-d` option, the `-p` option, or the `-h` option are omitted, the defaults for the database, port, and host are determined according to the same rules as other Advanced Server utility programs such as `edb-psql`, for example.

Any parameter listed in the preceding syntax diagram except for the `-d` option, `-p` option, `-h` option, and the `PARFILE` parameter may be specified in a *parameter file*. The parameter file is specified on the command line when `edbldr` is invoked using `PARFILE=param_file`. Some parameters may be specified in the `OPTIONS` clause in the control file. For more information on the control file, see [Building the EDB*Loader Control File](#).

The specification of `control_file`, `data_file`, `bad_file`, `discard_file`, `log_file`, and `param_file` may include the full directory path or a relative directory path to the file name. If the file name is specified alone or with a relative directory path, the file is assumed to exist (in the case of `control_file`, `data_file`, or `param_file`), or to be created (in the case of `bad_file`, `discard_file`, or `log_file`) relative to the current working directory from which `edbldr` is invoked.

Note

The control file must exist in the character set encoding of the client where `edbldr` is invoked. If the client is in a different encoding than the database encoding, then the `PGCLIENTENCODING` environment variable must be set on the client to the client's encoding prior to invoking `edbldr`. This must be done to ensure character set conversion is properly done between the client and the database server.

The operating system account used to invoke `edbldr` must have read permission on the directories and files specified by `control_file`, `data_file`, and `param_file`.

The operating system account `enterprisedb` must have write permission on the directories where `bad_file`, `discard_file`, and `log_file` are to be written.

Note

The file names for `control_file`, `data_file`, `bad_file`, `discard_file`, and `log_file` should include extensions of `.ctl`, `.dat`, `.bad`, `.dsc`, and `.log`, respectively. If the provided file name does not contain an extension, EDB*Loader assumes the actual file name includes the appropriate aforementioned extension.

Parameters

dbname

Name of the database containing the tables to be loaded.

port

Port number on which the database server is accepting connections.

host

IP address of the host on which the database server is running.

USERID={ username/password | username/ | username | / }

EDB*Loader connects to the database with `username`. `username` must be a superuser. `password` is the password for `username`.

If the `USERID` parameter is omitted, EDB*Loader prompts for `username` and `password`. If `USERID=username/` is specified, then EDB*Loader 1) uses the password file specified by environment variable `PGPASSFILE` if `PGPASSFILE` is set, or 2) uses the `.pqpass` password file (`pgpass.conf` on Windows systems) if `PGPASSFILE` is not set. If `USERID=username` is specified, then EDB*Loader prompts for `password`. If `USERID=/` is specified, the connection is attempted using the operating system account as the user name.

Note

The Advanced Server connection environment variables `PGUSER` and `PGPASSWORD` are ignored by EDB*Loader. See the PostgreSQL core documentation for information on the `PGPASSFILE` environment variable and the password file.

-c CONNECTION_STRING

connstr=CONNECTION_STRING

The `-c` or `connstr=` option allows you to specify all the connection parameters supported by libpq. With this option, SSL connection parameters or other connection parameters supported by libpq can also be specified. If connection options such as `-d`, `-h`, `-p` or `userid=dbuser/dbpass` are provided separately, they may override the values provided via the `-c` or `connstr=` option.

CONTROL=control_file

`control_file` specifies the name of the control file containing EDB*Loader directives. If a file extension is not specified, an extension of `.ctl` is assumed.

For more information on the control file, see [Building the EDB*Loader Control File](#).

DATA=data_file

`data_file` specifies the name of the file containing the data to be loaded into the target table. If a file extension is not specified, an extension of `.dat` is assumed. Specifying a `data_file` on the command line overrides the `INFILE` clause specified in the control file.

For more information about `data_file`, see [Building the EDB*Loader Control File](#).

BAD=bad_file

`bad_file` specifies the name of a file that receives input data records that cannot be loaded due to errors. Specifying a `bad_file` on the command line overrides any `BADFILE` clause specified in the control file.

For more information about `bad_file`, see [Building the EDB*Loader Control File](#).

DISCARD=discard_file

`discard_file` is the name of the file that receives input data records that do not meet any table's selection criteria. Specifying a `discard_file` on the command line overrides the `DISCARDFILE` clause in the control file.

For more information about `discard_file`, see [Building the EDB*Loader Control File](#).

DISCARDMAX=max_discard_recs

`max_discard_recs` is the maximum number of discarded records that may be encountered from the input data records before terminating the EDB*Loader session. Specifying `max_discard_recs` on the command line overrides the `DISCARDMAX` or `DISCARDS` clause in the control file.

For more information about `max_discard_recs`, see [Building the EDB*Loader Control File](#).

HANDLE_CONFLICTS={ FALSE | TRUE }

If any record insertion fails due to a unique constraint violation, EDB*Loader will abort the entire operation. You can instruct EDB*Loader to instead move the duplicate record to the `BAD` file and continue processing by setting `HANDLE CONFLICTS` to `TRUE`. This behavior will only apply if indexes are present. By default, `HANDLE_CONFLICTS` is set to `FALSE`.

`HANDLE_CONFLICTS` set to `TRUE` is not supported with direct path loading; if set to `TRUE` when direct path loading, EDB*Loader will throw an error.

LOG=log_file

`log_file` specifies the name of the file in which EDB*Loader records the results of the EDB*Loader session.

If the `LOG` parameter is omitted, EDB*Loader creates a log file with the name `control_file_base.log` in the directory from which `edbldr` is invoked. `control_file_base` is the base name of the control file used in the EDB*Loader session. The operating system account `enterprisedb` must have write permission on the directory where the log file is to be written.

PARFILE=param_file

`param_file` specifies the name of the file that contains command line parameters for the EDB*Loader session. Any command line parameter listed in this section except for the `-d`, `-p`, and `-h` options, and the `PARFILE` parameter itself, can be specified in `param_file` instead of on the command line.

Any parameter given in `param_file` overrides the same parameter supplied on the command line before the `PARFILE` option. Any parameter given on the command line that appears after the `PARFILE` option overrides the same parameter given in `param_file`.

Note

Unlike other EDB*Loader files, there is no default file name or extension assumed for `param_file`, though by Oracle SQL*Loader convention, `.par` is typically used, but not required, as an extension.

DIRECT= { FALSE | TRUE }

If `DIRECT` is set to `TRUE` EDB*Loader performs a direct path load instead of a conventional path load. The default value of `DIRECT` is `FALSE`.

You should not set `DIRECT=true` when loading the data into a replicated table. If you are using EDB*Loader to load data into a replicated table and set `DIRECT=true`, indexes may omit rows that are in a table or may potentially contain references to rows that have been deleted. EnterpriseDB does not support direct inserts to load data into replicated tables.

For information about direct path loads, see [Direct Path Load](#).

FREEZE= { FALSE | TRUE }

Set **FREEZE** to **TRUE** to indicate that the data should be copied with the rows **frozen**. A tuple guaranteed to be visible to all current and future transactions is marked as frozen to prevent transaction ID wrap-around. For more information about frozen tuples, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/routine-vacuuming.html>

You must specify a data-loading type of **TRUNCATE** in the control file when using the **FREEZE** option. **FREEZE** is not supported for direct loading.

By default, **FREEZE** is **FALSE**.

ERRORS=error_count

error_count specifies the number of errors permitted before aborting the EDB*Loader session. The default is **50**.

PARALLEL= { FALSE | TRUE }

Set **PARALLEL** to **TRUE** to indicate that this EDB*Loader session is one of a number of concurrent EDB*Loader sessions participating in a parallel direct path load. The default value of **PARALLEL** is **FALSE**.

When **PARALLEL** is **TRUE**, the **DIRECT** parameter must also be set to **TRUE**.

For more information about parallel direct path loads, see [Parallel Direct Path Load](#).

ROWS=n

n specifies the number of rows that EDB*Loader will commit before loading the next set of **n** rows.

SKIP=skip_count

Number of records at the beginning of the input data file that should be skipped before loading begins. The default is **0**.

SKIP_INDEX_MAINTENANCE= { FALSE | TRUE }

If set to **TRUE**, index maintenance is not performed as part of a direct path load, and indexes on the loaded table are marked as invalid. The default value of **SKIP_INDEX_MAINTENANCE** is **FALSE**.

During a parallel direct path load, target table indexes are not updated, and are marked as invalid after the load is complete.

You can use the **REINDEX** command to rebuild an index. For more information about the **REINDEX** command, see the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-reindex.html>

edb_resource_group=group_name

group_name specifies the name of an EDB Resource Manager resource group to which the EDB*Loader session is to be assigned.

Any default resource group that may have been assigned to the session (for example, a database user running the EDB*Loader session who had been assigned a default resource group with the **ALTER ROLE ... SET edb_resource_group** command) is overridden by the resource group given by the **edb_resource_group** parameter specified on the **edbldr** command line.

Examples

In the following example EDB*Loader is invoked using a control file named **emp.ctl** located in the current working

directory to load a table in database `edb`:

```
$ /usr/edb/as13/bin/edbldr -d edb USERID=enterprisedb/password
CONTROL=emp.ctl
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

Successfully loaded (4) records

In the following example, EDB*Loader prompts for the user name and password since they are omitted from the command line. In addition, the files for the bad file and log file are specified with the `BAD` and `LOG` command line parameters.

```
$ /usr/edb/as13/bin/edbldr -d edb CONTROL=emp.ctl BAD=/tmp/emp.bad
LOG=/tmp/emp.log
Enter the user name : enterprisedb
Enter the password :
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

Successfully loaded (4) records

The following example runs EDB*Loader with the same parameters as shown in the preceding example, but using a parameter file located in the current working directory. The `SKIP` and `ERRORS` parameters are altered from their defaults in the parameter file as well. The parameter file, `emp.par`, contains the following:

```
CONTROL=emp.ctl
BAD=/tmp/emp.bad
LOG=/tmp/emp.log
SKIP=1
ERRORS=10
```

EDB*Loader is invoked with the parameter file as shown by the following:

```
$ /usr/edb/as13/bin/edbldr -d edb PARFILE=emp.par
Enter the user name : enterprisedb
Enter the password :
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

Successfully loaded (3) records

In the following example EDB*Loader is invoked using a `connstr=` option using `emp.ctl` control file located in the current working directory to load a table in a database named `edb`:

```
$ /usr/edb/as13/bin/edbldr connstr=\"sslmode=verify-ca sslcompression=0
host=127.0.0.1 dbname=edb port=5444 user=enterprisedb\" CONTROL=emp.ctl
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

Successfully loaded (4) records

Exit Codes

When EDB*Loader exits, it will return one of the following codes:

| Exit Code | Description |
|-----------|-------------|
|-----------|-------------|

| Exit Code | Description |
|-----------|--|
| 0 | Indicates that all rows loaded successfully. |
| 1 | Indicates that EDB*Loader encountered command line or syntax errors, or aborted the load operation due to an unrecoverable error. |
| 2 | Indicates that the load completed, but some (or all) rows were rejected or discarded. |
| 3 | Indicates that EDB*Loader encountered fatal errors (such as OS errors). This class of errors is equivalent to the FATAL or PANIC severity levels of PostgreSQL errors. |

Direct Path Load

During a direct path load, EDB*Loader writes the data directly to the database pages, which is then synchronized to disk. The insert processing associated with a conventional path load is bypassed, thereby resulting in a performance improvement.

Bypassing insert processing reduces the types of constraints that may exist on the target table. The following types of constraints are permitted on the target table of a direct path load:

- Primary key
- Not null constraints
- Indexes (unique or non-unique)

The restrictions on the target table of a direct path load are the following:

- Triggers are not permitted
- Check constraints are not permitted
- Foreign key constraints on the target table referencing another table are not permitted
- Foreign key constraints on other tables referencing the target table are not permitted
- The table must not be partitioned
- Rules may exist on the target table, but they are not executed

Note

Currently, a direct path load in EDB*Loader is more restrictive than in Oracle SQL*Loader. The preceding restrictions do not apply to Oracle SQL*Loader in most cases. The following restrictions apply to a control file used in a direct path load:

- Multiple table loads are not supported. That is, only one `INTO TABLE` clause may be specified in the control file.
- SQL expressions may not be used in the data field definitions of the `INTO TABLE` clause.
- The `FREEZE` option is not supported for direct path loading.

To run a direct path load, add the `DIRECT=TRUE` option as shown by the following example:

```
$ /usr/edb/as13/bin/edbldr -d edb USERID=enterprisedb/password
CONTROL=emp.ctl DIRECT=TRUE
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

Successfully loaded (4) records

Parallel Direct Path Load

The performance of a direct path load can be further improved by distributing the loading process over two or more sessions running concurrently. Each session runs a direct path load into the same table.

Since the same table is loaded from multiple sessions, the input records to be loaded into the table must be divided amongst several data files so that each EDB*Loader session uses its own data file and the same record is not loaded more than once into the table.

The target table of a parallel direct path load is under the same restrictions as a direct path load run in a single session.

The restrictions on the target table of a direct path load are the following:

- Triggers are not permitted
- Check constraints are not permitted
- Foreign key constraints on the target table referencing another table are not permitted
- Foreign key constraints on other tables referencing the target table are not permitted
- The table must not be partitioned
- Rules may exist on the target table, but they are not executed

In addition, the `APPEND` clause must be specified in the control file used by each EDB*Loader session.

To run a parallel direct path load, run EDB*Loader in a separate session for each participant of the parallel direct path load. Invocation of each such EDB*Loader session must include the `DIRECT=TRUE` and `PARALLEL=TRUE` parameters.

Each EDB*Loader session runs as an independent transaction so if one of the parallel sessions aborts and rolls back its changes, the loading done by the other parallel sessions are not affected.

Note: In a parallel direct path load, each EDB*Loader session reserves a fixed number of blocks in the target table in a round-robin fashion. Some of the blocks in the last allocated chunk may not be used, and those blocks remain uninitialized. A subsequent use of the `VACUUM` command on the target table may show warnings regarding these uninitialized blocks such as the following:

```
WARNING: relation "emp" page 98264 is uninitialized --- fixing
```

```
WARNING: relation "emp" page 98265 is uninitialized --- fixing
```

```
WARNING: relation "emp" page 98266 is uninitialized --- fixing
```

This is an expected behavior and does not indicate data corruption.

Indexes on the target table are not updated during a parallel direct path load and are therefore marked as invalid after the load is complete. You must use the `REINDEX` command to rebuild the indexes.

The following example shows the use of a parallel direct path load on the `emp` table.

Note

If you attempt a parallel direct path load on the sample `emp` table provided with Advanced Server, you must first remove the triggers and constraints referencing the `emp` table. In addition the primary key column, `empno`, was expanded from `NUMBER(4)` to `NUMBER` in this example to allow for the insertion of a larger number of rows.

The following is the control file used in the first session:

```
LOAD DATA
INFILE  '/home/user/loader/emp_parallel_1.dat'
APPEND
INTO TABLE emp
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY ""
TRAILING NULLCOLS
(
    empno,
    ename,
```

```
job,
mgr,
hiredate,
sal,
deptno,
comm
)
```

The **APPEND** clause must be specified in the control file for a parallel direct path load.

The following shows the invocation of EDB*Loader in the first session. The **DIRECT=TRUE** and **PARALLEL=TRUE** parameters must be specified.

```
$ /usr/edb/as13/bin/edbldr -d edb USERID=enterprisedb/password
CONTROL=emp_parallel_1.ctl DIRECT=TRUE PARALLEL=TRUE
WARNING: index maintenance will be skipped with PARALLEL load
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

The control file used for the second session appears as follows. Note that it is the same as the one used in the first session, but uses a different data file.

```
LOAD DATA
INFILE '/home/user/loader/emp_parallel_2.dat'
APPEND
INTO TABLE emp
FIELDS TERMINATED BY ';' OPTIONALLY ENCLOSED BY ""
TRAILING NULLCOLS
(
empno,
ename,
job,
mgr,
hiredate,
sal,
deptno,
comm
)
```

The preceding control file is used in a second session as shown by the following:

```
$ /usr/edb/as13/bin/edbldr -d edb USERID=enterprisedb/password
CONTROL=emp_parallel_2.ctl DIRECT=TRUE PARALLEL=TRUE
WARNING: index maintenance will be skipped with PARALLEL load
EDB Loader: Copyright (c) 2007-2020, EnterpriseDB Corporation.
```

EDB*Loader displays the following message in each session when its respective load operation completes:

```
Successfully loaded (10000) records
```

The following query shows that the index on the emp table has been marked as **INVALID**:

```
SELECT index_name, status FROM user_indexes WHERE table_name = 'EMP';
index_name | status
-----+-----
EMP_PK    | INVALID
```

(1 row)

Note

`user_indexes` is the view of indexes compatible with Oracle databases owned by the current user.

Queries on the `emp` table will not utilize the index unless it is rebuilt using the `REINDEX` command as shown by the following:

```
REINDEX INDEX emp_pk;
```

A subsequent query on `user_indexes` shows that the index is now marked as `VALID`:

```
SELECT index_name, status FROM user_indexes WHERE table_name = 'EMP';
```

| index_name | status |
|------------|--------|
| EMP_PK | VALID |

(1 row)

Remote Loading

EDB*Loader supports a feature called *remote loading*. In remote loading, the database containing the table to be loaded is running on a database server on a different host than from where EDB*Loader is invoked with the input data source.

This feature is useful if you have a large amount of data to be loaded, and you do not want to create a large data file on the host running the database server.

In addition, you can use the standard input feature to pipe the data from the data source such as another program or script, directly to EDB*Loader, which then loads the table in the remote database. This bypasses the process of having to create a data file on disk for EDB*Loader.

Performing remote loading along with using standard input requires the following:

- The `edbldr` program must be installed on the client host on which it is to be invoked with the data source for the EDB*Loader session.
- The control file must contain the clause `INFILE 'stdin'` so you can pipe the data directly into EDB*Loader's standard input. For more information, see [Building the EDB*Loader Control File](#) for information on the `INFILE` clause and the EDB*Loader control file.
- All files used by EDB*Loader such as the control file, bad file, discard file, and log file must reside on, or are created on, the client host on which `edbldr` is invoked.
- When invoking EDB*Loader, use the `-h` option to specify the IP address of the remote database server. For more information, see [Invoking EDB*Loader](#) for information on invoking EDB*Loader.

- Use the operating system pipe operator `(|)` or input redirection operator `(<)` to supply the input data to EDB*Loader.

The following example loads a database running on a database server at `192.168.1.14` using data piped from a source named `datasource`.

```
datasource | ./edbldr -d edb -h 192.168.1.14 USERID=enterprisedb/password
CONTROL=remote.ctl
```

The following is another example of how standard input can be used:

```
./edbldr -d edb -h 192.168.1.14 USERID=enterprisedb/password
CONTROL=remote.ctl < datasource
```

Updating a Table with a Conventional Path Load

You can use EDB*Loader with a conventional path load to update the rows within a table, merging new data with the existing data. When you invoke EDB*Loader to perform an update, the server searches the table for an existing row with a matching primary key:

- If the server locates a row with a matching key, it replaces the existing row with the new row.
- If the server does not locate a row with a matching key, it adds the new row to the table.

To use EDB*Loader to update a table, the table must have a primary key. Please note that you cannot use EDB*Loader to [UPDATE](#) a partitioned table.

To perform an [UPDATE](#), use the same steps as when performing a conventional path load:

1. Create a data file that contains the rows you wish to [UPDATE](#) or [INSERT](#).
2. Define a control file that uses the [INFILE](#) keyword to specify the name of the data file. For information about building the EDB*Loader control file, see [Building the EDB*Loader Control File](#).
3. Invoke EDB*Loader, specifying the database name, connection information, and the name of the control file. For information about invoking EDB*Loader, see [Invoking EDB*Loader](#).

The following example uses the [emp](#) table that is distributed with the Advanced Server sample data. By default, the table contains:

```
edb=# select * from emp;
empno|ename | job | mgr |      hiredate | sal | comm | deptno
-----+-----+-----+-----+-----+-----+
7369 |SMITH |CLERK | 7902 | 17-DEC-80 00:00:00 | 800.00 |   | 20
7499 |ALLEN |SALESMAN | 7698 | 20-FEB-81 00:00:00 | 1600.00 |300.00 | 30
7521 |WARD |SALESMAN | 7698 | 22-FEB-81 00:00:00 | 1250.00 |500.00 | 30
7566 |JONES |MANAGER | 7839 | 02-APR-81 00:00:00 | 2975.00 |   | 20
7654 |MARTIN|SALESMAN | 7698 | 28-SEP-81 00:00:00 | 1250.00 |1400.00| 30
7698 |BLAKE |MANAGER | 7839 | 01-MAY-81 00:00:00 | 2850.00 |   | 30
7782 |CLARK |MANAGER | 7839 | 09-JUN-81 00:00:00 | 2450.00 |   | 10
7788 |SCOTT |ANALYST | 7566 | 19-APR-87 00:00:00 | 3000.00 |   | 20
7839 |KING |PRESIDENT|  | 17-NOV-81 00:00:00 | 5000.00 |   | 10
7844 |TURNER|SALESMAN | 7698 | 08-SEP-81 00:00:00 | 1500.00 | 0.00 | 30
7876 |ADAMS |CLERK | 7788 | 23-MAY-87 00:00:00 | 1100.00 |   | 20
7900 |JAMES |CLERK | 7698 | 03-DEC-81 00:00:00 | 950.00 |   | 30
7902 |FORD |ANALYST | 7566 | 03-DEC-81 00:00:00 | 3000.00 |   | 20
7934 |MILLER|CLERK | 7782 | 23-JAN-82 00:00:00 | 1300.00 |   | 10
(14 rows)
```

The following control file ([emp_update.ctl](#)) specifies the fields in the table in a comma-delimited list. The control file performs an [UPDATE](#) on the [emp](#) table:

```
LOAD DATA
INFILE 'emp_update.dat'
BADFILE 'emp_update.bad'
DISCARDFILE 'emp_update.dsc'
UPDATE INTO TABLE emp
FIELDS TERMINATED BY ","
(empno, ename, job, mgr, hiredate, sal, comm, deptno)
```

The data that is being updated or inserted is saved in the [emp_update.dat](#) file. [emp_update.dat](#) contains:

```
7521,WARD,MANAGER,7839,22-FEB-81 00:00:00,3000.00,0.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,3500.00,0.00,20
7903,BAKER,SALESMAN,7521,10-JUN-13 00:00:00,1800.00,500.00,20
```

```
7904,MILLS,SALESMAN,7839,13-JUN-13 00:00:00,1800.00,500.00,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1500.00,400.00,30
```

Invoke EDB*Loader, specifying the name of the database (`edb`), the name of a database superuser (and their associated password) and the name of the control file (`emp_update.ctl`):

```
| edbldr -d edb userid=user_name/password control=emp_update.ctl
```

After performing the update, the `emp` table contains:

```
edb=# select * from emp;
empno|ename | job | mgr |    hiredate   | sal | comm | deptno
-----+-----+-----+-----+-----+-----+
7369 |SMITH |CLERK | 7902 | 17-DEC-80 00:00:00 | 800.00 |      | 20
7499 |ALLEN |SALESMAN | 7698 | 20-FEB-81 00:00:00 | 1600.00 | 300.00 | 30
7521 |WARD |MANAGER | 7839 | 22-FEB-81 00:00:00 | 3000.00 | 0.00 | 30
7566 |JONES |MANAGER | 7839 | 02-APR-81 00:00:00 | 3500.00 | 0.00 | 20
7654 |MARTIN|SALESMAN | 7698 | 28-SEP-81 00:00:00 | 1500.00 | 400.00 | 30
7698 |BLAKE |MANAGER | 7839 | 01-MAY-81 00:00:00 | 2850.00 |      | 30
7782 |CLARK |MANAGER | 7839 | 09-JUN-81 00:00:00 | 2450.00 |      | 10
7788 |SCOTT |ANALYST | 7566 | 19-APR-87 00:00:00 | 3000.00 |      | 20
7839 |KING |PRESIDENT|      | 17-NOV-81 00:00:00 | 5000.00 |      | 10
7844 |TURNER|SALESMAN | 7698 | 08-SEP-81 00:00:00 | 1500.00 | 0.00 | 30
7876 |ADAMS |CLERK | 7788 | 23-MAY-87 00:00:00 | 1100.00 |      | 20
7900 |JAMES |CLERK | 7698 | 03-DEC-81 00:00:00 | 950.00 |      | 30
7902 |FORD |ANALYST | 7566 | 03-DEC-81 00:00:00 | 3000.00 |      | 20
7903 |BAKER |SALESMAN | 7521 | 10-JUN-13 00:00:00 | 1800.00 | 500.00 | 20
7904 |MILLS |SALESMAN | 7839 | 13-JUN-13 00:00:00 | 1800.00 | 500.00 | 20
7934 |MILLER|CLERK | 7782 | 23-JAN-82 00:00:00 | 1300.00 |      | 10
(16 rows)
```


The rows containing information for the three employees that are currently in the `emp` table are updated, while rows are added for the new employees (`BAKER` and `MILLS`)

8.3 EDB Wrap

The EDB Wrap utility protects proprietary source code and programs (functions, stored procedures, triggers, and packages) from unauthorized scrutiny. The EDB Wrap program translates a file that contains SPL or PL/pgSQL source code (the plaintext) into a file that contains the same code in a form that is nearly impossible to read. Once you have the obfuscated form of the code, you can send that code to the PostgreSQL server and the server will store those programs in obfuscated form. While EDB Wrap does obscure code, table definitions are still exposed.

Everything you wrap is stored in obfuscated form. If you wrap an entire package, the package body source, as well as the prototypes contained in the package header and the functions and procedures contained in the package body are stored in obfuscated form.

If you wrap a `CREATE PACKAGE` statement, you hide the package API from other developers. You may want to wrap the package body, but not the package header so users can see the package prototypes and other public variables that are defined in the package body. To allow users to see what prototypes the package contains, use EDBWrap to obfuscate only the `CREATE PACKAGE BODY` statement in the `edbwrap` input file, omitting the `CREATE PACKAGE` statement. The package header source will be stored plaintext, while the package body source and package functions and procedures will be stored obfuscated.



Once wrapped, source code and programs cannot be unwrapped or debugged. Reverse engineering is possible, but would be very difficult.

The entire source file is wrapped into one unit. Any `psql` meta-commands included in the wrapped file will not be recognized when the file is executed; executing an obfuscated file that contains a `psql` meta-command will cause a syntax error. `edbwrap` does not validate SQL source code -if the plaintext form contains a syntax error, `edbwrap` will not complain. Instead, the server will report an error and abort the entire file when you try to execute the obfuscated form.

Using EDB*Wrap to Obfuscate Source Code

EDB Wrap is a command line utility; it accepts a single input source file, obfuscates the contents and returns a single output file. When you invoke the `edbwrap` utility, you must provide the name of the file that contains the source code to obfuscate. You may also specify the name of the file where `edbwrap` will write the obfuscated form of the code. `edbwrap` offers three different command-line styles. The first style is compatible with Oracle's `wrap` utility:

```
edbwrap iname=<input_file> [oname= <output_file>]
```

The `iname=input_file` argument specifies the name of the input file; if `input_file` does not contain an extension, `edbwrap` will search for a file named `input_file.sql`

The `oname=output_file` argument (which is optional) specifies the name of the output file; if `output_file` does not contain an extension, `edbwrap` will append `.plb` to the name.

If you do not specify an output file name, `edbwrap` writes to a file whose name is derived from the input file name: `edbwrap` strips the suffix (typically `.sql`) from the input file name and adds `.plb`.

`edbwrap` offers two other command-line styles that may feel more familiar:

```
edbwrap --iname <input_file> [--oname <output_file>]  
edbwrap -i <input_file> [-o <output_file>]
```

You may mix command-line styles; the rules for deriving input and output file names are identical regardless of which style you use.

Once `edbwrap` has produced a file that contains obfuscated code, you typically feed that file into the PostgreSQL server using a client application such as `edb-psql`. The server executes the obfuscated code line by line and stores the source code for SPL and PL/pgSQL programs in wrapped form.

In summary, to obfuscate code with EDB Wrap, you:

1. Create the source code file.
2. Invoke EDB Wrap to obfuscate the code.
3. Import the file as if it were in plaintext form.

The following sequence demonstrates `edbwrap` functionality.

First, create the source code for the `list_emp` procedure (in plaintext form):

```
[bash] cat listemp.sql  
CREATE OR REPLACE PROCEDURE list_emp
```

```

IS
v_empno      NUMBER(4);
v_ename       VARCHAR2(10);
CURSOR emp_cur IS
  SELECT empno, ename FROM emp ORDER BY empno;
BEGIN
  OPEN emp_cur;
  DBMS_OUTPUT.PUT_LINE('EMPNO  ENAME');
  DBMS_OUTPUT.PUT_LINE('----- -----');
  LOOP
    FETCH emp_cur INTO v_empno, v_ename;
    EXIT WHEN emp_cur%NOTFOUND;
    DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || v_ename);
  END LOOP;
  CLOSE emp_cur;
END;
/

```

You can import the `list_emp` procedure with a client application such as `edb-psql`:

```

[bash] edb-psql edb
Welcome to edb-psql 8.4.3.2, the EnterpriseDB interactive terminal.
Type: \copyright for distribution terms
  \h for help with SQL commands
  \? for help with edb-psql commands
  \g or terminate with semicolon to execute query
  \q to quit
edb=# \i listemp.sql
CREATE PROCEDURE

```

You can view the plaintext source code (stored in the server) by examining the `pg_proc` system table:

```

edb=# SELECT prosrc FROM pg_proc WHERE proname = 'list_emp';
      prosrc
-----
```

```

v_empno      NUMBER(4);
v_ename       VARCHAR2(10);
CURSOR emp_cur IS
  SELECT empno, ename FROM emp ORDER BY empno;
BEGIN
  OPEN emp_cur;
  DBMS_OUTPUT.PUT_LINE('EMPNO  ENAME');
  DBMS_OUTPUT.PUT_LINE('----- -----');
  LOOP
    FETCH emp_cur INTO v_empno, v_ename;
    EXIT WHEN emp_cur%NOTFOUND;
    DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || v_ename);
  END LOOP;
  CLOSE emp_cur;
END
(1 row)

```

```
edb=# quit
```

Next, obfuscate the plaintext file with EDB Wrap:

```
[bash] edbwrap -i listemp.sql
EDB*Wrap Utility: Release 8.4.3.2
```

Copyright (c) 2004-2020 EnterpriseDB Corporation. All Rights Reserved.

Using encoding UTF8 for input
Processing listemp.sql to listemp.plb

Examining the contents of the output file (listemp.plb) file reveals that the code is obfuscated:

```
[bash] cat listemp.plb
$__EDBwrapped__$  
UTF8  
d+6DL30RVaGjYMIzkuoSzAQgtBw7MhYFuAFkBsfYfhdJ0rjwBv+bHr1FCyH6j9SgH  
movU+bYI+jR+hR2jbzq3sovHKEyZIp9y3/GckbQgualRhIlGpyWfE0dltDUpkYRLN  
/OUXmk0/P4H6EI98sAHevGDhOWI+58DjJ44qhZ+l5NNEVxbWDztpb/s5sdx4660qQ  
Ozx3/gh8VkjS2JbcxYMpjmrwVr6fAXfb68MI9mW2Hl7fNtxcb5kjSzXvfWR2XYzJf  
KFNrEhbL1DTVISEC5wE6IGlwhYvXOf22m1R2IFns0MtF9fwcnBWAs1YqjR00j6+fc  
er/f/efAFh4=
$__EDBwrapped__$
```

You may notice that the second line of the wrapped file contains an encoding name (in this case, the encoding is UTF8). When you obfuscate a file, edbwrap infers the encoding of the input file by examining the locale. For example, if you are running edbwrap while your locale is set to en_US.utf8, edbwrap assumes that the input file is encoded in UTF8. Be sure to examine the output file after running edbwrap; if the locale contained in the wrapped file does not match the encoding of the input file, you should change your locale and rewrap the input file.

You can import the obfuscated code into the PostgreSQL server using the same tools that work with plaintext code:

```
[bash] edb-psql edb
Welcome to edb-psql 8.4.3.2, the EnterpriseDB interactive terminal.
Type: \copyright for distribution terms
      \h for help with SQL commands
      \? for help with edb-psql commands
      \g or terminate with semicolon to execute query
      \q to quit
```

```
edb=# \i listemp.plb
CREATE PROCEDURE
```

Now, the pg_proc system table contains the obfuscated code:

```
edb=# SELECT prosrc FROM pg_proc WHERE proname = 'list_emp';
      prosrc
-----
$__EDBwrapped__$  
UTF8  
dw4B9Tz69J3WOsy0GgYJQa+G2sLZ3IOyxS8pDyuOTFuiYe/ExiEatwwG3h3tdJk  
ea+Alp35dS/4idbN8wpegM3s994dQ3R97NgNHfvTQnO2vt4wQtsQ/Zc4v4Lhfj  
nIV+A4UpHl5oQEnXeAch2LcRD87hkU0uo1ESeQV8IrXaj9BsZr+ueROnwhGs/Ec  
pva/tRV4m9RusFn0wyr38u4Z8w4dfnPW184Y3o6lt4b3aH07WxTkWrMLmOZW1jJ  
Nu6u4o+ezO64G9QKPazgehslv4JB9NQnuocActfDSPMY7R7anmgw
$__EDBwrapped__$  
(1 row)
```

Invoke the obfuscated code in the same way that you would invoke the plaintext form:

```
edb=# exec list_emp;
```

| EMPNO | ENAME |
|-------|--------|
| 7369 | SMITH |
| 7499 | ALLEN |
| 7521 | WARD |
| 7566 | JONES |
| 7654 | MARTIN |
| 7698 | BLAKE |
| 7782 | CLARK |
| 7788 | SCOTT |
| 7839 | KING |
| 7844 | TURNER |
| 7876 | ADAMS |
| 7900 | JAMES |
| 7902 | FORD |
| 7934 | MILLER |

When you use `pg_dump` to back up a database, wrapped programs remain obfuscated in the archive file.

Be aware that audit logs produced by the Postgres server will show wrapped programs in plaintext form. Source code is also displayed in plaintext in SQL error messages generated during the execution of a program.

Note

At this time, the bodies of the objects created by the following statements will not be stored in obfuscated form:

```
CREATE [OR REPLACE] TYPE type_name AS OBJECT
CREATE [OR REPLACE] TYPE type_name UNDER type_name
CREATE [OR REPLACE] TYPE BODY type_name
```

8.4 Dynamic Runtime Instrumentation Tools Architecture (DRITA)

The Dynamic Runtime Instrumentation Tools Architecture (DRITA) allows a DBA to query catalog views to determine the *wait events* that affect the performance of individual sessions or the system as a whole. DRITA records the number of times each event occurs as well as the time spent waiting; you can use this information to diagnose performance problems. DRITA offers this functionality, while consuming minimal system resources.

DRITA compares *snapshots* to evaluate the performance of a system. A snapshot is a saved set of system performance data at a given point in time. Each snapshot is identified by a unique ID number; you can use snapshot ID numbers with DRITA reporting functions to return system performance statistics.

Configuring and Using DRITA

Advanced Server's `postgresql.conf` file includes a configuration parameter named `timed_statistics` that controls the collection of timing data. The valid parameter values are `TRUE` or `FALSE`; the default value is `FALSE`.

This is a dynamic parameter which can be modified in the `postgresql.conf` file, or while a session is in progress. To enable DRITA, you must either:

Modify the `postgresql.conf` file, setting the `timed_statistics` parameter to `TRUE`.

or

Connect to the server with the EDB-PSQL client, and invoke the command:

```
SET timed_statistics = TRUE
```

After modifying the `timed_statistics` parameter, take a starting snapshot. A snapshot captures the current state of each timer and event counter. The server will compare the starting snapshot to a later snapshot to gauge system performance.

Use the `edbssnap()` function to take the beginning snapshot:

```
edb=# SELECT * FROM edbsnap();
edbssnap
```

```
-----  
Statement processed.  
(1 row)
```

Then, run the workload that you would like to evaluate; when the workload has completed (or at a strategic point during the workload), take another snapshot:

```
edb=# SELECT * FROM edbsnap();
edbssnap
```

```
-----  
Statement processed.  
(1 row)
```

You can capture multiple snapshots during a session. Then, use the DRITA functions and reports to manage and compare the snapshots to evaluate performance information.

DRITA Functions

You can use DRITA functions to gather wait information and manage snapshots. DRITA functions are fully supported by Advanced Server 13 whether your installation is made compatible with Oracle databases or is made in PostgreSQL-compatible mode.

`get_snaps()`

The `get_snaps()` function returns a list of the current snapshots. The signature is:

```
get_snaps()
```

The following example demonstrates using the `get_snaps()` function to display a list of snapshots:

```
SELECT * FROM get_snaps();
get_snaps
```

```

1 25-JUL-18 09:49:04.224597
2 25-JUL-18 09:49:09.310395
3 25-JUL-18 09:49:14.378728
4 25-JUL-18 09:49:19.448875
5 25-JUL-18 09:49:24.52103
6 25-JUL-18 09:49:29.586889
7 25-JUL-18 09:49:34.65529
8 25-JUL-18 09:49:39.723095
9 25-JUL-18 09:49:44.788392
10 25-JUL-18 09:49:49.855821
11 25-JUL-18 09:49:54.919954
12 25-JUL-18 09:49:59.987707
(12 rows)

```

The first column in the result list displays the snapshot identifier; the second column displays the date and time that the snapshot was captured.

sys_rpt()

The `sys_rpt()` function returns system wait information. The signature is:

```
sys_rpt(<beginning_id>, <ending_id>, <top_n>)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

This example demonstrates a call to the `sys_rpt()` function:

```

SELECT * FROM sys_rpt(9, 10, 10);
      sys_rpt
-----
WAIT NAME      COUNT    WAIT TIME      % WAIT
-----
wal flush      8359    1.357593      30.62
wal write       8358    1.349153      30.43
wal file sync   8358    1.286437      29.02
query plan     33439    0.439324      9.91
db file extend   54     0.000585      0.01
db file read     31     0.000307      0.01
other lwlock acquire  0     0.000000      0.00
ProcArrayLock     0     0.000000      0.00
CLogControlLock   0     0.000000      0.00
(11 rows)

```

The information displayed in the result set includes:

| Column Name | Description |
|-------------|---|
| WAIT NAME | The name of the wait. |
| COUNT | The number of times that the wait event occurred. |
| WAIT TIME | The time of the wait event in seconds. |
| % WAIT | The percentage of the total wait time used by this wait for this session. |

sess_rpt()

The `sess_rpt()` function returns session wait information. The signature is:

`sess_rpt(<beginning_id>, <ending_id>, <top_n>)`

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

The following example demonstrates a call to the `sess_rpt()` function:

```
SELECT * FROM sess_rpt(8, 9, 10);
```

| sess_rpt | | | | | | |
|-----------|------------|-----------------|-------|----------|------------|------------|
| ID | USER | WAIT NAME | COUNT | TIME | % WAIT SES | % WAIT ALL |
| 3501 | enterprise | wal flush | 8354 | 1.354958 | 30.61 | 30.61 |
| 3501 | enterprise | wal write | 8354 | 1.348192 | 30.46 | 30.46 |
| 3501 | enterprise | wal file sync | 8354 | 1.285607 | 29.04 | 29.04 |
| 3501 | enterprise | query plan | 33413 | 0.436901 | 9.87 | 9.87 |
| 3501 | enterprise | db file extend | 54 | 0.000578 | 0.01 | 0.01 |
| 3501 | enterprise | db file read | 56 | 0.000541 | 0.01 | 0.01 |
| 3501 | enterprise | ProcArrayLock | 0 | 0.000000 | 0.00 | 0.00 |
| 3501 | enterprise | CLogControlLock | 0 | 0.000000 | 0.00 | 0.00 |
| (10 rows) | | | | | | |

The information displayed in the result set includes:

| Column Name | Description |
|-------------------------|---|
| <code>ID</code> | The processID of the session. |
| <code>USER</code> | The name of the user incurring the wait. |
| <code>WAIT NAME</code> | The name of the wait event. |
| <code>COUNT</code> | The number of times that the wait event occurred. |
| <code>TIME</code> | The length of the wait event in seconds. |
| <code>% WAIT SES</code> | The percentage of the total wait time used by this wait for this session. |

% WAIT ALL The percentage of the total wait time used by this wait (for all sessions).

sessid_rpt()

The `sessid_rpt()` function returns session ID information for a specified backend. The signature is:

| `sessid_rpt(<beginning_id>, <ending_id>, <backend_id>)`

Parameters

`beginning_id`

| `beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

| `ending_id` is an integer value that represents the ending session identifier.

`backend_id`

| `backend_id` is an integer value that represents the backend identifier.

The following code sample demonstrates a call to `sessid_rpt()`:

```
SELECT * FROM sessid_rpt(8, 9, 3501);
```

| sessid_rpt | | | | | | |
|------------|------------|-----------------|-------|----------|------------|------------|
| ID | USER | WAIT NAME | COUNT | TIME | % WAIT SES | % WAIT ALL |
| 3501 | enterprise | CLogControlLock | 0 | 0.000000 | 0.00 | 0.00 |
| 3501 | enterprise | ProcArrayLock | 0 | 0.000000 | 0.00 | 0.00 |
| 3501 | enterprise | db file read | 56 | 0.000541 | 0.01 | 0.01 |
| 3501 | enterprise | db file extend | 54 | 0.000578 | 0.01 | 0.01 |
| 3501 | enterprise | query plan | 33413 | 0.436901 | 9.87 | 9.87 |
| 3501 | enterprise | wal file sync | 8354 | 1.285607 | 29.04 | 29.04 |
| 3501 | enterprise | wal write | 8354 | 1.348192 | 30.46 | 30.46 |
| 3501 | enterprise | wal flush | 8354 | 1.354958 | 30.61 | 30.61 |
| (10 rows) | | | | | | |

The information displayed in the result set includes:

| Column Name | Description |
|-------------------------|---|
| <code>ID</code> | The process ID of the wait. |
| <code>USER</code> | The name of the user that owns the session. |
| <code>WAIT NAME</code> | The name of the wait event. |
| <code>COUNT</code> | The number of times that the wait event occurred. |
| <code>TIME</code> | The length of the wait in seconds. |
| <code>% WAIT SES</code> | The percentage of the total wait time used by this wait for this session. |
| <code>% WAIT ALL</code> | The percentage of the total wait time used by this wait (for all sessions). |

sesshist_rpt()

The `sesshist_rpt()` function returns session wait information for a specified backend. The signature is:

`sesshist_rpt(<snapshot_id>, <session_id>)`

Parameters

`snapshot_id`

`snapshot_id` is an integer value that identifies the snapshot.

`session_id`

`session_id` is an integer value that represents the session.

The following example demonstrates a call to the `sesshist_rpt()` function:

Note

The following example has been shortened; over 1300 rows were actually generated.

```
SELECT * FROM sesshist_rpt (9, 3501);
          sesshist_rpt
-----
ID  USER    SEQ  WAIT NAME      ELAPSED File  Name          #
of Blk Sum of Blks

-----
3501 enterprise 1  query plan   13    0    N/A
0     0
3501 enterprise 1  query plan   13    0    edb_password_history
0     0
3501 enterprise 1  query plan   13    0    edb_password_history
0     0
3501 enterprise 1  query plan   13    0    edb_password_history
0     0
3501 enterprise 1  query plan   13    0    edb_profile
0     0
3501 enterprise 1  query plan   13    0    edb_profile_name_ind
0     0
3501 enterprise 1  query plan   13    0    edb_profile_oid_inde
0     0
3501 enterprise 1  query plan   13    0    edb_profile_password
0     0
3501 enterprise 1  query plan   13    0    edb_resource_group
0     0
3501 enterprise 1  query plan   13    0    edb_resource_group_n
0     0
3501 enterprise 1  query plan   13    0    edb_resource_group_o
0     0
3501 enterprise 1  query plan   13    0    pg_attribute
0     0
3501 enterprise 1  query plan   13    0    pg_attribute_relid_a
0     0
3501 enterprise 1  query plan   13    0    pg_attribute_relid_a
0     0
3501 enterprise 1  query plan   13    0    pg_auth_members
0     0
```

| | | | | | |
|------|--------------|------------|-----|---|----------------------|
| 3501 | enterprise 1 | query plan | 13 | 0 | pg_auth_members_memb |
| 0 | 0 | | | | |
| 3501 | enterprise 1 | query plan | 13 | 0 | pg_auth_members_role |
| 0 | 0 | | | | |
| | | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | N/A |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_password_history |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_password_history |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_password_history |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_profile |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_profile_name_ind |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_profile_oid_inde |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_profile_password |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_resource_group |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_resource_group_n |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | edb_resource_group_o |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | pg_attribute |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | pg_attribute_relid_a |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | pg_attribute_relid_a |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | pg_auth_members |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | pg_auth_members_memb |
| 0 | 0 | | | | |
| 3501 | enterprise 2 | wal flush | 149 | 0 | pg_auth_members_role |
| 0 | 0 | | | | |
| | | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | N/A |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_password_history |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_password_history |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_password_history |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_profile |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_profile_name_ind |

| | | | | | |
|------|---------------|-----------|-----|---|-----------------------|
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_profile_oid_index |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_profile_password |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_resource_group |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_resource_group_n |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | edb_resource_group_o |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | pg_attribute |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | pg_attribute_relid_a |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | pg_attribute_relid_a |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | pg_auth_members |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | pg_auth_members_memb |
| 0 | 0 | | | | |
| 3501 | enterprise 3 | wal write | 148 | 0 | pg_auth_members_role |
| 0 | 0 | | | | |
| | | . | . | . | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_1255 |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_1255_index |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_2396 |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_2396_index |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_2964 |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_2964_index |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_3592 |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_toast_3592_index |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_type |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_type_oid_index |
| 0 | 0 | | | | |
| 3501 | enterprise 24 | wal write | 130 | 0 | pg_type_typname_nsp_ |
| 0 | 0 | | | | |

(1304 rows)

The information displayed in the result set includes:

| Column Name | Description |
|-------------|---|
| ID | The system-assigned identifier of the wait. |

| | |
|-------------|---|
| USER | The name of the user that incurred the wait. |
| SEQ | The sequence number of the wait event. |
| WAIT NAME | The name of the wait event. |
| ELAPSED | The length of the wait event in microseconds. |
| File | The relfilenode number of the file. |
| Name | If available, the name of the file name related to the wait event. |
| # of Blk | The block number read or written for a specific instance of the event . |
| Sum of Blks | The number of blocks read. |

purgesnap()

The `purgesnap()` function purges a range of snapshots from the snapshot tables. The signature is:

`| purgesnap(<beginning_id>, <ending_id>)`

Parameters

`beginning_id`

`| beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`| ending_id` is an integer value that represents the ending session identifier.

`purgesnap()` removes all snapshots between `beginning_id` and `ending_id` (inclusive):

```
SELECT * FROM purgesnap(6, 9);
```

```
    purgesnap
```

```
-----  
Sno|beginning_id|ending_id|  
---|-----|-----|  
 1 |       6     |      9     |  
(1 row)
```

A call to the `get_snaps()` function after executing the example shows that snapshots 6 through 9 have been purged from the snapshot tables:

```
SELECT * FROM get_snaps();  
    get_snaps
```

```
-----  
Sno|beginning_id|ending_id|  
---|-----|-----|  
 1 |       1     |      5     |  
 2 |       6     |      9     |  
 3 |       10    |     12    |  
 4 |       13    |     15    |  
 5 |       16    |     18    |  
 6 |       19    |     21    |  
 7 |       22    |     24    |  
 8 |       25    |     27    |  
 9 |       28    |     30    |  
10 |       29    |     31    |  
11 |       32    |     34    |  
12 |       35    |     37    |  
13 |       38    |     40    |  
14 |       41    |     43    |  
15 |       44    |     46    |  
16 |       47    |     49    |  
17 |       50    |     52    |  
18 |       53    |     55    |  
19 |       56    |     58    |  
20 |       59    |     61    |  
21 |       62    |     64    |  
22 |       65    |     67    |  
23 |       68    |     70    |  
24 |       71    |     73    |  
25 |       74    |     76    |  
26 |       77    |     79    |  
27 |       80    |     82    |  
28 |       83    |     85    |  
29 |       86    |     88    |  
30 |       89    |     91    |  
31 |       92    |     94    |  
32 |       95    |     97    |  
33 |       98    |     100   |  
(8 rows)
```

truncsnap()

Use the `truncsnap()` function to delete all records from the snapshot table. The signature is:

`| truncsnap()`

For example:

```
SELECT * FROM truncsnap();
```

```
truncsnap
-----
```

```
Snapshots truncated.
(1 row)
```

A call to the `get_snaps()` function after calling the `truncsnap()` function shows that all records have been removed from the snapshot tables:

```
SELECT * FROM get_snaps();
get_snaps
-----
(0 rows)
```

Simulating Statspack AWR Reports

The functions described in this section return information comparable to the information contained in an Oracle Statspack/AWR (Automatic Workload Repository) report. When taking a snapshot, performance data from system catalog tables is saved into history tables. The reporting functions listed below report on the differences between two given snapshots.

- `stat db rpt()`
- `stat tables rpt()`
- `statio tables rpt()`
- `stat indexes rpt()`
- `statio_indexes_rpt()`

The reporting functions can be executed individually or you can execute all five functions by calling the `edbreport()` function.

edbreport()

The `edbreport()` function includes data from the other reporting functions, plus additional system information. The signature is:

`|edbreport(<beginning_id>, <ending_id>)`

Parameters

`beginning_id`

`| beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`| ending_id` is an integer value that represents the ending session identifier.

The call to the `edbreport()` function returns a composite report that contains system information and the reports returned by the other statspack functions.

```
SELECT * FROM edbreport(9, 10);
edbreport
```

EnterpriseDB Report for database acctg 25-JUL-18
 Version: PostgreSQL 12.0 (EnterpriseDB Advanced Server 12.0.2) on x86_64-pc-linux-gnu,
 compiled by gcc (GCC) 4.4.7 20120313 (Red Hat 4.4.7-18), 64-bit

Begin snapshot: 9 at 25-JUL-18 09:49:44.788392

End snapshot: 10 at 25-JUL-18 09:49:49.855821

Size of database acctg is 173 MB

Tablespace: pg_default Size: 231 MB Owner: enterpriseedb
 Tablespace: pg_global Size: 719 kB Owner: enterpriseedb

Schema: pg_toast_temp_1 Size: 0 bytes Owner: enterpriseedb

Schema: public Size: 158 MB Owner: enterpriseedb

The information displayed in the report introduction includes the database name and version, the current date, the beginning and ending snapshot date and times, database and tablespace details and schema information.

Top 10 Relations by pages

| TABLE | RELPAGES |
|-------------------------|----------|
| pgbench_accounts | 16394 |
| pgbench_history | 391 |
| pg_proc | 145 |
| pg_attribute | 92 |
| pg_depend | 81 |
| pg_collation | 60 |
| edb\$stat_all_indexes | 46 |
| edb\$statio_all_indexes | 46 |
| pg_description | 44 |
| edb\$stat_all_tables | 29 |

The information displayed in the [Top 10 Relations by pages](#) section includes:

| Column Name | Description |
|-------------|-----------------------------------|
| TABLE | The name of the table. |
| RELPAGES | The number of pages in the table. |

Top 10 Indexes by pages

| INDEX | RELPAGES |
|---------------------------------|----------|
| pgbench_accounts_pkey | 2745 |
| pg_depend_reference_index | 68 |
| pg_depend_depender_index | 63 |
| pg_procproname_args_nsp_index | 53 |
| pg_attribute_relid_attnam_index | 25 |
| pg_description_o_c_o_index | 24 |
| pg_attribute_relid_attnum_index | 17 |
| pg_proc_oid_index | 14 |
| pg_collation_name_enc_nsp_index | 12 |

```
edb$stat_idx_pk
```

```
10
```

The information displayed in the **Top 10 Indexes by pages** section includes:

| Column Name | Description |
|-------------|-----------------------------------|
| INDEX | The name of the index. |
| RELPAGES | The number of pages in the index. |

Top 10 Relations by DML

| SCHEMA | RELATION | UPDATES | DELETES | INSERTS |
|--------|------------------|---------|---------|---------|
| public | pgbench_accounts | 117209 | 0 | 1000000 |
| public | pgbench_tellers | 117209 | 0 | 100 |
| public | pgbench_branches | 117209 | 0 | 10 |
| public | pgbench_history | 0 | 0 | 117209 |

The information displayed in the **Top 10 Relations by DML** section includes:

| Column Name | Description |
|-------------|--|
| SCHEMA | The name of the schema in which the table resides. |
| RELATION | The name of the table. |
| UPDATES | The number of UPDATES performed on the table. |
| DELETES | The number of DELETES performed on the table. |
| INSERTS | The number of INSERTS performed on the table. |

DATA from pg_stat_database

| DATABASE | NUMBACKENDS | XACT COMMIT | XACT ROLLBACK | BLKS READ | BLKS HIT | HIT RATIO |
|----------|-------------|-------------|---------------|-----------|----------|-----------|
| acctg | 0 | 8261 | 0 | 117 | 127985 | 99.91 |

The information displayed in the DATA from **pg_stat_database** section of the report includes:

| Column Name | Description |
|---------------|--|
| DATABASE | The name of the database. |
| NUMBACKENDS | Number of backends currently connected to this database. This is the only column in this view that returns a value reflecting current state; all other columns return the accumulated values since the last reset. |
| XACT COMMIT | Number of transactions in this database that have been committed. |
| XACT ROLLBACK | Number of transactions in this database that have been rolled back. |
| BLKS READ | Number of disk blocks read in this database. |
| BLKS HIT | Number of times disk blocks were found already in the buffer cache (when a read was not necessary). |
| HIT RATIO | The percentage of times that a block was found in the shared buffer cache. |

DATA from pg_buffercache

| RELATION | BUFFERS |
|----------|---------|
| | |

| | |
|---------------------------|-------|
| pgbench_accounts | 16665 |
| pgbench_accounts_pkey | 2745 |
| pgbench_history | 751 |
| edb\$statio_all_indexes | 94 |
| edb\$stat_all_indexes | 94 |
| edb\$stat_all_tables | 60 |
| edb\$statio_all_tables | 56 |
| edb\$session_wait_history | 34 |
| edb\$statio_idx_pk | 17 |
| pg_depend | 17 |

The information displayed in the **DATA from pg_buffercache** section of the report includes:

| Column Name | Description |
|-----------------|--|
| RELATION | The name of the table. |
| BUFFERS | The number of shared buffers used by the relation. |

Note

In order to obtain the report for **DATA from pg_buffercache**, the **pg_buffercache** module must have been installed in the database. Perform the installation with the **CREATE EXTENSION** command.

For more information on the **CREATE EXTENSION** command, see the PostgreSQL Core documentation at:

<https://www.postgresql.org/docs/current/static/sql-createextension.html>

DATA from pg_stat_all_tables ordered by seq scan

| SCHEMA | RELATION | SEQ SCAN REL TUP READ IDX SCAN | | | | |
|------------|---------------------------|--------------------------------|--------|------|-----|-----|
| | | IDX | TUP | READ | INS | UPD |
| <hr/> | | | | | | |
| public | pgbench_branches | 8258 | 82580 | 0 | 0 | 0 |
| 0 | 0 | 8258 | 0 | 0 | 0 | 0 |
| public | pgbench_tellers | 8258 | 825800 | 0 | 0 | 0 |
| 0 | 0 | 8258 | 0 | 0 | 0 | 0 |
| pg_catalog | pg_class | 7 | 3969 | 92 | 0 | 0 |
| 80 | 0 | 0 | 0 | 0 | 0 | 0 |
| pg_catalog | pg_index | 5 | 950 | 31 | 0 | 0 |
| 38 | 0 | 0 | 0 | 0 | 0 | 0 |
| pg_catalog | pg_namespace | 4 | 144 | 5 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 |
| pg_catalog | pg_database | 2 | 12 | 7 | 0 | 0 |
| 7 | 0 | 0 | 0 | 0 | 0 | 0 |
| pg_catalog | pg_am | 1 | 1 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| pg_catalog | pg_authid | 1 | 10 | 2 | 0 | 0 |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 |
| sys | callback_queue_table | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| sys | edb\$session_wait_history | 0 | 0 | 0 | 0 | 0 |
| 0 | 125 | 0 | 0 | 0 | 0 | 0 |

The information displayed in the **DATA from pg_stat_all_tables ordered by seq scan** section includes:

| Column Name | Description |
|--------------|---|
| SCHEMA | The name of the schema in which the table resides. |
| RELATION | The name of the table. |
| SEQ SCAN | The number of sequential scans initiated on this table. |
| REL TUP READ | The number of tuples read in the table. |
| IDX SCAN | The number of index scans initiated on the table. |
| IDX TUP READ | The number of index tuples read. |
| INS | The number of rows inserted. |
| UPD | The number of rows updated. |
| DEL | The number of rows deleted. |

DATA from pg_stat_all_tables ordered by rel tup read

| SCHEMA | RELATION | SEQ SCAN | REL TUP READ | IDX SCAN |
|--------------|---------------------------|----------|--------------|----------|
| IDX TUP READ | INS | UPD | DEL | |
| <hr/> | | | | |
| public | pgbench_tellers | 8258 | 825800 | 0 |
| 0 | 0 | 8258 | 0 | |
| public | pgbench_branches | 8258 | 82580 | 0 |
| 0 | 0 | 8258 | 0 | |
| pg_catalog | pg_class | 7 | 3969 | 92 |
| 80 | 0 | 0 | 0 | |
| pg_catalog | pg_index | 5 | 950 | 31 |
| 38 | 0 | 0 | 0 | |
| pg_catalog | pg_namespace | 4 | 144 | 5 |
| 4 | 0 | 0 | 0 | |
| pg_catalog | pg_database | 2 | 12 | 7 |
| 7 | 0 | 0 | 0 | |
| pg_catalog | pg_authid | 1 | 10 | 2 |
| 2 | 0 | 0 | 0 | |
| pg_catalog | pg_am | 1 | 1 | 0 |
| 0 | 0 | 0 | 0 | |
| sys | callback_queue_table | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | |
| sys | edb\$session_wait_history | 0 | 0 | 0 |
| 0 | 125 | 0 | 0 | |

The information displayed in the [DATA from pg_stat_all_tables ordered by rel tup read](#) section includes:

| Column Name | Description |
|--------------|--|
| SCHEMA | The name of the schema in which the table resides. |
| RELATION | The name of the table. |
| SEQ SCAN | The number of sequential scans performed on the table. |
| REL TUP READ | The number of tuples read from the table. |
| IDX SCAN | The number of index scans performed on the table. |
| IDX TUP READ | The number of index tuples read. |
| INS | The number of rows inserted. |
| UPD | The number of rows updated. |
| DEL | The number of rows deleted. |

DATA from pg_statio_all_tables

| SCHEMA TOAST | RELATION TOAST | HEAP TIDX | IDX TIDX | READ | | HEAP | | IDX | |
|-----------------|-------------------|--------------|-------------|------|-------|------|-----|------|-----|
| | | | | READ | | HEAP | | IDX | |
| | | | | READ | HIT | READ | HIT | READ | HIT |
| <hr/> | | | | | | | | | |
| public | pgbench_accounts | 32 | 25016 | 0 | 49913 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| public | pgbench_tellers | 0 | 24774 | 0 | 0 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| public | pgbench_branches | 0 | 16516 | 0 | 0 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| public | pgbench_history | 53 | 8364 | 0 | 0 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| pg_catalog | pg_class | 0 | 199 | 0 | 187 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| pg_catalog | pg_attribute | 0 | 198 | 0 | 395 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| pg_catalog | pg_proc | 0 | 75 | 0 | 153 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| pg_catalog | pg_index | 0 | 56 | 0 | 33 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| pg_catalog | pg_amop | 0 | 48 | 0 | 56 | | | | |
| 0 0 | 0 0 | | | | | | | | |
| pg_catalog | pg_namespace | 0 | 28 | 0 | 7 | | | | |
| 0 0 | 0 0 | | | | | | | | |

The information displayed in the `DATA from pg_statio_all_tables` section includes:

| Column Name | Description |
|-------------------------|--|
| <code>SCHEMA</code> | The name of the schema in which the table resides. |
| <code>RELATION</code> | The name of the table. |
| <code>HEAP READ</code> | The number of heap blocks read. |
| <code>HEAP HIT</code> | The number of heap blocks hit. |
| <code>IDX READ</code> | The number of index blocks read. |
| <code>IDX HIT</code> | The number of index blocks hit. |
| <code>TOAST READ</code> | The number of toast blocks read. |

TOAST HIT

The number of toast blocks hit.

TIDX READ

The number of toast index blocks read.

TIDX HIT

The number of toast index blocks hit.

DATA from pg_stat_all_indexes

| SCHEMA | RELATION | INDEX |
|------------|---------------------------------|------------------------------|
| IDX SCAN | IDX TUP READ | IDX TUP FETCH |
| <hr/> | | |
| public | pgbench_accounts | pgbench_accounts_pkey |
| 16516 | 16679 | 16516 |
| pg_catalog | pg_attribute | |
| | pg_attribute_relid_attnum_index | 196 402 402 |
| pg_catalog | pg_proc | pg_proc_oid_index |
| 70 | 70 | |
| pg_catalog | pg_class | pg_class_oid_index |
| 61 | 61 | |
| pg_catalog | pg_class | pg_class_relname_nsp_index |
| 31 19 19 | | |
| pg_catalog | pg_type | pg_type_oid_index |
| 22 22 | 22 | |
| pg_catalog | edb_policy | edb_policy_object_name_index |
| 21 0 | 0 | |
| pg_catalog | pg_amop | pg_amop_fam_strat_index |
| 16 16 | 16 | |
| pg_catalog | pg_index | pg_index_indexrelid_index |
| 16 16 | 16 | |
| pg_catalog | pg_index | pg_index_indrelid_index |
| 15 22 | 22 | |

The information displayed in the **DATA from pg_stat_all_indexes** section includes:

Column Name**Description****SCHEMA**

The name of the schema in which the index resides.

RELATION

The name of the table on which the index is defined.

INDEX

The name of the index.

IDX SCAN

The number of indexes scans initiated on this index.

IDX TUP READ Number of index entries returned by scans on this index

IDX TUP FETCH Number of live table rows fetched by simple index scans using this index.

DATA from pg_statio_all_indexes

| SCHEMA | RELATION | INDEX |
|---------------------------------|---------------------------|----------------------------|
| IDX BLKS READ | IDX BLKS HIT | |
| <hr/> | | |
| public | pgbench_accounts | pgbench_accounts_pkey |
| 0 | 49913 | |
| pg_catalog | pg_attribute | |
| pg_attribute_relid_attnum_index | 0 | 395 |
| sys | edb\$stat_all_indexes | edb\$stat_idx_pk |
| 1 | 382 | |
| sys | edb\$statio_all_indexes | edb\$statio_idx_pk |
| 1 | 382 | |
| sys | edb\$statio_all_tables | edb\$statio_tab_pk |
| 2 | 262 | |
| sys | edb\$stat_all_tables | edb\$stat_tab_pk |
| 0 | 259 | |
| sys | edb\$session_wait_history | session_waits_hist_pk |
| 0 | 251 | |
| pg_catalog | pg_proc | pg_proc_oid_index |
| 0 | 142 | |
| pg_catalog | pg_class | pg_class_oid_index |
| 0 | 123 | |
| pg_catalog | pg_class | pg_class_relname_nsp_index |
| 0 | 63 | |

The information displayed in the **DATA from pg_statio_all_indexes** section includes:

| Column Name | Description |
|----------------------|--|
| SCHEMA | The name of the schema in which the index resides. |
| RELATION | The name of the table on which the index is defined. |
| INDEX | The name of the index. |
| IDX BLKS READ | The number of index blocks read. |
| IDX BLKS HIT | The number of index blocks hit. |

System Wait Information

| WAIT NAME | COUNT | WAIT TIME | % WAIT |
|---------------------|-------|-----------|--------|
| <hr/> | | | |
| wal flush | 8359 | 1.357593 | 30.62 |
| wal write | 8358 | 1.349153 | 30.43 |
| wal file sync | 8358 | 1.286437 | 29.02 |
| query plan | 33439 | 0.439324 | 9.91 |
| db file extend | 54 | 0.000585 | 0.01 |
| db file read | 31 | 0.000307 | 0.01 |
| other llock acquire | 0 | 0.000000 | 0.00 |

| | | | |
|-----------------|---|----------|------|
| ProcArrayLock | 0 | 0.000000 | 0.00 |
| CLogControlLock | 0 | 0.000000 | 0.00 |

The information displayed in the [System Wait Information](#) section includes:

| Column Name | Description |
|-------------|---|
| WAIT NAME | The name of the wait. |
| COUNT | The number of times that the wait event occurred. |
| WAIT TIME | The length of the wait time in seconds. |
| % WAIT | The percentage of the total wait time used by this wait for this session. |

Database Parameters from postgresql.conf

| PARAMETER | | |
|-------------------------------------|------------|------------|
| CONTEXT | MINVAL | MAXVAL |
| <hr/> | | |
| <hr/> | | |
| allow_system_table_mods | off | |
| postmaster | | |
| application_name | psql.bin | |
| user | | |
| archive_command | (disabled) | |
| sighup | | |
| archive_mode | off | |
| postmaster | | |
| archive_timeout | 0 | |
| sighup | 0 | 1073741823 |
| array_nulls | on | |
| user | | |
| authentication_timeout | 60 | |
| sighup | 1 | 600 |
| autovacuum | on | |
| sighup | | |
| autovacuum_analyze_scale_factor | 0.1 | |
| sighup | 0 | 100 |
| autovacuum_analyze_threshold | 50 | |
| sighup | 0 | 2147483647 |
| autovacuum_freeze_max_age | 200000000 | |
| postmaster | 100000 | 2000000000 |
| autovacuum_max_workers | 3 | |
| postmaster | 1 | 262143 |
| autovacuum_multixact_freeze_max_age | 400000000 | |
| postmaster | 10000 | 2000000000 |
| autovacuum_naptime | 60 | |
| sighup | 1 | 2147483 |
| autovacuum_vacuum_cost_delay | 20 | |
| sighup | -1 | 100 |
| <hr/> | | |
| <hr/> | | |

The information displayed in the [Database Parameters from postgresql.conf](#) section includes:

| Column Name | Description |
|-------------|-------------|
|-------------|-------------|

| | |
|------------------|--|
| PARAMETER | The name of the parameter. |
| SETTING | The current value assigned to the parameter. |
| CONTEXT | The context required to set the parameter value. |
| MINVAL | The minimum value allowed for the parameter. |
| MAXVAL | The maximum value allowed for the parameter. |

stat_db_rpt()

The signature is:

```
stat_db_rpt(<beginning_id>, <ending_id>)
```

Parameters

beginning_id

beginning_id is an integer value that represents the beginning session identifier.

ending_id

ending_id is an integer value that represents the ending session identifier.

The following example demonstrates the `stat_db_rpt()` function:

```
SELECT * FROM stat_db_rpt(9, 10);
```

```
stat_db_rpt
```

DATA from pg_stat_database

| DATABASE | NUMBACKENDS | XACT COMMIT | XACT ROLLBACK | BLKS READ | BLKS HIT | HIT RATIO |
|----------|-------------|-------------|---------------|-----------|----------|-----------|
| acctg | 0 | 8261 | 0 | 117 | 127985 | 99.91 |

(5 rows)

The information displayed in the DATA from `pg_stat_database` section of the report includes:

| Column Name | Description |
|----------------------|--|
| DATABASE | The name of the database. |
| NUMBACKENDS | Number of backends currently connected to this database. This is the only column in this view that returns a value reflecting current state; all other columns return the accumulated values since the last reset. |
| XACT COMMIT | The number of transactions in this database that have been committed. |
| XACT ROLLBACK | The number of transactions in this database that have been rolled back. |
| BLKS READ | The number of blocks read. |
| BLKS HIT | The number of blocks hit. |
| HIT RATIO | The percentage of times that a block was found in the shared buffer cache. |

stat_tables_rpt()

The signature is:

```
| function_name(<beginning_id>, <ending_id>, <top_n>, <scope>)
```

Parameters

`beginning_id`

| `beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

| `ending_id` is an integer value that represents the ending session identifier.

`top_n`

| `top_n` represents the number of rows to return

`scope`

scope determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- `SYS` indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, or `sys`.
- `USER` indicates that the function should return information about user-defined tables.
- `ALL` specifies that the function should return information about all tables.

The `stat_tables_rpt()` function returns a two-part report. The first portion of the report contains:

```
SELECT * FROM stat_tables_rpt(8, 9, 10, 'ALL');
```

`stat_tables_rpt`

DATA from `pg_stat_all_tables` ordered by seq scan

| SCHEMA IDX | RELATION TUP | SEQ INS | SCAN UPD | REL TUP READ IDX SCAN | | |
|---------------|---------------------------|------------|-------------|-----------------------|-----|------|
| | | | | READ | DEL | SCAN |
| public | pgbench_branches | 8249 | 82490 | 0 | 0 | |
| 0 | 0 | 8249 | 0 | | | |
| public | pgbench_tellers | 8249 | 824900 | 0 | 0 | |
| 0 | 0 | 8249 | 0 | | | |
| pg_catalog | pg_class | 7 | 3969 | 92 | 0 | |
| 80 | 0 | 0 | 0 | | | |
| pg_catalog | pg_index | 5 | 950 | 31 | 0 | |
| 38 | 0 | 0 | 0 | | | |
| pg_catalog | pg_namespace | 4 | 144 | 5 | 0 | |
| 4 | 0 | 0 | 0 | | | |
| pg_catalog | pg_am | 1 | 1 | 0 | 0 | |
| 0 | 0 | 0 | 0 | | | |
| pg_catalog | pg_authid | 1 | 10 | 2 | 0 | |
| 2 | 0 | 0 | 0 | | | |
| pg_catalog | pg_database | 1 | 6 | 3 | 0 | |
| 3 | 0 | 0 | 0 | | | |
| sys | callback_queue_table | 0 | 0 | 0 | 0 | |
| 0 | 0 | 0 | 0 | | | |
| sys | edb\$session_wait_history | 0 | 0 | 0 | 0 | |

| | | |
|---|-------|---|
| 0 | 125 0 | 0 |
|---|-------|---|

The information displayed in the **DATA from pg_stat_all_tables ordered by seq scan** section includes:

| Column Name | Description |
|--------------|--|
| SCHEMA | The name of the schema in which the table resides. |
| RELATION | The name of the table. |
| SEQ SCAN | The number of sequential scans on the table. |
| REL TUP READ | The number of tuples read from the table. |
| IDX SCAN | The number of index scans performed on the table. |
| IDX TUP READ | The number of index tuples read from the table. |
| INS | The number of rows inserted. |
| UPD | The number of rows updated. |
| DEL | The number of rows deleted. |

The second portion of the report contains:

DATA from pg_stat_all_tables ordered by rel tup read

| SCHEMA | RELATION | SEQ SCAN | REL TUP READ | IDX SCAN |
|--------------|---------------------------|----------|--------------|----------|
| IDX TUP READ | INS | UPD | DEL | |
| <hr/> | | | | |
| public | pgbench_tellers | 8249 | 824900 | 0 |
| 0 | 0 | 8249 0 | | |
| public | pgbench_branches | 8249 | 82490 | 0 |
| 0 | 0 | 8249 0 | | |
| pg_catalog | pg_class | 7 | 3969 | 92 |
| 80 | 0 | 0 0 | | |
| pg_catalog | pg_index | 5 | 950 | 31 |
| 38 | 0 | 0 0 | | |
| pg_catalog | pg_namespace | 4 | 144 | 5 |
| 4 | 0 | 0 0 | | |
| pg_catalog | pg_authid | 1 | 10 | 2 |
| 2 | 0 | 0 0 | | |
| pg_catalog | pg_database | 1 | 6 | 3 |
| 3 | 0 | 0 0 | | |
| pg_catalog | pg_am | 1 | 1 | 0 |
| 0 | 0 | 0 0 | | |
| sys | callback_queue_table | 0 | 0 | 0 |
| 0 | 0 | 0 0 | | |
| sys | edb\$session_wait_history | 0 | 0 | 0 |
| 0 | 125 | 0 0 | | |
| (29 rows) | | | | |

The information displayed in the **DATA from pg_stat_all_tables ordered by rel tup read** section includes:

| Column Name | Description |
|--------------|--|
| SCHEMA | The name of the schema in which the table resides. |
| RELATION | The name of the table. |
| SEQ SCAN | The number of sequential scans performed on the table. |
| REL TUP READ | The number of tuples read from the table. |

| | |
|--------------|---|
| IDX SCAN | The number of index scans performed on the table. |
| IDX TUP READ | The number of live rows fetched by index scans. |
| INS | The number of rows inserted. |
| UPD | The number of rows updated. |
| DEL | The number of rows deleted. |

statio_tables_rpt()

The signature is:

```
statio_tables_rpt(<beginning_id>, <ending_id>, <top_n>, <scope>)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- `SYS` indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, or `sys`.
- `USER` indicates that the function should return information about user-defined tables.
- `ALL` specifies that the function should return information about all tables.

The `statio_tables_rpt()` function returns a report that contains:

```
SELECT * FROM statio_tables_rpt(9, 10, 10, 'SYS');
```

```
statio_tables_rpt
```

```
-----
```

```
DATA from pg_statio_all_tables
```

| SCHEMA TOAST | RELATION TIDX | HEAP TIDX | IDX | | | TOAST |
|-----------------|------------------|----------------------------|------|-----|------|-------|
| | | | READ | HIT | READ | |
| HIT | READ | HIT | | | | |
| sys 0 | 0 | edb\$stat_all_indexes 0 | 18 | 1 | 382 | 0 |
| sys 0 | 0 | edb\$statio_all_index 0 | 18 | 1 | 382 | 0 |
| sys | | edb\$statio_all_table 5 | 12 | 2 | 262 | 0 |

```

0    0    0
sys      edb$stat_all_tables 4    10    0    259  0
0    0    0
sys      edb$session_wait_his 2    6    0    251  0
0    0    0
sys      edb$session_waits   1    4    0    12   0
0    0    0
sys      callback_queue_table 0    0    0    0    0
0    0    0
sys      dual        0    0    0    0    0
0    0    0
sys      edb$snap       0    1    0    2    0
0    0    0
sys      edb$stat_database 0    2    0    7    0
0    0    0
(15 rows)

```

The information displayed in the `DATA from pg_statio_all_tables` section includes:

| Column Name | Description |
|-------------------------|---|
| <code>SCHEMA</code> | The name of the schema in which the relation resides. |
| <code>RELATION</code> | The name of the relation. |
| <code>HEAP READ</code> | The number of heap blocks read. |
| <code>HEAP HIT</code> | The number of heap blocks hit. |
| <code>IDX READ</code> | The number of index blocks read. |
| <code>IDX HIT</code> | The number of index blocks hit. |
| <code>TOAST READ</code> | The number of toast blocks read. |
| <code>TOAST HIT</code> | The number of toast blocks hit. |
| <code>TIDX READ</code> | The number of toast index blocks read. |
| <code>TIDX HIT</code> | The number of toast index blocks hit. |

stat_indexes_rpt()

The signature is:

```
stat_indexes_rpt(<beginning_id>, <ending_id>, <top_n>, <scope>)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- **SYS** indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, or `sys`.
- **USER** indicates that the function should return information about user-defined tables.
- **ALL** specifies that the function should return information about all tables.

The `stat_indexes_rpt()` function returns a report that contains:

```
edb=# SELECT * FROM stat_indexes_rpt(9, 10, 10, 'ALL');
```

stat_indexes_rpt

DATA from pg_stat_all_indexes

| SCHEMA | RELATION | INDEX |
|------------|---------------------------------|------------------------------|
| IDX SCAN | IDX TUP READ | IDX TUP FETCH |
| <hr/> | | |
| public | pgbench_accounts | pgbench_accounts_pkey |
| 16516 | 16679 | 16516 |
| pg_catalog | pg_attribute | |
| | pg_attribute_relid_attnum_index | 196 402 402 |
| pg_catalog | pg_proc | pg_proc_oid_index |
| 70 | 70 | 70 |
| pg_catalog | pg_class | pg_class_oid_index |
| 61 | 61 | 61 |
| pg_catalog | pg_class | pg_class_relname_nsp_index |
| 31 | 19 | 19 |
| pg_catalog | pg_type | pg_type_oid_index |
| 22 | 22 | 22 |
| pg_catalog | edb_policy | edb_policy_object_name_index |
| 21 | 0 | 0 |
| pg_catalog | pg_amop | pg_amop_fam_strat_index |
| 16 | 16 | 16 |
| pg_catalog | pg_index | pg_index_indexrelid_index |
| 16 | 16 | 16 |
| pg_catalog | pg_index | pg_index_indrelid_index |
| 15 | 22 | 22 |
| (14 rows) | | |

The information displayed in the `DATA from pg_stat_all_indexes` section includes:

| Column Name | Description |
|---------------|---|
| SCHEMA | The name of the schema in which the relation resides. |
| RELATION | The name of the relation. |
| INDEX | The name of the index. |
| IDX SCAN | The number of indexes scanned. |
| IDX TUP READ | The number of index tuples read. |
| IDX TUP FETCH | The number of index tuples fetched. |

statio_indexes_rpt()

The signature is:

```
statio_indexes_rpt(<beginning_id>, <ending_id>, <top_n>, <scope>)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- `SYS` indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, or `sys`.
- `USER` indicates that the function should return information about user-defined tables.
- `ALL` specifies that the function should return information about all tables.

The `statio_indexes_rpt()` function returns a report that contains:

```
edb=# SELECT * FROM statio_indexes_rpt(9, 10, 10, 'SYS');
```

```
statio_indexes_rpt
```

DATA from `pg_statio_all_indexes`

| SCHEMA IDX BLKS READ | RELATION IDX BLKS HIT | INDEX BLKS HIT |
|---------------------------------|--------------------------|----------------------------|
| pg_catalog pg_attribute | | |
| pg_attribute_relid_attnum_index | 0 | 395 |
| sys edb\$stat_all_indexes | | edb\$stat_idx_pk |
| 1 382 | | |
| sys edb\$statio_all_indexes | | edb\$statio_idx_pk |
| 1 382 | | |
| sys edb\$statio_all_tables | | edb\$statio_tab_pk |
| 2 262 | | |
| sys edb\$stat_all_tables | | edb\$stat_tab_pk |
| 0 259 | | |
| sys edb\$session_wait_history | | session_waits_hist_pk |
| 0 251 | | |
| pg_catalog pg_proc | | pg_proc_oid_index |
| 0 142 | | |
| pg_catalog pg_class | | pg_class_oid_index |
| 0 123 | | |
| pg_catalog pg_class | | pg_class_relname_nsp_index |
| 0 63 | | |

| pg_catalog | pg_type | pg_type_oid_index |
|------------|---------|-------------------|
| 0 | 45 | |
| (14 rows) | | |

The information displayed in the [DATA from pg_statio_all_indexes](#) report includes:

| Column Name | Description |
|---------------|---|
| SCHEMA | The name of the schema in which the relation resides. |
| RELATION | The name of the table on which the index is defined. |
| INDEX | The name of the index. |
| IDX BLKS READ | The number of index blocks read. |
| IDX BLKS HIT | The number of index blocks hit. |

Performance Tuning Recommendations

To use DRITA reports for performance tuning, review the top five events in a given report, looking for any event that takes a disproportionately large percentage of resources. In a streamlined system, user I/O will probably make up the largest number of waits. Waits should be evaluated in the context of CPU usage and total time; an event may not be significant if it takes 2 minutes out of a total measurement interval of 2 hours, if the rest of the time is consumed by CPU time. The component of response time (CPU "work" time or other "wait" time) that consumes the highest percentage of overall time should be evaluated.

When evaluating events, watch for:

| Event type | Description |
|----------------------------|---|
| Checkpoint waits | Checkpoint waits may indicate that checkpoint parameters need to be adjusted, (checkpoint_segments and checkpoint_timeout). |
| WAL-related waits | WAL-related waits may indicate wal_buffers are under-sized. |
| SQL Parse waits | If the number of waits is high, try to use prepared statements. |
| db file random reads | If high, check that appropriate indexes and statistics exist. |
| db file random writes | If high, may need to decrease bgwriter_delay . |
| btree random lock acquires | May indicate indexes are being rebuilt. Schedule index builds during less active time. |

Performance reviews should also include careful scrutiny of the hardware, the operating system, the network and the application SQL statements.

Event Descriptions

The following table lists the basic wait events that are displayed by DRITA.

| Event Name | Description |
|------------------------------|--|
| add in shmem | Obsolete/unused |
| lock acquire | The bgwriter (background writer) process has waited for the short-term lock that synchronizes messages between the bgwriter and a backend process. |

| Event Name | Description |
|--------------------------------|---|
| btree vacuum lock acquire | The server has waited for the short-term lock that synchronizes access to the next available vacuum cycle ID. |
| buffer free list lock acquire | The server has waited for the short-term lock that synchronizes access to the list of free buffers (in shared memory). |
| checkpoint lock acquire | A server process has waited for the short-term lock that prevents simultaneous checkpoints. |
| checkpoint start lock acquire | The server has waited for the short-term lock that synchronizes access to the bgwriter checkpoint schedule. |
| clog control lock acquire | The server has waited for the short-term lock that synchronizes access to the commit log. |
| control file lock acquire | The server has waited for the short-term lock that synchronizes write access to the control file (this should usually be a low number). |
| db file extend | A server process has waited for the operating system while adding a new page to the end of a file. |
| db file read | A server process has waited for the completion of a read (from disk). |
| db file write | A server process has waited for the completion of a write (to disk). |
| db file sync | A server process has waited for the operating system to flush all changes to disk. |
| first buf mapping lock acquire | The server has waited for a short-term lock that synchronizes access to the shared-buffer mapping table. |
| freespace lock acquire | The server has waited for the short-term lock that synchronizes access to the freespace map. |
| lwlock acquire | The server has waited for a short-term lock that has not been described elsewhere in this section. |
| multi xact gen lock acquire | The server has waited for the short-term lock that synchronizes access to the next available multi-transaction ID (when a SELECT...FOR SHARE statement executes). |
| multi xact member lock acquire | The server has waited for the short-term lock that synchronizes access to the multi-transaction member file (when a SELECT...FOR SHARE statement executes). |
| multi xact offset lock acquire | The server has waited for the short-term lock that synchronizes access to the multi-transaction offset file (when a SELECT...FOR SHARE statement executes). |
| oid gen lock acquire | The server has waited for the short-term lock that synchronizes access to the next available OID (object ID). |
| query plan | The server has computed the execution plan for a SQL statement. |
| rel cache init lock acquire | The server has waited for the short-term lock that prevents simultaneous relation-cache loads/unloads. |
| shmem index lock acquire | The server has waited for the short-term lock that synchronizes access to the shared-memory map. |
| sinvil lock acquire | The server has waited for the short-term lock that synchronizes access to the cache invalidation state. |
| sql parse | The server has parsed a SQL statement. |
| subtrans control lock acquire | The server has waited for the short-term lock that synchronizes access to the subtransaction log. |
| tablespace create lock acquire | The server has waited for the short-term lock that prevents simultaneous CREATE TABLESPACE or DROP TABLESPACE commands. |
| two phase state lock acquire | The server has waited for the short-term lock that synchronizes access to the list of prepared transactions. |

| Event Name | Description |
|-------------------------|--|
| wal insert lock acquire | The server has waited for the short-term lock that synchronizes write access to the write-ahead log. A high number may indicate that WAL buffers are sized too small. |
| wal write lock acquire | The server has waited for the short-term lock that synchronizes write-ahead log flushes. |
| wal file sync | The server has waited for the write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to open_sync). |
| wal flush | The server has waited for the write-ahead log to flush to disk. |
| wal write | The server has waited for a write to the write-ahead log buffer (expect this value to be high). |
| xid gen lock acquire | The server has waited for the short-term lock that synchronizes access to the next available transaction ID. |

When wait events occur for lightweight locks, they are displayed by DRITA as well. A *lightweight lock* is used to protect a particular data structure in shared memory.

Certain wait events can be due to the server process waiting for one of a group of related lightweight locks, which is referred to as a *lightweight lock tranche*. Individual lightweight lock tranches are not displayed by DRITA, but their summation is displayed by a single event named `other lwlock acquire`.

For a list and description of lightweight locks displayed by DRITA, see Section 28.2, `The Statistics Collector` in the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/monitoring-stats.html>

Under Section 28.2.2. `Viewing Statistics`, the lightweight locks are listed in Table 28-4 `wait_event` Description where the `Wait Event Type` column designates `LWLock`.

The following example displays lightweight locks `ProcArrayLock`, `CLogControlLock`, `WALBufMappingLock`, and `XidGenLock`.

```
postgres=# select * from sys_rpt(40,70,20);
          sys_rpt
-----
```

| WAIT NAME | COUNT | WAIT TIME | % WAIT |
|----------------------|--------|-----------|--------|
| wal flush | 56107 | 44.456494 | 47.65 |
| db file read | 66123 | 19.543968 | 20.95 |
| wal write | 32886 | 12.780866 | 13.70 |
| wal file sync | 32933 | 11.792972 | 12.64 |
| query plan | 223576 | 4.539186 | 4.87 |
| db file extend | 2339 | 0.087038 | 0.09 |
| other lwlock acquire | 402 | 0.066591 | 0.07 |
| ProcArrayLock | 135 | 0.012942 | 0.01 |
| CLogControlLock | 212 | 0.010333 | 0.01 |
| WALBufMappingLock | 47 | 0.006068 | 0.01 |
| XidGenLock | 53 | 0.005296 | 0.01 |
| (13 rows) | | | |

DRITA also displays wait events that occur that are related to certain Advanced Server product features.

These Advanced Server feature specific wait events and the `other lwlock acquire` event are listed in the following table.

| Event Name | Description |
|--------------|---|
| BulkLoadLock | The server has waited for access related to EDB*Loader. |

| Event Name | Description |
|------------------------|---|
| EDBResourceManagerLock | The server has waited for access related to EDB Resource Manager. |
| other lwlock acquire | Summation of waits for lightweight lock tranches. |

9 ECPGPlus Guide

9.1 Introduction

EnterpriseDB has enhanced ECPG (the PostgreSQL pre-compiler) to create ECPGPlus. ECPGPlus allows you to include Pro*C compatible embedded SQL commands in C applications when connected to an EDB Postgres Advanced Server (Advanced Server) database. When you use ECPGPlus to compile an application, the SQL code is syntax-checked and translated into C.

ECPGPlus supports:

- Oracle Dynamic SQL – Method 4 (ODS-M4).
- Pro*C compatible anonymous blocks.
- A `CALL` statement compatible with Oracle databases.

As part of ECPGPlus's Pro*C compatibility, you do not need to include the `BEGIN DECLARE SECTION` and `END DECLARE SECTION` directives.


PostgreSQL Compatibility

While most ECPGPlus statements will work with community PostgreSQL, the `CALL` statement, and the `EXECUTE...END EXEC` statement work only when the client application is connected to EDB Postgres Advanced Server.

9.2 ECPGPlus - Overview

EnterpriseDB has enhanced ECPG (the PostgreSQL pre-compiler) to create ECPGPlus. ECPGPlus is a Pro*C-compatible version of the PostgreSQL C pre-compiler. ECPGPlus translates a program that combines C code and embedded SQL statements into an equivalent C program. As it performs the translation, ECPGPlus verifies that the syntax of each SQL construct is correct.

The following diagram charts the path of a program containing embedded SQL statements as it is compiled into an executable:



To produce an executable from a C program that contains embedded SQL statements, pass the program (`my_program.pgc` in the diagram above) to the ECPGPlus pre-compiler. ECPGPlus translates each SQL statement in `my_program.pgc` into C code that calls the `ecplib` API, and produces a C program (`my_program.c`). Then, pass the C program to a C compiler; the C compiler generates an object file (`my_program.o`). Finally, pass the object file (`my_program.o`), as well as the `ecplib` library file, and any other required libraries to the linker, which in turn produces the executable (`my_program`).

While the ECPGPlus preprocessor validates the *syntax* of each SQL statement, it cannot validate the *semantics*. For example, the preprocessor will confirm that an `INSERT` statement is syntactically correct, but it cannot confirm that the table mentioned in the `INSERT` statement actually exists.

Behind the Scenes

A client application contains a mix of C code and SQL code comprised of the following elements:

- C preprocessor directives
- C declarations (variables, types, functions, ...)
- C definitions (variables, types, functions, ...)
- SQL preprocessor directives
- SQL statements

For example:

```

1 #include <stdio.h>
2 EXEC SQL INCLUDE sqlca;
3
4 extern void printInt(char *label, int val);
5 extern void printStr(char *label, char *val);
6 extern void printFloat(char *label, float val);
7
8 void displayCustomer(int custNumber)
9 {
10 EXEC SQL BEGIN DECLARE SECTION;
11   VARCHAR custName[50];
12   float custBalance;
13   int custID = custNumber;
14 EXEC SQL END DECLARE SECTION;
15
16 EXEC SQL SELECT name, balance
17   INTO :custName, :custBalance
  
```

```

18 FROM customer
19 WHERE id = :custID;
20
21 printInt("ID", custID);
22 printStr("Name", custName);
23 printFloat("Balance", custBalance);
24 }
```

In the above code fragment:

- Line 1 specifies a directive to the C preprocessor.

C preprocessor directives may be interpreted or ignored; the option is controlled by a command line option (-C PROC) entered when you invoke ECPGPlus. In either case, ECPGPlus copies each C preprocessor directive to the output file (4) without change; any C preprocessor directive found in the source file will appear in the output file.

- Line 2 specifies a directive to the SQL preprocessor.

SQL preprocessor directives are interpreted by the ECPGPlus preprocessor, and are not copied to the output file.

- Lines 4 through 6 contain C declarations.

C declarations are copied to the output file without change, except that each VARCHAR declaration is translated into an equivalent struct declaration.

- Lines 10 through 14 contain an embedded-SQL declaration section.

C variables that you refer to within SQL code are known as host variables. If you invoke the ECPGPlus preprocessor in Pro*C mode (-C PROC), you may refer to any C variable within a SQL statement; otherwise you must declare each host variable within a BEGIN/END DECLARATION SECTION pair.

- Lines 16 through 19 contain a SQL statement.

SQL statements are translated into calls to the ECPGPlus run-time library.

- Lines 21 through 23 contain C code.

C code is copied to the output file without change.

Any SQL statement must be prefixed with EXEC SQL and extends to the next (unquoted) semicolon. For example:

```

printf("Updating employee salaries\n");

EXEC SQL UPDATE emp SET sal = sal * 1.25;
EXEC SQL COMMIT;

printf("Employee salaries updated\n");
```

When the preprocessor encounters the code fragment shown above, it passes the C code (the first line and the last line) to the output file without translation and converts each EXEC SQL statement into a call to an ecpglib function. The result would appear similar to the following:

```

printf("Updating employee salaries\n");

{
    ECPGdo( __LINE__, 0, 1, NULL, 0, ECPGst_normal,
            "update emp set sal = sal * 1.25",
            ECPGt_EOIT, ECPGt_EORT);
```

```

}

{
    ECPGtrans(__LINE__, NULL, "commit");
}

printf("Employee salaries updated\n");

```

Installation and Configuration

On Windows, ECPGPlus is installed by the Advanced Server installation wizard as part of the [Database Server](#) component. On Linux, install with the `edb-asxx-server-devel` RPM package where `xx` is the Advanced Server version number. By default, the executable is located in:

On Windows:

```
C:\Program Files\edb\as13\bin
```

On Linux:

```
/usr/edb/as13/bin
```

When invoking the ECPGPlus compiler, the executable must be in your search path (`%PATH%` on Windows, `$PATH` on Linux). For example, the following commands set the search path to include the directory that holds the ECPGPlus executable file `ecpg`.

On Windows:

```
set EDB_PATH=C:\Program Files\edb\as13\bin
set PATH=%EDB_PATH%;%PATH%
```

On Linux:

```
export EDB_PATH=/usr/edb/as13/bin
export PATH=$EDB_PATH:$PATH
```

Constructing a Makefile

A [makefile](#) contains a set of instructions that tell the `make` utility how to transform a program written in C (that contains embedded SQL) into a C program. To try the examples in this guide, you will need:

- a C compiler (and linker)
- the `make` utility
- ECPGPlus preprocessor and library
- a [makefile](#) that contains instructions for ECPGPlus

The following code is an example of a [makefile](#) for the samples included in this guide. To use the sample code, save it in a file named `makefile` in the directory that contains the source code file.

```
INCLUDES = -I$(shell pg_config --includedir)
LIBPATH = -L $(shell pg_config --libdir)
CFLAGS += $(INCLUDES) -g
LDFLAGS += -g
LDLIBS += $(LIBPATH) -lecpg -lpq
```

```
.SUFFIXES: .pgc,.pc
```

```
.pgc.c:  
    ecpg -c $(INCLUDES) $?
```

```
.pc.c:  
    ecpg -C PROC -c $(INCLUDES) $?
```

The first two lines use the `pg_config` program to locate the necessary header files and library directories:

```
INCLUDES = -I$(shell pg_config --includedir)  
LIBPATH = -L $(shell pg_config --libdir)
```

The `pg_config` program is shipped with Advanced Server.

`make` knows that it should use the `CFLAGS` variable when running the C compiler and `LDFLAGS` and `LDLIBS` when invoking the linker. ECPG programs must be linked against the ECPG run-time library (`-lecpq`) and the libpq library (`-lpq`)

```
CFLAGS += $(INCLUDES) -g  
LDFLAGS += -g  
LDLIBS += $(LIBPATH) -lecpq -lpq
```

The sample `makefile` instructs make how to translate a `.pgc` or a `.pc` file into a C program. Two lines in the `makefile` specify the mode in which the source file will be compiled. The first compile option is:

```
.pgc.c:  
    ecpg -c $(INCLUDES) $?
```

The first option tells `make` how to transform a file that ends in `.pgc` (presumably, an ECPG source file) into a file that ends in `.c` (a C program), using community ECPG (without the ECPGPlus enhancements). It invokes the ECPG pre-compiler with the `-c` flag (instructing the compiler to convert SQL code into C), using the value of the `INCLUDES` variable and the name of the `.pgc` file.

```
.pc.c:  
    ecpg -C PROC -c $(INCLUDES) $?
```

The second option tells make how to transform a file that ends in `.pg` (an ECPG source file) into a file that ends in `.c` (a C program), using the ECPGPlus extensions. It invokes the ECPG pre-compiler with the `-c` flag (instructing the compiler to convert SQL code into C), as well as the `-C PROC` flag (instructing the compiler to use ECPGPlus in Pro*C-compatibility mode), using the value of the `INCLUDES` variable and the name of the `.pgc` file.

When you run `make`, pass the name of the ECPG source code file you wish to compile. For example, to compile an ECPG source code file named `customer_list.pgc`, use the command:

```
| make customer_list
```

The `make` utility consults the `makefile` (located in the current directory), discovers that the `makefile` contains a rule that will compile `customer_list.pgc` into a C program (`customer_list.c`), and then uses the rules built into `make` to compile `customer_list.c` into an executable program.

ECPGPlus Command Line Options

In the sample `makefile` shown above, `make` includes the `-C` option when invoking ECPGPlus to specify that ECPGPlus should be invoked in Pro*C compatible mode.

If you include the `-C PROC` keywords on the command line, in addition to the ECPG syntax, you may use Pro*C command line syntax; for example:

```
$ ecpg -C PROC INCLUDE=/usr/edb/as13/include acct_update.c
```

To display a complete list of the other ECPGPlus options available, navigate to the ECPGPlus installation directory, and enter:

```
./ecpg --help
```

The command line options are:

| Option | Description |
|-----------------|---|
| <code>-c</code> | Automatically generate C code from embedded SQL code. |

Use the `-C` option to specify a compatibility mode:

`INFORMIX`

`-C mode`

`INFORMIX_SE`

`PROC`

Define a preprocessor *symbol*.

`-D symbol`

The `-D` keyword is not supported when compiling in *PROC mode*. Instead, use the Oracle-style '`DEFINE=`' clause.

`-h`

Parse a header file, this option includes option `c`.

`-i`

Parse system, include files as well.

`-I directory`

Search *directory* for include files.

`-o outfile`

Write the result to *outfile*.

Specify run-time behavior; *option* can be:

`no_indicator` - Do not use indicators, but instead use special values to represent NULL values.

`-r option`

`prepare` - Prepare all statements before using them.

`questionmarks` - Allow use of a question mark as a placeholder.

`usebulk` - Enable bulk processing for INSERT, UPDATE and DELETE statements that operate on host variable arrays.

--

`regression` Run in regression testing mode.

`-t` Turn on `autocommit` of transactions.

`-l` Disable `#line` directives.

`--help` Display the help options.

`--version` Output version information.

Note

If you do not specify an output file name when invoking ECPGPlus, the output file name is created by stripping off the `.pgc` filename extension, and appending `.c` to the file name.

9.3 Using Embedded SQL

Each of the following sections leads with a code sample, followed by an explanation of each section within the code sample.

Example - A Simple Query

The first code sample demonstrates how to execute a `SELECT` statement (which returns a single row), storing the results in a group of host variables. After declaring host variables, it connects to the `edb` sample database using a hard-coded role name and the associated password, and queries the `emp` table. The query returns the values into the declared host variables; after checking the value of the `NULL` indicator variable, it prints a simple result set onscreen and closes the connection.

```
*****
* print_emp.pgc
*
*/
#include <stdio.h>

int main(void)
{
    EXEC SQL BEGIN DECLARE SECTION;
    int v_empno;
    char v_ename[40];
    double v_sal;
    double v_comm;
    short v_comm_ind;
    EXEC SQL END DECLARE SECTION;

    EXEC SQL WHENEVER SQLERROR sqlprint;

    EXEC SQL CONNECT TO edb
        USER 'alice' IDENTIFIED BY '1safePWD';

    EXEC SQL
        SELECT
            empno, ename, sal, comm
        INTO
            :v_empno, :v_ename, :v_sal, :v_comm INDICATOR:v_comm_ind
        FROM
            emp
        WHERE
            empno = 7369;

    if (v_comm_ind)
        printf("empno(%d), ename(%s), sal(%f) comm(NULL)\n",
               v_empno, v_ename, v_sal);
    else
        printf("empno(%d), ename(%s), sal(%f) comm(%f)\n",
               v_empno, v_ename, v_sal, v_comm);
    EXEC SQL DISCONNECT;
```

```
}
```

The code sample begins by including the prototypes and type definitions for the C `stdio` library, and then declares the `main` function:

```
#include <stdio.h>

int main(void)
{
```

Next, the application declares a set of host variables used to interact with the database server:

```
EXEC SQL BEGIN DECLARE SECTION;
int v_empno;
char v_ename[40];
double v_sal;
double v_comm;
short v_comm_ind;
EXEC SQL END DECLARE SECTION;
```

Please note that if you plan to pre-compile the code in `PROC` mode, you may omit the `BEGIN DECLARE...END DECLARE` section. For more information about declaring host variables, refer to the [Declaring Host Variables](#).

The data type associated with each variable within the declaration section is a C data type. Data passed between the server and the client application must share a compatible data type; for more information about data types, see the [Supported C Data Types](#).

The next statement instructs the server how to handle an error:

```
| EXEC SQL WHENEVER SQLERROR sqlprint;
```

If the client application encounters an error in the SQL code, the server will print an error message to `stderr` (standard error), using the `sqlprint()` function supplied with `ecpglib`. The next `EXEC SQL` statement establishes a connection with Advanced Server:

```
| EXEC SQL CONNECT TO edb
```

```
| | USER 'alice' IDENTIFIED BY '1safePWD';
```

In our example, the client application connects to the `edb` database, using a role named `alice` with a password of `1safePWD`.

The code then performs a query against the `emp` table:

```
EXEC SQL
SELECT
    empno, ename, sal, comm
INTO
    :v_empno, :v_ename, :v_sal, :v_comm INDICATOR :v_comm_ind
FROM
    emp
WHERE
    empno = 7369;
```

The query returns information about employee number `7369`.

The `SELECT` statement uses an `INTO` clause to assign the retrieved values (from the `empno`, `ename`, `sal` and `comm` columns) into the `:v_empno`, `:v_ename`, `:v_sal` and `:v_comm` host variables (and the `:v_comm_ind` null

indicator). The first value retrieved is assigned to the first variable listed in the `INTO` clause, the second value is assigned to the second variable, and so on.

The `comm` column contains the commission values earned by an employee, and could potentially contain a `NULL` value. The statement includes the `INDICATOR` keyword, and a host variable to hold a null indicator.

The code checks the null indicator, and displays the appropriate on-screen results:

```
if (v_comm_ind)
    printf("empno(%d), ename(%s), sal(%f) comm(%c)\n",
        v_empno, v_ename, v_sal);
else
    printf("empno(%d), ename(%s), sal(%f) comm(%f)\n",
        v_empno, v_ename, v_sal, v_comm);
```

If the null indicator is `0` (that is, `false`), the `comm` column contains a meaningful value, and the `printf` function displays the commission. If the null indicator contains a non-zero value, `comm` is `NULL`, and `printf` displays a value of `NULL`. Please note that a host variable (other than a null indicator) contains no meaningful value if you fetch a `NULL` into that host variable; you must use null indicators to identify any value which may be `NULL`.

The final statement in the code sample closes the connection to the server:

```
EXEC SQL DISCONNECT;
}
```

Using Indicator Variables

The previous example included an *indicator variable* that identifies any row in which the value of the `comm` column (when returned by the server) was `NULL`. An indicator variable is an extra host variable that denotes if the content of the preceding variable is `NULL` or truncated. The indicator variable is populated when the contents of a row are stored. An indicator variable may contain the following values:

| Indicator Value | Denotes |
|---|---|
| If an indicator variable is less than 0. | The value returned by the server was <code>NULL</code> . |
| If an indicator variable is equal to 0. | The value returned by the server was not <code>NULL</code> , and was not truncated. |
| If an indicator variable is greater than 0. | The value returned by the server was truncated when stored in the host variable. |

When including an indicator variable in an `INTO` clause, you are not required to include the optional `INDICATOR` keyword.

You may omit an indicator variable if you are certain that a query will never return a `NULL` value into the corresponding host variable. If you omit an indicator variable and a query returns a `NULL` value, `ecpglib` will raise a run-time error.

Declaring Host Variables

You can use a *host variable* in a SQL statement at any point that a value may appear within that statement. A host variable is a C variable that you can use to pass data values from the client application to the server, and return data from the server to the client application. A host variable can be:

- an array
- a `typedef`

- a pointer
- a `struct`
- any scalar C data type

The code fragments that follow demonstrate using host variables in code compiled in `PROC` mode, and in non-`PROC` mode. The SQL statement adds a row to the `dept` table, inserting the values returned by the variables `v_deptno`, `v_dname` and `v_loc` into the `deptno` column, the `dname` column and the `loc` column, respectively.

If you are compiling in `PROC` mode, you may omit the `EXEC SQL BEGIN DECLARE SECTION` and `EXEC SQL END DECLARE SECTION` directives. `PROC` mode permits you to use C function parameters as host variables:

```
void addDept(int v_deptno, char v_dname, char v_loc)
{
    EXEC SQL INSERT INTO dept VALUES( :v_deptno, :v_dname, :v_loc);
}
```

If you are not compiling in `PROC` mode, you must wrap embedded variable declarations with the `EXEC SQL BEGIN DECLARE SECTION` and the `EXEC SQL END DECLARE SECTION` directives, as shown below:

```
void addDept(int v_deptno, char v_dname, char v_loc)
{
    EXEC SQL BEGIN DECLARE SECTION;
        int v_deptno_copy = v_deptno;
        char v_dname_copy[14+1] = v_dname;
        char v_loc_copy[13+1] = v_loc;
    EXEC SQL END DECLARE SECTION;

    EXEC SQL INSERT INTO dept VALUES( :v_deptno, :v_dname, :v_loc);
}
```

You can also include the `INTO` clause in a `SELECT` statement to use the host variables to retrieve information:

```
EXEC SQL SELECT deptno, dname, loc
    INTO :v_deptno, :v_dname, v_loc FROM dept;
```

Each column returned by the `SELECT` statement must have a type-compatible target variable in the `INTO` clause. This is a simple example that retrieves a single row; to retrieve more than one row, you must define a cursor, as demonstrated in the next example.

Example - Using a Cursor to Process a Result Set

The code sample that follows demonstrates using a cursor to process a result set. There are four basic steps involved in creating and using a cursor:

1. Use the `DECLARE CURSOR` statement to define a cursor.
2. Use the `OPEN CURSOR` statement to open the cursor.
3. Use the `FETCH` statement to retrieve data from a cursor.
4. Use the `CLOSE CURSOR` statement to close the cursor.

After declaring host variables, our example connects to the `edb` database using a user-supplied role name and password, and queries the `emp` table. The query returns the values into a cursor named `employees`. The code sample then opens the cursor, and loops through the result set a row at a time, printing the result set. When the sample detects the end of the result set, it closes the connection.

```
*****
* print_emps.pgc
*
```

```
/*
#include <stdio.h>

int main(int argc, char *argv[])
{
    EXEC SQL BEGIN DECLARE SECTION;
    char *username = argv[1];
    char *password = argv[2];
    int v_empno;
    char v_ename[40];
    double v_sal;
    double v_comm;
    short v_comm_ind;
    EXEC SQL END DECLARE SECTION;

    EXEC SQL WHENEVER SQLERROR sqlprint;

    EXEC SQL CONNECT TO edb USER :username IDENTIFIED BY :password;

    EXEC SQL DECLARE employees CURSOR FOR
        SELECT
            empno, ename, sal, comm
        FROM
            emp;

    EXEC SQL OPEN employees;

    EXEC SQL WHENEVER NOT FOUND DO break;

    for (;;)
    {
        EXEC SQL FETCH NEXT FROM employees
        INTO
            :v_empno, :v_ename, :v_sal, :v_comm INDICATOR :v_comm_ind;

        if (v_comm_ind)
            printf("empno(%d), ename(%s), sal(%f) comm(%f)\n",
                v_empno, v_ename, v_sal);
        else
            printf("empno(%d), ename(%s), sal(%f) comm(%f)\n",
                v_empno, v_ename, v_sal, v_comm);
    }
    EXEC SQL CLOSE employees;
    EXEC SQL DISCONNECT;
}
*****
```

The code sample begins by including the prototypes and type definitions for the C `stdio` library, and then declares the `main` function:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
```

Next, the application declares a set of host variables used to interact with the database server:

```
EXEC SQL BEGIN DECLARE SECTION;
char *username = argv[1];
char *password = argv[2];
int v_empno;
char v_ename[40];
double v_sal;
double v_comm;
short v_comm_ind;
EXEC SQL END DECLARE SECTION;
```

`argv[]` is an array that contains the command line arguments entered when the user runs the client application. `argv[1]` contains the first command line argument (in this case, a `username`), and `argv[2]` contains the second command line argument (a `password`); please note that we have omitted the error-checking code you would normally include in a real-world application. The declaration initializes the values of `username` and `password`, setting them to the values entered when the user invoked the client application.

You may be thinking that you could refer to `argv[1]` and `argv[2]` in a SQL statement (instead of creating a separate copy of each variable); that will not work. All host variables must be declared within a `BEGIN/END DECLARE SECTION` (unless you are compiling in `PROC` mode). Since `argv` is a function *parameter* (not an automatic variable), it cannot be declared within a `BEGIN/END DECLARE SECTION`. If you are compiling in `PROC` mode, you can refer to *any* C variable within a SQL statement.

The next statement instructs the server to respond to an SQL error by printing the text of the error message returned by ECPGPlus or the database server:

```
EXEC SQL WHENEVER SQLERROR sqlprint;
```

Then, the client application establishes a connection with Advanced Server:

```
EXEC SQL CONNECT TO edb USER :username IDENTIFIED BY :password;
```

The `CONNECT` statement creates a connection to the `edb` database, using the values found in the `:username` and `:password` host variables to authenticate the application to the server when connecting.

The next statement declares a cursor named `employees`:

```
EXEC SQL DECLARE employees CURSOR FOR
SELECT
    empno, ename, sal, comm
FROM
    emp;
```

`employees` will contain the result set of a `SELECT` statement on the `emp` table. The query returns employee information from the following columns: `empno`, `ename`, `sal` and `comm`. Notice that when you declare a cursor, you do not include an `INTO` clause - instead, you specify the target variables (or descriptors) when you `FETCH` from the cursor.

Before fetching rows from the cursor, the client application must `OPEN` the cursor:

```
EXEC SQL OPEN employees;
```

In the subsequent `FETCH` section, the client application will loop through the contents of the cursor; the client application includes a `WHENEVER` statement that instructs the server to `break` (that is, terminate the loop) when it reaches the end of the cursor:

```
EXEC SQL WHENEVER NOT FOUND DO break;
```

The client application then uses a `FETCH` statement to retrieve each row from the cursor `INTO` the previously declared host variables:

```
for (;;) {
    EXEC SQL FETCH NEXT FROM employees
    INTO
        :v_empno, :v_ename, :v_sal, :v_comm INDICATOR :v_comm_ind;
```

The `FETCH` statement uses an `INTO` clause to assign the retrieved values into the `:v_empno`, `:v_ename`, `:v_sal` and `:v_comm` host variables (and the `:v_comm_ind` null indicator). The first value in the cursor is assigned to the first variable listed in the `INTO` clause, the second value is assigned to the second variable, and so on.

The `FETCH` statement also includes the `INDICATOR` keyword and a host variable to hold a null indicator. If the `comm` column for the retrieved record contains a `NULL` value, `v_comm_ind` is set to a non-zero value, indicating that the column is `NULL`.

The code then checks the null indicator, and displays the appropriate on-screen results:

```
if (v_comm_ind)
    printf("empno(%d), ename(%s), sal(%,.2f) comm(NULL)\n",
           v_empno, v_ename, v_sal);
else
    printf("empno(%d), ename(%s), sal(%,.2f) comm(%,.2f)\n",
           v_empno, v_ename, v_sal, v_comm);
}
```

If the null indicator is `0` (that is, `false`), `v_comm` contains a meaningful value, and the `printf` function displays the commission. If the null indicator contains a non-zero value, `comm` is `NULL`, and `printf` displays the string '`NULL`'. Please note that a host variable (other than a null indicator) contains no meaningful value if you fetch a `NULL` into that host variable; you must use null indicators for any value which may be `NULL`.

The final statements in the code sample close the cursor (`employees`), and the connection to the server:

```
EXEC SQL CLOSE employees;
EXEC SQL DISCONNECT;
```

9.4 Using Descriptors

Dynamic SQL allows a client application to execute SQL statements that are composed at runtime. This is useful when you don't know the content or form a statement will take when you are writing a client application.

ECPGPlus does *not* allow you to use a host variable in place of an identifier (such as a table name, column name or index name); instead, you should use dynamic SQL statements to build a string that includes the information, and then execute that string. The string is passed between the client and the server in the form of a *descriptor*. A descriptor is a data structure that contains both the data and the information about the shape of the data.

A client application must use a `GET DESCRIPTOR` statement to retrieve information from a descriptor. The following steps describe the basic flow of a client application using dynamic SQL:

1. Use an `ALLOCATE DESCRIPTOR` statement to allocate a descriptor for the result set (select list).
2. Use an `ALLOCATE DESCRIPTOR` statement to allocate a descriptor for the input parameters (bind variables).
3. Obtain, assemble or compute the text of an SQL statement.
4. Use a `PREPARE` statement to parse and syntax-check the SQL statement.
5. Use a `DESCRIBE` statement to describe the select list into the select-list descriptor.
6. Use a `DESCRIBE` statement to describe the input parameters into the bind-variables descriptor.
7. Prompt the user (if required) for a value for each input parameter. Use a `SET DESCRIPTOR` statement to assign the values into a descriptor.

8. Use a **DECLARE CURSOR** statement to define a cursor for the statement.
9. Use an **OPEN CURSOR** statement to open a cursor for the statement.
10. Use a **FETCH** statement to fetch each row from the cursor, storing each row in select-list descriptor.
11. Use a **GET DESCRIPTOR** command to interrogate the select-list descriptor to find the value of each column in the current row.
12. Use a **CLOSE CURSOR** statement to close the cursor and free any cursor resources.

A descriptor may contain the attributes listed in the table below:

| Field | Type | Attribute Description |
|-----------------------------|---------|--|
| CARDINALITY | integer | The number of rows in the result set. |
| DATA | N/A | The data value. |
| | | If TYPE is 9: |
| DATETIME_INTERVAL_CODE | integer | 1 - DATE 2 - TIME 3 - TIMESTAMP 4 - TIME WITH TIMEZONE 5 - TIMESTAMP WITH TIMEZONE |
| DATETIME_INTERVAL_PRECISION | integer | Unused. |
| INDICATOR | integer | Indicates a NULL or truncated value. |
| KEY_MEMBER | integer | Unused (returns FALSE). |
| LENGTH | integer | The data length (as stored on server). |
| NAME | string | The name of the column in which the data resides. |
| NULLABLE | integer | Unused (returns TRUE). |
| OCTET_LENGTH | integer | The data length (in bytes) as stored on server. |
| PRECISION | integer | The data precision (if the data is of numeric type). |
| RETURNED_LENGTH | integer | Actual length of data item. |
| RETURNED_OCTET_LENGTH | integer | Actual length of data item. |
| SCALE | integer | The data scale (if the data is of numeric type). |

| Field | Type | Attribute Description |
|-------|---------|--|
| TYPE | integer | A numeric code that represents the data type of the column: 1 - SQL3_CHARACTER 2 - SQL3_NUMERIC 3 - SQL3_DECIMAL 4 - SQL3_INTEGER 5 - SQL3_SMALLINT 6 - SQL3_FLOAT 7 - SQL3_REAL 8 - SQL3_DOUBLE_PRECISION 9 - SQL3_DATE_TIME_TIMESTAMP 10 - SQL3_INTERVAL 12 - SQL3_CHARACTER_VARYING 13 - SQL3_ENUMERATED 14 - SQL3_BIT 15 - SQL3_BIT_VARYING 16 - SQL3_BOOLEAN |

Example - Using a Descriptor to Return Data

The following simple application executes an SQL statement entered by an end user. The code sample demonstrates:

- how to use a SQL descriptor to execute a `SELECT` statement.
- how to find the data and metadata returned by the statement.

The application accepts an SQL statement from an end user, tests the statement to see if it includes the `SELECT` keyword, and executes the statement.

When invoking the application, an end user must provide the name of the database on which the SQL statement will be performed, and a string that contains the text of the query.

For example, a user might invoke the sample with the following command:

```
./exec_stmt edb "SELECT * FROM emp"
```

```
*****
/* exec_stmt.pgc
*
*/
```

```

#include <stdio.h>
#include <stdlib.h>
#include <sql3types.h>
#include <sqlca.h>

EXEC SQL WHENEVER SQLERROR SQLPRINT;
static void print_meta_data( char * desc_name );

char *md1 = "col field      data      ret";
char *md2 = "num name      type      len";
char *md3 = "----- ----- ----- -----";

int main( int argc, char *argv[] )
{
EXEC SQL BEGIN DECLARE SECTION;
char *db = argv[1];
char *stmt = argv[2];
int col_count;
EXEC SQL END DECLARE SECTION;

EXEC SQL CONNECT TO :db;

EXEC SQL ALLOCATE DESCRIPTOR parse_desc;
EXEC SQL PREPARE query FROM :stmt;
EXEC SQL DESCRIBE query INTO SQL DESCRIPTOR parse_desc;
EXEC SQL GET DESCRIPTOR 'parse_desc' :col_count = COUNT;

if( col_count == 0 )
{
EXEC SQL EXECUTE IMMEDIATE :stmt;

if( sqlca.sqlcode >= 0 )
  EXEC SQL COMMIT;
}
else
{
int row;

EXEC SQL ALLOCATE DESCRIPTOR row_desc;
EXEC SQL DECLARE my_cursor CURSOR FOR query;
EXEC SQL OPEN my_cursor;

for( row = 0; ; row++ )
{
EXEC SQL BEGIN DECLARE SECTION;
int col;
EXEC SQL END DECLARE SECTION;
EXEC SQL FETCH IN my_cursor
  INTO SQL DESCRIPTOR row_desc;

if( sqlca.sqlcode != 0 )
  break;

if( row == 0 )

```

```

print_meta_data( "row_desc" );

printf("[RECORD %d]\n", row+1);

for( col = 1; col <= col_count; col++ )
{
    EXEC SQL BEGIN DECLARE SECTION;
    short ind;
    varchar val[40+1];
    varchar name[20+1];
    EXEC SQL END DECLARE SECTION;

    EXEC SQL GET DESCRIPTOR 'row_desc'
        VALUE :col
        :val = DATA, :ind = INDICATOR, :name = NAME;

    if( ind == -1 )
        printf( " %-20s : <null>\n", name.arr );
    else if( ind > 0 )
        printf( " %-20s : <truncated>\n", name.arr );
    else
        printf( " %-20s : %s\n", name.arr, val.arr );
}

printf( "\n" );

}

printf( "%d rows\n", row );
}

exit( 0 );
}

static void print_meta_data( char *desc_name )
{
EXEC SQL BEGIN DECLARE SECTION;
char *desc = desc_name;
int col_count;
int col;
EXEC SQL END DECLARE SECTION;

static char *types[] =
{
"unused      ",
"CHARACTER   ",
"NUMERIC    ",
"DECIMAL     ",
"INTEGER     ",
"SMALLINT   ",
"FLOAT       ",
"REAL        ",
"DOUBLE      ",
"DATE_TIME   ",
"INTERVAL   ",
"unused      ",
}

```

```

"CHARACTER_VARYING",
"ENUMERATED    ",
"BIT        ",
"BIT_VARYING   ",
"BOOLEAN      ",
"abstract     "
};

EXEC SQL GET DESCRIPTOR :desc :col_count = count;

printf( "%s\n", md1 );
printf( "%s\n", md2 );
printf( "%s\n", md3 );

for( col = 1; col <= col_count; col++ )
{

EXEC SQL BEGIN DECLARE SECTION;
int   type;
int   ret_len;
varchar name[21];
EXEC SQL END DECLARE SECTION;
char *type_name;

EXEC SQL GET DESCRIPTOR :desc
VALUE :col
:name = NAME,
:type = TYPE,
:ret_len = RETURNED_OCTET_LENGTH;

if( type > 0 && type < SQL3_abstract )
  type_name = types[type];
else
  type_name = "unknown";

printf( "%02d: %-20s %-17s %04d\n",
  col, name.arr, type_name, ret_len );
}
printf( "\n");
}

*****

```

The code sample begins by including the prototypes and type definitions for the C `stdio` and `stdlib` libraries, SQL data type symbols, and the `SQLCA` (SQL communications area) structure:

```

#include <stdio.h>
#include <stdlib.h>
#include <sql3types.h>
#include <sqlca.h>

```

The sample provides minimal error handling; when the application encounters an SQL error, it prints the error message to screen:

```
EXEC SQL WHENEVER SQLERROR SQLPRINT;
```

The application includes a forward-declaration for a function named `print_meta_data()` that will print the metadata found in a descriptor:

```
static void print_meta_data( char * desc_name );
```

The following code specifies the column header information that the application will use when printing the metadata:

```
char *md1 = "col field data ret";
char *md2 = "num name type len";
char *md3 = "----- ---";
```

```
int main( int argc, char *argv[] )
{
```

The following declaration section identifies the host variables that will contain the name of the database to which the application will connect, the content of the SQL Statement, and a host variable that will hold the number of columns in the result set (if any).

```
EXEC SQL BEGIN DECLARE SECTION;
char *db = argv[1];
char *stmt = argv[2];
int col_count;
EXEC SQL END DECLARE SECTION;
```

The application connects to the database (using the default credentials):

```
EXEC SQL CONNECT TO :db;
```

Next, the application allocates an SQL descriptor to hold the metadata for a statement:

```
EXEC SQL ALLOCATE DESCRIPTOR parse_desc;
```

The application uses a `PREPARE` statement to syntax check the string provided by the user:

```
EXEC SQL PREPARE query FROM :stmt;
```

and a `DESCRIBE` statement to move the metadata for the query into the SQL descriptor.

```
EXEC SQL DESCRIBE query INTO SQL DESCRIPTOR parse_desc;
```

Then, the application interrogates the descriptor to discover the number of columns in the result set, and stores that in the host variable `col_count`.

```
EXEC SQL GET DESCRIPTOR parse_desc :col_count = COUNT;
```

If the column count is zero, the end user did not enter a `SELECT` statement; the application uses an `EXECUTE IMMEDIATE` statement to process the contents of the statement:

```
if( col_count == 0 )
{
    EXEC SQL EXECUTE IMMEDIATE :stmt;
```

If the statement executes successfully, the application performs a `COMMIT`:

```
if( sqlca.sqlcode >= 0 )
    EXEC SQL COMMIT;
```

```

}
else
{

```

If the statement entered by the user is a **SELECT** statement (which we know because the column count is non-zero), the application declares a variable named **row**.

```
int row;
```

Then, the application allocates another descriptor that holds the description and the values of a specific row in the result set:

```
EXEC SQL ALLOCATE DESCRIPTOR row_desc;
```

The application declares and opens a cursor for the prepared statement:

```
EXEC SQL DECLARE my_cursor CURSOR FOR query;
EXEC SQL OPEN my_cursor;
```

Loops through the rows in result set:

```

for( row = 0; ; row++ )
{
    EXEC SQL BEGIN DECLARE SECTION;
    int col;
    EXEC SQL END DECLARE SECTION;

```

Then, uses a **FETCH** to retrieve the next row from the cursor into the descriptor:

```
EXEC SQL FETCH IN my_cursor INTO SQL DESCRIPTOR row_desc;
```

The application confirms that the **FETCH** did not fail; if the **FETCH** fails, the application has reached the end of the result set, and breaks the loop:

```
if( sqlca.sqlcode != 0 )
    break;
```

The application checks to see if this is the first row of the cursor; if it is, the application prints the metadata for the row.

```
if( row == 0 )
    print_meta_data( "row_desc" );
```

Next, it prints a record header containing the row number:

```
printf("[RECORD %d]\n", row+1);
```

Then, it loops through each column in the row:

```

for( col = 1; col <= col_count; col++ )
{
    EXEC SQL BEGIN DECLARE SECTION;
    short ind;
    varchar val[40+1];
    varchar name[20+1];
    EXEC SQL END DECLARE SECTION;

```

The application interrogates the row descriptor (`row_desc`) to copy the column value (`:val`), null indicator (`:ind`) and column name (`:name`) into the host variables declared above. Notice that you can retrieve multiple items from a descriptor using a comma-separated list.

```
EXEC SQL GET DESCRIPTOR row_desc
VALUE :col
:val = DATA, :ind = INDICATOR, :name = NAME;
```

If the null indicator (`ind`) is negative, the column value is `NULL`; if the null indicator is greater than `0`, the column value is too long to fit into the `val` host variable (so we print `<truncated>`); otherwise, the null indicator is `0` (meaning `NOT NULL`) so we print the value. In each case, we prefix the value (or `<null>` or `<truncated>`) with the name of the column.

```
if( ind == -1 )
printf( " %-20s : <null>\n", name.arr );
else if( ind > 0 )
printf( " %-20s : <truncated>\n", name.arr );
else
printf( " %-20s : %s\n", name.arr, val.arr );
}

printf( "\n" );
}
```

When the loop terminates, the application prints the number of rows fetched, and exits:

```
printf( "%d rows\n", row );
}

exit( 0 );
}
```

The `print_meta_data()` function extracts the metadata from a descriptor and prints the name, data type, and length of each column:

```
static void print_meta_data( char *desc_name )
{
```

The application declares host variables:

```
EXEC SQL BEGIN DECLARE SECTION;
char *desc = desc_name;
int col_count;
int col;
EXEC SQL END DECLARE SECTION;
```

The application then defines an array of character strings that map data type values (`numeric`) into data type names. We use the numeric value found in the descriptor to index into this array. For example, if we find that a given column is of type `2`, we can find the name of that type (`NUMERIC`) by writing `types[2]`.

```
static char *types[] =
{
"unused ",
"CHARACTER ",
"NUMERIC ",
"DECIMAL ",
"INTEGER ",
```

```

"SMALLINT",
"FLOAT",
"REAL",
"DOUBLE",
"DATE_TIME",
"INTERVAL",
"unused",
"CHARACTER_VARYING",
"ENUMERATED",
"BIT",
"BIT_VARYING",
"BOOLEAN",
"abstract"
};


```

The application retrieves the column count from the descriptor. Notice that the program refers to the descriptor using a host variable (`desc`) that contains the name of the descriptor. In most scenarios, you would use an identifier to refer to a descriptor, but in this case, the caller provided the descriptor name, so we can use a host variable to refer to the descriptor.

```
EXEC SQL GET DESCRIPTOR :desc :col_count = count;
```

The application prints the column headers (defined at the beginning of this application):

```

printf( "%s\n", md1 );
printf( "%s\n", md2 );
printf( "%s\n", md3 );

```

Then, loops through each column found in the descriptor, and prints the name, type and length of each column.

```

for( col = 1; col <= col_count; col++ )
{
    EXEC SQL BEGIN DECLARE SECTION;
    int type;
    int ret_len;
    varchar name[21];
    EXEC SQL END DECLARE SECTION;
    char *type_name;

```

It retrieves the name, type code, and length of the current column:

```

EXEC SQL GET DESCRIPTOR :desc
VALUE :col
:name = NAME,
:type = TYPE,
:ret_len = RETURNED_OCTET_LENGTH;

```

If the numeric type code matches a 'known' type code (that is, a type code found in the `types[]` array), it sets `type_name` to the name of the corresponding type; otherwise, it sets `type_name` to "unknown".

```

if( type > 0 && type < SQL3_abstract )
    type_name = types[type];
else
    type_name = "unknown";

```

and prints the column number, name, type name, and length:

```

printf( "%02d: %-20s %-17s %04d\n",
    col, name.arr, type_name, ret_len );
}
printf( "\n" );
}

```

If you invoke the sample application with the following command:

```
./exec_stmt test "SELECT * FROM emp WHERE empno IN(7902, 7934)"
```

The application returns:

| col field num name | data type | ret len |
|-----------------------|-------------------|---------|
| 01: empno | NUMERIC | 0004 |
| 02: ename | CHARACTER_VARYING | 0004 |
| 03: job | CHARACTER_VARYING | 0007 |
| 04: mgr | NUMERIC | 0004 |
| 05: hiredate | DATE_TIME | 0018 |
| 06: sal | NUMERIC | 0007 |
| 07: comm | NUMERIC | 0000 |
| 08: deptno | NUMERIC | 0002 |

[RECORD 1]

```

empno      : 7902
ename       : FORD
job        : ANALYST
mgr         : 7566
hiredate    : 03-DEC-81 00:00:00
sal         : 3000.00
comm        : <null>
deptno     : 20

```

[RECORD 2]

```

empno      : 7934
ename       : MILLER
job        : CLERK
mgr         : 7782
hiredate    : 23-JAN-82 00:00:00
sal         : 1300.00
comm        : <null>
deptno     : 10

```

2 rows

9.5 Building and Executing Dynamic SQL Statements

The following examples demonstrate four techniques for building and executing dynamic SQL statements. Each example demonstrates processing a different combination of statement and input types:

- The first example demonstrates processing and executing a SQL statement that does not contain a `SELECT` statement and does not require input variables. This example corresponds to the techniques used by Oracle Dynamic SQL Method 1.
- The second example demonstrates processing and executing a SQL statement that does not contain a `SELECT` statement, and contains a known number of input variables. This example corresponds to the techniques used by Oracle Dynamic SQL Method 2.
- The third example demonstrates processing and executing a SQL statement that may contain a `SELECT` statement, and includes a known number of input variables. This example corresponds to the techniques used by Oracle Dynamic SQL Method 3.
- The fourth example demonstrates processing and executing a SQL statement that may contain a `SELECT` statement, and includes an unknown number of input variables. This example corresponds to the techniques used by Oracle Dynamic SQL Method 4.

Example - Executing a Non-query Statement Without Parameters

The following example demonstrates how to use the `EXECUTE IMMEDIATE` command to execute a SQL statement where the text of the statement is not known until you run the application. You cannot use `EXECUTE IMMEDIATE` to execute a statement that returns a result set. You cannot use `EXECUTE IMMEDIATE` to execute a statement that contains parameter placeholders.

The `EXECUTE IMMEDIATE` statement parses and plans the SQL statement each time it executes, which can have a negative impact on the performance of your application. If you plan to execute the same statement repeatedly, consider using the `PREPARE/EXECUTE` technique described in the next example.

```
*****  

#include <stdio.h>
#include <string.h>
#include <stdlib.h>

static void handle_error(void);

int main(int argc, char *argv[])
{
    char *insertStmt;

    EXEC SQL WHENEVER SQLERROR DO handle_error();

    EXEC SQL CONNECT :argv[1];

    insertStmt = "INSERT INTO dept VALUES(50, 'ACCTG', 'SEATTLE')";

    EXEC SQL EXECUTE IMMEDIATE :insertStmt;

    fprintf(stderr, "ok\n");

    EXEC SQL COMMIT RELEASE;

    exit(EXIT_SUCCESS);
}

static void handle_error(void)
{
    fprintf(stderr, "%s\n", sqlca.sqlerrm.sqlerrmc);

    EXEC SQL WHENEVER SQLERROR CONTINUE;
```

```

EXEC SQL ROLLBACK RELEASE;

exit(EXIT_FAILURE);
}

/*****

```

The code sample begins by including the prototypes and type definitions for the C `stdio`, `string`, and `stdlib` libraries, and providing basic infrastructure for the program:

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>

static void handle_error(void);
int main(int argc, char *argv[])
{
    char *insertStmt;

```

The example then sets up an error handler; ECPGPlus calls the `handle_error()` function whenever a SQL error occurs:

```
EXEC SQL WHENEVER SQLERROR DO handle_error();
```

Then, the example connects to the database using the credentials specified on the command line:

```
EXEC SQL CONNECT :argv[1];
```

Next, the program uses an `EXECUTE IMMEDIATE` statement to execute a SQL statement, adding a row to the `dept` table:

```
insertStmt = "INSERT INTO dept VALUES(50, 'ACCTG', 'SEATTLE');
```

```
EXEC SQL EXECUTE IMMEDIATE :insertStmt;
```

If the `EXECUTE IMMEDIATE` command fails for any reason, ECPGPlus will invoke the `handle_error()` function (which terminates the application after displaying an error message to the user). If the `EXECUTE IMMEDIATE` command succeeds, the application displays a message (`ok`) to the user, commits the changes, disconnects from the server, and terminates the application.

```

fprintf(stderr, "ok\n");

EXEC SQL COMMIT RELEASE;

exit(EXIT_SUCCESS);
}
```

ECPGPlus calls the `handle_error()` function whenever it encounters a SQL error. The `handle_error()` function prints the content of the error message, resets the error handler, rolls back any changes, disconnects from the database, and terminates the application.

```

static void handle_error(void)
{
    fprintf(stderr, "%s\n", sqlca.sqlerrm.sqlerrmc);

EXEC SQL WHENEVER SQLERROR CONTINUE;
EXEC SQL ROLLBACK RELEASE;
```

```

    exit(EXIT_FAILURE);
}

```

Example - Executing a Non-query Statement with a Specified Number of Placeholders

To execute a non-query command that includes a known number of parameter placeholders, you must first **PREPARE** the statement (providing a *statement handle*), and then **EXECUTE** the statement using the statement handle. When the application executes the statement, it must provide a *value* for each placeholder found in the statement.

When an application uses the **PREPARE/EXECUTE** mechanism, each SQL statement is parsed and planned once, but may execute many times (providing different *values* each time).

ECPGPlus will convert each parameter value to the type required by the SQL statement, if possible; if not possible, ECPGPlus will report an error.

```

/*************************************************/
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <sqlca.h>

static void handle_error(void);

int main(int argc, char *argv[])
{
    char *stmtText;

    EXEC SQL WHENEVER SQLERROR DO handle_error();

    EXEC SQL CONNECT :argv[1];

    stmtText = "INSERT INTO dept VALUES(?, ?, ?)";

    EXEC SQL PREPARE stmtHandle FROM :stmtText;

    EXEC SQL EXECUTE stmtHandle USING :argv[2], :argv[3], :argv[4];

    fprintf(stderr, "ok\n");

    EXEC SQL COMMIT RELEASE;

    exit(EXIT_SUCCESS);
}

static void handle_error(void)
{
    printf("%s\n", sqlca.sqlerrm.sqlerrmc);
    EXEC SQL WHENEVER SQLERROR CONTINUE;
    EXEC SQL ROLLBACK RELEASE;

    exit(EXIT_FAILURE);
}
/*************************************************/

```

The code sample begins by including the prototypes and type definitions for the C `stdio`, `string`, `stdlib`, and `sqlca` libraries, and providing basic infrastructure for the program:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <sqlca.h>

static void handle_error(void);

int main(int argc, char *argv[])
{
    char *stmtText;
```

The example then sets up an error handler; ECPGPlus calls the `handle_error()` function whenever a SQL error occurs:

```
EXEC SQL WHENEVER SQLERROR DO handle_error();
```

Then, the example connects to the database using the credentials specified on the command line:

```
EXEC SQL CONNECT :argc[1];
```

Next, the program uses a `PREPARE` statement to parse and plan a statement that includes three parameter markers - if the `PREPARE` statement succeeds, it will create a statement handle that you can use to execute the statement (in this example, the statement handle is named `stmtHandle`). You can execute a given statement multiple times using the same statement handle.

```
stmtText = "INSERT INTO dept VALUES(?, ?, ?);
```

```
EXEC SQL PREPARE stmtHandle FROM :stmtText;
```

After parsing and planning the statement, the application uses the `EXECUTE` statement to execute the statement associated with the statement handle, substituting user-provided values for the parameter markers:

```
EXEC SQL EXECUTE stmtHandle USING :argc[2], :argc[3], :argc[4];
```

If the `EXECUTE` command fails for any reason, ECPGPlus will invoke the `handle_error()` function (which terminates the application after displaying an error message to the user). If the `EXECUTE` command succeeds, the application displays a message `(ok)` to the user, commits the changes, disconnects from the server, and terminates the application.

```
fprintf(stderr, "ok\n");
```

```
EXEC SQL COMMIT RELEASE;
```

```
exit(EXIT_SUCCESS);
}
```

ECPGPlus calls the `handle_error()` function whenever it encounters a SQL error. The `handle_error()` function prints the content of the error message, resets the error handler, rolls back any changes, disconnects from the database, and terminates the application.

```
static void handle_error(void)
{
    printf("%s\n", sqlca.sqlerrm.sqlerrmc);
```

```

EXEC SQL WHENEVER SQLERROR CONTINUE;
EXEC SQL ROLLBACK RELEASE;
exit(EXIT_FAILURE);
}

```

Example - Executing a Query With a Known Number of Placeholders

This example demonstrates how to execute a *query* with a known number of input parameters, and with a known number of columns in the result set. This method uses the `PREPARE` statement to parse and plan a query, before opening a cursor and iterating through the result set.

```

/*********************************/
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <stdbool.h>
#include <sqlca.h>

static void handle_error(void);

int main(int argc, char *argv[])
{
    VARCHAR empno[10];
    VARCHAR ename[20];

    EXEC SQL WHENEVER SQLERROR DO handle_error();

    EXEC SQL CONNECT :argv[1];

    EXEC SQL PREPARE queryHandle
        FROM "SELECT empno, ename FROM emp WHERE deptno = ?";

    EXEC SQL DECLARE empCursor CURSOR FOR queryHandle;

    EXEC SQL OPEN empCursor USING :argv[2];

    EXEC SQL WHENEVER NOT FOUND DO break;

    while(true)
    {
        EXEC SQL FETCH empCursor INTO :empno, :ename;
        printf("%-10s %s\n", empno.arr, ename.arr);
    }

    EXEC SQL CLOSE empCursor;

    EXEC SQL COMMIT RELEASE;

    exit(EXIT_SUCCESS);
}

static void handle_error(void)

```

```
{
printf("%s\n", sqlca.sqlerrm.sqlerrmc);

EXEC SQL WHENEVER SQLERROR CONTINUE;
EXEC SQL ROLLBACK RELEASE;

exit(EXIT_FAILURE);
}

******/
```

The code sample begins by including the prototypes and type definitions for the C `stdio`, `string`, `stdlib`, `stdbool`, and `sqlca` libraries, and providing basic infrastructure for the program:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <stdbool.h>
#include <sqlca.h>

static void handle_error(void);

int main(int argc, char *argv[])
{
    VARCHAR empno[10];
    VARCHAR ename[20];
```

The example then sets up an error handler; ECPGPlus calls the `handle_error()` function whenever a SQL error occurs:

```
EXEC SQL WHENEVER SQLERROR DO handle_error();
```

Then, the example connects to the database using the credentials specified on the command line:

```
EXEC SQL CONNECT :argv[1];
```

Next, the program uses a `PREPARE` statement to parse and plan a query that includes a single parameter marker - if the `PREPARE` statement succeeds, it will create a statement handle that you can use to execute the statement (in this example, the statement handle is named `stmtHandle`). You can execute a given statement multiple times using the same statement handle.

```
EXEC SQL PREPARE stmtHandle
  FROM "SELECT empno, ename FROM emp WHERE deptno = ?";
```

The program then declares and opens the cursor, `empCursor`, substituting a user-provided value for the parameter marker in the prepared `SELECT` statement. Notice that the `OPEN` statement includes a `USING` clause: the `USING` clause must provide a *value* for each placeholder found in the query:

```
EXEC SQL DECLARE empCursor CURSOR FOR stmtHandle;
```

```
EXEC SQL OPEN empCursor USING :argv[2];
```

```
EXEC SQL WHENEVER NOT FOUND DO break;
```

```
while(true)
{
```

The program iterates through the cursor, and prints the employee number and name of each employee in the selected department:

```
EXEC SQL FETCH empCursor INTO :empno, :ename;
printf("%-10s %s\n", empno.arr, ename.arr);
}
```

The program then closes the cursor, commits any changes, disconnects from the server, and terminates the application.

```
EXEC SQL CLOSE empCursor;
EXEC SQL COMMIT RELEASE;
exit(EXIT_SUCCESS);
}
```

The application calls the `handle_error()` function whenever it encounters a SQL error. The `handle_error()` function prints the content of the error message, resets the error handler, rolls back any changes, disconnects from the database, and terminates the application.

```
static void handle_error(void)
{
    printf("%s\n", sqlca.sqlerrm.sqlerrmc);

    EXEC SQL WHENEVER SQLERROR CONTINUE;
    EXEC SQL ROLLBACK RELEASE;

    exit(EXIT_FAILURE);
}
```

Example - Executing a Query With an Unknown Number of Variables

The next example demonstrates executing a query with an unknown number of input parameters and/or columns in the result set. This type of query may occur when you prompt the user for the text of the query, or when a query is assembled from a form on which the user chooses from a number of conditions (i.e., a filter).

```
/****************************************/
#include <stdio.h>
#include <stdlib.h>
#include <sqlda.h>
#include <sqlcpr.h>

SQLDA *params;
SQLDA *results;

static void allocateDescriptors(int count,
                               int varNameLength,
                               int indNameLenth);
static void bindParams(void);
static void displayResultSet(void);

int main(int argc, char *argv[])
{
```

```

EXEC SQL BEGIN DECLARE SECTION;
char *username = argv[1];
char *password = argv[2];
char *stmtText = argv[3];
EXEC SQL END DECLARE SECTION;

EXEC SQL WHENEVER SQLERROR sqlprint;

EXEC SQL CONNECT TO test
USER :username
IDENTIFIED BY :password;

params = sqlald(20, 64, 64);
results = sqlald(20, 64, 64);

EXEC SQL PREPARE stmt FROM :stmtText;

EXEC SQL DECLARE dynCursor CURSOR FOR stmt;

bindParams();

EXEC SQL OPEN dynCursor USING DESCRIPTOR params;

displayResultSet(20);
}

static void bindParams(void)
{
EXEC SQL DESCRIBE BIND VARIABLES FOR stmt INTO params;

if (params->F < 0)
fprintf(stderr, "Too many parameters required\n");
else
{
int i;

params->N = params->F;

for (i = 0; i < params->F; i++)
{
char *paramName = params->S[i];
int nameLen = params->C[i];
char paramValue[255];

printf("Enter value for parameter %.*s: ",
nameLen, paramName);

fgets(paramValue, sizeof(paramValue), stdin);

params->T[i] = 1; /* Data type = Character (1) */
params->L[i] = strlen(paramValue) - 1;
params->V[i] = strdup(paramValue);
}
}
}
}

```

```

static void displayResultSet(void)
{
    EXEC SQL DESCRIBE SELECT LIST FOR stmt INTO results;

    if (results->F < 0)
        fprintf(stderr, "Too many columns returned by query\n");
    else if (results->F == 0)
        return;
    else
    {
        int col;

        results->N = results->F;

        for (col = 0; col < results->F; col++)
        {
            int null_permitted, length;

            sqlnul(&results->T[col],
                   &results->T[col],
                   &null_permitted);

            switch (results->T[col])
            {
                case 2: /* NUMERIC */
                {
                    int precision, scale;

                    sqlprc(&results->L[col], &precision, &scale);

                    if (precision == 0)
                        precision = 38;

                    length = precision + 3;
                    break;
                }

                case 12: /* DATE */
                {
                    length = 30;
                    break;
                }

                default: /* Others */
                {
                    length = results->L[col] + 1;
                    break;
                }
            }

            results->V[col] = realloc(results->V[col], length);
            results->L[col] = length;
            results->T[col] = 1;
        }
    }
}

```

```

EXEC SQL WHENEVER NOT FOUND DO break;

while (1)
{
    const char *delimiter = "";

    EXEC SQL FETCH dynCursor USING DESCRIPTOR results;

    for (col = 0; col < results->F; col++)
    {
        if (*results->l[col] == -1)
            printf("%s%s", delimiter, "<null>");
        else
            printf("%s%s", delimiter, results->V[col]);
        delimiter = ", ";
    }

    printf("\n");
}
}
}

/*****

```

The code sample begins by including the prototypes and type definitions for the C `stdio` and `stdlib` libraries. In addition, the program includes the `sqlida.h` and `sqlcpr.h` header files. `sqlida.h` defines the SQLDA structure used throughout this example. `sqlcpr.h` defines a small set of functions used to interrogate the metadata found in an SQLDA structure.

```
#include <stdio.h>
#include <stdlib.h>
#include <sqlida.h>
#include <sqlcpr.h>
```

Next, the program declares pointers to two SQLDA structures. The first SQLDA structure (`params`) will be used to describe the metadata for any parameter markers found in the dynamic query text. The second SQLDA structure (`results`) will contain both the metadata and the result set obtained by executing the dynamic query.

```
SQLDA *params;
SQLDA *results;
```

The program then declares two helper functions (defined near the end of the code sample):

```
static void bindParams(void);
static void displayResultSet(void);
```

Next, the program declares three host variables; the first two (`username` and `password`) are used to connect to the database server; the third host variable (`stmtTxt`) is a NULL-terminated C string containing the text of the query to execute. Notice that the values for these three host variables are derived from the command-line arguments. When the program begins execution, it sets up an error handler and then connects to the database server:

```
int main(int argc, char *argv[])
{
    EXEC SQL BEGIN DECLARE SECTION;
    char *username = argv[1];
```

```

char *password = argv[2];
char *stmtText = argv[3];
EXEC SQL END DECLARE SECTION;

EXEC SQL WHENEVER SQLERROR sqlprint;
EXEC SQL CONNECT TO test
USER :username
IDENTIFIED BY :password;

```

Next, the program calls the `sqlald()` function to allocate the memory required for each descriptor. Each descriptor contains (among other things):

- a pointer to an array of column names
- a pointer to an array of indicator names
- a pointer to an array of data types
- a pointer to an array of lengths
- a pointer to an array of data values.

When you allocate an `SQLDA` descriptor, you specify the maximum number of columns you expect to find in the result set (for `SELECT`-list descriptors) or the maximum number of parameters you expect to find the dynamic query text (for bind-variable descriptors) - in this case, we specify that we expect no more than 20 columns and 20 parameters. You must also specify a maximum length for each column (or parameter) name and each indicator variable name - in this case, we expect names to be no more than 64 bytes long.

See [SQLDA Structure](#) section for a complete description of the `SQLDA` structure.

```

params = sqlald(20, 64, 64);
results = sqlald(20, 64, 64);

```

After allocating the `SELECT`-list and bind descriptors, the program prepares the dynamic statement and declares a cursor over the result set.

```
EXEC SQL PREPARE stmt FROM :stmtText;
```

```
EXEC SQL DECLARE dynCursor CURSOR FOR stmt;
```

Next, the program calls the `bindParams()` function. The `bindParams()` function examines the bind descriptor (`params`) and prompt the user for a value to substitute in place of each parameter marker found in the dynamic query.

```
bindParams();
```

Finally, the program opens the cursor (using the parameter values supplied by the user, if any) and calls the `displayResultSet()` function to print the result set produced by the query.

```
EXEC SQL OPEN dynCursor USING DESCRIPTOR params;
```

```
displayResultSet();
}
```

The `bindParams()` function determines whether the dynamic query contains any parameter markers, and, if so, prompts the user for a value for each parameter and then binds that value to the corresponding marker. The `DESCRIBE BIND VARIABLE` statement populates the `params` `SQLDA` structure with information describing each parameter marker.

```

static void bindParams(void)
{
    EXEC SQL DESCRIBE BIND VARIABLES FOR stmt INTO params;
}

```

If the statement contains no parameter markers, `params->F` will contain 0. If the statement contains more parameters than will fit into the descriptor, `params->F` will contain a negative number (in this case, the absolute value of `params->F` indicates the number of parameter markers found in the statement). If `params->F` contains a positive number, that number indicates how many parameter markers were found in the statement.

```
if (params->F < 0)
    fprintf(stderr, "Too many parameters required\n");
else
{
    int i;

    params->N = params->F;
```

Next, the program executes a loop that prompts the user for a value, iterating once for each parameter marker found in the statement.

```
for (i = 0; i < params->F; i++)
{
    char *paramName = params->S[i];
    int nameLen = params->C[i];
    char paramValue[255];

    printf("Enter value for parameter %.*s: ",
           nameLen, paramName);

    fgets(paramValue, sizeof(paramValue), stdin);
```

After prompting the user for a value for a given parameter, the program *binds* that value to the parameter by setting `params->T[i]` to indicate the data type of the value (see [Type Codes](#) for a list of type codes), `params->L[i]` to the length of the value (we subtract one to trim off the trailing new-line character added by `fgets()`), and `params->V[i]` to point to a copy of the NULL-terminated string provided by the user.

```
params->T[i] = 1;      /* Data type = Character (1) */
params->L[i] = strlen(paramValue) + 1;
params->V[i] = strdup(paramValue);
}
}
}
```

The `displayResultSet()` function loops through each row in the result set and prints the value found in each column. `displayResultSet()` starts by executing a `DESCRIBE SELECT LIST` statement - this statement populates an SQLDA descriptor (`results`) with a description of each column in the result set.

```
static void displayResultSet(void)
{
    EXEC SQL DESCRIBE SELECT LIST FOR stmt INTO results;
```

If the dynamic statement returns no columns (that is, the dynamic statement is not a `SELECT` statement), `results->F` will contain 0. If the statement returns more columns than will fit into the descriptor, `results->F` will contain a negative number (in this case, the absolute value of `results->F` indicates the number of columns returned by the statement). If `results->F` contains a positive number, that number indicates how many columns where returned by the query.

```
if (results->F < 0)
    fprintf(stderr, "Too many columns returned by query\n");
else if (results->F == 0)
    return;
```

```

else
{
    int col;

    results->N = results->F;
}

```

Next, the program enters a loop, iterating once for each column in the result set:

```

for (col = 0; col < results->F; col++)
{
    int null_permitted, length;
}

```

To decode the type code found in `results->T`, the program invokes the `sqlnul()` function (see the description of the `T` member of the SQLDA structure in the [The SQLDA Structure](#)). This call to `sqlnul()` modifies `results->T[col]` to contain only the type code (the nullability flag is copied to `null_permitted`). This step is necessary because the [DESCRIBE SELECT LIST](#) statement encodes the type of each column and the nullability of each column into the `T` array.

```

sqlnul(&results->T[col]
       &results->T[col]
       &null_permitted);
}

```

After decoding the actual data type of the column, the program modifies the results descriptor to tell ECPGPlus to return each value in the form of a NULL-terminated string. Before modifying the descriptor, the program must compute the amount of space required to hold each value. To make this computation, the program examines the maximum length of each column (`results->V[col]`) and the data type of each column (`results->T[col]`).

For numeric values (where `results->T[col] = 2`), the program calls the `sqlprc()` function to extract the precision and scale from the column length. To compute the number of bytes required to hold a numeric value in string form, `displayResultSet()` starts with the precision (that is, the maximum number of digits) and adds three bytes for a sign character, a decimal point, and a NULL terminator.

```

switch (results->T[col])
{
    case 2: /* NUMERIC */
    {
        int precision, scale;

        sqlprc(&results->L[col], &precision, &scale);

        if (precision == 0)
            precision = 38;
        length = precision + 3;
        break;
    }
}

```

For date values, the program uses a somewhat arbitrary, hard-coded length of 30. In a real-world application, you may want to more carefully compute the amount of space required.

```

case 12: /* DATE */
{
    length = 30;
    break;
}

```

For a value of any type other than date or numeric, `displayResultSet()` starts with the maximum column width reported by [DESCRIBE SELECT LIST](#) and adds one extra byte for the NULL terminator. Again, in a real-world application you may want to include more careful calculations for other data types.

```

default: /* Others */
{
    length = results->L[col] + 1;
    break;
}
}

```

After computing the amount of space required to hold a given column, the program allocates enough memory to hold the value, sets `results->L[col]` to indicate the number of bytes found at `results->V[col]`, and set the type code for the column (`results->T[col]`) to `1` to instruct the upcoming `FETCH` statement to return the value in the form of a NULL-terminated string.

```

results->V[col] = malloc(length);
results->L[col] = length;
results->T[col] = 1;
}

```

At this point, the results descriptor is configured such that a `FETCH` statement can copy each value into an appropriately sized buffer in the form of a NULL-terminated string.

Next, the program defines a new error handler to break out of the upcoming loop when the cursor is exhausted.

```

EXEC SQL WHENEVER NOT FOUND DO break;

while (1)
{
    const char *delimiter = "";

```

The program executes a `FETCH` statement to fetch the next row in the cursor into the `results` descriptor. If the `FETCH` statement fails (because the cursor is exhausted), control transfers to the end of the loop because of the `EXEC SQL WHENEVER` directive found before the top of the loop.

| `EXEC SQL FETCH dynCursor USING DESCRIPTOR results;`

The `FETCH` statement will populate the following members of the results descriptor:

- `*results->I[col]` will indicate whether the column contains a NULL value (`-1`) or a non-NULL value (`0`). If the value non-NULL but too large to fit into the space provided, the value is truncated and `*results->I[col]` will contain a positive value.
- `results->V[col]` will contain the value fetched for the given column (unless `*results->I[col]` indicates that the column value is NULL).
- `results->L[col]` will contain the length of the value fetched for the given column

Finally, `displayResultSet()` iterates through each column in the result set, examines the corresponding NULL indicator, and prints the value. The result set is not aligned - instead, each value is separated from the previous value by a comma.

```

for (col = 0; col < results->F; col++)
{
    if (*results->I[col] == -1)
        printf("%s%s", delimiter, "<null>");
    else
        printf("%s%s", delimiter, results->V[col]);
    delimiter = ", ";
}

printf("\n");
}

```

```

}
}
/*****

```

9.6 Error Handling

ECPGPlus provides two methods to detect and handle errors in embedded SQL code:

- A client application can examine the `sqlca` data structure for error messages, and supply customized error handling for your client application.
- A client application can include `EXEC SQL WHENEVER` directives to instruct the ECPGPlus compiler to add error-handling code.

Error Handling with `sqlca`

`sqlca` (SQL communications area) is a global variable used by `ecpglib` to communicate information from the server to the client application. After executing a SQL statement (for example, an `INSERT` or `SELECT` statement) you can inspect the contents of `sqlca` to determine if the statement has completed successfully or if the statement has failed.

`sqlca` has the following structure:

```

struct
{
    char sqlcaid[8];
    long sqlabc;
    long sqlcode;
    struct
    {
        int sqlerrml;
        char sqlerrmc[SQLERRMC_LEN];
    } sqlerrm;
    char sqlerrp[8];
    long sqlerrd[6];
    char sqlwarn[8];
    char sqlstate[5];
} sqlca;

```

Use the following directive to implement `sqlca` functionality:

```
EXEC SQL INCLUDE sqlca;
```

If you include the `ecpg` directive, you do not need to `#include` the `sqlca.h` file in the client application's header declaration.

The Advanced Server `sqlca` structure contains the following members:

`sqlcaid`

`sqlcaid` contains the string: "SQLCA".

sqlabc

sqlabc contains the size of the **sqlca** structure.

sqlcode

The **sqlcode** member has been deprecated with SQL 92; Advanced Server supports **sqlcode** for backward compatibility, but you should use the **sqlstate** member when writing new code.

sqlcode is an integer value; a positive **sqlcode** value indicates that the client application has encountered a harmless processing condition, while a negative value indicates a warning or error.

If a statement processes without error, **sqlcode** will contain a value of **0**. If the client application encounters an error (or warning) during a statement's execution, **sqlcode** will contain the last code returned.

The SQL standard defines only a positive value of 100, which indicates that the most recent SQL statement processed returned/affected no rows. Since the SQL standard does not define other **sqlcode** values, please be aware that the values assigned to each condition may vary from database to database.

sqlerrm is a structure embedded within **sqlca**, composed of two members:

sqlerrml

sqlerrml contains the length of the error message currently stored in **sqlerrmc**.

sqlerrmc

sqlerrmc contains the null-terminated message text associated with the code stored in **sqlstate**. If a message exceeds 149 characters in length, **ecpglib** will truncate the error message.

sqlerrp

sqlerrp contains the string "NOT SET".

sqlerrd is an array that contains six elements:

sqlerrd[1] contains the OID of the processed row (if applicable).

sqlerrd[2] contains the number of processed or returned rows.

sqlerrd[0], **sqlerrd[3]**, **sqlerrd[4]** and **sqlerrd[5]** are unused.

sqlwarn is an array that contains 8 characters:

sqlwarn[0] contains a value of '**W**' if any other element within **sqlwarn** is set to '**W**'.

sqlwarn[1] contains a value of '**W**' if a data value was truncated when it was stored in a host variable.

sqlwarn[2] contains a value of '**W**' if the client application encounters a non-fatal warning.

sqlwarn[3], **sqlwarn[4]**, **sqlwarn[5]**, **sqlwarn[6]**, and **sqlwarn[7]** are unused.

sqlstate

sqlstate is a 5 character array that contains a SQL-compliant status code after the execution of a statement from the client application. If a statement processes without error, **sqlstate** will contain a value of **00000**. Please note that **sqlstate** is *not* a null-terminated string.

sqlstate codes are assigned in a hierarchical scheme:

- The first two characters of **sqlstate** indicate the general class of the condition.
- The last three characters of **sqlstate** indicate a specific status within the class.

If the client application encounters multiple errors (or warnings) during an SQL statement's execution `sqlstate` will contain the last code returned.

The following table lists the `sqlstate` and `sqlcode` values, as well as the symbolic name and error description for the related condition:

| sqlstate | sqlcode (Deprecated) | Symbolic Name | Description |
|--------------------|-------------------------|-------------------------|--|
| YE001 | -12 | ECPG_OUT_OF_MEMORY | Virtual memory is exhausted. |
| YE002 | -200 | ECPG_UNSUPPORTED | The preprocessor has generated an unsupported command. |
| 07001, or 07002 | -201 | ECPG_TOO_MANY_ARGUMENTS | The program specifies more variables than were declared. |
| 07001, or 07002 | -202 | ECPG_TOO_FEW_ARGUMENTS | The program specified fewer variables than were declared. |
| 21000 | -203 | ECPG_TOO_MANY_MATCHES | The SQL command has returned too many matches. |
| 42804 | -204 | ECPG_INT_FORMAT | The host variable (defined in the application) is not in integer format. |
| 42804 | -205 | ECPG_UINT_FORMAT | The host variable (defined in the application) is not in unsigned integer format. |
| 42804 | -206 | ECPG_FLOAT_FORMAT | The host variable (defined in the application) is not in floating point format. |
| 42804 | -211 | ECPG_CONVERT_BOOL | The host variable (defined in the application) is not in boolean format. |
| YE002 | -2-1 | ECPG_EMPTY | The statement sent to the server was empty. |
| 22002 | -213 | ECPG_MISSING_INDICATOR | A NULL indicator variable has been passed to the server. |
| 42804 | -214 | ECPG_NO_ARRAY | The server has returned an array indicator. |
| 42804 | -215 | ECPG_DATA_NOT_ARRAY | The server has returned a value that is not an array. |
| 08003 | -220 | ECPG_NO_CONN | The client application has attempted to connect to a non-existent database. |
| YE002 | -221 | ECPG_NOT_CONN | The client application has attempted to connect to a database that is not connected. |

| | | | |
|-------|------|-------------------------------|---------------------------------|
| 26000 | -230 | ECPG_INVALID_STMT | The statement has not been p |
| 33000 | -240 | ECPG_UNKNOWN_DESCRIPTOR | The specified descriptor is not |
| 07009 | -241 | ECPG_INVALID_DESCRIPTOR_INDEX | The descriptor index is out-of- |
| YE002 | -242 | ECPG_UNKNOWN_DESCRIPTOR_ITEM | The client application has req |
| 07006 | -243 | ECPG_VAR_NOT_NUMERIC | A dynamic statement has retu |
| 07006 | -244 | ECPG_VAR_NOT_CHAR | A dynamic SQL statement has |
| | -400 | ECPG_PGSQL | The server has returned an er |
| 08007 | -401 | ECPG_TRANS | The server cannot start, comm |
| 08001 | -402 | ECPG_CONNECT | The client application's attemp |
| 02000 | 100 | ECPG_NOT_FOUND | The last command retrieved o |

EXEC SQL WHENEVER

Use the **EXEC SQL WHENEVER** directive to implement simple error handling for client applications compiled with ECPGPlus. The syntax of the directive is:

```
EXEC SQL WHENEVER <condition> <action>;
```

This directive instructs the ECPG compiler to insert error-handling code into your program.

The code instructs the client application that it should perform a specified action if the client application detects a given condition. The *condition* may be one of the following:

SQLERROR

A **SQLERROR** condition exists when `sqlca.sqlcode` is less than zero.

SQLWARNING

A **SQLWARNING** condition exists when `sqlca.sqlwarn[0]` contains a 'W'.

NOT FOUND

A **NOT FOUND** condition exists when `sqlca.sqlcode` is **ECPG_NOT_FOUND** (when a query returns no data).

You can specify that the client application perform one of the following *actions* if it encounters one of the previous conditions:

CONTINUE

Specify `CONTINUE` to instruct the client application to continue processing, ignoring the current `condition`. `CONTINUE` is the default action.

DO CONTINUE

An action of `DO CONTINUE` will generate a `CONTINUE` statement in the emitted C code that if it encounters the condition, skips the rest of the code in the loop and continues with the next iteration. You can only use it within a loop.

GOTO label

or

GO TO label

Use a C `goto` statement to jump to the specified `label`.

SQLPRINT

Print an error message to `stderr` (standard error), using the `sqlprint()` function. The `sqlprint()` function prints `sql error`, followed by the contents of `sqlca.sqlerrm.sqlerrmc`.

STOP

Call `exit(1)` to signal an error, and terminate the program.

DO BREAK

Execute the C `break` statement. Use this action in loops, or `switch` statements.

CALL name(args)

or

DO name(args)

Invoke the C function specified by the name `parameter`, using the parameters specified in the `args` parameter.

Example:

The following code fragment prints a message if the client application encounters a warning, and aborts the application if it encounters an error:

```
EXEC SQL WHENEVER SQLWARNING SQLPRINT;
EXEC SQL WHENEVER SQLERROR STOP;
```

Note

The ECPGPlus compiler processes your program from top to bottom, even though the client application may not execute from top to bottom. The compiler directive is applied to each line in order, and remains in effect until the compiler encounters another directive. If the control of the flow within your program is not top-to-bottom, you should consider adding error-handling directives to any parts of the program that may be inadvertently missed during compilation.

9.7 Reference

The sections that follow describe ecpgPlus language elements:

- C-Preprocessor Directives
- Supported C Data Types
- Type Codes
- The SQLDA Structure
- ECPGPlus Statements

C-preprocessor Directives

The ECPGPlus C-preprocessor enforces two behaviors that are dependent on the mode in which you invoke ECPGPlus:

- **PROC** mode
- non-**PROC** mode

Compiling in PROC Mode

In **PROC** mode, ECPGPlus allows you to:

- Declare host variables outside of an **EXEC SQL BEGIN/END DECLARE SECTION**.
- Use any C variable as a host variable as long as it is of a data type compatible with ECPG.

When you invoke ECPGPlus in **PROC** mode (by including the **-C PROC** keywords), the ECPG compiler honors the following C-preprocessor directives:

```
#include
#if expression
#define symbolName
#ifndef symbolName
#else
#elif expression
#endif
#define symbolName expansion
#define symbolName([macro arguments]) expansion
#undef symbolName
#define(symbolName)
```

Pre-processor directives are used to effect or direct the code that is received by the compiler. For example, using the following code sample:

```
#if HAVE_LONG_LONG == 1
#define BALANCE_TYPE long long
#else
#define BALANCE_TYPE double
#endif
...
BALANCE_TYPE customerBalance;
```

If you invoke ECPGPlus with the following command-line arguments:

```
ecpg -C PROC -D HAVE_LONG_LONG=1
```

ECPGPlus will copy the entire fragment (without change) to the output file, but will only send the following tokens to the ECPG parser:

```
long long customerBalance;
```

On the other hand, if you invoke ECPGPlus with the following command-line arguments:

```
ecpg -C PROC -D HAVE_LONG_LONG=0
```

The ECPG parser will receive the following tokens:

```
double customerBalance;
```

If your code uses preprocessor directives to filter the code that is sent to the compiler, the complete code is retained in the original code, while the ECPG parser sees only the processed token stream.

You can also use compatible syntax when executing the following preprocessor directives with an **EXEC** directive:

```
EXEC ORACLE DEFINE
EXEC ORACLE UNDEF
EXEC ORACLE INCLUDE
EXEC ORACLE IFDEF
EXEC ORACLE IFNDEF
EXEC ORACLE ELIF
EXEC ORACLE ELSE
EXEC ORACLE ENDIF
EXEC ORACLE OPTION
```

For example, if your code includes the following:

```
EXEC ORACLE IFDEF HAVE_LONG_LONG;
#define BALANCE_TYPE long long
EXEC ORACLE ENDIF;
BALANCE_TYPE customerBalance;
```

If you invoke ECPGPlus with the following command-line arguments:

```
ecpg -C PROC DEFINE=HAVE_LONG_LONG=1
```

ECPGPlus will send the following tokens to the output file, and the ECPG parser:

```
long long customerBalance;
```

Note

The **EXEC ORACLE** pre-processor directives only work if you specify **-C PROC** on the ECPG command line.

Using the SELECT_ERROR Precompiler Option

When using ECPGPlus in compatible mode, you can use the **SELECT_ERROR** precompiler option to instruct your program how to handle result sets that contain more rows than the host variable can accommodate. The syntax is:

```
SELECT_ERROR={YES|NO}
```

The default value is **YES**; a **SELECT** statement will return an error message if the result set exceeds the capacity of the host variable. Specify **NO** to instruct the program to suppress error messages when a **SELECT** statement returns more rows than a host variable can accommodate.

Use **SELECT_ERROR** with the **EXEC ORACLE OPTION** directive.

Compiling in non-PROC Mode

If you do not include the `-C PROC` command-line option:

- C preprocessor directives are copied to the output file without change.
- You must declare the type and name of each C variable that you intend to use as a host variable within an `EXEC SQL BEGIN/END DECLARE` section.

When invoked in non-`PROC` mode, ECPG implements the behavior described in the PostgreSQL Core documentation.

Supported C Data Types

An ECPGPlus application must deal with two sets of data types: SQL data types (such as `SMALLINT`, `DOUBLE PRECISION` and `CHARACTER VARYING`) and C data types (like `short`, `double` and `varchar[n]`). When an application fetches data from the server, ECPGPlus will map each SQL data type to the type of the C variable into which the data is returned.

In general, ECPGPlus can convert most SQL server types into similar C types, but not all combinations are valid. For example, ECPGPlus will try to convert a SQL character value into a C integer value, but the conversion may fail (at execution time) if the SQL character value contains non-numeric characters. The reverse is also true; when an application sends a value to the server, ECPGPlus will try to convert the C data type into the required SQL type. Again, the conversion may fail (at execution time) if the C value cannot be converted into the required SQL type.

ECPGPlus can convert any SQL type into C character values (`char[n]` or `varchar[n]`). Although it is safe to convert any SQL type to/from `char[n]` or `varchar[n]`, it is often convenient to use more natural C types such as `int`, `double`, or `float`.

The supported C data types are:

- `short`
- `int`
- `unsigned int`
- `long long int`
- `float`
- `double`
- `char[n+1]`
- `varchar[n+1]`
- `bool`
- and any equivalent created by a `typedef`

In addition to the numeric and character types supported by C, the `pgtypeslib` run-time library offers custom data types (and functions to operate on those types) for dealing with date/time and exact numeric values:

- `timestamp`
- `interval`
- `date`
- `decimal`
- `numeric`

To use a data type supplied by `pgtypeslib`, you must `#include` the proper header file.

Type Codes

The following table contains the type codes for *external* data types. An external data type is used to indicate the type of a C host variable. When an application binds a value to a parameter or binds a buffer to a `SELECT`-list item, the type code in the corresponding SQLDA descriptor (`descriptor->T[column]`) should be set to one of the following values:

| Type Code | Host Variable Type (C Data Type) |
|---|---|
| 1, 2, 8, 11, 12, 15, 23, 24, 91, 94, 95, 96, 97 | char[] |
| 3 | int |
| 4, 7, 21 | float |
| 5, 6 | null-terminated string (char[length+1]) |
| 9 | varchar |
| 22 | double |
| 68 | unsigned int |

The following table contains the type codes for *internal* data types. An internal type code is used to indicate the type of a value as it resides in the database. The **DESCRIBE SELECT LIST** statement populates the data type array (`descriptor->T[column]`) using the following values.

| Internal Type Code | Server Type |
|--------------------|------------------------|
| 1 | VARCHAR2 |
| 2 | NUMBER |
| 8 | LONG |
| 11 | ROWID |
| 12 | DATE |
| 23 | RAW |
| 24 | LONG RAW |
| 96 | CHAR |
| 100 | BINARY FLOAT |
| 101 | BINARY DOUBLE |
| 104 | UROWID |
| 187 | TIMESTAMP |
| 188 | TIMESTAMP W/TIMEZONE |
| 189 | INTERVAL YEAR TO MONTH |
| 190 | INTERVAL DAY TO SECOND |
| 232 | TIMESTAMP LOCAL_TZ |

The SQLDA Structure

Oracle Dynamic SQL method 4 uses the SQLDA data structure to hold the data and metadata for a dynamic SQL statement. A SQLDA structure can describe a set of input parameters corresponding to the parameter markers found in the text of a dynamic statement or the result set of a dynamic statement. The layout of the SQLDA structure is:

```
struct SQLDA
{
    int    N; /* Number of entries      */
    char **V; /* Variables            */
    int   *L; /* Variable lengths     */
    short *T; /* Variable types       */
    short **I; /* Indicators          */
    int    F; /* Count of variables discovered by DESCRIBE */
    char **S; /* Variable names       */
    short *M; /* Variable name maximum lengths */
    short *C; /* Variable name actual lengths */
    char **X; /* Indicator names     */
}
```

```
short *Y; /* Indicator name maximum lengths */
short *Z; /* Indicator name actual lengths */
};
```

Parameters

N - maximum number of entries

The N structure member contains the maximum number of entries that the SQLDA may describe. This member is populated by the `sqlald()` function when you allocate the SQLDA structure. Before using a descriptor in an `OPEN` or `FETCH` statement, you must set N to the *actual* number of values described.

V - data values

The V structure member is a pointer to an array of data values.

- For a `SELECT`-list descriptor, V points to an array of values returned by a `FETCH` statement (each member in the array corresponds to a column in the result set).
- For a bind descriptor, V points to an array of parameter values (you must populate the values in this array before opening a cursor that uses the descriptor).

Your application must allocate the space required to hold each value. Refer to `displayResultSet` function for an example of how to allocate space for `SELECT`-list values.

L - length of each data value

The L structure member is a pointer to an array of lengths. Each member of this array must indicate the amount of memory available in the corresponding member of the V array. For example, if V[5] points to a buffer large enough to hold a 20-byte NULL-terminated string, L[5] should contain the value 21 (20 bytes for the characters in the string plus 1 byte for the NULL-terminator). Your application must set each member of the L array.

T - data types

The T structure member points to an array of data types, one for each column (or parameter) described by the descriptor.

- For a bind descriptor, you must set each member of the T array to tell ECPGPlus the data type of each parameter.
- For a `SELECT`-list descriptor, the `DESCRIBE SELECT LIST` statement sets each member of the T array to reflect the type of data found in the corresponding column.

You may change any member of the T array before executing a `FETCH` statement to force ECPGPlus to convert the corresponding value to a specific data type. For example, if the `DESCRIBE SELECT LIST` statement indicates that a given column is of type `DATE`, you may change the corresponding T member to request that the next `FETCH` statement return that value in the form of a NULL-terminated string. Each member of the T array is a numeric type code (see `Type Codes` for a list of type codes). The type codes returned by a `DESCRIBE SELECT LIST` statement differ from those expected by a `FETCH` statement. After executing a `DESCRIBE SELECT LIST` statement, each member of T encodes a data type and a flag indicating whether the corresponding column is nullable. You can use the `sqlnul()` function to extract the type code and nullable flag from a member of the T array. The signature of the `sqlnul()` function is as follows:

```
void sqlnul(unsigned short *valType,
            unsigned short *typeCode,
            int           *isNull)
```

For example, to find the type code and nullable flag for the third column of a descriptor named results, you would invoke `sqlnul()` as follows:

```
sqlnul(&results->T[2], &typeCode, &isNull);
```

I - indicator variables

The **I** structure member points to an array of indicator variables. This array is allocated for you when your application calls the `sqlald()` function to allocate the descriptor.

- For a `SELECT`-list descriptor, each member of the **I** array indicates whether the corresponding column contains a NULL (non-zero) or non-NUL (zero) value.
- For a bind parameter, your application must set each member of the **I** array to indicate whether the corresponding parameter value is NULL.

F - number of entries

The **F** structure member indicates how many values are described by the descriptor (the **N** structure member indicates the *maximum* number of values which may be described by the descriptor; **F** indicates the actual number of values). The value of the **F** member is set by ECPGPlus when you execute a `DESCRIBE` statement. **F** may be positive, negative, or zero.

- For a `SELECT`-list descriptor, **F** will contain a positive value if the number of columns in the result set is equal to or less than the maximum number of values permitted by the descriptor (as determined by the **N** structure member); 0 if the statement is *not* a `SELECT` statement, or a negative value if the query returns more columns than allowed by the **N** structure member.
- For a bind descriptor, **F** will contain a positive number if the number of parameters found in the statement is less than or equal to the maximum number of values permitted by the descriptor (as determined by the **N** structure member); 0 if the statement contains no parameters markers, or a negative value if the statement contains more parameter markers than allowed by the **N** structure member.

If **F** contains a positive number (after executing a `DESCRIBE` statement), that number reflects the count of columns in the result set (for a `SELECT`-list descriptor) or the number of parameter markers found in the statement (for a bind descriptor). If **F** contains a negative value, you may compute the absolute value of **F** to discover how many values (or parameter markers) are required. For example, if **F** contains **-24** after describing a `SELECT` list, you know that the query returns 24 columns.

S - column/parameter names

The **S** structure member points to an array of NULL-terminated strings.

- For a `SELECT`-list descriptor, the `DESCRIBE SELECT LIST` statement sets each member of this array to the name of the corresponding column in the result set.
- For a bind descriptor, the `DESCRIBE BIND VARIABLES` statement sets each member of this array to the name of the corresponding bind variable.

In this release, the name of each bind variable is determined by the left-to-right order of the parameter marker within the query - for example, the name of the first parameter is always **?0**, the name of the second parameter is always **?1**, and so on.

M - maximum column/parameter name length

The **M** structure member points to an array of lengths. Each member in this array specifies the *maximum* length of the corresponding member of the **S** array (that is, **M[0]** specifies the maximum length of the column/parameter name found at **S[0]**). This array is populated by the `sqlald()` function.

C - actual column/parameter name length

The **C** structure member points to an array of lengths. Each member in this array specifies the *actual* length of the corresponding member of the **S** array (that is, **C[0]** specifies the actual length of the column/parameter name found at **S[0]**).

This array is populated by the `DESCRIBE` statement.

X - indicator variable names

The `X` structure member points to an array of NULL-terminated strings -each string represents the name of a NULL indicator for the corresponding value.

This array is not used by ECPGPlus, but is provided for compatibility with Pro*C applications.

`Y - maximum indicator name length`

The `Y` structure member points to an array of lengths. Each member in this array specifies the *maximum* length of the corresponding member of the `X` array (that is, `Y[0]` specifies the maximum length of the indicator name found at `X[0]`).

This array is not used by ECPGPlus, but is provided for compatibility with Pro*C applications.

`Z - actual indicator name length`

The `Z` structure member points to an array of lengths. Each member in this array specifies the *actual* length of the corresponding member of the `X` array (that is, `Z[0]` specifies the actual length of the indicator name found at `X[0]`).

This array is not used by ECPGPlus, but is provided for compatibility with Pro*C applications.

ECPGPlus Statements

An embedded SQL statement allows your client application to interact with the server, while an embedded directive is an instruction to the ECPGPlus compiler.

You can embed any Advanced Server SQL statement in a C program. Each statement should begin with the keywords `EXEC SQL`, and must be terminated with a semi-colon (;). Within the C program, a SQL statement takes the form:

```
EXEC SQL <sql_command_body>;
```

Where `sql_command_body` represents a standard SQL statement. You can use a host variable anywhere that the SQL statement expects a value expression. For more information about substituting host variables for value expressions, refer to [Declaring Host Variables](#).

ECPGPlus extends the PostgreSQL server-side syntax for some statements; for those statements, syntax differences are outlined in the following reference sections. For a complete reference to the supported syntax of other SQL commands, refer to the *PostgreSQL Core Documentation* available at:

<https://www.postgresql.org/docs/current/static/sql-commands.html>

ALLOCATE DESCRIPTOR

Use the `ALLOCATE DESCRIPTOR` statement to allocate an SQL descriptor area:

```
EXEC SQL [FOR <array_size>] ALLOCATE DESCRIPTOR <descriptor_name>
[WITH MAX <variable_count>];
```

Where:

`array_size` is a variable that specifies the number of array elements to allocate for the descriptor. `array_size` may be an `INTEGER` value or a host variable.

`descriptor_name` is the host variable that contains the name of the descriptor, or the name of the descriptor. This value may take the form of an identifier, a quoted string literal, or of a host variable.

`variable_count` specifies the maximum number of host variables in the descriptor. The default value of `variable_count` is 100.

The following code fragment allocates a descriptor named `emp_query` that may be processed as an array (`emp_array`):

```
EXEC SQL FOR :emp_array ALLOCATE DESCRIPTOR emp_query;
```

CALL

Use the `CALL` statement to invoke a procedure or function on the server. The `CALL` statement works only on Advanced Server. The `CALL` statement comes in two forms; the first form is used to call a *function*:

```
EXEC SQL CALL <program_name> '(['<actual_arguments>'])'  
INTO [:<ret_variable>][:<ret_indicator>];
```

The second form is used to call a *procedure*:

```
EXEC SQL CALL <program_name> '(['<actual_arguments>'])';
```

Where:

`program_name` is the name of the stored procedure or function that the `CALL` statement invokes. The program name may be schema-qualified or package-qualified (or both); if you do not specify the schema or package in which the program resides, ECPGPlus will use the value of `search_path` to locate the program.

`actual_arguments` specifies a comma-separated list of arguments required by the program. Note that each `actual_argument` corresponds to a formal argument expected by the program. Each formal argument may be an `IN` parameter, an `OUT` parameter, or an `INOUT` parameter.

`:ret_variable` specifies a host variable that will receive the value returned if the program is a function.

`:ret_indicator` specifies a host variable that will receive the indicator value returned, if the program is a function.

For example, the following statement invokes the `get_job_desc` function with the value contained in the `:ename` host variable, and captures the value returned by that function in the `:job` host variable:

```
EXEC SQL CALL get_job_desc(:ename)  
INTO :job;
```

CLOSE

Use the `CLOSE` statement to close a cursor, and free any resources currently in use by the cursor. A client application cannot fetch rows from a closed cursor. The syntax of the `CLOSE` statement is:

```
EXEC SQL CLOSE [<cursor_name>];
```

Where:

`cursor_name` is the name of the cursor closed by the statement. The cursor name may take the form of an identifier or of a host variable.

The `OPEN` statement initializes a cursor. Once initialized, a cursor result set will remain unchanged unless the cursor is re-opened. You do not need to `CLOSE` a cursor before re-opening it.

To manually close a cursor named `emp_cursor`, use the command:

```
EXEC SQL CLOSE emp_cursor;
```

A cursor is automatically closed when an application terminates.

COMMIT

Use the **COMMIT** statement to complete the current transaction, making all changes permanent and visible to other users. The syntax is:

```
EXEC SQL [AT <database_name>] COMMIT [WORK]
[COMMENT <'text'>] [COMMENT <'text'> RELEASE];
```

Where:

database_name is the name of the database (or host variable that contains the name of the database) in which the work resides. This value may take the form of an unquoted string literal, or of a host variable.

For compatibility, ECPGPlus accepts the **COMMENT** clause without error but does *not* store any text included with the **COMMENT** clause.

Include the **RELEASE** clause to close the current connection after performing the commit.

For example, the following command commits all work performed on the **dept** database and closes the current connection:

```
EXEC SQL AT dept COMMIT RELEASE;
```

By default, statements are committed only when a client application performs a **COMMIT** statement. Include the **-t** option when invoking ECPGPlus to specify that a client application should invoke **AUTOCOMMIT** functionality. You can also control **AUTOCOMMIT** functionality in a client application with the following statements:

```
EXEC SQL SET AUTOCOMMIT TO ON
```

and

```
EXEC SQL SET AUTOCOMMIT TO OFF
```

CONNECT

Use the **CONNECT** statement to establish a connection to a database. The **CONNECT** statement is available in two forms - one form is compatible with Oracle databases, the other is not.

The first form is compatible with Oracle databases:

```
EXEC SQL CONNECT
{{:<user_name> IDENTIFIED BY :<password>} | :<connection_id>}
[AT <database_name>]
[USING :database_string]
[ALTER AUTHORIZATION :new_password];
```

Where:

user_name is a host variable that contains the role that the client application will use to connect to the server.

password is a host variable that contains the password associated with that role.

`connection_id` is a host variable that contains a slash-delimited user name and password used to connect to the database.

Include the `AT` clause to specify the database to which the connection is established. `database_name` is the name of the database to which the client is connecting; specify the value in the form of a variable, or as a string literal.

Include the `USING` clause to specify a host variable that contains a null-terminated string identifying the database to which the connection will be established.

The `ALTER AUTHORIZATION` clause is supported for syntax compatibility only; ECPGPlus parses the `ALTER AUTHORIZATION` clause, and reports a warning.

Using the first form of the `CONNECT` statement, a client application might establish a connection with a host variable named `user` that contains the identity of the connecting role, and a host variable named `password` that contains the associated password using the following command:

```
EXEC SQL CONNECT :user IDENTIFIED BY :password;
```

A client application could also use the first form of the `CONNECT` statement to establish a connection using a single host variable named `:connection_id`. In the following example, `connection_id` contains the slash-delimited role name and associated password for the user:

```
EXEC SQL CONNECT :connection_id;
```

The syntax of the second form of the `CONNECT` statement is:

```
EXEC SQL CONNECT TO <database_name>
[AS <connection_name>] [<credentials>];
```

Where `credentials` is one of the following:

```
USER user_name password
USER user_name IDENTIFIED BY password
USER user_name USING password
```

In the second form:

`database_name` is the name or identity of the database to which the client is connecting. Specify `database_name` as a variable, or as a string literal, in one of the following forms:

```
<database_name>[@<hostname>][:<port>]
tcp:postgresql://<hostname>[:<port>]/<database_name>[options]
unix:postgresql://<hostname>[:<port>]/<database_name>[options]
```

Where:

`hostname` is the name or IP address of the server on which the database resides.

`port` is the port on which the server listens.

You can also specify a value of `DEFAULT` to establish a connection with the default database, using the default role name. If you specify `DEFAULT` as the target database, do not include a `connection_name` or `credentials`.

`connection_name` is the name of the connection to the database. `connection_name` should take the form of an identifier (that is, not a string literal or a variable). You can open multiple connections, by providing a unique `connection_name` for each connection.

If you do not specify a name for a connection, `ecpglib` assigns a name of `DEFAULT` to the connection. You can refer to the connection by name (`DEFAULT`) in any `EXEC SQL` statement.

`CURRENT` is the most recently opened or the connection mentioned in the most-recent `SET CONNECTION TO` statement. If you do not refer to a connection by name in an `EXEC SQL` statement, ECPG assumes the name of the connection to be `CURRENT`.

`user_name` is the role used to establish the connection with the Advanced Server database. The privileges of the specified role will be applied to all commands performed through the connection.

`password` is the password associated with the specified `user_name`.

The following code fragment uses the second form of the `CONNECT` statement to establish a connection to a database named `edb`, using the role `alice` and the password associated with that role, `1safePWD`:

```
EXEC SQL CONNECT TO edb AS acctg_conn
USER 'alice' IDENTIFIED BY '1safePWD';
```

The name of the connection is `acctg_conn`; you can use the connection name when changing the connection name using the `SET CONNECTION` statement.

DEALLOCATE DESCRIPTOR

Use the `DEALLOCATE DESCRIPTOR` statement to free memory in use by an allocated descriptor. The syntax of the statement is:

```
EXEC SQL DEALLOCATE DESCRIPTOR <descriptor_name>
```

Where:

`descriptor_name` is the name of the descriptor. This value may take the form of a quoted string literal, or of a host variable.

The following example deallocates a descriptor named `emp_query`:

```
EXEC SQL DEALLOCATE DESCRIPTOR emp_query;
```

DECLARE CURSOR

Use the `DECLARE CURSOR` statement to define a cursor. The syntax of the statement is:

```
EXEC SQL [AT <database_name>] DECLARE <cursor_name> CURSOR FOR
(<select_statement> | <statement_name>);
```

Where:

`database_name` is the name of the database on which the cursor operates. This value may take the form of an identifier or of a host variable. If you do not specify a database name, the default value of `database_name` is the default database.

`cursor_name` is the name of the cursor.

`select_statement` is the text of the `SELECT` statement that defines the cursor result set; the `SELECT` statement cannot contain an `INTO` clause.

`statement_name` is the name of a SQL statement or block that defines the cursor result set.

The following example declares a cursor named `employees`:

```
EXEC SQL DECLARE employees CURSOR FOR
SELECT
  empno, ename, sal, comm
FROM
  emp;
```

The cursor generates a result set that contains the employee number, employee name, salary and commission for each employee record that is stored in the `emp` table.

DECLARE DATABASE

Use the `DECLARE DATABASE` statement to declare a database identifier for use in subsequent SQL statements (for example, in a `CONNECT` statement). The syntax is:

```
EXEC SQL DECLARE <database_name> DATABASE;
```

Where:

`database_name` specifies the name of the database.

The following example demonstrates declaring an identifier for the `acctg` database:

```
EXEC SQL DECLARE acctg DATABASE;
```

After invoking the command declaring `acctg` as a database identifier, the `acctg` database can be referenced by name when establishing a connection or in `AT` clauses.

This statement has no effect and is provided for Pro*C compatibility only.

DECLARE STATEMENT

Use the `DECLARE STATEMENT` directive to declare an identifier for an SQL statement. Advanced Server supports two versions of the `DECLARE STATEMENT` directive:

```
EXEC SQL [<database_name>] DECLARE <statement_name> STATEMENT;
```

and

```
EXEC SQL DECLARE STATEMENT <statement_name>;
```

Where:

`statement_name` specifies the identifier associated with the statement.

`database_name` specifies the name of the database. This value may take the form of an identifier or of a host variable that contains the identifier.

A typical usage sequence that includes the `DECLARE STATEMENT` directive might be:

```
EXEC SQL DECLARE give_raise STATEMENT; // give_raise is now a statement
handle (not prepared)
EXEC SQL PREPARE give_raise FROM :stmtText; // give_raise is now associated
with a statement
```

```
EXEC SQL EXECUTE give_raise;
```

This statement has no effect and is provided for Pro*C compatibility only.

DELETE

Use the `DELETE` statement to delete one or more rows from a table. The syntax for the ECPGPlus `DELETE` statement is the same as the syntax for the SQL statement, but you can use parameter markers and host variables any place that an expression is allowed. The syntax is:

```
[FOR <exec_count>] DELETE FROM [ONLY] <table> [[AS] <alias>]
[USING <using_list>]
[WHERE <condition> | WHERE CURRENT OF <cursor_name>]
[&{RETURNING|RETURN} * | <output_expression> [[ AS] <output_name>]
[, ...] INTO <host_variable_list> ]
```

Where:

Include the `FOR exec_count` clause to specify the number of times the statement will execute; this clause is valid only if the `VALUES` clause references an array or a pointer to an array.

`table` is the name (optionally schema-qualified) of an existing table. Include the `ONLY` clause to limit processing to the specified table; if you do not include the `ONLY` clause, any tables inheriting from the named table are also processed.

`alias` is a substitute name for the target table.

`using_list` is a list of table expressions, allowing columns from other tables to appear in the `WHERE` condition.

Include the `WHERE` clause to specify which rows should be deleted. If you do not include a `WHERE` clause in the statement, `DELETE` will delete all rows from the table, leaving the table definition intact.

`condition` is an expression, host variable or parameter marker that returns a value of type `BOOLEAN`. Those rows for which `condition` returns true will be deleted.

`cursor_name` is the name of the cursor to use in the `WHERE CURRENT OF` clause; the row to be deleted will be the one most recently fetched from this cursor. The cursor must be a non-grouping query on the `DELETE` statements target table. You cannot specify `WHERE CURRENT OF` in a `DELETE` statement that includes a Boolean condition.

The `RETURN/RETURNING` clause specifies an `output_expression` or `host_variable_list` that is returned by the `DELETE` command after each row is deleted:

`output_expression` is an expression to be computed and returned by the `DELETE` command after each row is deleted. `output_name` is the name of the returned column; include * to return all columns.

`host_variable_list` is a comma-separated list of host variables and optional indicator variables. Each host variable receives a corresponding value from the `RETURNING` clause.

For example, the following statement deletes all rows from the `emp` table where the `sal` column contains a value greater than the value specified in the host variable, `:max_sal`:

```
DELETE FROM emp WHERE sal > :max_sal;
```

For more information about using the `DELETE` statement, see the PostgreSQL Core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-delete.html>

DESCRIBE

Use the `DESCRIBE` statement to find the number of input values required by a prepared statement or the number of output values returned by a prepared statement. The `DESCRIBE` statement is used to analyze a SQL statement whose shape is unknown at the time you write your application.

The `DESCRIBE` statement populates an `SQLDA` descriptor; to populate a SQL descriptor, use the `ALLOCATE DESCRIPTOR` and `DESCRIBE...DESCRIPTOR` statements.

```
EXEC SQL DESCRIBE BIND VARIABLES FOR <statement_name> INTO
<descriptor>;
```

or

```
EXEC SQL DESCRIBE SELECT LIST FOR <statement_name> INTO <descriptor>;
```

Where:

`statement_name` is the identifier associated with a prepared SQL statement or PL/SQL block.

`descriptor` is the name of C variable of type `SQLDA*`. You must allocate the space for the descriptor by calling `sqlald()` (and initialize the descriptor) before executing the `DESCRIBE` statement.

When you execute the first form of the `DESCRIBE` statement, ECPG populates the given descriptor with a description of each input variable *required* by the statement. For example, given two descriptors:

```
SQLDA *query_values_in;
SQLDA *query_values_out;
```

You might prepare a query that returns information from the `emp` table:

```
EXEC SQL PREPARE get_emp FROM
"SELECT ename, empno, sal FROM emp WHERE empno = ?";
```

The command requires one input variable (for the parameter marker (?)).

```
EXEC SQL DESCRIBE BIND VARIABLES
FOR get_emp INTO query_values_in;
```

After describing the bind variables for this statement, you can examine the descriptor to find the number of variables required and the type of each variable.

When you execute the second form, ECPG populates the given descriptor with a description of each value *returned* by the statement. For example, the following statement returns three values:

```
EXEC SQL DESCRIBE SELECT LIST
FOR get_emp INTO query_values_out;
```

After describing the select list for this statement, you can examine the descriptor to find the number of returned values and the name and type of each value.

Before *executing* the statement, you must bind a variable for each input value and a variable for each output value. The variables that you bind for the input values specify the actual values used by the statement. The variables that you bind for the output values tell ECPGPlus where to put the values when you execute the statement.

This is alternate Pro*C compatible syntax for the `DESCRIBE DESCRIPTOR` statement.

DESCRIBE DESCRIPTOR

Use the **DESCRIBE DESCRIPTOR** statement to retrieve information about a SQL statement, and store that information in a SQL descriptor. Before using **DESCRIBE DESCRIPTOR**, you must allocate the descriptor with the **ALLOCATE DESCRIPTOR** statement. The syntax is:

```
EXEC SQL DESCRIBE [INPUT | OUTPUT] <statement_identifier>
  USING [SQL] DESCRIPTOR <descriptor_name>;
```

Where:

statement_name is the name of a prepared SQL statement.

descriptor_name is the name of the descriptor. **descriptor_name** can be a quoted string value or a host variable that contains the name of the descriptor.

If you include the **INPUT** clause, ECPGPlus populates the given descriptor with a description of each input variable *required* by the statement.

For example, given two descriptors:

```
EXEC SQL ALLOCATE DESCRIPTOR query_values_in;
EXEC SQL ALLOCATE DESCRIPTOR query_values_out;
```

You might prepare a query that returns information from the **emp** table:

```
EXEC SQL PREPARE get_emp FROM
  "SELECT ename, empno, sal FROM emp WHERE empno = ?";
```

The command requires one input variable (for the parameter marker (?)).

```
EXEC SQL DESCRIBE INPUT get_emp USING 'query_values_in';
```

After describing the bind variables for this statement, you can examine the descriptor to find the number of variables required and the type of each variable.

If you do not specify the **INPUT** clause, **DESCRIBE DESCRIPTOR** populates the specified descriptor with the values returned by the statement.

If you include the **OUTPUT** clause, ECPGPlus populates the given descriptor with a description of each value *returned* by the statement.

For example, the following statement returns three values:

```
EXEC SQL DESCRIBE OUTPUT FOR get_emp USING 'query_values_out';
```

After describing the select list for this statement, you can examine the descriptor to find the number of returned values and the name and type of each value.

DISCONNECT

Use the **DISCONNECT** statement to close the connection to the server. The syntax is:

```
EXEC SQL DISCONNECT [<connection_name>][CURRENT][DEFAULT][ALL];
```

Where:

`connection_name` is the connection name specified in the `CONNECT` statement used to establish the connection. If you do not specify a connection name, the current connection is closed.

Include the `CURRENT` keyword to specify that ECPGPlus should close the most-recently used connection.

Include the `DEFAULT` keyword to specify that ECPGPlus should close the connection named `DEFAULT`. If you do not specify a name when opening a connection, ECPGPlus assigns the name, `DEFAULT`, to the connection.

Include the `ALL` keyword to instruct ECPGPlus to close all active connections.

The following example creates a connection (named `hr_connection`) that connects to the `hr` database, and then disconnects from the connection:

```
/* client.pgc*/
int main()
{
    EXEC SQL CONNECT TO hr AS connection_name;
    EXEC SQL DISCONNECT connection_name;
    return(0);
}
```

EXECUTE

Use the `EXECUTE` statement to execute a statement previously prepared using an `EXEC SQL PREPARE` statement. The syntax is:

```
EXEC SQL [FOR <array_size>] EXECUTE <statement_name>
[USING {DESCRIPTOR <SQLDA_descriptor>
|: <host_variable> [|INDICATOR] :<indicator_variable>]}];
```

Where:

`array_size` is an integer value or a host variable that contains an integer value that specifies the number of rows to be processed. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`statement_name` specifies the name assigned to the statement when the statement was created (using the `EXEC SQL PREPARE` statement).

Include the `USING` clause to supply values for parameters within the prepared statement:

Include the `DESCRIPTOR` `SQLDA_descriptor` clause to provide an SQLDA descriptor value for a parameter.

Use a `host_variable` (and an optional `indicator_variable`) to provide a user-specified value for a parameter.

The following example creates a prepared statement that inserts a record into the `emp` table:

```
EXEC SQL PREPARE add_emp (numeric, text, text, numeric) AS
    INSERT INTO emp VALUES($1, $2, $3, $4);
```

Each time you invoke the prepared statement, provide fresh parameter values for the statement:

```
EXEC SQL EXECUTE add_emp USING 8000, 'DAWSON', 'CLERK', 7788;
EXEC SQL EXECUTE add_emp USING 8001, 'EDWARDS', 'ANALYST', 7698;
```

EXECUTE DESCRIPTOR

Use the `EXECUTE` statement to execute a statement previously prepared by an `EXEC SQL PREPARE` statement, using an SQL descriptor. The syntax is:

```
EXEC SQL [FOR <array_size>] EXECUTE <statement_identifier>
[USING [SQL] DESCRIPTOR <descriptor_name>]
[INTO [SQL] DESCRIPTOR <descriptor_name>];
```

Where:

`array_size` is an integer value or a host variable that contains an integer value that specifies the number of rows to be processed. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`statement_identifier` specifies the identifier assigned to the statement with the `EXEC SQL PREPARE` statement.

Include the `USING` clause to specify values for any input parameters required by the prepared statement.

Include the `INTO` clause to specify a descriptor into which the `EXECUTE` statement will write the results returned by the prepared statement.

`descriptor_name` specifies the name of a descriptor (as a single-quoted string literal), or a host variable that contains the name of a descriptor.

The following example executes the prepared statement, `give_raise`, using the values contained in the descriptor `stmtText`:

```
EXEC SQL PREPARE give_raise FROM :stmtText;
EXEC SQL EXECUTE give_raise USING DESCRIPTOR :stmtText;
```

EXECUTE...END EXEC

Use the `EXECUTE...END-EXEC` statement to embed an anonymous block into a client application. The syntax is:

```
EXEC SQL [AT <database_name>] EXECUTE <anonymous_block> END-EXEC;
```

Where:

`database_name` is the database identifier or a host variable that contains the database identifier. If you omit the `AT` clause, the statement will be executed on the current default database.

`anonymous_block` is an inline sequence of PL/pgSQL or SPL statements and declarations. You may include host variables and optional indicator variables within the block; each such variable is treated as an `IN/OUT` value.

The following example executes an anonymous block:

```
EXEC SQL EXECUTE
BEGIN
  IF (current_user = :admin_user_name) THEN
    DBMS_OUTPUT.PUT_LINE('You are an administrator');
  END IF;
END-EXEC;
```

Note

The `EXECUTE...END EXEC` statement is supported only by Advanced Server.

EXECUTE IMMEDIATE

Use the `EXECUTE IMMEDIATE` statement to execute a string that contains a SQL command. The syntax is:

```
EXEC SQL [AT <database_name>] EXECUTE IMMEDIATE <command_text>;
```

Where:

`database_name` is the database identifier or a host variable that contains the database identifier. If you omit the `AT` clause, the statement will be executed on the current default database.

`command_text` is the command executed by the `EXECUTE IMMEDIATE` statement.

This dynamic SQL statement is useful when you don't know the text of an SQL statement (ie., when writing a client application). For example, a client application may prompt a (trusted) user for a statement to execute. After the user provides the text of the statement as a string value, the statement is then executed with an `EXECUTE IMMEDIATE` command.

The statement text may not contain references to host variables. If the statement may contain parameter markers or returns one or more values, you must use the `PREPARE` and `DESCRIBE` statements.

The following example executes the command contained in the `:command_text` host variable:

```
EXEC SQL EXECUTE IMMEDIATE :command_text;
```

FETCH

Use the `FETCH` statement to return rows from a cursor into an SQLDA descriptor or a target list of host variables. Before using a `FETCH` statement to retrieve information from a cursor, you must prepare the cursor using `DECLARE` and `OPEN` statements. The statement syntax is:

```
EXEC SQL [FOR <array_size>] FETCH <cursor>
{ USING DESCRIPTOR <SQLDA_descriptor> }{| INTO <target_list> };
```

Where:

`array_size` is an integer value or a host variable that contains an integer value specifying the number of rows to fetch. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`cursor` is the name of the cursor from which rows are being fetched, or a host variable that contains the name of the cursor.

If you include a `USING` clause, the `FETCH` statement will populate the specified SQLDA descriptor with the values returned by the server.

If you include an `INTO` clause, the `FETCH` statement will populate the host variables (and optional indicator variables) specified in the `target_list`.

The following code fragment declares a cursor named `employees` that retrieves the `employee number`, `name` and `salary` from the `emp` table:

```
EXEC SQL DECLARE employees CURSOR
  SELECT empno, ename, esal FROM emp
EXEC SQL OPEN emp_cursor
EXEC SQL FETCH emp_cursor INTO :emp_no, :emp_name, :emp_sal;
```

FETCH DESCRIPTOR

Use the `FETCH DESCRIPTOR` statement to retrieve rows from a cursor into an SQL descriptor. The syntax is:

```
EXEC SQL [FOR <array_size>] FETCH <cursor>
INTO [SQL] DESCRIPTOR <descriptor_name>;
```

Where:

`array_size` is an integer value or a host variable that contains an integer value specifying the number of rows to fetch. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`cursor` is the name of the cursor from which rows are fetched, or a host variable that contains the name of the cursor. The client must `DECLARE` and `OPEN` the cursor before calling the `FETCH DESCRIPTOR` statement.

Include the `INTO` clause to specify an SQL descriptor into which the `EXECUTE` statement will write the results returned by the prepared statement. `descriptor_name` specifies the name of a descriptor (as a single-quoted string literal), or a host variable that contains the name of a descriptor. Prior to use, the descriptor must be allocated using an `ALLOCATE DESCRIPTOR` statement.

The following example allocates a descriptor named `row_desc` that will hold the description and the values of a specific row in the result set. It then declares and opens a cursor for a prepared statement (`my_cursor`), before looping through the rows in result set, using a `FETCH` to retrieve the next row from the cursor into the descriptor:

```
EXEC SQL ALLOCATE DESCRIPTOR 'row_desc';
EXEC SQL DECLARE my_cursor CURSOR FOR query;
EXEC SQL OPEN my_cursor;

for( row = 0; ; row++)
{
    EXEC SQL BEGIN DECLARE SECTION;
    int col;
    EXEC SQL END DECLARE SECTION;
    EXEC SQL FETCH my_cursor INTO SQL DESCRIPTOR 'row_desc';
```

GET DESCRIPTOR

Use the `GET DESCRIPTOR` statement to retrieve information from a descriptor. The `GET DESCRIPTOR` statement comes in two forms. The first form returns the number of values (or columns) in the descriptor.

```
EXEC SQL GET DESCRIPTOR <descriptor_name>
:<host_variable> = COUNT;
```

The second form returns information about a specific value (specified by the `VALUE column_number` clause).

```
EXEC SQL [FOR <array_size>] GET DESCRIPTOR <descriptor_name>
VALUE <column_number> {:<host_variable> = <descriptor_item> {...}};
```

Where:

`array_size` is an integer value or a host variable that contains an integer value that specifies the number of rows to be processed. If you specify an `array_size`, the `host_variable` must be an array of that size; for example, if `array_size` is `10`, `:host_variable` must be a 10-member array of `host_variables`. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`descriptor_name` specifies the name of a descriptor (as a single-quoted string literal), or a host variable that contains the name of a descriptor.

Include the `VALUE` clause to specify the information retrieved from the descriptor.

`column_number` identifies the position of the variable within the descriptor.

`host_variable` specifies the name of the host variable that will receive the value of the item.

`descriptor_item` specifies the type of the retrieved descriptor item.

ECPGPlus implements the following `descriptor_item` types:

- `TYPE`
- `LENGTH`
- `OCTET LENGTH`
- `RETURNED LENGTH`
- `RETURNED_OCTET_LENGTH`
- `PRECISION`
- `SCALE`
- `NULLABLE`
- `INDICATOR`
- `DATA`
- `NAME`

The following code fragment demonstrates using a `GET DESCRIPTOR` statement to obtain the number of columns entered in a user-provided string:

```
EXEC SQL ALLOCATE DESCRIPTOR parse_desc;
EXEC SQL PREPARE query FROM :stmt;
EXEC SQL DESCRIBE query INTO SQL DESCRIPTOR parse_desc;
EXEC SQL GET DESCRIPTOR parse_desc :col_count = COUNT;
```

The example allocates an SQL descriptor (named `parse_desc`), before using a `PREPARE` statement to syntax check the string provided by the user (`:stmt`). A `DESCRIBE` statement moves the user-provided string into the descriptor, `parse_desc`. The call to `EXEC SQL GET DESCRIPTOR` interrogates the descriptor to discover the number of columns (`:col_count`) in the result set.

INSERT

Use the `INSERT` statement to add one or more rows to a table. The syntax for the ECPGPlus `INSERT` statement is the same as the syntax for the SQL statement, but you can use parameter markers and host variables any place that a value is allowed. The syntax is:

```
[FOR <exec_count>] INSERT INTO <table> [(<column> [, ...])]
{DEFAULT VALUES |
 VALUES ({<expression> | DEFAULT} [, ...])[, ...] | <query>}
[RETURNING * | <output_expression> [[ AS ] <output_name>] [, ...]]
```

Where:

Include the `FOR exec_count` clause to specify the number of times the statement will execute; this clause is valid only if the `VALUES` clause references an array or a pointer to an array.

`table` specifies the (optionally schema-qualified) name of an existing table.

`column` is the name of a column in the table. The column name may be qualified with a subfield name or array subscript. Specify the `DEFAULT VALUES` clause to use default values for all columns.

`expression` is the expression, value, host variable or parameter marker that will be assigned to the corresponding column. Specify `DEFAULT` to fill the corresponding column with its default value.

`query` specifies a `SELECT` statement that supplies the row(s) to be inserted.

`output_expression` is an expression that will be computed and returned by the `INSERT` command after each row is inserted. The expression can refer to any column within the table. Specify `*` to return all columns of the inserted row(s).

`output_name` specifies a name to use for a returned column.

The following example adds a row to the `employees` table:

```
INSERT INTO emp (empno, ename, job, hiredate)
VALUES ('8400', :ename, 'CLERK', '2011-10-31');
```

Note

The `INSERT` statement uses a host variable (`:ename`) to specify the value of the `ename` column.

For more information about using the `INSERT` statement, see the PostgreSQL Core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-insert.html>

OPEN

Use the `OPEN` statement to open a cursor. The syntax is:

```
EXEC SQL [FOR <array_size>] OPEN <cursor> [USING <parameters>];
```

Where `parameters` is one of the following:

`DESCRIPTOR <SQLDA_descriptor>`

or

`<host_variable> [[INDICATOR] <indicator_variable>, ...]`

Where:

`array_size` is an integer value or a host variable that contains an integer value specifying the number of rows to fetch. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`cursor` is the name of the cursor being opened.

`parameters` is either `DESCRIPTOR SQLDA_descriptor` or a comma-separated list of `host variables` (and optional `indicator variables`) that initialize the cursor. If specifying an `SQLDA_descriptor`, the descriptor must be initialized with a `DESCRIBE` statement.

The `OPEN` statement initializes a cursor using the values provided in `parameters`. Once initialized, the cursor result set will remain unchanged unless the cursor is closed and re-opened. A cursor is automatically closed when an application terminates.

The following example declares a cursor named `employees`, that queries the `emp` table, returning the `employee number`, `name`, `salary` and `commission` of an employee whose name matches a user-supplied value (stored in the host variable, `:emp_name`).

```
EXEC SQL DECLARE employees CURSOR FOR
SELECT
    empno, ename, sal, comm
FROM
    emp
WHERE ename = :emp_name;
```

```
EXEC SQL OPEN employees;
```

...

After declaring the cursor, the example uses an `OPEN` statement to make the contents of the cursor available to a client application.

OPEN DESCRIPTOR

Use the `OPEN DESCRIPTOR` statement to open a cursor with a SQL descriptor. The syntax is:

```
EXEC SQL [FOR <array_size>] OPEN <cursor>
[USING [SQL] DESCRIPTOR <descriptor_name>]
[INTO [SQL] DESCRIPTOR <descriptor_name>];
```

Where:

`array_size` is an integer value or a host variable that contains an integer value specifying the number of rows to fetch. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`cursor` is the name of the cursor being opened.

`descriptor_name` specifies the name of an SQL descriptor (in the form of a single-quoted string literal) or a host variable that contains the name of an SQL descriptor that contains the query that initializes the cursor.

For example, the following statement opens a cursor (named `emp_cursor`), using the host variable, `:employees`:

```
EXEC SQL OPEN emp_cursor USING DESCRIPTOR :employees;
```

PREPARE

Prepared statements are useful when a client application must perform a task multiple times; the statement is parsed, written and planned only once, rather than each time the statement is executed, saving repetitive processing time.

Use the `PREPARE` statement to prepare an SQL statement or PL/pgSQL block for execution. The statement is available in two forms; the first form is:

```
EXEC SQL [AT <database_name>] PREPARE <statement_name>
FROM <sql_statement>;
```

The second form is:

```
EXEC SQL [AT <database_name>] PREPARE <statement_name>
AS <sql_statement>;
```

Where:

`database_name` is the database identifier or a host variable that contains the database identifier against which the statement will execute. If you omit the `AT` clause, the statement will execute against the current default database.

`statement_name` is the identifier associated with a prepared SQL statement or PL/SQL block.

`sql_statement` may take the form of a `SELECT` statement, a single-quoted string literal or host variable that contains the text of an SQL statement.

To include variables within a prepared statement, substitute placeholders (\$1, \$2, \$3, etc.) for statement values that might change when you **PREPARE** the statement. When you **EXECUTE** the statement, provide a value for each parameter. The values must be provided in the order in which they will replace placeholders.

The following example creates a prepared statement (named `add_emp`) that inserts a record into the `emp` table:

```
EXEC SQL PREPARE add_emp (int, text, text, numeric) AS
  INSERT INTO emp VALUES($1, $2, $3, $4);
```

Each time you invoke the statement, provide fresh parameter values for the statement:

```
EXEC SQL EXECUTE add_emp(8003, 'Davis', 'CLERK', 2000.00);
EXEC SQL EXECUTE add_emp(8004, 'Myer', 'CLERK', 2000.00);
```

Note

A client application must issue a **PREPARE** statement within each session in which a statement will be executed; prepared statements persist only for the duration of the current session.

ROLLBACK

Use the **ROLLBACK** statement to abort the current transaction, and discard any updates made by the transaction. The syntax is:

```
EXEC SQL [AT <database_name>] ROLLBACK [WORK]
[ { TO [SAVEPOINT] <savepoint> } | RELEASE ]
```

Where:

`database_name` is the database identifier or a host variable that contains the database identifier against which the statement will execute. If you omit the `AT` clause, the statement will execute against the current default database.

Include the `TO` clause to abort any commands that were executed after the specified `savepoint`; use the `SAVEPOINT` statement to define the `savepoint`. If you omit the `TO` clause, the **ROLLBACK** statement will abort the transaction, discarding all updates.

Include the `RELEASE` clause to cause the application to execute an **EXEC SQL COMMIT RELEASE** and close the connection.

Use the following statement to rollback a complete transaction:

```
EXEC SQL ROLLBACK;
```

Invoking this statement will abort the transaction, undoing all changes, erasing any savepoints, and releasing all transaction locks. If you include a savepoint (`my_savepoint` in the following example):

```
EXEC SQL ROLLBACK TO SAVEPOINT my_savepoint;
```

Only the portion of the transaction that occurred after the `my_savepoint` is rolled back; `my_savepoint` is retained, but any savepoints created after `my_savepoint` will be erased.

Rolling back to a specified savepoint releases all locks acquired after the savepoint.

SAVEPOINT

Use the `SAVEPOINT` statement to define a `savepoint`; a savepoint is a marker within a transaction. You can use a `ROLLBACK` statement to abort the current transaction, returning the state of the server to its condition prior to the specified savepoint. The syntax of a `SAVEPOINT` statement is:

```
EXEC SQL [AT <database_name>] SAVEPOINT <savepoint_name>
```

Where:

`database_name` is the database identifier or a host variable that contains the database identifier against which the savepoint resides. If you omit the `AT` clause, the statement will execute against the current default database.

`savepoint_name` is the name of the savepoint. If you re-use a `savepoint_name`, the original savepoint is discarded.

Savepoints can only be established within a transaction block. A transaction block may contain multiple savepoints.

To create a savepoint named `my_savepoint`, include the statement:

```
EXEC SQL SAVEPOINT my_savepoint;
```

SELECT

ECPGPlus extends support of the `SQL SELECT` statement by providing the `INTO host_variables` clause. The clause allows you to select specified information from an Advanced Server database into a host variable. The syntax for the `SELECT` statement is:

```
EXEC SQL [AT <database_name>]
SELECT
[ <hint> ]
[ ALL | DISTINCT [ ON( <expression>, ... ) ] ]
select_list INTO <host_variables>
[ FROM from_item [, from_item ]...]
[ WHERE condition ]
[ hierarchical_query_clause ]
[ GROUP BY expression [, ...]]
[ HAVING condition ]
[ { UNION [ ALL ] | INTERSECT | MINUS } (subquery) ]
[ ORDER BY expression> [order_by_options]]
[ LIMIT { count | ALL }]
[ OFFSET start [ ROW | ROWS ] ]
[ FETCH { FIRST | NEXT } [ count ] { ROW | ROWS } ONLY ]
[ FOR { UPDATE | SHARE } [OF table_name [, ...]][NOWAIT ][...]]
```

Where:

`database_name` is the name of the database (or host variable that contains the name of the database) in which the table resides. This value may take the form of an unquoted string literal, or of a host variable.

`host_variables` is a list of host variables that will be populated by the `SELECT` statement. If the `SELECT` statement returns more than a single row, `host_variables` must be an array.

ECPGPlus provides support for the additional clauses of the SQL `SELECT` statement as documented in the PostgreSQL Core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-select.html>

To use the `INTO host_variables` clause, include the names of defined host variables when specifying the `SELECT` statement. For example, the following `SELECT` statement populates the `:emp_name` and `:emp_sal` host variables with a list of `employee names` and `salaries`:

```
EXEC SQL SELECT ename, sal
  INTO :emp_name, :emp_sal
  FROM emp
 WHERE empno = 7988;
```

The enhanced `SELECT` statement also allows you to include parameter markers (question marks) in any clause where a value would be permitted. For example, the following query contains a parameter marker in the `WHERE` clause:

```
SELECT * FROM emp WHERE dept_no = ?;
```

This `SELECT` statement allows you to provide a value at run-time for the `dept_no` parameter marker.

SET CONNECTION

There are (at least) three reasons you may need more than one connection in a given client application:

- You may want different privileges for different statements;
- You may need to interact with multiple databases within the same client.
- Multiple threads of execution (within a client application) cannot share a connection concurrently.

The syntax for the `SET CONNECTION` statement is:

```
EXEC SQL SET CONNECTION <connection_name>;
```

Where:

`connection_name` is the name of the connection to the database.

To use the `SET CONNECTION` statement, you should open the connection to the database using the second form of the `CONNECT` statement; include the `AS` clause to specify a `connection_name`.

By default, the current thread uses the current connection; use the `SET CONNECTION` statement to specify a default connection for the current thread to use. The default connection is only used when you execute an `EXEC SQL` statement that does not explicitly specify a connection name. For example, the following statement will use the default connection because it does not include an `AT connection_name` clause. :

```
EXEC SQL DELETE FROM emp;
```

This statement will not use the default connection because it specifies a connection name using the `AT connection_name` clause:

```
EXEC SQL AT acctg_conn DELETE FROM emp;
```

For example, a client application that creates and maintains multiple connections (such as):

```
EXEC SQL CONNECT TO edb AS acctg_conn
  USER 'alice' IDENTIFIED BY 'acctpwd';
```

and

```
EXEC SQL CONNECT TO edb AS hr_conn
  USER 'bob' IDENTIFIED BY 'hrpwd';
```

Can change between the connections with the `SET CONNECTION` statement:

```
SET CONNECTION acctg_conn;
```

or

```
SET CONNECTION hr_conn;
```

The server will use the privileges associated with the connection when determining the privileges available to the connecting client. When using the `acctg_conn` connection, the client will have the privileges associated with the role, `alice`; when connected using `hr_conn`, the client will have the privileges associated with `bob`.

SET DESCRIPTOR

Use the `SET DESCRIPTOR` statement to assign a value to a descriptor area using information provided by the client application in the form of a host variable or an integer value. The statement comes in two forms; the first form is:

```
EXEC SQL [FOR <array_size>] SET DESCRIPTOR <descriptor_name>
  VALUE <column_number> <descriptor_item> = <host_variable>;
```

The second form is:

```
EXEC SQL [FOR <array_size>] SET DESCRIPTOR <descriptor_name>
  COUNT = integer;
```

Where:

`array_size` is an integer value or a host variable that contains an integer value specifying the number of rows to fetch. If you omit the `FOR` clause, the statement is executed once for each member of the array.

`descriptor_name` specifies the name of a descriptor (as a single-quoted string literal), or a host variable that contains the name of a descriptor.

Include the `VALUE` clause to describe the information stored in the descriptor.

`column_number` identifies the position of the variable within the descriptor.

`descriptor_item` specifies the type of the descriptor item.

`host_variable` specifies the name of the host variable that contains the value of the item.

ECPGPlus implements the following `descriptor_item` types:

- `TYPE`
- `LENGTH`
- `[REF] INDICATOR`
- `[REF] DATA`
- `[REF] RETURNED LENGTH`

For example, a client application might prompt a user for a dynamically created query:

```
query_text = promptUser("Enter a query");
```

To execute a dynamically created query, you must first `prepare` the query (parsing and validating the syntax of the query), and then `describe` the `input` parameters found in the query using the `EXEC SQL DESCRIBE INPUT` statement.

```
EXEC SQL ALLOCATE DESCRIPTOR query_params;
EXEC SQL PREPARE emp_query FROM :query_text;

EXEC SQL DESCRIBE INPUT emp_query
  USING SQL DESCRIPTOR 'query_params';
```

After describing the query, the `query_params` descriptor contains information about each parameter required by the query.

For this example, we'll assume that the user has entered:

```
SELECT ename FROM emp WHERE sal > ? AND job = ?;
```

In this case, the descriptor describes two parameters, one for `sal > ?` and one for `job = ?`.

To discover the number of parameter markers (question marks) in the query (and therefore, the number of values you must provide before executing the query), use:

```
EXEC SQL GET DESCRIPTOR ... :host_variable = COUNT;
```

Then, you can use `EXEC SQL GET DESCRIPTOR` to retrieve the name of each parameter. You can also use `EXEC SQL GET DESCRIPTOR` to retrieve the type of each parameter (along with the number of parameters) from the descriptor, or you can supply each `value` in the form of a character string and ECPG will convert that string into the required data type.

The data type of the first parameter is `numeric`; the type of the second parameter is `varchar`. The name of the first parameter is `sal`; the name of the second parameter is `job`.

Next, loop through each parameter, prompting the user for a value, and store those values in host variables. You can use `GET DESCRIPTOR ... COUNT` to find the number of parameters in the query.

```
EXEC SQL GET DESCRIPTOR 'query_params'
:param_count = COUNT;

for(param_number = 1;
    param_number <= param_count;
    param_number++)
{
```

Use `GET DESCRIPTOR` to copy the name of the parameter into the `param_name` host variable:

```
EXEC SQL GET DESCRIPTOR 'query_params'
  VALUE :param_number :param_name = NAME;

reply = promptUser(param_name);
if (reply == NULL)
  reply_ind = 1; /* NULL */
else
  reply_ind = 0; /* NOT NULL */
```

To associate a `value` with each parameter, you use the `EXEC SQL SET DESCRIPTOR` statement. For example:

```
EXEC SQL SET DESCRIPTOR 'query_params'
  VALUE :param_number DATA = :reply;
EXEC SQL SET DESCRIPTOR 'query_params'
  VALUE :param_number INDICATOR = :reply_ind;
{
```

Now, you can use the `EXEC SQL EXECUTE DESCRIPTOR` statement to execute the prepared statement on the server.

UPDATE

Use an `UPDATE` statement to modify the data stored in a table. The syntax is:

```
EXEC SQL [AT <database_name>][FOR <exec_count>]
  UPDATE [ ONLY ] table [ [ AS ] alias ]
    SET {column = { expression | DEFAULT } |
          (column [, ...]) = ({ expression|DEFAULT } [, ...])} [, ...]
    [ FROM from_list ]
    [ WHERE condition | WHERE CURRENT OF cursor_name ]
    [ RETURNING * | output_expression [[ AS ] output_name] [, ...] ]
```

Where:

`database_name` is the name of the database (or host variable that contains the name of the database) in which the table resides. This value may take the form of an unquoted string literal, or of a host variable.

Include the `FOR exec_count` clause to specify the number of times the statement will execute; this clause is valid only if the `SET` or `WHERE` clause contains an array.

ECPGPlus provides support for the additional clauses of the SQL `UPDATE` statement as documented in the PostgreSQL Core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-update.html>

A host variable can be used in any clause that specifies a value. To use a host variable, simply substitute a defined variable for any value associated with any of the documented `UPDATE` clauses.

The following `UPDATE` statement changes the job description of an employee (identified by the `:ename` host variable) to the value contained in the `:new_job` host variable, and increases the employees salary, by multiplying the current salary by the value in the `:increase` host variable:

```
EXEC SQL UPDATE emp
  SET job = :new_job, sal = sal * :increase
  WHERE ename = :ename;
```

The enhanced `UPDATE` statement also allows you to include parameter markers (question marks) in any clause where an input value would be permitted. For example, we can write the same update statement with a parameter marker in the `WHERE` clause:

```
EXEC SQL UPDATE emp
  SET job = ?, sal = sal * ?
  WHERE ename = :ename;
```

This `UPDATE` statement could allow you to prompt the user for a new value for the `job` column and provide the amount by which the `sal` column is incremented for the employee specified by `:ename`.

WHENEVER

Use the `WHENEVER` statement to specify the action taken by a client application when it encounters an SQL error or warning. The syntax is:

```
EXEC SQL WHENEVER <condition> <action>;
```

The following table describes the different conditions that might trigger an [action](#):

| Condition | Description |
|------------|---|
| NOT FOUND | The server returns a NOT FOUND condition when it encounters a SELECT that returns no rows, or when a FETCH reaches the end of a result set. |
| SQLERROR | The server returns an SQLERROR condition when it encounters a serious error returned by an SQL statement. |
| SQLWARNING | The server returns an SQLWARNING condition when it encounters a non-fatal warning returned by an SQL statement. |

The following table describes the actions that result from a client encountering a [condition](#):

| Action | Description |
|---------------------------|---|
| CALL function [(args)] | Instructs the client application to call the named function. |
| CONTINUE | Instructs the client application to proceed to the next statement. |
| DO BREAK | Instructs the client application to a C break statement. A break statement may appear in a loop or a switch statement. If executed, the break statement terminate the loop or the switch statement.. |
| DO CONTINUE | Instructs the client application to emit a C continue statement. A continue statement may only exist within a loop, and if executed, will cause the flow of control to return to the top of the loop. |
| DO function (args) | Instructs the client application to call the named function. |
| GOTO label or GO TO label | Instructs the client application to proceed to the statement that contains the label. |
| SQLPRINT | Instructs the client application to print a message to standard error. |
| STOP | Instructs the client application to stop execution. |

The following code fragment prints a message if the client application encounters a warning, and aborts the application if it encounters an error:

```
EXEC SQL WHENEVER SQLWARNING SQLPRINT;
EXEC SQL WHENEVER SQLERROR STOP;
```

Include the following code to specify that a client should continue processing after warning a user of a problem:

```
EXEC SQL WHENEVER SQLWARNING SQLPRINT;
```

Include the following code to call a function if a query returns no rows, or when a cursor reaches the end of a result set:

```
EXEC SQL WHENEVER NOT FOUND CALL error_handler(__LINE__);
```

10 Language Pack Guide

Language pack installers contain supported languages that may be used with EDB Postgres Advanced Server and

EnterpriseDB PostgreSQL database installers. The language pack installer allows you to install Perl, TCL/TK, and Python without installing supporting software from third party vendors. The Language Pack installer includes:

- TCL with TK version 8.6
- Perl version 5.26
- Python version 3.7

The Perl package contains the `cpan` package manager, and Python contains `pip` and `easy_install` package managers. There is no package manager for TCL/TK.

In previous Postgres releases, `plpython` was statically linked with ActiveState's python library. The Language Pack Installer dynamically links with our shared object for python. In ActiveState Linux installers for Python, there is no dynamic library. As a result of these changes, `plpython` will no longer work with ActiveState installers.

This document uses the term *Postgres* to mean either EDB Postgres Advanced Server or EDB PostgreSQL. For more information about using EDB Postgres products, please visit the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

StackBuilder Plus is distributed with Advanced Server; Stack Builder (distributed with PostgreSQL) provides comparable functionality. This document uses the term *StackBuilder Plus* to mean either StackBuilder Plus or Stack Builder.

10.1 Supported Database Server Versions

Language Pack installers are version and platform specific; select the Language Pack installer that corresponds to your Advanced Server or PostgreSQL server version:

Linux:

| Advanced Server or PostgreSQL Version | Language Pack Installer Version | Installs Language Versions |
|---------------------------------------|---------------------------------|--------------------------------|
| 9.4 | 9.4 | Perl 5.16, Python 3.3, Tcl 8.5 |
| 9.5 | 9.5 | Perl 5.20, Python 3.3, Tcl 8.5 |
| 9.6 through 10 | 1.0 | Perl 5.26, Python 3.7, Tcl 8.6 |

For detailed information about using an RPM package to add Language Pack, please see the EDB Postgres Advanced Server Installation Guide for Linux, available at:

<https://www.enterprisedb.com/edb-docs/p/edb-postgres-advanced-server>

Mac OS:

| Advanced Server or PostgreSQL Version | Language Pack Installer Version | Installs Language Version |
|---------------------------------------|---------------------------------|--------------------------------|
| 9.4 through 12 | 1.0 | Perl 5.26, Python 3.7, Tcl 8.6 |

Windows 32:

| Advanced Server or PostgreSQL Version | Language Pack Installer Version | Installs Language Version |
|---------------------------------------|---------------------------------|--------------------------------|
| 9.4 | 9.4 | Perl 5.16, Python 3.3, Tcl 8.5 |
| 9.5 | 9.5 | Perl 5.20, Python 3.3, Tcl 8.5 |

| Advanced Server or PostgreSQL Version | Language Pack Installer Version | Installs Language Version |
|---------------------------------------|---------------------------------|--------------------------------|
| 9.6 through 12 | 1.0 | Perl 5.26, Python 3.7, Tcl 8.6 |

Windows 64:

| Advanced Server or PostgreSQL Version | Language Pack Installer Version | Installs Language Version |
|---------------------------------------|---------------------------------|--------------------------------|
| 9.4 | 9.4 | Perl 5.16, Python 3.3, Tcl 8.5 |
| PostgreSQL 9.5 | 9.5 | Perl 5.20, Python 3.3, Tcl 8.5 |
| PostgreSQL 9.6 through 12 | 1.0 | Perl 5.26, Python 3.7, Tcl 8.8 |
| Advanced Server 12 | 1.0 | Perl 5.26, Python 3.7, Tcl 8.6 |


10.2 Installing Language Pack

The graphical installer is available from the EnterpriseDB website or via StackBuilder Plus.


Invoking the Graphical Installer

Assume Administrator privileges, and double-click the installer icon; if prompted, provide the password associated with the Administrator account. When prompted, select an installation language, and click OK.

The installer **Welcome** window opens.



Click **Next** to continue.




The **Ready to Install** window displays the Language Pack installation directory:


On Windows 64: `C:\edb\languagepack/v1`

On OSX: `/Library/edb/languagepack/v1`

You cannot modify the installation directory. Click **Next** to continue.




A progress bar marks installation progress; click **Next** to continue.



The installer will inform you that the Language Pack installation has completed; click **Finish** to exit the installer.

Installing Language Pack with StackBuilder Plus

You can use StackBuilder Plus to download and invoke the Language Pack graphical installer. To open StackBuilder Plus, select the StackBuilder Plus menu item from the version-specific EDB Postgres sub-menu.



Select your server from the drop-down menu on the StackBuilder Plus Welcome window and click **Next** to continue.

Expand the **Add-ons, tools and utilities** node of the **Categories** tree control, and check the box to the left of **EDB Language Pack**; click **Next** to continue.

StackBuilder Plus will confirm your package selection before downloading the installer. When the download completes, StackBuilder Plus will offer to invoke the installer for you, or will delay the installation until a more convenient time.

For details about using the graphical installer, see [Invoking the Graphical Installer](#).

Configuring Language Pack on an Advanced Server Host

After installing Language Pack on an Advanced Server host, you must configure the installation.

Configuring Language Pack on Windows

On Windows, the Language Pack installer places the languages in:

```
C:\edb\languagepack\v1
```

After installing Language Pack, you must set the following variables:

```
set PYTHONHOME=C:\edb\languagepack\v1\Python-3.7
```

Use the following commands to add Python, Perl and Tcl to your search path:

```
set PATH=C:\edb\languagepack\v1\Python-3.7;
C:\edb\languagepack\v1\Perl-5.26\bin;
C:\edb\languagepack\v1\Tcl-8.6\bin;%PATH%
```

After performing the steps required to configure Language Pack on Windows, use the Windows [Services](#) applet to restart the Advanced Server database server.

Configuring Language Pack on a PostgreSQL Host

After installing Language Pack on a PostgreSQL host, you must configure the installation.

Configuring Language Pack on Windows

After installing Language Pack, you must tell the Python interpreter where to find Python:

```
set PYTHONHOME=C:\edb\languagepack\v1\Python-3.7
```

Then, use the following commands to add Language Pack to your search path:

```
set PATH=C:\edb\languagepack\v1\Python-3.7;
C:\edb\languagepack\v1\Perl-5.26\bin;
C:\edb\languagepack\v1\Tcl-8.6\bin;%PATH%
```

After setting the system-specific steps required to configure Language Pack on Windows, restart the database server.

Configuring Language Pack on OSX

To simplify setting the value of `PATH` or `LD_LIBRARY_PATH`, you can create environment variables that identify the installation location:

```
PERLHOME=/Library/edb/languagepack/v1/Perl-5.26
PYTHONHOME=/Library/edb/languagepack/v1/Python-3.7
TCLHOME=/Library/edb/languagepack/v1/Tcl-8.6
```

Then, instruct the Python interpreter where to find Python:

```
export PYTHONHOME
```

You can use the same environment variables when setting the value of `PATH`:

```
export PATH=$PYTHONHOME/bin:
$PERLHOME/bin:
$TCLHOME/bin:$PATH
```

Lastly, use the variables to tell OSX where to find the shared libraries:

```
export DYLD_LIBRARY_PATH=$PYTHONHOME/lib:  
$PERLHOME/lib/CORE:$TCLHOME/lib:  
$DYLD_LIBRARY_PATH
```

10.3 Using the Procedural Languages

The Postgres procedural languages (PL/Perl, PL/Python, and PL/Java) are installed by the Language Pack installer. You can also use an RPM package to add procedural language functionality to your Advanced Server installation. For more information about using an RPM package, please see the EDB Advanced Server Installation Guide, available at:

<https://www.enterprisedb.com/edb-docs>

PL/Perl

The PL/Perl procedural language allows you to use Perl functions in Postgres applications.

You must install PL/Perl in each database (or in a template database) before creating a PL/Perl function. Use the `\CREATE LANGUAGE` command at the EDB-PSQL command line to install PL/Perl. Open the EDB-PSQL client, establish a connection to the database in which you wish to install PL/Perl, and enter the command:

`\CREATE EXTENSION plperl;`

You can now use a Postgres client application to access the features of the PL/Perl language. The following PL/Perl example creates a function named `perl_max` that returns the larger of two integer values:

```
CREATE OR REPLACE FUNCTION perl_max(integer, integer) RETURNS integer  
AS  
$$  
if ($_[0] > $_[1])  
{ return $_[0]; }  
return $_[1];  
$$ LANGUAGE plperl;
```

Pass two values when calling the function:

`\SELECT perl_max(1, 2);`

The server returns:

```
perl_max  
-----  
 2  
(1 row)
```

For more information about using the Perl procedural language, consult the official Postgres documentation available at:

<https://www.postgresql.org/docs/12/static/plperl.html>

PL/Python

The PL/Python procedural language allows you to create and execute functions written in Python within Postgres applications. The version of PL/Python used by Advanced Server and PostgreSQL is untrusted (`plpython3u`); it offers no restrictions on users to prevent potential security risks.

Install PL/Python in each database (or in a template database) before creating a PL/Python function. You can use the `CREATE LANGUAGE` command at the EDB-PSQL command line to install PL/Python. Use EDB-PSQL to connect to the database in which you wish to install PL/Python, and enter the command:

```
| CREATE EXTENSION plpython3u;
```

After installing PL/Python in your database, you can use the features of the PL/Python language.

Please note: The indentation shown in the following example must be included as you enter the sample function in EDB-PSQL. The following PL/Python example creates a function named `pymax` that returns the larger of two integer values:

```
CREATE OR REPLACE FUNCTION pymax (a integer, b integer) RETURNS
integer AS
$$
if a > b:
return a
return b
$$ LANGUAGE plpython3u;
```

When calling the `pymax` function, pass two values as shown below:

```
| SELECT pymax(12, 3);
```

The server returns:

```
pymax
-----
 12
(1 row)
```

For more information about using the Python procedural language, consult the official PostgreSQL documentation available at:

```
| https://www.postgresql.org/docs/12/static/plpython.html
```

PL/Tcl

The PL/Tcl procedural language allows you to use Tcl/Tk functions in applications.

You must install PL/Tcl in each database (or in a template database) before creating a PL/Tcl function. Use the `CREATE LANGUAGE` command at the EDB-PSQL command line to install PL/Tcl. Use the `psql` client to connect to the database in which you wish to install PL/Tcl, and enter the command:

```
| CREATE EXTENSION pltcl;
```

After creating the `pltcl` language, you can use the features of the PL/Tcl language from within your Postgres server.

The following PL/Tcl example creates a function named `tcl_max` that returns the larger of two integer values:

```

CREATE OR REPLACE FUNCTION tcl_max(integer, integer) RETURNS integer
AS $$$
if {[argisnull 1]} {
if {[argisnull 2]} { return_null }
return $2
}
if {[argisnull 2]} { return $1 }
if {$1 > $2} {return $1}
return $2
$$ LANGUAGE pltcl;

```

Pass two values when calling the function:

```
SELECT tcl_max(1, 2);
```

The server returns:

```
tcl_max
-----
2
(1 row)
```

For more information about using the Tcl procedural language, consult the official Postgres documentation available at:

<https://www.postgresql.org/docs/12/static/pltcl.html>

11 Advanced Server User Guide

Introduction

Advanced Server adds extended functionality to the open-source PostgreSQL database. The extended functionality supports database administration, enhanced SQL capabilities, database and application security, performance monitoring and analysis, and application development utilities. This guide documents those features that are exclusive to Advanced Server.

- **Enhanced Compatibility Features.** This section provides an overview of compatibility features supported by Advanced Server.
- **Database Administration.** This section contains information about features and tools that are helpful to the database administrator.
 - [Index Advisor](#) helps to determine the additional indexes needed on tables to improve application performance.
 - [SQL Profiler](#) locates and diagnoses poorly running SQL queries in applications.
 - [pgsnmpd](#) is an SNMP agent that returns hierarchical monitoring information regarding the current state of Advanced Server.
- **Security.** This section contains information about security features supported by Advanced Server.
 - [SQL/Protect](#) provides protection against SQL injection attacks.
 - [Virtual Private Database](#) provides fine-grained, row level access.
 - [sslutils](#) provides SSL certificate generation functions.
 - [Data redaction](#) provides protection against sensitive data exposure.

- **EDB Resource Manager.** This section contains information about the EDB Resource Manager feature, which provides the capability to control system resource usage by Advanced Server processes.
 - **Resource Groups** shows how to create and maintain the groups on which resource limits can be defined.
 - **CPU Usage Throttling** provides a method to control CPU usage by Advanced Server processes.
 - **Dirty Buffer Throttling** provides a method to control the dirty rate of shared buffers by Advanced Server processes.
- **The libpq C Library.** The **libpq C library** is the C application programming interface (API) language for Advanced Server.
- **The PL Debugger.** The **PL Debugger** is a graphically oriented debugging tool for PL/pgSQL.
- **Performance Analysis and Tuning.** This section contains the various tools for analyzing and improving application and database server performance.
 - **Dynatune** provides a quick and easy means for configuring Advanced Server depending upon the type of application usage.
 - **EDB wait states** provides a way to capture wait events and other data for performance diagnosis.
- **EDB Clone Schema.** This section contains information about the EDB Clone Schema feature, which provides the capability to copy a schema and its database objects within a single database or from one database to another database.
- **Enhanced SQL and Other Miscellaneous Features.** This section contains information on enhanced SQL functionality and other features that provide additional flexibility and convenience.
- **System Catalog Tables.** This section contains additional *system catalog tables* added for Advanced Server specific database objects.
- **Advanced Server Keywords.** This section contains information about the words that Advanced Server recognizes as keywords.

For information about the features that are shared by Advanced Server and PostgreSQL, see the PostgreSQL core documentation, available at:

<https://www.postgresql.org/docs/current/index.html>

What's New

The following features have been changed in EDB Postgres Advanced Server 12 to create Advanced Server 13:

- Advanced Server now supports the **PARALLEL | NOPARALLEL** clause in the **CREATE TABLE, ALTER TABLE, CREATE INDEX**, and **ALTER INDEX** commands to enable or disable parallelism on an index or a table. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- Advanced Server now supports **USING INDEX ... create_index_statement** clause in the **CREATE TABLE**, and **ALTER TABLE** commands to explicitly create an index on a table. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- Advanced Server introduces **AUTOMATIC LIST PARTITIONING**, which allows a database to automatically create a new partition if the partitioning key value does not correspond to an existing partition. For information, see the *Database Compatibility Table Partitioning Guide*.
- Advanced Server has added the **STATS_MODE** function that takes an argument a set of values to return the most frequent input value. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added the **SYSDATE** function that returns the timestamp without time zone. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added **smallint, init, bigint** and **numeric** variants to the **MEDIAN** function. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added support for the **FM** format in the **TO_NUMBER** function. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server now supports statement logging and records the number of rows processed during bulk execution. For information, see **edb_log_every_bulk_value** under [Reporting and Logging / What to Log](#).
- Advanced Server has added a new GUC to support Oracle-style display output. For information, see the *Database Compatibility for Oracle Developer's Built-in Package Guide*.
- EDB*Loader now supports **HANDLE_CONFLICTS** parameter to move duplicate records to the **BAD** file. For information, see the *Database Compatibility for Oracle Developer's Tools and Utilities Guide*.
- Advanced Server has added **ALTER DIRECTORY ...OWNER TO** command to change the owner of a

- directory. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- Advanced Server has added `ALTER TRIGGER ...ON AUTHORIZATION` command to change an owner of the trigger's implicit objects. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
 - EDB*Loader now supports `-c CONNECTION STRING` or `connstr=CONNECTION_STRING` parameter, which allows you to specify the SSL parameters and the connection parameters supported by libpq. For information, see the *Database Compatibility for Oracle Developer's Tools and Utilities Guide*.
 - Advanced Server has added the `TO_TIMESTAMP_TZ` function that converts a timestamp formatted string to a `TIMESTAMPTZ` datatype. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
 - Advanced Server has added additional fields to make logs consistent across both CSV and XML audit logs. For information, see [EDB Audit Logging](#).
 - Advanced Server has added a new GUC `edb_dblink_oci.rescan` to control the scan type for the remote statement. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.

Conventions

The following is a list of conventions used throughout this document.

- This guide applies to both Linux and Windows systems. Directory paths are presented in the Linux format with forward slashes. When working on Windows systems, start the directory path with the drive letter followed by a colon and substitute back slashes for forward slashes.
- Some of the information in this document may apply interchangeably to the PostgreSQL and EDB Postgres Advanced Server database systems. The term *Advanced Server* is used to refer to EDB Postgres Advanced Server. The term *Postgres* is used to generically refer to both PostgreSQL and Advanced Server. When a distinction needs to be made between these two database systems, the specific names, PostgreSQL or Advanced Server are used.
- The installation directory path of the PostgreSQL or Advanced Server products is referred to as `POSTGRES_INSTALL_HOME`.
 - For PostgreSQL Linux installations, this defaults to `/opt/PostgreSQL/x.x` for version 10 and earlier. For later versions, use the PostgreSQL community packages.
 - For Advanced Server Linux installations accomplished using the interactive installer for version 10 and earlier, this defaults to `/opt/edb/asx.x`.
 - For Advanced Server Linux installations accomplished using an RPM package, this defaults to `/usr/edb/asxx`.
 - For Advanced Server Windows installations, this defaults to `C:\Program Files\edb\as\xx`. The product version number is represented by `x.x` or by `xx` for version 10 and later.

Examples used

The examples in this guide are shown in the type and background illustrated below.

Examples and output from examples are shown in fixed-width, black font on a white background.

The examples use the sample tables, `dept`, `emp`, and `jobhist`, created and loaded when Advanced Server is installed.

The tables and programs in the sample database can be re-created at any time by executing the following script:

```
/usr/edb/asxx/share/pg-sample.sql
```

where `xx` is the Advanced Server version number.

In addition there is a script in the same directory containing the database objects created using syntax compatible with Oracle databases. This script file is `edb-sample.sql`.

The script:

- Creates the sample tables and programs in the currently connected database.

- Grants all permissions on the tables to the **PUBLIC** group.

The tables and programs will be created in the first schema of the search path in which the current user has permission to create tables and procedures. You can display the search path by issuing the command:

```
| SHOW SEARCH_PATH;
```

You can use PSQL commands to modify the search path.

Sample Database Description

The sample database represents employees in an organization. It contains three types of records: employees, departments, and historical records of employees.

Each employee has an identification number, name, hire date, salary, and manager. Some employees earn a commission in addition to their salary. All employee-related information is stored in the **emp** table.

The sample company is regionally diverse, so it tracks the locations of its departments. Each company employee is assigned to a department. Each department is identified by a unique department number and a short name. Each department is associated with one location. All department-related information is stored in the **dept** table.

The company also tracks information about jobs held by the employees. Some employees have been with the company for a long time and have held different positions, received raises, switched departments, etc. When a change in employee status occurs, the company records the end date of the former position. A new job record is added with the start date and the new job title, department, salary, and the reason for the status change. All employee history is maintained in the **jobhist** table.

The following is the **pg-sample.sql** script:

```
SET datestyle TO 'iso, dmy';

-- 
-- Script that creates the 'sample' tables, views
-- functions, triggers, etc.
-- 
-- Start new transaction - commit all or nothing
-- 
BEGIN;

-- 
-- Create and load tables used in the documentation examples.
-- 
-- Create the 'dept' table
-- 
CREATE TABLE dept (
    deptno NUMERIC(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname VARCHAR(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc VARCHAR(13)
);
-- 
-- Create the 'emp' table
-- 
CREATE TABLE emp (
    empno NUMERIC(4) NOT NULL CONSTRAINT emp_pk PRIMARY KEY,
    ename VARCHAR(10),
    job VARCHAR(9),
    mgr NUMERIC(4),
```

```

hiredate DATE,
sal NUMERIC(7,2) CONSTRAINT emp_sal_ck CHECK (sal > 0),
comm NUMERIC(7,2),
deptno NUMERIC(2) CONSTRAINT emp_ref_dept_fk
    REFERENCES dept(deptno)
);
-- 
-- Create the 'jobhist' table
-- 
CREATE TABLE jobhist (
    empno NUMERIC(4) NOT NULL,
    startdate TIMESTAMP(0) NOT NULL,
    enddate TIMESTAMP(0),
    job VARCHAR(9),
    sal NUMERIC(7,2),
    comm NUMERIC(7,2),
    deptno NUMERIC(2),
    chgdesc VARCHAR(80),
    CONSTRAINT jobhist_pk PRIMARY KEY (empno, startdate),
    CONSTRAINT jobhist_ref_emp_fk FOREIGN KEY (empno)
        REFERENCES emp(empno) ON DELETE CASCADE,
    CONSTRAINT jobhist_ref_dept_fk FOREIGN KEY (deptno)
        REFERENCES dept (deptno) ON DELETE SET NULL,
    CONSTRAINT jobhist_date_chk CHECK (startdate <= enddate)
);
-- 
-- Create the 'salesemp' view
-- 
CREATE OR REPLACE VIEW salesemp AS
    SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job = 'SALESMAN';
-- 
-- Sequence to generate values for function 'new_empno'.
-- 
CREATE SEQUENCE next_empno START WITH 8000 INCREMENT BY 1;
-- 
-- Issue PUBLIC grants
-- 
--GRANT ALL ON emp TO PUBLIC;
--GRANT ALL ON dept TO PUBLIC;
--GRANT ALL ON jobhist TO PUBLIC;
--GRANT ALL ON salesemp TO PUBLIC;
--GRANT ALL ON next_empno TO PUBLIC;
-- 
-- Load the 'dept' table
-- 
INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT INTO dept VALUES (20,'RESEARCH','DALLAS');
INSERT INTO dept VALUES (30,'SALES','CHICAGO');
INSERT INTO dept VALUES (40,'OPERATIONS','BOSTON');
-- 
-- Load the 'emp' table
-- 
INSERT INTO emp VALUES
(7369,'SMITH','CLERK',7902,'17-DEC-80',800,NULL,20);
INSERT INTO emp VALUES

```

```

(7499,'ALLEN','SALESMAN',7698,'20-FEB-81',1600,300,30);
INSERT INTO emp VALUES
(7521,'WARD','SALESMAN',7698,'22-FEB-81',1250,500,30);
INSERT INTO emp VALUES
(7566,'JONES','MANAGER',7839,'02-APR-81',2975,NULL,20);
INSERT INTO emp VALUES
(7654,'MARTIN','SALESMAN',7698,'28-SEP-81',1250,1400,30);
INSERT INTO emp VALUES
(7698,'BLAKE','MANAGER',7839,'01-MAY-81',2850,NULL,30);
INSERT INTO emp VALUES
(7782,'CLARK','MANAGER',7839,'09-JUN-81',2450,NULL,10);
INSERT INTO emp VALUES
(7788,'SCOTT','ANALYST',7566,'19-APR-87',3000,NULL,20);
INSERT INTO emp VALUES
(7839,'KING','PRESIDENT',NULL,'17-NOV-81',5000,NULL,10);
INSERT INTO emp VALUES
(7844,'TURNER','SALESMAN',7698,'08-SEP-81',1500,0,30);
INSERT INTO emp VALUES
(7876,'ADAMS','CLERK',7788,'23-MAY-87',1100,NULL,20);
INSERT INTO emp VALUES
(7900,'JAMES','CLERK',7698,'03-DEC-81',950,NULL,30);
INSERT INTO emp VALUES
(7902,'FORD','ANALYST',7566,'03-DEC-81',3000,NULL,20);
INSERT INTO emp VALUES
(7934,'MILLER','CLERK',7782,'23-JAN-82',1300,NULL,10);

-- 
-- Load the 'jobhist' table
-- 

INSERT INTO jobhist VALUES
(7369,'17-DEC-80',NULL,'CLERK',800,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7499,'20-FEB-81',NULL,'SALESMAN',1600,300,30,'New Hire');
INSERT INTO jobhist VALUES
(7521,'22-FEB-81',NULL,'SALESMAN',1250,500,30,'New Hire');
INSERT INTO jobhist VALUES
(7566,'02-APR-81',NULL,'MANAGER',2975,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7654,'28-SEP-81',NULL,'SALESMAN',1250,1400,30,'New Hire');
INSERT INTO jobhist VALUES
(7698,'01-MAY-81',NULL,'MANAGER',2850,NULL,30,'New Hire');
INSERT INTO jobhist VALUES
(7782,'09-JUN-81',NULL,'MANAGER',2450,NULL,10,'New Hire');
INSERT INTO jobhist VALUES
(7788,'19-APR-87','12-APR-88','CLERK',1000,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7788,'13-APR-88','04-MAY-89','CLERK',1040,NULL,20,'Raise');
INSERT INTO jobhist VALUES
(7788,'05-MAY-90',NULL,'ANALYST',3000,NULL,20,'Promoted to Analyst');
INSERT INTO jobhist VALUES
(7839,'17-NOV-81',NULL,'PRESIDENT',5000,NULL,10,'New Hire');
INSERT INTO jobhist VALUES
(7844,'08-SEP-81',NULL,'SALESMAN',1500,0,30,'New Hire');
INSERT INTO jobhist VALUES
(7876,'23-MAY-87',NULL,'CLERK',1100,NULL,20,'New Hire');
INSERT INTO jobhist VALUES

```

```

(7900,'03-DEC-81','14-JAN-83','CLERK',950,NULL,10,'New Hire');
INSERT INTO jobhist VALUES
(7900,'15-JAN-83',NULL,'CLERK',950,NULL,30,'Changed to Dept 30');
INSERT INTO jobhist VALUES
(7902,'03-DEC-81',NULL,'ANALYST',3000,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7934,'23-JAN-82',NULL,'CLERK',1300,NULL,10,'New Hire');

-- 
-- Populate statistics table and view (pg_statistic/pg_stats)
-- 

ANALYZE dept;
ANALYZE emp;
ANALYZE jobhist;

-- 
-- Function that lists all employees' numbers and names
-- from the 'emp' table using a cursor.

-- 
CREATE OR REPLACE FUNCTION list_emp() RETURNS VOID
AS $$$
DECLARE
    v_empno NUMERIC(4);
    v_ename VARCHAR(10);
    emp_cur CURSOR FOR
        SELECT empno, ename FROM emp ORDER BY empno;
BEGIN
    OPEN emp_cur;
    RAISE INFO 'EMPNO ENAME';
    RAISE INFO '----- -----';
    LOOP
        FETCH emp_cur INTO v_empno, v_ename;
        EXIT WHEN NOT FOUND;
        RAISE INFO '% %', v_empno, v_ename;
    END LOOP;
    CLOSE emp_cur;
    RETURN;
END;
$$ LANGUAGE 'plpgsql';

-- 
-- Function that selects an employee row given the employee
-- number and displays certain columns.

-- 
CREATE OR REPLACE FUNCTION select_emp (
    p_empno NUMERIC
) RETURNS VOID
AS $$$
DECLARE
    v_ename emp.ename%TYPE;
    v_hiredate emp.hiredate%TYPE;
    v_sal emp.sal%TYPE;
    v_comm emp.comm%TYPE;
    v_dname dept.dname%TYPE;
    v_disp_date VARCHAR(10);
BEGIN
    SELECT INTO
        v_ename, v_hiredate, v_sal, v_comm, v_dname

```

```

ename, hiredate, sal, COALESCE(comm, 0), dname
FROM emp e, dept d
WHERE empno = p_empno
AND e.deptno = d.deptno;
IF NOT FOUND THEN
  RAISE INFO 'Employee % not found', p_empno;
  RETURN;
END IF;
v_disp_date := TO_CHAR(v_hiredate, 'MM/DD/YYYY');
RAISE INFO 'Number : %', p_empno;
RAISE INFO 'Name : %', v_ename;
RAISE INFO 'Hire Date : %', v_disp_date;
RAISE INFO 'Salary : %', v_sal;
RAISE INFO 'Commission: %', v_comm;
RAISE INFO 'Department: %', v_dname;
RETURN;
EXCEPTION
WHEN OTHERS THEN
  RAISE INFO 'The following is SQLERRM : %', SQLERRM;
  RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
  RETURN;
END;
$$ LANGUAGE 'plpgsql';

-- 
-- A RECORD type used to format the return value of
-- function, 'emp_query'.
-- 

CREATE TYPE emp_query_type AS (
  empno NUMERIC,
  ename VARCHAR(10),
  job VARCHAR(9),
  hiredate DATE,
  sal NUMERIC
);
-- 
-- Function that queries the 'emp' table based on
-- department number and employee number or name. Returns
-- employee number and name as INOUT parameters and job,
-- hire date, and salary as OUT parameters. These are
-- returned in the form of a record defined by
-- RECORD type, 'emp_query_type'.
-- 

CREATE OR REPLACE FUNCTION emp_query (
  IN p_deptno NUMERIC,
  INOUT p_empno NUMERIC,
  INOUT p_ename VARCHAR,
  OUT p_job VARCHAR,
  OUT p_hiredate DATE,
  OUT p_sal NUMERIC
)
AS $$

BEGIN
  SELECT INTO
    p_empno, p_ename, p_job, p_hiredate, p_sal
    empno, ename, job, hiredate, sal

```

```

FROM emp
WHERE deptno = p_deptno
  AND (empno = p_empno
    OR ename = UPPER(p_ename)));
END;
$$ LANGUAGE 'plpgsql';
-- 
-- Function to call 'emp_query_caller' with IN and INOUT
-- parameters. Displays the results received from INOUT and
-- OUT parameters.
-- 
CREATE OR REPLACE FUNCTION emp_query_caller() RETURNS VOID
AS $$$
DECLARE
  v_deptno NUMERIC;
  v_empno NUMERIC;
  v_ename VARCHAR;
  v_rows INTEGER;
  r_emp_query EMP_QUERY_TYPE;
BEGIN
  v_deptno := 30;
  v_empno := 0;
  v_ename := 'Martin';
  r_emp_query := emp_query(v_deptno, v_empno, v_ename);
  RAISE INFO 'Department : %', v_deptno;
  RAISE INFO 'Employee No: %', (r_emp_query).empno;
  RAISE INFO 'Name : %', (r_emp_query).ename;
  RAISE INFO 'Job : %', (r_emp_query).job;
  RAISE INFO 'Hire Date : %', (r_emp_query).hiredate;
  RAISE INFO 'Salary : %', (r_emp_query).sal;
  RETURN;
EXCEPTION
  WHEN OTHERS THEN
    RAISE INFO 'The following is SQLERRM : %', SQLERRM;
    RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
    RETURN;
END;
$$ LANGUAGE 'plpgsql';
-- 
-- Function to compute yearly compensation based on semimonthly
-- salary.
-- 
CREATE OR REPLACE FUNCTION emp_comp (
  p_sal NUMERIC,
  p_comm NUMERIC
) RETURNS NUMERIC
AS $$$
BEGIN
  RETURN (p_sal + COALESCE(p_comm, 0)) \* 24;
END;
$$ LANGUAGE 'plpgsql';
-- 
-- Function that gets the next number from sequence, 'next_empno',
-- and ensures it is not already in use as an employee number.
-- 

```

```

CREATE OR REPLACE FUNCTION new_empno() RETURNS INTEGER
AS $$$
DECLARE
    v_cnt INTEGER := 1;
    v_new_empno INTEGER;
BEGIN
    WHILE v_cnt > 0 LOOP
        SELECT INTO v_new_empno nextval('next_empno');
        SELECT INTO v_cnt COUNT(*) FROM emp WHERE empno = v_new_empno;
    END LOOP;
    RETURN v_new_empno;
END;
$$ LANGUAGE 'plpgsql';
-- 
-- Function that adds a new clerk to table 'emp'.
-- 
CREATE OR REPLACE FUNCTION hire_clerk (
    p_ename VARCHAR,
    p_deptno NUMERIC
) RETURNS NUMERIC
AS $$$
DECLARE
    v_empno NUMERIC(4);
    v_ename VARCHAR(10);
    v_job VARCHAR(9);
    v_mgr NUMERIC(4);
    v_hiredate DATE;
    v_sal NUMERIC(7,2);
    v_comm NUMERIC(7,2);
    v_deptno NUMERIC(2);
BEGIN
    v_empno := new_empno();
    INSERT INTO emp VALUES (v_empno, p_ename, 'CLERK', 7782,
                           CURRENT_DATE, 950.00, NULL, p_deptno);
    SELECT INTO
        v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno
        empno, ename, job, mgr, hiredate, sal, comm, deptno
    FROM emp WHERE empno = v_empno;
    RAISE INFO 'Department : %', v_deptno;
    RAISE INFO 'Employee No: %', v_empno;
    RAISE INFO 'Name : %', v_ename;
    RAISE INFO 'Job : %', v_job;
    RAISE INFO 'Manager : %', v_mgr;
    RAISE INFO 'Hire Date : %', v_hiredate;
    RAISE INFO 'Salary : %', v_sal;
    RAISE INFO 'Commission : %', v_comm;
    RETURN v_empno;
EXCEPTION
    WHEN OTHERS THEN
        RAISE INFO 'The following is SQLERRM : %', SQLERRM;
        RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
        RETURN -1;
END;
$$ LANGUAGE 'plpgsql';
-- 

```

```
-- Function that adds a new salesman to table 'emp'.
--
CREATE OR REPLACE FUNCTION hire_salesman (
    p_ename VARCHAR,
    p_sal NUMERIC,
    p_comm NUMERIC
) RETURNS NUMERIC
AS $$$
DECLARE
    v_empno NUMERIC(4);
    v_ename VARCHAR(10);
    v_job VARCHAR(9);
    v_mgr NUMERIC(4);
    v_hiredate DATE;
    v_sal NUMERIC(7,2);
    v_comm NUMERIC(7,2);
    v_deptno NUMERIC(2);
BEGIN
    v_empno := new_empno();
    INSERT INTO emp VALUES (v_empno, p_ename, 'SALESMAN', 7698,
        CURRENT_DATE, p_sal, p_comm, 30);
    SELECT INTO
        v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno
        empno, ename, job, mgr, hiredate, sal, comm, deptno
    FROM emp WHERE empno = v_empno;
    RAISE INFO 'Department : %', v_deptno;
    RAISE INFO 'Employee No: %', v_empno;
    RAISE INFO 'Name : %', v_ename;
    RAISE INFO 'Job : %', v_job;
    RAISE INFO 'Manager : %', v_mgr;
    RAISE INFO 'Hire Date : %', v_hiredate;
    RAISE INFO 'Salary : %', v_sal;
    RAISE INFO 'Commission : %', v_comm;
    RETURN v_empno;
EXCEPTION
    WHEN OTHERS THEN
        RAISE INFO 'The following is SQLERRM : %', SQLERRM;
        RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
        RETURN -1;
END;
$$ LANGUAGE 'plpgsql';

-- Rule to INSERT into view 'salesemp'
--
CREATE OR REPLACE RULE salesemp_i AS ON INSERT TO salesemp
DO INSTEAD
    INSERT INTO emp VALUES (NEW.empno, NEW.ename, 'SALESMAN', 7698,
        NEW.hiredate, NEW.sal, NEW.comm, 30);

-- Rule to UPDATE view 'salesemp'
--
CREATE OR REPLACE RULE salesemp_u AS ON UPDATE TO salesemp
DO INSTEAD
    UPDATE emp SET empno = NEW.empno,
        ename = NEW.ename,
```

```

        hiredate = NEW.hiredate,
        sal     = NEW.sal,
        comm    = NEW.comm
    WHERE empno = OLD.empno;

-- 
-- Rule to DELETE from view 'salesemp'
--

CREATE OR REPLACE RULE salesemp_d AS ON DELETE TO salesemp
DO INSTEAD
    DELETE FROM emp WHERE empno = OLD.empno;

-- 
-- After statement-level trigger that displays a message after
-- an insert, update, or deletion to the 'emp' table. One message
-- per SQL command is displayed.
--

CREATE OR REPLACE FUNCTION user_audit_trig() RETURNS TRIGGER
AS $$$
DECLARE
    v_action VARCHAR(24);
    v_text TEXT;
BEGIN
    IF TG_OP = 'INSERT' THEN
        v_action := ' added employee(s) on ';
    ELSIF TG_OP = 'UPDATE' THEN
        v_action := ' updated employee(s) on ';
    ELSIF TG_OP = 'DELETE' THEN
        v_action := ' deleted employee(s) on ';
    END IF;
    v_text := 'User ' || USER || v_action || CURRENT_DATE;
    RAISE INFO '%', v_text;
    RETURN NULL;
END;
$$ LANGUAGE 'plpgsql';
CREATE TRIGGER user_audit_trig
    AFTER INSERT OR UPDATE OR DELETE ON emp
    FOR EACH STATEMENT EXECUTE PROCEDURE user_audit_trig();

-- 
-- Before row-level trigger that displays employee number and
-- salary of an employee that is about to be added, updated,
-- or deleted in the 'emp' table.
-- 

CREATE OR REPLACE FUNCTION emp_sal_trig() RETURNS TRIGGER
AS $$$
DECLARE
    sal_diff NUMERIC(7,2);
BEGIN
    IF TG_OP = 'INSERT' THEN
        RAISE INFO 'Inserting employee %', NEW.empno;
        RAISE INFO '..New salary: %', NEW.sal;
        RETURN NEW;
    END IF;
    IF TG_OP = 'UPDATE' THEN
        sal_diff := NEW.sal - OLD.sal;
        RAISE INFO 'Updating employee %', OLD.empno;
        RAISE INFO '..Old salary: %', OLD.sal;
    END IF;
END;
$$ LANGUAGE 'plpgsql';

```

```

RAISE INFO '..New salary: %', NEW.sal;
RAISE INFO '..Raise : %', sal_diff;
RETURN NEW;
END IF;
IF TG_OP = 'DELETE' THEN
  RAISE INFO 'Deleting employee %', OLD.empno;
  RAISE INFO '..Old salary: %', OLD.sal;
  RETURN OLD;
END IF;
END;
$$ LANGUAGE 'plpgsql';
CREATE TRIGGER emp_sal_trig
  BEFORE DELETE OR INSERT OR UPDATE ON emp
  FOR EACH ROW EXECUTE PROCEDURE emp_sal_trig();
COMMIT;

```

11.1 Introduction

Advanced Server adds extended functionality to the open-source PostgreSQL database. The extended functionality supports database administration, enhanced SQL capabilities, database and application security, performance monitoring and analysis, and application development utilities. This guide documents those features that are exclusive to Advanced Server.

- **Enhanced Compatibility Features.** This section provides an overview of compatibility features supported by Advanced Server.
- **Database Administration.** This section contains information about features and tools that are helpful to the database administrator.
 - [Index Advisor](#) helps to determine the additional indexes needed on tables to improve application performance.
 - [SQL Profiler](#) locates and diagnoses poorly running SQL queries in applications.
 - [pgsnmpd](#) is an SNMP agent that returns hierarchical monitoring information regarding the current state of Advanced Server.
- **Security.** This section contains information about security features supported by Advanced Server.
 - [SQL/Protect](#) provides protection against SQL injection attacks.
 - [Virtual Private Database](#) provides fine-grained, row level access.
 - [sslutils](#) provides SSL certificate generation functions.
 - [Data redaction](#) provides protection against sensitive data exposure.
- **EDB Resource Manager.** This section contains information about the EDB Resource Manager feature, which provides the capability to control system resource usage by Advanced Server processes.
 - [Resource Groups](#) shows how to create and maintain the groups on which resource limits can be defined.
 - [CPU Usage Throttling](#) provides a method to control CPU usage by Advanced Server processes.
 - [Dirty Buffer Throttling](#) provides a method to control the dirty rate of shared buffers by Advanced Server processes.
- **The libpq C Library.** The [libpq C library](#) is the C application programming interface (API) language for Advanced Server.
- **The PL Debugger.** The [PL Debugger](#) is a graphically oriented debugging tool for PL/pgSQL.
- **Performance Analysis and Tuning.** This section contains the various tools for analyzing and improving application and database server performance.
 - [Dynatune](#) provides a quick and easy means for configuring Advanced Server depending upon the type of application usage.
 - [EDB wait states](#) provides a way to capture wait events and other data for performance diagnosis.

- **EDB Clone Schema.** This section contains information about the EDB Clone Schema feature, which provides the capability to copy a schema and its database objects within a single database or from one database to another database.
- **Enhanced SQL and Other Miscellaneous Features.** This section contains information on enhanced SQL functionality and other features that provide additional flexibility and convenience.
- **System Catalog Tables.** This section contains additional system catalog tables added for Advanced Server specific database objects.
- **Advanced Server Keywords.** This section contains information about the words that Advanced Server recognizes as keywords.

For information about the features that are shared by Advanced Server and PostgreSQL, see the PostgreSQL core documentation, available at:

<https://www.postgresql.org/docs/current/index.html>

11.1.1 What's New

The following features have been changed in EDB Postgres Advanced Server 12 to create Advanced Server 13:

- Advanced Server now supports the `PARALLEL | NOPARALLEL` clause in the `CREATE TABLE, ALTER TABLE, CREATE INDEX`, and `ALTER INDEX` commands to enable or disable parallelism on an index or a table. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- Advanced Server now supports `USING INDEX ... create_index_statement` clause in the `CREATE TABLE`, and `ALTER TABLE` commands to explicitly create an index on a table. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- Advanced Server introduces `AUTOMATIC LIST PARTITIONING`, which allows a database to automatically create a new partition if the partitioning key value does not correspond to an existing partition. For information, see the *Database Compatibility Table Partitioning Guide*.
- Advanced Server has added the `STATS_MODE` function that takes an argument a set of values to return the most frequent input value. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added the `SYSDATE` function that returns the timestamp without time zone. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added `smallint, init, bigint` and `numeric` variants to the `MEDIAN` function. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added support for the `FM` format in the `TO_NUMBER` function. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server now supports statement logging and records the number of rows processed during bulk execution. For information, see `edb_log_every_bulk_value` under [Reporting and Logging / What to Log](#).
- Advanced Server has added a new GUC to support Oracle-style display output. For information, see the *Database Compatibility for Oracle Developer's Built-in Package Guide*.
- EDB*Loader now supports `HANDLE_CONFLICTS` parameter to move duplicate records to the `BAD` file. For information, see the *Database Compatibility for Oracle Developer's Tools and Utilities Guide*.
- Advanced Server has added `ALTER DIRECTORY ...OWNER TO` command to change the owner of a directory. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- Advanced Server has added `ALTER TRIGGER ...ON AUTHORIZATION` command to change an owner of the trigger's implicit objects. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.
- EDB*Loader now supports `-c CONNECTION STRING` or `connstr=CONNECTION_STRING` parameter, which allows you to specify the SSL parameters and the connection parameters supported by libpq. For information, see the *Database Compatibility for Oracle Developer's Tools and Utilities Guide*.
- Advanced Server has added the `TO_TIMESTAMP_TZ` function that converts a timestamp formatted string to a `TIMESTAMPTZ` datatype. For information, see the *Database Compatibility for Oracle Developer's Reference Guide*.
- Advanced Server has added additional fields to make logs consistent across both CSV and XML audit logs. For information, see [EDB Audit Logging](#).
- Advanced Server has added a new GUC `edb_dblink_oci.rescan` to control the scan type for the remote

statement. For information, see the *Database Compatibility for Oracle Developers SQL Guide*.

11.1.2 Conventions Used in this Guide

The following is a list of conventions used throughout this document.

- This guide applies to both Linux and Windows systems. Directory paths are presented in the Linux format with forward slashes. When working on Windows systems, start the directory path with the drive letter followed by a colon and substitute back slashes for forward slashes.
- Some of the information in this document may apply interchangeably to the PostgreSQL and EDB Postgres Advanced Server database systems. The term *Advanced Server* is used to refer to EDB Postgres Advanced Server. The term *Postgres* is used to generically refer to both PostgreSQL and Advanced Server. When a distinction needs to be made between these two database systems, the specific names, PostgreSQL or Advanced Server are used.
- The installation directory path of the PostgreSQL or Advanced Server products is referred to as `POSTGRES_INSTALL_HOME`.
 - For PostgreSQL Linux installations, this defaults to `/opt/PostgreSQL/x.x` for version 10 and earlier. For later versions, use the PostgreSQL community packages.
 - For Advanced Server Linux installations accomplished using the interactive installer for version 10 and earlier, this defaults to `/opt/edb/asx.x`.
 - For Advanced Server Linux installations accomplished using an RPM package, this defaults to `/usr/edb/asxx`.
 - For Advanced Server Windows installations, this defaults to `C:\Program Files\edb\as\xx`. The product version number is represented by `x.x` or by `xx` for version 10 and later.

11.1.3 About the Examples Used in this Guide

The examples in this guide are shown in the type and background illustrated below.

Examples and output from examples are shown in fixed-width, black font on a white background.

The examples use the sample tables, `dept`, `emp`, and `jobhist`, created and loaded when Advanced Server is installed.

The tables and programs in the sample database can be re-created at any time by executing the following script:

```
/usr/edb/asxx/share/pg-sample.sql
```

where `xx` is the Advanced Server version number.

In addition there is a script in the same directory containing the database objects created using syntax compatible with Oracle databases. This script file is `edb-sample.sql`.

The script:

- Creates the sample tables and programs in the currently connected database.
- Grants all permissions on the tables to the `PUBLIC` group.

The tables and programs will be created in the first schema of the search path in which the current user has permission to create tables and procedures. You can display the search path by issuing the command:

```
| SHOW SEARCH_PATH;
```

You can use PSQL commands to modify the search path.

Sample Database Description

The sample database represents employees in an organization. It contains three types of records: employees, departments, and historical records of employees.

Each employee has an identification number, name, hire date, salary, and manager. Some employees earn a commission in addition to their salary. All employee-related information is stored in the `emp` table.

The sample company is regionally diverse, so it tracks the locations of its departments. Each company employee is assigned to a department. Each department is identified by a unique department number and a short name. Each department is associated with one location. All department-related information is stored in the `dept` table.

The company also tracks information about jobs held by the employees. Some employees have been with the company for a long time and have held different positions, received raises, switched departments, etc. When a change in employee status occurs, the company records the end date of the former position. A new job record is added with the start date and the new job title, department, salary, and the reason for the status change. All employee history is maintained in the `jobhist` table.

The following is the `pg-sample.sql` script:

```
SET datestyle TO 'iso, dmy';

-- Script that creates the 'sample' tables, views
-- functions, triggers, etc.

-- Start new transaction - commit all or nothing
--

BEGIN;

-- Create and load tables used in the documentation examples.
--

-- Create the 'dept' table
--

CREATE TABLE dept (
    deptno NUMERIC(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname VARCHAR(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc VARCHAR(13)
);
-- Create the 'emp' table
--

CREATE TABLE emp (
    empno NUMERIC(4) NOT NULL CONSTRAINT emp_pk PRIMARY KEY,
    ename VARCHAR(10),
    job VARCHAR(9),
    mgr NUMERIC(4),
    hiredate DATE,
    sal NUMERIC(7,2) CONSTRAINT emp_sal_ck CHECK (sal > 0),
    comm NUMERIC(7,2),
    deptno NUMERIC(2) CONSTRAINT emp_ref_dept_fk
        REFERENCES dept(deptno)
```

```

);
-- 
-- Create the 'jobhist' table
-- 

CREATE TABLE jobhist (
    empno NUMERIC(4) NOT NULL,
    startdate TIMESTAMP(0) NOT NULL,
    enddate TIMESTAMP(0),
    job VARCHAR(9),
    sal NUMERIC(7,2),
    comm NUMERIC(7,2),
    deptno NUMERIC(2),
    chgdesc VARCHAR(80),
    CONSTRAINT jobhist_pk PRIMARY KEY (empno, startdate),
    CONSTRAINT jobhist_ref_emp_fk FOREIGN KEY (empno)
        REFERENCES emp(empno) ON DELETE CASCADE,
    CONSTRAINT jobhist_ref_dept_fk FOREIGN KEY (deptno)
        REFERENCES dept (deptno) ON DELETE SET NULL,
    CONSTRAINT jobhist_date_chk CHECK (startdate <= enddate)
);

-- 
-- Create the 'salesemp' view
-- 

CREATE OR REPLACE VIEW salesemp AS
SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job = 'SALESMAN';
-- 
-- Sequence to generate values for function 'new_empno'.
-- 

CREATE SEQUENCE next_empno START WITH 8000 INCREMENT BY 1;
-- 
-- Issue PUBLIC grants
-- 

--GRANT ALL ON emp TO PUBLIC;
--GRANT ALL ON dept TO PUBLIC;
--GRANT ALL ON jobhist TO PUBLIC;
--GRANT ALL ON salesemp TO PUBLIC;
--GRANT ALL ON next_empno TO PUBLIC;
-- 
-- Load the 'dept' table
-- 

INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT INTO dept VALUES (20,'RESEARCH','DALLAS');
INSERT INTO dept VALUES (30,'SALES','CHICAGO');
INSERT INTO dept VALUES (40,'OPERATIONS','BOSTON');
-- 
-- Load the 'emp' table
-- 

INSERT INTO emp VALUES
(7369,'SMITH','CLERK',7902,'17-DEC-80',800,NULL,20);
INSERT INTO emp VALUES
(7499,'ALLEN','SALESMAN',7698,'20-FEB-81',1600,300,30);
INSERT INTO emp VALUES
(7521,'WARD','SALESMAN',7698,'22-FEB-81',1250,500,30);
INSERT INTO emp VALUES
(7566,'JONES','MANAGER',7839,'02-APR-81',2975,NULL,20);

```

```

INSERT INTO emp VALUES
(7654,'MARTIN','SALESMAN',7698,'28-SEP-81',1250,1400,30);
INSERT INTO emp VALUES
(7698,'BLAKE','MANAGER',7839,'01-MAY-81',2850,NULL,30);
INSERT INTO emp VALUES
(7782,'CLARK','MANAGER',7839,'09-JUN-81',2450,NULL,10);
INSERT INTO emp VALUES
(7788,'SCOTT','ANALYST',7566,'19-APR-87',3000,NULL,20);
INSERT INTO emp VALUES
(7839,'KING','PRESIDENT',NULL,'17-NOV-81',5000,NULL,10);
INSERT INTO emp VALUES
(7844,'TURNER','SALESMAN',7698,'08-SEP-81',1500,0,30);
INSERT INTO emp VALUES
(7876,'ADAMS','CLERK',7788,'23-MAY-87',1100,NULL,20);
INSERT INTO emp VALUES
(7900,'JAMES','CLERK',7698,'03-DEC-81',950,NULL,30);
INSERT INTO emp VALUES
(7902,'FORD','ANALYST',7566,'03-DEC-81',3000,NULL,20);
INSERT INTO emp VALUES
(7934,'MILLER','CLERK',7782,'23-JAN-82',1300,NULL,10);

-- 
-- Load the 'jobhist' table
-- 

INSERT INTO jobhist VALUES
(7369,'17-DEC-80',NULL,'CLERK',800,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7499,'20-FEB-81',NULL,'SALESMAN',1600,300,30,'New Hire');
INSERT INTO jobhist VALUES
(7521,'22-FEB-81',NULL,'SALESMAN',1250,500,30,'New Hire');
INSERT INTO jobhist VALUES
(7566,'02-APR-81',NULL,'MANAGER',2975,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7654,'28-SEP-81',NULL,'SALESMAN',1250,1400,30,'New Hire');
INSERT INTO jobhist VALUES
(7698,'01-MAY-81',NULL,'MANAGER',2850,NULL,30,'New Hire');
INSERT INTO jobhist VALUES
(7782,'09-JUN-81',NULL,'MANAGER',2450,NULL,10,'New Hire');
INSERT INTO jobhist VALUES
(7788,'19-APR-87','12-APR-88','CLERK',1000,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7788,'13-APR-88','04-MAY-89','CLERK',1040,NULL,20,'Raise');
INSERT INTO jobhist VALUES
(7788,'05-MAY-90',NULL,'ANALYST',3000,NULL,20,'Promoted to Analyst');
INSERT INTO jobhist VALUES
(7839,'17-NOV-81',NULL,'PRESIDENT',5000,NULL,10,'New Hire');
INSERT INTO jobhist VALUES
(7844,'08-SEP-81',NULL,'SALESMAN',1500,0,30,'New Hire');
INSERT INTO jobhist VALUES
(7876,'23-MAY-87',NULL,'CLERK',1100,NULL,20,'New Hire');
INSERT INTO jobhist VALUES
(7900,'03-DEC-81','14-JAN-83','CLERK',950,NULL,10,'New Hire');
INSERT INTO jobhist VALUES
(7900,'15-JAN-83',NULL,'CLERK',950,NULL,30,'Changed to Dept 30');
INSERT INTO jobhist VALUES
(7902,'03-DEC-81',NULL,'ANALYST',3000,NULL,20,'New Hire');

```

```

INSERT INTO jobhist VALUES
(7934,'23-JAN-82',NULL,'CLERK',1300,NULL,10,'New Hire');
-- 
-- Populate statistics table and view (pg_statistic/pg_stats)
-- 

ANALYZE dept;
ANALYZE emp;
ANALYZE jobhist;

-- 
-- Function that lists all employees' numbers and names
-- from the 'emp' table using a cursor.
-- 

CREATE OR REPLACE FUNCTION list_emp() RETURNS VOID
AS $$$
DECLARE
    v_empno NUMERIC(4);
    v_ename VARCHAR(10);
    emp_cur CURSOR FOR
        SELECT empno, ename FROM emp ORDER BY empno;
BEGIN
    OPEN emp_cur;
    RAISE INFO 'EMPNO ENAME';
    RAISE INFO '----- -----';
    LOOP
        FETCH emp_cur INTO v_empno, v_ename;
        EXIT WHEN NOT FOUND;
        RAISE INFO '% %', v_empno, v_ename;
    END LOOP;
    CLOSE emp_cur;
    RETURN;
END;
$$ LANGUAGE 'plpgsql';

-- 
-- Function that selects an employee row given the employee
-- number and displays certain columns.
-- 

CREATE OR REPLACE FUNCTION select_emp (
    p_empno NUMERIC
) RETURNS VOID
AS $$$
DECLARE
    v_ename emp.ename%TYPE;
    v_hiredate emp.hiredate%TYPE;
    v_sal emp.sal%TYPE;
    v_comm emp.comm%TYPE;
    v_dname dept.dname%TYPE;
    v_disp_date VARCHAR(10);
BEGIN
    SELECT INTO
        v_ename, v_hiredate, v_sal, v_comm, v_dname
        ename, hiredate, sal, COALESCE(comm, 0), dname
    FROM emp e, dept d
    WHERE empno = p_empno
        AND e.deptno = d.deptno;
    IF NOT FOUND THEN

```

```

        RAISE INFO 'Employee % not found', p_empno;
        RETURN;
    END IF;
    v_disp_date := TO_CHAR(v_hiredate, 'MM/DD/YYYY');
    RAISE INFO 'Number : %', p_empno;
    RAISE INFO 'Name : %', v_ename;
    RAISE INFO 'Hire Date : %', v_disp_date;
    RAISE INFO 'Salary : %', v_sal;
    RAISE INFO 'Commission: %', v_comm;
    RAISE INFO 'Department: %', v_dname;
    RETURN;
EXCEPTION
WHEN OTHERS THEN
    RAISE INFO 'The following is SQLERRM : %', SQLERRM;
    RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
    RETURN;
END;
$$ LANGUAGE 'plpgsql';

-- 
-- A RECORD type used to format the return value of
-- function, 'emp_query'.

-- 
CREATE TYPE emp_query_type AS (
    empno NUMERIC,
    ename VARCHAR(10),
    job VARCHAR(9),
    hiredate DATE,
    sal NUMERIC
);
-- 
-- Function that queries the 'emp' table based on
-- department number and employee number or name. Returns
-- employee number and name as INOUT parameters and job,
-- hire date, and salary as OUT parameters. These are
-- returned in the form of a record defined by
-- RECORD type, 'emp_query_type'.
-- 
CREATE OR REPLACE FUNCTION emp_query (
    IN p_deptno NUMERIC,
    INOUT p_empno NUMERIC,
    INOUT p_ename VARCHAR,
    OUT p_job VARCHAR,
    OUT p_hiredate DATE,
    OUT p_sal NUMERIC
)
AS $$$
BEGIN
    SELECT INTO
        p_empno, p_ename, p_job, p_hiredate, p_sal
        empno, ename, job, hiredate, sal
    FROM emp
    WHERE deptno = p_deptno
        AND (empno = p_empno
            OR ename = UPPER(p_ename));
END;

```

```

$$ LANGUAGE 'plpgsql';
-- 
-- Function to call 'emp_query_caller' with IN and INOUT
-- parameters. Displays the results received from INOUT and
-- OUT parameters.
-- 
CREATE OR REPLACE FUNCTION emp_query_caller() RETURNS VOID
AS $$$
DECLARE
    v_deptno NUMERIC;
    v_empno NUMERIC;
    v_ename VARCHAR;
    v_rows INTEGER;
    r_emp_query EMP_QUERY_TYPE;
BEGIN
    v_deptno := 30;
    v_empno := 0;
    v_ename := 'Martin';
    r_emp_query := emp_query(v_deptno, v_empno, v_ename);
    RAISE INFO 'Department : %', v_deptno;
    RAISE INFO 'Employee No: %', (r_emp_query).empno;
    RAISE INFO 'Name : %', (r_emp_query).ename;
    RAISE INFO 'Job : %', (r_emp_query).job;
    RAISE INFO 'Hire Date : %', (r_emp_query).hiredate;
    RAISE INFO 'Salary : %', (r_emp_query).sal;
    RETURN;
EXCEPTION
    WHEN OTHERS THEN
        RAISE INFO 'The following is SQLERRM : %', SQLERRM;
        RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
        RETURN;
END;
$$ LANGUAGE 'plpgsql';
-- 
-- Function to compute yearly compensation based on semimonthly
-- salary.
-- 
CREATE OR REPLACE FUNCTION emp_comp (
    p_sal NUMERIC,
    p_comm NUMERIC
) RETURNS NUMERIC
AS $$$
BEGIN
    RETURN (p_sal + COALESCE(p_comm, 0)) \* 24;
END;
$$ LANGUAGE 'plpgsql';
-- 
-- Function that gets the next number from sequence, 'next_empno',
-- and ensures it is not already in use as an employee number.
-- 
CREATE OR REPLACE FUNCTION new_empno() RETURNS INTEGER
AS $$$
DECLARE
    v_cnt INTEGER := 1;
    v_new_empno INTEGER;

```

```

BEGIN
  WHILE v_cnt > 0 LOOP
    SELECT INTO v_new_empno nextval('next_empno');
    SELECT INTO v_cnt COUNT(*) FROM emp WHERE empno = v_new_empno;
  END LOOP;
  RETURN v_new_empno;
END;
$$ LANGUAGE 'plpgsql';

-- Function that adds a new clerk to table 'emp'.
-- CREATE OR REPLACE FUNCTION hire_clerk (
  p_ename VARCHAR,
  p_deptno NUMERIC
) RETURNS NUMERIC
AS $$

DECLARE
  v_empno NUMERIC(4);
  v_ename VARCHAR(10);
  v_job VARCHAR(9);
  v_mgr NUMERIC(4);
  v_hiredate DATE;
  v_sal NUMERIC(7,2);
  v_comm NUMERIC(7,2);
  v_deptno NUMERIC(2);

BEGIN
  v_empno := new_empno();
  INSERT INTO emp VALUES (v_empno, p_ename, 'CLERK', 7782,
    CURRENT_DATE, 950.00, NULL, p_deptno);
  SELECT INTO
    v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno
    empno, ename, job, mgr, hiredate, sal, comm, deptno
    FROM emp WHERE empno = v_empno;
  RAISE INFO 'Department : %', v_deptno;
  RAISE INFO 'Employee No: %', v_empno;
  RAISE INFO 'Name : %', v_ename;
  RAISE INFO 'Job : %', v_job;
  RAISE INFO 'Manager : %', v_mgr;
  RAISE INFO 'Hire Date : %', v_hiredate;
  RAISE INFO 'Salary : %', v_sal;
  RAISE INFO 'Commission : %', v_comm;
  RETURN v_empno;
EXCEPTION
  WHEN OTHERS THEN
    RAISE INFO 'The following is SQLERRM : %', SQLERRM;
    RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
    RETURN -1;
END;
$$ LANGUAGE 'plpgsql';

-- Function that adds a new salesman to table 'emp'.
-- CREATE OR REPLACE FUNCTION hire_salesman (
  p_ename VARCHAR,
  p_sal NUMERIC,

```

```

p_comm NUMERIC
) RETURNS NUMERIC
AS $$

DECLARE
    v_empno NUMERIC(4);
    v_ename VARCHAR(10);
    v_job VARCHAR(9);
    v_mgr NUMERIC(4);
    v_hiredate DATE;
    v_sal NUMERIC(7,2);
    v_comm NUMERIC(7,2);
    v_deptno NUMERIC(2);

BEGIN
    v_empno := new_empno();
    INSERT INTO emp VALUES (v_empno, p_ename, 'SALESMAN', 7698,
                           CURRENT_DATE, p_sal, p_comm, 30);
    SELECT INTO
        v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno
        empno, ename, job, mgr, hiredate, sal, comm, deptno
    FROM emp WHERE empno = v_empno;
    RAISE INFO 'Department : %', v_deptno;
    RAISE INFO 'Employee No: %', v_empno;
    RAISE INFO 'Name : %', v_ename;
    RAISE INFO 'Job : %', v_job;
    RAISE INFO 'Manager : %', v_mgr;
    RAISE INFO 'Hire Date : %', v_hiredate;
    RAISE INFO 'Salary : %', v_sal;
    RAISE INFO 'Commission : %', v_comm;
    RETURN v_empno;
EXCEPTION
    WHEN OTHERS THEN
        RAISE INFO 'The following is SQLERRM : %', SQLERRM;
        RAISE INFO 'The following is SQLSTATE: %', SQLSTATE;
        RETURN -1;
END;
$$ LANGUAGE 'plpgsql';

-- Rule to INSERT into view 'salesemp'
-- CREATE OR REPLACE RULE salesemp_i AS ON INSERT TO salesemp
DO INSTEAD
    INSERT INTO emp VALUES (NEW.empno, NEW.ename, 'SALESMAN', 7698,
                           NEW.hiredate, NEW.sal, NEW.comm, 30);

-- Rule to UPDATE view 'salesemp'
-- CREATE OR REPLACE RULE salesemp_u AS ON UPDATE TO salesemp
DO INSTEAD
    UPDATE emp SET empno = NEW.empno,
                  ename = NEW.ename,
                  hiredate = NEW.hiredate,
                  sal = NEW.sal,
                  comm = NEW.comm
    WHERE empno = OLD.empno;

```

```
-- Rule to DELETE from view 'salesemp'
--
CREATE OR REPLACE RULE salesemp_d AS ON DELETE TO salesemp
DO INSTEAD
    DELETE FROM emp WHERE empno = OLD.empno;
--
-- After statement-level trigger that displays a message after
-- an insert, update, or deletion to the 'emp' table. One message
-- per SQL command is displayed.
--
CREATE OR REPLACE FUNCTION user_audit_trig() RETURNS TRIGGER
AS $$$
DECLARE
    v_action VARCHAR(24);
    v_text TEXT;
BEGIN
    IF TG_OP = 'INSERT' THEN
        v_action := ' added employee(s) on ';
    ELSIF TG_OP = 'UPDATE' THEN
        v_action := ' updated employee(s) on ';
    ELSIF TG_OP = 'DELETE' THEN
        v_action := ' deleted employee(s) on ';
    END IF;
    v_text := 'User ' || USER || v_action || CURRENT_DATE;
    RAISE INFO '%', v_text;
    RETURN NULL;
END;
$$ LANGUAGE 'plpgsql';
CREATE TRIGGER user_audit_trig
AFTER INSERT OR UPDATE OR DELETE ON emp
FOR EACH STATEMENT EXECUTE PROCEDURE user_audit_trig();
--
-- Before row-level trigger that displays employee number and
-- salary of an employee that is about to be added, updated,
-- or deleted in the 'emp' table.
--
CREATE OR REPLACE FUNCTION emp_sal_trig() RETURNS TRIGGER
AS $$$
DECLARE
    sal_diff NUMERIC(7,2);
BEGIN
    IF TG_OP = 'INSERT' THEN
        RAISE INFO 'Inserting employee %', NEW.empno;
        RAISE INFO '..New salary: %', NEW.sal;
        RETURN NEW;
    END IF;
    IF TG_OP = 'UPDATE' THEN
        sal_diff := NEW.sal - OLD.sal;
        RAISE INFO 'Updating employee %', OLD.empno;
        RAISE INFO '..Old salary: %', OLD.sal;
        RAISE INFO '..New salary: %', NEW.sal;
        RAISE INFO '..Raise : %', sal_diff;
        RETURN NEW;
    END IF;
    IF TG_OP = 'DELETE' THEN

```

```

RAISE INFO 'Deleting employee %', OLD.empno;
RAISE INFO '..Old salary: %', OLD.sal;
RETURN OLD;
END IF;
END;
$$ LANGUAGE 'plpgsql';
CREATE TRIGGER emp_sal_trig
  BEFORE DELETE OR INSERT OR UPDATE ON emp
  FOR EACH ROW EXECUTE PROCEDURE emp_sal_trig();
COMMIT;

```

11.2 Enhanced Compatibility Features

Advanced Server includes extended functionality that provides compatibility for syntax supported by Oracle applications. Detailed information about the compatibility features supported by Advanced Server is provided in the Database Compatibility for Oracle Developers Guides; the version-specific guides are available at:

<https://www.enterprisedb.com/edb-docs>

The following sections highlight some of the compatibility features supported by Advanced Server.

Enabling Compatibility Features

There are several ways to install Advanced Server that will allow you to take advantage of compatibility features:

- Use the `INITDBOPTS` variable (in the Advanced Server service configuration file) to specify `--redwood-like` before initializing your cluster.
- When invoking `initdb` to initialize your cluster, include the `--redwood-like` option.

For more information about the installation options supported by the Advanced Server installers, see the *EDB Postgres Advanced Server Installation Guide*, available from the EDB website at:

<https://www.enterprisedb.com/edb-docs>

Stored Procedural Language

Advanced Server supports a highly productive procedural language that allows you to write custom procedures, functions, triggers and packages. The procedural language:

- complements the SQL language and built-in packages.
- provides a seamless development and testing environment.
- allows you to create reusable code.

For information about using the Stored Procedural Language, see the *Database Compatibility for Stored Procedural Language Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Optimizer Hints

When you invoke a `DELETE`, `INSERT`, `SELECT`, or `UPDATE` command, the server generates a set of execution plans; after analyzing those execution plans, the server selects a plan that will (generally) return the result set in the least amount of time. The server's choice of plan is dependent upon several factors:

- The estimated execution cost of data handling operations.
- Parameter values assigned to parameters in the Query Tuning section of the `postgresql.conf` file.
- Column statistics that have been gathered by the `ANALYZE` command.

As a rule, the query planner will select the least expensive plan. You can use an optimizer hint to influence the server as it selects a query plan.

An optimizer hint is a directive (or multiple directives) embedded in a comment-like syntax that immediately follows a `DELETE`, `INSERT`, `SELECT` or `UPDATE` command. Keywords in the comment instruct the server to employ or avoid a specific plan when producing the result set. For information about using optimizer hints, see the *Database Compatibility for Oracle Developers Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Data Dictionary Views

Advanced Server includes a set of views that provide information about database objects in a manner compatible with the Oracle data dictionary views. For detailed information about the views available with Advanced Server, see the *Database Compatibility for Oracle Developers Catalog Views Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

dblink_ora

`dblink_ora` provides an OCI-based database link that allows you to `SELECT`, `INSERT`, `UPDATE` or `DELETE` data stored on an Oracle system from within Advanced Server. For detailed information about using `dblink_ora`, and the supported functions and procedures, see the *Database Compatibility for Oracle Developers Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Profile Management

Advanced Server supports compatible SQL syntax for profile management. Profile management commands allow a database superuser to create and manage named *profiles*. Each profile defines rules for password management that augment password and md5 authentication. The rules in a profile can:

- count failed login attempts
- lock an account due to excessive failed login attempts
- mark a password for expiration
- define a grace period after a password expiration
- define rules for password complexity
- define rules that limit password re-use

A profile is a named set of attributes that allow you to easily manage a group of roles that share comparable authentication requirements. If password requirements change, you can modify the profile to have the new requirements applied to each user that is associated with that profile.

After creating the profile, you can associate the profile with one or more users. When a user connects to the

server, the server enforces the profile that is associated with their login role. Profiles are shared by all databases within a cluster, but each cluster may have multiple profiles. A single user with access to multiple databases will use the same profile when connecting to each database within the cluster.

For information about using profile management commands, see the *Database Compatibility for Oracle Developers Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Built-In Packages

Advanced Server supports a number of built-in packages that provide compatibility with Oracle procedures and functions.

| Package Name | Description |
|----------------|---|
| DBMS_ALERT | The DBMS_ALERT package provides the capability to register for, send, and receive alerts. |
| DBMS_AQ | The DBMS_AQ package provides message queueing and processing for Advanced Server. |
| DBMS_AQADM | The DBMS_AQADM package provides supporting procedures for Advanced Queueing functionality. |
| DBMS_CRYPTO | The DBMS_CRYPTO package provides functions and procedures that allow you to encrypt or decrypt RAW, BLOB or CLOB data. You can also use DBMS_CRYPTO functions to generate cryptographically strong random values. |
| DBMS_JOB | The DBMS_JOB package provides for the creation, scheduling, and managing of jobs. |
| DBMS_LOB | The DBMS_LOB package provides the capability to operate on large objects. |
| DBMS_LOCK | Advanced Server provides support for the DBMS_LOCK.SLEEP procedure. |
| DBMS_MVIEW | Use procedures in the DBMS_MVIEW package to manage and refresh materialized views and their dependencies. |
| DBMS_OUTPUT | The DBMS_OUTPUT package provides the capability to send messages to a message buffer, or get messages from the message buffer. |
| DBMS_PIPE | The DBMS_PIPE package provides the capability to send messages through a pipe within or between sessions connected to the same database cluster. |
| DBMS_PROFILER | The DBMS_PROFILER package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a performance profiling session. |
| DBMS_RANDOM | The DBMS_RANDOM package provides a number of methods to generate random values. The procedures and functions available in the DBMS_RANDOM package are listed in the following table. |
| DBMS_REDACT | The DBMS_REDACT package enables the redacting or masking of data that is returned by a query. |
| DBMS_RLS | The DBMS_RLS package enables the implementation of Virtual Private Database on certain Advanced Server database objects. |
| DBMS_SCHEDULER | The DBMS_SCHEDULER package provides a way to create and manage jobs, programs and job schedules. |
| DBMS_SESSION | Advanced Server provides support for the DBMS_SESSION.SET_ROLE procedure. |
| DBMS_SQL | The DBMS_SQL package provides an application interface to the EnterpriseDB dynamic SQL functionality. |
| DBMS.Utility | The DBMS.Utility package provides various utility programs. |
| UTL_ENCODE | The UTL_ENCODE package provides a way to encode and decode data. |
| UTL_FILE | The UTL_FILE package provides the capability to read from, and write to files on the operating system's file system. |

| Package Name | Description |
|--------------|--|
| UTL_HTTP | The UTL_HTTP package provides a way to use the HTTP or HTTPS protocol to retrieve information found at an URL. |
| UTL_MAIL | The UTL_MAIL package provides the capability to manage e-mail. |
| UTL_RAW | The UTL_RAW package allows you to manipulate or retrieve the length of raw data types. |
| UTL_SMTP | The UTL_SMTP package provides the capability to send e-mails over the Simple Mail Transfer Protocol (SMTP). |
| UTL_URL | The UTL_URL package provides a way to escape illegal and reserved characters within an URL. |


For detailed information about the procedures and functions available within each package, see the *Database Compatibility for Oracle Developers Built-In Package Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Open Client Library

The Open Client Library provides application interoperability with the Oracle Call Interface – an application that was formerly “locked in” can now work with either an Advanced Server or an Oracle database with minimal to no changes to the application code. The EnterpriseDB implementation of the Open Client Library is written in C.

The following diagram compares the Open Client Library and Oracle Call Interface application stacks.



For detailed information about the functions supported by the Open Client Library, see the *EDB Postgres Advanced Server OCL Connector Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

Utilities

For detailed information about the compatible syntax supported by the utilities listed below, see the *Database Compatibility for Oracle Developer's Tools and Utilities Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

EDB*Plus

EDB*Plus is a utility program that provides a command line user interface to the Advanced Server that will be familiar to Oracle developers and users. EDB*Plus accepts SQL commands, SPL anonymous blocks, and EDB*Plus commands.

EDB*Plus allows you to:

- Query certain database objects
- Execute stored procedures
- Format output from SQL commands
- Execute batch scripts
- Execute OS commands
- Record output

For detailed information about EDB*Plus, see the *EDB Plus User's Guide* available at:

<https://www.enterprisedb.com/edb-docs/p/edbplus>

EDB*Loader

EDB*Loader is a high-performance bulk data loader that provides an interface compatible with Oracle databases for Advanced Server. The EDB*Loader command line utility loads data from an input source, typically a file, into one or more tables using a subset of the parameters offered by Oracle SQL*Loader.

EDB*Loader features include:

- Support for the Oracle SQL*Loader data loading methods - conventional path load, direct path load, and parallel direct path load
- Oracle SQL*Loader compatible syntax for control file directives
- Input data with delimiter-separated or fixed-width fields
- Bad file for collecting rejected records
- Loading of multiple target tables
- Discard file for collecting records that do not meet the selection criteria of any target table
- Data loading from standard input and remote loading

EDB Wrap

The EDB Wrap utility protects proprietary source code and programs (functions, stored procedures, triggers, and packages) from unauthorized scrutiny. The EDB Wrap program translates a file that contains SPL or PL/pgSQL source code (the plaintext) into a file that contains the same code in a form that is nearly impossible to read. Once you have the obfuscated form of the code, you can send that code to Advanced Server and it will store those programs in obfuscated form. While EDB Wrap does obscure code, table definitions are still exposed.

Everything you wrap is stored in obfuscated form. If you wrap an entire package, the package body source, as well as the prototypes contained in the package header and the functions and procedures contained in the package body are stored in obfuscated form.

Dynamic Runtime Instrumentation Tools Architecture (DRITA)

The Dynamic Runtime Instrumentation Tools Architecture (DRITA) allows a DBA to query catalog views to determine the *wait events* that affect the performance of individual sessions or the system as a whole. DRITA records the number of times each event occurs as well as the time spent waiting; you can use this information to diagnose performance problems. DRITA offers this functionality, while consuming minimal system resources.

DRITA compares *snapshots* to evaluate the performance of a system. A snapshot is a saved set of system performance data at a given point in time. Each snapshot is identified by a unique ID number; you can use snapshot ID numbers with DRITA reporting functions to return system performance statistics.

ECPGPlus

EnterpriseDB has enhanced ECPG (the PostgreSQL pre-compiler) to create ECPGPlus. ECPGPlus allows you to

include embedded SQL commands in C applications; when you use ECPGPlus to compile an application that contains embedded SQL commands, the SQL code is syntax-checked and translated into C.

ECPGPlus supports Pro*C syntax in C programs when connected to an Advanced Server database. ECPGPlus supports:

- Oracle Dynamic SQL – Method 4 (ODS-M4)
- Pro*C compatible anonymous blocks
- A `CALL` statement compatible with Oracle databases

For information about using ECPGPlus, see the *EDB Postgres Advanced Server ECPG Connector Guide*, available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

Table Partitioning

In a partitioned table, one logically large table is broken into smaller physical pieces. Partitioning can provide several benefits:

- Query performance can be improved dramatically in certain situations, particularly when most of the heavily accessed rows of the table are in a single partition or a small number of partitions. Partitioning allows you to omit the partition column from the front of an index, reducing index size and making it more likely that the heavily used parts of the index fits in memory.
- When a query or update accesses a large percentage of a single partition, performance may improve because the server will perform a sequential scan of the partition instead of using an index and random access reads scattered across the whole table.
- A bulk load (or unload) can be implemented by adding or removing partitions, if you plan that requirement into the partitioning design. `ALTER TABLE` is far faster than a bulk operation. It also entirely avoids the `VACUUM` overhead caused by a bulk `DELETE`.
- Seldom-used data can be migrated to less-expensive (or slower) storage media.

Table partitioning is worthwhile only when a table would otherwise be very large. The exact point at which a table will benefit from partitioning depends on the application; a good rule of thumb is that the size of the table should exceed the physical memory of the database server.

For information about database compatibility features supported by Advanced Server see the *Database Compatibility for Table Partitioning Guide*, available at:

<https://www.enterprisedb.com/edb-docs>

11.3 Database Administration

This chapter describes the features that aid in the management and administration of Advanced Server databases.

11.3.1 Configuration Parameters

This section describes the database server configuration parameters of Advanced Server. These parameters control various aspects of the database server's behavior and environment such as data file and log file locations, connection, authentication, and security settings, resource allocation and consumption, archiving and replication settings, error logging and statistics gathering, optimization and performance tuning, locale and formatting settings, and so on.

Configuration parameters that apply only to Advanced Server are noted in the [Summary of Configuration Parameters](#) section.

Additional information about configuration parameters can be found in the PostgreSQL Core Documentation available at:

<https://www.postgresql.org/docs/current/static/runtime-config.html>

11.3.1.1 Setting Configuration Parameters

This section provides an overview of how configuration parameters are specified and set.

Each configuration parameter is set using a `name/value` pair. Parameter names are case-insensitive. The parameter name is typically separated from its value by an optional equals sign (`=`).

The following is an example of some configuration parameter settings in the `postgresql.conf` file:

```
### This is a comment
log_connections = yes
log_destination = 'syslog'
search_path = "$user", public'
shared_buffers = 128MB
```

Parameter values are specified as one of five types:

- **Boolean.** Acceptable values can be written as `on`, `off`, `true`, `false`, `yes`, `no`, `1`, `0`, or any unambiguous prefix of these.
- **Integer.** Number without a fractional part.
- **Floating Point.** Number with an optional fractional part separated by a decimal point.
- **String.** Text value. Enclose in single quotes if the value is not a simple identifier or number (that is, the value contains special characters such as spaces or other punctuation marks).
- **Enum.** Specific set of string values. The allowed values can be found in the system view `pg_settings.enumvals`. Enum values are case-insensitive.

Some settings specify a memory or time value. Each of these has an implicit unit, which is kilobytes, blocks (typically 8 kilobytes), milliseconds, seconds, or minutes. Default units can be found by referencing the system view `pg_settings.unit`. A different unit can be specified explicitly.

Valid memory units are `KB` (kilobytes), `MB` (megabytes), and `GB` (gigabytes). Valid time units are `ms` (milliseconds), `s` (seconds), `min` (minutes), `h` (hours), and `d` (days). The multiplier for memory units is 1024.

The configuration parameter settings can be established in a number of different ways:

- There is a number of parameter settings that are established when the Advanced Server database product is built. These are read-only parameters, and their values cannot be changed. There are also a couple of parameters that are permanently set for each database when the database is created. These parameters are read-only as well and cannot be subsequently changed for the database.
- The initial settings for almost all configurable parameters across the entire database cluster are listed in the

configuration file, `postgresql.conf`. These settings are put into effect upon database server start or restart. Some of these initial parameter settings can be overridden as discussed in the following bullet points. All configuration parameters have built-in default settings that are in effect if not explicitly overridden.

- Configuration parameters in the `postgresql.conf` file are overridden when the same parameters are included in the `postgresql.auto.conf` file. The `ALTER SYSTEM` command is used to manage the configuration parameters in the `postgresql.auto.conf` file.
- Parameter settings can be modified in the configuration file while the database server is running. If the configuration file is then reloaded (meaning a SIGHUP signal is issued), for certain parameter types, the changed parameters settings immediately take effect. For some of these parameter types, the new settings are available in a currently running session immediately after the reload. For other of these parameter types, a new session must be started to use the new settings. And yet for other parameter types, modified settings do not take effect until the database server is stopped and restarted. See the following section of the PostgreSQL Core Documentation for information on how to reload the configuration file:

<https://www.postgresql.org/docs/current/config-setting.html>

- The SQL commands `ALTER DATABASE`, `ALTER ROLE`, or `ALTER ROLE IN DATABASE` can be used to modify certain parameter settings. The modified parameter settings take effect for new sessions after the command is executed. `ALTER DATABASE` affects new sessions connecting to the specified database. `ALTER ROLE` affects new sessions started by the specified role. `ALTER ROLE IN DATABASE` affects new sessions started by the specified role connecting to the specified database. Parameter settings established by these SQL commands remain in effect indefinitely, across database server restarts, overriding settings established by the methods discussed in the second and third bullet points. Parameter settings established using the `ALTER DATABASE`, `ALTER ROLE`, or `ALTER ROLE IN DATABASE` commands can only be changed by: a) re-issuing these commands with a different parameter value, or
- b) issuing these commands using either of the `SET parameter TO DEFAULT` clause or the `RESET parameter` clause. These clauses change the parameter back to using the setting established by the methods set forth in the prior bullet points. See the “SQL Commands” section of Chapter VI “Reference” in the *PostgreSQL Core Documentation* for the exact syntax of these SQL commands:

<https://www.postgresql.org/docs/current/sql-commands.html>

- Changes can be made for certain parameter settings for the duration of individual sessions using the `PGOPTIONS` environment variable or by using the `SET` command within the `EDB-PSQL` or `PSQL` command line terminal programs. Parameter settings made in this manner override settings established using any of the methods described by the second, third, and fourth bullet points, but only for the duration of the session.

11.3.1.2 Summary of Configuration Parameters

This section contains a summary table listing all Advanced Server configuration parameters along with a number of key attributes of the parameters.

These attributes are described by the following columns of the summary table:

- **Parameter.** Configuration parameter name.
- **Scope of Effect.** Scope of effect of the configuration parameter setting.
 - `Cluster` – Setting affects the entire database cluster (that is, all databases managed by the database server instance).
 - `Database` – Setting can vary by database and is established when the database is created. Applies to a small number of parameters related to locale settings.
 - `Session` – Setting can vary down to the granularity of individual sessions.

In other words, different settings can be made for the following entities whereby the latter settings in this list override prior ones: a) the entire database cluster, b) specific databases in the database cluster, c) specific roles, d) specific roles when connected to specific databases, e) a specific session.

- **When Takes Effect.** When a changed parameter setting takes effect.

- **Preset** – Established when the Advanced Server product is built or a particular database is created. This is a read-only parameter and cannot be changed.
- **Restart** – Database server must be restarted.
- **Reload** – Configuration file must be reloaded (or the database server can be restarted).
- **Immediate** – Immediately effective in a session if the **PGOPTIONS** environment variable or the **SET** command is used to change the setting in the current session. Effective in new sessions if **ALTER DATABASE**, **ALTER ROLE**, or **ALTER ROLE IN DATABASE** commands are used to change the setting.

- **Authorized User.** Type of operating system account or database role that must be used to put the parameter setting into effect.

- **EPAS service account** – EDB Postgres Advanced Server service account (**enterprisedb** for an installation compatible with Oracle databases, **postgres** for a PostgreSQL compatible mode installation).
- **Superuser** – Database role with superuser privileges.
- **User** – Any database role with permissions on the affected database object (the database or role to be altered with the **ALTER** command).
- **n/a** – Parameter setting cannot be changed by any user.

- **Description.** Brief description of the configuration parameter.

- **EPAS Only.** ‘X’ – Configuration parameter is applicable to EDB Postgres Advanced Server only. No entry in this column indicates the configuration parameter applies to PostgreSQL as well.

Note

There are a number of parameters that should never be altered. These are designated as “**Note: For internal use only**” in the Description column.

The following table shows the summary of configuration parameters:

| Parameter | Scope of Effect | When Takes Effect | Authorized User | Description | EPAS Only |
|--------------------------------|-----------------|-------------------|----------------------|--|-----------|
| custom_variable_classes | Cluster | Reload | EPAS service account | Deprecated in Advanced Server 9.2. | X |
| db_dialect | Session | Immediate | User | Sets the precedence of built-in namespaces. | X |
| dbms_alert.max_alerts | Cluster | Restart | EPAS service account | Sets maximum number of alerts. | X |
| dbms_pipe.total_message_buffer | Cluster | Restart | EPAS service account | Specifies the total size of the buffer used for the DBMS_PIPE package. | X |
| default_heap_fillfactor | Session | Immediate | User | Create new tables with this heap fillfactor by default. | X |
| default_with_rowids | Session | Immediate | User | Create new tables with ROWID support (ROWIDs with indexes) by default. | X |

| | | | | | |
|----------------------------|---------|-----------|----------------------|--|---|
| edb_audit | Cluster | Reload | EPAS service account | Enable EDB Auditing to create audit reports in XML or CSV format. | X |
| edb_audit_connect | Cluster | Reload | EPAS service account | Audits each successful connection. | X |
| edb_audit_destination | Cluster | Reload | EPAS service account | Sets <code>edb_audit_directory</code> or <code>syslog</code> as the destination directory for audit files. The <code>syslog</code> setting is only valid for a Linux system. | X |
| edb_audit_directory | Cluster | Reload | EPAS service account | Sets the destination directory for audit files. | X |
| edb_audit_disconnect | Cluster | Reload | EPAS service account | Audits end of a session. | X |
| edb_audit_filename | Cluster | Reload | EPAS service account | Sets the file name pattern for audit files. | X |
| edb_audit_rotation_day | Cluster | Reload | EPAS service account | Automatic rotation of log files based on day of week. | X |
| edb_audit_rotation_seconds | Cluster | Reload | EPAS service account | Automatic log file rotation will occur after N seconds. | X |
| edb_audit_rotation_size | Cluster | Reload | EPAS service account | Automatic log file rotation will occur after N Megabytes. | X |
| edb_audit_statement | Cluster | Reload | EPAS service account | Sets the type of statements to audit. | X |
| edb_audit_tag | Session | Immediate | User | Specify a tag to be included in the audit log. | X |
| edb_connectby_order | Session | Immediate | User | Sort results of CONNECT BY queries with no ORDER BY to depth-first order. Note: For internal use only. | X |
| edb_data_redaction | Session | Immediate | User | Enable data redaction. | X |
| edb_dynatune | Cluster | Restart | EPAS service account | Sets the edb utilization percentage. | X |
| edb_dynatune_profile | Cluster | Restart | EPAS service account | Sets the workload profile for dynatune. | X |
| edb_enable_pruning | Session | Immediate | User | Enables the planner to early-prune partitioned tables. | X |
| edb_log_every_bulk_value | Session | Immediate | Superuser | Sets the statements logged for bulk processing. | X |

| | | | | | |
|---|---------|-----------|----------------------|--|---|
| edb_max_resource_groups | Cluster | Restart | EPAS service account | Specifies the maximum number of resource groups for simultaneous use. | X |
| edb_max_spins_per_delay | Cluster | Restart | EPAS service account | Specifies the number of times a session will spin while waiting for a lock. | X |
| edb_redwood_date | Session | Immediate | User | Determines whether DATE should behave like a TIMESTAMP or not. | X |
| edb_redwood_greatest_least | Session | Immediate | User | Determines how GREATEST and LEAST functions should handle NULL parameters. | X |
| edb_redwood_raw_names | Session | Immediate | User | Return the unmodified name stored in the PostgreSQL system catalogs from Redwood interfaces. | X |
| edb_redwood_strings | Session | Immediate | User | Treat NULL as an empty string when concatenated with a text value. | X |
| edb_resource_group | Session | Immediate | User | Specifies the resource group to be used by the current process. | X |
| edb_sql_protect.enabled | Cluster | Reload | EPAS service account | Defines whether SQL/Protect should track queries or not. | X |
| edb_sql_protect.level | Cluster | Reload | EPAS service account | Defines the behavior of SQL/Protect when an event is found. | X |
| edb_sql_protect.max_protected_relations | Cluster | Restart | EPAS service account | Sets the maximum number of relations protected by SQL/Protect per role. | X |
| edb_sql_protect.max_protected_roles | Cluster | Restart | EPAS service account | Sets the maximum number of roles protected by SQL/Protect. | X |
| edb_sql_protect.max_queries_to_save | Cluster | Restart | EPAS service account | Sets the maximum number of offending queries to save by SQL/Protect. | X |
| edb_stmt_level_tx | Session | Immediate | User | Allows continuing on errors instead of requiring a transaction abort. | X |
| edb_wait_states.directory | Cluster | Restart | EPAS service account | The EDB wait states logs are stored in this directory. | X |
| edb_wait_states.retention_period | Cluster | Reload | EPAS service account | EDB wait state log files will be automatically deleted after retention period. | X |

| | | | | | |
|-----------------------------------|----------|-----------|----------------------|--|---|
| edb_wait_states.sampling_interval | Cluster | Reload | EPAS service account | The interval between two EDB wait state sampling cycles. | X |
| edbldr.empty_csv_field | Session | Immediate | Superuser | Specifies how EDB Loader handles empty strings. | X |
| enable_hints | Session | Immediate | User | Enable optimizer hints in SQL statements. | X |
| icu_short_form | Database | Preset | n/a | Shows the ICU collation order configuration. | X |
| index_advisor.enabled | Session | Immediate | User | Enable Index Advisor plugin. | X |
| nls_length_semantics | Session | Immediate | Superuser | Sets the semantics to use for char, varchar, varchar2 and long columns. | X |
| odbc_lib_path | Cluster | Restart | EPAS service account | Sets the path for ODBC library. | X |
| optimizer_mode | Session | Immediate | User | Default optimizer mode. | X |
| oracle_home | Cluster | Restart | EPAS service account | Sets the path for the Oracle home directory. | X |
| pg_prewarm.autoprewarm | Cluster | Restart | EPAS service account | Enables the autoprewarm background worker. | X |
| pg_prewarm.autoprewarm_interval | Cluster | Reload | EPAS service account | Sets the minimum number of seconds after which autoprewarm dumps shared buffers. | X |
| qreplace_function | Session | Immediate | Superuser | The function to be used by Query Replace feature. Note: For internal use only. | X |
| query_rewrite_enabled | Session | Immediate | User | Child table scans will be skipped if their constraints guarantee that no rows match the query. | X |
| query_rewrite_integrity | Session | Immediate | Superuser | Sets the degree to which query rewriting must be enforced. | X |
| timed_statistics | Session | Immediate | User | Enables the recording of timed statistics. | X |
| trace_hints | Session | Immediate | User | Emit debug info about hints being honored. | X |
| util_encode.uudecode_redwood | Session | Immediate | User | Allows decoding of Oracle-created uuencoded data. | X |
| util_file.umask | Session | Immediate | User | Umask used for files created through the UTL_FILE package. | X |

11.3.1.3 Configuration Parameters by Functionality

This section provides more detail for certain groups of configuration parameters.

The section heading for each parameter is followed by a list of attributes:

- **Parameter Type.** Type of values the parameter can accept. See [Setting Configuration Parameters](#) section for a discussion of parameter type values.
- **Default Value.** Default setting if a value is not explicitly set in the configuration file.
- **Range.** Permitted range of values.
- **Minimum Scope of Effect.** Smallest scope for which a distinct setting can be made. Generally, the minimal scope of a distinct setting is either the entire `cluster` (the setting is the same for all databases and sessions thereof, in the cluster), or `per session` (the setting may vary by role, database, or individual session). (This attribute has the same meaning as the `Scope of Effect` column in the table of [Summary of Configuration Parameters](#) section).
- **When Value Changes Take Effect.** Least invasive action required to activate a change to a parameter's value. All parameter setting changes made in the configuration file can be put into effect with a restart of the database server; however certain parameters require a database server `restart`. Some parameter setting changes can be put into effect with a `reload` of the configuration file without stopping the database server. Finally, other parameter setting changes can be put into effect with some client side action whose result is `immediate`. (This attribute has the same meaning as the `When Takes Effect` column in the table of [Summary of Configuration Parameters](#) section).
- **Required Authorization to Activate.** The type of user authorization to activate a change to a parameter's setting. If a database server restart or a configuration file reload is required, then the user must be a EPAS service account (`enterprisedb` or `postgres` depending upon the Advanced Server compatibility installation mode). (This attribute has the same meaning as the `Authorized User` column in the table of [Summary of Configuration Parameters](#) section).

11.3.1.3.1 Top Performance Related Parameters

This section discusses the configuration parameters that have the most immediate impact on performance.

11.3.1.3.1.1 `shared_buffers`

Parameter Type: Integer

Default Value: 32MB

Range: 128kB to system dependent

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart**Required Authorization to Activate:** EPAS service account

Sets the amount of memory the database server uses for shared memory buffers. The default is typically 32 megabytes (32MB), but might be less if your kernel settings will not support it (as determined during `initdb`). This setting must be at least 128 kilobytes. (Non-default values of `BLCKSZ` change the minimum.) However, settings significantly higher than the minimum are usually needed for good performance.

If you have a dedicated database server with 1GB or more of RAM, a reasonable starting value for `shared_buffers` is 25% of the memory in your system. There are some workloads where even large settings for `shared_buffers` are effective, but because Advanced Server also relies on the operating system cache, it is unlikely that an allocation of more than 40% of RAM to `shared_buffers` will work better than a smaller amount.

On systems with less than 1GB of RAM, a smaller percentage of RAM is appropriate, so as to leave adequate space for the operating system (15% of memory is more typical in these situations). Also, on Windows, large values for `shared_buffers` aren't as effective. You may find better results keeping the setting relatively low and using the operating system cache more instead. The useful range for `shared_buffers` on Windows systems is generally from 64MB to 512MB.

Increasing this parameter might cause Advanced Server to request more System V shared memory than your operating system's default configuration allows. See the section [Shared Memory and Semaphores](#) in the *PostgreSQL Core Documentation* for information on how to adjust those parameters, if necessary.

11.3.1.3.1.2 work_mem

Parameter Type: Integer**Default Value:** 1MB**Range:** 64kB to 2097151kB**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Specifies the amount of memory to be used by internal sort operations and hash tables before writing to temporary disk files. The value defaults to one megabyte (1MB). Note that for a complex query, several sort or hash operations might be running in parallel; each operation will be allowed to use as much memory as this value specifies before it starts to write data into temporary files. Also, several running sessions could be doing such operations concurrently. Therefore, the total memory used could be many times the value of `work_mem`; it is necessary to keep this fact in mind when choosing the value. Sort operations are used for `ORDER BY`, `DISTINCT`, and merge joins. Hash tables are used in hash joins, hash-based aggregation, and hash-based processing of `IN` subqueries.

Reasonable values are typically between 4MB and 64MB, depending on the size of your machine, how many concurrent connections you expect (determined by `max_connections`), and the complexity of your queries.

11.3.1.3.1.3 maintenance_work_mem

Parameter Type: Integer**Default Value:** 16MB**Range:** 1024kB to 2097151kB**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Specifies the maximum amount of memory to be used by maintenance operations, such as `VACUUM`, `CREATE INDEX`, and `ALTER TABLE ADD FOREIGN KEY`. It defaults to 16 megabytes (16MB). Since only one of these operations can be executed at a time by a database session, and an installation normally doesn't have many of them running concurrently, it's safe to set this value significantly larger than `work_mem`. Larger settings might improve performance for vacuuming and for restoring database dumps.

Note that when autovacuum runs, up to `autovacuum_max_workers` times this memory may be allocated, so be careful not to set the default value too high.

A good rule of thumb is to set this to about 5% of system memory, but not more than about 512MB. Larger values won't necessarily improve performance.

11.3.1.3.1.4 wal_buffers

Parameter Type: Integer**Default Value:** 64kB**Range:** 32kB to system dependent**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Restart**Required Authorization to Activate:** EPAS service account

The amount of memory used in shared memory for WAL data. The default is 64 kilobytes (64kB). The setting need only be large enough to hold the amount of WAL data generated by one typical transaction, since the data is written out to disk at every transaction commit.

Increasing this parameter might cause Advanced Server to request more System V shared memory than your operating system's default configuration allows. See the section "Shared Memory and Semaphores" in the *PostgreSQL Core Documentation* for information on how to adjust those parameters, if necessary.

Although even this very small setting does not always cause a problem, there are situations where it can result in extra `fsync` calls, and degrade overall system throughput. Increasing this value to 1MB or so can alleviate this problem. On very busy systems, an even higher value may be needed, up to a maximum of about 16MB. Like `shared_buffers`, this parameter increases Advanced Server's initial shared memory allocation, so if increasing it causes an Advanced Server start failure, you will need to increase the operating system limit.

11.3.1.3.1.5 checkpoint_segments

Now deprecated; this parameter is not supported by server versions 9.5 or later.

11.3.1.3.1.6 checkpoint_completion_target

Parameter Type: Floating point

Default Value: 0.5

Range: 0 to 1

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the target of checkpoint completion as a fraction of total time between checkpoints. This spreads out the checkpoint writes while the system starts working towards the next checkpoint.

The default of 0.5 means aim to finish the checkpoint writes when 50% of the next checkpoint is ready. A value of 0.9 means aim to finish the checkpoint writes when 90% of the next checkpoint is done, thus throttling the checkpoint writes over a larger amount of time and avoiding spikes of performance bottlenecks.

11.3.1.3.1.7 checkpoint_timeout

Parameter Type: Integer

Default Value: 5min

Range: 30s to 3600s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Maximum time between automatic WAL checkpoints, in seconds. The default is five minutes (5min). Increasing this parameter can increase the amount of time needed for crash recovery.

Increasing `checkpoint_timeout` to a larger value, such as 15 minutes, can reduce the I/O load on your system, especially when using large values for `shared_buffers`.

The downside of making the aforementioned adjustments to the checkpoint parameters is that your system will use a modest amount of additional disk space, and will take longer to recover in the event of a crash. However, for most users, this is a small price to pay for a significant performance improvement.

11.3.1.3.1.8 max_wal_size

Parameter Type: Integer

Default Value: 1 GB

Range: 2 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

`max_wal_size` specifies the maximum size that the WAL will reach between automatic WAL checkpoints. This is a soft limit; WAL size can exceed `max_wal_size` under special circumstances (when under a heavy load, a failing archive_command, or a high `wal_keep_segments` setting).

Increasing this parameter can increase the amount of time needed for crash recovery. This parameter can only be set in the `postgresql.conf` file or on the server command line.

11.3.1.3.1.9 min_wal_size

Parameter Type: Integer

Default Value: 80 MB

Range: 2 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

If WAL disk usage stays below the value specified by `min_wal_size`, old WAL files are recycled for future use at a checkpoint, rather than removed. This ensures that enough WAL space is reserved to handle spikes in WAL usage (like when running large batch jobs). This parameter can only be set in the `postgresql.conf` file or on the server command line.

11.3.1.3.1.10 bgwriter_delay

Parameter Type: Integer

Default Value: 200ms

Range: 10ms to 10000ms**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

Specifies the delay between activity rounds for the background writer. In each round the writer issues writes for some number of dirty buffers (controllable by the `bgwriter_lru_maxpages` and `bgwriter_lru_multiplier` parameters). It then sleeps for `bgwriter_delay` milliseconds, and repeats.

The default value is 200 milliseconds (200ms). Note that on many systems, the effective resolution of sleep delays is 10 milliseconds; setting `bgwriter_delay` to a value that is not a multiple of 10 might have the same results as setting it to the next higher multiple of 10.

Typically, when tuning `bgwriter_delay`, it should be reduced from its default value. This parameter is rarely increased, except perhaps to save on power consumption on a system with very low utilization.

11.3.1.3.1.11 seq_page_cost

Parameter Type: Floating point**Default Value:** 1**Range:** 0 to 1.79769e+308**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Sets the planner's estimate of the cost of a disk page fetch that is part of a series of sequential fetches. The default is 1.0. This value can be overridden for a particular tablespace by setting the tablespace parameter of the same name. (Refer to the [ALTER TABLESPACE](#) command in the *PostgreSQL Core Documentation*.)

The default value assumes very little caching, so it's frequently a good idea to reduce it. Even if your database is significantly larger than physical memory, you might want to try setting this parameter to less than 1 (rather than its default value of 1) to see whether you get better query plans that way. If your database fits entirely within memory, you can lower this value much more, perhaps to 0.1.

11.3.1.3.1.12 random_page_cost

Parameter Type: Floating point**Default Value:** 4**Range:** 0 to 1.79769e+308

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the planner's estimate of the cost of a non-sequentially-fetched disk page. The default is 4.0. This value can be overridden for a particular tablespace by setting the tablespace parameter of the same name. (Refer to the [ALTER TABLESPACE](#) command in the *PostgreSQL Core Documentation*.)

Reducing this value relative to `seq_page_cost` will cause the system to prefer index scans; raising it will make index scans look relatively more expensive. You can raise or lower both values together to change the importance of disk I/O costs relative to CPU costs, which are described by the `cpu_tuple_cost` and `cpu_index_tuple_cost` parameters.

The default value assumes very little caching, so it's frequently a good idea to reduce it. Even if your database is significantly larger than physical memory, you might want to try setting this parameter to 2 (rather than its default of 4) to see whether you get better query plans that way. If your database fits entirely within memory, you can lower this value much more, perhaps to 0.1.

Although the system will let you do so, never set `random_page_cost` less than `seq_page_cost`. However, setting them equal (or very close to equal) makes sense if the database fits mostly or entirely within memory, since in that case there is no penalty for touching pages out of sequence. Also, in a heavily-cached database you should lower both values relative to the CPU parameters, since the cost of fetching a page already in RAM is much smaller than it would normally be.

11.3.1.3.1.13 effective_cache_size

Parameter Type: Integer

Default Value: 128MB

Range: 8kB to 17179869176kB

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the planner's assumption about the effective size of the disk cache that is available to a single query. This is factored into estimates of the cost of using an index; a higher value makes it more likely index scans will be used, a lower value makes it more likely sequential scans will be used. When setting this parameter you should consider both Advanced Server's shared buffers and the portion of the kernel's disk cache that will be used for Advanced Server data files. Also, take into account the expected number of concurrent queries on different tables, since they will have to share the available space. This parameter has no effect on the size of shared memory allocated by Advanced Server, nor does it reserve kernel disk cache; it is used only for estimation purposes. The default is 128 megabytes ([128MB](#)).

If this parameter is set too low, the planner may decide not to use an index even when it would be beneficial to do so. Setting `effective_cache_size` to 50% of physical memory is a normal, conservative setting. A more aggressive setting would be approximately 75% of physical memory.

11.3.1.3.1.14 synchronous_commit

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Specifies whether transaction commit will wait for WAL records to be written to disk before the command returns a [success](#) indication to the client. The default, and safe, setting is on. When off, there can be a delay between when success is reported to the client and when the transaction is really guaranteed to be safe against a server crash. (The maximum delay is three times [wal_writer_delay](#).)

Unlike [fsync](#), setting this parameter to off does not create any risk of database inconsistency: an operating system or database crash might result in some recent allegedly-committed transactions being lost, but the database state will be just the same as if those transactions had been aborted cleanly.

So, turning [synchronous_commit](#) off can be a useful alternative when performance is more important than exact certainty about the durability of a transaction. See the section [Asynchronous Commit](#) in the *PostgreSQL Core Documentation* for information.

This parameter can be changed at any time; the behavior for any one transaction is determined by the setting in effect when it commits. It is therefore possible, and useful, to have some transactions commit synchronously and others asynchronously. For example, to make a single multistatement transaction commit asynchronously when the default is the opposite, issue [SET LOCAL synchronous_commit TO OFF](#) within the transaction.

11.3.1.3.1.15 edb_max_spins_per_delay

Parameter Type: Integer

Default Value: 1000

Range: 10 to 1000

Minimum Scope of Effect: Per cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Use [edb_max_spins_per_delay](#) to specify the maximum number of times that a session will spin while waiting for a spin-lock. If a lock is not acquired, the session will sleep. If you do not specify an alternative value for [edb_max_spins_per_delay](#), the server will enforce the default value of [1000](#).

This may be useful for systems that use [NUMA](#) (non-uniform memory access) architecture.

11.3.1.3.1.16 pg_prewarm.autoprewarm

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

This parameter controls whether or not the database server should run `autoprewarm`, which is a background worker process that automatically dumps shared buffers to disk before a shutdown. It then `prewams` the shared buffers the next time the server is started, meaning it loads blocks from the disk back into the buffer pool.

The advantage is that it shortens the warm up times after the server has been restarted by loading the data that has been dumped to disk before shutdown.

If `pg_prewarm.autoprewarm` is set to on, the `autoprewarm` worker is enabled. If the parameter is set to off, `autoprewarm` is disabled. The parameter is on by default.

Before `autoprewarm` can be used, you must add `$libdir/pg_prewarm` to the libraries listed in the `shared_preload_libraries` configuration parameter of the `postgresql.conf` file as shown by the following example:

```
shared_preload_libraries =
'$libdir/dbms_pipe,$libdir/edb_gen,$libdir/dbms_aq,$libdir/pg_prewarm'
```

After modifying the `shared_preload_libraries` parameter, restart the database server after which the `autoprewarm` background worker is launched immediately after the server has reached a consistent state.

The `autoprewarm` process will start loading blocks that were previously recorded in `$PGDATA/autoprewarm` blocks until there is no free buffer space left in the buffer pool. In this manner, any new blocks that were loaded either by the recovery process or by the querying clients, are not replaced.

Once the `autoprewarm` process has finished loading buffers from disk, it will periodically dump shared buffers to disk at the interval specified by the `pg_prewarm.autoprewarm_interval` parameter. See the `pg_prewarm.autoprewarm_interval` for information on the `autoprewarm` background worker. Upon the next server restart, the `autoprewarm` process will prewarm shared buffers with the blocks that were last dumped to disk.

11.3.1.3.1.17 pg_prewarm.autoprewarm_interval

Parameter Type: Integer

Default Value: 300s

Range: 0s to 2147483s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

This is the minimum number of seconds after which the `autoprewarm` background worker dumps shared buffers to disk. The default is 300 seconds. If set to 0, shared buffers are not dumped at regular intervals, but only when the server is shut down.

See the [pg_prewarm.autoprewarm](#) for information on the `autoprewarm` background worker.

11.3.1.3.2 Resource Usage / Memory

The configuration parameters in this section control resource usage pertaining to memory.

`edb_dynatune`

Parameter Type: Integer

Default Value: 0

Range: 0 to 100

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Determines how much of the host system's resources are to be used by the database server based upon the host machine's total available resources and the intended usage of the host machine.

When Advanced Server is initially installed, the `edb_dynatune` parameter is set in accordance with the selected usage of the host machine on which it was installed (i.e., development machine, mixed use machine, or dedicated server). For most purposes, there is no need for the database administrator to adjust the various configuration parameters in the `postgresql.conf` file in order to improve performance.

The `edb_dynatune` parameter can be set to any integer value between 0 and 100, inclusive. A value of 0, turns off the dynamic tuning feature thereby leaving the database server resource usage totally under the control of the other configuration parameters in the `postgresql.conf` file.

A low non-zero, value (e.g., 1 - 33) dedicates the least amount of the host machine's resources to the database server. This setting would be used for a development machine where many other applications are being used.

A value in the range of 34 - 66 dedicates a moderate amount of resources to the database server. This setting might be used for a dedicated application server that may have a fixed number of other applications running on the same machine as Advanced Server.

The highest values (e.g., 67 - 100) dedicate most of the server's resources to the database server. This setting would be used for a host machine that is totally dedicated to running Advanced Server.

Once a value of `edb_dynatune` is selected, database server performance can be further fine-tuned by adjusting the other configuration parameters in the `postgresql.conf` file. Any adjusted setting overrides the corresponding value chosen by `edb_dynatune`. You can change the value of a parameter by un-commenting the configuration parameter, specifying the desired value, and restarting the database server.

`edb_dynatune_profile`

Parameter Type: Enum

Default Value: oltp

Range: {oltp | reporting | mixed}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

This parameter is used to control tuning aspects based upon the expected workload profile on the database server.

The following are the possible values:

- `oltp`. Recommended when the database server is processing heavy online transaction processing workloads.
 - `reporting`. Recommended for database servers used for heavy data reporting.
 - `mixed`. Recommended for servers that provide a mix of transaction processing and data reporting.
-

11.3.1.3.3 Resource Usage / EDB Resource Manager

The configuration parameters in this section control resource usage through EDB Resource Manager.

`edb_max_resource_groups`

Parameter Type: Integer

Default Value: 16

Range: 0 to 65536

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

This parameter controls the maximum number of resource groups that can be used simultaneously by EDB Resource Manager. More resource groups can be created than the value specified by `edb_max_resource_groups`, however, the number of resource groups in active use by processes in these groups cannot exceed this value.

Parameter `edb_max_resource_groups` should be set comfortably larger than the number of groups you expect to maintain so as not to run out.

`edb_resource_group`

Parameter Type: String

Default Value: none

Range: n/a

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Set the `edb_resource_group` parameter to the name of the resource group to which the current session is to be controlled by EDB Resource Manager according to the group's resource type settings.

If the parameter is not set, then the current session does not utilize EDB Resource Manager.

11.3.1.3.4 Query Tuning

This section describes the configuration parameters used for optimizer hints.

`enable_hints`

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Optimizer hints embedded in SQL commands are utilized when `enable_hints` is on. Optimizer hints are ignored when this parameter is off.

11.3.1.3.5 Query Tuning / Planner Method Configuration

This section describes the configuration parameters used for planner method configuration.

`edb_enable_pruning`

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When set to `TRUE`, `edb_enable_pruning` allows the query planner to early-prune partitioned tables. **Early-pruning** means that the query planner can `prune` (i.e., ignore) partitions that would not be searched in a query **before** generating query plans. This helps improve performance time as it eliminates the generation of query plans of partitions that would not be searched.

Conversely, **late-pruning** means that the query planner prunes partitions **after** generating query plans for each partition. (The `constraint_exclusion` configuration parameter controls late-pruning.)

The ability to early-prune depends upon the nature of the query in the `WHERE` clause. Early-pruning can be utilized in only simple queries with constraints of the type `WHERE column = literal` (e.g., `WHERE deptno = 10`).

Early-pruning is not used for more complex queries such as `WHERE column = expression` (e.g., `WHERE deptno = 10 + 5`).

11.3.1.3.6 Reporting and Logging / What to Log

The configuration parameters in this section control reporting and logging.

trace_hints

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Use with the optimizer hints feature to provide more detailed information regarding whether or not a hint was used by the planner. Set the `client_min_messages` and `trace_hints` configuration parameters as follows:

```
SET client_min_messages TO info;
SET trace_hints TO true;
```

The `SELECT` command with the `NO_INDEX` hint shown below illustrates the additional information produced when the aforementioned configuration parameters are set.

```
EXPLAIN SELECT /*+ NO_INDEX(accounts accounts_pkey) */ * FROM accounts
WHERE aid = 100;
```

INFO: [HINTS] Index Scan of [accounts].[accounts_pkey] rejected because of NO_INDEX hint.

INFO: [HINTS] Bitmap Heap Scan of [accounts].[accounts_pkey] rejected because of NO_INDEX hint.

QUERY PLAN

Seq Scan on accounts (cost=0.00..14461.10 rows=1 width=97)

 Filter: (aid = 100)

 (2 rows)

`edb_log_every_bulk_value`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Superuser

Bulk processing logs the resulting statements into both the Advanced Server log file and the EDB Audit log file. However, logging each and every statement in bulk processing is costly. This can be controlled by the `edb_log_every_bulk_value` configuration parameter. When set to `true`, each and every statement in bulk processing is logged. During bulk execution, when `edb_log_every_bulk_value` is set to `false`, a log message is recorded once per bulk processing along with the number of rows processed. In addition, the duration is emitted once per bulk processing. Default is set to `false`.

11.3.1.3.7 Auditing Settings

This section describes configuration parameters used by the Advanced Server database auditing feature.

11.3.1.3.7.1 `edb_audit`

Parameter Type: Enum

Default Value: none

Range: {none | csv | xml}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Enables or disables database auditing. The values `xml` or `csv` will enable database auditing. These values represent the file format in which auditing information will be captured. `none` will disable database auditing and is also the default.

11.3.1.3.7.2 `edb_audit_directory`

Parameter Type: String

Default Value: `edb_audit`

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the directory where the audit log files will be created. The path of the directory can be absolute or relative to the Advanced Server `data` directory.

11.3.1.3.7.3 `edb_audit_filename`

Parameter Type: String

Default Value: `audit-%Y%m%d_%H%M%S`

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the file name of the audit file where the auditing information will be stored. The default file name will be `audit-%Y%m%d_%H%M%S`. The escape sequences, `%Y`, `%m` etc., will be replaced by the appropriate current values according to the system date and time.

11.3.1.3.7.4 `edb_audit_rotation_day`

Parameter Type: String

Default Value: every

Range: {none | every | sun | mon | tue | wed | thu | fri | sat} ...

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the day of the week on which to rotate the audit files. Valid values are `sun`, `mon`, `tue`, `wed`, `thu`, `fri`, `sat`, `every`, and `none`. To disable rotation, set the value to `none`. To rotate the file every day, set the `edb_audit_rotation_day` value to `every`. To rotate the file on a specific day of the week, set the value to the desired day of the week.

11.3.1.3.7.5 `edb_audit_rotation_size`

Parameter Type: Integer

Default Value: 0MB

Range: 0MB to 5000MB

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies a file size threshold in megabytes when file rotation will be forced to occur. The default value is 0MB. If the parameter is commented out or set to 0, rotation of the file on a size basis will not occur.

11.3.1.3.7.6 `edb_audit_rotation_seconds`

Parameter Type: Integer

Default Value: 0s

Range: 0s to 2147483647s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the rotation time in seconds when a new log file should be created. To disable this feature, set this parameter to 0.

11.3.1.3.7.7 `edb_audit_connect`

Parameter Type: Enum

Default Value: failed

Range: {none | failed | all}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Enables auditing of database connection attempts by users. To disable auditing of all connection attempts, set `edb_audit_connect` to `none`. To audit all failed connection attempts, set the value to `failed`. To audit all connection attempts, set the value to `all`.

11.3.1.3.7.8 `edb_audit_disconnect`

Parameter Type: Enum

Default Value: none

Range: {none | all}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Enables auditing of database disconnections by connected users. To enable auditing of disconnections, set the value to `all`. To disable, set the value to `none`.

11.3.1.3.7.9 `edb_audit_statement`

Parameter Type: String

Default Value: ddl, error**Range:** {none| ddl | dml | insert | update | delete | truncate | select | error | create | drop | alter | grant | revoke | rollback | all} ...**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

This configuration parameter is used to specify auditing of different categories of SQL statements as well as those statements related to specific SQL commands. To log errors, set the parameter value to `error`. To audit all DDL statements such as `CREATE TABLE`, `ALTER TABLE`, etc., set the parameter value to `ddl`. To audit specific types of DDL statements, the parameter values can include those specific SQL commands (`create`, `drop`, or `alter`). In addition, the object type may be specified following the command such as `create table`, `create view`, `drop role`, etc. All modification statements such as `INSERT`, `UPDATE`, `DELETE` or `TRUNCATE` can be audited by setting `edb_audit_statement` to `dml`. To audit specific types of DML statements, the parameter values can include the specific SQL commands, `insert`, `update`, `delete`, or `truncate`. Include parameter values `select`, `grant`, `revoke`, or `rollback` to audit statements regarding those SQL commands. Setting the value to `all` will audit every statement while `none` disables this feature.

11.3.1.3.7.10 `edb_audit_tag`

Parameter Type: String**Default Value:** none**Minimum Scope of Effect:** Session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** User

Use `edb audit tag` to specify a string value that will be included in the audit log when the `edb_audit` parameter is set to `csv` or `xml`.

11.3.1.3.7.11 `edb_audit_destination`

Parameter Type: Enum**Default Value:** file**Range:** {file | syslog}**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

Specifies whether the audit log information is to be recorded in the directory as given by the `edb_audit_directory` parameter or to the directory and file managed by the `syslog` process. Set to `file` to use the directory specified by `edb_audit_directory` (the default setting).

Set to `syslog` to use the syslog process and its location as configured in the `/etc/syslog.conf` file. The `syslog` setting is valid only for Advanced Server running on a Linux host, and is not supported on Windows systems.

Note: In recent Linux versions, `syslog` has been replaced by `rsyslog` and the configuration file is in `/etc/rsyslog.conf`.

11.3.1.3.7.12 `edb_log_every_bulk_value`

For information on `edb_log_every_bulk_value`, see the [edb_log_every_bulk_value](#).

11.3.1.3.8 Client Connection Defaults / Locale and Formatting

This section describes configuration parameters affecting locale and formatting.

`icu_short_form`

Parameter Type: String

Default Value: none

Range: n/a

Minimum Scope of Effect: Database

When Value Changes Take Effect: n/a

Required Authorization to Activate: n/a

The configuration parameter `icu_short_form` is a parameter containing the default ICU short form name assigned to a database or to the Advanced Server instance. See [Unicode Collation Algorithm](#) for general information about the ICU short form and the Unicode Collation Algorithm.

This configuration parameter is set either when the `CREATE DATABASE` command is used with the `ICU_SHORT_FORM` parameter in which case the specified short form name is set and appears in the `icu_short_form` configuration parameter when connected to this database, or when an Advanced Server instance is created with the `initdb` command used with the `--icu_short_form` option in which case the specified short form name is set and appears in the `icu_short_form` configuration parameter when connected to a database in that Advanced Server instance, and the database does not override it with its own `ICU_SHORT_FORM` parameter with a different short form.

Once established in the manner described, the `icu_short_form` configuration parameter cannot be changed.

11.3.1.3.9 Client Connection Defaults / Statement Behavior

This section describes configuration parameters affecting statement behavior.

`default_heap_fillfactor`

Parameter Type: Integer

Default Value: 100

Range: 10 to 100

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the fillfactor for a table when the `FILLFACTOR` storage parameter is omitted from a `CREATE TABLE` command.

The fillfactor for a table is a percentage between 10 and 100. 100 (complete packing) is the default. When a smaller `fillfactor` is specified, `INSERT` operations pack table pages only to the indicated percentage; the remaining space on each page is reserved for updating rows on that page. This gives `UPDATE` a chance to place the updated copy of a row on the same page as the original, which is more efficient than placing it on a different page. For a table whose entries are never updated, complete packing is the best choice, but in heavily updated tables smaller fillfactors are appropriate.

`edb_data_redaction`

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Data redaction is the support of policies to limit the exposure of certain sensitive data to certain users by altering the displayed information.

The default setting is `TRUE` so the data redaction is applied to all users except for superusers and the table owner:

- Superusers and table owner bypass data redaction.
- All other users get the redaction policy applied and see the reformatted data.

If the parameter is disabled by setting it to `FALSE`, then the following occurs:

- Superusers and table owner still bypass data redaction.

- All other users will get an error.

For information on data redaction, see the *EDB Postgres Advanced Server Security Features Guide*.

11.3.1.3.10 Client Connection Defaults / Other Defaults

The parameters in this section set other miscellaneous client connection defaults.

`oracle_home`

Parameter Type: String

Default Value: none

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Before creating an Oracle Call Interface (OCI) database link to an Oracle server, you must direct Advanced Server to the correct Oracle home directory. Set the `LD_LIBRARY_PATH` environment variable on Linux (or `PATH` on Windows) to the `lib` directory of the Oracle client installation directory.

For Windows only, you can instead set the value of the `oracle_home` configuration parameter in the `postgresql.conf` file. The value specified in the `oracle_home` configuration parameter will override the Windows `PATH` environment variable.

The `LD_LIBRARY_PATH` environment variable on Linux (`PATH` environment variable or `oracle_home` configuration parameter on Windows) must be set properly each time you start Advanced Server.

For Windows only: To set the `oracle_home` configuration parameter in the `postgresql.conf` file, edit the file, adding the following line:

```
| oracle_home = 'lib_directory'
```

Substitute the name of the Windows directory that contains `oci.dll` for `lib_directory`.

After setting the `oracle_home` configuration parameter, you must restart the server for the changes to take effect. Restart the server from the Windows Services console.

`odbc_lib_path`

Parameter Type: String

Default Value: none

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

If you will be using an ODBC driver manager, and if it is installed in a non-standard location, you specify the location by setting the `odbc_lib_path` configuration parameter in the `postgresql.conf` file:

```
odbc_lib_path = 'complete_path_to_libodbc.so'
```

The configuration file must include the complete pathname to the driver manager shared library (typically `libodbc.so`).

11.3.1.3.11 Compatibility Options

The configuration parameters described in this section control various database compatibility features.

edb_redwood_date

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When `DATE` appears as the data type of a column in the commands, it is translated to `TIMESTAMP` at the time the table definition is stored in the database if the configuration parameter `edb_redwood_date` is set to `TRUE`. Thus, a time component will also be stored in the column along with the date.

If `edb_redwood_date` is set to `FALSE` the column's data type in a `CREATE TABLE` or `ALTER TABLE` command remains as a native PostgreSQL `DATE` data type and is stored as such in the database. The PostgreSQL `DATE` data type stores only the date without a time component in the column.

Regardless of the setting of `edb_redwood_date`, when `DATE` appears as a data type in any other context such as the data type of a variable in an SPL declaration section, or the data type of a formal parameter in an SPL procedure or SPL function, or the return type of an SPL function, it is always internally translated to a `TIMESTAMP` and thus, can handle a time component if present.

edb_redwood_greatest_least

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate**Required Authorization to Activate:** Session user

The **GREATEST** function returns the parameter with the greatest value from its list of parameters. The **LEAST** function returns the parameter with the least value from its list of parameters.

When `edb_redwood_greatest_least` is set to `TRUE`, the **GREATEST** and **LEAST** functions return null when at least one of the parameters is null.

```
SETedb_redwood_greatest_least TO on;
```

```
SELECTGREATEST(1, 2, NULL, 3);
```

```
greatest
```

```
-----
```

```
(1 row)
```

When `edb_redwood_greatest_least` is set to `FALSE`, null parameters are ignored except when all parameters are null in which case null is returned by the functions.

```
SETedb_redwood_greatest_least TO off;
```

```
SELECTGREATEST(1, 2, NULL, 3);
```

```
greatest
```

```
-----
```

```
3
```

```
(1 row)
```

```
SELECTGREATEST(NULL, NULL, NULL);
```

```
greatest
```

```
-----
```

```
(1 row)
```

edb_redwood_raw_names

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When `edb_redwood_raw_names` is set to its default value of `FALSE`, database object names such as table names, column names, trigger names, program names, user names, etc. appear in uppercase letters when viewed from Redwood catalogs (that is, system catalogs prefixed by `ALL_`, `DBA_`, or `USER_`). In addition, quotation marks enclose names that were created with enclosing quotation marks.

When `edb_redwood_raw_names` is set to `TRUE`, the database object names are displayed exactly as they are stored in the PostgreSQL system catalogs when viewed from the Redwood catalogs. Thus, names created without enclosing quotation marks appear in lowercase as expected in PostgreSQL. Names created with enclosing quotation marks appear exactly as they were created, but without the quotation marks.

For example, the following user name is created, and then a session is started with that user.

```
CREATE USER reduser IDENTIFIED BY password;
edb=# \c - reduser
Password for user reduser:
You are now connected to database "edb" as user "reduser".
```

When connected to the database as reduser, the following tables are created.

```
CREATE TABLE all_lower (col INTEGER);
CREATE TABLE ALL_UPPER (COL INTEGER);
CREATE TABLE "Mixed_Case" ("Col" INTEGER);
```

When viewed from the Redwood catalog, `USER_TABLES`, with `edb_redwood_raw_names` set to the default value `FALSE`, the names appear in uppercase except for the `Mixed_Case` name, which appears as created and also with enclosing quotation marks.

```
edb=> SELECT * FROM USER_TABLES;
schema_name | table_name | tablespace_name | status | temporary
-----+-----+-----+-----+
REDUSER   | ALL_LOWER |           | VALID  | N
REDUSER   | ALL_UPPER |           | VALID  | N
REDUSER   | "Mixed_Case" |         | VALID  | N
(3 rows)
```

When viewed with `edb_redwood_raw_names` set to `TRUE`, the names appear in lowercase except for the `Mixed_Case` name, which appears as created, but now without the enclosing quotation marks.

```
edb=> SET edb_redwood_raw_names TO true;
SET
edb=> SELECT * FROM USER_TABLES;
schema_name | table_name | tablespace_name | status | temporary
-----+-----+-----+-----+
reduser   | all_lower |           | VALID  | N
reduser   | all_upper |           | VALID  | N
reduser   | Mixed_Case |         | VALID  | N
(3 rows)
```

These names now match the case when viewed from the PostgreSQL `pg_tables` catalog.

```
edb=> SELECT schemaname, tablename, tableowner FROM pg_tables WHERE
tableowner = 'reduser';
schemaname | tablename | tableowner
-----+-----+-----
reduser   | all_lower | reduser
reduser   | all_upper | reduser
reduser   | Mixed_Case | reduser
(3 rows)
```

edb_redwood_strings

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

If the `edb_redwood_strings` parameter is set to `TRUE`, when a string is concatenated with a null variable or null column, the result is the original string. If `edb_redwood_strings` is set to `FALSE`, the native PostgreSQL behavior is maintained, which is the concatenation of a string with a null variable or null column gives a null result.

The following example illustrates the difference.

The sample application contains a table of employees. This table has a column named `comm` that is null for most employees. The following query is run with `edb_redwood_string` set to `FALSE`. The concatenation of a null column with non-empty strings produces a final result of null, so only employees that have a commission appear in the query result. The output line for all other employees is null.

```
SET edb_redwood_strings TO off;

SELECT RPAD(ename,10) || ' ' || TO_CHAR(sal,'99,999.99') || ' ' ||
TO_CHAR(comm,'99,999.99') "EMPLOYEE COMPENSATION" FROM emp;
```

EMPLOYEE COMPENSATION

| | | |
|--------|----------|----------|
| ALLEN | 1,600.00 | 300.00 |
| WARD | 1,250.00 | 500.00 |
| MARTIN | 1,250.00 | 1,400.00 |

| | | |
|--------|----------|-----|
| TURNER | 1,500.00 | .00 |
|--------|----------|-----|

(14 rows)

The following is the same query executed when `edb_redwood_strings` is set to `TRUE`. Here, the value of a null column is treated as an empty string. The concatenation of an empty string with a non-empty string produces the non-empty string.

```
SET edb_redwood_strings TO on;

SELECT RPAD(ename,10) || ' ' || TO_CHAR(sal,'99,999.99') || ' ' ||
TO_CHAR(comm,'99,999.99') "EMPLOYEE COMPENSATION" FROM emp;
```

EMPLOYEE COMPENSATION

| | |
|-------|--------|
| SMITH | 800.00 |
|-------|--------|

```

ALLEN 1,600.00 300.00
WARD 1,250.00 500.00
JONES 2,975.00
MARTIN 1,250.00 1,400.00
BLAKE 2,850.00
CLARK 2,450.00
SCOTT 3,000.00
KING 5,000.00
TURNER 1,500.00 .00
ADAMS 1,100.00
JAMES 950.00
FORD 3,000.00
MILLER 1,300.00
(14 rows)

```

`edb_stmt_level_tx`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

The term **statement level transaction isolation** describes the behavior whereby when a runtime error occurs in a SQL command, all the updates on the database caused by that single command are rolled back. For example, if a single **UPDATE** command successfully updates five rows, but an attempt to update a sixth row results in an exception, the updates to all six rows made by this **UPDATE** command are rolled back. The effects of prior SQL commands that have not yet been committed or rolled back are pending until a **COMMIT** or **ROLLBACK** command is executed.

In Advanced Server, if an exception occurs while executing a SQL command, all the updates on the database since the start of the transaction are rolled back. In addition, the transaction is left in an aborted state and either a **COMMIT** or **ROLLBACK** command must be issued before another transaction can be started.

If `edb_stmt_level_tx` is set to `TRUE`, then an exception will not automatically roll back prior uncommitted database updates. If `edb_stmt_level_tx` is set to `FALSE`, then an exception will roll back uncommitted database updates.

Note

Use `edb_stmt_level_tx` set to `TRUE` only when absolutely necessary, as this may cause a negative performance impact.

The following example run in PSQL shows that when `edb_stmt_level_tx` is `FALSE`, the abort of the second **INSERT** command also rolls back the first **INSERT** command. Note that in PSQL, the command `\set AUTOCOMMIT off` must be issued, otherwise every statement commits automatically defeating the purpose of this demonstration of the effect of `edb_stmt_level_tx`.

```

\set AUTOCOMMIT off
SET edb_stmt_level_tx TO off;

INSERT INTO emp (empno,ename,deptno) VALUES (9001, 'JONES', 40);

```

```
INSERT INTO emp (empno,ename,deptno) VALUES (9002, 'JONES', 00);
ERROR: insert or update on table "emp" violates foreign key constraint
"emp_ref_dept_fk"
DETAIL: Key (deptno)=(0) is not present in table "dept".
```

```
COMMIT;
SELECT empno, ename, deptno FROM emp WHERE empno > 9000;
```

| empno | ename | deptno |
|-------|-------|--------|
| | | |

(0 rows)

In the following example, with `edb_stmt_level_tx` set to `TRUE`, the first `INSERT` command has not been rolled back after the error on the second `INSERT` command. At this point, the first `INSERT` command can either be committed or rolled back.

```
\set AUTOCOMMIT off
SET edb_stmt_level_tx TO on;
```

```
INSERT INTO emp (empno,ename,deptno) VALUES (9001, 'JONES', 40);
INSERT INTO emp (empno,ename,deptno) VALUES (9002, 'JONES', 00);
ERROR: insert or update on table "emp" violates foreign key constraint
"emp_ref_dept_fk"
DETAIL: Key (deptno)=(0) is not present in table "dept"
```

```
SELECT empno, ename, deptno FROM emp WHERE empno > 9000;
```

| empno | ename | deptno |
|-------|-------|--------|
| 9001 | JONES | 40 |

(1 row)

```
COMMIT;
```

A `ROLLBACK` command could have been issued instead of the `COMMIT` command in which case the insert of employee number `9001` would have been rolled back as well.

db_dialect

Parameter Type: Enum

Default Value: postgres

Range: {postgres | redwood}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

In addition to the native PostgreSQL system catalog, `pg_catalog`, Advanced Server contains an extended catalog view. This is the `sys` catalog for the expanded catalog view. The `db_dialect` parameter controls the order in which these catalogs are searched for name resolution.

When set to `postgres`, the namespace precedence is `pg_catalog` then `sys`, giving the PostgreSQL catalog the

highest precedence. When set to `redwood`, the namespace precedence is `sys` then `pg_catalog`, giving the expanded catalog views the highest precedence.

`default_with_rowids`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When set to `on`, `CREATE TABLE` includes a `ROWID` column in newly created tables, which can then be referenced in SQL commands. In earlier versions of Advanced Server `ROWIDs` were mapped to `OIDs`, but from Advanced Server version 12 onwards, the `ROWID` is an auto-incrementing value based on a sequence that starts with 1 and assigned to each row of a table created with `ROWIDs` option. By default, a unique index is created on a `ROWID` column.

The `ALTER` and `DROP` operations are restricted on a `ROWID` column.

To restore a database with `ROWIDs` from Advanced Server 11 or an earlier version, you must perform the following:

- `pg_dump`: If a table includes `OIDs` then specify `--convert-oids-to-rowids` to dump a database. Otherwise, ignore the `OIDs` to continue table creation on Advanced Server version 12 onwards.
- `pg_upgrade`: Errors out. But if a table includes `OIDs` or `ROWIDs`, then you must perform the following steps:
 1. Take a dump of the tables by specifying `--convert-oids-to-rowids` option.
 2. Drop the tables and then perform the upgrade.
 3. Restore the dump after the upgrade is successful into a new cluster that contains the dumped tables into a target database.

`optimizer_mode`

Parameter Type: Enum

Default Value: choose

Range: {choose | ALL_ROWS | FIRST_ROWS | FIRST_ROWS_10 | FIRST_ROWS_100 | FIRST_ROWS_1000}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the default optimization mode for analyzing optimizer hints.

The following table shows the possible values:

| Hint | Description |
|-----------------------|--|
| <code>ALL_ROWS</code> | Optimizes for retrieval of all rows of the result set. |

| | |
|-----------------|--|
| CHOOSE | Does no default optimization based on assumed number of rows to be retrieved from the result set. This is the default. |
| FIRST_ROWS | Optimizes for retrieval of only the first row of the result set. |
| FIRST_ROWS_10 | Optimizes for retrieval of the first 10 rows of the results set. |
| FIRST_ROWS_100 | Optimizes for retrieval of the first 100 rows of the result set. |
| FIRST_ROWS_1000 | Optimizes for retrieval of the first 1000 rows of the result set. |

These optimization modes are based upon the assumption that the client submitting the SQL command is interested in viewing only the first `n` rows of the result set and will then abandon the remainder of the result set. Resources allocated to the query are adjusted as such.

11.3.1.3.12 Customized Options

In previous releases of Advanced Server, the `custom_variable_classes` was required by those parameters not normally known to be added by add-on modules (such as procedural languages).

`custom_variable_classes`

The `custom_variable_classes` parameter is deprecated in Advanced Server 9.2; parameters that previously depended on this parameter no longer require its support.

`dbms_alert.max_alerts`

Parameter Type: Integer

Default Value: 100

Range: 0 to 500

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Specifies the maximum number of concurrent alerts allowed on a system using the `DBMS_ALERTS` package.

`dbms_pipe.total_message_buffer`

Parameter Type: Integer

Default Value: 30 Kb

Range: 30 Kb to 256 Kb

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Specifies the total size of the buffer used for the `DBMS_PIPE` package.

index_advisor.enabled

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Provides the capability to temporarily suspend Index Advisor in an EDB-PSQL or PSQL session. The Index Advisor plugin, `index advisor`, must be loaded in the EDB-PSQL or PSQL session in order to use the `index_advisor.enabled` configuration parameter.

The Index Advisor plugin can be loaded as shown by the following example:

```
$ psql -d edb -U enterpriseedb
Password for user enterpriseedb:
psql (13.0.0)
Type "help" for help.
```

```
edb=# LOAD 'index_advisor';
LOAD
```

Use the `SET` command to change the parameter setting to control whether or not Index Advisor generates an alternative query plan as shown by the following example:

```
edb=# SET index_advisor.enabled TO off;
SET
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
      QUERY PLAN
```

```
-----  
Seq Scan on t (cost=0.00..1693.00 rows=9864 width=8)  
  Filter: (a < 10000)  
(2 rows)
```

```
edb=# SET index_advisor.enabled TO on;
SET
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
      QUERY PLAN
```

```
-----  
Seq Scan on t (cost=0.00..1693.00 rows=9864 width=8)  
  Filter: (a < 10000)  
Result (cost=0.00..327.88 rows=9864 width=8)  
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]===:text  
-> Index Scan using "<hypothetical-index>:1" on t (cost=0.00..327.88  
rows=9864 width=8)
```

Index Cond: (a < 10000)
 (6 rows)

`edb_sql_protect.enabled`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Controls whether or not SQL/Protect is actively monitoring protected roles by analyzing SQL statements issued by those roles and reacting according to the setting of `edb_sql_protect.level`. When you are ready to begin monitoring with SQL/Protect set this parameter to on.

`edb_sql_protect.level`

Parameter Type: Enum

Default Value: passive

Range: {learn | passive | active}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Sets the action taken by SQL/Protect when a SQL statement is issued by a protected role.

The `edb_sql_protect.level` configuration parameter can be set to one of the following values to use either learn mode, passive mode, or active mode:

- **learn.** Tracks the activities of protected roles and records the relations used by the roles. This is used when initially configuring SQL/Protect so the expected behaviors of the protected applications are learned.
- **passive.** Issues warnings if protected roles are breaking the defined rules, but does not stop any SQL statements from executing. This is the next step after SQL/Protect has learned the expected behavior of the protected roles. This essentially behaves in intrusion detection mode and can be run in production when properly monitored.
- **active.** Stops all invalid statements for a protected role. This behaves as a SQL firewall preventing dangerous queries from running. This is particularly effective against early penetration testing when the attacker is trying to determine the vulnerability point and the type of database behind the application. Not only does SQL/Protect close those vulnerability points, but it tracks the blocked queries allowing administrators to be alerted before the attacker finds an alternate method of penetrating the system.

If you are using SQL/Protect for the first time, set `edb_sql_protect.level` to `learn`.

`edb_sql_protect.max_protected_relations`

Parameter Type: Integer

Default Value: 1024

Range: 1 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the maximum number of relations that can be protected per role. Please note the total number of protected relations for the server will be the number of protected relations times the number of protected roles. Every protected relation consumes space in shared memory. The space for the maximum possible protected relations is reserved during database server startup.

If the server is started when `edb_sql_protect.max_protected_relations` is set to a value outside of the valid range (for example, a value of 2,147,483,648), then a warning message is logged in the database server log file:

```
2014-07-18 16:04:12 EDT WARNING: invalid value for parameter
"edb_sql_protect.max_protected_relations": "2147483648"
2014-07-18 16:04:12 EDT HINT: Value exceeds integer range.
```

The database server starts successfully, but with `edb_sql_protect.max_protected_relations` set to the default value of 1024.

Though the upper range for the parameter is listed as the maximum value for an integer data type, the practical setting depends on how much shared memory is available and the parameter value used during database server startup.

As long as the space required can be reserved in shared memory, the value will be acceptable. If the value is such that the space in shared memory cannot be reserved, the database server startup fails with an error message such as the following:

```
2014-07-18 15:22:17 EDT FATAL: could not map anonymous shared memory:
Cannot allocate memory
2014-07-18 15:22:17 EDT HINT: This error usually means that PostgreSQL's
request for a shared memory segment exceeded available memory, swap
space or huge pages. To reduce the request size (currently 2070118400
bytes), reduce PostgreSQL's shared memory usage, perhaps by reducing
shared_buffers or max_connections.
```

In such cases, reduce the parameter value until the database server can be started successfully.

`edb_sql_protect.max_protected_roles`

Parameter Type: Integer

Default Value: 64

Range: 1 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the maximum number of roles that can be protected.

Every protected role consumes space in shared memory. Please note that the server will reserve space for the number of protected roles times the number of protected relations (`edb_sql_protect.max_protected_relations`). The space for the maximum possible protected roles is reserved during database server startup.

If the database server is started when `edb_sql_protect.max_protected_roles` is set to a value outside of the valid range (for example, a value of 2,147,483,648), then a warning message is logged in the database server log file:

```
2014-07-18 16:04:12 EDT WARNING: invalid value for parameter
"edb_sql_protect.max_protected_roles": "2147483648"
2014-07-18 16:04:12 EDT HINT: Value exceeds integer range.
```

The database server starts successfully, but with `edb_sql_protect.max_protected_roles` set to the default value of 64.

Though the upper range for the parameter is listed as the maximum value for an integer data type, the practical setting depends on how much shared memory is available and the parameter value used during database server startup.

As long as the space required can be reserved in shared memory, the value will be acceptable. If the value is such that the space in shared memory cannot be reserved, the database server startup fails with an error message such as the following:

```
2014-07-18 15:22:17 EDT FATAL: could not map anonymous shared memory:
Cannot allocate memory
2014-07-18 15:22:17 EDT HINT: This error usually means that PostgreSQL's
request for a shared memory segment exceeded available memory, swap
space or huge pages. To reduce the request size (currently 2070118400
bytes), reduce PostgreSQL's shared memory usage, perhaps by reducing
shared_buffers or max_connections.
```

In such cases, reduce the parameter value until the database server can be started successfully.

`edb_sql_protect.max_queries_to_save`

Parameter Type: Integer

Default Value: 5000

Range: 100 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the maximum number of offending queries to save in view `edb_sql_protect_queries`.

Every query that is saved consumes space in shared memory. The space for the maximum possible queries that can be saved is reserved during database server startup.

If the database server is started when `edb_sql_protect.max_queries_to_save` is set to a value outside of the valid range (for example, a value of 10), then a warning message is logged in the database server log file:

```
2014-07-18 13:05:31 EDT WARNING: 10 is outside the valid range for
parameter "edb_sql_protect.max_queries_to_save" (100 .. 2147483647)
```

The database server starts successfully, but with `edb_sql_protect.max_queries_to_save` set to the default value of 5000.

Though the upper range for the parameter is listed as the maximum value for an integer data type, the practical setting depends on how much shared memory is available and the parameter value used during database server startup.

As long as the space required can be reserved in shared memory, the value will be acceptable. If the value is such that the space in shared memory cannot be reserved, the database server startup fails with an error message such as the following:

```
2014-07-18 15:22:17 EDT FATAL: could not map anonymous shared memory:
```

Cannot allocate memory

```
2014-07-18 15:22:17 EDT HINT: This error usually means that PostgreSQL's request for a shared memory segment exceeded available memory, swap space or huge pages. To reduce the request size (currently 2070118400 bytes), reduce PostgreSQL's shared memory usage, perhaps by reducing shared_buffers or max_connections.
```

In such cases, reduce the parameter value until the database server can be started successfully.

`edb_wait_states.directory`

Parameter Type: String

Default Value: `edb_wait_states`

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the directory path where the EDB wait states log files are stored. The specified path should be a full, absolute path and not a relative path. However, the default setting is `edb_wait_states`, which makes `$PGDATA/edb_wait_states` the default directory location. For information on EDB wait states see [EDB Wait States](#) under [Performance Analysis and Tuning](#).

`edb_wait_states.retention_period`

Parameter Type: Integer

Default Value: 604800s

Range: 86400s to 2147483647s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Sets the period of time after which the log files for EDB wait states will be deleted. The default is 604800 seconds, which is 7 days. For information on EDB wait states see [EDB Wait States](#) under [Performance Analysis and Tuning](#).

`edb_wait_states.sampling_interval`

Parameter Type: Integer

Default Value: 1s

Range: 1s to 2147483647s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Sets the timing interval between two sampling cycles for EDB wait states. The default setting is 1 second. For information on EDB wait states see [EDB Wait States](#) under [Performance Analysis and Tuning](#).

`edbldr.empty_csv_field`

Parameter Type: Enum

Default Value: NULL

Range: {NULL | empty_string | pgsql}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Use the `edbldr.empty_csv_field` parameter to specify how EDB*Loader will treat an empty string. The valid values for the `edbldr.empty_csv_field` parameter are:

| Parameter Setting | EDB*Loader Behavior |
|-------------------|---|
| NULL | An empty field is treated as NULL. |
| empty_string | An empty field is treated as a string of length zero. |
| pgsql | An empty field is treated as a NULL if it does not contain quotes and as an empty string if it contains quotes. |

For more information about the `edbldr.empty_csv_field` parameter in EDB*Loader, see the *Database Compatibility for Oracle Developers Tools and Utilities Guide* at the EnterpriseDB website:

<https://www.enterprisedb.com/edb-docs>

`utl_encode.uudecode_redwood`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When set to `TRUE`, Advanced Server's `UTL_ENCODE.UUDECODE` function can decode uuencoded data that was created by the Oracle implementation of the `UTL_ENCODE.UUENCODE` function.

When set to the default setting of `FALSE`, Advanced Server's `UTL_ENCODE.UUDECODE` function can decode uuencoded data created by Advanced Server's `UTL_ENCODE.UUENCODE` function.

utl_file.umask

Parameter Type: String

Default Value: 0077

Range: Octal digits for umask settings

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

The `utl_file.umask` parameter sets the *file mode creation mask* or simply, the *mask*, in a manner similar to the Linux `umask` command. This is for usage only within the Advanced Server `UTL_FILE` package.

Note

The `utl_file.umask` parameter is not supported on Windows systems.

The value specified for `utl_file.umask` is a 3 or 4-character octal string that would be valid for the Linux `umask` command. The setting determines the permissions on files created by the `UTL_FILE` functions and procedures. (Refer to any information source regarding Linux or Unix systems for information on file permissions and the usage of the `umask` command.)

The following shows the results of the default `utl_file.umask` setting of 0077. Note that all permissions are denied on users belonging to the `enterprisedb` group as well as all other users. Only user `enterprisedb` has read and write permissions on the file.

```
-rw----- 1 enterprisedb enterprisedb 21 Jul 24 16:08 utlfile
```

11.3.1.3.13 Ungrouped

Configuration parameters in this section apply to Advanced Server only and are for a specific, limited purpose.

nls_length_semantics

Parameter Type: Enum

Default Value: byte

Range: {byte | char}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Superuser

This parameter has no effect in Advanced Server.

For example, the form of the `ALTER SESSION` command is accepted in Advanced Server without throwing a syntax error, but does not alter the session environment:

```
ALTER SESSION SET nls_length_semantics = char;
```

Note

Since the setting of this parameter has no effect on the server environment, it does not appear in the system view `pg_settings`.

query_rewrite_enabled

Parameter Type: Enum

Default Value: false

Range: {true | false | force}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

This parameter has no effect in Advanced Server.

For example, the following form of the `ALTER SESSION` command is accepted in Advanced Server without throwing a syntax error, but does not alter the session environment:

```
ALTER SESSION SET query_rewrite_enabled = force;
```

Note

Since the setting of this parameter has no effect on the server environment, it does not appear in the system view `pg_settings`.

query_rewrite_integrity

Parameter Type: Enum

Default Value: enforced

Range: {enforced | trusted | stale_tolerated}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Superuser

This parameter has no effect in Advanced Server.

For example, the following form of the `ALTER SESSION` command is accepted in Advanced Server without throwing a syntax error, but does not alter the session environment:

```
ALTER SESSION SET query_rewrite_integrity = stale_tolerated;
```

Note

Since the setting of this parameter has no effect on the server environment, it does not appear in the system view `pg_settings`.

timed_statistics

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Controls the collection of timing data for the Dynamic Runtime Instrumentation Tools Architecture (DRITA) feature. When set to on, timing data is collected.

Note

When Advanced Server is installed, the `postgresql.conf` file contains an explicit entry setting `timed_statistics` to off. If this entry is commented out letting `timed_statistics` to default, and the configuration file is reloaded, timed statistics collection would be turned on.

11.3.2 Index Advisor

The Index Advisor utility helps determine which columns you should index to improve performance in a given workload. Index Advisor considers B-tree (single-column or composite) index types, and does not identify other index types (GIN, GiST, Hash) that may improve performance. Index Advisor is installed with EDB Postgres Advanced Server.

Index Advisor works with Advanced Server's query planner by creating *hypothetical indexes* that the query planner uses to calculate execution costs as if such indexes were available. Index Advisor identifies the indexes by analyzing SQL queries supplied in the workload.

There are three ways to use Index Advisor to analyze SQL queries:

- Invoke the Index Advisor utility program, supplying a text file containing the SQL queries that you wish to analyze; Index Advisor will generate a text file with `CREATE INDEX` statements for the recommended indexes.
- Provide queries at the EDB-PSQL command line that you want Index Advisor to analyze.

- Access Index Advisor through the Postgres Enterprise Manager client. When accessed via the PEM client, Index Advisor works with SQL Profiler, providing indexing recommendations on code captured in SQL traces. For more information about using SQL Profiler and Index Advisor with PEM, see the *PEM Getting Started Guide* available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

Index Advisor will attempt to make indexing recommendations on `INSERT`, `UPDATE`, `DELETE` and `SELECT` statements. When invoking Index Advisor, you supply the workload in the form of a set of queries (if you are providing the command in an SQL file) or an `EXPLAIN` statement (if you are specifying the SQL statement at the `psql` command line). Index Advisor displays the query plan and estimated execution cost for the supplied query, but does not actually execute the query.

During the analysis, Index Advisor compares the query execution costs with and without hypothetical indexes. If the execution cost using a hypothetical index is less than the execution cost without it, both plans are reported in the `EXPLAIN` statement output, metrics that quantify the improvement are calculated, and Index Advisor generates the `CREATE INDEX` statement needed to create the index.

If no hypothetical index can be found that reduces the execution cost, Index Advisor displays only the original query plan output of the `EXPLAIN` statement.

Index Advisor does not actually create indexes on the tables. Use the `CREATE INDEX` statements supplied by Index Advisor to add any recommended indexes to your tables.

A script supplied with Advanced Server creates the table in which Index Advisor stores the indexing recommendations generated by the analysis; the script also creates a function and a view of the table to simplify the retrieval and interpretation of the results.

If you choose to forego running the script, Index Advisor will log recommendations in a temporary table that is available only for the duration of the Index Advisor session.

11.3.2.1 Index Advisor Components

The Index Advisor shared library interacts with the query planner to make indexing recommendations. On Windows, the Advanced Server installer creates the following shared library in the `libdir` subdirectory of your Advanced Server home directory. For Linux, install the `edb-asxx-server-indexadvisor` RPM package where `xx` is the Advanced Server version number.

On Linux:

`index_advisor.so`

On Windows:

`index_advisor.dll`

Please note that libraries in the `libdir` directory can only be loaded by a superuser. A database administrator can allow a non-superuser to use Index Advisor by manually copying the Index Advisor file from the `libdir` directory into the `libdir/plugins` directory (under your Advanced Server home directory). Only a trusted non-superuser should be allowed access to the plugin; this is an unsafe practice in a production environment.

The installer also creates the Index Advisor utility program and setup script:

`pg_advise_index`

`pg_advise_index` is a utility program that reads a user-supplied input file containing SQL queries and

produces a text file containing `CREATE INDEX` statements that can be used to create the indexes recommended by the Index Advisor. The `pg_advise_index` program is located in the `bin` subdirectory of the Advanced Server home directory.

`index_advisor.sql`

`index_advisor.sql` is a script that creates a permanent Index Advisor log table along with a function and view to facilitate reporting of recommendations from the log table. The script is located in the `share/contrib` subdirectory of the Advanced Server directory.

The `index_advisor.sql` script creates the `index_advisor_log` table, the `show_index_recommendations()` function and the `index_recommendations` view. These database objects must be created in a schema that is accessible by, and included in the search path of the role that will invoke Index Advisor.

`index_advisor_log`

Index Advisor logs indexing recommendations in the `index_advisor_log` table. If Index Advisor does not find the `index_advisor_log` table in the user's search path, Index Advisor will store any indexing recommendations in a temporary table of the same name. The temporary table exists only for the duration of the current session.

`show_index_recommendations()`

`show_index_recommendations()` is a PL/pgSQL function that interprets and displays the recommendations made during a specific Index Advisor session (as identified by its backend process ID).

`index_recommendations`

Index Advisor creates the `index_recommendations` view based on information stored in the `index_advisor_log` table during a query analysis. The view produces output in the same format as the `show_index_recommendations()` function, but contains Index Advisor recommendations for all stored sessions, while the result set returned by the `show_index_recommendations()` function are limited to a specified session.

11.3.2.2 Index Advisor Configuration

Index Advisor does not require configuration to generate recommendations that are available only for the duration of the current session; to store the results of multiple sessions, you must create the `index_advisor_log` table. To create the `index_advisor_log` table, you must run the `index_advisor.sql` script.

When selecting a storage schema for the Index Advisor table, function and view, keep in mind that all users that invoke Index Advisor (and query the result set) must have `USAGE` privileges on the schema. The schema must be in the search path of all users that are interacting with the Index Advisor.

1. Place the selected schema at the start of your `search_path` parameter. For example, if your search path is currently:

```
search_path=public, accounting
```

and you want the Index Advisor objects to be created in a schema named `advisor`, use the command:

```
SET search_path = advisor, public, accounting;
```

2. Run the `index_advisor.sql` script to create the database objects. If you are running the `psql` client, you can use the command:

`\i full.pathname/index_advisor.sql`

Specify the pathname to the `index_advisor.sql` script in place of `full.pathname`.

3. Grant privileges on the `index_advisor_log` table to all Index Advisor users; this step is not necessary if the Index Advisor user is a superuser, or the owner of these database objects.

- Grant `SELECT` and `INSERT` privileges on the `index_advisor_log` table to allow a user to invoke Index Advisor.
- Grant `DELETE` privileges on the `index_advisor_log` table to allow the specified user to delete the table contents.
- Grant `SELECT` privilege on the `index_recommendations` view.

The following example demonstrates the creation of the Index Advisor database objects in a schema named `ia`, which will then be accessible to an Index Advisor user with user name `ia_user`:

```
$ edb-psql -d edb -U enterpriseedb
psql.bin (13.0.0, server 13.0.0)
Type "help" for help.

edb=# CREATE SCHEMA ia;
CREATE SCHEMA
edb=# SET search_path TO ia;
SET
edb=# \i /usr/edb/as13/share/contrib/index_advisor.sql
CREATE TABLE
CREATE INDEX
CREATE INDEX
CREATE FUNCTION
CREATE FUNCTION
CREATE VIEW
edb=# GRANT USAGE ON SCHEMA ia TO ia_user;
GRANT
edb=# GRANT SELECT, INSERT, DELETE ON index_advisor_log TO ia_user;
GRANT
edb=# GRANT SELECT ON index_recommendations TO ia_user;
GRANT
```

While using Index Advisor, the specified schema (`ia`) must be included in `ia_user`'s `search_path` parameter.

11.3.2.3 Using Index Advisor

When you invoke Index Advisor, you must supply a workload; the workload is either a query (specified at the command line), or a file that contains a set of queries (executed by the `pg_advise_index()` function). After analyzing the workload, Index Advisor will either store the result set in a temporary table, or in a permanent table. You can review the indexing recommendations generated by Index Advisor and use the `CREATE INDEX` statements generated by Index Advisor to create the recommended indexes.

Note

You should not run Index Advisor in read-only transactions.

The following examples assume that superuser `enterpriseedb` is the Index Advisor user, and the Index Advisor

database objects have been created in a schema in the `search_path` of superuser `enterprisedb`.

The examples in the following sections use the table created with the statement shown below:

```
CREATE TABLE t( a INT, b INT );
INSERT INTO t SELECT s, 99999 - s FROM generate_series(0,99999) AS s;
ANALYZE t;
```

The resulting table contains the following rows:

| a | b |
|-------|-------|
| 0 | 99999 |
| 1 | 99998 |
| 2 | 99997 |
| 3 | 99996 |
| . | . |
| . | . |
| 99997 | 2 |
| 99998 | 1 |
| 99999 | 0 |

Using the pg_advise_index Utility

When invoking the `pg_advise_index` utility, you must include the name of a file that contains the queries that will be executed by `pg_advise_index`; the queries may be on the same line, or on separate lines, but each query must be terminated by a semicolon. Queries within the file should not begin with the `EXPLAIN` keyword.

The following example shows the contents of a sample `workload.sql` file:

```
SELECT * FROM t WHERE a = 500;
SELECT * FROM t WHERE b < 1000;
```

Run the `pg_advise_index` program as shown in the code sample below:

```
$ pg_advise_index -d edb -h localhost -U enterprisedb -s 100M -o
advisory.sql workload.sql
poolsize = 102400 KB
ad workload from file 'workload.sql'
Analyzing queries .. done.
size = 2184 KB, benefit = 1684.720000
size = 2184 KB, benefit = 1655.520000
/* 1. t(a): size=2184 KB, benefit=1684.72 */
/* 2. t(b): size=2184 KB, benefit=1655.52 */
/* Total size = 4368KB */
```

In the code sample, the `-d`, `-h`, and `-U` options are psql connection options.

`-S`

`-S` is an optional parameter that limits the maximum size of the indexes recommended by Index Advisor. If Index Advisor does not return a result set, `-S` may be set too low.

`-O`

The recommended indexes are written to the file specified after the `-o` option.

The information displayed by the `pg_advise_index` program is logged in the `index_advisor_log` table. In response to the command shown in the example, Index Advisor writes the following `CREATE INDEX` statements to the `advisory.sql` output file.

```
create index idx_t_1 on t (a);
create index idx_t_2 on t (b);
```

You can create the recommended indexes at the psql command line with the `CREATE INDEX` statements in the file, or create the indexes by executing the `advisory.sql` script.

```
$ edb-psql -d edb -h localhost -U enterprise -e -f advisory.sql
create index idx_t_1 on t (a);
CREATE INDEX
create index idx_t_2 on t (b);
CREATE INDEX
```

Using Index Advisor at the psql Command Line

You can use Index Advisor to analyze SQL statements entered at the `edb-psql` (or `psql`) command line; the following steps detail loading the Index Advisor plugin and using Index Advisor:

1. Connect to the server with the `edb-psql` command line utility, and load the Index Advisor plugin:

```
$ edb-psql -d edb -U enterprise
...
edb=# LOAD 'index_advisor';
LOAD
```

1. Use the `edb-psql` command line to invoke each SQL command that you would like Index Advisor to analyze. Index Advisor stores any recommendations for the queries in the `index_advisor_log` table. If the `index_advisor_log` table does not exist in the user's `search_path`, a temporary table is created with the same name. This temporary table exists only for the duration of the user's session.

After loading the Index Advisor plugin, Index Advisor will analyze all SQL statements and log any indexing recommendations for the duration of the session.

If you would like Index Advisor to analyze a query (and make indexing recommendations) without actually executing the query, preface the SQL statement with the `EXPLAIN` keyword.

If you do not preface the statement with the `EXPLAIN` keyword, Index Advisor will analyze the statement while the statement executes, writing the indexing recommendations to the `index_advisor_log` table for later review.

In the example that follows, the `EXPLAIN` statement displays the normal query plan, followed by the query plan of the same query, if the query were using the recommended hypothetical index:

```
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
```

QUERY PLAN

```
Seq Scan on t (cost=0.00..1693.00 rows=10105 width=8)
  Filter: (a < 10000)
Result (cost=0.00..337.10 rows=10105 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==':text
-> Index Scan using "<hypothetical-index>:1" on t
```

```
(cost=0.00..337.10 rows=10105 width=8)
```

```
Index Cond: (a < 10000)
```

```
(6 rows)
```

```
edb=# EXPLAIN SELECT * FROM t WHERE a = 100;
      QUERY PLAN
```

```
Seq Scan on t (cost=0.00..1693.00 rows=1 width=8)
```

```
Filter: (a = 100)
```

```
Result (cost=0.00..8.28 rows=1 width=8)
```

```
One-Time Filter: '===[ HYPOTHETICAL PLAN ]==='::text
```

```
-> Index Scan using "<hypothetical-index>:3" on t
```

```
(cost=0.00..8.28 rows=1 width=8)
```

```
Index Cond: (a = 100)
```

```
(6 rows)
```

After loading the Index Advisor plugin, the default value of `index_advisor.enabled` is `on`. The Index Advisor plugin must be loaded to use a `SET` or `SHOW` command to display the current value of `index_advisor.enabled`.

You can use the `index_advisor.enabled` parameter to temporarily disable Index Advisor without interrupting the `psql` session:

```
edb=# SET index_advisor.enabled TO off;
SET
```

To enable Index Advisor, set the parameter to `on`:

```
edb=# SET index_advisor.enabled TO on;
SET
```

11.3.2.4 Reviewing the Index Advisor Recommendations

There are several ways to review the index recommendations generated by Index Advisor. You can:

- Query the `index_advisor.log` table.
- Run the `show_index_recommendations` function.
- Query the `index_recommendations` view.

Using the `show_index_recommendations()` Function

To review the recommendations of the Index Advisor utility using the `show_index_recommendations()` function, call the function, specifying the process ID of the session:

```
| SELECT show_index_recommendations( pid );
```

Where `pid` is the process ID of the current session. If you do not know the process ID of your current session, passing a value of `NULL` will also return a result set for the current session.

The following code fragment shows an example of a row in a result set:

```
edb=# SELECT show_index_recommendations(null);
```

```
show_index_recommendations
```

```
create index idx_t_a on t(a);/* size: 2184 KB, benefit: 3040.62,
gain: 1.39222666981456 */
(1 row)
```

In the example, `create index idx_t_a on t(a)` is the SQL statement needed to create the index suggested by Index Advisor. Each row in the result set shows:

- The command required to create the recommended index.
- The maximum estimated size of the index.
- The calculated benefit of using the index.
- The estimated gain that will result from implementing the index.

You can display the results of all Index Advisor sessions from the following view:

```
| SELECT * FROM index_recommendations;
```

Querying the `index_advisor_log` Table

Index Advisor stores indexing recommendations in a table named `index_advisor_log`. Each row in the `index_advisor_log` table contains the result of a query where Index Advisor determines it can recommend a hypothetical index to reduce the execution cost of that query.

| Column | Type | Description |
|-------------|-----------|--|
| reloid | oid | OID of the base table for the index |
| relname | name | Name of the base table for the index |
| attrs | integer[] | Recommended index columns (identified by column number) |
| benefit | real | Calculated benefit of the index for this query |
| index_size | integer | Estimated index size in disk-pages |
| backend_pid | integer | Process ID of the process generating this recommendation |
| timestamp | timestamp | Date/Time when the recommendation was generated |

You can query the `index_advisor_log` table at the psql command line. The following example shows the `index_advisor_log` table entries resulting from two Index Advisor sessions. Each session contains two queries, and can be identified (in the table below) by a different `backend_pid` value. For each session, Index Advisor generated two index recommendations.

```
edb=# SELECT * FROM index_advisor_log;
reloid | relname | attrs | benefit | index_size | backend_pid | timestamp
-----+-----+-----+-----+-----+-----+
16651 | t | {1} | 1684.72 | 2184 | 3442 | 22-MAR-11
16:44:32.712638 -04:00
16651 | t | {2} | 1655.52 | 2184 | 3442 | 22-MAR-11
16:44:32.759436 -04:00
16651 | t | {1} | 1355.9 | 2184 | 3506 | 22-MAR-11
16:48:28.317016 -04:00
16651 | t | {1} | 1684.72 | 2184 | 3506 | 22-MAR-11
16:51:45.927906 -04:00
(4 rows)
```

Index Advisor added the first two rows to the table after analyzing the following two queries executed by the `pg_advise_index` utility:

```
SELECT * FROM t WHERE a = 500;
SELECT * FROM t WHERE b < 1000;
```

The value of `3442` in column `backend_pid` identifies these results as coming from the session with process ID `3442`.

The value of `1` in column `attrs` in the first row indicates that the hypothetical index is on the first column of the table (column `a` of table `t`).

The value of `2` in column `attrs` in the second row indicates that the hypothetical index is on the second column of the table (column `b` of table `t`).

Index Advisor added the last two rows to the table after analyzing the following two queries (executed at the psql command line):

```
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
      QUERY PLAN
-----
-----  

Seq Scan on t (cost=0.00..1693.00 rows=10105 width=8)
  Filter: (a < 10000)
Result (cost=0.00..337.10 rows=10105 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==='::text
    -> Index Scan using "<hypothetical-index>:1" on t (cost=0.00..337.10
rows=10105 width=8)
      Index Cond: (a < 10000)
(6 rows)
```

```
edb=# EXPLAIN SELECT * FROM t WHERE a = 100;
      QUERY PLAN
-----
```

```
Seq Scan on t (cost=0.00..1693.00 rows=1 width=8)
  Filter: (a = 100)
Result (cost=0.00..8.28 rows=1 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==='::text
    -> Index Scan using "<hypothetical-index>:3" on t (cost=0.00..8.28
rows=1 width=8)
      Index Cond: (a = 100)
(1 rows)
```

The values in the `benefit` column of the `index_advisor_log` table are calculated using the following formula:

`benefit = (normal execution cost) - (execution cost with hypothetical index)`

The value of the `benefit` column for the last row of the `index_advisor_log` table (shown in the example) is calculated using the query plan for the following SQL statement:

`EXPLAIN SELECT * FROM t WHERE a = 100;`

The execution costs of the different execution plans are evaluated and compared:

`benefit = (Seq Scan on t cost) - (Index Scan using <hypothetical-index>)`

and the benefit is added to the table:

`benefit = 1693.00 - 8.28`

`benefit = 1684.72`

You can delete rows from the `index_advisor_log` table when you no longer have the need to review the results of the queries stored in the row.

Querying the `index_recommendations` View

The `index_recommendations` view contains the calculated metrics and the `CREATE INDEX` statements to create the recommended indexes for all sessions whose results are currently in the `index_advisor_log` table. You can display the results of all stored Index Advisor sessions by querying the `index_recommendations` view as shown below:

```
| SELECT * FROM index_recommendations;
```

Using the example shown in the previous section ([Querying the `index_advisor_log` Table](#)), the `index_recommendations` view displays the following:

```
edb=# SELECT * FROM index_recommendations;
backend_pid | show_index_recommendations
-----+-----
3442 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
1684.72, gain: 0.771392654586624 */
3442 | create index idx_t_b on t(b);/* size: 2184 KB, benefit:
1655.52, gain: 0.758021539820856 */
3506 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
3040.62, gain: 1.39222666981456 */
(3 rows)
```

Within each session, the results of all queries that benefit from the same recommended index are combined to produce one set of metrics per recommended index, reflected in the fields named `benefit` and `gain`.

The formulas for the fields are as follows:

```
size = MAX(index size of all queries)
benefit = SUM(benefit of each query)
gain = SUM(benefit of each query) / MAX(index size of all queries)
```

So for example, using the following query results from the process with a `backend_pid` of `3506`:

```
reloid | relname | attrs | benefit | index_size | backend_pid | timestamp
-----+-----+-----+-----+-----+
16651 | t | {1} | 1355.9 | 2184 | 3506 | 22-MAR-11
16:48:28.317016 -04:00
16651 | t | {1} | 1684.72 | 2184 | 3506 | 22-MAR-11
16:51:45.927906 -04:00
```

The metrics displayed from the `index_recommendations` view for `backend_pid` `3506` are:

```
backend_pid | show_index_recommendations
-----+-----
3506 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
3040.62, gain: 1.39222666981456 */
```

The metrics from the view are calculated as follows:

```
benefit = (benefit from 1st query) + (benefit from 2nd query)
```

```
benefit = 1355.9 + 1684.72
```

```
benefit = 3040.62
```

and

```
gain = ((benefit from 1st query) + (benefit from 2nd query)) / MAX(index  
size of all queries)
```

```
gain = (1355.9 + 1684.72) / MAX(2184, 2184)
```

```
gain = 3040.62 / 2184
```

```
gain = 1.39223
```

The gain metric is useful when comparing the relative advantage of the different recommended indexes derived during a given session. The larger the gain value, the better the cost effectiveness derived from the index weighed against the possible disk space consumption of the index.

11.3.2.5 Limitations

Index Advisor does not consider Index Only scans; it does consider Index scans when making recommendations.

Index Advisor ignores any computations found in the `WHERE` clause. Effectively, the index field in the recommendations will not be any kind of expression; the field will be a simple column name.

Index Advisor does not consider inheritance when recommending hypothetical indexes. If a query references a parent table, Index Advisor does not make any index recommendations on child tables.

Restoration of a `pg_dump` backup file that includes the `index_advisor_log` table or any tables for which indexing recommendations were made and stored in the `index_advisor_log` table, may result in "broken links" between the `index_advisor_log` table and the restored tables referenced by rows in the `index_advisor_log` table because of changes in object identifiers (OIDs).

If it is necessary to display the recommendations made prior to the backup, you can replace the old OIDs in the `reloid` column of the `index_advisor_log` table with the new OIDs of the referenced tables using the SQL `UPDATE` statement:

```
UPDATE index_advisor_log SET reloid = new_oid WHERE reloid = old_oid;
```

11.3.3 SQL Profiler

Inefficient SQL code is one of, if not the leading cause of database performance problems. The challenge for database administrators and developers is locating and then optimizing this code in large, complex systems.

SQL Profiler helps you locate and optimize poorly running SQL code.

Specific features and benefits of SQL Profiler include the following:

- **On-Demand Traces.** You can capture SQL traces at any time by manually setting up your parameters and starting the trace.
- **Scheduled Traces.** For inconvenient times, you can also specify your trace parameters and schedule them to

run at some later time.

- **Save Traces.** Execute your traces and save them for later review.
- **Trace Filters.** Selectively filter SQL captures by database and by user, or capture every SQL statement sent by all users against all databases.
- **Trace Output Analyzer.** A graphical table lets you quickly sort and filter queries by duration or statement, and a graphical or text based **EXPLAIN** plan lays out your query paths and joins.
- **Index Advisor Integration.** Once you have found your slow queries and optimized them, you can also let the Index Advisor recommend the creation of underlying table indices to further improve performance.

The following describes the installation process.

Step 1: Install SQL Profiler

SQL Profiler is installed by the Advanced Server installer on Windows or from the **edb-asxx-server-sqlprofiler** RPM package on Linux where **xx** is the Advanced Server version number.

Step 2: Add the SQL Profiler library

Modify the **postgresql.conf** parameter file for the instance to include the SQL Profiler library in the **shared_preload_libraries** configuration parameter.

For Linux installations, the parameter value should include:

```
$libdir/sql-profiler
```

On Windows, the parameter value should include:

```
$libdir\sql-profiler.dll
```

Step 3: Create the functions used by SQL Profiler

The SQL Profiler installation program places a SQL script (named **sql-profiler.sql**) in:

On Linux:

```
/usr/edb/as13/share/contrib/
```

On Windows:

```
C:\Program Files\edb\as13\share\contrib\
```

Use the **psql** command line interface to run the **sql-profiler.sql** script in the database specified as the Maintenance Database on the server you wish to profile. If you are using Advanced Server, the default maintenance database is named **edb**. If you are using a PostgreSQL instance, the default maintenance database is named **postgres**.

The following command uses the **psql** command line to invoke the **sql-profiler.sql** script on a Linux system:

```
$ /usr/edb/as13/bin/psql -U user_name database_name <
/usr/edb/as13/share/contrib/sql-profiler.sql
```

Step 4: Stop and restart the server for the changes to take effect.

After configuring SQL Profiler, it is ready to use with all databases that reside on the server. You can take advantage of SQL Profiler functionality with EDB Postgres Enterprise Manager; for more information about Postgres Enterprise Manager, visit the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

Troubleshooting

If (after performing an upgrade to a newer version of SQL Profiler) you encounter an error that contains the

following text:

An error has occurred:

ERROR: function return row and query-specified return row do not match.

DETAIL: Returned row contains 11 attributes, but the query expects 10.

To correct this error, you must replace the existing query set with a new query set. First, uninstall SQL Profiler by invoking the `uninstall-sql-profiler.sql` script, and then reinstall SQL Profiler by invoking the `sql-profiler.sql` script.

11.3.4 pgsnmpd

`pgsnmpd` is an SNMP agent that can return hierarchical information about the current state of Advanced Server on a Linux system. `pgsnmpd` is distributed and installed using the `edb-asxx-pgsnmpd` RPM package where `xx` is the Advanced Server version number. The `pgsnmpd` agent can operate as a stand-alone SNMP agent, as a pass-through sub-agent, or as an AgentX sub-agent.

After installing Advanced Server, you will need to update the `LD_LIBRARY_PATH` variable. Use the command:

```
$ export LD_LIBRARY_PATH=/usr/edb/as13/lib:$LD_LIBRARY_PATH
```

This command does not persistently alter the value of `LD_LIBRARY_PATH`. Consult the documentation for your distribution of Linux for information about persistently setting the value of `LD_LIBRARY_PATH`.

The examples that follow demonstrate the simplest usage of `pgsnmpd`, implementing read only access. `pgsnmpd` is based on the net-snmp library; for more information about net-snmp, visit:

<http://net-snmp.sourceforge.net/>

Configuring pgsnmpd

The `pgsnmpd` configuration file is named `snmpd.conf`. For information about the directives that you can specify in the configuration file, please review the `snmpd.conf` man page (`man snmpd.conf`).

You can create the configuration file by hand, or you can use the `snmpconf` perl script to create the configuration file. The perl script is distributed with net-snmp package.

net-snmp is an open-source package available from:

<http://www.net-snmp.org/>

To use the `snmpconf` configuration file wizard, download and install net-snmp. When the installation completes, open a command line and enter:

`snmpconf`

When the configuration file wizard opens, it may prompt you to read in an existing configuration file. Enter `none` to generate a new configuration file (not based on a previously existing configuration file).

`snmpconf` is a menu-driven wizard. Select menu item `1: snmpd.conf` to start the configuration wizard. As you select each top-level menu option displayed by `snmpconf`, the wizard walks through a series of questions, prompting you for information required to build the configuration file. When you have provided information in each of the category relevant to your system, enter `Finished` to generate a configuration file named `snmpd.conf`. Copy the file to:

/usr/edb/as13/share/

Setting the Listener Address

By default, `pgsnmpd` listens on port `161`. If the listener port is already being used by another service, you may receive the following error:

Error opening specified endpoint "udp:161".

You can specify an alternate listener port by adding the following line to your `snmpd.conf` file:

`agentaddress $host_address:2000`

The example instructs `pgsnmpd` to listen on UDP port `2000`, where `$host_address` is the IP address of the server (e.g., `127.0.0.1`).

Invoking pgsnmpd

Ensure that an instance of Advanced Server is up and running (`pgsnmpd` will connect to this server). Open a command line and assume superuser privileges, before invoking `pgsnmpd` with a command that takes the following form:

```
POSTGRES_INSTALL_HOME/bin/pgsnmpd -b  
-c POSTGRES_INSTALL_HOME/share/snmpd.conf  
-C "user=enterprisedb dbname=edb password=safe_password  
port=5444"
```

Where `POSTGRES_INSTALL_HOME` specifies the Advanced Server installation directory.

Include the `-b` option to specify that `pgsnmpd` should run in the background.

Include the `-c` option, specifying the path and name of the `pgsnmpd` configuration file.

Include connection information for your installation of Advanced Server (in the form of a `libpq` connection string) after the `-C` option.

Viewing pgsnmpd Help

Include the `--help` option when invoking the `pgsnmpd` utility to view other `pgsnmpd` command line options:

```
pgsnmpd --help  
Version PGSQ-L-SNMP-Ver1.0  
usage: pgsnmpd [-s] [-b] [-c FILE] [-x address] [-g] [-C "Connect  
String"]  
-s : run as AgentX sub-agent of an existing snmpd process  
-b : run in the background  
-c : configuration file name  
-g : use syslog  
-C : libpq connection string  
-x : address:port of a network interface  
-V : display version strings
```

Requesting Information from pgsnmpd

You can use `net-snmp` commands to query the `pgsnmpd` service. For example:

```
snmpgetnext -v 2c -c public localhost .1.3.6.1.4.1.5432.1.4.2.1.1.0
```

In the above example:

`-v 2c` option instructs the `snmpgetnext` client to send the request in SNMP version 2c format.

`-c public` specifies the community name.

`localhost` indicates the host machine running the `pgsnmpd` server.

`.1.3.6.1.4.1.5432.1.4.2.1.1.0` is the identity of the requested object. To see a list of all databases, increment the last digit by one (e.g. .1.1, .1.2, .1.3 etc.).

The encodings required to query any given object are defined in the MIB (Management Information Base). An SNMP client can monitor a variety of servers; the server type determines the information exposed by a given server. Each SNMP server describes the exposed data in the form of a MIB (Management information base). By default, `pgsnmpd` searches for MIB's in the following locations:

`/usr/share/snmp/mibs`

`$HOME/.snmp/mibs`

11.3.5 EDB Audit Logging

Advanced Server allows database and security administrators, auditors, and operators to track and analyze database activities using the *EDB Audit Logging* functionality. EDB Audit Logging generates audit log files, which contains all of the relevant information. The audit logs can be configured to record information such as:

- When a role establishes a connection to an Advanced Server database
- What database objects a role creates, modifies, or deletes when connected to Advanced Server
- When any failed authentication attempts occur

The parameters specified in the configuration files, `postgresql.conf` or `postgresql.auto.conf`, control the information included in the audit logs.

11.3.5.1 Audit Logging Configuration Parameters

Use the following configuration parameters to control database auditing. See [Summary of Configuration Parameters](#) to determine if a change to the configuration parameter takes effect immediately, or if the configuration needs to be reloaded, or if the database server needs to be restarted.

`edb_audit`

Enables or disables database auditing. The values `xml` or `csv` will enable database auditing. These values represent the file format in which auditing information will be captured. `none` will disable database auditing

and is also the default.

`edb_audit_directory`

Specifies the directory where the log files will be created. The path of the directory can be relative or absolute to the data folder. The default is the `PGDATA/edb_audit` directory.

`edb_audit_filename`

Specifies the file name of the audit file where the auditing information will be stored. The default file name will be `audit-%Y%m%d_%H%M%S`. The escape sequences, `%Y`, `%m` etc., will be replaced by the appropriate current values according to the system date and time.

`edb_audit_rotation_day`

Specifies the day of the week on which to rotate the audit files. Valid values are `sun`, `mon`, `tue`, `wed`, `thu`, `fri`, `sat`, `every`, and `none`. To disable rotation, set the value to `none`. To rotate the file every day, set the `edb_audit_rotation_day` value to `every`. To rotate the file on a specific day of the week, set the value to the desired day of the week. `every` is the default value.

`edb_audit_rotation_size`

Specifies a file size threshold in megabytes when file rotation will be forced to occur. The default value is 0 MB. If the parameter is commented out or set to 0, rotation of the file on a size basis will not occur.

`edb_audit_rotation_seconds`

Specifies the rotation time in seconds when a new log file should be created. To disable this feature, set this parameter to 0, which is the default.

`edb_audit_connect`

Enables auditing of database connection attempts by users. To disable auditing of all connection attempts, set `edb_audit_connect` to `none`. To audit all failed connection attempts, set the value to `failed`, which is the default. To audit all connection attempts, set the value to `all`.

`edb_audit_disconnect`

Enables auditing of database disconnections by connected users. To enable auditing of disconnections, set the value to `all`. To disable, set the value to `none`, which is the default.

`edb_audit_statement`

This configuration parameter is used to specify auditing of different categories of SQL statements. Various combinations of the following values may be specified: `none`, `dml`, `insert`, `update`, `delete`, `truncate`, `select`, `error`, `rollback`, `ddl`, `create`, `drop`, `alter`, `grant`, `revoke`, and `all`. The default is `ddl` and `error`. See [Selecting SQL Statements to Audit](#) for information regarding the setting of this parameter.

`edb_audit_tag`

Use this configuration parameter to specify a string value that will be included in the audit log file for each entry as a tracking tag.

`edb_log_every_bulk_value`

Bulk processing logs the resulting statements into both the Advanced Server log file and the EDB Audit log file. However, logging each and every statement in bulk processing is costly. This can be controlled by the `edb_log_every_bulk_value` configuration parameter. When set to `true`, each and every statement in bulk processing is logged. During bulk execution, when `edb_log_every_bulk_value` is set to `false`, a log message is recorded once per bulk processing along with the number of rows processed. In addition, the duration is emitted once per bulk processing. Default is set to `false`.

edb_audit_destination

Specifies whether the audit log information is to be recorded in the directory as given by the `edb_audit_directory` parameter or to the directory and file managed by the `syslog` process. Set to `file` to use the directory specified by `edb_audit_directory`, which is the default setting.

Set to `syslog` to use the `syslog` process and its location as configured in the `/etc/syslog.conf` file. The `syslog` setting is valid for Advanced Server running on a Linux host and is not supported on Windows systems. **Note:** In recent Linux versions, `syslog` has been replaced by `rsyslog` and the configuration file is in `/etc/rsyslog.conf`.

The following section describes selection of specific SQL statements for auditing using the `edb_audit_statement` parameter.

11.3.5.2 Selecting SQL Statements to Audit

The `edb_audit_statement` permits inclusion of one or more, comma-separated values to control which SQL statements are to be audited. The following is the general format:

```
edb_audit_statement = 'value_1[, value_2]...'
```

The comma-separated values may include or omit space characters following the comma. The values can be specified in any combination of lowercase or uppercase letters.

The basic parameter values are the following:

- `all` – Results in the auditing and logging of every statement including any error messages on statements.
- `none` – Disables all auditing and logging. A value of `none` overrides any other value included in the list.
- `ddl` – Results in the auditing of all data definition language (DDL) statements (`CREATE`, `ALTER`, and `DROP`) as well as `GRANT` and `REVOKE` data control language (DCL) statements.
- `dml` – Results in the auditing of all data manipulation language (DML) statements (`INSERT`, `UPDATE`, `DELETE`, and `TRUNCATE`).
- `select` – Results in the auditing of `SELECT` statements.
- `rollback` – Results in the auditing of `ROLLBACK` statements.
- `error` – Results in the logging of all error messages that occur. Unless the `error` value is included, no error messages are logged regarding any errors that occur on SQL statements related to any of the other preceding parameter values except when `all` is used.

[Data Definition Language and Data Control Language Statements](#) describes additional parameter values for selecting particular DDL or DCL statements for auditing.

[Data Manipulation Language Statements](#) describes additional parameter values for selecting particular DML statements for auditing.

If an unsupported value is included in the `edb_audit_statement` parameter, then an error occurs when applying the setting to the database server. See the database server log file for the error such as in the following example where `create materialized vw` results in the error. (The correct value is `create materialized view`.)

```
2017-07-16 11:20:39 EDT LOG: invalid value for parameter
"edb_audit_statement": "create materialized vw, create sequence, grant"
2017-07-16 11:20:39 EDT FATAL: configuration file "/var/lib/edb/as13/data/
postgresql.conf" contains errors
```

The following sections describe the values for the SQL language types DDL, DCL, and DML.

Data Definition Language and Data Control Language Statements

This section describes values that can be included in the `edb_audit_statement` parameter to audit DDL and DCL statements.

The following general rules apply:

- If the `edb_audit_statement` parameter includes either `ddl` or `all`, then all DDL statements are audited. In addition, the DCL statements `GRANT` and `REVOKE` are audited as well.
- If the `edb_audit_statement` is set to `none`, then no DDL nor DCL statements are audited.
- Specific types of DDL and DCL statements can be chosen for auditing by including a combination of values within the `edb_audit_statement` parameter.

Use the following syntax to specify an `edb_audit_statement` parameter value for DDL statements:

```
> { create | alter | drop } [ <object_type> ]
```

`object_type` is any of the following:

ACCESS METHOD

AGGREGATE

CAST

COLLATION

CONVERSION

DATABASE

EVENT TRIGGER

EXTENSION

FOREIGN TABLE

FUNCTION

INDEX

LANGUAGE

LARGE OBJECT

MATERIALIZED VIEW

OPERATOR

OPERATOR CLASS

OPERATOR FAMILY

POLICY

PUBLICATION

ROLE

RULE

SCHEMA

SEQUENCE**SERVER****SUBSCRIPTION****TABLE****TABLESPACE****TEXT SEARCH CONFIGURATION****TEXT SEARCH DICTIONARY****TEXT SEARCH PARSER****TEXT SEARCH TEMPLATE****TRANSFORM****TRIGGER****TYPE****USER MAPPING****VIEW**

Descriptions of object types as used in SQL commands can be found in the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-commands.html>

If `object_type` is omitted from the parameter value, then all of the specified command statements (either `create`, `alter`, or `drop`) are audited.

Use the following syntax to specify an `edb_audit_statement` parameter value for DCL statements:

> { grant | revoke }

The following are some DDL and DCL examples.

Example 1

The following is an example where `edb_audit_connect` and `edb_audit_statement` are set with the following non-default values:

```
logging_collector = 'on'
edb_audit_connect = 'all'
edb_audit_statement = 'create, alter, error'
```

Thus, only SQL statements invoked by the `CREATE` and `ALTER` commands are audited. Error messages are also included in the audit log.

The database session that occurs is the following:

```
$ psql edb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.
```

```

edb=# SHOW edb_audit_connect;
edb_audit_connect
-----
all
(1 row)

edb=# SHOW edb_audit_statement;
edb_audit_statement
-----
create, alter, error
(1 row)

edb=# CREATE ROLE adminuser;
CREATE ROLE
edb=# ALTER ROLE adminuser WITH LOGIN, SUPERUSER, PASSWORD 'password';
ERROR: syntax error at or near ","
LINE 1: ALTER ROLE adminuser WITH LOGIN, SUPERUSER, PASSWORD 'passwo...
          ^
edb=# ALTER ROLE adminuser WITH LOGIN SUPERUSER PASSWORD 'password';
ALTER ROLE
edb=# CREATE DATABASE auditdb;
CREATE DATABASE
edb=# ALTER DATABASE auditdb OWNER TO adminuser;
ALTER DATABASE
edb=# \c auditdb adminuser
Password for user adminuser:
You are now connected to database "auditdb" as user "adminuser".
auditdb=# CREATE SCHEMA edb;
CREATE SCHEMA
auditdb=# SET search_path TO edb;
SET
auditdb=# CREATE TABLE department (
auditdb(# deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
auditdb(# dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
auditdb(# loc        VARCHAR2(13)
auditdb(# );
CREATE TABLE
auditdb=# DROP TABLE department;
DROP TABLE
auditdb=# CREATE TABLE dept (
auditdb(# deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
auditdb(# dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
auditdb(# loc        VARCHAR2(13)
auditdb(# );
CREATE TABLE

```

The resulting audit log file contains the following.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-25 12:32:22.799 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,1,
"authentication",2020-05-25 12:32:22 IST,4/19,0,AUDIT,00000,"connection authorized:
user=enterprisedb database=edb",,,,,,,,"client backend","","connect"

```

```

2020-05-25 12:34:05.843 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,2,

```

```
"idle",2020-05-25 12:32:22 IST,4/23,0,AUDIT,00000,"statement: CREATE ROLE adminuser;",
,,,,,,,"pgsql","client backend","CREATE ROLE","create"

2020-05-25 12:34:16.617 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,3,
"idle",2020-05-25 12:32:22 IST,4/24,0,ERROR,42601,"syntax error at or near ""","",
,"ALTER ROLE adminuser WITH LOGIN, SUPERUSER, PASSWORD 'password';",32,,,"pgsql",
"client backend","","error"

2020-05-25 12:34:29.954 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,4,
"idle",2020-05-25 12:32:22 IST,4/25,0,AUDIT,00000,"statement: ALTER ROLE adminuser
WITH LOGIN SUPERUSER PASSWORD 'password';",,,,,,,,"pgsql","client backend","ALTER ROLE
","", "alter"

2020-05-25 12:34:40.114 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,5,
"idle",2020-05-25 12:32:22 IST,4/26,0,AUDIT,00000,"statement: CREATE DATABASE auditdb;
",,,,,,,,"pgsql","client backend","CREATE DATABASE","create"

2020-05-25 12:34:50.591 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,6,
"idle",2020-05-25 12:32:22 IST,4/27,0,AUDIT,00000,"statement: ALTER DATABASE auditdb
OWNER TO adminuser;","",,,,"pgsql","client backend","ALTER DATABASE","alter"

2020-05-25 12:35:01.554 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,1,
"authentication",2020-05-25 12:35:01 IST,5/11,0,AUDIT,00000,"connection authorized:
user=adminuser database=auditdb","",,,,"client backend","","connect"

2020-05-25 12:35:12.931 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,2,
"idle",2020-05-25 12:35:01 IST,5/13,0,AUDIT,00000,"statement: CREATE SCHEMA edb;","",,
,,,"pgsql","client backend","CREATE SCHEMA","create"

2020-05-25 12:37:18.547 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,3,
"idle",2020-05-25 12:35:01 IST,5/15,0,AUDIT,00000,"statement: CREATE TABLE department
(
    deptno NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname  VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc    VARCHAR2(13)
);","",,,,"pgsql","client backend","CREATE TABLE","create"

2020-05-25 12:39:09.065 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,4,
"idle",2020-05-25 12:35:01 IST,5/17,0,AUDIT,00000,"statement: CREATE TABLE dept (
    deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc         VARCHAR2(13)
);","",,,,"pgsql","client backend","CREATE TABLE","create"
```

The `CREATE` and `ALTER` statements for the `adminuser` role and `auditdb` database are audited. The error for the `ALTER ROLE adminuser` statement is also logged since error is included in the `edb_audit_statement` parameter.

Similarly, the `CREATE` statements for schema `edb` and tables `department` and `dept` are audited.

Note that the `DROP TABLE department` statement is not in the audit log since there is no `edb_audit_statement` setting that would result in the auditing of successfully processed `DROP` statements such as `ddl`, `all`, or `drop`.

Example 2

The following is an example where `edb_audit_connect` and `edb_audit_statement` are set with the following non-default values:

```
logging_collector = 'on'
edb_audit_connect = 'all'
edb_audit_statement = create view,create materialized view,create
sequence,grant'
```

Thus, only SQL statements invoked by the **CREATE VIEW**, **CREATE MATERIALIZED VIEW**, **CREATE SEQUENCE** and **GRANT** commands are audited.

The database session that occurs is the following:

```
$ psql auditdb adminuser
Password for user adminuser:
psql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_connect;
edb_audit_connect
-----
all
(1 row)

auditdb=# SHOW edb_audit_statement;
edb_audit_statement
-----
create view,create materialized view,create sequence,grant
(1 row)

auditdb=# SET search_path TO edb;
SET
auditdb=# CREATE TABLE emp (
auditdb(#   empno      NUMBER(4) NOT NULL CONSTRAINT emp_pk PRIMARY KEY,
auditdb(#   ename      VARCHAR2(10),
auditdb(#   job        VARCHAR2(9),
auditdb(#   mgr        NUMBER(4),
auditdb(#   hiredate    DATE,
auditdb(#   sal        NUMBER(7,2) CONSTRAINT emp_sal_ck CHECK (sal >
0),
auditdb(#   comm       NUMBER(7,2),
auditdb(#   deptno     NUMBER(2) CONSTRAINT emp_ref_dept_fk
auditdb(#                   REFERENCES dept(deptno)
auditdb(# );
CREATE TABLE
auditdb=# CREATE VIEW salesemp AS
auditdb-#   SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job =
'SALESMAN';
CREATE VIEW
auditdb=# CREATE MATERIALIZED VIEW managers AS
auditdb-#   SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job =
'MANAGER';
SELECT 0
auditdb=# CREATE SEQUENCE next_empno START WITH 8000 INCREMENT BY 1;
CREATE SEQUENCE
auditdb=# GRANT ALL ON dept TO PUBLIC;
GRANT
auditdb=# GRANT ALL ON emp TO PUBLIC;
GRANT
```

The resulting audit log file contains the following.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```
2020-05-25 14:05:29.163 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,1,
"authentication",2020-05-25 14:05:29 IST,4/28,0,AUDIT,00000,"connection authorized:
user=adminuser database=auditdb",",,,,"client backend",",","connect"
```

```
2020-05-25 14:12:06.318 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,2,
"idle",2020-05-25 14:05:29 IST,4/34,0,AUDIT,00000,"statement: CREATE VIEW salesemp AS
SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job = 'SALESMAN';",",,,",
"psql","client backend","CREATE VIEW","create"
```

```
2020-05-25 14:13:26.657 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,3,
"idle",2020-05-25 14:05:29 IST,4/36,0,AUDIT,00000,"statement: CREATE MATERIALIZED
VIEW managers AS SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job =
'MANAGER';",",,,,"psql","client backend","CREATE MATERIALIZED VIEW","create"
```

```
2020-05-25 14:13:38.928 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,4,
"idle",2020-05-25 14:05:29 IST,4/37,0,AUDIT,00000,"statement: CREATE SEQUENCE next_
empno START WITH 8000 INCREMENT BY 1;","",,"psql","client backend","CREATE
SEQUENCE","create"
```

```
2020-05-25 14:13:51.434 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,5,
"idle",2020-05-25 14:05:29 IST,4/38,0,AUDIT,00000,"statement: GRANT ALL ON dept TO
PUBLIC;","",,"psql","client backend","GRANT","grant"
```

```
2020-05-25 14:14:03.737 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,6,
"idle",2020-05-25 14:05:29 IST,4/39,0,AUDIT,00000,"statement: GRANT ALL ON emp TO
PUBLIC;","",,"psql","client backend","GRANT","grant"
```

The `CREATE VIEW` and `CREATE MATERIALIZED VIEW` statements are audited. Note that the prior `CREATE TABLE` `emp` statement is not audited since none of the values `create`, `create table`, `ddl`, nor `all` are included in the `edb_audit_statement` parameter.

The `CREATE SEQUENCE` and `GRANT` statements are audited since those values are included in the `edb_audit_statement` parameter.

Data Manipulation Language Statements

This section describes the values that can be included in the `edb_audit_statement` parameter to audit DML statements.

The following general rules apply:

- If the `edb_audit_statement` parameter includes either `dml` or `all`, then all DML statements are audited.
- If the `edb_audit_statement` is set to `none`, then no DML statements are audited.
- Specific types of DML statements can be chosen for auditing by including a combination of values within the `edb_audit_statement` parameter.

Use the following syntax to specify an `edb_audit_statement` parameter value for `SQL INSERT`, `UPDATE`, `DELETE`, or `TRUNCATE` statements:

```
> { insert | update | delete | truncate }
```

Example

The following is an example where `edb_audit_connect` and `edb_audit_statement` are set with the following non-default values:

```
logging_collector = 'on'
edb_audit_connect = 'all'
edb_audit_statement = 'UPDATE, DELETE, error'
```

Thus, only SQL statements invoked by the `UPDATE` and `DELETE` commands are audited. All errors are also included in the audit log (even errors not related to the `UPDATE` and `DELETE` commands).

The database session that occurs is the following:

```
$ psql auditdb adminuser
Password for user adminuser:
psql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_connect;
edb_audit_connect
-----
all
(1 row)

auditdb=# SHOW edb_audit_statement;
edb_audit_statement
-----
UPDATE, DELETE, error
(1 row)

auditdb=# SET search_path TO edb;
SET
auditdb=# INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT 0 1
auditdb=# INSERT INTO dept VALUES (20,'RESEARCH','DALLAS');
INSERT 0 1
auditdb=# INSERT INTO dept VALUES (30,'SALES','CHICAGO');
INSERT 0 1
auditdb=# INSERT INTO dept VALUES (40,'OPERATIONS','BOSTON');
INSERT 0 1
auditdb=# INSERT INTO emp VALUES (7369,'SMITH','CLERK',7902,'17-DEC-80',800,NULL,20);
INSERT 0 1
auditdb=# INSERT INTO emp VALUES (7499,'ALLEN','SALESMAN',7698,'20-FEB-81',1600,300,30);
INSERT 0 1
auditdb=# INSERT INTO emp VALUES (7521,'WARD','SALESMAN',7698,'22-FEB-81',1250,500,30);
INSERT 0 1
.

auditdb=# INSERT INTO emp VALUES (7934,'MILLER','CLERK',7782,'23-JAN-82',1300,NULL,10);
INSERT 0 1
auditdb=# UPDATE dept SET loc = 'BEDFORD' WHERE deptno = 40;
UPDATE 1
```

```
auditdb=# SELECT * FROM dept;
deptno | dname   | loc
-----+-----+
 10 | ACCOUNTING | NEW YORK
 20 | RESEARCH   | DALLAS
 30 | SALES     | CHICAGO
 40 | OPERATIONS | BEDFORD
(4 rows)
```

```
auditdb=# DELETE FROM emp WHERE deptno = 10;
DELETE 3
auditdb=# TRUNCATE employee;
ERROR: relation "employee" does not exist
auditdb=# TRUNCATE emp;
TRUNCATE TABLE
auditdb=# \q
```

The resulting audit log file contains the following.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```
2020-05-25 15:17:16.392 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,1,
"authentication",2020-05-25 15:17:16 IST,5/3,0,AUDIT,00000,"connection authorized:
user=adminuser database=auditdb",,,,,,,,"client backend","",connect"
```

```
2020-05-25 15:19:18.066 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,2,
"idle",2020-05-25 15:17:16 IST,5/14,0,AUDIT,00000,"statement: UPDATE dept SET loc =
'BEDFORD' WHERE deptno = 40;,,,,,,,"psql","client backend","UPDATE","update"
```

```
2020-05-25 15:19:38.524 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,3,
"idle",2020-05-25 15:17:16 IST,5/16,0,AUDIT,00000,"statement: DELETE FROM emp WHERE
deptno = 10;,,,,,,,"psql","client backend","DELETE","delete"
```

```
2020-05-25 15:19:48.149 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,4,
"TRUNCATE TABLE",2020-05-25 15:17:16 IST,5/17,0,ERROR,42P01,"relation ""employee"""
does not exist",,,,,"TRUNCATE employee;","","psql","client backend","TRUNCATE TABLE",
"error"
```

The `UPDATE dept` and `DELETE FROM emp` statements are audited. Note that all of the prior `INSERT` statements are not audited since none of the values `insert`, `dml`, nor `all` are included in the `edb_audit_statement` parameter.

The `SELECT * FROM dept` statement is not audited as well since neither `select` nor `all` is included in the `edb_audit_statement` parameter.

Since `error` is specified in the `edb_audit_statement` parameter, but not the `truncate` value, the error on the `TRUNCATE employee` statement is logged in the audit file, but not the successful `TRUNCATE emp` statement.

11.3.5.3 Enabling Audit Logging

The following steps describe how to configure Advanced Server to log all connections, disconnections, DDL

statements, DCL statements, DML statements, and any statements resulting in an error.

1. Enable auditing by the setting the `edb_audit` parameter to `xml` or `csv`.
2. Set the file rotation day when the new file will be created by setting the parameter `edb_audit_rotation_day` to the desired value.
3. To audit all connections, set the parameter, `edb_audit_connect`, to `all`.
4. To audit all disconnections, set the parameter, `edb_audit_disconnect`, to `all`.
5. To audit DDL, DCL, DML and other statements, set the parameter, `edb_audit_statement` according to the instructions in [Selecting SQL Statements to Audit](#).

The setting of the `edb_audit_statement` parameter in the configuration file affects the entire database cluster.

The type of statements that are audited as controlled by the `edb_audit_statement` parameter can be further refined according to the database in use as well as the database role running the session:

- The `edb_audit_statement` parameter can be set as an attribute of a specified database with the `ALTER DATABASE dbname SET edb_audit_statement` command. This setting overrides the `edb_audit_statement` parameter in the configuration file for statements executed when connected to database `dbname`.
- The `edb_audit_statement` parameter can be set as an attribute of a specified role with the `ALTER ROLE rolename SET edb_audit_statement` command. This setting overrides the `edb_audit_statement` parameter in the configuration file as well as any setting assigned to the database by the previously described `ALTER DATABASE` command when the specified role is running the current session.
- The `edb_audit_statement` parameter can be set as an attribute of a specified role when using a specified database with the `ALTER ROLE rolename IN DATABASE dbname SET edb_audit_statement` command. This setting overrides the `edb_audit_statement` parameter in the configuration file as well as any setting assigned to the database by the previously described `ALTER DATABASE` command as well as any setting assigned to the role with the `ALTER ROLE` command without the `IN DATABASE` clause as previously described.

The following are examples of this technique.

The database cluster is established with `edb_audit_statement` set to `all` as shown in its `postgresql.conf` file:

```
logging_collector = 'on'
edb_audit_statement = 'all'    # Statement type to be audited:
                                # none, dml, insert, update, delete,
                                # truncate,
                                # select, error, rollback, ddl, create, drop,
                                # alter, grant, revoke, all
```

A database and role are established with the following settings for the `edb_audit_statement` parameter:

- Database `auditdb` with `ddl`, `insert`, `update`, and `delete`
- Role `admin` with `select` and `truncate`
- Role `admin` in database `auditdb` with `create table`, `insert`, and `update`

Creation and alteration of the database and role are shown by the following:

```
$ psql edb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

edb=# SHOW edb_audit_statement;
edb_audit_statement
-----
all
(1 row)

edb=# CREATE DATABASE auditdb;
CREATE DATABASE
```

```

edb=# ALTER DATABASE auditdb SET edb_audit_statement TO 'ddl, insert,
update, delete';
ALTER DATABASE
edb=# CREATE ROLE admin WITH LOGIN SUPERUSER PASSWORD 'password';
CREATE ROLE
edb=# ALTER ROLE admin SET edb_audit_statement TO 'select, truncate';
ALTER ROLE
edb=# ALTER ROLE admin IN DATABASE auditdb SET edb_audit_statement TO
'create table, insert, update';
ALTER ROLE

```

The following demonstrates the changes made and the resulting audit log file for three cases.

Case 1: Changes made in database `auditdb` by role `enterprisedb`. Only `ddl, insert, update`, and `delete` statements are audited:

```

$ psql auditdb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_statement;
  edb_audit_statement
-----
 ddl, insert, update, delete
(1 row)

auditdb=# CREATE TABLE audit_tbl (f1 INTEGER PRIMARY KEY, f2 TEXT);
CREATE TABLE
auditdb=# INSERT INTO audit_tbl VALUES (1, 'Row 1');
INSERT 0 1
auditdb=# UPDATE audit_tbl SET f2 = 'Row A' WHERE f1 = 1;
UPDATE 1
auditdb=# SELECT * FROM audit_tbl;           <== Should not be audited
   f1 |   f2
  ----+
  1 | Row A
(1 row)

auditdb=# TRUNCATE audit_tbl;           <== Should not be audited
TRUNCATE TABLE

```

The following audit log file shows entries only for the `CREATE TABLE`, `INSERT INTO audit_tbl`, and `UPDATE audit_tbl` statements. The `SELECT * FROM audit_tbl` and `TRUNCATE audit_tbl` statements were not audited.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-25 15:59:12.332 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,2,
"idle",2020-05-25 15:59:02 IST,5/7,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl
(f1 INTEGER PRIMARY KEY, f2 TEXT);",,,,,,,,"psql","client backend","CREATE TABLE","",",
"create"

```

```

2020-05-25 15:59:22.419 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,3,
"idle",2020-05-25 15:59:02 IST,5/8,0,AUDIT,00000,"statement: INSERT INTO audit_tbl
VALUES (1, 'Row 1');",,,,,,,,"psql","client backend","INSERT","",,"insert"

```

```
2020-05-25 15:59:32.180 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,4,
"idle",2020-05-25 15:59:02 IST,5/9,0,AUDIT,00000,"statement: UPDATE audit_tbl SET f2
= 'Row A' WHERE f1 = 1;","",,"pgsql","client backend","UPDATE","",,"update"
```

Case 2: Changes made in database `edb` by role `admin`. Only `select` and `truncate` statements are audited:

```
$ psql edb admin
Password for user admin:
psql.bin (13.0.0)
Type "help" for help.

edb=# SHOW edb_audit_statement;
edb_audit_statement
-----
select, truncate
(1 row)

edb=# CREATE TABLE edb_tbl (f1 INTEGER PRIMARY KEY, f2 TEXT); <== Should not
be audited
CREATE TABLE
edb=# INSERT INTO edb_tbl VALUES (1, 'Row 1'); <== Should not
be audited
INSERT 0 1
edb=# SELECT * FROM edb_tbl;
f1 | f2
----+---
 1 | Row 1
(1 row)

edb=# TRUNCATE edb_tbl;
TRUNCATE TABLE
```

Continuation of the audit log file now appears as follows. The last two entries representing the second case show only the `SELECT * FROM edb_tbl` and `TRUNCATE edb_tbl` statements. The `CREATE TABLE edb_tbl` and `INSERT INTO edb_tbl` statements were not audited.

```
2020-05-25 15:59:12.332 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,2,
"idle",2020-05-25 15:59:02 IST,5/7,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl
(f1 INTEGER PRIMARY KEY, f2 TEXT);","",,"pgsql","client backend","CREATE TABLE","",,
"create"

2020-05-25 15:59:22.419 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,3,
"idle",2020-05-25 15:59:02 IST,5/8,0,AUDIT,00000,"statement: INSERT INTO audit_tbl
VALUES (1, 'Row 1');","",,"pgsql","client backend","INSERT","",,"insert"

2020-05-25 15:59:32.180 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,4,
"idle",2020-05-25 15:59:02 IST,5/9,0,AUDIT,00000,"statement: UPDATE audit_tbl SET f2
= 'Row A' WHERE f1 = 1;","",,"pgsql","client backend","UPDATE","",,"update"
```

```
2020-05-25 16:25:37.509 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,2,"idle",
2020-05-25 16:24:54 IST,4/33,0,AUDIT,00000,"statement: SELECT * FROM edb_tbl;","",,
,"pgsql","client backend","SELECT","",,"select"
```

```
2020-05-25 16:25:46.131 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,3,"idle",
2020-05-25 16:24:54 IST,4/34,0,AUDIT,00000,"statement: TRUNCATE edb_tbl;","",,
```

```
"psql","client backend","TRUNCATE TABLE","","truncate"
```

Case 3: Changes made in database `auditdb` by role `admin`. Only `create table`, `insert`, and `update` statements are audited:

```
$ psql auditdb admin
Password for user admin:
psql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_statement;
      edb_audit_statement
-----
create table, insert, update
(1 row)

auditdb=# CREATE TABLE audit_tbl_2 (f1 INTEGER PRIMARY KEY, f2 TEXT);
CREATE TABLE
auditdb=# INSERT INTO audit_tbl_2 VALUES (1, 'Row 1');
INSERT 0 1
auditdb=# SELECT * FROM audit_tbl_2;           <== Should not be audited
   f1 |   f2
   ---+-----
   1 | Row 1
(1 row)

auditdb=# TRUNCATE audit_tbl_2;               <== Should not be audited
TRUNCATE TABLE
```

Continuation of the audit log file now appears as follows. The next to last two entries representing the third case show only `CREATE TABLE audit_tbl_2` and `INSERT INTO audit_tbl_2` statements. The `SELECT * FROM audit_tbl_2` and `TRUNCATE audit_tbl_2` statements were not audited.

```
2020-05-25 15:59:12.332 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,2,
"idle",2020-05-25 15:59:02 IST,5/7,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl
(f1 INTEGER PRIMARY KEY, f2 TEXT);",,,,,,,,"psql","client backend","CREATE TABLE","","
"create"
```

```
2020-05-25 15:59:22.419 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,3,
"idle",2020-05-25 15:59:02 IST,5/8,0,AUDIT,00000,"statement: INSERT INTO audit_tbl
VALUES (1, 'Row 1);",,,,,,,,"psql","client backend","INSERT","","
"insert"
```

```
2020-05-25 15:59:32.180 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,4,
"idle",2020-05-25 15:59:02 IST,5/9,0,AUDIT,00000,"statement: UPDATE audit_tbl SET f2
= 'Row A' WHERE f1 = 1;",,,,,,,,"psql","client backend","UPDATE","","
"update"
```

```
2020-05-25 16:25:37.509 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,2,"idle",
2020-05-25 16:24:54 IST,4/33,0,AUDIT,00000,"statement: SELECT * FROM edb_tbl;","",",
,"psql","client backend","SELECT","","
"select"
```

```
2020-05-25 16:25:46.131 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,3,"idle",
2020-05-25 16:24:54 IST,4/34,0,AUDIT,00000,"statement: TRUNCATE edb_tbl;","",",
,"psql","client backend","TRUNCATE TABLE","","
"truncate"
```

```
2020-05-25 17:30:59.057 IST,"admin","auditdb",122093,"[local]",5ecbb370.1dced,2,"idle
",2020-05-25 17:30:48 IST,5/11,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl_2 (f1
```

```
INTEGER PRIMARY KEY, f2 TEXT);",",,""",""psql",""client backend",""CREATE TABLE","","",
"create"
```

```
2020-05-25 17:31:08.866 IST,"admin","auditdb",122093,"[local]",5ecbb370.1dced,3,"idle
",2020-05-25 17:30:48 IST,5/12,0,AUDIT,00000,"statement: INSERT INTO audit_tbl_2
VALUES (1, 'Row 1');",",,""psql",""client backend",""INSERT","",""insert"
```

11.3.5.4 Audit Log File

The audit log file can be generated in either CSV or XML format depending upon the setting of the `edb_audit` configuration parameter.

The information in the audit log is based on the logging performed by PostgreSQL as described in the section “Using CSV-Format Log Output” within Section “Error Reporting and Logging” in the PostgreSQL core documentation, available at:

<https://www.postgresql.org/docs/current/static/runtime-config-logging.html>

The following table lists the fields in the order they appear in the XML audit log format. The table contains the following information:

- **Field.** Name of the field as shown in the sample table definition in the PostgreSQL documentation as previously referenced.
- **XML Element/Attribute.** For the XML format, name of the XML element and its attribute (if used), referencing the value. **Note:** `n/a` indicates that there is no XML representation for this field.
- **Data Type.** Data type of the field as given by the PostgreSQL sample table definition.
- **Description.** Description of the field. For certain fields, no output is generated in the audit log as those fields are not supported by auditing. Those fields are indicated by “Not supported”.

The fields that do not have any values for logging appear as consecutive commas (,,) in the CSV format.

| Field | XML Element/Attribute | Data Type | Description |
|------------------------|------------------------------|--------------------------|--|
| user_name | event/user | text | Database user who executed the statement. |
| database_name | event/database | text | Database in which the statement was executed. |
| process_id | event/process_id | integer | Operating system process ID in which the statement was executed. |
| connection_from | event/remote_host | text | Host and port location from where the statement was executed. |
| session_id | event/session_id | text | Session ID in which the statement was executed. |
| session_line_num | event/session_line_num | bigint | Order of the statement within the session. |
| process_status | event/process_status | text | Processing status. |
| session_start_time | event/session_start_time | timestamp with time zone | Date/time when the session was started. |
| log_time | event/log_time | timestamp with time zone | Log date/time of the statement. |
| virtual_transaction_id | event/virtual_transaction_id | text | Virtual transaction ID of the statement. |
| transaction_id | event/transaction_id | bigint | Regular transaction ID of the statement. |

| Field | XML Element/Attribute | Data Type | Description |
|--------------------|------------------------|-----------|--|
| type | event/type | text | Determines the audit event_type to identify messages in the log. |
| sql_state_code | event/sql_state_code | text | SQL state code returned for the statement. The sql_state_code is not logged when its value is 00000 for XML log format. |
| command_tag | event/command_tag | text | SQL command of the statement. |
| audit_tag | event/audit_tag | text | Value specified by the audit_tag parameter in the configuration file. |
| application_name | event/application_name | text | Name of the application from which the statement was executed. (for example, psql.bin). |
| backend_type | event/backend_type | text | The backend_type corresponds to what pg_stat_activity.bac kend_type shows and is added as a column to the csv log. |
| error_severity | error_severity | text | Statement severity. Values are AUDIT for audited statements and ERROR for any resulting error messages. |
| message | message | text | The SQL statement that was attempted for execution. |
| detail | detail | text | Error message detail. |
| hint | hint | text | Hint for error. |
| internal_query | internal_query | text | Internal query that led to the error, if any. |
| internal_query_pos | internal_query_pos | integer | Character count of the error position therein. |
| context | context | text | Error context. |
| query | query | text | User query that led to the error. For (For errors only) |
| query_pos | query_pos | integer | Character count of the error position therein. (For errors only) |
| location | location | text | Location of the error in the source code. The location field will be populated if log_error_verbosity is set to verbose. |

The following examples are generated in the CSV and XML formats.

The non-default audit settings in the `postgresql.conf` file are as follows:

```
logging_collector = 'on'
edb_audit = 'csv'
edb_audit_connect = 'all'
edb_audit_disconnect = 'all'
edb_audit_statement = 'ddl, dml, select, error'
edb_audit_tag = 'edbaudit'
```

The `edb_audit` parameter is changed to `xml` when generating the XML format.

The audited session is the following:

```
$ psql edb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

edb=# CREATE SCHEMA edb;
CREATE SCHEMA
```

```

edb=# SET search_path TO edb;
SET
edb=# CREATE TABLE dept (
edb(#   deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
edb(#   dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
edb(#   loc        VARCHAR2(13)
edb(# );
CREATE TABLE
edb=# INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT 0 1
edb=# UPDATE department SET loc = 'BOSTON' WHERE deptno = 10;
ERROR: relation "department" does not exist
LINE 1: UPDATE department SET loc = 'BOSTON' WHERE deptno = 10;
          ^
edb=# UPDATE dept SET loc = 'BOSTON' WHERE deptno = 10;
UPDATE 1
edb=# SELECT * FROM dept;
deptno | dname | loc
-----+-----+
 10 | ACCOUNTING | BOSTON
(1 row)

edb=# \q

```

CSV Audit Log File

The following is the CSV format of the audit log file.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-22 16:53:37.817 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,1,
"authentication",2020-05-22 16:53:37 IST,4/21,0,AUDIT,00000,"connection authorized:
user=enterprisedb database=edb",,,,,,,,"client backend","","edbaudit","connect"

```

```

2020-05-22 16:53:42.279 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,2,"idle
",2020-05-22 16:53:37 IST,4/23,0,AUDIT,00000,"statement: CREATE SCHEMA edb;",
"psql","client backend","CREATE SCHEMA","edbaudit","create"

```

```

2020-05-22 16:54:07.896 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,3,"idle
",2020-05-22 16:53:37 IST,4/25,0,AUDIT,00000,"statement: CREATE TABLE dept (
  deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
  dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
  loc        VARCHAR2(13)
);",,,,,,,,"psql","client backend","CREATE TABLE","edbaudit","create"

```

```

2020-05-22 16:54:20.500 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,4,"idle
",2020-05-22 16:53:37 IST,4/26,0,AUDIT,00000,"statement: INSERT INTO dept VALUES
(10,'ACCOUNTING','NEW YORK');",,,,,,,,"psql","client backend","INSERT","edbaudit",
"insert"

```

```

2020-05-22 16:54:34.821 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,5,"idle
",2020-05-22 16:53:37 IST,4/27,0,AUDIT,00000,"statement: UPDATE department SET loc =
'BOSTON' WHERE deptno = 10;",,,,,,,,"psql","client backend","UPDATE","edbaudit",
"update"

```

2020-05-22 16:54:34.821 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,6,
 "UPDATE",2020-05-22 16:53:37 IST,4/27,0,ERROR,42P01,"relation ""department"" does not
 exist",,,,,"UPDATE department SET loc = 'BOSTON' WHERE deptno = 10;",8,,"psql",
 "client backend","UPDATE","edbaudit","error"

2020-05-22 16:54:51.308 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,7,"idle
 ",2020-05-22 16:53:37 IST,4/28,0,AUDIT,00000,"statement: UPDATE dept SET loc = 'BOSTON
 ' WHERE deptno = 10;,,,,,,,"psql","client backend","UPDATE","edbaudit","update"

2020-05-22 16:55:00.774 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,8,"idle
 ",2020-05-22 16:53:37 IST,4/29,0,AUDIT,00000,"statement: SELECT * FROM dept;,,,,,,,"
 "psql","client backend","SELECT","edbaudit","select"

2020-05-22 16:55:06.548 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,9,"idle
 ",2020-05-22 16:53:37 IST,,0,AUDIT,00000,"disconnection: session time: 0:01:28.732
 user=enterprisedb database=edb host=[local],,,,"psql","client backend","","
 "edbaudit","disconnect"

XML Audit Log File

The following is the XML format of the audit log file. The output has been formatted for more clarity in the appearance in the example.

```
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"  

  session_id="5ec7ac4d.1735" session_line_num="1" process_status="authentication"  

  session_start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:11:17.806  

  IST" virtual_transaction_id="4/19" type="connect" audit_tag="edbaudit" backend_  

  type="client backend">  

  <error_severity>AUDIT</error_severity>  

  <message>connection authorized: user=enterprisedb database=edb</message>  

</event>  

<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"  

  session_id="5ec7ac4d.1735" session_line_num="2" process_status="idle" session_  

  start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:11:32.558 IST"  

  virtual_transaction_id="4/21" type="create" command_tag="CREATE SCHEMA" audit_  

  tag="edbaudit" application_name="psql" backend_type="client backend">  

  <error_severity>AUDIT</error_severity>  

  <message>statement: CREATE SCHEMA edb;</message>  

</event>  

<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"  

  session_id="5ec7ac4d.1735" session_line_num="3" process_status="idle" session_  

  start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:00.199 IST"  

  virtual_transaction_id="4/23" type="create" command_tag="CREATE TABLE" audit_  

  tag="edbaudit" application_name="psql" backend_type="client backend">  

  <error_severity>AUDIT</error_severity>  

  <message>statement: CREATE TABLE dept (  

    deptno NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,  

    dname VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,  

    loc VARCHAR2(13));  

  </message>  

</event>  

<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"  

  session_id="5ec7ac4d.1735" session_line_num="4" process_status="idle" session_  

  start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:10.992 IST"  

  virtual_transaction_id="4/24" type="insert" command_tag="INSERT" audit_tag=  

  "edbaudit" application_name="psql" backend_type="client backend">
```

```

<error_severity>AUDIT</error_severity>
<message>statement: INSERT INTO dept VALUES
          (10,&'ACCOUNTING'&,'NEW YORK');
</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
       session_id="5ec7ac4d.1735" session_line_num="5" process_status="idle" session_
       start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:21.764 IST"
       virtual_transaction_id="4/25" type="update" command_tag="UPDATE" audit_tag=
       "edbaudit" application_name="psql" backend_type="client backend">
    <error_severity>AUDIT</error_severity>
    <message>statement: UPDATE department SET loc = &'BOSTON'&;
               WHERE deptno = 10;
    </message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
       session_id="5ec7ac4d.1735" session_line_num="6" process_status="UPDATE"
       session_start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:21.765
       IST" virtual_transaction_id="4/25" type="error" sql_state_code="42P01" command_
       tag="UPDATE" audit_tag="edbaudit" application_name="psql" backend_type="client
       backend">
    <error_severity>ERROR</error_severity>
    <message>relation "department" does not exist</message>
    <query>UPDATE department SET loc = &'BOSTON'&; WHERE deptno =
       10;</query>
    <query_pos>8</query_pos>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
       session_id="5ec7ac4d.1735" session_line_num="7" process_status="idle" session_
       start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:34.878 IST"
       virtual_transaction_id="4/26" type="update" command_tag="UPDATE" audit_tag=
       "edbaudit" application_name="psql" backend_type="client backend">
    <error_severity>AUDIT</error_severity>
    <message>statement: UPDATE dept SET loc = &'BOSTON'&; WHERE
               deptno = 10;
    </message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
       session_id="5ec7ac4d.1735" session_line_num="8" process_status="idle" session_
       start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:45.471 IST"
       virtual_transaction_id="4/27" type="select" command_tag="SELECT" audit_tag=
       "edbaudit" application_name="psql" backend_type="client backend">
    <error_severity>AUDIT</error_severity>
    <message>statement: SELECT * FROM dept;</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
       session_id="5ec7ac4d.1735" session_line_num="9" process_status="idle" session_
       start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:53.048 IST"
       type="disconnect" audit_tag="edbaudit" application_name="psql" backend_type=
       "client backend">
    <error_severity>AUDIT</error_severity>
    <message>disconnection: session time: 0:01:35.243 user=enterprisedb
               database=edb host=[local]</message>
</event>
```

11.3.5.5 Using Error Codes to Filter Audit Logs

Advanced Server includes an extension that you can use to exclude log file entries that include a user-specified error code from the Advanced Server log files. To filter audit log entries, you must first enable the extension by modifying the `postgresql.conf` file, adding the following value to the values specified in the `shared_preload_libraries` parameter:

```
$libdir/edb_filter_log
```

Then, use the `edb_filter_log.errcodes` parameter to specify any error codes you wish to omit from the log files:

```
edb_filter_log.errcode = 'error_code'
```

Where `error_code` specifies one or more error codes that you wish to omit from the log file. Provide multiple error codes in a comma-delimited list.

For example, if `edb_filter_log` is enabled, and `edb_filter_log.errcode` is set to '`23505,23502,22012`', any log entries that return one of the following `SQLSTATE` errors:

`23505` (for violating a unique constraint)

`23502` (for violating a not-null constraint)

`22012` (for dividing by zero)

will be omitted from the log file.

For a complete list of the error codes supported by Advanced Server audit log filtering, see the core documentation at:

<https://www.postgresql.org/docs/current/static/errcodes-appendix.html>

11.3.5.6 Using Command Tags to Filter Audit Logs

Each entry in the log file except for those displaying an error message contains a *command tag*. A command tag is the SQL command executed for that particular log entry. The command tag makes it possible to use subsequent tools to scan the log file to find entries related to a particular SQL command.

The following is an example in XML form. The example has been formatted for easier review. The command tag is displayed as the `command_tag` attribute of the event element with values `CREATE ROLE`, `ALTER ROLE`, and `DROP ROLE` in the example.

```
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"
       session_id="5ecbc7e6.faea" session_line_num="2" process_status="idle" session_
       start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:21.147 IST"
       virtual_transaction_id="4/30" type="create" command_tag="CREATE ROLE"
       application_name="psql" backend_type="client backend">
    <error_severity>AUDIT</error_severity>
    <message>statement: CREATE ROLE newuser WITH LOGIN;</message>
</event>
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]">
```

```

session_id="5ecbc7e6.faea" session_line_num="3" process_status="idle" session_
start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:34.142 IST"
virtual_transaction_id="4/31" type="error" sql_state_code="42601" application_
name="psql" backend_type="client backend">
<error_severity>ERROR</error_severity>
<message>unrecognized role option &quot;super&quot;</message>
<query>ALTER ROLE newuser WITH SUPER USER;</query>
<query_pos>25</query_pos>
</event>
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"
session_id="5ecbc7e6.faea" session_line_num="4" process_status="idle" session_
start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:44.680 IST"
virtual_transaction_id="4/32" type="alter" command_tag="ALTER ROLE"
application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: ALTER ROLE newuser WITH SUPERUSER;</message>
</event>
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"
session_id="5ecbc7e6.faea" session_line_num="5" process_status="idle" session_
start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:58.173 IST"
virtual_transaction_id="4/33" type="drop" command_tag="DROP ROLE" application_
name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: DROP ROLE newuser;</message>
</event>

```

The following is the same example in CSV form. The command tag is the next to last column of each entry. In the listing, the column that appears empty (""), would be the value of the `edb_audit_tag` parameter if provided.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-25 19:09:32.105 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,2,
"idle",2020-05-25 19:09:07 IST,4/30,0,AUDIT,00000,"statement: CREATE ROLE newuser
WITH LOGIN;,,,,,,,"psql","client backend","CREATE ROLE","","","create"

```

```

2020-05-25 19:09:50.975 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,3,
"idle",2020-05-25 19:09:07 IST,4/31,0,ERROR,42601,"unrecognized role option ""super"""
",,,,"ALTER ROLE newuser WITH SUPER USER;";25,"psql","client backend","","","","error"

```

```

2020-05-25 19:10:04.128 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,4,
"idle",2020-05-25 19:09:07 IST,4/32,0,AUDIT,00000,"statement: ALTER ROLE newuser WITH
SUPERUSER;,,,,,,,"psql","client backend","ALTER ROLE","","","alter"

```

```

2020-05-25 19:10:15.959 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,5,
"idle",2020-05-25 19:09:07 IST,4/33,0,AUDIT,00000,"statement: DROP ROLE newuser;,,,,,
",,,,"psql","client backend","DROP ROLE","","","drop"

```

11.3.5.7 Redacting Passwords from Audit Logs

You can use the `edb_filter_log.redact_password_commands` extension to instruct the server to redact stored

passwords from the log file. Note that the module only recognizes the following syntax:

```
{CREATE|ALTER} {USER|ROLE|GROUP} identifier { [WITH] [ENCRYPTED] PASSWORD
'nonempty_string_literal' | IDENTIFIED BY { 'nonempty_string_literal' |
bareword } } [ REPLACE { 'nonempty_string_literal' | bareword } ]
```

When such a statement is logged by `log_statement`, the server will redact the old and new passwords to 'x'. For example, the command:

```
| ALTER USER carol PASSWORD '1safePWD' REPLACE 'old_pwd';
```

Will be added to log files as:

```
| statement: ALTER USER carol PASSWORD 'x' REPLACE 'x';
```

When a statement that includes a redacted password is logged, the server redacts the statement text. When the statement is logged as context for some other message, the server omits the statement from the context.

To enable password redaction, you must first enable the extension by modifying the `postgresql.conf` file, adding the following value to the values specified in the `shared_preload_libraries` parameter:

```
| $libdir/edb_filter_log
```

Then, set `edb_filter_log.redact_password_commands` to `true`:

```
| edb_filter_log.redact_password_commands = true
```

After modifying the `postgresql.conf` file, you must restart the server for the changes to take effect.

11.3.6 Unicode Collation Algorithm

The *Unicode Collation Algorithm* (UCA) is a specification (*Unicode Technical Report #10*) that defines a customizable method of collating and comparing Unicode data. *Collation* means how data is sorted as with a `SELECT ... ORDER BY` clause. *Comparison* is relevant for searches that use ranges with less than, greater than, or equal to operators.

Customizability is an important factor for various reasons such as the following.

- Unicode supports a vast number of languages. Letters that may be common to several languages may be expected to collate in different orders depending upon the language.
- Characters that appear with letters in certain languages such as accents or umlauts have an impact on the expected collation depending upon the language.
- In some languages, combinations of several consecutive characters should be treated as a single character with regards to its collation sequence.
- There may be certain preferences as to the collation of letters according to case. For example, should the lowercase form of a letter collate before the uppercase form of the same letter or vice versa.
- There may be preferences as to whether punctuation marks such as underscore characters, hyphens, or space characters should be considered in the collating sequence or should they simply be ignored as if they did not exist in the string.

Given all of these variations with the vast number of languages supported by Unicode, there is a necessity for a method to select the specific criteria for determining a collating sequence. This is what the Unicode Collation Algorithm defines.

Note

In addition, another advantage for using ICU collations (the implementation of the Unicode Collation Algorithm) is for performance. Sorting tasks, including B-tree index creation, can complete in less than half the time it takes with a non-ICU collation. The exact performance gain depends on your operating system version, the language of your text data, and other factors.

The following sections provide a brief, simplified explanation of the Unified Collation Algorithm concepts. As the algorithm and its usage are quite complex with numerous variations, refer to the official documents cited in these sections for complete details.

Basic Unicode Collation Algorithm Concepts

The official information for the Unicode Collation Algorithm is specified in *Unicode Technical Report #10*, which can be found on The Unicode Consortium website at:

<http://www.unicode.org/reports/tr10/>

The ICU – International Components for Unicode also provides much useful information. An explanation of the collation concepts can be found on their website located at:

<http://userguide.icu-project.org/collation/concepts>

The basic concept behind the Unicode Collation Algorithm is the use of multilevel comparison. This means that a number of levels are defined, which are listed as level 1 through level 5 in the following bullet points. Each level defines a type of comparison. Strings are first compared using the primary level, also called level 1.

If the order can be determined based on the primary level, then the algorithm is done. If the order cannot be determined based on the primary level, then the secondary level, level 2, is applied. If the order can be determined based on the secondary level, then the algorithm is done, otherwise the tertiary level is applied, and so on. There is typically, a final tie-breaking level to determine the order if it cannot be resolved by the prior levels.

- **Level 1 – Primary Level for Base Characters.** The order of basic characters such as letters and digits determines the difference such as A < B.
- **Level 2 – Secondary Level for Accents.** If there are no primary level differences, then the presence or absence of accents and other such characters determine the order such as a < á.
- **Level 3 – Tertiary Level for Case.** If there are no primary level or secondary level differences, then a difference in case determines the order such as a < A.
- **Level 4 – Quaternary Level for Punctuation.** If there are no primary, secondary, or tertiary level differences, then the presence or absence of white space characters, control characters, and punctuation determine the order such as -A < A.
- **Level 5 – Identical Level for Tie-Breaking.** If there are no primary, secondary, tertiary, or quaternary level differences, then some other difference such as the code point values determines the order.

International Components for Unicode

The Unicode Collation Algorithm is implemented by open source software provided by the *International Components for Unicode* (ICU). The software is a set of C/C++ and Java libraries.

When Advanced Server is used to create a collation that invokes the ICU components to produce the collation, the result is referred to as an *ICU collation*.

Locale Collations

When creating a collation for a locale, a predefined ICU short form name for the given locale is typically provided.

An *ICU short form* is a method of specifying *collation attributes*, which are the properties of a collation. [Collation Attributes](#) provides additional information on collation attributes.

There are predefined ICU short forms for locales. The ICU short form for a locale incorporates the collation attribute settings typically used for the given locale. This simplifies the collation creation process by eliminating the need to specify the entire list of collation attributes for that locale.

The system catalog `pg_catalog.pg_icu_collate_names` contains a list of the names of the ICU short forms for locales. The ICU short form name is listed in column `icu_short_form`.

```
edb=# SELECT icu_short_form, valid_locale FROM pg_icu_collate_names ORDER BY
valid_locale;
          icu_short_form | valid_locale
-----+-----
LAF      | af
LAR      | ar
LAS      | as
LAZ      | az
LBE      | be
LBG      | bg
LBN      | bn
LBS      | bs
LBS_ZCYRL | bs_Cyril
LROOT    | ca
LROOT    | chr
LCS      | cs
LCY      | cy
LDA      | da
LROOT    | de
LROOT    | dz
LEE      | ee
LEL      | el
LROOT    | en
LROOT    | en_US
LEN_RUS_VPOSIX | en_US_POSIX
LEO      | eo
LES      | es
LET      | et
LFA      | fa
LFA_RAF  | fa_AF
.
.
```

If needed, the default characteristics of an ICU short form for a given locale can be overridden by specifying the collation attributes to override that property. This is discussed in the next section.

Collation Attributes

When creating an ICU collation, the desired characteristics of the collation must be specified. As discussed in [Locale Collations](#), this can typically be done with an ICU short form for the desired locale. However, if more specific information is required, the specification of the collation properties can be done by using *collation attributes*.

Collation attributes define the rules of how characters are to be compared for determining the collation sequence of text strings. As Unicode covers a vast set of languages in numerous variations according to country, territory and culture, these collation attributes are quite complex.

For the complete, precise meaning and usage of collation attributes, see Section “Collator Naming Scheme” on

the ICU – International Components for Unicode website at:

| <http://userguide.icu-project.org/collation/concepts>

The following is a brief summary of the collation attributes and how they are specified using the ICU short form method

Each collation attribute is represented by an uppercase letter, which are listed in the following bullet points. The possible valid values for each attribute are given by codes shown within the parentheses. Some codes have general meanings for all attributes. **X** means to set the attribute off. **O** means to set the attribute on. **D** means to set the attribute to its default value.

- **A – Alternate (N, S, D)**. Handles treatment of *variable* characters such as white spaces, punctuation marks, and symbols. When set to non-ignorable (N), differences in variable characters are treated with the same importance as differences in letters. When set to shifted (S), then differences in variable characters are of minor importance (that is, the variable character is ignored when comparing base characters).
- **C – Case First (X, L, U, D)**. Controls whether a lowercase letter sorts before the same uppercase letter (L), or the uppercase letter sorts before the same lowercase letter (U). Off (X) is typically specified when lowercase first (L) is desired.
- **E – Case Level (X, O, D)**. Set in combination with the Strength attribute, the Case Level attribute is used when accents are to be ignored, but not case.
- **F – French Collation (X, O, D)**. When set to on, secondary differences (presence of accents) are sorted from the back of the string as done in the French Canadian locale.
- **H – Hiragana Quaternary (X, O, D)**. Introduces an additional level to distinguish between the Hiragana and Katakana characters for compatibility with the JIS X 4061 collation of Japanese character strings.
- **N – Normalization Checking (X, O, D)**. Controls whether or not text is thoroughly normalized for comparison. Normalization deals with the issue of canonical equivalence of text whereby different code point sequences represent the same character, which then present issues when sorting or comparing such characters. Languages such as Arabic, ancient Greek, Hebrew, Hindi, Thai, or Vietnamese should be used with Normalization Checking set to on.
- **S – Strength (1, 2, 3, 4, I, D)**. Maximum collation level used for comparison. Influences whether accents or case are taken into account when collating or comparing strings. Each number represents a level. A setting of I represents identical strength (that is, level 5).
- **T – Variable Top (hexadecimal digits)**. Applicable only when the Alternate attribute is not set to non-ignorable (N). The hexadecimal digits specify the highest character sequence that is to be considered ignorable. For example, if white space is to be ignorable, but visible variable characters are not to be ignorable, then Variable Top set to 0020 would be specified along with the Alternate attribute set to S and the Strength attribute set to 3. (The space character is hexadecimal 0020. Other non-visible variable characters such as backspace, tab, line feed, carriage return, etc. have values less than 0020. All visible punctuation marks have values greater than 0020.)

A set of collation attributes and their values is represented by a text string consisting of the collation attribute letter concatenated with the desired attribute value. Each attribute/value pair is joined to the next pair with an underscore character as shown by the following example.

| **AN_CX_EX_FX_HX_NO_S3**

Collation attributes can be specified along with a locale's ICU short form name to override those default attribute settings of the locale.

The following is an example where the ICU short form named **LROOT** is modified with a number of other collation attribute/value pairs.

| **AN_CX_EX_LROOT_NO_S3**

In the preceding example, the Alternate attribute (**A**) is set to non-ignorable (**N**). The Case First attribute (**C**) is set to off (**X**). The Case Level attribute (**E**) is set to off (**X**). The Normalization attribute (**N**) is set to on (**O**). The Strength attribute (**S**) is set to the tertiary level **3**. **LROOT** is the ICU short form to which these other attributes are applying modifications.

Using a Collation

A newly defined ICU collation can be used anywhere the `COLLATION "collation_name"` clause can be used in a SQL command such as in the column specifications of the `CREATE TABLE` command or appended to an expression in the `ORDER BY` clause of a `SELECT` command.

The following are some examples of the creation and usage of ICU collations based on the English language in the United States (`en_US.UTF8`).

In these examples, ICU collations are created with the following characteristics.

Collation `icu_collate_lowercase` forces the lowercase form of a letter to sort before its uppercase counterpart (`CL`).

Collation `icu_collate_uppercase` forces the uppercase form of a letter to sort before its lowercase counterpart (`CU`).

Collation `icu_collate_ignore_punct` causes variable characters (white space and punctuation marks) to be ignored during sorting (`AS`).

Collation `icu_collate_ignore_white_sp` causes white space and other non-visible variable characters to be ignored during sorting, but visible variable characters (punctuation marks) are not ignored (`AS`, `T0020`).

```
CREATE COLLATION icu_collate_lowercase (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AN_CL_EX_NX_LROOT'
);

CREATE COLLATION icu_collate_uppercase (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AN_CU_EX_NX_LROOT'
);

CREATE COLLATION icu_collate_ignore_punct (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AS_CX_EX_NX_LROOT_L3'
);

CREATE COLLATION icu_collate_ignore_white_sp (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AS_CX_EX_NX_LROOT_L3_T0020'
);
```

Note

When creating collations, ICU may generate notice and warning messages when attributes are given to modify the `LROOT` collation.

The following `psql` command lists the collations.

```
edb=# \dO
          List of collations
 Schema |      Name       | Collate | Ctype |      ICU
-----+-----+-----+-----+
enterprisedb | icu_collate_ignore_punct | en_US.UTF8 | en_US.UTF8 |
AS_CX_EX_NX_LROOT_L3
enterprisedb | icu_collate_ignore_white_sp | en_US.UTF8 | en_US.UTF8 |
AS_CX_EX_NX_LROOT_L3_T0020
enterprisedb | icu_collate_lowercase | en_US.UTF8 | en_US.UTF8 |
```

```

AN_CL_EX_NX_LROOT
enterprisedb | icu_collate_uppercase      | en_US.UTF8 | en_US.UTF8 |
AN_CU_EX_NX_LROOT
(4 rows)

```

The following table is created and populated.

```

CREATE TABLE collate_tbl (
    id      INTEGER,
    c2      VARCHAR(2)
);

INSERT INTO collate_tbl VALUES (1, 'A');
INSERT INTO collate_tbl VALUES (2, 'B');
INSERT INTO collate_tbl VALUES (3, 'C');
INSERT INTO collate_tbl VALUES (4, 'a');
INSERT INTO collate_tbl VALUES (5, 'b');
INSERT INTO collate_tbl VALUES (6, 'c');
INSERT INTO collate_tbl VALUES (7, '1');
INSERT INTO collate_tbl VALUES (8, '2');
INSERT INTO collate_tbl VALUES (9, '.B');
INSERT INTO collate_tbl VALUES (10, '-B');
INSERT INTO collate_tbl VALUES (11, ' B');

```

The following query sorts on column `c2` using the default collation. Note that variable characters (white space and punctuation marks) with id column values of `9`, `10`, and `11` are ignored and sort with the letter `B`.

```

edb=# SELECT * FROM collate_tbl ORDER BY c2;
id | c2
----+-
7  | 1
8  | 2
4  | a
1  | A
5  | b
2  | B
11 | B
10 | -B
9   | .B
6   | c
3   | C
(11 rows)

```

The following query sorts on column `c2` using collation `icu_collate_lowercase`, which forces the lowercase form of a letter to sort before the uppercase form of the same base letter. Also note that the `AN` attribute forces variable characters to be included in the sort order at the same level when comparing base characters so rows with `id` values of `9`, `10`, and `11` appear at the beginning of the sort list before all letters and numbers.

```

edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE "icu_collate_lowercase";
id | c2
----+-
11 | B
10 | -B
9  | .B
7  | 1
8  | 2

```

```
4 | a
1 | A
5 | b
2 | B
6 | c
3 | C
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_uppercase`, which forces the uppercase form of a letter to sort before the lowercase form of the same base letter.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE "icu_collate_uppercase";
id | c2
----+-
11 | B
10 | -B
9 | .B
7 | 1
8 | 2
1 | A
4 | a
2 | B
5 | b
3 | C
6 | c
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_ignore_punct`, which causes variable characters to be ignored so rows with id values of `9`, `10`, and `11` sort with the letter `B` as that is the character immediately following the ignored variable character.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE
"icu_collate_ignore_punct";
id | c2
----+-
7 | 1
8 | 2
4 | a
1 | A
5 | b
11 | B
2 | B
9 | .B
10 | -B
6 | c
3 | C
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_ignore_white_sp`. The `AS` and `T0020` attributes of the collation cause variable characters with code points less than or equal to hexadecimal `0020` to be ignored while variable characters with code points greater than hexadecimal `0020` are included in the sort.

The row with id value of `11`, which starts with a space character (hexadecimal `0020`) sorts with the letter `B`. The rows with id values of `9` and `10`, which start with visible punctuation marks greater than hexadecimal `0020`, appear at the beginning of the sort list as these particular variable characters are included in the sort order at the same level when comparing base characters.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE
"icu_collate_ignore_white_sp";
id | c2
----+-
10 | -B
9 | .B
7 | 1
8 | 2
4 | a
1 | A
5 | b
11 | B
2 | B
6 | c
3 | C
(11 rows)
```

11.4 Configuration Parameters

This section describes the database server configuration parameters of Advanced Server. These parameters control various aspects of the database server's behavior and environment such as data file and log file locations, connection, authentication, and security settings, resource allocation and consumption, archiving and replication settings, error logging and statistics gathering, optimization and performance tuning, locale and formatting settings, and so on.

Configuration parameters that apply only to Advanced Server are noted in the [Summary of Configuration Parameters](#) section.

Additional information about configuration parameters can be found in the PostgreSQL Core Documentation available at:

<https://www.postgresql.org/docs/current/static/runtime-config.html>

11.4.1 Setting Configuration Parameters

This section provides an overview of how configuration parameters are specified and set.

Each configuration parameter is set using a `name/value` pair. Parameter names are case-insensitive. The parameter name is typically separated from its value by an optional equals sign (`=`).

The following is an example of some configuration parameter settings in the `postgresql.conf` file:

```
### This is a comment
log_connections = yes
log_destination = 'syslog'
search_path = "$user", public
shared_buffers = 128MB
```

Parameter values are specified as one of five types:

- **Boolean.** Acceptable values can be written as `on`, `off`, `true`, `false`, `yes`, `no`, `1`, `0`, or any unambiguous prefix of these.
- **Integer.** Number without a fractional part.
- **Floating Point.** Number with an optional fractional part separated by a decimal point.
- **String.** Text value. Enclose in single quotes if the value is not a simple identifier or number (that is, the value contains special characters such as spaces or other punctuation marks).
- **Enum.** Specific set of string values. The allowed values can be found in the system view `pg_settings.enumvals`. Enum values are case-insensitive.

Some settings specify a memory or time value. Each of these has an implicit unit, which is kilobytes, blocks (typically 8 kilobytes), milliseconds, seconds, or minutes. Default units can be found by referencing the system view `pg_settings.unit`. A different unit can be specified explicitly.

Valid memory units are `KB` (kilobytes), `MB` (megabytes), and `GB` (gigabytes). Valid time units are `ms` (milliseconds), `s` (seconds), `min` (minutes), `h` (hours), and `d` (days). The multiplier for memory units is 1024.

The configuration parameter settings can be established in a number of different ways:

- There is a number of parameter settings that are established when the Advanced Server database product is built. These are read-only parameters, and their values cannot be changed. There are also a couple of parameters that are permanently set for each database when the database is created. These parameters are read-only as well and cannot be subsequently changed for the database.
- The initial settings for almost all configurable parameters across the entire database cluster are listed in the configuration file, `postgresql.conf`. These settings are put into effect upon database server start or restart. Some of these initial parameter settings can be overridden as discussed in the following bullet points. All configuration parameters have built-in default settings that are in effect if not explicitly overridden.
- Configuration parameters in the `postgresql.conf` file are overridden when the same parameters are included in the `postgresql.auto.conf` file. The `ALTER SYSTEM` command is used to manage the configuration parameters in the `postgresql.auto.conf` file.
- Parameter settings can be modified in the configuration file while the database server is running. If the configuration file is then reloaded (meaning a SIGHUP signal is issued), for certain parameter types, the changed parameters settings immediately take effect. For some of these parameter types, the new settings are available in a currently running session immediately after the reload. For other of these parameter types, a new session must be started to use the new settings. And yet for other parameter types, modified settings do not take effect until the database server is stopped and restarted. See the following section of the PostgreSQL Core Documentation for information on how to reload the configuration file:

<https://www.postgresql.org/docs/current/config-setting.html>

- The SQL commands `ALTER DATABASE`, `ALTER ROLE`, or `ALTER ROLE IN DATABASE` can be used to modify certain parameter settings. The modified parameter settings take effect for new sessions after the command is executed. `ALTER DATABASE` affects new sessions connecting to the specified database. `ALTER ROLE` affects new sessions started by the specified role. `ALTER ROLE IN DATABASE` affects new sessions started by the specified role connecting to the specified database. Parameter settings established by these SQL commands remain in effect indefinitely, across database server restarts, overriding settings established by the methods discussed in the second and third bullet points. Parameter settings established using the `ALTER DATABASE`, `ALTER ROLE`, or `ALTER ROLE IN DATABASE` commands can only be changed by: a) re-issuing these commands with a different parameter value, or
- b) issuing these commands using either of the `SET parameter TO DEFAULT` clause or the `RESET parameter` clause. These clauses change the parameter back to using the setting established by the methods set forth in the prior bullet points. See the “SQL Commands” section of Chapter VI “Reference” in the PostgreSQL Core Documentation for the exact syntax of these SQL commands:

<https://www.postgresql.org/docs/current/sql-commands.html>

- Changes can be made for certain parameter settings for the duration of individual sessions using the

PGOPTIONS environment variable or by using the **SET** command within the **EDB-PSQL** or **PSQL** command line terminal programs. Parameter settings made in this manner override settings established using any of the methods described by the second, third, and fourth bullet points, but only for the duration of the session.

11.4.2 Summary of Configuration Parameters

This section contains a summary table listing all Advanced Server configuration parameters along with a number of key attributes of the parameters.

These attributes are described by the following columns of the summary table:

- **Parameter.** Configuration parameter name.
- **Scope of Effect.** Scope of effect of the configuration parameter setting.
 - **Cluster** – Setting affects the entire database cluster (that is, all databases managed by the database server instance).
 - **Database** – Setting can vary by database and is established when the database is created. Applies to a small number of parameters related to locale settings.
 - **Session** – Setting can vary down to the granularity of individual sessions.
 In other words, different settings can be made for the following entities whereby the latter settings in this list override prior ones: a) the entire database cluster, b) specific databases in the database cluster, c) specific roles, d) specific roles when connected to specific databases, e) a specific session.
- **When Takes Effect.** When a changed parameter setting takes effect.
 - **Preset** – Established when the Advanced Server product is built or a particular database is created. This is a read-only parameter and cannot be changed.
 - **Restart** – Database server must be restarted.
 - **Reload** – Configuration file must be reloaded (or the database server can be restarted).
 - **Immediate** – Immediately effective in a session if the **PGOPTIONS** environment variable or the **SET** command is used to change the setting in the current session. Effective in new sessions if **ALTER DATABASE**, **ALTER ROLE**, or **ALTER ROLE IN DATABASE** commands are used to change the setting.
- **Authorized User.** Type of operating system account or database role that must be used to put the parameter setting into effect.
 - **EPAS service account** – EDB Postgres Advanced Server service account (**enterprisedb** for an installation compatible with Oracle databases, **postgres** for a PostgreSQL compatible mode installation).
 - **Superuser** – Database role with superuser privileges.
 - **User** – Any database role with permissions on the affected database object (the database or role to be altered with the **ALTER** command).
 - **n/a** – Parameter setting cannot be changed by any user.
- **Description.** Brief description of the configuration parameter.
- **EPAS Only.** ‘X’ – Configuration parameter is applicable to EDB Postgres Advanced Server only. No entry in this column indicates the configuration parameter applies to PostgreSQL as well.

Note

There are a number of parameters that should never be altered. These are designated as “**Note: For internal use only**” in the Description column.

The following table shows the summary of configuration parameters:

| Parameter | Scope of Effect | When Takes Effect | Authorized User | Description | EPAS Only |
|--------------------------------|-----------------|-------------------|----------------------|---|-----------|
| custom_variable_classes | Cluster | Reload | EPAS service account | Deprecated in Advanced Server 9.2. | X |
| db_dialect | Session | Immediate | User | Sets the precedence of built-in namespaces. | X |
| dbms_alert.max_alerts | Cluster | Restart | EPAS service account | Sets maximum number of alerts. | X |
| dbms_pipe.total_message_buffer | Cluster | Restart | EPAS service account | Specifies the total size of the buffer used for the DBMS_PIPE package. | X |
| default_heap_fillfactor | Session | Immediate | User | Create new tables with this heap fillfactor by default. | X |
| default_with_rowids | Session | Immediate | User | Create new tables with ROWID support (ROWIDs with indexes) by default. | X |
| edb_audit | Cluster | Reload | EPAS service account | Enable EDB Auditing to create audit reports in XML or CSV format. | X |
| edb_audit_connect | Cluster | Reload | EPAS service account | Audits each successful connection. | X |
| edb_audit_destination | Cluster | Reload | EPAS service account | Sets edb_audit_directory or syslog as the destination directory for audit files. The syslog setting is only valid for a Linux system. | X |
| edb_audit_directory | Cluster | Reload | EPAS service account | Sets the destination directory for audit files. | X |
| edb_audit_disconnect | Cluster | Reload | EPAS service account | Audits end of a session. | X |
| edb_audit_filename | Cluster | Reload | EPAS service account | Sets the file name pattern for audit files. | X |
| edb_audit_rotation_day | Cluster | Reload | EPAS service account | Automatic rotation of log files based on day of week. | X |
| edb_audit_rotation_seconds | Cluster | Reload | EPAS service account | Automatic log file rotation will occur after N seconds. | X |
| edb_audit_rotation_size | Cluster | Reload | EPAS service account | Automatic log file rotation will occur after N Megabytes. | X |
| edb_audit_statement | Cluster | Reload | EPAS service account | Sets the type of statements to audit. | X |

| | | | | | |
|---|---------|-----------|----------------------|--|---|
| edb_audit_tag | Session | Immediate | User | Specify a tag to be included in the audit log. | X |
| edb_connectby_order | Session | Immediate | User | Sort results of CONNECT BY queries with no ORDER BY to depth-first order. Note: For internal use only. | X |
| edb_data_redaction | Session | Immediate | User | Enable data redaction. | X |
| edb_dynatune | Cluster | Restart | EPAS service account | Sets the edb utilization percentage. | X |
| edb_dynatune_profile | Cluster | Restart | EPAS service account | Sets the workload profile for dynatune. | X |
| edb_enable_pruning | Session | Immediate | User | Enables the planner to early-prune partitioned tables. | X |
| edb_log_every_bulk_value | Session | Immediate | Superuser | Sets the statements logged for bulk processing. | X |
| edb_max_resource_groups | Cluster | Restart | EPAS service account | Specifies the maximum number of resource groups for simultaneous use. | X |
| edb_max_spins_per_delay | Cluster | Restart | EPAS service account | Specifies the number of times a session will spin while waiting for a lock. | X |
| edb_redwood_date | Session | Immediate | User | Determines whether DATE should behave like a TIMESTAMP or not. | X |
| edb_redwood_greatest_least | Session | Immediate | User | Determines how GREATEST and LEAST functions should handle NULL parameters. | X |
| edb_redwood_raw_names | Session | Immediate | User | Return the unmodified name stored in the PostgreSQL system catalogs from Redwood interfaces. | X |
| edb_redwood_strings | Session | Immediate | User | Treat NULL as an empty string when concatenated with a text value. | X |
| edb_resource_group | Session | Immediate | User | Specifies the resource group to be used by the current process. | X |
| edb_sql_protect.enabled | Cluster | Reload | EPAS service account | Defines whether SQL/Protect should track queries or not. | X |
| edb_sql_protect.level | Cluster | Reload | EPAS service account | Defines the behavior of SQL/Protect when an event is found. | X |
| edb_sql_protect.max_protected_relations | Cluster | Restart | EPAS service account | Sets the maximum number of relations protected by SQL/Protect per role. | X |

| | | | | | |
|-------------------------------------|----------|-----------|----------------------|--|---|
| edb_sql_protect.max_protected_roles | Cluster | Restart | EPAS service account | Sets the maximum number of roles protected by SQL/Protect. | X |
| edb_sql_protect.max_queries_to_save | Cluster | Restart | EPAS service account | Sets the maximum number of offending queries to save by SQL/Protect. | X |
| edb_stmt_level_tx | Session | Immediate | User | Allows continuing on errors instead of requiring a transaction abort. | X |
| edb_wait_states.directory | Cluster | Restart | EPAS service account | The EDB wait states logs are stored in this directory. | X |
| edb_wait_states.retention_period | Cluster | Reload | EPAS service account | EDB wait state log files will be automatically deleted after retention period. | X |
| edb_wait_states.sampling_interval | Cluster | Reload | EPAS service account | The interval between two EDB wait state sampling cycles. | X |
| edbldr.empty_csv_field | Session | Immediate | Superuser | Specifies how EDB Loader handles empty strings. | X |
| enable_hints | Session | Immediate | User | Enable optimizer hints in SQL statements. | X |
| icu_short_form | Database | Preset | n/a | Shows the ICU collation order configuration. | X |
| index_advisor.enabled | Session | Immediate | User | Enable Index Advisor plugin. | X |
| nls_length_semantics | Session | Immediate | Superuser | Sets the semantics to use for char, varchar, varchar2 and long columns. | X |
| odbc_lib_path | Cluster | Restart | EPAS service account | Sets the path for ODBC library. | X |
| optimizer_mode | Session | Immediate | User | Default optimizer mode. | X |
| oracle_home | Cluster | Restart | EPAS service account | Sets the path for the Oracle home directory. | X |
| pg_prewarm.autoprewarm | Cluster | Restart | EPAS service account | Enables the autoprewarm background worker. | X |
| pg_prewarm.autoprewarm_interval | Cluster | Reload | EPAS service account | Sets the minimum number of seconds after which autoprewarm dumps shared buffers. | X |
| qreplace_function | Session | Immediate | Superuser | The function to be used by Query Replace feature. Note: For internal use only. | X |

| | | | | | |
|------------------------------|---------|-----------|-----------|--|---|
| query_rewrite_enabled | Session | Immediate | User | Child table scans will be skipped if their constraints guarantee that no rows match the query. | X |
| query_rewrite_integrity | Session | Immediate | Superuser | Sets the degree to which query rewriting must be enforced. | X |
| timed_statistics | Session | Immediate | User | Enables the recording of timed statistics. | X |
| trace_hints | Session | Immediate | User | Emit debug info about hints being honored. | X |
| util_encode.uudecode_redwood | Session | Immediate | User | Allows decoding of Oracle-created uuencoded data. | X |
| util_file.umask | Session | Immediate | User | Umask used for files created through the UTL_FILE package. | X |

11.4.3 Configuration Parameters by Functionality

This section provides more detail for certain groups of configuration parameters.

The section heading for each parameter is followed by a list of attributes:

- **Parameter Type.** Type of values the parameter can accept. See [Setting Configuration Parameters](#) section for a discussion of parameter type values.
- **Default Value.** Default setting if a value is not explicitly set in the configuration file.
- **Range.** Permitted range of values.
- **Minimum Scope of Effect.** Smallest scope for which a distinct setting can be made. Generally, the minimal scope of a distinct setting is either the entire `cluster` (the setting is the same for all databases and sessions thereof, in the cluster), or `per session` (the setting may vary by role, database, or individual session). (This attribute has the same meaning as the `Scope of Effect` column in the table of [Summary of Configuration Parameters](#) section).
- **When Value Changes Take Effect.** Least invasive action required to activate a change to a parameter's value. All parameter setting changes made in the configuration file can be put into effect with a restart of the database server; however certain parameters require a database server `restart`. Some parameter setting changes can be put into effect with a `reload` of the configuration file without stopping the database server. Finally, other parameter setting changes can be put into effect with some client side action whose result is `immediate`. (This attribute has the same meaning as the `When Takes Effect` column in the table of [Summary of Configuration Parameters](#) section).
- **Required Authorization to Activate.** The type of user authorization to activate a change to a parameter's setting. If a database server restart or a configuration file reload is required, then the user must be a EPAS service account (`enterprisedb` or `postgres` depending upon the Advanced Server compatibility installation mode). (This attribute has the same meaning as the `Authorized User` column in the table of [Summary of Configuration Parameters](#) section).

11.4.3.1 Top Performance Related Parameters

This section discusses the configuration parameters that have the most immediate impact on performance.

11.4.3.1.1 shared_buffers

Parameter Type: Integer

Default Value: 32MB

Range: 128kB to system dependent

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the amount of memory the database server uses for shared memory buffers. The default is typically 32 megabytes ([32MB](#)), but might be less if your kernel settings will not support it (as determined during [initdb](#)). This setting must be at least 128 kilobytes. (Non-default values of [BLCKSZ](#) change the minimum.) However, settings significantly higher than the minimum are usually needed for good performance.

If you have a dedicated database server with 1GB or more of RAM, a reasonable starting value for [shared_buffers](#) is 25% of the memory in your system. There are some workloads where even large settings for [shared_buffers](#) are effective, but because Advanced Server also relies on the operating system cache, it is unlikely that an allocation of more than 40% of RAM to [shared_buffers](#) will work better than a smaller amount.

On systems with less than 1GB of RAM, a smaller percentage of RAM is appropriate, so as to leave adequate space for the operating system (15% of memory is more typical in these situations). Also, on Windows, large values for [shared_buffers](#) aren't as effective. You may find better results keeping the setting relatively low and using the operating system cache more instead. The useful range for [shared_buffers](#) on Windows systems is generally from 64MB to 512MB.

Increasing this parameter might cause Advanced Server to request more System V shared memory than your operating system's default configuration allows. See the section [Shared Memory and Semaphores](#) in the [PostgreSQL Core Documentation](#) for information on how to adjust those parameters, if necessary.

11.4.3.1.2 work_mem

Parameter Type: Integer

Default Value: 1MB

Range: 64kB to 2097151kB

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Specifies the amount of memory to be used by internal sort operations and hash tables before writing to temporary disk files. The value defaults to one megabyte (**1MB**). Note that for a complex query, several sort or hash operations might be running in parallel; each operation will be allowed to use as much memory as this value specifies before it starts to write data into temporary files. Also, several running sessions could be doing such operations concurrently. Therefore, the total memory used could be many times the value of **work_mem**; it is necessary to keep this fact in mind when choosing the value. Sort operations are used for **ORDER BY**, **DISTINCT**, and merge joins. Hash tables are used in hash joins, hash-based aggregation, and hash-based processing of **IN** subqueries.

Reasonable values are typically between 4MB and 64MB, depending on the size of your machine, how many concurrent connections you expect (determined by **max_connections**), and the complexity of your queries.

11.4.3.1.3 maintenance_work_mem

Parameter Type: Integer

Default Value: 16MB

Range: 1024kB to 2097151kB

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Specifies the maximum amount of memory to be used by maintenance operations, such as **VACUUM**, **CREATE INDEX**, and **ALTER TABLE ADD FOREIGN KEY**. It defaults to 16 megabytes (**16MB**). Since only one of these operations can be executed at a time by a database session, and an installation normally doesn't have many of them running concurrently, it's safe to set this value significantly larger than **work_mem**. Larger settings might improve performance for vacuuming and for restoring database dumps.

Note that when autovacuum runs, up to **autovacuum_max_workers** times this memory may be allocated, so be careful not to set the default value too high.

A good rule of thumb is to set this to about 5% of system memory, but not more than about 512MB. Larger values won't necessarily improve performance.

11.4.3.1.4 wal_buffers

Parameter Type: Integer

Default Value: 64kB

Range: 32kB to system dependent

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

The amount of memory used in shared memory for WAL data. The default is 64 kilobytes ([64kB](#)). The setting need only be large enough to hold the amount of WAL data generated by one typical transaction, since the data is written out to disk at every transaction commit.

Increasing this parameter might cause Advanced Server to request more System V shared memory than your operating system's default configuration allows. See the section "Shared Memory and Semaphores" in the *PostgreSQL Core Documentation* for information on how to adjust those parameters, if necessary.

Although even this very small setting does not always cause a problem, there are situations where it can result in extra `fsync` calls, and degrade overall system throughput. Increasing this value to 1MB or so can alleviate this problem. On very busy systems, an even higher value may be needed, up to a maximum of about 16MB. Like [shared_buffers](#), this parameter increases Advanced Server's initial shared memory allocation, so if increasing it causes an Advanced Server start failure, you will need to increase the operating system limit.

11.4.3.1.5 `checkpoint_segments`

Now deprecated; this parameter is not supported by server versions 9.5 or later.

11.4.3.1.6 `checkpoint_completion_target`

Parameter Type: Floating point

Default Value: 0.5

Range: 0 to 1

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the target of checkpoint completion as a fraction of total time between checkpoints. This spreads out the checkpoint writes while the system starts working towards the next checkpoint.

The default of 0.5 means aim to finish the checkpoint writes when 50% of the next checkpoint is ready. A value of 0.9 means aim to finish the checkpoint writes when 90% of the next checkpoint is done, thus throttling the checkpoint writes over a larger amount of time and avoiding spikes of performance bottlenecks.

11.4.3.1.7 checkpoint_timeout

Parameter Type: Integer

Default Value: 5min

Range: 30s to 3600s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Maximum time between automatic WAL checkpoints, in seconds. The default is five minutes (5min). Increasing this parameter can increase the amount of time needed for crash recovery.

Increasing `checkpoint_timeout` to a larger value, such as 15 minutes, can reduce the I/O load on your system, especially when using large values for `shared_buffers`.

The downside of making the aforementioned adjustments to the checkpoint parameters is that your system will use a modest amount of additional disk space, and will take longer to recover in the event of a crash. However, for most users, this is a small price to pay for a significant performance improvement.

11.4.3.1.8 max_wal_size

Parameter Type: Integer

Default Value: 1 GB

Range: 2 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

`max_wal_size` specifies the maximum size that the WAL will reach between automatic WAL checkpoints. This is a soft limit; WAL size can exceed `max_wal_size` under special circumstances (when under a heavy load, a failing archive_command, or a high `wal_keep_segments` setting).

Increasing this parameter can increase the amount of time needed for crash recovery. This parameter can only be set in the `postgresql.conf` file or on the server command line.

11.4.3.1.9 min_wal_size

Parameter Type: Integer

Default Value: 80 MB**Range:** 2 to 2147483647**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

If WAL disk usage stays below the value specified by `min_wal_size`, old WAL files are recycled for future use at a checkpoint, rather than removed. This ensures that enough WAL space is reserved to handle spikes in WAL usage (like when running large batch jobs). This parameter can only be set in the `postgresql.conf` file or on the server command line.

11.4.3.1.10 bgwriter_delay

Parameter Type: Integer**Default Value:** 200ms**Range:** 10ms to 10000ms**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

Specifies the delay between activity rounds for the background writer. In each round the writer issues writes for some number of dirty buffers (controllable by the `bgwriter_lru_maxpages` and `bgwriter_lru_multiplier` parameters). It then sleeps for `bgwriter_delay` milliseconds, and repeats.

The default value is 200 milliseconds (200ms). Note that on many systems, the effective resolution of sleep delays is 10 milliseconds; setting `bgwriter_delay` to a value that is not a multiple of 10 might have the same results as setting it to the next higher multiple of 10.

Typically, when tuning `bgwriter_delay`, it should be reduced from its default value. This parameter is rarely increased, except perhaps to save on power consumption on a system with very low utilization.

11.4.3.1.11 seq_page_cost

Parameter Type: Floating point**Default Value:** 1**Range:** 0 to 1.79769e+308**Minimum Scope of Effect:** Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the planner's estimate of the cost of a disk page fetch that is part of a series of sequential fetches. The default is 1.0. This value can be overridden for a particular tablespace by setting the tablespace parameter of the same name. (Refer to the [ALTER TABLESPACE](#) command in the *PostgreSQL Core Documentation*.)

The default value assumes very little caching, so it's frequently a good idea to reduce it. Even if your database is significantly larger than physical memory, you might want to try setting this parameter to less than 1 (rather than its default value of 1) to see whether you get better query plans that way. If your database fits entirely within memory, you can lower this value much more, perhaps to 0.1.

11.4.3.1.12 random_page_cost

Parameter Type: Floating point

Default Value: 4

Range: 0 to 1.79769e+308

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the planner's estimate of the cost of a non-sequentially-fetched disk page. The default is 4.0. This value can be overridden for a particular tablespace by setting the tablespace parameter of the same name. (Refer to the [ALTER TABLESPACE](#) command in the *PostgreSQL Core Documentation*.)

Reducing this value relative to [seq_page_cost](#) will cause the system to prefer index scans; raising it will make index scans look relatively more expensive. You can raise or lower both values together to change the importance of disk I/O costs relative to CPU costs, which are described by the [cpu_tuple_cost](#) and [cpu_index_tuple_cost](#) parameters.

The default value assumes very little caching, so it's frequently a good idea to reduce it. Even if your database is significantly larger than physical memory, you might want to try setting this parameter to 2 (rather than its default of 4) to see whether you get better query plans that way. If your database fits entirely within memory, you can lower this value much more, perhaps to 0.1.

Although the system will let you do so, never set [random_page_cost](#) less than [seq_page_cost](#). However, setting them equal (or very close to equal) makes sense if the database fits mostly or entirely within memory, since in that case there is no penalty for touching pages out of sequence. Also, in a heavily-cached database you should lower both values relative to the CPU parameters, since the cost of fetching a page already in RAM is much smaller than it would normally be.

11.4.3.1.13 effective_cache_size

Parameter Type: Integer

Default Value: 128MB**Range:** 8kB to 17179869176kB**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Sets the planner's assumption about the effective size of the disk cache that is available to a single query. This is factored into estimates of the cost of using an index; a higher value makes it more likely index scans will be used, a lower value makes it more likely sequential scans will be used. When setting this parameter you should consider both Advanced Server's shared buffers and the portion of the kernel's disk cache that will be used for Advanced Server data files. Also, take into account the expected number of concurrent queries on different tables, since they will have to share the available space. This parameter has no effect on the size of shared memory allocated by Advanced Server, nor does it reserve kernel disk cache; it is used only for estimation purposes. The default is 128 megabytes (128MB).

If this parameter is set too low, the planner may decide not to use an index even when it would be beneficial to do so. Setting `effective_cache_size` to 50% of physical memory is a normal, conservative setting. A more aggressive setting would be approximately 75% of physical memory.

11.4.3.1.14 `synchronous_commit`

Parameter Type: Boolean**Default Value:** true**Range:** {true | false}**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Specifies whether transaction commit will wait for WAL records to be written to disk before the command returns a `success` indication to the client. The default, and safe, setting is on. When off, there can be a delay between when success is reported to the client and when the transaction is really guaranteed to be safe against a server crash. (The maximum delay is three times `wal_writer_delay`.)

Unlike `fsync`, setting this parameter to off does not create any risk of database inconsistency: an operating system or database crash might result in some recent allegedly-committed transactions being lost, but the database state will be just the same as if those transactions had been aborted cleanly.

So, turning `synchronous_commit` off can be a useful alternative when performance is more important than exact certainty about the durability of a transaction. See the section [Asynchronous Commit](#) in the *PostgreSQL Core Documentation* for information.

This parameter can be changed at any time; the behavior for any one transaction is determined by the setting in effect when it commits. It is therefore possible, and useful, to have some transactions commit synchronously and others asynchronously. For example, to make a single multistatement transaction commit asynchronously when the default is the opposite, issue `SET LOCAL synchronous_commit TO OFF` within the transaction.

11.4.3.1.15 edb_max_spins_per_delay

Parameter Type: Integer

Default Value: 1000

Range: 10 to 1000

Minimum Scope of Effect: Per cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Use `edb_max_spins_per_delay` to specify the maximum number of times that a session will spin while waiting for a spin-lock. If a lock is not acquired, the session will sleep. If you do not specify an alternative value for `edb_max_spins_per_delay`, the server will enforce the default value of 1000.

This may be useful for systems that use NUMA (non-uniform memory access) architecture.

11.4.3.1.16 pg_prewarm.autoprewarm

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

This parameter controls whether or not the database server should run `autoprewarm`, which is a background worker process that automatically dumps shared buffers to disk before a shutdown. It then `prewams` the shared buffers the next time the server is started, meaning it loads blocks from the disk back into the buffer pool.

The advantage is that it shortens the warm up times after the server has been restarted by loading the data that has been dumped to disk before shutdown.

If `pg_prewarm.autoprewarm` is set to on, the `autoprewarm` worker is enabled. If the parameter is set to off, `autoprewarm` is disabled. The parameter is on by default.

Before `autoprewarm` can be used, you must add `$libdir/pg_prewarm` to the libraries listed in the `shared_preload_libraries` configuration parameter of the `postgresql.conf` file as shown by the following example:

```
shared_preload_libraries =
'$libdir/dbms_pipe,$libdir/edb_gen,$libdir/dbms_aq,$libdir/pg_prewarm'
```

After modifying the `shared_preload_libraries` parameter, restart the database server after which the `autoprewarm` background worker is launched immediately after the server has reached a consistent state.

The `autoprewarm` process will start loading blocks that were previously recorded in `$PGDATA/autoprewarm` blocks until there is no free buffer space left in the buffer pool. In this manner, any new blocks that were loaded either by the recovery process or by the querying clients, are not replaced.

Once the `autoprewarm` process has finished loading buffers from disk, it will periodically dump shared buffers to disk at the interval specified by the `pg_prewarm.autoprewarm_interval` parameter. See the `pg_prewarm.autoprewarm_interval` for information on the `autoprewarm` background worker. Upon the next server restart, the `autoprewarm` process will prewarm shared buffers with the blocks that were last dumped to disk.

11.4.3.1.17 pg_prewarm.autoprewarm_interval

Parameter Type: Integer

Default Value: 300s

Range: 0s to 2147483s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

This is the minimum number of seconds after which the `autoprewarm` background worker dumps shared buffers to disk. The default is 300 seconds. If set to 0, shared buffers are not dumped at regular intervals, but only when the server is shut down.

See the `pg_prewarm.autoprewarm` for information on the `autoprewarm` background worker.

11.4.3.2 Resource Usage / Memory

The configuration parameters in this section control resource usage pertaining to memory.

edb_dynatune

Parameter Type: Integer

Default Value: 0

Range: 0 to 100

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Determines how much of the host system's resources are to be used by the database server based upon the host

machine's total available resources and the intended usage of the host machine.

When Advanced Server is initially installed, the `edb_dynatune` parameter is set in accordance with the selected usage of the host machine on which it was installed (i.e., development machine, mixed use machine, or dedicated server). For most purposes, there is no need for the database administrator to adjust the various configuration parameters in the `postgresql.conf` file in order to improve performance.

The `edb_dynatune` parameter can be set to any integer value between 0 and 100, inclusive. A value of 0, turns off the dynamic tuning feature thereby leaving the database server resource usage totally under the control of the other configuration parameters in the `postgresql.conf` file.

A low non-zero, value (e.g., 1 - 33) dedicates the least amount of the host machine's resources to the database server. This setting would be used for a development machine where many other applications are being used.

A value in the range of 34 - 66 dedicates a moderate amount of resources to the database server. This setting might be used for a dedicated application server that may have a fixed number of other applications running on the same machine as Advanced Server.

The highest values (e.g., 67 - 100) dedicate most of the server's resources to the database server. This setting would be used for a host machine that is totally dedicated to running Advanced Server.

Once a value of `edb_dynatune` is selected, database server performance can be further fine-tuned by adjusting the other configuration parameters in the `postgresql.conf` file. Any adjusted setting overrides the corresponding value chosen by `edb_dynatune`. You can change the value of a parameter by un-commenting the configuration parameter, specifying the desired value, and restarting the database server.

`edb_dynatune_profile`

Parameter Type: Enum

Default Value: oltp

Range: {oltp | reporting | mixed}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

This parameter is used to control tuning aspects based upon the expected workload profile on the database server.

The following are the possible values:

- `oltp`. Recommended when the database server is processing heavy online transaction processing workloads.
- `reporting`. Recommended for database servers used for heavy data reporting.
- `mixed`. Recommended for servers that provide a mix of transaction processing and data reporting.

11.4.3.3 Resource Usage / EDB Resource Manager

The configuration parameters in this section control resource usage through EDB Resource Manager.

`edb_max_resource_groups`

Parameter Type: Integer

Default Value: 16

Range: 0 to 65536

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

This parameter controls the maximum number of resource groups that can be used simultaneously by EDB Resource Manager. More resource groups can be created than the value specified by `edb_max_resource_groups`, however, the number of resource groups in active use by processes in these groups cannot exceed this value.

Parameter `edb_max_resource_groups` should be set comfortably larger than the number of groups you expect to maintain so as not to run out.

`edb_resource_group`

Parameter Type: String

Default Value: none

Range: n/a

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Set the `edb_resource_group` parameter to the name of the resource group to which the current session is to be controlled by EDB Resource Manager according to the group's resource type settings.

If the parameter is not set, then the current session does not utilize EDB Resource Manager.

11.4.3.4 Query Tuning

This section describes the configuration parameters used for optimizer hints.

`enable_hints`

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Optimizer hints embedded in SQL commands are utilized when `enable_hints` is on. Optimizer hints are ignored when this parameter is off.

11.4.3.5 Query Tuning / Planner Method Configuration

This section describes the configuration parameters used for planner method configuration.

edb_enable_pruning

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When set to `TRUE`, `edb_enable_pruning` allows the query planner to early-prune partitioned tables. **Early-pruning** means that the query planner can `prune` (i.e., ignore) partitions that would not be searched in a query `before` generating query plans. This helps improve performance time as it eliminates the generation of query plans of partitions that would not be searched.

Conversely, `late-pruning` means that the query planner prunes partitions `after` generating query plans for each partition. (The `constraint_exclusion` configuration parameter controls late-pruning.)

The ability to early-prune depends upon the nature of the query in the `WHERE` clause. Early-pruning can be utilized in only simple queries with constraints of the type `WHERE column = literal` (e.g., `WHERE deptno = 10`).

Early-pruning is not used for more complex queries such as `WHERE column = expression` (e.g., `WHERE deptno = 10 + 5`).

11.4.3.6 Reporting and Logging / What to Log

The configuration parameters in this section control reporting and logging.

trace_hints

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Use with the optimizer hints feature to provide more detailed information regarding whether or not a hint was used by the planner. Set the `client_min_messages` and `trace_hints` configuration parameters as follows:

```
SET client_min_messages TO info;
SET trace_hints TO true;
```

The `SELECT` command with the `NO_INDEX` hint shown below illustrates the additional information produced when the aforementioned configuration parameters are set.

```
EXPLAIN SELECT /*+ NO_INDEX(accounts accounts_pkey) */ * FROM accounts
WHERE aid = 100;
```

```
INFO: [HINTS] Index Scan of [accounts].[accounts_pkey] rejected because
of NO_INDEX hint.
```

```
INFO: [HINTS] Bitmap Heap Scan of [accounts].[accounts_pkey] rejected
because of NO_INDEX hint.
```

QUERY PLAN

```
Seq Scan on accounts (cost=0.00..14461.10 rows=1 width=97)
  Filter: (aid = 100)
(2 rows)
```

edb_log_every_bulk_value

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Superuser

Bulk processing logs the resulting statements into both the Advanced Server log file and the EDB Audit log file. However, logging each and every statement in bulk processing is costly. This can be controlled by the `edb_log_every_bulk_value` configuration parameter. When set to `true`, each and every statement in bulk

processing is logged. During bulk execution, when `edb_log_every_bulk_value` is set to `false`, a log message is recorded once per bulk processing along with the number of rows processed. In addition, the duration is emitted once per bulk processing. Default is set to `false`.

11.4.3.7 Auditing Settings

This section describes configuration parameters used by the Advanced Server database auditing feature.

11.4.3.7.1 `edb_audit`

Parameter Type: Enum

Default Value: none

Range: {none | csv | xml}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Enables or disables database auditing. The values `xml` or `csv` will enable database auditing. These values represent the file format in which auditing information will be captured. `none` will disable database auditing and is also the default.

11.4.3.7.2 `edb_audit_directory`

Parameter Type: String

Default Value: `edb_audit`

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the directory where the audit log files will be created. The path of the directory can be absolute or relative to the Advanced Server `data` directory.

11.4.3.7.3 edb_audit_filename

Parameter Type: String

Default Value: audit-%Y%m%d_%H%M%S

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the file name of the audit file where the auditing information will be stored. The default file name will be `audit-%Y%m%d_%H%M%S`. The escape sequences, `%Y`, `%m` etc., will be replaced by the appropriate current values according to the system date and time.

11.4.3.7.4 edb_audit_rotation_day

Parameter Type: String

Default Value: every

Range: {none | every | sun | mon | tue | wed | thu | fri | sat} ...

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the day of the week on which to rotate the audit files. Valid values are `sun`, `mon`, `tue`, `wed`, `thu`, `fri`, `sat`, `every`, and `none`. To disable rotation, set the value to `none`. To rotate the file every day, set the `edb_audit_rotation_day` value to `every`. To rotate the file on a specific day of the week, set the value to the desired day of the week.

11.4.3.7.5 edb_audit_rotation_size

Parameter Type: Integer

Default Value: 0MB

Range: 0MB to 5000MB

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies a file size threshold in megabytes when file rotation will be forced to occur. The default value is 0MB. If the parameter is commented out or set to 0, rotation of the file on a size basis will not occur.

11.4.3.7.6 `edb_audit_rotation_seconds`

Parameter Type: Integer

Default Value: 0s

Range: 0s to 2147483647s

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies the rotation time in seconds when a new log file should be created. To disable this feature, set this parameter to 0.

11.4.3.7.7 `edb_audit_connect`

Parameter Type: Enum

Default Value: failed

Range: {none | failed | all}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Enables auditing of database connection attempts by users. To disable auditing of all connection attempts, set `edb_audit_connect` to `none`. To audit all failed connection attempts, set the value to `failed`. To audit all connection attempts, set the value to `all`.

11.4.3.7.8 `edb_audit_disconnect`

Parameter Type: Enum**Default Value:** none**Range:** {none | all}**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

Enables auditing of database disconnections by connected users. To enable auditing of disconnections, set the value to `all`. To disable, set the value to `none`.

11.4.3.7.9 `edb_audit_statement`

Parameter Type: String**Default Value:** ddl, error

Range: {none| ddl | dml | insert | update | delete | truncate | select | error | create | drop | alter | grant | revoke | rollback | all} ...

Minimum Scope of Effect: Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

This configuration parameter is used to specify auditing of different categories of SQL statements as well as those statements related to specific SQL commands. To log errors, set the parameter value to `error`. To audit all DDL statements such as `CREATE TABLE`, `ALTER TABLE`, etc., set the parameter value to `ddl`. To audit specific types of DDL statements, the parameter values can include those specific SQL commands (`create`, `drop`, or `alter`). In addition, the object type may be specified following the command such as `create table`, `create view`, `drop role`, etc. All modification statements such as `INSERT`, `UPDATE`, `DELETE` or `TRUNCATE` can be audited by setting `edb_audit_statement` to `dml`. To audit specific types of DML statements, the parameter values can include the specific SQL commands, `insert`, `update`, `delete`, or `truncate`. Include parameter values `select`, `grant`, `revoke`, or `rollback` to audit statements regarding those SQL commands. Setting the value to `all` will audit every statement while `none` disables this feature.

11.4.3.7.10 `edb_audit_tag`

Parameter Type: String**Default Value:** none**Minimum Scope of Effect:** Session**When Value Changes Take Effect:** Immediate

Required Authorization to Activate: User

Use `edb_audit_tag` to specify a string value that will be included in the audit log when the `edb_audit` parameter is set to `csv` or `xml`.

11.4.3.7.11 `edb_audit_destination`

Parameter Type: Enum

Default Value: file

Range: {file | syslog}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Specifies whether the audit log information is to be recorded in the directory as given by the `edb_audit_directory` parameter or to the directory and file managed by the `syslog` process. Set to `file` to use the directory specified by `edb_audit_directory` (the default setting).

Set to `syslog` to use the syslog process and its location as configured in the `/etc/syslog.conf` file. The `syslog` setting is valid only for Advanced Server running on a Linux host, and is not supported on Windows systems.

Note: In recent Linux versions, `syslog` has been replaced by `rsyslog` and the configuration file is in `/etc/rsyslog.conf`.

11.4.3.7.12 `edb_log_every_bulk_value`

For information on `edb_log_every_bulk_value`, see the [edb_log_every_bulk_value](#).

11.4.3.8 Client Connection Defaults / Locale and Formatting

This section describes configuration parameters affecting locale and formatting.

`icu_short_form`

Parameter Type: String

Default Value: none

Range: n/a**Minimum Scope of Effect:** Database**When Value Changes Take Effect:** n/a**Required Authorization to Activate:** n/a

The configuration parameter `icu_short_form` is a parameter containing the default ICU short form name assigned to a database or to the Advanced Server instance. See [Unicode Collation Algorithm](#) for general information about the ICU short form and the Unicode Collation Algorithm.

This configuration parameter is set either when the `CREATE DATABASE` command is used with the `ICU_SHORT_FORM` parameter in which case the specified short form name is set and appears in the `icu_short_form` configuration parameter when connected to this database, or when an Advanced Server instance is created with the `initdb` command used with the `--icu_short_form` option in which case the specified short form name is set and appears in the `icu_short_form` configuration parameter when connected to a database in that Advanced Server instance, and the database does not override it with its own `ICU_SHORT_FORM` parameter with a different short form.

Once established in the manner described, the `icu_short_form` configuration parameter cannot be changed.

11.4.3.9 Client Connection Defaults / Statement Behavior

This section describes configuration parameters affecting statement behavior.

`default_heap_fillfactor`

Parameter Type: Integer**Default Value:** 100**Range:** 10 to 100**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Sets the fillfactor for a table when the `FILLFACTOR` storage parameter is omitted from a `CREATE TABLE` command.

The fillfactor for a table is a percentage between 10 and 100. 100 (complete packing) is the default. When a smaller `fillfactor` is specified, `INSERT` operations pack table pages only to the indicated percentage; the remaining space on each page is reserved for updating rows on that page. This gives `UPDATE` a chance to place the updated copy of a row on the same page as the original, which is more efficient than placing it on a different page. For a table whose entries are never updated, complete packing is the best choice, but in heavily updated tables smaller fillfactors are appropriate.

`edb_data_redaction`

Parameter Type: Boolean**Default Value:** true**Range:** {true | false}**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Data redaction is the support of policies to limit the exposure of certain sensitive data to certain users by altering the displayed information.

The default setting is `TRUE` so the data redaction is applied to all users except for superusers and the table owner:

- Superusers and table owner bypass data redaction.
- All other users get the redaction policy applied and see the reformatted data.

If the parameter is disabled by setting it to `FALSE`, then the following occurs:

- Superusers and table owner still bypass data redaction.
- All other users will get an error.

For information on data redaction, see the *EDB Postgres Advanced Server Security Features Guide*.

11.4.3.10 Client Connection Defaults / Other Defaults

The parameters in this section set other miscellaneous client connection defaults.

`oracle_home`

Parameter Type: String**Default Value:** none**Range:** n/a**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Restart**Required Authorization to Activate:** EPAS service account

Before creating an Oracle Call Interface (OCI) database link to an Oracle server, you must direct Advanced Server to the correct Oracle home directory. Set the `LD_LIBRARY_PATH` environment variable on Linux (or `PATH` on Windows) to the `lib` directory of the Oracle client installation directory.

For Windows only, you can instead set the value of the `oracle_home` configuration parameter in the `postgresql.conf` file. The value specified in the `oracle_home` configuration parameter will override the Windows `PATH` environment variable.

The `LD_LIBRARY_PATH` environment variable on Linux (PATH environment variable or `oracle_home` configuration parameter on Windows) must be set properly each time you start Advanced Server.

For Windows only: To set the `oracle_home` configuration parameter in the `postgresql.conf` file, edit the file, adding the following line:

```
oracle_home = 'lib_directory'
```

Substitute the name of the Windows directory that contains `oci.dll` for `lib_directory`.

After setting the `oracle_home` configuration parameter, you must restart the server for the changes to take effect. Restart the server from the Windows Services console.

odbc_lib_path

Parameter Type: String

Default Value: none

Range: n/a

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

If you will be using an ODBC driver manager, and if it is installed in a non-standard location, you specify the location by setting the `odbc_lib_path` configuration parameter in the `postgresql.conf` file:

```
odbc_lib_path = 'complete_path_to_libodbc.so'
```

The configuration file must include the complete pathname to the driver manager shared library (typically `libodbc.so`).

11.4.3.11 Compatibility Options

The configuration parameters described in this section control various database compatibility features.

edb_redwood_date

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When `DATE` appears as the data type of a column in the commands, it is translated to `TIMESTAMP` at the time the table definition is stored in the database if the configuration parameter `edb_redwood_date` is set to `TRUE`. Thus, a time component will also be stored in the column along with the date.

If `edb_redwood_date` is set to `FALSE` the column's data type in a `CREATE TABLE` or `ALTER TABLE` command remains as a native PostgreSQL `DATE` data type and is stored as such in the database. The PostgreSQL `DATE` data type stores only the date without a time component in the column.

Regardless of the setting of `edb_redwood_date`, when `DATE` appears as a data type in any other context such as the data type of a variable in an SPL declaration section, or the data type of a formal parameter in an SPL procedure or SPL function, or the return type of an SPL function, it is always internally translated to a `TIMESTAMP` and thus, can handle a time component if present.

`edb_redwood_greatest_least`

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

The `GREATEST` function returns the parameter with the greatest value from its list of parameters. The `LEAST` function returns the parameter with the least value from its list of parameters.

When `edb_redwood_greatest_least` is set to `TRUE`, the `GREATEST` and `LEAST` functions return null when at least one of the parameters is null.

```
SETedb_redwood_greatest_leastTOon;
```

```
SELECTGREATEST(1,2,NULL,3);
```

```
greatest
```

```
-----
```

```
(1row)
```

When `edb_redwood_greatest_least` is set to `FALSE`, null parameters are ignored except when all parameters are null in which case null is returned by the functions.

```
SETedb_redwood_greatest_leastTOoff;
```

```
SELECTGREATEST(1,2,NULL,3);
```

```
greatest
```

```
-----
```

```
3
```

```
(1row)
```

```
SELECTGREATEST(NULL,NULL,NULL);
```

```
greatest
-----
```

```
(1 row)
```

edb_redwood_raw_names

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When `edb_redwood_raw_names` is set to its default value of `FALSE`, database object names such as table names, column names, trigger names, program names, user names, etc. appear in uppercase letters when viewed from Redwood catalogs (that is, system catalogs prefixed by `ALL_`, `DBA_`, or `USER_`). In addition, quotation marks enclose names that were created with enclosing quotation marks.

When `edb_redwood_raw_names` is set to `TRUE`, the database object names are displayed exactly as they are stored in the PostgreSQL system catalogs when viewed from the Redwood catalogs. Thus, names created without enclosing quotation marks appear in lowercase as expected in PostgreSQL. Names created with enclosing quotation marks appear exactly as they were created, but without the quotation marks.

For example, the following user name is created, and then a session is started with that user.

```
CREATE USER reduser IDENTIFIED BY password;
edb=# \c - reduser
Password for user reduser:
You are now connected to database "edb" as user "reduser".
```

When connected to the database as reduser, the following tables are created.

```
CREATE TABLE all_lower (col INTEGER);
CREATE TABLE ALL_UPPER (COL INTEGER);
CREATE TABLE "Mixed_Case" ("Col" INTEGER);
```

When viewed from the Redwood catalog, `USER_TABLES`, with `edb_redwood_raw_names` set to the default value `FALSE`, the names appear in uppercase except for the `Mixed_Case` name, which appears as created and also with enclosing quotation marks.

```
edb=> SELECT * FROM USER_TABLES;
schema_name | table_name | tablespace_name | status | temporary
-----+-----+-----+-----+
REDUSER    | ALL_LOWER   |           | VALID  | N
REDUSER    | ALL_UPPER   |           | VALID  | N
REDUSER    | "Mixed_Case" |           | VALID  | N
(3 rows)
```

When viewed with `edb_redwood_raw_names` set to `TRUE`, the names appear in lowercase except for the `Mixed_Case` name, which appears as created, but now without the enclosing quotation marks.

```
edb=> SET edb_redwood_raw_names TO true;
SET
edb=> SELECT * FROM USER_TABLES;
schema_name | table_name | tablespace_name | status | temporary
-----+-----+-----+
reduser | all_lower | VALID | N
reduser | all_upper | VALID | N
reduser | Mixed_Case | VALID | N
(3 rows)
```

These names now match the case when viewed from the PostgreSQL `pg_tables` catalog.

```
edb=> SELECT schemaname, tablename, tableowner FROM pg_tables WHERE
tableowner = 'reduser';
schemaname | tablename | tableowner
-----+-----+-----
reduser | all_lower | reduser
reduser | all_upper | reduser
reduser | Mixed_Case | reduser
(3 rows)
```

edb_redwood_strings

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

If the `edb_redwood_strings` parameter is set to `TRUE`, when a string is concatenated with a null variable or null column, the result is the original string. If `edb_redwood_strings` is set to `FALSE`, the native PostgreSQL behavior is maintained, which is the concatenation of a string with a null variable or null column gives a null result.

The following example illustrates the difference.

The sample application contains a table of employees. This table has a column named `comm` that is null for most employees. The following query is run with `edb_redwood_string` set to `FALSE`. The concatenation of a null column with non-empty strings produces a final result of null, so only employees that have a commission appear in the query result. The output line for all other employees is null.

```
SET edb_redwood_strings TO off;

SELECT RPAD(ename,10) || ' ' || TO_CHAR(sal,'99,999.99') || ' ' ||
TO_CHAR(comm,'99,999.99') "EMPLOYEE COMPENSATION" FROM emp;

EMPLOYEE COMPENSATION
-----
ALLEN    1,600.00  300.00
WARD     1,250.00  500.00
```

```
MARTIN    1,250.00 1,400.00
```

```
TURNER    1,500.00 .00
```

(14 rows)

The following is the same query executed when `edb_redwood_strings` is set to `TRUE`. Here, the value of a null column is treated as an empty string. The concatenation of an empty string with a non-empty string produces the non-empty string.

```
SET edb_redwood_strings TO on;

SELECT RPAD(ename,10) || '' || TO_CHAR(sal,'99,999.99') || '' ||
TO_CHAR(comm,'99,999.99') "EMPLOYEE COMPENSATION" FROM emp;
```

EMPLOYEE COMPENSATION

```
SMITH 800.00
ALLEN 1,600.00 300.00
WARD 1,250.00 500.00
JONES 2,975.00
MARTIN 1,250.00 1,400.00
BLAKE 2,850.00
CLARK 2,450.00
SCOTT 3,000.00
KING 5,000.00
TURNER 1,500.00 .00
ADAMS 1,100.00
JAMES 950.00
FORD 3,000.00
MILLER 1,300.00
```

(14 rows)

`edb_stmt_level_tx`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

The term `statement level transaction isolation` describes the behavior whereby when a runtime error occurs in a SQL command, all the updates on the database caused by that single command are rolled back. For example, if a

single `UPDATE` command successfully updates five rows, but an attempt to update a sixth row results in an exception, the updates to all six rows made by this `UPDATE` command are rolled back. The effects of prior SQL commands that have not yet been committed or rolled back are pending until a `COMMIT` or `ROLLBACK` command is executed.

In Advanced Server, if an exception occurs while executing a SQL command, all the updates on the database since the start of the transaction are rolled back. In addition, the transaction is left in an aborted state and either a `COMMIT` or `ROLLBACK` command must be issued before another transaction can be started.

If `edb_stmt_level_tx` is set to `TRUE`, then an exception will not automatically roll back prior uncommitted database updates. If `edb_stmt_level_tx` is set to `FALSE`, then an exception will roll back uncommitted database updates.

Note

Use `edb_stmt_level_tx` set to `TRUE` only when absolutely necessary, as this may cause a negative performance impact.

The following example run in PSQL shows that when `edb_stmt_level_tx` is `FALSE`, the abort of the second `INSERT` command also rolls back the first `INSERT` command. Note that in PSQL, the command `\set AUTOCOMMIT off` must be issued, otherwise every statement commits automatically defeating the purpose of this demonstration of the effect of `edb_stmt_level_tx`.

```
\set AUTOCOMMIT off
SET edb_stmt_level_tx TO off;

INSERT INTO emp (empno,ename,deptno) VALUES (9001, 'JONES', 40);
INSERT INTO emp (empno,ename,deptno) VALUES (9002, 'JONES', 00);
ERROR: insert or update on table "emp" violates foreign key constraint
"emp_ref_dept_fk"
DETAIL: Key (deptno)=(0) is not present in table "dept".

COMMIT;
SELECT empno, ename, deptno FROM emp WHERE empno > 9000;

empno | ename | deptno
-----+-----+
(0 rows)
```

In the following example, with `edb_stmt_level_tx` set to `TRUE`, the first `INSERT` command has not been rolled back after the error on the second `INSERT` command. At this point, the first `INSERT` command can either be committed or rolled back.

```
\set AUTOCOMMIT off
SET edb_stmt_level_tx TO on;

INSERT INTO emp (empno,ename,deptno) VALUES (9001, 'JONES', 40);
INSERT INTO emp (empno,ename,deptno) VALUES (9002, 'JONES', 00);
ERROR: insert or update on table "emp" violates foreign key constraint
"emp_ref_dept_fk"
DETAIL: Key (deptno)=(0) is not present in table "dept"

SELECT empno, ename, deptno FROM emp WHERE empno > 9000;

empno | ename | deptno
-----+-----+
 9001 | JONES |    40
(1 row)
```

```
COMMIT;
```

A `ROLLBACK` command could have been issued instead of the `COMMIT` command in which case the insert of employee number `9001` would have been rolled back as well.

`db_dialect`

Parameter Type: Enum

Default Value: postgres

Range: {postgres | redwood}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

In addition to the native PostgreSQL system catalog, `pg_catalog`, Advanced Server contains an extended catalog view. This is the `sys` catalog for the expanded catalog view. The `db_dialect` parameter controls the order in which these catalogs are searched for name resolution.

When set to `postgres`, the namespace precedence is `pg_catalog` then `sys`, giving the PostgreSQL catalog the highest precedence. When set to `redwood`, the namespace precedence is `sys` then `pg_catalog`, giving the expanded catalog views the highest precedence.

`default_with_rowids`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When set to `on`, `CREATE TABLE` includes a `ROWID` column in newly created tables, which can then be referenced in SQL commands. In earlier versions of Advanced Server `ROWIDs` were mapped to `OIDs`, but from Advanced Server version 12 onwards, the `ROWID` is an auto-incrementing value based on a sequence that starts with 1 and assigned to each row of a table created with `ROWIDs` option. By default, a unique index is created on a `ROWID` column.

The `ALTER` and `DROP` operations are restricted on a `ROWID` column.

To restore a database with `ROWIDs` from Advanced Server 11 or an earlier version, you must perform the following:

- `pg_dump`: If a table includes `OIDs` then specify `--convert-oids-to-rowids` to dump a database. Otherwise, ignore the `OIDs` to continue table creation on Advanced Server version 12 onwards.
- `pg_upgrade`: Errors out. But if a table includes `OIDs` or `ROWIDs`, then you must perform the following steps:
 1. Take a dump of the tables by specifying `--convert-oids-to-rowids` option.
 2. Drop the tables and then perform the upgrade.

3. Restore the dump after the upgrade is successful into a new cluster that contains the dumped tables into a target database.

optimizer_mode

Parameter Type: Enum

Default Value: choose

Range: {choose | ALL_ROWS | FIRST_ROWS | FIRST_ROWS_10 | FIRST_ROWS_100 | FIRST_ROWS_1000}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Sets the default optimization mode for analyzing optimizer hints.

The following table shows the possible values:

| Hint | Description |
|-----------------|--|
| ALL_ROWS | Optimizes for retrieval of all rows of the result set. |
| CHOOSE | Does no default optimization based on assumed number of rows to be retrieved from the result set. This is the default. |
| FIRST_ROWS | Optimizes for retrieval of only the first row of the result set. |
| FIRST_ROWS_10 | Optimizes for retrieval of the first 10 rows of the results set. |
| FIRST_ROWS_100 | Optimizes for retrieval of the first 100 rows of the result set. |
| FIRST_ROWS_1000 | Optimizes for retrieval of the first 1000 rows of the result set. |

These optimization modes are based upon the assumption that the client submitting the SQL command is interested in viewing only the first **n** rows of the result set and will then abandon the remainder of the result set. Resources allocated to the query are adjusted as such.

11.4.3.12 Customized Options

In previous releases of Advanced Server, the **custom_variable_classes** was required by those parameters not normally known to be added by add-on modules (such as procedural languages).

custom_variable_classes

The **custom_variable_classes** parameter is deprecated in Advanced Server 9.2; parameters that previously depended on this parameter no longer require its support.

dbms_alert.max_alerts

Parameter Type: Integer

Default Value: 100

Range: 0 to 500

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Specifies the maximum number of concurrent alerts allowed on a system using the `DBMS_ALERTS` package.

dbms_pipe.total_message_buffer

Parameter Type: Integer

Default Value: 30 Kb

Range: 30 Kb to 256 Kb

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Specifies the total size of the buffer used for the `DBMS_PIPE` package.

index_advisor.enabled

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Provides the capability to temporarily suspend Index Advisor in an EDB-PSQL or PSQL session. The Index Advisor plugin, `index_advisor`, must be loaded in the EDB-PSQL or PSQL session in order to use the `index_advisor.enabled` configuration parameter.

The Index Advisor plugin can be loaded as shown by the following example:

```
$ psql -d edb -U enterpriseedb
Password for user enterpriseedb:
psql (13.0.0)
Type "help" for help.

edb=# LOAD 'index_advisor';
LOAD
```

Use the `SET` command to change the parameter setting to control whether or not Index Advisor generates an alternative query plan as shown by the following example:

```
edb=# SET index_advisor.enabled TO off;
SET
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
      QUERY PLAN
```

```
-----  
Seq Scan on t (cost=0.00..1693.00 rows=9864 width=8)  
  Filter: (a < 10000)  
(2 rows)
```

```
edb=# SET index_advisor.enabled TO on;
SET
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
      QUERY PLAN
```

```
-----  
Seq Scan on t (cost=0.00..1693.00 rows=9864 width=8)  
  Filter: (a < 10000)  
Result (cost=0.00..327.88 rows=9864 width=8)  
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==':text  
    -> Index Scan using "<hypothetical-index>:1" on t (cost=0.00..327.88  
rows=9864 width=8)  
      Index Cond: (a < 10000)  
(6 rows)
```

`edb_sql_protect.enabled`

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload

Required Authorization to Activate: EPAS service account

Controls whether or not SQL/Protect is actively monitoring protected roles by analyzing SQL statements issued by those roles and reacting according to the setting of `edb_sql_protect.level`. When you are ready to begin monitoring with SQL/Protect set this parameter to on.

`edb_sql_protect.level`

Parameter Type: Enum

Default Value: passive

Range: {learn | passive | active}

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Reload**Required Authorization to Activate:** EPAS service account

Sets the action taken by SQL/Protect when a SQL statement is issued by a protected role.

The `edb_sql_protect.level` configuration parameter can be set to one of the following values to use either learn mode, passive mode, or active mode:

- **learn.** Tracks the activities of protected roles and records the relations used by the roles. This is used when initially configuring SQL/Protect so the expected behaviors of the protected applications are learned.
- **passive.** Issues warnings if protected roles are breaking the defined rules, but does not stop any SQL statements from executing. This is the next step after SQL/Protect has learned the expected behavior of the protected roles. This essentially behaves in intrusion detection mode and can be run in production when properly monitored.
- **active.** Stops all invalid statements for a protected role. This behaves as a SQL firewall preventing dangerous queries from running. This is particularly effective against early penetration testing when the attacker is trying to determine the vulnerability point and the type of database behind the application. Not only does SQL/Protect close those vulnerability points, but it tracks the blocked queries allowing administrators to be alerted before the attacker finds an alternate method of penetrating the system.

If you are using SQL/Protect for the first time, set `edb_sql_protect.level` to `learn`.

`edb_sql_protect.max_protected_relations`

Parameter Type: Integer

Default Value: 1024

Range: 1 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the maximum number of relations that can be protected per role. Please note the total number of protected relations for the server will be the number of protected relations times the number of protected roles. Every protected relation consumes space in shared memory. The space for the maximum possible protected relations is reserved during database server startup.

If the server is started when `edb_sql_protect.max_protected_relations` is set to a value outside of the valid range (for example, a value of 2,147,483,648), then a warning message is logged in the database server log file:

```
2014-07-18 16:04:12 EDT WARNING: invalid value for parameter
"edb_sql_protect.max_protected_relations": "2147483648"
2014-07-18 16:04:12 EDT HINT: Value exceeds integer range.
```

The database server starts successfully, but with `edb_sql_protect.max_protected_relations` set to the default value of 1024.

Though the upper range for the parameter is listed as the maximum value for an integer data type, the practical setting depends on how much shared memory is available and the parameter value used during database server startup.

As long as the space required can be reserved in shared memory, the value will be acceptable. If the value is such that the space in shared memory cannot be reserved, the database server startup fails with an error message such as the following:

2014-07-18 15:22:17 EDT FATAL: could not map anonymous shared memory:

Cannot allocate memory

2014-07-18 15:22:17 EDT HINT: This error usually means that PostgreSQL's request for a shared memory segment exceeded available memory, swap space or huge pages. To reduce the request size (currently 2070118400 bytes), reduce PostgreSQL's shared memory usage, perhaps by reducing shared_buffers or max_connections.

In such cases, reduce the parameter value until the database server can be started successfully.

edb_sql_protect.max_protected_roles

Parameter Type: Integer

Default Value: 64

Range: 1 to 2147483647

Minimum Scope of Effect: Cluster

When Value Changes Take Effect: Restart

Required Authorization to Activate: EPAS service account

Sets the maximum number of roles that can be protected.

Every protected role consumes space in shared memory. Please note that the server will reserve space for the number of protected roles times the number of protected relations (`edb_sql_protect.max_protected_relations`). The space for the maximum possible protected roles is reserved during database server startup.

If the database server is started when `edb_sql_protect.max_protected_roles` is set to a value outside of the valid range (for example, a value of 2,147,483,648), then a warning message is logged in the database server log file:

2014-07-18 16:04:12 EDT WARNING: invalid value for parameter

"`edb_sql_protect.max_protected_roles`": "2147483648"

2014-07-18 16:04:12 EDT HINT: Value exceeds integer range.

The database server starts successfully, but with `edb_sql_protect.max_protected_roles` set to the default value of 64.

Though the upper range for the parameter is listed as the maximum value for an integer data type, the practical setting depends on how much shared memory is available and the parameter value used during database server startup.

As long as the space required can be reserved in shared memory, the value will be acceptable. If the value is such that the space in shared memory cannot be reserved, the database server startup fails with an error message such as the following:

2014-07-18 15:22:17 EDT FATAL: could not map anonymous shared memory:

Cannot allocate memory

2014-07-18 15:22:17 EDT HINT: This error usually means that PostgreSQL's request for a shared memory segment exceeded available memory, swap space or huge pages. To reduce the request size (currently 2070118400 bytes), reduce PostgreSQL's shared memory usage, perhaps by reducing shared_buffers or max_connections.

In such cases, reduce the parameter value until the database server can be started successfully.

edb_sql_protect.max_queries_to_save**Parameter Type:** Integer**Default Value:** 5000**Range:** 100 to 2147483647**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Restart**Required Authorization to Activate:** EPAS service account

Sets the maximum number of offending queries to save in view `edb_sql_protect_queries`.

Every query that is saved consumes space in shared memory. The space for the maximum possible queries that can be saved is reserved during database server startup.

If the database server is started when `edb_sql_protect.max_queries_to_save` is set to a value outside of the valid range (for example, a value of 10), then a warning message is logged in the database server log file:

```
2014-07-18 13:05:31 EDT WARNING: 10 is outside the valid range for
parameter "edb_sql_protect.max_queries_to_save" (100 .. 2147483647)
```

The database server starts successfully, but with `edb_sql_protect.max_queries_to_save` set to the default value of 5000.

Though the upper range for the parameter is listed as the maximum value for an integer data type, the practical setting depends on how much shared memory is available and the parameter value used during database server startup.

As long as the space required can be reserved in shared memory, the value will be acceptable. If the value is such that the space in shared memory cannot be reserved, the database server startup fails with an error message such as the following:

```
2014-07-18 15:22:17 EDT FATAL: could not map anonymous shared memory:
Cannot allocate memory
2014-07-18 15:22:17 EDT HINT: This error usually means that PostgreSQL's
request for a shared memory segment exceeded available memory, swap
space or huge pages. To reduce the request size (currently 2070118400
bytes), reduce PostgreSQL's shared memory usage, perhaps by reducing
shared_buffers or max_connections.
```

In such cases, reduce the parameter value until the database server can be started successfully.

edb_wait_states.directory**Parameter Type:** String**Default Value:** `edb_wait_states`**Range:** n/a**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Restart

Required Authorization to Activate: EPAS service account

Sets the directory path where the EDB wait states log files are stored. The specified path should be a full, absolute path and not a relative path. However, the default setting is `edb_wait_states`, which makes `$PGDATA/edb_wait_states` the default directory location. For information on EDB wait states see [EDB Wait States](#) under [Performance Analysis and Tuning](#).

`edb_wait_states.retention_period`**Parameter Type:** Integer**Default Value:** 604800s**Range:** 86400s to 2147483647s**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

Sets the period of time after which the log files for EDB wait states will be deleted. The default is 604800 seconds, which is 7 days. For information on EDB wait states see [EDB Wait States](#) under [Performance Analysis and Tuning](#).

`edb_wait_states.sampling_interval`**Parameter Type:** Integer**Default Value:** 1s**Range:** 1s to 2147483647s**Minimum Scope of Effect:** Cluster**When Value Changes Take Effect:** Reload**Required Authorization to Activate:** EPAS service account

Sets the timing interval between two sampling cycles for EDB wait states. The default setting is 1 second. For information on EDB wait states see [EDB Wait States](#) under [Performance Analysis and Tuning](#).

`edbldr.empty_csv_field`**Parameter Type:** Enum**Default Value:** NULL**Range:** {NULL | empty_string | pgsql}**Minimum Scope of Effect:** Per session**When Value Changes Take Effect:** Immediate**Required Authorization to Activate:** Session user

Use the `edbldr.empty_csv_field` parameter to specify how EDB*Loader will treat an empty string. The valid values for the `edbldr.empty_csv_field` parameter are:

| Parameter Setting | EDB*Loader Behavior |
|-------------------|---|
| NULL | An empty field is treated as NULL. |
| empty_string | An empty field is treated as a string of length zero. |
| pgsql | An empty field is treated as a NULL if it does not contain quotes and as an empty string if it contains quotes. |

For more information about the `edbldr.empty_csv_field` parameter in EDB*Loader, see the *Database Compatibility for Oracle Developers Tools and Utilities Guide* at the EnterpriseDB website:

<https://www.enterprisedb.com/edb-docs>

utl_encode.uudecode_redwood

Parameter Type: Boolean

Default Value: false

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

When set to `TRUE`, Advanced Server's `UTL_ENCODE.UUDECODE` function can decode uuencoded data that was created by the Oracle implementation of the `UTL_ENCODE.UUENCODE` function.

When set to the default setting of `FALSE`, Advanced Server's `UTL_ENCODE.UUDECODE` function can decode uuencoded data created by Advanced Server's `UTL_ENCODE.UUENCODE` function.

utl_file.umask

Parameter Type: String

Default Value: 0077

Range: Octal digits for umask settings

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

The `utl_file.umask` parameter sets the *file mode creation mask* or simply, the *mask*, in a manner similar to the Linux `umask` command. This is for usage only within the Advanced Server `UTL_FILE` package.

Note

The `utl_file.umask` parameter is not supported on Windows systems.

The value specified for `utl_file.umask` is a 3 or 4-character octal string that would be valid for the Linux `umask` command. The setting determines the permissions on files created by the `UTL_FILE` functions and procedures. (Refer to any information source regarding Linux or Unix systems for information on file permissions and the usage of the `umask` command.)

The following shows the results of the default `utl_file.umask` setting of 0077. Note that all permissions are denied on users belonging to the `enterprisedb` group as well as all other users. Only user `enterprisedb` has read and write permissions on the file.

```
-rw----- 1 enterprisedb enterprisedb 21 Jul 24 16:08 utlfile
```

11.4.3.13 Ungrouped

Configuration parameters in this section apply to Advanced Server only and are for a specific, limited purpose.

`nls_length_semantics`

Parameter Type: Enum

Default Value: byte

Range: {byte | char}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Superuser

This parameter has no effect in Advanced Server.

For example, the form of the `ALTER SESSION` command is accepted in Advanced Server without throwing a syntax error, but does not alter the session environment:

```
ALTER SESSION SET nls_length_semantics = char;
```

Note

Since the setting of this parameter has no effect on the server environment, it does not appear in the system view `pg_settings`.

`query_rewrite_enabled`

Parameter Type: Enum

Default Value: false

Range: {true | false | force}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

This parameter has no effect in Advanced Server.

For example, the following form of the `ALTER SESSION` command is accepted in Advanced Server without throwing a syntax error, but does not alter the session environment:

```
ALTER SESSION SET query_rewrite_enabled = force;
```

Note

Since the setting of this parameter has no effect on the server environment, it does not appear in the system view `pg_settings`.

query_rewrite_integrity

Parameter Type: Enum

Default Value: enforced

Range: {enforced | trusted | stale_tolerated}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Superuser

This parameter has no effect in Advanced Server.

For example, the following form of the `ALTER SESSION` command is accepted in Advanced Server without throwing a syntax error, but does not alter the session environment:

```
ALTER SESSION SET query_rewrite_integrity = stale_tolerated;
```

Note

Since the setting of this parameter has no effect on the server environment, it does not appear in the system view `pg_settings`.

timed_statistics

Parameter Type: Boolean

Default Value: true

Range: {true | false}

Minimum Scope of Effect: Per session

When Value Changes Take Effect: Immediate

Required Authorization to Activate: Session user

Controls the collection of timing data for the Dynamic Runtime Instrumentation Tools Architecture (DRITA) feature.

When set to on, timing data is collected.

Note

When Advanced Server is installed, the `postgresql.conf` file contains an explicit entry setting `timed_statistics` to off. If this entry is commented out letting `timed_statistics` to default, and the configuration file is reloaded, timed statistics collection would be turned on.

11.5 Index Advisor

The Index Advisor utility helps determine which columns you should index to improve performance in a given workload. Index Advisor considers B-tree (single-column or composite) index types, and does not identify other index types (GIN, GiST, Hash) that may improve performance. Index Advisor is installed with EDB Postgres Advanced Server.

Index Advisor works with Advanced Server's query planner by creating *hypothetical indexes* that the query planner uses to calculate execution costs as if such indexes were available. Index Advisor identifies the indexes by analyzing SQL queries supplied in the workload.

There are three ways to use Index Advisor to analyze SQL queries:

- Invoke the Index Advisor utility program, supplying a text file containing the SQL queries that you wish to analyze; Index Advisor will generate a text file with `CREATE INDEX` statements for the recommended indexes.
- Provide queries at the EDB-PSQL command line that you want Index Advisor to analyze.
- Access Index Advisor through the Postgres Enterprise Manager client. When accessed via the PEM client, Index Advisor works with SQL Profiler, providing indexing recommendations on code captured in SQL traces. For more information about using SQL Profiler and Index Advisor with PEM, see the *PEM Getting Started Guide* available from the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

Index Advisor will attempt to make indexing recommendations on `INSERT`, `UPDATE`, `DELETE` and `SELECT` statements. When invoking Index Advisor, you supply the workload in the form of a set of queries (if you are providing the command in an SQL file) or an `EXPLAIN` statement (if you are specifying the SQL statement at the psql command line). Index Advisor displays the query plan and estimated execution cost for the supplied query, but does not actually execute the query.

During the analysis, Index Advisor compares the query execution costs with and without hypothetical indexes. If the execution cost using a hypothetical index is less than the execution cost without it, both plans are reported in the `EXPLAIN` statement output, metrics that quantify the improvement are calculated, and Index Advisor generates the `CREATE INDEX` statement needed to create the index.

If no hypothetical index can be found that reduces the execution cost, Index Advisor displays only the original query plan output of the `EXPLAIN` statement.

Index Advisor does not actually create indexes on the tables. Use the `CREATE INDEX` statements supplied by Index Advisor to add any recommended indexes to your tables.

A script supplied with Advanced Server creates the table in which Index Advisor stores the indexing recommendations generated by the analysis; the script also creates a function and a view of the table to simplify the retrieval and interpretation of the results.

If you choose to forego running the script, Index Advisor will log recommendations in a temporary table that is available only for the duration of the Index Advisor session.

11.5.1 Index Advisor Components

The Index Advisor shared library interacts with the query planner to make indexing recommendations. On Windows, the Advanced Server installer creates the following shared library in the `libdir` subdirectory of your Advanced Server home directory. For Linux, install the `edb-asxx-server-indexadvisor` RPM package where `xx` is the Advanced Server version number.

On Linux:

`index_advisor.so`

On Windows:

`index_advisor.dll`

Please note that libraries in the `libdir` directory can only be loaded by a superuser. A database administrator can allow a non-superuser to use Index Advisor by manually copying the Index Advisor file from the `libdir` directory into the `libdir/plugins` directory (under your Advanced Server home directory). Only a trusted non-superuser should be allowed access to the plugin; this is an unsafe practice in a production environment.

The installer also creates the Index Advisor utility program and setup script:

`pg_advise_index`

`pg_advise_index` is a utility program that reads a user-supplied input file containing SQL queries and produces a text file containing `CREATE INDEX` statements that can be used to create the indexes recommended by the Index Advisor. The `pg_advise_index` program is located in the `bin` subdirectory of the Advanced Server home directory.

`index_advisor.sql`

`index_advisor.sql` is a script that creates a permanent Index Advisor log table along with a function and view to facilitate reporting of recommendations from the log table. The script is located in the `share/contrib` subdirectory of the Advanced Server directory.

The `index_advisor.sql` script creates the `index_advisor_log` table, the `show_index_recommendations()` function and the `index_recommendations` view. These database objects must be created in a schema that is accessible by, and included in the search path of the role that will invoke Index Advisor.

`index_advisor_log`

Index Advisor logs indexing recommendations in the `index_advisor_log` table. If Index Advisor does not find the `index_advisor_log` table in the user's search path, Index Advisor will store any indexing recommendations in a temporary table of the same name. The temporary table exists only for the duration of the current session.

`show_index_recommendations()`

`show_index_recommendations()` is a PL/pgSQL function that interprets and displays the recommendations made during a specific Index Advisor session (as identified by its backend process ID).

`index_recommendations`

Index Advisor creates the `index_recommendations` view based on information stored in the `index_advisor_log` table during a query analysis. The view produces output in the same format as the `show_index_recommendations()` function, but contains Index Advisor recommendations for all stored sessions, while the result set returned by the `show_index_recommendations()` function are limited to a specified session.

11.5.2 Index Advisor Configuration

Index Advisor does not require configuration to generate recommendations that are available only for the duration of the current session; to store the results of multiple sessions, you must create the `index_advisor_log` table. To create the `index_advisor_log` table, you must run the `index_advisor.sql` script.

When selecting a storage schema for the Index Advisor table, function and view, keep in mind that all users that invoke Index Advisor (and query the result set) must have `USAGE` privileges on the schema. The schema must be in the search path of all users that are interacting with the Index Advisor.

1. Place the selected schema at the start of your `search_path` parameter. For example, if your search path is currently:

```
search_path=public, accounting
```

and you want the Index Advisor objects to be created in a schema named `advisor`, use the command:

```
SET search_path = advisor, public, accounting;
```

2. Run the `index_advisor.sql` script to create the database objects. If you are running the `psql` client, you can use the command:

```
\i full.pathname/index_advisor.sql
```

Specify the pathname to the `index_advisor.sql` script in place of `full.pathname`.

3. Grant privileges on the `index_advisor_log` table to all Index Advisor users; this step is not necessary if the Index Advisor user is a superuser, or the owner of these database objects.

- Grant `SELECT` and `INSERT` privileges on the `index_advisor_log` table to allow a user to invoke Index Advisor.
- Grant `DELETE` privileges on the `index_advisor_log` table to allow the specified user to delete the table contents.
- Grant `SELECT` privilege on the `index_recommendations` view.

The following example demonstrates the creation of the Index Advisor database objects in a schema named `ia`, which will then be accessible to an Index Advisor user with user name `ia_user`:

```
$ edb-psql -d edb -U enterpriseedb
psql.bin (13.0.0, server 13.0.0)
Type "help" for help.
```

```
edb=# CREATE SCHEMA ia;
CREATE SCHEMA
edb=# SET search_path TO ia;
SET
edb=# \i /usr/edb/as13/share/contrib/index_advisor.sql
CREATE TABLE
CREATE INDEX
CREATE INDEX
CREATE FUNCTION
CREATE FUNCTION
CREATE VIEW
edb=# GRANT USAGE ON SCHEMA ia TO ia_user;
GRANT
edb=# GRANT SELECT, INSERT, DELETE ON index_advisor_log TO ia_user;
```

```
GRANT
edb=# GRANT SELECT ON index_recommendations TO ia_user;
GRANT
```

While using Index Advisor, the specified schema (`ia`) must be included in `ia_user`'s `search_path` parameter.

11.5.3 Using Index Advisor

When you invoke Index Advisor, you must supply a workload; the workload is either a query (specified at the command line), or a file that contains a set of queries (executed by the `pg_advise_index()` function). After analyzing the workload, Index Advisor will either store the result set in a temporary table, or in a permanent table. You can review the indexing recommendations generated by Index Advisor and use the `CREATE INDEX` statements generated by Index Advisor to create the recommended indexes.

Note

You should not run Index Advisor in read-only transactions.

The following examples assume that superuser `enterprisedb` is the Index Advisor user, and the Index Advisor database objects have been created in a schema in the `search_path` of superuser `enterprisedb`.

The examples in the following sections use the table created with the statement shown below:

```
CREATE TABLE t( a INT, b INT );
INSERT INTO t SELECT s, 99999 - s FROM generate_series(0,99999) AS s;
ANALYZE t;
```

The resulting table contains the following rows:

| a | b |
|-------|-------|
| 0 | 99999 |
| 1 | 99998 |
| 2 | 99997 |
| 3 | 99996 |
| . | . |
| 99997 | 2 |
| 99998 | 1 |
| 99999 | 0 |

Using the `pg_advise_index` Utility

When invoking the `pg_advise_index` utility, you must include the name of a file that contains the queries that will be executed by `pg_advise_index`; the queries may be on the same line, or on separate lines, but each query must be terminated by a semicolon. Queries within the file should not begin with the `EXPLAIN` keyword.

The following example shows the contents of a sample `workload.sql` file:

```
SELECT * FROM t WHERE a = 500;
```

```
SELECT * FROM t WHERE b < 1000;
```

Run the `pg_advise_index` program as shown in the code sample below:

```
$ pg_advise_index -d edb -h localhost -U enterprisedb -s 100M -o
advisory.sql workload.sql
poolsize = 102400 KB
ad workload from file 'workload.sql'
Analyzing queries .. done.
size = 2184 KB, benefit = 1684.720000
size = 2184 KB, benefit = 1655.520000
/* 1. t(a): size=2184 KB, benefit=1684.72 */
/* 2. t(b): size=2184 KB, benefit=1655.52 */
/* Total size = 4368KB */
```

In the code sample, the `-d`, `-h`, and `-U` options are psql connection options.

`-s`

`-s` is an optional parameter that limits the maximum size of the indexes recommended by Index Advisor. If Index Advisor does not return a result set, `-s` may be set too low.

`-o`

The recommended indexes are written to the file specified after the `-o` option.

The information displayed by the `pg_advise_index` program is logged in the `index_advisor_log` table. In response to the command shown in the example, Index Advisor writes the following `CREATE INDEX` statements to the `advisory.sql` output file.

```
create index idx_t_1 on t (a);
create index idx_t_2 on t (b);
```

You can create the recommended indexes at the psql command line with the `CREATE INDEX` statements in the file, or create the indexes by executing the `advisory.sql` script.

```
$ edb-psql -d edb -h localhost -U enterprisedb -e -f advisory.sql
create index idx_t_1 on t (a);
CREATE INDEX
create index idx_t_2 on t (b);
CREATE INDEX
```

Using Index Advisor at the psql Command Line

You can use Index Advisor to analyze SQL statements entered at the `edb-psql` (or `psql`) command line; the following steps detail loading the Index Advisor plugin and using Index Advisor:

1. Connect to the server with the `edb-psql` command line utility, and load the Index Advisor plugin:

```
$ edb-psql -d edb -U enterprisedb
...
edb=# LOAD 'index_advisor';
LOAD
```

1. Use the `edb-psql` command line to invoke each SQL command that you would like Index Advisor to analyze. Index Advisor stores any recommendations for the queries in the `index_advisor_log` table. If the

`index_advisor_log` table does not exist in the user's `search_path`, a temporary table is created with the same name. This temporary table exists only for the duration of the user's session.

After loading the Index Advisor plugin, Index Advisor will analyze all SQL statements and log any indexing recommendations for the duration of the session.

If you would like Index Advisor to analyze a query (and make indexing recommendations) without actually executing the query, preface the SQL statement with the `EXPLAIN` keyword.

If you do not preface the statement with the `EXPLAIN` keyword, Index Advisor will analyze the statement while the statement executes, writing the indexing recommendations to the `index_advisor_log` table for later review.

In the example that follows, the `EXPLAIN` statement displays the normal query plan, followed by the query plan of the same query, if the query were using the recommended hypothetical index:

```
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
```

QUERY PLAN

```
Seq Scan on t (cost=0.00..1693.00 rows=10105 width=8)
  Filter: (a < 10000)
Result (cost=0.00..337.10 rows=10105 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==='::text
    -> Index Scan using "<hypothetical-index>:1" on t
      (cost=0.00..337.10 rows=10105 width=8)
        Index Cond: (a < 10000)
(6 rows)
```

```
edb=# EXPLAIN SELECT * FROM t WHERE a = 100;
```

QUERY PLAN

```
Seq Scan on t (cost=0.00..1693.00 rows=1 width=8)
  Filter: (a = 100)
Result (cost=0.00..8.28 rows=1 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==='::text
    -> Index Scan using "<hypothetical-index>:3" on t
      (cost=0.00..8.28 rows=1 width=8)
        Index Cond: (a = 100)
(6 rows)
```

After loading the Index Advisor plugin, the default value of `index_advisor.enabled` is `on`. The Index Advisor plugin must be loaded to use a `SET` or `SHOW` command to display the current value of `index_advisor.enabled`.

You can use the `index_advisor.enabled` parameter to temporarily disable Index Advisor without interrupting the psql session:

```
edb=# SET index_advisor.enabled TO off;
SET
```

To enable Index Advisor, set the parameter to `on`:

```
edb=# SET index_advisor.enabled TO on;
SET
```

11.5.4 Reviewing the Index Advisor Recommendations

There are several ways to review the index recommendations generated by Index Advisor. You can:

- Query the `index_advisor_log` table.
- Run the `show_index_recommendations` function.
- Query the `index_recommendations` view.

Using the `show_index_recommendations()` Function

To review the recommendations of the Index Advisor utility using the `show_index_recommendations()` function, call the function, specifying the process ID of the session:

```
| SELECT show_index_recommendations( pid );
```

Where `pid` is the process ID of the current session. If you do not know the process ID of your current session, passing a value of `NULL` will also return a result set for the current session.

The following code fragment shows an example of a row in a result set:

```
edb=# SELECT show_index_recommendations(null);
      show_index_recommendations
-----
create index idx_t_a on t(a);/* size: 2184 KB, benefit: 3040.62,
gain: 1.39222666981456 */
(1 row)
```

In the example, `create index idx_t_a on t(a)` is the SQL statement needed to create the index suggested by Index Advisor. Each row in the result set shows:

- The command required to create the recommended index.
- The maximum estimated size of the index.
- The calculated benefit of using the index.
- The estimated gain that will result from implementing the index.

You can display the results of all Index Advisor sessions from the following view:

```
| SELECT * FROM index_recommendations;
```

Querying the `index_advisor_log` Table

Index Advisor stores indexing recommendations in a table named `index_advisor_log`. Each row in the `index_advisor_log` table contains the result of a query where Index Advisor determines it can recommend a hypothetical index to reduce the execution cost of that query.

| Column | Type | Description |
|-------------|-----------|--|
| reloid | oid | OID of the base table for the index |
| relname | name | Name of the base table for the index |
| attrs | integer[] | Recommended index columns (identified by column number) |
| benefit | real | Calculated benefit of the index for this query |
| index_size | integer | Estimated index size in disk-pages |
| backend_pid | integer | Process ID of the process generating this recommendation |

| Column | Type | Description |
|-----------|-----------|---|
| timestamp | timestamp | Date/Time when the recommendation was generated |

You can query the `index_advisor_log` table at the psql command line. The following example shows the `index_advisor_log` table entries resulting from two Index Advisor sessions. Each session contains two queries, and can be identified (in the table below) by a different `backend_pid` value. For each session, Index Advisor generated two index recommendations.

```
edb=# SELECT * FROM index_advisor_log;
reloid | relname | attrs | benefit | index_size | backend_pid | timestamp
-----+-----+-----+-----+-----+
16651 | t | {1} | 1684.72 | 2184 | 3442 | 22-MAR-11
16:44:32.712638 -04:00
16651 | t | {2} | 1655.52 | 2184 | 3442 | 22-MAR-11
16:44:32.759436 -04:00
16651 | t | {1} | 1355.9 | 2184 | 3506 | 22-MAR-11
16:48:28.317016 -04:00
16651 | t | {1} | 1684.72 | 2184 | 3506 | 22-MAR-11
16:51:45.927906 -04:00
(4 rows)
```

Index Advisor added the first two rows to the table after analyzing the following two queries executed by the `pg_advise_index` utility:

```
SELECT * FROM t WHERE a = 500;
SELECT * FROM t WHERE b < 1000;
```

The value of `3442` in column `backend_pid` identifies these results as coming from the session with process ID `3442`.

The value of `1` in column `attrs` in the first row indicates that the hypothetical index is on the first column of the table (column `a` of table `t`).

The value of `2` in column `attrs` in the second row indicates that the hypothetical index is on the second column of the table (column `b` of table `t`).

Index Advisor added the last two rows to the table after analyzing the following two queries (executed at the psql command line):

```
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
          QUERY PLAN
-----
Seq Scan on t (cost=0.00..1693.00 rows=10105 width=8)
  Filter: (a < 10000)
Result (cost=0.00..337.10 rows=10105 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]===:text'
    -> Index Scan using "<hypothetical-index>:1" on t (cost=0.00..337.10
      rows=10105 width=8)
        Index Cond: (a < 10000)
(6 rows)
```

```
edb=# EXPLAIN SELECT * FROM t WHERE a = 100;
          QUERY PLAN
-----
```

```

Seq Scan on t (cost=0.00..1693.00 rows=1 width=8)
  Filter: (a = 100)
Result (cost=0.00..8.28 rows=1 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==='::text
    -> Index Scan using "<hypothetical-index>:3" on t (cost=0.00..8.28
rows=1 width=8)
      Index Cond: (a = 100)
(6 rows)

```

The values in the `benefit` column of the `index_advisor_log` table are calculated using the following formula:

`benefit = (normal execution cost) - (execution cost with hypothetical index)`

The value of the `benefit` column for the last row of the `index_advisor_log` table (shown in the example) is calculated using the query plan for the following SQL statement:

`EXPLAIN SELECT * FROM t WHERE a = 100;`

The execution costs of the different execution plans are evaluated and compared:

`benefit = (Seq Scan on t cost) - (Index Scan using <hypothetical-index>)`

and the benefit is added to the table:

`benefit = 1693.00 - 8.28`

`benefit = 1684.72`

You can delete rows from the `index_advisor_log` table when you no longer have the need to review the results of the queries stored in the row.

Querying the `index_recommendations` View

The `index_recommendations` view contains the calculated metrics and the `CREATE INDEX` statements to create the recommended indexes for all sessions whose results are currently in the `index advisor log` table. You can display the results of all stored Index Advisor sessions by querying the `index_recommendations` view as shown below:

`SELECT * FROM index_recommendations;`

Using the example shown in the previous section ([Querying the `index_advisor_log` Table](#)), the `index_recommendations` view displays the following:

```

edb=# SELECT * FROM index_recommendations;
backend_pid | show_index_recommendations
-----+
-----+
 3442 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
1684.72, gain: 0.771392654586624 */
 3442 | create index idx_t_b on t(b);/* size: 2184 KB, benefit:
1655.52, gain: 0.758021539820856 */
 3506 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
3040.62, gain: 1.39222666981456 */
(3 rows)

```

Within each session, the results of all queries that benefit from the same recommended index are combined to produce one set of metrics per recommended index, reflected in the fields named `benefit` and `gain`.

The formulas for the fields are as follows:

```
size = MAX(index size of all queries)
benefit = SUM(benefit of each query)
gain = SUM(benefit of each query) / MAX(index size of all queries)
```

So for example, using the following query results from the process with a `backend_pid` of `3506`:

| reloid | relname | attrs | benefit | index_size | backend_pid | timestamp |
|--------|---------|-------|---------|------------|-------------|-------------------------------------|
| 16651 | t | {1} | 1355.9 | 2184 | 3506 | 22-MAR-11 16:48:28.317016 -04:00 |
| 16651 | t | {1} | 1684.72 | 2184 | 3506 | 22-MAR-11 16:51:45.927906 -04:00 |

The metrics displayed from the `index_recommendations` view for `backend_pid 3506` are:

| backend_pid | show_index_recommendations |
|-------------|--|
| 3506 | create index idx_t_a on t(a);/* size: 2184 KB, benefit: 3040.62, gain: 1.39222666981456 */ |

The metrics from the view are calculated as follows:

```
benefit = (benefit from 1st query) + (benefit from 2nd query)
benefit = 1355.9 + 1684.72
benefit = 3040.62
```

and

```
gain = ((benefit from 1st query) + (benefit from 2nd query)) / MAX(index
size of all queries)
gain = (1355.9 + 1684.72) / MAX(2184, 2184)
gain = 3040.62 / 2184
gain = 1.39223
```

The gain metric is useful when comparing the relative advantage of the different recommended indexes derived during a given session. The larger the gain value, the better the cost effectiveness derived from the index weighed against the possible disk space consumption of the index.

11.5.5 Limitations

Index Advisor does not consider Index Only scans; it does consider Index scans when making recommendations.

Index Advisor ignores any computations found in the `WHERE` clause. Effectively, the index field in the recommendations will not be any kind of expression; the field will be a simple column name.

Index Advisor does not consider inheritance when recommending hypothetical indexes. If a query references a parent table, Index Advisor does not make any index recommendations on child tables.

Restoration of a `pg_dump` backup file that includes the `index_advisor_log` table or any tables for which indexing recommendations were made and stored in the `index_advisor_log` table, may result in "broken links" between the `index_advisor_log` table and the restored tables referenced by rows in the `index_advisor_log` table because of changes in object identifiers (OIDs).

If it is necessary to display the recommendations made prior to the backup, you can replace the old OIDs in the `reloid` column of the `index_advisor_log` table with the new OIDs of the referenced tables using the SQL `UPDATE` statement:

```
UPDATE index_advisor_log SET reloid = new_oid WHERE reloid = old_oid;
```

11.6 SQL Profiler

Inefficient SQL code is one of, if not the leading cause of database performance problems. The challenge for database administrators and developers is locating and then optimizing this code in large, complex systems.

SQL Profiler helps you locate and optimize poorly running SQL code.

Specific features and benefits of SQL Profiler include the following:

- **On-Demand Traces.** You can capture SQL traces at any time by manually setting up your parameters and starting the trace.
- **Scheduled Traces.** For inconvenient times, you can also specify your trace parameters and schedule them to run at some later time.
- **Save Traces.** Execute your traces and save them for later review.
- **Trace Filters.** Selectively filter SQL captures by database and by user, or capture every SQL statement sent by all users against all databases.
- **Trace Output Analyzer.** A graphical table lets you quickly sort and filter queries by duration or statement, and a graphical or text based `EXPLAIN` plan lays out your query paths and joins.
- **Index Advisor Integration.** Once you have found your slow queries and optimized them, you can also let the Index Advisor recommend the creation of underlying table indices to further improve performance.

The following describes the installation process.

Step 1: Install SQL Profiler

SQL Profiler is installed by the Advanced Server installer on Windows or from the `edb-asxx-server-sqlprofiler` RPM package on Linux where `xx` is the Advanced Server version number.

Step 2: Add the SQL Profiler library

Modify the `postgresql.conf` parameter file for the instance to include the SQL Profiler library in the `shared_preload_libraries` configuration parameter.

For Linux installations, the parameter value should include:

```
$libdir/sql-profiler
```

On Windows, the parameter value should include:

```
$libdir\sql-profiler.dll
```

Step 3: Create the functions used by SQL Profiler

The SQL Profiler installation program places a SQL script (named `sql-profiler.sql`) in:

On Linux:

```
/usr/edb/as13/share/contrib/
```

On Windows:

```
C:\Program Files\edb\as13\share\contrib\
```

Use the `psql` command line interface to run the `sql-profiler.sql` script in the database specified as the Maintenance Database on the server you wish to profile. If you are using Advanced Server, the default maintenance database is named `edb`. If you are using a PostgreSQL instance, the default maintenance database is named `postgres`.

The following command uses the `psql` command line to invoke the `sql-profiler.sql` script on a Linux system:

```
$ /usr/edb/as13/bin/psql -U user_name database_name <
/usr/edb/as13/share/contrib/sql-profiler.sql
```

Step 4: Stop and restart the server for the changes to take effect.

After configuring SQL Profiler, it is ready to use with all databases that reside on the server. You can take advantage of SQL Profiler functionality with EDB Postgres Enterprise Manager; for more information about Postgres Enterprise Manager, visit the EnterpriseDB website at:

<https://www.enterprisedb.com/edb-docs>

Troubleshooting

If (after performing an upgrade to a newer version of SQL Profiler) you encounter an error that contains the following text:

An error has occurred:

```
ERROR: function return row and query-specified return row do not match.
DETAIL: Returned row contains 11 attributes, but the query expects 10.
```

To correct this error, you must replace the existing query set with a new query set. First, uninstall SQL Profiler by invoking the `uninstall-sql-profiler.sql` script, and then reinstall SQL Profiler by invoking the `sql-profiler.sql` script.

11.7 pgsnmpd

`pgsnmpd` is an SNMP agent that can return hierarchical information about the current state of Advanced Server on a Linux system. `pgsnmpd` is distributed and installed using the `edb-asxx-pgsnmpd` RPM package where `xx` is the Advanced Server version number. The `pgsnmpd` agent can operate as a stand-alone SNMP agent, as a pass-through sub-agent, or as an AgentX sub-agent.

After installing Advanced Server, you will need to update the `LD_LIBRARY_PATH` variable. Use the command:

```
$ export LD_LIBRARY_PATH=/usr/edb/as13/lib:$LD_LIBRARY_PATH
```

This command does not persistently alter the value of `LD_LIBRARY_PATH`. Consult the documentation for your distribution of Linux for information about persistently setting the value of `LD_LIBRARY_PATH`.

The examples that follow demonstrate the simplest usage of `pgsnmpd`, implementing read only access. `pgsnmpd` is based on the net-snmp library; for more information about net-snmp, visit:

| <http://net-snmp.sourceforge.net/>

Configuring pgsqlsnmpd

The `pgsqlsnmpd` configuration file is named `snmpd.conf`. For information about the directives that you can specify in the configuration file, please review the `snmpd.conf` man page (`man snmpd.conf`).

You can create the configuration file by hand, or you can use the `snmpconf` perl script to create the configuration file. The perl script is distributed with net-snmp package.

net-snmp is an open-source package available from:

| <http://www.net-snmp.org/>

To use the `snmpconf` configuration file wizard, download and install net-snmp. When the installation completes, open a command line and enter:

| `snmpconf`

When the configuration file wizard opens, it may prompt you to read in an existing configuration file. Enter `none` to generate a new configuration file (not based on a previously existing configuration file).

`snmpconf` is a menu-driven wizard. Select menu item `1: snmpd.conf` to start the configuration wizard. As you select each top-level menu option displayed by `snmpconf`, the wizard walks through a series of questions, prompting you for information required to build the configuration file. When you have provided information in each of the category relevant to your system, enter `Finished` to generate a configuration file named `snmpd.conf`. Copy the file to:

| `/usr/edb/as13/share/`

Setting the Listener Address

By default, `pgsqlsnmpd` listens on port `161`. If the listener port is already being used by another service, you may receive the following error:

| Error opening specified endpoint "udp:161".

You can specify an alternate listener port by adding the following line to your `snmpd.conf` file:

| `agentaddress $host_address:2000`

The example instructs `pgsqlsnmpd` to listen on UDP port `2000`, where `$host_address` is the IP address of the server (e.g., `127.0.0.1`).

Invoking pgsqlsnmpd

Ensure that an instance of Advanced Server is up and running (`pgsqlsnmpd` will connect to this server). Open a command line and assume superuser privileges, before invoking pgsqlsnmpd with a command that takes the following form:

```
POSTGRES_INSTALL_HOME/bin/pgsqlsnmpd -b
-c POSTGRES_INSTALL_HOME/share/snmpd.conf
-C "user=enterprisedb dbname=edb password=safe_password
port=5444"
```

Where `POSTGRES_INSTALL_HOME` specifies the Advanced Server installation directory.

Include the `-b` option to specify that `pgsnmpd` should run in the background.

Include the `-c` option, specifying the path and name of the `pgsnmpd` configuration file.

Include connection information for your installation of Advanced Server (in the form of a `libpq` connection string) after the `-C` option.

Viewing pgsnmpd Help

Include the `--help` option when invoking the `pgsnmpd` utility to view other `pgsnmpd` command line options:

```
pgsnmpd --help
Version PGSQL-SNMP-Ver1.0
usage: pgsnmpd [-s] [-b] [-c FILE] [-x address] [-g] [-C "Connect
String"]
-s : run as AgentX sub-agent of an existing snmpd process
-b : run in the background
-c : configuration file name
-g : use syslog
-C : libpq connection string
-x : address:port of a network interface
-V : display version strings
```

Requesting Information from pgsnmpd

You can use `net-snmp` commands to query the `pgsnmpd` service. For example:

```
snmpgetnext -v 2c -c public localhost .1.3.6.1.4.1.5432.1.4.2.1.1.0
```

In the above example:

`-v 2c` option instructs the `snmpgetnext` client to send the request in SNMP version 2c format.

`-c public` specifies the community name.

`localhost` indicates the host machine running the `pgsnmpd` server.

`.1.3.6.1.4.1.5432.1.4.2.1.1.0` is the identity of the requested object. To see a list of all databases, increment the last digit by one (e.g. .1.1, .1.2, .1.3 etc.).

The encodings required to query any given object are defined in the MIB (Management Information Base). An SNMP client can monitor a variety of servers; the server type determines the information exposed by a given server. Each SNMP server describes the exposed data in the form of a MIB (Management information base). By default, `pgsnmpd` searches for MIB's in the following locations:

`/usr/share/snmp/mibs`

`$HOME/.snmp/mibs`

11.8 EDB Audit Logging

Advanced Server allows database and security administrators, auditors, and operators to track and analyze database activities using the *EDB Audit Logging* functionality. EDB Audit Logging generates audit log files, which contains all of the relevant information. The audit logs can be configured to record information such as:

- When a role establishes a connection to an Advanced Server database
- What database objects a role creates, modifies, or deletes when connected to Advanced Server
- When any failed authentication attempts occur

The parameters specified in the configuration files, `postgresql.conf` or `postgresql.auto.conf`, control the information included in the audit logs.

11.8.1 Audit Logging Configuration Parameters

Use the following configuration parameters to control database auditing. See [Summary of Configuration Parameters](#) to determine if a change to the configuration parameter takes effect immediately, or if the configuration needs to be reloaded, or if the database server needs to be restarted.

`edb_audit`

Enables or disables database auditing. The values `xml` or `csv` will enable database auditing. These values represent the file format in which auditing information will be captured. `none` will disable database auditing and is also the default.

`edb_audit_directory`

Specifies the directory where the log files will be created. The path of the directory can be relative or absolute to the data folder. The default is the `PGDATA/edb_audit` directory.

`edb_audit_filename`

Specifies the file name of the audit file where the auditing information will be stored. The default file name will be `audit-%Y%m%d_%H%M%S`. The escape sequences, `%Y`, `%m` etc., will be replaced by the appropriate current values according to the system date and time.

`edb_audit_rotation_day`

Specifies the day of the week on which to rotate the audit files. Valid values are `sun`, `mon`, `tue`, `wed`, `thu`, `fri`, `sat`, `every`, and `none`. To disable rotation, set the value to `none`. To rotate the file every day, set the `edb_audit_rotation_day` value to `every`. To rotate the file on a specific day of the week, set the value to the desired day of the week. `every` is the default value.

`edb_audit_rotation_size`

Specifies a file size threshold in megabytes when file rotation will be forced to occur. The default value is 0 MB. If the parameter is commented out or set to 0, rotation of the file on a size basis will not occur.

`edb_audit_rotation_seconds`

Specifies the rotation time in seconds when a new log file should be created. To disable this feature, set this parameter to 0, which is the default.

`edb_audit_connect`

Enables auditing of database connection attempts by users. To disable auditing of all connection attempts, set `edb_audit_connect` to `none`. To audit all failed connection attempts, set the value to `failed`, which is the default. To audit all connection attempts, set the value to `all`.

`edb_audit_disconnect`

Enables auditing of database disconnections by connected users. To enable auditing of disconnections, set the value to `all`. To disable, set the value to `none`, which is the default.

`edb_audit_statement`

This configuration parameter is used to specify auditing of different categories of SQL statements. Various combinations of the following values may be specified: `none`, `dml`, `insert`, `update`, `delete`, `truncate`, `select`, `error`, `rollback`, `ddl`, `create`, `drop`, `alter`, `grant`, `revoke`, and `all`. The default is `ddl` and `error`. See [Selecting SQL Statements to Audit](#) for information regarding the setting of this parameter.

`edb_audit_tag`

Use this configuration parameter to specify a string value that will be included in the audit log file for each entry as a tracking tag.

`edb_log_every_bulk_value`

Bulk processing logs the resulting statements into both the Advanced Server log file and the EDB Audit log file. However, logging each and every statement in bulk processing is costly. This can be controlled by the `edb_log_every_bulk_value` configuration parameter. When set to `true`, each and every statement in bulk processing is logged. During bulk execution, when `edb_log_every_bulk_value` is set to `false`, a log message is recorded once per bulk processing along with the number of rows processed. In addition, the duration is emitted once per bulk processing. Default is set to `false`.

`edb_audit_destination`

Specifies whether the audit log information is to be recorded in the directory as given by the `edb_audit_directory` parameter or to the directory and file managed by the `syslog` process. Set to `file` to use the directory specified by `edb_audit_directory`, which is the default setting.

Set to `syslog` to use the `syslog` process and its location as configured in the `/etc/syslog.conf` file. The `syslog` setting is valid for Advanced Server running on a Linux host and is not supported on Windows systems. **Note:** In recent Linux versions, `syslog` has been replaced by `rsyslog` and the configuration file is in `/etc/rsyslog.conf`.

The following section describes selection of specific SQL statements for auditing using the `edb_audit_statement` parameter.

11.8.2 Selecting SQL Statements to Audit

The `edb_audit_statement` permits inclusion of one or more, comma-separated values to control which SQL statements are to be audited. The following is the general format:

```
edb_audit_statement = 'value_1[, value_2]...'
```

The comma-separated values may include or omit space characters following the comma. The values can be specified in any combination of lowercase or uppercase letters.

The basic parameter values are the following:

- `all` – Results in the auditing and logging of every statement including any error messages on statements.
- `none` – Disables all auditing and logging. A value of `none` overrides any other value included in the list.
- `ddl` – Results in the auditing of all data definition language (DDL) statements (`CREATE`, `ALTER`, and `DROP`) as well as `GRANT` and `REVOKE` data control language (DCL) statements.
- `dml` – Results in the auditing of all data manipulation language (DML) statements (`INSERT`, `UPDATE`, `DELETE`, and `TRUNCATE`).
- `select` – Results in the auditing of `SELECT` statements.
- `rollback` – Results in the auditing of `ROLLBACK` statements.
- `error` – Results in the logging of all error messages that occur. Unless the `error` value is included, no error messages are logged regarding any errors that occur on SQL statements related to any of the other preceding parameter values except when `all` is used.

[Data Definition Language and Data Control Language Statements](#) describes additional parameter values for selecting particular DDL or DCL statements for auditing.

[Data Manipulation Language Statements](#) describes additional parameter values for selecting particular DML statements for auditing.

If an unsupported value is included in the `edb_audit_statement` parameter, then an error occurs when applying the setting to the database server. See the database server log file for the error such as in the following example where `create materialized vw` results in the error. (The correct value is `create materialized view`.)

```
2017-07-16 11:20:39 EDT LOG: invalid value for parameter
"edb_audit_statement": "create materialized vw, create sequence, grant"
2017-07-16 11:20:39 EDT FATAL: configuration file "/var/lib/edb/as13/data/
postgresql.conf" contains errors
```

The following sections describe the values for the SQL language types DDL, DCL, and DML.

Data Definition Language and Data Control Language Statements

This section describes values that can be included in the `edb_audit_statement` parameter to audit DDL and DCL statements.

The following general rules apply:

- If the `edb_audit_statement` parameter includes either `ddl` or `all`, then all DDL statements are audited. In addition, the DCL statements `GRANT` and `REVOKE` are audited as well.
- If the `edb_audit_statement` is set to `none`, then no DDL nor DCL statements are audited.
- Specific types of DDL and DCL statements can be chosen for auditing by including a combination of values within the `edb_audit_statement` parameter.

Use the following syntax to specify an `edb_audit_statement` parameter value for DDL statements:

```
> { create | alter | drop } [ <object_type> ]
```

`object_type` is any of the following:

`ACCESS METHOD`

`AGGREGATE`

`CAST`

`COLLATION`

`CONVERSION`

`DATABASE`

[EVENT TRIGGER](#)[EXTENSION](#)[FOREIGN TABLE](#)[FUNCTION](#)[INDEX](#)[LANGUAGE](#)[LARGE OBJECT](#)[MATERIALIZED VIEW](#)[OPERATOR](#)[OPERATOR CLASS](#)[OPERATOR FAMILY](#)[POLICY](#)[PUBLICATION](#)[ROLE](#)[RULE](#)[SCHEMA](#)[SEQUENCE](#)[SERVER](#)[SUBSCRIPTION](#)[TABLE](#)[TABLESPACE](#)[TEXT SEARCH CONFIGURATION](#)[TEXT SEARCH DICTIONARY](#)[TEXT SEARCH PARSER](#)[TEXT SEARCH TEMPLATE](#)[TRANSFORM](#)[TRIGGER](#)[TYPE](#)[USER MAPPING](#)[VIEW](#)

Descriptions of object types as used in SQL commands can be found in the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-commands.html>

If `object_type` is omitted from the parameter value, then all of the specified command statements (either `create`, `alter`, or `drop`) are audited.

Use the following syntax to specify an `edb_audit_statement` parameter value for DCL statements:

```
> { grant | revoke }
```

The following are some DDL and DCL examples.

Example 1

The following is an example where `edb_audit_connect` and `edb_audit_statement` are set with the following non-default values:

```
logging_collector = 'on'
edb_audit_connect = 'all'
edb_audit_statement = 'create, alter, error'
```

Thus, only SQL statements invoked by the `CREATE` and `ALTER` commands are audited. Error messages are also included in the audit log.

The database session that occurs is the following:

```
$ psql edb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

edb=# SHOW edb_audit_connect;
edb_audit_connect
-----
all
(1 row)

edb=# SHOW edb_audit_statement;
edb_audit_statement
-----
create, alter, error
(1 row)

edb=# CREATE ROLE adminuser;
CREATE ROLE
edb=# ALTER ROLE adminuser WITH LOGIN, SUPERUSER, PASSWORD 'password';
ERROR: syntax error at or near ","
LINE 1: ALTER ROLE adminuser WITH LOGIN, SUPERUSER, PASSWORD 'passwo...
          ^
edb=# ALTER ROLE adminuser WITH LOGIN SUPERUSER PASSWORD 'password';
ALTER ROLE
edb=# CREATE DATABASE auditdb;
CREATE DATABASE
edb=# ALTER DATABASE auditdb OWNER TO adminuser;
ALTER DATABASE
edb=# \c auditdb adminuser
Password for user adminuser:
You are now connected to database "auditdb" as user "adminuser".
auditdb=# CREATE SCHEMA edb;
```

```

CREATE SCHEMA
auditdb=# SET search_path TO edb;
SET
auditdb=# CREATE TABLE department (
auditdb(# deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
auditdb(# dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
auditdb(# loc        VARCHAR2(13)
auditdb(# );
CREATE TABLE
auditdb=# DROP TABLE department;
DROP TABLE
auditdb=# CREATE TABLE dept (
auditdb(# deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
auditdb(# dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
auditdb(# loc        VARCHAR2(13)
auditdb(# );
CREATE TABLE

```

The resulting audit log file contains the following.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-25 12:32:22.799 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,1,
"authentication",2020-05-25 12:32:22 IST,4/19,0,AUDIT,00000,"connection authorized:
user=enterprisedb database=edb",,,,,,,,"client backend","","connect"

```

```

2020-05-25 12:34:05.843 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,2,
"idle",2020-05-25 12:32:22 IST,4/23,0,AUDIT,00000,"statement: CREATE ROLE adminuser:",
,,,,,,,"psql","client backend","CREATE ROLE","create"

```

```

2020-05-25 12:34:16.617 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,3,
"idle",2020-05-25 12:32:22 IST,4/24,0,ERROR,42601,"syntax error at or near ""","",
,"ALTER ROLE adminuser WITH LOGIN, SUPERUSER, PASSWORD 'password';",32,,,"psql",
"client backend","","error"

```

```

2020-05-25 12:34:29.954 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,4,
"idle",2020-05-25 12:32:22 IST,4/25,0,AUDIT,00000,"statement: ALTER ROLE adminuser
WITH LOGIN SUPERUSER PASSWORD 'password';",,,,,,,,"psql","client backend","ALTER ROLE
","alter"

```

```

2020-05-25 12:34:40.114 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,5,
"idle",2020-05-25 12:32:22 IST,4/26,0,AUDIT,00000,"statement: CREATE DATABASE auditdb;
",,,,,,,,"psql","client backend","CREATE DATABASE","create"

```

```

2020-05-25 12:34:50.591 IST,"enterprisedb","edb",72518,"[local]",5ecb6d7e.11b46,6,
"idle",2020-05-25 12:32:22 IST,4/27,0,AUDIT,00000,"statement: ALTER DATABASE auditdb
OWNER TO adminuser;",,,,,,,,"psql","client backend","ALTER DATABASE","alter"

```

```

2020-05-25 12:35:01.554 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,1,
"authentication",2020-05-25 12:35:01 IST,5/11,0,AUDIT,00000,"connection authorized:
user=adminuser database=auditdb",,,,,,,,"client backend","","connect"

```

```

2020-05-25 12:35:12.931 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,2,
"idle",2020-05-25 12:35:01 IST,5/13,0,AUDIT,00000,"statement: CREATE SCHEMA edb;",
,,,,"psql","client backend","CREATE SCHEMA","create"

```

```
2020-05-25 12:37:18.547 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,3,
"idle",2020-05-25 12:35:01 IST,5/15,0,AUDIT,00000,"statement: CREATE TABLE department
(
    deptno  NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname   VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc     VARCHAR2(13)
);",,,,,,,,"psql","client backend","CREATE TABLE","create"
```

```
2020-05-25 12:39:09.065 IST,"adminuser","auditdb",75531,"[local]",5ecb6e1d.1270b,4,
"idle",2020-05-25 12:35:01 IST,5/17,0,AUDIT,00000,"statement: CREATE TABLE dept (
    deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc         VARCHAR2(13)
);",,,,,,,,"psql","client backend","CREATE TABLE","create"
```

The `CREATE` and `ALTER` statements for the `adminuser` role and `auditdb` database are audited. The error for the `ALTER ROLE adminuser` statement is also logged since error is included in the `edb_audit_statement` parameter.

Similarly, the `CREATE` statements for schema `edb` and tables `department` and `dept` are audited.

Note that the `DROP TABLE department` statement is not in the audit log since there is no `edb_audit_statement` setting that would result in the auditing of successfully processed `DROP` statements such as `ddl`, `all`, or `drop`.

Example 2

The following is an example where `edb_audit_connect` and `edb_audit_statement` are set with the following non-default values:

```
logging_collector = 'on'
edb_audit_connect = 'all'
edb_audit_statement = create view,create materialized view,create
sequence,grant'
```

Thus, only SQL statements invoked by the `CREATE VIEW`, `CREATE MATERIALIZED VIEW`, `CREATE SEQUENCE` and `GRANT` commands are audited.

The database session that occurs is the following:

```
$ psql auditdb adminuser
Password for user adminuser:
psql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_connect;
edb_audit_connect
-----
all
(1 row)

auditdb=# SHOW edb_audit_statement;
edb_audit_statement
-----
create view,create materialized view,create sequence,grant
(1 row)
```

```

auditdb=# SET search_path TO edb;
SET
auditdb=# CREATE TABLE emp (
auditdb(#   empno      NUMBER(4) NOT NULL CONSTRAINT emp_pk PRIMARY KEY,
auditdb(#   ename      VARCHAR2(10),
auditdb(#   job        VARCHAR2(9),
auditdb(#   mgr        NUMBER(4),
auditdb(#   hiredate    DATE,
auditdb(#   sal        NUMBER(7,2) CONSTRAINT emp_sal_ck CHECK (sal >
0),
auditdb(#   comm      NUMBER(7,2),
auditdb(#   deptno    NUMBER(2) CONSTRAINT emp_ref_dept_fk
auditdb(#                   REFERENCES dept(deptno)
auditdb(# );
CREATE TABLE
auditdb=# CREATE VIEW salesemp AS
auditdb#   SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job =
'SALESMAN';
CREATE VIEW
auditdb=# CREATE MATERIALIZED VIEW managers AS
auditdb#   SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job =
'MANAGER';
SELECT 0
auditdb=# CREATE SEQUENCE next_empno START WITH 8000 INCREMENT BY 1;
CREATE SEQUENCE
auditdb=# GRANT ALL ON dept TO PUBLIC;
GRANT
auditdb=# GRANT ALL ON emp TO PUBLIC;
GRANT

```

The resulting audit log file contains the following.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-25 14:05:29.163 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,1,
"authentication",2020-05-25 14:05:29 IST,4/28,0,AUDIT,00000,"connection authorized:
user=adminuser database=auditdb",,,,,,,,"client backend","","connect"

```

```

2020-05-25 14:12:06.318 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,2,
"idle",2020-05-25 14:05:29 IST,4/34,0,AUDIT,00000,"statement: CREATE VIEW salesemp AS
SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job = 'SALESMAN';",,,,,,,,
"psql","client backend","CREATE VIEW","create"

```

```

2020-05-25 14:13:26.657 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,3,
"idle",2020-05-25 14:05:29 IST,4/36,0,AUDIT,00000,"statement: CREATE MATERIALIZED
VIEW managers AS SELECT empno, ename, hiredate, sal, comm FROM emp WHERE job =
'MANAGER';",,,,,,,,"psql","client backend","CREATE MATERIALIZED VIEW","create"

```

```

2020-05-25 14:13:38.928 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,4,
"idle",2020-05-25 14:05:29 IST,4/37,0,AUDIT,00000,"statement: CREATE SEQUENCE next_
empno START WITH 8000 INCREMENT BY 1;","",,"psql","client backend","CREATE
SEQUENCE","create"

```

```

2020-05-25 14:13:51.434 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,5,
"idle",2020-05-25 14:05:29 IST,4/38,0,AUDIT,00000,"statement: GRANT ALL ON dept TO

```

```
PUBLIC;,,,,,,,"psql","client backend","GRANT","grant"
```

```
2020-05-25 14:14:03.737 IST,"adminuser","auditdb",40810,"[local]",5ecb8351.9f6a,6,
"idle",2020-05-25 14:05:29 IST,4/39,0,AUDIT,00000,"statement: GRANT ALL ON emp TO
PUBLIC;,,,,,,,"psql","client backend","GRANT","grant"
```

The `CREATE VIEW` and `CREATE MATERIALIZED VIEW` statements are audited. Note that the prior `CREATE TABLE emp` statement is not audited since none of the values `create`, `create table`, `ddl`, nor `all` are included in the `edb_audit_statement` parameter.

The `CREATE SEQUENCE` and `GRANT` statements are audited since those values are included in the `edb_audit_statement` parameter.

Data Manipulation Language Statements

This section describes the values that can be included in the `edb_audit_statement` parameter to audit DML statements.

The following general rules apply:

- If the `edb_audit_statement` parameter includes either `dml` or `all`, then all DML statements are audited.
- If the `edb_audit_statement` is set to `none`, then no DML statements are audited.
- Specific types of DML statements can be chosen for auditing by including a combination of values within the `edb_audit_statement` parameter.

Use the following syntax to specify an `edb_audit_statement` parameter value for `SQL INSERT`, `UPDATE`, `DELETE`, or `TRUNCATE` statements:

```
> { insert | update | delete | truncate }
```

Example

The following is an example where `edb_audit_connect` and `edb_audit_statement` are set with the following non-default values:

```
logging_collector = 'on'
edb_audit_connect = 'all'
edb_audit_statement = 'UPDATE, DELETE, error'
```

Thus, only SQL statements invoked by the `UPDATE` and `DELETE` commands are audited. All errors are also included in the audit log (even errors not related to the `UPDATE` and `DELETE` commands).

The database session that occurs is the following:

```
$ psql auditdb adminuser
Password for user adminuser:
pgsql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_connect;
edb_audit_connect
-----
all
(1 row)

auditdb=# SHOW edb_audit_statement;
edb_audit_statement
```

 UPDATE, DELETE, error
 (1 row)

```
auditdb=# SET search_path TO edb;
SET
auditdb=# INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT 0 1
auditdb=# INSERT INTO dept VALUES (20,'RESEARCH','DALLAS');
INSERT 0 1
auditdb=# INSERT INTO dept VALUES (30,'SALES','CHICAGO');
INSERT 0 1
auditdb=# INSERT INTO dept VALUES (40,'OPERATIONS','BOSTON');
INSERT 0 1
auditdb=# INSERT INTO emp VALUES (7369,'SMITH','CLERK',7902,'17-DEC-80',800,NULL,20);
INSERT 0 1
auditdb=# INSERT INTO emp VALUES (7499,'ALLEN','SALESMAN',7698,'20-FEB-81',1600,300,30);
INSERT 0 1
auditdb=# INSERT INTO emp VALUES (7521,'WARD','SALESMAN',7698,'22-FEB-81',1250,500,30);
INSERT 0 1
.
.
.
auditdb=# INSERT INTO emp VALUES (7934,'MILLER','CLERK',7782,'23-JAN-82',1300,NULL,10);
INSERT 0 1
auditdb=# UPDATE dept SET loc = 'BEDFORD' WHERE deptno = 40;
UPDATE 1
auditdb=# SELECT * FROM dept;
deptno | dname   | loc
-----+-----+
 10 | ACCOUNTING | NEW YORK
 20 | RESEARCH   | DALLAS
 30 | SALES     | CHICAGO
 40 | OPERATIONS | BEDFORD
(4 rows)
```

```
auditdb=# DELETE FROM emp WHERE deptno = 10;
DELETE 3
auditdb=# TRUNCATE employee;
ERROR: relation "employee" does not exist
auditdb=# TRUNCATE emp;
TRUNCATE TABLE
auditdb=# \q
```

The resulting audit log file contains the following.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```
2020-05-25 15:17:16.392 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,1,
"authentication",2020-05-25 15:17:16 IST,5/3,0,AUDIT,00000,"connection authorized:
user=adminuser database=auditdb","","","","client backend","","connect"
```

```
2020-05-25 15:19:18.066 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,2,
"idle",2020-05-25 15:17:16 IST,5/14,0,AUDIT,00000,"statement: UPDATE dept SET loc =
'BEDFORD' WHERE deptno = 40;","","pgsql","client backend","UPDATE","update"
```

```
2020-05-25 15:19:38.524 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,3,
"idle",2020-05-25 15:17:16 IST,5/16,0,AUDIT,00000,"statement: DELETE FROM emp WHERE
deptno = 10;","","pgsql","client backend","DELETE","delete"
```

```
2020-05-25 15:19:48.149 IST,"adminuser","auditdb",123420,"[local]",5ecb9424.1e21c,4,
"TRUNCATE TABLE",2020-05-25 15:17:16 IST,5/17,0,ERROR,42P01,"relation ""employee"""
does not exist","","TRUNCATE employee;","","pgsql","client backend","TRUNCATE TABLE",
"error"
```

The `UPDATE dept` and `DELETE FROM emp` statements are audited. Note that all of the prior `INSERT` statements are not audited since none of the values `insert`, `dml`, nor `all` are included in the `edb_audit_statement` parameter.

The `SELECT * FROM dept` statement is not audited as well since neither `select` nor `all` is included in the `edb_audit_statement` parameter.

Since `error` is specified in the `edb_audit_statement` parameter, but not the `truncate` value, the error on the `TRUNCATE employee` statement is logged in the audit file, but not the successful `TRUNCATE emp` statement.

11.8.3 Enabling Audit Logging

The following steps describe how to configure Advanced Server to log all connections, disconnections, DDL statements, DCL statements, DML statements, and any statements resulting in an error.

1. Enable auditing by the setting the `edb_audit` parameter to `xml` or `csv`.
2. Set the file rotation day when the new file will be created by setting the parameter `edb_audit_rotation_day` to the desired value.
3. To audit all connections, set the parameter, `edb_audit_connect`, to `all`.
4. To audit all disconnections, set the parameter, `edb_audit_disconnect`, to `all`.
5. To audit DDL, DCL, DML and other statements, set the parameter, `edb_audit_statement` according to the instructions in [Selecting SQL Statements to Audit](#).

The setting of the `edb_audit_statement` parameter in the configuration file affects the entire database cluster.

The type of statements that are audited as controlled by the `edb_audit_statement` parameter can be further refined according to the database in use as well as the database role running the session:

- The `edb_audit_statement` parameter can be set as an attribute of a specified database with the `ALTER DATABASE dbname SET edb_audit_statement` command. This setting overrides the `edb_audit_statement` parameter in the configuration file for statements executed when connected to database `dbname`.
- The `edb_audit_statement` parameter can be set as an attribute of a specified role with the `ALTER ROLE rolename SET edb_audit_statement` command. This setting overrides the `edb_audit_statement` parameter in the configuration file as well as any setting assigned to the database by the previously described `ALTER DATABASE` command when the specified role is running the current session.
- The `edb_audit_statement` parameter can be set as an attribute of a specified role when using a specified database with the `ALTER ROLE rolename IN DATABASE dbname SET edb_audit_statement` command. This setting overrides the `edb_audit_statement` parameter in the configuration file as well as any setting assigned to the database by the previously described `ALTER DATABASE` command as well as any setting assigned to the role with the `ALTER ROLE` command without the `IN DATABASE` clause as previously described.

The following are examples of this technique.

The database cluster is established with `edb_audit_statement` set to `all` as shown in its `postgresql.conf` file:

```
logging_collector = 'on'
edb_audit_statement = 'all'    # Statement type to be audited:
                                # none, dml, insert, update, delete,
                                # truncate,
                                # select, error, rollback, ddl, create, drop,
                                # alter, grant, revoke, all
```

A database and role are established with the following settings for the `edb_audit_statement` parameter:

- Database `auditdb` with `ddl`, `insert`, `update`, and `delete`
- Role `admin` with `select` and `truncate`
- Role `admin` in database `auditdb` with `create table`, `insert`, and `update`

Creation and alteration of the database and role are shown by the following:

```
$ psql edb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

edb=# SHOW edb_audit_statement;
edb_audit_statement
-----
all
(1 row)

edb=# CREATE DATABASE auditdb;
CREATE DATABASE
edb=# ALTER DATABASE auditdb SET edb_audit_statement TO 'ddl, insert,
update, delete';
ALTER DATABASE
edb=# CREATE ROLE admin WITH LOGIN SUPERUSER PASSWORD 'password';
CREATE ROLE
edb=# ALTER ROLE admin SET edb_audit_statement TO 'select, truncate';
ALTER ROLE
edb=# ALTER ROLE admin IN DATABASE auditdb SET edb_audit_statement TO
'create table, insert, update';
ALTER ROLE
```

The following demonstrates the changes made and the resulting audit log file for three cases.

Case 1: Changes made in database `auditdb` by role `enterprisedb`. Only `ddl`, `insert`, `update`, and `delete` statements are audited:

```
$ psql auditdb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

auditdb=# SHOW edb_audit_statement;
edb_audit_statement
-----
ddl, insert, update, delete
```

(1 row)

```

auditdb=# CREATE TABLE audit_tbl (f1 INTEGER PRIMARY KEY, f2 TEXT);
CREATE TABLE
auditdb=# INSERT INTO audit_tbl VALUES (1, 'Row 1');
INSERT 0 1
auditdb=# UPDATE audit_tbl SET f2 = 'Row A' WHERE f1 = 1;
UPDATE 1
auditdb=# SELECT * FROM audit_tbl;           <== Should not be audited
f1 | f2
-----+
 1 | Row A
(1 row)

```

```

auditdb=# TRUNCATE audit_tbl;           <== Should not be audited
TRUNCATE TABLE

```

The following audit log file shows entries only for the `CREATE TABLE`, `INSERT INTO audit_tbl`, and `UPDATE audit_tbl` statements. The `SELECT * FROM audit_tbl` and `TRUNCATE audit_tbl` statements were not audited.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```

2020-05-25 15:59:12.332 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,2,
"idle",2020-05-25 15:59:02 IST,5/7,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl
(f1 INTEGER PRIMARY KEY, f2 TEXT);",,,,,,,,"psql","client backend","CREATE TABLE","",",
"create"

```

```

2020-05-25 15:59:22.419 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,3,
"idle",2020-05-25 15:59:02 IST,5/8,0,AUDIT,00000,"statement: INSERT INTO audit_tbl
VALUES (1, 'Row 1');",,,,,,,,"psql","client backend","INSERT","",,"insert"

```

```

2020-05-25 15:59:32.180 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,4,
"idle",2020-05-25 15:59:02 IST,5/9,0,AUDIT,00000,"statement: UPDATE audit_tbl SET f2
= 'Row A' WHERE f1 = 1;",,,,,,,,"psql","client backend","UPDATE","",,"update"

```

Case 2: Changes made in database `edb` by role `admin`. Only `select` and `truncate` statements are audited:

```

$ psql edb admin
Password for user admin:
psql.bin (13.0.0)
Type "help" for help.

```

```

edb=# SHOW edb_audit_statement;
edb_audit_statement
-----
select, truncate
(1 row)

```

```

edb=# CREATE TABLE edb_tbl (f1 INTEGER PRIMARY KEY, f2 TEXT); <== Should not
be audited
CREATE TABLE
edb=# INSERT INTO edb_tbl VALUES (1, 'Row 1');           <== Should not
be audited
INSERT 0 1
edb=# SELECT * FROM edb_tbl;

```

```
f1 | f2
----+-----
 1 | Row 1
(1 row)
```

```
edb=# TRUNCATE edb_tbl;
TRUNCATE TABLE
```

Continuation of the audit log file now appears as follows. The last two entries representing the second case show only the `SELECT * FROM edb_tbl` and `TRUNCATE edb_tbl` statements. The `CREATE TABLE edb_tbl` and `INSERT INTO edb_tbl` statements were not audited.

```
2020-05-25 15:59:12.332 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,2,
"idle",2020-05-25 15:59:02 IST,5/7,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl
(f1 INTEGER PRIMARY KEY, f2 TEXT);",,,,,,,,"psql","client backend","CREATE TABLE","","
"create"
```

```
2020-05-25 15:59:22.419 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,3,
"idle",2020-05-25 15:59:02 IST,5/8,0,AUDIT,00000,"statement: INSERT INTO audit_tbl
VALUES (1, 'Row 1');",,,,,,,,"psql","client backend","INSERT","", "insert"
```

```
2020-05-25 15:59:32.180 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,4,
"idle",2020-05-25 15:59:02 IST,5/9,0,AUDIT,00000,"statement: UPDATE audit_tbl SET f2
= 'Row A' WHERE f1 = 1; ",,,,,,,,"psql","client backend","UPDATE","", "update"
```

```
2020-05-25 16:25:37.509 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,2,"idle",
2020-05-25 16:24:54 IST,4/33,0,AUDIT,00000,"statement: SELECT * FROM edb_tbl; ",,,,
,"psql","client backend","SELECT","", "select"
```

```
2020-05-25 16:25:46.131 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,3,"idle",
2020-05-25 16:24:54 IST,4/34,0,AUDIT,00000,"statement: TRUNCATE edb_tbl; ",,,,
,"psql","client backend","TRUNCATE TABLE","", "truncate"
```

Case 3: Changes made in database `auditdb` by role `admin`. Only `create table`, `insert`, and `update` statements are audited:

```
$ psql auditdb admin
Password for user admin:
psql.bin (13.0.0)
Type "help" for help.
```

```
auditdb=# SHOW edb_audit_statement;
edb_audit_statement
-----
create table, insert, update
(1 row)

auditdb=# CREATE TABLE audit_tbl_2 (f1 INTEGER PRIMARY KEY, f2 TEXT);
CREATE TABLE
auditdb=# INSERT INTO audit_tbl_2 VALUES (1, 'Row 1');
INSERT 0 1
auditdb=# SELECT * FROM audit_tbl_2;           <== Should not be audited
f1 | f2
----+-----
 1 | Row 1
(1 row)
```

```
auditdb=# TRUNCATE audit_tbl_2;           <== Should not be audited
TRUNCATE TABLE
```

Continuation of the audit log file now appears as follows. The next to last two entries representing the third case show only `CREATE TABLE audit_tbl_2` and `INSERT INTO audit_tbl_2` statements. The `SELECT * FROM audit_tbl_2` and `TRUNCATE audit_tbl_2` statements were not audited.

```
2020-05-25 15:59:12.332 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,2,
"idle",2020-05-25 15:59:02 IST,5/7,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl
(f1 INTEGER PRIMARY KEY, f2 TEXT);",,,,,,,,"psql","client backend","CREATE TABLE","",",
"create"
```

```
2020-05-25 15:59:22.419 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,3,
"idle",2020-05-25 15:59:02 IST,5/8,0,AUDIT,00000,"statement: INSERT INTO audit_tbl
VALUES (1, 'Row 1');",,,,,,,,"psql","client backend","INSERT","",,"insert"
```

```
2020-05-25 15:59:32.180 IST,"enterprisedb","auditdb",41837,"[local]",5ecb9dee.a36d,4,
"idle",2020-05-25 15:59:02 IST,5/9,0,AUDIT,00000,"statement: UPDATE audit_tbl SET f2
= 'Row A' WHERE f1 = 1;",,,,,,,,"psql","client backend","UPDATE","",,"update"
```

```
2020-05-25 16:25:37.509 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,2,"idle",
2020-05-25 16:24:54 IST,4/33,0,AUDIT,00000,"statement: SELECT * FROM edb_tbl;","",,
,"psql","client backend","SELECT","",,"select"
```

```
2020-05-25 16:25:46.131 IST,"admin","edb",71920,"[local]",5ecba3fe.118f0,3,"idle",
2020-05-25 16:24:54 IST,4/34,0,AUDIT,00000,"statement: TRUNCATE edb_tbl;","",,
,"psql","client backend","TRUNCATE TABLE","",,"truncate"
```

```
2020-05-25 17:30:59.057 IST,"admin","auditdb",122093,"[local]",5ecbb370.1dced,2,"idle
",2020-05-25 17:30:48 IST,5/11,0,AUDIT,00000,"statement: CREATE TABLE audit_tbl_2 (f1
INTEGER PRIMARY KEY, f2 TEXT);",,,,,,,,"psql","client backend","CREATE TABLE","",",
"create"
```

```
2020-05-25 17:31:08.866 IST,"admin","auditdb",122093,"[local]",5ecbb370.1dced,3,"idle
",2020-05-25 17:30:48 IST,5/12,0,AUDIT,00000,"statement: INSERT INTO audit_tbl_2
VALUES (1, 'Row 1');",,,,,,,,"psql","client backend","INSERT","",,"insert"
```

11.8.4 Audit Log File

The audit log file can be generated in either CSV or XML format depending upon the setting of the `edb_audit` configuration parameter.

The information in the audit log is based on the logging performed by PostgreSQL as described in the section “Using CSV-Format Log Output” within Section “Error Reporting and Logging” in the PostgreSQL core documentation, available at:

<https://www.postgresql.org/docs/current/static/runtime-config-logging.html>

The following table lists the fields in the order they appear in the XML audit log format. The table contains the following information:

- Field.** Name of the field as shown in the sample table definition in the PostgreSQL documentation as previously referenced.
- XML Element/Attribute.** For the XML format, name of the XML element and its attribute (if used), referencing the value. **Note:** *n/a* indicates that there is no XML representation for this field.
- Data Type.** Data type of the field as given by the PostgreSQL sample table definition.
- Description.** Description of the field. For certain fields, no output is generated in the audit log as those fields are not supported by auditing. Those fields are indicated by "Not supported".

The fields that do not have any values for logging appear as consecutive commas (,,) in the CSV format.

| Field | XML Element/Attribute | Data Type | Description |
|------------------------|------------------------------|--------------------------|---|
| user_name | event/user | text | Database user who executed the statement. |
| database_name | event/database | text | Database in which the statement was executed. |
| process_id | event/process_id | integer | Operating system process ID in which the statement was executed. |
| connection_from | event/remote_host | text | Host and port location from where the statement was executed. |
| session_id | event/session_id | text | Session ID in which the statement was executed. |
| session_line_num | event/session_line_num | bigint | Order of the statement within the session. |
| process_status | event/process_status | text | Processing status. |
| session_start_time | event/session_start_time | timestamp with time zone | Date/time when the session was started. |
| log_time | event/log_time | timestamp with time zone | Log date/time of the statement. |
| virtual_transaction_id | event/virtual_transaction_id | text | Virtual transaction ID of the statement. |
| transaction_id | event/transaction_id | bigint | Regular transaction ID of the statement. |
| type | event/type | text | Determines the audit event_type to identify messages in the log. |
| sql_state_code | event/sql_state_code | text | SQL state code returned for the statement. The sql_state_code is not logged when its value is 00000 for XML log format. |
| command_tag | event/command_tag | text | SQL command of the statement. |
| audit_tag | event/audit_tag | text | Value specified by the audit_tag parameter in the configuration file. |
| application_name | event/application_name | text | Name of the application from which the statement was executed. (for example, psql.bin). |
| backend_type | event/backend_type | text | The backend_type corresponds to what pg_stat_activity.backend_type shows and is added as a column to the csv log. |
| error_severity | error_severity | text | Statement severity. Values are AUDIT for audited statements and ERROR for any resulting error messages. |
| message | message | text | The SQL statement that was attempted for execution. |
| detail | detail | text | Error message detail. |
| hint | hint | text | Hint for error. |
| internal_query | internal_query | text | Internal query that led to the error, if any. |
| internal_query_pos | internal_query_pos | integer | Character count of the error position therein. |
| context | context | text | Error context. |

| Field | XML Element/Attribute | Data Type | Description |
|-----------|-----------------------|-----------|--|
| query | query | text | User query that led to the error. For (For errors only) |
| query_pos | query_pos | integer | Character count of the error position therein. (For errors only) |
| location | location | text | Location of the error in the source code. The location field will be populated if log_error_verbosity is set to verbose. |

The following examples are generated in the CSV and XML formats.

The non-default audit settings in the `postgresql.conf` file are as follows:

```
logging_collector = 'on'
edb_audit = 'csv'
edb_audit_connect = 'all'
edb_audit_disconnect = 'all'
edb_audit_statement = 'ddl, dml, select, error'
edb_audit_tag = 'edbaudit'
```

The `edb_audit` parameter is changed to `xml` when generating the XML format.

The audited session is the following:

```
$ psql edb enterprisedb
Password for user enterprisedb:
psql.bin (13.0.0)
Type "help" for help.

edb=# CREATE SCHEMA edb;
CREATE SCHEMA
edb=# SET search_path TO edb;
SET
edb=# CREATE TABLE dept (
edb(#   deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
edb(#   dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
edb(#   loc         VARCHAR2(13)
edb(# );
CREATE TABLE
edb=# INSERT INTO dept VALUES (10,'ACCOUNTING','NEW YORK');
INSERT 0 1
edb=# UPDATE department SET loc = 'BOSTON' WHERE deptno = 10;
ERROR: relation "department" does not exist
LINE 1: UPDATE department SET loc = 'BOSTON' WHERE deptno = 10;
^
edb=# UPDATE dept SET loc = 'BOSTON' WHERE deptno = 10;
UPDATE 1
edb=# SELECT * FROM dept;
deptno | dname  | loc
-----+-----+
 10 | ACCOUNTING | BOSTON
(1 row)

edb=# \q
```

CSV Audit Log File

The following is the CSV format of the audit log file.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```
2020-05-22 16:53:37.817 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,1,
"authentication",2020-05-22 16:53:37 IST,4/21,0,AUDIT,00000,"connection authorized:
user=enterprisedb database=edb",,,,,,,,"client backend","",,"edbaudit","connect"
```

```
2020-05-22 16:53:42.279 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,2,"idle
",2020-05-22 16:53:37 IST,4/23,0,AUDIT,00000,"statement: CREATE SCHEMA edb;",
"psql","client backend","CREATE SCHEMA","edbaudit","create"
```

```
2020-05-22 16:54:07.896 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,3,"idle
",2020-05-22 16:53:37 IST,4/25,0,AUDIT,00000,"statement: CREATE TABLE dept (
    deptno      NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname       VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc         VARCHAR2(13)
);",,,,,,,,"psql","client backend","CREATE TABLE","edbaudit","create"
```

```
2020-05-22 16:54:20.500 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,4,"idle
",2020-05-22 16:53:37 IST,4/26,0,AUDIT,00000,"statement: INSERT INTO dept VALUES
(10,'ACCOUNTING','NEW YORK');",,,,,,,,"psql","client backend","INSERT","edbaudit",
"insert"
```

```
2020-05-22 16:54:34.821 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,5,"idle
",2020-05-22 16:53:37 IST,4/27,0,AUDIT,00000,"statement: UPDATE department SET loc =
'BOSTON' WHERE deptno = 10;",,,,,,,,"psql","client backend","UPDATE","edbaudit",
"update"
```

```
2020-05-22 16:54:34.821 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,6,
"UPDATE",2020-05-22 16:53:37 IST,4/27,0,ERROR,42P01,"relation ""department"" does not
exist",,,,,"UPDATE department SET loc = 'BOSTON' WHERE deptno = 10;",8,,,"psql",
"client backend","UPDATE","edbaudit","error"
```

```
2020-05-22 16:54:51.308 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,7,"idle
",2020-05-22 16:53:37 IST,4/28,0,AUDIT,00000,"statement: UPDATE dept SET loc = 'BOSTON
' WHERE deptno = 10;",,,,,,,,"psql","client backend","UPDATE","edbaudit","update"
```

```
2020-05-22 16:55:00.774 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,8,"idle
",2020-05-22 16:53:37 IST,4/29,0,AUDIT,00000,"statement: SELECT * FROM dept;",
"psql","client backend","SELECT","edbaudit","select"
```

```
2020-05-22 16:55:06.548 IST,"enterprisedb","edb",55279,"[local]",5ec7b639.d7ef,9,"idle
",2020-05-22 16:53:37 IST,,0,AUDIT,00000,"disconnection: session time: 0:01:28.732
user=enterprisedb database=edb host=[local]",,,,,,,,"psql","client backend","",",
"edbaudit","disconnect"
```

XML Audit Log File

The following is the XML format of the audit log file. The output has been formatted for more clarity in the appearance in the example.

```
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="1" process_status="authentication"
session_start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:11:17.806
```

```

IST" virtual_transaction_id="4/19" type="connect" audit_tag="edbaudit" backend_
type="client backend">
<error_severity>AUDIT</error_severity>
<message>connection authorized: user=enterprisedb database=edb</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="2" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:11:32.558 IST"
virtual_transaction_id="4/21" type="create" command_tag="CREATE SCHEMA" audit_
tag="edbaudit" application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: CREATE SCHEMA edb;</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="3" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:00.199 IST"
virtual_transaction_id="4/23" type="create" command_tag="CREATE TABLE" audit_
tag="edbaudit" application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: CREATE TABLE dept (
    deptno  NUMBER(2) NOT NULL CONSTRAINT dept_pk PRIMARY KEY,
    dname   VARCHAR2(14) CONSTRAINT dept_dname_uq UNIQUE,
    loc     VARCHAR2(13));
</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="4" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:10.992 IST"
virtual_transaction_id="4/24" type="insert" command_tag="INSERT" audit_tag=
"edbaudit" application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: INSERT INTO dept VALUES
    (10,'ACCOUNTING','NEW YORK');
</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="5" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:21.764 IST"
virtual_transaction_id="4/25" type="update" command_tag="UPDATE" audit_tag=
"edbaudit" application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: UPDATE department SET loc = 'BOSTON'
    WHERE deptno = 10;
</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="6" process_status="UPDATE"
session_start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:21.765
IST" virtual_transaction_id="4/25" type="error" sql_state_code="42P01" command_
tag="UPDATE" audit_tag="edbaudit" application_name="psql" backend_type="client
backend">
<error_severity>ERROR</error_severity>
<message>relation "department" does not exist</message>
<query>UPDATE department SET loc = 'BOSTON' WHERE deptno =
10;</query>

```

```

<query_pos>8</query_pos>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="7" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:34.878 IST"
virtual_transaction_id="4/26" type="update" command_tag="UPDATE" audit_tag=
"edbaudit" application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: UPDATE dept SET loc = &apos;BOSTON&apos; WHERE
deptno = 10;
</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="8" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:45.471 IST"
virtual_transaction_id="4/27" type="select" command_tag="SELECT" audit_tag=
"edbaudit" application_name="psql" backend_type="client backend">
<error_severity>AUDIT</error_severity>
<message>statement: SELECT * FROM dept;</message>
</event>
<event user="enterprisedb" database="edb" process_id="5941" remote_host="[local]"
session_id="5ec7ac4d.1735" session_line_num="9" process_status="idle" session_
start_time="2020-05-22 16:11:17 IST" log_time="2020-05-22 16:12:53.048 IST"
type="disconnect" audit_tag="edbaudit" application_name="psql" backend_type=
"client backend">
<error_severity>AUDIT</error_severity>
<message>disconnection: session time: 0:01:35.243 user=enterprisedb
database=edb host=[local]</message>
</event>
```

11.8.5 Using Error Codes to Filter Audit Logs

Advanced Server includes an extension that you can use to exclude log file entries that include a user-specified error code from the Advanced Server log files. To filter audit log entries, you must first enable the extension by modifying the `postgresql.conf` file, adding the following value to the values specified in the `shared_preload_libraries` parameter:

`$libdir/edb_filter_log`

Then, use the `edb_filter_log.errcodes` parameter to specify any error codes you wish to omit from the log files:

`edb_filter_log.errcode = 'error_code'`

Where `error_code` specifies one or more error codes that you wish to omit from the log file. Provide multiple error codes in a comma-delimited list.

For example, if `edb_filter_log` is enabled, and `edb_filter_log.errcode` is set to '`23505,23502,22012`', any log entries that return one of the following `SQLSTATE` errors:

`23505` (for violating a unique constraint)

`23502` (for violating a not-null constraint)

| 22012 (for dividing by zero)

will be omitted from the log file.

For a complete list of the error codes supported by Advanced Server audit log filtering, see the core documentation at:

| <https://www.postgresql.org/docs/current/static/errcodes-appendix.html>

11.8.6 Using Command Tags to Filter Audit Logs

Each entry in the log file except for those displaying an error message contains a *command tag*. A command tag is the SQL command executed for that particular log entry. The command tag makes it possible to use subsequent tools to scan the log file to find entries related to a particular SQL command.

The following is an example in XML form. The example has been formatted for easier review. The command tag is displayed as the `command_tag` attribute of the event element with values `CREATE ROLE`, `ALTER ROLE`, and `DROP ROLE` in the example.

```
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"  
    session_id="5ecbc7e6.faea" session_line_num="2" process_status="idle" session_  
    start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:21.147 IST"  
    virtual_transaction_id="4/30" type="create" command_tag="CREATE ROLE"  
    application_name="psql" backend_type="client backend">  
    <error_severity>AUDIT</error_severity>  
    <message>statement: CREATE ROLE newuser WITH LOGIN;</message>  
</event>  
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"  
    session_id="5ecbc7e6.faea" session_line_num="3" process_status="idle" session_  
    start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:34.142 IST"  
    virtual_transaction_id="4/31" type="error" sql_state_code="42601" application_  
    name="psql" backend_type="client backend">  
    <error_severity>ERROR</error_severity>  
    <message>unrecognized role option &quot;super&quot;</message>  
    <query>ALTER ROLE newuser WITH SUPER USER;</query>  
    <query_pos>25</query_pos>  
</event>  
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"  
    session_id="5ecbc7e6.faea" session_line_num="4" process_status="idle" session_  
    start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:44.680 IST"  
    virtual_transaction_id="4/32" type="alter" command_tag="ALTER ROLE"  
    application_name="psql" backend_type="client backend">  
    <error_severity>AUDIT</error_severity>  
    <message>statement: ALTER ROLE newuser WITH SUPERUSER;</message>  
</event>  
<event user="enterprisedb" database="edb" process_id="64234" remote_host="[local]"  
    session_id="5ecbc7e6.faea" session_line_num="5" process_status="idle" session_  
    start_time="2020-05-25 18:58:06 IST" log_time="2020-05-25 18:58:58.173 IST"  
    virtual_transaction_id="4/33" type="drop" command_tag="DROP ROLE" application_  
    name="psql" backend_type="client backend">  
    <error_severity>AUDIT</error_severity>  
    <message>statement: DROP ROLE newuser;</message>
```

```
</event>
```

The following is the same example in CSV form. The command tag is the next to last column of each entry. In the listing, the column that appears empty (""), would be the value of the `edb_audit_tag` parameter if provided.

Each audit log entry has been split and displayed across multiple lines, and a blank line has been inserted between the audit log entries for more clarity in the appearance of the results.

```
2020-05-25 19:09:32.105 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,2,  
"idle",2020-05-25 19:09:07 IST,4/30,0,AUDIT,00000,"statement: CREATE ROLE newuser  
WITH LOGIN;","","pgsql","client backend","CREATE ROLE","","create"
```

```
2020-05-25 19:09:50.975 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,3,  
"idle",2020-05-25 19:09:07 IST,4/31,0,ERROR,42601,"unrecognized role option ""super""  
","",,"ALTER ROLE newuser WITH SUPER USER;","25,,"pgsql","client backend","","","error"
```

```
2020-05-25 19:10:04.128 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,4,  
"idle",2020-05-25 19:09:07 IST,4/32,0,AUDIT,00000,"statement: ALTER ROLE newuser WITH  
SUPERUSER;","","pgsql","client backend","ALTER ROLE","","alter"
```

```
2020-05-25 19:10:15.959 IST,"enterprisedb","edb",77212,"[local]",5ecbca7b.12d9c,5,  
"idle",2020-05-25 19:09:07 IST,4/33,0,AUDIT,00000,"statement: DROP ROLE newuser;","","  
","","pgsql","client backend","DROP ROLE","","drop"
```

11.8.7 Redacting Passwords from Audit Logs

You can use the `edb_filter_log.redact_password_commands` extension to instruct the server to redact stored passwords from the log file. Note that the module only recognizes the following syntax:

```
{CREATE|ALTER} {USER|ROLE|GROUP} identifier { [WITH] [ENCRYPTED] PASSWORD  
'nonempty_string_literal' | IDENTIFIED BY { 'nonempty_string_literal' |  
bareword } } [ REPLACE { 'nonempty_string_literal' | bareword } ]
```

When such a statement is logged by `log_statement`, the server will redact the old and new passwords to 'x'. For example, the command:

```
| ALTER USER carol PASSWORD '1safepwd' REPLACE 'old_pwd';
```

Will be added to log files as:

```
| statement: ALTER USER carol PASSWORD 'x' REPLACE 'x';
```

When a statement that includes a redacted password is logged, the server redacts the statement text. When the statement is logged as context for some other message, the server omits the statement from the context.

To enable password redaction, you must first enable the extension by modifying the `postgresql.conf` file, adding the following value to the values specified in the `shared_preload_libraries` parameter:

```
| $libdir/edb_filter_log
```

Then, set `edb_filter_log.redact_password_commands` to `true`:

```
| edb_filter_log.redact_password_commands = true
```

After modifying the `postgresql.conf` file, you must restart the server for the changes to take effect.

11.9 Unicode Collation Algorithm

The *Unicode Collation Algorithm* (UCA) is a specification (*Unicode Technical Report #10*) that defines a customizable method of collating and comparing Unicode data. *Collation* means how data is sorted as with a `SELECT ... ORDER BY` clause. *Comparison* is relevant for searches that use ranges with less than, greater than, or equal to operators.

Customizability is an important factor for various reasons such as the following.

- Unicode supports a vast number of languages. Letters that may be common to several languages may be expected to collate in different orders depending upon the language.
- Characters that appear with letters in certain languages such as accents or umlauts have an impact on the expected collation depending upon the language.
- In some languages, combinations of several consecutive characters should be treated as a single character with regards to its collation sequence.
- There may be certain preferences as to the collation of letters according to case. For example, should the lowercase form of a letter collate before the uppercase form of the same letter or vice versa.
- There may be preferences as to whether punctuation marks such as underscore characters, hyphens, or space characters should be considered in the collating sequence or should they simply be ignored as if they did not exist in the string.

Given all of these variations with the vast number of languages supported by Unicode, there is a necessity for a method to select the specific criteria for determining a collating sequence. This is what the Unicode Collation Algorithm defines.

Note

In addition, another advantage for using ICU collations (the implementation of the Unicode Collation Algorithm) is for performance. Sorting tasks, including B-tree index creation, can complete in less than half the time it takes with a non-ICU collation. The exact performance gain depends on your operating system version, the language of your text data, and other factors.

The following sections provide a brief, simplified explanation of the Unified Collation Algorithm concepts. As the algorithm and its usage are quite complex with numerous variations, refer to the official documents cited in these sections for complete details.

Basic Unicode Collation Algorithm Concepts

The official information for the Unicode Collation Algorithm is specified in *Unicode Technical Report #10*, which can be found on The Unicode Consortium website at:

<http://www.unicode.org/reports/tr10/>

The ICU – International Components for Unicode also provides much useful information. An explanation of the collation concepts can be found on their website located at:

<http://userguide.icu-project.org/collation/concepts>

The basic concept behind the Unicode Collation Algorithm is the use of multilevel comparison. This means that a number of levels are defined, which are listed as level 1 through level 5 in the following bullet points. Each level defines a type of comparison. Strings are first compared using the primary level, also called level 1.

If the order can be determined based on the primary level, then the algorithm is done. If the order cannot be determined based on the primary level, then the secondary level, level 2, is applied. If the order can be determined based on the secondary level, then the algorithm is done, otherwise the tertiary level is applied, and so on. There is typically, a final tie-breaking level to determine the order if it cannot be resolved by the prior levels.

- **Level 1 – Primary Level for Base Characters.** The order of basic characters such as letters and digits determines the difference such as `A < B`.
- **Level 2 – Secondary Level for Accents.** If there are no primary level differences, then the presence or absence of accents and other such characters determine the order such as `a < á`.
- **Level 3 – Tertiary Level for Case.** If there are no primary level or secondary level differences, then a difference in case determines the order such as `a < A`.
- **Level 4 – Quaternary Level for Punctuation.** If there are no primary, secondary, or tertiary level differences, then the presence or absence of white space characters, control characters, and punctuation determine the order such as `-A < A`.
- **Level 5 – Identical Level for Tie-Breaking.** If there are no primary, secondary, tertiary, or quaternary level differences, then some other difference such as the code point values determines the order.

International Components for Unicode

The Unicode Collation Algorithm is implemented by open source software provided by the *International Components for Unicode* (ICU). The software is a set of C/C++ and Java libraries.

When Advanced Server is used to create a collation that invokes the ICU components to produce the collation, the result is referred to as an *ICU collation*.

Locale Collations

When creating a collation for a locale, a predefined ICU short form name for the given locale is typically provided.

An *ICU short form* is a method of specifying *collation attributes*, which are the properties of a collation. [Collation Attributes](#) provides additional information on collation attributes.

There are predefined ICU short forms for locales. The ICU short form for a locale incorporates the collation attribute settings typically used for the given locale. This simplifies the collation creation process by eliminating the need to specify the entire list of collation attributes for that locale.

The system catalog `pg_catalog.pg_icu_collate_names` contains a list of the names of the ICU short forms for locales. The ICU short form name is listed in column `icu_short_form`.

```
edb=# SELECT icu_short_form, valid_locale FROM pg_icu_collate_names ORDER BY
valid_locale;
icu_short_form | valid_locale
-----+-----
LAF      | af
LAR      | ar
LAS      | as
LAZ      | az
LBE      | be
LBG      | bg
LBN      | bn
LBS      | bs
LBS_ZCYRL | bs_Cyril
LROOT    | ca
LROOT    | chr
LCS      | cs
LCY      | cy
```

```

LDA      | da
LROOT    | de
LROOT    | dz
LEE      | ee
LEL      | el
LROOT    | en
LROOT    | en_US
LEN_RUS_VPOSIX | en_US_POSIX
LEO      | eo
LES      | es
LET      | et
LFA      | fa
LFA_RAF  | fa_AF
.
.
```

If needed, the default characteristics of an ICU short form for a given locale can be overridden by specifying the collation attributes to override that property. This is discussed in the next section.

Collation Attributes

When creating an ICU collation, the desired characteristics of the collation must be specified. As discussed in [Locale Collations](#), this can typically be done with an ICU short form for the desired locale. However, if more specific information is required, the specification of the collation properties can be done by using *collation attributes*.

Collation attributes define the rules of how characters are to be compared for determining the collation sequence of text strings. As Unicode covers a vast set of languages in numerous variations according to country, territory and culture, these collation attributes are quite complex.

For the complete, precise meaning and usage of collation attributes, see Section “Collator Naming Scheme” on the ICU – International Components for Unicode website at:

<http://userguide.icu-project.org/collation/concepts>

The following is a brief summary of the collation attributes and how they are specified using the ICU short form method

Each collation attribute is represented by an uppercase letter, which are listed in the following bullet points. The possible valid values for each attribute are given by codes shown within the parentheses. Some codes have general meanings for all attributes. **X** means to set the attribute off. **O** means to set the attribute on. **D** means to set the attribute to its default value.

- **A – Alternate (N, S, D)**. Handles treatment of *variable* characters such as white spaces, punctuation marks, and symbols. When set to non-ignorable (N), differences in variable characters are treated with the same importance as differences in letters. When set to shifted (S), then differences in variable characters are of minor importance (that is, the variable character is ignored when comparing base characters).
- **C – Case First (X, L, U, D)**. Controls whether a lowercase letter sorts before the same uppercase letter (L), or the uppercase letter sorts before the same lowercase letter (U). Off (X) is typically specified when lowercase first (L) is desired.
- **E – Case Level (X, O, D)**. Set in combination with the Strength attribute, the Case Level attribute is used when accents are to be ignored, but not case.
- **F – French Collation (X, O, D)**. When set to on, secondary differences (presence of accents) are sorted from the back of the string as done in the French Canadian locale.
- **H – Hiragana Quaternary (X, O, D)**. Introduces an additional level to distinguish between the Hiragana and Katakana characters for compatibility with the JIS X 4061 collation of Japanese character strings.
- **N – Normalization Checking (X, O, D)**. Controls whether or not text is thoroughly normalized for comparison.

Normalization deals with the issue of canonical equivalence of text whereby different code point sequences represent the same character, which then present issues when sorting or comparing such characters. Languages such as Arabic, ancient Greek, Hebrew, Hindi, Thai, or Vietnamese should be used with Normalization Checking set to on.

- **S – Strength (1, 2, 3, 4, I, D).** Maximum collation level used for comparison. Influences whether accents or case are taken into account when collating or comparing strings. Each number represents a level. A setting of I represents identical strength (that is, level 5).
- **T – Variable Top (hexadecimal digits).** Applicable only when the Alternate attribute is not set to non-ignorable (N). The hexadecimal digits specify the highest character sequence that is to be considered ignorable. For example, if white space is to be ignorable, but visible variable characters are not to be ignorable, then Variable Top set to 0020 would be specified along with the Alternate attribute set to S and the Strength attribute set to 3. (The space character is hexadecimal 0020. Other non-visible variable characters such as backspace, tab, line feed, carriage return, etc. have values less than 0020. All visible punctuation marks have values greater than 0020.)

A set of collation attributes and their values is represented by a text string consisting of the collation attribute letter concatenated with the desired attribute value. Each attribute/value pair is joined to the next pair with an underscore character as shown by the following example.

AN_CX_EX_FX_HX_NO_S3

Collation attributes can be specified along with a locale's ICU short form name to override those default attribute settings of the locale.

The following is an example where the ICU short form named **LROOT** is modified with a number of other collation attribute/value pairs.

AN_CX_EX_LROOT_NO_S3

In the preceding example, the Alternate attribute **(A)** is set to non-ignorable **(N)**. The Case First attribute **(C)** is set to off **(X)**. The Case Level attribute **(E)** is set to off **(X)**. The Normalization attribute **(N)** is set to on **(O)**. The Strength attribute **(S)** is set to the tertiary level **3**. **LROOT** is the ICU short form to which these other attributes are applying modifications.

Using a Collation

A newly defined ICU collation can be used anywhere the **COLLATION "collation_name"** clause can be used in a SQL command such as in the column specifications of the **CREATE TABLE** command or appended to an expression in the **ORDER BY** clause of a **SELECT** command.

The following are some examples of the creation and usage of ICU collations based on the English language in the United States (**en_US.UTF8**).

In these examples, ICU collations are created with the following characteristics.

Collation **icu_collate_lowercase** forces the lowercase form of a letter to sort before its uppercase counterpart (**CL**).

Collation **icu_collate_uppercase** forces the uppercase form of a letter to sort before its lowercase counterpart (**CU**).

Collation **icu_collate_ignore_punct** causes variable characters (white space and punctuation marks) to be ignored during sorting (**AS**).

Collation **icu_collate_ignore_white_sp** causes white space and other non-visible variable characters to be ignored during sorting, but visible variable characters (punctuation marks) are not ignored (**AS**, **T0020**).

```
CREATE COLLATION icu_collate_lowercase (
  LOCALE = 'en_US.UTF8',
  ICU_SHORT_FORM = 'AN_CL_EX_NX_LROOT'
```

```
);

CREATE COLLATION icu_collate_uppercase (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AN_CU_EX_NX_LROOT'
);
```

```
CREATE COLLATION icu_collate_ignore_punct (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AS_CX_EX_NX_LROOT_L3'
);
```

```
CREATE COLLATION icu_collate_ignore_white_sp (
    LOCALE = 'en_US.UTF8',
    ICU_SHORT_FORM = 'AS_CX_EX_NX_LROOT_L3_T0020'
);
```

Note

When creating collations, ICU may generate notice and warning messages when attributes are given to modify the **LROOT** collation.

The following `psql` command lists the collations.

```
edb=# \dO
          List of collations
 Schema |      Name       | Collate | Ctype |      ICU
-----+-----+-----+-----+
-----+
enterprisedb | icu_collate_ignore_punct | en_US.UTF8 | en_US.UTF8 |
AS_CX_EX_NX_LROOT_L3
enterprisedb | icu_collate_ignore_white_sp | en_US.UTF8 | en_US.UTF8 |
AS_CX_EX_NX_LROOT_L3_T0020
enterprisedb | icu_collate_lowercase     | en_US.UTF8 | en_US.UTF8 |
AN_CL_EX_NX_LROOT
enterprisedb | icu_collate_uppercase     | en_US.UTF8 | en_US.UTF8 |
AN_CU_EX_NX_LROOT
(4 rows)
```

The following table is created and populated.

```
CREATE TABLE collate_tbl (
    id      INTEGER,
    c2      VARCHAR(2)
);

INSERT INTO collate_tbl VALUES (1, 'A');
INSERT INTO collate_tbl VALUES (2, 'B');
INSERT INTO collate_tbl VALUES (3, 'C');
INSERT INTO collate_tbl VALUES (4, 'a');
INSERT INTO collate_tbl VALUES (5, 'b');
INSERT INTO collate_tbl VALUES (6, 'c');
INSERT INTO collate_tbl VALUES (7, '1');
INSERT INTO collate_tbl VALUES (8, '2');
INSERT INTO collate_tbl VALUES (9, '.B');
INSERT INTO collate_tbl VALUES (10, '-B');
```

```
INSERT INTO collate_tbl VALUES (11, 'B');
```

The following query sorts on column `c2` using the default collation. Note that variable characters (white space and punctuation marks) with id column values of `9`, `10`, and `11` are ignored and sort with the letter `B`.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2;
id | c2
----+---
 7 | 1
 8 | 2
 4 | a
 1 | A
 5 | b
 2 | B
11 | B
10 | -B
 9 | .B
 6 | c
 3 | C
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_lowercase`, which forces the lowercase form of a letter to sort before the uppercase form of the same base letter. Also note that the `AN` attribute forces variable characters to be included in the sort order at the same level when comparing base characters so rows with `id` values of `9`, `10`, and `11` appear at the beginning of the sort list before all letters and numbers.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE "icu_collate_lowercase";
id | c2
----+---
11 | B
10 | -B
 9 | .B
 7 | 1
 8 | 2
 4 | a
 1 | A
 5 | b
 2 | B
 6 | c
 3 | C
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_uppercase`, which forces the uppercase form of a letter to sort before the lowercase form of the same base letter.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE "icu_collate_uppercase";
id | c2
----+---
11 | B
10 | -B
 9 | .B
 7 | 1
 8 | 2
 1 | A
 4 | a
 2 | B
```

```
5 | b
3 | C
6 | c
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_ignore_punct`, which causes variable characters to be ignored so rows with id values of `9`, `10`, and `11` sort with the letter `B` as that is the character immediately following the ignored variable character.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE
"icu_collate_ignore_punct";
id | c2
----+---
7 | 1
8 | 2
4 | a
1 | A
5 | b
11 | B
2 | B
9 | .B
10 | -B
6 | c
3 | C
(11 rows)
```

The following query sorts on column `c2` using collation `icu_collate_ignore_white_sp`. The `AS` and `T0020` attributes of the collation cause variable characters with code points less than or equal to hexadecimal `0020` to be ignored while variable characters with code points greater than hexadecimal `0020` are included in the sort.

The row with id value of `11`, which starts with a space character (hexadecimal `0020`) sorts with the letter `B`. The rows with id values of `9` and `10`, which start with visible punctuation marks greater than hexadecimal `0020`, appear at the beginning of the sort list as these particular variable characters are included in the sort order at the same level when comparing base characters.

```
edb=# SELECT * FROM collate_tbl ORDER BY c2 COLLATE
"icu_collate_ignore_white_sp";
id | c2
----+---
10 | -B
9 | .B
7 | 1
8 | 2
4 | a
1 | A
5 | b
11 | B
2 | B
6 | c
3 | C
(11 rows)
```

11.10 EDB Resource Manager

EDB Resource Manager is an Advanced Server feature that provides the capability to control the usage of operating system resources used by Advanced Server processes.

This capability allows you to protect the system from processes that may uncontrollably overuse and monopolize certain system resources.

The following are some key points about using EDB Resource Manager.

- The basic component of EDB Resource Manager is a resource group. A *resource group* is a named, global group, available to all databases in an Advanced Server instance, on which various resource usage limits can be defined. Advanced Server processes that are assigned as members of a given resource group are then controlled by EDB Resource Manager so that the aggregate resource usage of all processes in the group is kept near the limits defined on the group.
- Data definition language commands are used to create, alter, and drop resource groups. These commands can only be used by a database user with superuser privileges.
- The desired, aggregate consumption level of all processes belonging to a resource group is defined by *resource type parameters*. There are different resource type parameters for the different types of system resources currently supported by EDB Resource Manager.
- Multiple resource groups can be created, each with different settings for its resource type parameters, thus defining different consumption levels for each resource group.
- EDB Resource Manager throttles processes in a resource group to keep resource consumption near the limits defined by the resource type parameters. If there are multiple resource type parameters with defined settings in a resource group, the actual resource consumption may be significantly lower for certain resource types than their defined resource type parameter settings. This is because EDB Resource Manager throttles processes attempting to keep *all resources with defined resource type settings within their defined limits*.
- The definition of available resource groups and their resource type settings are stored in a shared global system catalog. Thus, resource groups can be utilized by all databases in a given Advanced Server instance.
- The `edb_max_resource_groups` configuration parameter sets the maximum number of resource groups that can be active simultaneously with running processes. The default setting is 16 resource groups. Changes to this parameter take effect on database server restart.
- Use the `SET edb_resource_group TO group_name` command to assign the current process to a specified resource group. Use the `RESET edb_resource_group` command or `SET edb_resource_group TO DEFAULT` to remove the current process from a resource group.
- A default resource group can be assigned to a role using the `ALTER ROLE ... SET` command, or to a database by the `ALTER DATABASE ... SET` command. The entire database server instance can be assigned a default resource group by setting the parameter in the `postgresql.conf` file.
- In order to include resource groups in a backup file of the database server instance, use the `pg_dumpall` backup utility with default settings (That is, do not specify any of the `--globals-only`, `--roles-only`, or `-- tablespaces-only` options.)

Creating and Managing Resource Groups

The data definition language commands described in this section provide for the creation and management of resource groups.

CREATE RESOURCE GROUP

Use the `CREATE RESOURCE GROUP` command to create a new resource group.

```
| CREATE RESOURCE GROUP <group_name>;
```

Description

The `CREATE RESOURCE GROUP` command creates a resource group with the specified name. Resource limits

can then be defined on the group with the **ALTER RESOURCE GROUP** command. The resource group is accessible from all databases in the Advanced Server instance.

To use the **CREATE RESOURCE GROUP** command you must have superuser privileges.

Parameters

group_name

The name of the resource group.

Example

The following example results in the creation of three resource groups named `resgrp_a`, `resgrp_b`, and `resgrp_c`.

```
edb=# CREATE RESOURCE GROUP resgrp_a;
CREATE RESOURCE GROUP
edb=# CREATE RESOURCE GROUP resgrp_b;
CREATE RESOURCE GROUP
edb=# CREATE RESOURCE GROUP resgrp_c;
CREATE RESOURCE GROUP
```

The following query shows the entries for the resource groups in the `edb_resource_group` catalog.

```
edb=# SELECT * FROM edb_resource_group;
rgrpname | rgrpcpuratelim | rgrpdirtyratelim
-----+-----+
resgrp_a |      0 | 0
resgrp_b |      0 | 0
resgrp_c |      0 | 0
(3 rows)
```

ALTER RESOURCE GROUP

Use the **ALTER RESOURCE GROUP** command to change the attributes of an existing resource group. The command syntax comes in three forms.

The first form renames the resource group:

```
ALTER RESOURCE GROUP <group_name> RENAME TO <new_name>;
```

The second form assigns a resource type to the resource group:

```
ALTER RESOURCE GROUP <group_name> SET
<resource_type> { TO | = } { <value> | DEFAULT };
```

The third form resets the assignment of a resource type to its default within the group:

```
ALTER RESOURCE GROUP <group_name> RESET <resource_type>;
```

Description

The **ALTER RESOURCE GROUP** command changes certain attributes of an existing resource group.

The first form with the **RENAME TO** clause assigns a new name to an existing resource group.

The second form with the **SET resource_type TO** clause either assigns the specified literal value to a resource

type, or resets the resource type when `DEFAULT` is specified. Resetting or setting a resource type to `DEFAULT` means that the resource group has no defined limit on that resource type.

The third form with the `RESET resource_type` clause resets the resource type for the group as described previously.

To use the `ALTER RESOURCE GROUP` command, you must have superuser privileges.

Parameters

`group_name`

The name of the resource group to be altered.

`new_name`

The new name to be assigned to the resource group.

`resource_type`

The resource type parameter specifying the type of resource to which a usage value is to be set.

`value | DEFAULT`

When `value` is specified, the literal value to be assigned to `resource_type`. When `DEFAULT` is specified, the assignment of `resource_type` is reset for the resource group.

Example

The following are examples of the `ALTER RESOURCE GROUP` command.

```
edb=# ALTER RESOURCE GROUP resgrp_a RENAME TO newgrp;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_b SET cpu_rate_limit = .5;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_b SET dirty_rate_limit = 6144;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_c RESET cpu_rate_limit;
ALTER RESOURCE GROUP
```

The following query shows the results of the `ALTER RESOURCE GROUP` commands to the entries in the `edb_resource_group` catalog.

```
edb=# SELECT * FROM edb_resource_group;
rgrpname | rgrpcpuratelim | rgrpdirtyratelim
-----+-----+
newgrp  | 0          | 0
resgrp_b | 0.5       | 6144
resgrp_c | 0          | 0
(3 rows)
```

DROP RESOURCE GROUP

Use the `DROP RESOURCE GROUP` command to remove a resource group.

`DROP RESOURCE GROUP [IF EXISTS] <group_name>;`

Description

The `DROP RESOURCE GROUP` command removes a resource group with the specified name.

To use the `DROP RESOURCE GROUP` command you must have superuser privileges.

Parameters

`group_name`

The name of the resource group to be removed.

`IF EXISTS`

Do not throw an error if the resource group does not exist. A notice is issued in this case.

Example

The following example removes resource group `newgrp`.

```
edb=# DROP RESOURCE GROUP newgrp
DROP RESOURCE GROUP
```

Assigning a Process to a Resource Group

Use the `SET edb_resource_group TO group_name` command to assign the current process to a specified resource group as shown by the following.

```
edb=# SET edb_resource_group TO resgrp_b;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
resgrp_b
(1 row)
```

The resource type settings of the group immediately take effect on the current process. If the command is used to change the resource group assigned to the current process, the resource type settings of the newly assigned group immediately take effect.

Processes can be included by default in a resource group by assigning a default resource group to roles, databases, or an entire database server instance.

A default resource group can be assigned to a role using the `ALTER ROLE ... SET` command. For more information about the `ALTER ROLE` command, refer to the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-alterrole.html>

A default resource group can be assigned to a database by the `ALTER DATABASE ... SET` command. For more information about the `ALTER DATABASE` command, refer to the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/static/sql-alterdatabase.html>

The entire database server instance can be assigned a default resource group by setting the `edb_resource_group` configuration parameter in the `postgresql.conf` file as shown by the following.

```
### - EDB Resource Manager -
#edb_max_resource_groups = 16          # 0-65536 (change requires restart)
edb_resource_group = 'resgrp_b'
```

A change to `edb_resource_group` in the `postgresql.conf` file requires a configuration file reload before it takes effect on the database server instance.

Removing a Process from a Resource Group

Set `edb_resource_group` to `DEFAULT` or use `RESET edb_resource_group` to remove the current process from a resource group as shown by the following.

```
edb=# SET edb_resource_group TO DEFAULT;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
(1 row)
```

For removing a default resource group from a role, use the `ALTER ROLE ... RESET` form of the `ALTER ROLE` command.

For removing a default resource group from a database, use the `ALTER DATABASE ... RESET` form of the `ALTER DATABASE` command.

For removing a default resource group from the database server instance, set the `edb_resource_group` configuration parameter to an empty string in the `postgresql.conf` file and reload the configuration file.

Monitoring Processes in Resource Groups

After resource groups have been created, the number of processes actively using these resource groups can be obtained from the view `edb_all_resource_groups`.

The columns in `edb_all_resource_groups` are the following:

- **group_name**. Name of the resource group.
- **active_processes**. Number of active processes in the resource group.
- **cpu_rate_limit**. The value of the CPU rate limit resource type assigned to the resource group.
- **per_process_cpu_rate_limit**. The CPU rate limit applicable to an individual, active process in the resource group.
- **dirty_rate_limit**. The value of the dirty rate limit resource type assigned to the resource group.
- **per_process_dirty_rate_limit**. The dirty rate limit applicable to an individual, active process in the resource group.

Note

Columns `per_process_cpu_rate_limit` and `per_process_dirty_rate_limit` do not show the *actual* resource consumption used by the processes, but indicate how `EDB Resource Manager` sets the resource limit for an individual process based upon the number of active processes in the resource group.

The following shows `edb_all_resource_groups` when resource group `resgrp_a` contains no active processes, resource group `resgrp_b` contains two active processes, and resource group `resgrp_c` contains one active process.

```
edb=# SELECT * FROM edb_all_resource_groups ORDER BY group_name;
-[ RECORD 1 ]-----+
group_name      | resgrp_a
active_processes | 0
cpu_rate_limit   | 0.5
per_process_cpu_rate_limit |
```

```

dirty_rate_limit      | 12288
per_process_dirty_rate_limit |
-[ RECORD 2 ]-----+
group_name           | resgrp_b
active_processes     | 2
cpu_rate_limit       | 0.4
per_process_cpu_rate_limit | 0.195694289022895
dirty_rate_limit     | 6144
per_process_dirty_rate_limit | 3785.92924684337
-[ RECORD 3 ]-----+
group_name           | resgrp_c
active_processes     | 1
cpu_rate_limit       | 0.3
per_process_cpu_rate_limit | 0.292342129631091
dirty_rate_limit     | 3072
per_process_dirty_rate_limit | 3072

```

The CPU rate limit and dirty rate limit settings that are assigned to these resource groups are as follows.

```

edb=# SELECT * FROM edb_resource_group;
rgrpname | rgrpcpuratelimit | rgrpdirtyratelimit
-----+-----+-----+
resgrp_a | 0.5      | 12288
resgrp_b | 0.4      | 6144
resgrp_c | 0.3      | 3072
(3 rows)

```

In the `edb_all_resource_groups` view, note that the `per_process_cpu_rate_limit` and `per_process_dirty_rate_limit` values are roughly the corresponding CPU rate limit and dirty rate limit divided by the number of active processes.

CPU Usage Throttling

CPU usage of a resource group is controlled by setting the `cpu_rate_limit` resource type parameter.

Set the `cpu_rate_limit` parameter to the fraction of CPU time over wall-clock time to which the combined, simultaneous CPU usage of all processes in the group should not exceed. Thus, the value assigned to `cpu_rate_limit` should typically be less than or equal to 1.

The valid range of the `cpu_rate_limit` parameter is 0 to `1.67772e+07`. A setting of 0 means no CPU rate limit has been set for the resource group.

When multiplied by 100, the `cpu_rate_limit` can also be interpreted as the CPU usage percentage for a resource group.

EDB Resource Manager utilizes *CPU throttling* to keep the aggregate CPU usage of all processes in the group within the limit specified by the `cpu_rate_limit` parameter. A process in the group may be interrupted and put into sleep mode for a short interval of time to maintain the defined limit. When and how such interruptions occur is defined by a proprietary algorithm used by EDB Resource Manager.

Setting the CPU Rate Limit for a Resource Group

The `ALTER RESOURCE GROUP` command with the `SET cpu_rate_limit` clause is used to set the CPU rate limit for a resource group.

In the following example the CPU usage limit is set to 50% for `resgrp_a`, 40% for `resgrp_b` and 30% for

`resgrp_c`. This means that the combined CPU usage of all processes assigned to `resgrp_a` is maintained at approximately 50%. Similarly, for all processes in `resgrp_b`, the combined CPU usage is kept to approximately 40%, etc.

```
edb=# ALTER RESOURCE GROUP resgrp_a SET cpu_rate_limit TO .5;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_b SET cpu_rate_limit TO .4;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_c SET cpu_rate_limit TO .3;
ALTER RESOURCE GROUP
```

The following query shows the settings of `cpu_rate_limit` in the catalog.

```
edb=# SELECT rgrpname, rgrpcpuratelim FROM edb_resource_group;
rgrpname | rgrpcpuratelim
-----+
resgrp_a |      0.5
resgrp_b |      0.4
resgrp_c |      0.3
(3 rows)
```

Changing the `cpu_rate_limit` of a resource group not only affects new processes that are assigned to the group, but any currently running processes that are members of the group are immediately affected by the change. That is, if the `cpu_rate_limit` is changed from .5 to .3, currently running processes in the group would be throttled downward so that the aggregate group CPU usage would be near 30% instead of 50%.

To illustrate the effect of setting the CPU rate limit for resource groups, the following examples use a CPU-intensive calculation of 20000 factorial (multiplication of 20000 * 19999 * 19998, etc.) performed by the query `SELECT 20000!`; run in the `psql` command line utility.

The resource groups with the CPU rate limit settings shown in the previous query are used in these examples.

Example – Single Process in a Single Group

The following shows that the current process is set to use resource group `resgrp_b`. The factorial calculation is then started.

```
edb=# SET edb_resource_group TO resgrp_b;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
resgrp_b
(1 row)
edb=# SELECT 20000!;
```

In a second session, the Linux `top` command is used to display the CPU usage as shown under the `%CPU` column. The following is a snapshot at an arbitrary point in time as the `top` command output periodically changes.

```
$ top
top - 16:37:03 up 4:15, 7 users, load average: 0.49, 0.20, 0.38
Tasks: 202 total, 1 running, 201 sleeping, 0 stopped, 0 zombie
Cpu(s): 42.7%us, 2.3%sy, 0.0%ni, 55.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0
Mem: 1025624k total, 791160k used, 234464k free, 23400k buffers
Swap: 103420k total, 13404k used, 90016k free, 373504k cached
```

```
PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
28915 enterpri 20 0 195m 5900 4212 S 39.9 0.6 3:36.98 edb-postgres
1033 root 20 0 171m 77m 2960 S 1.0 7.8 3:43.96 Xorg
3040 user 20 0 278m 22m 14m S 1.0 2.2 3:41.72 knotify4
```

The `psql` session performing the factorial calculation is shown by the row where `edb-postgres` appears under the `COMMAND` column. The CPU usage of the session shown under the `%CPU` column shows 39.9, which is close to the 40% CPU limit set for resource group `resgrp_b`.

By contrast, if the `psql` session is removed from the resource group and the factorial calculation is performed again, the CPU usage is much higher.

```
edb=# SET edb_resource_group TO DEFAULT;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
```

(1 row)

```
edb=# SELECT 20000!;
```

Under the `%CPU` column for `edb-postgres`, the CPU usage is now 93.6, which is significantly higher than the 39.9 when the process was part of the resource group.

```
$ top
top - 16:43:03 up 4:21, 7 users, load average: 0.66, 0.33, 0.37
Tasks: 202 total, 5 running, 197 sleeping, 0 stopped, 0 zombie
Cpu(s): 96.7%us, 3.3%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0
Mem: 1025624k total, 791228k used, 234396k free, 23560k buffers
Swap: 103420k total, 13404k used, 90016k free, 373508k cached
```

```
PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
28915 enterpri 20 0 195m 5900 4212 R 93.6 0.6 5:01.56 edb-postgres
1033 root 20 0 171m 77m 2960 S 1.0 7.8 3:48.15 Xorg
2907 user 20 0 98.7m 11m 9100 S 0.3 1.2 0:46.51 vmware-user-lo
```

Example – Multiple Processes in a Single Group

As stated previously, the CPU rate limit applies to the aggregate of all processes in the resource group. This concept is illustrated in the following example.

The factorial calculation is performed simultaneously in two separate `psql` sessions, each of which has been added to resource group `resgrp_b` that has `cpu_rate_limit` set to .4 (CPU usage of 40%).

Session 1:

```
edb=# SET edb_resource_group TO resgrp_b;
SET
```

```
edb=# SHOW edb_resource_group;
```

```
edb_resource_group
```

```
-----
```

```
resgrp_b
```

```
(1 row)
```

```
edb=# SELECT 20000!;
```

Session 2:

```
edb=# SET edb_resource_group TO resgrp_b;
```

```
SET
```

```
edb=# SHOW edb_resource_group;
```

```
edb_resource_group
```

```
-----
```

```
resgrp_b
```

```
(1 row)
```

```
edb=# SELECT 20000!;
```

A third session monitors the CPU usage.

```
$ top
```

```
top - 16:53:03 up 4:31, 7 users, load average: 0.31, 0.19, 0.27
```

```
Tasks: 202 total, 1 running, 201 sleeping, 0 stopped, 0 zombie
```

```
Cpu(s): 41.2%us, 3.0%sy, 0.0%ni, 55.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0
```

```
Mem: 1025624k total, 792020k used, 233604k free, 23844k buffers
```

```
Swap: 103420k total, 13404k used, 90016k free, 373508k cached
```

```
PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
```

```
29857 enterpri 20 0 195m 4708 3312 S 19.9 0.5 0:57.35 edb-postgres
```

```
28915 enterpri 20 0 195m 5900 4212 S 19.6 0.6 5:35.49 edb-postgres
```

```
3040 user 20 0 278m 22m 14m S 1.0 2.2 3:54.99 knotify4
```

```
1033 root 20 0 171m 78m 2960 S 0.3 7.8 3:55.71 Xorg
```

```
.
```

```
.
```

```
.
```

There are now two processes named `edb-postgres` with `%CPU` values of 19.9 and 19.6, whose sum is close to the 40% CPU usage set for resource group `resgrp_b`.

The following command sequence displays the sum of all `edb-postgres` processes sampled over half second time intervals. This shows how the total CPU usage of the processes in the resource group changes over time as EDB Resource Manager throttles the processes to keep the total resource group CPU usage near 40%.

```
$ while [[ 1 -eq 1 ]]; do top -d0.5 -b -n2 | grep edb-postgres | awk '{ SUM
```

```
+=$9} END { print SUM / 2 }'; done
```

```
37.2
```

```
39.1
```

```
38.9
```

```
38.3
```

```
44.7
```

```
39.2
```

```
42.5
```

```
39.1
```

```
39.2
```

```
39.2
41
42.85
46.1
.
```

Example – Multiple Processes in Multiple Groups

In this example, two additional `psql` sessions are used along with the previous two sessions. The third and fourth sessions perform the same factorial calculation within resource group `resgrp_c` with a `cpu_rate_limit` of `.3` (30% CPU usage).

Session 3:

```
edb=# SET edb_resource_group TO resgrp_c;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
resgrp_c
(1 row)

edb=# SELECT 20000!;
```

Session 4:

```
edb=# SET edb_resource_group TO resgrp_c;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
resgrp_c
(1 row)
edb=# SELECT 20000!;
```

The `top` command displays the following output.

```
$ top
top - 17:45:09 up 5:23, 8 users, load average: 0.47, 0.17, 0.26
Tasks: 203 total, 4 running, 199 sleeping, 0 stopped, 0 zombie
Cpu(s): 70.2%us, 0.0%sy, 0.0%ni, 29.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0
Mem: 1025624k total, 806140k used, 219484k free, 25296k buffers
Swap: 103420k total, 13404k used, 90016k free, 374092k cached

 PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
29857 enterpri 20 0 195m 4820 3324 S 19.9 0.5 4:25.02 edb-postgres
28915 enterpri 20 0 195m 5900 4212 R 19.6 0.6 9:07.50 edb-postgres
29023 enterpri 20 0 195m 4744 3248 R 16.3 0.5 4:01.73 edb-postgres
11019 enterpri 20 0 195m 4120 2764 R 15.3 0.4 0:04.92 edb-postgres
2907 user 20 0 98.7m 12m 9112 S 1.3 1.2 0:56.54 vmware-user-lo
3040 user 20 0 278m 22m 14m S 1.3 2.2 4:38.73 knotify4
```

The two resource groups in use have CPU usage limits of 40% and 30%. The sum of the %CPU column for the first two `edb-postgres` processes is 39.5 (approximately 40%, which is the limit for `resgrp_b`) and the sum of the %CPU column for the third and fourth `edb-postgres` processes is 31.6 (approximately 30%, which is the limit for `resgrp_c`).

The sum of the CPU usage limits of the two resource groups to which these processes belong is 70%. The following output shows that the sum of the four processes borders around 70%.

```
$ while [[ 1 -eq 1 ]]; do top -d0.5 -b -n2 | grep edb-postgres | awk '{ SUM  
+= $9} END { print SUM / 2 }'; done  
61.8  
76.4  
72.6  
69.55  
64.55  
79.95  
68.55  
71.25  
74.85  
62  
74.85  
76.9  
72.4  
65.9  
74.9  
68.25
```

By contrast, if three sessions are processing where two sessions remain in `resgrp_b`, but the third session does not belong to any resource group, the `top` command shows the following output.

```
$ top  
top - 17:24:55 up 5:03, 7 users, load average: 1.00, 0.41, 0.38  
Tasks: 199 total, 3 running, 196 sleeping, 0 stopped, 0 zombie  
Cpu(s): 99.7%us, 0.3%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0  
Mem: 1025624k total, 797692k used, 227932k free, 24724k buffers  
Swap: 103420k total, 13404k used, 90016k free, 374068k cached  
  
 PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND  
29023 enterpri 20 0 195m 4744 3248 R 58.6 0.5 2:53.75 edb-postgres  
28915 enterpri 20 0 195m 5900 4212 S 18.9 0.6 7:58.45 edb-postgres  
29857 enterpri 20 0 195m 4820 3324 S 18.9 0.5 3:14.85 edb-postgres  
1033 root 20 0 174m 81m 2960 S 1.7 8.2 4:26.50 Xorg  
3040 user 20 0 278m 22m 14m S 1.0 2.2 4:21.20 knotify4
```

The second and third `edb-postgres` processes belonging to the resource group where the CPU usage is limited to 40%, have a total CPU usage of 37.8. However, the first `edb-postgres` process has a 58.6% CPU usage as it is not within a resource group, and basically utilizes the remaining, available CPU resources on the system.

Likewise, the following output shows the sum of all three sessions is around 95% since one of the sessions has no set limit on its CPU usage.

```
$ while [[ 1 -eq 1 ]]; do top -d0.5 -b -n2 | grep edb-postgres | awk '{ SUM  
+= $9} END { print SUM / 2 }'; done  
96  
90.35  
92.55  
96.4
```

```

94.1
90.7
95.7
95.45
93.65
87.95
96.75
94.25
95.45
97.35
92.9
96.05
96.25
94.95
.
.
```

Dirty Buffer Throttling

Writing to shared buffers is controlled by setting the `dirty_rate_limit` resource type parameter.

Set the `dirty_rate_limit` parameter to the number of kilobytes per second for the combined rate at which all the processes in the group should write to or “dirty” the shared buffers. An example setting would be 3072 kilobytes per seconds.

The valid range of the `dirty_rate_limit` parameter is `0` to `1.67772e+07`. A setting of 0 means no dirty rate limit has been set for the resource group.

EDB Resource Manager utilizes *dirty buffer throttling* to keep the aggregate, shared buffer writing rate of all processes in the group near the limit specified by the `dirty_rate_limit` parameter. A process in the group may be interrupted and put into sleep mode for a short interval of time to maintain the defined limit. When and how such interruptions occur is defined by a proprietary algorithm used by EDB Resource Manager.

Setting the Dirty Rate Limit for a Resource Group

The `ALTER RESOURCE GROUP` command with the `SET dirty_rate_limit` clause is used to set the dirty rate limit for a resource group.

In the following example the dirty rate limit is set to 12288 kilobytes per second for `resgrp_a`, 6144 kilobytes per second for `resgrp_b` and 3072 kilobytes per second for `resgrp_c`. This means that the combined writing rate to the shared buffer of all processes assigned to `resgrp_a` is maintained at approximately 12288 kilobytes per second. Similarly, for all processes in `resgrp_b`, the combined writing rate to the shared buffer is kept to approximately 6144 kilobytes per second, etc.

```

edb=# ALTER RESOURCE GROUP resgrp_a SET dirty_rate_limit TO 12288;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_b SET dirty_rate_limit TO 6144;
ALTER RESOURCE GROUP
edb=# ALTER RESOURCE GROUP resgrp_c SET dirty_rate_limit TO 3072;
ALTER RESOURCE GROUP

```

The following query shows the settings of `dirty_rate_limit` in the catalog.

```
edb=# SELECT rgrpname, rgrpdirtyratelimit FROM edb_resource_group;
rgrpname | rgrpdirtyratelimit
-----+-----
resgrp_a |    12288
resgrp_b |     6144
resgrp_c |     3072
(3 rows)
```

Changing the `dirty_rate_limit` of a resource group not only affects new processes that are assigned to the group, but any currently running processes that are members of the group are immediately affected by the change. That is, if the `dirty_rate_limit` is changed from 12288 to 3072, currently running processes in the group would be throttled downward so that the aggregate group dirty rate would be near 3072 kilobytes per second instead of 12288 kilobytes per second.

To illustrate the effect of setting the dirty rate limit for resource groups, the following examples use the following table for intensive I/O operations.

```
CREATE TABLE t1 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLFACTOR = 10);
```

The `FILLFACTOR = 10` clause results in `INSERT` commands packing rows up to only 10% per page. This results in a larger sampling of dirty shared blocks for the purpose of these examples.

The `pg_stat_statements` module is used to display the number of shared buffer blocks that are dirtied by a SQL command and the amount of time the command took to execute. This provides the information to calculate the actual kilobytes per second writing rate for the SQL command, and thus compare it to the dirty rate limit set for a resource group.

In order to use the `pg_stat_statements` module, perform the following steps.

Step 1: In the `postgresql.conf` file, add `$libdir/pg_stat_statements` to the `shared_preload_libraries` configuration parameter as shown by the following.

```
shared_preload_libraries = '$libdir/dbms_pipe,$libdir/edb_gen,$libdir/
pg_stat_statements'
```

Step 2: Restart the database server.

Step 3: Use the `CREATE EXTENSION` command to complete the creation of the `pg_stat_statements` module.

```
edb=# CREATE EXTENSION pg_stat_statements SCHEMA public;
CREATE EXTENSION
```

The `pg_stat_statements_reset()` function is used to clear out the `pg_stat_statements` view for clarity of each example.

The resource groups with the dirty rate limit settings shown in the previous query are used in these examples.

Example – Single Process in a Single Group

The following sequence of commands shows the creation of table `t1`. The current process is set to use resource group `resgrp_b`. The `pg_stat_statements` view is cleared out by running the `pg_stat_statements_reset()` function.

Finally, the `INSERT` command generates a series of integers from 1 to 10,000 to populate the table, and dirty approximately 10,000 blocks.

```
edb=# CREATE TABLE t1 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLFACTOR = 10);
CREATE TABLE
```

```

edb=# SETedb_resource_group TO resgrp_b;
SET
edb=# SHOWedb_resource_group;
edb_resource_group
-----
resgrp_b
(1 row)

edb=# SELECT pg_stat_statements_reset();
pg_stat_statements_reset
-----
(1 row)

edb=# INSERT INTO t1 VALUES (generate_series (1,10000), 'aaa');
INSERT 0 10000

```

The following shows the results from the `INSERT` command.

```

edb=# SELECT query, rows, total_time, shared_blk_dirtied FROM
pg_stat_statements;
-[ RECORD 1 ]-----+
query      | INSERT INTO t1 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 13496.184
shared_blk_dirtied | 10003

```

The actual dirty rate is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 13496.184 ms, which yields *0.74117247 blocks per millisecond*.
- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *741.17247 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *6072 kilobytes per second*.

Note that the actual dirty rate of 6072 kilobytes per second is close to the dirty rate limit for the resource group, which is 6144 kilobytes per second.

By contrast, if the steps are repeated again without the process belonging to any resource group, the dirty buffer rate is much higher.

```

edb=# CREATE TABLE t1 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLFACTOR = 10);
CREATE TABLE
edb=# SHOWedb_resource_group;
edb_resource_group
-----
(1 row)

edb=# SELECT pg_stat_statements_reset();
pg_stat_statements_reset
-----
(1 row)

edb=# INSERT INTO t1 VALUES (generate_series (1,10000), 'aaa');

```

```
INSERT 0 10000
```

The following shows the results from the `INSERT` command without the usage of a resource group.

```
edb=# SELECT query, rows, total_time, shared_blks_dirtied FROM pg_stat_statements;
-[ RECORD 1 ]-----+
query      | INSERT INTO t1 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 2432.165
shared_blks_dirtied | 10003
```

First, note the total time was only 2432.165 milliseconds as compared to 13496.184 milliseconds when a resource group with a dirty rate limit set to 6144 kilobytes per second was used.

The actual dirty rate without the use of a resource group is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 2432.165 ms, which yields *4.112797 blocks per millisecond*.
- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *4112.797 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *33692 kilobytes per second*.

Note that the actual dirty rate of 33692 kilobytes per second is significantly higher than when the resource group with a dirty rate limit of 6144 kilobytes per second was used.

Example – Multiple Processes in a Single Group

As stated previously, the dirty rate limit applies to the aggregate of all processes in the resource group. This concept is illustrated in the following example.

For this example the inserts are performed simultaneously on two different tables in two separate `psql` sessions, each of which has been added to resource group `resgrp_b` that has a `dirty_rate_limit` set to 6144 kilobytes per second.

Session 1:

```
edb=# CREATE TABLE t1 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLCODE = 10);
CREATE TABLE
edb=# SET edb_resource_group TO resgrp_b;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
resgrp_b
(1 row)
```

```
edb=# INSERT INTO t1 VALUES (generate_series (1,10000), 'aaa');
INSERT 0 10000
```

Session 2:

```
edb=# CREATE TABLE t2 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLCODE = 10);
CREATE TABLE
edb=# SET edb_resource_group TO resgrp_b;
```

```
SET
edb=# SHOW edb_resource_group;
edb_resource_group
```

```
-----+
resgrp_b
(1 row)
```

```
edb=# SELECT pg_stat_statements_reset();
pg_stat_statements_reset
```

```
-----+
(1 row)
```

```
edb=# INSERT INTO t2 VALUES (generate_series (1,10000), 'aaa');
INSERT 0 10000
```

Note

The `INSERT` commands in session 1 and session 2 were started after the `SELECT pg_stat_statements_reset()` command in session 2 was run.

The following shows the results from the `INSERT` commands in the two sessions. **RECORD 3** shows the results from session 1. **RECORD 2** shows the results from session 2.

```
edb=# SELECT query, rows, total_time, shared_blk_dirtied FROM
pg_stat_statements;
-[ RECORD 1 ]-----+
query      | SELECT pg_stat_statements_reset();
rows       | 1
total_time | 0.43
shared_blk_dirtied | 0
-[ RECORD 2 ]-----+
query      | INSERT INTO t2 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 30591.551
shared_blk_dirtied | 10003
-[ RECORD 3 ]-----+
query      | INSERT INTO t1 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 33215.334
shared_blk_dirtied | 10003
```

First, note the total time was 33215.334 milliseconds for session 1 and 30591.551 milliseconds for session 2. When only one session was active in the same resource group as shown in the first example, the time was 13496.184 milliseconds. Thus more active processes in the resource group result in a slower dirty rate for each active process in the group. This is shown in the following calculations.

The actual dirty rate for session 1 is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 33215.334 ms, which yields *0.30115609 blocks per millisecond*.
- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *301.15609 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *2467 kilobytes per second*.

The actual dirty rate for session 2 is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 30591.551 ms, which yields *0.32698571*

blocks per millisecond.

- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields 326.98571 *blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately 2679 *kilobytes per second*.

The combined dirty rate from session 1 (2467 kilobytes per second) and from session 2 (2679 kilobytes per second) yields 5146 kilobytes per second, which is below the set dirty rate limit of the resource group (6144 kilobytes per seconds).

Example – Multiple Processes in Multiple Groups

In this example, two additional `psql` sessions are used along with the previous two sessions. The third and fourth sessions perform the same `INSERT` command in resource group `resgrp_c` with a `dirty_rate_limit` of 3072 kilobytes per second.

Sessions 1 and 2 are repeated as illustrated in the prior example using resource group `resgrp_b` with a `dirty_rate_limit` of 6144 kilobytes per second.

Session 3:

```
edb=# CREATE TABLE t3 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLFACTOR = 10);
CREATE TABLE
edb=# SET edb_resource_group TO resgrp_c;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
```

```
resgrp_c
(1 row)
```

```
edb=# INSERT INTO t3 VALUES (generate_series (1,10000), 'aaa');
INSERT 0 10000
```

Session 4:

```
edb=# CREATE TABLE t4 (c1 INTEGER, c2 CHARACTER(500)) WITH (FILLFACTOR =
10);
CREATE TABLE
edb=# SET edb_resource_group TO resgrp_c;
SET
edb=# SHOW edb_resource_group;
edb_resource_group
-----
```

```
resgrp_c
(1 row)
```

```
edb=# SELECT pg_stat_statements_reset();
pg_stat_statements_reset
-----
```

```
(1 row)
```

```
edb=# INSERT INTO t4 VALUES (generate_series (1,10000), 'aaa');
```

INSERT 0 10000

Note

The `INSERT` commands in all four sessions were started after the `SELECT pg_stat_statements_reset()` command in session 4 was run.

The following shows the results from the `INSERT` commands in the four sessions.

`RECORD 3` shows the results from session 1. `RECORD 2` shows the results from session 2.

`RECORD 4` shows the results from session 3. `RECORD 5` shows the results from session 4.

```
edb=# SELECT query, rows, total_time, shared_blks_dirtied FROM
pg_stat_statements;
-[ RECORD 1 ]-----+
query      | SELECT pg_stat_statements_reset();
rows       | 1
total_time | 0.467
shared_blks_dirtied | 0
-[ RECORD 2 ]-----+
query      | INSERT INTO t2 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 31343.458
shared_blks_dirtied | 10003
-[ RECORD 3 ]-----+
query      | INSERT INTO t1 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 28407.435
shared_blks_dirtied | 10003
-[ RECORD 4 ]-----+
query      | INSERT INTO t3 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 52727.846
shared_blks_dirtied | 10003
-[ RECORD 5 ]-----+
query      | INSERT INTO t4 VALUES (generate_series (?,?), ?);
rows       | 10000
total_time | 56063.697
shared_blks_dirtied | 10003
```

First note that the times of session 1 (28407.435) and session 2 (31343.458) are close to each other as they are both in the same resource group with `dirty_rate_limit` set to 6144, as compared to the times of session 3 (52727.846) and session 4 (56063.697), which are in the resource group with `dirty_rate_limit` set to 3072. The latter group has a slower dirty rate limit so the expected processing time is longer as is the case for sessions 3 and 4.

The actual dirty rate for session 1 is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 28407.435 ms, which yields *0.35212612 blocks per millisecond*.
- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *352.12612 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *2885 kilobytes per second*.

The actual dirty rate for session 2 is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 31343.458 ms, which yields *0.31914156*

blocks per millisecond.

- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *319.14156 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *2614 kilobytes per second*.

The combined dirty rate from session 1 (2885 kilobytes per second) and from session 2 (2614 kilobytes per second) yields 5499 kilobytes per second, which is near the set dirty rate limit of the resource group (6144 kilobytes per seconds).

The actual dirty rate for session 3 is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 52727.846 ms, which yields *0.18971001 blocks per millisecond*.
- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *189.71001 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *1554 kilobytes per second*.

The actual dirty rate for session 4 is calculated as follows.

- The number of blocks dirtied per millisecond (ms) is 10003 blocks / 56063.697 ms, which yields *0.17842205 blocks per millisecond*.
- Multiply the result by 1000 to give the number of shared blocks dirtied per second (1 second = 1000 ms), which yields *178.42205 blocks per second*.
- Multiply the result by 8.192 to give the number of kilobytes dirtied per second (1 block = 8.192 kilobytes), which yields approximately *1462 kilobytes per second*.

The combined dirty rate from session 3 (1554 kilobytes per second) and from session 4 (1462 kilobytes per second) yields 3016 kilobytes per second, which is near the set dirty rate limit of the resource group (3072 kilobytes per seconds).

Thus, this demonstrates how EDB Resource Manager keeps the aggregate dirty rate of the active processes in its groups close to the dirty rate limit set for each group.

System Catalogs

This section describes the system catalogs that store the resource group information used by EDB Resource Manager.

edb_all_resource_groups

The following table lists the information available in the `edb_all_resource_groups` catalog:

| Column | Type | Description |
|------------------------------|---------|--|
| group_name | name | The name of the resource group. |
| active_processes | integer | Number of currently active processes in the resource group. |
| cpu_rate_limit | float8 | Maximum CPU rate limit for the resource group. 0 means no limit. |
| per_process_cpu_rate_limit | float8 | Maximum CPU rate limit per currently active process in the resource group. |
| dirty_rate_limit | float8 | Maximum dirty rate limit for a resource group. 0 means no limit. |
| per_process_dirty_rate_limit | float8 | Maximum dirty rate limit per currently active process in the resource group. |

edb_resource_group

The following table lists the information available in the `edb_resource_group` catalog:

| Column | Type | Description |
|--------------------|--------|--|
| rgrpname | name | The name of the resource group. |
| rgrpcpuratelimit | float8 | Maximum CPU rate limit for a resource group. 0 means no limit. |
| rgrpdirtyratelimit | float8 | Maximum dirty rate limit for a resource group. 0 means no limit. |

11.11 libpq C Library

libpq is the C application programmer's interface to Advanced Server. libpq is a set of library functions that allow client programs to pass queries to the Advanced Server and to receive the results of these queries.

libpq is also the underlying engine for several other EnterpriseDB application interfaces including those written for C++, Perl, Python, Tcl and ECPG. So some aspects of libpq's behavior will be important to the user if one of those packages is used.

Client programs that use libpq must include the header file `libpq-fe.h` and must link with the `libpq` library.

Using libpq with EnterpriseDB SPL

The EnterpriseDB SPL language can be used with the libpq interface library, providing support for:

- Procedures, functions, packages
- Prepared statements
- `REFCURSORs`
- Static cursors
- `structs` and `typedefs`
- Arrays
- DML and DDL operations
- `IN/OUT/IN OUT` parameters

REFCURSOR Support

In earlier releases, Advanced Server provided support for REFCURSORs through the following libpq functions; these functions should now be considered deprecated:

- `POCursorResult()`
- `PQgetCursorResult()`
- `PQnCursor()`

You may now use `POexec()` and `PQgetvalue()` to retrieve a `REFCURSOR` returned by an SPL (or PL/pgSQL) function. A `REFCURSOR` is returned in the form of a null-terminated string indicating the name of the cursor. Once you have the name of the cursor, you can execute one or more `FETCH` statements to retrieve the values exposed through the cursor.

Please note that the samples that follow do not include error-handling code that would be required in a real-world client application.

Returning a Single REFCURSOR

The following example shows an SPL function that returns a value of type **REFCURSOR**:

```
CREATE OR REPLACE FUNCTION getEmployees(p_deptno NUMERIC)
RETURN REFCURSOR AS
result REFCURSOR;
BEGIN
OPEN result FOR SELECT * FROM emp WHERE deptno = p_deptno;

RETURN result;
END;
```

This function expects a single parameter, **p_deptno**, and returns a **REFCURSOR** that holds the result set for the **SELECT** query shown in the **OPEN** statement. The **OPEN** statement executes the query and stores the result set in a cursor. The server constructs a name for that cursor and stores the name in a variable (named **result**). The function then returns the name of the cursor to the caller.

To call this function from a C client using libpq, you can use **PQexec()** and **PQgetvalue()**:

```
#include <stdio.h>
#include <stdlib.h>
#include "libpq-fe.h"

static void fetchAllRows(PGconn *conn,
    const char *cursorName,
    const char *description);
static void fail(PGconn *conn, const char *msg);

int
main(int argc, char *argv[])
{
    PGconn    *conn = PQconnectdb(argv[1]);
    PGresult  *result;

    if (PQstatus(conn) != CONNECTION_OK)
        fail(conn, PQerrorMessage(conn));

    result = PQexec(conn, "BEGIN TRANSACTION");

    if (PQresultStatus(result) != PGRES_COMMAND_OK)
        fail(conn, PQerrorMessage(conn));

    PQclear(result);

    result = PQexec(conn, "SELECT * FROM getEmployees(10)");

    if (PQresultStatus(result) != PGRES_TUPLES_OK)
        fail(conn, PQerrorMessage(conn));

    fetchAllRows(conn, PQgetvalue(result, 0, 0), "employees");

    PQclear(result);

    PQexec(conn, "COMMIT");

    PQfinish(conn);
```

```

    exit(0);
}

static void
fetchAllRows(PGconn *conn,
    const char *cursorName,
    const char *description)
{
    size_t commandLength = strlen("FETCH ALL FROM ") +
        strlen(cursorName) + 3;

    char *commandText = malloc(commandLength);
    PGresult *result;
    int     row;

    sprintf(commandText, "FETCH ALL FROM \"%s\"", cursorName);

    result = PQexec(conn, commandText);

    if (PQresultStatus(result) != PGRES_TUPLES_OK)
        fail(conn, PQerrorMessage(conn));

    printf("-- %s --\n", description);

    for (row = 0; row < PQntuples(result); row++)
    {
        const char *delimiter = "\t";
        int      col;

        for (col = 0; col < PQnfields(result); col++)
        {
            printf("%s%s", delimiter, PQgetvalue(result, row, col));
            delimiter = ",";
        }

        printf("\n");
    }

    PQclear(result);
    free(commandText);
}

static void
fail(PGconn *conn, const char *msg)
{
    fprintf(stderr, "%s\n", msg);

    if (conn != NULL)
        PQfinish(conn);

    exit(-1);
}

```

The code sample contains a line of code that calls the `getEmployees()` function, and returns a result set that contains all of the employees in department `10`:

```
result = PQexec(conn, "SELECT * FROM getEmployees(10)");
```

The `PQexec()` function returns a result set handle to the C program. The result set will contain exactly one value; that value is the name of the cursor as returned by `getEmployees()`.

Once you have the name of the cursor, you can use the SQL `FETCH` statement to retrieve the rows in that cursor. The function `fetchAllRows()` builds a `FETCH ALL` statement, executes that statement, and then prints the result set of the `FETCH ALL` statement.

The output of this program is shown below:

```
-- employees --
7782,CLARK,MANAGER,7839,09-JUN-81 00:00:00,2450.00,,10
7839,KING,PRESIDENT,,17-NOV-81 00:00:00,5000.00,,10
7934,MILLER,CLERK,7782,23-JAN-82 00:00:00,1300.00,,10
```

Returning Multiple REFCURSORS

The next example returns two `REFCURSORs`:

- The first `REFCURSOR` contains the name of a cursor (`employees`) that contains all employees who work in a department within the range specified by the caller.
- The second `REFCURSOR` contains the name of a cursor (`departments`) that contains all of the departments in the range specified by the caller.

In this example, instead of returning a single `REFCURSOR`, the function returns a `SETOF REFCURSOR` (which means `0` or more `REFCURSORS`). One other important difference is that the libpq program should not expect a single `REFCURSOR` in the result set, but should expect two rows, each of which will contain a single value (the first row contains the name of the `employees` cursor, and the second row contains the name of the `departments` cursor).

```
CREATE OR REPLACE FUNCTION getEmpsAndDepts(p_min NUMERIC,
                                         p_max NUMERIC)
RETURNS SETOF REFCURSOR AS
  employees REFCURSOR;
  departments REFCURSOR;
BEGIN
  OPEN employees FOR
    SELECT * FROM emp WHERE deptno BETWEEN p_min AND p_max;
  RETURN NEXT employees;

  OPEN departments FOR
    SELECT * FROM dept WHERE deptno BETWEEN p_min AND p_max;
  RETURN NEXT departments;
END;
```

As in the previous example, you can use `PQexec()` and `PQgetvalue()` to call the SPL function:

```
#include <stdio.h>
#include <stdlib.h>
#include "libpq-fe.h"

static void fetchAllRows(PGconn *conn,
                       const char *cursorName,
                       const char *description);
static void fail(PGconn *conn, const char *msg);

int
```

```

main(int argc, char *argv[])
{
    PGconn *conn = PQconnectdb(argv[1]);
    PGresult *result;

    if (PQstatus(conn) != CONNECTION_OK)
        fail(conn, PQerrorMessage(conn));

    result = PQexec(conn, "BEGIN TRANSACTION");

    if (PQresultStatus(result) != PGRES_COMMAND_OK)
        fail(conn, PQerrorMessage(conn));

    PQclear(result);

    result = PQexec(conn, "SELECT * FROM getEmpsAndDepts(20, 30)");

    if (PQresultStatus(result) != PGRES_TUPLES_OK)
        fail(conn, PQerrorMessage(conn));

    fetchAllRows(conn, PQgetvalue(result, 0, 0), "employees");
    fetchAllRows(conn, PQgetvalue(result, 1, 0), "departments");

    PQclear(result);

    PQexec(conn, "COMMIT");

    PQfinish(conn);

    exit(0);
}

static void
fetchAllRows(PGconn *conn,
            const char *cursorName,
            const char *description)
{
    size_t    commandLength = strlen("FETCH ALL FROM ") +
                           strlen(cursorName) + 3;
    char     *commandText   = malloc(commandLength);
    PGresult *result;
    int      row;

    sprintf(commandText, "FETCH ALL FROM \"%s\"", cursorName);

    result = PQexec(conn, commandText);

    if (PQresultStatus(result) != PGRES_TUPLES_OK)
        fail(conn, PQerrorMessage(conn));

    printf("-- %s --\n", description);

    for (row = 0; row < PQntuples(result); row++)
    {
        const char *delimiter = "\t";

```

```

int      col;

for (col = 0; col < PQnfields(result); col++)
{
    printf("%s%s", delimiter, PQgetvalue(result, row, col));
    delimiter = ",";
}

printf("\n");
}

PQclear(result);
free(commandText);
}

static void
fail(PGconn *conn, const char *msg)
{
    fprintf(stderr, "%s\n", msg);

    if (conn != NULL)
        PQfinish(conn);

    exit(-1);
}

```

If you call `getEmpsAndDepts(20, 30)`, the server will return a cursor that contains all employees who work in department `20` or `30`, and a second cursor containing the description of departments `20` and `30`.

```

-- employees --
7369,SMITH,CLERK,7902,17-DEC-80 00:00:00,800.00,,20
7499,ALLEN,SALESMAN,7698,20-FEB-81 00:00:00,1600.00,300.00,30
7521,WARD,SALESMAN,7698,22-FEB-81 00:00:00,1250.00,500.00,30
7566,JONES,MANAGER,7839,02-APR-81 00:00:00,2975.00,,20
7654,MARTIN,SALESMAN,7698,28-SEP-81 00:00:00,1250.00,1400.00,30
7698,BLAKE,MANAGER,7839,01-MAY-81 00:00:00,2850.00,,30
7788,SCOTT,ANALYST,7566,19-APR-87 00:00:00,3000.00,,20
7844,TURNER,SALESMAN,7698,08-SEP-81 00:00:00,1500.00,0.00,30
7876,ADAMS,CLERK,7788,23-MAY-87 00:00:00,1100.00,,20
7900,JAMES,CLERK,7698,03-DEC-81 00:00:00,950.00,,30
7902,FORD,ANALYST,7566,03-DEC-81 00:00:00,3000.00,,20
-- departments --
20,RESEARCH,DALLAS
30,SALES,CHICAGO

```

Array Binding

Advanced Server's array binding functionality allows you to send an array of data across the network in a single round-trip. When the back end receives the bulk data, it can use the data to perform insert or update operations.

Perform bulk operations with a prepared statement; use the following function to prepare the statement:

```

PGresult *PQprepare(PGconn *conn,
                    const char *stmtName,

```

```
const char *query,
int nParams,
const Oid *paramTypes);
```

Details of [PQprepare\(\)](#) can be found in the prepared statement section.

The following functions can be used to perform bulk operations:

- [PQBulkStart](#)
- [PQexecBulk](#)
- [PQBulkFinish](#)
- [PQexecBulkPrepared](#)

PQBulkStart

[PQBulkStart\(\)](#) initializes bulk operations on the server. You must call this function before sending bulk data to the server. [PQBulkStart\(\)](#) initializes the prepared statement specified in stmtName to receive data in a format specified by [paramFmts](#).

API Definition

```
PGresult *PQBulkStart(PGconn *conn,
                      const char *Stmt_Name,
                      unsigned int nCol,
                      const int *paramFmts);
```

PQexecBulk

[PQexecBulk\(\)](#) is used to supply data ([paramValues](#)) for a statement that was previously initialized for bulk operation using [PQBulkStart\(\)](#).

This function can be used more than once after [PQBulkStart\(\)](#) to send multiple blocks of data. See the example for more details.

API Definition

```
PGresult *PQexecBulk(PGconn *conn,
                     unsigned int nRows,
                     const char **const paramValues,
                     const int *paramLengths);
```

PQBulkFinish

This function completes the current bulk operation. You can use the prepared statement again without re-preparing it.

API Definition

```
PGresult *PQBulkFinish(PGconn *conn);
```

PQexecBulkPrepared

Alternatively, you can use the [PQexecBulkPrepared\(\)](#) function to perform a bulk operation with a single function call. [PQexecBulkPrepared\(\)](#) sends a request to execute a prepared statement with the given parameters, and waits for the result. This function combines the functionality of [PQbulkStart\(\)](#), [PQexecBulk\(\)](#), and [PQbulkFinish\(\)](#). When using this function, you are not required to initialize or terminate the bulk operation; this function starts the bulk operation, passes the data to the server, and terminates the bulk operation.

Specify a previously prepared statement in the place of stmtName. Commands that will be used repeatedly will be parsed and planned just once, rather than each time they are executed.

API Definition

```
PGresult *PQexecBulkPrepared(PGconn *conn,
                           const char *stmtName,
                           unsigned int nCols,
                           unsigned int nRows,
                           const char *const *paramValues,
                           const int *paramLengths,
                           const int *paramFormats);
```

Example Code (Using PQbulkStart, PQexecBulk, PQbulkFinish)

The following example uses [PQBulkStart](#), [PQexecBulk](#), and [PQbulkFinish](#).

```
void InsertDataUsingBulkStyle( PGconn *conn )
{
    PGresult      *res;
    Oid           paramTypes[2];
    char          *paramVals[5][2];
    int           paramLens[5][2];
    int           paramFmts[2];
    int           i;

    int           a[5] = { 10, 20, 30, 40, 50 };
    char          b[5][10] = { "Test_1", "Test_2", "Test_3",
                            "Test_4", "Test_5" };

    paramTypes[0] = 23;
    paramTypes[1] = 1043;
    res = PQprepare( conn, "stmt_1", "INSERT INTO testtable1 values( $1, $2
)", 2, paramTypes );
    PQclear( res );

    paramFmts[0] = 1; /* Binary format */
    paramFmts[1] = 0;

    for( i = 0; i < 5; i++ )
    {
        a[i] = htonl( a[i] );
        paramVals[i][0] = &(a[i]);
        paramVals[i][1] = b[i];

        paramLens[i][0] = 4;
        paramLens[i][1] = strlen( b[i] );
    }
}
```

```

res = PQBulkStart(conn, "stmt_1", 2, paramFmts);
PQclear( res );
printf( "< -- PQBulkStart -- >\n" );

res = PQexecBulk(conn, 5, (const char *const *)paramVals, (const int
*)paramLens);
PQclear( res );
printf( "< -- PQexecBulk -- >\n" );

res = PQexecBulk(conn, 5, (const char *const *)paramVals, (const int
*)paramLens);
PQclear( res );
printf( "< -- PQexecBulk -- >\n" );

res = PQBulkFinish(conn);
PQclear( res );
printf( "< -- PQBulkFinish -- >\n" );
}

```

Example Code (Using PQexecBulkPrepared)

The following example uses `PQexecBulkPrepared`.

```

void InsertDataUsingBulkStyleCombinedVersion( PGconn *conn )
{
    PGresult      *res;
    Oid           paramTypes[2];
    char          *paramVals[5][2];
    int           paramLens[5][2];
    int           paramFmts[2];
    int           i;

    int           a[5] = { 10, 20, 30, 40, 50 };
    char          b[5][10] = { "Test_1", "Test_2", "Test_3",
                            "Test_4", "Test_5" };

    paramTypes[0] = 23;
    paramTypes[1] = 1043;
    res = PQprepare( conn, "stmt_2", "INSERT INTO testtable1 values( $1, $2
)", 2, paramTypes );
    PQclear( res );

    paramFmts[0] = 1; /* Binary format */
    paramFmts[1] = 0;

    for( i = 0; i < 5; i++ )
    {
        a[i] = htonl( a[i] );
        paramVals[i][0] = &(a[i]);
        paramVals[i][1] = b[i];

        paramLens[i][0] = 4;
        paramLens[i][1] = strlen( b[i] );
    }
}

```

```

    }

res = PQexecBulkPrepared(conn, "stmt_2", 2, 5, (const char *const
*)paramVals,(const int *)paramLens, (const int *)paramFmts);
PQclear( res );
}

```

11.12 Debugger

The Debugger gives developers and DBAs the ability to test and debug server-side programs using a graphical, dynamic environment. The types of programs that can be debugged are SPL stored procedures, functions, triggers, and packages as well as PL/pgSQL functions and triggers.

The Debugger is integrated with *pgAdmin 4* and *EDB Postgres Enterprise Manager*. If you have installed Advanced Server on a Windows host, pgAdmin 4 is automatically installed; you will find the pgAdmin 4 icon in the [Windows Start](#) menu. If your Advanced Server host is on a CentOS or Linux system, you can use [yum](#) to install pgAdmin4. Open a command line, assume superuser privileges, and enter:

```
| yum install edb-pgadmin4*
```

On Linux, the `edb-asxx-server-pldebugger` RPM package where `xx` is the Advanced Server version number, must be installed as well. Information about pgAdmin 4 is available at:

```
| https://www.pgadmin.org/
```

The RPM installation will add the pgAdmin4 icon to your Applications menu.

There are two basic ways the Debugger can be used to test programs:

- **Standalone Debugging.** The Debugger is used to start the program to be tested. You supply any input parameter values required by the program and you can immediately observe and step through the code of the program. Standalone debugging is the typical method used for new programs and for initial problem investigation.
- **In-Context Debugging.** The program to be tested is initiated by an application other than the Debugger. You first set a *global breakpoint* on the program to be tested. The application that makes the first call to the program encounters the global breakpoint. The application suspends execution at which point the Debugger takes control of the called program. You can then observe and step through the code of the called program as it runs within the context of the calling application. After you have completely stepped through the code of the called program in the Debugger, the suspended application resumes execution. In-context debugging is useful if it is difficult to reproduce a problem using standalone debugging due to complex interaction with the calling application.

The debugging tools and operations are the same whether using standalone or in-context debugging. The difference is in how the program to be debugged is invoked.

The following sections discuss the features and functionality of the Debugger using the standalone debugging method. The directions for starting the Debugger for in-context debugging are discussed in the [Setting Global Breakpoint for In-Context Debugging](#).

Configuring the Debugger

Before using the Debugger, edit the `postgresql.conf` file (located in the `data` subdirectory of your Advanced Server home directory), adding `$libdir/plugin_debugger` to the libraries listed in the `shared_preload_libraries`

configuration parameter:

```
shared_preload_libraries = '$libdir/dbms_pipe,$libdir/edb_gen,$libdir/
plugin_debugger'
```


- On Linux, the `postgresql.conf` file is located in: `/var/lib/edb/asxx/data`
- On Windows, the `postgresql.conf` file is located in: `C:\Program Files\edb\asxx\data`

Where `x` is the version of Advanced Server.


After modifying the `shared_preload_libraries` parameter, you must restart the database server.

Starting the Debugger

Use pgAdmin 4 to access the Debugger for standalone debugging. To open the Debugger, highlight the name of the stored procedure or function you wish to debug in the pgAdmin 4 `Browser` panel. Then, navigate through the `Object` menu to the `Debugging` menu and select `Debug` from the submenu.



You can also right-click on the name of the stored procedure or function in the pgAdmin 4 `Browser`, and select `Debugging`, and the `Debug` from the context menu.



Note that triggers cannot be debugged using standalone debugging. Triggers must be debugged using in-context debugging. See the [Setting Global Breakpoint for In-Context Debugging](#) for information on setting a global breakpoint for in-context debugging.

To debug a package, highlight the specific procedure or function under the package node of the package you wish to debug and follow the same directions as for stored procedures and functions.

The Debugger Window

You can use the **Debugger** window to pass parameter values when you are standalone-debugging a program that expects parameters. When you start the debugger, the **Debugger** window opens automatically to display any **IN** or **IN OUT** parameters expected by the program. If the program declares no **IN** or **IN OUT** parameters, the **Debugger** window does not open.

| Name | Type | Null? | Expression? | Value | Use Default? | Default value |
|--------|---------|-------------------------------------|-------------------------------------|-------|--------------------------|--------------------|
| p_sal | numeric | <input checked="" type="checkbox"/> | <input checked="" type="checkbox"/> | NaN | <input type="checkbox"/> | <No default value> |
| p_comm | numeric | <input checked="" type="checkbox"/> | <input checked="" type="checkbox"/> | NaN | <input type="checkbox"/> | <No default value> |

Use the fields on the **Debugger** window to provide a value for each parameter:

- The **Name** field contains the formal parameter name.
- The **Type** field contains the parameter data type.
- Check the **Null?** checkbox to indicate that the parameter is a **NULL** value.

- Check the **Expression?** checkbox if the **Value** field contains an expression.
- The **Value** field contains the parameter value that will be passed to the program.
- Check the **Use Default?** checkbox to indicate that the program should use the value in the **Default Value** field.
- The **Default Value** field contains the default value of the parameter.

Press the **Tab** key to select the next parameter in the list for data entry, or click on a **Value** field to select the parameter for data entry.

If you are debugging a procedure or function that is a member of a package that has an initialization section, check the **Debug Package Initializer** check box to instruct the Debugger to step into the package initialization section, allowing you to debug the initialization section code before debugging the procedure or function. If you do not select the check box, the Debugger executes the package initialization section without allowing you to see or step through the individual lines of code as they are executed.

After entering the desired parameter values, click the **Debug** button to start the debugging process. Click the **Cancel** button to terminate the Debugger.

Note

The **Debugger** window does not open during in-context debugging. Instead, the application calling the program to be debugged must supply any required input parameter values.

When you have completed a full debugging cycle by stepping through the program code, the **Debugger** window re-opens, allowing you to enter new parameter values and repeat the debugging cycle, or end the debugging session.

Main Debugger Window

The Main Debugger window contains two panels:

- The top **Program Body** panel displays the program source code.
- The bottom **Tabs** panel provides a set of tabs for different information.

Use the **Tool Bar** icons located at the top panel to access debugging functions.

The screenshot shows the pgAdmin 4 interface with the 'Program Body' panel open. The left sidebar displays a tree view of database objects, including Schemas, Functions, and Views. The main panel shows the source code for a function named 'select_emp'. The code includes declarations for variables like v_ename, v_hiredate, v_sal, v_comm, v_dname, and v_disp_date, followed by a SELECT INTO statement and a WHERE clause. Below the code, there is a table titled 'Parameters' with one entry: p_empno of type integer with value 7900.

```

1  DECLARE
2  v_ename emp.ename%TYPE;
3  v_hiredate emp.hiredate%TYPE;
4  v_sal emp.sal%TYPE;
5  v_comm emp.comm%TYPE;
6  v_dname dept.dname%TYPE;
7  v_disp_date VARCHAR(10);
8  BEGIN
9    SELECT INTO
10   v_ename, v_hiredate, v_sal, v_comm, v_dname
11   ename, hiredate, sal, COALESCE(comm, 0), dname
12   FROM emp e, dept d
13   WHERE empno = p_empno;
14 END;

```

| Parameters | Type | Value |
|------------|---------|-------|
| p_empno | integer | 7900 |

The two panels are described in the following sections.

The Program Body Panel

The **Program Body** panel displays the source code of the program that is being debugged.

The figure shows that the Debugger is about to execute the `SELECT` statement. The blue indicator in the program body highlights the next statement to execute.

The Tabs Panel

You can use the bottom `Tabs` panel to view or modify parameter values or local variables, or to view messages generated by `RAISE INFO` and function results.

The following is the information displayed by the tabs in the panel:

- The `Parameters` tab displays the current parameter values.
- The `Local variables` tab displays the value of any variables declared within the program.
- The `Messages` tab displays any results returned by the program as it executes.
- The `Results` tab displays program results (if applicable) such as the value from the `RETURN` statement of a function.
- The `Stack tab` displays the call stack.

The following figures show the results from the various tabs.

pgAdmin 4

File ▾ Object ▾ Tools ▾ Help ▾

Browser

- > Extensions (6)
- > Foreign Data Wrappers (2)
- > Languages (2)
- > Schemas (2)
 - > enterprisesdb
 - > public
 - > Collations
 - > Domains
 - > FTS Configurations
 - > FTS Dictionaries
 - > FTS Parsers
 - > FTS Templates
 - > Foreign Tables
 - > Functions (6)
 - < emp_comp(p_sa
 - < emp_query_calle
 - < hire_clerk(p_en
 - < hire_salesman(p
 - < new_empno()
 - < select_emp(p_en
- > Materialized Views
- > Packages
- > Procedures
- > Sequences
- > Synonyms
- > Tables
- > Trigger Functions
- > Types
- > Views

Dashboard Properties SQL Statistics Dependencies Dependents Debugger

```
select_emp(p_empno integer)
1 DECLARE
2 v_ename emp.ename%TYPE;
3 v_hiredate emp.hiredate%TYPE;
4 v_sal emp.sal%TYPE;
5 v_comm emp.comm%TYPE;
6 v_dname dept.dname%TYPE;
7 v_disp_date VARCHAR(10);
8 BEGIN
9 SELECT INTO
10 v_ename, v_hiredate, v_sal, v_comm, v_dname
11 ename, hiredate, sal, COALESCE(comm, 0), dname
12 FROM emp e, dept d
13 WHERE empno = p_empno;
```

Parameters Local variables Messages Results Stack

| Name | Type | Value |
|---------|---------|-------|
| p_empno | integer | 7900 |

pgAdmin 4

File ▾ Object ▾ Tools ▾ Help ▾

Browser

- > Extensions (6)
- > Foreign Data Wrappers (2)
- > Languages (2)
- > Schemas (2)
 - > enterprisesdb
 - > public
 - > Collation
 - > Domain
 - > FTS Configuration
 - > FTS Dictionary
 - > FTS Parser
 - > FTS Template
 - > Foreign Table
 - > Function
 - < emp_comp(p_sa
 - < emp_query_calle
 - < hire_clerk(p_en
 - < hire_salesman(p
 - < new_empno()
 - < select_emp(p_en
 - > Materialized View
 - > Package
 - > Procedure
 - > Sequence
 - > Synonym
 - > Table
 - > Trigger Function
 - > Type
 - > View

Dashboard Properties SQL Statistics Dependencies Dependents Debugger

```
emp_query_caller()
1 DECLARE
2 Create
3 no NUMERIC;
4 o NUMERIC;
5 e VARCHAR;
6 i INTEGER;
7 query TEXT;
```

Debugging > Debug

Set breakpoint

```
e := 'Martin';
query := emp_query(v_deptno, v_empno, v_ename);
INFO 'Department : %', v_deptno;
INFO 'Employee No: %', (r_emp_query).empno;
```

Parameters Local variables Messages Results Stack

| Name | Type | Value |
|-------------|-------------------|-------|
| v_deptno | numeric | NULL |
| v_empno | numeric | NULL |
| v_ename | character varying | NULL |
| v_rows | integer | NULL |
| r_emp_query | text | NULL |

pgAdmin 4

File ▾ Object ▾ Tools ▾ Help ▾

Browser

- > Extensions (6)
- > Foreign Data Wrappers (2)
- > Languages (2)
- > Schemas (2)
 - > enterprisesdb
 - > public
 - > Collations
 - > Domains
 - > FTS Configurations
 - > FTS Dictionaries
 - > FTS Parsers
 - > FTS Templates
 - > Foreign Tables
 - > Functions (6)
 - < emp_comp(p_ename, p_deptno)
 - < emp_query_call(p_ename, p_deptno)
 - < hire_clerk(p_ename, p_deptno)
 - < hire_salesman(p_ename, p_deptno)
 - < new_empno()
 - < select_emp(p_ename, p_deptno)
- > Materialized Views
- > Packages
- > Procedures
- > Sequences
- > Synonyms
- > Tables
- > Trigger Functions
- > Types
- > Views

Debugger

```
hire_clerk(p_ename character varying, p_deptno numeric)

1   v_empno      NUMBER(4);
2   v_ename       VARCHAR2(10);
3   v_job        VARCHAR2(9);
4   v_mgr         NUMBER(4);
5   v_hiredate    DATE;
6   v_sal         NUMBER(7,2);
7   v_comm        NUMBER(7,2);
8   v_deptno      NUMBER(2);

10  BEGIN
11    v_empno := new_empno;
12    INSERT INTO emp VALUES (v_empno, p_ename, 'CLERK', 7782,
13                           TRUNC(SYSDATE), 950.00, NULL, p_deptno);

Parameters Local variables Messages Results Stack
```

Messages tab content:

Inserting employee 8001
..New salary: 950.00
User enterprisesdb added employee(s) on 2019-04-01
Department :
Employee No: 8001
Name :
Job : CLERK
Manager : 7782
Hire Date : 2019-04-01 00:00:00
Salary : 950.00
Commission :

pgAdmin 4

File ▾ Object ▾ Tools ▾ Help ▾

Browser

- > Extensions (6)
- > Foreign Data Wrappers (2)
- > Languages (2)
- > Schemas (2)
 - > enterprisesdb
 - > public
 - > Collations
 - > Domains
 - > FTS Configurations
 - > FTS Dictionaries
 - > FTS Parsers
 - > FTS Templates
 - > Foreign Tables
 - > Functions (6)
 - < emp_comp(p_ename, p_deptno)
 - < emp_query_call(p_ename, p_deptno)
 - < hire_clerk(p_ename, p_deptno)
 - < hire_salesman(p_ename, p_deptno)
 - < new_empno()
 - < select_emp(p_ename, p_deptno)
- > Materialized Views
- > Packages
- > Procedures
- > Sequences
- > Synonyms
- > Tables
- > Trigger Functions
- > Types
- > Views

Debugger

```
hire_clerk(p_ename character varying, p_deptno numeric)

1   v_empno      NUMBER(4);
2   v_ename       VARCHAR2(10);
3   v_job        VARCHAR2(9);
4   v_mgr         NUMBER(4);
5   v_hiredate    DATE;
6   v_sal         NUMBER(7,2);
7   v_comm        NUMBER(7,2);
8   v_deptno      NUMBER(2);

10  BEGIN
11    v_empno := new_empno;
12    INSERT INTO emp VALUES (v_empno, p_ename, 'CLERK', 7782,
13                           TRUNC(SYSDATE), 950.00, NULL, p_deptno);

Parameters Local variables Messages Results Stack
```

Results tab content:

hire_clerk

| |
|------|
| 8001 |
|------|

The Stack Tab


The **Stack** tab displays a list of programs that are currently on the call stack (programs that have been invoked,

but which have not yet completed). When a program is called, the name of the program is added to the top of the list displayed in the **Stack** tab. When the program ends, its name is removed from the list.

The **Stack** tab also displays information about program calls. The information includes:

- The location of the call within the program
- The call arguments
- The name of the program being called

Reviewing the call stack can help you trace the course of execution through a series of nested programs.



The screenshot shows the pgAdmin 4 interface with the 'Debugger' tab selected. The left sidebar shows a tree view of database objects under 'Schemas'. In the main pane, the 'Program Body' panel displays the code for the 'emp_query_caller' program. The code is as follows:

```

1  DECLARE
2      v_deptno NUMERIC;
3      v_empno NUMERIC;
4      v_ename VARCHAR;
5      v_rows INTEGER;
6      r_emp_query TEXT;
7  BEGIN
8      v_deptno := 30;
9      v_empno := 0;
10     v_ename := 'Martin';
11     r_emp_query := emp_query(v_deptno, v_empno, v_ename);
12     RAISE INFO 'Department : %', v_deptno;
13     RAISE INFO 'Employee No: %', (r_emp_query).empno;
...

```

Below the code, there are tabs for 'Parameters', 'Local variables', 'Messages', 'Results', and 'Stack'. The 'Stack' tab is selected, showing a table with one row:

| Name | Value | Line No. |
|--------------------|-------|----------|
| emp_query_caller() | | 8 |

The above figure shows that `emp_query_caller` is about to call a subprogram named `emp_query`. `emp_query_caller` is currently at the top of the call stack.

After the call to `emp_query` executes, `emp_query` is displayed at the top of the **Stack** tab, and its code is displayed in the **Program Body** panel.

pgAdmin 4

File ▾ Object ▾ Tools ▾ Help ▾

Browser

- > Extensions (6)
- > Foreign Data Wrappers (2)
- > Languages (2)
- ▼ Schemas (2)
 - > enterprisesdb
 - < public
 - > Collations
 - > Domains
 - > FTS Configurations
 - > FTS Dictionaries
 - > FTS Parsers
 - > FTS Templates
 - > Foreign Tables
 - < Functions (6)
 - < emp_comp(p_s)
 - < emp_query_caller()
 - < hire_clerk(p_en)
 - < hire_salesman()
 - < new_empno()
 - < select_emp(p_en)
- > Materialized Views
- > Packages
- > Procedures
- > Sequences
- > Synonyms
- > Tables
- > Trigger Functions
- < Types
- > Views

Dashboard Properties SQL Statistics Dependencies Dependents Debugger

```
emp_query_caller()

1  DECLARE
2  v_deptno NUMERIC;
3  v.empno NUMERIC;
4  v.ename VARCHAR;
5  v.rows INTEGER;
6  r.emp_query TEXT;
7  BEGIN
8  v_deptno := 30;
9  v.empno := 0;
10 v.ename := 'Martin';
11 r.emp_query := emp_query(v_deptno, v.empno, v.ename);
12 RAISE INFO 'Department : %', v_deptno;
13 RAISE INFO 'Employee No: %', (r.emp_query).empno;
14 END;
```

Parameters Local variables Messages Results Stack

| Name | Value | Line No. |
|--------------------|-------|----------|
| emp_query_caller() | | 8 |

Upon completion of execution of the subprogram, control returns to the calling program (`emp_query_caller`), now displayed at the top of the `Stack` tab.

pgAdmin 4

File ▾ Object ▾ Tools ▾ Help ▾

Browser

- > Extensions (6)
- > Foreign Data Wrappers (2)
- > Languages (2)
- ▼ Schemas (2)
 - > enterprisesdb
 - < public
 - > Collations
 - > Domains
 - > FTS Configurations
 - > FTS Dictionaries
 - > FTS Parsers
 - > FTS Templates
 - > Foreign Tables
 - < Functions (6)
 - < emp_comp(p_s)
 - < emp_query_caller()
 - < hire_clerk(p_en)
 - < hire_salesman()
 - < new_empno()
 - < select_emp(p_en)
- > Materialized Views
- > Packages
- > Procedures
- > Sequences
- > Synonyms
- > Tables
- > Trigger Functions
- < Types
- > Views

Dashboard Properties SQL Statistics Dependencies Dependents Debugger

```
emp_query_caller()

1  DECLARE
2  v_deptno NUMERIC;
3  v.empno NUMERIC;
4  v.ename VARCHAR;
5  v.rows INTEGER;
6  r.emp_query TEXT;
7  BEGIN
8  v_deptno := 30;
9  v.empno := 0;
10 v.ename := 'Martin';
11 r.emp_query := emp_query(v_deptno, v.empno, v.ename);
12 RAISE INFO 'Department : %', v_deptno;
13 RAISE INFO 'Employee No: %', (r.emp_query).empno;
14 END;
```

Parameters Local variables Messages Results Stack

| Name | Value | Line No. |
|--------------------|-------|----------|
| emp_query_caller() | | 8 |

Debugging a Program

You can perform the following operations to debug a program:

- Step through the program one line at a time
- Execute the program until you reach a breakpoint
- View and change local variable values within the program

Stepping Through the Code

Use the tool bar icons to step through a program with the Debugger:



The icons serve the following purposes:

- **Step into.** Click the **Step into** icon to execute the currently highlighted line of code.
- **Step over.** Click the **Step over** icon to execute a line of code, stepping over any sub-functions invoked by the code. The sub-function executes, but is not debugged unless it contains a breakpoint.
- **Continue/Start.** Click the **Continue/Start** icon to execute the highlighted code, and continue until the program encounters a breakpoint or completes.
- **Stop.** Click the **Stop** icon to halt the execution of a program.

Using Breakpoints


As the Debugger executes a program, it pauses whenever it reaches a breakpoint. When the Debugger pauses, you can observe or change local variables, or navigate to an entry in the call stack to observe variables or set other breakpoints. The next step into, step over, or continue operation forces the debugger to resume execution with the next line of code following the breakpoint. There are two types of breakpoints:

Local Breakpoint - A local breakpoint can be set at any executable line of code within a program. The Debugger pauses execution when it reaches a line where a local breakpoint has been set.

Global Breakpoint - A global breakpoint will trigger when *any* session reaches that breakpoint. Set a global breakpoint if you want to perform in-context debugging of a program. When a global breakpoint is set on a program, the debugging session that set the global breakpoint waits until that program is invoked in another session. A global breakpoint can only be set by a superuser.

To create a local breakpoint, left-click within the grey shaded margin to the left of the line of code where you want the local breakpoint set. Where you click in the grey shaded margin should be close to the right side of the margin as in the spot where the breakpoint dot is shown on source code line 12.

When created, the Debugger displays a dark dot in the margin, indicating a breakpoint has been set at the selected line of code.



You can set as many local breakpoints as desired. Local breakpoints remain in effect for the duration of a debugging session until they are removed.

Removing a Local Breakpoint

To remove a local breakpoint, left-click the mouse on the breakpoint dot in the grey shaded margin of the **Program Body** panel. The dot disappears, indicating that the breakpoint has been removed.

You can remove all of the breakpoints from the program that currently appears in the **Program Body** frame by clicking the **Clear all breakpoints** icon.



Note

When you perform any of the preceding actions, only the breakpoints in the program that currently appears in the **Program Body** panel are removed. Breakpoints in called subprograms or breakpoints in programs that call the program currently appearing in the **Program Body** panel are not removed.

Setting a Global Breakpoint for In-Context Debugging

To set a global breakpoint for in-context debugging, highlight the stored procedure, function, or trigger on which you wish to set the breakpoint in the **Browser** panel. Navigate through the **Object** menu to select **Debugging**, and then **Set Breakpoint**.

The screenshot shows the pgAdmin 4 interface. In the left sidebar, under the 'Functions' node, a context menu is open for the 'select_emp(p_en)' function. The 'Debugging' option is selected, and its submenu shows 'Debug' and 'Set breakpoint' as options. The 'Set breakpoint' option is highlighted with a blue background. The main pane displays various monitoring charts: 'Transactions per second' (Transactions, Commits, Rollbacks), 'Tuples in' (Inserts, Updates, Deletes), 'Tuples out' (Fetched, Returned), and 'Block I/O' (Reads, Hits). Below these charts is a 'Server activity' section with tabs for 'Sessions', 'Locks', and 'Prepared Transactions'. The 'Sessions' tab is active, showing a table of backend sessions with columns: PID, User, Application, Client, Backend start, State, Wait Event, and Blocking.

| PID | User | Application | Client | Backend start | State | Wait Event | Blocking |
|------|--------------|--------------------------|--------|-------------------------|-------|--------------------|----------|
| 2568 | enterprisedb | pgAdmin 4 - CONN:6098332 | ::1 | 2019-04-01 15:04:56 IST | idle | Client: ClientRead | |
| 2992 | enterprisedb | pgAdmin 4 - CONN:1190633 | ::1 | 2019-04-01 15:04:56 IST | idle | Client: ClientRead | |
| 3732 | enterprisedb | pgAdmin 4 - CONN:2023868 | ::1 | 2019-04-01 15:04:57 IST | idle | Client: ClientRead | |
| 4200 | enterprisedb | pgAdmin 4 - CONN:2058600 | ::1 | 2019-04-01 15:04:57 IST | idle | Client: ClientRead | |

Alternatively, you can right-click on the name of the stored procedure, function, or trigger on which you wish to set a global breakpoint and select **Debugging**, then **Set Breakpoint** from the context menu as shown by the following.


This screenshot is similar to the one above but focuses on triggers. In the left sidebar, under the 'Triggers' node, a context menu is open for the 'select_emp(p_en)' trigger. The 'Debugging' option is selected, and its submenu shows 'Debug' and 'Set breakpoint' as options. The 'Set breakpoint' option is highlighted with a blue background. The main pane displays monitoring charts and a 'Server activity' section identical to the first screenshot.

| PID | User | Application | Client | Backend start | State | Wait Event | Blocking |
|------|--------------|--------------------------|--------|-------------------------|-------|--------------------|----------|
| 2568 | enterprisedb | pgAdmin 4 - CONN:6098332 | ::1 | 2019-04-01 15:04:56 IST | idle | Client: ClientRead | |
| 2992 | enterprisedb | pgAdmin 4 - CONN:1190633 | ::1 | 2019-04-01 15:04:56 IST | idle | Client: ClientRead | |
| 3732 | enterprisedb | pgAdmin 4 - CONN:2023868 | ::1 | 2019-04-01 15:04:57 IST | idle | Client: ClientRead | |
| 4200 | enterprisedb | pgAdmin 4 - CONN:2058600 | ::1 | 2019-04-01 15:04:57 IST | idle | Client: ClientRead | |

To set a global breakpoint on a trigger, expand the table node that contains the trigger, highlight the specific trigger you wish to debug, and follow the same directions as for stored procedures and functions.

To set a global breakpoint in a package, highlight the specific procedure or function under the package node of the package you wish to debug and follow the same directions as for stored procedures and functions.

After you choose **Set Breakpoint**, the Debugger window opens and waits for an application to call the program to be debugged.



The PSQL client invokes the `select_emp` function (on which a global breakpoint has been set).

```
$ psql edb enterprisedb
psql.bin (13.0.0, server 13.0.0)
Type "help" for help.
```

```
edb=# SELECT select_emp(7900);
```

The `select_emp` function does not complete until you step through the program in the Debugger.

The screenshot shows the pgAdmin 4 interface with the Debugger tab selected. On the left, the Browser pane lists various database objects like Extensions, Foreign Data Wrappers, Languages, Schemas, Functions, and Views. The Functions section is expanded, and the select_emp function is selected. The main pane displays the source code of the function:

```

1  DECLARE
2    v_ename emp.ename%TYPE;
3    v_hiredate emp.hiredate%TYPE;
4    v_sal emp.sal%TYPE;
5    v_comm emp.comm%TYPE;
6    v_dname dept.dname%TYPE;
7    v_disp_date VARCHAR(10);
8  BEGIN
9    SELECT INTO
10      v_ename, v_hiredate, v_sal, v_comm, v_dname
11      ename, hiredate, sal, COALESCE(comm, 0), dname
12    FROM emp e, dept d
13   WHERE empno = p_empno;

```

Below the code, there are several tabs: Parameters, Local variables, Messages, Results, and Stack. The Parameters tab is active, showing a table with one row:

| Name | Type | Value |
|---------|---------|-------|
| p_empno | integer | 7900 |

You can now debug the program using any of the previously discussed operations such as step into, step over, and continue, or set local breakpoints. When you have stepped through execution of the program, the calling application (PSQL) regains control and the `select_emp` function completes execution and its output is displayed.

```
$ psql edb enterprisedb
psql.bin (13.0.0, server 13.0.0)
Type "help" for help.
```

```
edb=# SELECT select_emp(7900);
INFO: Number : 7900
INFO: Name : JAMES
INFO: Hire Date : 12/03/1981
INFO: Salary : 950.00
INFO: Commission: 0.00
INFO: Department: SALES
select_emp
-----
(1 row)
```

At this point, you can end the Debugger session. If you do not end the Debugger session, the next application that invokes the program will encounter the global breakpoint and the debugging cycle will begin again.

Exiting the Debugger

To end a Debugger session and exit the Debugger, click on the close icon (x) located in the upper-right corner to close the tab.

The screenshot shows the EDB Postgres Advanced Server interface with the 'Debugger' tab selected. A script named 'hire_clerk' is displayed, containing PL/pgSQL code for inserting a new employee record into the 'emp' table. The code defines several local variables (v_empno, v_ename, v_job, v_mgr, v_hiredate, v_sal, v_comm, v_deptno) and performs an INSERT INTO statement. Below the code, the 'Messages' tab is active, showing the output of the script's execution: 'Inserting employee 8003' and '.New salary: 950.00'.

```

hire_clerk(p_ename character varying, p_deptno numeric)

1
2     v_empno      NUMBER(4);
3     v_ename      VARCHAR2(10);
4     v_job       VARCHAR2(9);
5     v_mgr       NUMBER(4);
6     v_hiredate    DATE;
7     v_sal        NUMBER(7,2);
8     v_comm        NUMBER(7,2);
9     v_deptno      NUMBER(2);

10    BEGIN
11        v_empno := new_empno;
12        INSERT INTO emp VALUES (v_empno, p_ename, 'CLERK', 7782,
13                                TRUNC(SYSDATE), 950.00, NULL, p_deptno);
14    END;

```

Parameters Local variables Messages Results Stack

Inserting employee 8003
..New salary: 950.00

11.13 Performance Analysis and Tuning

Advanced Server provides various tools for performance analysis and tuning. These features are described in this section.

Dynatune

Advanced Server supports dynamic tuning of the database server to make the optimal usage of the system resources available on the host machine on which it is installed. The two parameters that control this functionality are located in the [postgresql.conf](#) file. These parameters are:

- `edb_dynatune`
- `edb_dynatune_profile`

`edb_dynatune`

`edb_dynatune` determines how much of the host system's resources are to be used by the database server based upon the host machine's total available resources and the intended usage of the host machine.

When Advanced Server is initially installed, the `edb_dynatune` parameter is set in accordance with the selected usage of the host machine on which it was installed - i.e., development machine, mixed use machine, or dedicated server. For most purposes, there is no need for the database administrator to adjust the various configuration parameters in the [postgresql.conf](#) file in order to improve performance.

You can change the value of the `edb_dynatune` parameter after the initial installation of Advanced Server by editing the [postgresql.conf](#) file. The postmaster must be restarted in order for the new configuration to take effect.

The `edb_dynatune` parameter can be set to any integer value between `0` and `100`, inclusive. A value of `0`, turns off the dynamic tuning feature thereby leaving the database server resource usage totally under the control of the other configuration parameters in the `postgresql.conf` file.

A low non-zero, value (e.g., `1` - `33`) dedicates the least amount of the host machine's resources to the database server. This setting would be used for a development machine where many other applications are being used.

A value in the range of `34` - `66` dedicates a moderate amount of resources to the database server. This setting might be used for a dedicated application server that may have a fixed number of other applications running on the same machine as Advanced Server.

The highest values (e.g., `67` - `100`) dedicate most of the server's resources to the database server. This setting would be used for a host machine that is totally dedicated to running Advanced Server.

Once a value of `edb_dynatune` is selected, database server performance can be further fine-tuned by adjusting the other configuration parameters in the `postgresql.conf` file. Any adjusted setting overrides the corresponding value chosen by `edb_dynatune`. You can change the value of a parameter by un-commenting the configuration parameter, specifying the desired value, and restarting the database server.

`edb_dynatune_profile`

The `edb_dynatune_profile` parameter is used to control tuning aspects based upon the expected workload profile on the database server. This parameter takes effect upon startup of the database server.

The possible values for `edb_dynatune_profile` are:

| Value | Usage |
|------------------------|---|
| <code>oltp</code> | Recommended when the database server is processing heavy online transaction processing workloads. |
| <code>reporting</code> | Recommended for database servers used for heavy data reporting. |
| <code>mixed</code> | Recommended for servers that provide a mix of transaction processing and data reporting. |

EDB Wait States

The *EDB wait states* contrib module contains two main components.

EDB Wait States Background Worker (EWSBW)

When the wait states background worker is registered as one of the shared preload libraries, EWSBW probes each of the running sessions at regular intervals.

For every session it collects information such as the database to which it is connected, the logged in user of the session, the query running in that session, and the wait events on which it is waiting.

This information is saved in a set of files in a user-configurable path and directory folder given by the `edb_wait_states.directory` parameter to be added to the `postgresql.conf` file. The specified path must be a full, absolute path and not a relative path.

The following describes the installation process on a Linux system.

Step 1: EDB wait states is installed with the `edb-asxx-server-edb-modules` RPM package where `xx` is the Advanced Server version number.

Step 2: To launch the worker, it must be registered in the `postgresql.conf` file using the `shared_preload_libraries` parameter, for example:

```
| shared_preload_libraries = '$libdir/edb_wait_states'
```

Step 3: Restart the database server. After a successful restart, the background worker begins collecting data.

Step 4: To review the data, create the following extension:

```
| CREATE EXTENSION edb_wait_states;
```

Step 5: To terminate the EDB wait states worker, remove `$libdir/edb_wait_states` from the `shared_preload_libraries` parameter and restart the database server.

The following describes the installation process on a Windows system.

Step 1: EDB wait states module is installed with the `EDB Modules` installer by invoking StackBuilder Plus utility. Follow the onscreen instructions to complete the installation of the `EDB Modules`.

Step 2: To register the worker, modify the `postgresql.conf` file to include the wait states library in the `shared_preload_libraries` configuration parameter. The parameter value must include:

```
| shared_preload_libraries = '$libdir/edb_wait_states.dll'
```

The EDB wait states installation places the `edb_wait_states.dll` library file in the following path:

```
| C:\Program Files\edb\as13\lib\
```

Step 3: Restart the database server for the changes to take effect. After a successful restart, the background worker gets started and starts collecting the data.

Step 4: To view the data, create the following extension:

```
| CREATE EXTENSION edb_wait_states;
```

The installer places the `edb_wait_states.control` file in the following path:

```
| C:\Program Files\edb\as13\share\extension
```

Terminating the Wait States Worker

To terminate the EDB wait states worker, use the `DROP EXTENSION` command to drop the `edb_wait_states` extension; then modify the `postgresql.conf` file, removing `$libdir/edb_wait_states.dll` from the `shared_preload_libraries` parameter. Restart the database server after modifying the `postgresql.conf` file to apply your changes.

The Wait States Interface

The interface includes the functions listed in the following sections. Each of these functions has common input and output parameters. Those parameters are as follows:

- **start_ts and end_ts (IN).** Together these specify the time interval and the data within which is to be read. If only `start_ts` is specified, the data starting from `start_ts` is output. If only `end_ts` is provided, data up to `end_ts` is output. If none of those are provided, all the data is output. Every function outputs different data. The output of each function will be explained below.
- **query_id (OUT).** Identifies a normalized query. It is internal hash code computed from the query.
- **session_id (OUT).** Identifies a session.
- **ref_start_ts and ref_end_ts (OUT).** Provide the timestamps of a file containing a particular data point. A data point may be a wait event sample record or a query record or a session record.

The syntax of each function is given in the following sections.

Note: The examples shown in the following sections are based on the following three queries executed on four different sessions connected to different databases using different users, simultaneously:

```

SELECT schemaname FROM pg_tables, pg_sleep(15) WHERE schemaname <>
'pg_catalog'; /* ran on 2 sessions */
SELECT tablename FROM pg_tables, pg_sleep(10) WHERE schemaname <>
'pg_catalog';
SELECT tablename, schemaname FROM pg_tables, pg_sleep(10) WHERE schemaname
<> 'pg_catalog';

```

edb_wait_states_data

This function is used to read the data collected by EWSBW.

```

edb_wait_states_data(
    IN start_ts timestamptz default '-infinity'::timestamptz,
    IN end_ts timestamptz default 'infinity'::timestamptz,
    OUT session_id int4,
    OUT <dbname> text,
    OUT <username> text,
    OUT <query> text,
    OUT <query_start_time> timestamptz,
    OUT <sample_time> timestamptz,
    OUT <wait_event_type> text,
    OUT <wait_event> text
)

```

This function can be used to find out the following:

The queries running in the given duration (defined by `start_ts` and `end_ts`) in all the sessions, and the wait events, if any, they were waiting on. For example:

```

SELECT query, session_id, wait_event_type, wait_event
FROM edb_wait_states_data(start_ts, end_ts);

```

The progress of a session within a given duration (that is, the queries run in a session (`session_id = 100000`) and the wait events the queries waited on). For example:

```

SELECT query, wait_event_type, wait_event
FROM edb_wait_states_data(start_ts, end_ts)
WHERE session_id = 100000;

```

The duration for which the samples are available. For example:

```

SELECT min(sample_time), max(sample_time)
FROM edb_wait_states_data();

```

Parameters

In addition to the common parameters described previously, each row of the output gives the following:

`dbname`

| The session's database

`username`

| The session's logged in user

query

The query running in the session

query_start_time

The time when the query started

sample_time

The time when wait event data was collected

wait_event_type

The type of wait event the session (backend) is waiting on

wait_event

The wait event the session (backend) is waiting on

Example

The following is a sample output from the `edb_wait_states_data()` function.

```
edb=# SELECT * FROM edb_wait_states_data();
-[ RECORD 1 ]-----+
-----+-----+
session_id | 4398
dbname     | edb
username   | enterprise
query      | SELECT schemaname FROM pg_tables, pg_sleep($1) WHERE
schemaname <> $2
query_start_time | 17-AUG-18 11:56:05.271962 -04:00
sample_time   | 17-AUG-18 11:56:19.700236 -04:00
wait_event_type | Timeout
wait_event    | PgSleep
-[ RECORD 2 ]-----+
-----+-----+
session_id | 4398
dbname     | edb
username   | enterprise
query      | SELECT schemaname FROM pg_tables, pg_sleep($1) WHERE
schemaname <> $2
query_start_time | 17-AUG-18 11:56:05.271962 -04:00
sample_time   | 17-AUG-18 11:56:18.699938 -04:00
wait_event_type | Timeout
wait_event    | PgSleep
-[ RECORD 3 ]-----+
-----+-----+
session_id | 4398
dbname     | edb
username   | enterprise
query      | SELECT schemaname FROM pg_tables, pg_sleep($1) WHERE
schemaname <> $2
query_start_time | 17-AUG-18 11:56:05.271962 -04:00
sample_time   | 17-AUG-18 11:56:17.700253 -04:00
wait_event_type | Timeout
wait_event    | PgSleep
```

```

.
.
```

edb_wait_states_queries

This function gives information about the queries sampled by EWSBW.

```
edb_wait_states_queries(
    IN start_ts timestamptz default '-infinity'::timestamptz,
    IN end_ts timestamptz default 'infinity'::timestamptz,
    OUT query_id int8,
    OUT <query> text,
    OUT ref_start_ts timestamptz
    OUT ref_end_ts timestamptz
)
```

A new queries file is created periodically and thus, there can be multiple query files generated corresponding to specific intervals.

This function returns all the queries in query files that overlap with the given time interval. A query as shown below, gives all the queries in query files that contained queries sampled between `start_ts` and `end_ts`.

In other words, the function may output queries that did not run in the given interval. To exactly know that the user should use `edb_wait_states_data()`.

```
| SELECT query FROM edb_wait_states_queries(start_ts, end_ts);
```

Parameters

In addition to the common parameters described previously, each row of the output gives the following:

`query`

```
| Normalized query text
```

Example

The following is a sample output from the `edb_wait_states_queries()` function.

```
edb=# SELECT * FROM edb_wait_states_queries();
-[ RECORD 1 ]+-----
-----  

query_id | 4292540138852956818  

query    | SELECT schemaname FROM pg_tables, pg_sleep($1) WHERE  

schemaname <> $2  

ref_start_ts | 17-AUG-18 11:52:38.698793 -04:00  

ref_end_ts  | 18-AUG-18 11:52:38.698793 -04:00  

-[ RECORD 2 ]+-----
-----  

query_id | 3754591102365859187  

query    | SELECT tablename FROM pg_tables, pg_sleep($1) WHERE  

schemaname <> $2  

ref_start_ts | 17-AUG-18 11:52:38.698793 -04:00  

ref_end_ts  | 18-AUG-18 11:52:38.698793 -04:00  

-[ RECORD 3 ]+-----
```

```
-----
query_id | 349089096300352897
query    | SELECT tablename, schemaname FROM pg_tables, pg_sleep($1)
WHERE schemaname <> $2
ref_start_ts | 17-AUG-18 11:52:38.698793 -04:00
ref_end_ts  | 18-AUG-18 11:52:38.698793 -04:00
```

edb_wait_states_sessions

This function gives information about the sessions sampled by EWSBW.

```
edb_wait_states_sessions(
    IN start_ts timestamptz default '-infinity'::timestamptz,
    IN end_ts timestamptz default 'infinity'::timestamptz,
    OUT session_id int4,
    OUT <dbname> text,
    OUT <username> text,
    OUT ref_start_ts timestamptz
    OUT ref_end_ts timestamptz
)
```

This function can be used to identify the databases that were connected and/or which users started those sessions. For example:

```
SELECT dbname, username, session_id
  FROM edb_wait_states_sessions();
```

Similar to [edb_wait_states_queries\(\)](#), this function outputs all the sessions logged in session files that contain sessions sampled within the given interval and not necessarily only the sessions sampled within the given interval. To identify that one should use [edb_wait_states_data\(\)](#).

Parameters

In addition to the common parameters described previously, each row of the output gives the following:

dbname

The database to which the session is connected

username

Login user of the session

Example

The following is a sample output from the [edb_wait_states_sessions\(\)](#) function.

```
edb=# SELECT * FROM edb_wait_states_sessions();
-[ RECORD 1 ]+-----
session_id | 4340
dbname     | edb
username   | enterprisedb
ref_start_ts | 17-AUG-18 11:52:38.698793 -04:00
ref_end_ts  | 18-AUG-18 11:52:38.698793 -04:00
-[ RECORD 2 ]+-----
session_id | 4398
```

```

dbname |edb
username |enterprisedb
ref_start_ts |17-AUG-18 11:52:38.698793 -04:00
ref_end_ts |18-AUG-18 11:52:38.698793 -04:00
-[ RECORD 3 ]+-----
session_id |4410
dbname |db1
username |user1
ref_start_ts |17-AUG-18 11:52:38.698793 -04:00
ref_end_ts |18-AUG-18 11:52:38.698793 -04:00
-[ RECORD 4 ]+-----
session_id |4422
dbname |db2
username |user2
ref_start_ts |17-AUG-18 11:52:38.698793 -04:00
ref_end_ts |18-AUG-18 11:52:38.698793 -04:00

```

edb_wait_states_samples

This function gives information about wait events sampled by EWSBW.

```

edb_wait_states_samples(
    IN start_ts timestamptz default '-infinity'::timestamptz,
    IN end_ts timestamptz default 'infinity'::timestamptz,
    OUT query_id int8,
    OUT session_id int4,
    OUT <query_start_time> timestamptz,
    OUT <sample_time> timestamptz,
    OUT <wait_event_type> text,
    OUT <wait_event> text
)

```

Usually, a user would not be required to call this function directly.

Parameters

In addition to the common parameters described previously, each row of the output gives the following:

query_start_time

The time when the query started in this session

sample_time

The time when wait event data was collected

wait_event_type

The type of wait event on which the session is waiting

wait_event

The wait event on which the session (backend) is waiting

Example

The following is a sample output from the `edb_wait_states_samples()` function.

```
edb=# SELECT * FROM edb_wait_states_samples();
-[ RECORD 1 ]-----+
query_id      | 4292540138852956818
session_id    | 4340
query_start_time | 17-AUG-18 11:56:00.39421 -04:00
sample_time    | 17-AUG-18 11:56:00.699934 -04:00
wait_event_type | Timeout
wait_event     | PgSleep
-[ RECORD 2 ]-----+
query_id      | 4292540138852956818
session_id    | 4340
query_start_time | 17-AUG-18 11:56:00.39421 -04:00
sample_time    | 17-AUG-18 11:56:01.699003 -04:00
wait_event_type | Timeout
wait_event     | PgSleep
-[ RECORD 3 ]-----+
query_id      | 4292540138852956818
session_id    | 4340
query_start_time | 17-AUG-18 11:56:00.39421 -04:00
sample_time    | 17-AUG-18 11:56:02.70001 -04:00
wait_event_type | Timeout
wait_event     | PgSleep
-[ RECORD 4 ]-----+
query_id      | 4292540138852956818
session_id    | 4340
query_start_time | 17-AUG-18 11:56:00.39421 -04:00
sample_time    | 17-AUG-18 11:56:03.700081 -04:00
wait_event_type | Timeout
wait_event     | PgSleep
.
.
.
```

`edb_wait_states_purge`

The function deletes all the sampled data files (queries, sessions and wait event samples) that were created after `start_ts` and aged (rotated) before `end_ts`.

```
edb_wait_states_purge(
  IN start_ts timestamptz default '-infinity'::timestamptz,
  IN end_ts timestamptz default 'infinity'::timestamptz
)
```

Usually a user does not need to run this function. The backend should purge those according to the retention age, but in case, that doesn't happen for some reason, this function may be used.

In order to know the duration for which the samples have been retained, use `edb_wait_states_data()` as explained in the previous examples of that function.

Example

The `$PGDATA/edb_wait_states` directory before running `edb_wait_states_purge()`:

```
[root@localhost data]# pwd
/var/lib/edb/as13/data
[root@localhost data]# ls -l edb_wait_states
total 12
-rw----- 1 enterprise... 253 Aug 17 11:56
edb_ws_queries_587836358698793_587922758698793
-rw----- 1 enterprise... 1600 Aug 17 11:56
edb_ws_samples_587836358698793_587839958698793
-rw----- 1 enterprise... 94 Aug 17 11:56
edb_ws_sessions_587836358698793_587922758698793
```

The `$PGDATA/edb_wait_states` directory after running `edb_wait_states_purge()`:

```
edb=# SELECT * FROM edb_wait_states_purge();
edb_wait_states_purge
-----
```

(1 row)

```
[root@localhost data]# pwd
/var/lib/edb/as13/data
[root@localhost data]# ls -l edb_wait_states
total 0
```

11.14 EDB Clone Schema

EDB Clone Schema is an extension module for Advanced Server that allows you to copy a schema and its database objects from a local or remote database (the source database) to a receiving database (the target database).

The source and target databases can be the same physical database, or different databases within the same database cluster, or separate databases running under different database clusters on separate database server hosts.

Use the following functions with EDB Clone Schema:

- **localcopschema.** This function makes a copy of a schema and its database objects from a source database back into the same database (the target), but with a different schema name than the original. Use this function when the original source schema and the resulting copy are to reside within the same database. See `localcopschema` for information on the `localcopschema` function.
- **localcopschema_nb.** This function performs the same purpose as `localcopschema`, but as a background job, thus freeing up the terminal from which the function was initiated. This is referred to as a *non-blocking* function. See `localcopschema_nb` for information on the `localcopschema_nb` function.
- **remotecloneschema.** This function makes a copy of a schema and its database objects from a source database to a different target database. Use this function when the original source schema and the resulting copy are to reside in two, separate databases. The separate databases can reside in the same, or in different Advanced Server database clusters. See `remotecloneschema` for information on the `remotecloneschema` function.
- **remotecloneschema_nb.** This function performs the same purpose as `remotecloneschema`, but as a background job, thus freeing up the terminal from which the function was initiated. This is referred to as a *non-blocking* function. See `remotecloneschema_nb` for information on the `remotecloneschema_nb` function.
- **process_status_from_log.** This function displays the status of the cloning functions. The information is obtained from a log file that must be specified when a cloning function is invoked. See

[process_status_from_log](#) for information on the `process_status_from_log` function.

- **remove_log_file_and_job.** This function deletes the log file created by a cloning function. This function can also be used to delete a job created by the non-blocking form of the function. See [remove_log_file_and_job](#) for information on the `remove_log_file_and_job` function.

The database objects that can be cloned from one schema to another are the following:

- Data types
- Tables including partitioned tables, excluding foreign tables
- Indexes
- Constraints
- Sequences
- View definitions
- Materialized views
- Private synonyms
- Table triggers, but excluding event triggers
- Rules
- Functions
- Procedures
- Packages
- Comments for all supported object types
- Access control lists (ACLs) for all supported object types

The following database objects cannot be cloned:

- Large objects (Postgres `LOBs` and `BFILEs`)
- Logical replication attributes for a table
- Database links
- Foreign data wrappers
- Foreign tables
- Event triggers
- Extensions (For cloning objects that rely on extensions, see the third bullet point in the following limitations list.)
- Row level security
- Policies
- Operator class

In addition, the following limitations apply:

- EDB Clone Schema is supported on Advanced Server only when a dialect of `Compatible with Oracle` is specified on the Advanced Server `Dialect` dialog during installation, or when the `--redwood-like` keywords are included during a text mode installation or cluster initialization.
- The source code within functions, procedures, triggers, packages, etc., are not modified after being copied to the target schema. If such programs contain coded references to objects with schema names, the programs may fail upon invocation in the target schema if such schema names are no longer consistent within the target schema.
- Cross schema object dependencies are not resolved. If an object in the target schema depends upon an object in another schema, this dependency is not resolved by the cloning functions.
- For remote cloning, if an object in the source schema is dependent upon an extension, then this extension must be created in the public schema of the remote database before invoking the remote cloning function.
- At most, 16 copy jobs can run in parallel to clone schemas, whereas each job can have at most 16 worker processes to copy table data in parallel.
- Queries being run by background workers cannot be cancelled.

The following section describes how to set up EDB Clone Schema on the databases.

Setup Process

Several extensions along with the PL/Perl language must be installed on any database to be used as the source or target database by an EDB Clone Schema function.

In addition, some configuration parameters in the `postgresql.conf` file of the database servers may benefit from some modification.

The following is the setup instructions for these requirements.

Installing Extensions and PL/Perl

The following describes the steps to install the required extensions and the PL/Perl language.

These steps must be performed on any database to be used as the source or target database by an EDB Clone Schema function.

Step 1: The following extensions must be installed on the database:

- `postgres_fdw`
- `dblink`
- `adminpack`
- `pgagent`

Ensure that pgAgent is installed before creating the `pgagent` extension. On Linux, you can use the `edb-asxx-pgagent` RPM package where `xx` is the Advanced Server version number to install pgAgent. On Windows, use StackBuilder Plus to download and install pgAgent.

The previously listed extensions can be installed by the following commands if they do not already exist:

```
CREATE EXTENSION postgres_fdw SCHEMA public;
CREATE EXTENSION dblink SCHEMA public;
CREATE EXTENSION adminpack;
CREATE EXTENSION pgagent;
```

For more information about using the `CREATE EXTENSION` command, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/sql-createextension.html>

Step 2: Modify the `postgresql.conf` file.

Modify the `postgresql.conf` file by adding `$libdir/parallel_clone` to the `shared_preload_libraries` configuration parameter as shown by the following example:

```
shared_preload_libraries = '$libdir/dbms_pipe,$libdir/dbms_aq,$libdir/
parallel_clone'
```

Step 3: The Perl Procedural Language (PL/Perl) must be installed on the database and the `CREATE TRUSTED LANGUAGE plperl` command must be run. For Linux, install PL/Perl using the `edb-asxx-server-plperl` RPM package where `xx` is the Advanced Server version number. For Windows, use the EDB Postgres Language Pack. For information on EDB Language Pack, see the *EDB Postgres Language Pack Guide* available at:

<https://www.enterprisedb.com/edb-docs>

Step 4: Connect to the database as a superuser and run the following command:

```
CREATE TRUSTED LANGUAGE plperl;
```

For more information about using the `CREATE LANGUAGE` command, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/sql-createlanguage.html>

Setting Configuration Parameters

The following sections describe configuration parameters that may need to be altered in the `postgresql.conf` file.

Performance Configuration Parameters

You may need to tune the system for copying a large schema as part of one transaction.

Tuning of configuration parameters is for the source database server referenced in a cloning function.

The configuration parameters in the `postgresql.conf` file that may need to be tuned include the following:

- **work_mem**. Specifies the amount of memory to be used by internal sort operations and hash tables before writing to temporary disk files.
- **maintenance work mem**. Specifies the maximum amount of memory to be used by maintenance operations, such as `VACUUM`, `CREATE INDEX`, and `ALTER TABLE ADD FOREIGN KEY`.
- **max_worker_processes**. Sets the maximum number of background processes that the system can support.
- **checkpoint_timeout**. Maximum time between automatic WAL checkpoints, in seconds.
- **checkpoint_completion_target**. Specifies the target of checkpoint completion, as a fraction of total time between checkpoints.
- **checkpoint_flush_after**. Whenever more than `checkpoint_flush_after` bytes have been written while performing a checkpoint, attempt to force the OS to issue these writes to the underlying storage.
- **max_wal_size**. Maximum size to let the WAL grow to between automatic WAL checkpoints.
- **max_locks_per_transaction**. This parameter controls the average number of object locks allocated for each transaction; individual transactions can lock more objects as long as the locks of all transactions fit in the lock table.

For information about the configuration parameters, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/runtime-config.html>

Status Logging

Status logging by the cloning functions creates log files in the directory specified by the `log_directory` parameter in the `postgresql.conf` file for the database server to which you are connected when invoking the cloning function.

The default location is `PGDATA/log` as shown by the following:

```
#log_directory = 'log'      # directory where log files are written,
                           # can be absolute or relative to PGDATA
```

This directory must exist prior to running a cloning function.

The name of the log file is determined by what you specify in the parameter list when invoking the cloning function.

To display the status from a log file, use the `process_status_from_log` function.

To delete a log file, use the `remove_log_file_and_job` function, or simply navigate to the log directory and delete it manually.

Installing EDB Clone Schema

The following are the directions for installing EDB Clone Schema.

These steps must be performed on any database to be used as the source or target database by an EDB Clone Schema function.

Step 1: If you had previously installed an older version of the `edb_cloneschema` extension, then you must run the following command:

```
DROP EXTENSION parallel_clone CASCADE;
```

This command also drops the `edb_cloneschema` extension.

Step 2: Install the extensions using the following commands:

```
CREATE EXTENSION parallel_clone SCHEMA public;
```

```
CREATE EXTENSION edb_cloneschema;
```

Make sure you create the `parallel_clone` extension before creating the `edb_cloneschema` extension.

Creating the Foreign Servers and User Mappings

When using one of the local cloning functions, `localcopschema` or `localcopschema_nb`, one of the required parameters includes a single, foreign server for identifying the database server along with its database that is the source and the receiver of the cloned schema.

When using one of the remote cloning functions, `remotecopschema` or `remotecopschema_nb`, two of the required parameters include two foreign servers. The foreign server specified as the first parameter identifies the source database server along with its database that is the provider of the cloned schema. The foreign server specified as the second parameter identifies the target database server along with its database that is the receiver of the cloned schema.

For each foreign server, a user mapping must be created. When a selected database superuser invokes a cloning function, that database superuser who invokes the function must have been mapped to a database user name and password that has access to the foreign server that is specified as a parameter in the cloning function.

For general information about foreign data, foreign servers, and user mappings, see the PostgreSQL core documentation at:

```
https://www.postgresql.org/docs/current/static/ddl-foreign-data.html
```

The following two sections describe how these foreign servers and user mappings are defined.

Foreign Server and User Mapping for Local Cloning Functions

For the `localcopschema` and `localcopschema_nb` functions, the source and target schemas are both within the same database of the same database server. Thus, only one foreign server must be defined and specified for these functions. This foreign server is also referred to as the *local server*.

This server is referred to as the local server because this server is the one to which you must be connected when invoking the `localcopschema` or `localcopschema_nb` function.

The user mapping defines the connection and authentication information for the foreign server.

This foreign server and user mapping must be created within the database of the local server in which the cloning is to occur.

The database user for whom the user mapping is defined must be a superuser and the user connected to the local server when invoking an EDB Clone Schema function.

The following example creates the foreign server for the database containing the schema to be cloned, and to

receive the cloned schema as well.

```
CREATE SERVER local_server FOREIGN DATA WRAPPER postgres_fdw
OPTIONS(
  host 'localhost',
  port '5444',
  dbname 'edb'
);
```

For more information about using the `CREATE SERVER` command, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/sql-createserver.html>

The user mapping for this server is the following:

```
CREATE USER MAPPING FOR enterprisedb SERVER local_server
OPTIONS (
  user 'enterprisedb',
  password 'password'
);
```

For more information about using the `CREATE USER MAPPING` command, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/sql-createusermapping.html>

The following `psql` commands show the foreign server and user mapping:

```
edb=# \des+
List of foreign servers
-[ RECORD 1 ]-----+
Name      | local_server
Owner     | enterprisedb
Foreign-data wrapper | postgres_fdw
Access privileges  |
Type      |
Version   |
FDW options | (host 'localhost', port '5444', dbname 'edb')
Description |
```



```
edb=# \deu+
      List of user mappings
  Server | User name |          FDW options
-----+-----+
local_server | enterprisedb | ("user" 'enterprisedb', password 'password')
(1 row)
```

When database superuser `enterprisedb` invokes a cloning function, the database user `enterprisedb` with its password is used to connect to `local_server` on the `localhost` with port `5444` to database `edb`.

In this case, the mapped database user, `enterprisedb`, and the database user, `enterprisedb`, used to connect to the local `edb` database happen to be the same, identical database user, but that is not an absolute requirement.

For specific usage of these foreign server and user mapping examples, see the example given in `localcopschema`.

Foreign Server and User Mapping for Remote Cloning Functions

For the `remotecopyschema` and `remotecopyschema_nb` functions, the source and target schemas are in different databases of either the same or different database servers. Thus, two foreign servers must be defined and specified for these functions.

The foreign server defining the originating database server and its database containing the source schema to be cloned is referred to as the *source server* or the *remote server*.

The foreign server defining the database server and its database to receive the schema to be cloned is referred to as the *target server* or the *local server*.

The target server is also referred to as the local server because this server is the one to which you must be connected when invoking the `remotecopyschema` or `remotecopyschema_nb` function.

The user mappings define the connection and authentication information for the foreign servers.

All of these foreign servers and user mappings must be created within the target database of the target/local server.

The database user for whom the user mappings are defined must be a superuser and the user connected to the local server when invoking an EDB Clone Schema function.

The following example creates the foreign server for the local, target database that is to receive the cloned schema.

```
CREATE SERVER tgt_server FOREIGN DATA WRAPPER postgres_fdw
OPTIONS(
  host 'localhost',
  port '5444',
  dbname 'tgtdb'
);
```

The user mapping for this server is the following:

```
CREATE USER MAPPING FOR enterprisedb SERVER tgt_server
OPTIONS (
  user 'tgtuser',
  password 'tgtpassword'
);
```

The following example creates the foreign server for the remote, source database that is to be the source for the cloned schema.

```
CREATE SERVER src_server FOREIGN DATA WRAPPER postgres_fdw
OPTIONS(
  host '192.168.2.28',
  port '5444',
  dbname 'srcdb'
);
```

The user mapping for this server is the following:

```
CREATE USER MAPPING FOR enterprisedb SERVER src_server
OPTIONS (
  user 'srcuser',
  password 'srcpassword'
);
```

The following `psql` commands show the foreign servers and user mappings:

```
tgtdb=# \des+
List of foreign servers
-[ RECORD 1 ]+-
Name      | src_server
Owner     | tgtuser
Foreign-data wrapper | postgres_fdw
Access privileges |
Type      |
Version   |
FDW options | (host '192.168.2.28', port '5444', dbname 'srcdb')
Description |
-[ RECORD 2 ]+-
Name      | tgt_server
Owner     | tgtuser
Foreign-data wrapper | postgres_fdw
Access privileges |
Type      |
Version   |
FDW options | (host 'localhost', port '5444', dbname 'tgtdb')
Description |
```

```
tgtdb=# \deu+
          List of user mappings
  Server | User name |           FDW options
```

| Server | User name | FDW options |
|------------|--------------|--|
| src_server | enterprisedb | ("user" 'srcuser', password 'srcpassword') |
| tgt_server | enterprisedb | ("user" 'tgtuser', password 'tgtpassword') |

(2 rows)

When database superuser `enterprisedb` invokes a cloning function, the database user `tgtuser` with password `tgtpassword` is used to connect to `tgt_server` on the `localhost` with port `5444` to database `tgtdb`.

In addition, database user `srcuser` with password `srcpassword` connects to `src_server` on host `192.168.2.28` with port `5444` to database `srcdb`.

Note

Be sure the `pg_hba.conf` file of the database server running the source database `srcdb` has an appropriate entry permitting connection from the target server location (address `192.168.2.27` in the following example) connecting with the database user `srcuser` that was included in the user mapping for the foreign server `src_server` defining the source server and database.

```
### TYPE DATABASE USER ADDRESS METHOD
### "local" is for Unix domain socket connections only
local  all      all            md5
### IPv4 local connections:
host   srcdb    srcuser     192.168.2.27/32    md5
```

For specific usage of these foreign server and user mapping examples, see the example given [remotecopyschema](#).

EDB Clone Schema Functions

The EDB Clone Schema functions are created in the `edb_util` schema when the `parallel_clone` and

`edb_cloneschema` extensions are installed.

Verify the following conditions before using an EDB Clone Schema function:

- You are connected to the target or local database as the database superuser defined in the `CREATE USER MAPPING` command for the foreign server of the target or local database.
- The `edb_util` schema is in the search path, or the cloning function is to be invoked with the `edb_util` prefix.
- The target schema does not exist in the target database.
- When using the remote copy functions, if the `on_tblspace` parameter is to be set to `true`, then the target database cluster contains all tablespaces that are referenced by objects in the source schema, otherwise creation of the DDL statements for those database objects will fail in the target schema. This causes a failure of the cloning process.
- When using the remote copy functions, if the `copy_acls` parameter is to be set to `true`, then all roles that have `GRANT` privileges on objects in the source schema exist in the target database cluster, otherwise granting of privileges to those roles will fail in the target schema. This causes a failure of the cloning process.
- pgAgent is running against the target database if the non-blocking form of the function is to be used.

For information about pgAgent, see the pgAdmin documentation available at:

<https://www.pgadmin.org/docs/pgadmin4/dev/pgagent.html>

Note that pgAgent is provided as a component with Advanced Server.

localcopschema

The `localcopschema` function copies a schema and its database objects within a local database specified within the `source_fdw` foreign server from the source schema to the specified target schema within the same database.

```
localcopschema(
    <source_fdw> TEXT,
    <source_schema> TEXT,
    <target_schema> TEXT,
    <log_filename> TEXT
    [, <on_tblspace> BOOLEAN
    [, <verbose_on> BOOLEAN
    [, <copy_acls> BOOLEAN
    [, <worker_count> INTEGER ]]]]
)
```

A `BOOLEAN` value is returned by the function. If the function succeeds, then `true` is returned. If the function fails, then `false` is returned.

The `source_fdw`, `source_schema`, `target_schema`, and `log_filename` are required parameters while all other parameters are optional.

Parameters

`source_fdw`

Name of the foreign server managed by the `postgres_fdw` foreign data wrapper from which database objects are to be cloned.

`source_schema`

Name of the schema from which database objects are to be cloned.

`target_schema`

Name of the schema into which database objects are to be cloned from the source schema.

log_filename

Name of the log file in which information from the function is recorded. The log file is created under the directory specified by the `log_directory` configuration parameter in the `postgresql.conf` file.

on_tblspace

`BOOLEAN` value to specify whether or not database objects are to be created within their tablespaces. If `false` is specified, then the `TABLESPACE` clause is not included in the applicable `CREATE` DDL statement when added to the target schema. If `true` is specified, then the `TABLESPACE` clause is included in the `CREATE` DDL statement when added to the target schema. If the `on_tblspace` parameter is omitted, the default value is `false`.

verbose_on

`BOOLEAN` value to specify whether or not the DDLs are to be printed in `log_filename` when creating objects in the target schema. If `false` is specified, then DDLs are not printed. If `true` is specified, then DDLs are printed. If omitted, the default value is `false`.

copy_acls

`BOOLEAN` value to specify whether or not the access control list (ACL) is to be included while creating objects in the target schema. The access control list is the set of `GRANT` privilege statements. If `false` is specified, then the access control list is not included for the target schema. If `true` is specified, then the access control list is included for the target schema. If the `copy_acls` parameter is omitted, the default value is `false`.

worker_count

Number of background workers to perform the clone in parallel. If omitted, the default value is `1`.

Example

The following example shows the cloning of schema `edb` containing a set of database objects to target schema `edbcopy`, both within database `edb` as defined by `local_server`.

The example is for the following environment:

- Host on which the database server is running: `localhost`
- Port of the database server: `5444`
- Database source/target of the clone: `edb`
- Foreign server (`local_server`) and user mapping with the information of the preceding bullet points
- Source schema: `edb`
- Target schema: `edbcopy`
- Database superuser to invoke `localcopschema: enterprisedb`

Before invoking the function, the connection is made by database user `enterprisedb` to database `edb`.

```
edb=# SET search_path TO "$user",public,edb_util;
SET
edb=# SHOW search_path;
      search_path
-----
"$user", public, edb_util
(1 row)

edb=# SELECT localcopschema
('local_server','edb','edbcopy','clone_edb_edbcopy');
localcopschema
-----
t
```

(1 row)

The following displays the logging status using the `process_status_from_log` function:

```
edb=# SELECT process_status_from_log('clone_edb_edbcopy');
          process_status_from_log
-----
```

```
-----  
(FINISH,"2017-06-29 11:07:36.830783-04",3855,INFO,"STAGE:  
FINAL","successfully cloned schema")
```

(1 row)

After the clone has completed, the following shows some of the database objects copied to the `edbcopy` schema:

```
edb=# SET search_path TO edbcopy;
```

```
SET
```

```
edb=# \dt+
```

```
      List of relations
```

| Schema | Name | Type | Owner | Size | Description |
|--------|------|------|-------|------|-------------|
|--------|------|------|-------|------|-------------|

```
-----+-----+-----+-----+-----+-----  
edbcopy | dept | table | enterprisedb | 8192 bytes |  
edbcopy | emp | table | enterprisedb | 8192 bytes |  
edbcopy | jobhist | table | enterprisedb | 8192 bytes |  
(3 rows)
```

```
edb=# \dv
```

```
      List of relations
```

| Schema | Name | Type | Owner |
|--------|------|------|-------|
|--------|------|------|-------|

```
-----+-----+-----+-----  
edbcopy | salesemp | view | enterprisedb  
(1 row)
```

```
edb=# \di
```

```
      List of relations
```

| Schema | Name | Type | Owner | Table |
|--------|------|------|-------|-------|
|--------|------|------|-------|-------|

```
-----+-----+-----+-----+-----  
edbcopy | dept_dname_uq | index | enterprisedb | dept  
edbcopy | dept_pk | index | enterprisedb | dept  
edbcopy | emp_pk | index | enterprisedb | emp  
edbcopy | jobhist_pk | index | enterprisedb | jobhist  
(4 rows)
```

```
edb=# \ds
```

```
      List of relations
```

| Schema | Name | Type | Owner |
|--------|------|------|-------|
|--------|------|------|-------|

```
-----+-----+-----+-----  
edbcopy | next_empno | sequence | enterprisedb  
(1 row)
```

```
edb=# SELECT DISTINCT schema_name, name, type FROM user_source WHERE
schema_name = 'EDBCOPY' ORDER BY type, name;
```

```
schema_name | name | type
```

```
-----+-----+-----  
EDBCOPY | EMP_COMP | FUNCTION  
EDBCOPY | HIRE_CLERK | FUNCTION  
EDBCOPY | HIRE SALESMAN | FUNCTION
```

```

EDBCOPY | NEW_EMPNO           | FUNCTION
EDBCOPY | EMP_ADMIN            | PACKAGE
EDBCOPY | EMP_ADMIN            | PACKAGE BODY
EDBCOPY | EMP_QUERY             | PROCEDURE
EDBCOPY | EMP_QUERY_CALLER      | PROCEDURE
EDBCOPY | LIST_EMP              | PROCEDURE
EDBCOPY | SELECT_EMP             | PROCEDURE
EDBCOPY | EMP_SAL_TRIG          | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_a_19991" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_a_19992" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_a_19999" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_a_20000" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_a_20004" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_a_20005" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_c_19993" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_c_19994" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_c_20001" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_c_20002" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_c_20006" | TRIGGER
EDBCOPY | "RI_ConstraintTrigger_c_20007" | TRIGGER
EDBCOPY | USER_AUDIT_TRIG        | TRIGGER
(24 rows)

```

localcopyschema_nb

The `localcopyschema_nb` function copies a schema and its database objects within a local database specified within the `source_fdw` foreign server from the source schema to the specified target schema within the same database, but in a non-blocking manner as a job submitted to pgAgent.

```

localcopyschema_nb(
  <source_fdw> TEXT,
  <source> TEXT,
  <target> TEXT,
  <log_filename> TEXT
 [, <on_tblspace> BOOLEAN
 [, <verbose_on> BOOLEAN
 [, <copy_acls> BOOLEAN
 [, <worker_count> INTEGER ]]]]
)

```

An `INTEGER` value job ID is returned by the function for the job submitted to pgAgent. If the function fails, then null is returned.

The `source_fdw`, `source`, `target`, and `log_filename` are required parameters while all other parameters are optional.

After completion of the pgAgent job, remove the job with the `remove_log_file_and_job` function.

Parameters

`source_fdw`

Name of the foreign server managed by the `postgres_fdw` foreign data wrapper from which database objects are to be cloned.

`source`

Name of the schema from which database objects are to be cloned.

target

Name of the schema into which database objects are to be cloned from the source schema.

log_filename

Name of the log file in which information from the function is recorded. The log file is created under the directory specified by the `log_directory` configuration parameter in the `postgresql.conf` file.

on_tblspace

`BOOLEAN` value to specify whether or not database objects are to be created within their tablespaces. If `false` is specified, then the `TABLESPACE` clause is not included in the applicable `CREATE` DDL statement when added to the target schema. If `true` is specified, then the `TABLESPACE` clause is included in the `CREATE` DDL statement when added to the target schema. If the `on_tblspace` parameter is omitted, the default value is `false`.

verbose_on

`BOOLEAN` value to specify whether or not the DDLs are to be printed in `log_filename` when creating objects in the target schema. If `false` is specified, then DDLs are not printed. If `true` is specified, then DDLs are printed. If omitted, the default value is `false`.

copy_acls

`BOOLEAN` value to specify whether or not the access control list (ACL) is to be included while creating objects in the target schema. The access control list is the set of `GRANT` privilege statements. If `false` is specified, then the access control list is not included for the target schema. If `true` is specified, then the access control list is included for the target schema. If the `copy_acls` parameter is omitted, the default value is `false`.

worker_count

Number of background workers to perform the clone in parallel. If omitted, the default value is `1`.

Example

The same cloning operation is performed as the example in `localcopschema`, but using the non-blocking function `localcopschema_nb`.

The following command can be used to observe if pgAgent is running on the appropriate local database:

```
[root@localhost ~]# ps -ef | grep pgagent
root    4518     1  0 11:35 pts/1  00:00:00 pgagent -s /tmp/
pgagent_edb_log hostaddr=127.0.0.1 port=5444 dbname=edb user=enterprisedb
password=password
root    4525  4399  0 11:35 pts/1  00:00:00 grep --color=auto pgagent
```

If pgAgent is not running, it can be started as shown by the following. The `pgagent` program file is located in the `bin` subdirectory of the Advanced Server installation directory.

```
[root@localhost bin]# ./pgagent -l 2 -s /tmp/pgagent_edb_log
hostaddr=127.0.0.1 port=5444 dbname=edb user=enterprisedb password=password
```

Note

The `pgagent -l 2` option starts pgAgent in `DEBUG` mode, which logs continuous debugging information into the log file specified with the `-s` option. Use a lower value for the `-l` option, or omit it entirely to record less information.

The `localcopschema_nb` function returns the job ID shown as `4` in the example.

```
edb=# SELECTedb_util.localcopschema_nb
('local_server','edb','edbcopy','clone_edb_edbcopy');
localcopschema_nb
-----
4
(1 row)
```

The following displays the job status:

```
edb=# SELECTedb_util.process_status_from_log('clone_edb_edbcopy');
process_status_from_log
-----
(FINISH,"29-JUN-17 11:39:11.620093 -04:00",4618,INFO,"STAGE:
FINAL","successfully cloned schema")
(1 row)
```

The following removes the pgAgent job:

```
edb=# SELECTedb_util.remove_log_file_and_job (4);
remove_log_file_and_job
-----
t
(1 row)
```

remotecopschema

The `remotecopschema` function copies a schema and its database objects from a source schema in the remote source database specified within the `source_fdw` foreign server to a target schema in the local target database specified within the `target_fdw` foreign server.

```
remotecopschema(
  <source_fdw> TEXT,
  <target_fdw> TEXT,
  <source_schema> TEXT,
  <target_schema> TEXT,
  <log_filename> TEXT
 [, <on_tblspace> BOOLEAN
 [, <verbose_on> BOOLEAN
 [, <copy_acls> BOOLEAN
 [, <worker_count> INTEGER ]]]]
)
```

A `BOOLEAN` value is returned by the function. If the function succeeds, then `true` is returned. If the function fails, then `false` is returned.

The `source_fdw`, `target_fdw`, `source_schema`, `target_schema`, and `log_filename` are required parameters while all other parameters are optional.

Parameters

`source_fdw`

Name of the foreign server managed by the `postgres_fdw` foreign data wrapper from which database objects are to be cloned.

`target_fdw`

Name of the foreign server managed by the `postgres_fdw` foreign data wrapper to which database objects are to be cloned.

`source_schema`

Name of the schema from which database objects are to be cloned.

`target_schema`

Name of the schema into which database objects are to be cloned from the source schema.

`log_filename`

Name of the log file in which information from the function is recorded. The log file is created under the directory specified by the `log_directory` configuration parameter in the `postgresql.conf` file.

`on_tblspace`

`BOOLEAN` value to specify whether or not database objects are to be created within their tablespaces. If `false` is specified, then the `TABLESPACE` clause is not included in the applicable `CREATE` DDL statement when added to the target schema. If `true` is specified, then the `TABLESPACE` clause is included in the `CREATE` DDL statement when added to the target schema. If the `on_tblspace` parameter is omitted, the default value is `false`.

Note

If `true` is specified and a database object has a `TABLESPACE` clause, but that tablespace does not exist in the target database cluster, then the cloning function fails.

`verbose_on`

`BOOLEAN` value to specify whether or not the DDLs are to be printed in `log_filename` when creating objects in the target schema. If `false` is specified, then DDLs are not printed. If `true` is specified, then DDLs are printed. If omitted, the default value is `false`.

`copy_acls`

`BOOLEAN` value to specify whether or not the access control list (ACL) is to be included while creating objects in the target schema. The access control list is the set of `GRANT` privilege statements. If `false` is specified, then the access control list is not included for the target schema. If `true` is specified, then the access control list is included for the target schema. If the `copy_acls` parameter is omitted, the default value is `false`.

Note

If `true` is specified and a role with `GRANT` privilege does not exist in the target database cluster, then the cloning function fails.

`worker_count`

Number of background workers to perform the clone in parallel. If omitted, the default value is `1`.

Example

The following example shows the cloning of schema `srcschema` within database `srcdb` as defined by `src_server` to target schema `tgtschema` within database `tgtdb` as defined by `tgt_server`.

The source server environment:

- Host on which the source database server is running: 192.168.2.28
- Port of the source database server: 5444
- Database source of the clone: srctdb
- Foreign server (src_server) and user mapping with the information of the preceding bullet points
- Source schema: srccschema

The target server environment:

- Host on which the target database server is running: localhost
- Port of the target database server: 5444
- Database target of the clone: tgttgt
- Foreign server (tgt_server) and user mapping with the information of the preceding bullet points
- Target schema: tgtschema
- Database superuser to invoke remotecopyschema: enterprisedb

Before invoking the function, the connection is made by database user enterprisedb to database tgttgt. A worker_count of 4 is specified for this function.

```
tgttgt=# SELECTedb_util.remotecopyschema
('src_server','tgt_server','srccschema','tgtschema','clone_rmt_src_tgt',worker
_count >= 4);
remotecopyschema
-----
t
(1 row)
```

The following displays the status from the log file during various points in the cloning process:

```
tgttgt=# SELECTedb_util.process_status_from_log('clone_rmt_src_tgt');
process_status_from_log
-----
-- 
(RUNNING,"28-JUN-17 current:18:05.299953 -04:00",4021,INFO,"STAGE: DATA-
COPY","[0][0] successfully copied data in [tgtschema.pgbench_tellers]
")
(1 row)
```

```
tgttgt=# SELECTedb_util.process_status_from_log('clone_rmt_src_tgt');
process_status_from_log
-----
-- 
(RUNNING,"28-JUN-17 current:18:06.634364 -04:00",4022,INFO,"STAGE: DATA-
COPY","[0][1] successfully copied data in [tgtschema.pgbench_history]
")
(1 row)
```

```
tgttgt=# SELECTedb_util.process_status_from_log('clone_rmt_src_tgt');
process_status_from_log
-----
--
```

```
(RUNNING,"28-JUN-17 current:18:10.550393 -04:00",4039,INFO,"STAGE: POST-  
DATA","CREATE PRIMARY KEY CONSTRAINT pgbench_tellers_pkey successful"  
)  
(1 row)
```

```
tgtmdb=# SELECTedb_util.process_status_from_log('clone_rmt_src_tgt');  
process_status_from_log
```

```
-----  
(FINISH,"28-JUN-17 current:18:12.019627 -04:00",4039,INFO,"STAGE:  
FINAL","successfully clone schema into tgtschema")  
(1 row)
```

The following shows the cloned tables:

```
tgtmdb=# \dt+  
          List of relations  
 Schema |     Name      | Type | Owner | Size | Description  
-----+-----+-----+-----+-----+  
--  
tgtschema | pgbench_accounts | table | enterprisedb | 256 MB |  
tgtschema | pgbench_branches | table | enterprisedb | 8192 bytes |  
tgtschema | pgbench_history | table | enterprisedb | 25 MB |  
tgtschema | pgbench_tellers | table | enterprisedb | 16 kB |  
(4 rows)
```

When the `remotecopyschema` function was invoked, four background workers were specified.

The following portion of the log file `clone_rmt_src_tgt` shows the status of the parallel data copying operation using four background workers:

```
Wed Jun 28 current:18:05.232949 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0]  
table count [4]  
Wed Jun 28 current:18:05.233321 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0][0]  
worker started to copy data  
Wed Jun 28 current:18:05.233640 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0][1]  
worker started to copy data  
Wed Jun 28 current:18:05.233919 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0][2]  
worker started to copy data  
Wed Jun 28 current:18:05.234231 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0][3]  
worker started to copy data  
Wed Jun 28 current:18:05.298174 2017 EDT: [4024] INFO: [STAGE: DATA-COPY] [0][3]  
successfully copied data in [tgtschema.pgbench_branches]  
Wed Jun 28 current:18:05.2999current 2017 EDT: [4021] INFO: [STAGE: DATA-COPY] [0][0]  
successfully copied data in [tgtschema.pgbench_tellers]  
Wed Jun 28 current:18:06.634310 2017 EDT: [4022] INFO: [STAGE: DATA-COPY] [0][1]  
successfully copied data in [tgtschema.pgbench_history]  
Wed Jun 28 current:18:10.477333 2017 EDT: [4023] INFO: [STAGE: DATA-COPY] [0][2]  
successfully copied data in [tgtschema.pgbench_accounts]  
Wed Jun 28 current:18:10.477609 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0] all  
workers finished [4]  
Wed Jun 28 current:18:10.477654 2017 EDT: [4019] INFO: [STAGE: DATA-COPY] [0] copy  
done [4] tables  
Wed Jun 28 current:18:10.493938 2017 EDT: [4019] INFO: [STAGE: DATA-COPY]  
successfully copied data into tgtschema
```

Note that the `DATA-COPY` log message includes two, square bracket numbers (for example, `[0][3]`).

The first number is the job index whereas the second number is the worker index. The worker index values range from `0` to `3` for the four background workers.

In case two clone schema jobs are running in parallel, the first log file will have `0` as the job index whereas the second will have `1` as the job index.

`remotecopyschema_nb`

The `remotecopyschema_nb` function copies a schema and its database objects from a source schema in the remote source database specified within the `source_fdw` foreign server to a target schema in the local target database specified within the `target_fdw` foreign server, but in a non-blocking manner as a job submitted to pgAgent.

```
remotecopyschema_nb(
    <source_fdw> TEXT,
    <target_fdw> TEXT,
    <source> TEXT,
    <target> TEXT,
    <log_filename> TEXT
[, <on_tblspace> BOOLEAN
[, <verbose_on> BOOLEAN
[, <copy_acls> BOOLEAN
[, <worker_count> INTEGER ]]]]
)
```

An `INTEGER` value job ID is returned by the function for the job submitted to pgAgent. If the function fails, then null is returned.

The `source_fdw`, `target_fdw`, `source`, `target`, and `log_filename` are required parameters while all other parameters are optional.

After completion of the pgAgent job, remove the job with the `remove_log_file_and_job` function.

Parameters

`source_fdw`

Name of the foreign server managed by the `postgres_fdw` foreign data wrapper from which database objects are to be cloned.

`target_fdw`

Name of the foreign server managed by the `postgres_fdw` foreign data wrapper to which database objects are to be cloned.

`source`

Name of the schema from which database objects are to be cloned.

`target`

Name of the schema into which database objects are to be cloned from the source schema.

`log_filename`

Name of the log file in which information from the function is recorded. The log file is created under the

directory specified by the `log_directory` configuration parameter in the `postgresql.conf` file.

`on_tblspace`

`BOOLEAN` value to specify whether or not database objects are to be created within their tablespaces. If `false` is specified, then the `TABLESPACE` clause is not included in the applicable `CREATE` DDL statement when added to the target schema. If `true` is specified, then the `TABLESPACE` clause is included in the `CREATE` DDL statement when added to the target schema. If the `on_tblspace` parameter is omitted, the default value is `false`.

Note

If `true` is specified and a database object has a `TABLESPACE` clause, but that tablespace does not exist in the target database cluster, then the cloning function fails.

`verbose_on`

`BOOLEAN` value to specify whether or not the DDLs are to be printed in `log_filename` when creating objects in the target schema. If `false` is specified, then DDLs are not printed. If `true` is specified, then DDLs are printed. If omitted, the default value is `false`.

`copy_acls`

`BOOLEAN` value to specify whether or not the access control list (ACL) is to be included while creating objects in the target schema. The access control list is the set of `GRANT` privilege statements. If `false` is specified, then the access control list is not included for the target schema. If `true` is specified, then the access control list is included for the target schema. If the `copy_acls` parameter is omitted, the default value is `false`.

Note

If `true` is specified and a role with `GRANT` privilege does not exist in the target database cluster, then the cloning function fails.

`worker_count`

Number of background workers to perform the clone in parallel. If omitted, the default value is `1`.

Example

The same cloning operation is performed as the example in `remotecopyschema`, but using the non-blocking function `remotecopyschema_nb`.

The following command starts pgAgent on the target database `tgtdb`. The `pgagent` program file is located in the `bin` subdirectory of the Advanced Server installation directory.

```
[root@localhost bin]# ./pgagent -I 1 -s /tmp/pgagent_tgtdb_log
hostaddr=127.0.0.1 port=5444 user=enterprisedb dbname=tgtdb password=password
```

The `remotecopyschema_nb` function returns the job ID shown as `2` in the example.

```
tgtdb=# SELECTedb_util.remotecopyschema_nb
('src_server','tgt_server','srcschema','tgtschema','clone_rmt_src_tgt',worker
_count >= 4);
remotecopyschema_nb
-----
2
(1 row)
```

The completed status of the job is shown by the following:

```
tgtmdb=# SELECT edb_util.process_status_from_log('clone_rmt_src_tgt');
          process_status_from_log
-----
-----  

(FINISH,"29-JUN-17 current:16:00.100284 -04:00",3849,INFO,"STAGE:
FINAL","successfully clone schema into tgtschema")
(1 row)
```

The following removes the log file and the pgAgent job:

```
tgtmdb=# SELECT edb_util.remove_log_file_and_job ('clone_rmt_src_tgt',2);
remove_log_file_and_job
-----
t
(1 row)
```

process_status_from_log

The `process_status_from_log` function provides the status of a cloning function from its log file.

```
process_status_from_log (
  <log_file> TEXT
)
```

The function returns the following fields from the log file:

| Field Name | Description |
|----------------|---|
| status | Displays either STARTING, RUNNING, FINISH, or FAILED. |
| execution_time | When the command was executed. Displayed in timestamp format. |
| pid | Session process ID in which clone schema is getting called. |
| level | Displays either INFO, ERROR, or SUCCESSFUL. |
| stage | Displays either STARTUP, INITIAL, DDL-COLLECTION, PRE-DATA, DATA-COPY, POST-DATA, or FINAL. |
| message | Information respective to each command or failure. |

Parameters

log_file

Name of the log file recording the cloning of a schema as specified when the cloning function was invoked.

Example

The following shows usage of the `process_status_from_log` function:

```
edb=# SELECT edb_util.process_status_from_log('clone_edb_edbcopy');
          process_status_from_log
-----
-----  

(FINISH,"26-JUN-17 11:57:03.214458 -04:00",3691,INFO,"STAGE:
FINAL","successfully cloned schema")
(1 row)
```

remove_log_file_and_job

The `remove_log_file_and_job` function performs cleanup tasks by removing the log files created by the schema cloning functions and the jobs created by the non-blocking functions.

```
remove_log_file_and_job (
{ <log_file> TEXT |
<job_id> INTEGER |
<log_file> TEXT, <job_id> INTEGER
}
)
```

Values for any or both of the two parameters may be specified when invoking the `remove_log_file_and_job` function:

- If only `log_file` is specified, then the function will only remove the log file.
- If only `job_id` is specified, then the function will only remove the job.
- If both are specified, then the function will remove the log file and the job.

Parameters

`log_file`

Name of the log file to be removed.

`job_id`

Job ID of the job to be removed.

Example

The following examples removes only the log file, given the log filename.

```
edb=# SELECTedb_util.remove_log_file_and_job ('clone_edb_edbcopy');
remove_log_file_and_job
-----
t
(1 row)
```

The following example removes only the job, given the job ID.

```
edb=# SELECTedb_util.remove_log_file_and_job (3);
remove_log_file_and_job
-----
t
(1 row)
```

The following example removes the log file and the job, given both values:

```
tgtedb=# SELECTedb_util.remove_log_file_and_job ('clone_rmt_src_tgt',2);
remove_log_file_and_job
-----
t
(1 row)
```

11.15 Enhanced SQL and Other Miscellaneous Features

Advanced Server includes enhanced SQL functionality and various other features that provide additional flexibility and convenience. This chapter discusses some of these additions.

COMMENT

In addition to commenting on objects supported by the PostgreSQL `COMMENT` command, Advanced Server supports comments on additional object types. The complete supported syntax is:

```
COMMENT ON
{
    AGGREGATE <aggregate_name> ( <aggregate_signature> ) |
    CAST (<source_type> AS <target_type>) |
    COLLATION <object_name> |
    COLUMN <relation_name>. <column_name> |
    CONSTRAINT <constraint_name> ON <table_name> |
    CONSTRAINT <constraint_name> ON DOMAIN <domain_name> |
    CONVERSION <object_name> |
    DATABASE <object_name> |
    DOMAIN <object_name> |
    EXTENSION <object_name> |
    EVENT TRIGGER <object_name> |
    FOREIGN DATA WRAPPER <object_name> |
    FOREIGN TABLE <object_name> |
    FUNCTION <func_name> ([<argmode>] [<argname>] <argtype> [...]) |
    INDEX <object_name> |
    LARGE OBJECT <large_object_oid> |
    MATERIALIZED VIEW <object_name> |
    OPERATOR <operator_name> (left_type, right_type) |
    OPERATOR CLASS <object_name> USING <index_method> |
    OPERATOR FAMILY <object_name> USING <index_method> |
    PACKAGE <object_name> |
    POLICY <policy_name> ON <table_name> |
    [ PROCEDURAL ] LANGUAGE <object_name> |
    PROCEDURE <proc_name> ([<argmode>] [<argname>] <argtype> [, ...]) |
    PUBLIC SYNONYM <object_name> |
    ROLE <object_name> |
    RULE <rule_name> ON <table_name> |
    SCHEMA <object_name> |
    SEQUENCE <object_name> |
    SERVER <object_name> |
    TABLE <object_name> |
    TABLESPACE <object_name> |
    TEXT SEARCH CONFIGURATION <object_name> |
    TEXT SEARCH DICTIONARY <object_name> |
    TEXT SEARCH PARSER <object_name> |
    TEXT SEARCH TEMPLATE <object_name> |
    TRANSFORM FOR <type_name> LANGUAGE <lang_name> |
    TRIGGER <trigger_name> ON <table_name> |
    TYPE <object_name> |
    VIEW <object_name>
} IS '<text>'
```

where `aggregate_signature` is:

```
* |
[ <argmode> ] [ <argname> ] <argtype> [ , ... ] |
[ [ <argmode> ] [ <argname> ] <argtype> [ , ... ] ]
ORDER BY [ <argmode> ] [ <argname> ] <argtype> [ , ... ]
```

Parameters

`object_name`

The name of the object on which you are commenting.

`AGGREGATE aggregate_name (aggregate_signature)`

Include the `AGGREGATE` clause to create a comment about an aggregate. `aggregate_name` specifies the name of an aggregate, and `aggregate_signature` specifies the associated signature in one of the following forms:

```
* |
[ <argmode> ] [ <argname> ] <argtype> [ , ... ] |
[ [ <argmode> ] [ <argname> ] <argtype> [ , ... ] ]
ORDER BY [ <argmode> ] [ <argname> ] <argtype> [ , ... ]
```

Where ```argmode``` is the mode of a function, procedure, or aggregate argument, `argmode` may be ```IN```, ```OUT```, ```INOUT```, or ```VARIADIC```. If omitted, the default is ```IN```.

```argname``` is the name of an aggregate argument.

```argtype``` is the data type of an aggregate argument.

`CAST (source_type AS target_type)`

Include the `CAST` clause to create a comment about a cast. When creating a comment about a cast, `source_type` specifies the source data type of the cast, and `target_type` specifies the target data type of the cast.

`COLUMN relation_name.column_name`

Include the `COLUMN` clause to create a comment about a column. `column_name` specifies name of the column to which the comment applies. `relation_name` is the table, view, composite type, or foreign table in which a column resides.

`CONSTRAINT constraint_name ON table_name`

`CONSTRAINT constraint_name ON DOMAIN domain_name`

Include the `CONSTRAINT` clause to add a comment about a constraint. When creating a comment about a constraint, `constraint_name` specifies the name of the constraint. `table_name` or `domain_name` specifies the name of the table or domain on which the constraint is defined.

`FUNCTION func_name ([[argmode] [argname] argtype [, ...]])`

Include the `FUNCTION` clause to add a comment about a function. `func_name` specifies the name of the function. `argmode` specifies the mode of the function. `argmode` may be `IN`, `OUT`, `INOUT`, or `VARIADIC`. If omitted, the default is `IN`.

`argname` specifies the name of a function, procedure, or aggregate argument. `argtype` specifies the data

type of a function, procedure, or aggregate argument.

`large_object_oid`

`large_object_oid` is the system-assigned OID of the large object about which you are commenting.

`OPERATOR operator_name (left_type, right_type)`

Include the `OPERATOR` clause to add a comment about an operator. `operator_name` specifies the (optionally schema-qualified) name of an operator on which you are commenting. `left_type` and `right_type` are the (optionally schema-qualified) data type(s) of the operator's arguments.

`OPERATOR CLASS object_name USING index_method`

Include the `OPERATOR CLASS` clause to add a comment about an operator class. `object_name` specifies the (optionally schema-qualified) name of an operator on which you are commenting. `index_method` specifies the associated index method of the operator class.

`OPERATOR FAMILY object_name USING index_method`

Include the `OPERATOR FAMILY` clause to add a comment about an operator family. `object_name` specifies the (optionally schema-qualified) name of an operator family on which you are commenting. `index_method` specifies the associated index method of the operator family.

`POLICY policy_name ON table_name`

Include the `POLICY` clause to add a comment about a policy. `policy_name` specifies the name of the policy, and `table_name` specifies the table that the policy is associated with.

`PROCEDURE proc_name [([argmode] [argname] argtype [, ...]])]`

Include the `PROCEDURE` clause to add a comment about a procedure. `proc_name` specifies the name of the procedure. `argmode` specifies the mode of the procedure. `argmode` may be `IN`, `OUT`, `INOUT`, or `VARIADIC`. If omitted, the default is `IN`.

`argname` specifies the name of a function, procedure, or aggregate argument. `argtype` specifies the data type of a function, procedure, or aggregate argument.

`RULE rule_name ON table_name`

Include the `RULE` clause to specify a `COMMENT` on a rule. `rule_name` specifies the name of the rule, and `table_name` specifies the name of the table on which the rule is defined.

`TRANSFORM FOR type_name LANGUAGE lang_name |`

Include the `TRANSFORM FOR` clause to specify a `COMMENT` on a `TRANSFORM`.

`type_name` specifies the name of the data type of the transform and `lang_name` specifies the name of the language of the transform.

`TRIGGER trigger_name ON table_name`

Include the `TRIGGER` clause to specify a `COMMENT` on a trigger. `trigger_name` specifies the name of the trigger, and `table_name` specifies the name of the table on which the trigger is defined.

`text`

The comment, written as a string literal or `NULL` to drop the comment.

Notes:

Names of tables, aggregates, collations, conversions, domains, foreign tables, functions, indexes, operators, operator classes, operator families, packages, procedures, sequences, text search objects, types, and views can

be schema-qualified.

Example:

The following example adds a comment to a table named `new_emp`:

```
COMMENT ON TABLE new_emp IS 'This table contains information about new employees.';
```

For more information about using the `COMMENT` command, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/sql-comment.html>

Output of Function `version()`

The text string output of the `version()` function displays the name of the product, its version, and the host system on which it has been installed.

For Advanced Server, the `version()` output is in a format similar to the PostgreSQL community version in that the first text word is *PostgreSQL* instead of *EnterpriseDB* as in Advanced Server version 10 and earlier.

The general format of the `version()` output is the following:

```
PostgreSQL $PG_VERSION_EXT (EnterpriseDB Advanced Server $PG_VERSION) on $host
```

So for the current Advanced Server the version string appears as follows:

```
edb@45032=#select version();
version
-----
```

```
-----  
PostgreSQL 13.0 (EnterpriseDB Advanced Server 13.0.0) on x86_64-pc-linux-gnu, compiled by gcc (GCC) 4.8.5 20150623 (Red Hat 4.8.5-11), 64-bit  
(1 row)
```

In contrast, for Advanced Server 10, the version string was the following:

```
edb=# select version();
version
-----  
-----  
EnterpriseDB 10.4.9 on x86_64-pc-linux-gnu, compiled by gcc (GCC) 4.4.7  
20120313 (Red Hat 4.4.7-18), 64-bit  
(1 row)
```

Logical Decoding on Standby

Logical decoding on a standby server allows you to create a logical replication slot on a standby server that can respond to API operations such as `get`, `peek`, `advance`, etc..

For more information about the `LOGICAL DECODING`, please refer to the PostgreSQL core documentation available at:

<https://www.postgresql.org/docs/current/logicaldecoding-explanation.html>

For a logical slot on a standby server to work, the `hot_standby_feedback` parameter must be set to `ON` on the standby. The `hot_standby_feedback` parameter prevents `VACUUM` from removing recently-dead rows that are required by an existing logical replication slot on the standby server. If a slot conflict occurs on the standby, the slots will be dropped.

For logical decoding on a standby to work, `wal_level` must be set to `logical` on both the primary and standby server. If `wal_level` is set to a value other than `logical`, then slots are not created. If you set `wal_level` to a value other than `logical` on primary and if there are existing logical slots on standby, such slots are dropped and new slots cannot be created.

When transactions are written to the primary server, the activity will trigger the creation of a logical slot on the standby server. If a primary server is idle, creating a logical slot on a standby server may take noticeable time.

For more information about functions that support replication and logical decoding example, see to the PostgreSQL documentation available at:

<https://www.postgresql.org/docs/current/functions-admin.html#FUNCTIONS-REPLICATION>

<https://www.postgresql.org/docs/current/logicaldecoding-example.html>

11.16 System Catalog Tables

The following system catalog tables contain definitions of database objects. The layout of the system tables is subject to change, if you are writing an application that depends on information stored in the system tables, it would be prudent to use an existing catalog view, or create a catalog view to isolate the application from changes to the system table.

edb_dir

The `edb_dir` table contains one row for each alias that points to a directory created with the `CREATE DIRECTORY` command. A directory is an alias for a pathname that allows a user limited access to the host file system.

You can use a directory to fence a user into a specific directory tree within the file system. For example, the `UTL_FILE` package offers functions that permit a user to read and write files and directories in the host file system, but only allows access to paths that the database administrator has granted access to via a `CREATE DIRECTORY` command.

| Column | Type | Modifiers | Description |
|----------|-----------|-----------|---|
| dirname | "name" | not null | The name of the alias. |
| dirowner | oid | | |
| dirpath | text | | |
| diracl | aclitem[] | not null | The OID of the user that owns the alias. The directory name to which access is granted. The access control list that determines which users may access the alias. |

edb_all_resource_groups

The `edb_all_resource_groups` table contains one row for each resource group created with the `CREATE RESOURCE GROUP` command and displays the number of active processes in each resource group.

| Column | Type | Modifiers | Description |
|------------------------------|---------|-----------|-------------|
| group_name | "name" | | |
| active_processes | integer | | |
| cpu_rate_limit | float8 | | |
| per_process_cpu_rate_limit | float8 | | |
| dirty_rate_limit | float8 | | |
| per_process_dirty_rate_limit | float8 | | |

edb_policy

The `edb_policy` table contains one row for each policy.

| Column | Type | Modifiers | Description |
|-------------------|------------|-----------|---|
| policyname | name | not null | The policy name. |
| policygroup | oid | not null | Currently unused. |
| policyobject | oid | not null | The OID of the table secured by this policy (the object_schema plus the object_name). |
| | | | The kind of object secured by this policy: 'r' for a table |
| policykind | char | not null | 'v' for a view = for a synonym Currently always 'r'. |
| | | | |
| policyproc | oid | not null | The OID of the policy function (function_schema plus policy_function). |
| policyinsert | boolean | not null | True if the policy is enforced by INSERT statements. |
| policyselect | boolean | not null | True if the policy is enforced by SELECT statements. |
| policydelete | boolean | not null | True if the policy is enforced by DELETE statements. |
| policyupdate | boolean | not null | True if the policy is enforced by UPDATE statements. |
| policyindex | boolean | not null | Currently unused. |
| policyenabled | boolean | not null | True if the policy is enabled. |
| policyupdatecheck | boolean | not null | True if rows updated by an UPDATE statement must satisfy the policy. |
| policystatic | boolean | not null | Currently unused. |
| policytype | integer | not null | Currently unused. |
| policyopts | integer | not null | Currently unused. |
| policyseccols | int2vector | not null | The column numbers for columns listed in sec_relevant_cols. |

edb_profile

The `edb_profile` table stores information about the available profiles. `edb_profiles` is shared across all databases within a cluster.

| Column | Type | References | Description |
|--------|------|------------|-------------|
| | | | |

| Column | Type | References | Description |
|--|---|-----------------|--|
| oid prfname prffailedloginattempts prfpasswordlocktime prfpasswordlifetime prfpasswordgracetime prfpasswordreusetime prfpasswordreusemax | oid name integer integer integer integer integer integer | | Row identifier (hidden attribute, must be explicitly selected). The name of the profile. The number of failed login attempts allowed by the profile. -1 indicates that the value from the default profile should be used. -2 indicates no limit on failed login attempts. The password lock time associated with the profile (in seconds). -1 indicates that the value from the default profile should be used. -2 indicates that the account should be locked permanently. The password life time associated with the profile (in seconds). -1 indicates that the value from the default profile should be used. -2 indicates that the password never expires. The password grace time associated with the profile (in seconds). -1 indicates that the value from the default profile should be used. -2 indicates that the password never expires. The number of seconds that a user must wait before reusing a password. -1 indicates that the value from the default profile should be used. -2 indicates that the old passwords can never be reused. The number of password changes that have to occur before a password can be reused. -1 indicates that the value from the default profile should be used. -2 indicates that the old passwords can never be reused. |
| prfpasswordverifyfuncdb | oid | pg_database.oid | The OID of the database in which the password verify function exists. |
| prfpasswordverifyfunc | oid | pg_proc.oid | The OID of the password verify function associated with the profile. |

edb_redaction_column

The catalog `edb_redaction_column` stores information of data redaction policy attached to the columns of the table.

| Column | Type | References | Description |
|---|--------------------------------|--------------------------|--|
| oid | oid | | Row identifier (hidden attribute, must be explicitly selected) |
| rdpolicyid | oid | edb_redaction_policy.oid | The data redaction policy applies to the described column |
| rdrelid | oid | pg_class.oid | The table that the described column belongs to |
| rdattnum rdscope rdexception rfuncexpr | int2 int2 int2 pg_node_tree | pg_attribute.attnum | The number of the described column The redaction scope: 1 = query, 2 = top_tlist, 4 = top_tlist_or_error The redaction exception: 8 = none, 16 = equal, 32 = leakproof Data redaction function expression |

Note

The described column will be redacted if the redaction policy `edb_redaction_column.rdpolicyid` on the table is enabled and the redaction policy expression `edb_redaction_policy.rdexpr` evaluates to true.

edb_redaction_policy

The catalog `edb_redaction_policy` stores information of the redaction policies for tables.

| Column | Type | References | Description |
|-------------------------------|--------------------------|--------------|---|
| oid rdname | oid name | | Row identifier (hidden attribute, must be explicitly selected) The name of the data redaction policy |
| rdrelid rdenable rdexpr | oid boolean pg_node_tree | pg_class.oid | The table to which the data redaction policy applies Is the data redaction policy enabled? The data redaction policy expression |

Note

The data redaction policy applies for the table if it is enabled and the expression ever evaluated true.

edb_resource_group

The `edb_resource_group` table contains one row for each resource group created with the `CREATE RESOURCE GROUP` command.

| Column | Type | Modifiers | Description |
|--------------------|--------|-----------|--|
| rgrpname | "name" | not null | The name of the resource group. |
| rgrpcpuratelimit | float8 | not null | Maximum CPU rate limit for a resource group. 0 means no limit. |
| rgrpdirtyratelimit | float8 | not null | Maximum dirty rate limit for a resource group. 0 means no limit. |

edb_variable

The `edb_variable` table contains one row for each package level variable (each variable declared within a package).

| Column | Type | Modifiers | Description |
|------------|----------|-----------|---|
| varname | "name" | not null | The name of the variable. |
| varpackage | oid | not null | The OID of the pg_namespace row that stores the package. |
| vartype | oid | not null | The OID of the pg_type row that defines the type of the variable. |
| | | | + if the variable is visible outside of the package. - if the variable is only visible within the package. |
| varaccess | "char" | not null | |
| varsrcc | text | | Note: Public variables are declared within the package header, private variables are declared within the package body. Contains the source of the variable declaration, including any default value expressions for the variable. |
| varseq | smallint | not null | The order in which the variable was declared in the package. |

pg_synonym

The `pg_synonym` table contains one row for each synonym created with the `CREATE SYNONYM` command or `CREATE PUBLIC SYNONYM` command.

| Column | Type | Modifiers | Description |
|---------|--------|-----------|--------------------------|
| synname | "name" | not null | The name of the synonym. |

| Column | Type | Modifiers | Description |
|--------------|--------|-----------|--|
| synnamespace | oid | not null | Replaces synowner. Contains the OID of the pg_namespace row where the synonym is stored |
| synowner | oid | not null | The OID of the user that owns the synonym. |
| synobjschema | "name" | not null | The schema in which the referenced object is defined. |
| synobjname | "name" | | The name of the referenced object. The (optional) name of the database link in which the referenced object is defined. |
| synlink | text | not null | |

product_component_version

The `product_component_version` table contains information about feature compatibility, an application can query this table at installation or run time to verify that features used by the application are available with this deployment.

| Column | Type | Description |
|---------|------------------------|------------------------------------|
| product | character varying (74) | The name of the product. |
| version | character varying (74) | The version number of the product. |
| status | character varying (74) | The status of the release. |

11.17 Advanced Server Keywords

A keyword is a word that is recognized by the Advanced Server parser as having a special meaning or association. You can use the `pg_get_keywords()` function to retrieve an up-to-date list of the Advanced Server keywords:

```
acctg=# acctg=# SELECT * FROM pg_get_keywords();
      word      | catcode |      catdesc
-----+-----+-----
    abort     | U      | unreserved
  absolute   | U      | unreserved
   access    | U      | unreserved
...
...
```

`pg_get_keywords` returns a table containing the keywords recognized by Advanced Server:

- The `word` column displays the keyword.
- The `catcode` column displays a category code.
- The `catdesc` column displays a brief description of the category to which the keyword belongs.

Note that any character can be used in an identifier if the name is enclosed in double quotes. You can selectively query the `pg_get_keywords()` function to retrieve an up-to-date list of the Advanced Server keywords that belong to a specific category:

```
| SELECT * FROM pg_get_keywords() WHERE catcode = 'code';
```

Where `code` is:

| R - The word is reserved. Reserved keywords may never be used as an identifier, they are reserved for use

by the server.

U - The word is unreserved. Unreserved words are used internally in some contexts, but may be used as a name for a database object.

T - The word is used internally, but may be used as a name for a function or type.

C - The word is used internally, and may not be used as a name for a function or type.

For more information about Advanced Server identifiers and keywords, see to the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/current/static/sql-syntax-lexical.html>

12 EDB Plus User Guide

This guide describes how to connect to an Advanced Server database using EDB*Plus. EDB*Plus provides a command line user interface to EDB Postgres Advanced Server that accepts SQL commands that allow you to:

- Query certain database objects
- Execute stored procedures
- Format output from SQL commands
- Execute batch scripts
- Execute OS commands
- Record output

For detailed information about the features supported by Advanced Server, please consult the complete library of Advanced Server guides available at:

<https://www.enterprisedb.com/edb-docs>

12.1 What's New

The following features have been added to EDB*Plus version 36:

- Advanced Server supports use of Internet Protocol version 6 (IPv6) to log into EDB*Plus. For more information, see [Starting EDB*Plus](#).
 - Advanced Server supports use of a Secure Sockets Layer (SSL) connection from EDB*Plus to the Advanced Server database. For more information, see [Using a Secure Sockets Layer \(SSL\) Connection](#).
-

12.2 EDB Plus

EDB*Plus is a utility program that provides a command line user interface to EDB Postgres Advanced Server. EDB*Plus accepts SQL commands, SPL anonymous blocks, and EDB*Plus commands.

EDB*Plus commands are compatible with Oracle SQL*Plus commands and provide various capabilities including:

- Querying certain database objects
- Executing stored procedures
- Formatting output from SQL commands
- Executing batch scripts
- Executing OS commands
- Recording output

12.3 Installing EDB Plus

You can use an RPM installer or a graphical installer to add EDB*Plus to your Advanced Server installation.

Installation Pre-Requisites

Before installing EDB*Plus, you must first install Java (version 1.7 or later). On a Linux system, you can use the yum package manager to install Java. Open a terminal window, assume superuser privileges, and enter:

```
# yum install java
```

If you are using Windows, Java installers and instructions are available online at:

```
http://www.java.com/en/download/manual.jsp
```

You must also have credentials that allow access to the EnterpriseDB repository. For information about requesting credentials, visit:

```
https://info.enterprisedb.com/rs/069-ALB-339/images/Repository%20Access%2004-09-2019.pdf
```

After receiving your repository credentials you can:

1. Create the repository configuration file.
2. Modify the file, providing your user name and password.
3. Install EDB*Plus.

Performing an RPM Installation

For detailed information about creating and using EnterpriseDB repositories to install Advanced Server or its supporting components, please see the *EDB Postgres Advanced Server Installation Guide*, available at:

```
https://www.enterprisedb.com/edb-docs
```

Creating a Repository Configuration File

To create the repository configuration file, assume superuser privileges, and invoke the following command:

```
yum -y install https://yum.enterprisedb.com/edb-repo-rpms/edb-repo-latest.noarch.rpm
```

The repository configuration file is named `edb.repo`. The file resides in `/etc/yum.repos.d`.

After creating the `edb.repo` file, use your choice of editor to ensure that the value of the enabled parameter is `1`, and replace the `username` and `password` placeholders in the `baseurl` specification with the name and password

of a registered EnterpriseDB user.

```
[edb]
name=EnterpriseDB RPMs $releasever - $basearch
baseurl=https://<username>:<password>@yum.enterprisedb.com/edb/redhat/rhel-$releasever-$basearch
enabled=1
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/ENTERPRISEDB-GPG-KEY
```

After saving your changes to the configuration file, you can use the `yum install` command to install EDB*Plus:

`| edb-as11-edbplus`

When you install an RPM package that is signed by a source that is not recognized by your system, yum may ask for your permission to import the key to your local server. If prompted, and you are satisfied that the packages come from a trustworthy source, enter `y`, and press `Return` to continue.

During the installation, yum may encounter a dependency that it cannot resolve. If it does, it will provide a list of the required dependencies that you must manually resolve.

Configuring an RPM Installation

After performing an RPM installation of EDB*Plus, you must set the values of environment variables that allow EDB*Plus to locate your Java installation. Use the following commands to set variable values:

`| export JAVA_HOME=<path_to_java> export PATH=<path_to_java>/bin:$PATH`

By default, the `pg_hba.conf` file for the RPM installer enforces IDENT authentication. Before invoking EDB*Plus, you must either modify the `pg_hba.conf` file, changing the authentication method to a form other than IDENT (and restarting the server), or perform the following steps to ensure that an IDENT server is accessible:

You must confirm that an `identd` server is installed and running. You can use the yum package manager to install an `identd` server by invoking the command:

`| yum install xinetd authd`

The command should create a file named `/etc/xinetd.d/auth` that contains:

```
service auth
{
disable = yes
socket_type = stream
wait =no
user = ident
cps = 4096 10
instances = UNLIMITED
server = /usr/sbin/in.authd server_args = -t60 --xerror --os
}
```

Note

If the file includes a `-E` argument at the end of the server arguments, please erase `-E`.

Then, to start the `identd` server, invoke the following commands:

`| systemctl enable xinetd`

`| systemctl start xinetd`

Open the `pg_ident.conf` file and create a user mapping:

```
### map_name    system_username    postgres_username
edbas        enterprisedb      enterprisedb
```

Where:

- The name specified in the `map_name` column is a user-defined name that will identify the mapping in the `pg_hba.conf` file.
- The name specified in the `system_username` column is `enterprisedb`.
- The name specified in the `postgres_username` column is `enterprisedb`.

Then, open the `pg_hba.conf` file and modify the IDENT entries:

- If you are using an IPv4 local connection, modify the file entry to read:

```
| host all all 127.0.0.0/0 ident map=edbas
```

- If you are using an IPv6 local connection, modify the file entry to read:


```
| host all all ::1/128 ident map=edbas
```

You must restart the Advanced Server service before invoking EDB*Plus. For detailed information about controlling the Advanced Server service, see the *EDB Postgres Advanced Server Installation Guide*, available at:


```
| https://www.enterprisedb.com/resources/product-documentation
```

Using the Graphical Installer


Graphical installers for EDB*Plus are available via StackBuilder Plus; you can access StackBuilder Plus through your Windows or Linux start menu. After opening StackBuilder Plus and selecting the installation for which you wish to install EDB*Plus, expand the component selection screen tree control to select and download the EDB*Plus installer.



The EDB*Plus installer welcomes you to the setup wizard, as shown in the figure below.



Use the **Installation Directory** field to specify the directory in which you wish to install the EDB*Plus software. Then, click **Next** to continue.



Use fields on the **EDB Postgres Advanced Server Installation Details** window to identify the location of the Advanced Server host:

- Use the Host field to identify the system on which Advanced Server resides.
- Use the Port field to identify the listener port that Advanced Server monitors for client connections.

Then, click **Next** to continue.



The Ready to Install window notifies you when the installer has all of the information needed to install EDB*Plus on your system. Click **Next** to install EDB*Plus.



The installer notifies you when the setup wizard has completed the EDB*Plus installation. Click **Finish** to exit the installer.

12.4 Using EDB Plus

To open an EDB*Plus command line, navigate through the **Applications** or **Start** menu to the **Advanced Server** menu, to the **Run SQL Command Line** menu, and select the **EDB Plus** option. You can also invoke **EDB Plus** from the operating system command line with the following command:

```
> edbplus [ -S[ILENT] ] [ <login> | NOLOG ] [ @<scriptfile> [.<ext>] ] -SILENT
```

If specified, the EDB*Plus sign-on banner is suppressed along with all prompts.

`login`

Login information for connecting to the database server and database. `login` takes the following form; please note that there must be no white space within the login information:

```
> <username>/[<password>][@{<connectstring>} | <variable>]
```

Where:

- `username` is a database username with which to connect to the database.
- `password` is the password associated with the specified `username`. If a `password` is not provided, but a password is required for authentication, a password file is used if available. If there is no password file or no entry in the password file with the matching connection parameters, then EDB*Plus will prompt for the password.
- `connectstring` is the database connection string with the following format:

```
<host>[:<port>]//<dbname>[?ssl={true | false}]
```

Where:

- `host` is the hostname or IP address on which the database server resides. If neither `@ connectstring` nor `@ variable` nor `/NOLOG` is specified, the default host is assumed to be the localhost.
- `port` is the port number receiving connections on the database server. If not specified, the default is `5444`.
- `dbname` is the name of the database to connect to. If not specified the default is `edb`.

- If **Internet Protocol version 6** (IPv6) is used for the connection instead of IPv4, then the IP address must be enclosed within square brackets (that is, [**ipv6_address**]). The following is an example that includes an IPv6 connection:

```
edbplus.sh enterprisedb/password@[fe80::20c:29ff:fe7c:78b2]:5444/edb
```

The **pg_hba.conf** file for the database server must contain an appropriate entry for the IPv6 connection. The following example shows an entry that allows all addresses:

```
### TYPE DATABASE USER ADDRESS METHOD
host all all ::/0/0 md5
```

For more information about the **pg_hba.conf** file, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/Current/static/auth-pg-hba-conf.html>

If an SSL connection is desired, then include the **?ssl=true** parameter in the connection string. In such a case, the connection string must minimally include *host:port*, with or without */dbname*. If the *ssl* parameter is not specified, the default is **false**. See Section 2.3 for instructions on setting up an SSL connection.

- variable** is a variable defined in the **login.sql** file that contains a database connection string. The **login.sql** file can be found in the **edbplus** subdirectory of the Advanced Server home directory.

/NOLOG

Specify **/NOLOG** to start EDB*Plus without establishing a database connection. SQL commands and EDB*Plus commands that require a database connection cannot be used in this mode. The CONNECT command can be subsequently given to connect to a database after starting EDB*Plus with the **/NOLOG** option.

scriptfile\ [<ext>]

scriptfile is the name of a file residing in the current working directory, containing SQL and/or EDB*Plus commands that will be automatically executed after startup of EDB*Plus. *ext* is the filename extension. If the filename extension is *sql*, then the **.sql** extension may be omitted when specifying **scriptfile**. When creating a script file, always name the file with an extension, otherwise it will not be accessible by EDB*Plus. (EDB*Plus will always assume a **.sql** extension on filenames that are specified with no extension.)

The following example shows user **enterprisedb** with password **password**, connecting to database **edb** running on a database server on the **localhost** at port **5444**.

```
C:\Program Files\edb\as10\edbplus>edbplus enterprisedb/password
Connected to EnterpriseDB 10.0.1 (localhost:5444/edb) AS enterprisedb
EDB Plus: Release 10 (Build 36.0.0)
Copyright (c) 2008-2017, EnterpriseDB Corporation. All rights reserved.
SQL>
```

The following example shows user **enterprisedb** with password, **password**, connecting to database **edb** running on a database server on the **localhost** at port **5445**.

```
C:\Program Files\edb\as10\edbplus>edbplus enterprisedb/password@localhost:5445/edb
Connected to EnterpriseDB 10.0.1 (localhost:5445/edb) AS enterprisedb

EDB Plus: Release 10 (Build 36.0.0)
Copyright (c) 2008-2017, EnterpriseDB Corporation. All rights reserved.

SQL>
```

The following example demonstrates using variable **hr_5445** in the **login.sql** file, and connecting to database **hr** on **localhost** at port **5445**.

```
C:\Program Files\edb\as10\edbplus>edbplus enterprisedb/password@hr_5445
Connected to EnterpriseDB 10.0.1 (localhost:5445/hr) AS enterprisedb
```

EDB Plus: Release 10 (Build 36.0.0)
 Copyright (c) 2008-2017, EnterpriseDB Corporation. All rights reserved.

SQL>

The following is the content of the `login.sql` file used in the previous example.

```
define edb="localhost:5445/edb"
define hr_5445="localhost:5445/hr"
```

The following example executes a script file, `dept_query.sql` after connecting to database `edb` on server `localhost` at port `5444`.

```
C:\Program Files\edb\as10\edbplus>edbplus enterprisedb/password @dept_query
Connected to EnterpriseDB 10.0.1 (localhost:5444/edb) AS enterprisedb
```

SQL> SELECT * FROM dept;

| DEPTNO | DNAME | LOC |
|--------|------------|----------|
| 10 | ACCOUNTING | NEW YORK |
| 20 | RESEARCH | DALLAS |
| 30 | SALES | CHICAGO |
| 40 | OPERATIONS | BOSTON |

SQL> EXIT
 Disconnected from EnterpriseDB Database.

The following is the content of file `dept_query.sql` used in the previous example.

```
SET PAGESIZE 9999
SET ECHO ON
SELECT * FROM dept;
EXIT
```

12.5 Using a Secure Sockets Layer (SSL) Connection

An EDB*Plus connection to the Advanced Server database can be accomplished using secure sockets layer (SSL) connectivity.

Using SSL requires various prerequisite configuration steps performed on the database server involved with the SSL connection as well as creation of the Java truststore and keystore on the host that will run EDB*Plus.

The Java *truststore* is the file containing the Certificate Authority (CA) certificates with which the Java client (EDB*Plus) uses to verify the authenticity of the server to which it is initiating an SSL connection.

The Java *keystore* is the file containing private and public keys and their corresponding certificates. The keystore is required for client authentication to the server, which is used for the EDB*Plus connection.

The following is material to which you can refer to for guidance in setting up the SSL connections:

- For information on setting up SSL connectivity to the Advanced Server database, see the section about secure TCP connections with SSL in (*Server Setup and Operation*) in the *PostgreSQL Core Documentation* located at:
<https://www.postgresql.org/docs/11/ssl-tcp.html>
- For information on JDBC client connectivity using SSL, see the section on configuring the client ([Using SSL](#)) in *The PostgreSQL JDBC Interface* located at:
<https://jdbc.postgresql.org/documentation/head/ssl-client.html>

The following sections provide information for the configuration steps of using SSL.

- Configuring SSL on Advanced Server
- Configuring SSL for the EDB*Plus client
- Requesting SSL connection to the Advanced Server database

Configuring SSL on Advanced Server

This section provides an example of configuring SSL on a database server to demonstrate the use of SSL with EDB*Plus. A self-signed certificate is used for this purpose.

Step 1: Create the certificate signing request (CSR).

In the following example the generated certificate signing request file is `server.csr`. The private key is generated as file `server.key`.

```
$ openssl req -new -nodes -text -out server.csr \
> -keyout server.key -subj "/CN=enterprisedb"
Generating a 2048 bit RSA private key
.....+++
.....+++
writing new private key to 'server.key'
---
```

Note

When creating the certificate, the value specified for the common name field (designated as CN=enterprisedb in this example) must be the database user name that is specified when connecting to EDB*Plus.

In addition, user name maps can be used as defined in the `pg_ident.conf` file to permit more flexibility for the common name and database user name. Steps 8 and 9 describe the use of user name `maps`.

Step 2: Generate the self-signed certificate.

The following generates a self-signed certificate to file `server.crt` using the certificate signing request file, `server.csr`, and the private key, `server.key`, as input.

```
$ openssl x509 -req -days 365 -in server.csr -signkey server.key \
> -out server.crt
Signature ok
subject=/CN=enterprisedb
Getting Private key
```

Step 3: Make a copy of the server certificate (`server.crt`) to be used as the root Certificate Authority (CA) file (`root.crt`).

```
$ cp server.crt root.crt
```

Step 4: Delete the now redundant certificate signing request (`server.csr`).

```
$ rm -f server.csr
```

Step 5: Move or copy the certificate and private key files to the Advanced Server data directory (for example, `/opt/edb/as10/data`).

```
$ mv root.crt /opt/edb/as10/data
$ mv server.crt /opt/edb/as10/data
$ mv server.key /opt/edb/as10/data
```

Step 6: Set the file ownership and permissions on the certificate files and private key file.

Set the ownership to the operating system account that owns the data sub-directory of the database server. Set the permissions so that no other groups or accounts other than the owner can access these files.

```
$ chown enterprise db root.crt server.crt server.key
$ chgrp enterprise db root.crt server.crt server.key
$ chmod 600 root.crt server.crt server.key
$ ls -l
total 152
.
.
.
-rw----- 1 enterprise db 985 Aug 22 11:00 root.crt
-rw----- 1 enterprise db 985 Aug 22 10:59 server.crt
-rw----- 1 enterprise db 1704 Aug 22 10:58 server.key
```

Step 7: In the `postgresql.conf` file, make the following modifications.

```
ssl = on
ssl_cert_file = 'server.crt'
ssl_key_file = 'server.key'
ssl_ca_file = 'root.crt'
```

Step 8: Modify the `pg_hba.conf` file to enable SSL usage on the desired database to which EDB*Plus is to make the SSL connection.

In the `pg_hba.conf` file, the `hostssl` type indicates the entry is used to validate SSL connection attempts from the client (EDB*Plus).

The authentication method is set to cert with the option `clientcert=1` in order to require an SSL certificate from the client against which authentication is performed using the common name of the certificate (enterprise db in this example).

The `map=sslusers` option specifies that a mapping named `sslusers` defined in the `pg_ident.conf` file is to be used for authentication. This mapping allows a connection to the database if the common name from the certificate and the database user name attempting the connection match the SYSTEM-USERNAME/PG-USERNAME pair listed in the `pg_ident.conf` file.

The following is an example of the settings in the `pg_hba.conf` file if the database (edb) must use SSL connections.

```
## TYPE DATABASE USER ADDRESS METHOD
## "local" is for Unix domain socket connections only
local all all md5
```

```
### IPv4 local connections:
```

```
hostssl edb all 192.168.2.0/24 cert clientcert=1 map=sslusers
```

Step 9: The following shows the user name maps in the `pg_ident.conf` file related to the `pg_hba.conf` file by the `map=sslusers` option. These user name maps permit you to specify database user names `edbuser`, `postgres`, or `enterprisedb` when connecting with EDB*Plus.

```
### MAPNAME SYSTEM-USERNAME PG-USERNAME
sslusers enterprisedb edbuser
sslusers enterprisedb postgres
sslusers enterprisedb enterprisedb
```

Step 10: Restart the database server after you have made the changes to the configuration files.

Configuring SSL for the EDB*Plus Client

After you have configured SSL on the database server, the following steps provide an example of generating certificate and keystore files for EDB*Plus (the JDBC client).

Step 1: Using files `server.crt` and `server.key` located under the database server data sub-directory, create copies of these files and move them to the host where EDB*Plus is to be running.

Store these files in the desired directory to contain the trusted certificate and keystore files to be generated in the following steps. The suggested location is to create a `.postgresql` sub-directory under the home user account that will invoke EDB*Plus. Thus, these files will be under the `~/.postgresql` directory of the user account that will run EDB*Plus.

For this example, assume file `edb.crt` is a copy of `server.crt` and `edb.key` is a copy of `server.key`.

Step 2: Create an additional copy of `edb.crt`.

```
$ cp edb.crt edb_root.crt
$ ls -l
total 12
-rw-r--r-- 1 user user 985 Aug 22 14:17 edb.crt
-rw-r--r-- 1 user user 1704 Aug 22 14:18 edb.key
-rw-r--r-- 1 user user 985 Aug 22 14:19 edb_root.crt
```

Step 3: Create a Distinguished Encoding Rules (DER) format of file `edb_root.crt`. The generated DER format of this file is `edb_root.crt.der`. The DER format of the file is required for the `keytool` program used in Step 4.

```
$ openssl x509 -in edb_root.crt -out edb_root.crt.der -outform der
$ ls -l
total 16
-rw-r--r-- 1 user user 985 Aug 22 14:17 edb.crt
-rw-r--r-- 1 user user 1704 Aug 22 14:18 edb.key
-rw-r--r-- 1 user user 985 Aug 22 14:19 edb_root.crt
-rw-rw-r-- 1 user user 686 Aug 22 14:21 edb_root.crt.der
```

Step 4: Use the `keytool` program to create a keystore file (`postgresql.keystore`) using `edb_root.crt.der` as the input. This process adds the certificate of the Postgres database server to the keystore file.

Note: The file name `postgresql.keystore` is recommended so that it can be accessed in its default directory location `~/.postgresql/postgresql.keystore`, which is under the home directory of the user account invoking EDB*Plus. Also note that the file name suffix can be `.jks` instead of `.keystore` (thus, file name `postgresql.jks`). However, in the remainder of these examples, file name `postgresql.keystore` is used.

For Windows only: The path is %APPDATA%\postgresql\postgresql.keystore

The `keytool` program can be found under the `bin` subdirectory of the Java Runtime Environment installation.

You will be prompted for a new password. **Save this password as it must be specified with the PGSSLCERTPASS environment variable.**

```
$ /usr/java/jdk1.8.0_131/jre/bin/keytool -keystore postgresql.keystore
> -alias postgresqlstore -import -file edb_root.crt.der
Enter keystore password:
Re-enter new password:
Owner: CN=enterprisedb
Issuer: CN=enterprisedb
Serial number: c60f40256b0e8d53
Valid from: Tue Aug 22 10:59:25 EDT 2017 until: Wed Aug 22 10:59:25 EDT
2018
Certificate fingerprints:
MD5: 85:0B:E9:A7:6E:4F:7C:B0:9B:D6:3A:44:55:E2:E9:8E
SHA1: DD:A6:71:24:0B:6C:F8:BC:7A:4C:89:9B:DC:22:6A:6C:B0:F5:3F:7C
SHA256:
DC:02:64:E2:B0:E9:6F:1C:FC:4F:AE:E6:18:85:0B:79:57:43:C3:C5:AE:43:0D:37:49:53:6D:11:69:06:46:48
Signature algorithm name: SHA1withRSA
Version: 1
Trust this certificate? [no]: yes
Certificate was added to keystore
```

Step 5: Create a `PKCS #12` format of the keystore file (`postgresql.p12`) using files `edb.crt` and `edb.key` as input.

Note: The file name `postgresql.p12` is recommended so that it can be accessed in its default directory location `~/.postgresql/postgresql.p12`, which is under the home directory of the user account invoking EDB*Plus.

For Windows only: The path is %APPDATA%.postgresql\postgresql.p12

You will be prompted for a new password. **Save this password as it must be specified with the PGSSLKEYPASS environment variable.**

```
$ openssl pkcs12 -export -in edb.crt -inkey edb.key -out postgresql.p12

Enter Export Password:

Verifying - Enter Export Password:

$ ls -l

total 24

-rw-rw-r-- 1 user user 985 Aug 24 12:18 edb.crt

-rw-rw-r-- 1 user user 1704 Aug 24 12:18 edb.key

-rw-rw-r-- 1 user user 985 Aug 24 12:20 edb_root.crt

-rw-rw-r-- 1 user user 686 Aug 24 12:20 edb_root.crt.der

-rw-rw-r-- 1 user user 758 Aug 24 12:26 postgresql.keystore

-rw-rw-r-- 1 user user 2285 Aug 24 12:28 postgresql.p12
```

Step 6: If the `postgresql.keystore` and `postgresql.p12` files are not already in the `~/.postgresql` directory, move or copy them to that location.

For Windows only: The directory is `%APPDATA%\postgresql`.

Step 7: If the default location `~/.postgresql` is not used, then the full path (including the file name) to the `postgresql.keystore` file must be set with the `PGSSLCERT` environment variable, and the full path (including the file name) to file `postgresql.p12` must be set with the `PGSSLKEY` environment variable before invoking EDB*Plus.

In addition, if the generated file from Step 4 was not named `postgresql.keystore` or `postgresql.jks` then, use the `PGSSLCERT` environment variable to designate the file name and its location. Similarly, if the generated file from Step 5 was not named `postgresql.p12` then, use the `PGSSLKEY` environment variable to designate the file name and its location.

Requesting an SSL Connection between EDB*Plus and the Advanced Server Database

Be sure the following topics have been addressed in order to perform an SSL connection:

- The trusted certificate and keystore files have been generated for both the database server and the client host to be invoking EDB*Plus.
- The `postgresql.conf` file for the database server contains the updated configuration parameters.
- The `pg_hba.conf` file for the database server contains the required entry for permitting the SSL connection.
- For the client host, either the client's certificate and keystore files have been placed in the user account's `~/.postgresql` directory or the environment variables `PGSSLCERT` and `PGSSLKEY` are set before invoking EDB*Plus.
- The `PGSSLCERTPASS` environment variable is set with a password.
- The `PGSSLKEYPASS` environment variable is set with a password

When invoking EDB*Plus, include the `?ssl=true` parameter in the database connection string as shown for the `connectstring` option.

The following example demonstrates EDB*Plus invoked from a host that is remote to the database server.

The `postgresql.conf` file of the database server contains the following modified parameters:

```
ssl = on
ssl_cert_file = 'server.crt'
ssl_key_file = 'server.key'
ssl_ca_file = 'root.crt'
```

The `pg_hba.conf` file of the database server contains the following entry for connecting from EDB*Plus on the remote host:

```
### TYPE DATABASE USER ADDRESS METHOD
### "local" is for Unix domain socket connections only
local all all md5
### IPv4 local connections:
hostssl edb all 192.168.2.24/32 cert clientcert=1
```

On the remote host where EDB*Plus is to be invoked, the Linux user account named `user` contains the certificate

and keystore files in its `~/.postgresql` directory:

```
[user@localhost ~]$ whoami
user
[user@localhost ~]$ cd .postgresql
[user@localhost .postgresql]$ pwd
/home/user/.postgresql
[user@localhost .postgresql]$ ls -l
total 8
-rw-rw-r-- 1 user user 758 Aug 24 12:37 postgresql.keystore
-rw-rw-r-- 1 user user 2285 Aug 24 12:37 postgresql.p12
```

Logged into Linux with the account named user, EDB Plus is successfully invoked with the `ssl=true` parameter:

```
$ export PGSSLCERTPASS=keypass
$ export PGSSLKEYPASS=exppass
$ cd /opt/edb/as10/edbplus
$ ./edbplus.sh enterprisedb/password@192.168.2.22:5444/edb?ssl=true
```

Connected to EnterpriseDB 10.0.1 (192.168.2.22:5444/edb) AS enterprisedb

EDB Plus: Release 10 (Build 36.0.0)

Copyright (c) 2008-2017, EnterpriseDB Corporation. All rights reserved.

SQL>

Alternatively, without placing the certificate and keystore files in `~/.postgresql`, but in a different directory, EDB*Plus can be invoked in the following manner:

```
$ export PGSSLCERT=/home/user/ssl/postgresql.keystore
$ export PGSSLKEY=/home/user/ssl/postgresql.p12
$ export PGSSLCERTPASS=keypass
$ export PGSSLKEYPASS=exppass
$ cd /opt/edb/as10/edbplus
$ ./edbplus.sh enterprisedb/password@192.168.2.22:5444/edb?ssl=true
```

Connected to EnterpriseDB 10.0.1 (192.168.2.22:5444/edb) AS enterprisedb

EDB Plus: Release 10 (Build 36.0.0)

Copyright (c) 2008-2017, EnterpriseDB Corporation. All rights reserved.

SQL>

Note that in both cases the database user name used to log into EDB*Plus is enterpriseDb as this is the user specified for the common name field when creating the certificate. Other database user names can be used if the `pg_hba.conf` file with the `map` option and the `pg_ident.conf` file are used.

12.6 Command Summary

The following sections contains a summary of EDB*Plus commands.

ACCEPT

The ACCEPT command displays a prompt and waits for the user's keyboard input. The value input by the user is placed in the specified variable.

`ACC[EPT] variable`

The following example creates a new variable named `my_name`, accepts a value of John Smith, then displays the value using the DEFINE command.

```
SQL> ACCEPT my_name
Enter value for my_name: John Smith
SQL> DEFINE my_name
DEFINE MY_NAME = "John Smith"
```

APPEND

APPEND is a line editor command that appends the given text to the end of the current line in the SQL buffer.

`A[PPEND] text`

In the following example, a SELECT command is built-in the SQL buffer using the APPEND command. Note that two spaces are placed between the APPEND command and the WHERE clause in order to separate dept and WHERE by one space in the SQL buffer.

```
SQL> APPEND SELECT * FROM dept
SQL> LIST
1 SELECT * FROM dept
SQL> APPEND WHERE deptno = 10
SQL> LIST
1 SELECT * FROM dept WHERE deptno = 10
```

CHANGE

CHANGE is a line editor command performs a search-and-replace on the current line in the SQL buffer.

C[HANGE] FROM [TO]

If **TO** / is specified, the first occurrence of text **FROM** in the current line is changed to text **TO** . If **TO** / is omitted, the first occurrence of text **FROM** in the current line is deleted.

The following sequence of commands makes line 3 the current line, then changes the department number in the WHERE clause from 20 to 30.

```
SQL> LIST
1 SELECT empno, ename, job, sal, comm
2 FROM emp
3 WHERE deptno = 20
4* ORDER BY empno
SQL> 3
3* WHERE deptno = 20
SQL> CHANGE /20/30/
3* WHERE deptno = 30
SQL> LIST
1 SELECT empno, ename, job, sal, comm
2 FROM emp
3 WHERE deptno = 30
4* ORDER BY empno
```

CLEAR

The CLEAR command removes the contents of the SQL buffer, deletes all column definitions set with the COLUMN command, or clears the screen.

CL[EAR] [BUFF[ER] | SQL | COL[UMNS] | SCR[EEN]] BUFFER | SQL

Clears the SQL buffer.

COLUMNS

Removes column definitions.

SCREEN

Clears the screen. This is the default if no options are specified.

COLUMN

The COLUMN command controls output formatting. The formatting attributes set by using the COLUMN command remain in effect only for the duration of the current session.

```
COL[UMN ]
[ column
{ CLE[AR ] |
{ FOR[MAT ] spec |
HEA[DING ] text |
{ OFF | ON }
} [...]
```

```
}
```

If the `COLUMN` command is specified with no subsequent options, formatting options for current columns in effect for the session are displayed.

If the `COLUMN` command is followed by a column name, then the column name may be followed by one of the following:

1. No other options
2. `CLEAR`
3. Any combination of `FORMAT`, `HEADING`, and one of `OFF` or `ON`

`column`

Name of a column in a table to which subsequent column formatting options are to apply. If no other options follow `column`, then the current column formatting options if any, of `column` are displayed.

`CLEAR`

The `CLEAR` option reverts all formatting options back to their defaults for `column`. If the `CLEAR` option is specified, it must be the only option specified.

`spec`

Format specification to be applied to `column`. For character columns, `spec` takes the following format:

`n`

`n` is a positive integer that specifies the column width in characters within which to display the data. Data in excess of `n` will wrap around with the specified column width.

For numeric columns, `spec` is comprised of the following elements.

Table - Numeric Column Format Elements

| Element | Description |
|---------|--|
| \$ | Display a leading dollar sign. |
| , | Display a comma in the indicated position. |
| . | Marks the location of the decimal point. |
| 0 | Display leading zeros. |
| 9 | Number of significant digits to display. |

If loss of significant digits occurs due to overflow of the format, then all #'s are displayed.

`text`

Text to be used for the column heading of `column`.

`OFF` | `ON`

If OFF is specified, formatting options are reverted back to their defaults, but are still available within the session. If ON is specified, the formatting options specified by previous COLUMN commands for `column` within the session are re-activated.

The following example shows the effect of changing the display width of the job column.

```
SQL> SET PAGESIZE 9999
SQL> COLUMN job FORMAT A5
```

```
SQL> COLUMN job
COLUMN JOB ON
FORMAT A5
wrapped
SQL> SELECT empno, ename, job FROM emp;
```

| EMPNO | ENAME | JOB |
|-------|--------|---------------|
| 7369 | SMITH | CLERK |
| 799 | ALLEN | SALES MAN |
| 7521 | WARD | SALES MAN |
| 7566 | JONES | MANAG ER |
| 7654 | MARTIN | SALES MAN |
| 7698 | BLAKE | MANAG ER |
| 7782 | CLARK | MANAG ER |
| 7788 | SCOTT | ANALY ST |
| 7839 | KING | PRESI DENT |
| 7844 | TURNER | SALES MAN |
| 7876 | ADAMS | CLERK |
| 7900 | JAMES | CLERK |
| 7902 | FORD | ANALY ST |
| 7934 | MILLER | CLERK |

| | | |
|------|--------|---------------|
| 7369 | SMITH | CLERK |
| 799 | ALLEN | SALES MAN |
| 7521 | WARD | SALES MAN |
| 7566 | JONES | MANAG ER |
| 7654 | MARTIN | SALES MAN |
| 7698 | BLAKE | MANAG ER |
| 7782 | CLARK | MANAG ER |
| 7788 | SCOTT | ANALY ST |
| 7839 | KING | PRESI DENT |
| 7844 | TURNER | SALES MAN |
| 7876 | ADAMS | CLERK |
| 7900 | JAMES | CLERK |
| 7902 | FORD | ANALY ST |
| 7934 | MILLER | CLERK |

14 rows retrieved.

The following example applies a format to the `sal` column.

```
SQL> COLUMN sal FORMAT $99,999.00
SQL> COLUMN JOB ON
FORMAT A5
wrapped
COLUMN SAL ON
FORMAT $99,999.00
wrapped
SQL> SELECT empno, ename, job, sal FROM emp;
```

| EMPNO | ENAME | JOB | SAL |
|-------|-------|-------|-------------------|
| 7369 | SMITH | CLERK | \$800.00 |
| 7499 | ALLEN | SALES | \$1,600.00 MAN |
| 7521 | WARD | SALES | \$1,250.00 MAN |
| 7566 | JONES | MANAG | \$2,975.00 |

| | | | |
|------|-------|-------|-------------------|
| 7369 | SMITH | CLERK | \$800.00 |
| 7499 | ALLEN | SALES | \$1,600.00 MAN |
| 7521 | WARD | SALES | \$1,250.00 MAN |
| 7566 | JONES | MANAG | \$2,975.00 |

```

        ER
7654 MARTIN    SALES $1,250.00
        MAN
7698 BLAKE     MANAG $2,850.00
        ER
7782 CLARK     MANAG $2,450.00
        ER
7788 SCOTT      ANALY $3,000.00
        ST
7839 KING       PRESI $5,000.00
        DENT
7844 TURNER     SALES $1,500.00
        MAN
7876 ADAMS      CLERK $1,100.00
7900 JAMES      CLERK $950.00
7902 FORD       ANALY $3,000.00
        ST
7934 MILLER     CLERK $1,300.00

```

14 rows retrieved.

CONNECT

Change the database connection to a different user and/or connect to a different database. There must be no white space between any of the parameters following the CONNECT command.

`CON[NECT] username [/password][@{ connectstring | variable }]`

Where:

`username` is a database username with which to connect to the database.

`password` is the password associated with the specified `username`. If a `password` is not provided, but a password is required for authentication, a search is made for a password file, first in the home directory of the Linux operating system account invoking EDB*Plus (or in the `%APPDATA%\postgresql` directory for Windows) and then at the location specified by the PGPASSFILE environment variable. The password file is `.pgpass` on Linux hosts and `pgpass.conf` on Windows hosts. The following is an example on a Windows host:

`C:\Users\Administrator\AppData\Roaming\postgresql\pgpass.conf`

If a password file cannot be located, or it does not have an entry matching the EDB*Plus connection parameters, then EDB*Plus will prompt for the password. For more information about password files, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/10/static/libpq-pgpass.html>

Note: When a password is not required, EDB*Plus does not prompt for a password such as when the trust authentication method is specified in the `pg_hba.conf` file. For more information about the `pg_hba.conf` file and authentication methods, see the PostgreSQL core documentation at:

<https://www.postgresql.org/docs/10/static/auth-pg-hba-conf.html>

`connectstring` is the database connection string. See Section 2.2 for further information on the database connection string.

`variable` is a variable defined in the `login.sql` file that contains a database connection string. The `login.sql` file can be found in the `edbplus` subdirectory of the Advanced Server home directory.

In the following example, the database connection is changed to database `edb` on the `localhost` at port `5445` with username `smith`.

```
SQL> CONNECT smith/mypassword@localhost:5445/edb
Disconnected from EnterpriseDB Database.
Connected to EnterpriseDB 10.0.1 (localhost:5445/edb) AS smith
```

From within the session shown above, the connection is changed to username `enterprisedb`. Also note that the host defaults to the `localhost`, the port defaults to `5444` (which is not the same as the port previously used), and the database defaults to `edb`.

```
SQL> CONNECT enterprisedb/password
Disconnected from EnterpriseDB Database.
Connected to EnterpriseDB 10.0.1 (localhost:5444/edb) AS enterprisedb
```

DEFINE

The `DEFINE` command creates or replaces the value of a *user variable* (also called a *substitution variable*).

`DEF[INE] [<variable> [= <text>]]`

If the `DEFINE` command is given without any parameters, all current variables and their values are displayed.

If `DEFINE variable` is given, only `variable` is displayed with its value.

`DEFINE variable = text` assigns `text` to `variable`. `text` may be optionally enclosed within single or double quotation marks. Quotation marks must be used if `text` contains space characters.

The following example defines two variables, `dept` and `name`.

```
SQL> DEFINE dept = 20
SQL> DEFINE name = 'John Smith'
SQL> DEFINE
DEFINE EDB = "localhost:5445/edb"
DEFINE DEPT = "20"
DEFINE NAME = "John Smith"
```

Note: The variable `EDB` is read from the `login.sql` file located in the `edbplus` subdirectory of the Advanced Server home directory.

DEL

`DEL` is a line editor command that deletes one or more lines from the SQL buffer.

`DEL [n | n m | n * | n L[AST] | n | L[AST] | L[AST]]`

The parameters specify which lines are to be deleted from the SQL buffer. Two parameters specify the start and end of a range of lines to be deleted. If the `DEL` command is given with no parameters, the current line is deleted.

`n`

`n` is an integer representing the `n`th line

`n m`

`n` and `m` are integers where `m` is greater than `n` representing the `n`th through the `m`th lines

- Current line

LAST

Last line

In the following example, the fifth and sixth lines containing columns `sal` and `comm`, respectively, are deleted from the SELECT command in the SQL buffer.

```
SQL> LIST
1 SELECT
2 empno
3 ,ename
4 ,job
5 ,sal
6 ,comm
7 ,deptno
8* FROM emp
SQL> DEL 5 6
SQL> LIST
1 SELECT
2 empno
3 ,ename
4 ,job
5 ,deptno
6\* FROM emp
```

DESCRIBE

The DESCRIBE command displays:

- A list of columns, column data types, and column lengths for a table or view
- A list of parameters for a procedure or function
- A list of procedures and functions and their respective parameters for a package.

The DESCRIBE command will also display the structure of the database object referred to by a synonym. The syntax is:

`DESC[RIBE] [<schema>.] <object>`

schema

Name of the schema containing the object to be described.

object

Name of the table, view, procedure, function, or package to be displayed, or the synonym of an object.

DISCONNECT

The DISCONNECT command closes the current database connection, but does not terminate EDB*Plus.

DISC[ONNECT]

EDIT

The EDIT command invokes an external editor to edit the contents of an operating system file or the SQL buffer.

| ED[IT] [<filename> [.<ext>]]

<filename> [.<ext>]

filename is the name of the file to open with an external editor. **ext** is the filename extension. If the filename extension is sql, then the .sql extension may be omitted when specifying **filename**. EDIT always assumes a .sql extension on filenames that are specified with no extension. If the filename parameter is omitted from the EDIT command, the contents of the SQL buffer are brought into the editor.

EXECUTE

The EXECUTE command executes an SPL procedure from EDB*Plus.

| EXEC[UTE] <spl_procedure> [([<parameters>])]

spl_procedure

The name of the SPL procedure to be executed.

parameters

Comma-delimited list of parameters. If there are no parameters, then a pair of empty parentheses may optionally be specified.

EXIT

The EXIT command terminates the EDB*Plus session and returns control to the operating system. QUIT is a synonym for EXIT. Specifying no parameters is equivalent to EXIT SUCCESS COMMIT.

```
{ EXIT | QUIT }
[ SUCCESS | FAILURE | WARNING | <value> | <variable> ]
[ COMMIT | ROLLBACK ]SUCCESS | FAILURE |WARNING]
```

Returns an operating system dependent return code indicating successful operation, failure, or warning for SUCCESS, FAILURE, and WARNING, respectively. The default is SUCCESS.

value

An integer value that is returned as the return code.

variable

A variable created with the DEFINE command whose value is returned as the return code.

COMMIT | ROLLBACK

If COMMIT is specified, uncommitted updates are committed upon exit. If ROLLBACK is specified, uncommitted updates are rolled back upon exit. The default is COMMIT.

GET

The GET command loads the contents of the given file to the SQL buffer.

```
> GET filename[.<ext>] [ LIST | NOLIST ]
```

`filename` [.<ext>]

`filename` is the name of the file to load into the SQL buffer. `ext` is the filename extension. If the filename extension is `sql`, then the `.sql` extension may be omitted when specifying `filename`. GET always assumes a `.sql` extension on filenames that are specified with no extension.

`LIST | NOLIST`

If LIST is specified, the content of the SQL buffer is displayed after the file is loaded. If NOLIST is specified, no listing is displayed. The default is LIST.

HELP

The HELP command obtains an index of topics or help on a specific topic. The question mark (?) is synonymous with specifying HELP.

```
> { HELP | ? } { INDEX | <topic> }
```

`INDEX`

Displays an index of available topics.

`topic`

The name of a specific topic – e.g., an EDB*Plus command, for which help is desired.

HOST

The HOST command executes an operating system command from EDB*Plus.

```
| HO[ST] [<os_command>]
```

`os_command`

The operating system command to be executed. If you do not provide an operating system command, EDB*Plus pauses execution and opens a new shell prompt. When the shell exits, EDB*Plus resumes execution.

INPUT

The INPUT line editor command adds a line of text to the SQL buffer after the current line.

```
| I[INPUT] <text>
```

The following sequence of INPUT commands constructs a SELECT command.

```
SQL> INPUT SELECT empno, ename, job, sal, comm
SQL> INPUT FROM emp
SQL> INPUT WHERE deptno = 20
SQL> INPUT ORDER BY empno
SQL> LIST
1 SELECT empno, ename, job, sal, comm
```

```
2 FROM emp
3 WHERE deptno = 20
4* ORDER BY empno
```

LIST

LIST is a line editor command that displays the contents of the SQL buffer.

L[IST] [n | n m | n * | n L[AST] | * | * n | * L[AST] | L[AST]]

The buffer does not include a history of the EDB*Plus commands.

n

| n represents the buffer line number.

n m

| n m displays a list of lines between n and m.

n *

| n * displays a list of lines that range between line n and the current line.

n L[AST]

| n L[AST] displays a list of lines that range from line n through the last line in the buffer.

| * displays the current line.

*** n**

| * n displays a list of lines that range from the current line through line n.

*** L[AST]**

| * L[AST] displays a list of lines that range from the current line through the last line.

L[AST]

| L[AST] displays the last line.

PASSWORD

Use the PASSWORD command to change your database password.

PASSW[ORD] [<user_name>]

You must have sufficient privileges to use the PASSWORD command to change another user's password. The following example demonstrates using the PASSWORD command to change the password for a user named **acctg**:

```
SQL> PASSWORD acctg
Changing password for acctg
New password:
```

New password again:
Password successfully changed.

PAUSE

The PAUSE command displays a message, and waits for the user to press ENTER.

PAU[SE] [<optional_text>]

`optional_text` specifies the text that will be displayed to the user. If the `optional_text` is omitted, Advanced Server will display two blank lines. If you double quote the `optional_text` string, the quotes will be included in the output.

PROMPT

The PROMPT command displays a message to the user before continuing.

PRO[MPT] [<message_text>]

`message_text` specifies the text displayed to the user. Double quote the string to include quotes in the output.

QUIT

The QUIT command terminates the session and returns control to the operating system. QUIT is a synonym for EXIT.

```
QUIT
[SUCCESS | FAILURE | WARNING | value> | <sub_variable>]
[COMMIT | ROLLBACK]``
```

The default value is QUIT SUCCESS COMMIT.

REMARK

Use REMARK to include comments in a script.

| **REM[ARK] [<optional_text>]**

You may also use the following convention to include a comment:

```
*
```

* This is an example of a three line comment.

```
*
```

SAVE

Use the SAVE command to write the SQL Buffer to an operating system file.

SAV[E] <file_name> [CRE[ATE] | REP[LACE] | APP[END]]

file_name

`file_name` specifies the name of the file (including the path) where the buffer contents are written. If you do not provide a file extension, .sql is appended to the end of the file name.

CREATE

Include the CREATE keyword to create a new file. A new file is created *only* if a file with the specified name does not already exist. This is the default.

REPLACE

Include the REPLACE keyword to specify that Advanced Server should overwrite an existing file.

APPEND

Include the APPEND keyword to specify that Advanced Server should append the contents of the SQL buffer to the end of the specified file.

The following example saves the contents of the SQL buffer to a file named `example.sql`, located in the temp directory:

```
SQL> SAVE C:example.sql CREATE File "example.sql" written.
```

SET

Use the SET command to specify a value for a session level variable that controls EDB*Plus behavior. The following forms of the SET command are valid:

SET AUTOCOMMIT

Use the SET AUTOCOMMIT command to specify commit behavior for Advanced Server transactions.

SET AUTO[COMMIT] {ON | OFF | IMMEDIATE | <statement_count>}

Please note that EDB*Plus always automatically commits DDL statements.

ON

Specify ON to turn autocommit behavior on.

OFF

Specify OFF to turn autocommit behavior off.

IMMEDIATE

IMMEDIATE has the same effect as ON.

statement_count

Include a value for `statement_count` to instruct EDB*Plus to issue a commit after the specified count of successful SQL statements.

SET COLUMN SEPARATOR

Use the SET COLUMN SEPARATOR command to specify the text that Advanced Server displays between columns.

SET COLSEP <column_separator>

The default value of `column_separator` is a single space.

SET ECHO

Use the SET ECHO command to specify if SQL and EDB*Plus script statements should be displayed onscreen as they are executed.

SET ECHO {ON | OFF}

The default value is OFF.

SET FEEDBACK

The SET FEEDBACK command controls the display of interactive information after a SQL statement executes.

SET FEED[BACK] {ON | OFF | row_threshold}

`row_threshold`

Specify an integer value for `row_threshold`. Setting `row_threshold` to 0 is same as setting FEEDBACK to OFF. Setting `row_threshold` equal 1 effectively sets FEEDBACK to ON.

SET FLUSH

Use the SET FLUSH command to control display buffering.

SET FLU[SH] {ON | OFF}

Set FLUSH to OFF to enable display buffering. If you enable buffering, messages bound for the screen may not appear until the script completes. Please note that setting FLUSH to OFF will offer better performance.

Set FLUSH to ON to disable display buffering. If you disable buffering, messages bound for the screen appear immediately.

SET HEADING

Use the SET HEADING variable to specify if Advanced Server should display column headings for SELECT statements.

SET HEA[DING] {ON | OFF}

SET HEAD SEPARATOR

The SET HEADSEP command sets the new heading separator character used by the COLUMN HEADING command. The default is '|'.

SET HEADS[EP]

SET LINESIZE

Use the SET LINESIZE command to specify the width of a line in characters.

SET LIN[ESIZE] <width_of_line>

`width_of_line`

The default value of `width_of_line` is 132.

SET NEWPAGE

Use the SET NEWPAGE command to specify how many blank lines are printed after a page break.

SET NEWP[AGE] <lines_per_page>

`lines_per_page`

The default value of `lines_per_page` is 1.

SET NULL

Use the SET NULL command to specify a string that is displayed to the user when a NULL column value is displayed in the output buffer.

`SET NULL <null_string>`

SET PAGESIZE

Use the SET PAGESIZE command to specify the number of printed lines that fit on a page.

`SET PAGESIZE <line_count>`

Use the `line_count` parameter to specify the number of lines per page.

SET SQLCASE

The SET SQLCASE command specifies if SQL statements transmitted to the server should be converted to upper or lower case.

`SET SQLCASE {MIX[ED] | UP[PER] | LO[WER]}`

UPPER

Specify UPPER to convert the command text to uppercase.

LOWER

Specify LOWER to convert the command text to lowercase.

MIXED

Specify MIXED to leave the case of SQL commands unchanged. The default is MIXED.

SET PAUSE

The SET PAUSE command is most useful when included in a script; the command displays a prompt and waits for the user to press Return.

`SET PAUSE {ON | OFF}`

If SET PAUSE is ON, the message Hit ENTER to continue... will be displayed before each command is executed.

SET SPACE

Use the SET SPACE command to specify the number of spaces to display between columns:

`SET SPACE <number_of_spaces>`

SET SQLPROMPT

Use SET SQLPROMPT to set a value for a user-interactive prompt:

`SET SQLPROMPT "<prompt>"`

By default, SQLPROMPT is set to "SQL> "

SET TERMOUT

Use the SET TERMOUT command to specify if command output should be displayed onscreen.

SET TERM[OUT] {ON | OFF}

SET TIMING

The SET TIMING command specifies if Advanced Server should display the execution time for each SQL statement after it is executed.

SET TIMI[NG] {ON | OFF}

SET TRIMSPOOL

Use the SET TRIMSPOOL command to remove trailing spaces from each line in the output file specified by the SPOOL command.

SET TRIMS[POOL] {ON | OFF}

The default value is OFF.

SET VERIFY

Specifies if both the old and new values of a SQL statement are displayed when a substitution variable is encountered.

SET VER[IFY] { ON | OFF }

SHOW

Use the SHOW command to display current parameter values.

> **SHO[W] {ALL | parameter_name}**

Display the current parameter settings by including the ALL keyword:

```
SQL> SHOW ALL
autocommit OFF
colsep ""
define "&"
echo OFF
FEEDBACK ON for 6 row(s).
flush ON
heading ON
headsep "|"
linesize 78
newpage 1
null ""
pagesize 14
pause OFF
serveroutput OFF
spool OFF
sqlcase MIXED
sqlprompt "SQL> "
sqlterminator ";"
suffix ".sql"
```

```
termout ON
timing OFF
verify ON
USER is "enterprisedb"
HOST is "localhost"
PORT is "5444"
DATABASE is "edb"
VERSION is "10.0.0"
```

Or display a specific parameter setting by including the parameter_name in the SHOW command:

```
SQL> SHOW VERSION
VERSION is "10.0.0"
```

SPOOL

The SPOOL command sends output from the display to a file.

```
> SP[OOL] <output_file> | OFF
```

Use the `output_file` parameter to specify a path name for the output file.

START

Use the START command to run an EDB*Plus script file; START is an alias for @ command.

```
| STA[RT] <script_file>
```

Specify the name of a script file in the `script_file` parameter.

UNDEFINE

The UNDEFINE command erases a user variable created by the DEFINE command.

```
| UNDEF[INE] <variable_name> [ <variable_name>...]
```

Use the `variable_name` parameter to specify the name of a variable or variables.

WHENEVER SQLERROR

The WHENEVER SQLERROR command provides error handling for SQL errors or PL/SQL block errors. The syntax is:

```
WHENEVER SQLERROR
{CONTINUE|COMMIT|ROLLBACK|NONE}
|EXIT[SUCCESS|FAILURE|WARNING|<n>|<sub_variable>]
[COMMIT|ROLLBACK]}
```

If Advanced Server encounters an error during the execution of a SQL command or PL/SQL block, EDB*Plus performs the action specified in the WHENEVER SQLERROR command:

Include the **CONTINUE** clause to instruct EDB*Plus to perform the specified action before continuing.

Include the **COMMIT** clause to instruct EDB*Plus to COMMIT the current transaction before exiting or continuing.

Include the **ROLLBACK** clause to instruct EDB*Plus to ROLLBACK the current transaction before exiting or continuing.

Include the **NONE** clause to instruct EDB*Plus to continue without committing or rolling back the transaction.

Include the **EXIT** clause to instruct EDB*Plus to perform the specified action and exit if it encounters an error.

Use the following options to specify a status code that EDB*Plus will return before exiting:

[SUCCESS|FAILURE|WARNING|<n>|<sub_variable>]

Please note that EDB*Plus supports substitution variables, but does not support bind variables.